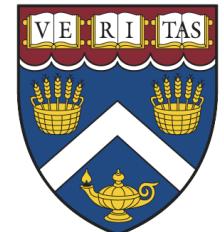
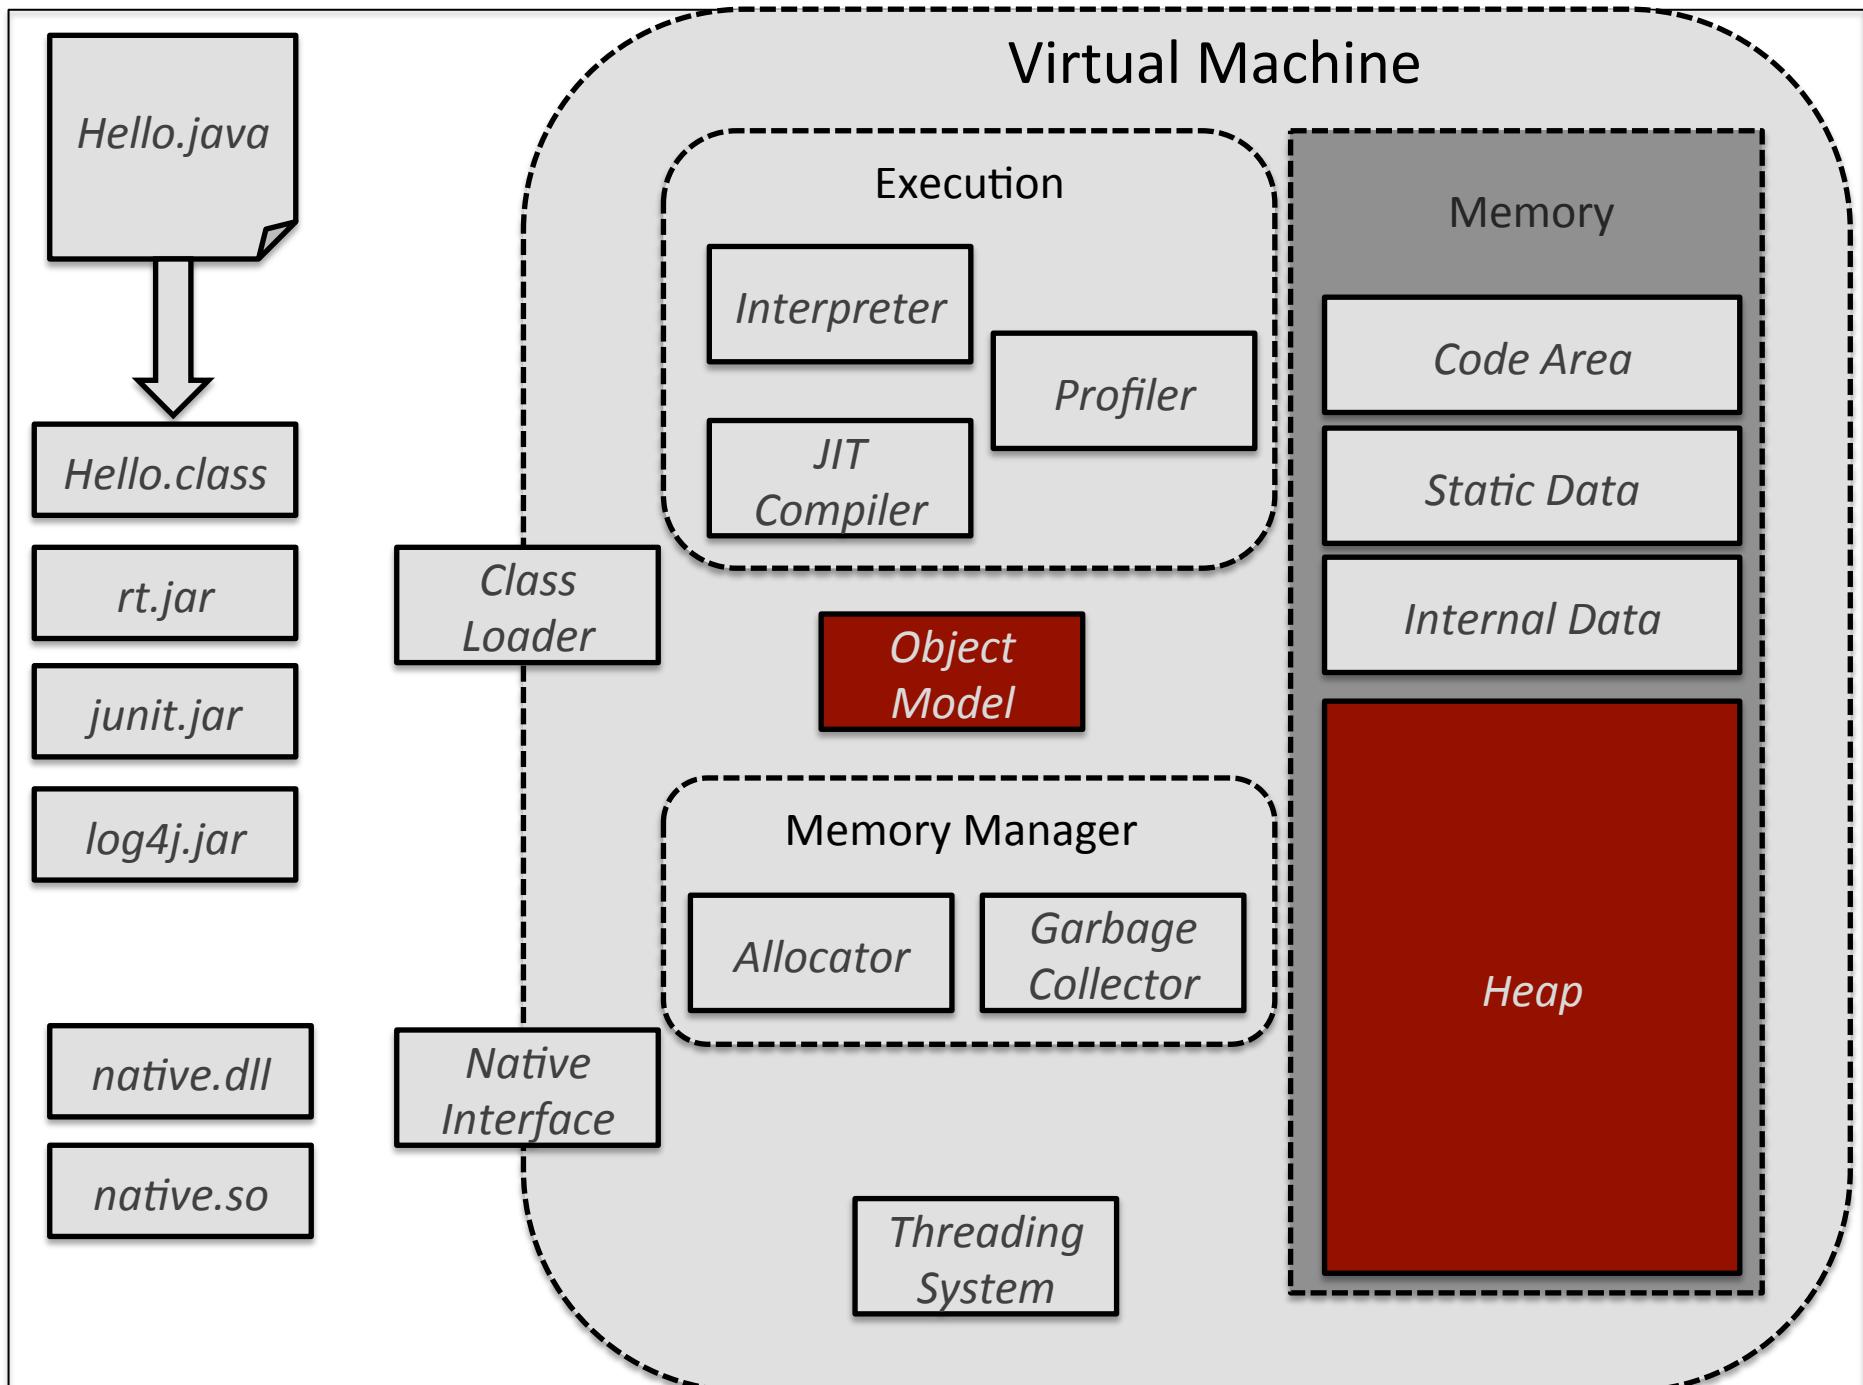


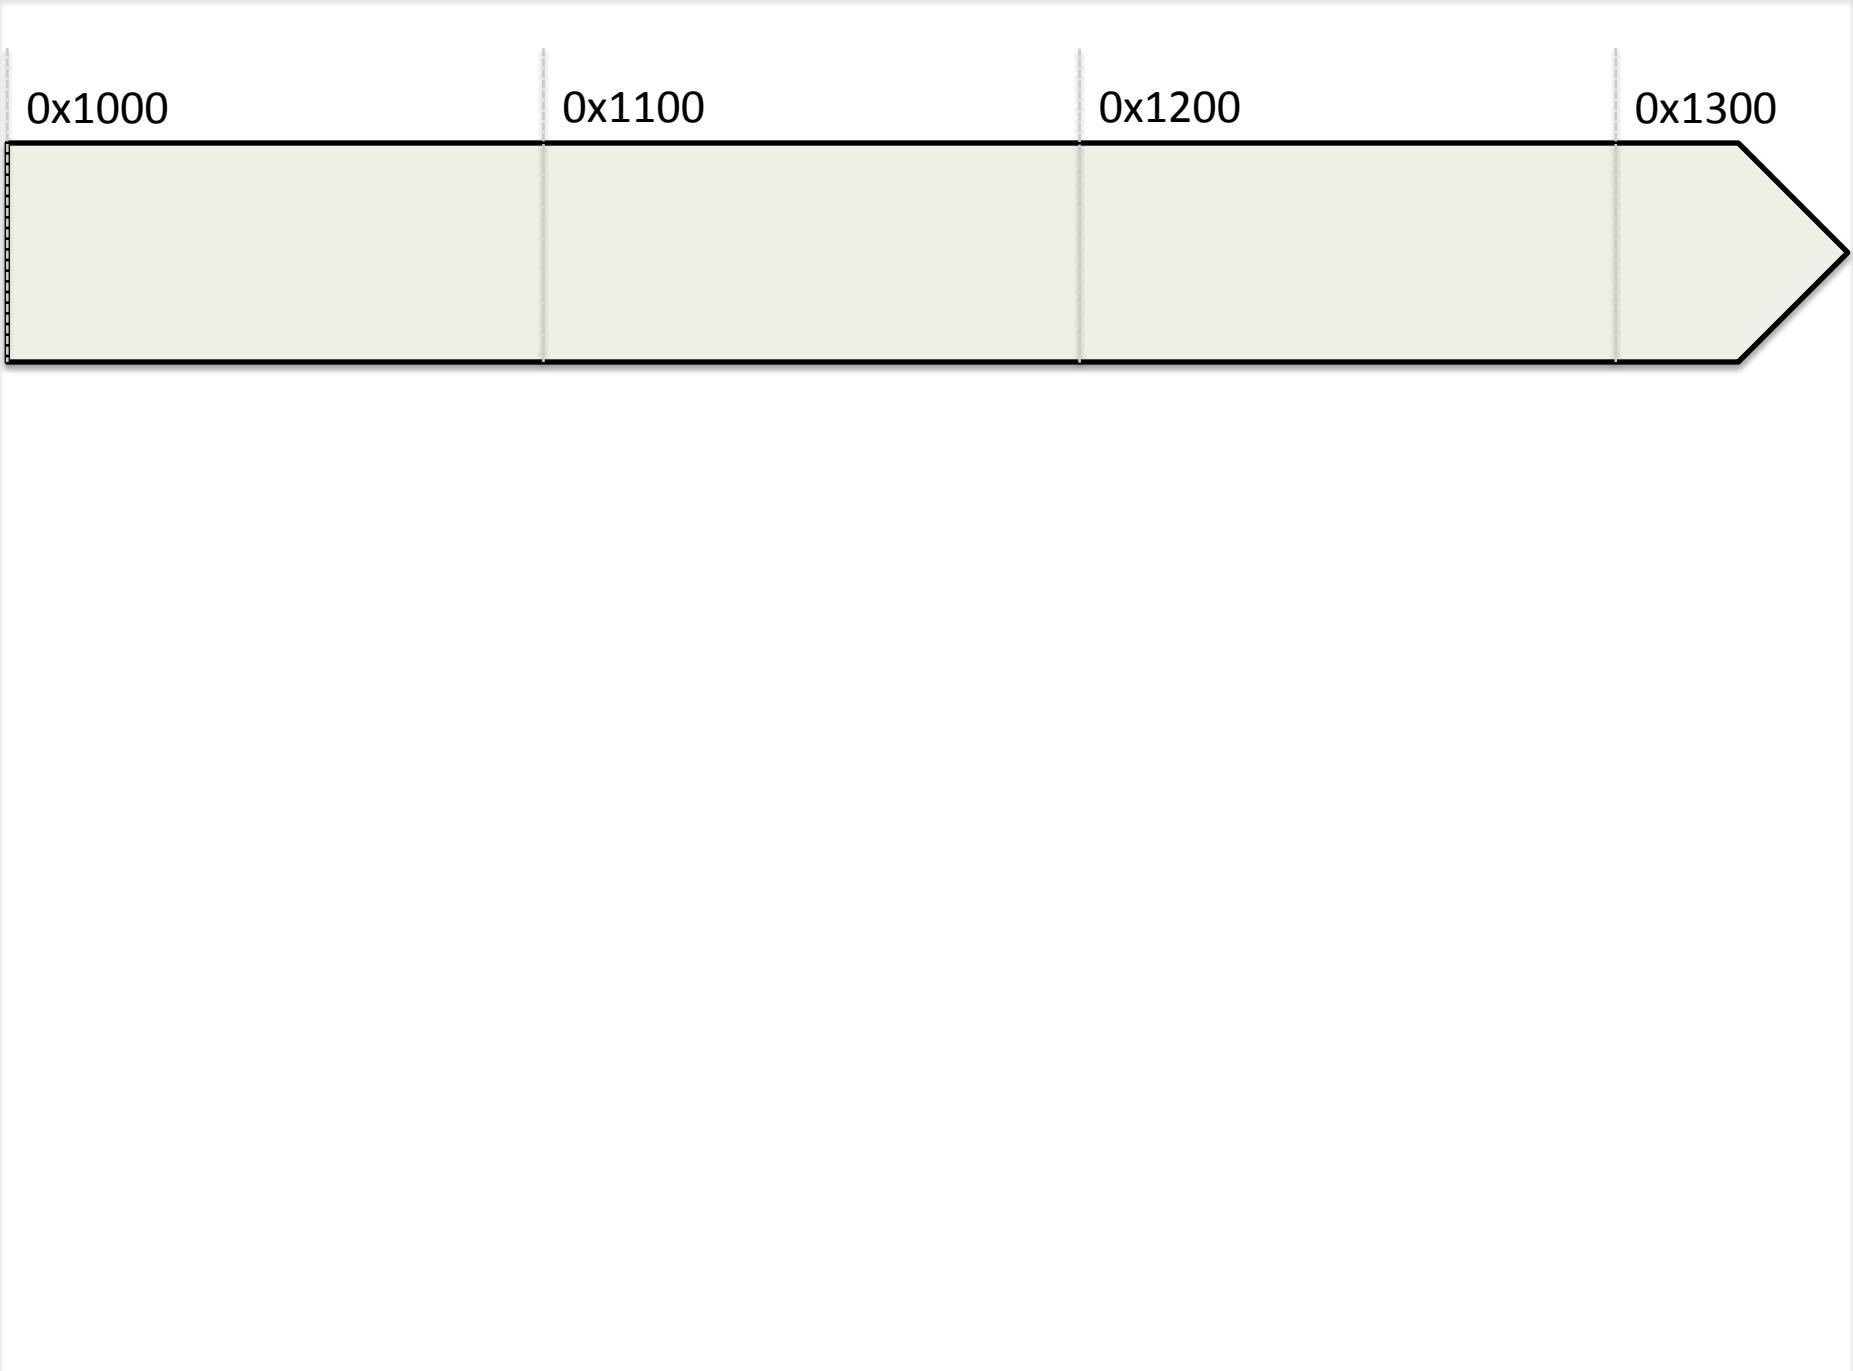
Object Representations

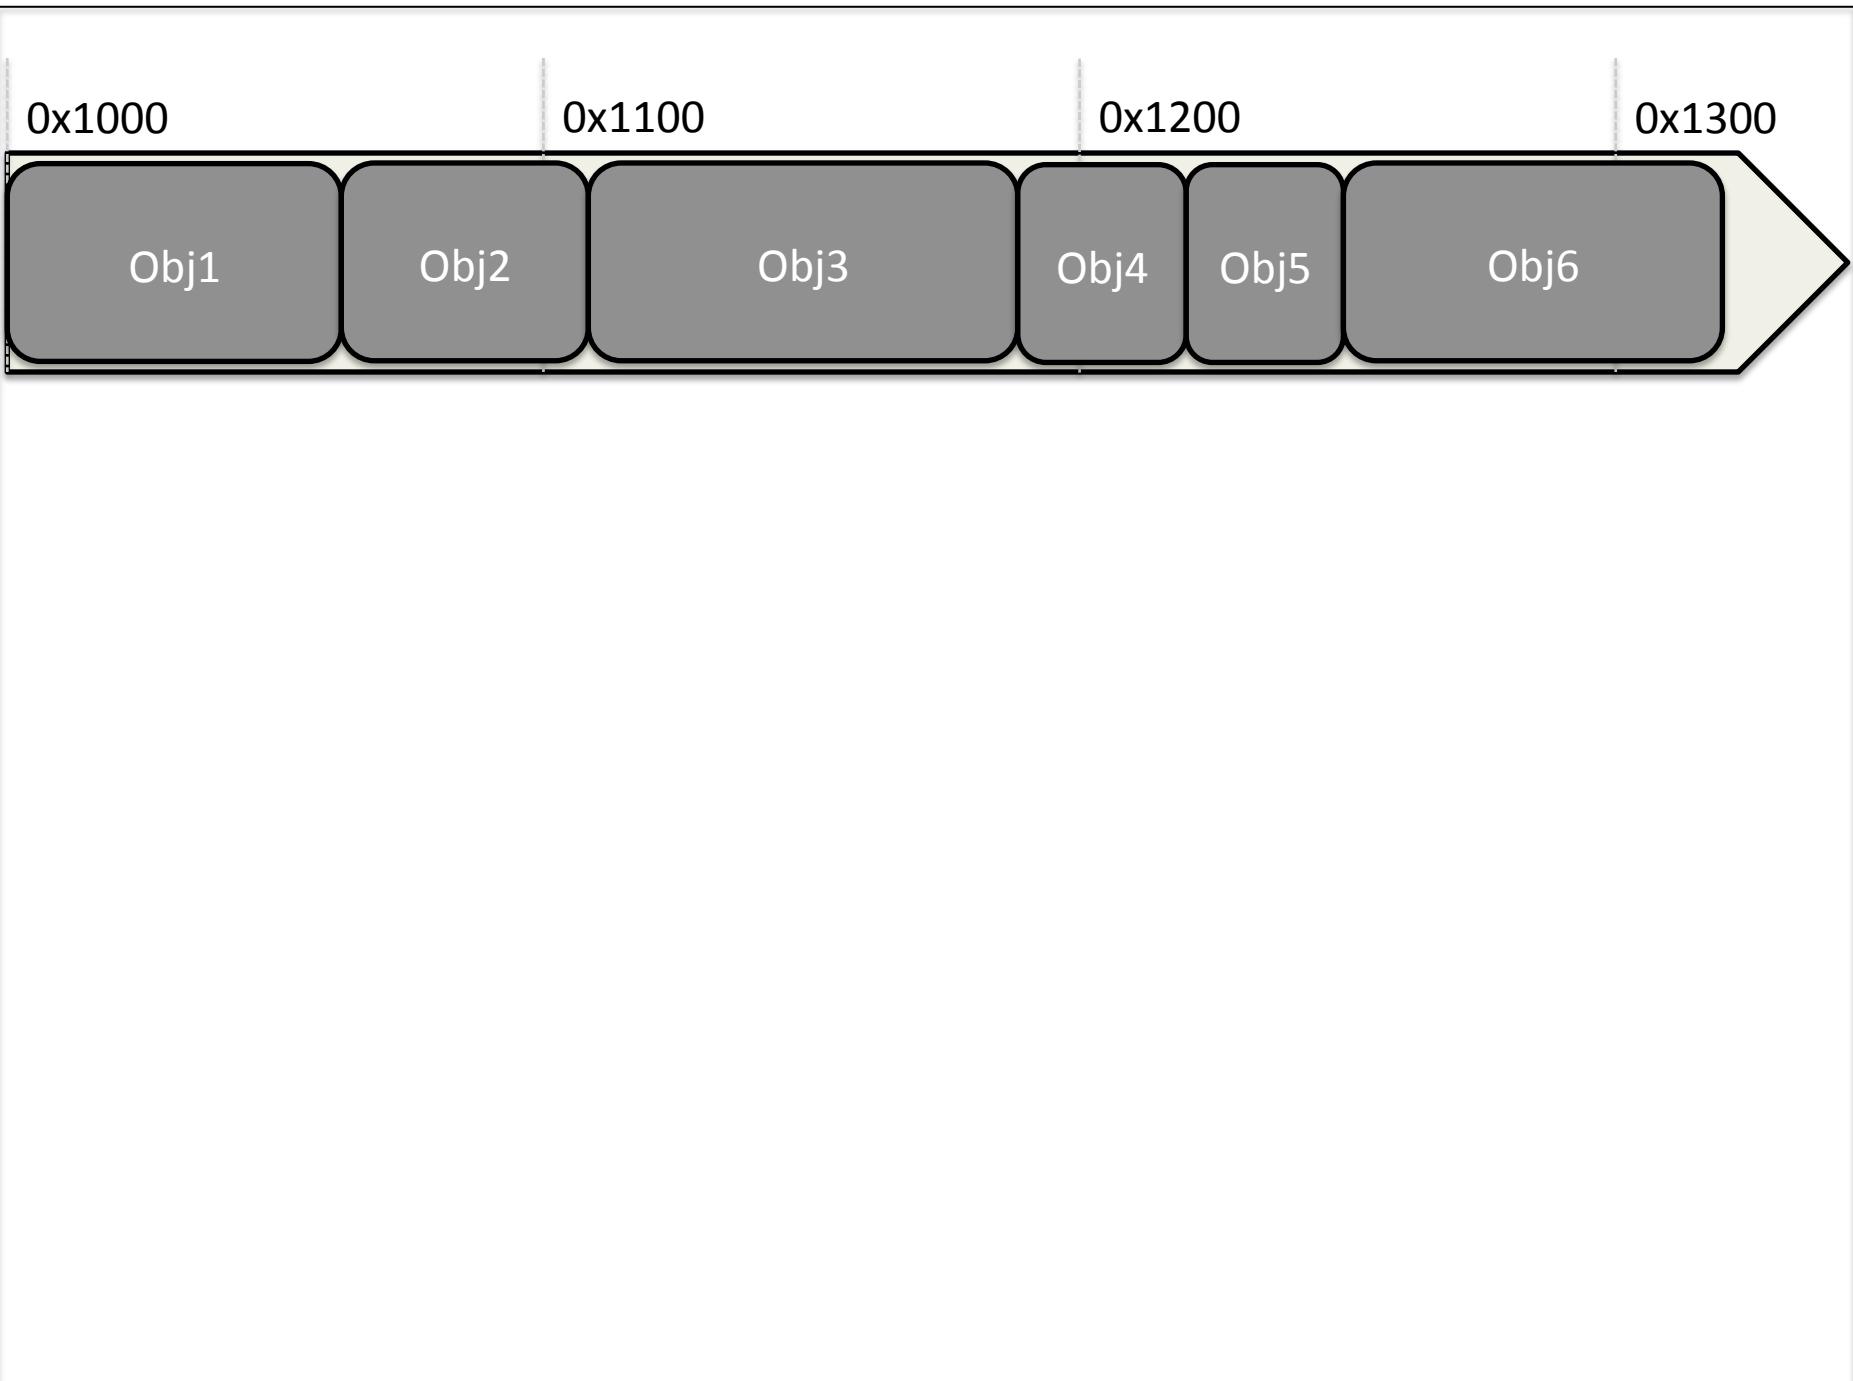


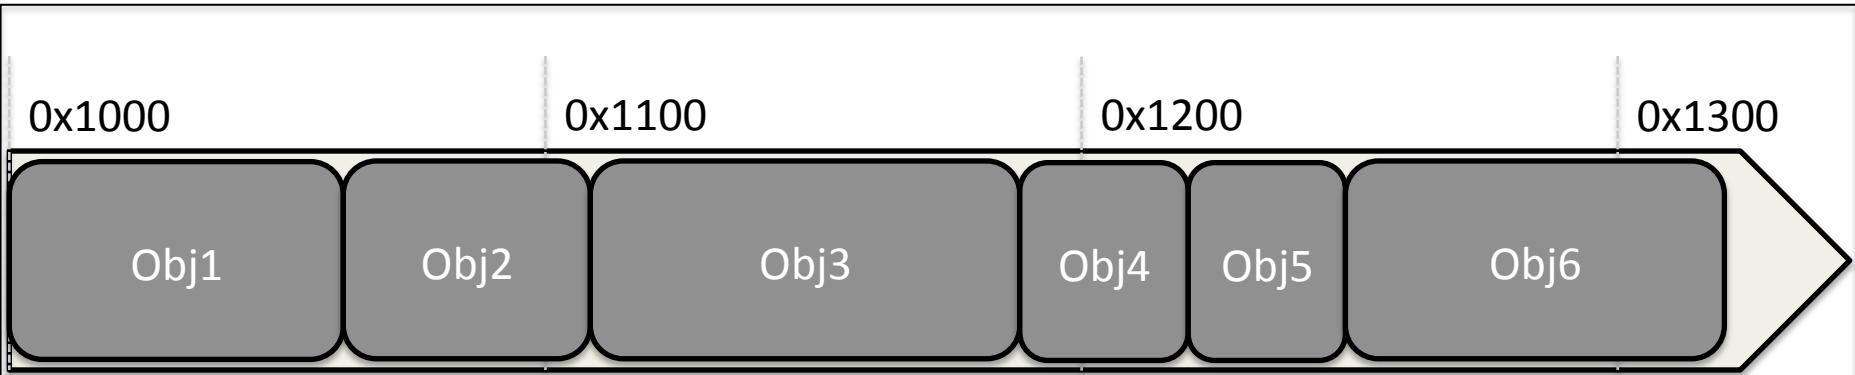


Primitives and Objects

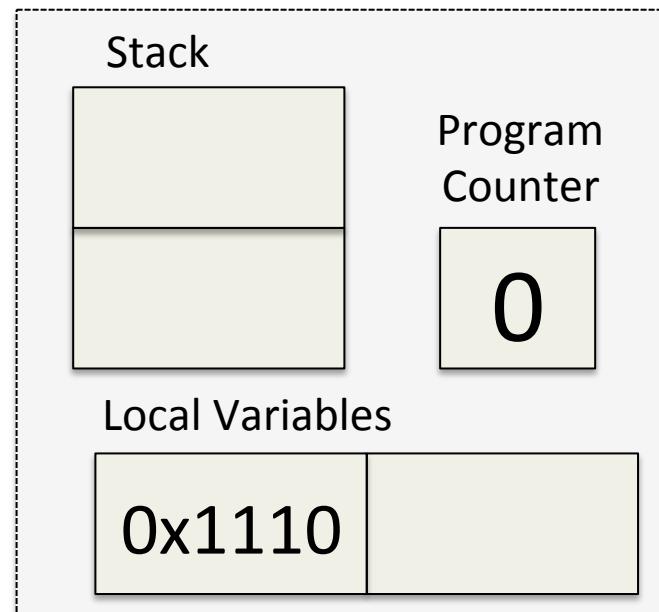
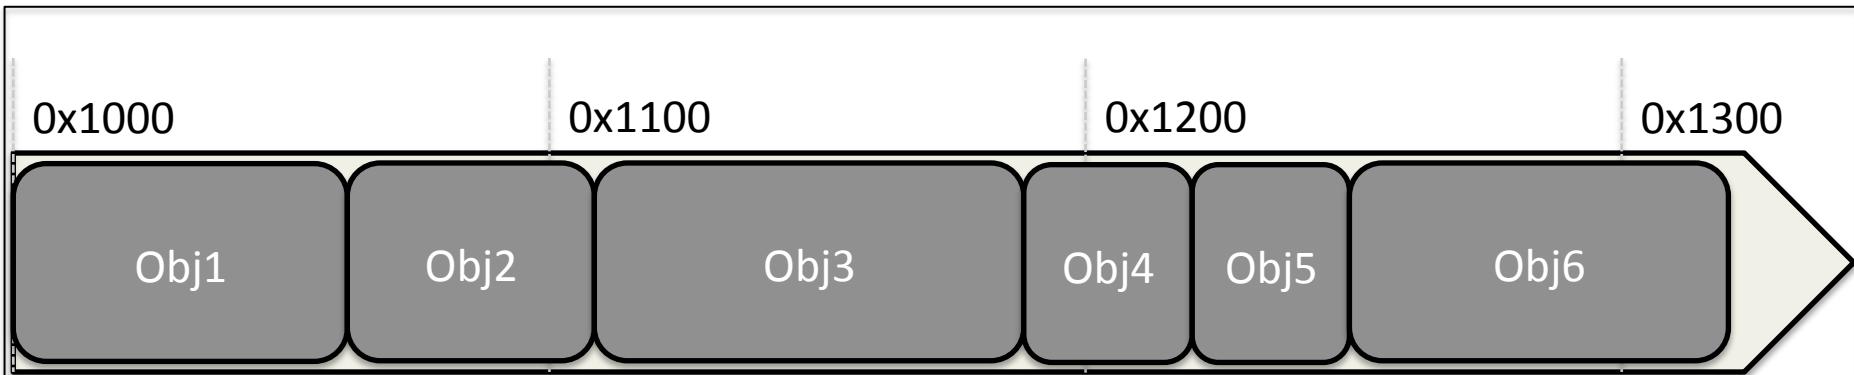
- Primitives
 - Built-in
 - Immutable
 - Stored on the Java stack or in fields
- Objects
 - Classes defined by application or library
 - Mutable
 - Stored on the heap



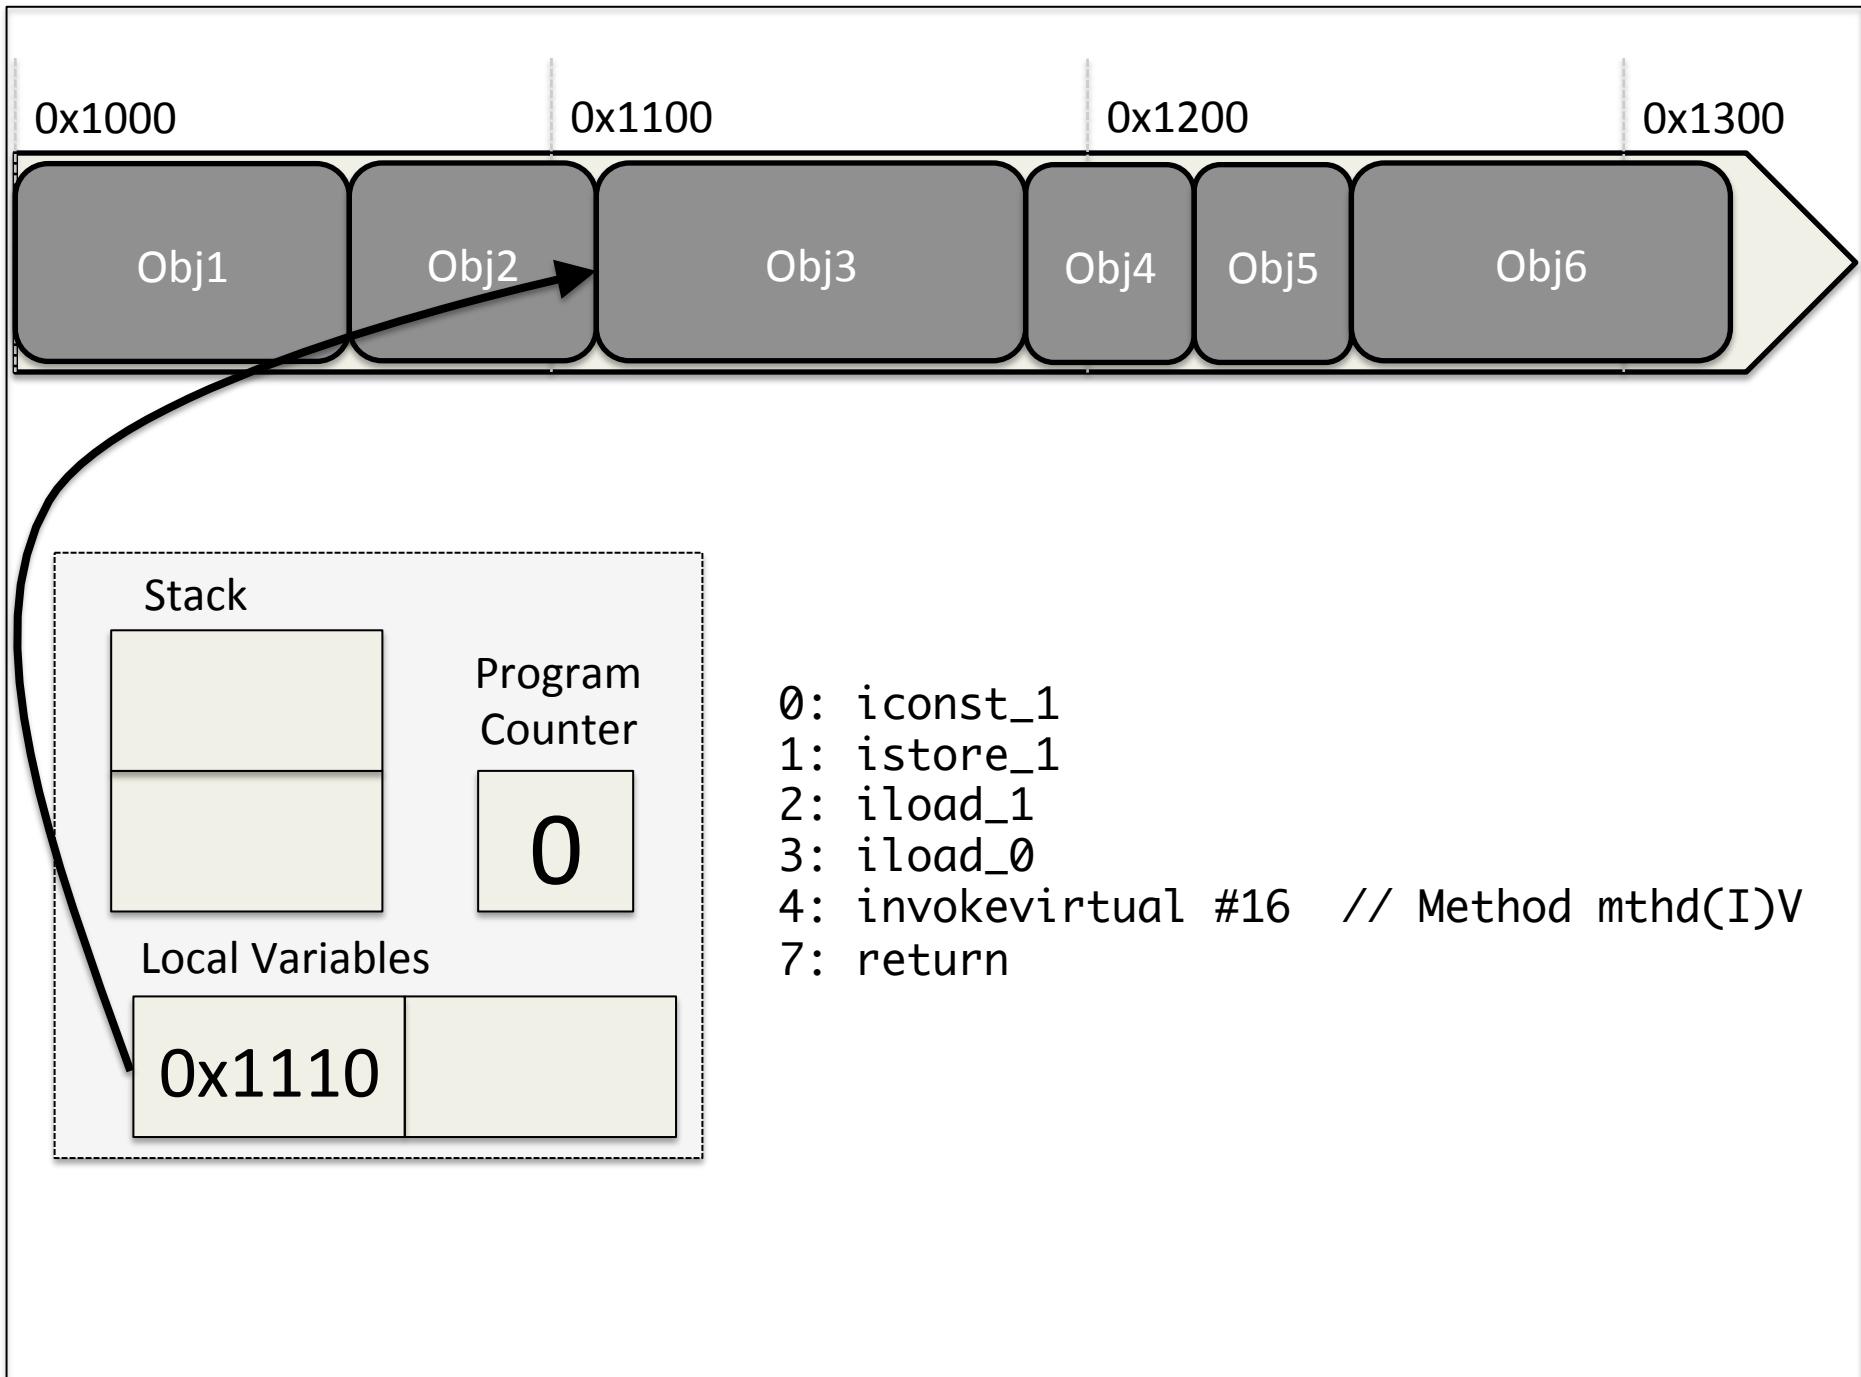


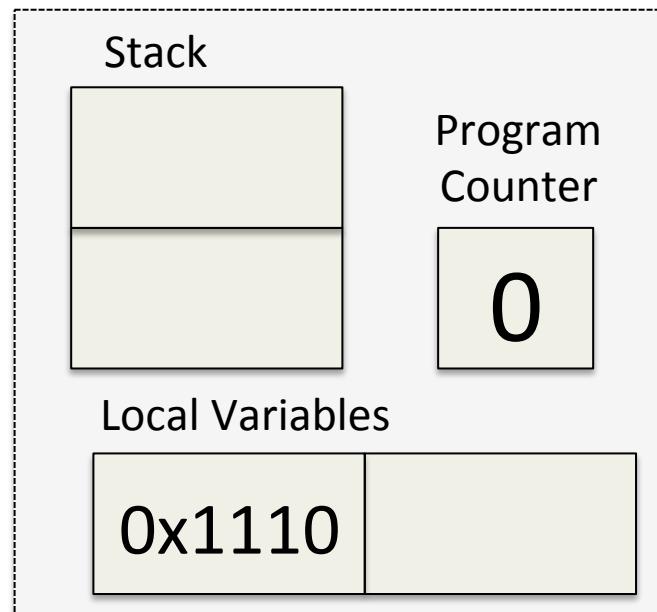
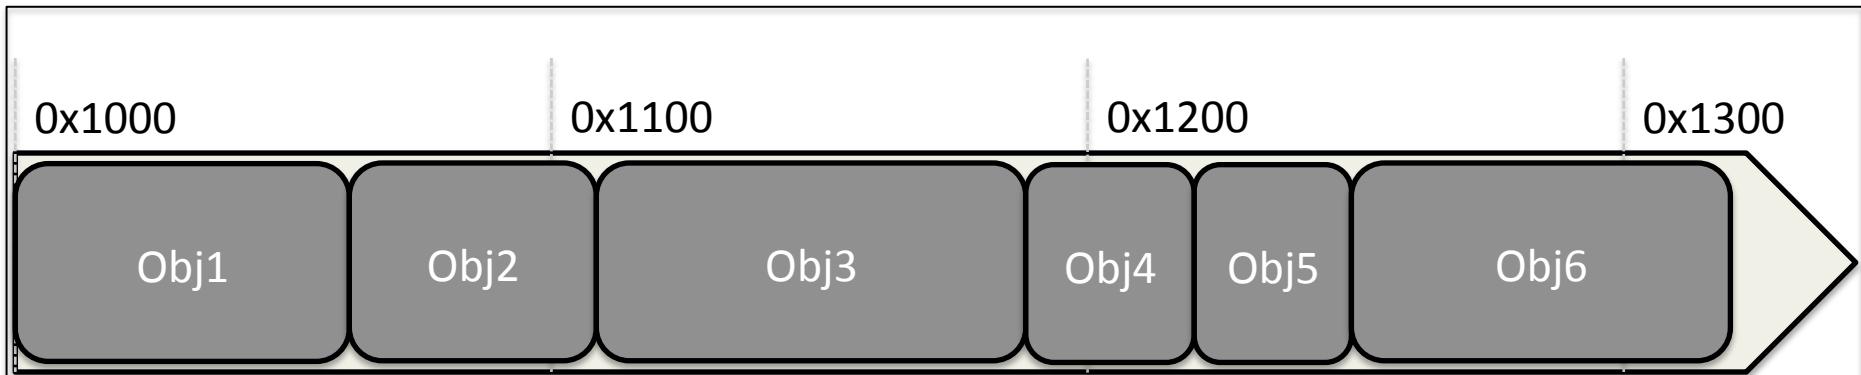


```
0:  iconst_1
1:  istore_1
2:  iload_1
3:  iload_0
4:  invokevirtual #16 // Method mthd(I)V
7:  return
```

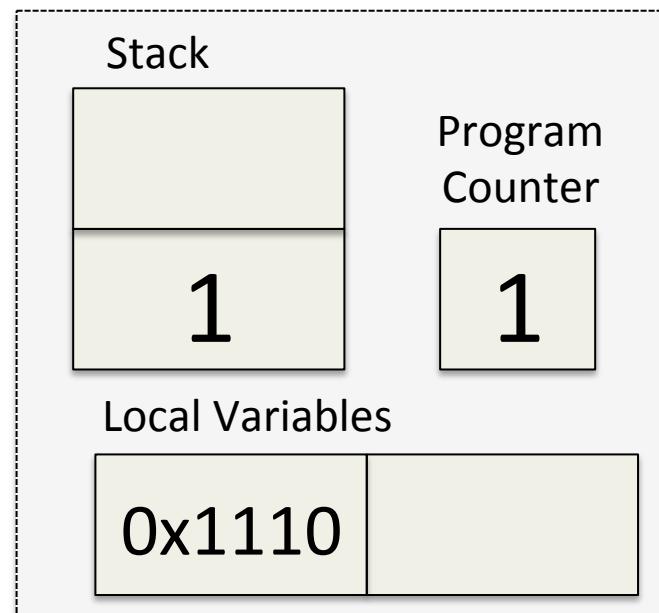
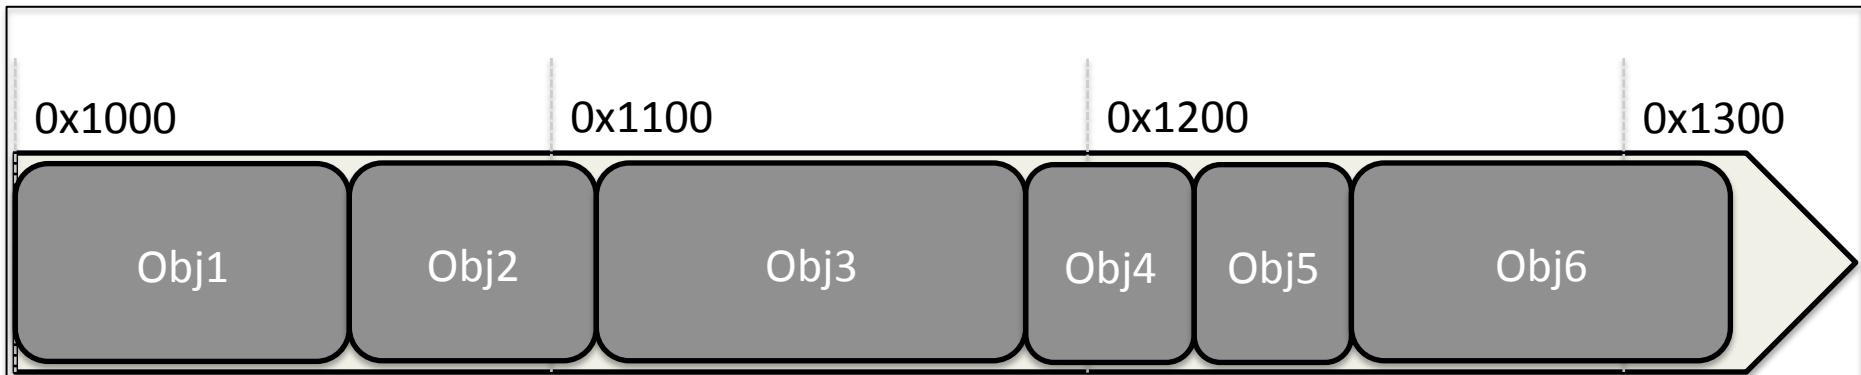


```
0:  iconst_1
1:  istore_1
2:  iload_1
3:  iload_0
4:  invokevirtual #16 // Method mthd(I)V
7:  return
```

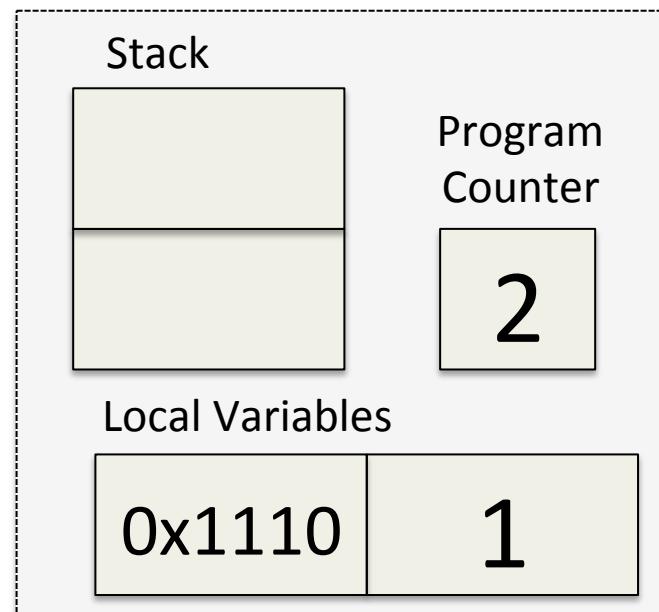
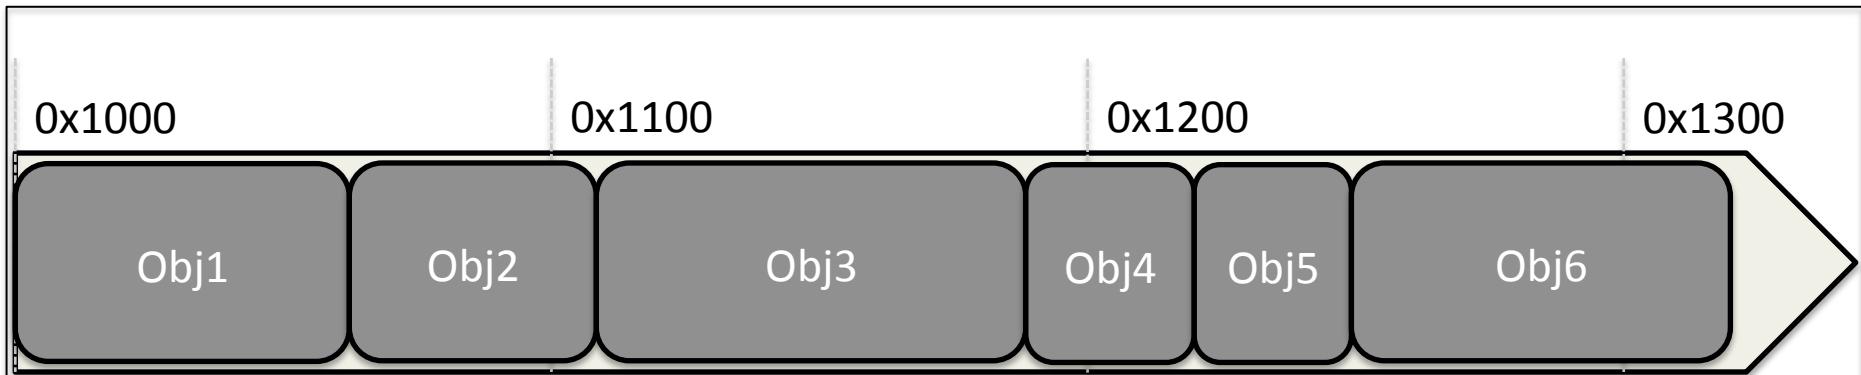




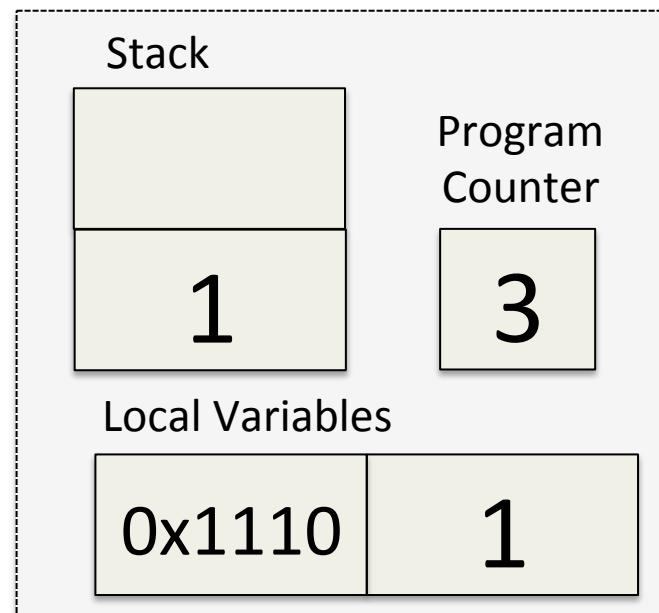
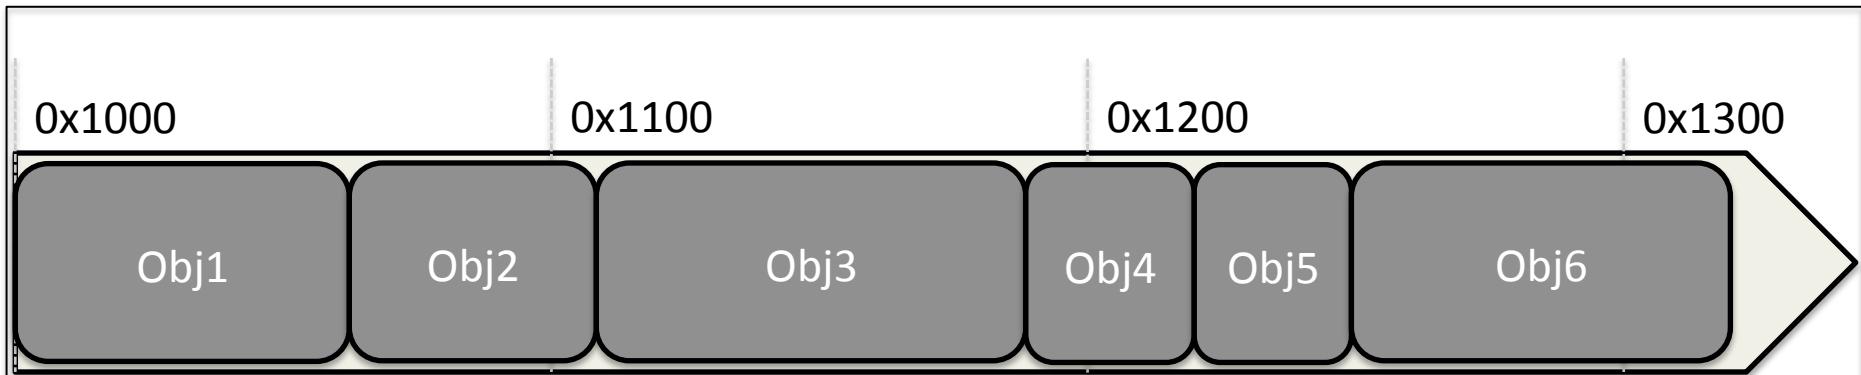
```
0: iconst_1
1: istore_1
2: iload_1
3: iload_0
4: invokevirtual #16 // Method mthd(I)V
7: return
```



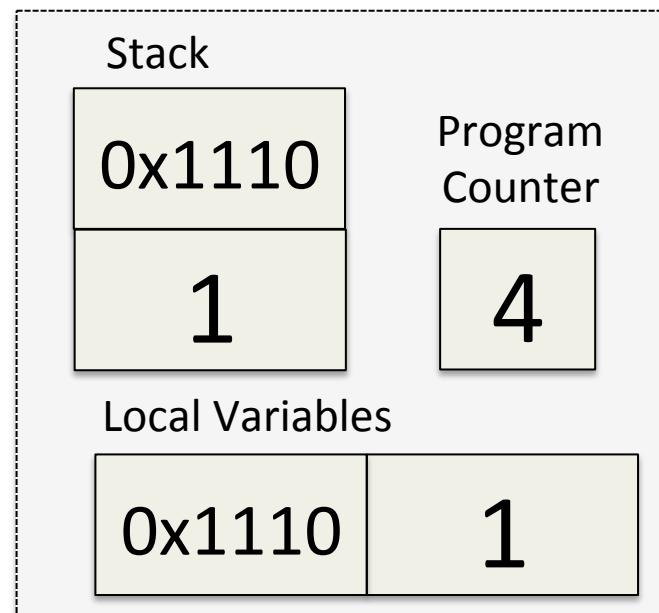
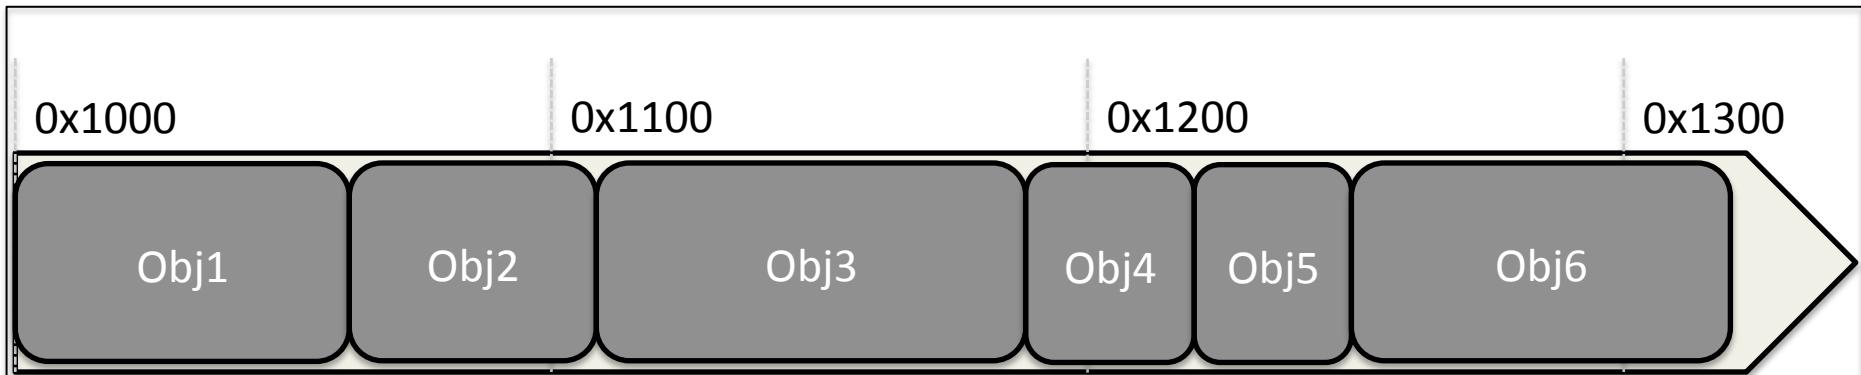
```
0: iconst 1
1: istore_1
2: iload_1
3: iload_0
4: invokevirtual #16 // Method mthd(I)V
7: return
```



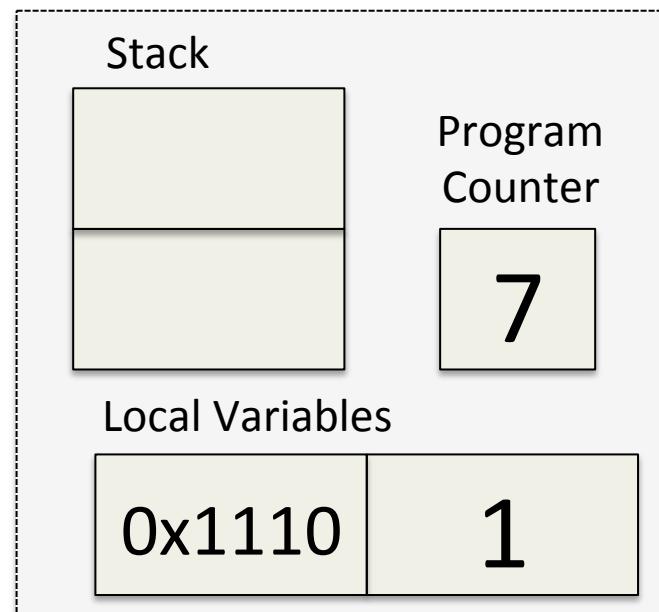
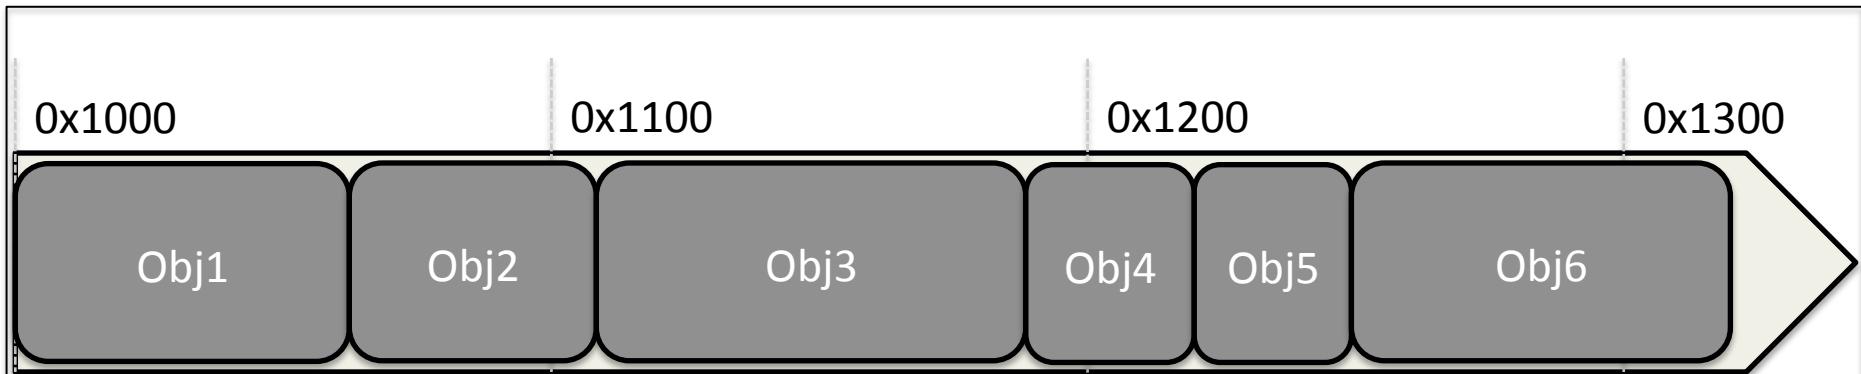
```
0:  iconst_1
1:  istore_1
2:  iload_1
3:  iload_0
4:  invokevirtual #16 // Method mthd(I)V
7:  return
```



```
0:  iconst_1
1:  istore_1
2:  iload_1
3:  iload_0 // The instruction at index 3 is highlighted in red.
4:  invokevirtual #16   // Method mthd(I)V
7:  return
```



```
0:  iconst_1
1:  istore_1
2:  iload_1
3:  iload_0
4:  invokevirtual #16 // Method mthd(I)V
7:  return
```



```
0:  iconst_1
1:  istore_1
2:  iload_1
3:  iload_0
4:  invokevirtual #16 // Method mthd(I)V
7:  return
```

References in Java

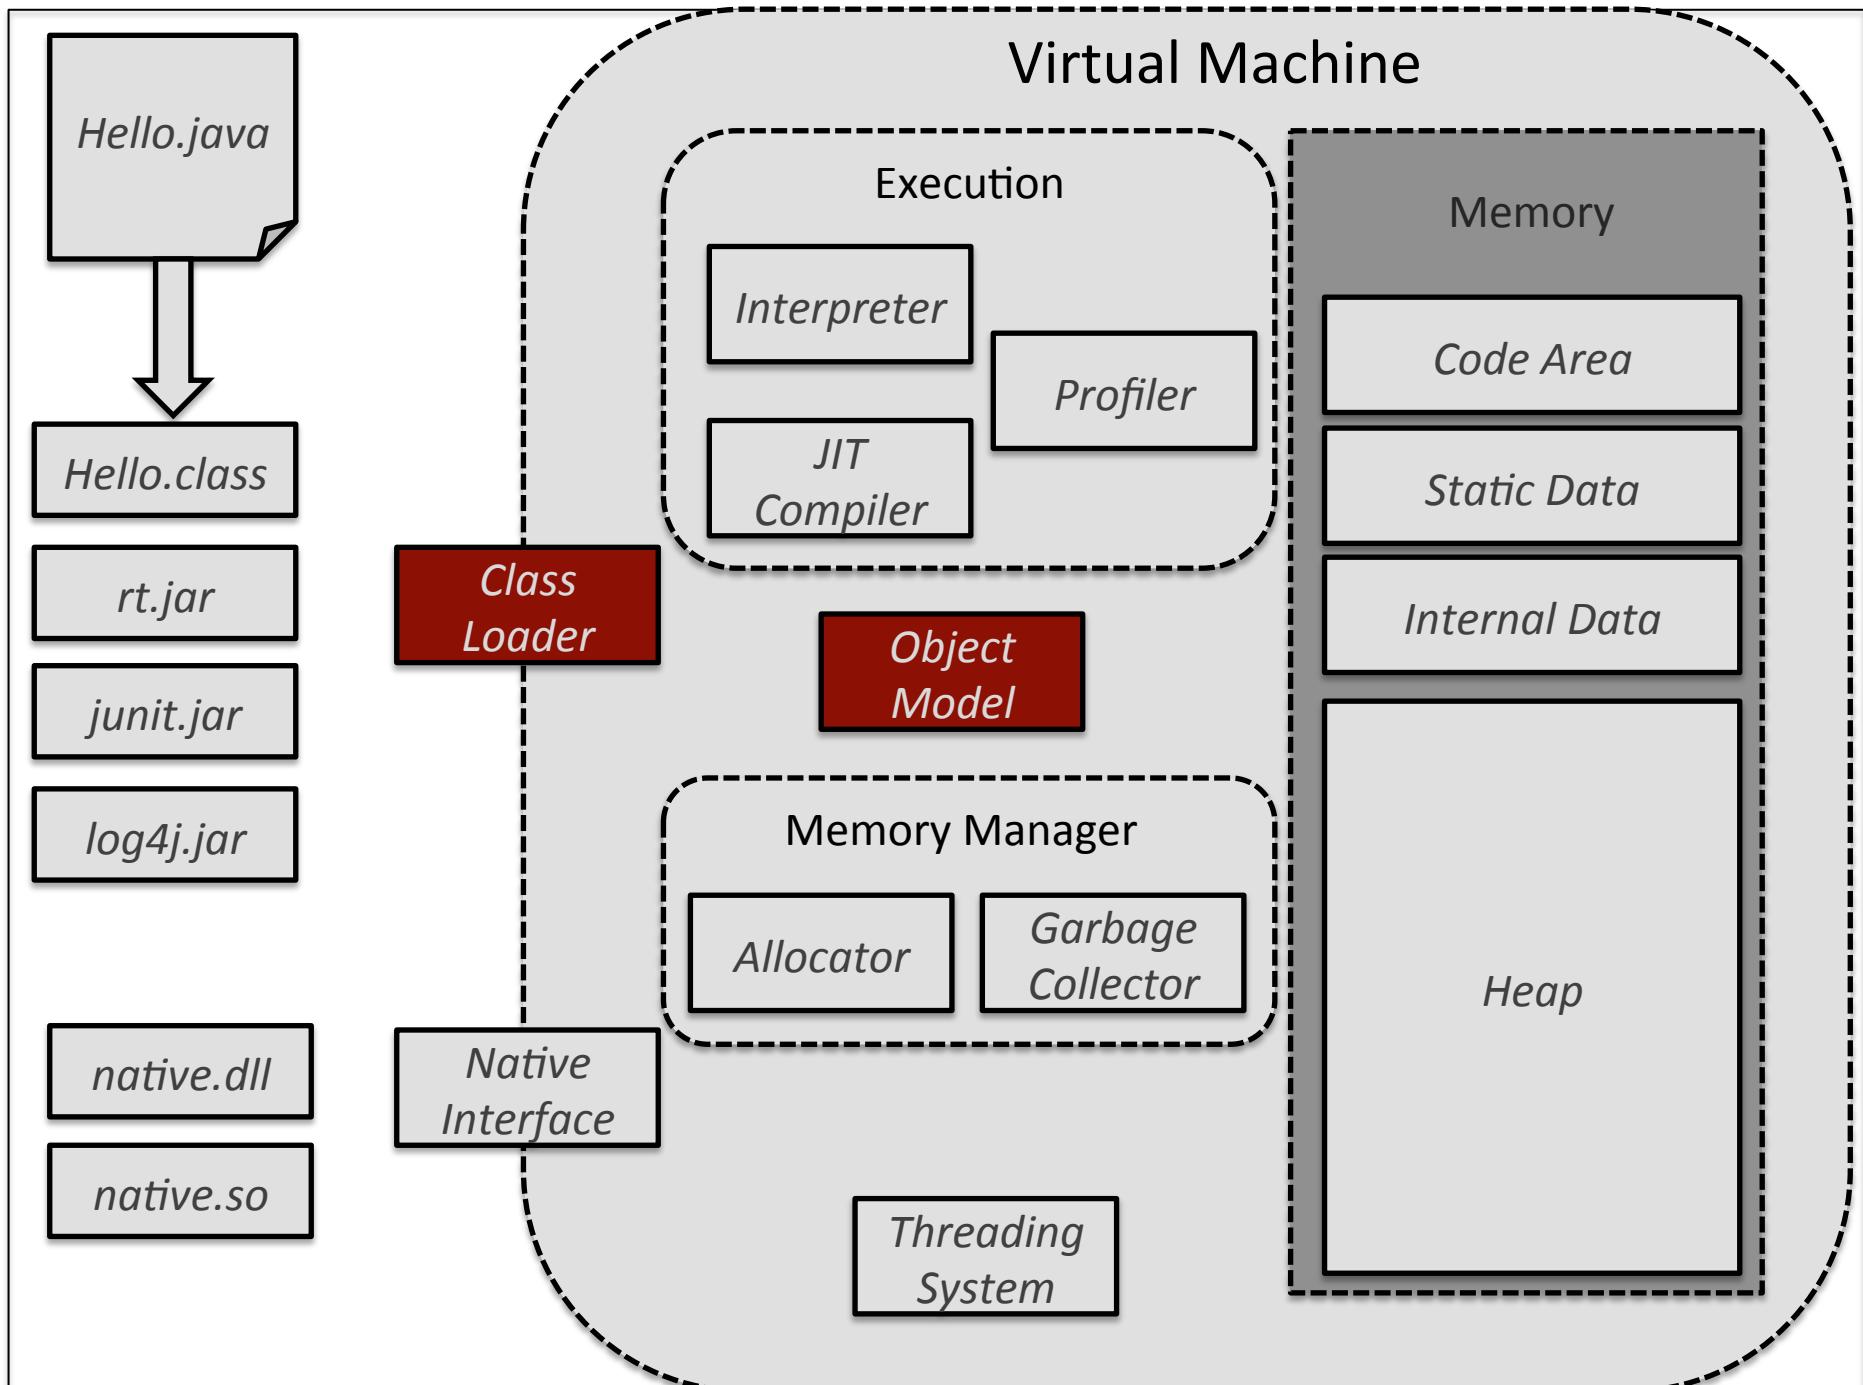
- Java hides references from developers
 - All a developer knows is that they have an object
 - As VM developers, we need two different concepts
- References can be treated similarly to primitives
 - Exist on the stack and local variables
 - Category 1 value
 - They can be stored in object fields
 - They are passed by value during method calls

Pointers vs References

- References
 - Guaranteed to point to a valid heap object
 - Dereference a reference to get an object
- Pointers
 - Arbitrary location in the heap
 - May point to the middle of an object
 - Not possible in Java

Arrays

- Largely similar to Objects
 - Stored on the heap
 - Mutable
- Some differences
 - All fields are the same type
 - VM needs to track array size
 - Accessed using different bytecodes



Object Representation in C

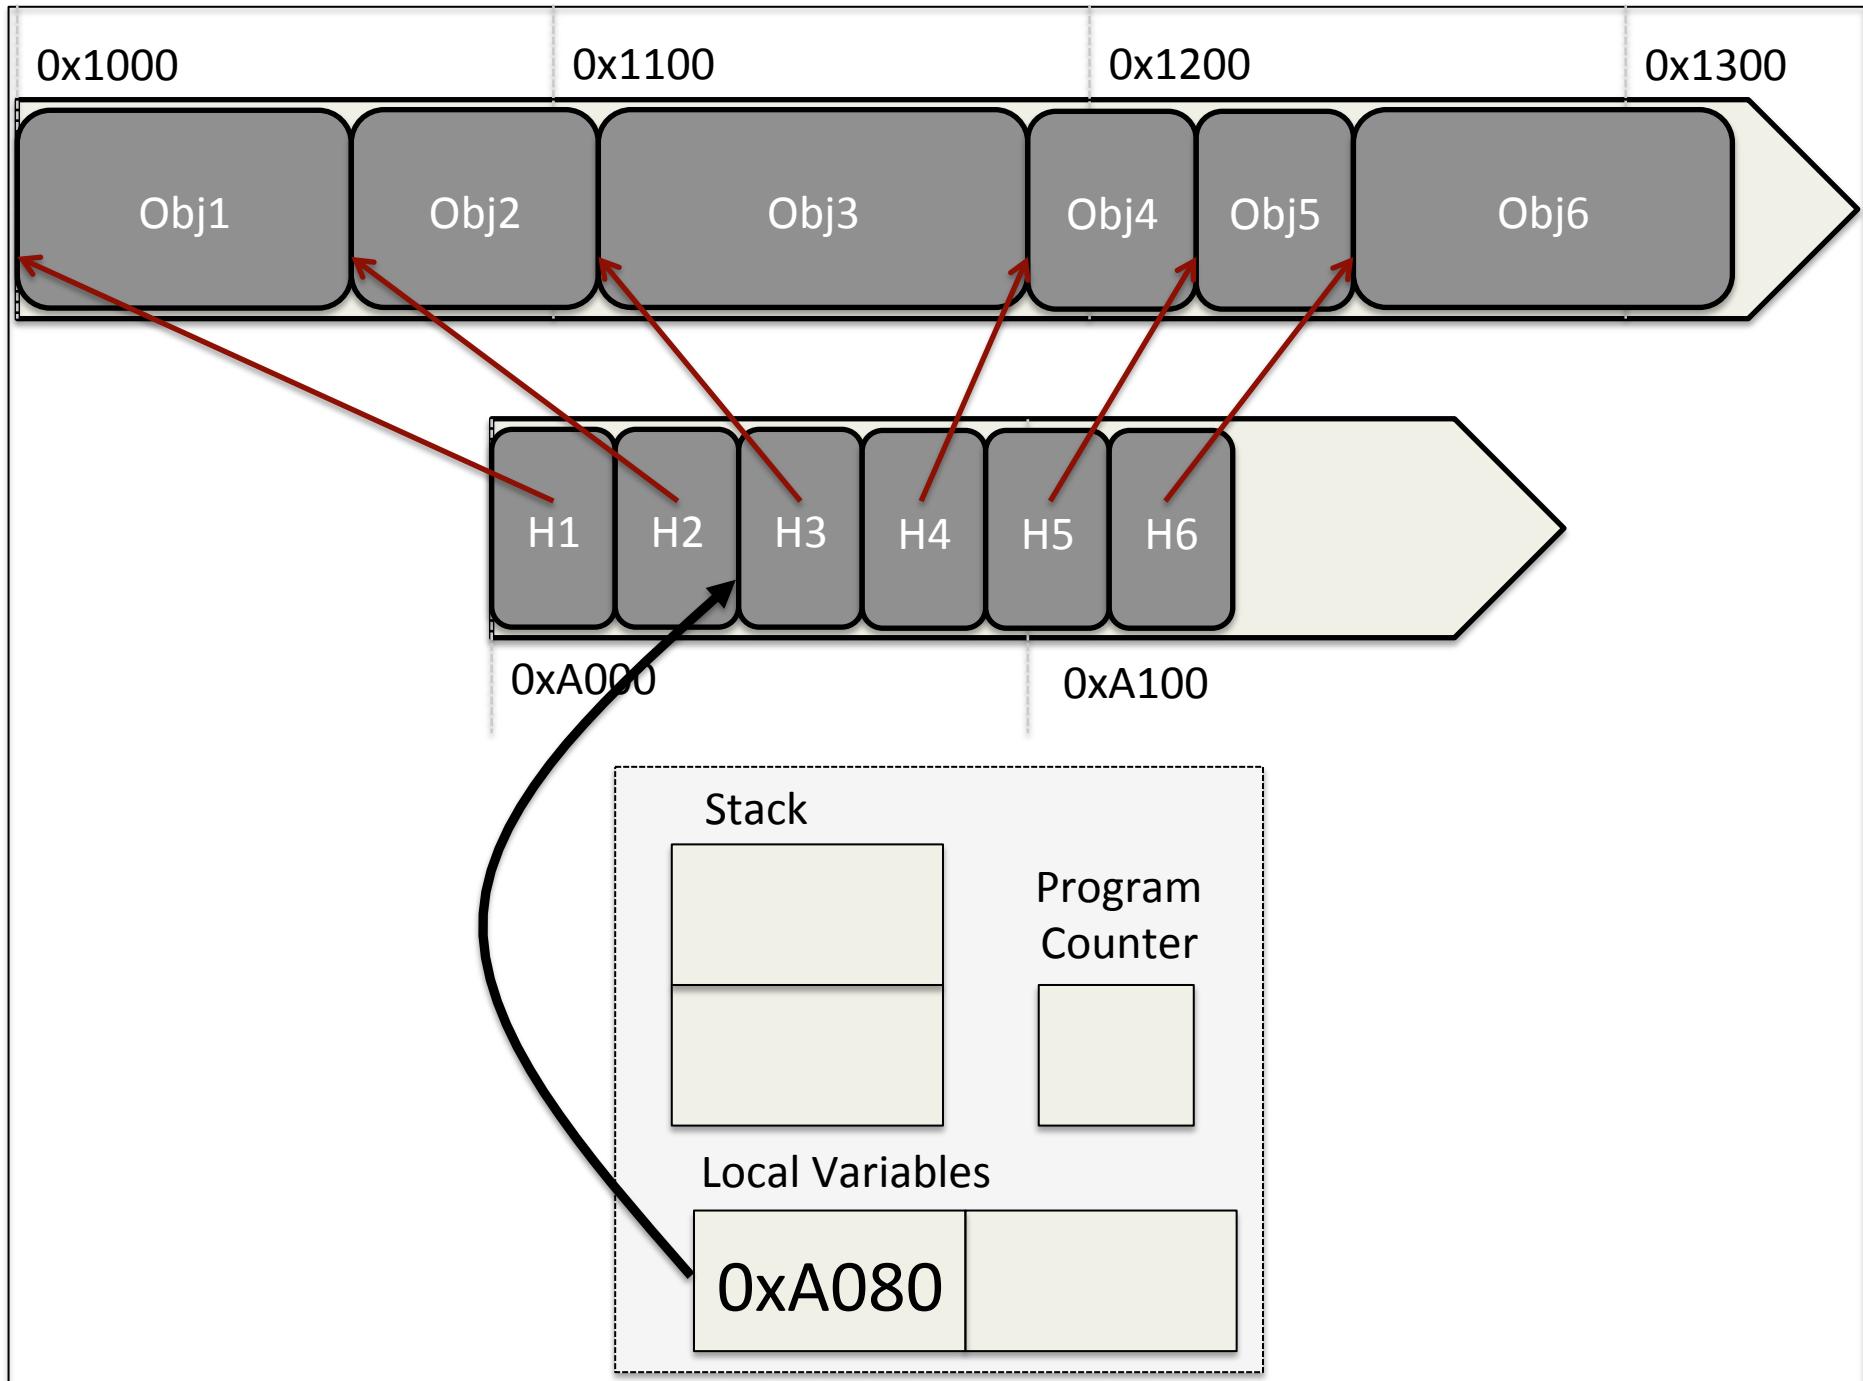
- Bits is bits
 - Object is a series of bits on the heap
- The meaning of the object comes from its use
 - Unions
 - Casting char array as a struct
- Developer or compiler tracks types

Managed Objects

- Runtime needs to know more about the object
 - Type of object
 - Class hierarchy for dynamic dispatch
 - Layout of fields, location of references
 - Lock state
 - GC state
 - Hash code
- Two kinds of information
 - Per-Class
 - Per-Object

Handles

- Object references point to metadata block
 - Metadata is stored off to the side
 - Keeps track of where the object is on the heap

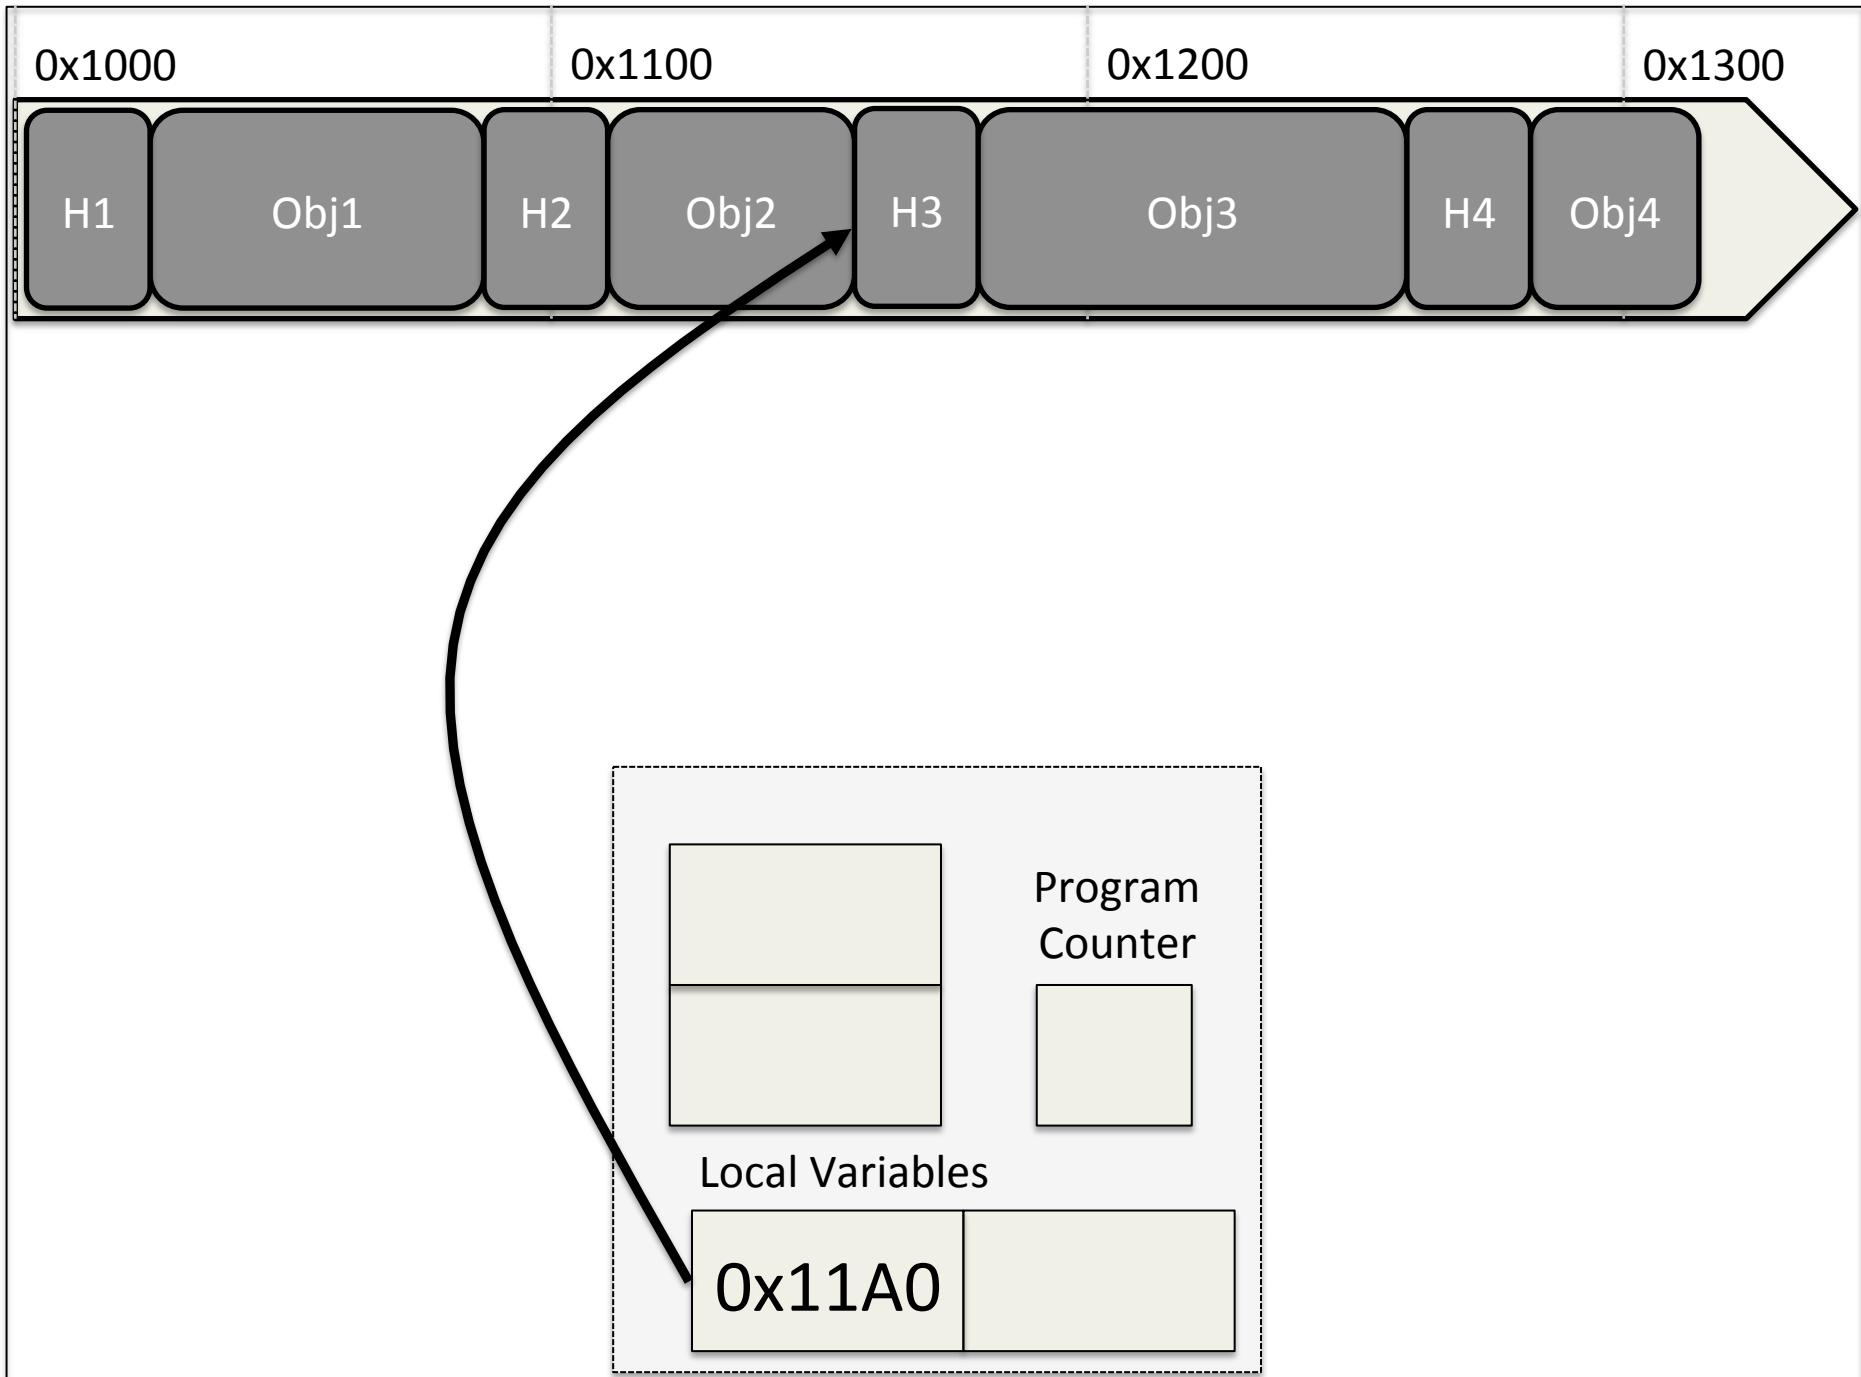


Handles

- Object references point to metadata block
 - Metadata is stored off to the side
 - Keeps track of where the object is on the heap
- All object dereferences go through indirection
 - Two memory references instead of one
- Handle locations independent of heap layout
 - Useful for copying garbage collection

Object Headers

- Every object has a header block



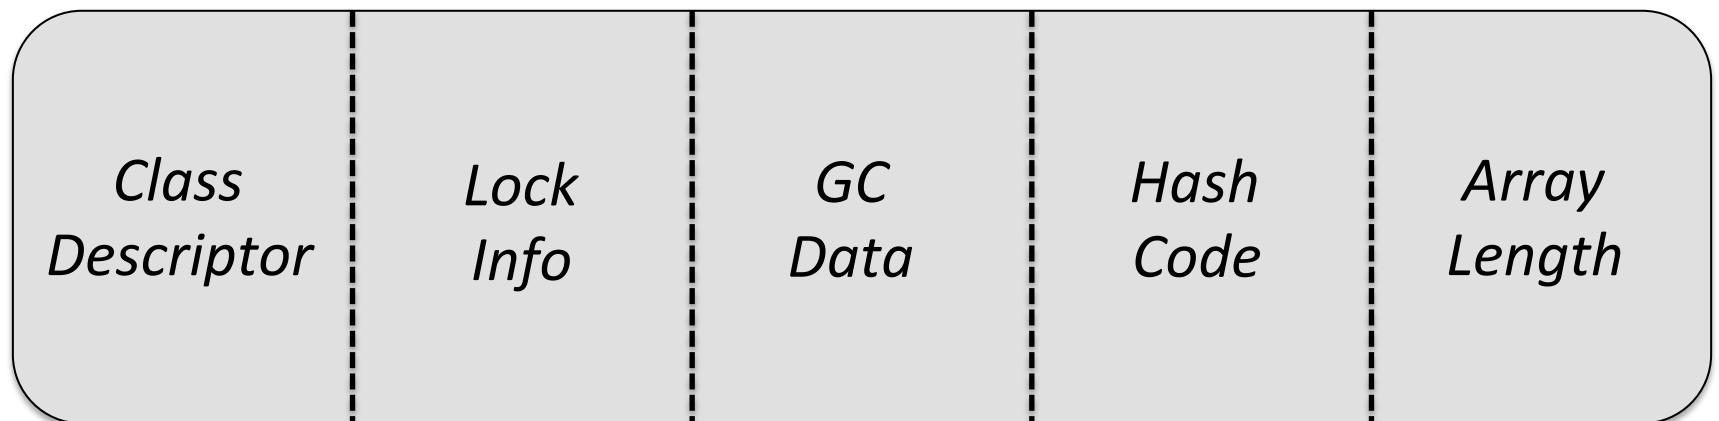
Object Headers

- Every object has a header block
- Keep metadata close to the object
 - Only one access per dereference
 - Makes better use of the cache
- Takes up more space on the heap

Header Format

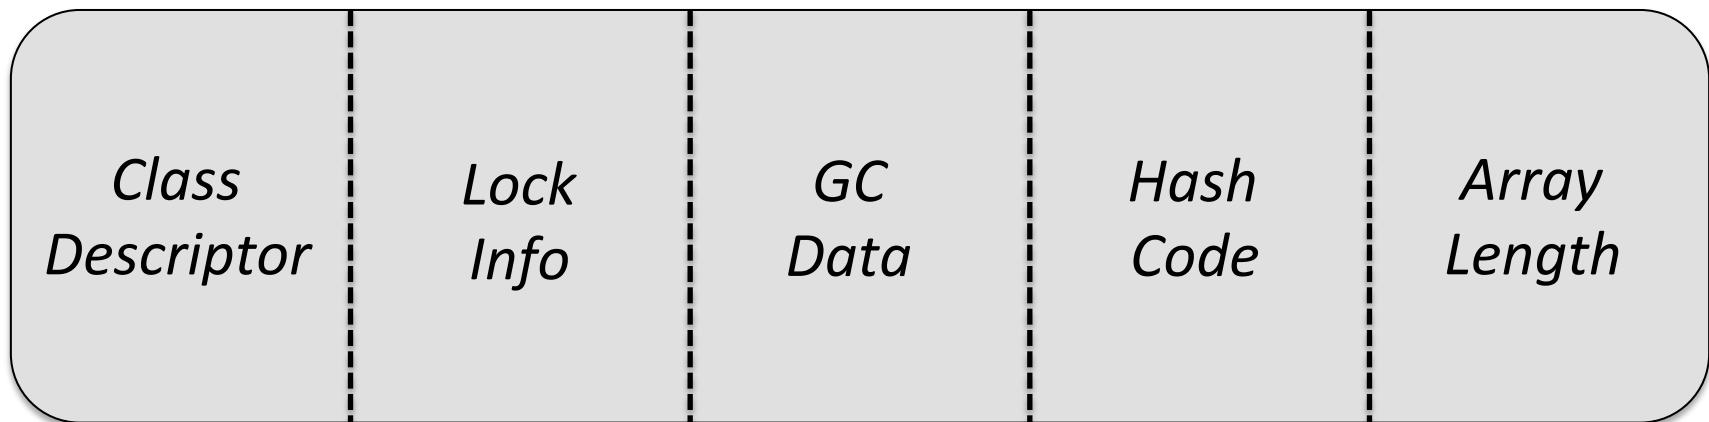
- Fixed header format
 - Implementation specific
 - Simplifies generated access code
- Slight difference between arrays and objects
 - Array length field
 - Different bytecodes to handle each

Header Data



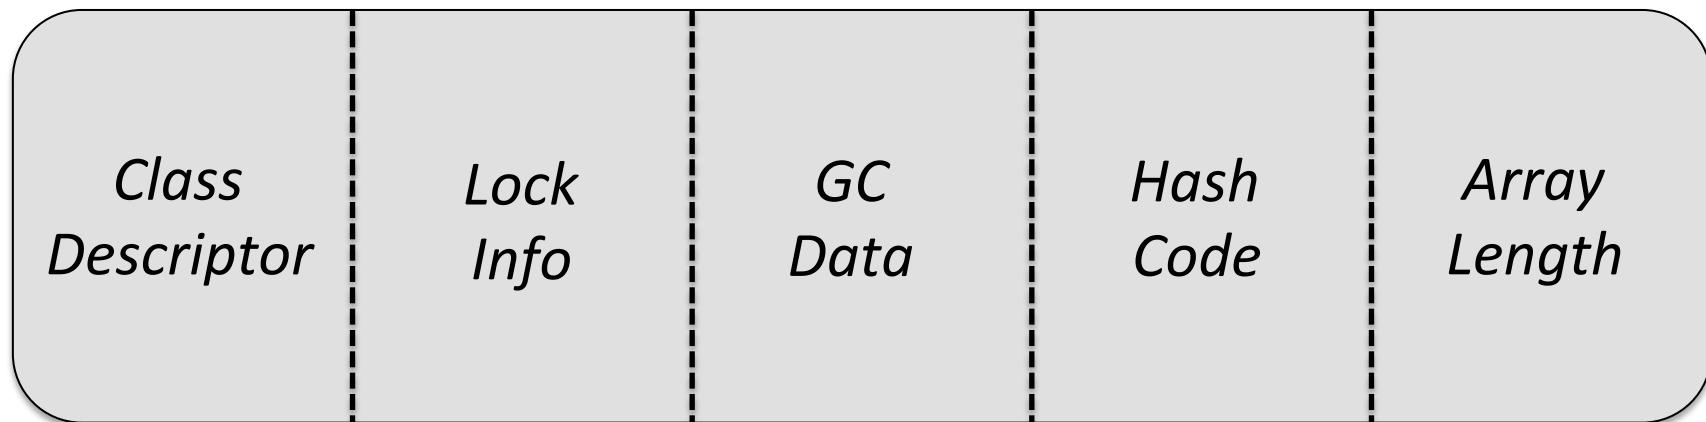
Class Descriptor

- One instance per class
- Identifies the runtime type of the object
 - Includes the class hierarchy
- Gives access to static data



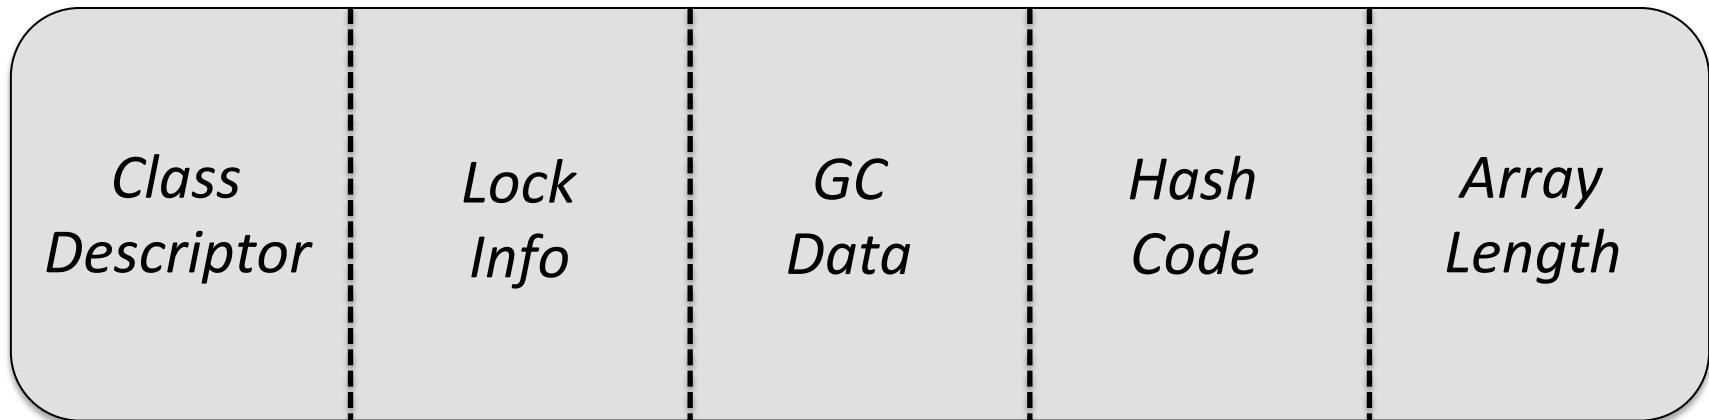
Lock Word

- Any object can be used as a monitor
- Java locks are re-entrant
 - You can take the a lock twice with the same thread
 - We'll talk more about this later in the semester



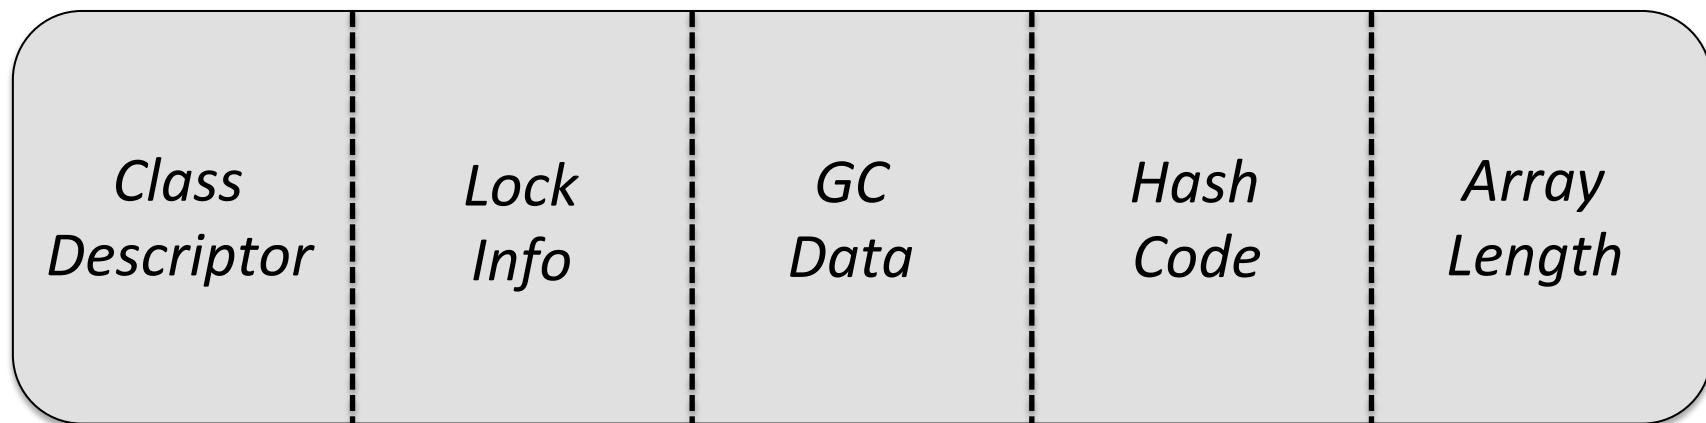
Garbage Collection Data

- Most GC algorithms need some per-object data
 - Mark bits
 - Forwarding pointers
 - We'll talk a lot more about this later



Hash Code

- Any object can be hashed
 - Used when stored in maps or sets
- Need to track hash code when assigned
 - Can't just use the object address



The SimpleJava Header

- SimpleJava has a two-word header
 - Single header structure for objects or arrays
 - No locking or hashing
- Both words contain `java.lang.Object` type
 - Word 0 is the class descriptor
 - Word 1 is the array size
- `Object` type lets us overwrite words during GC
- See `ObjectHeader` interface for more details

Header Compression

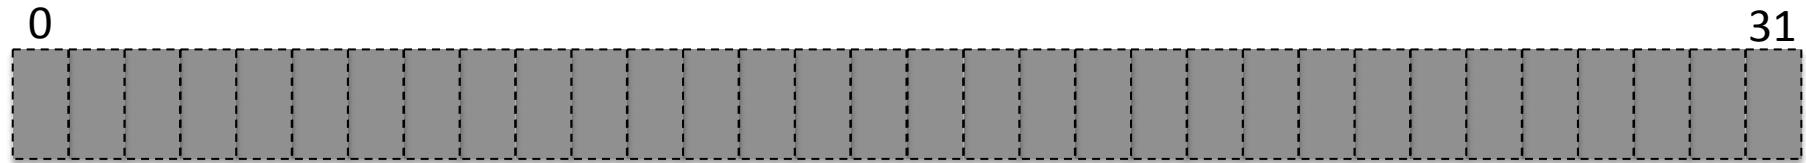
- Five words per object is excessive
 - More header than object data
- A few observations let us reduce the overhead
 - Most objects aren't locked or hashed
 - Most of the time the GC isn't running
- Objects can be aligned on byte boundaries

Addressing

- Hardware is generally word-addressed
 - 32-bit memory can address 4GB of memory
- The VM controls allocation
 - It can force objects to align on particular boundaries

Addressing

- Word aligned
 - Requires all 32 bits to address



Addressing

- Word aligned
 - Requires all 32 bits to address
- Double word aligned
 - We know that the least significant bit will be a zero



Addressing

- Word aligned
 - Requires all 32 bits to address
- Double word aligned
 - We know that the least significant bit will be a zero
- Four-word aligned
 - The two least significant bits will always be zero



Bit Stealing

- We can align on any power of two value
 - Guarantees a number of bits to be zero
- Those bits can then be re-used for other things
 - Known as bit stealing
- Need to mask the stolen bits on read
 - Normally a cheap operation
 - Sometimes done in hardware

Compressed Headers

- Hotspot uses two words per object
 - Three words for arrays
 - Could go smaller, but with higher overhead
- One word is reserved for the class descriptor
 - Frequently used
- Second word is the mark word
 - Encodes all the other information

Mark Word

- Bottom two or three bits are used as a tag
 - Two bits in most cases, third for a corner case
 - Remainder of the word determined by their value
- Five separate states
 - Unlocked
 - Lightweight locked
 - Heavyweight locked
 - Marked by GC
 - Biased Locks
- We'll talk about each state later in the semester

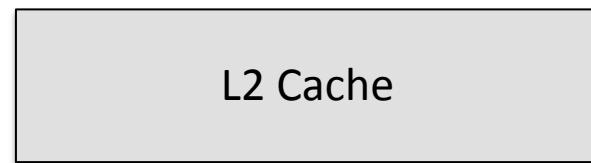
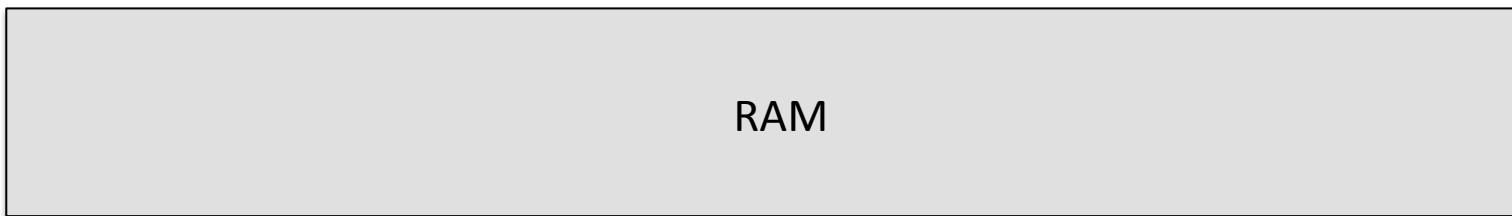
Compression Trade-offs

- We can save a lot of space
- A lot depends on the assumptions
 - We win if uncommon cases don't arise
 - If they do, we have extra memory accesses
- Code sequences can be complicated
 - Instructions have to check the header's mode

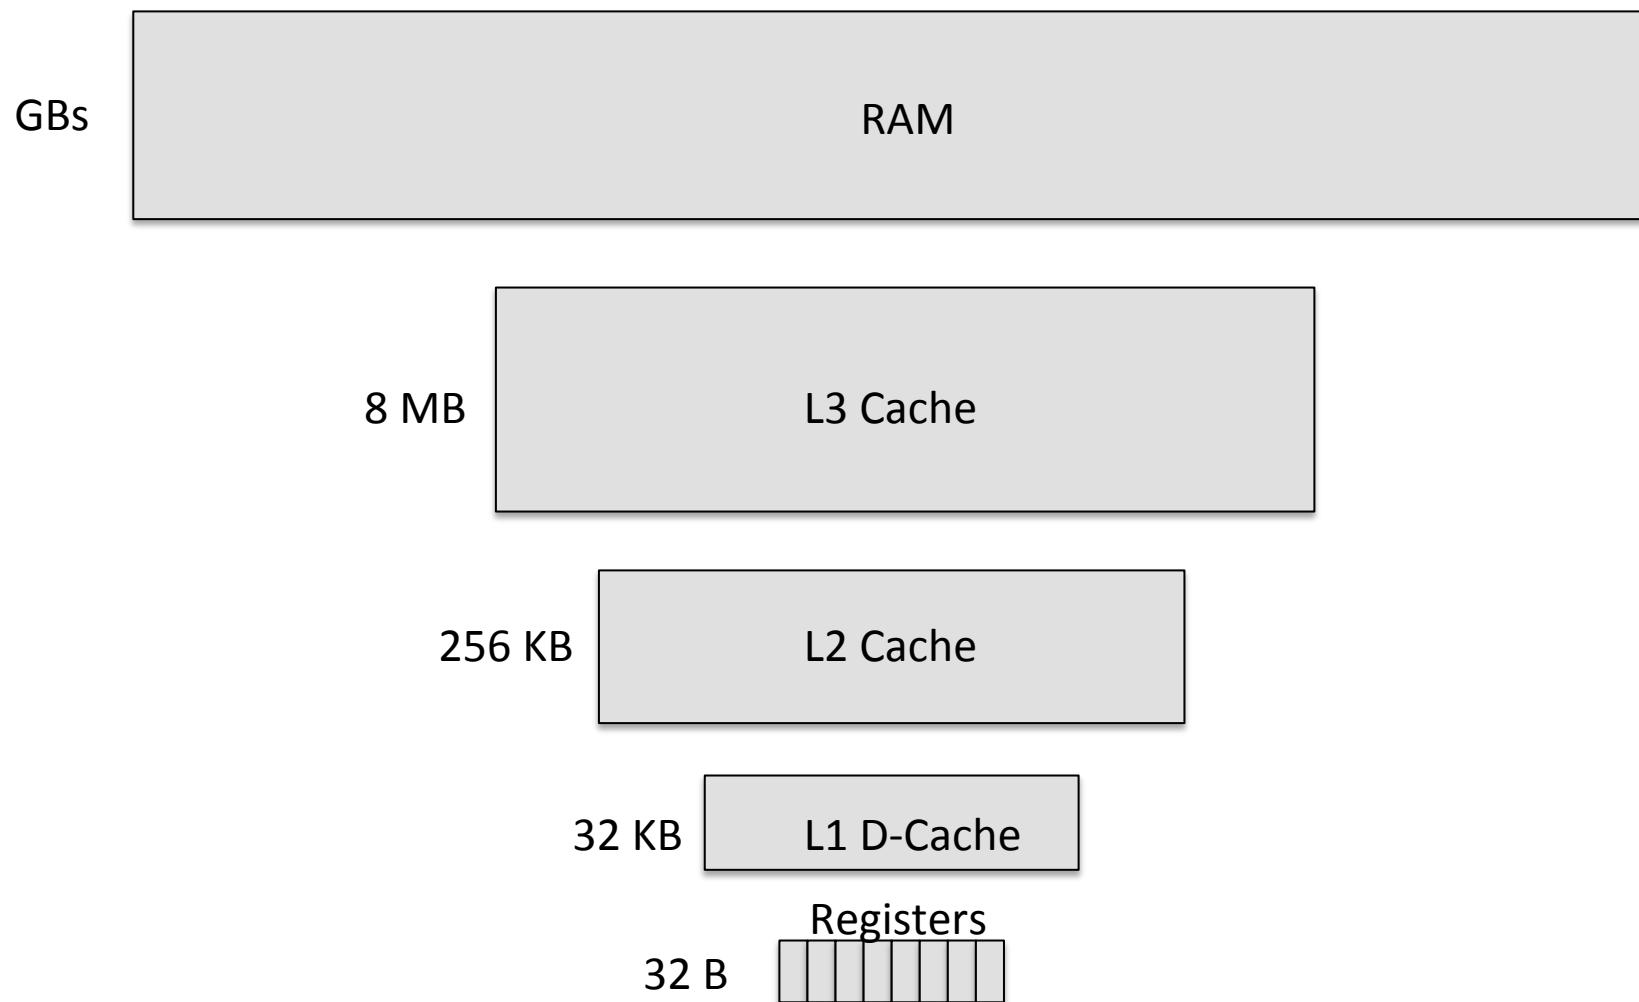
Cache Digression

- Why do we care about large headers?
 - Memory is cheap
- Larger objects require more garbage collection
 - GC can have a per-byte overhead
 - Large objects harder to fit in a fragmented heap
- Large objects take up more space in the cache

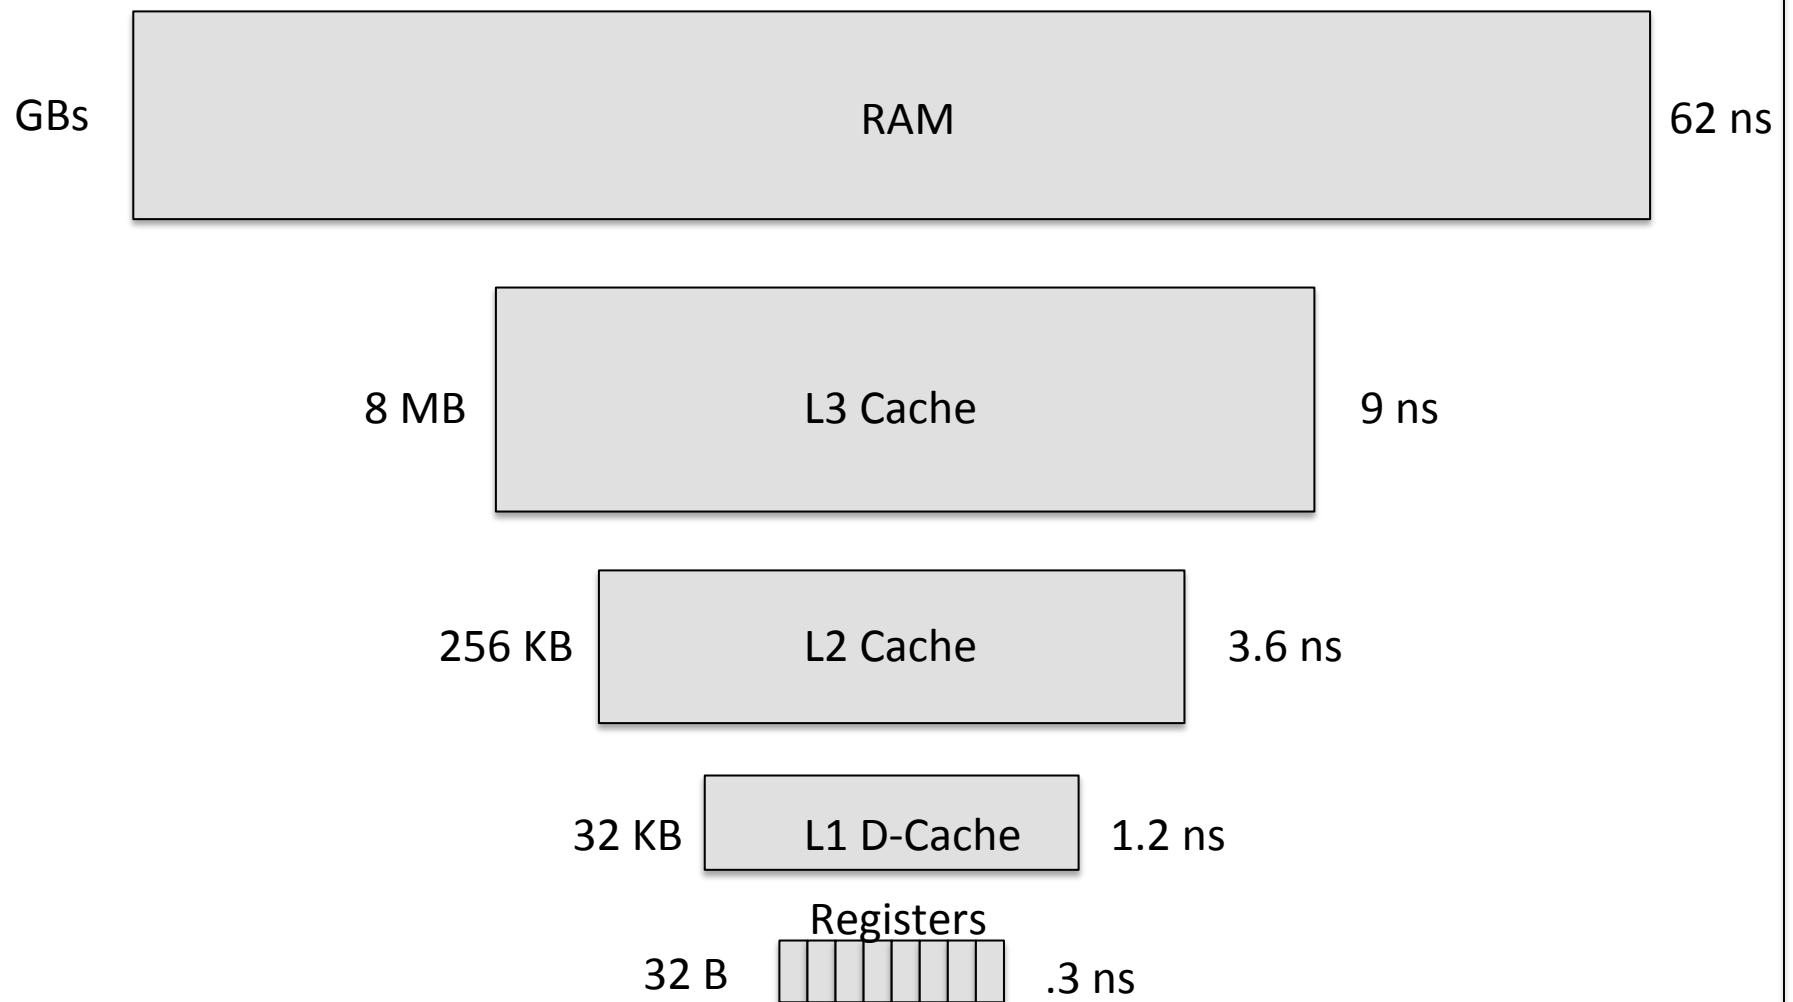
Intel I-7 Memory Hierarchy



Intel I-7 Memory Hierarchy



Intel I-7 Memory Hierarchy



Caching 101

- Caches are made up of lines
 - 64 bytes per line, in Intel Ivy Bridge L1 D-Cache
- Main memory access is by cache line
 - You don't just fetch a single object
 - Recall that objects are word-aligned
- The cache line includes header + data
 - Smaller header means more data
 - More data means fewer memory accesses

Caching 101

- Caching will be increasingly important to us
 - A big part of memory management
 - Implications on JIT compilation strategies
- We'll be covering it on-demand
- Caching is a fascinating topic
 - Its effects can be massive and counter-intuitive
 - Consider taking CSCI E-93

Object Layout

- Objects contain non-static field data
 - Values can be different for every instance
- Includes all fields defined by super classes
 - Remember that fields do not override
- Method data is not stored with the object
 - Loaded once per class

Field Layout

- Fields start at a fixed offset from the header
 - Remember that the header is a known size
- Object header can be at the end of the object
 - Fields laid out backwards
 - Pushes null pointer checks into hardware
- Same effect by guarding first few memory pages

Field Layout

- Layout strategy is implementation dependent
 - Ordered by defining class
 - Ordered by definition in the class file
 - Ordered by predicted access patterns
 - Ordered by type, keeping references together
- All have different cache characteristics