

In the last two weeks

- Assignment 2 was due yesterday
 - Feedback survey now posted
- Reading on conservative GC
- Questions on allocation algorithms
 - How would you choose which to use

Assignment 2

- More conceptual than code intensive
 - Deeper dive into some of the issues from class
 - More open in terms of implementation
- Quite a lot of questions and comments
 - Not many on the same topic
- A couple of bugs showed up in SimpleJava
 - Inline code reversed method parameters
 - Leaky abstraction in the profiler interface

Conservative GC

- Managed languages have precise type information
 - Many other languages don't
 - Can add a conservative GC without compiler support
- Conservative GC works around imprecise types
 - If it looks like a pointer, treat it as one
 - Can lead to false retention in some cases

Garbage Collecting C++

- Useful when the language doesn't support GC
 - Some comparisons with reference counting
 - Benefits in large, multi-person projects
- Bohem-Demers-Weiser collector
 - Allows conservative collection in C/C++
 - How watertight is the abstraction?
- How expensive is type accuracy?

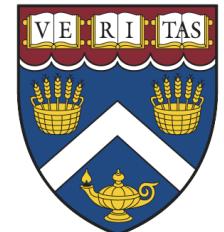
False Retention Rates

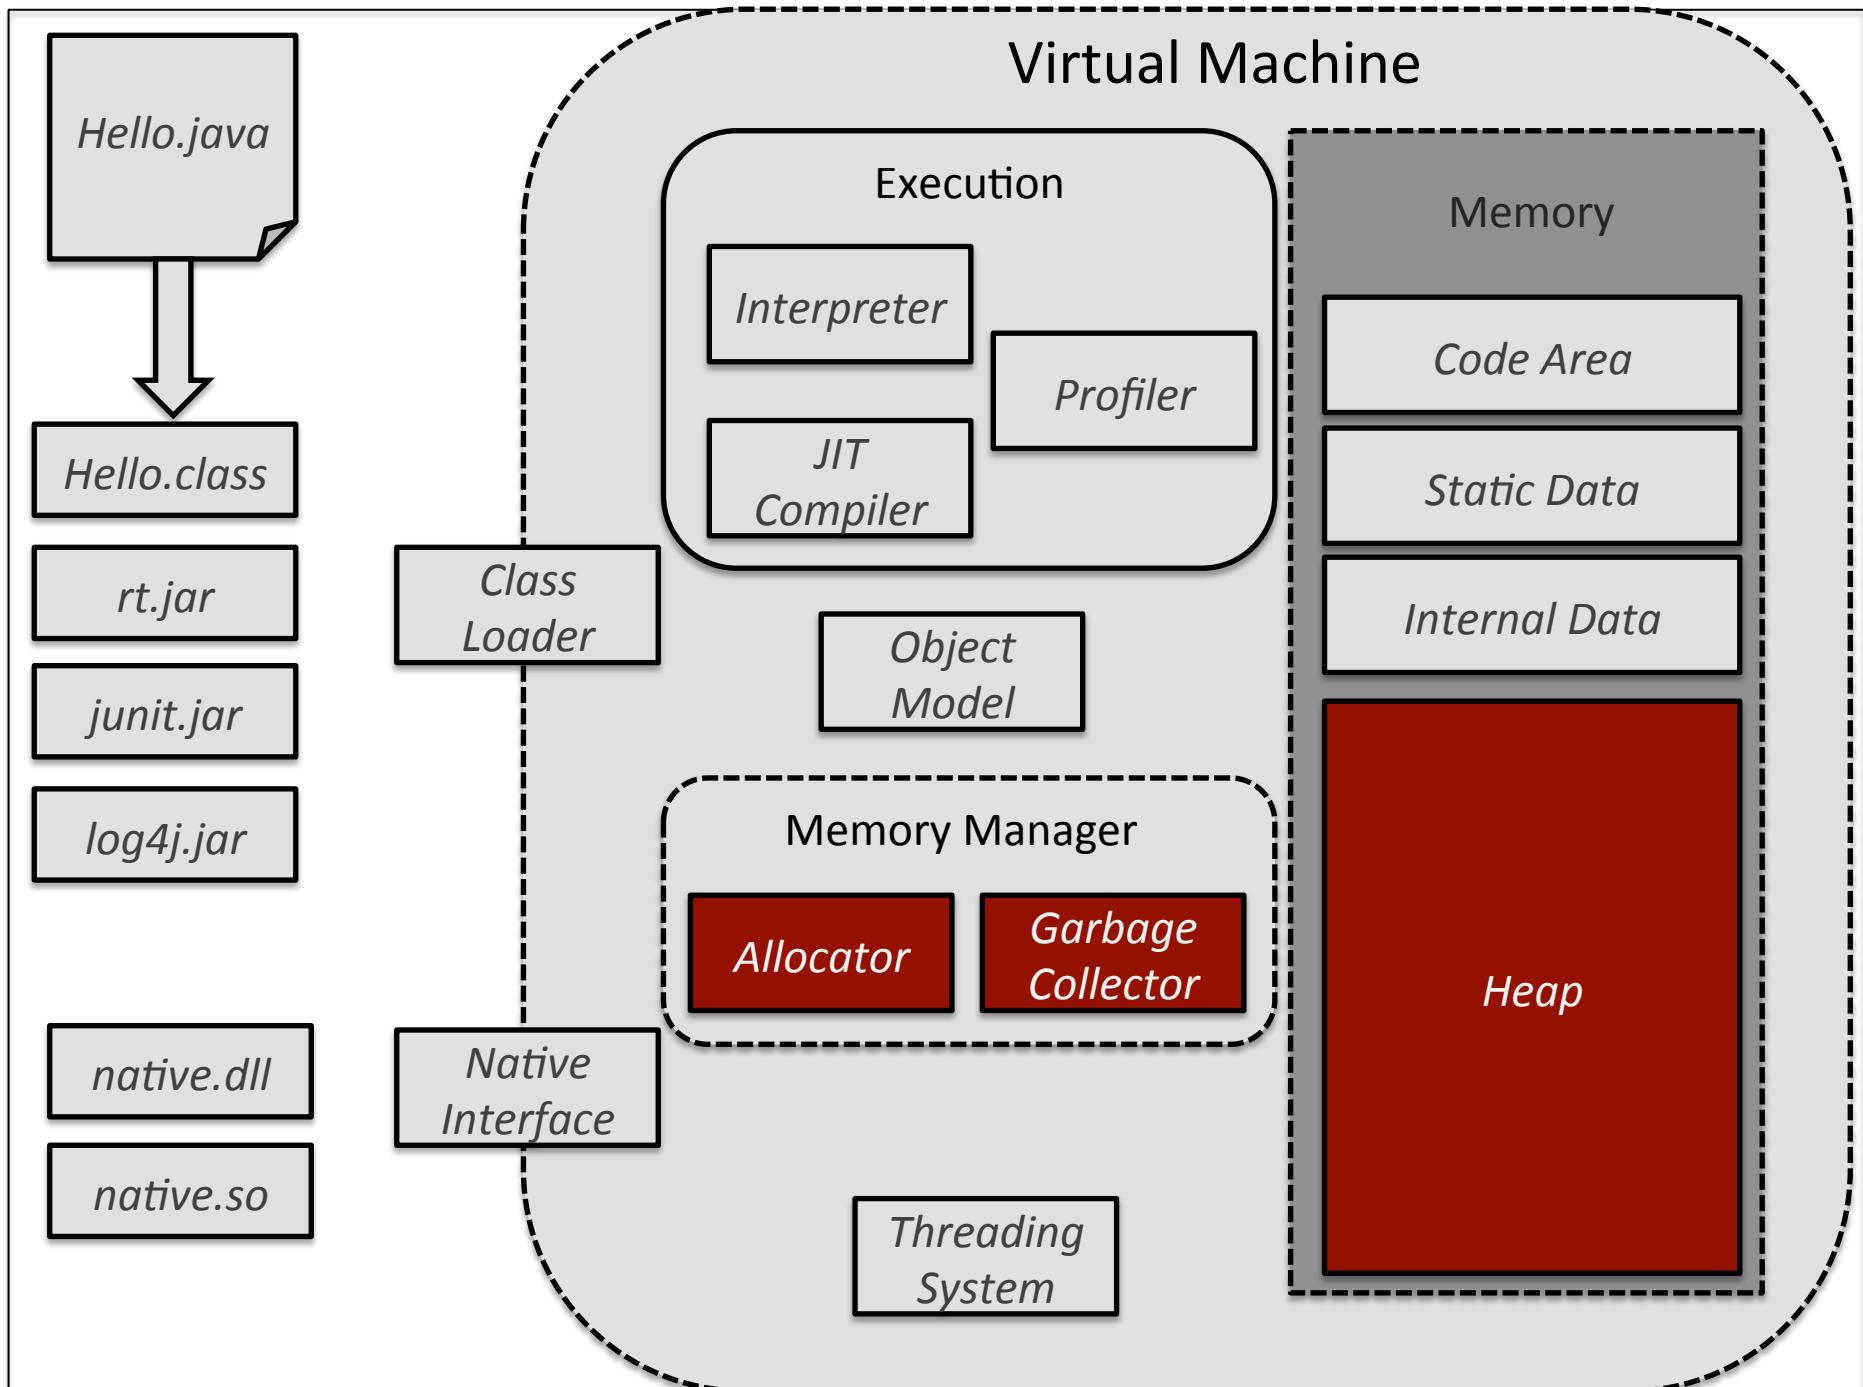
- Excess retention very low
 - On the order of 0.01% in Jikes RVM
- Does that add up over time?
- May still have to do some maintenance
 - Not good for a server environment

Target Applications

- Conservative GC good for small programs
 - Flash and Javascript programs tend to be small
 - Memory leaks don't have time to become serious
 - Is that still the case?
- The same languages tend to be user-facing
 - Pause times are important
 - Argument for reference counting?

Copying Garbage Collectors





Garbage Collection

- Data used by the application stored on the heap
 - Also some VM internal structures
- Objects are garbage when they're no longer used
 - Approximate using reachability
- Garbage collector frees unreachable objects

Reference Counting

- Keep track of all references to an object
- Some nice features
 - Incremental
 - Simple to implement and understand
 - Doesn't require runtime support
- One big problem
 - Garbage cycles

Stop The World Collectors

- Pause all mutator threads
 - Do garbage collection work
 - Restart mutator threads
-
- Can take a big-picture view of the heap
 - Know that it won't change during collection
-
- Leads to large pause times
 - Useful when studying algorithms

Mark and Sweep Collection

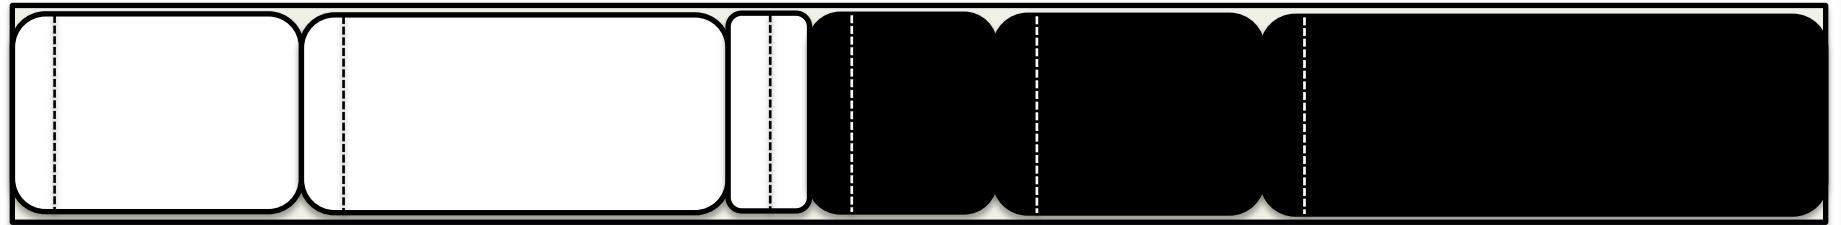
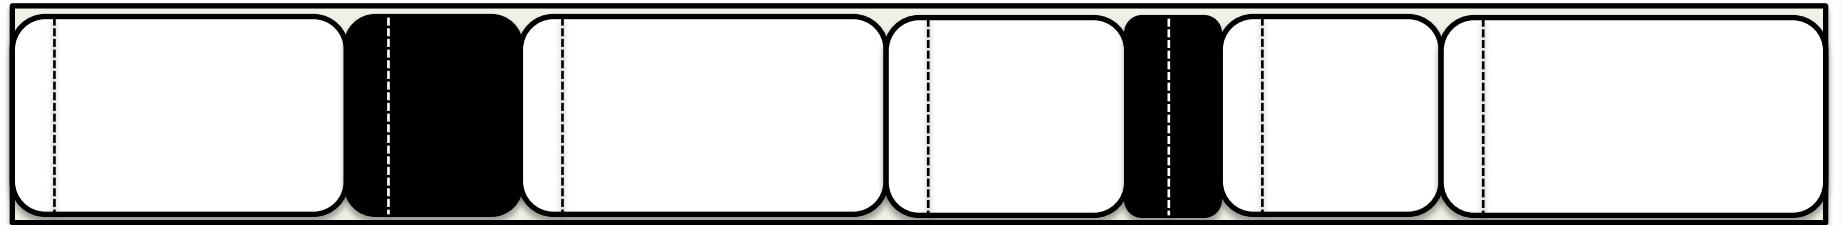
- Works out what objects are live
 - Traces the object graph from the system roots
 - Marks any objects that are reachable
- Scans linearly over the heap
- Returns unmarked objects to the free pool
- Introduces fragmentation to the heap

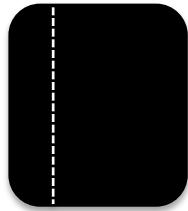
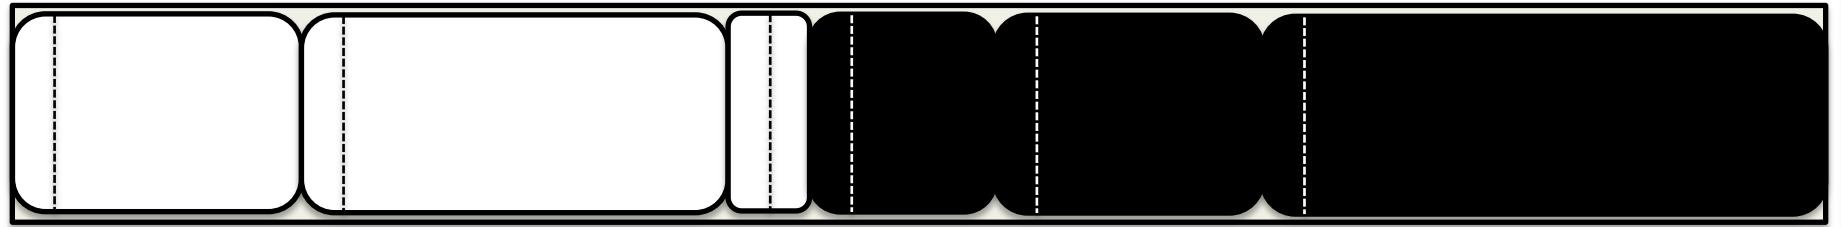
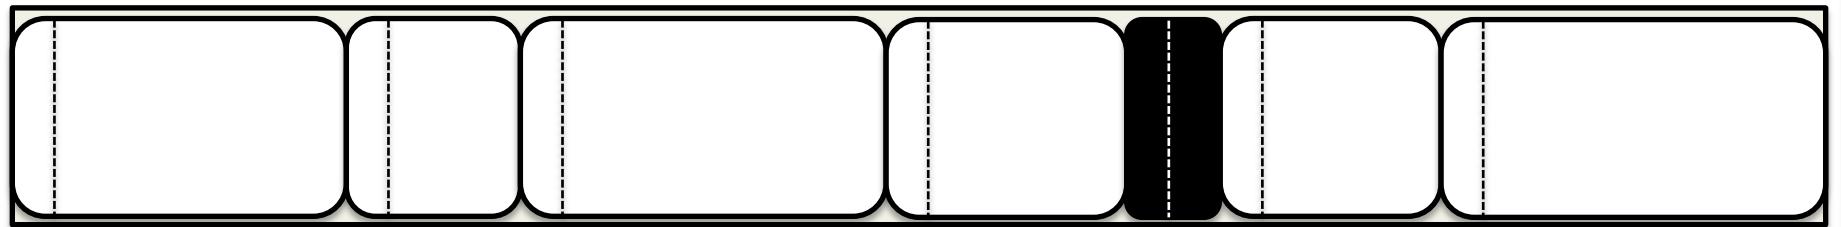
Fragmentation

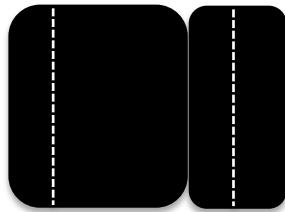
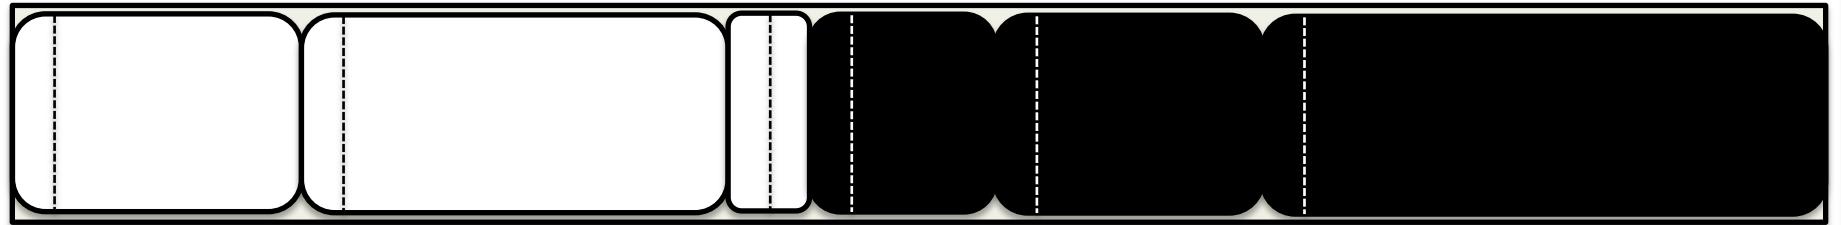
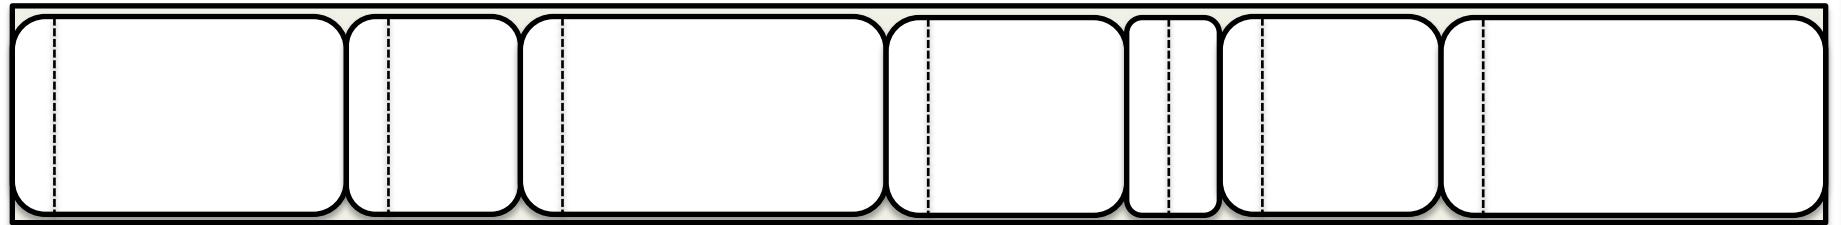
- Any collector that frees objects creates holes
 - Spaces where freed objects used to be
- Fragmentation can lead to wasted memory
 - There may not be an object of the correct size
- Allocation to fragmented memory is expensive
 - Need to figure out where to put objects
 - May not have a perfect fit

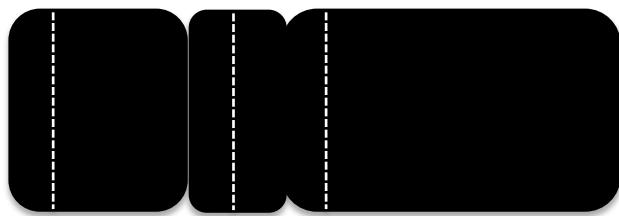
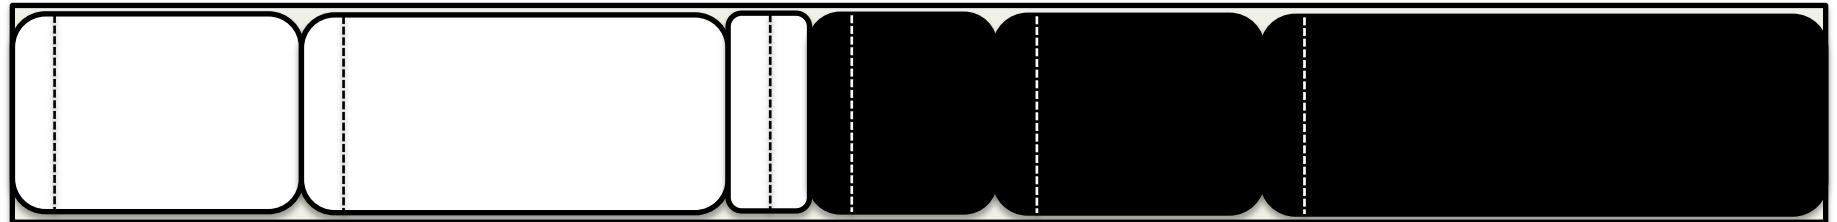
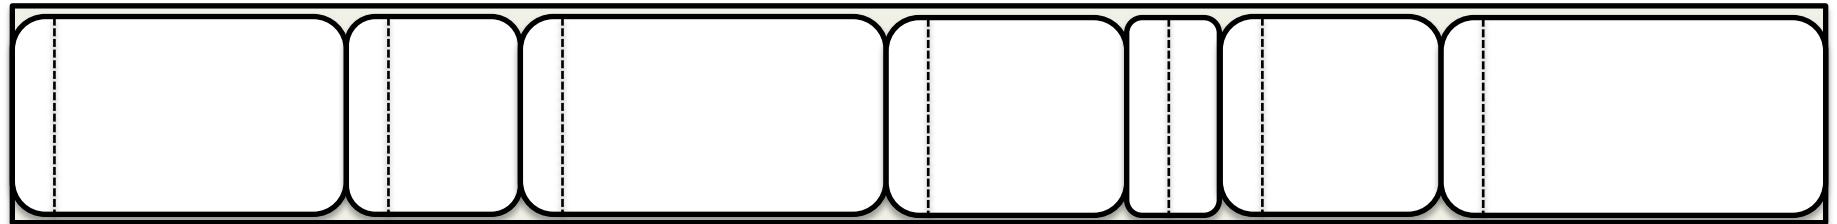
Eliminating Fragmentation

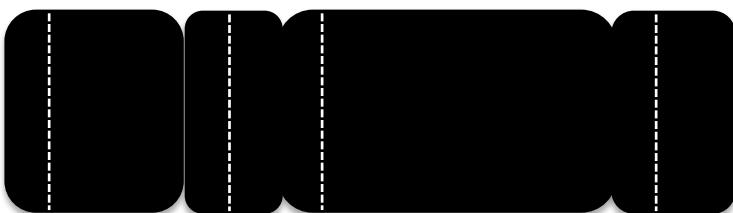
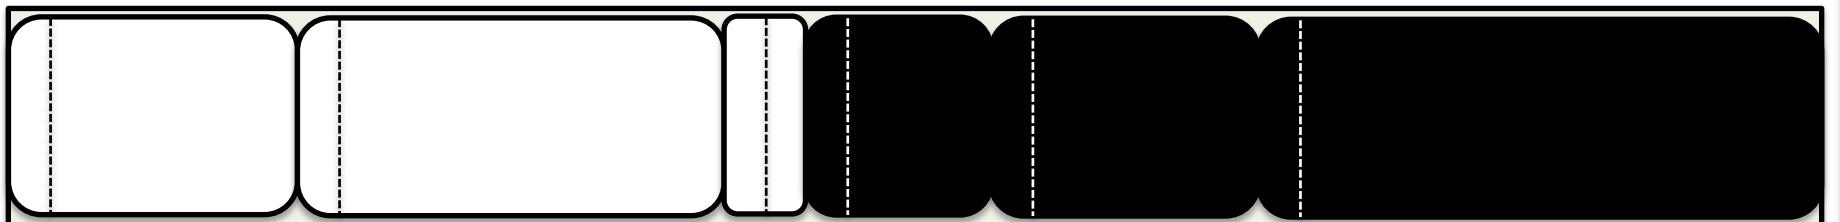
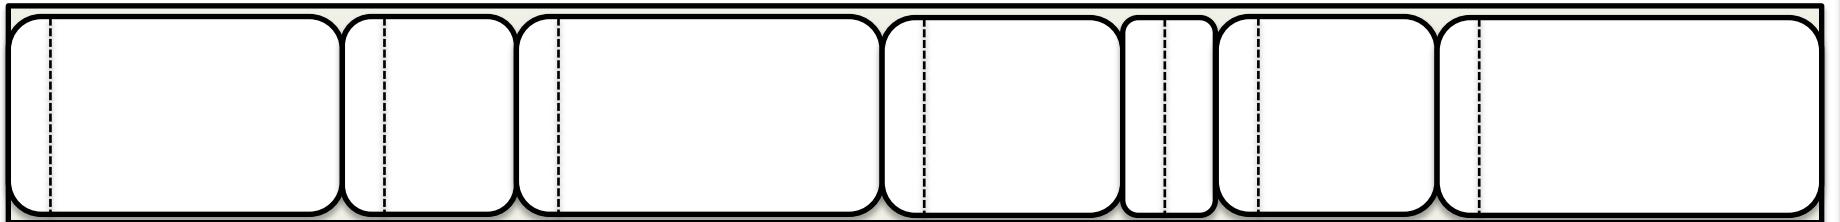
- Fragmentation is caused by freeing objects
- Garbage collector can compact the heap
 - Move all objects to a contiguous block
- MS and RC collectors free unreachable objects
- Copying collectors evacuate reachable objects
 - Everything left is garbage

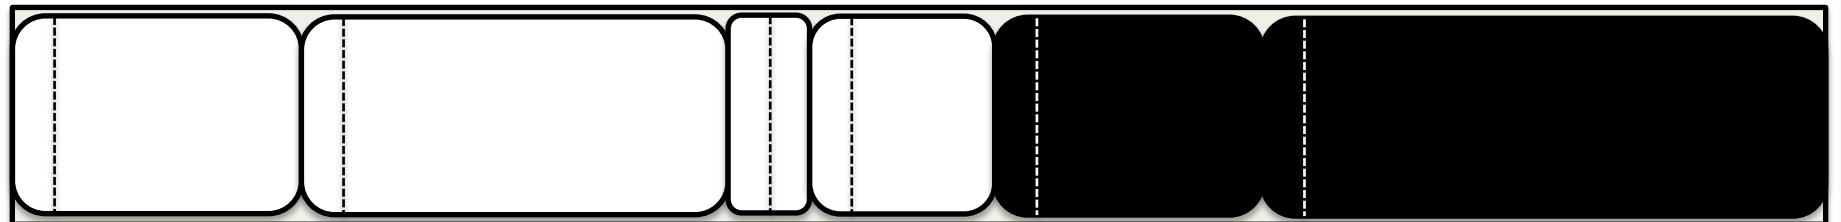
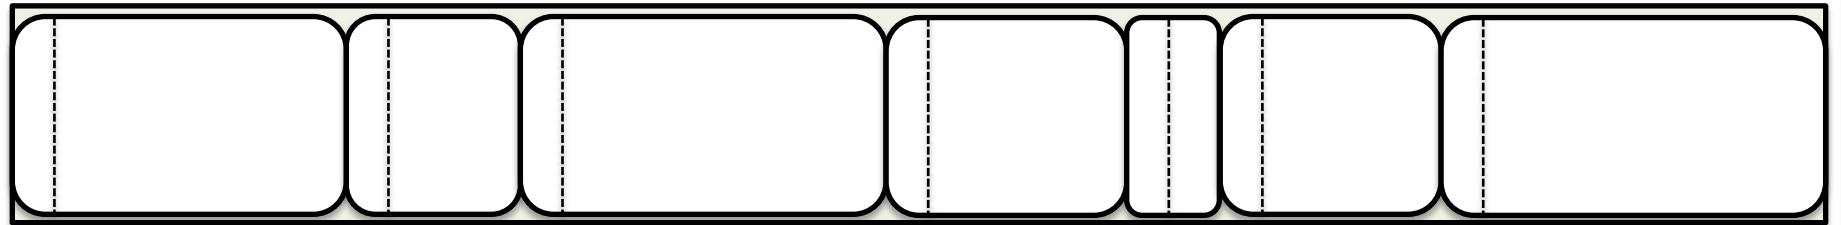


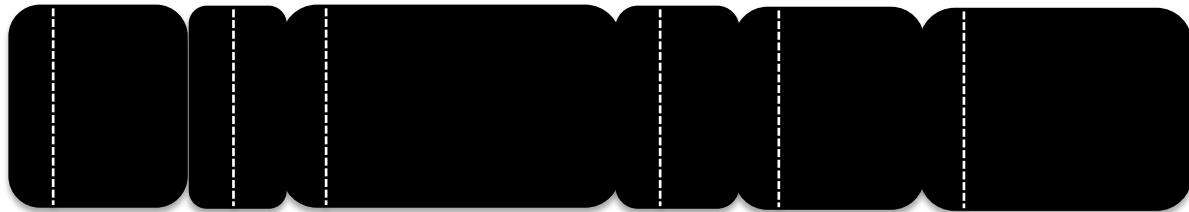
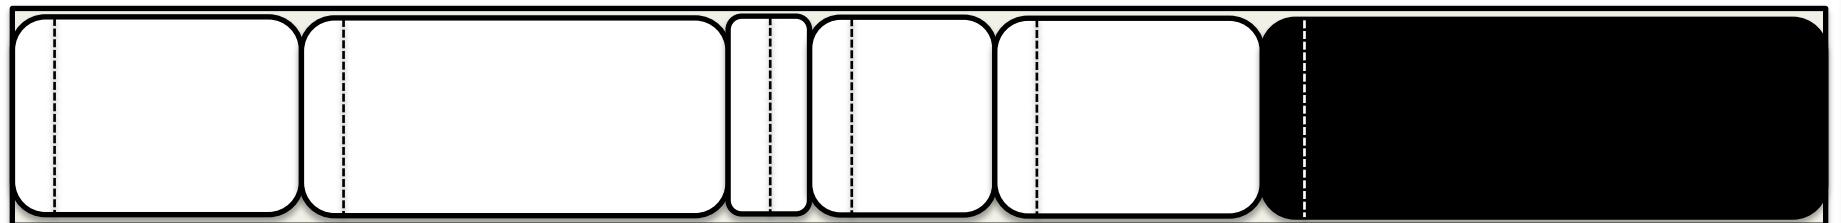
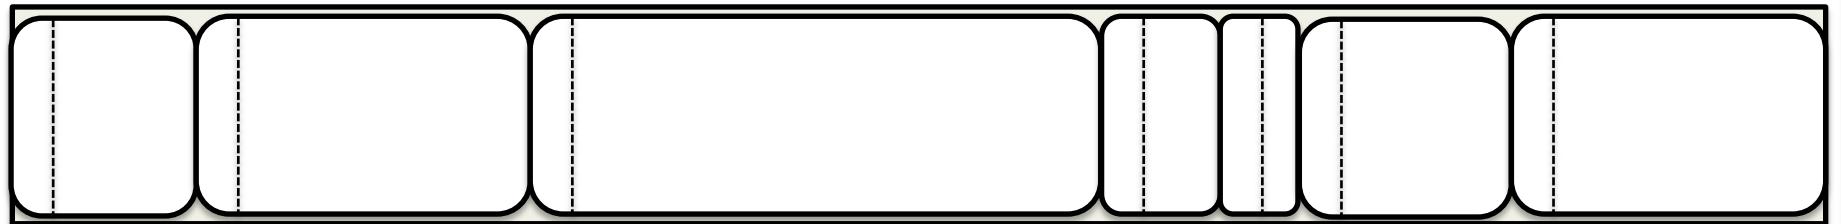
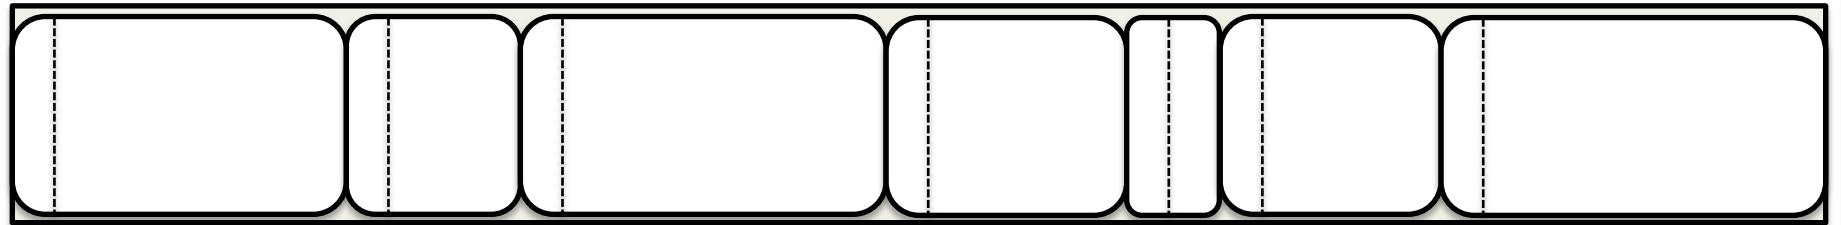


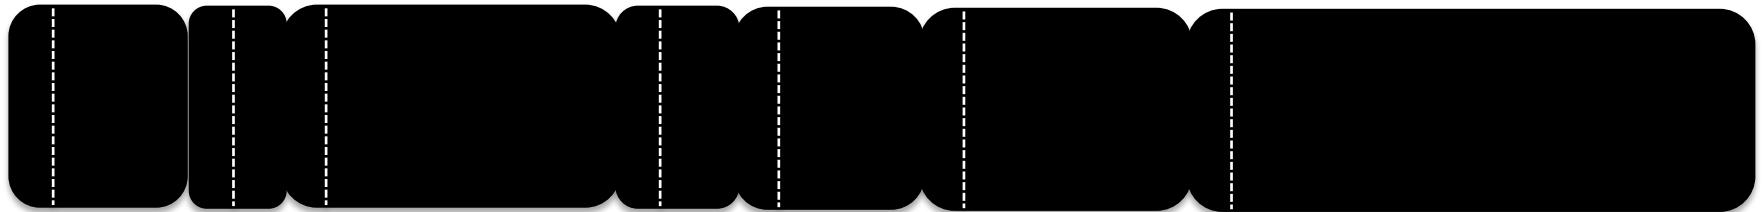
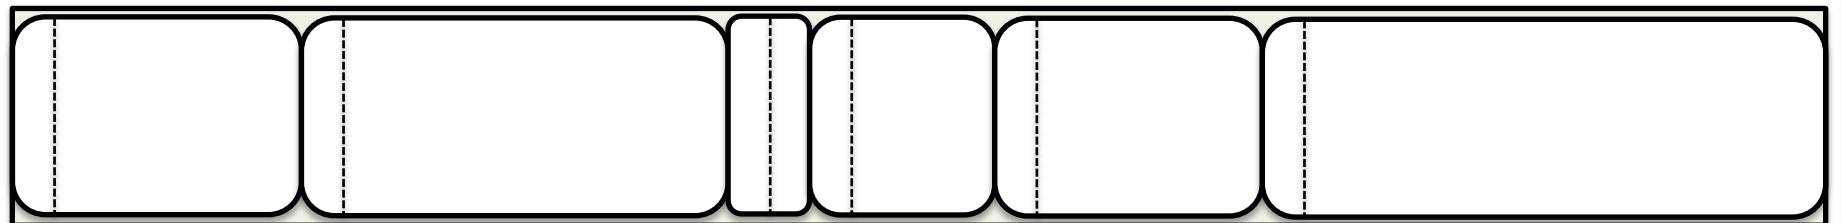
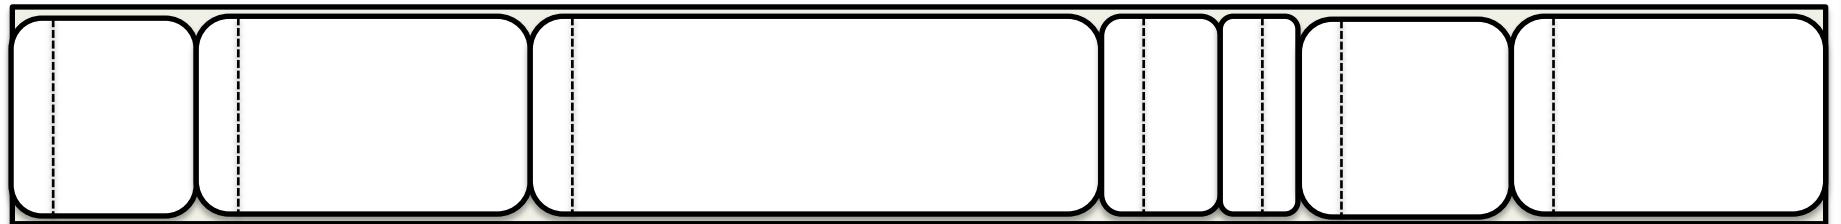
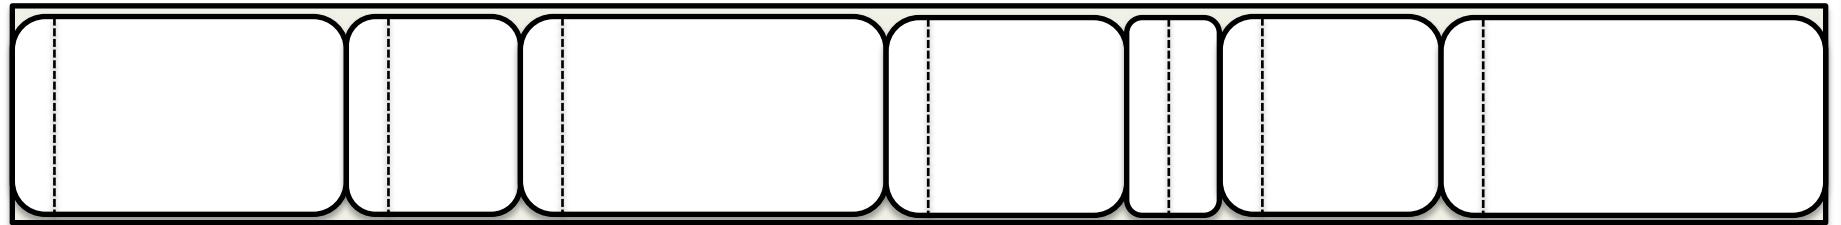


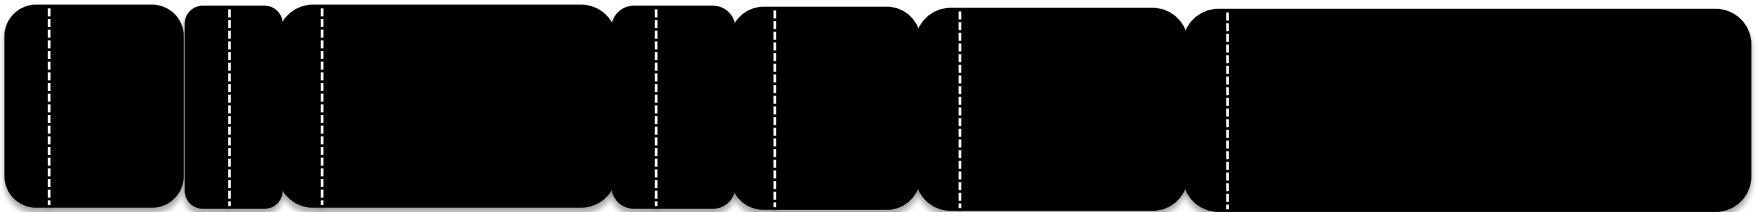
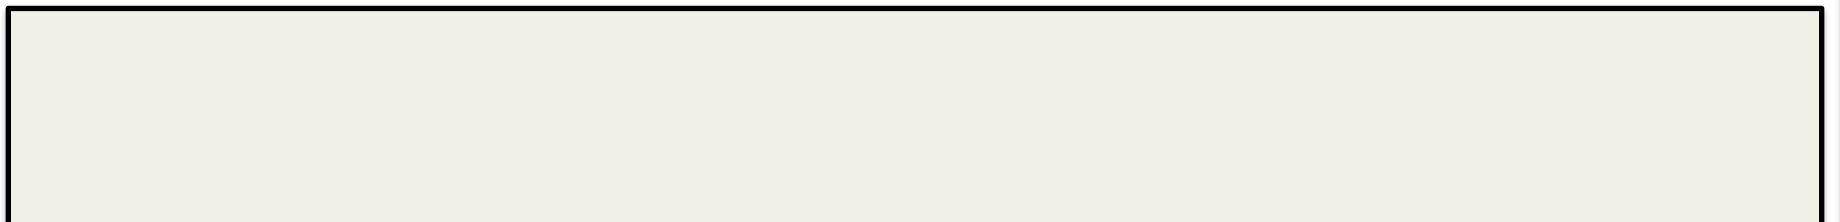
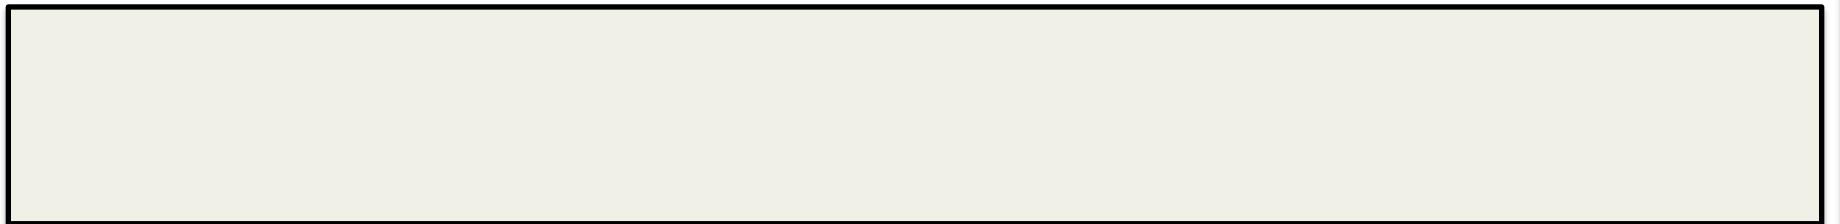










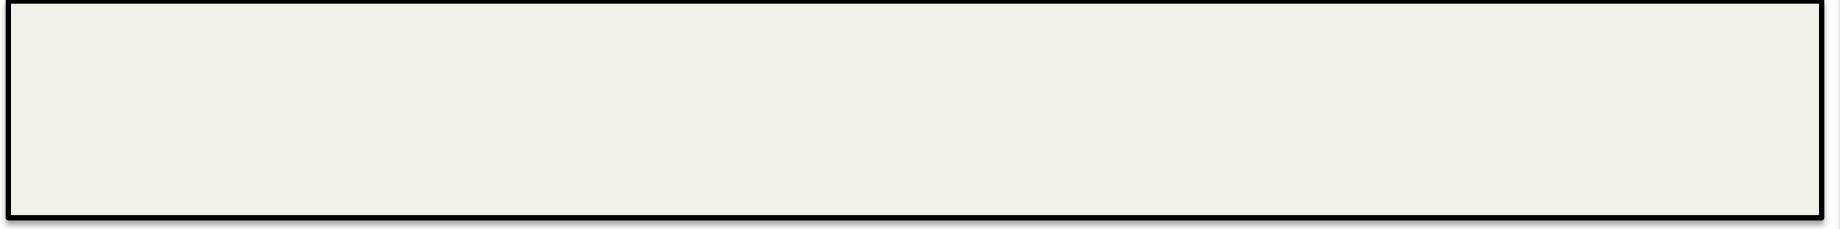
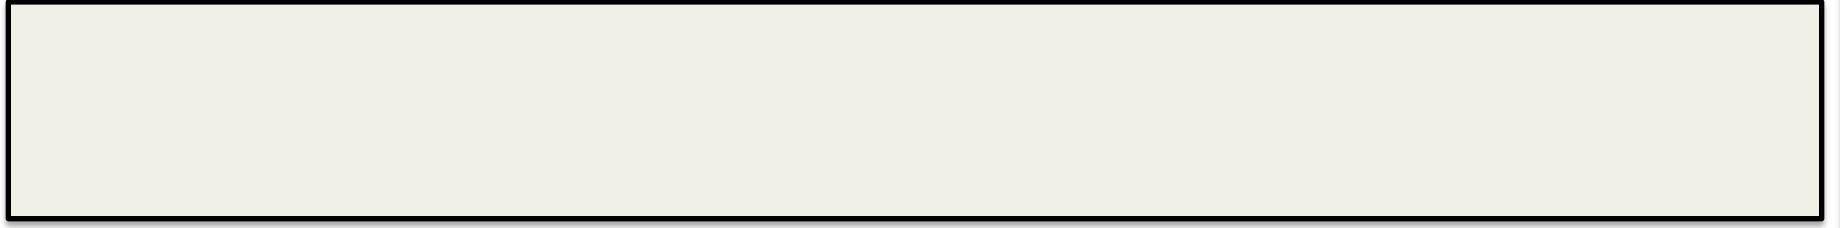


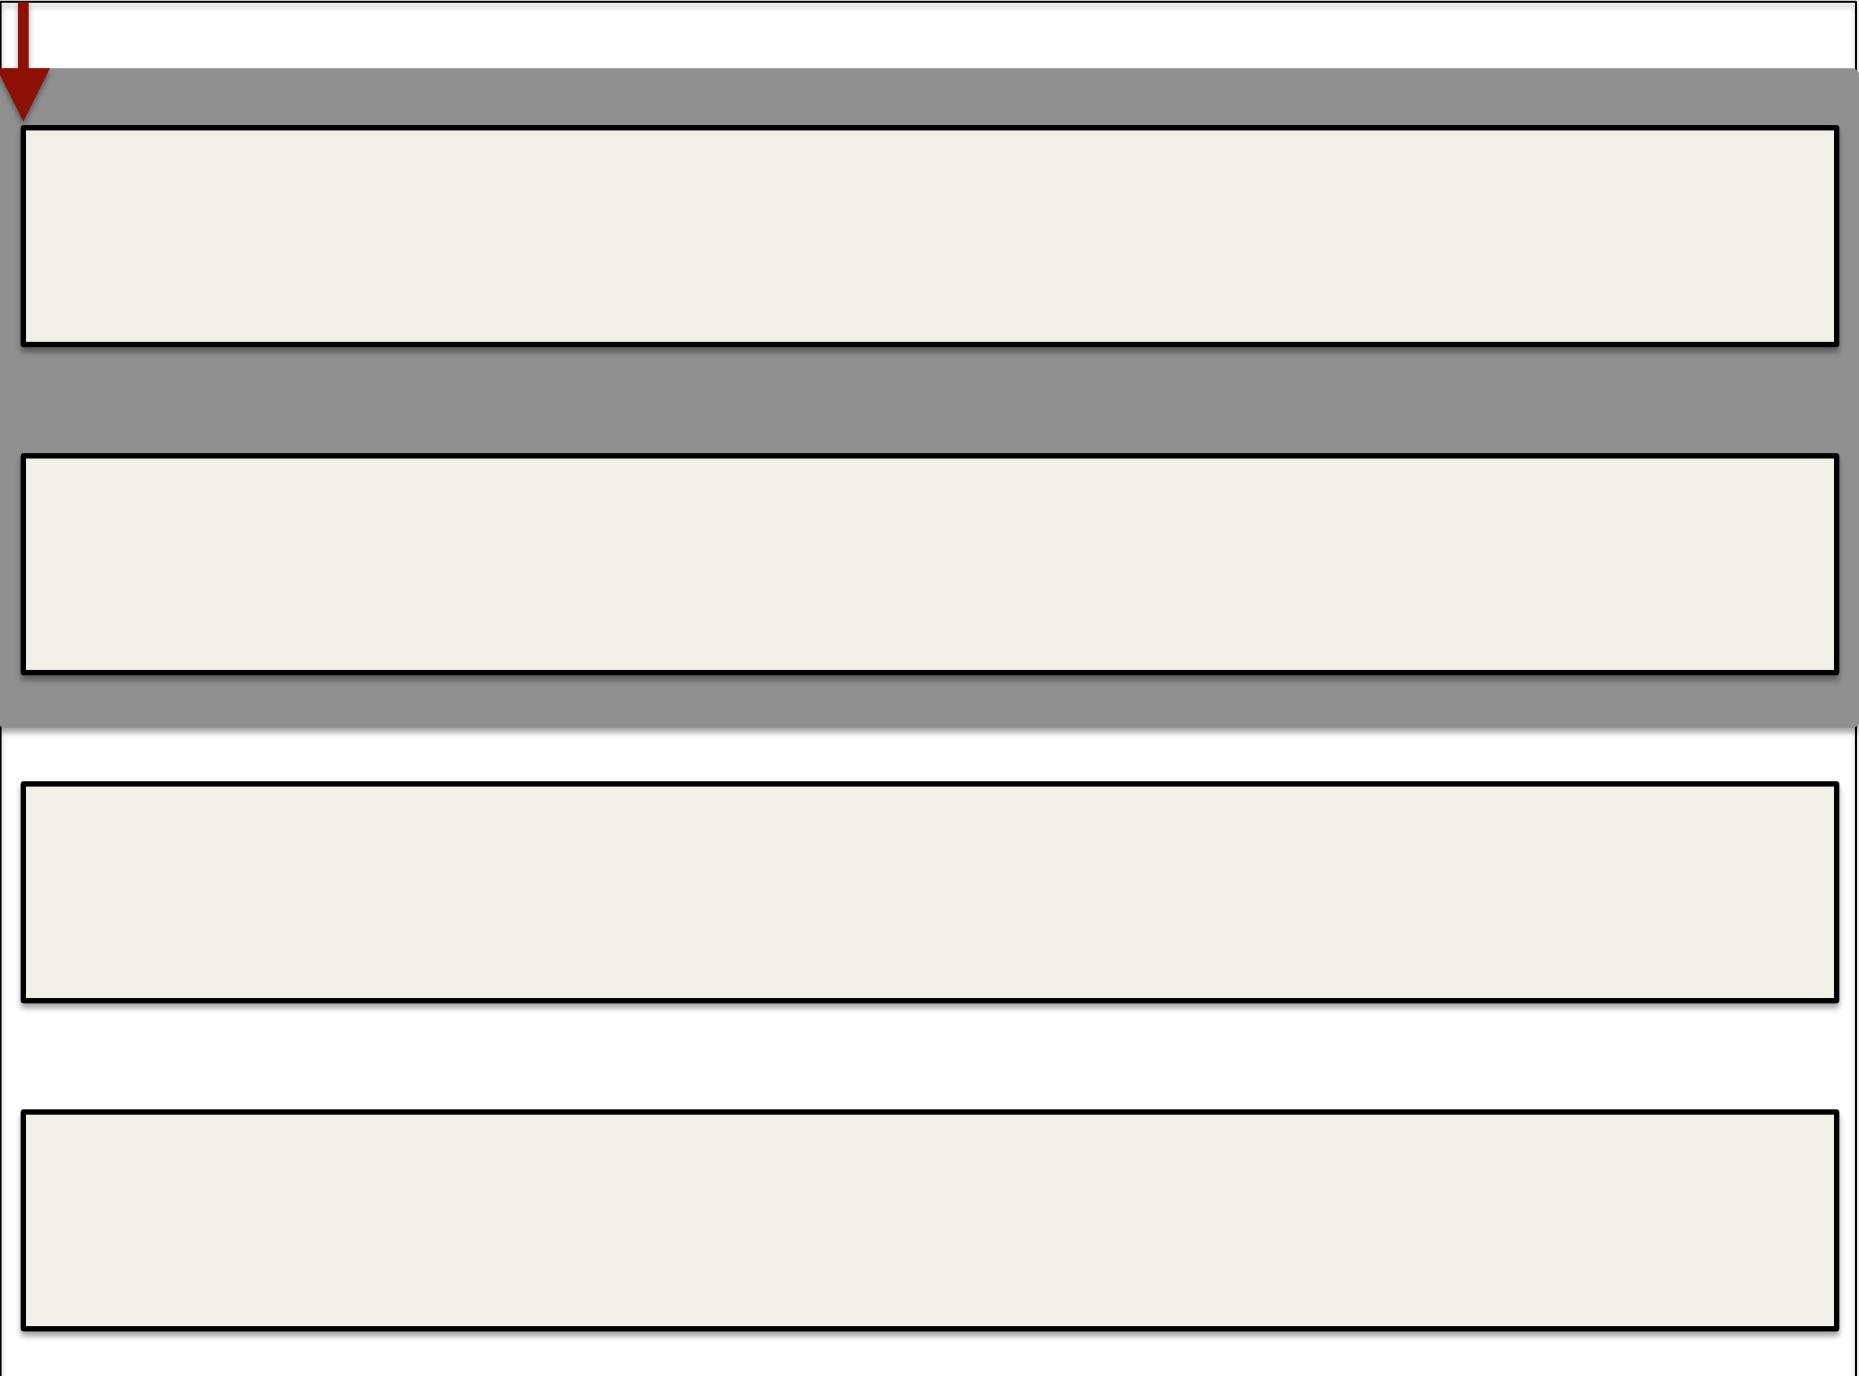
Copy Reserve

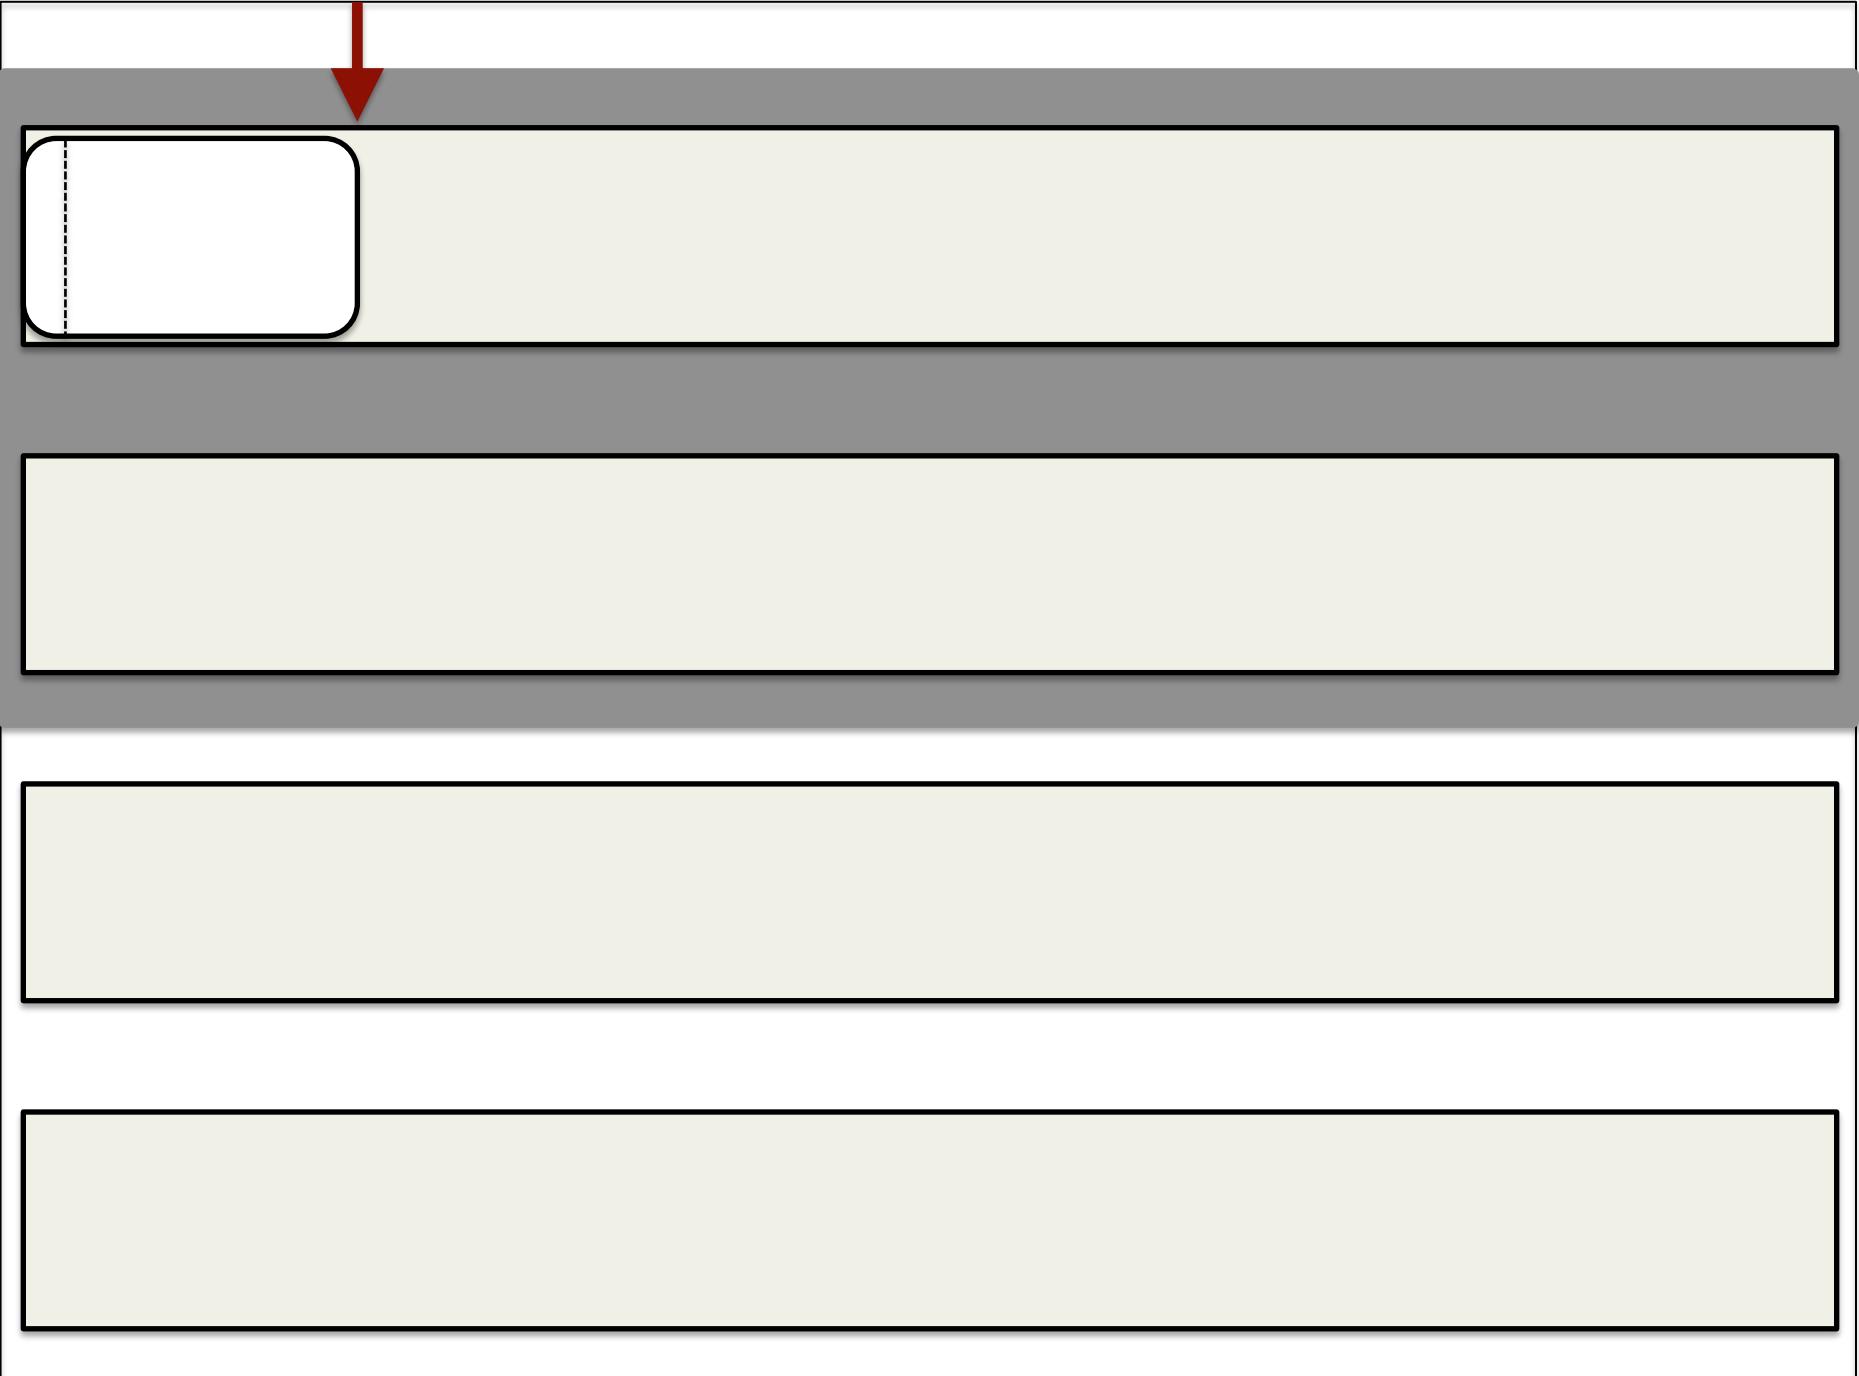
- Where do the copied objects go?
- We need to set aside space to copy them
 - In the worst case, all objects survive a collection
 - Need to set aside half the heap
 - We'll see later how to reduce this
- The extra space is called the copy reserve

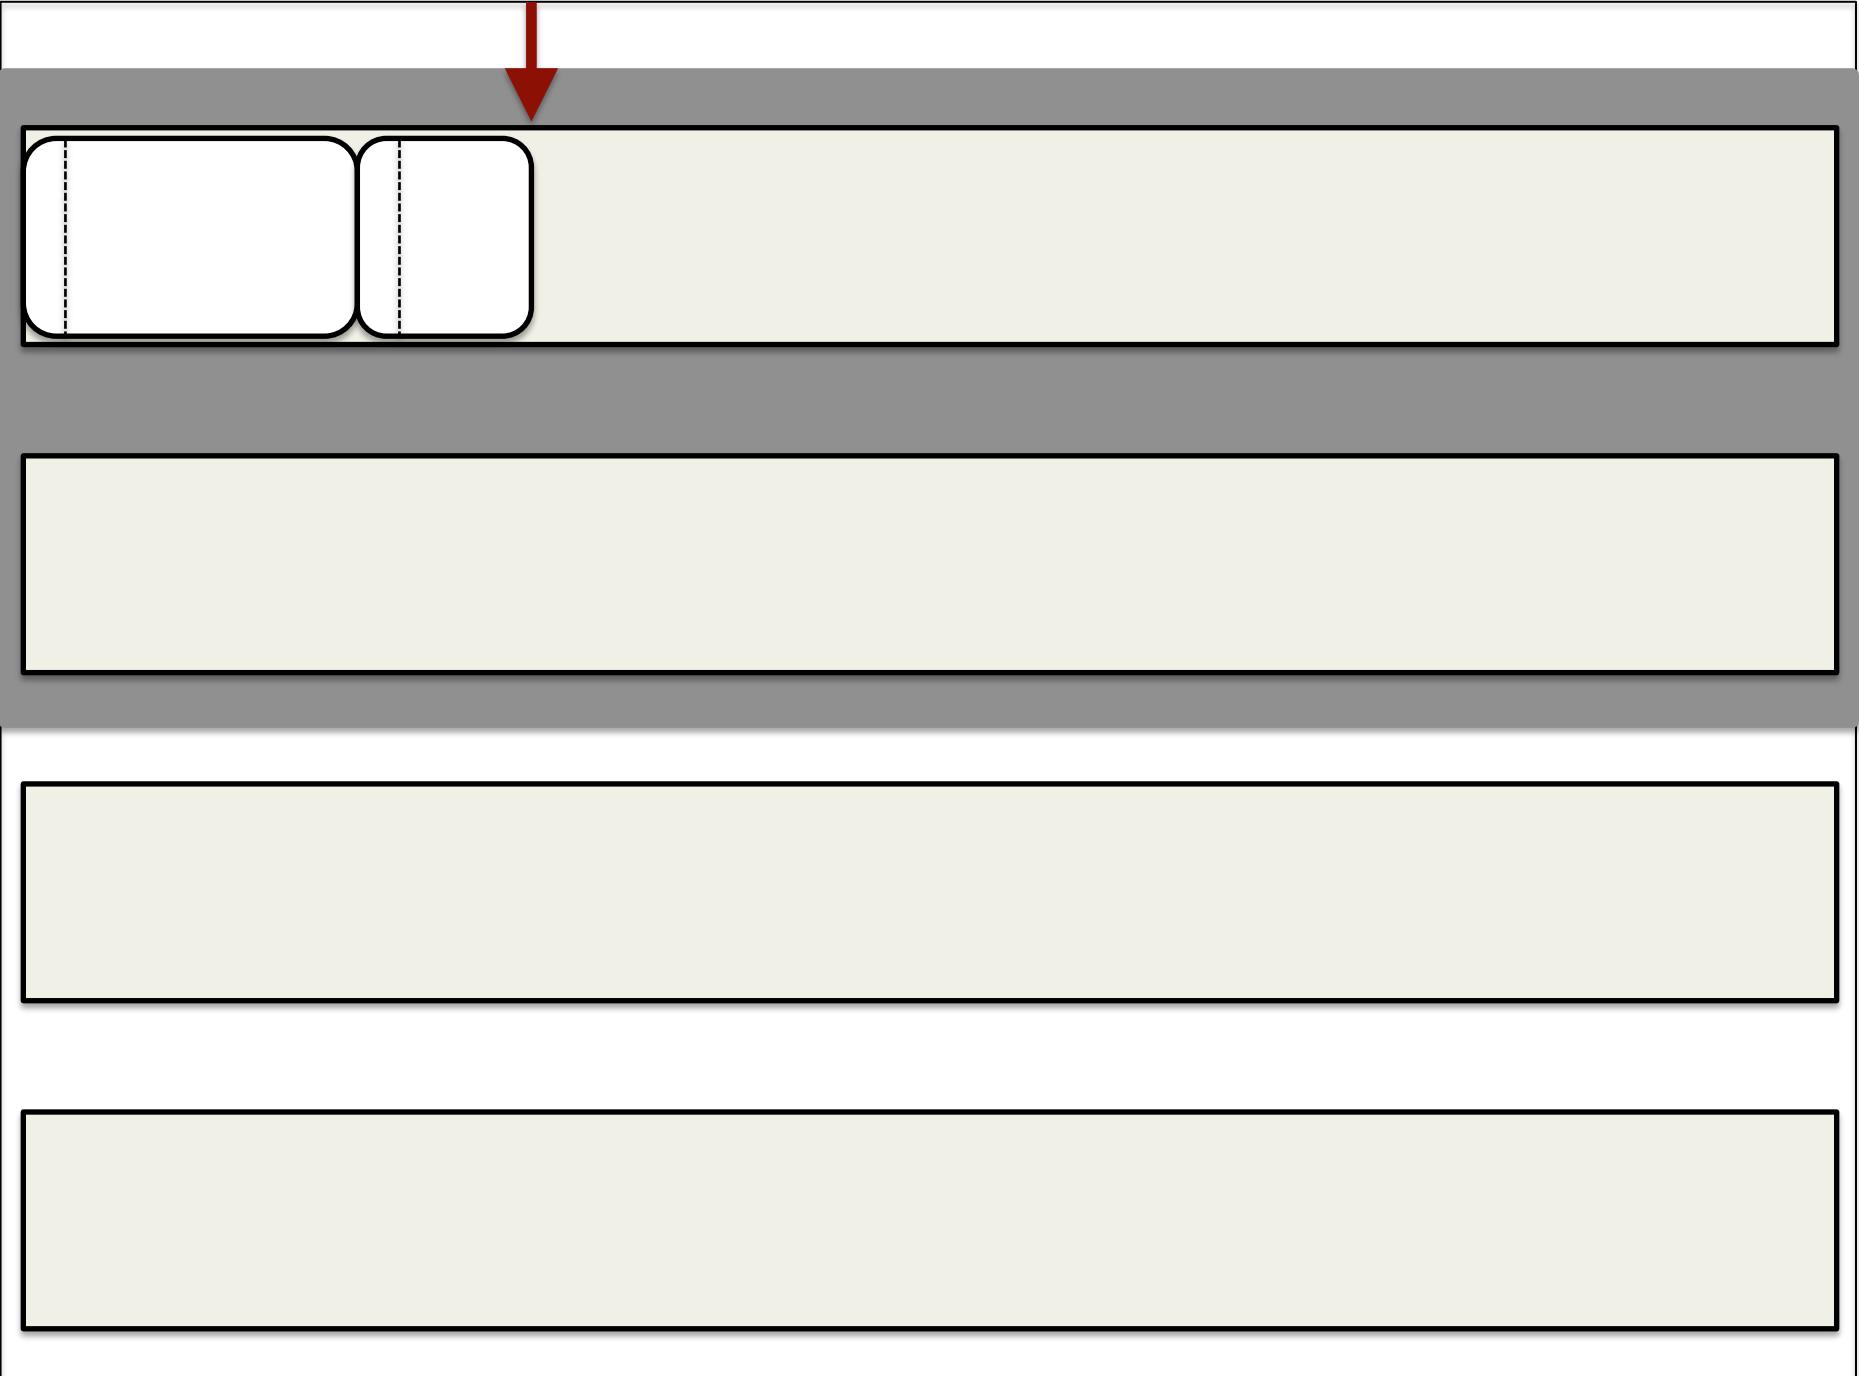
Semi-Space Collectors

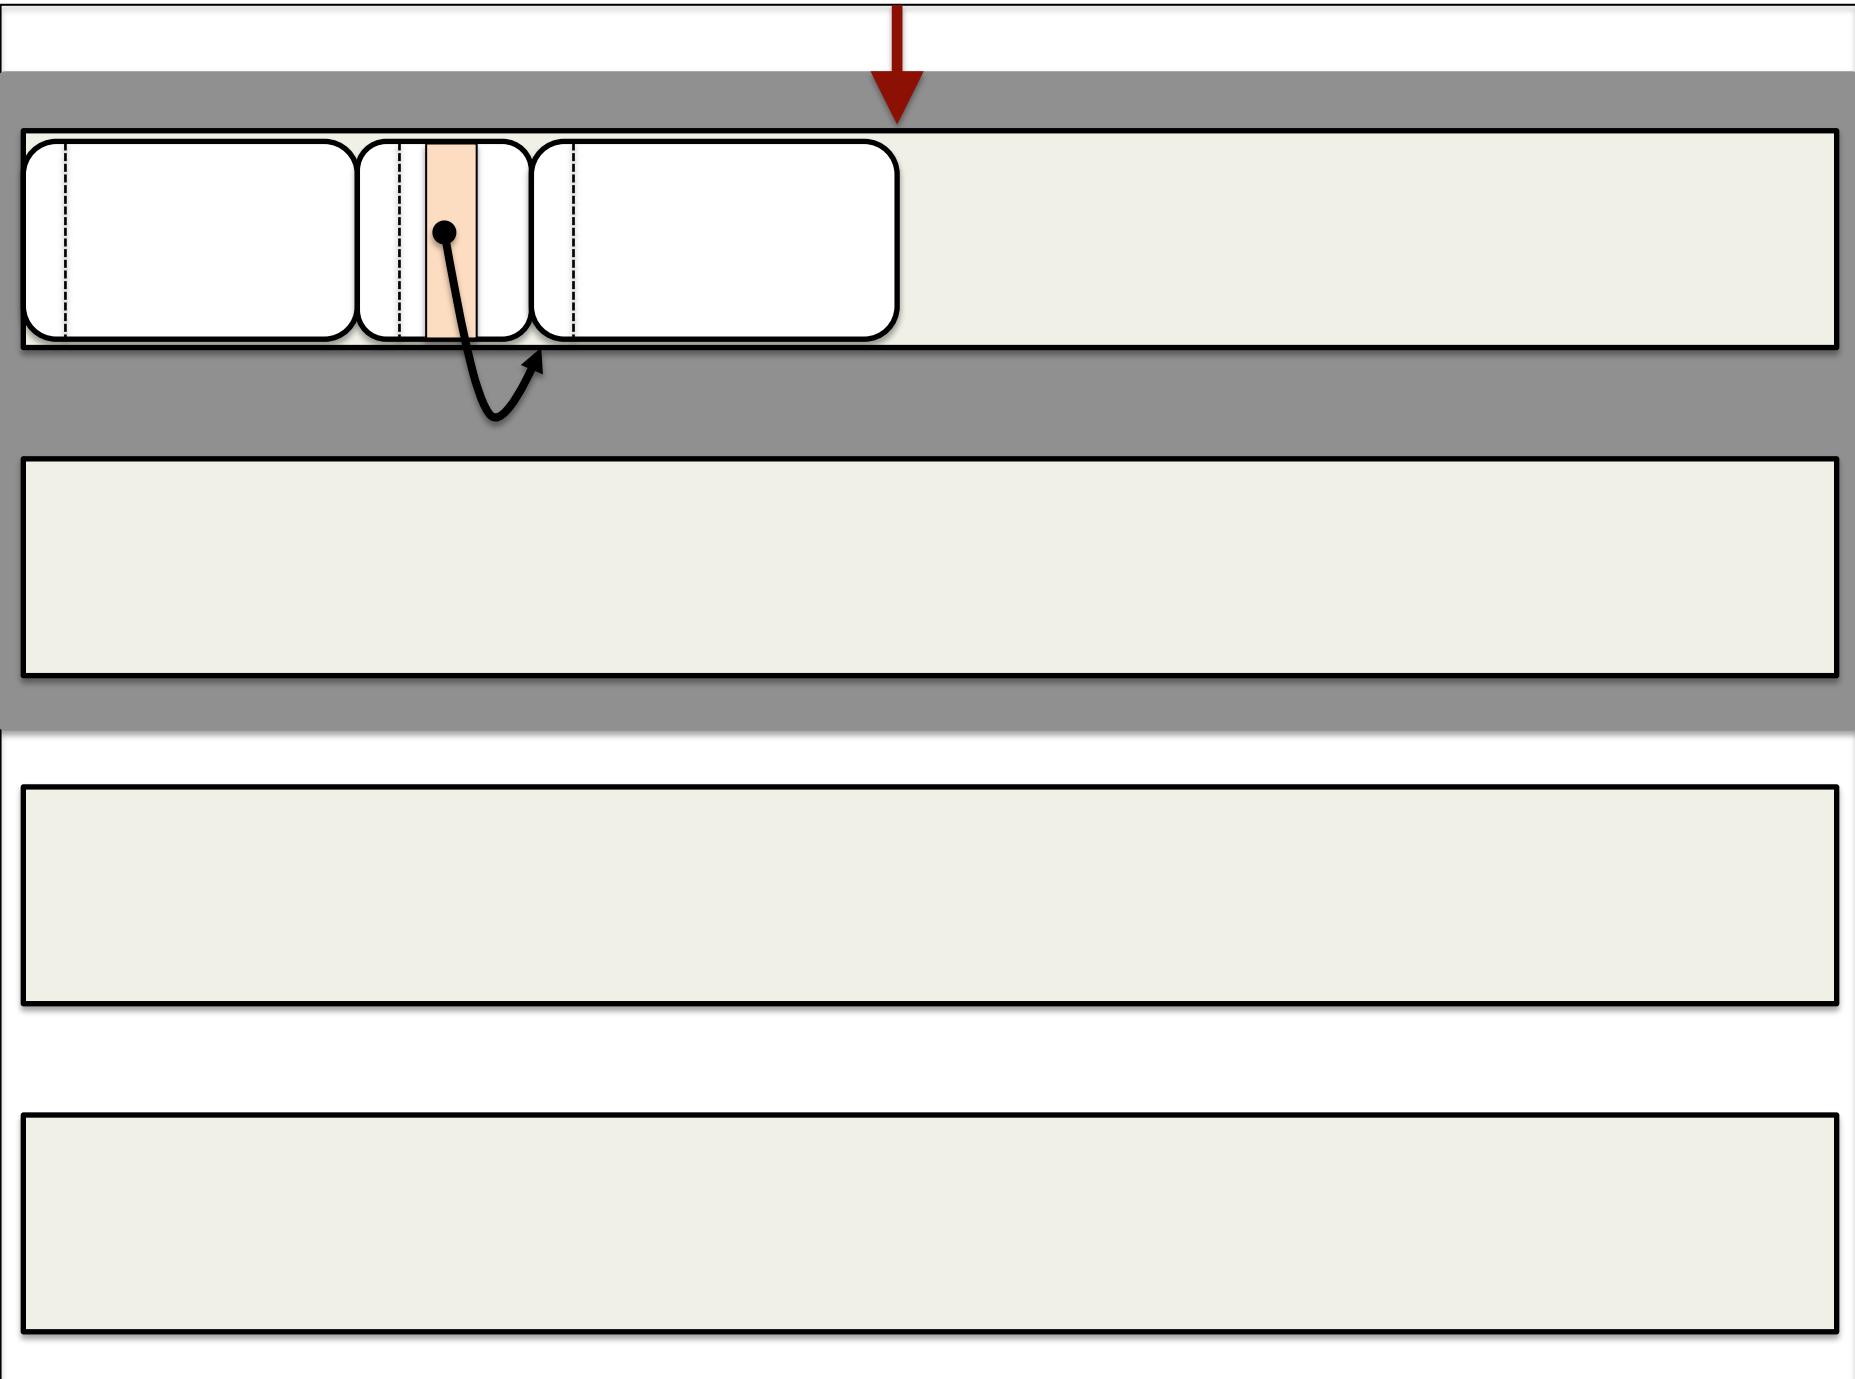
- Simplest form of copying collector
 - Divide the heap into two
 - Allocate to the first half
- Collector triggered when the heap is half full
 - Trace the heap to find reachable objects
 - Copy all reachable objects to the second half
- Allocation continues in the second semi-space

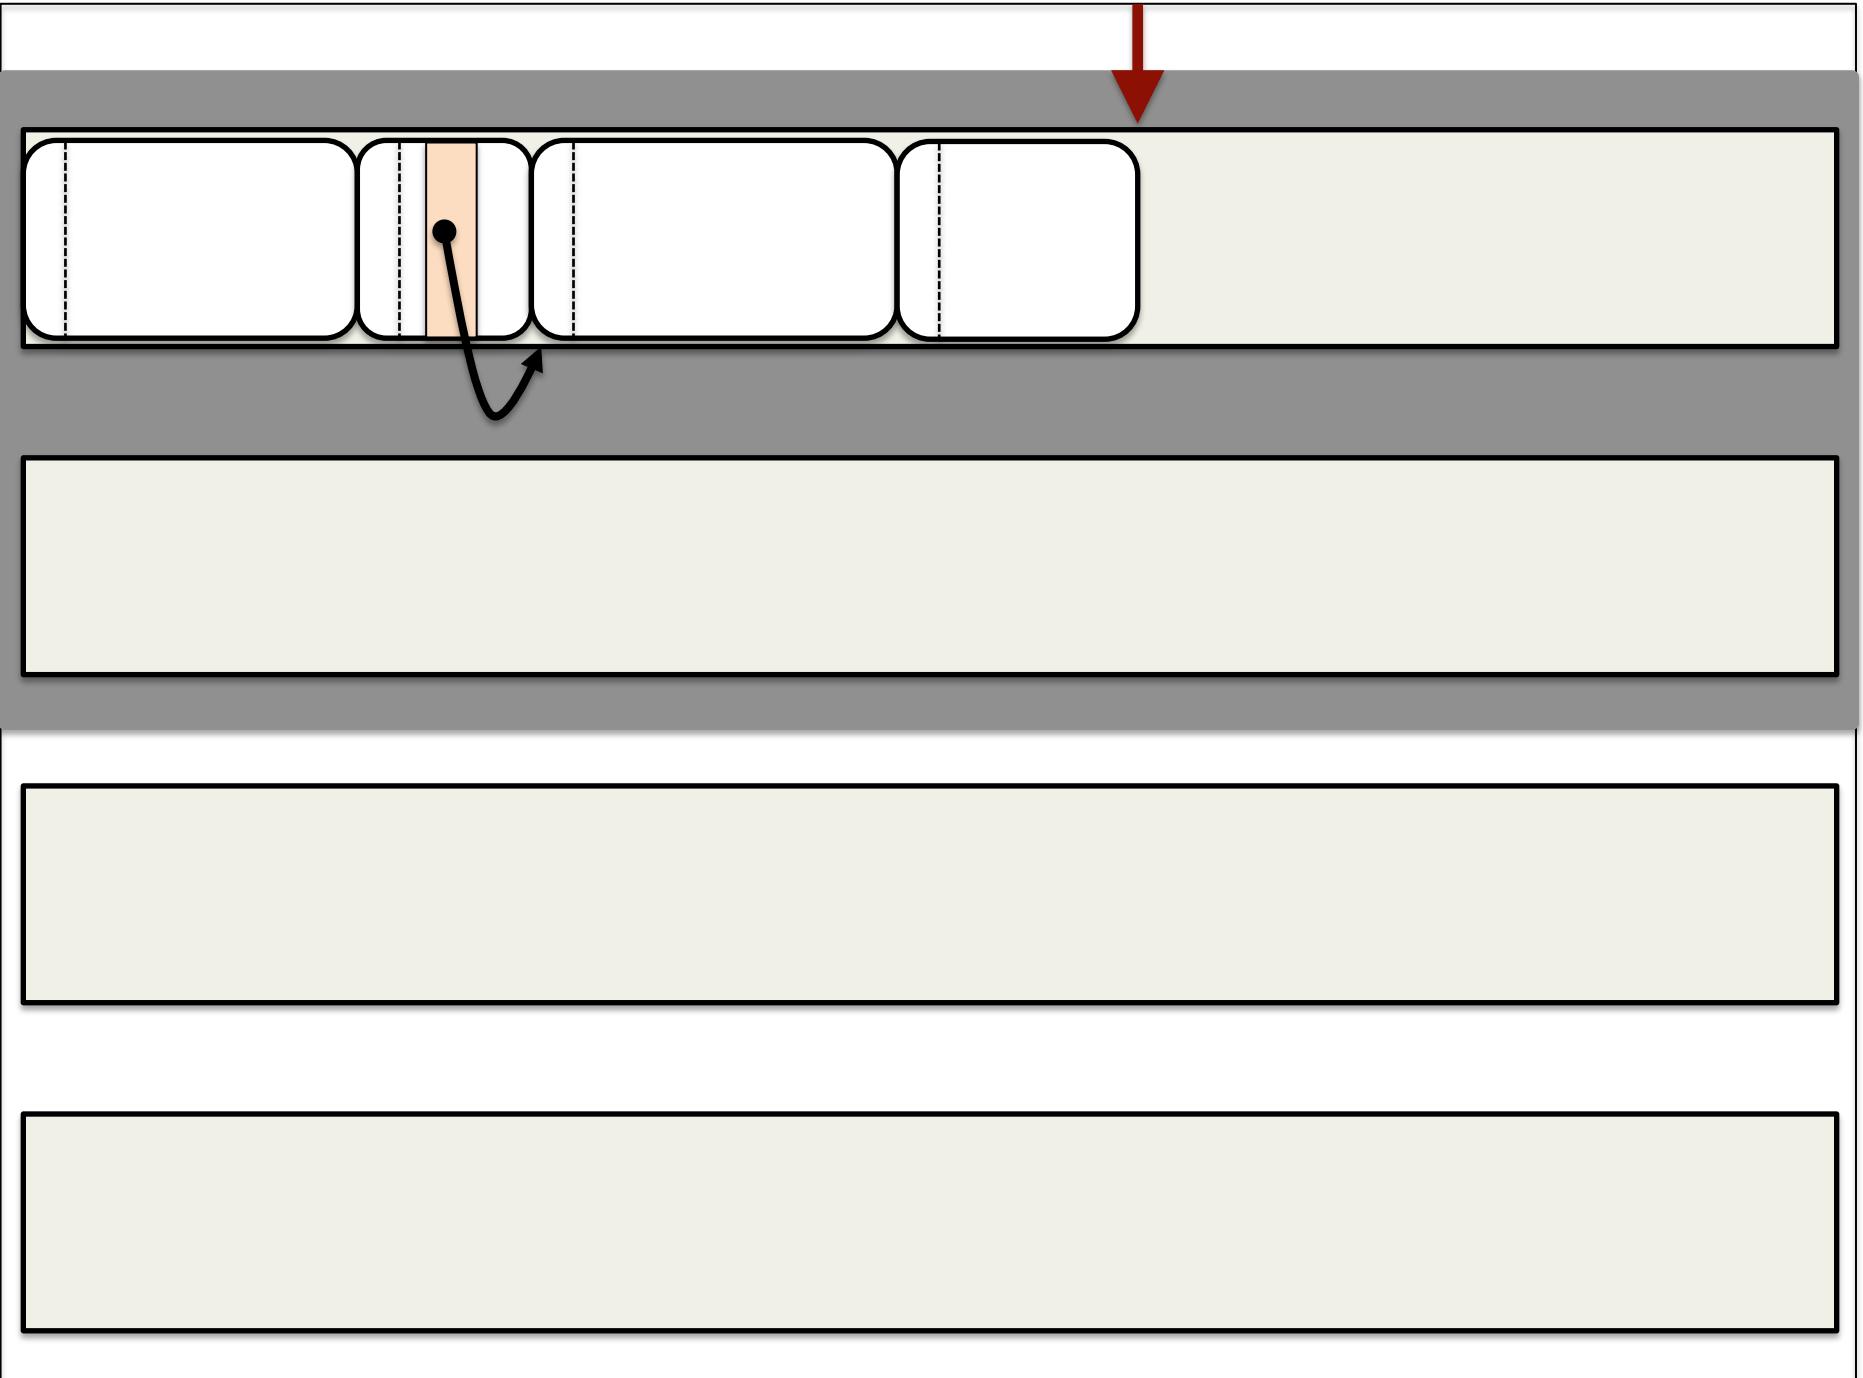


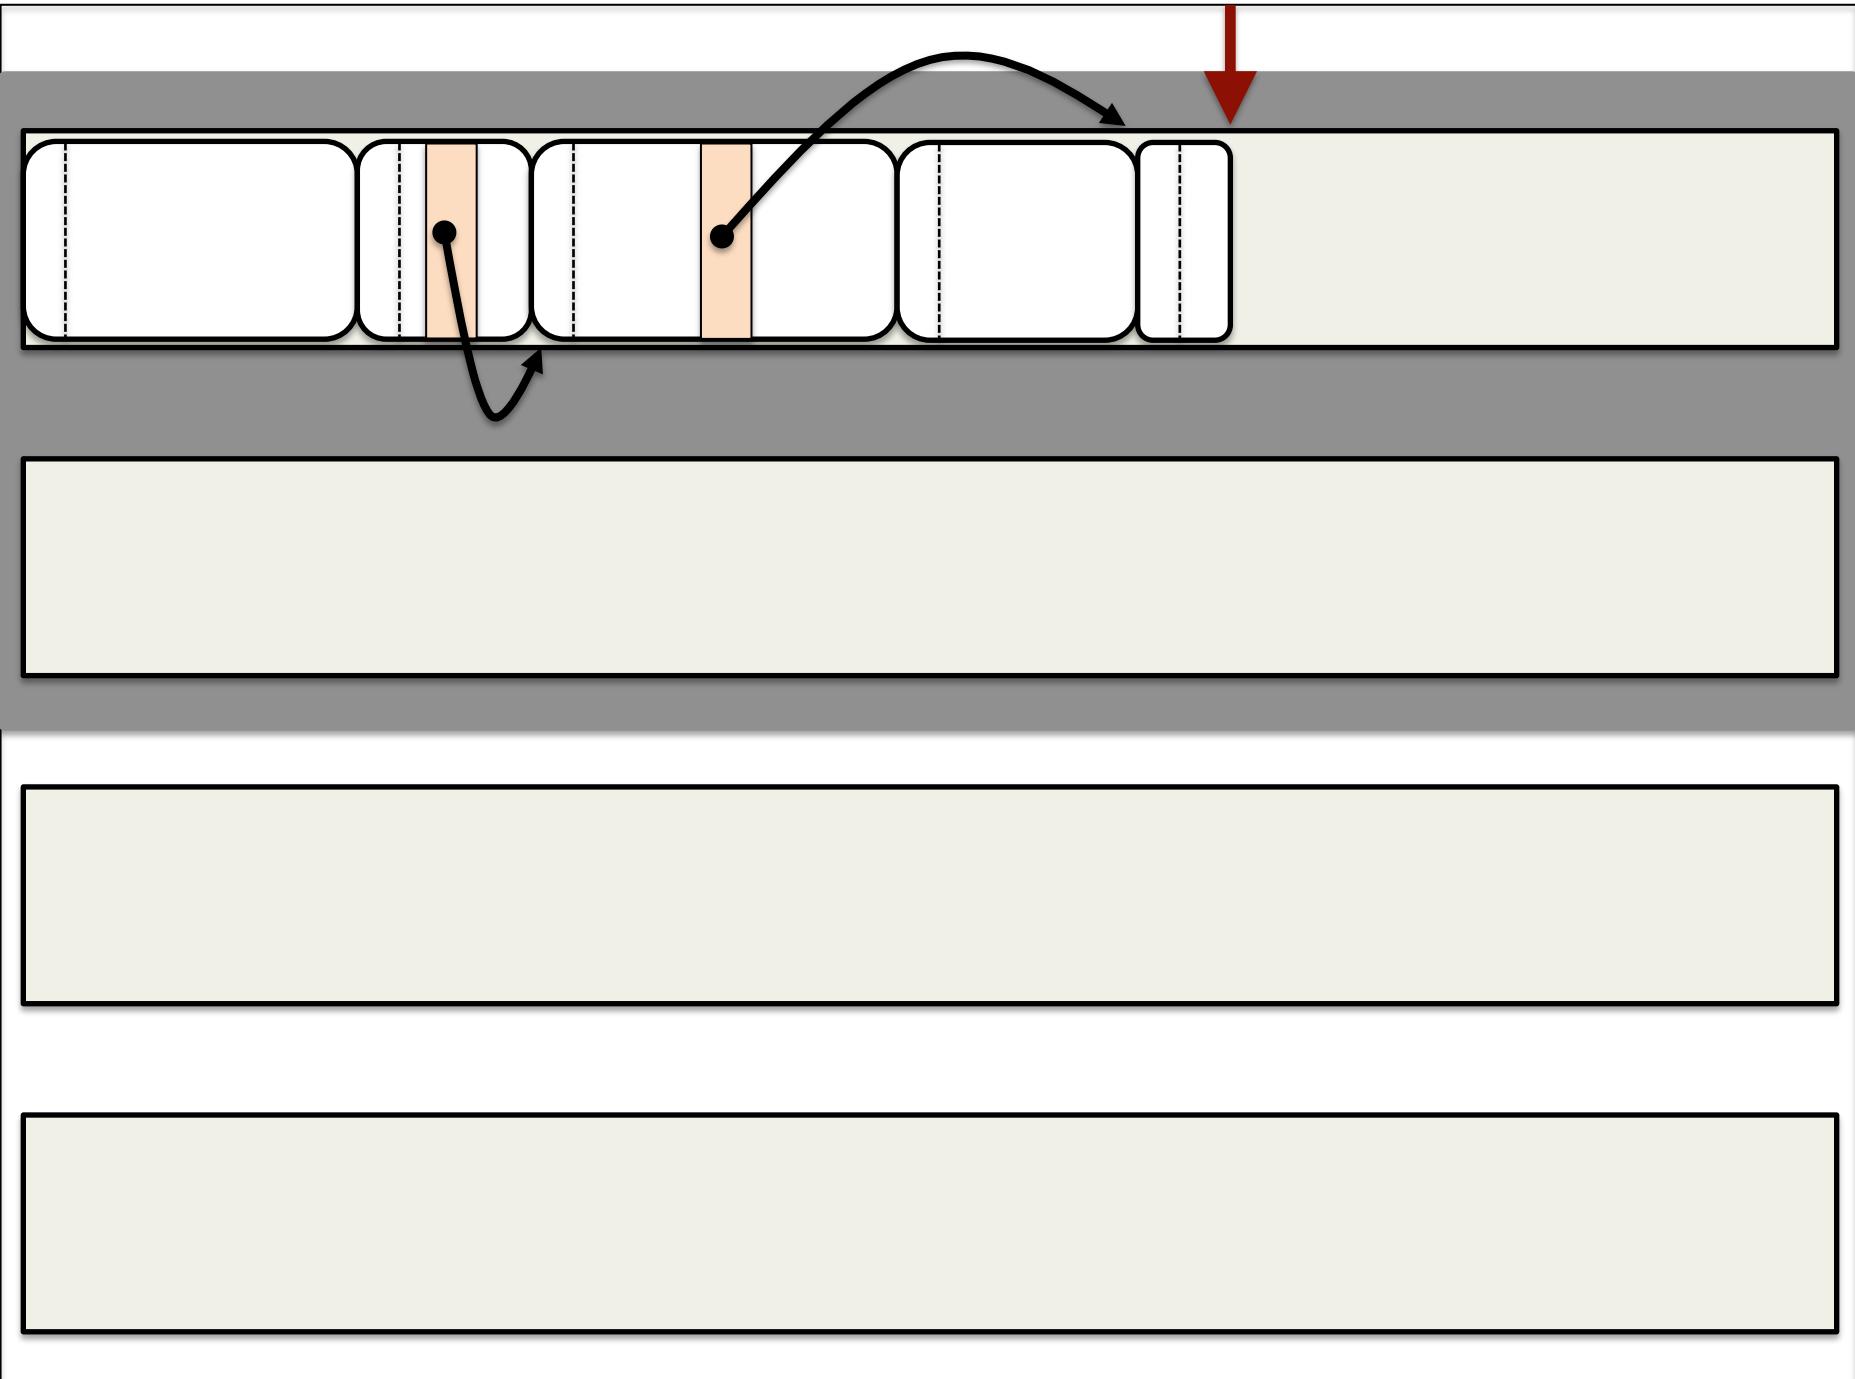


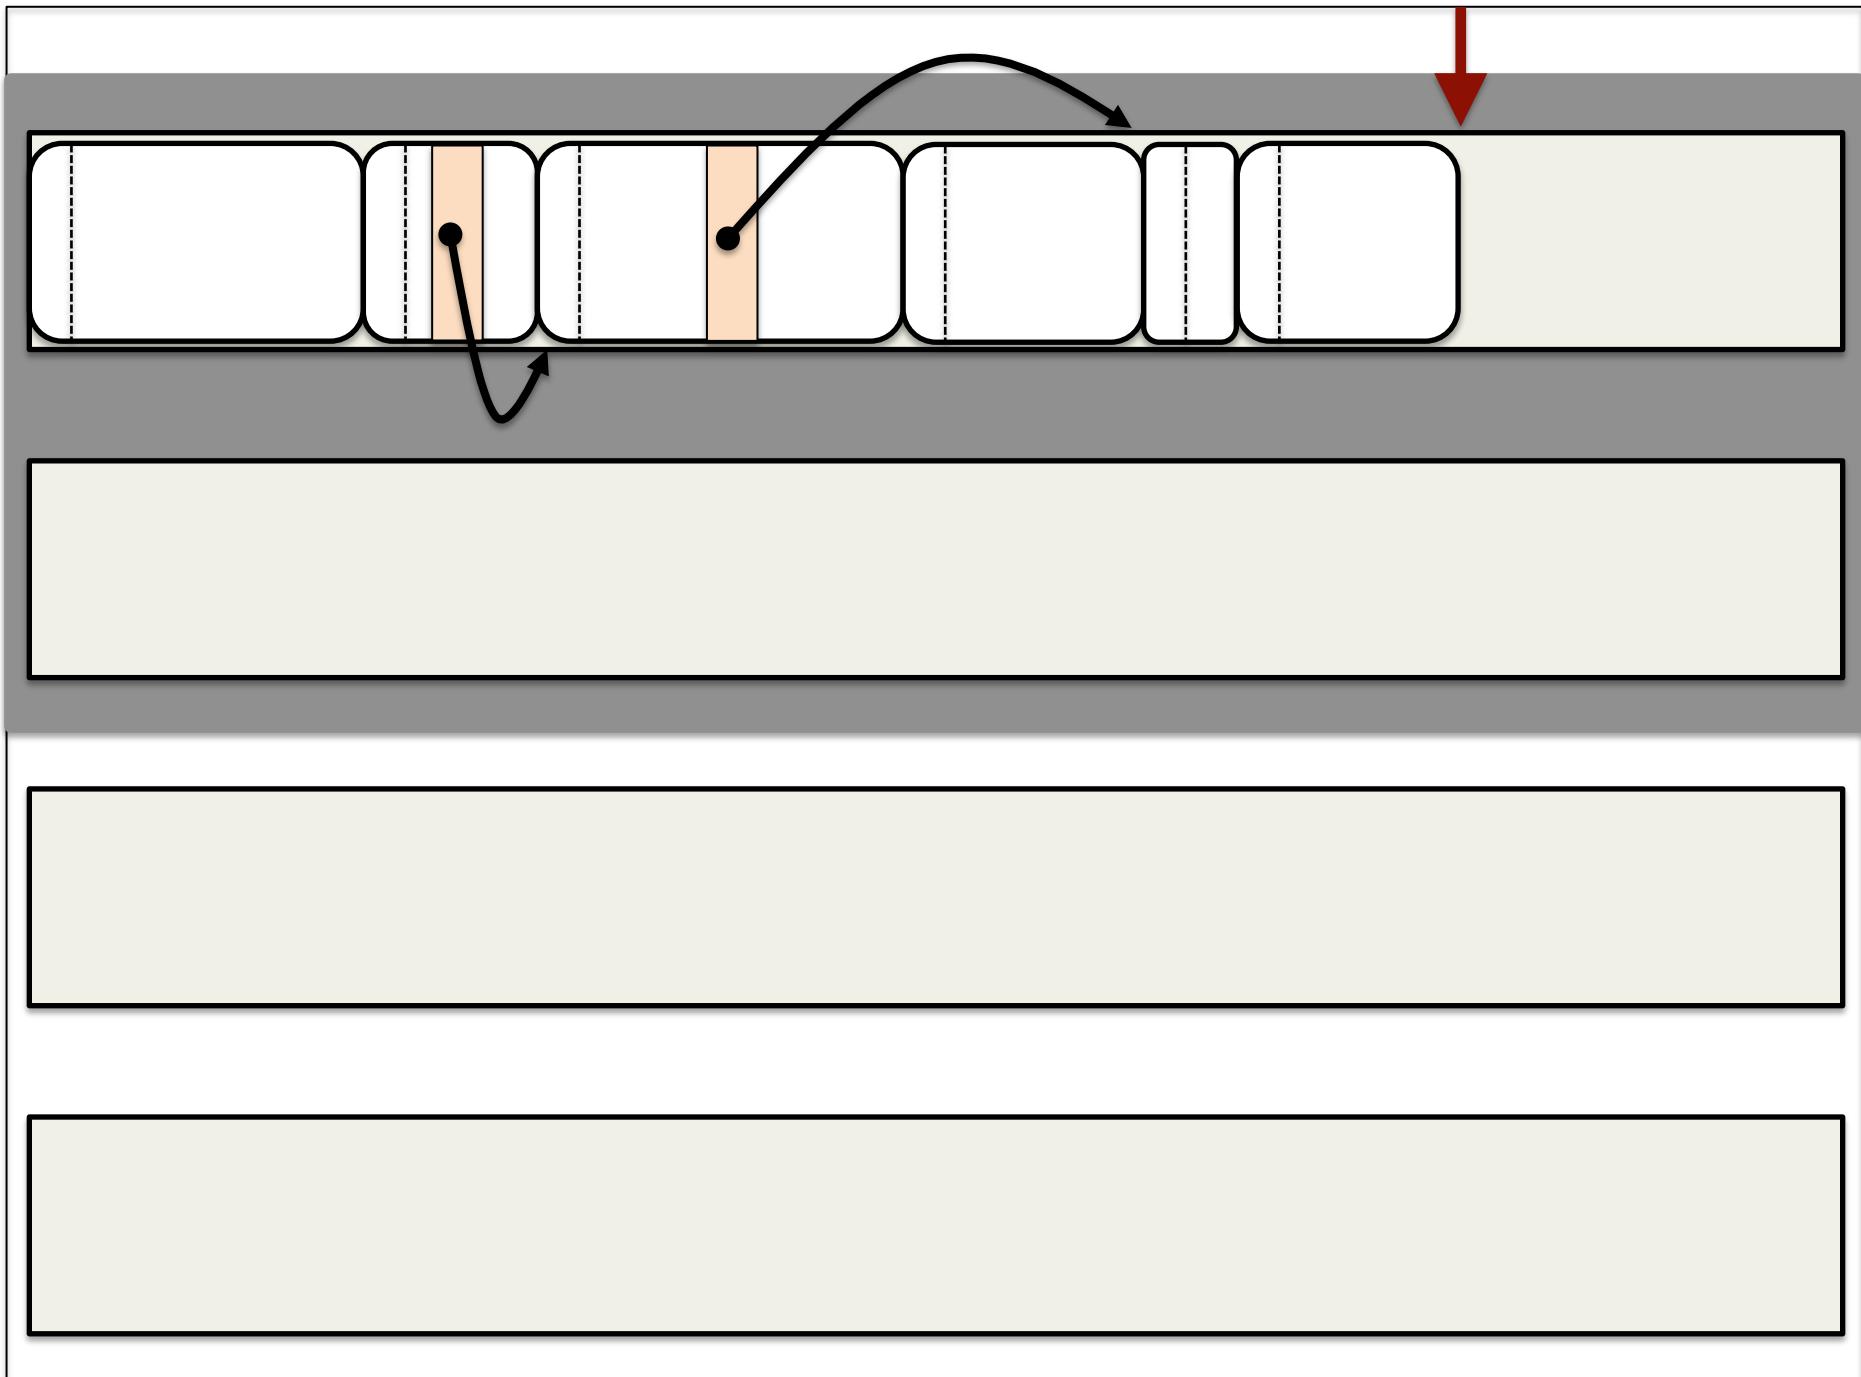


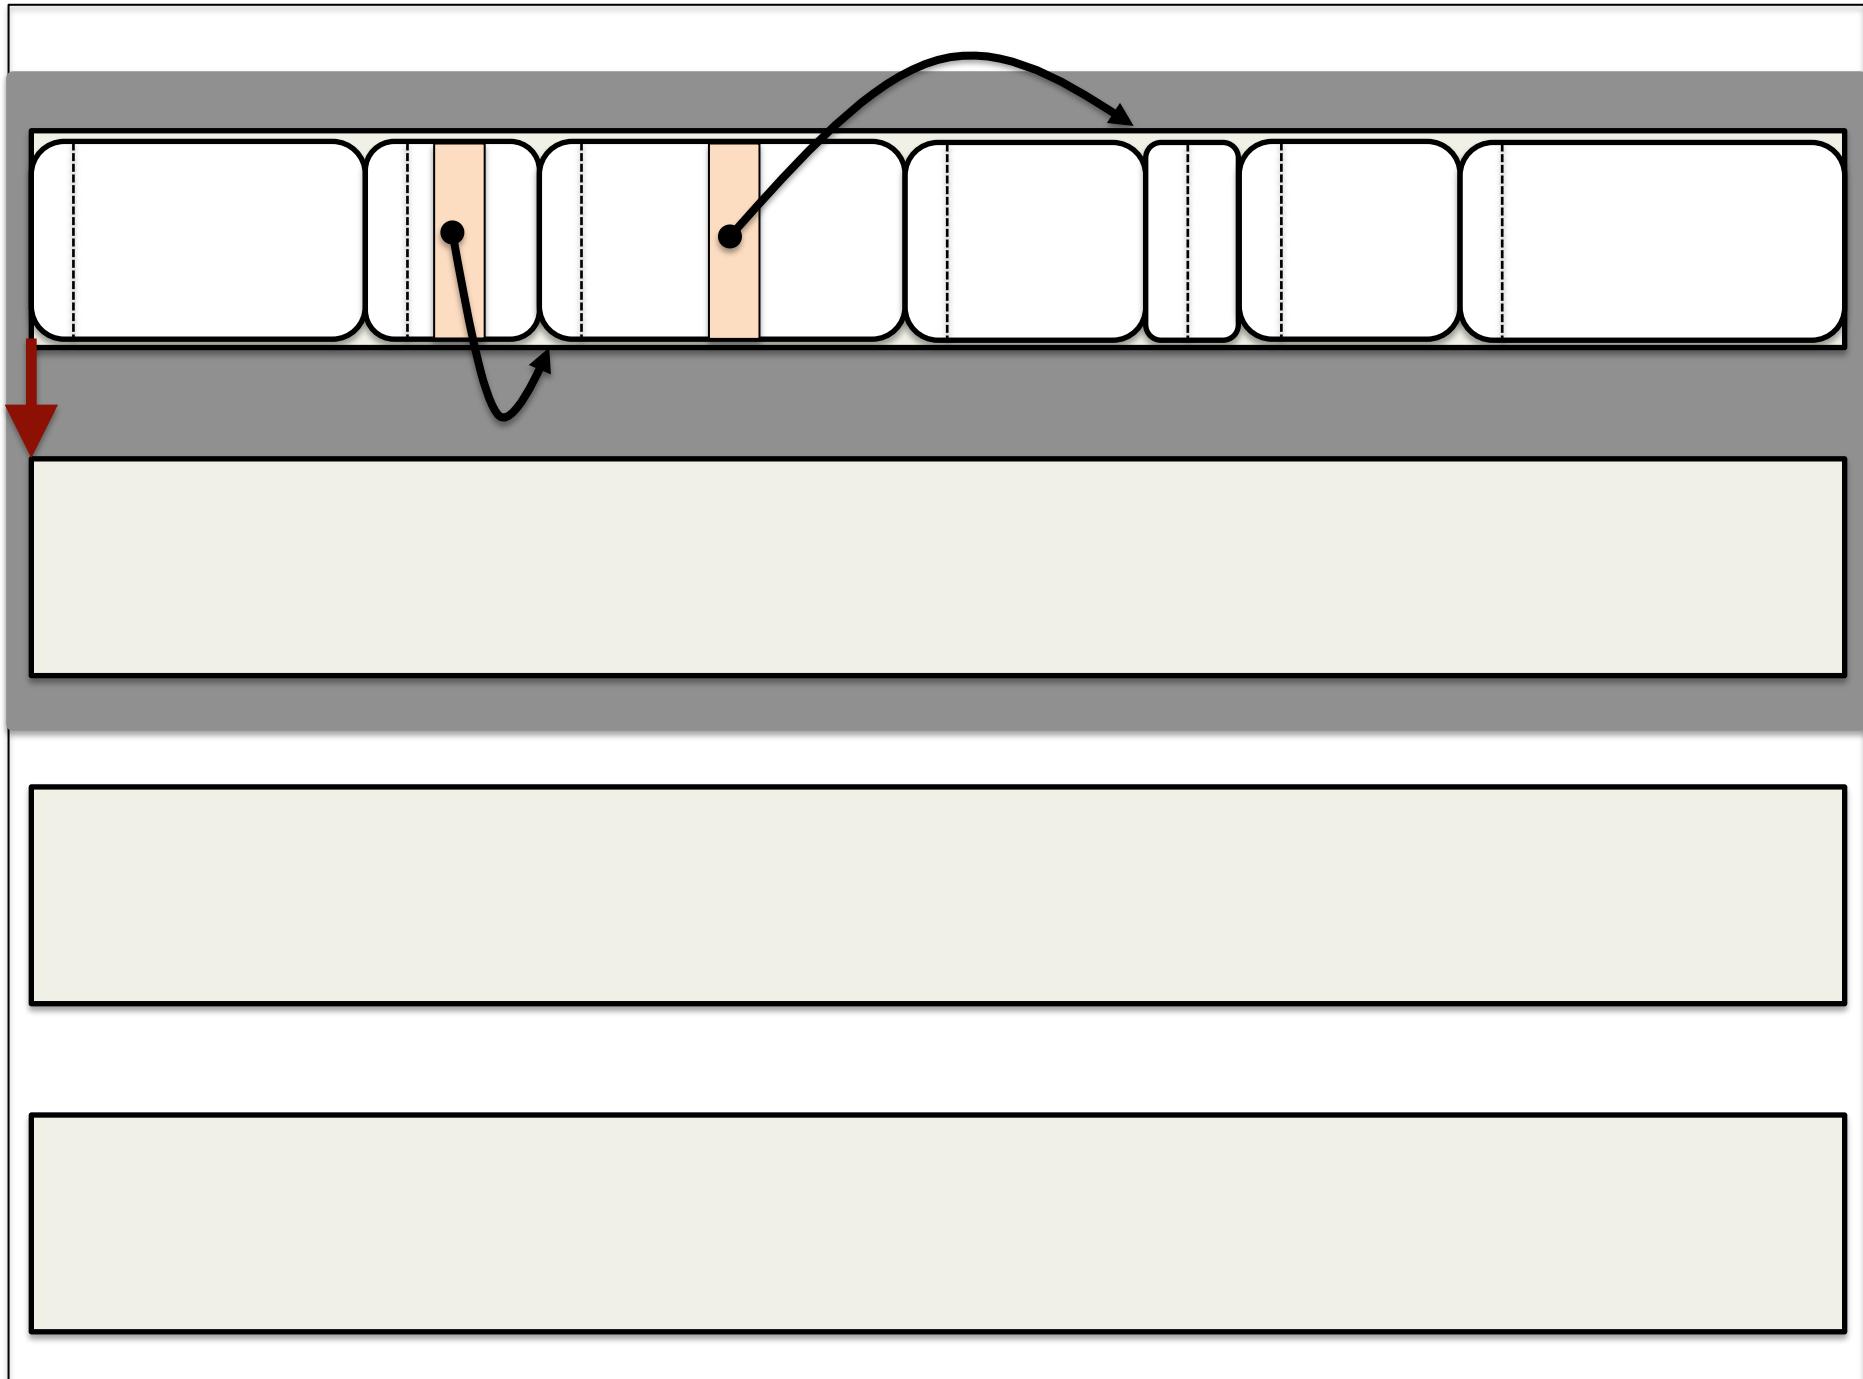










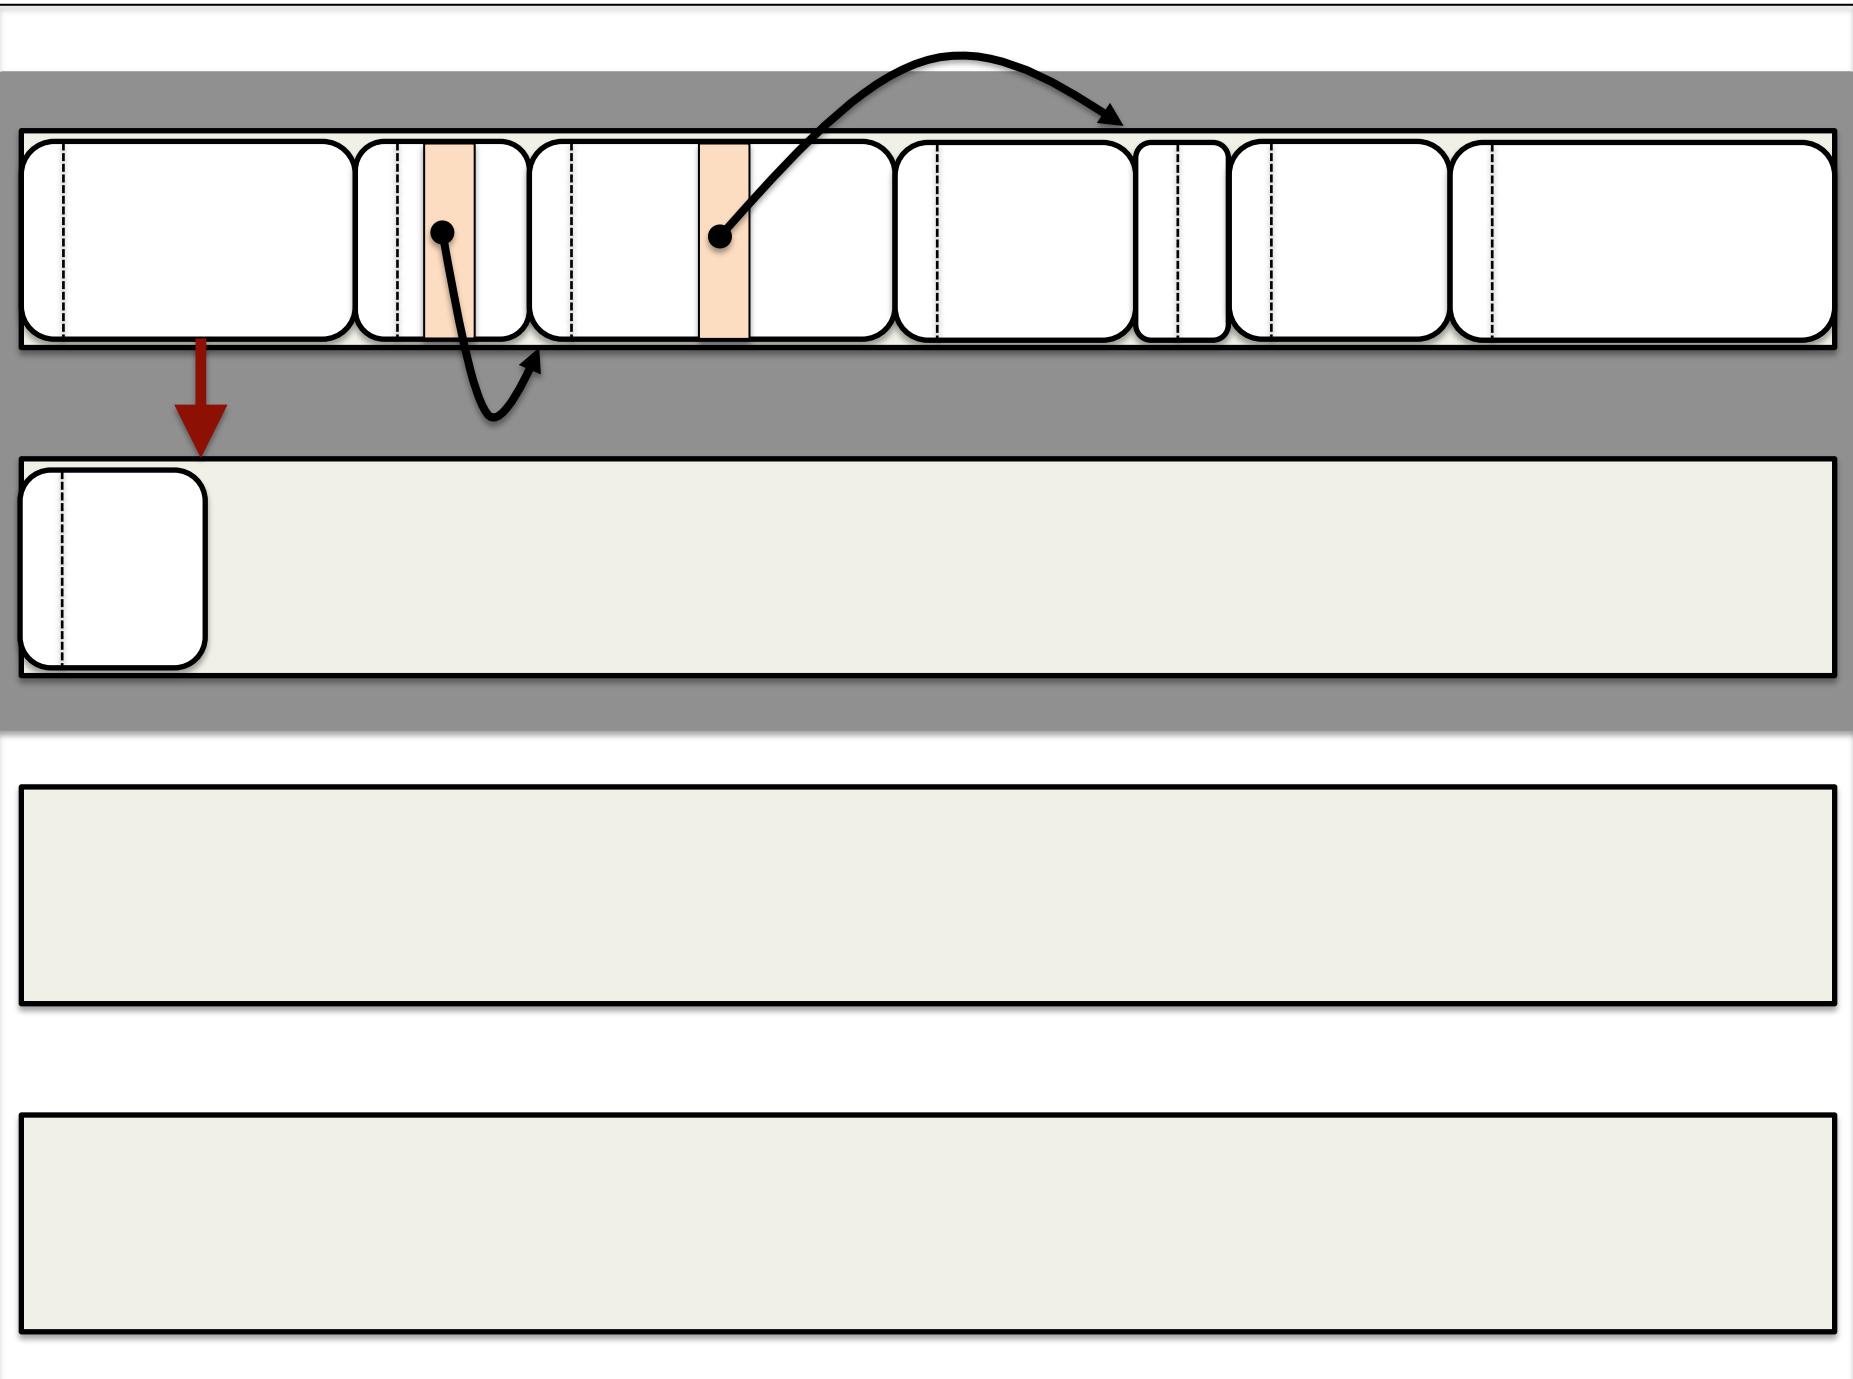


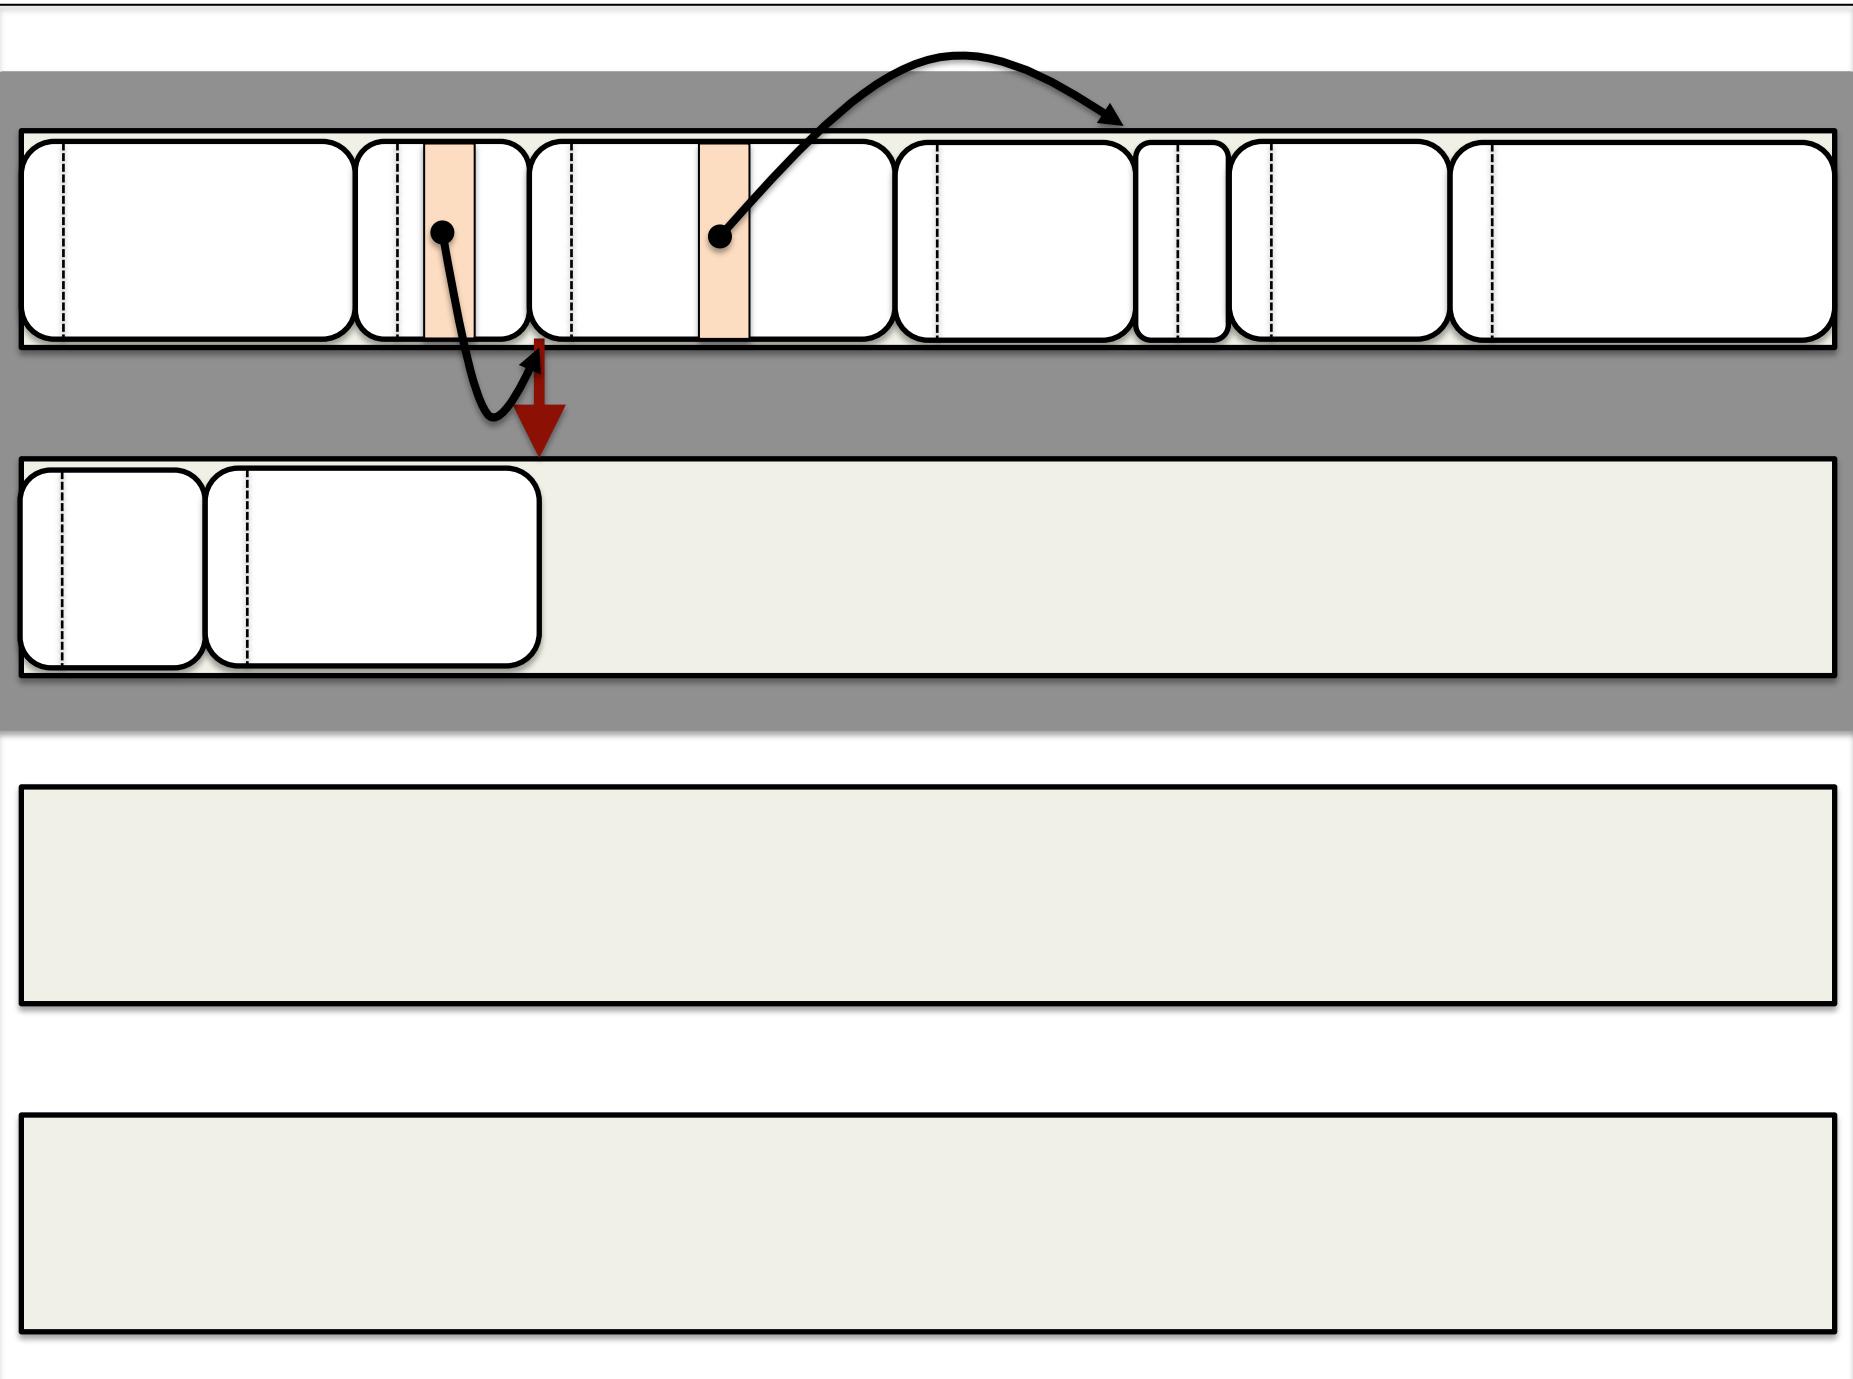
Bump Pointer Allocation

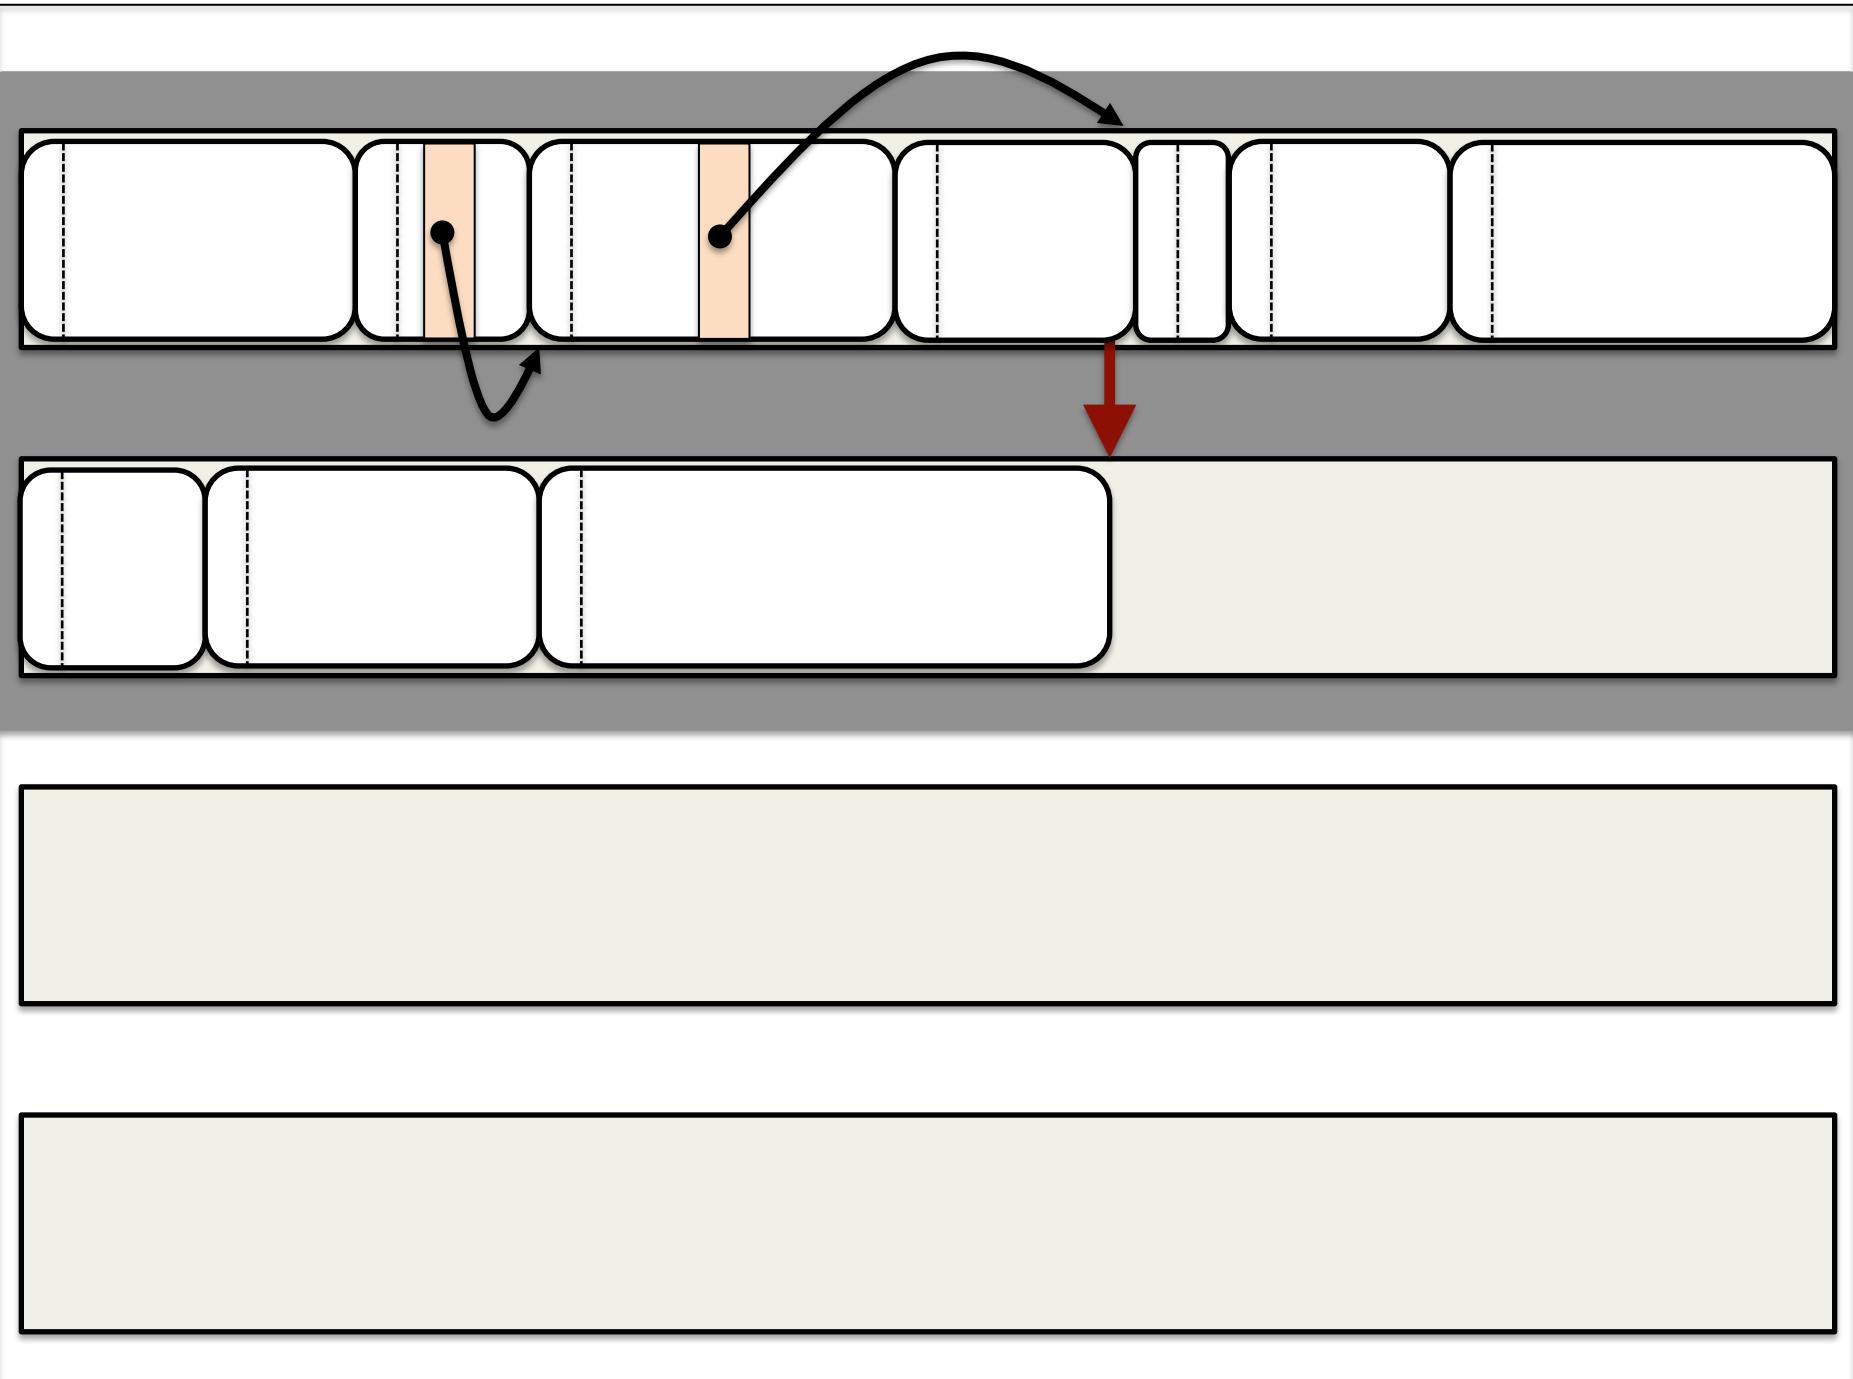
- Each object allocated directly after the previous
 - Only works if there is no fragmentation
- Straightforward algorithm
 - No heuristics as in previous allocators
- Allocation is very cheap

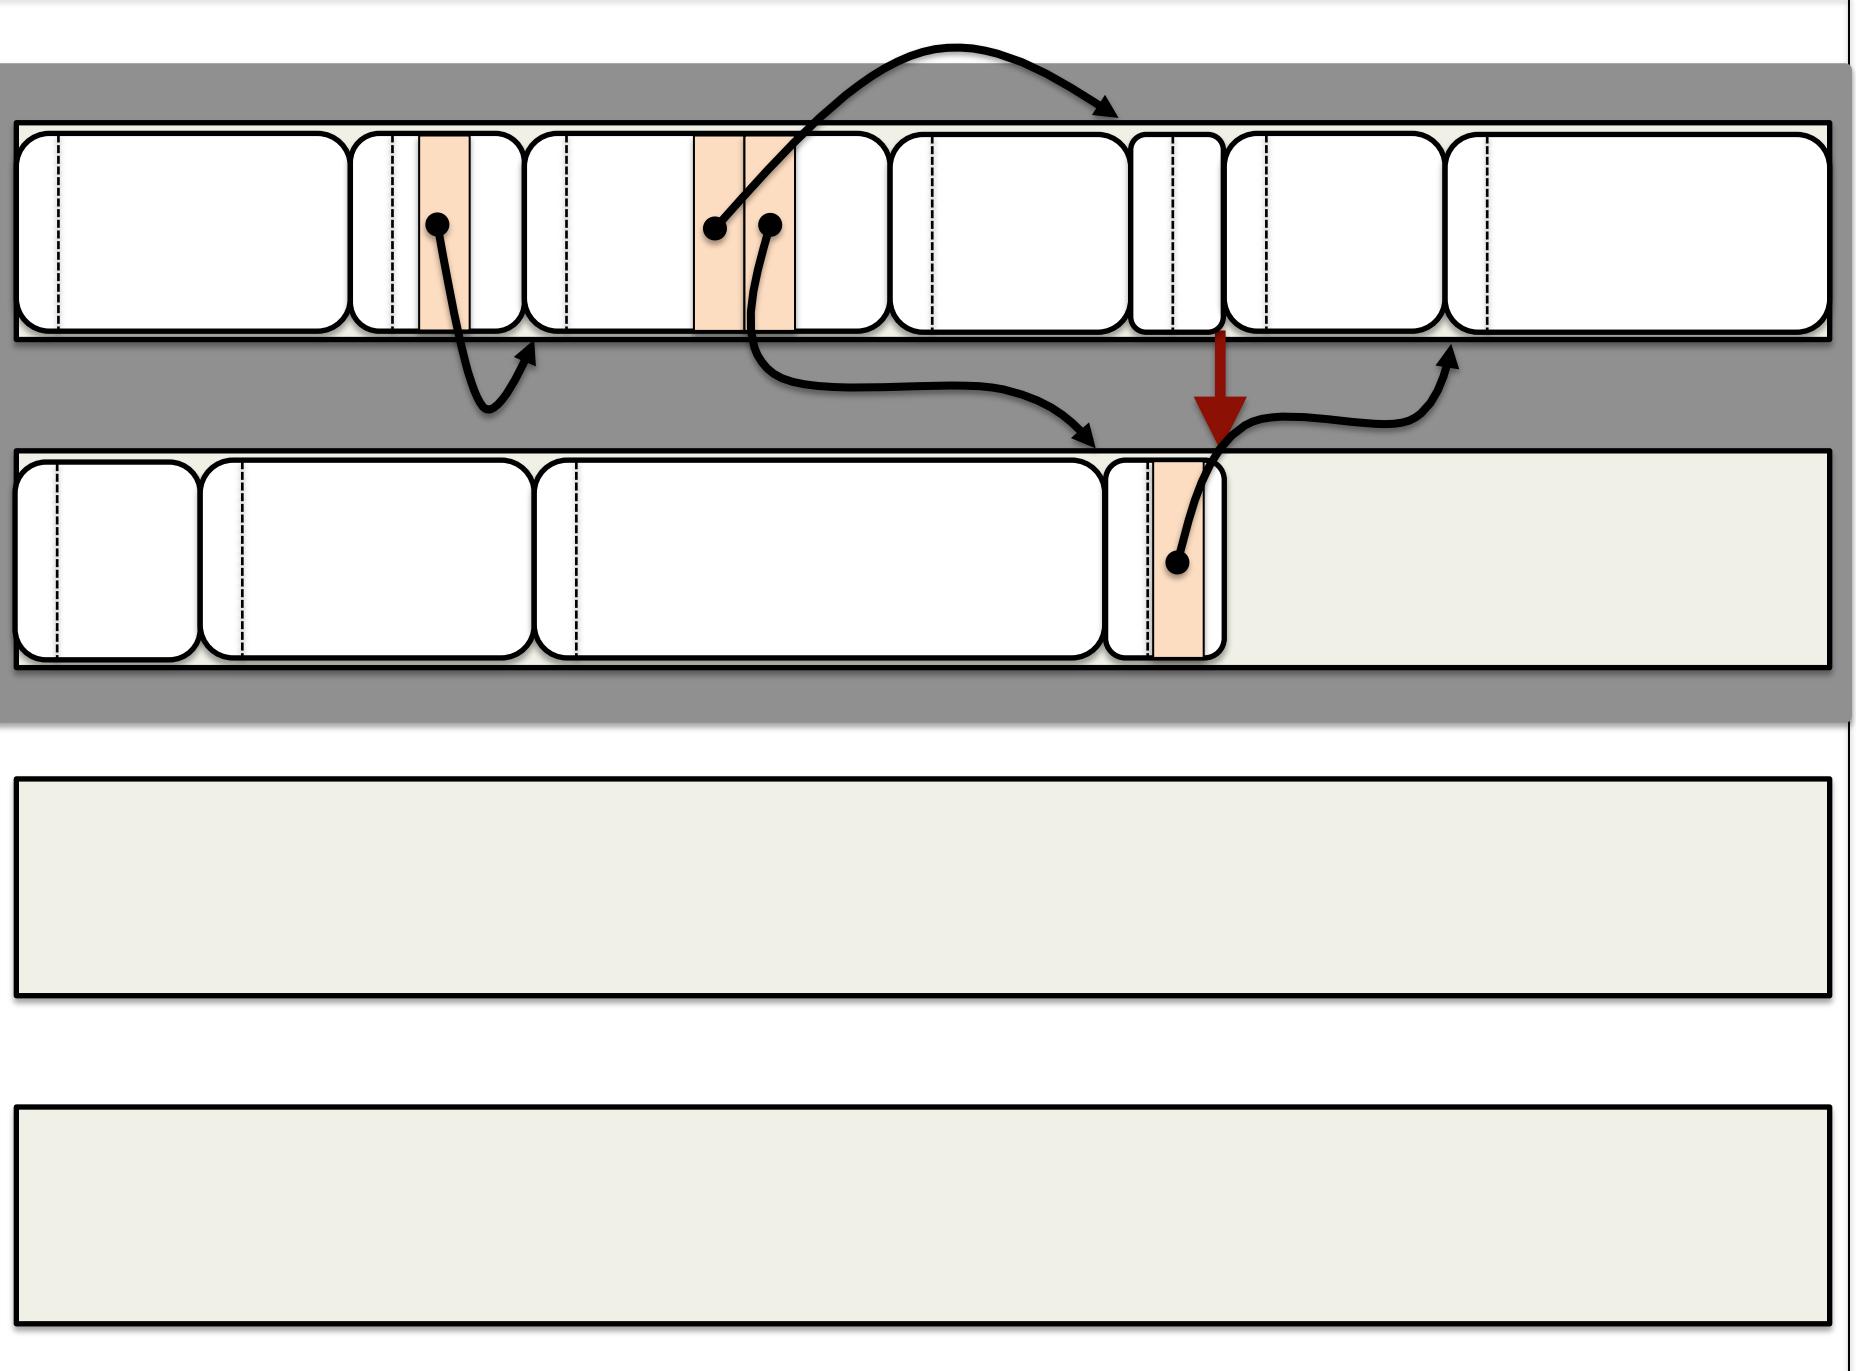
Bump Pointer Characteristics

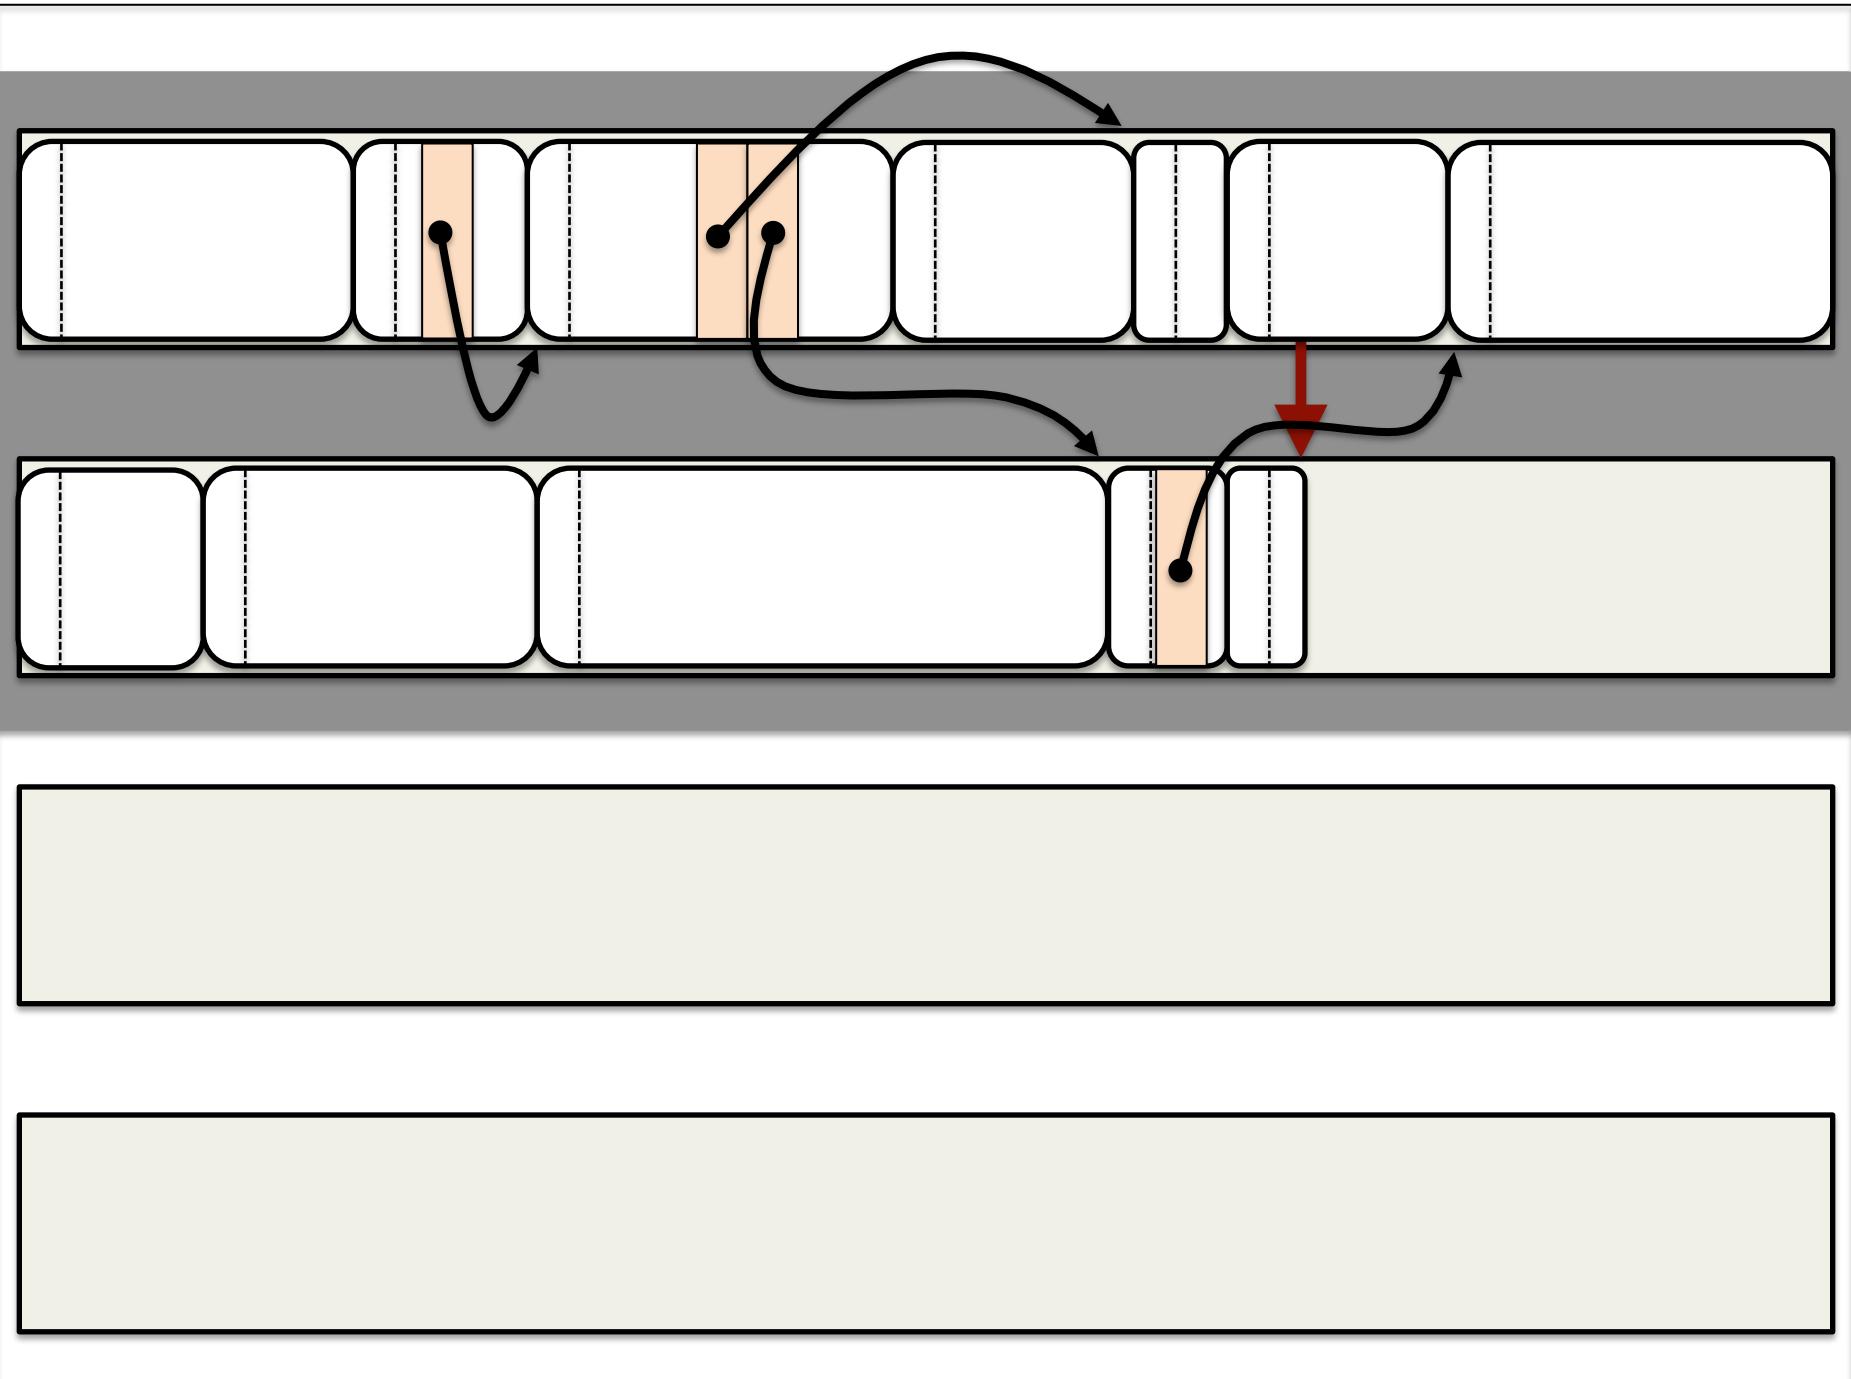
- Fast path
 - Taken when there is enough memory available
 - Allocation is an increment and comparison
 - Easily inlined and optimized
- Slow path
 - Taken when comparison fails
 - Get new memory region from heap
 - Trigger garbage collection if none available

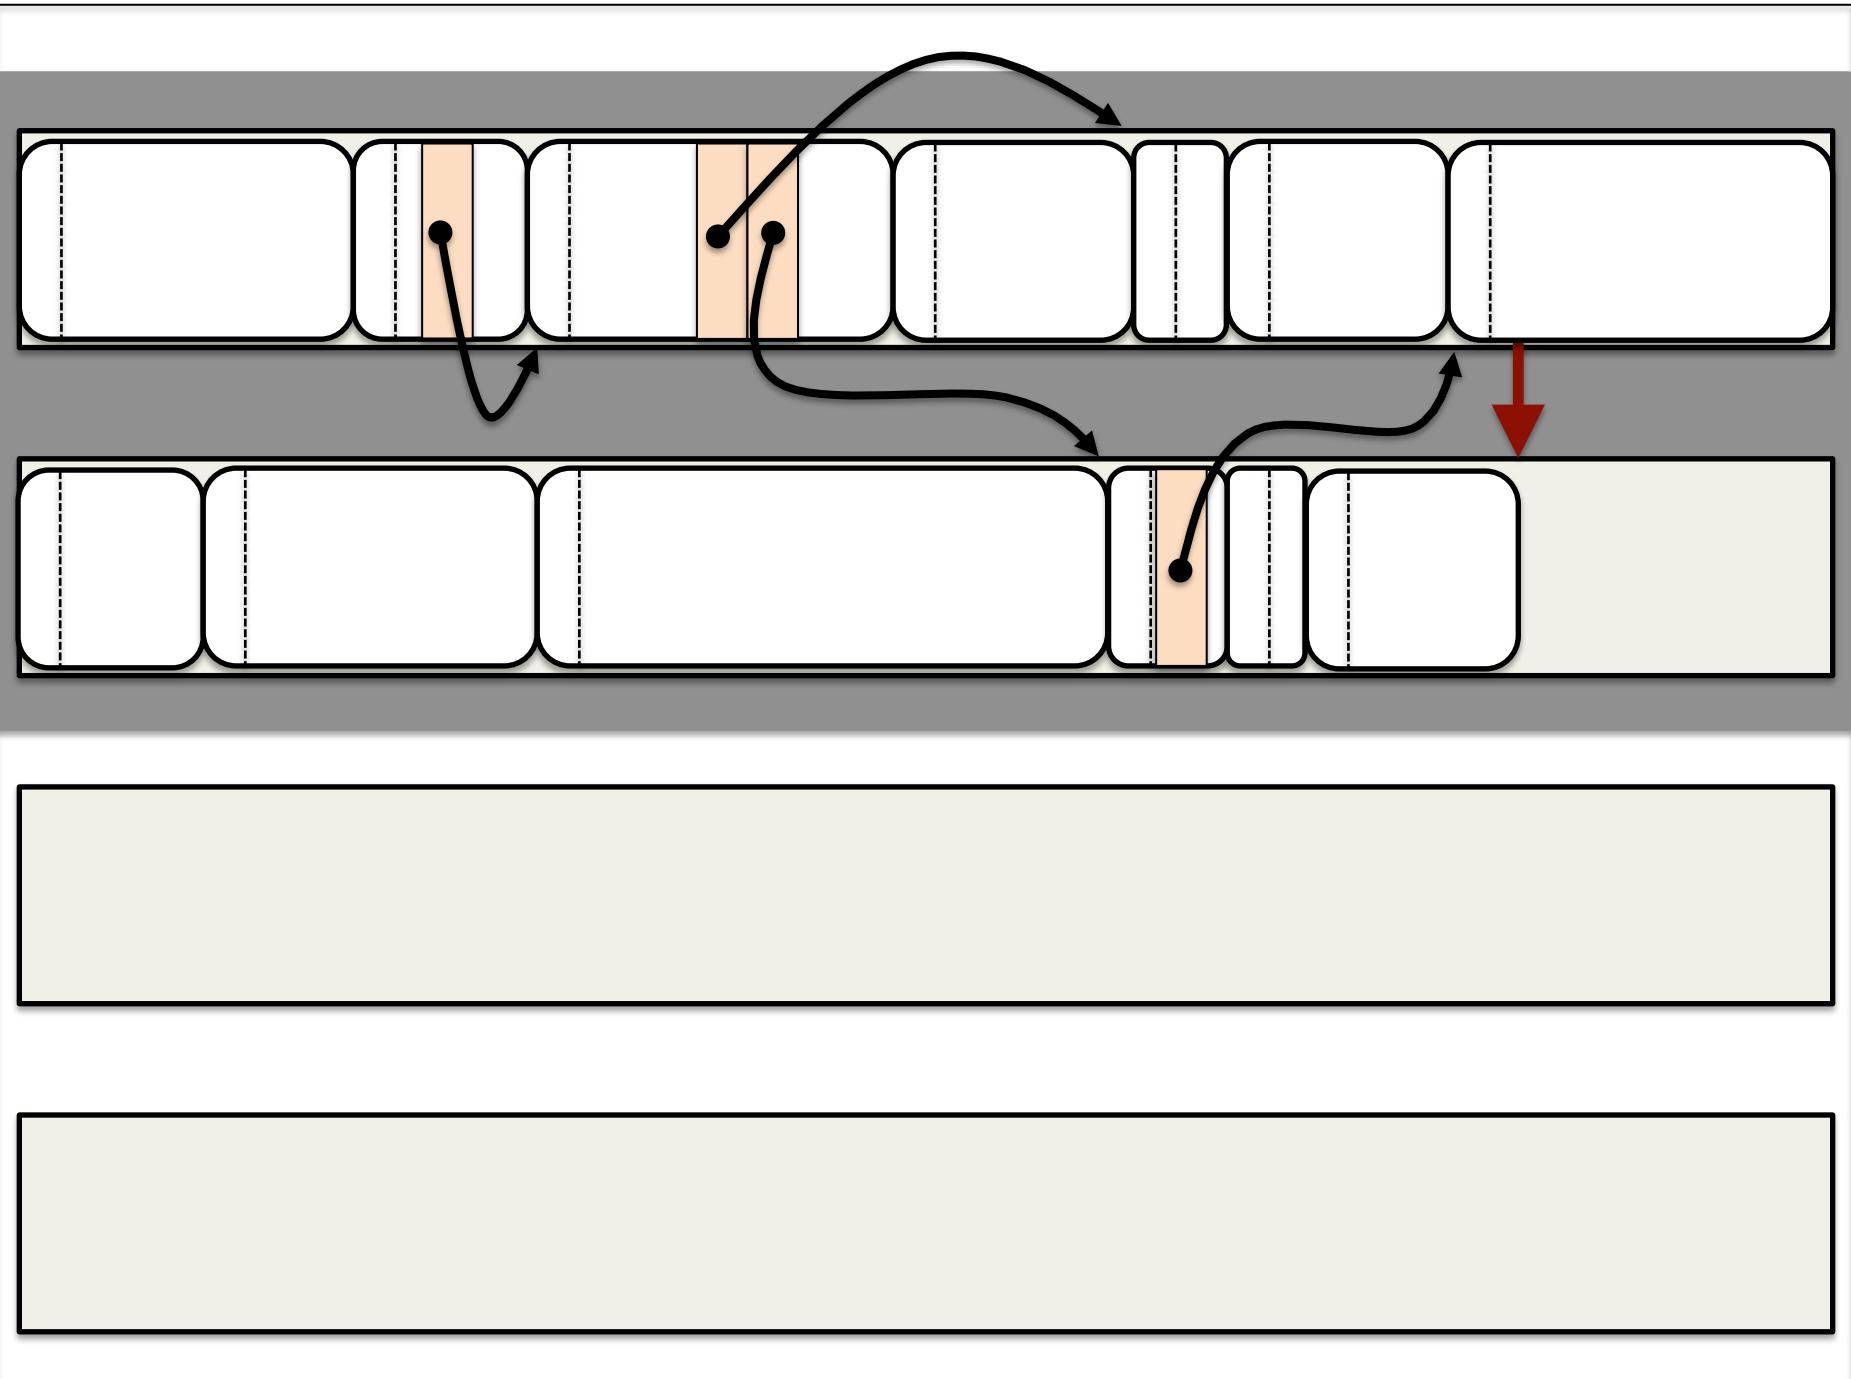


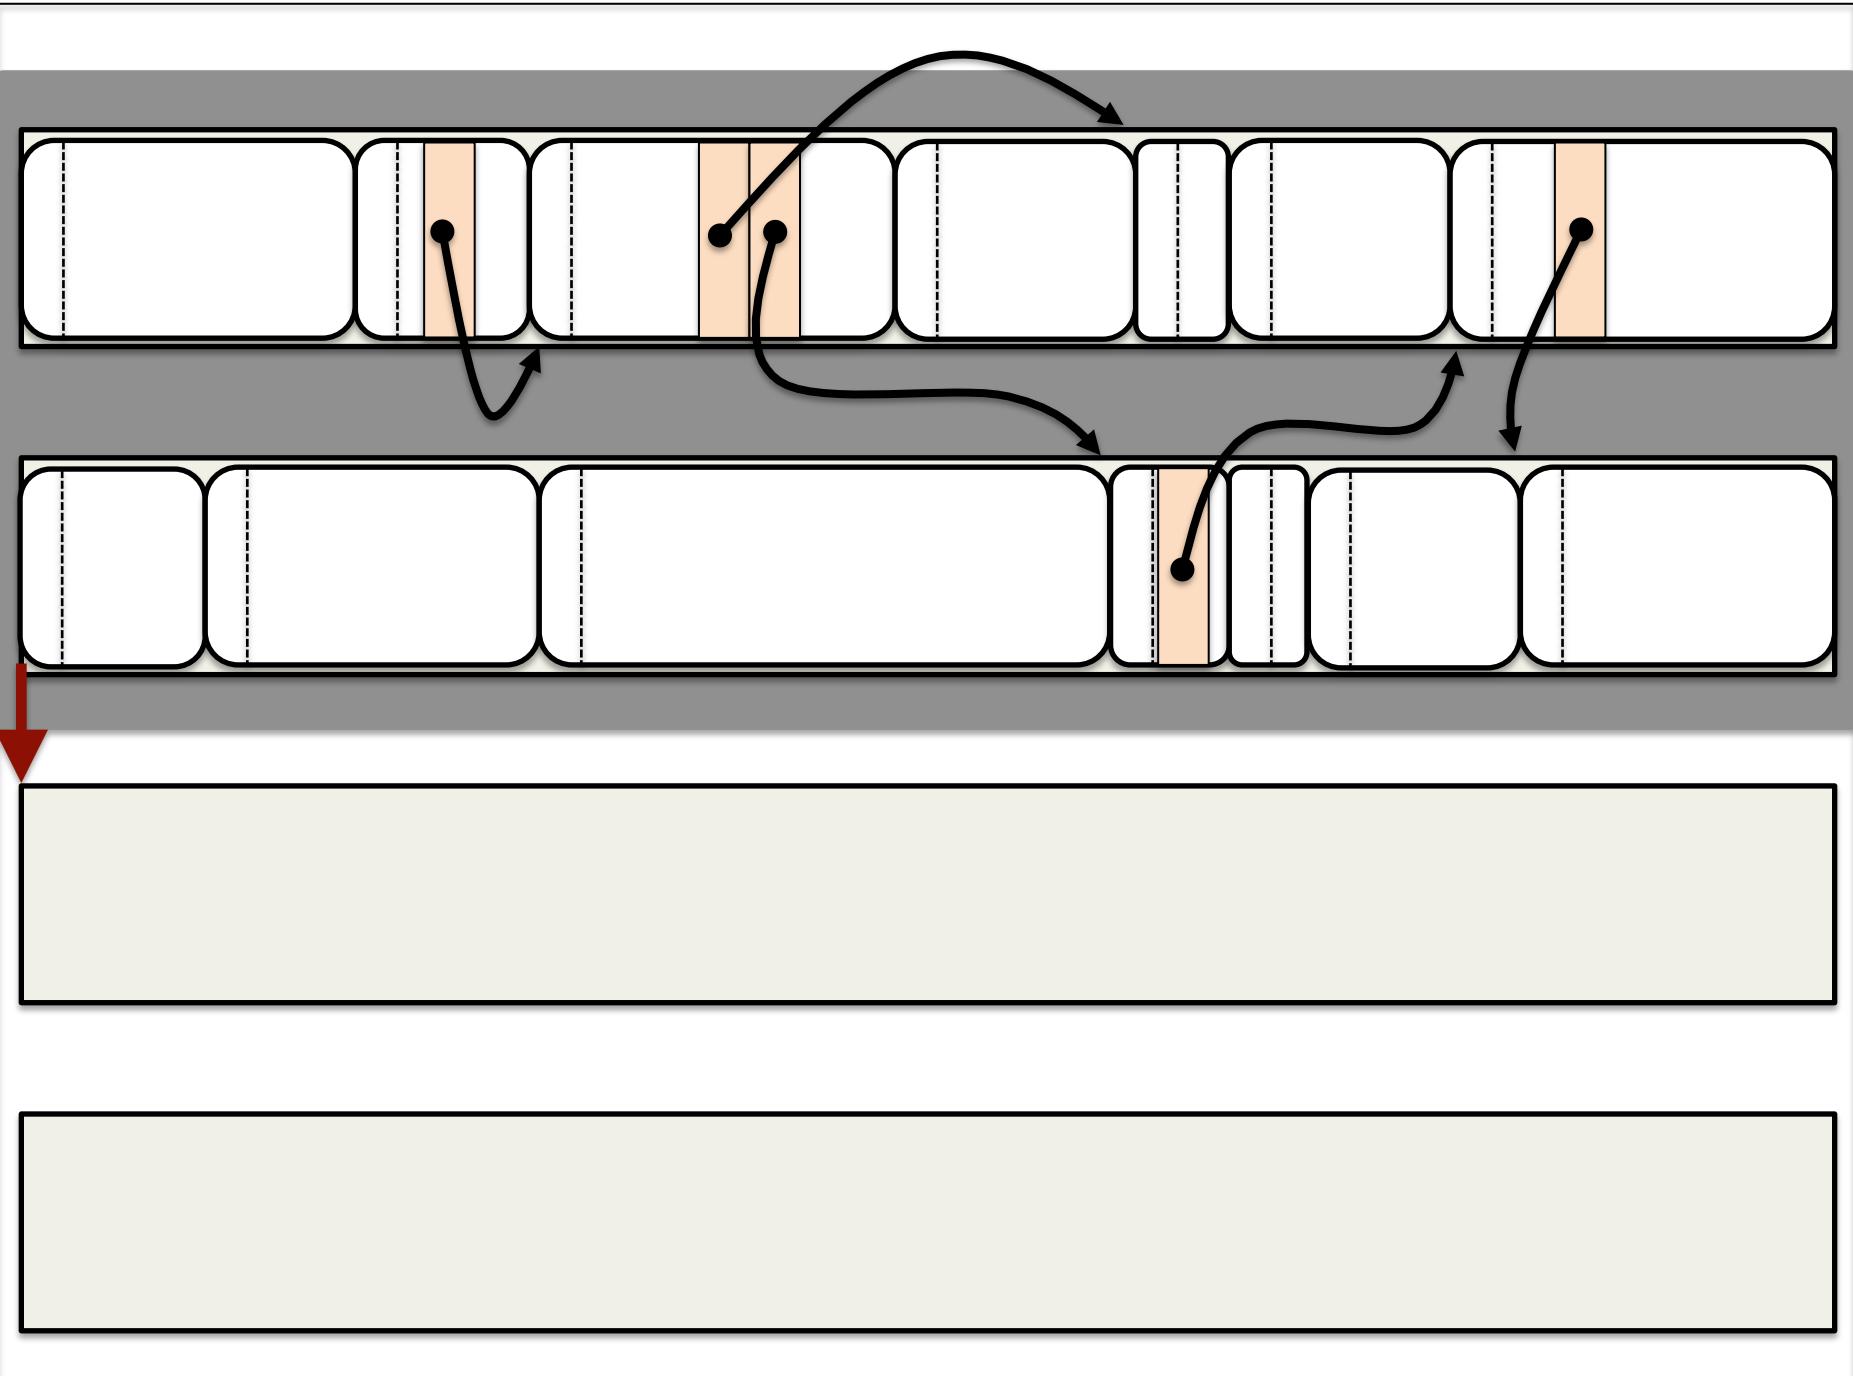






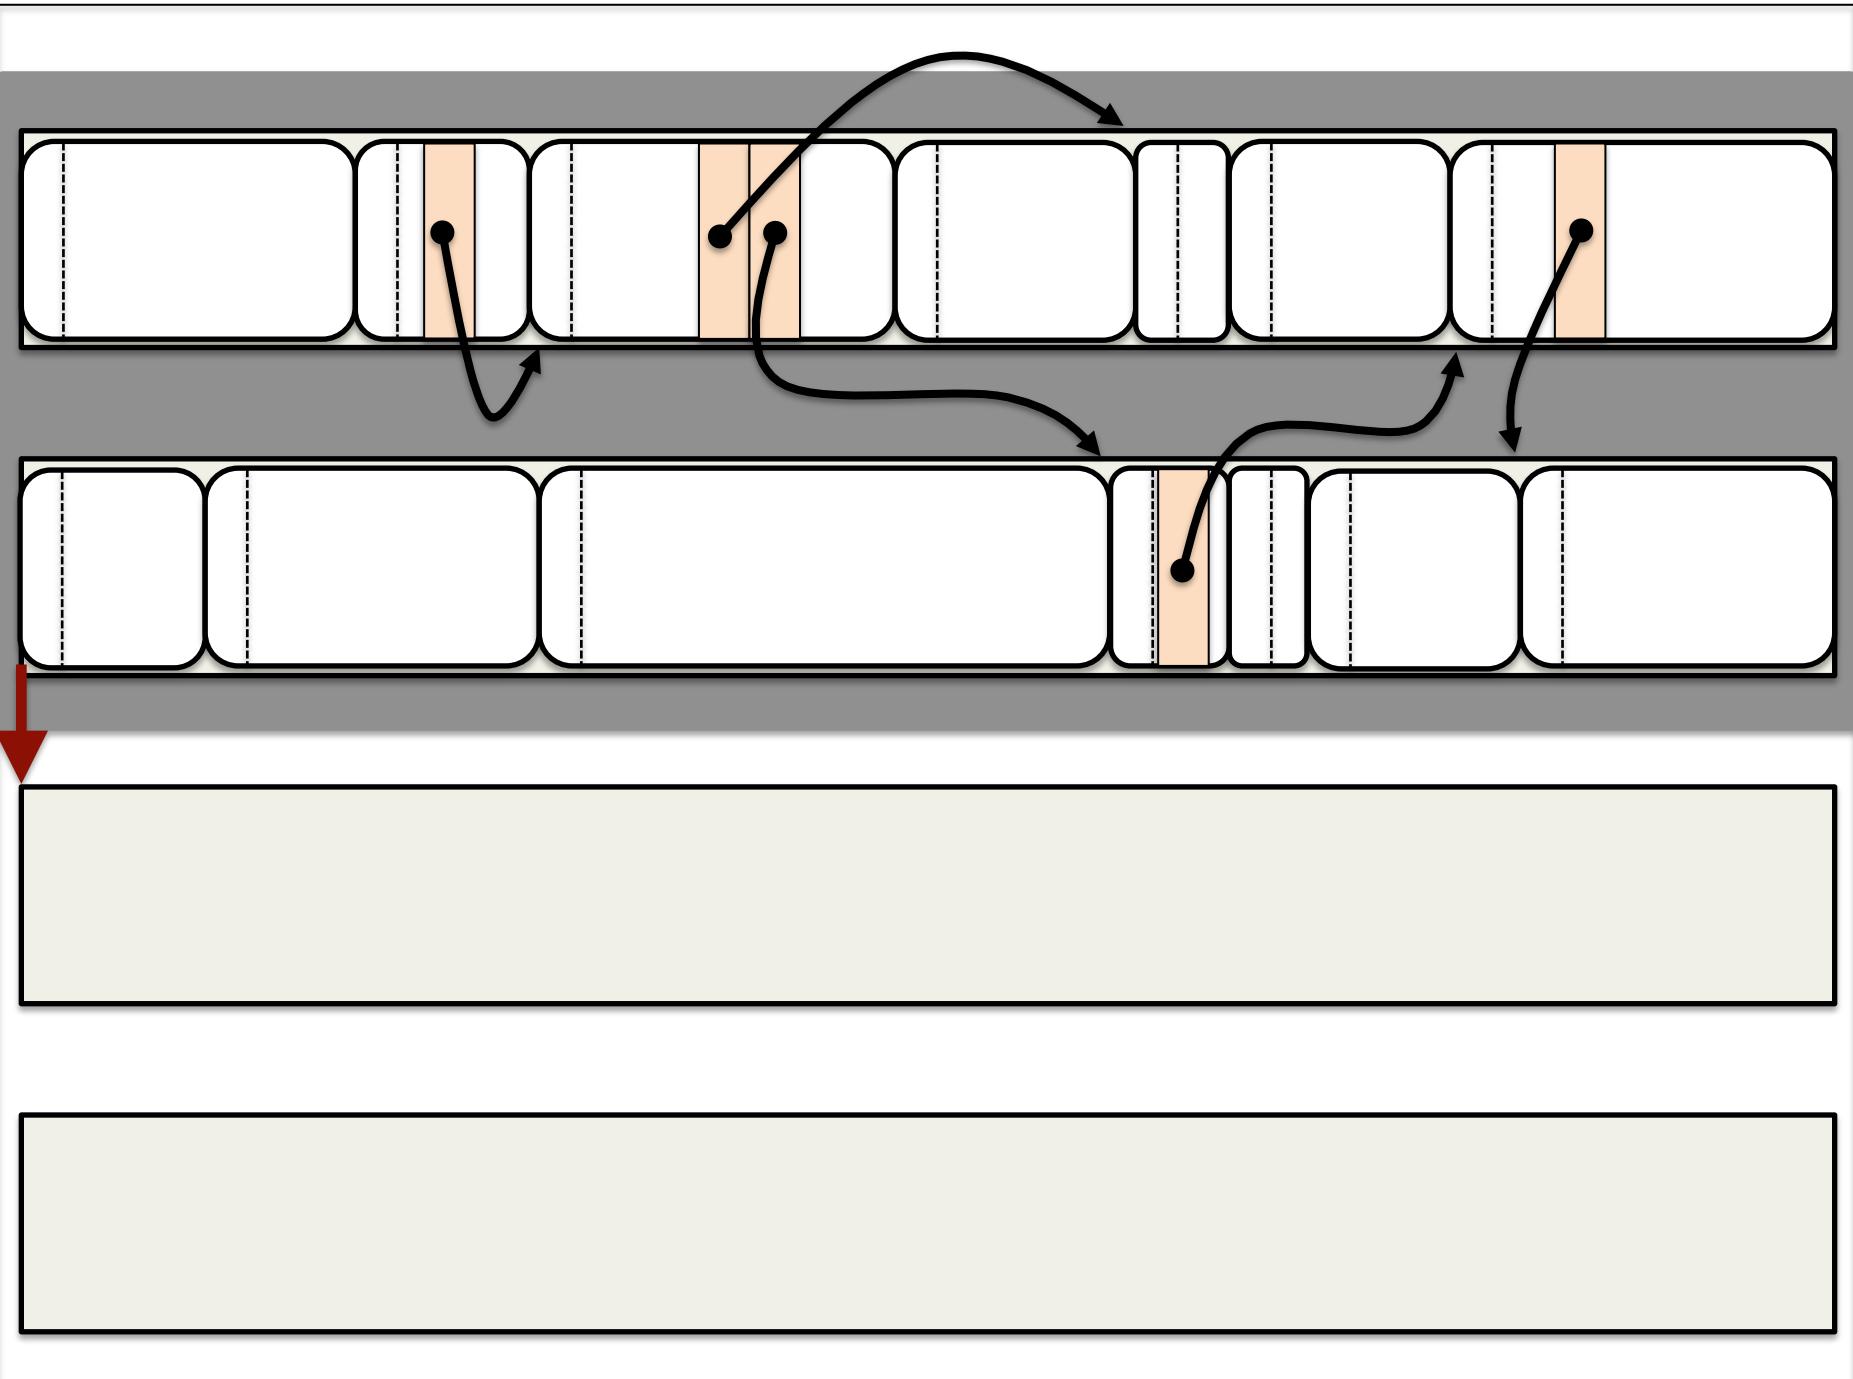


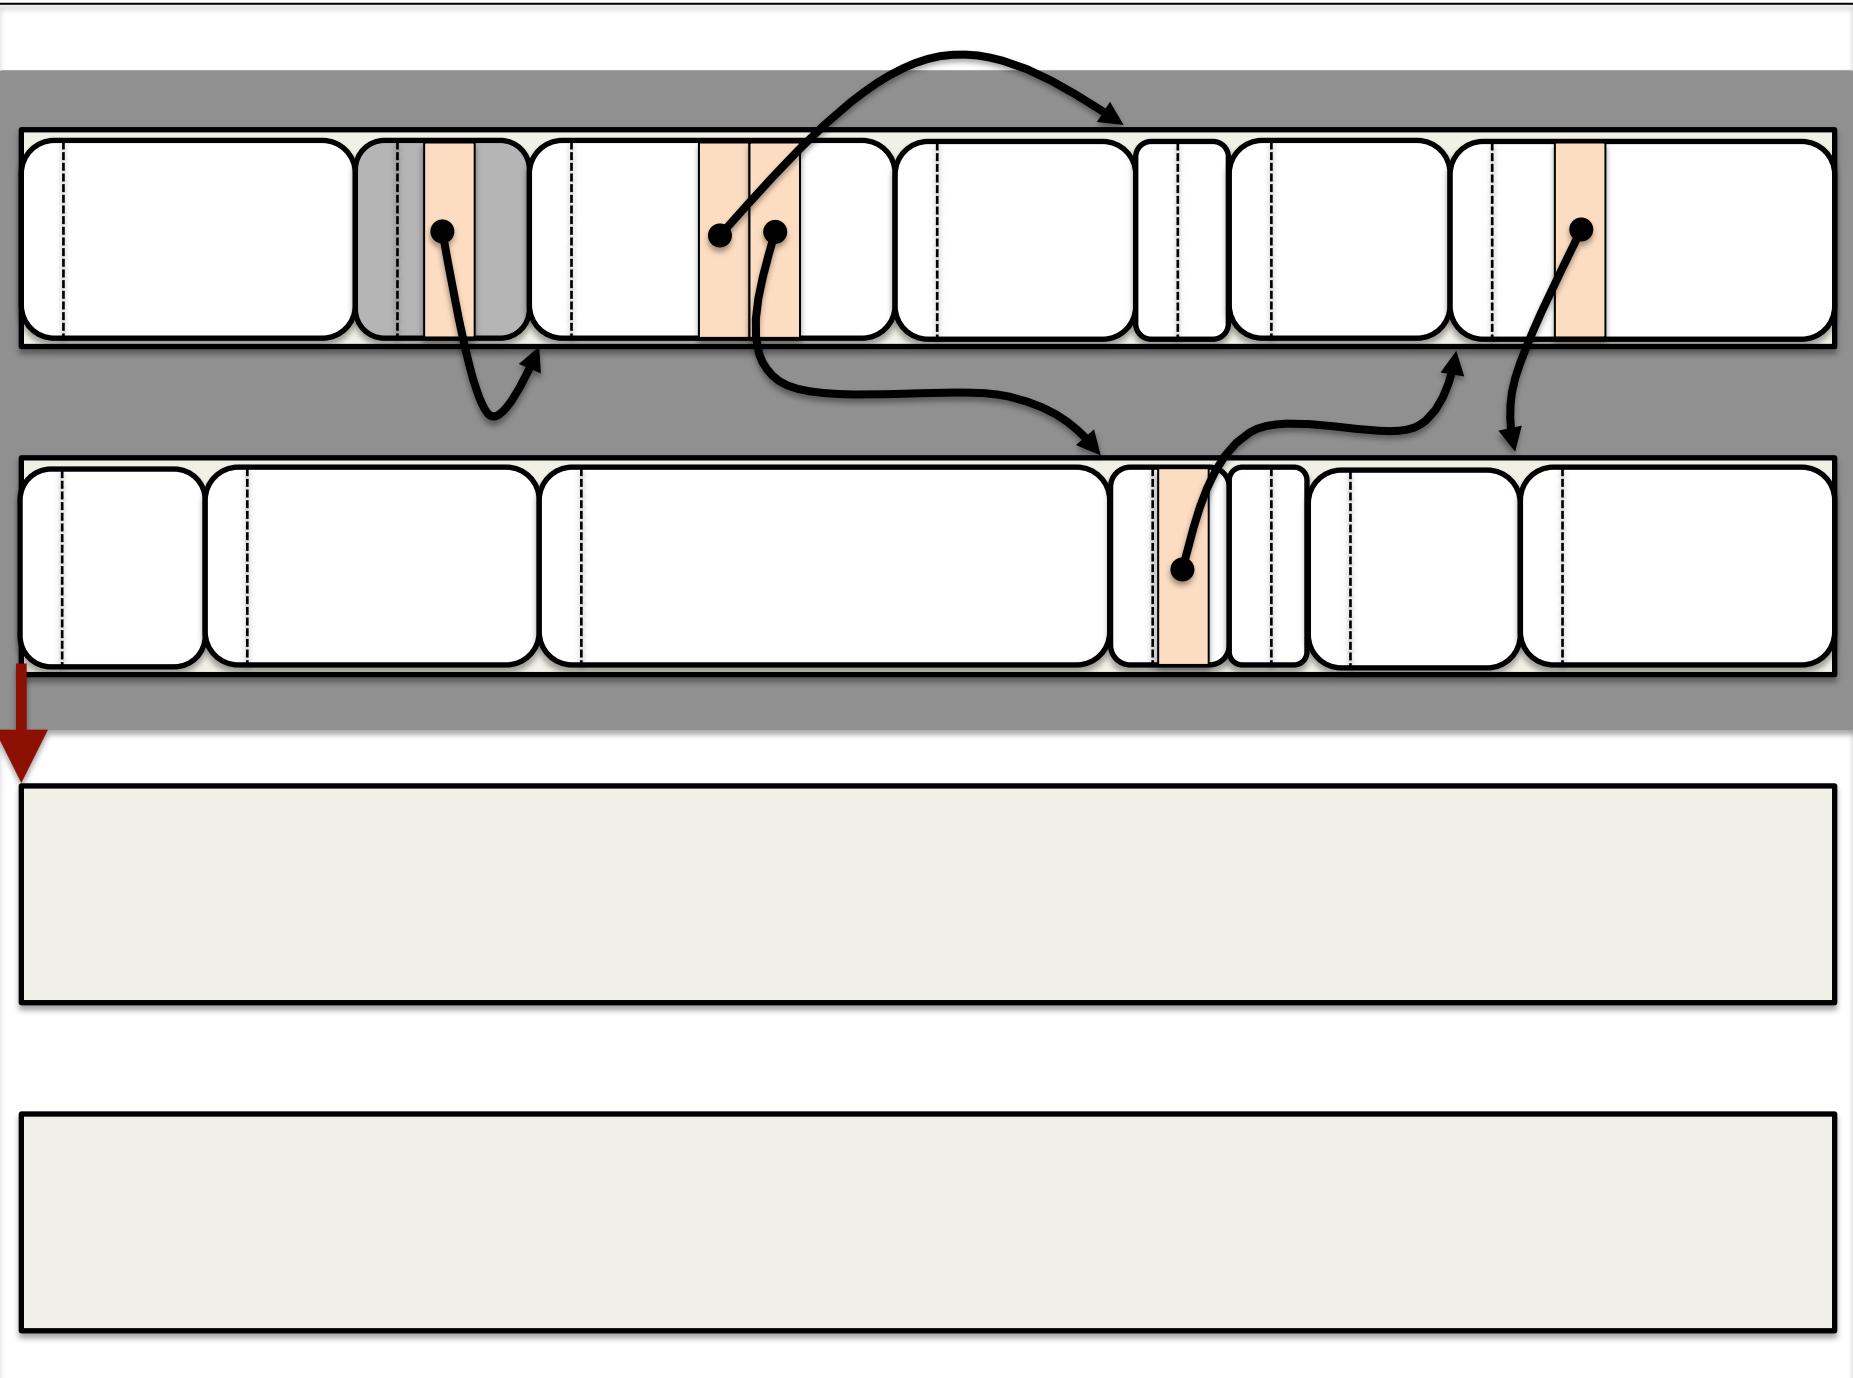


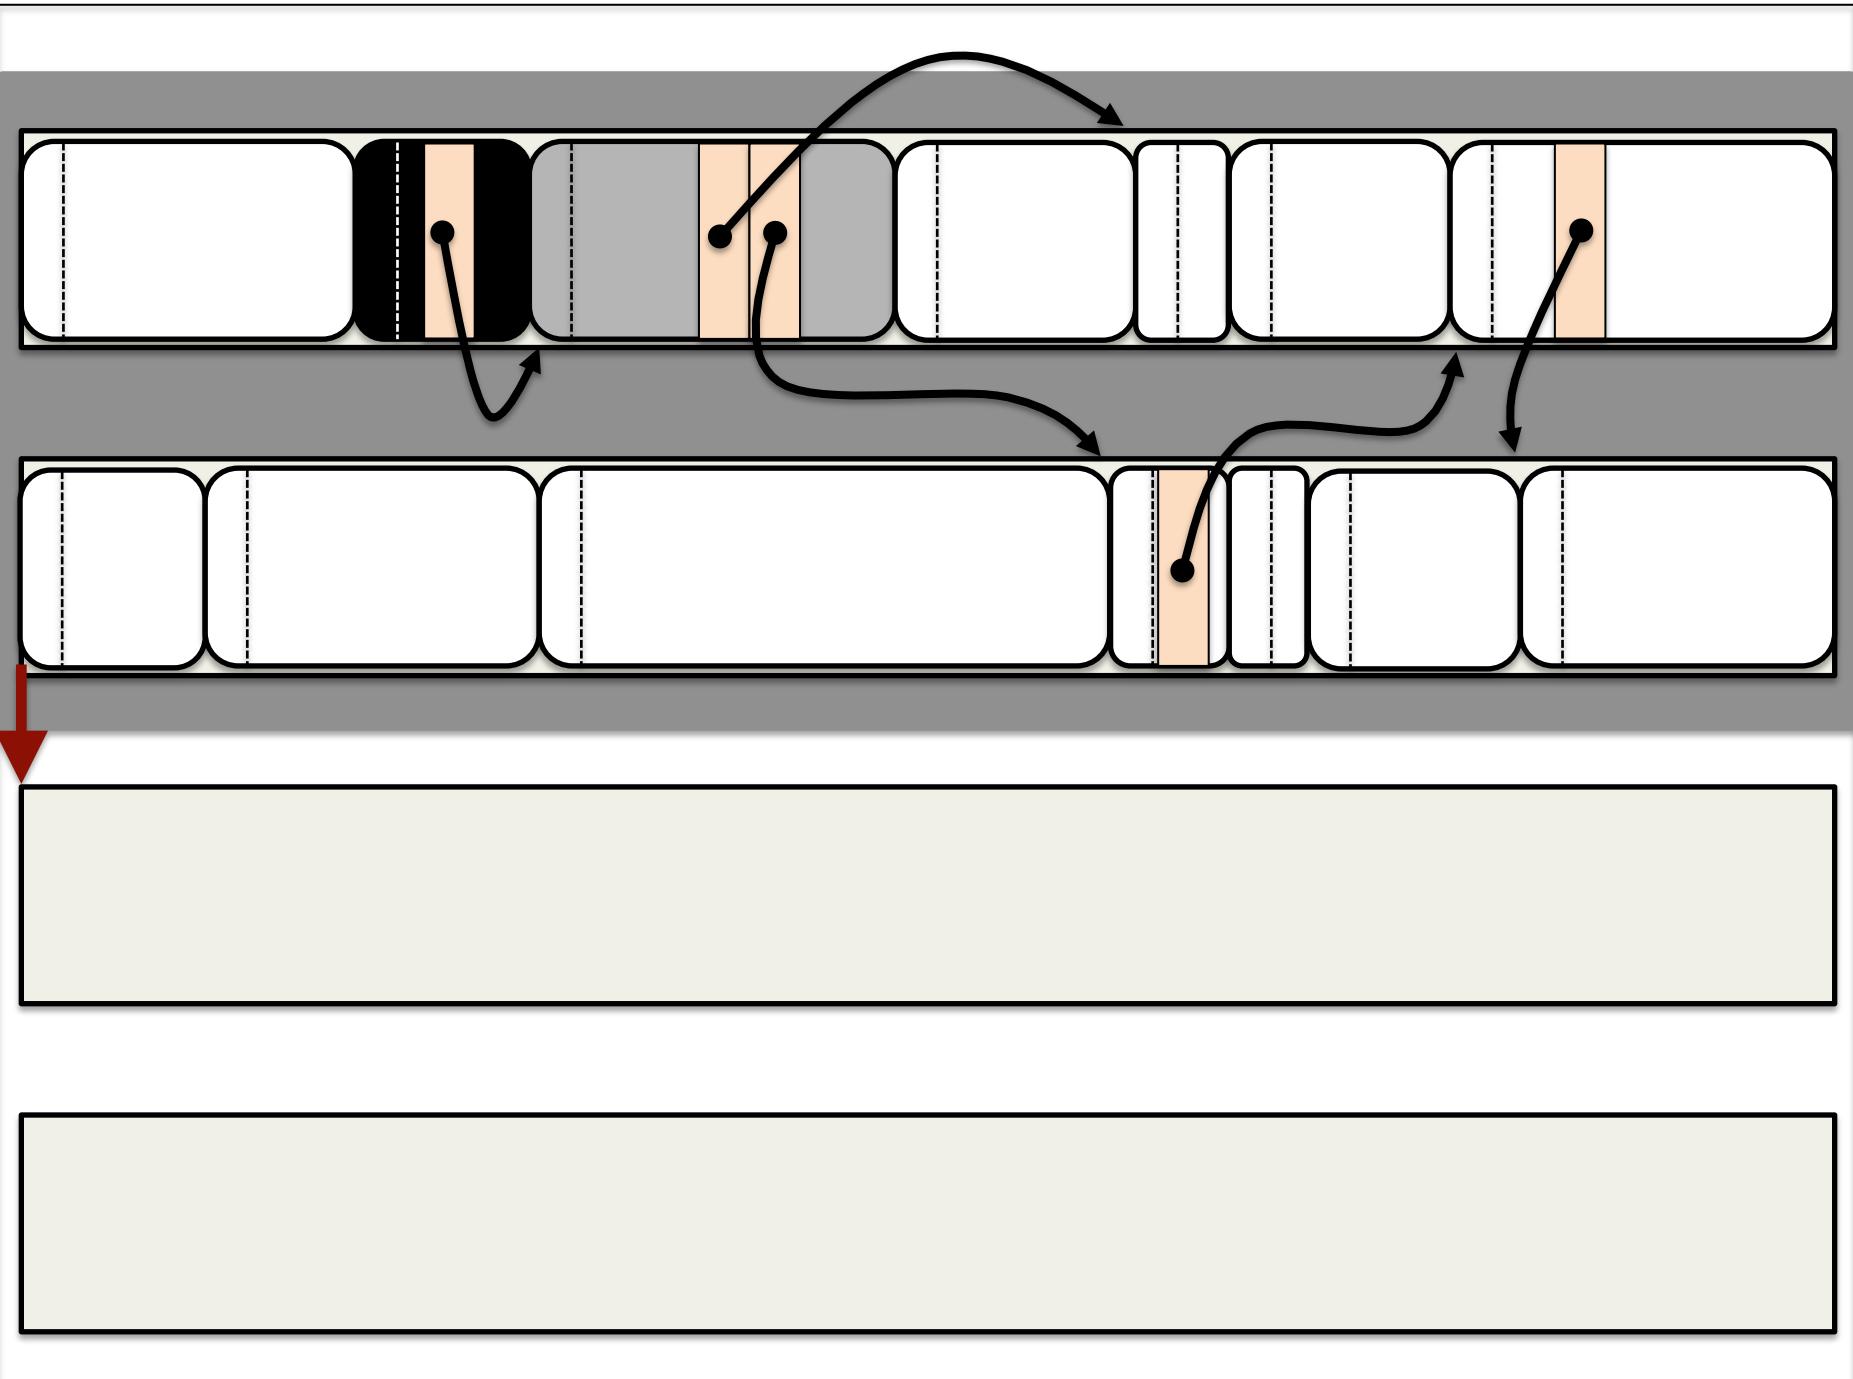


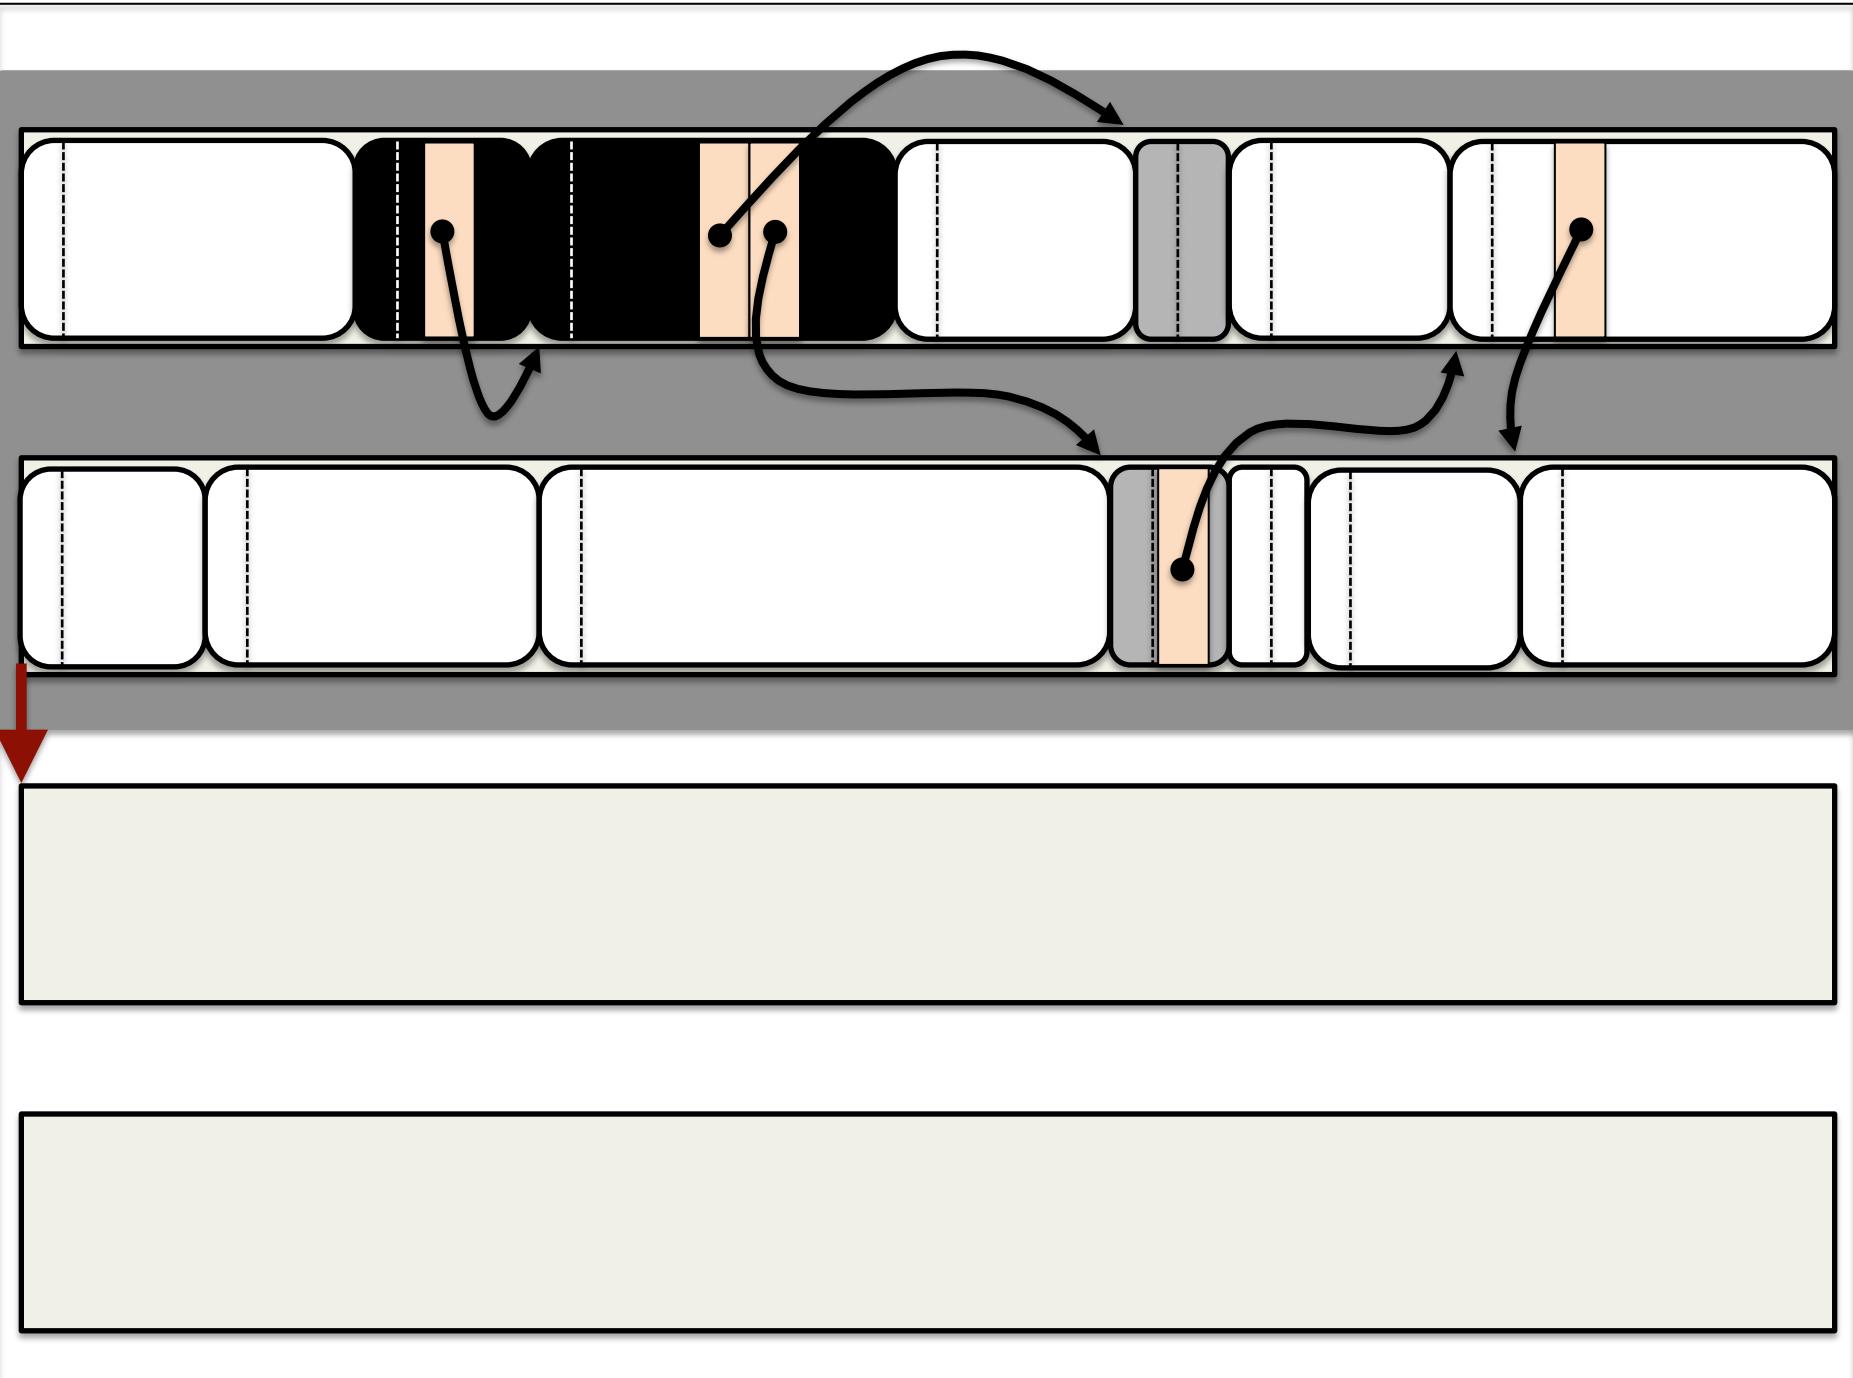
Tracing

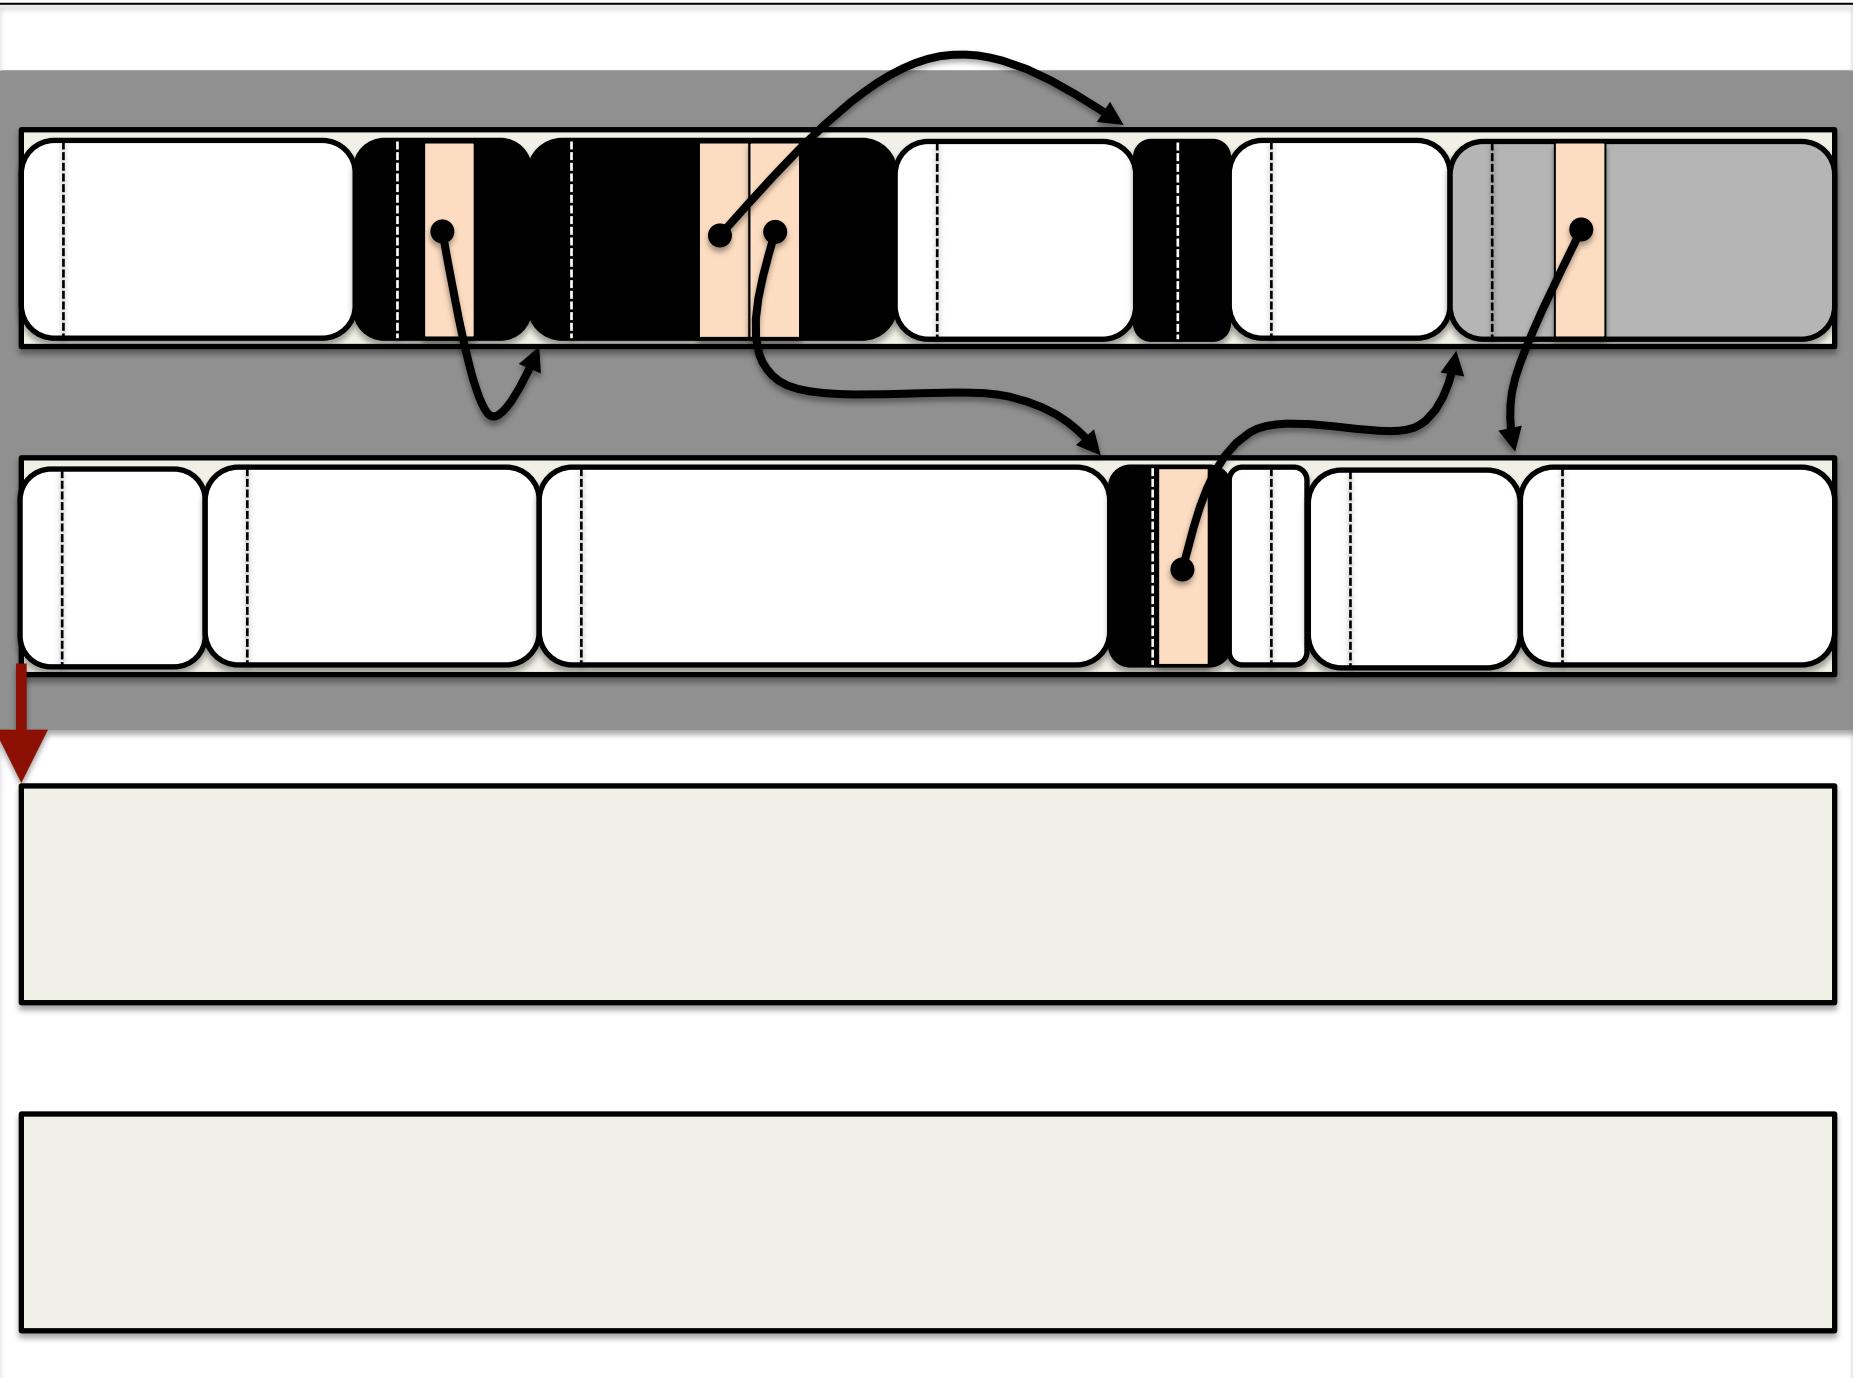
- Trace process the same as in mark and sweep
 - We'll see a more specialized algorithm later
- Start with the system roots
- Use tricolor algorithm to find all live objects
- Result is that reachable objects are black

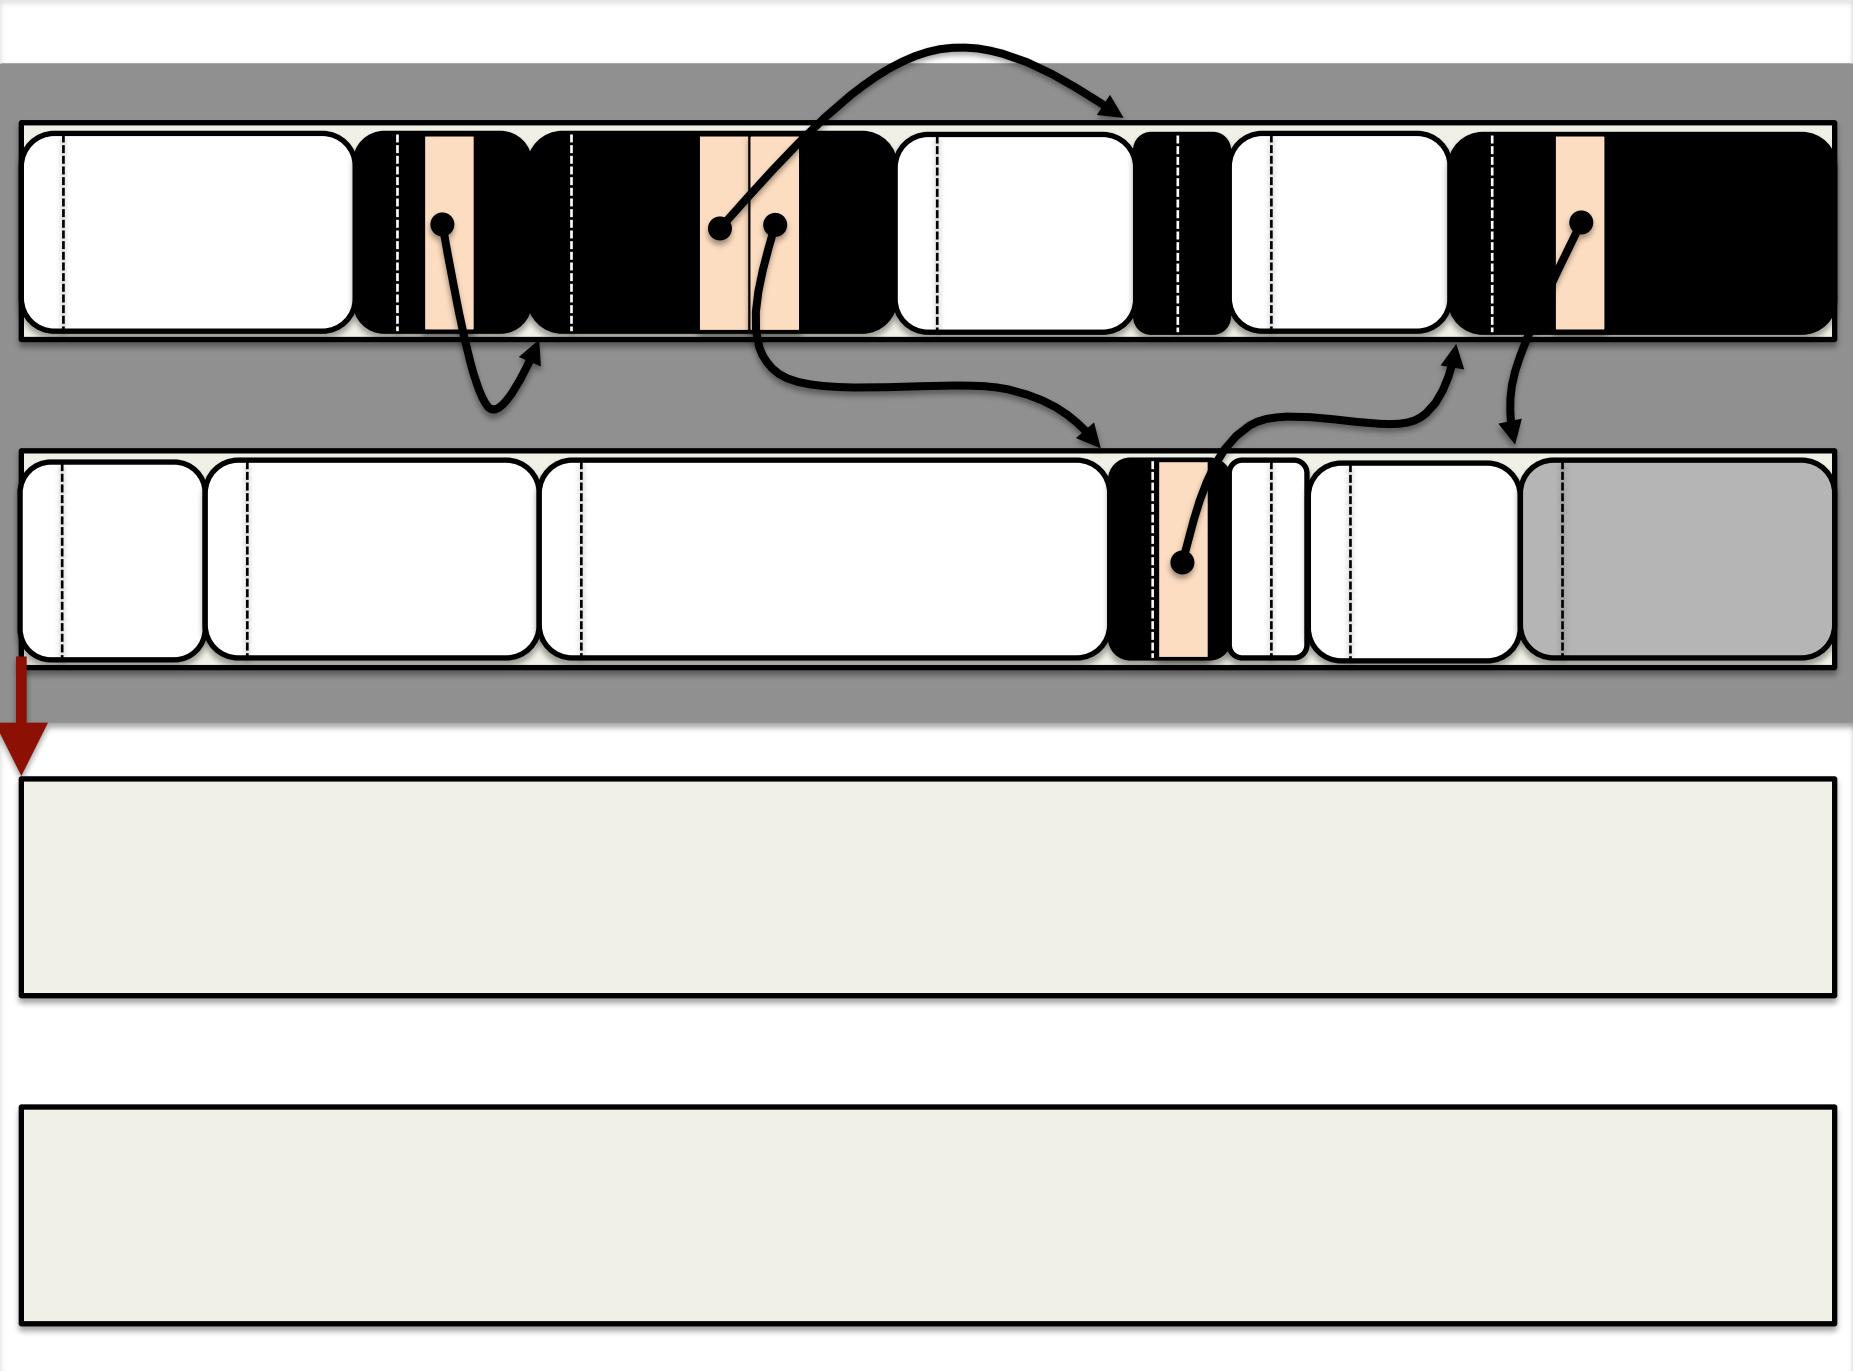


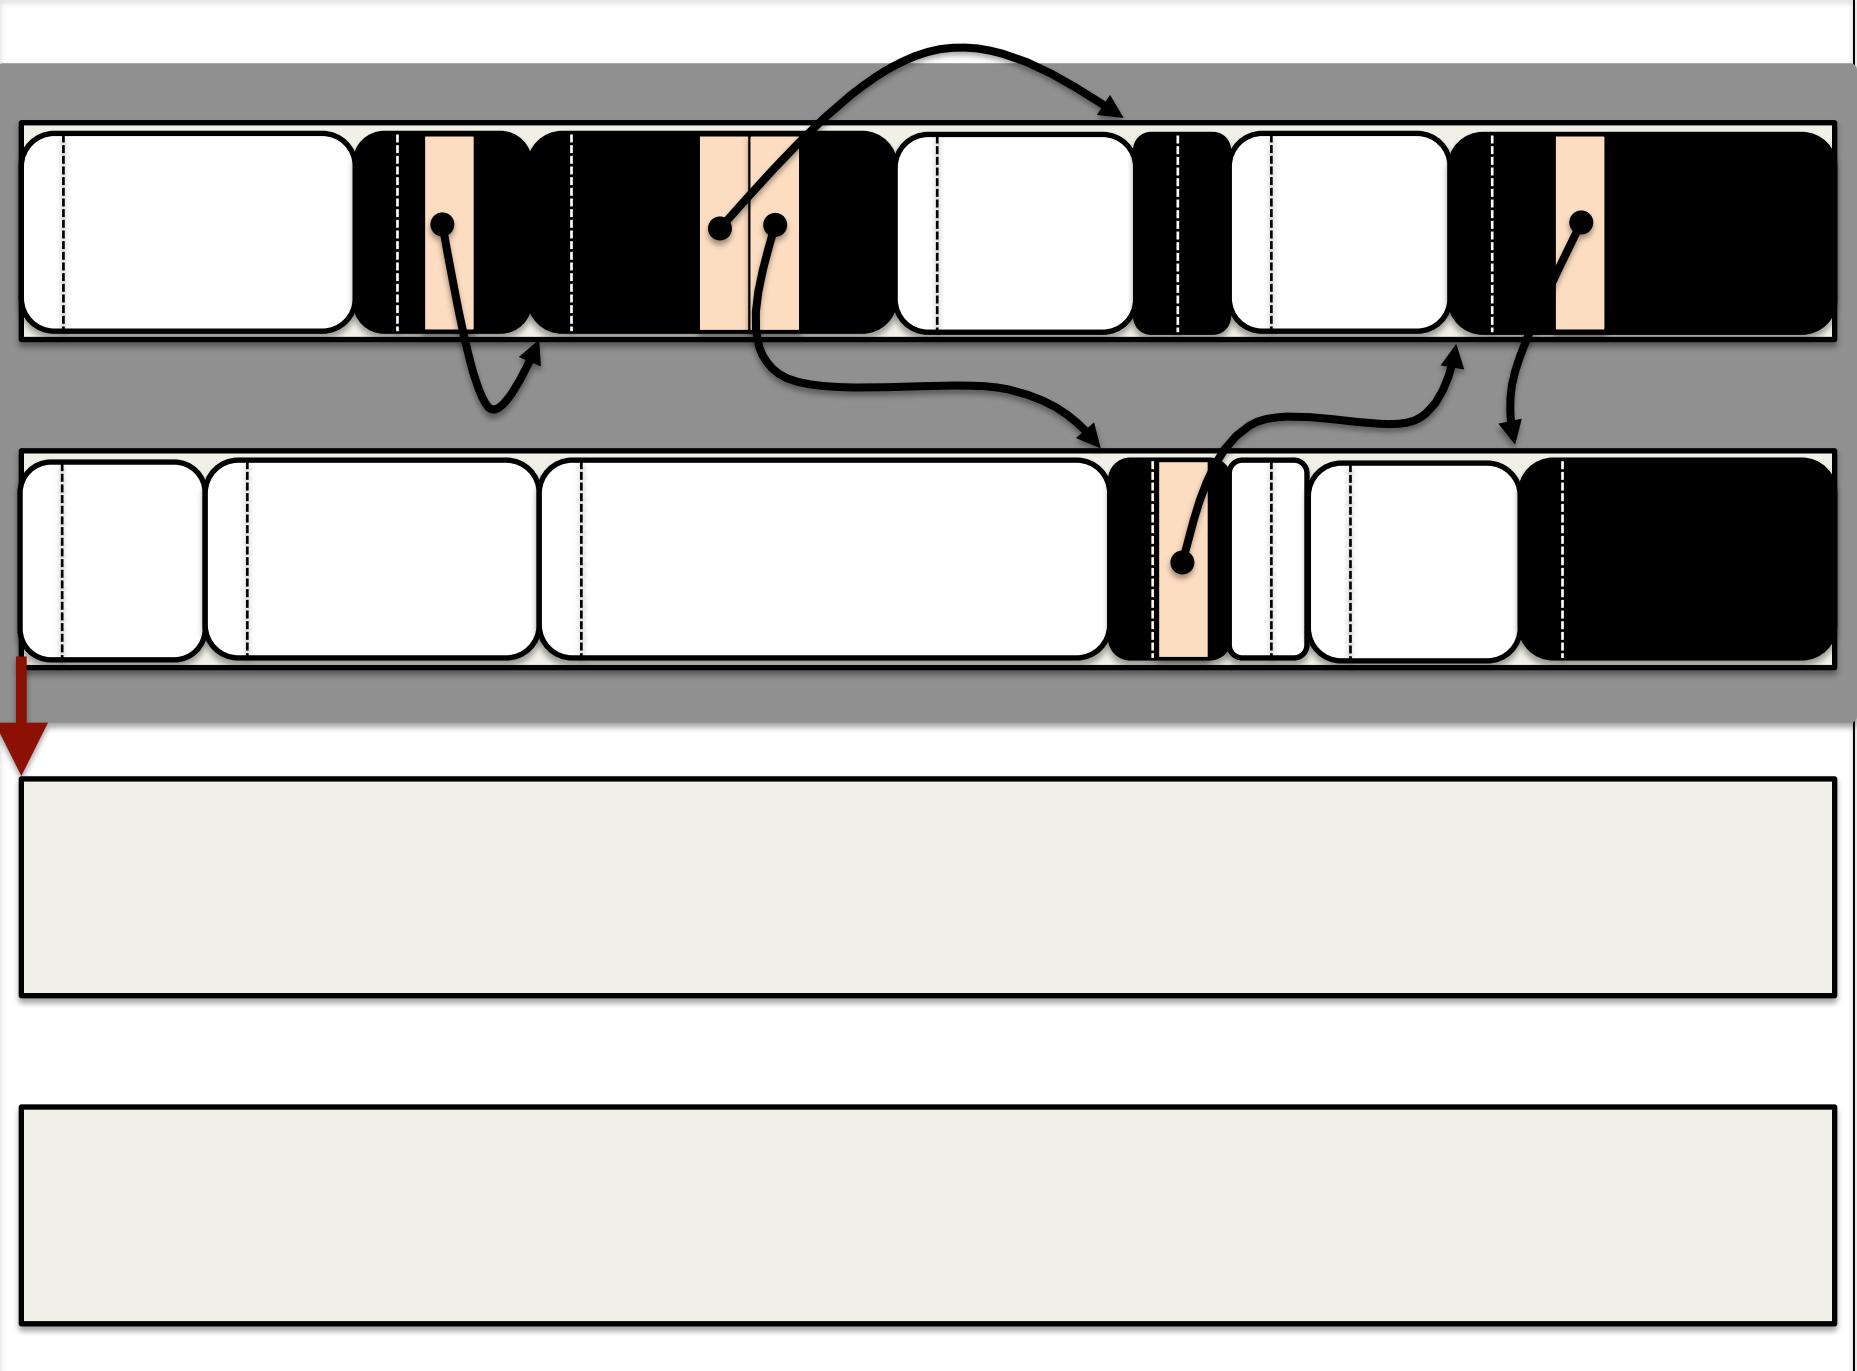






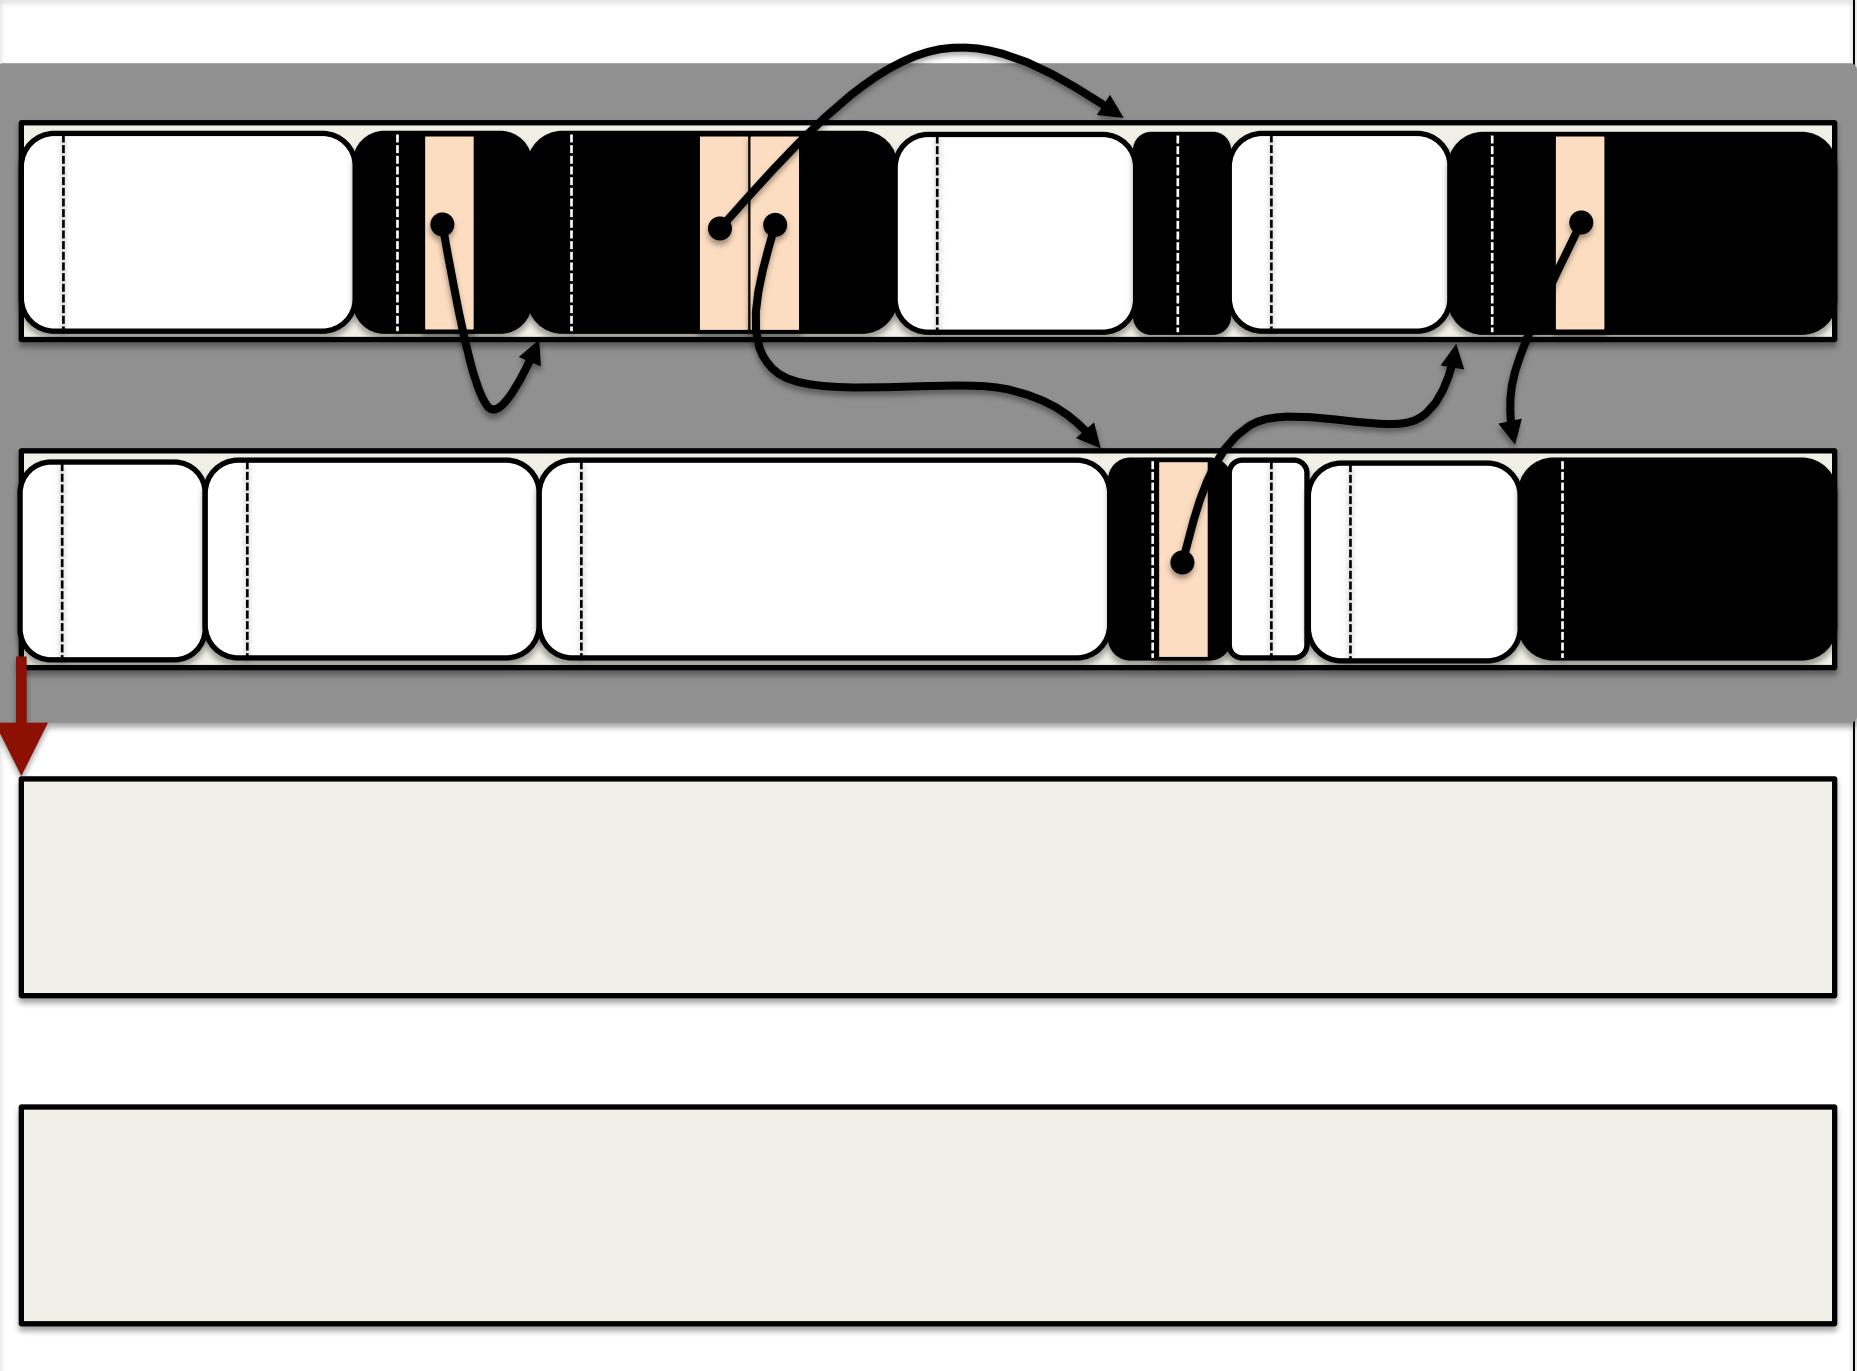


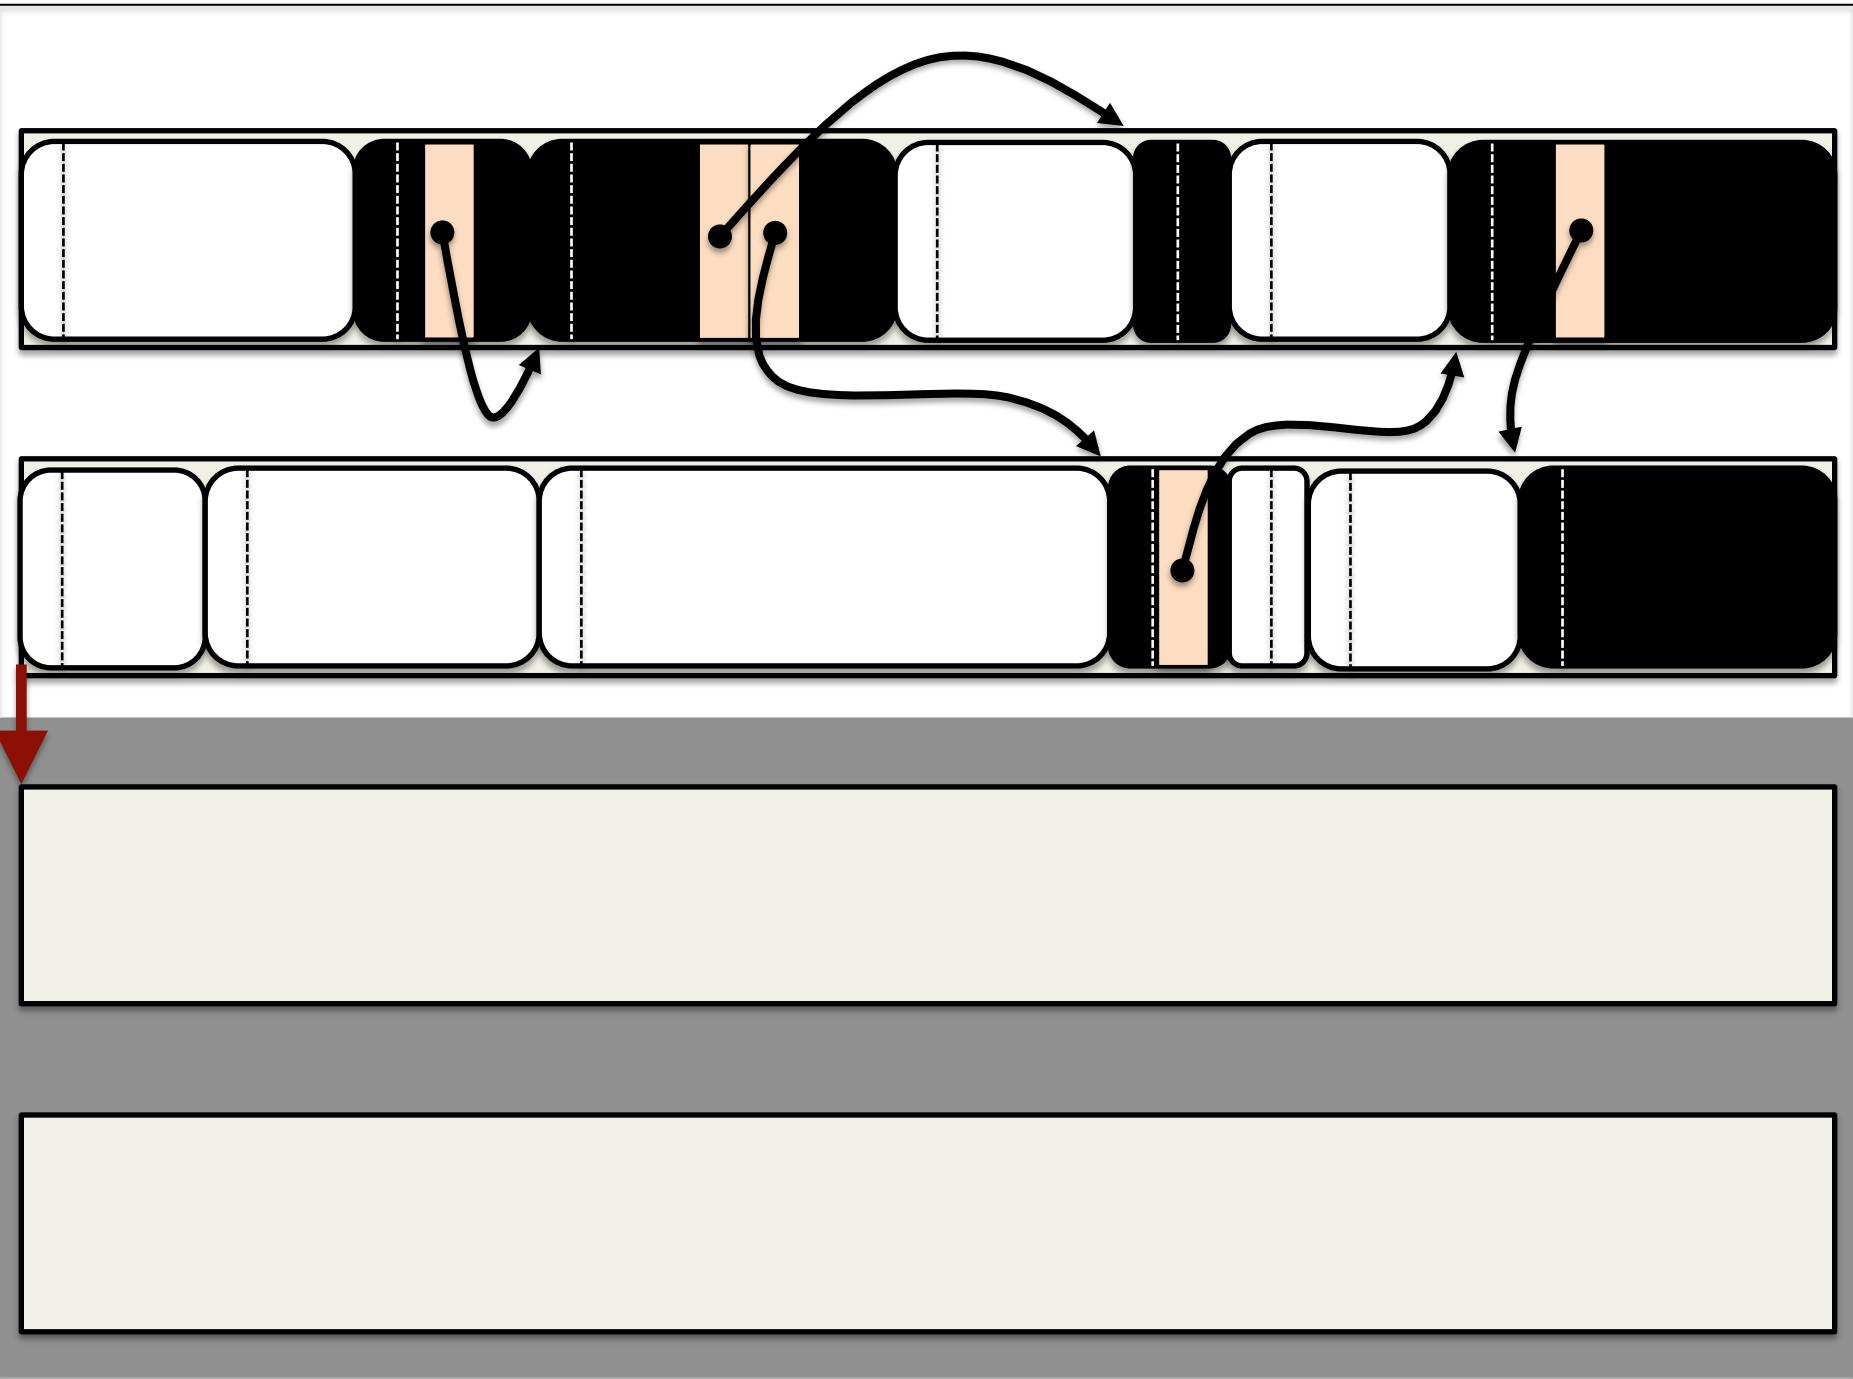


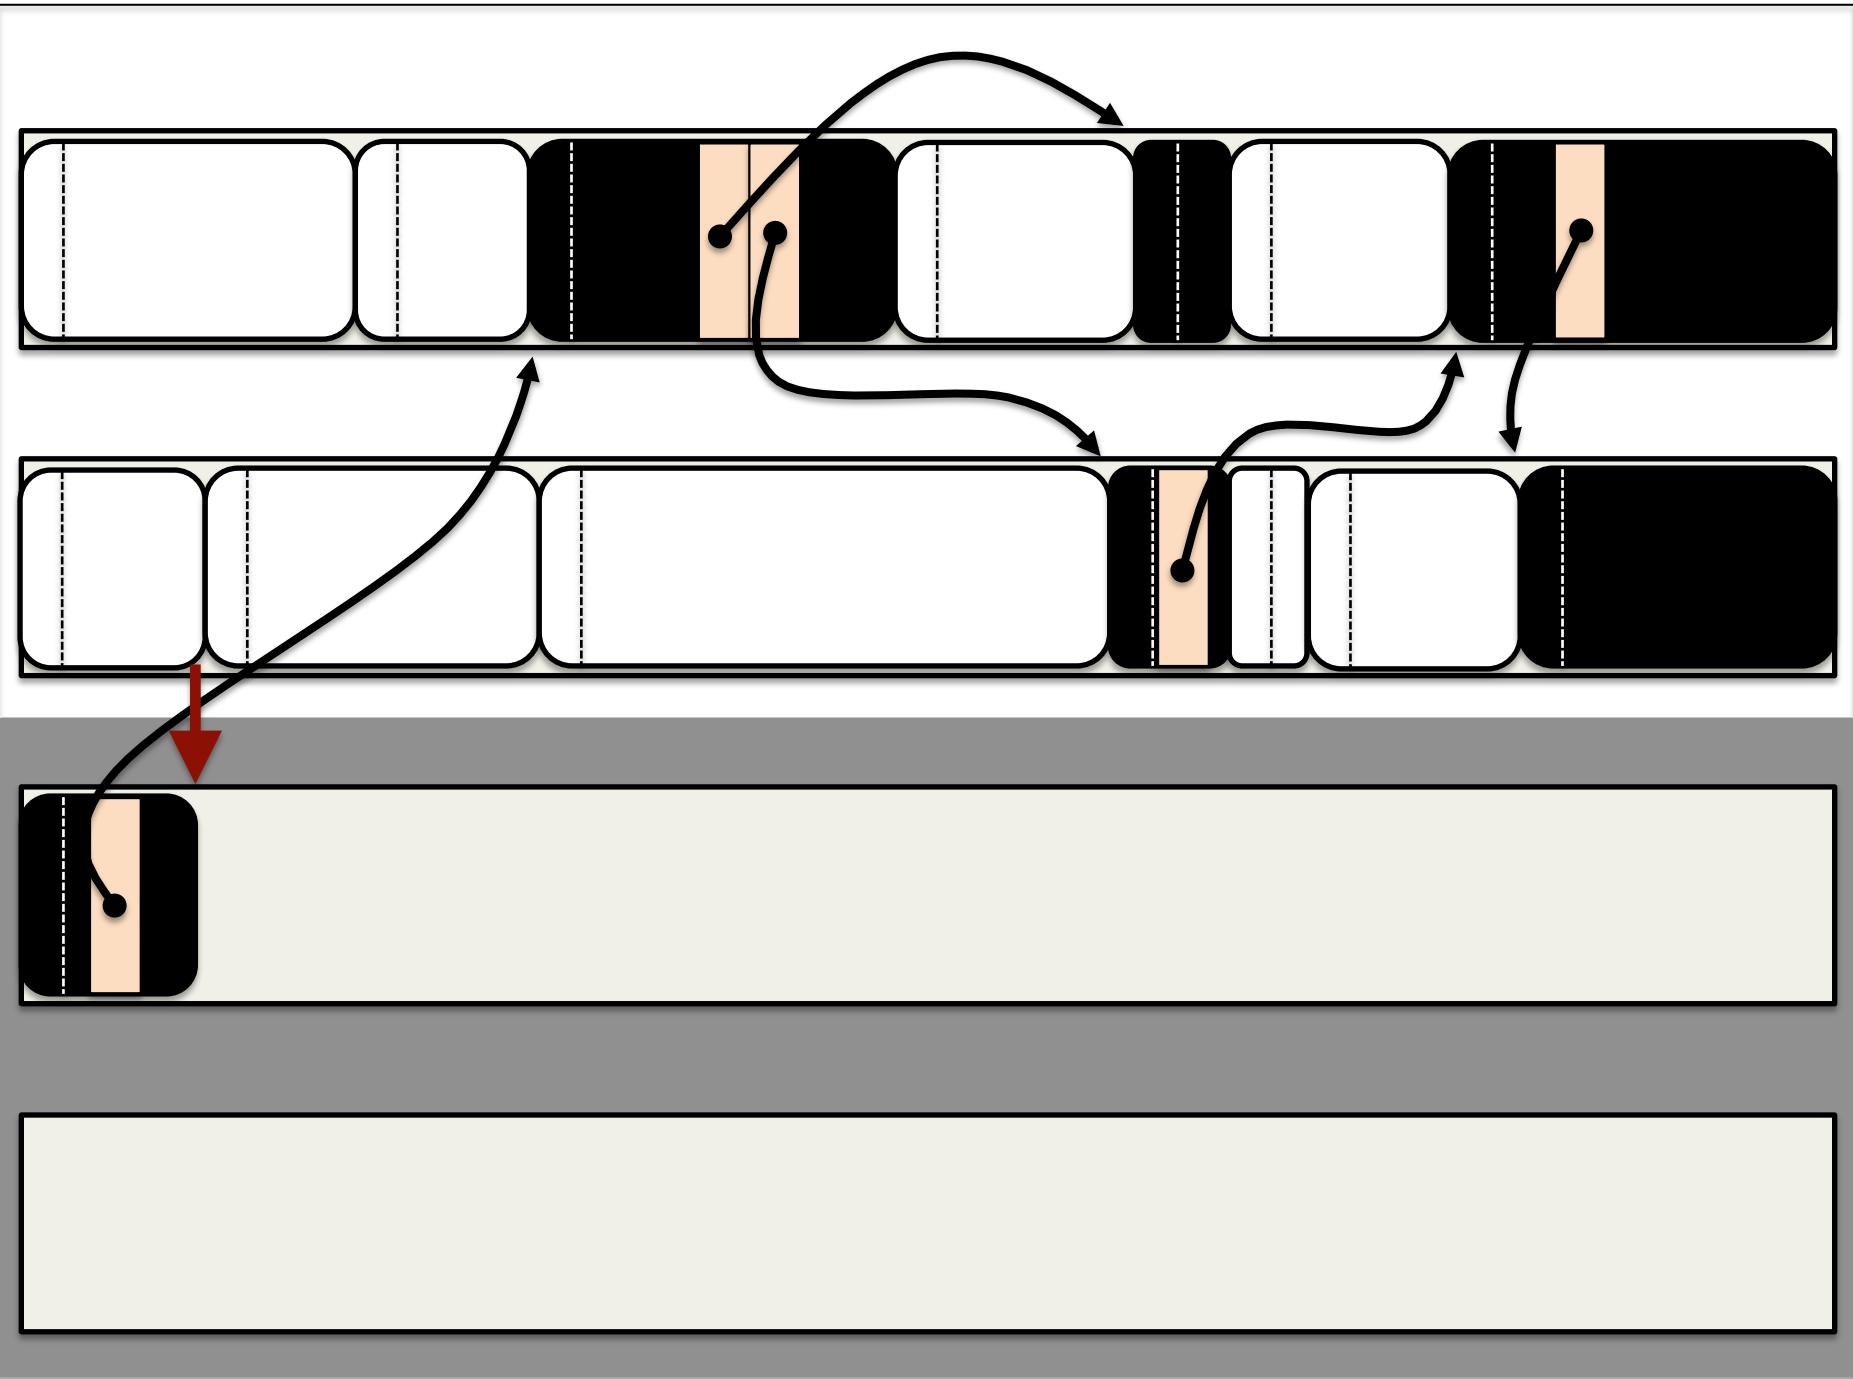


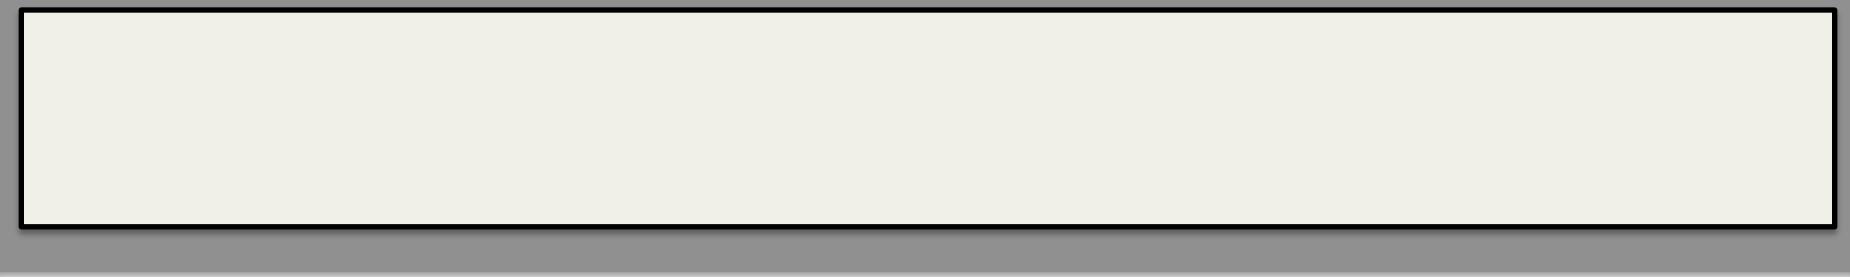
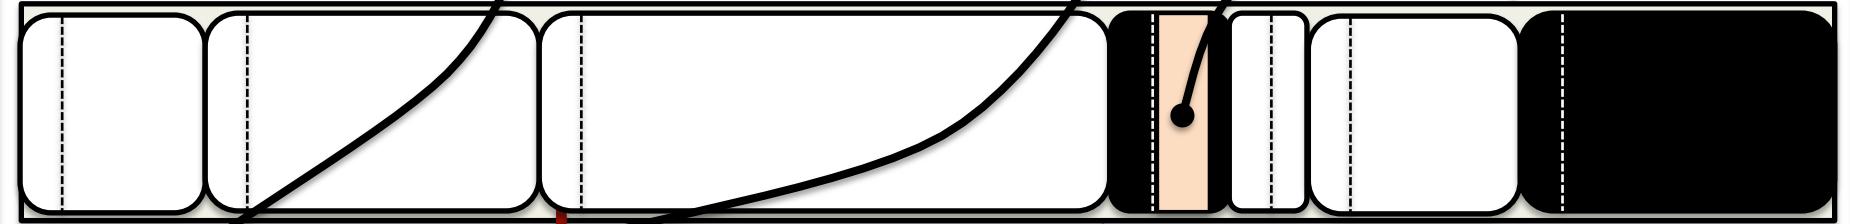
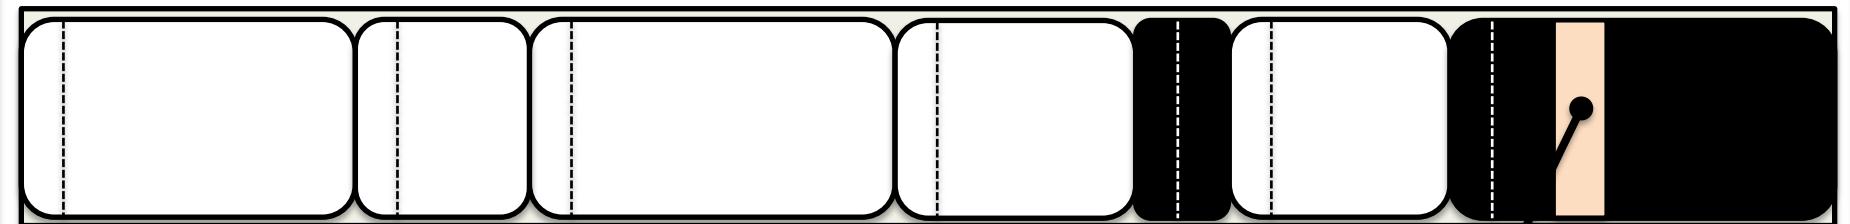
Old and New Space

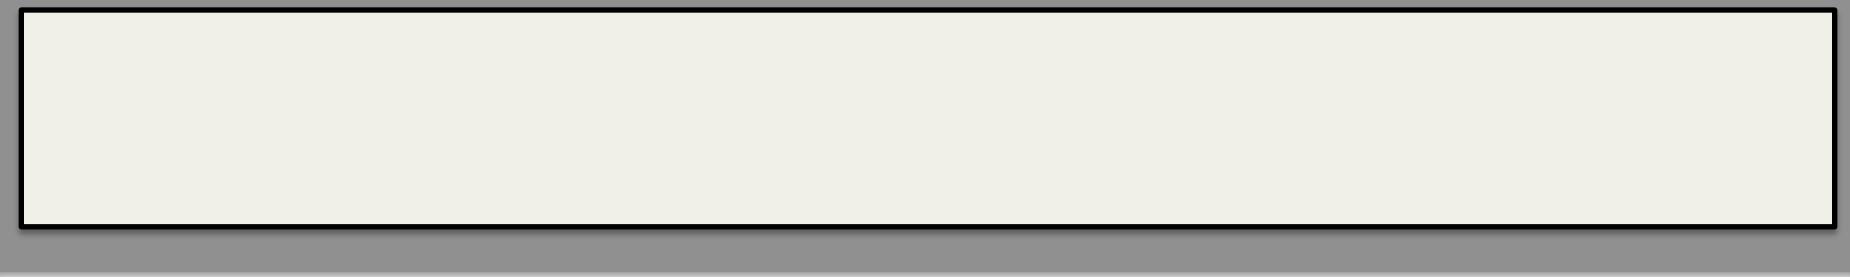
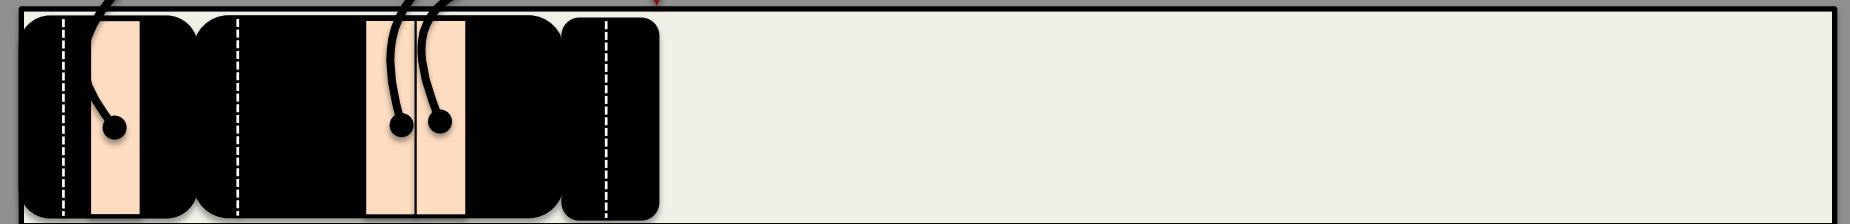
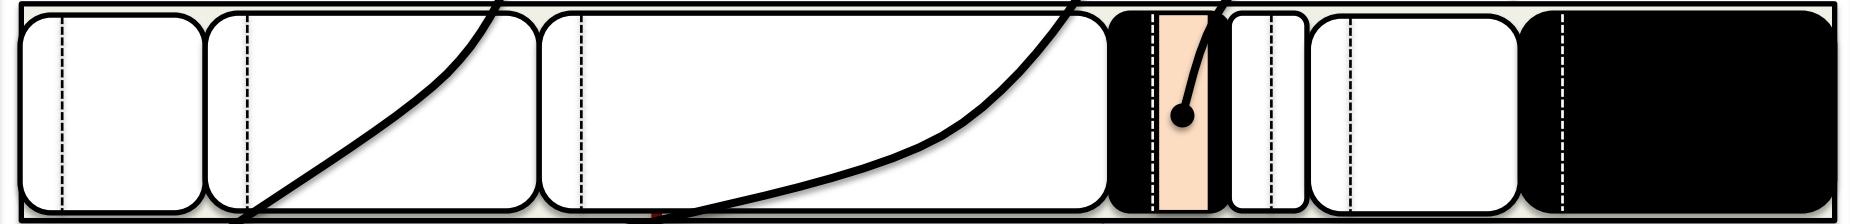
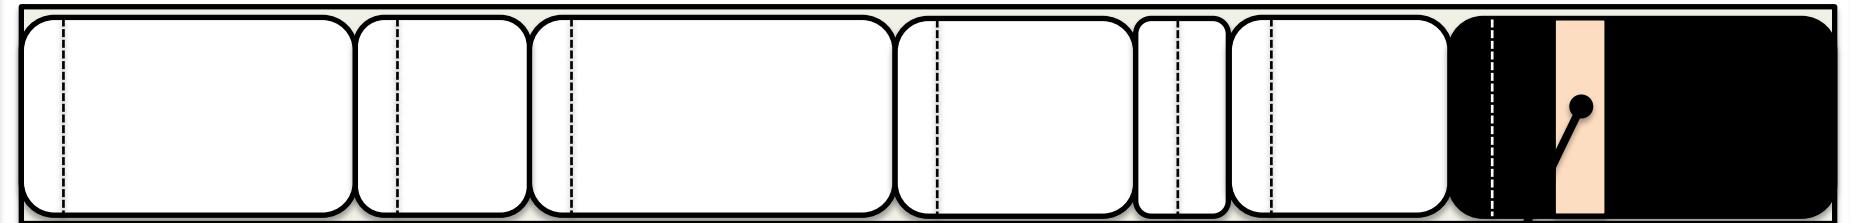
- Copy all live objects from one space to the other
- Spaces referred to by different names
 - Old space or From space
 - New space or To space
- Which one is used depends on convention
 - Both mean the same thing

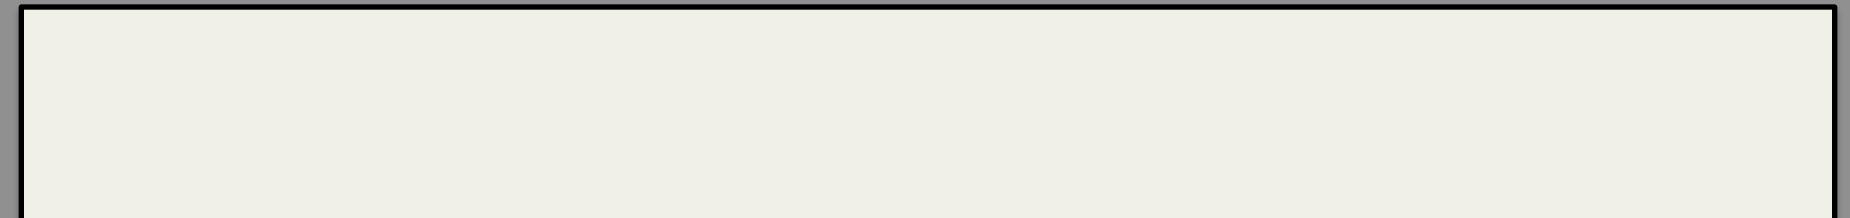
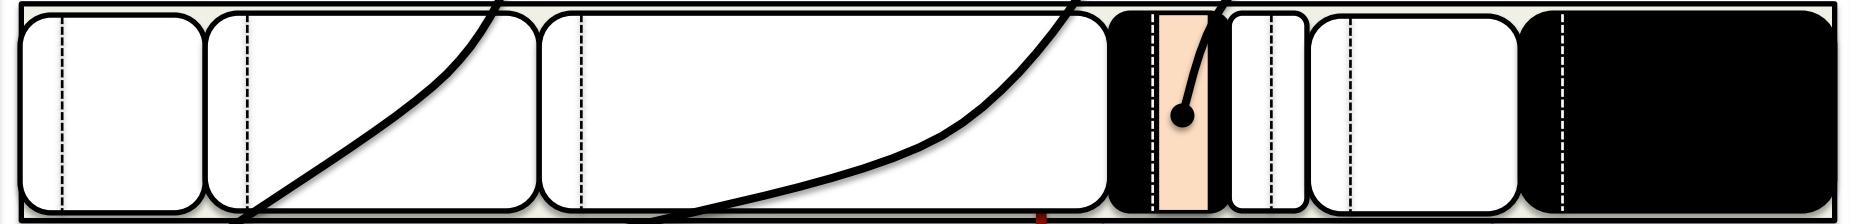
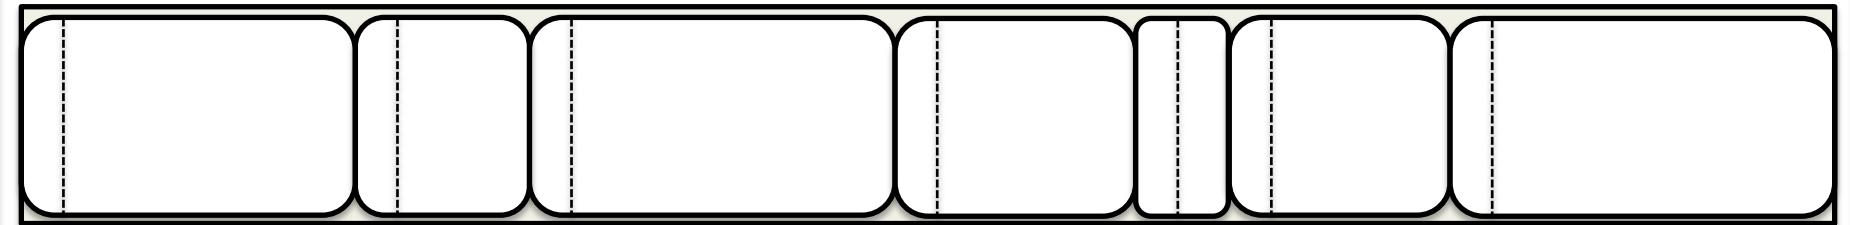


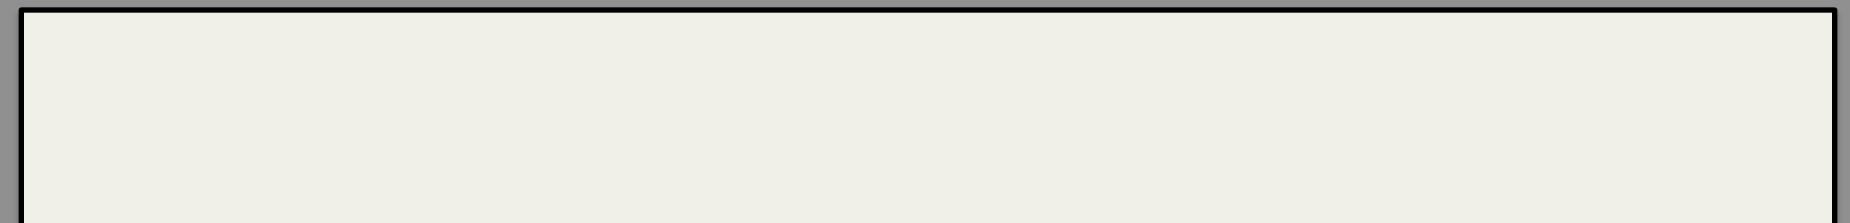
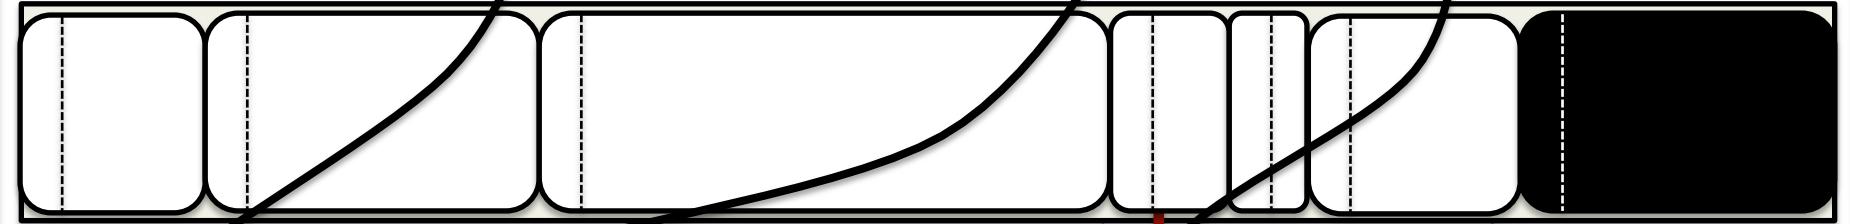
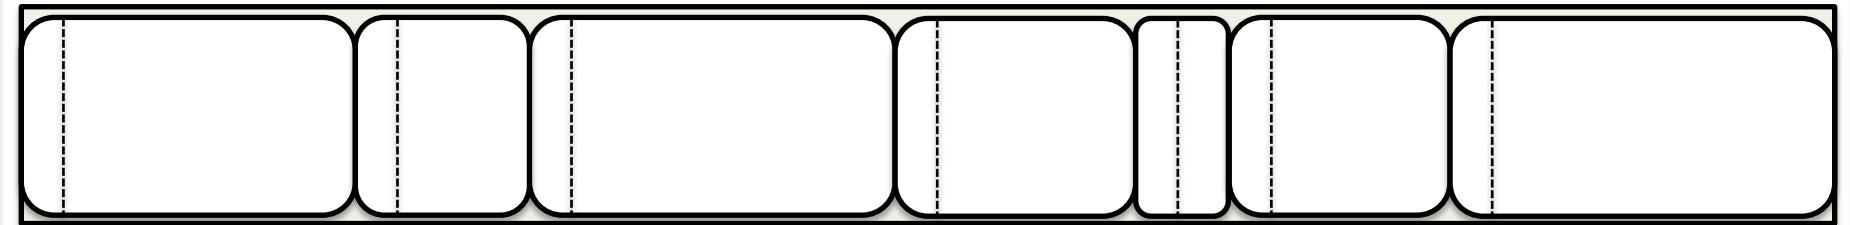


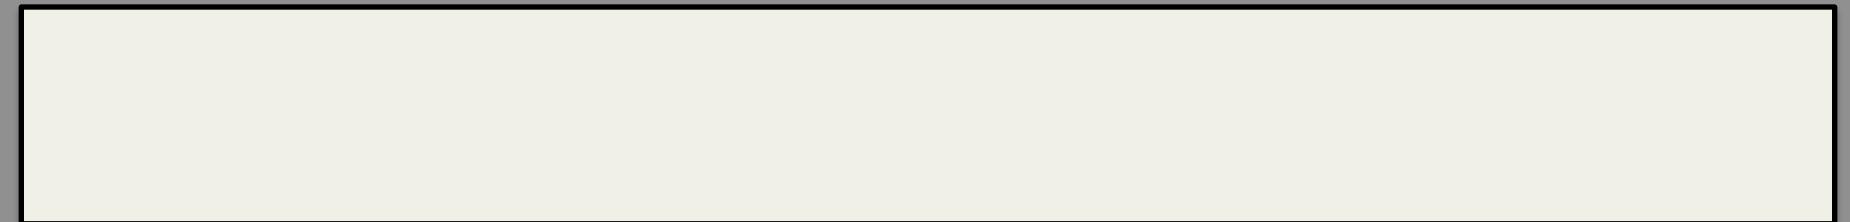
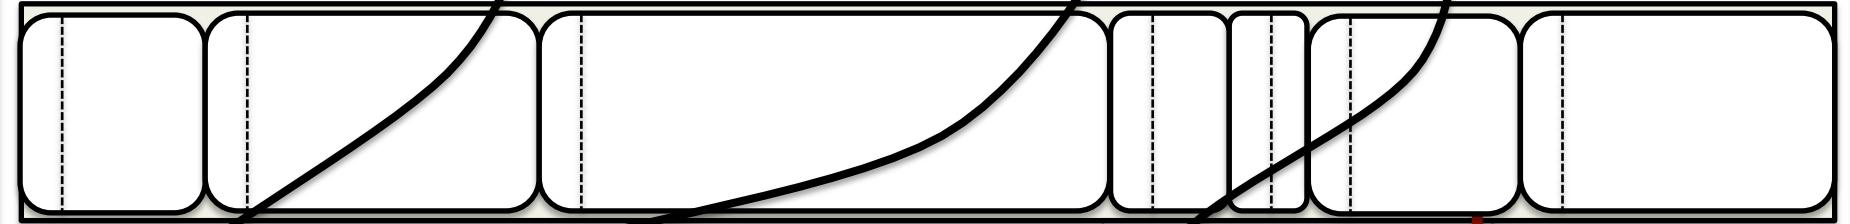
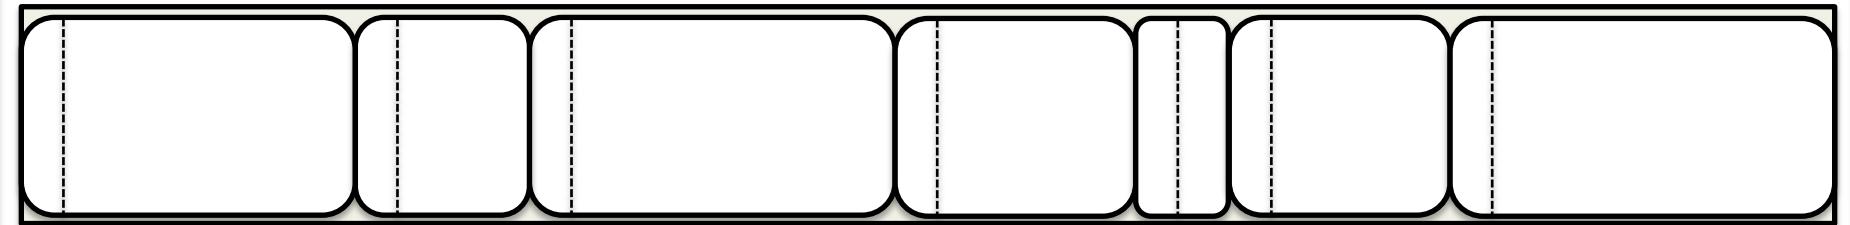












Updating References

- Copy into new space is bit-for-bit
 - Generally use memcpy
 - Pointers are just patterns of bits
- Copied object refers to old space
 - But old space is garbage
 - Need to have one canonical version of the object
- We have to update references when we copy

Handles

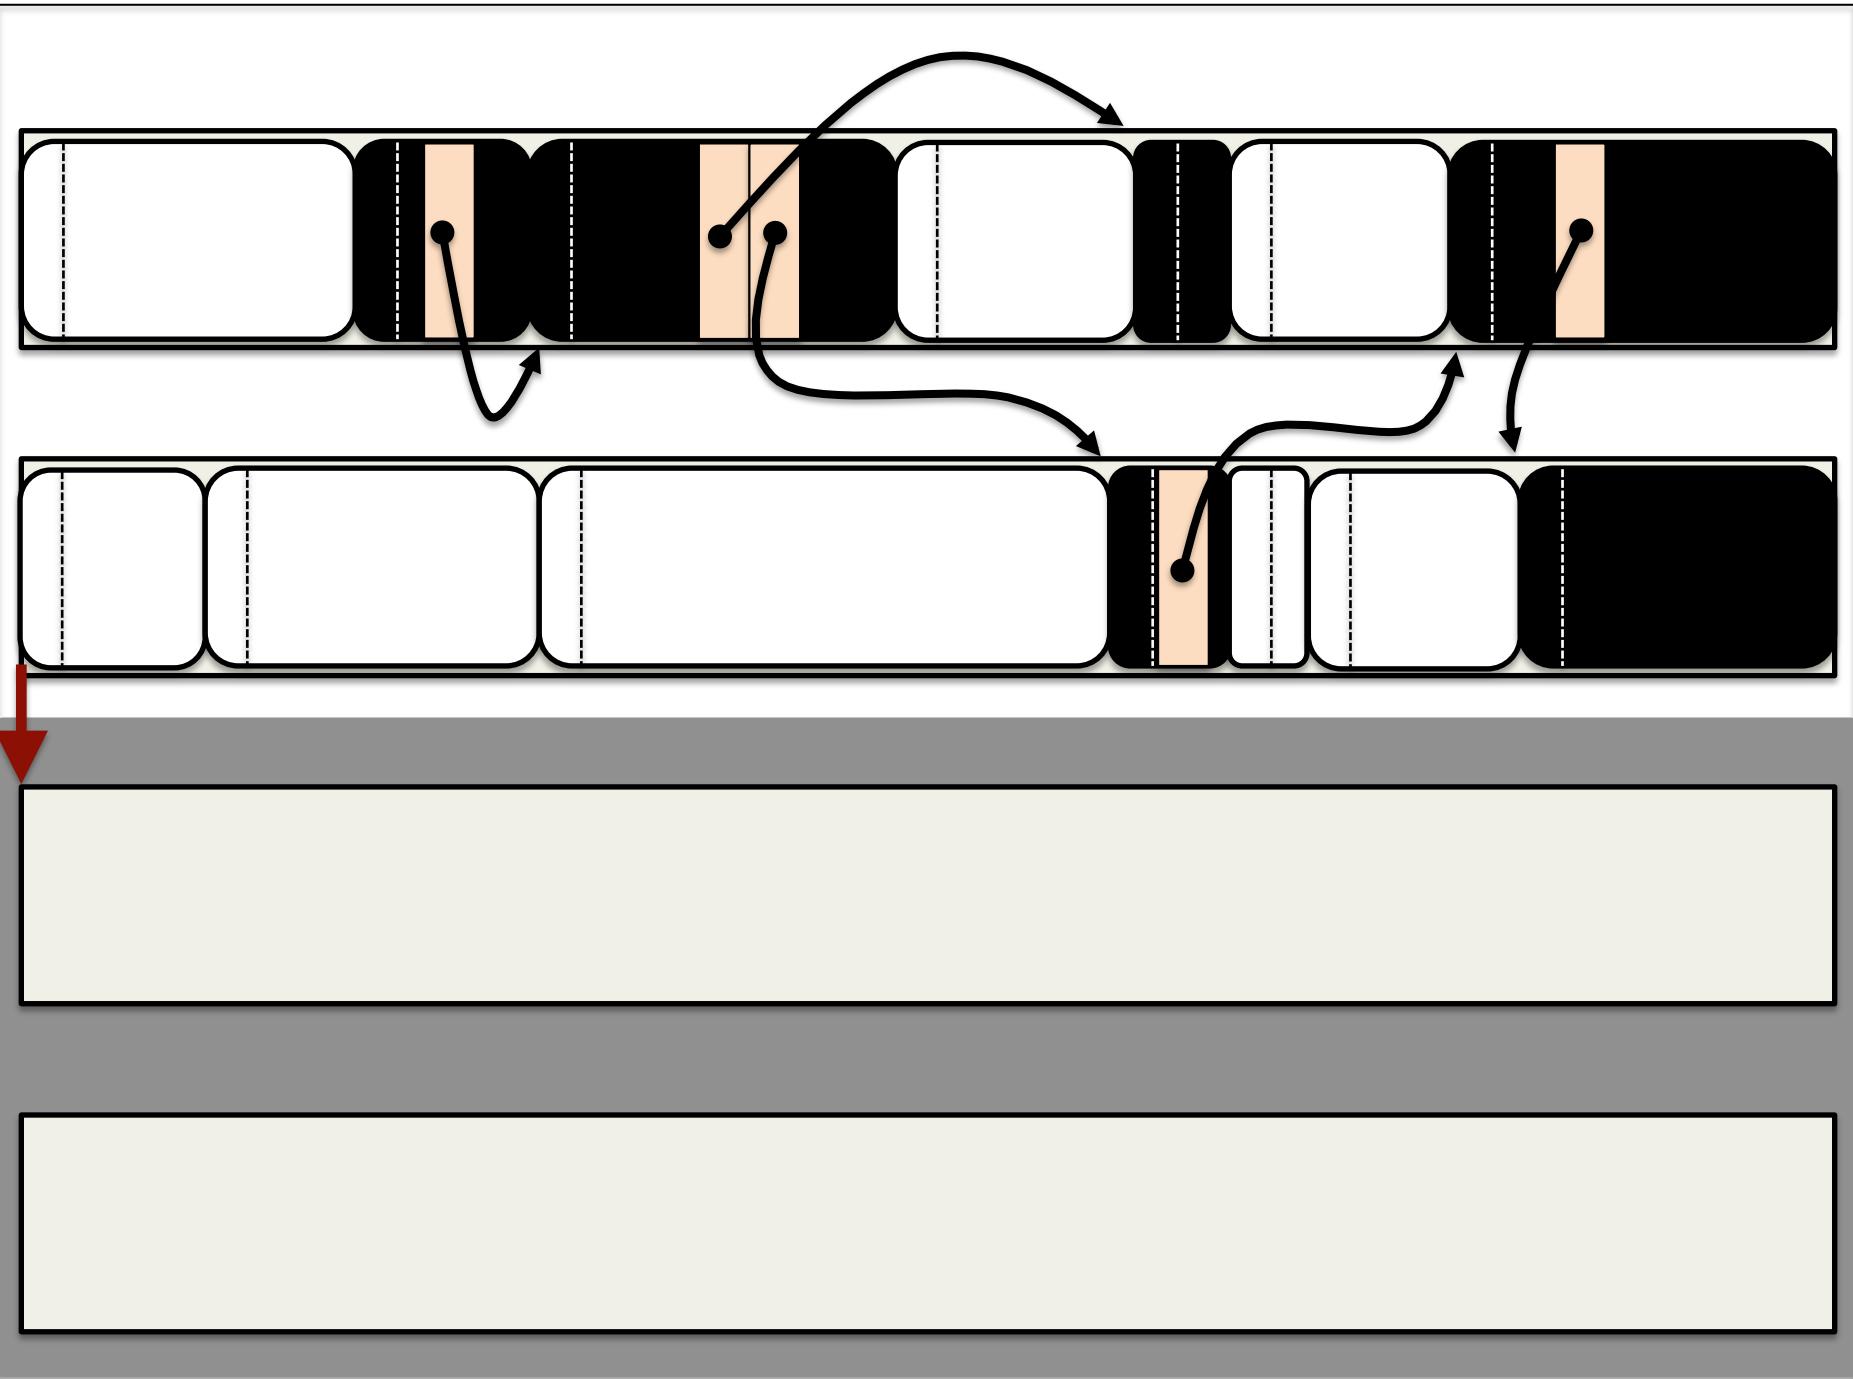
- Remember the two ways of storing metadata
 - Handles
 - Object headers
- Handles give us a single reference to an object
 - All other references go through the handle
- Solves the forwarding pointer problem
- Still other issues with handles

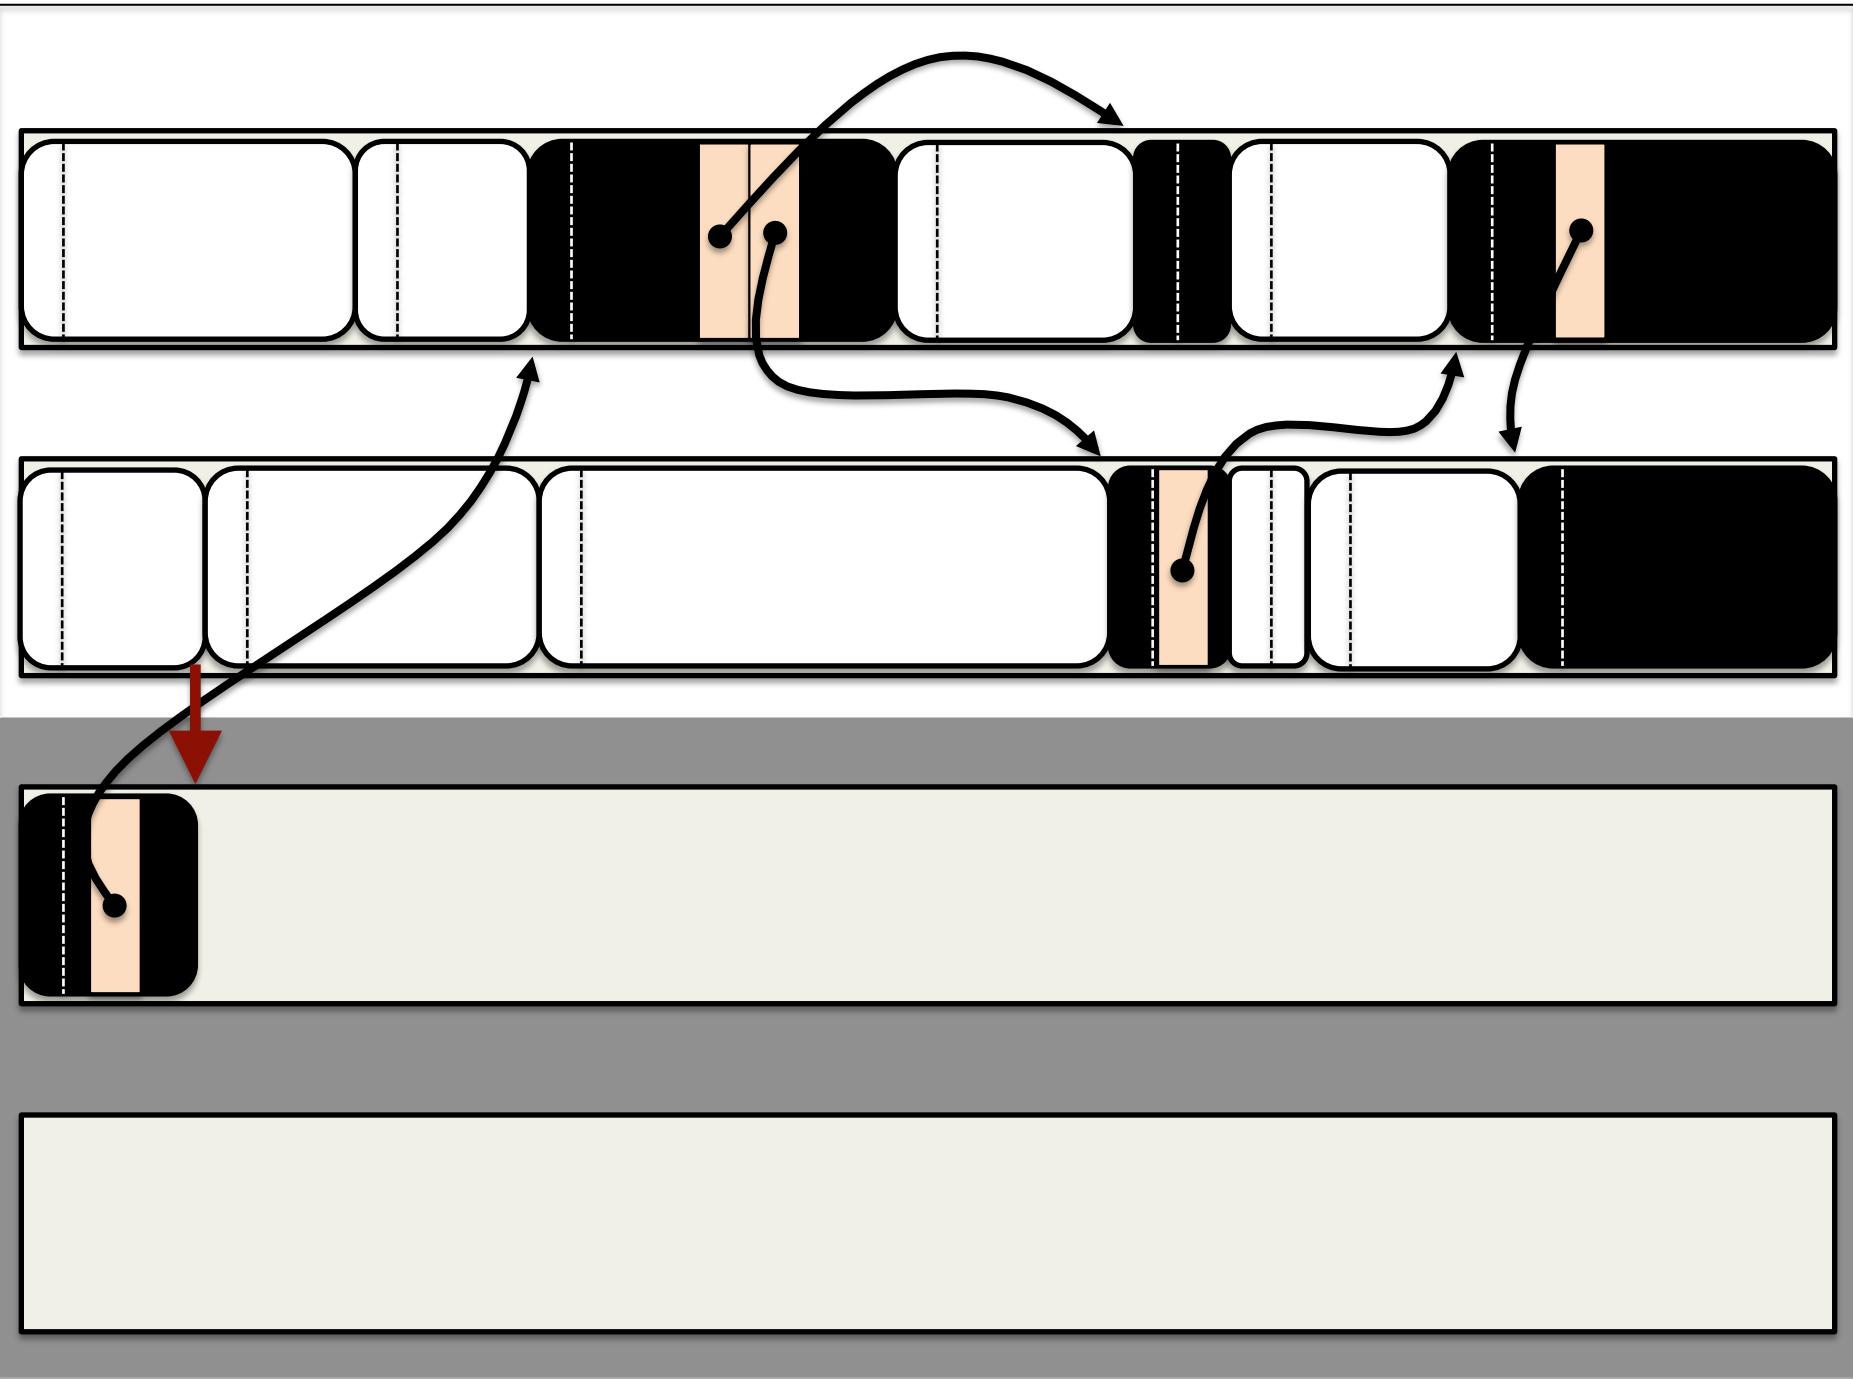
Forwarding Pointers

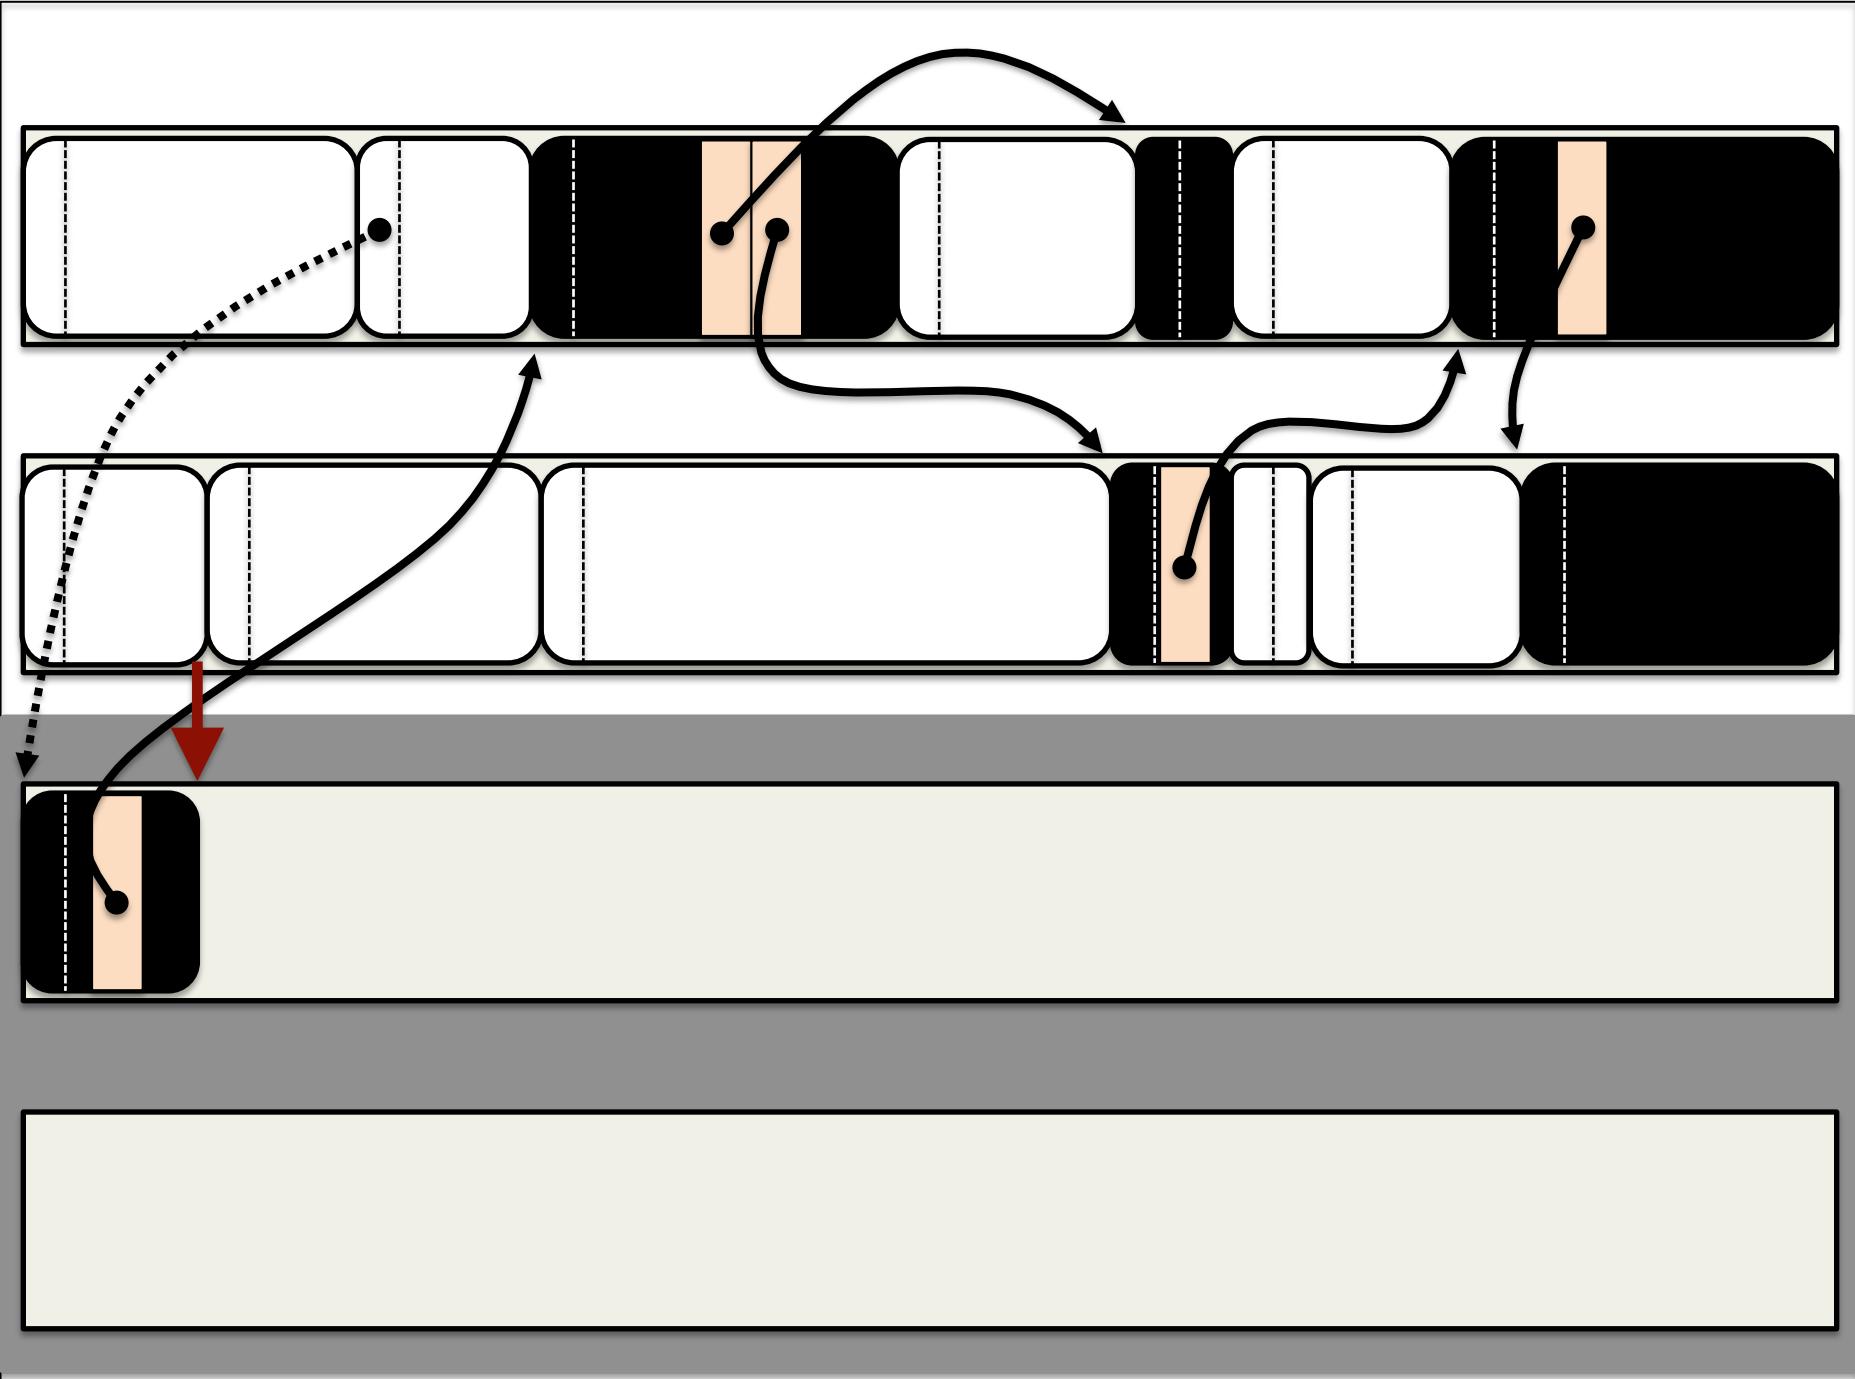
- We can't update all references when we move
 - Object graph is directional
 - We don't know where all the references are
- Instead we track where the object went
 - Fix up pointers in another pass
- Need some per-object metadata for pointer

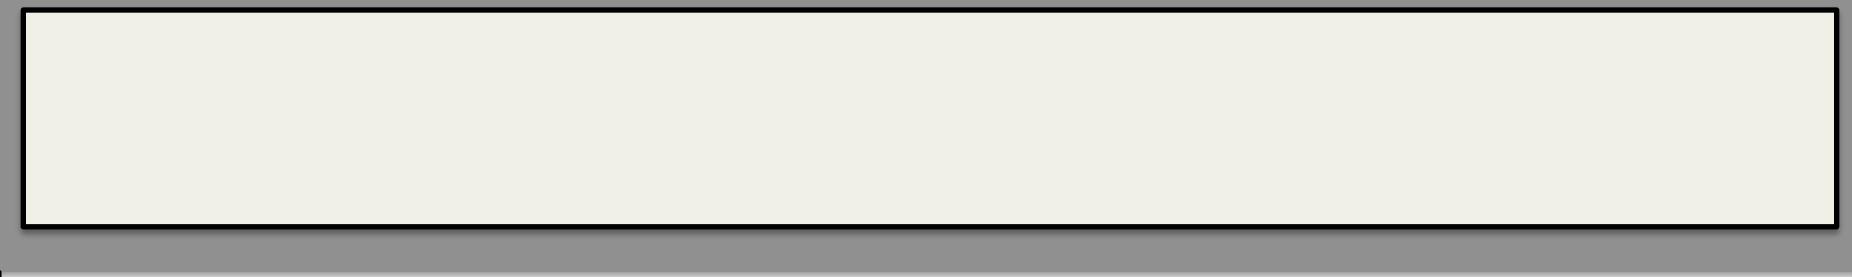
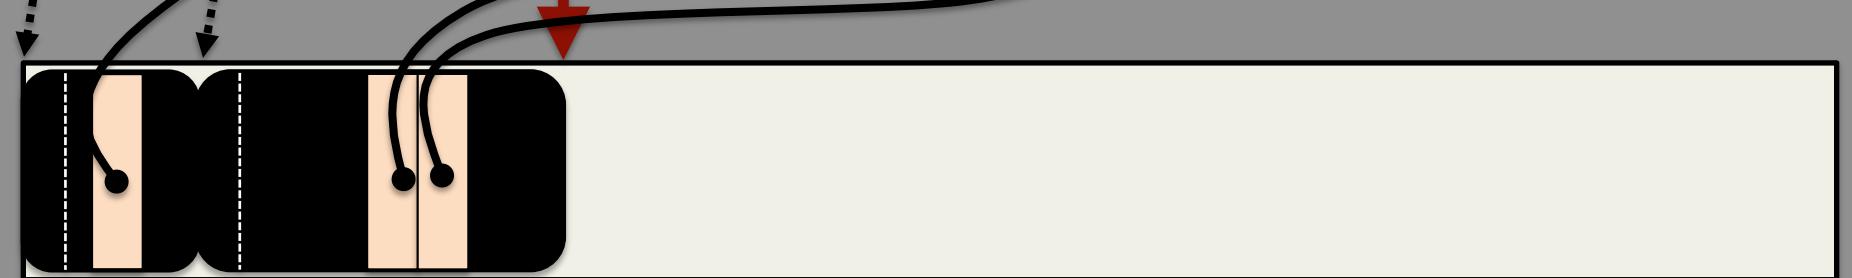
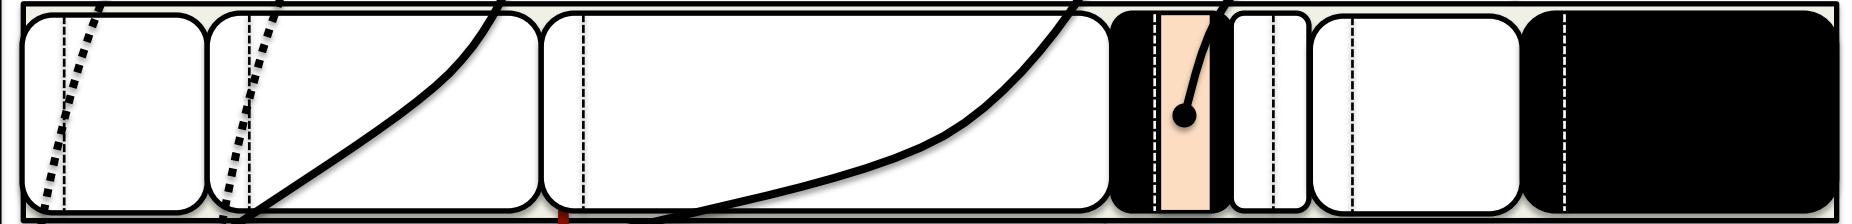
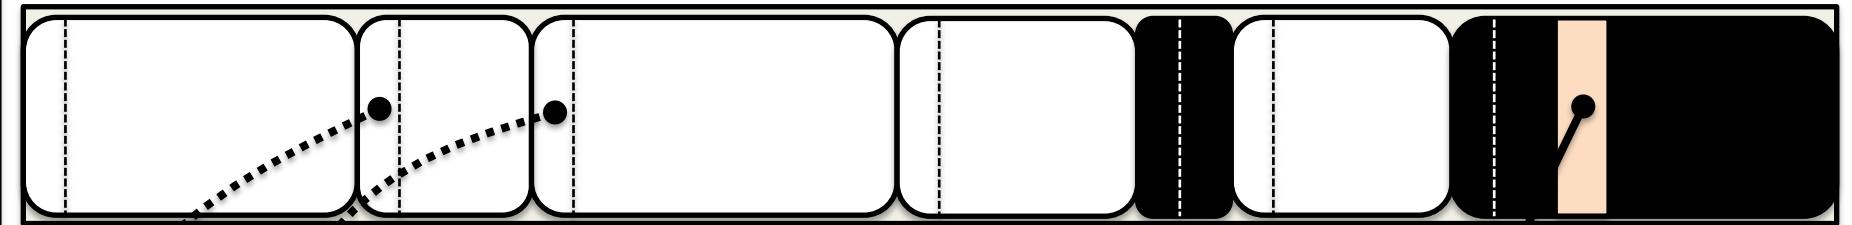
Forwarding Pointers

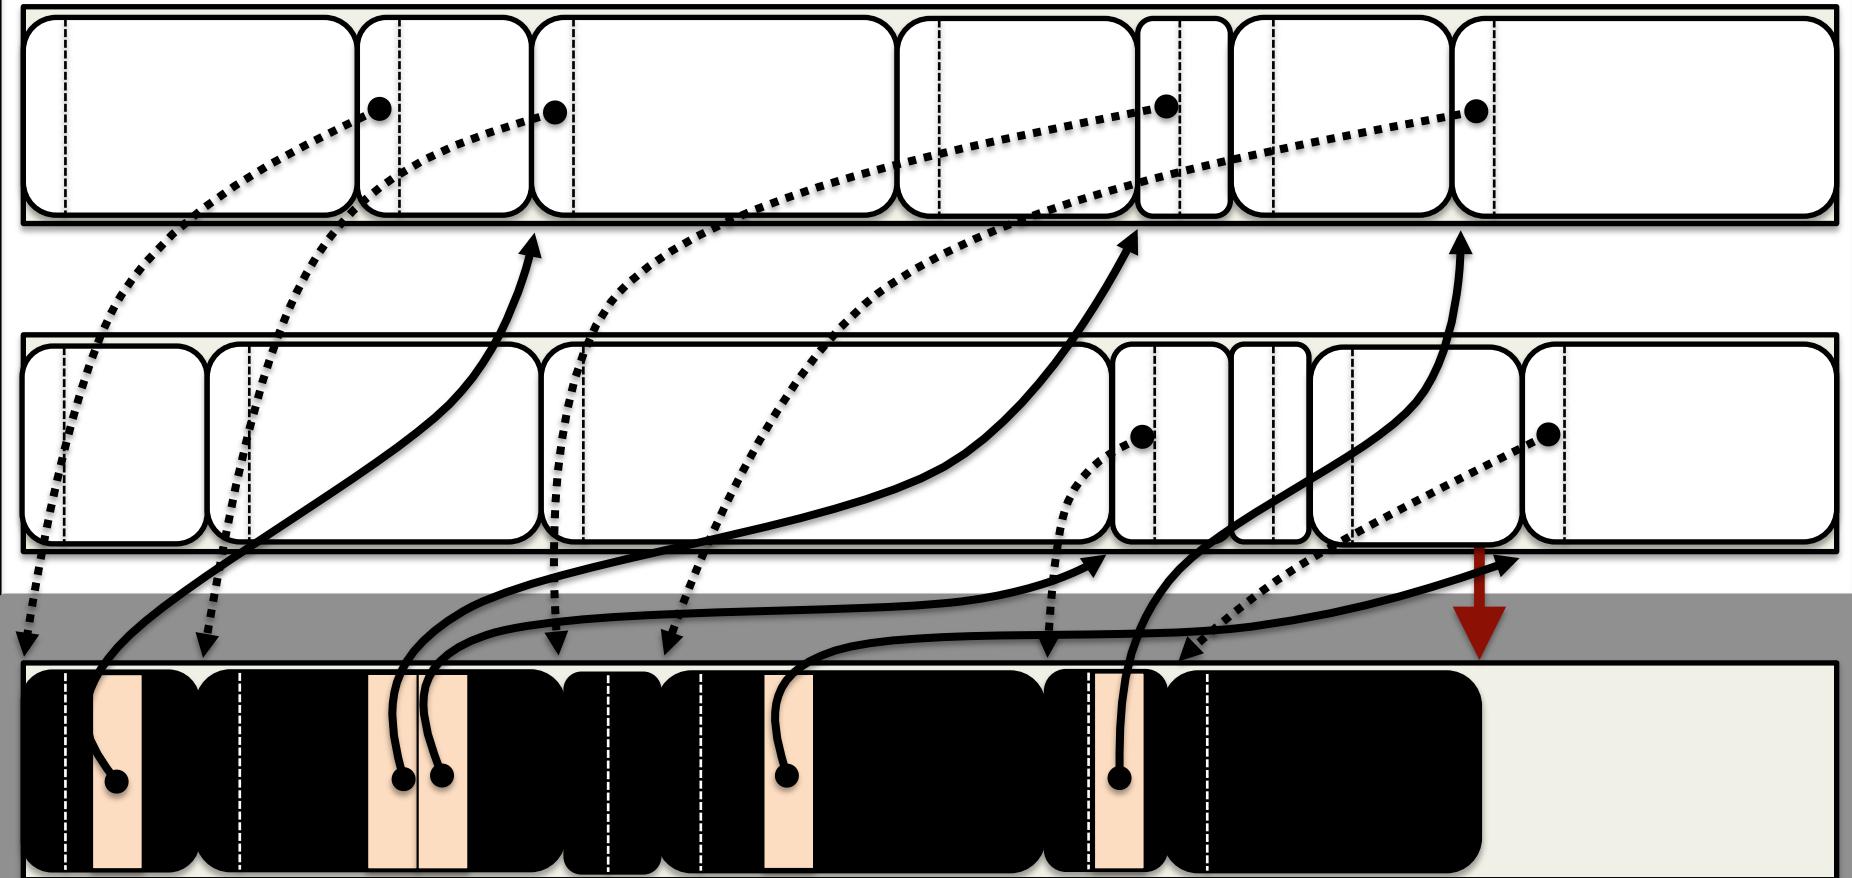
- We know that the old version is invalid
 - All objects in the old space are garbage
- We can overwrite all or part of that object
 - Install a forwarding pointer
 - Generally use a word in the header
- The GC has to recognize a forwarding pointer
 - Steal the low-order bit to use as a flag

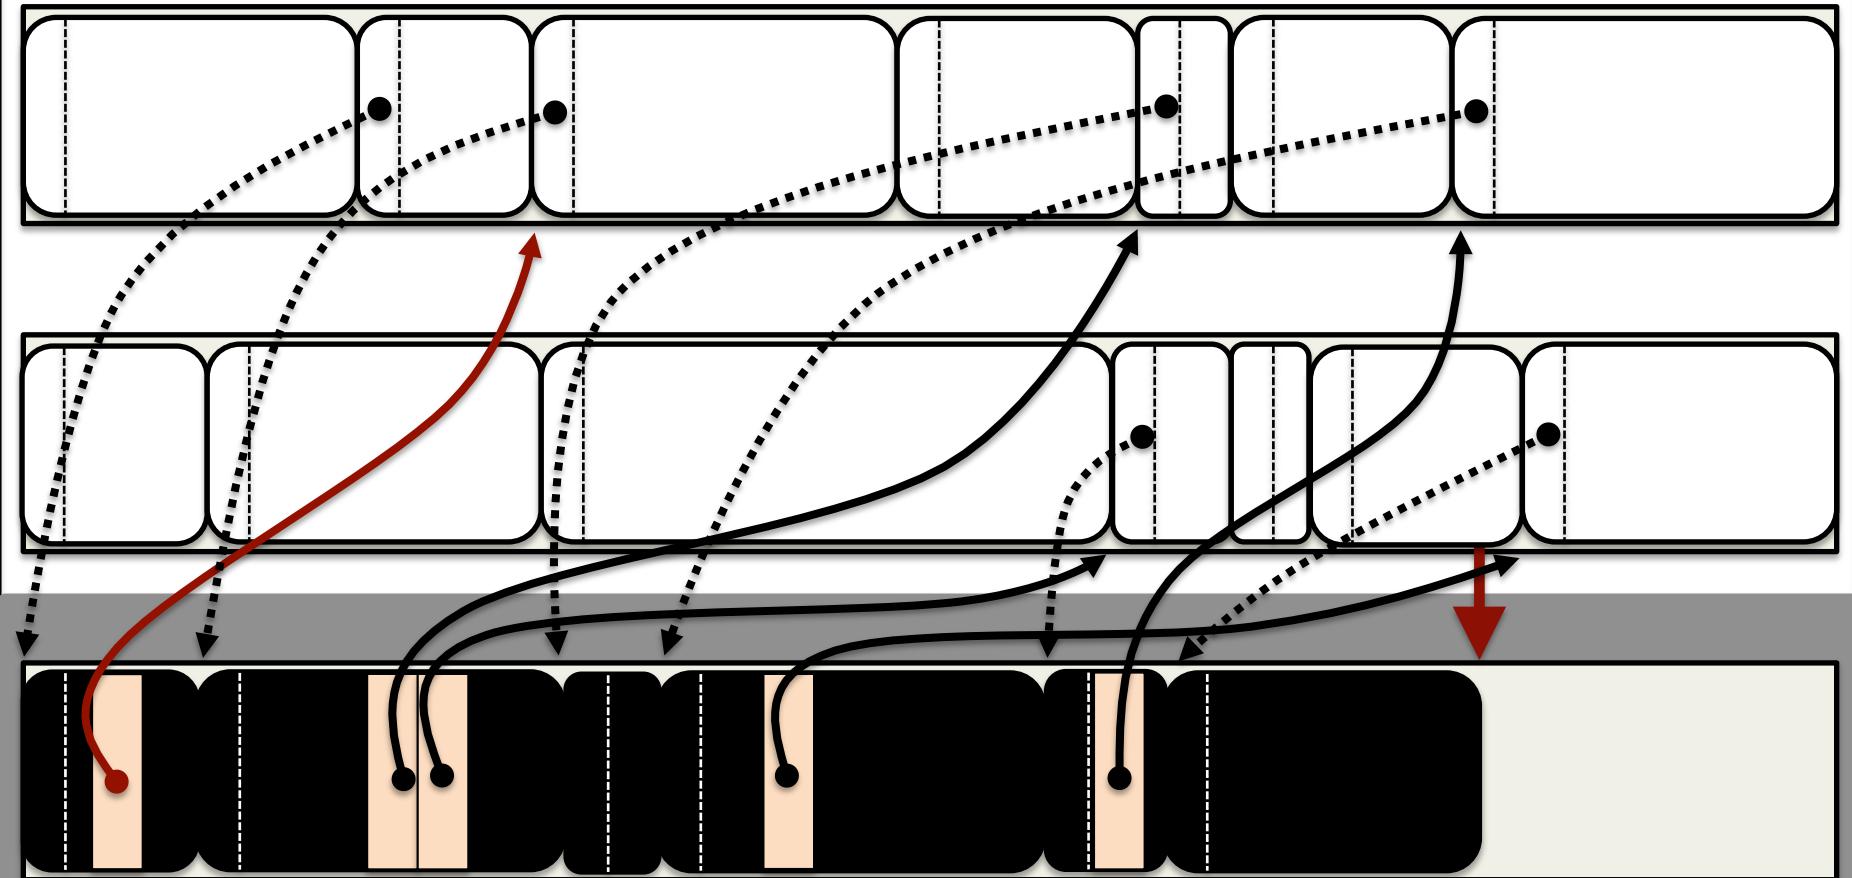


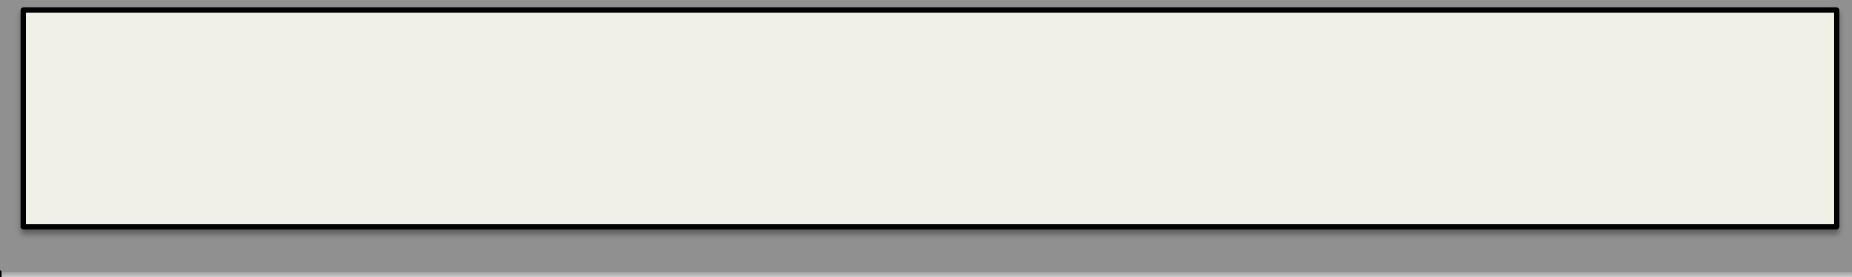
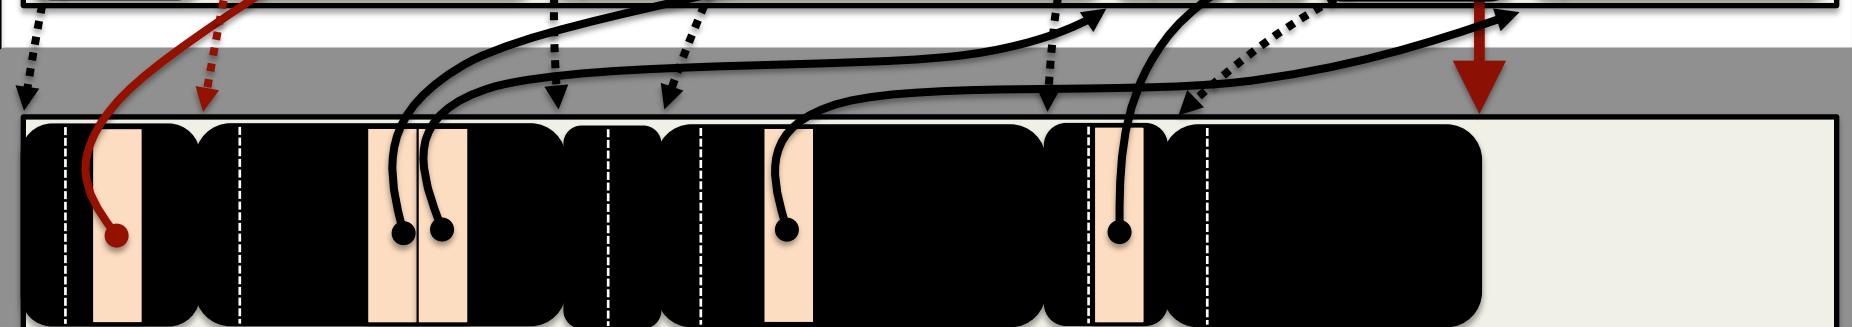
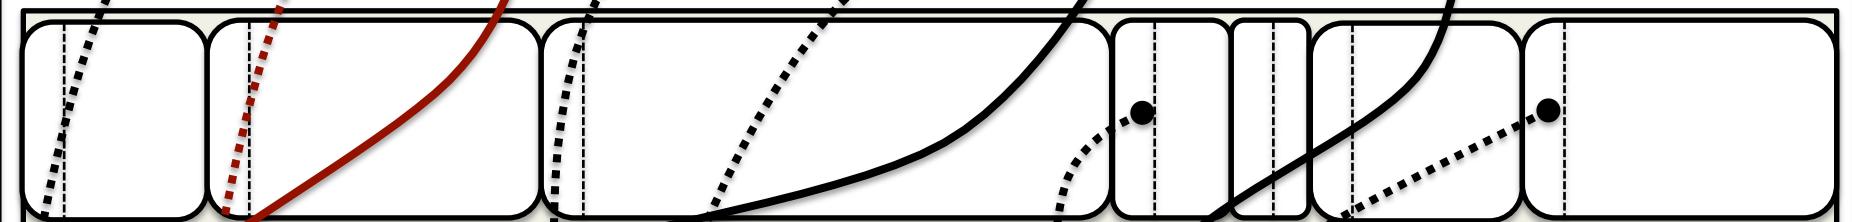
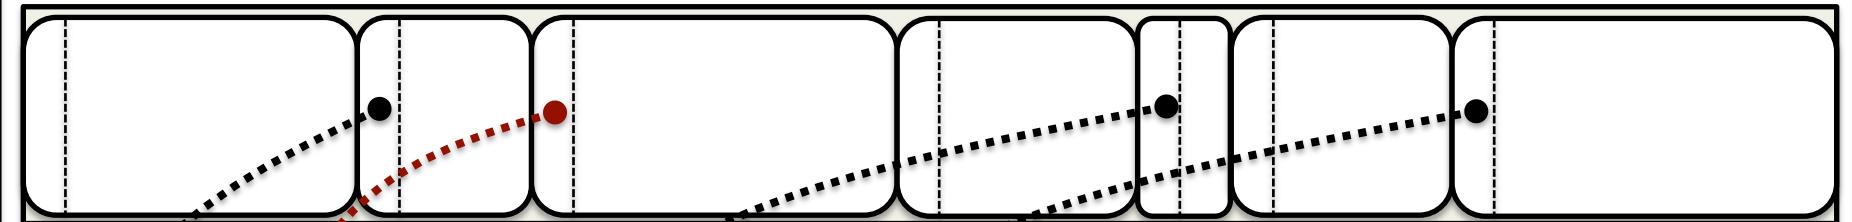


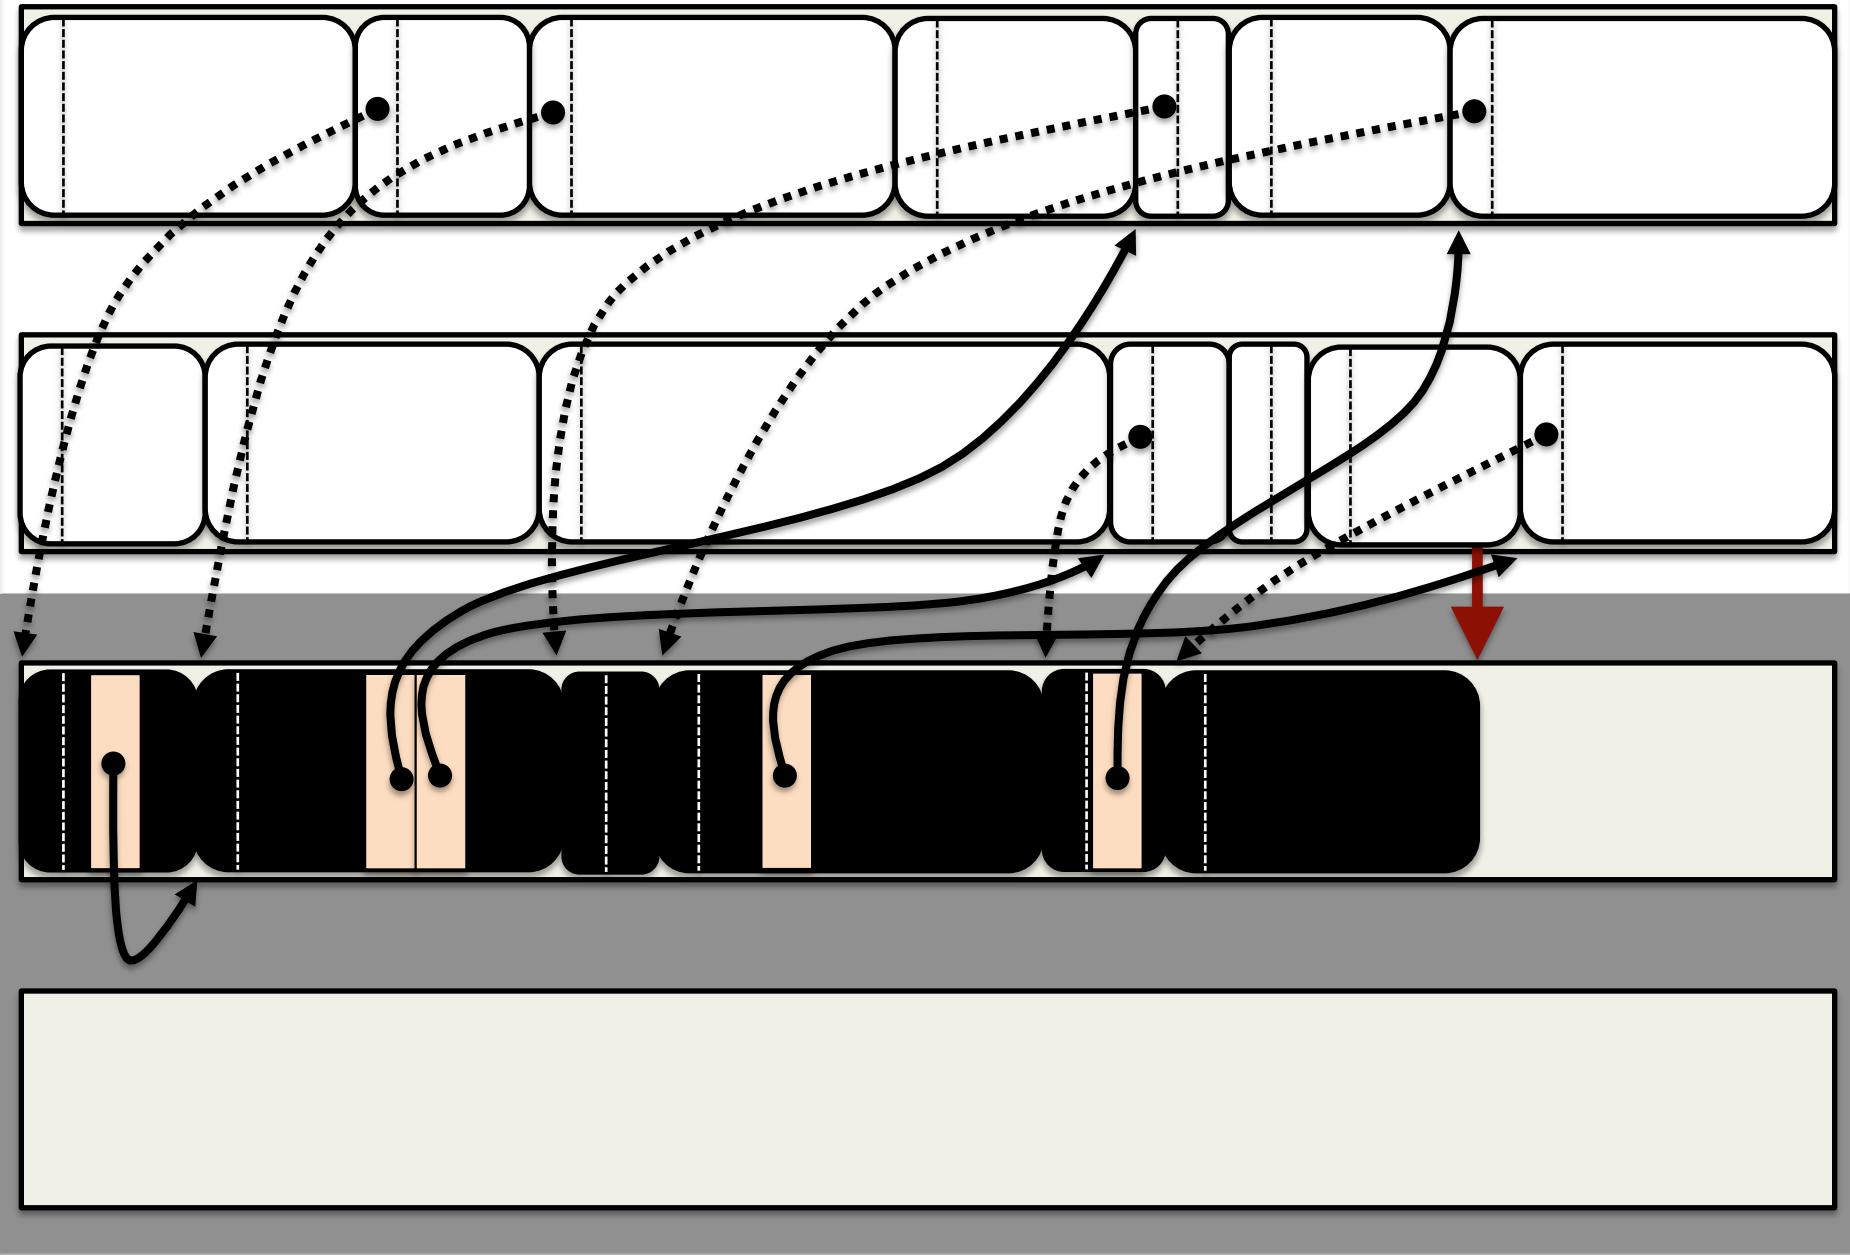


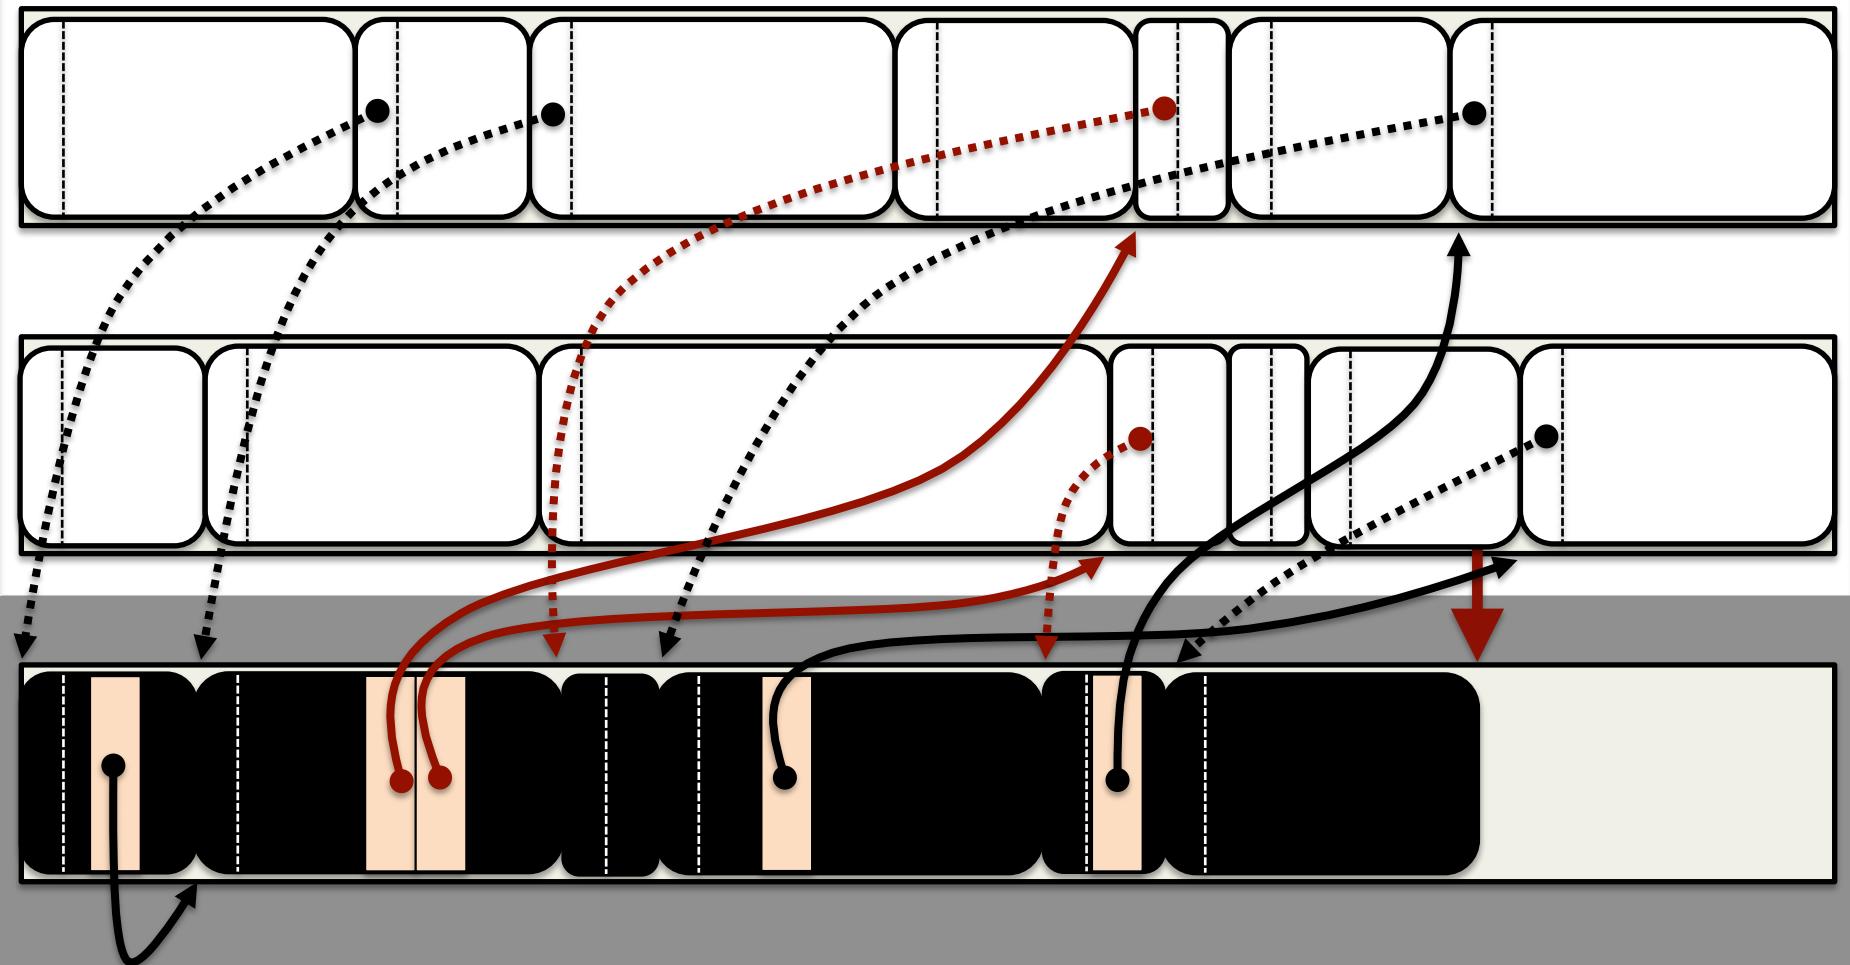


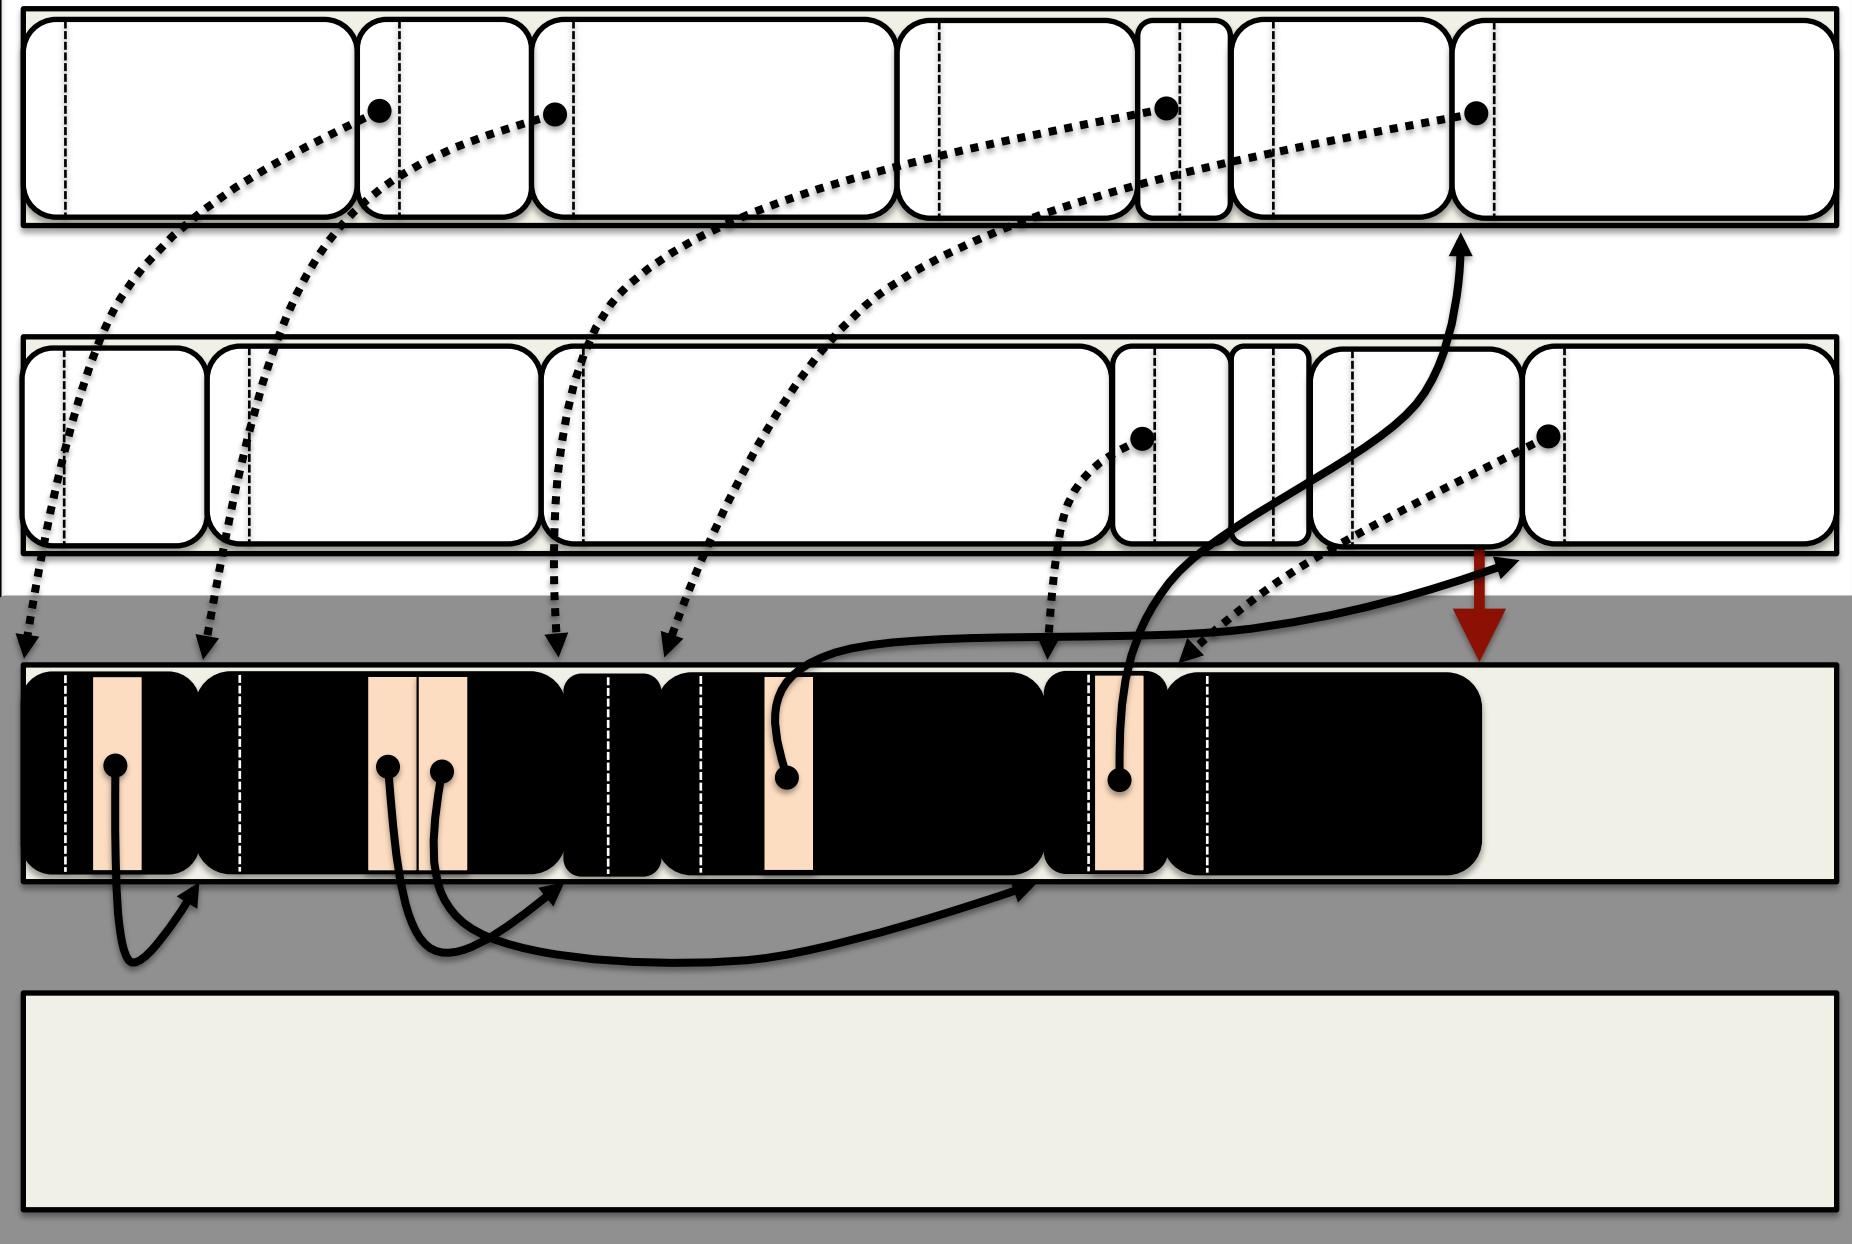


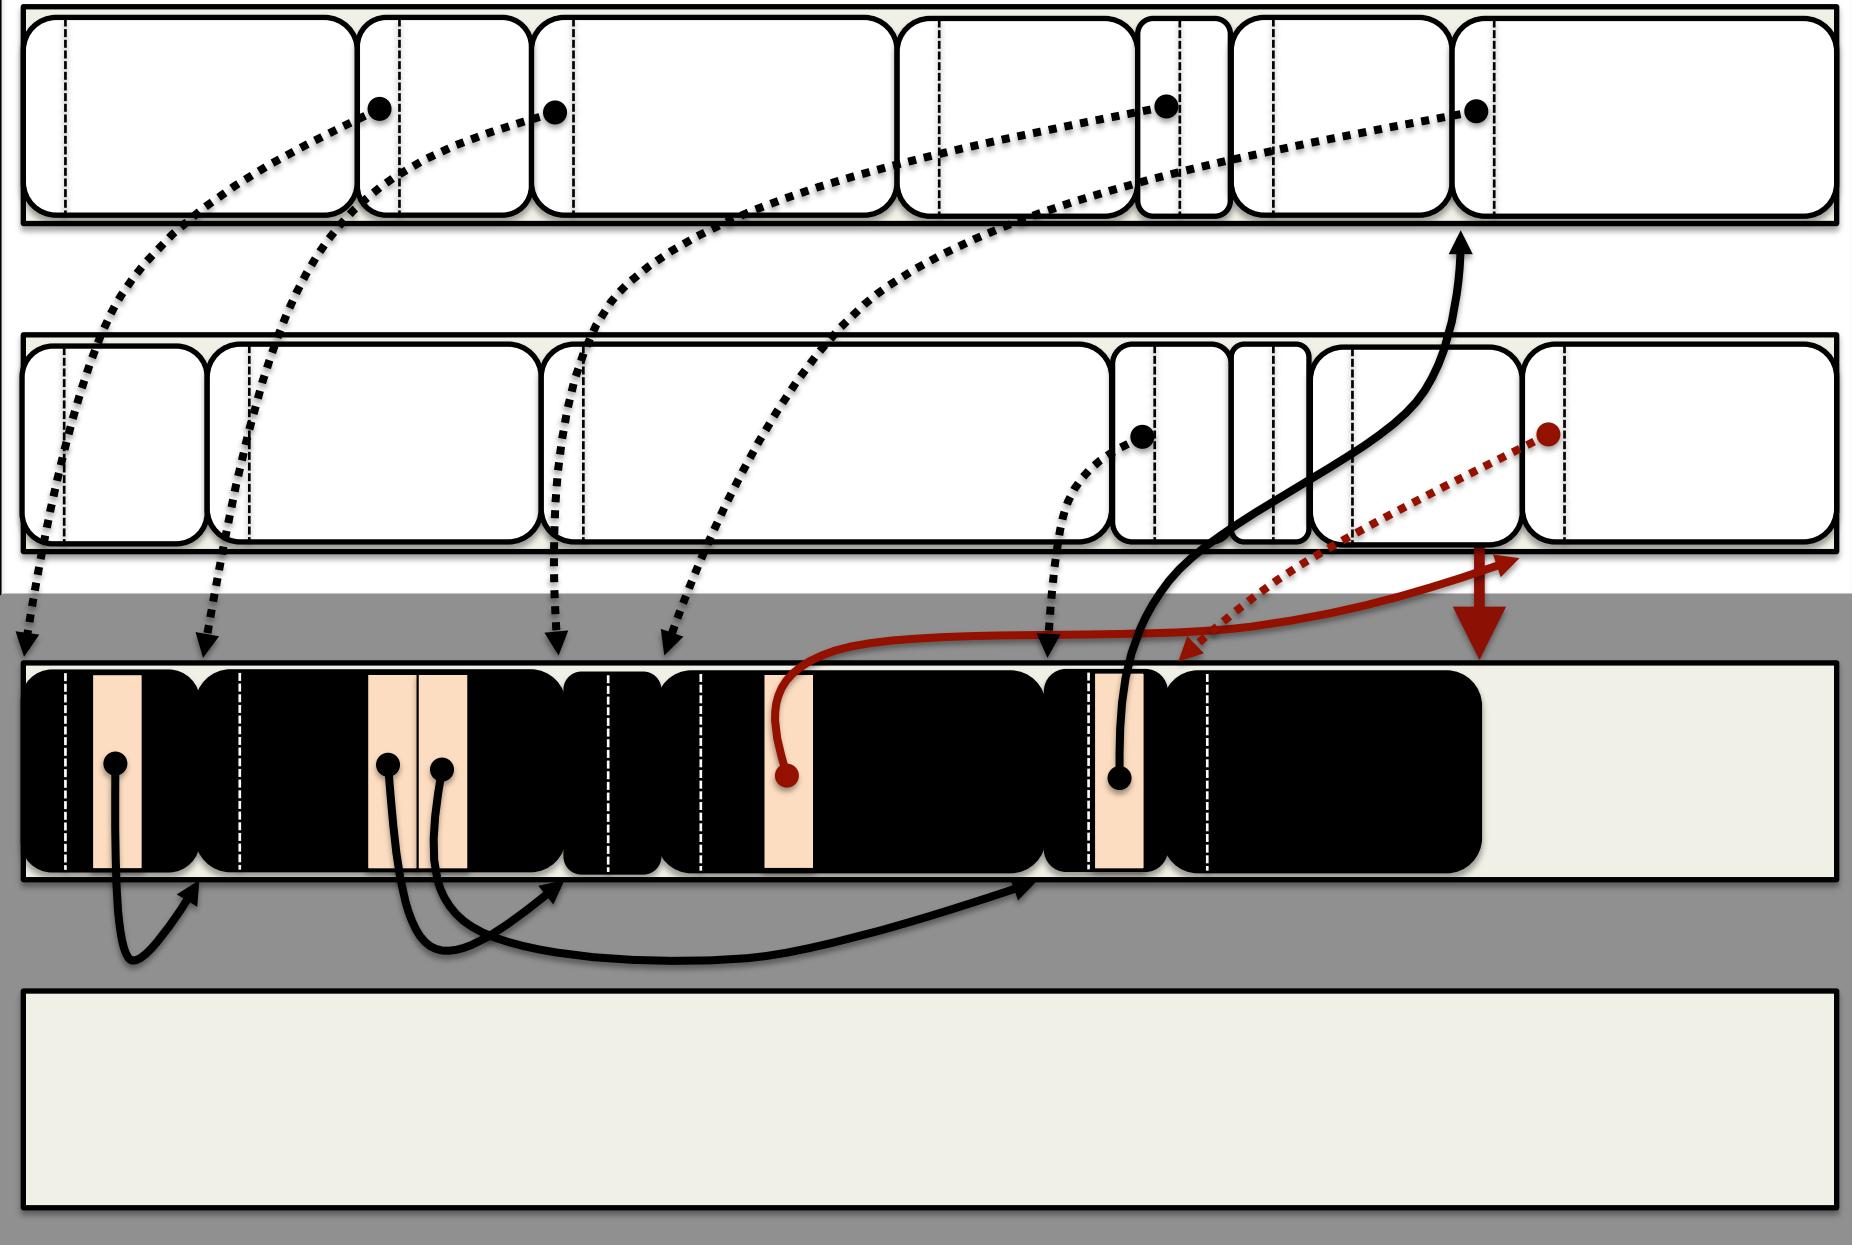


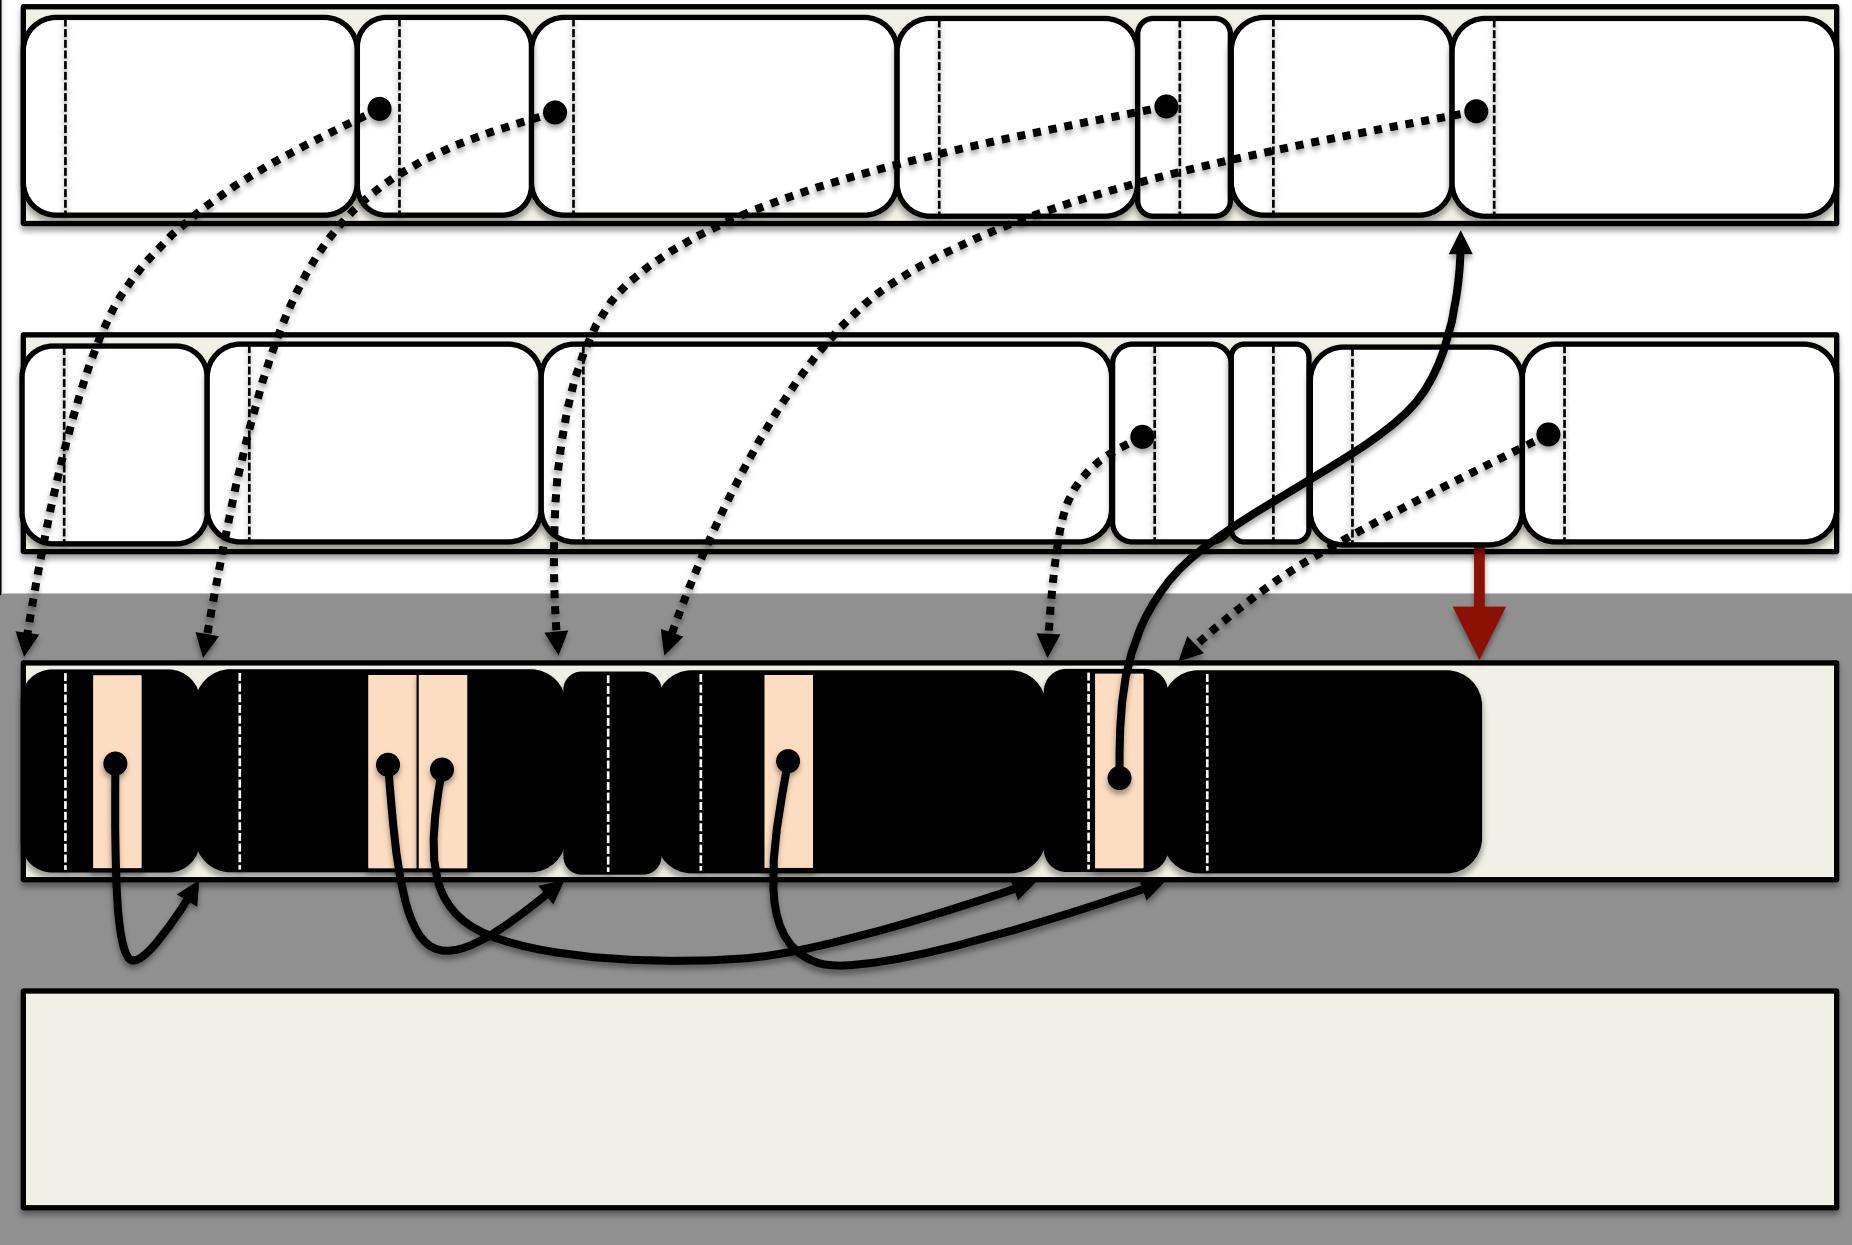


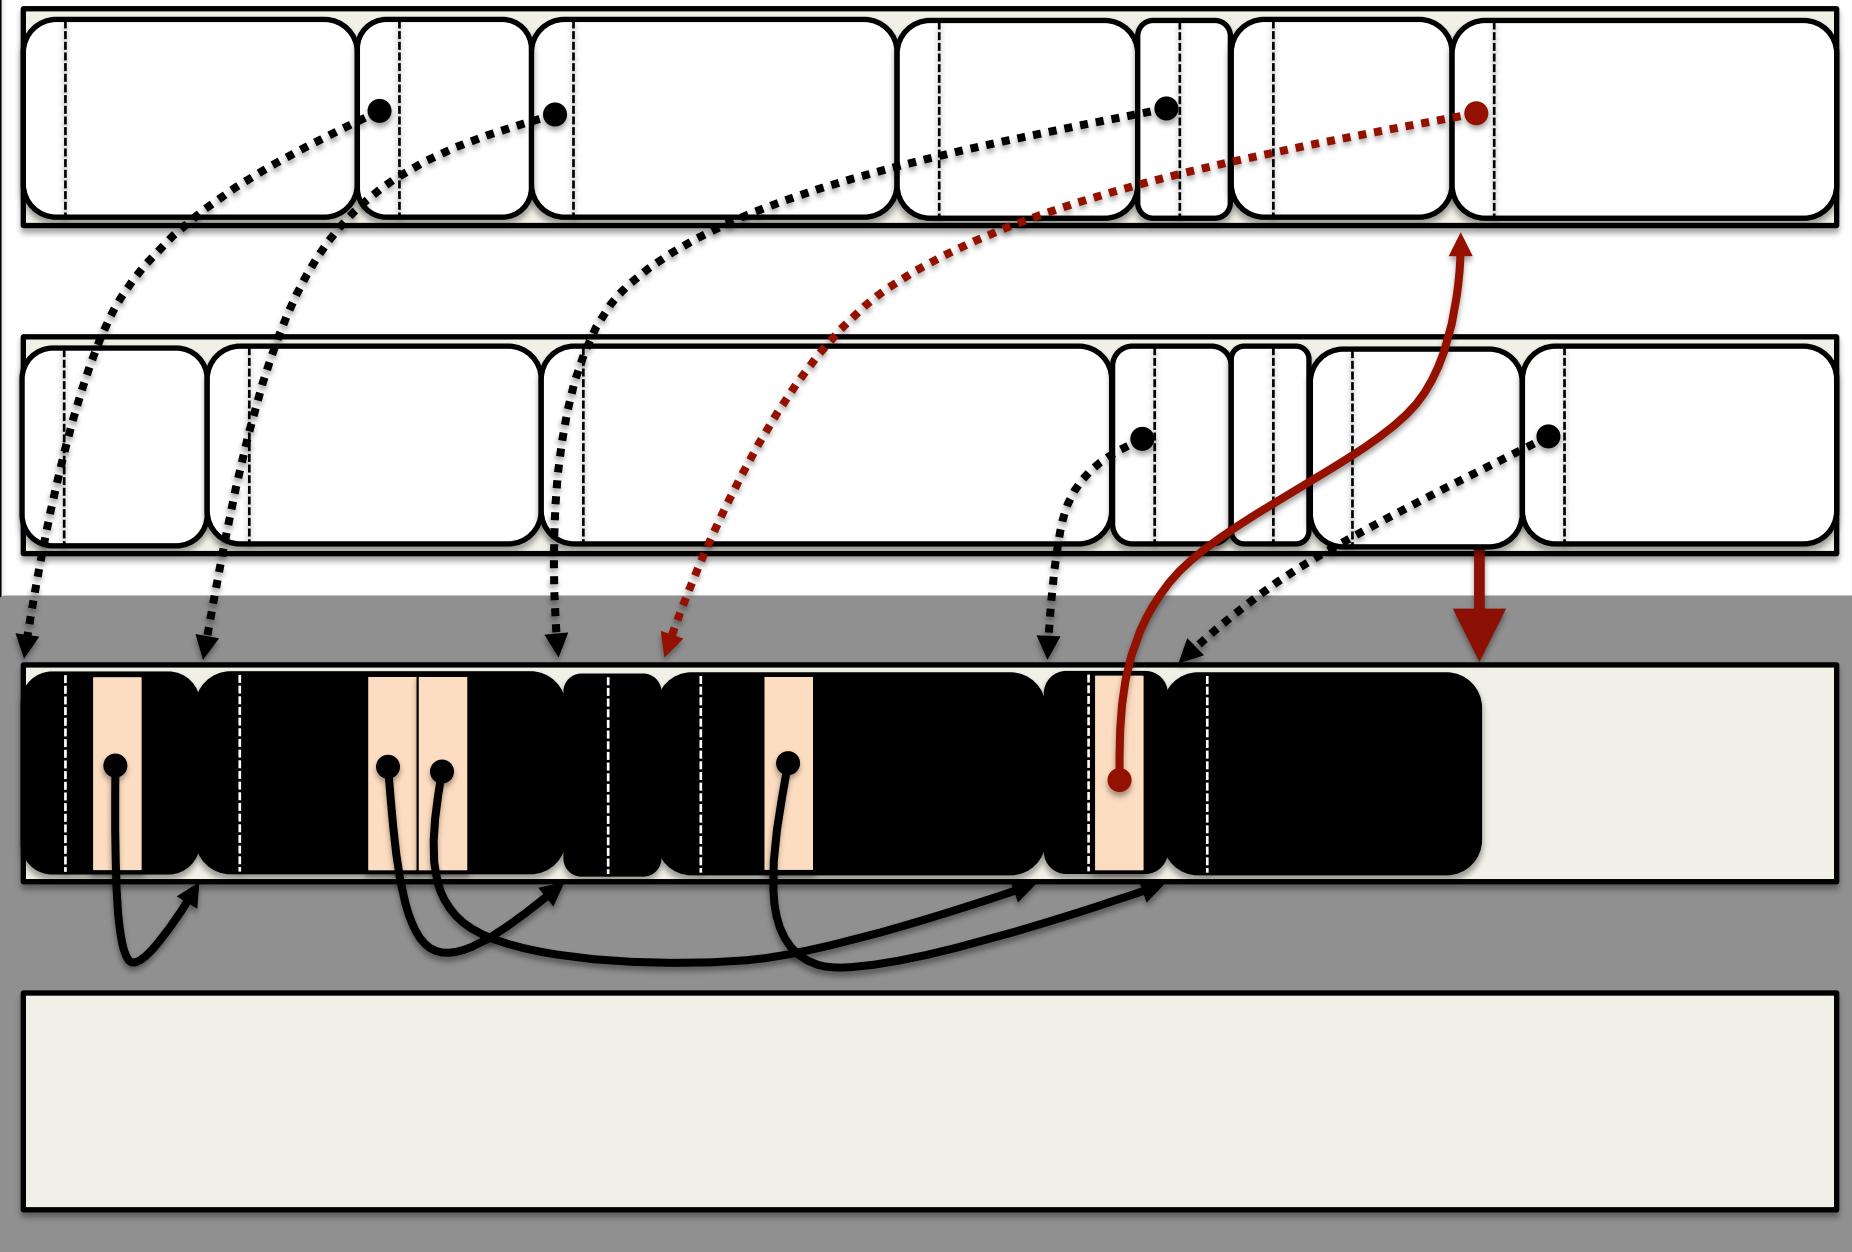


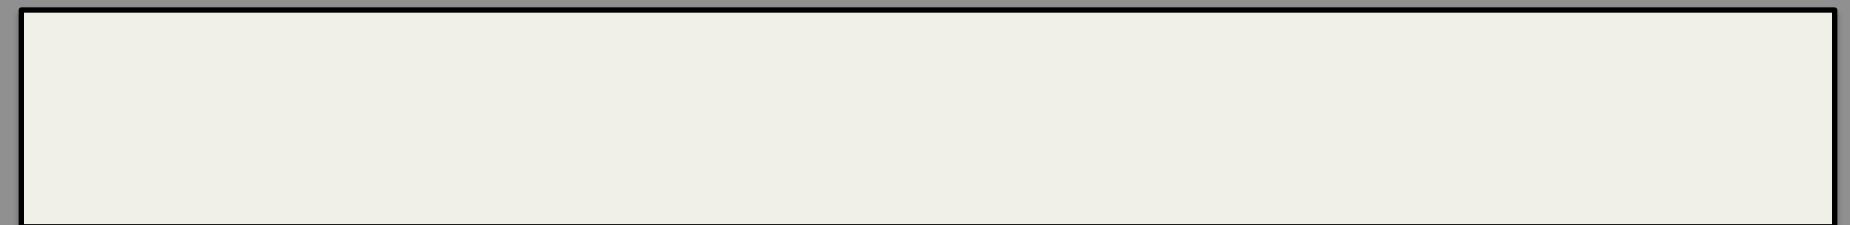
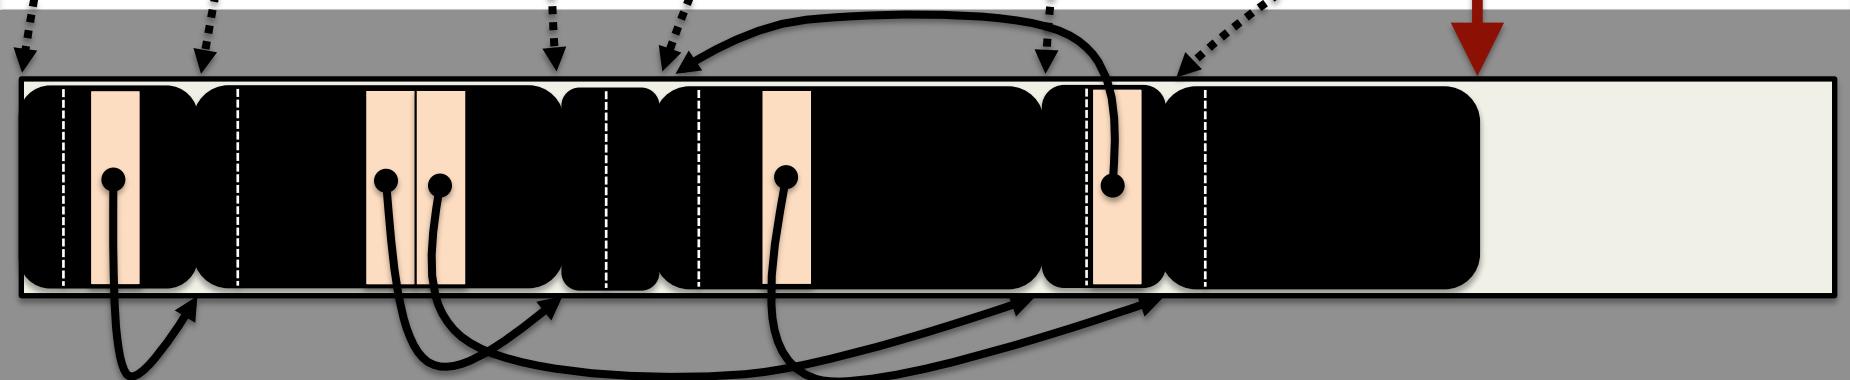
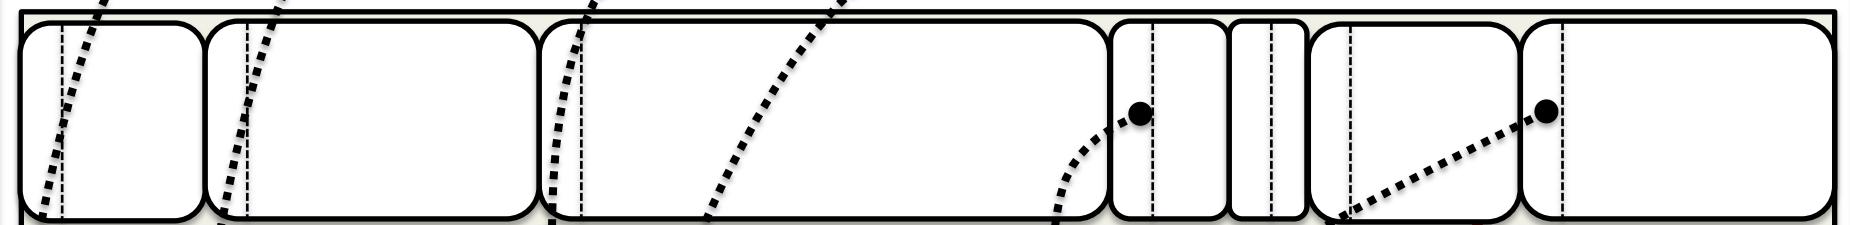
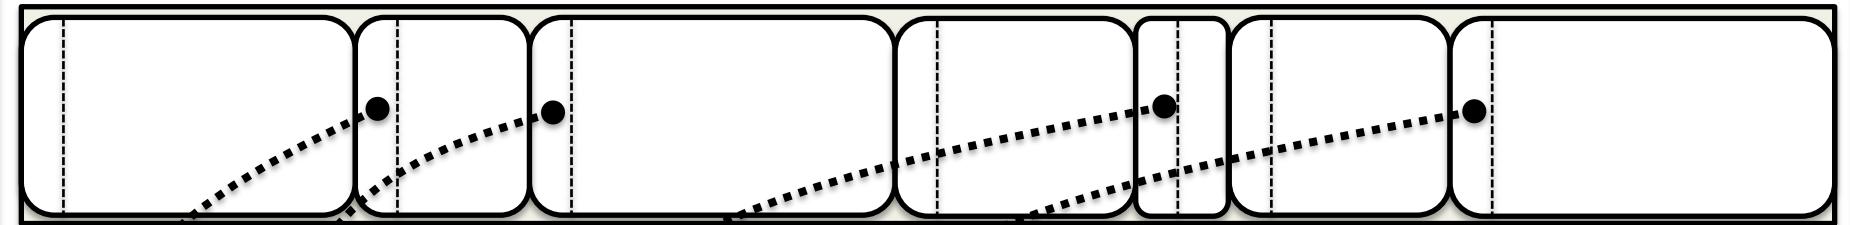






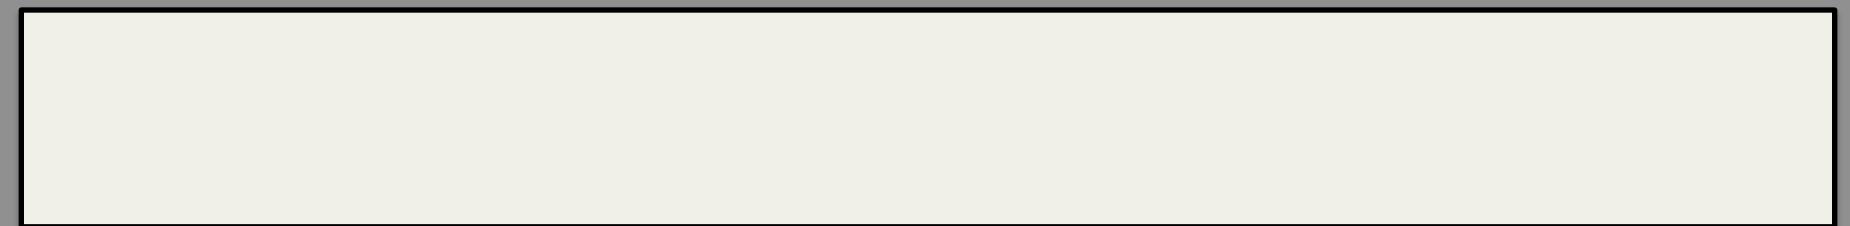
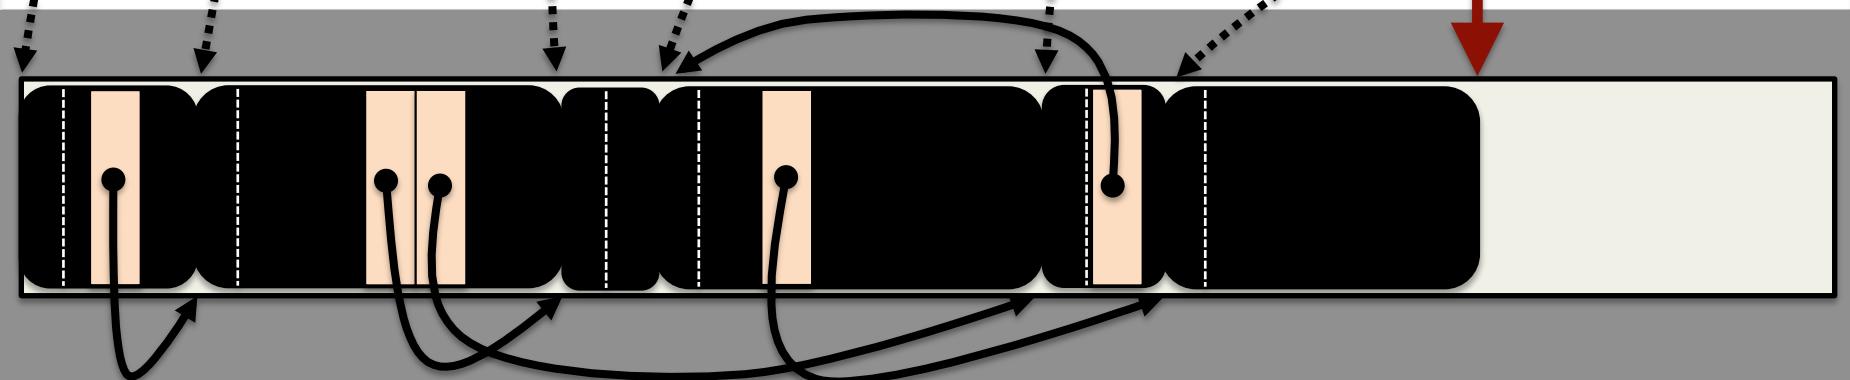
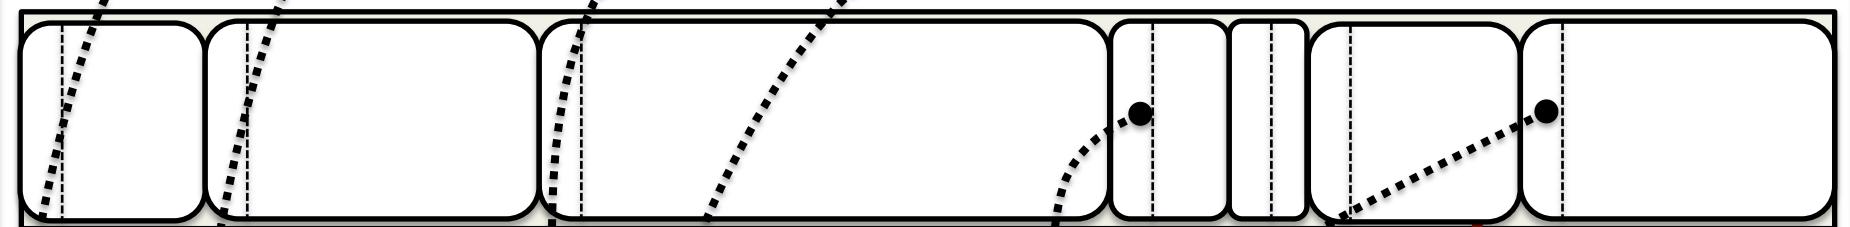
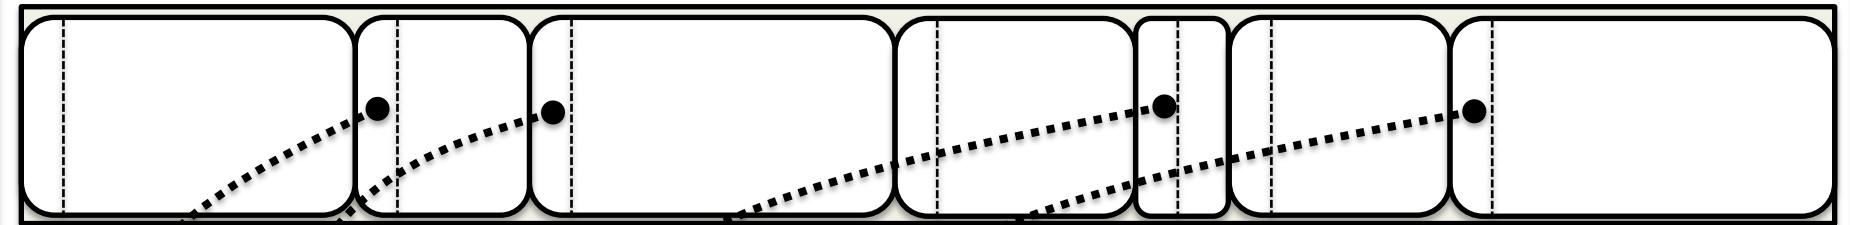


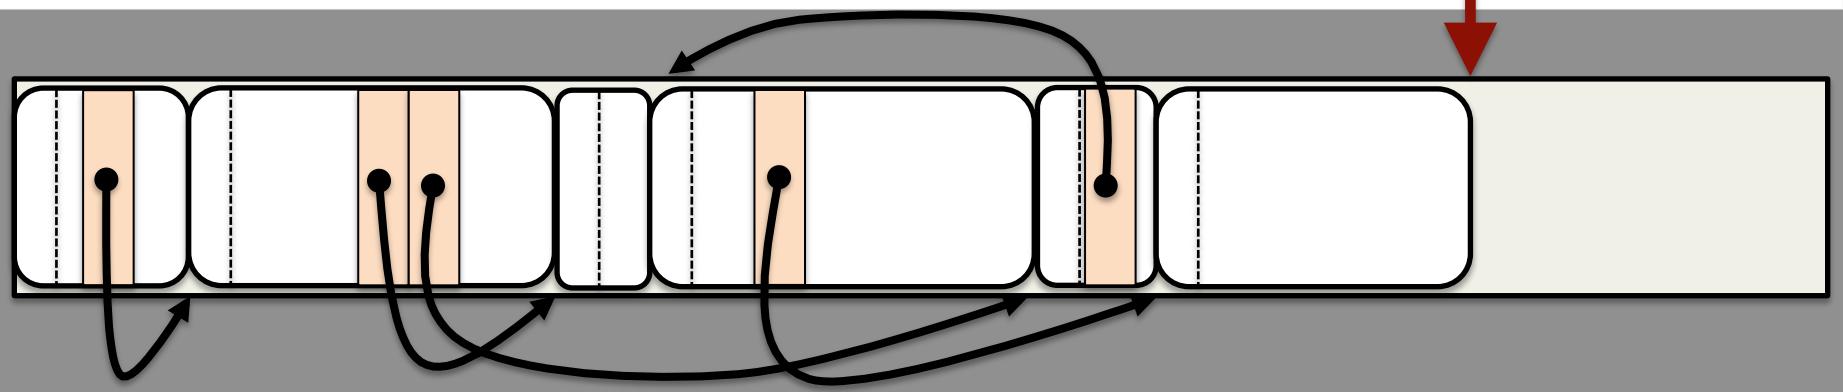


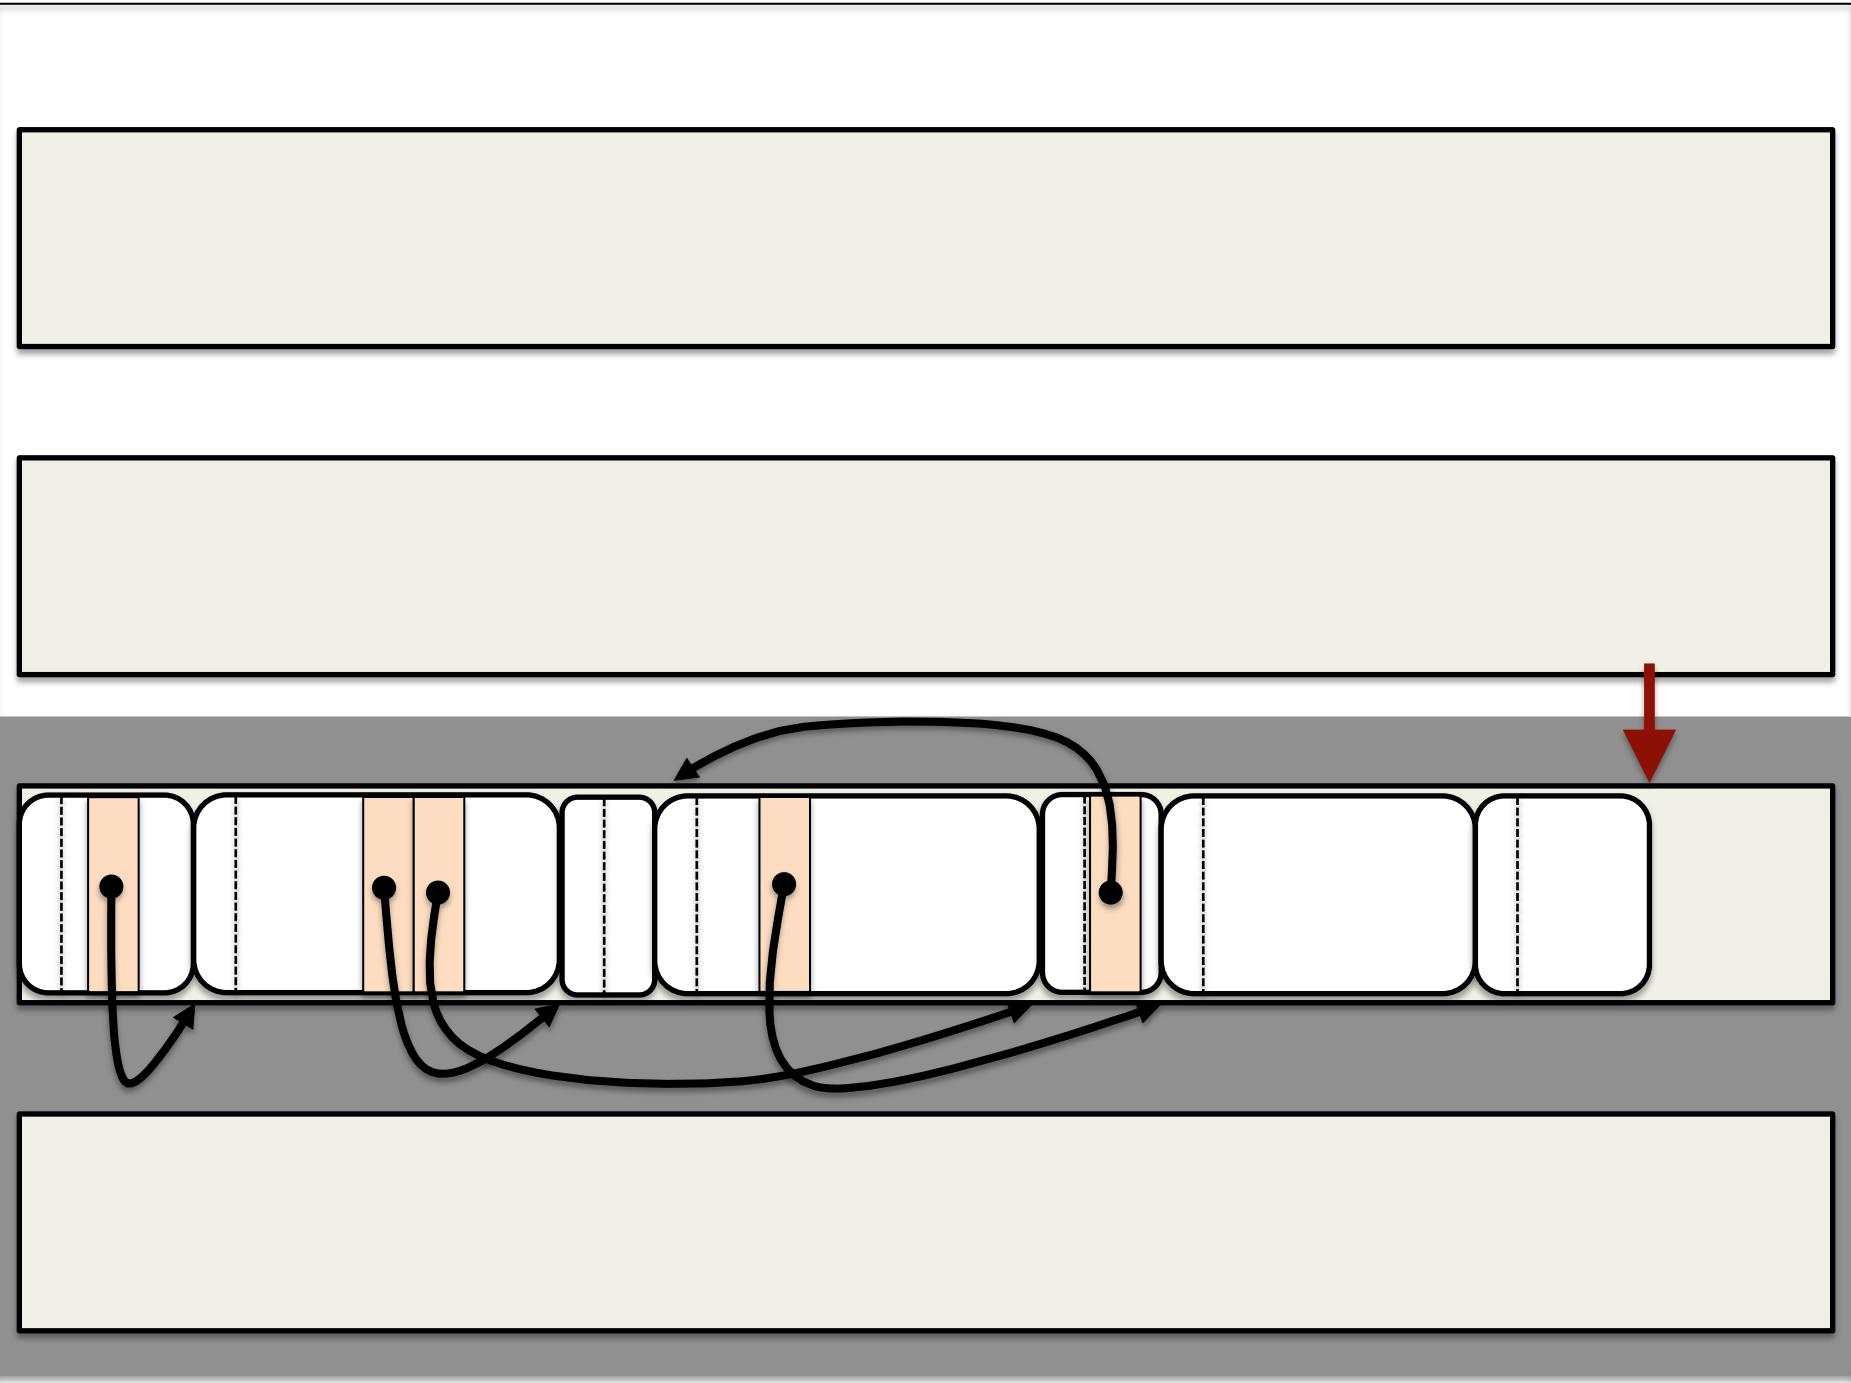


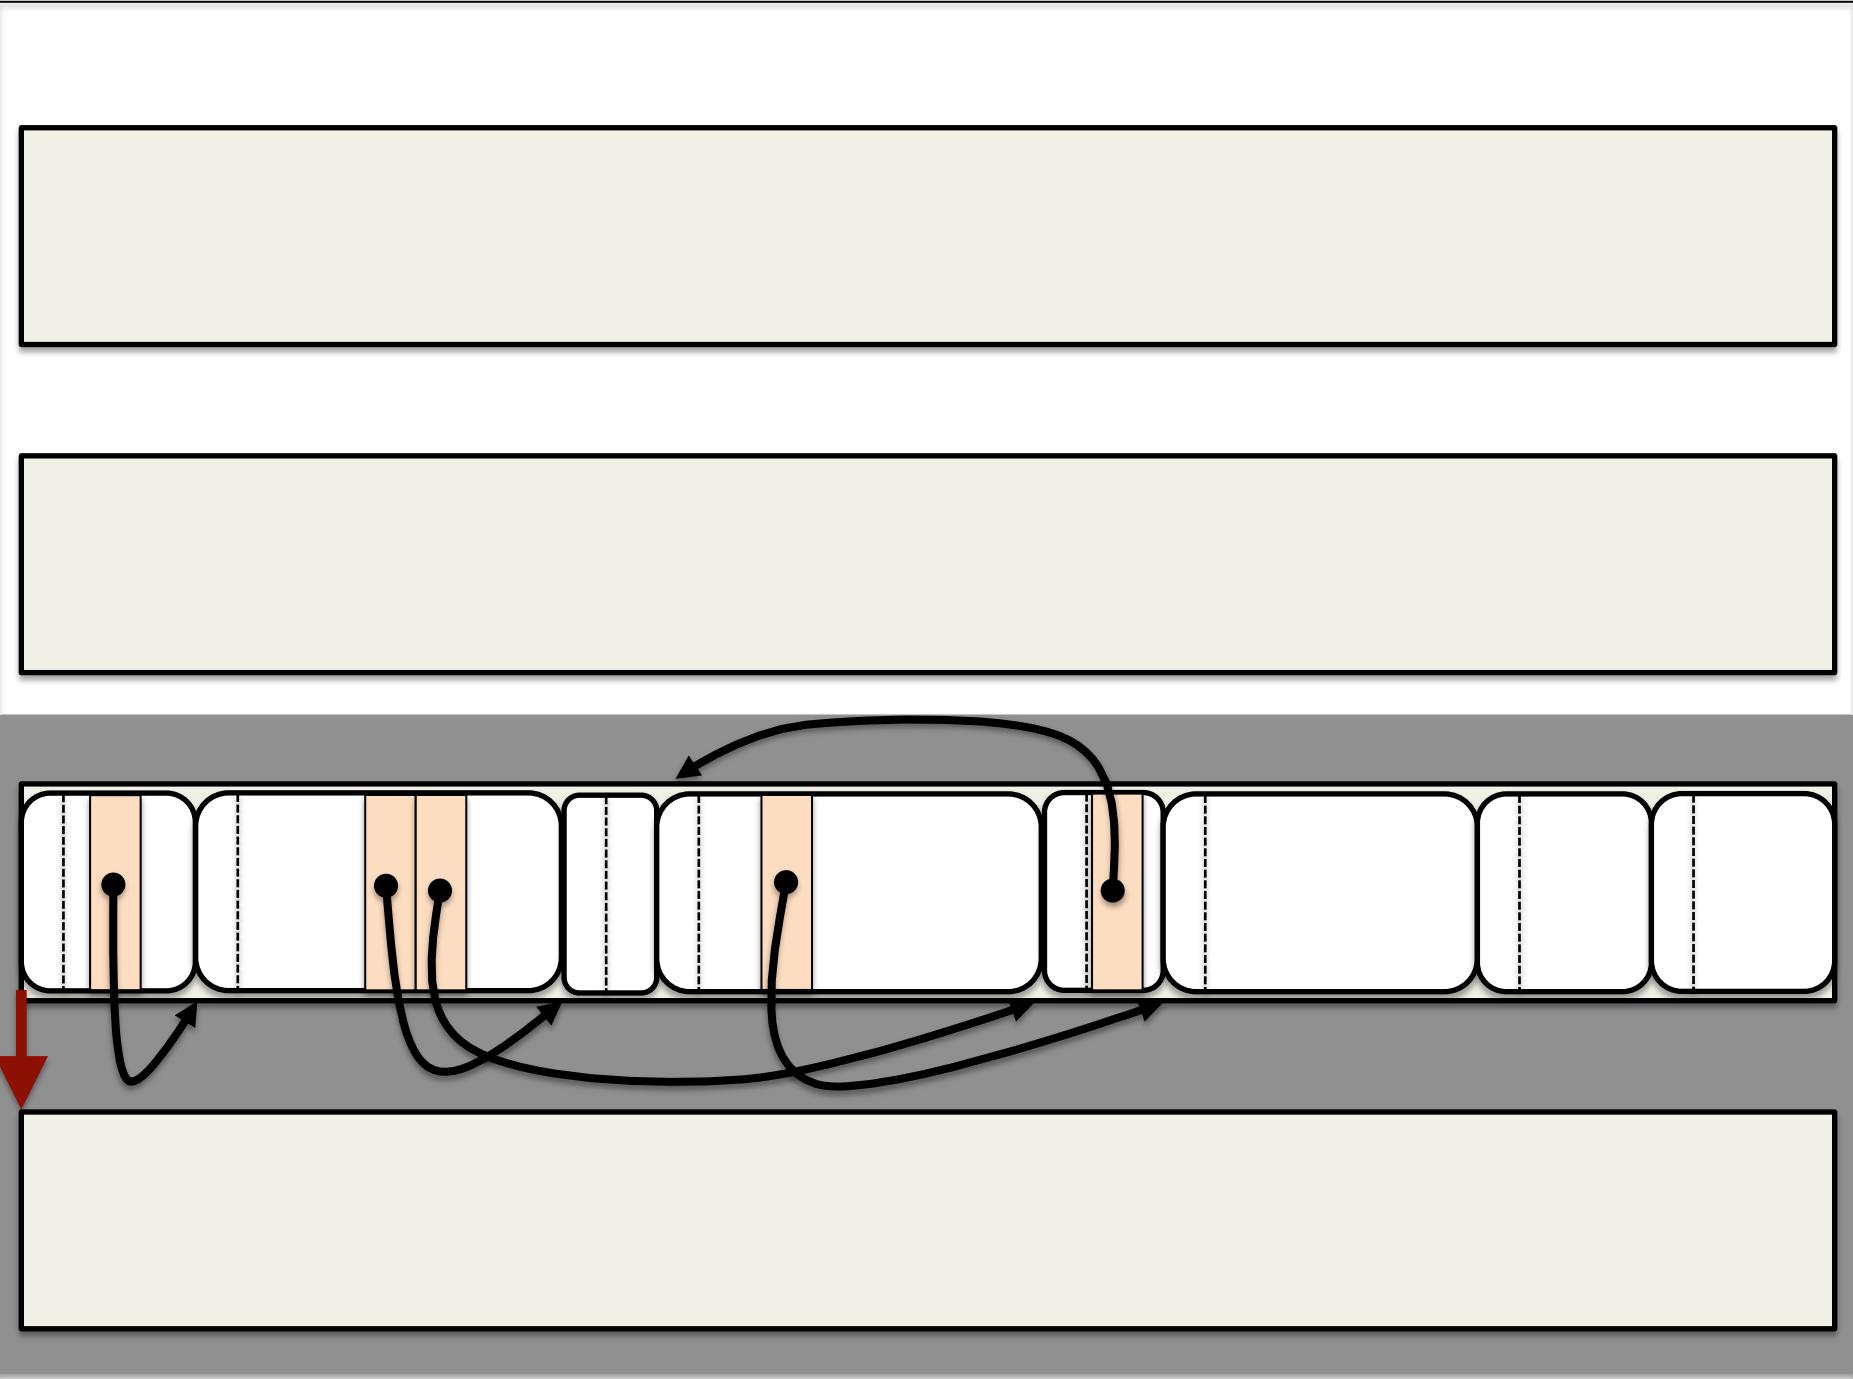
Finishing Up

- Remember to update the roots as well
 - There can be no references to the old space
- The old space can now be reused
 - Reset the bump pointer to the start of the space
- Allocation continues to the new space









Copying Collectors

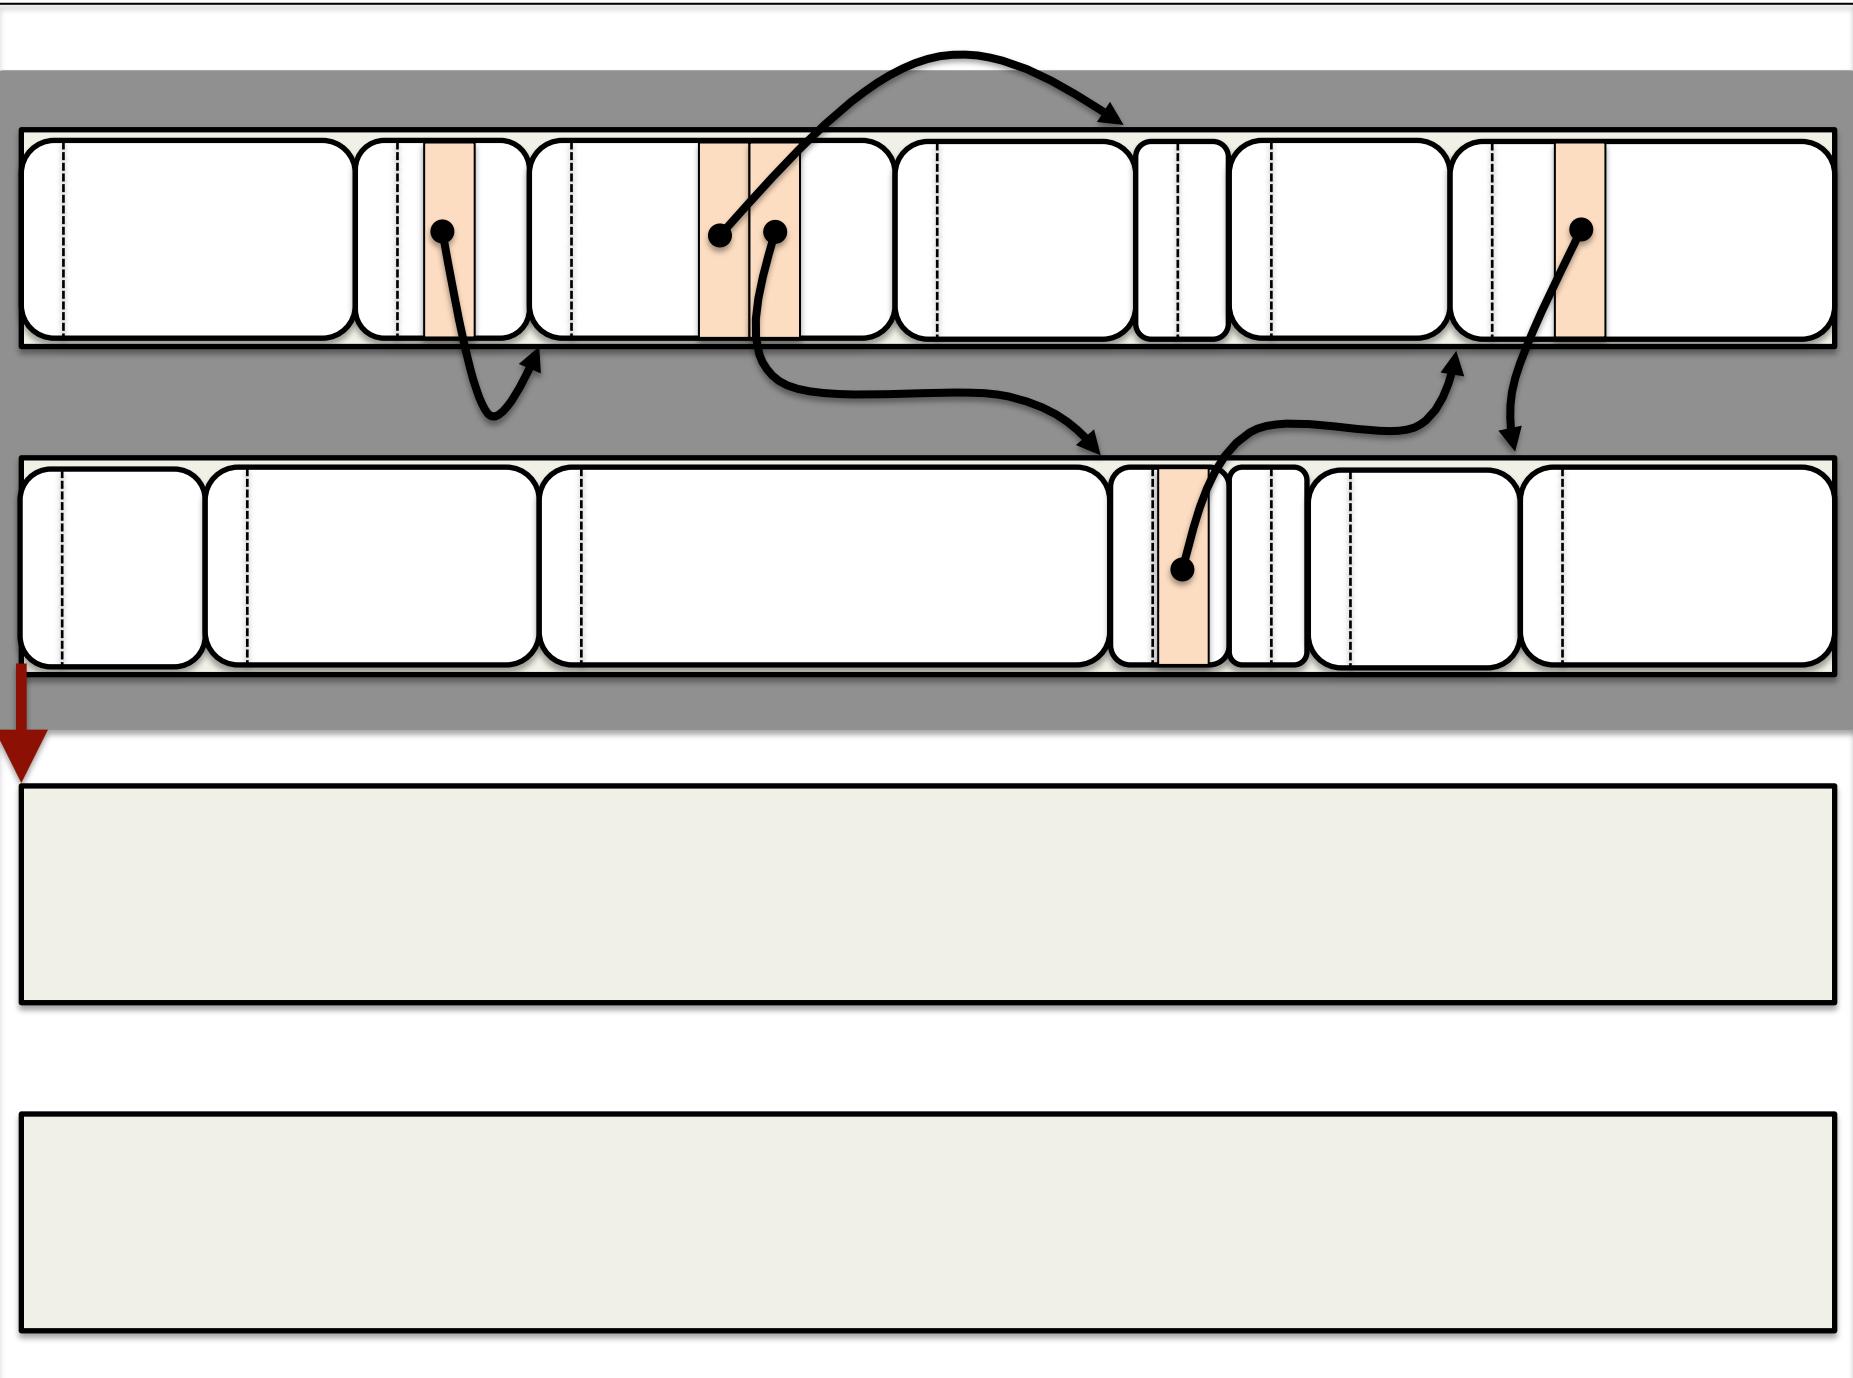
- Moving objects gives us some nice features
 - No fragmentation
 - Much better cache behavior
- Comes with some large overheads
 - Copying can be expensive
 - Wastes half of the heap
- We'll see how to reduce these over time

Object Mobility

- For copying to work, all objects must be mobile
- Any pinned objects create holes in the heap
 - Bump-pointer allocation can't use the holes
 - Scanning algorithms become much harder
 - Basic assumption is that heap is contiguous
- Imposes limitations throughout the VM
 - Particularly when working with native code

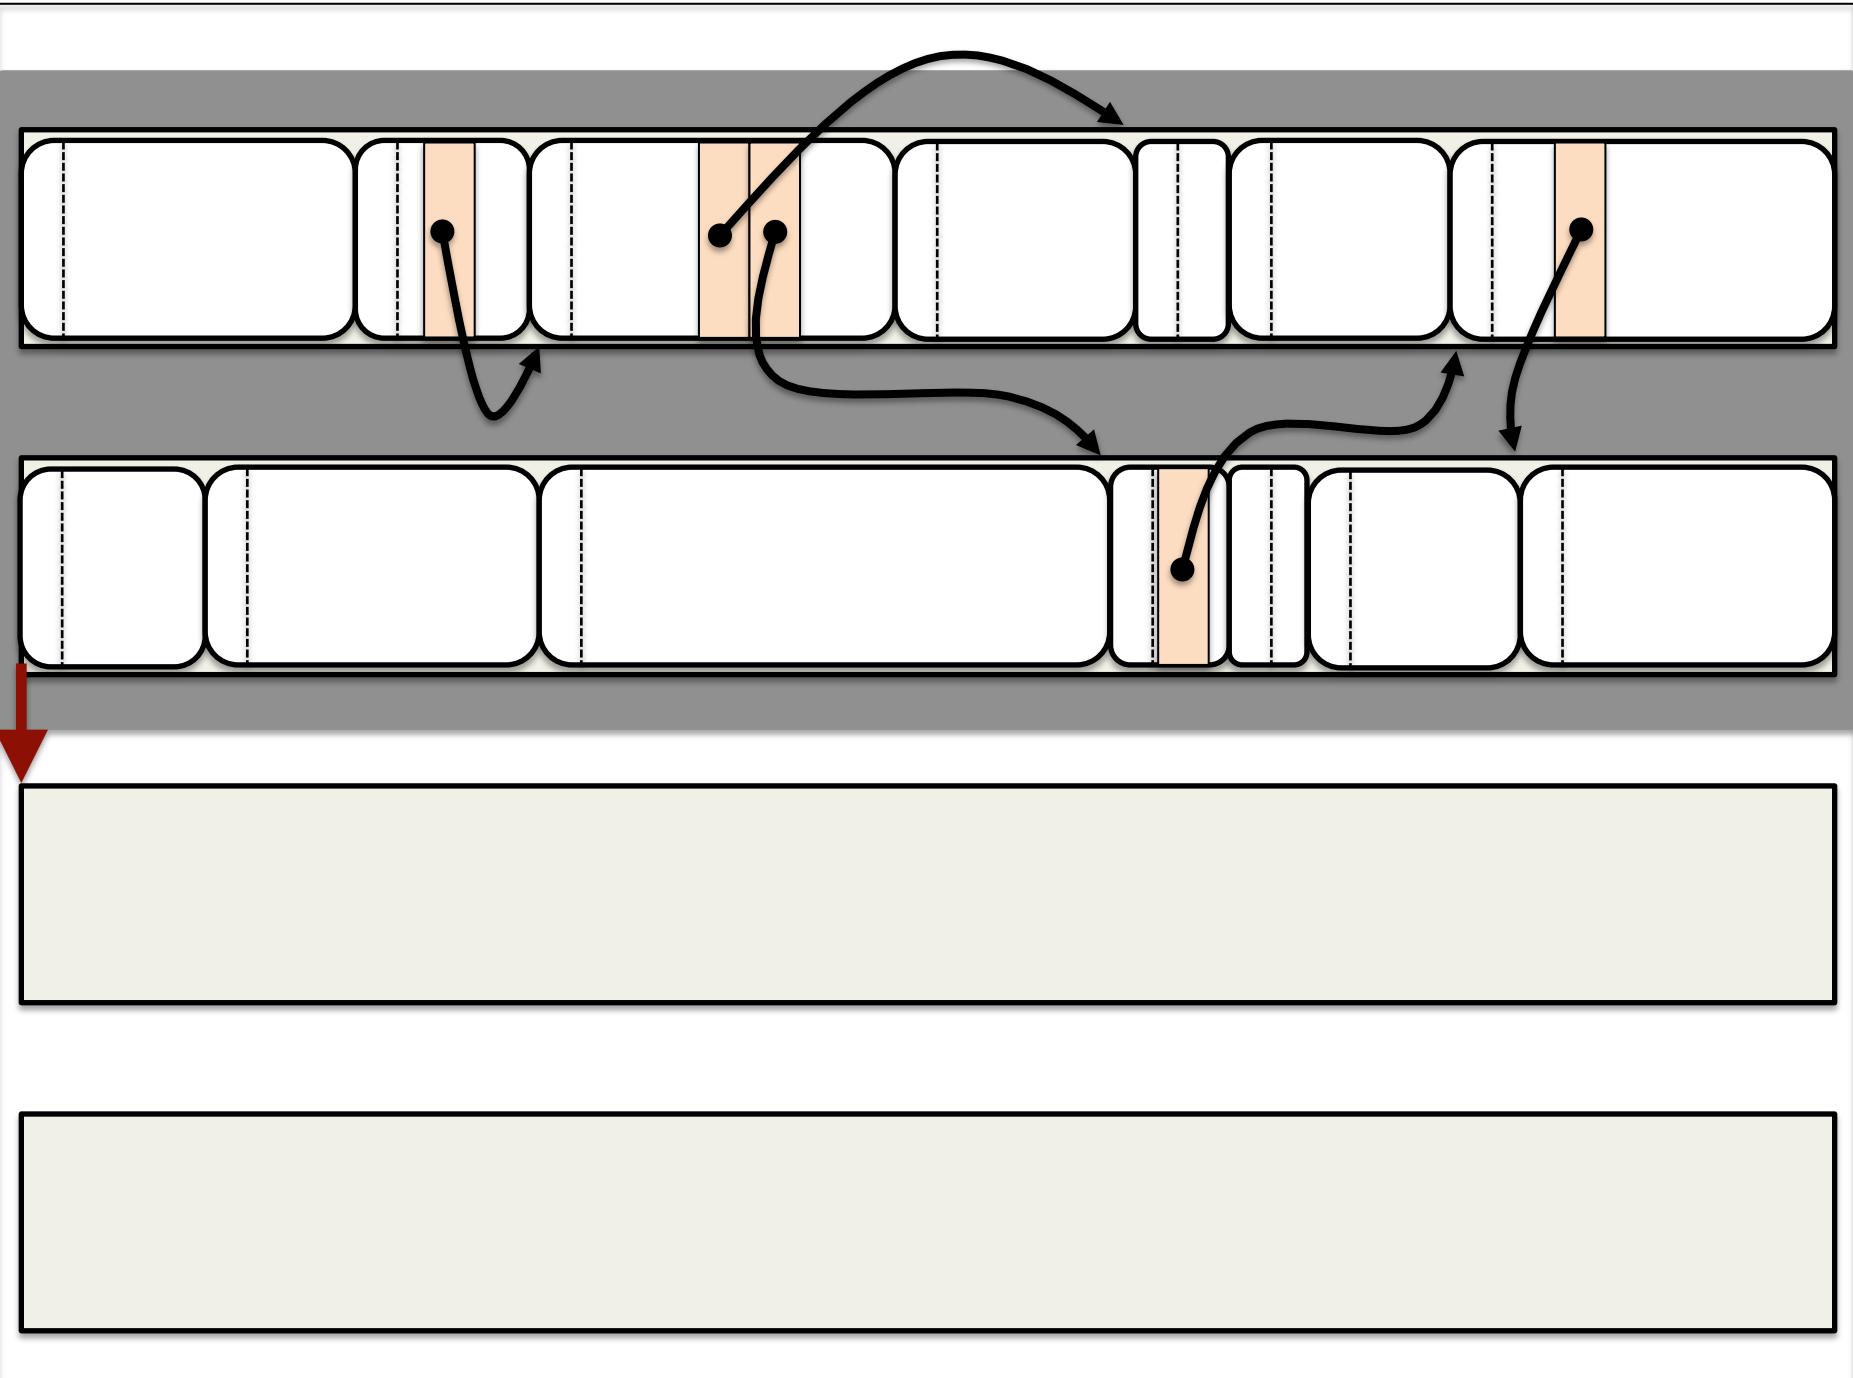
Cheney's Algorithm

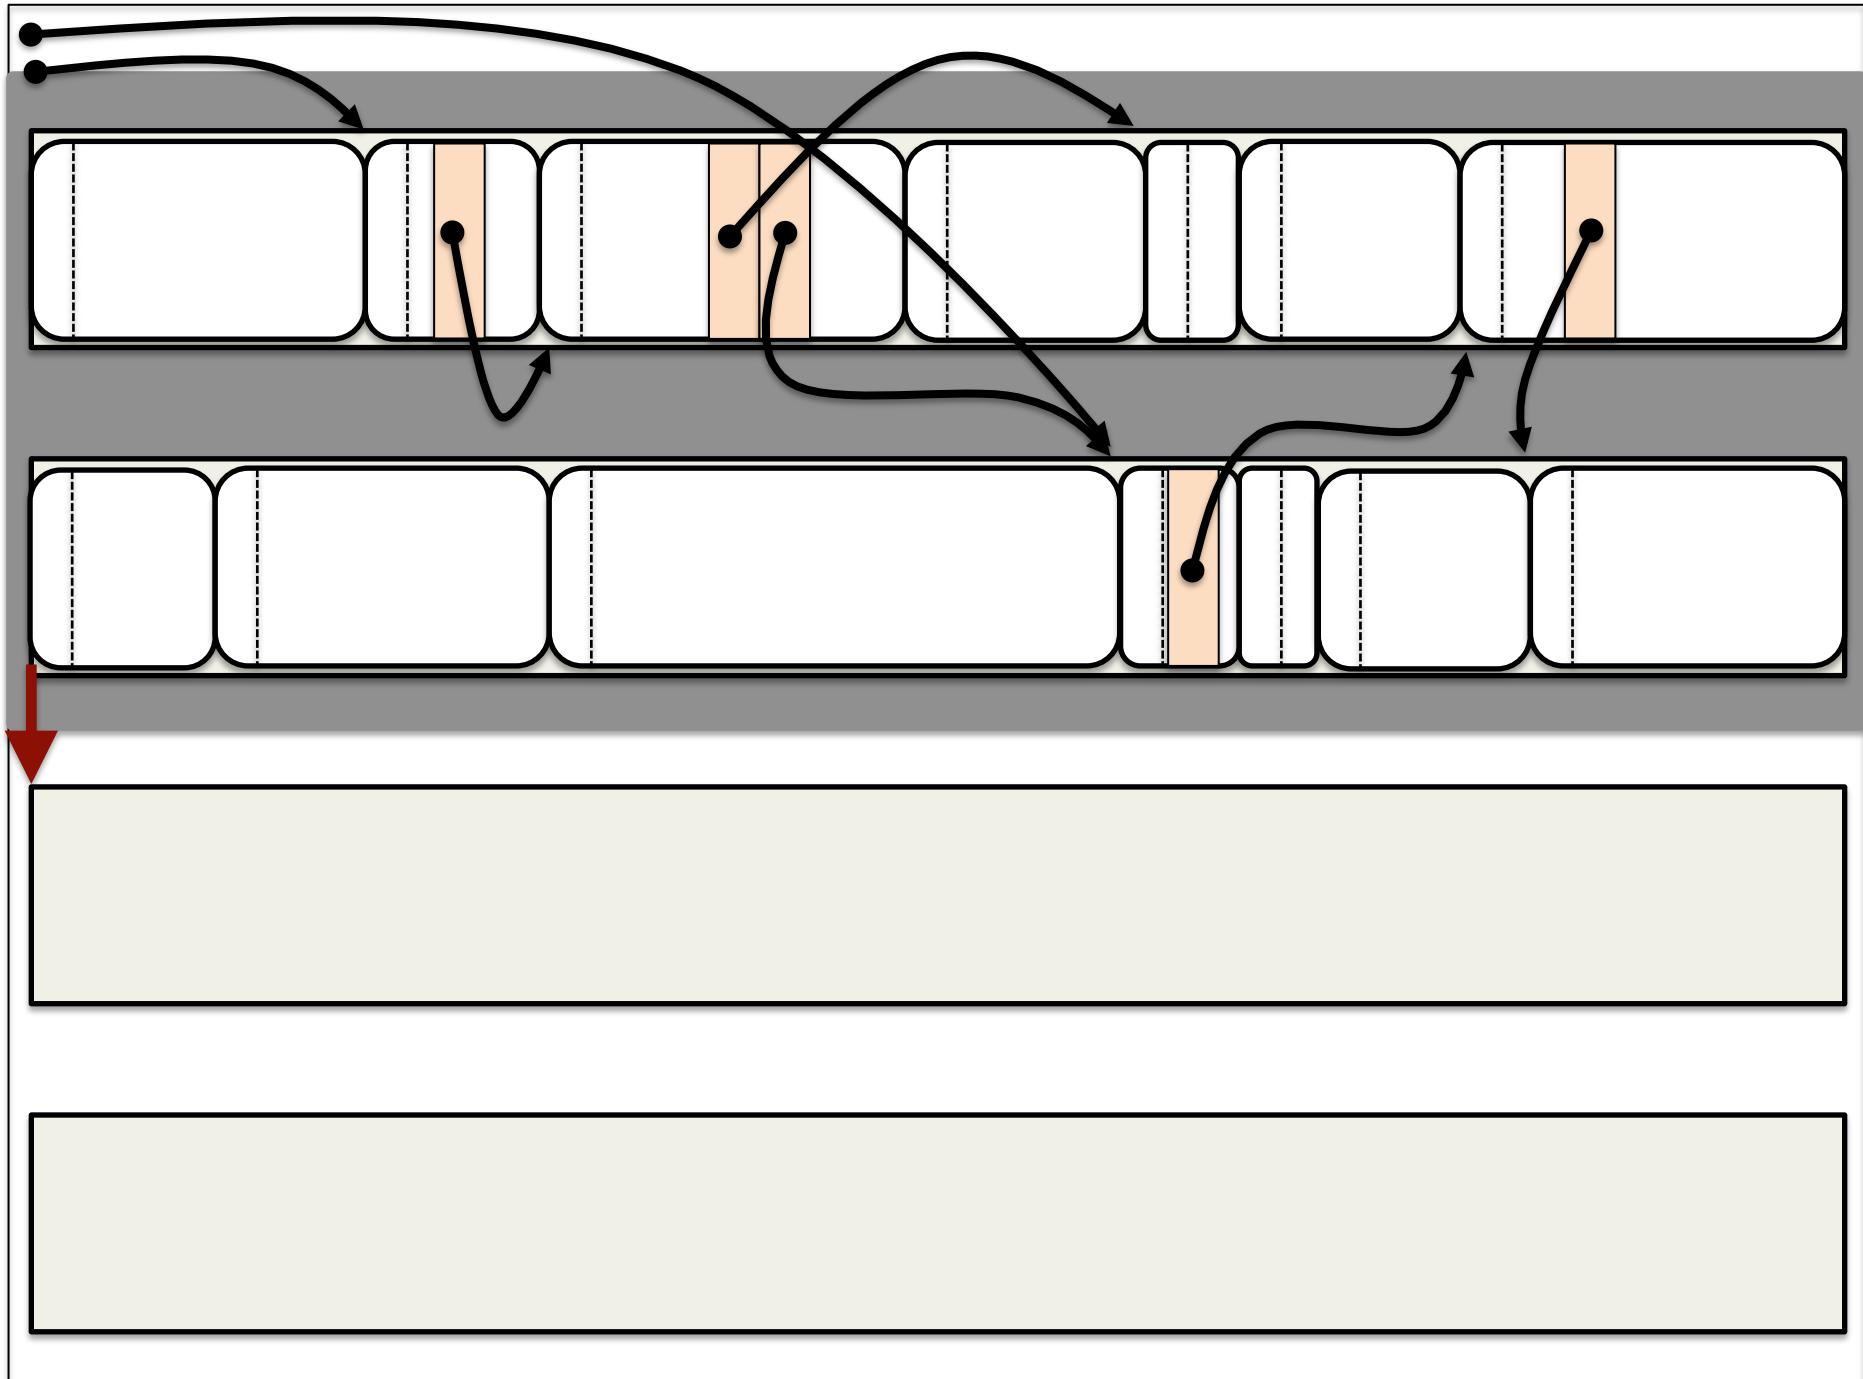
- Proposed in 1970
 - Tends to be the canonical copying algorithm
- Uses a single pass of the heap
 - Merges marking and copying phases
 - Updates references as it goes
- You will implement this algorithm in SimpleJava
 - Assignment 3 is now available

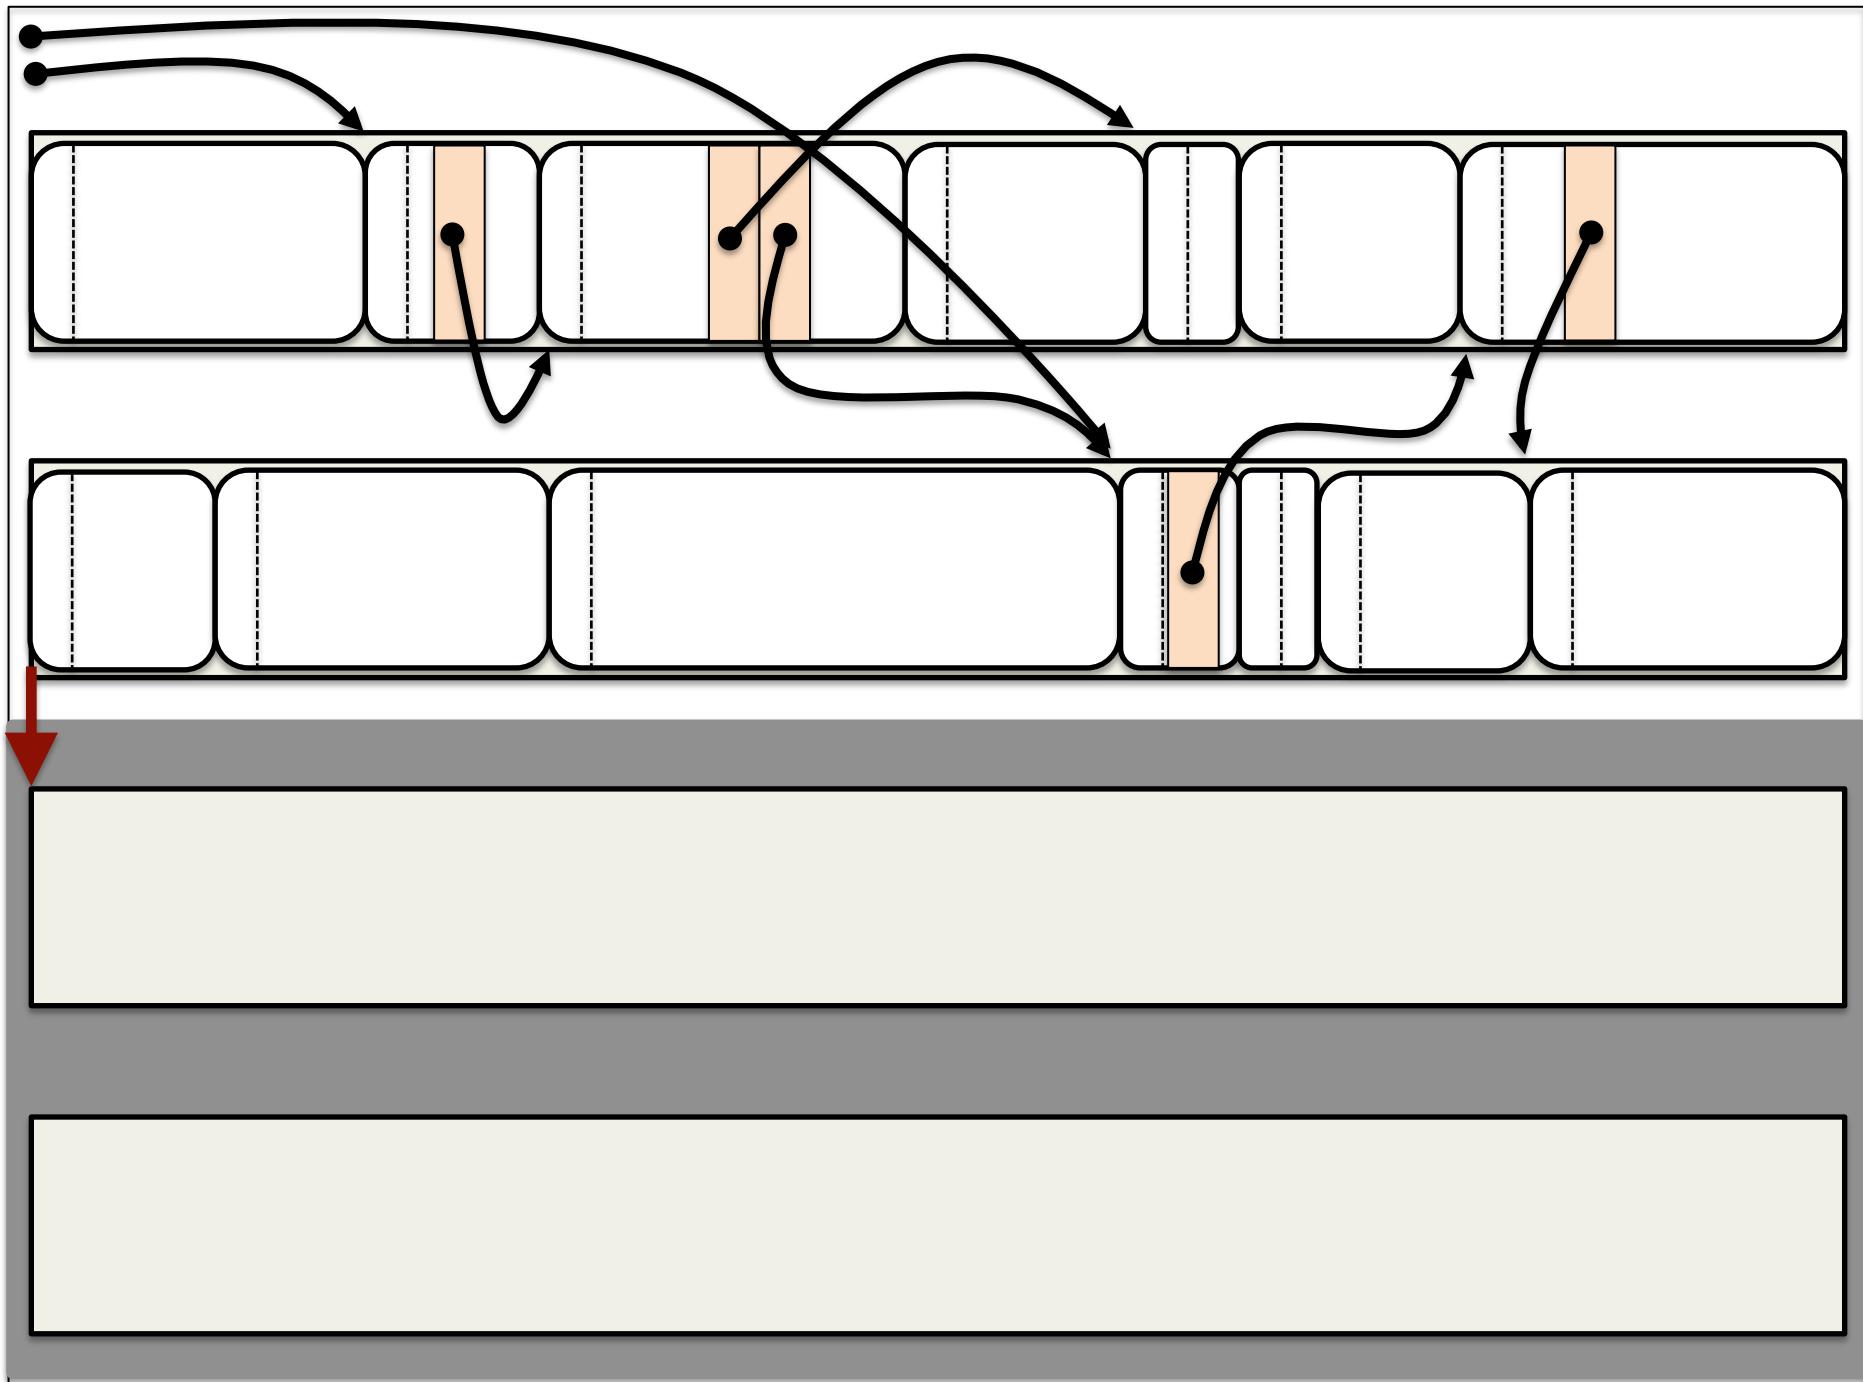


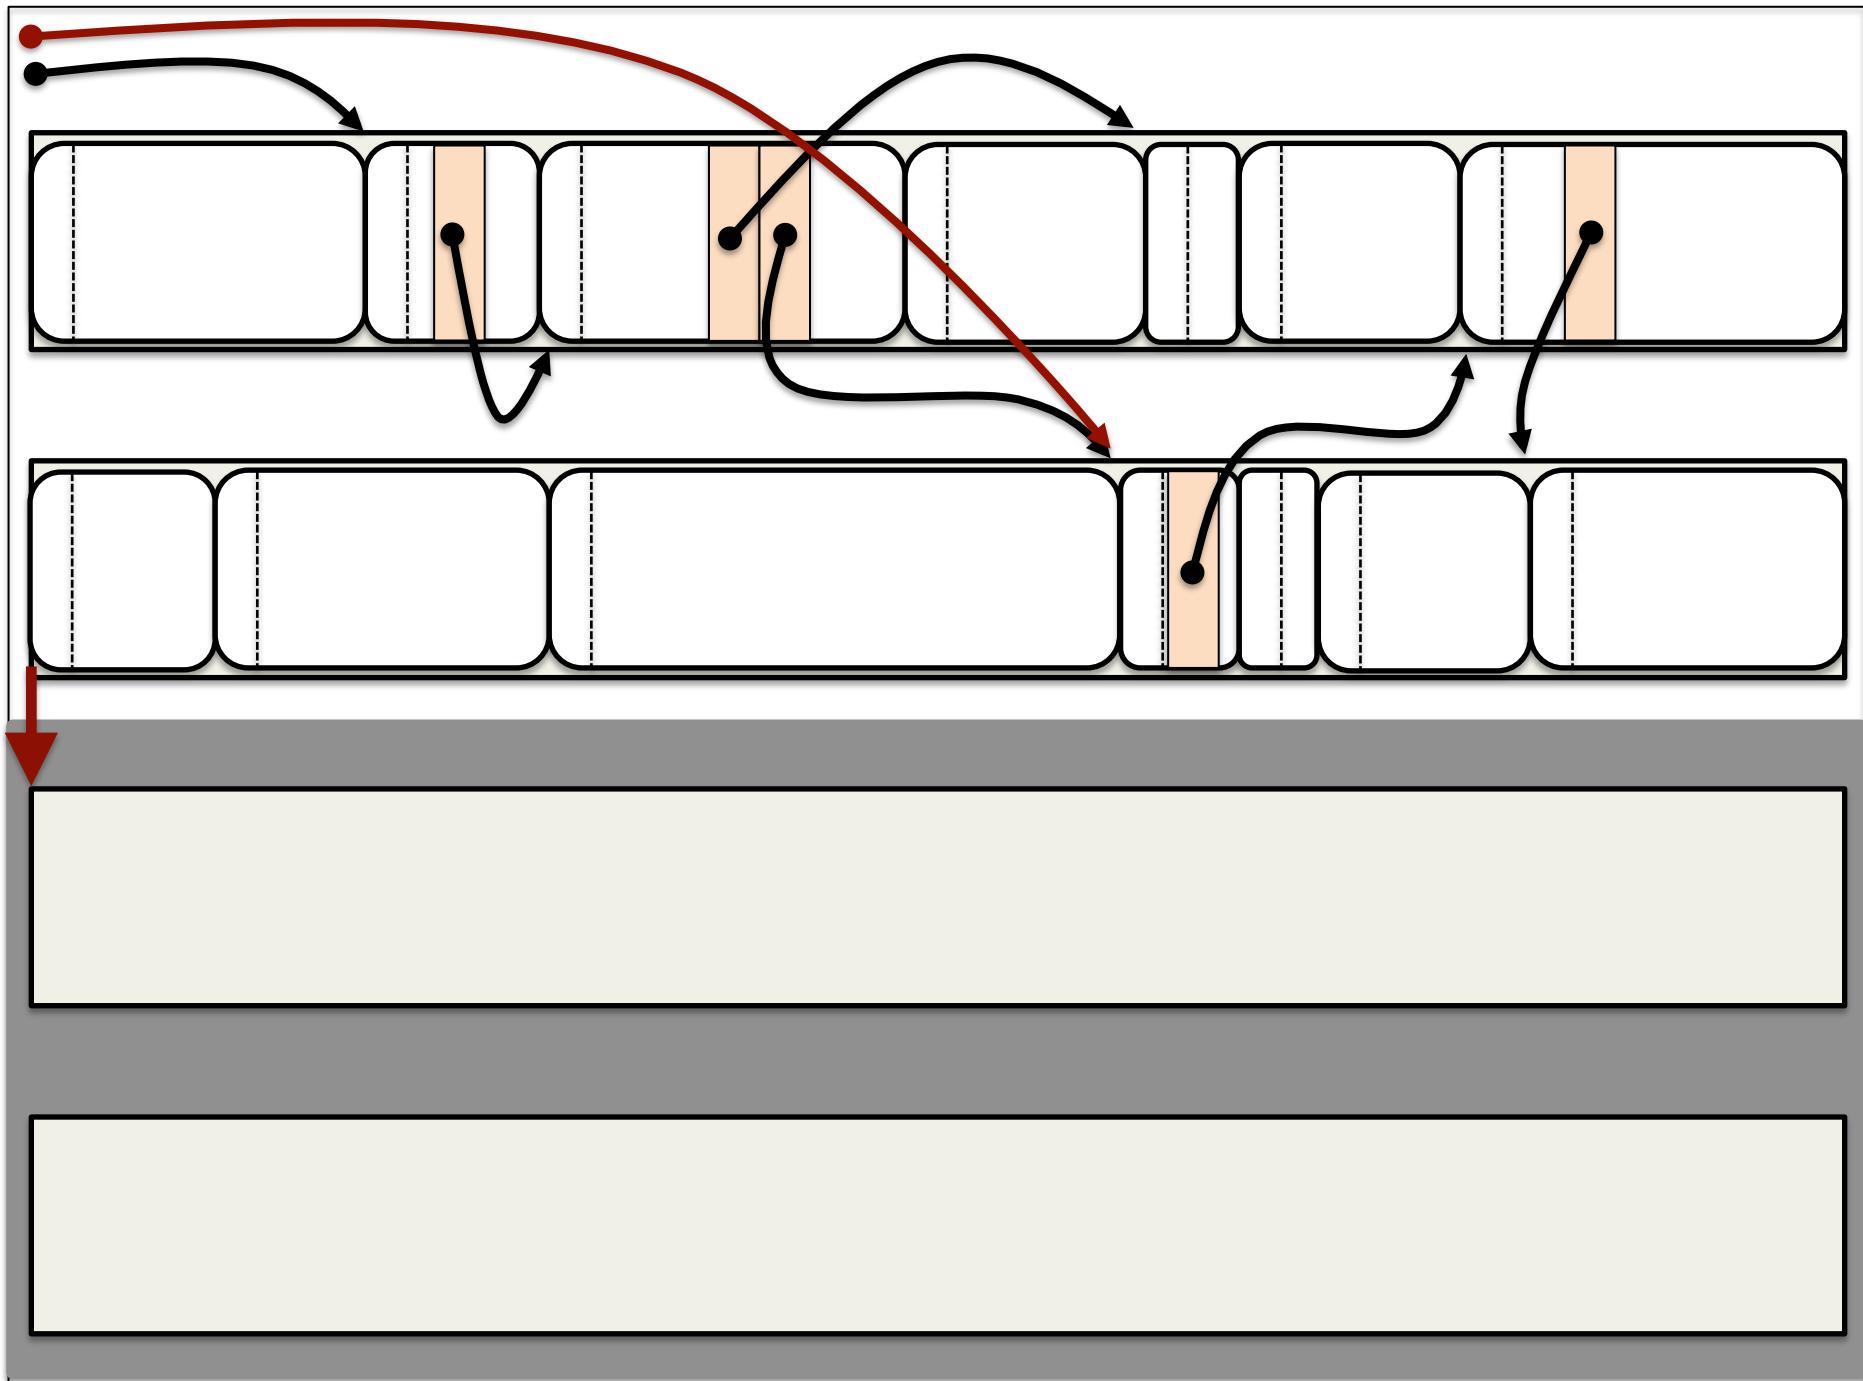
Root Scanning

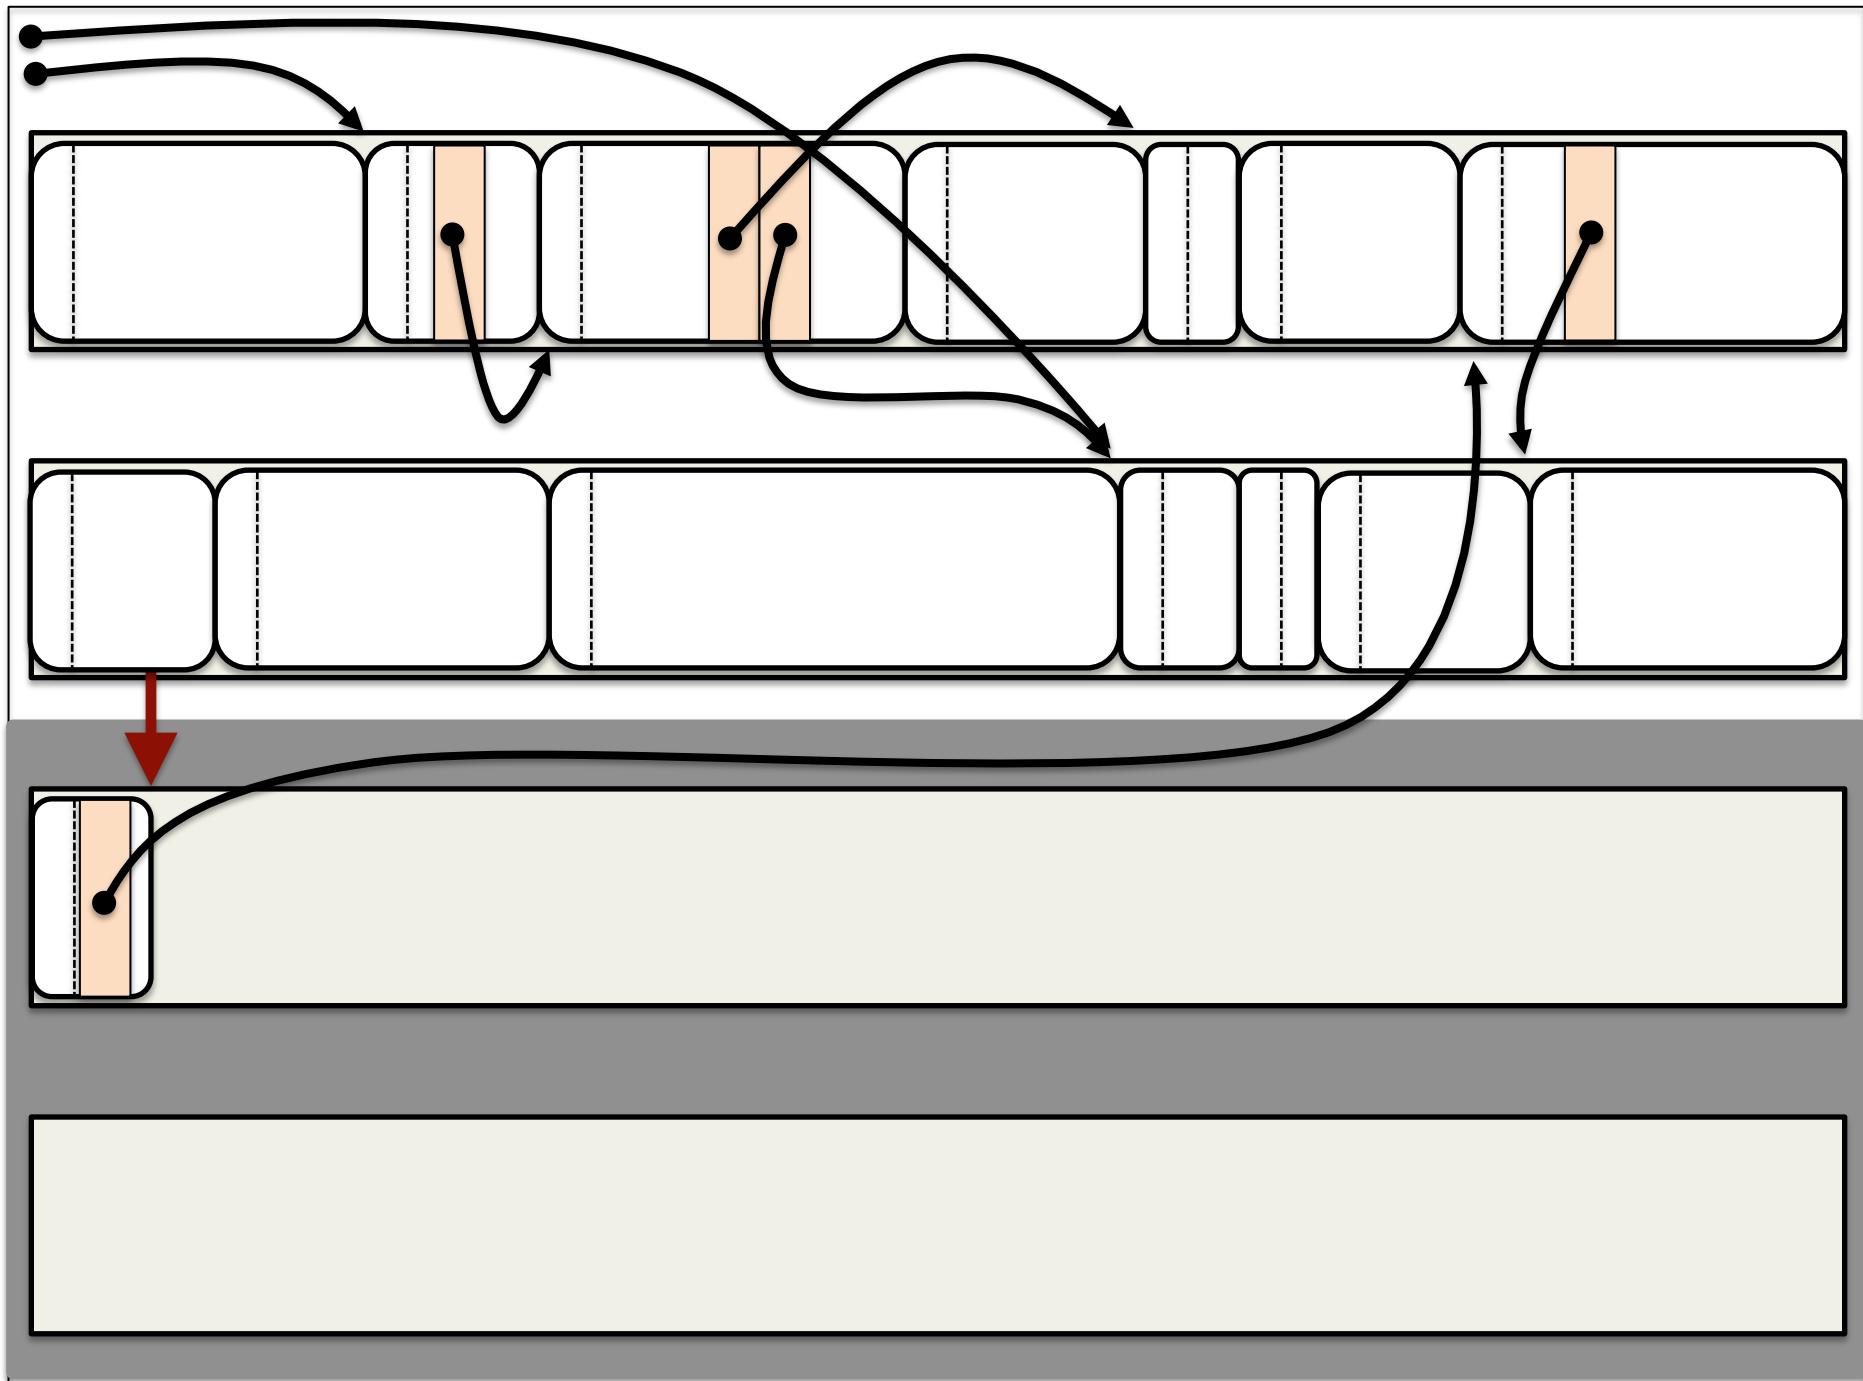
- First step of any tracing algorithm
 - Copy all objects pointed to by a root
- Root set of a VM is implementation dependent
- SimpleJava has three root locations
 - Thread stack and local references
 - Class static references
 - String intern table

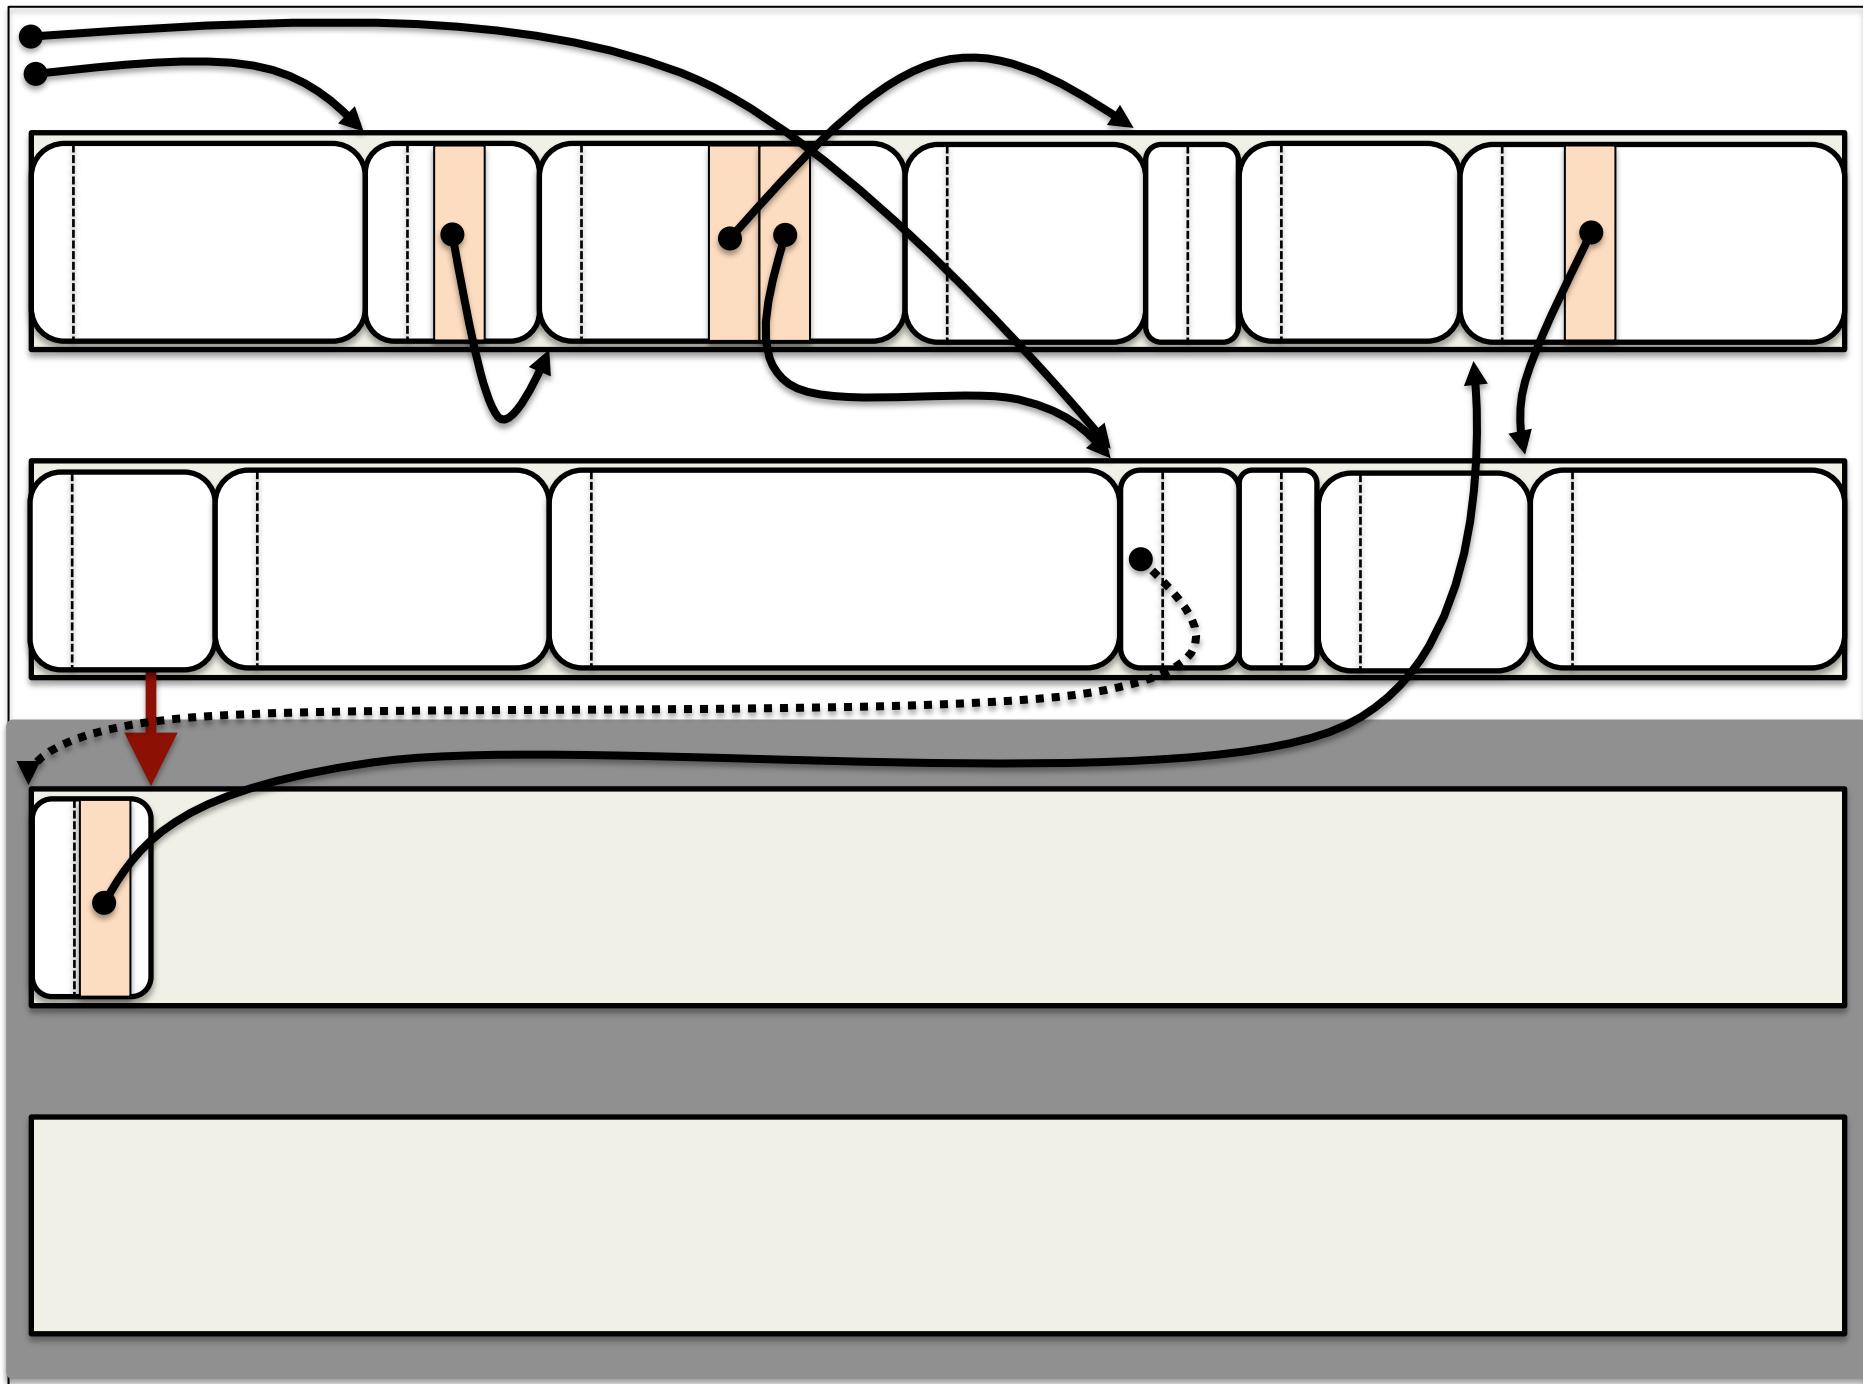


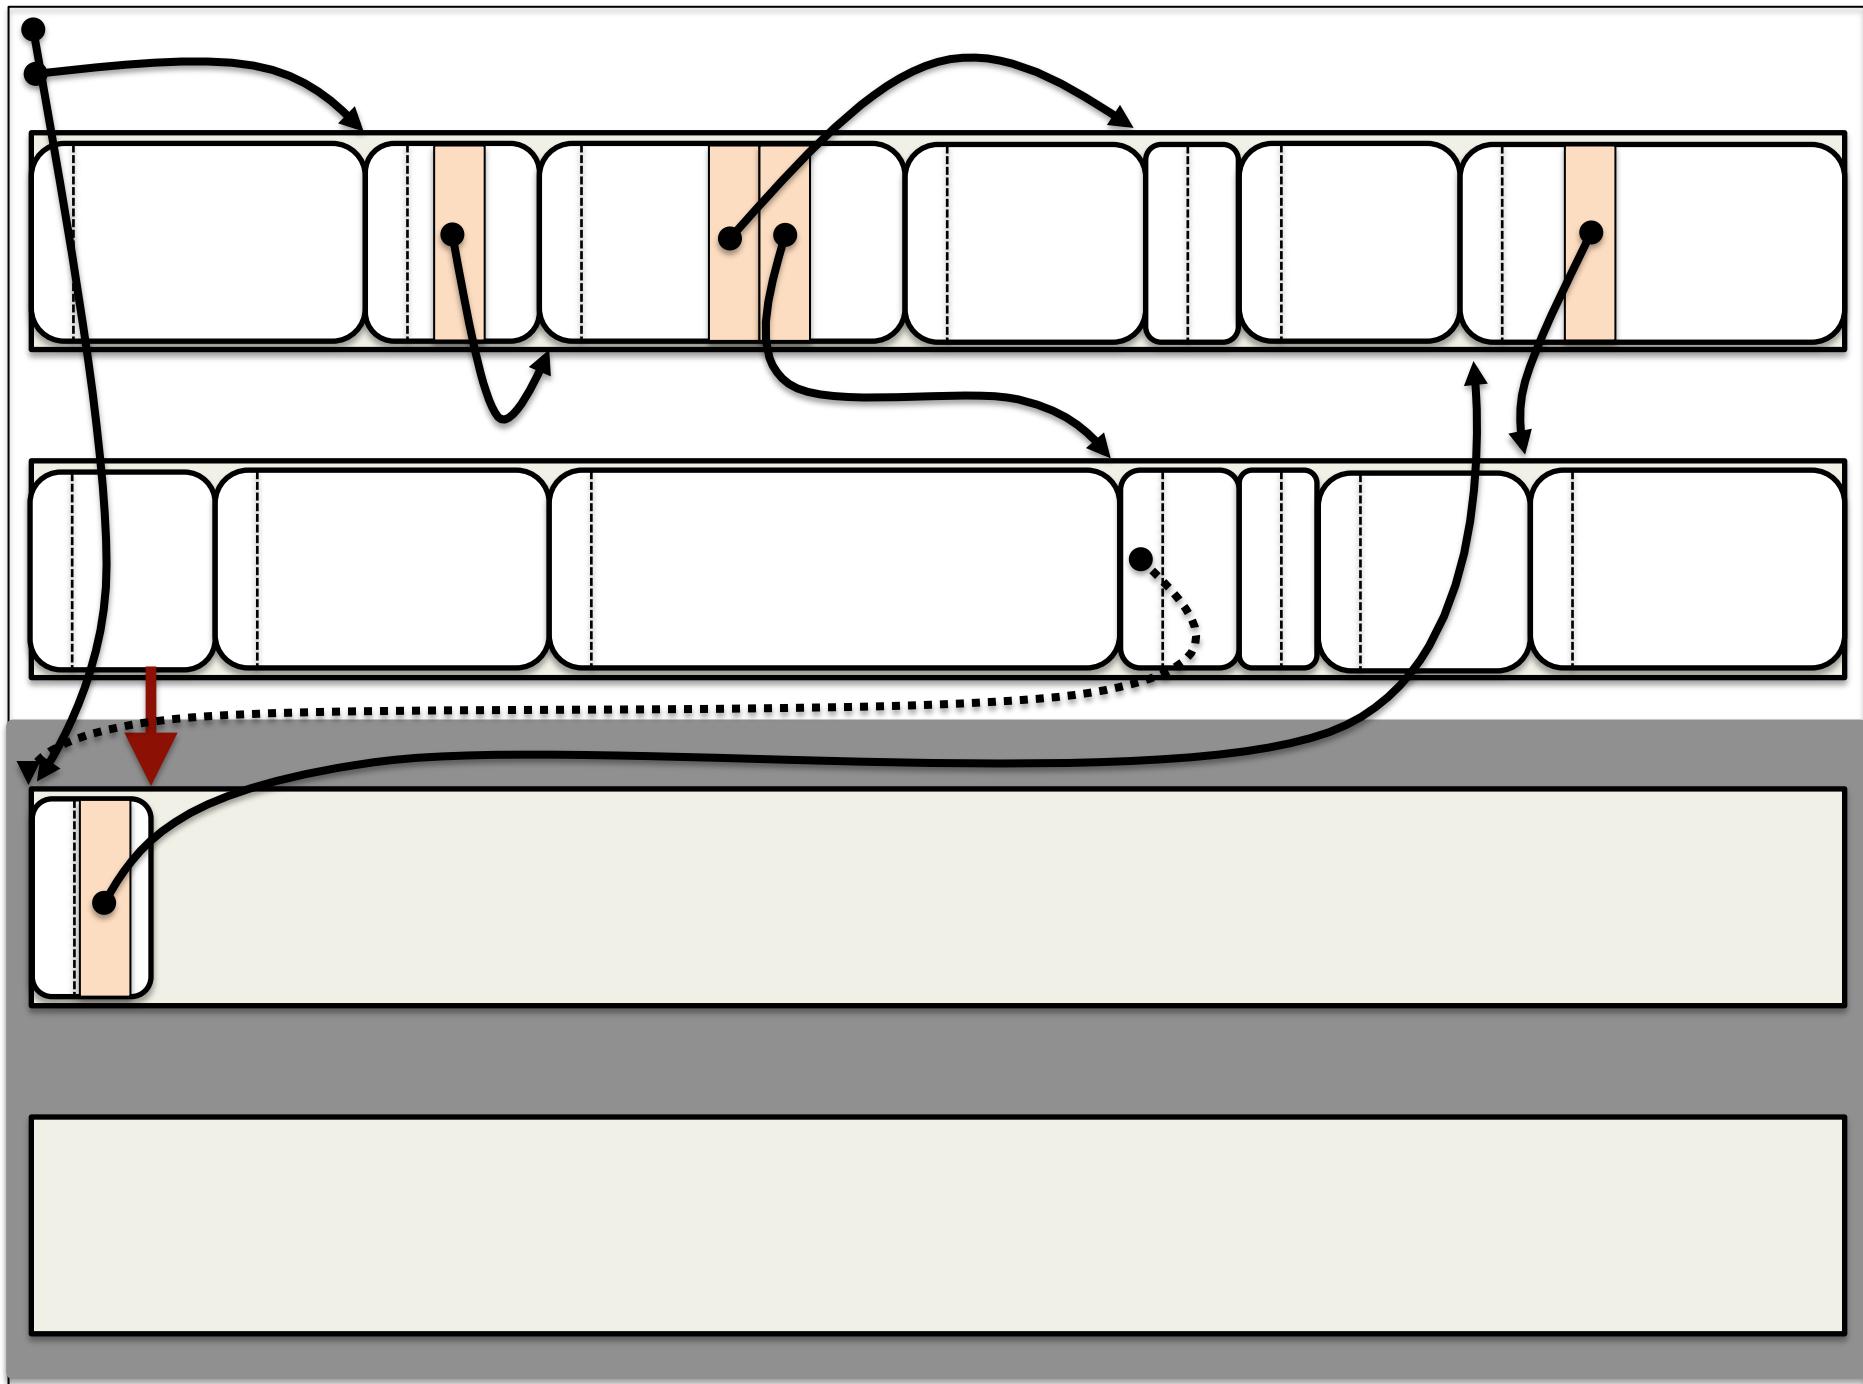


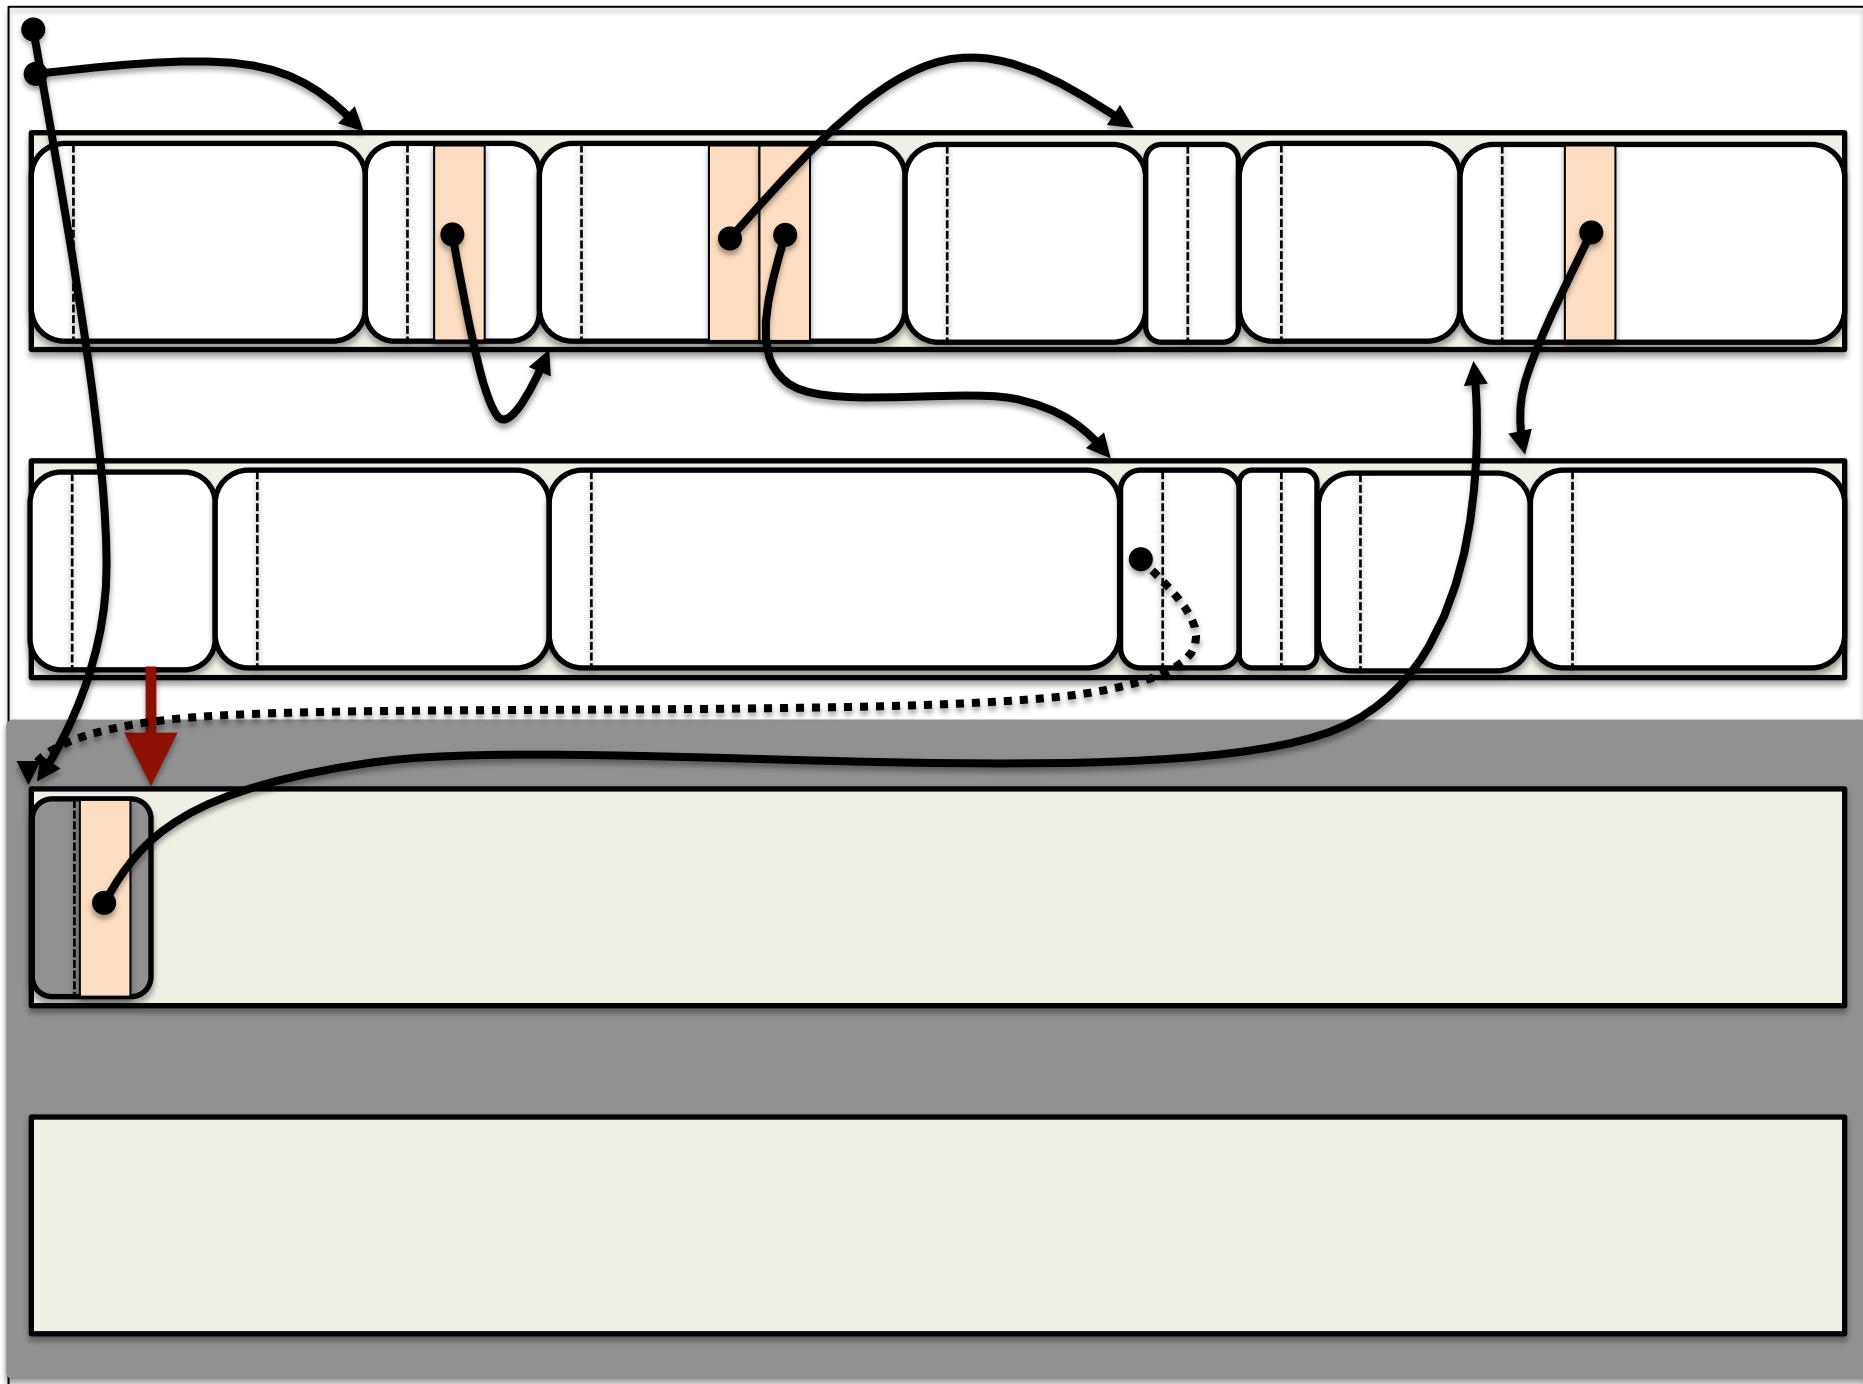


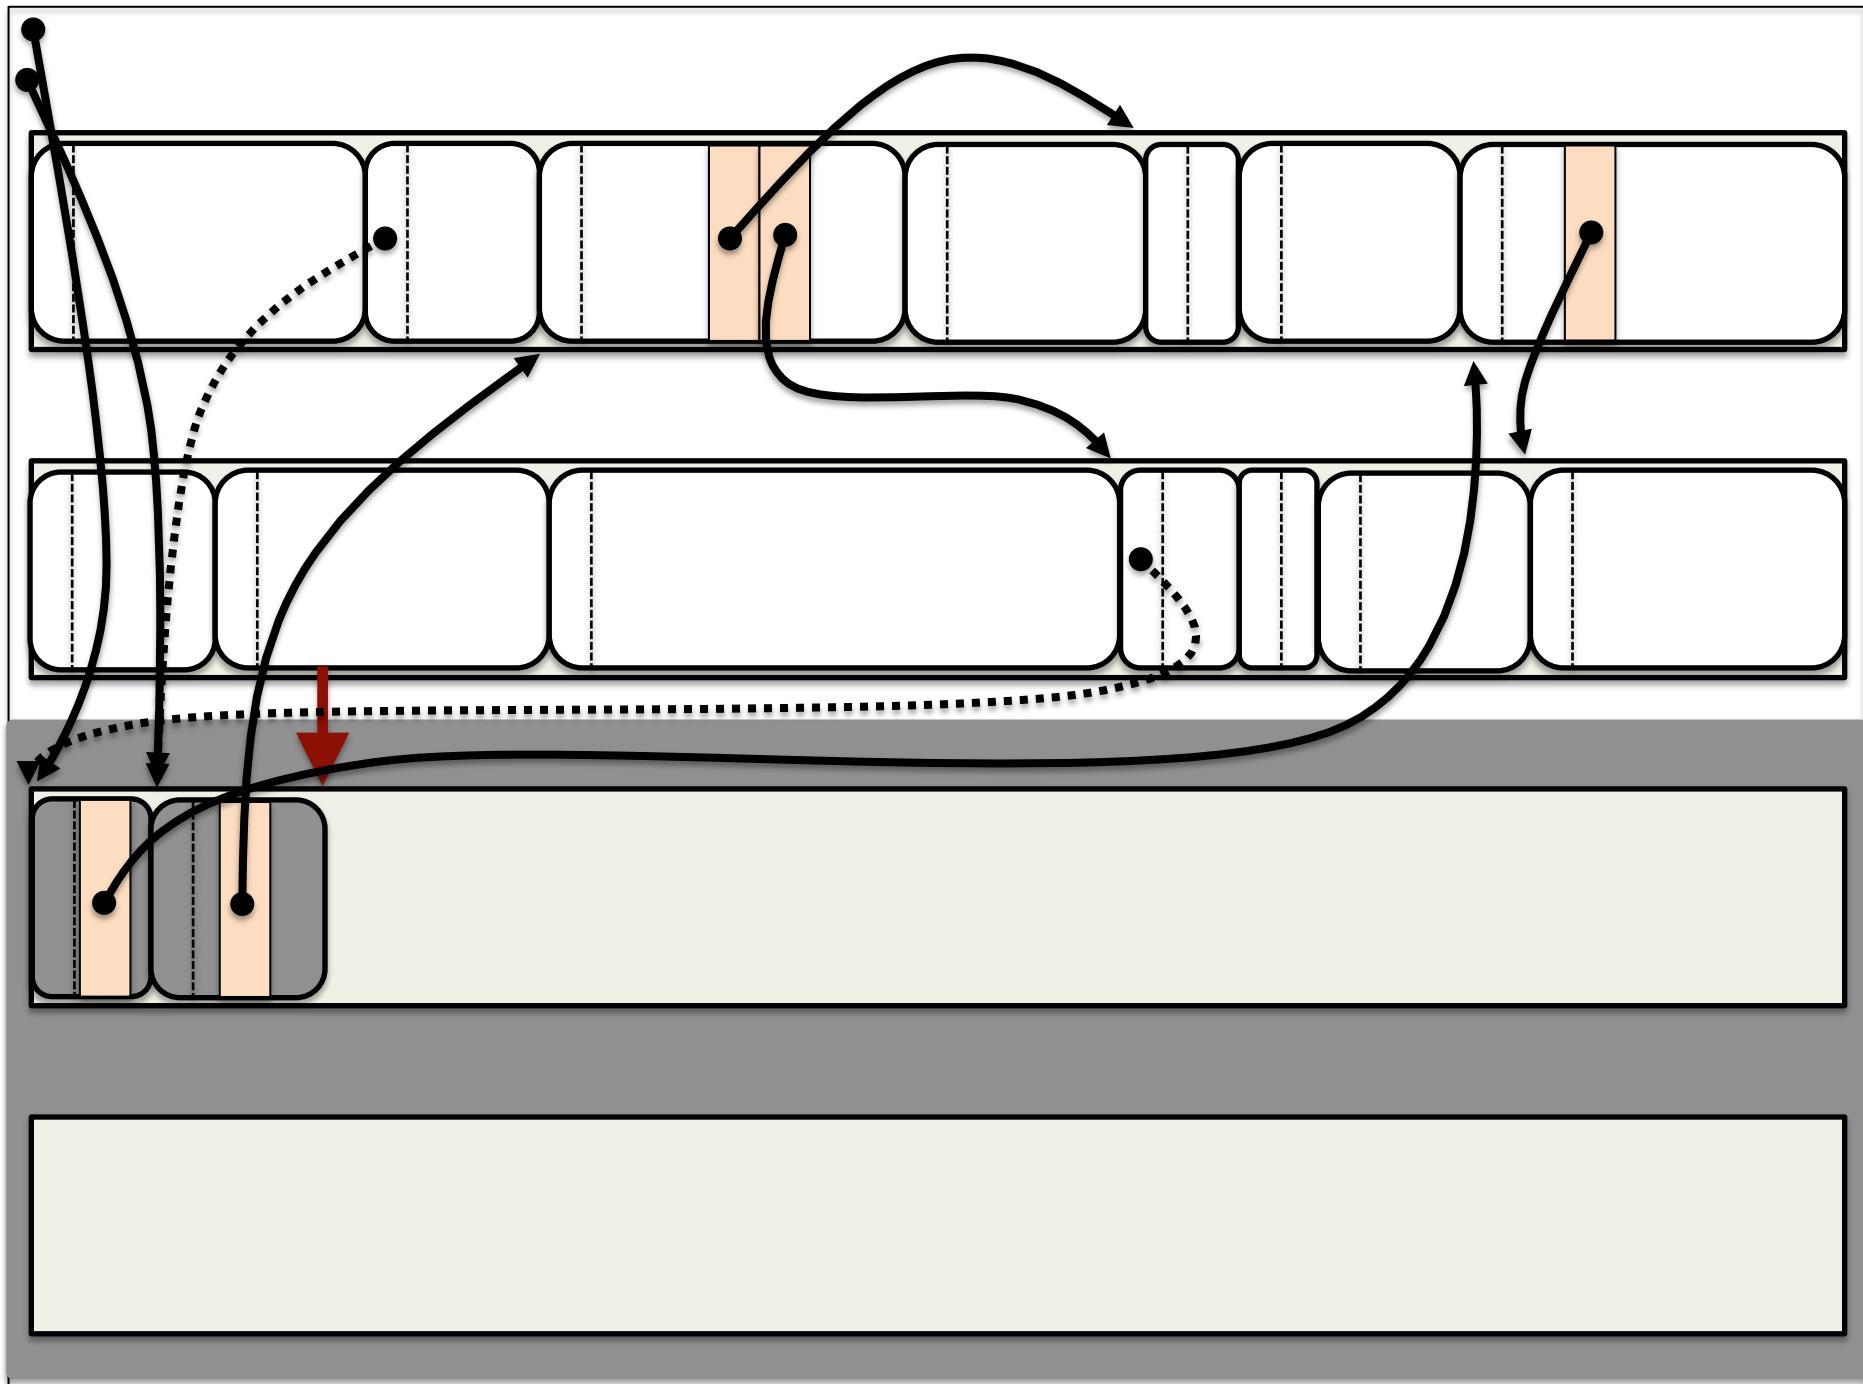






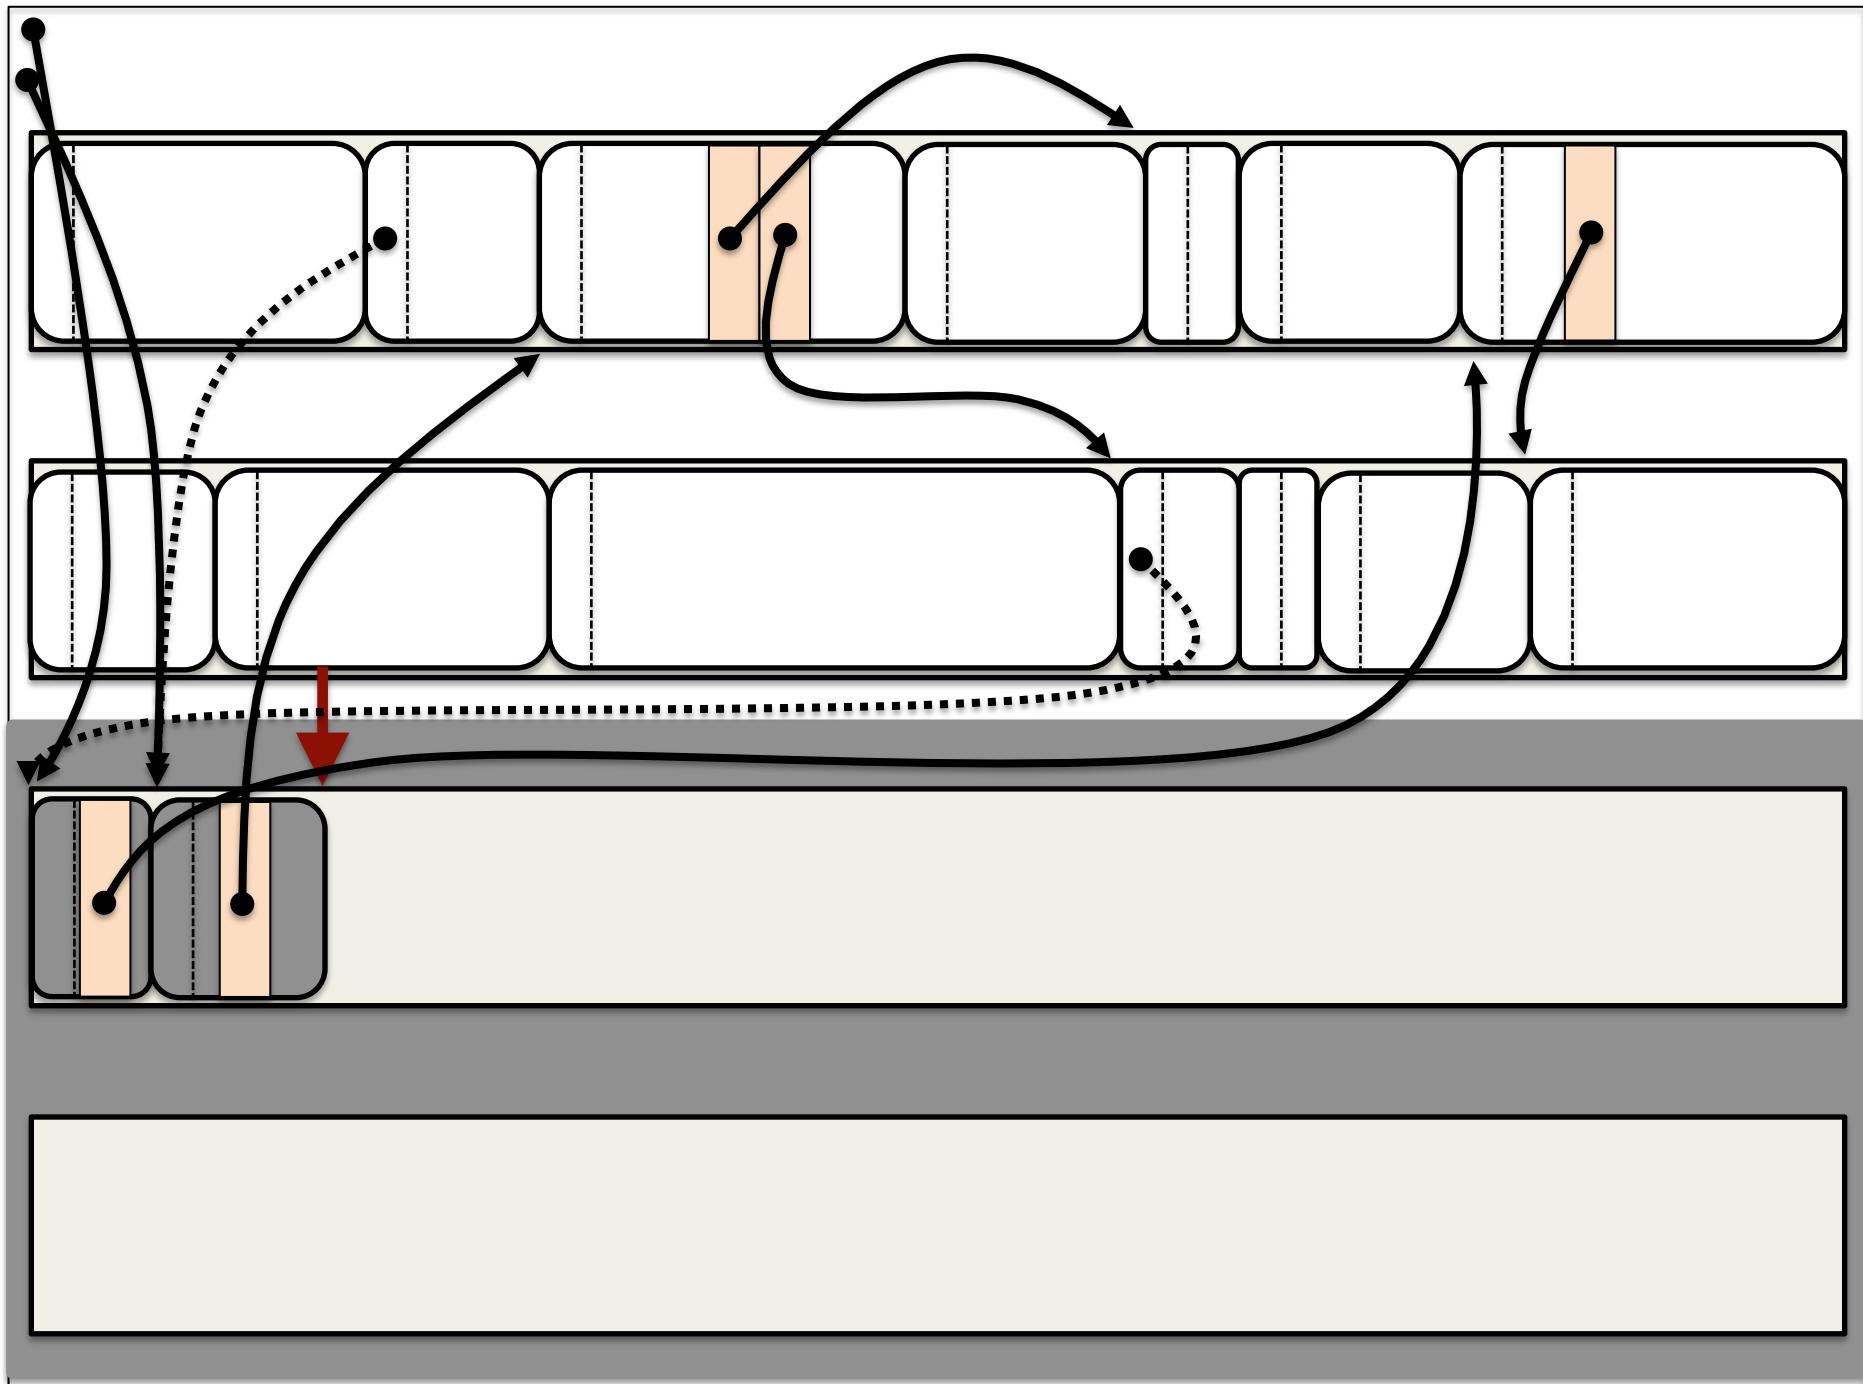


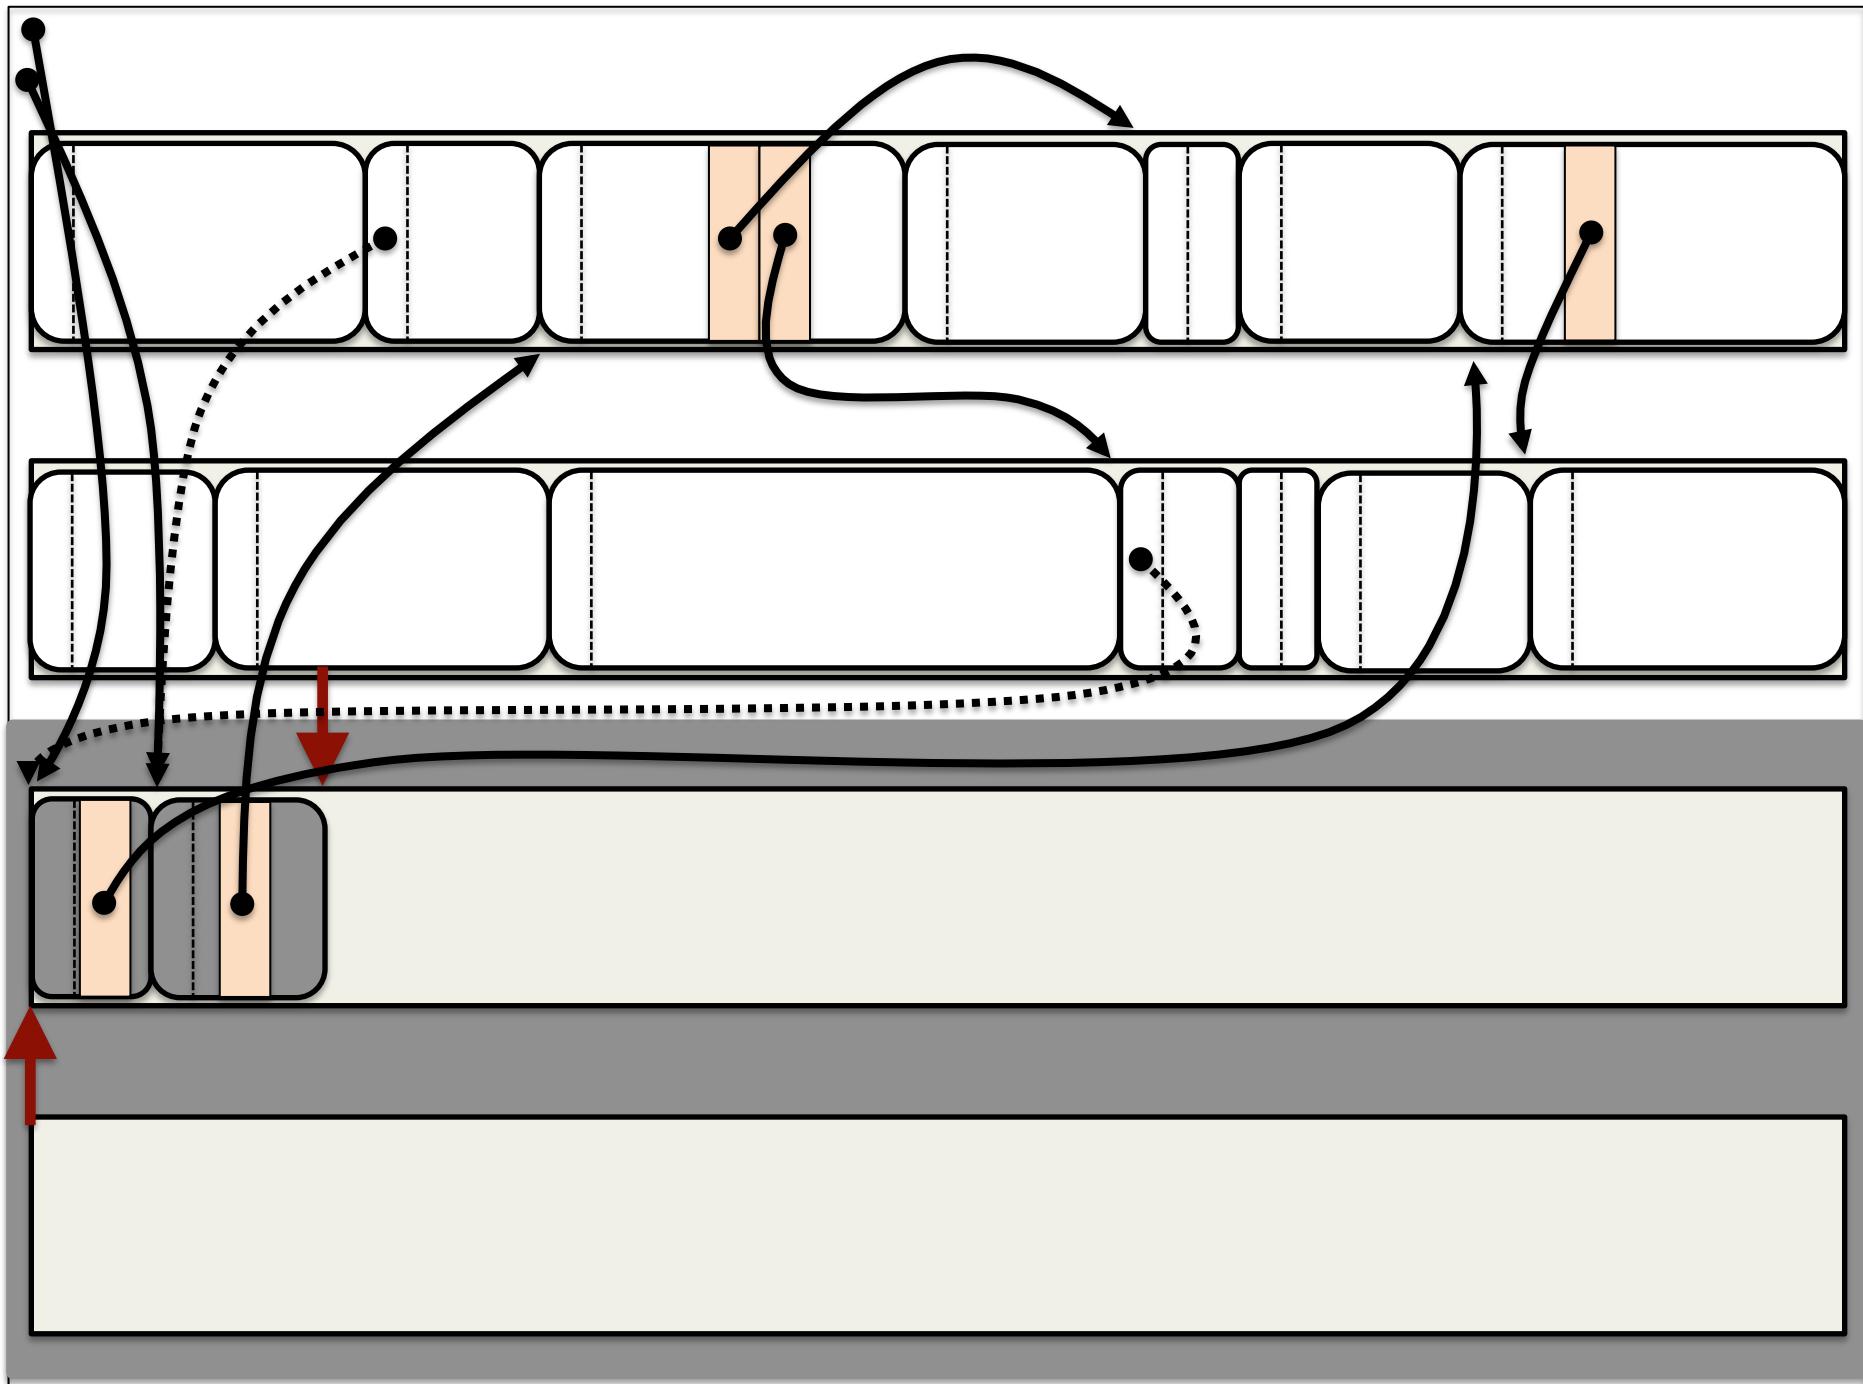


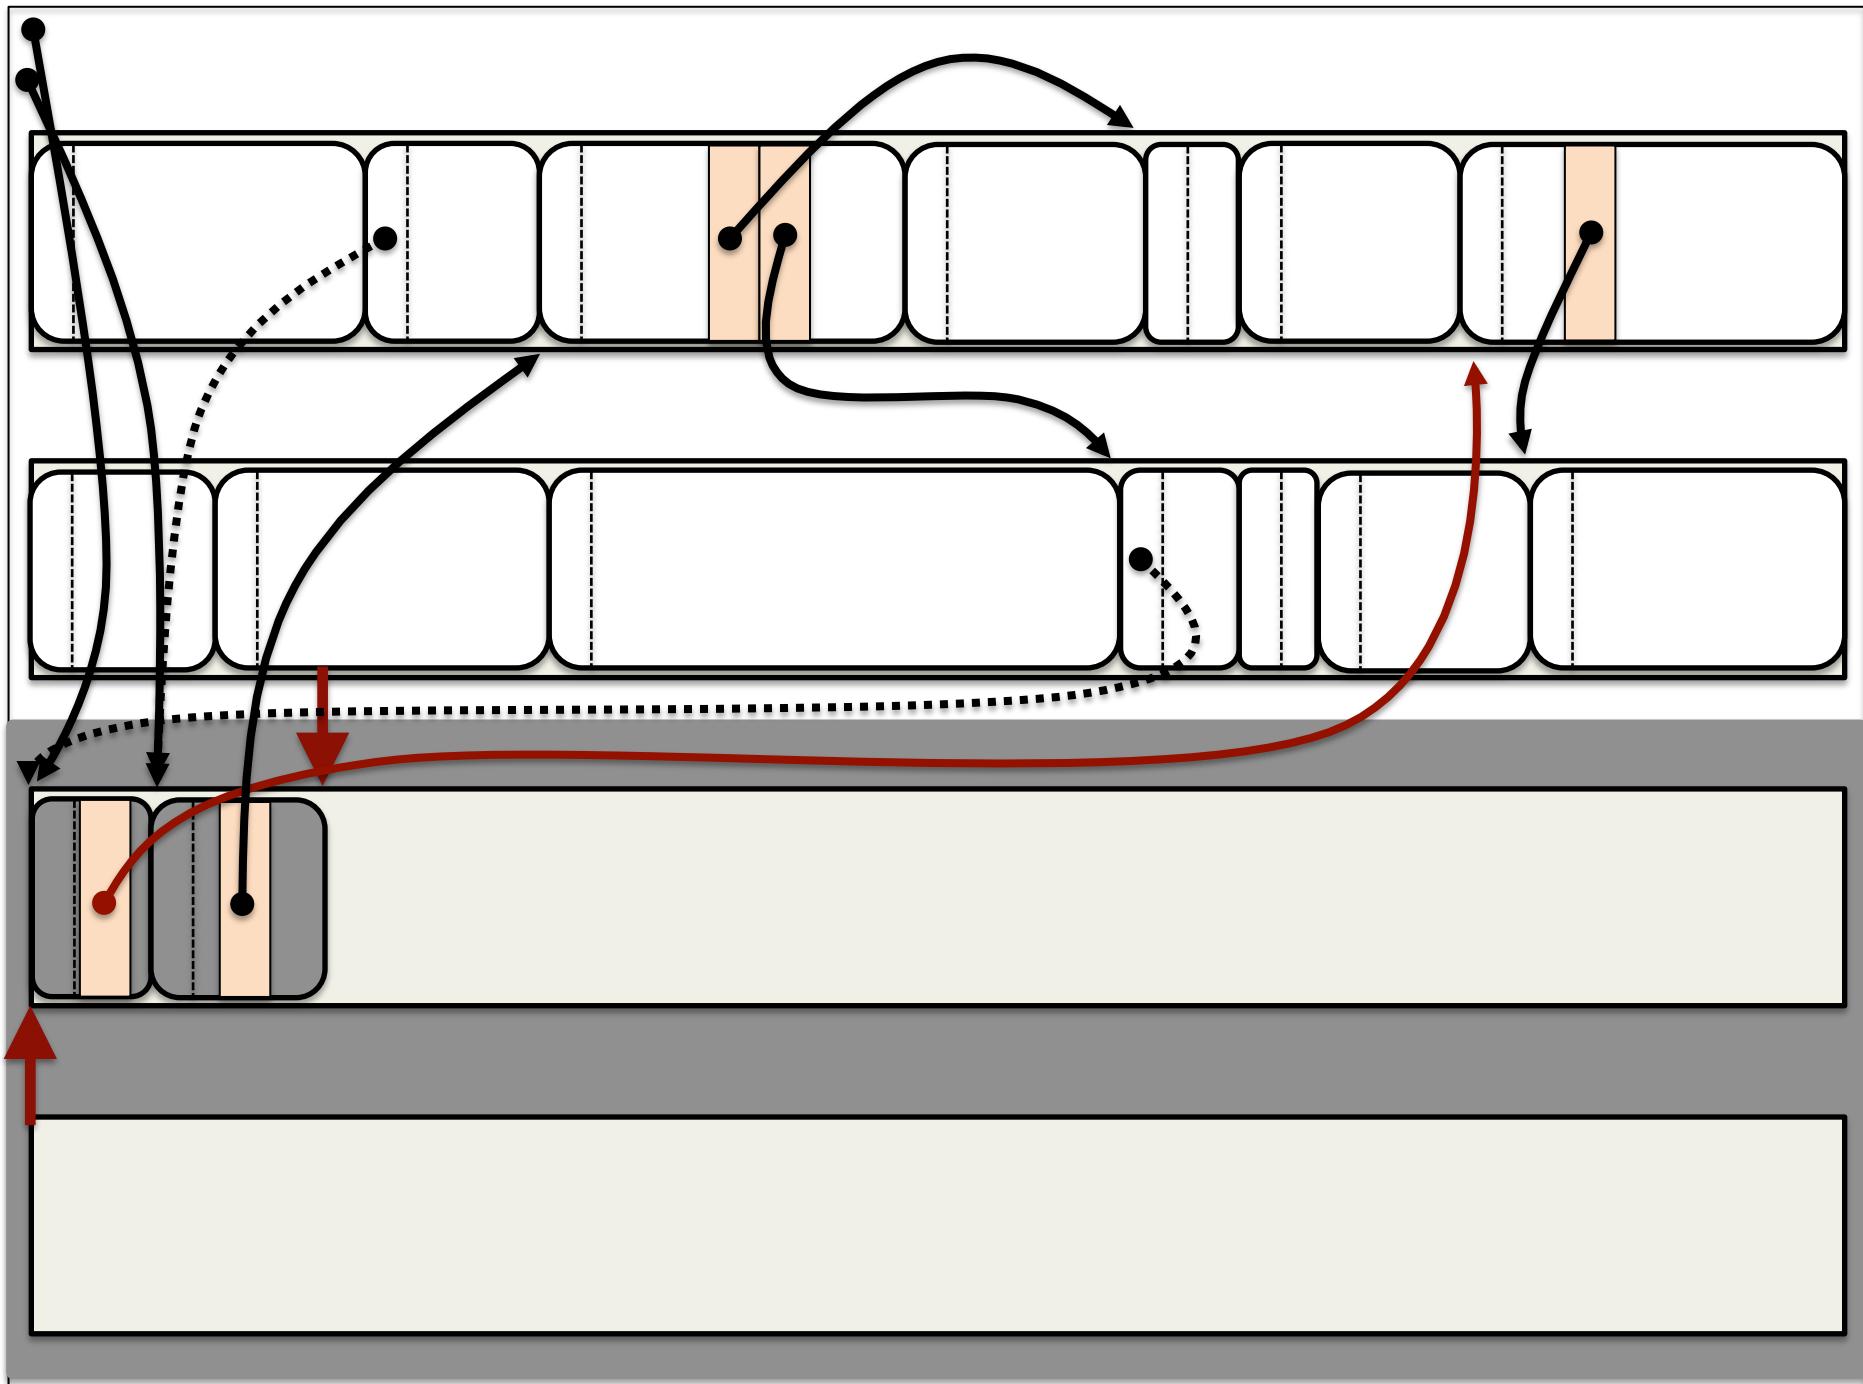


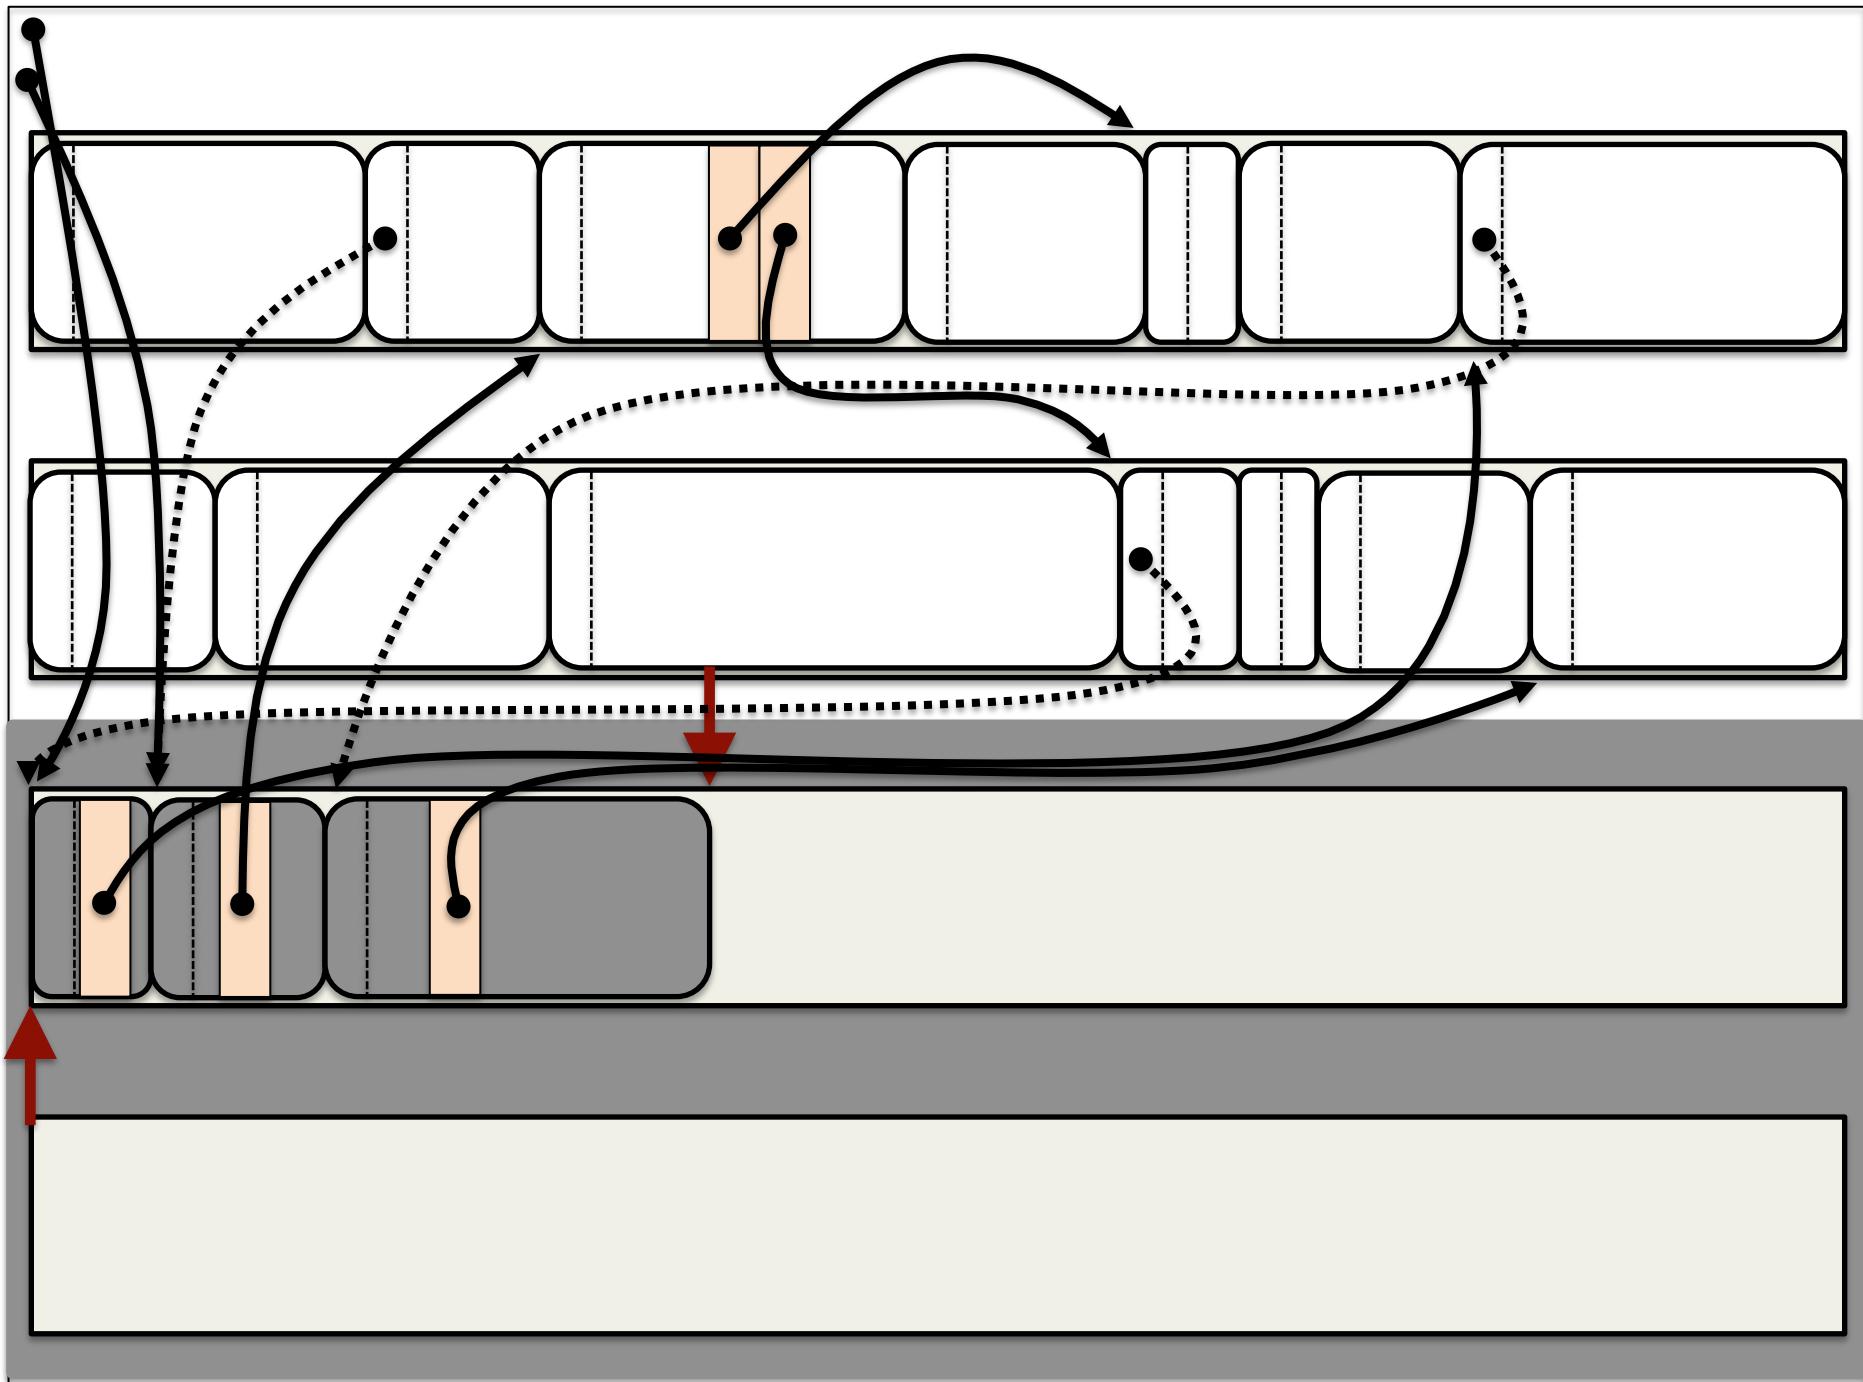
Two Finger Algorithm

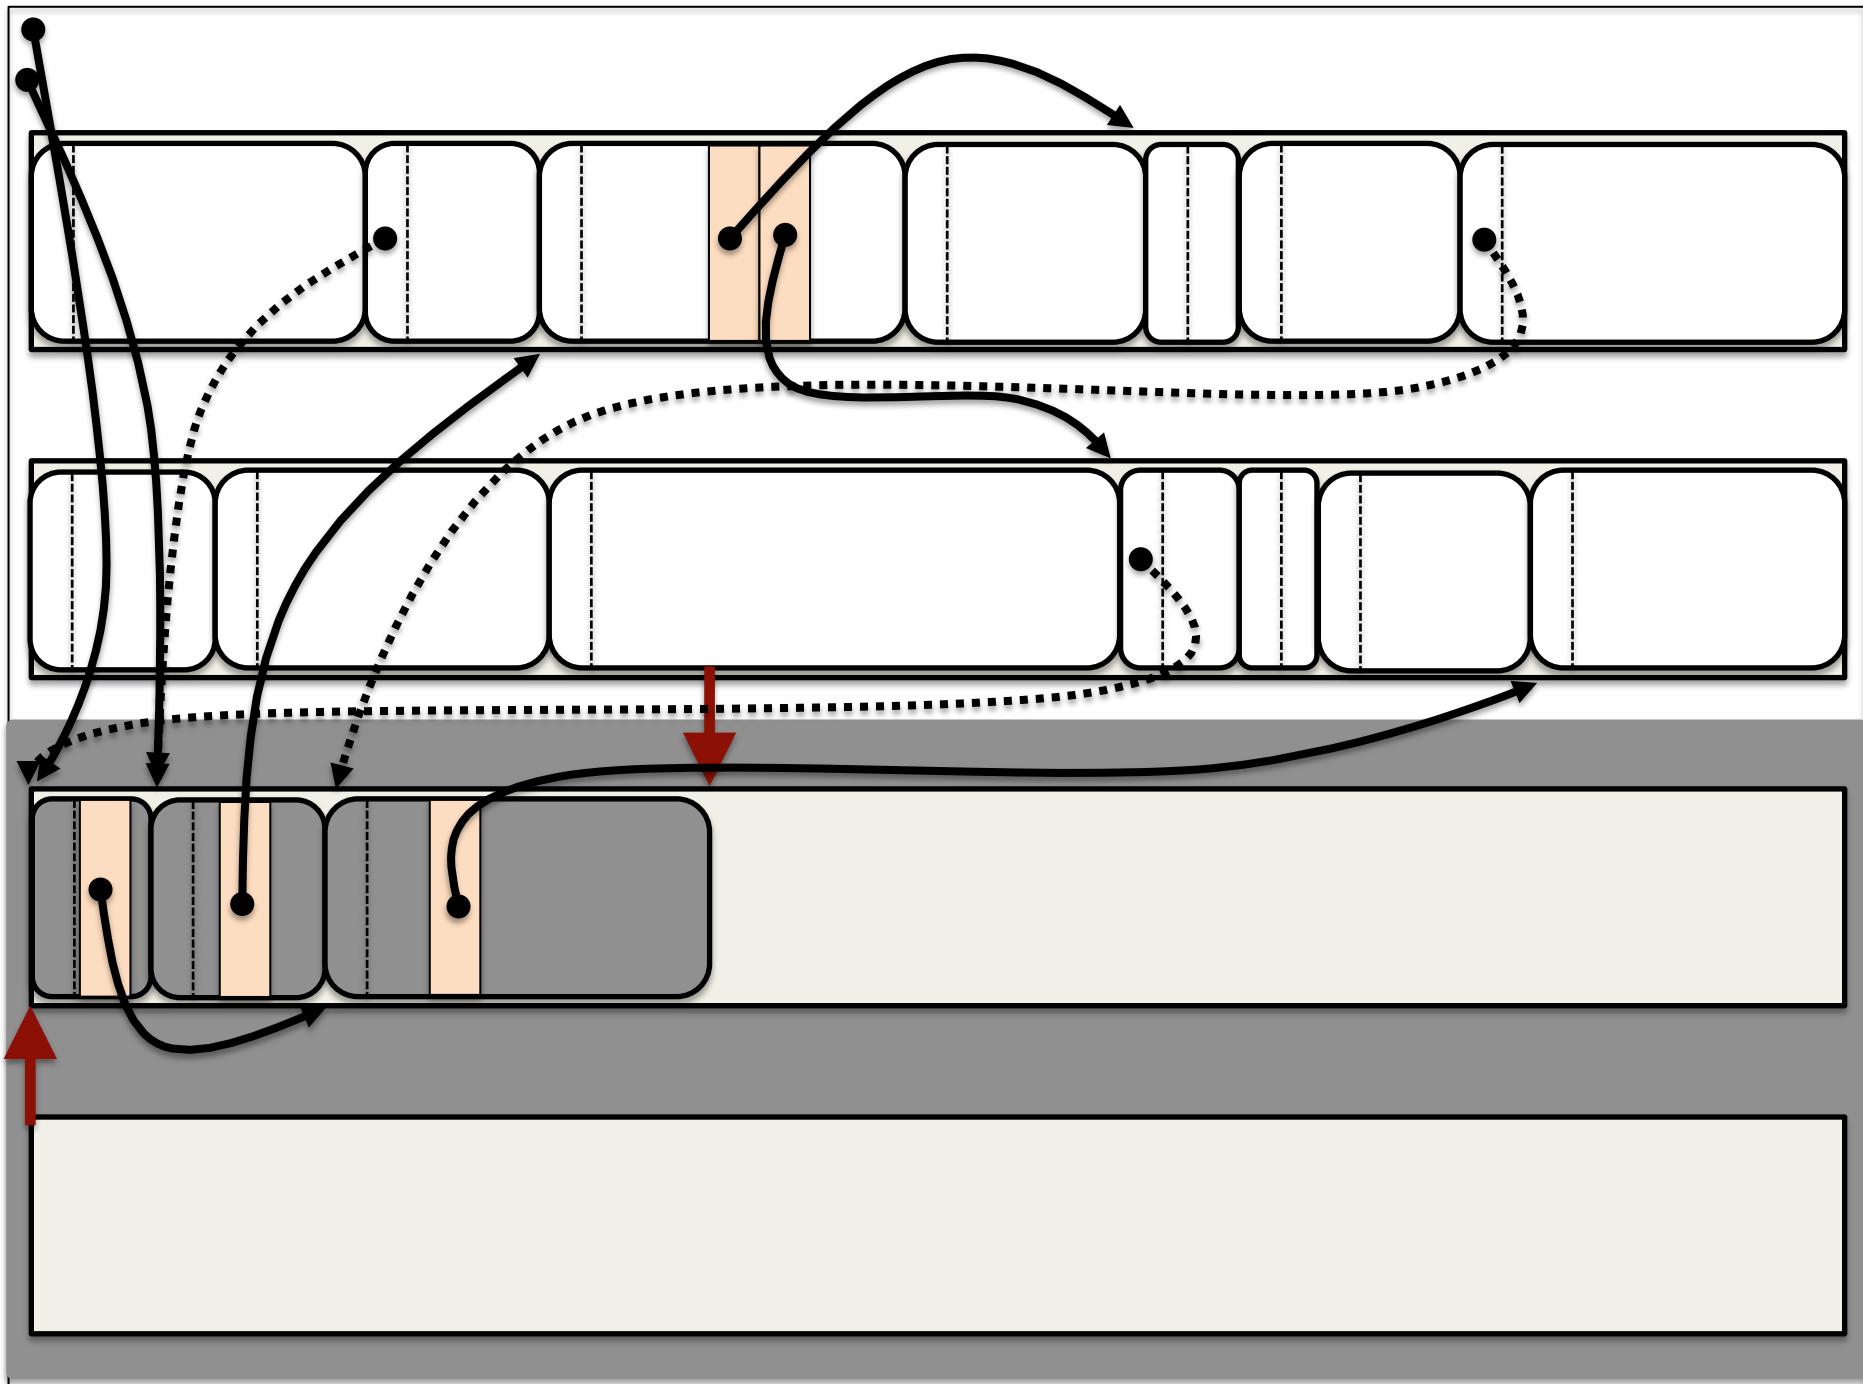
- Cheney's algorithm doesn't use a gray stack
- Objects are copied immediately
- We can use the new space as an implicit stack
 - Note that objects are allocated sequentially
 - Scan objects in the same order
- We're done when the two pointers meet

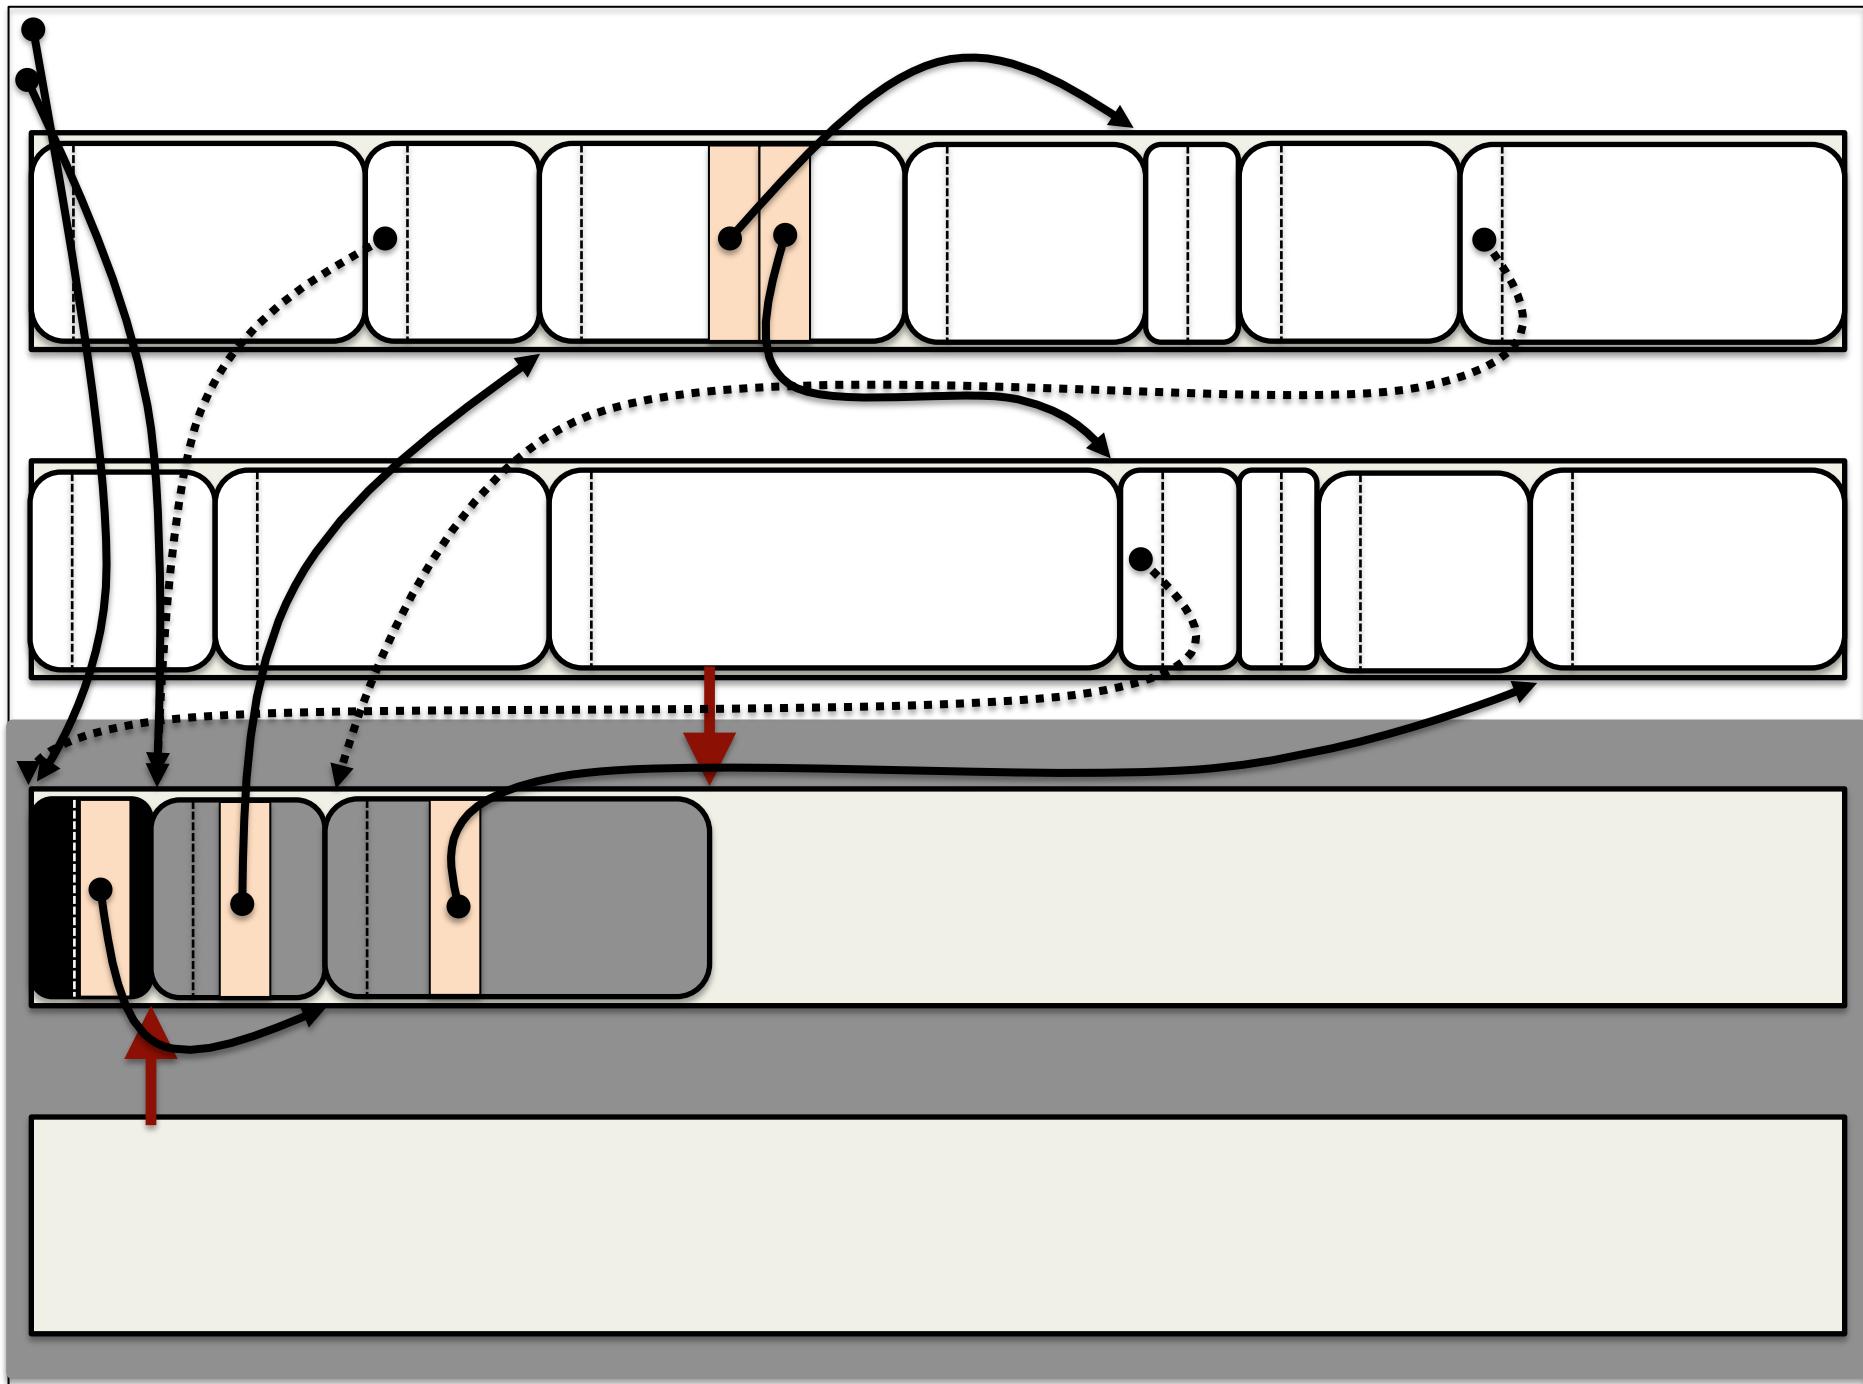


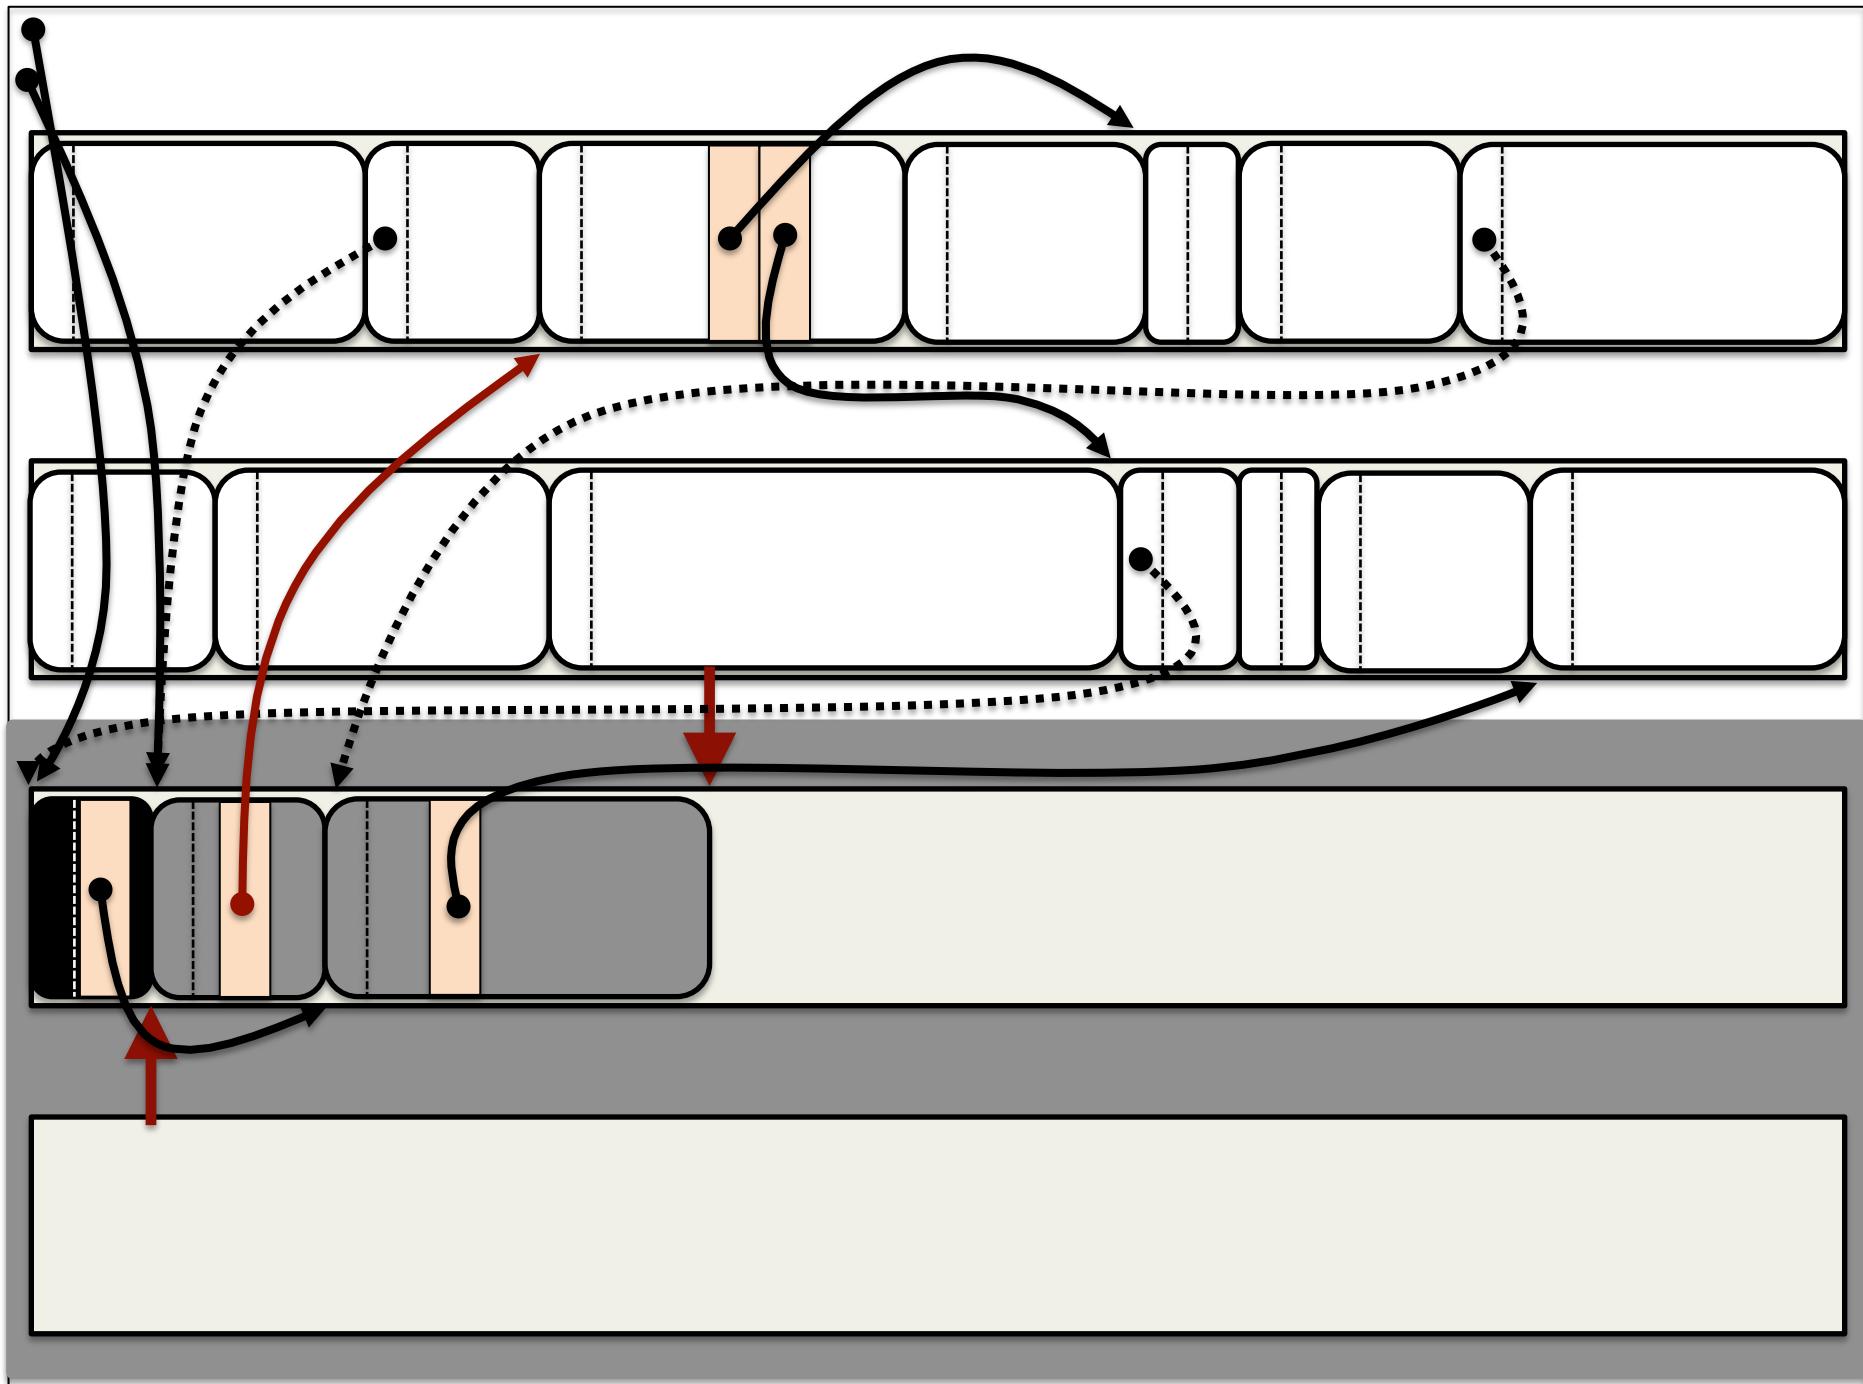


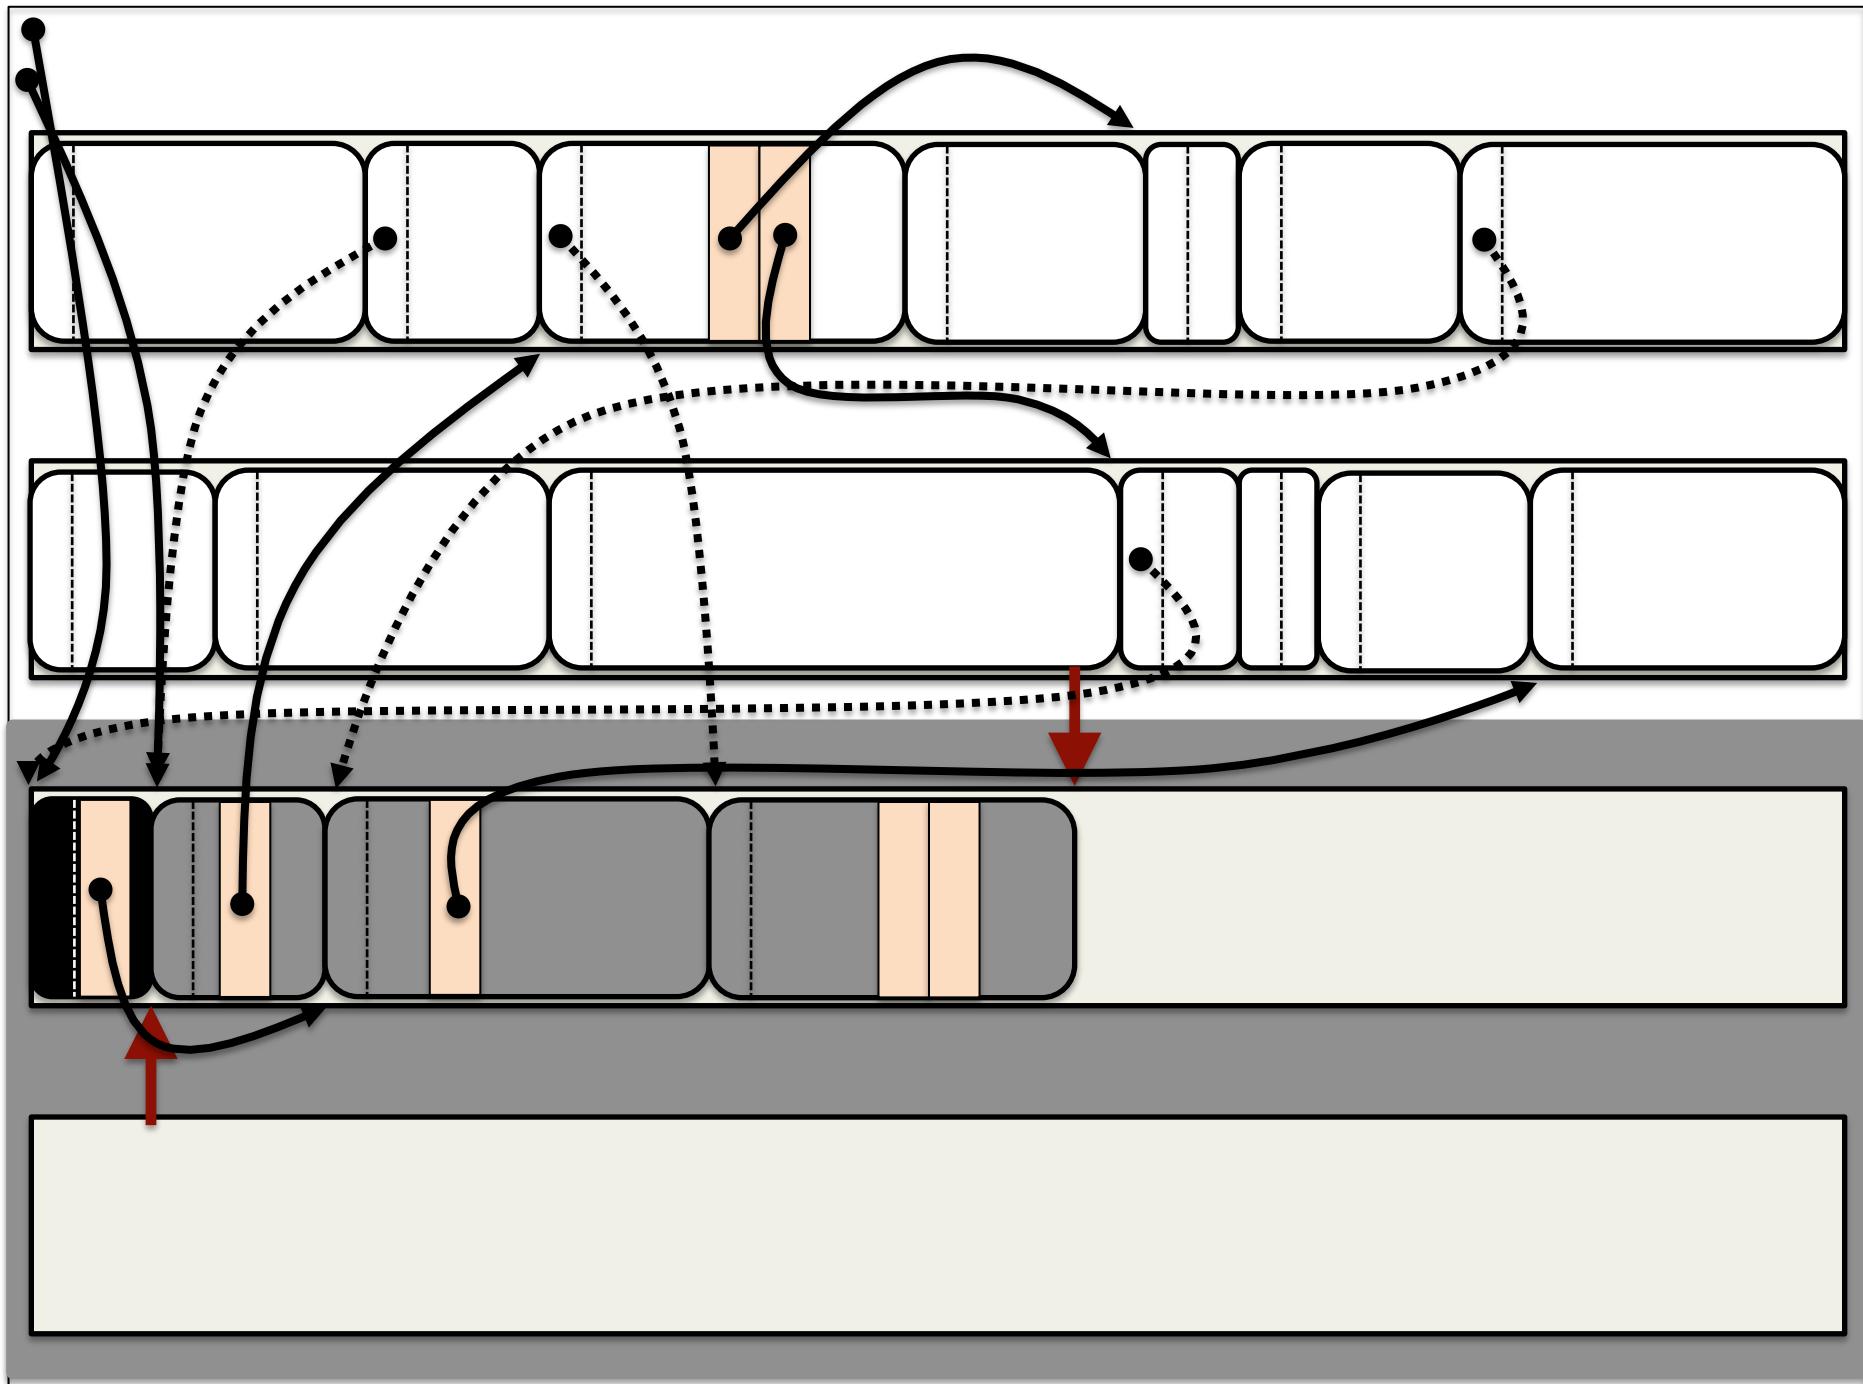


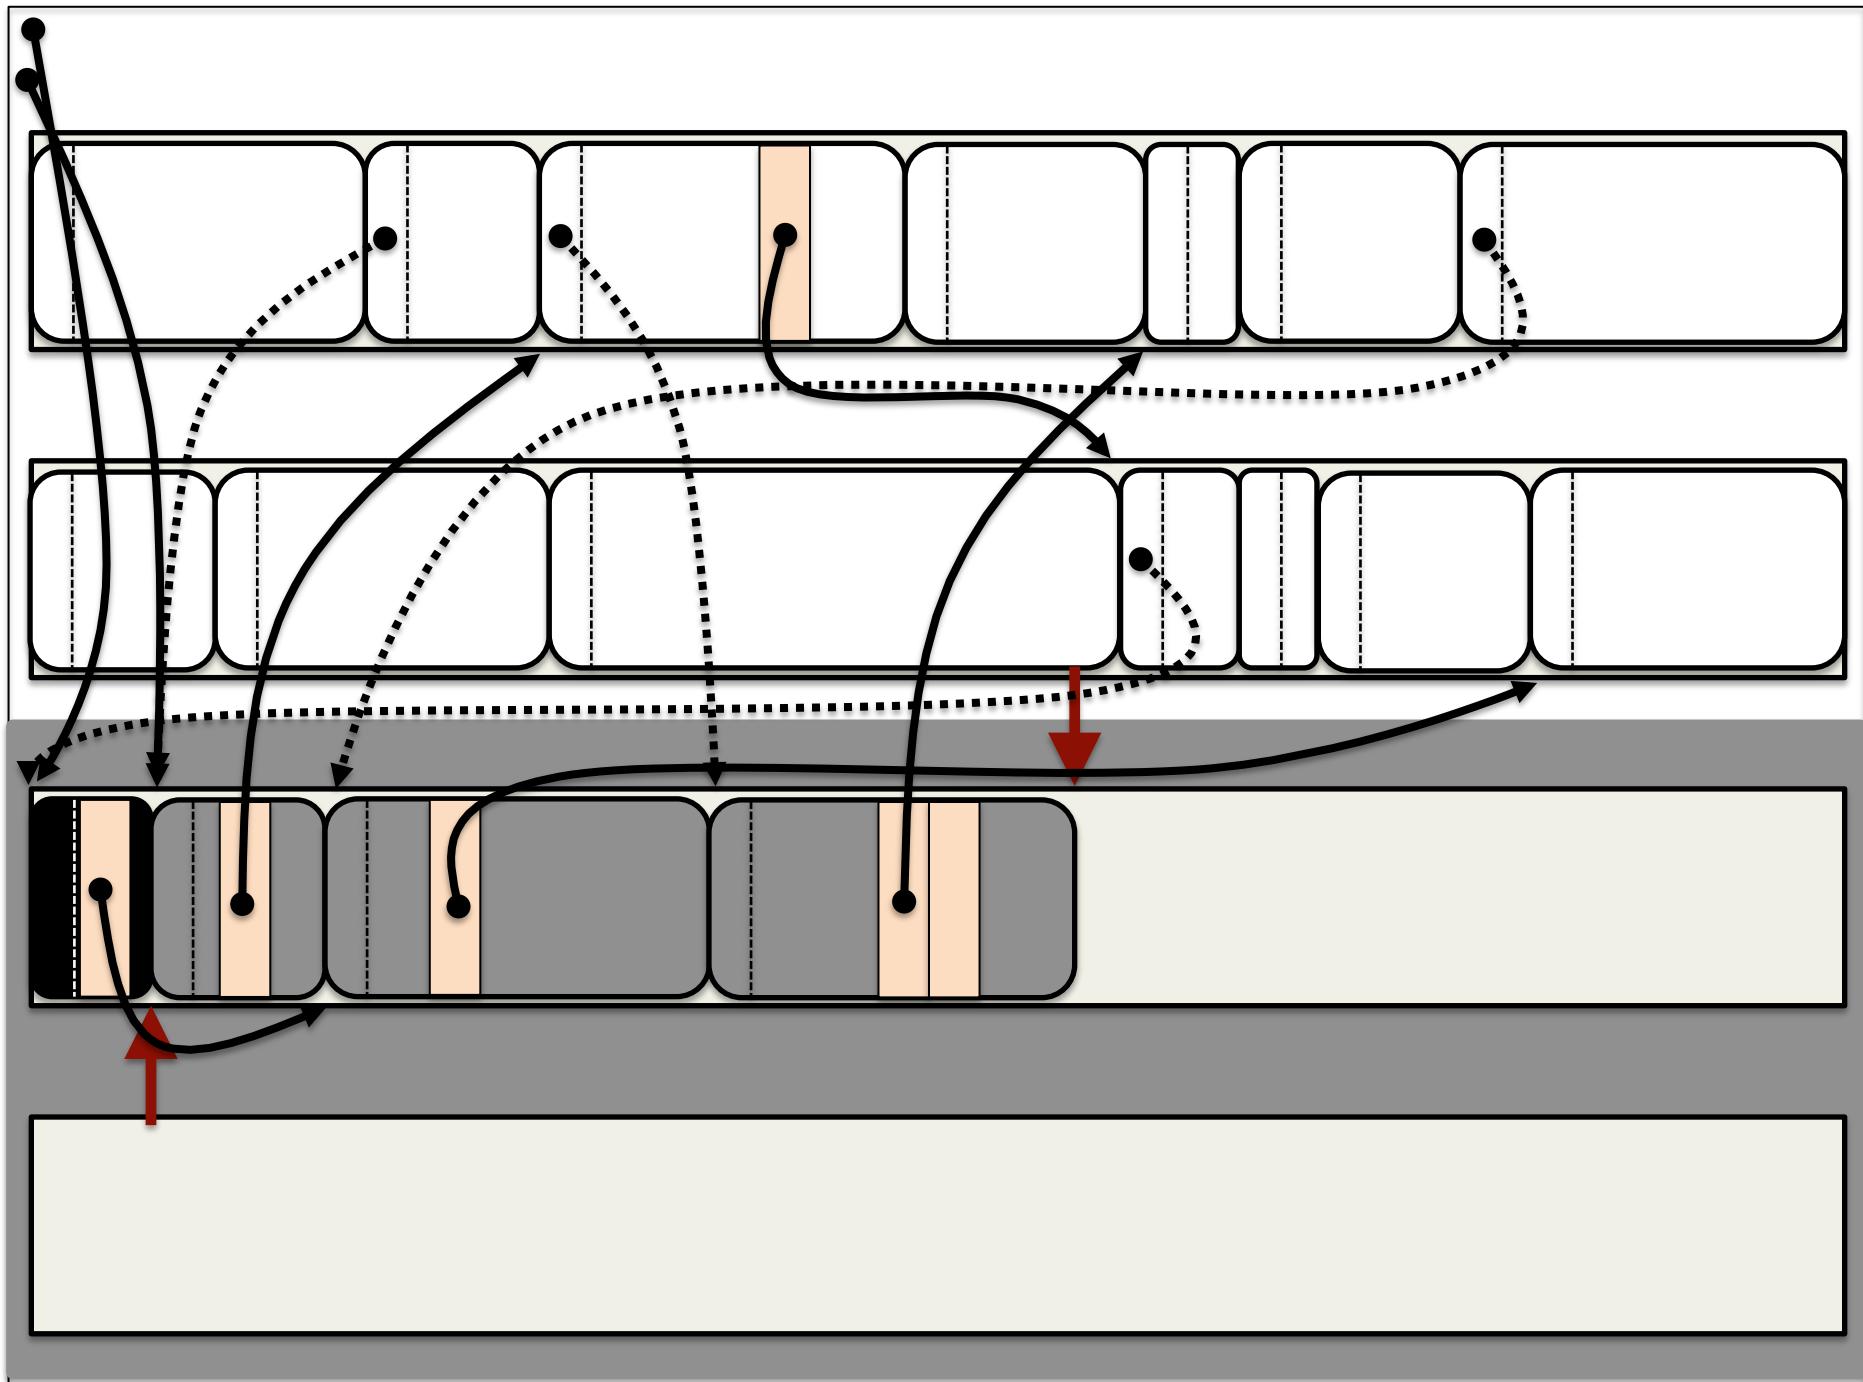


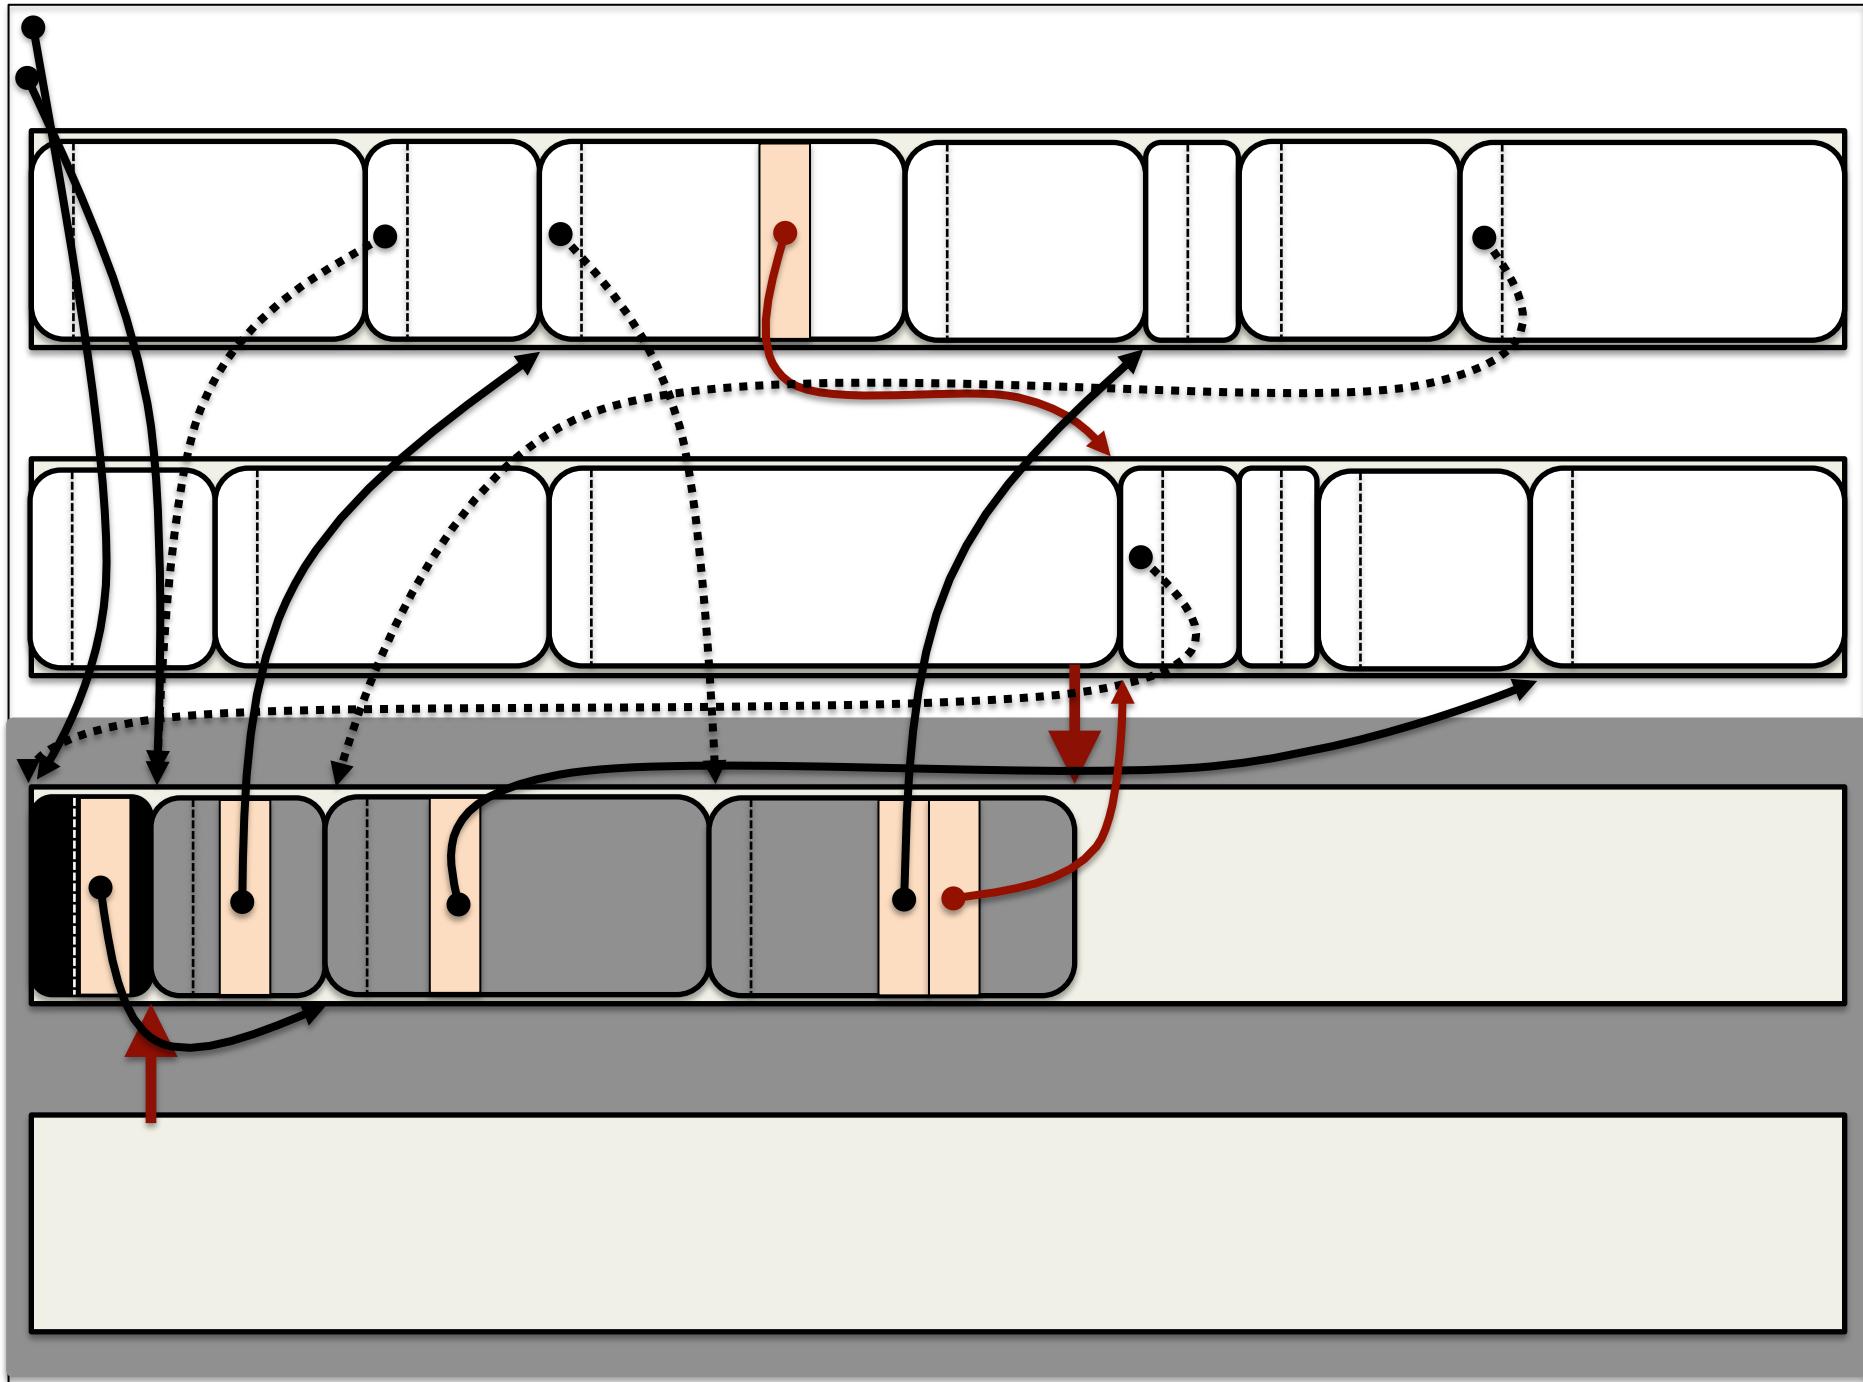


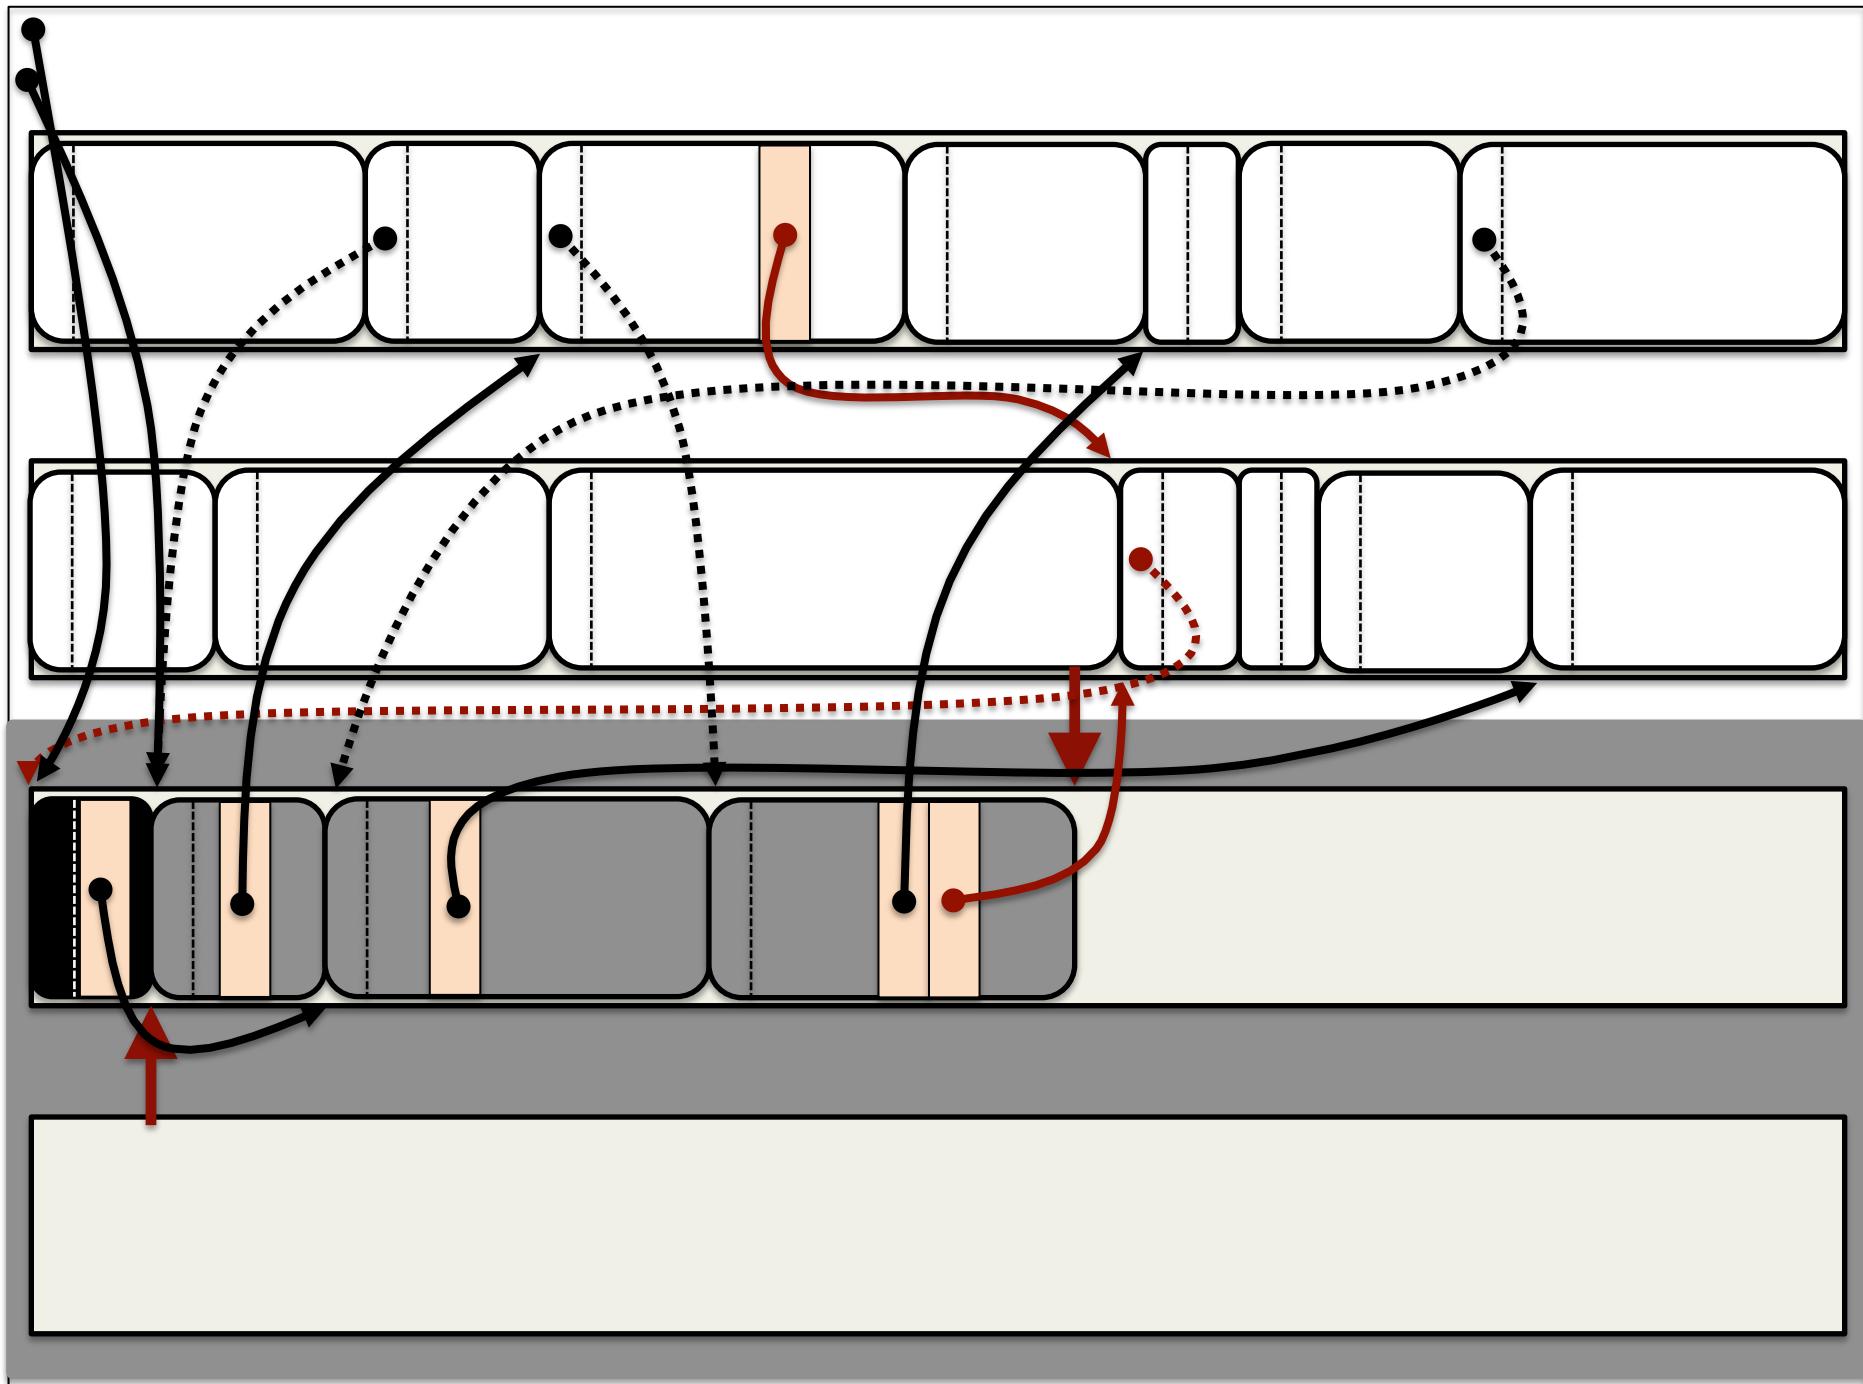


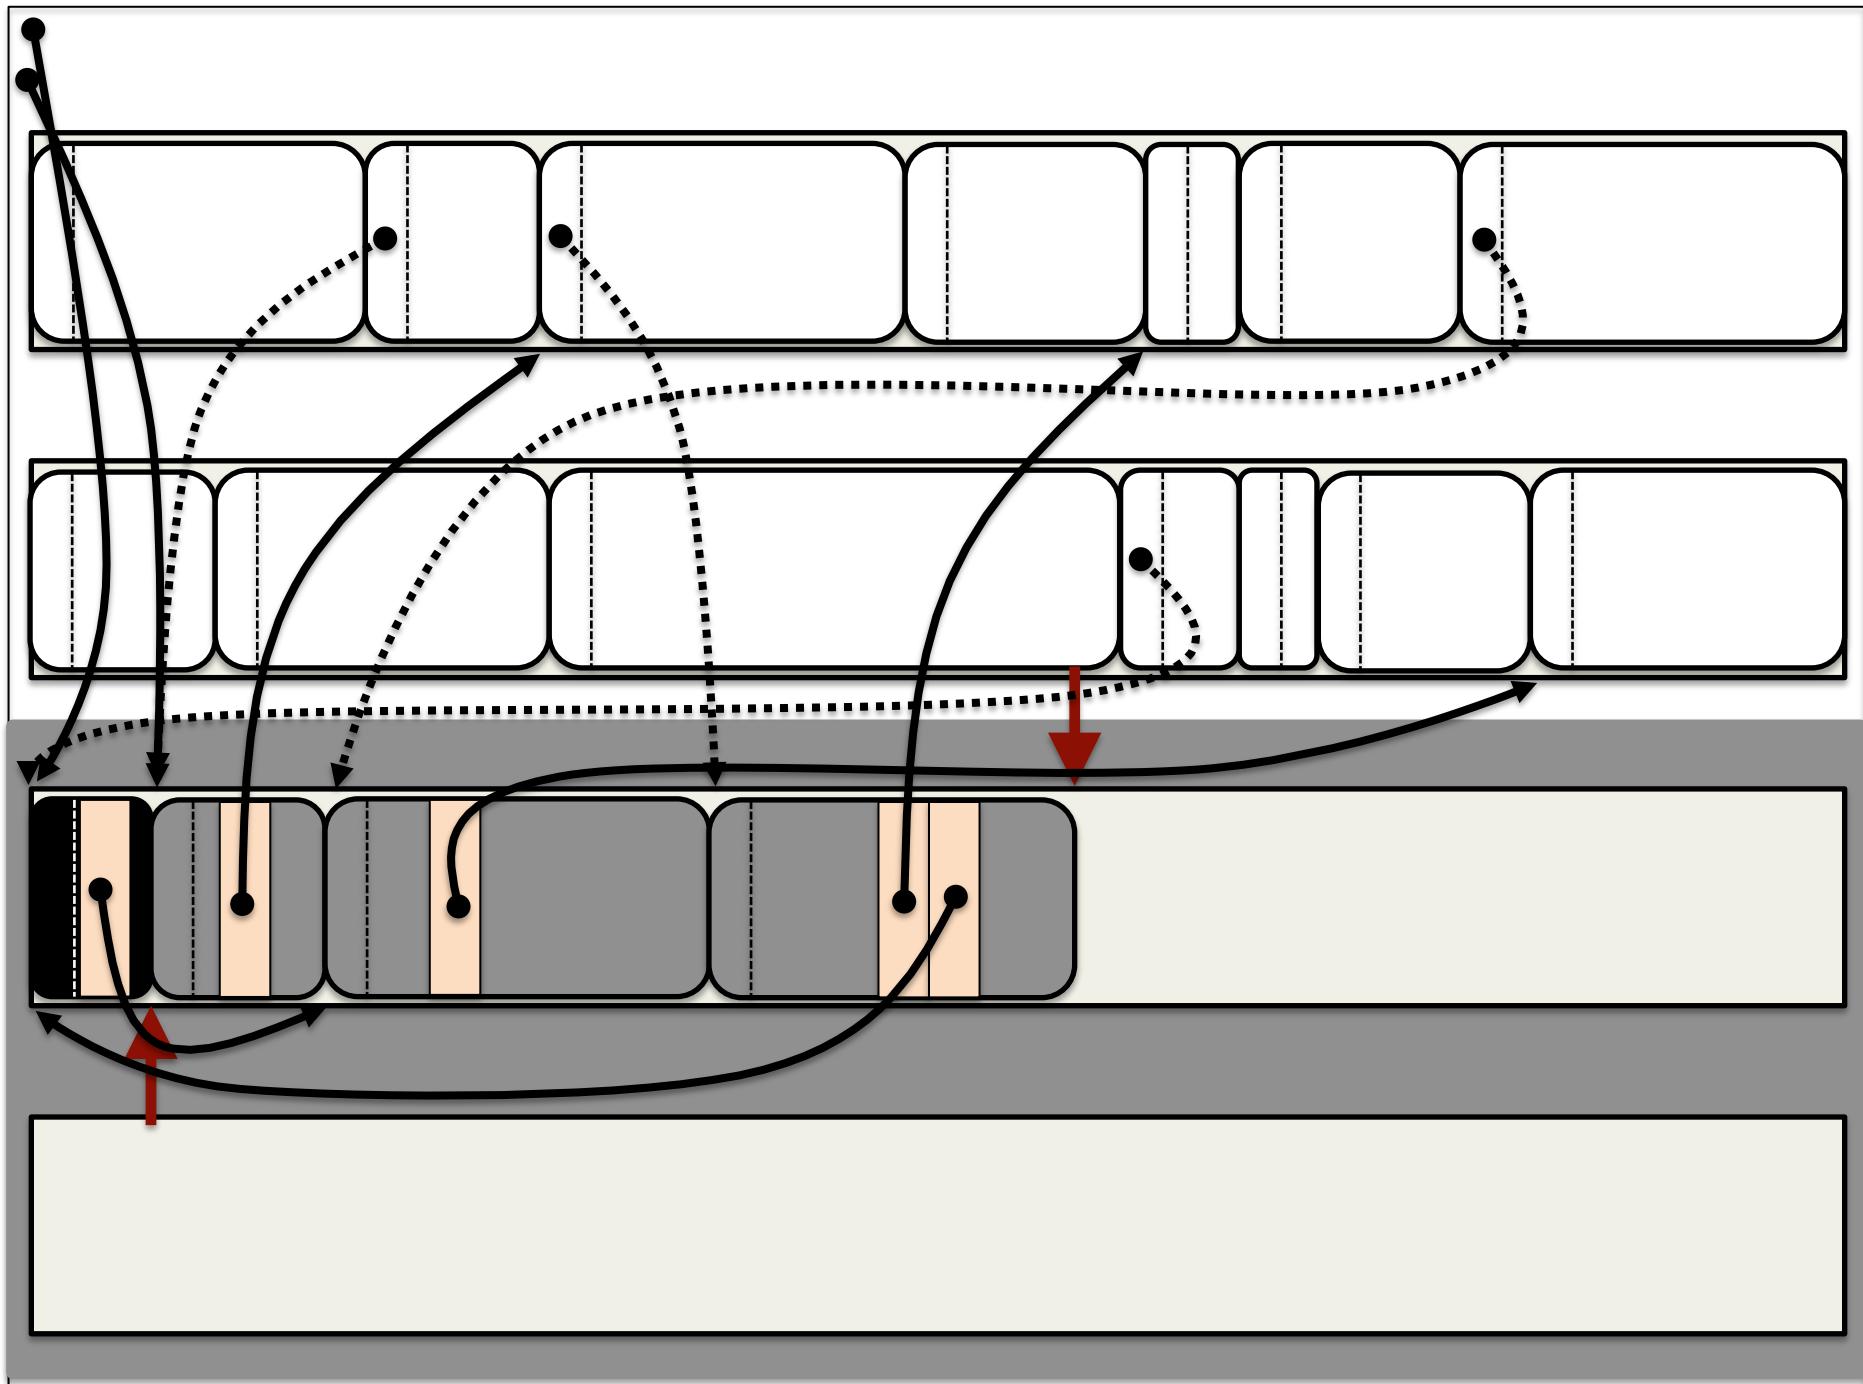


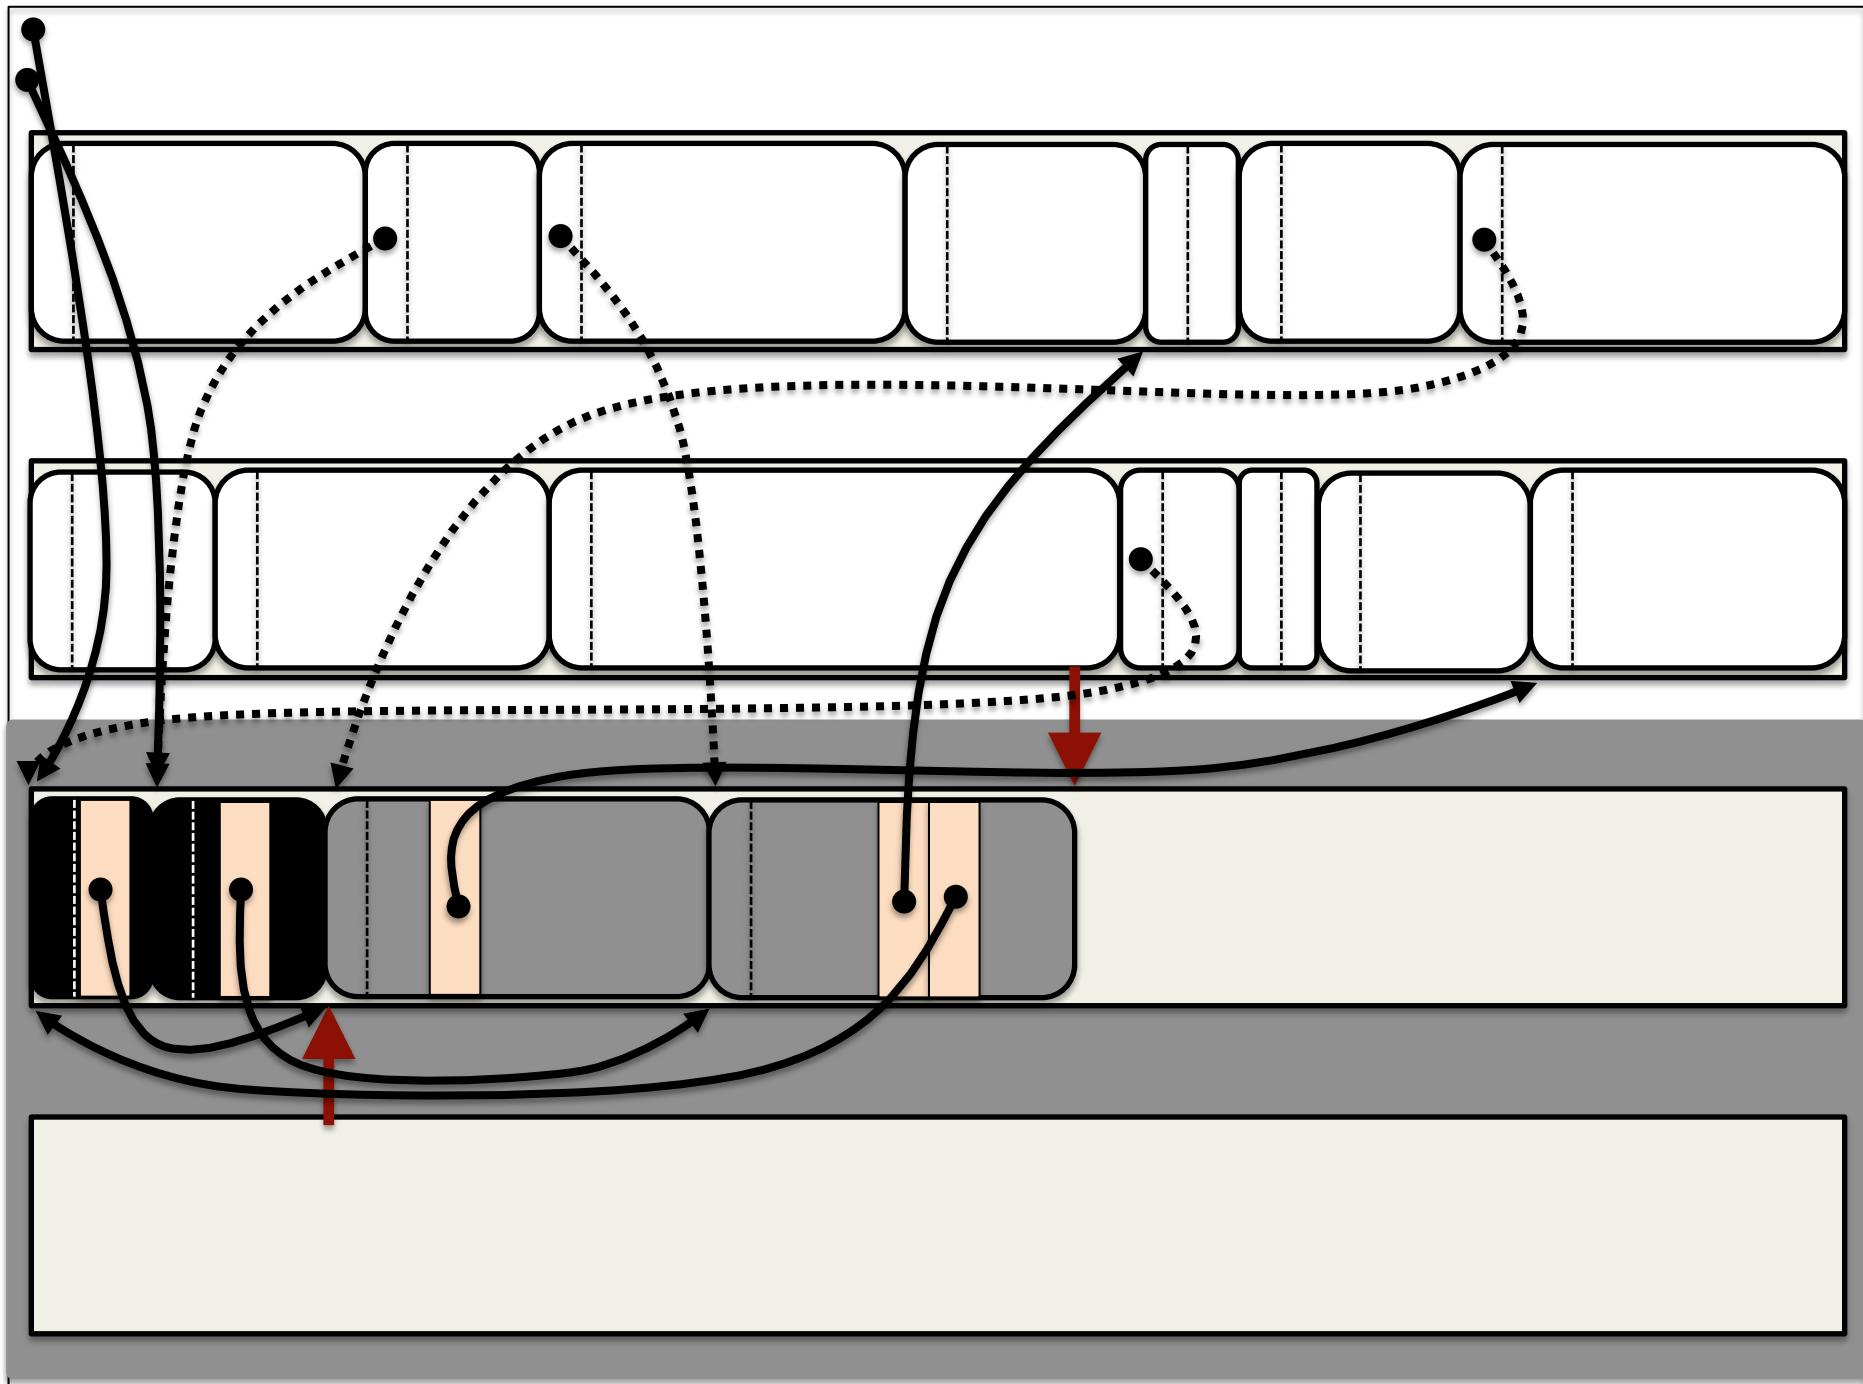


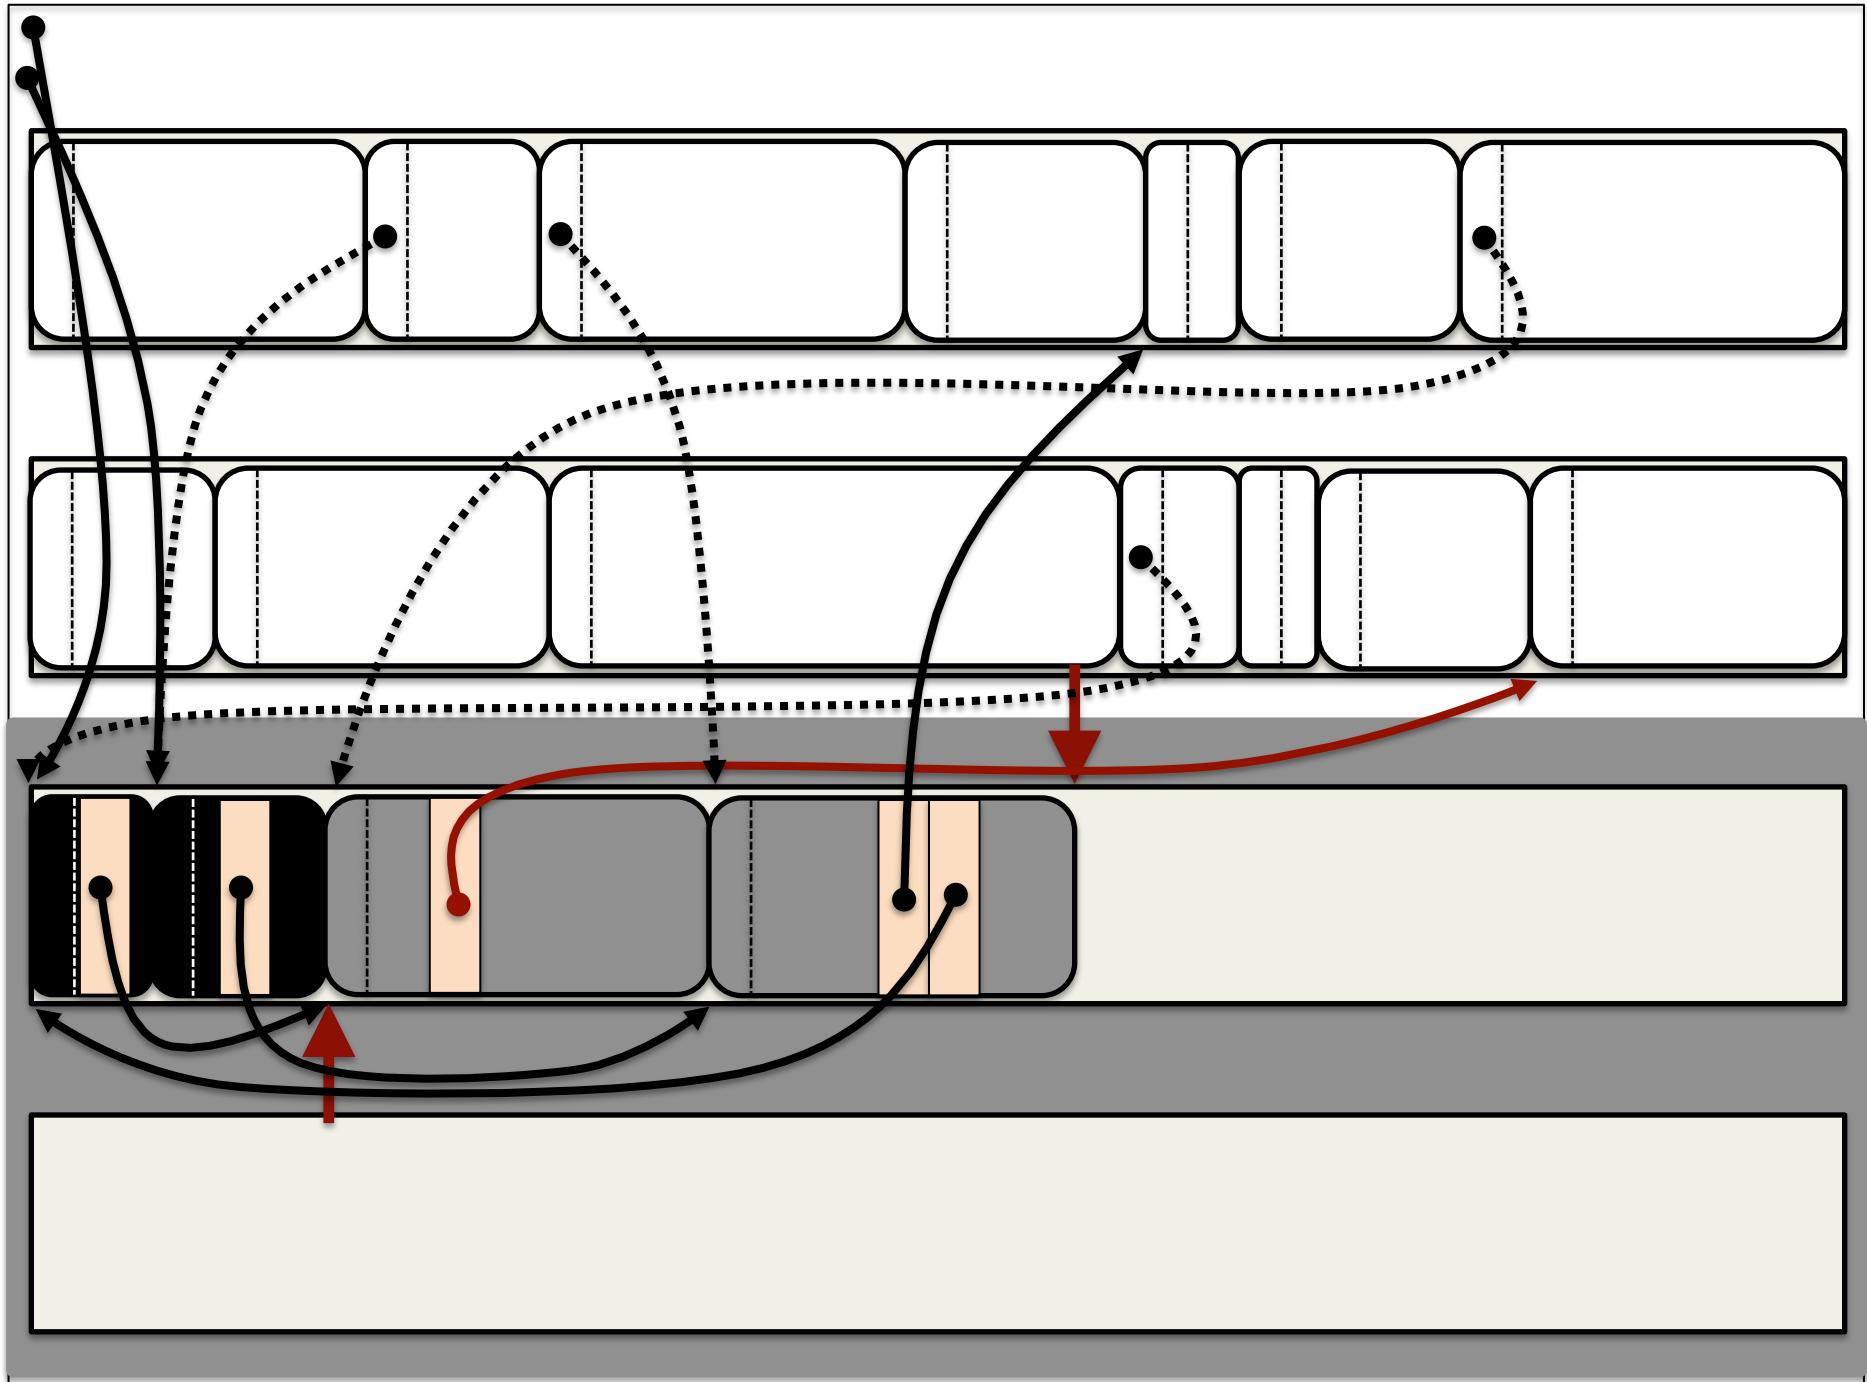


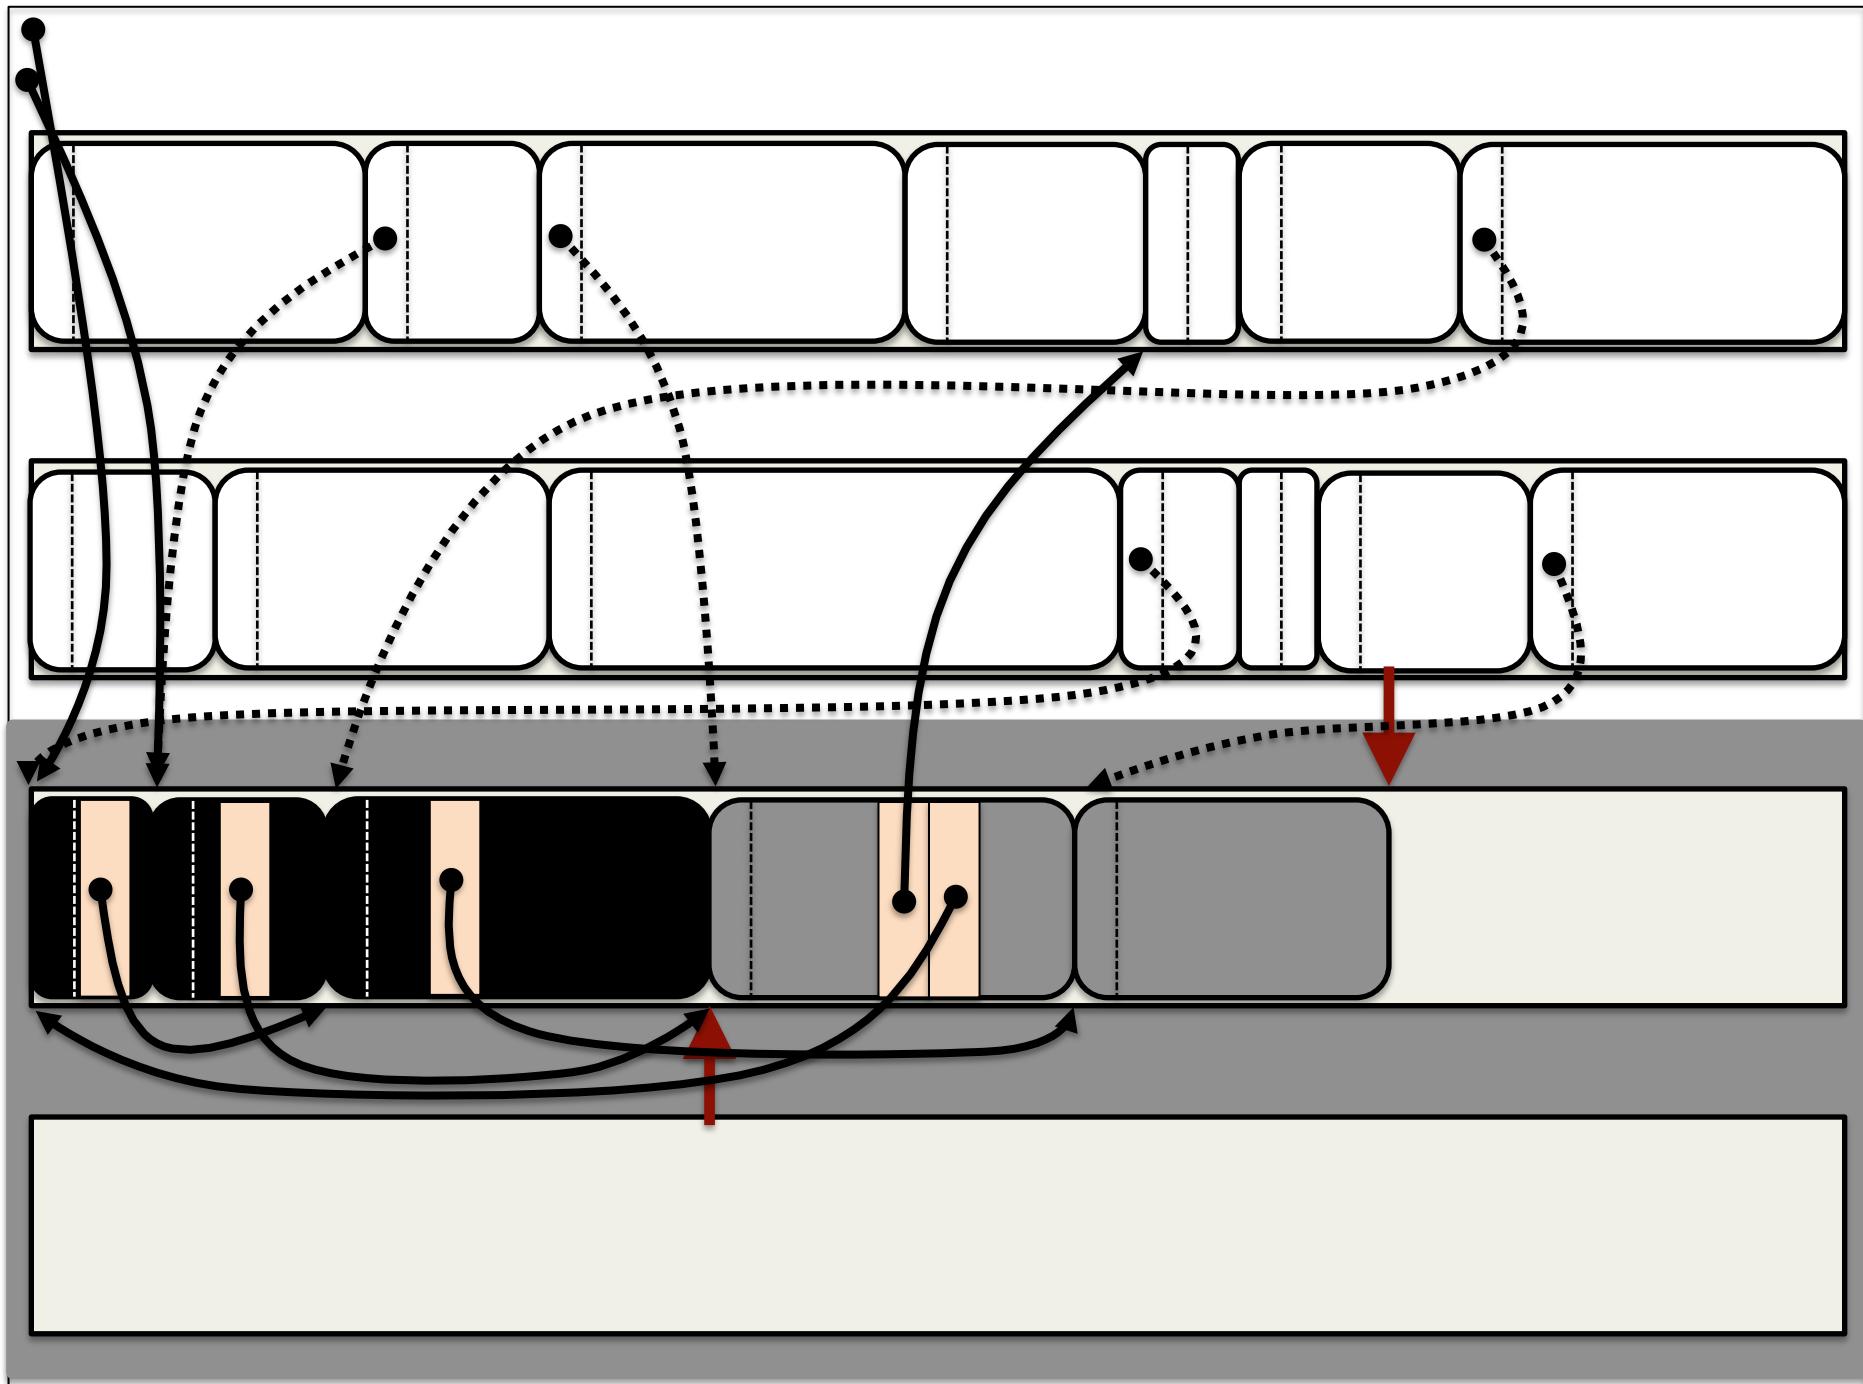






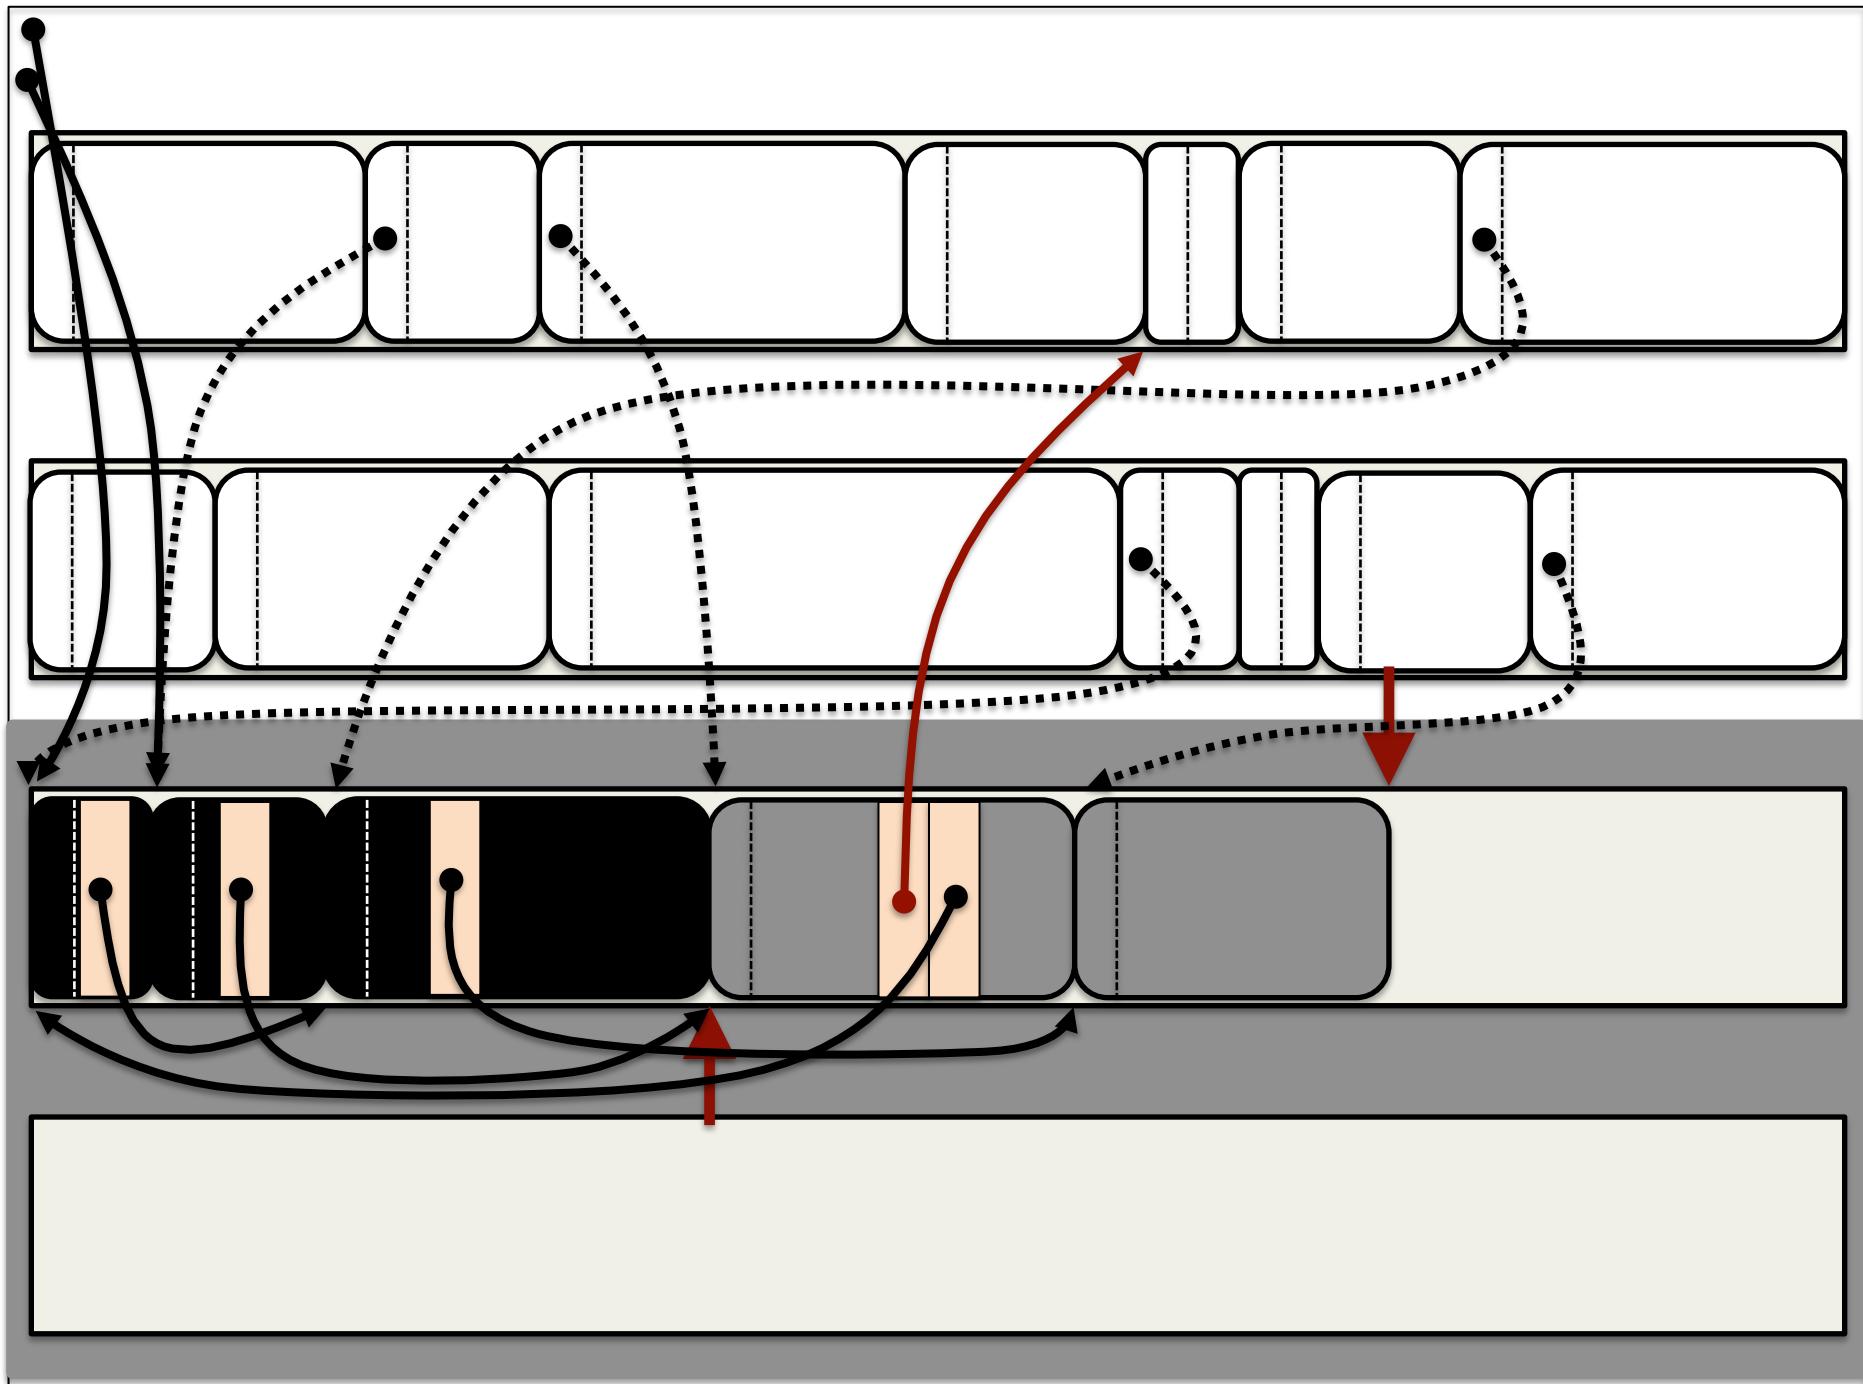


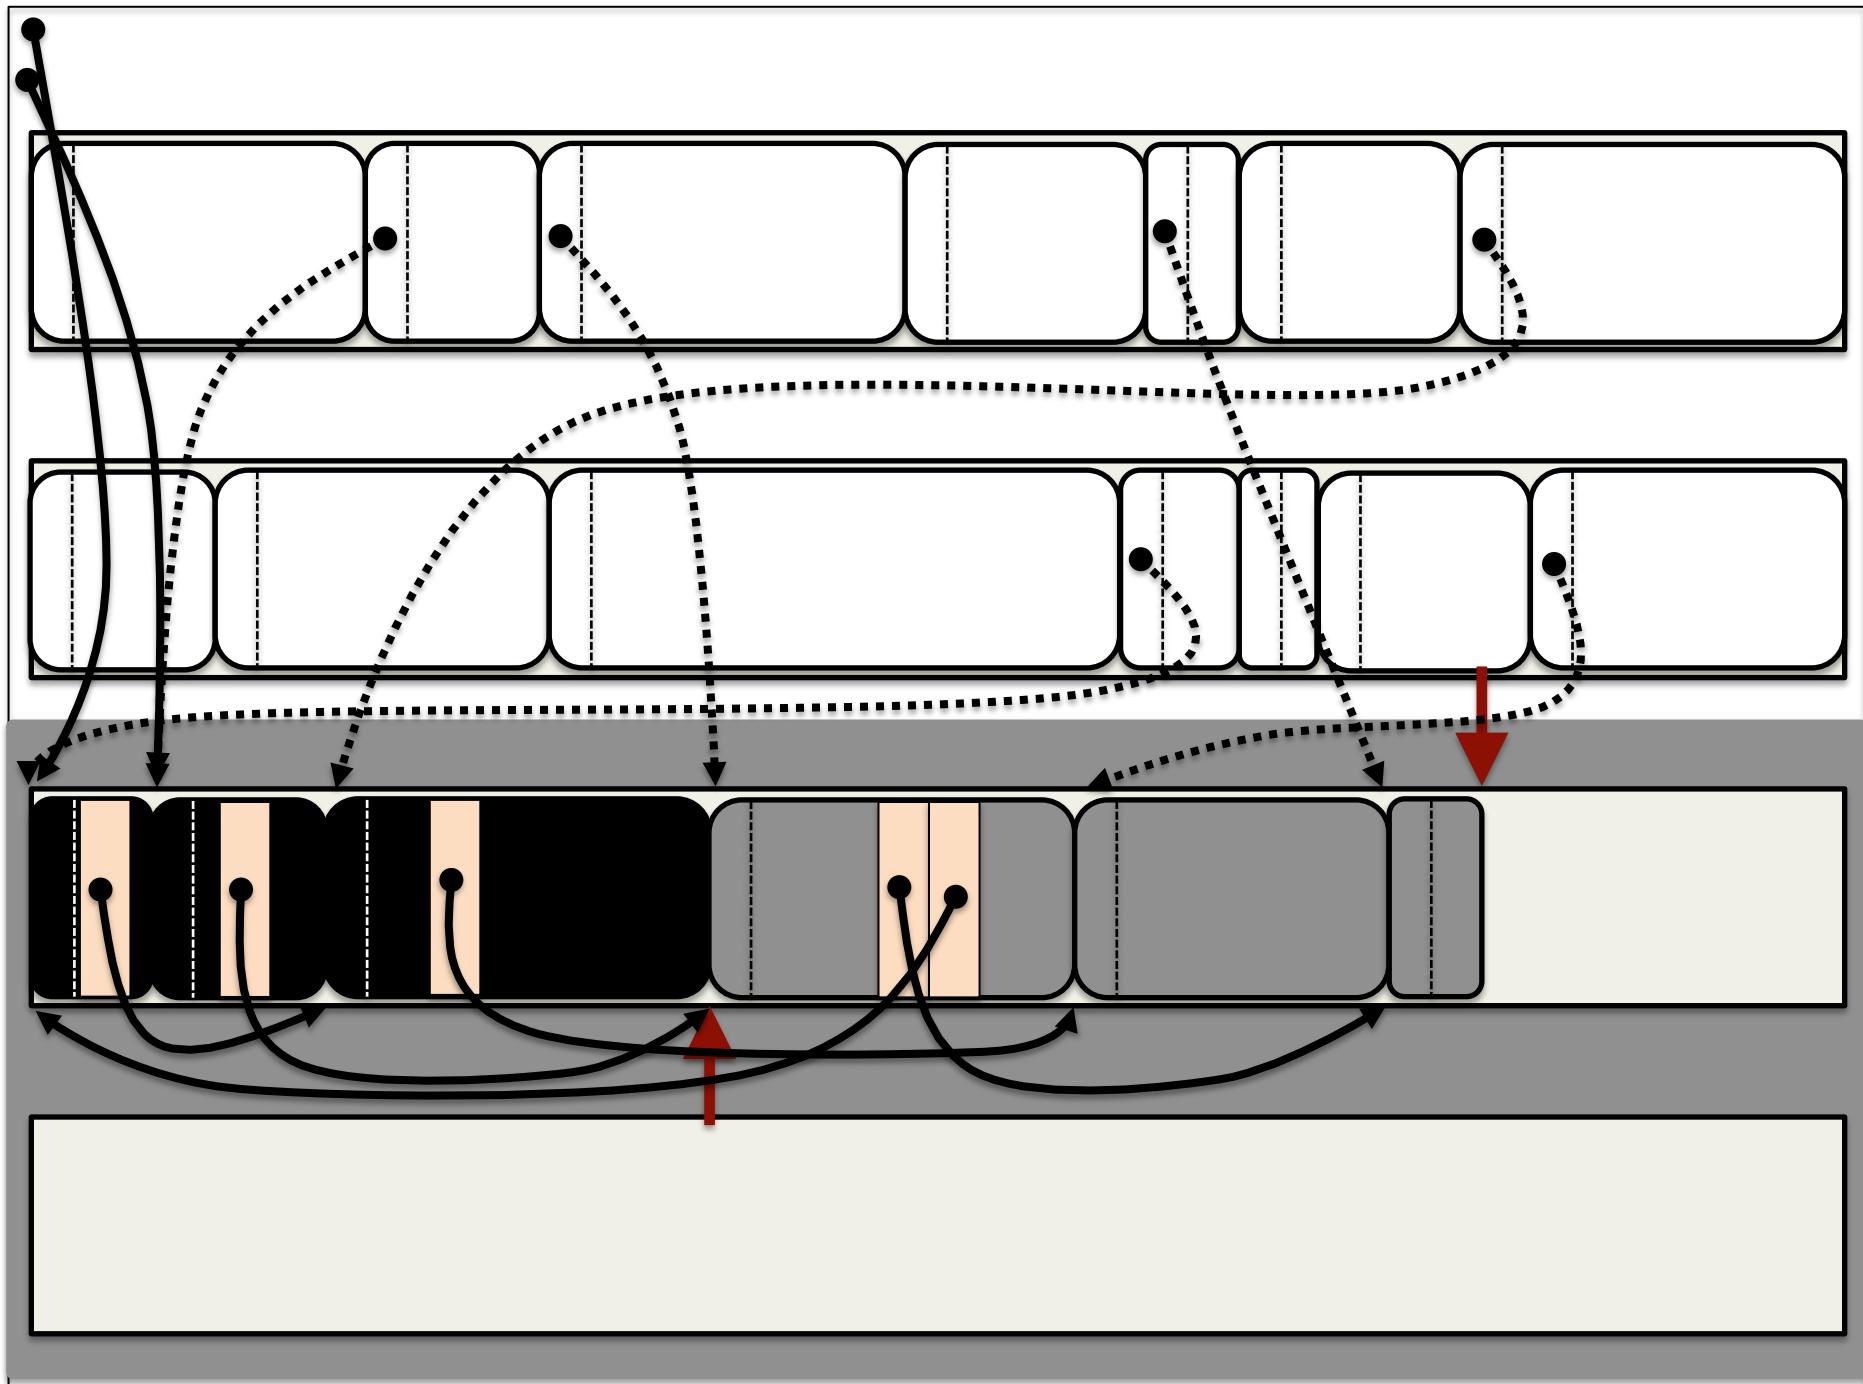


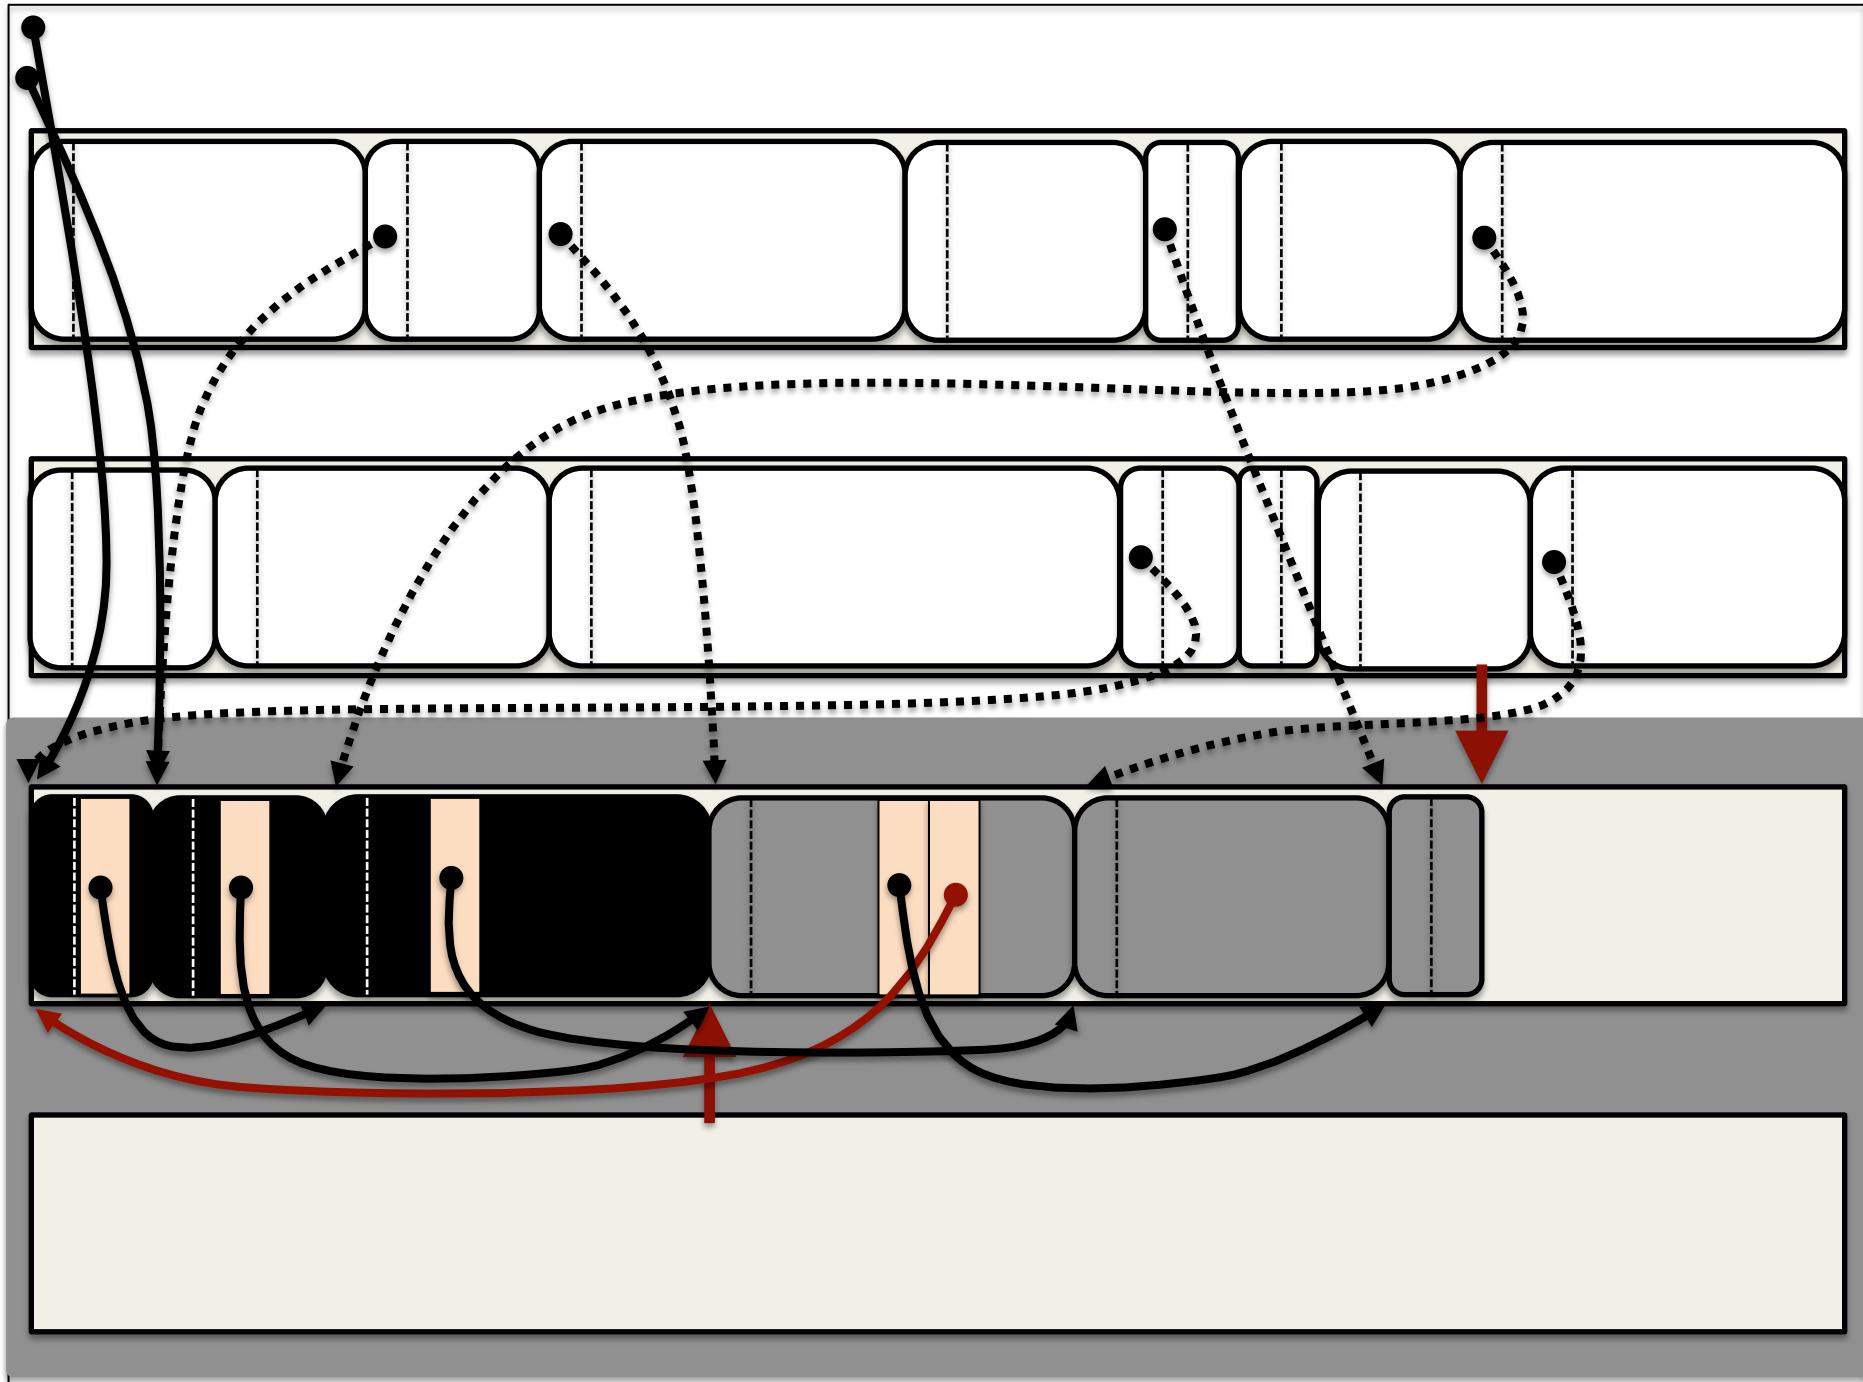


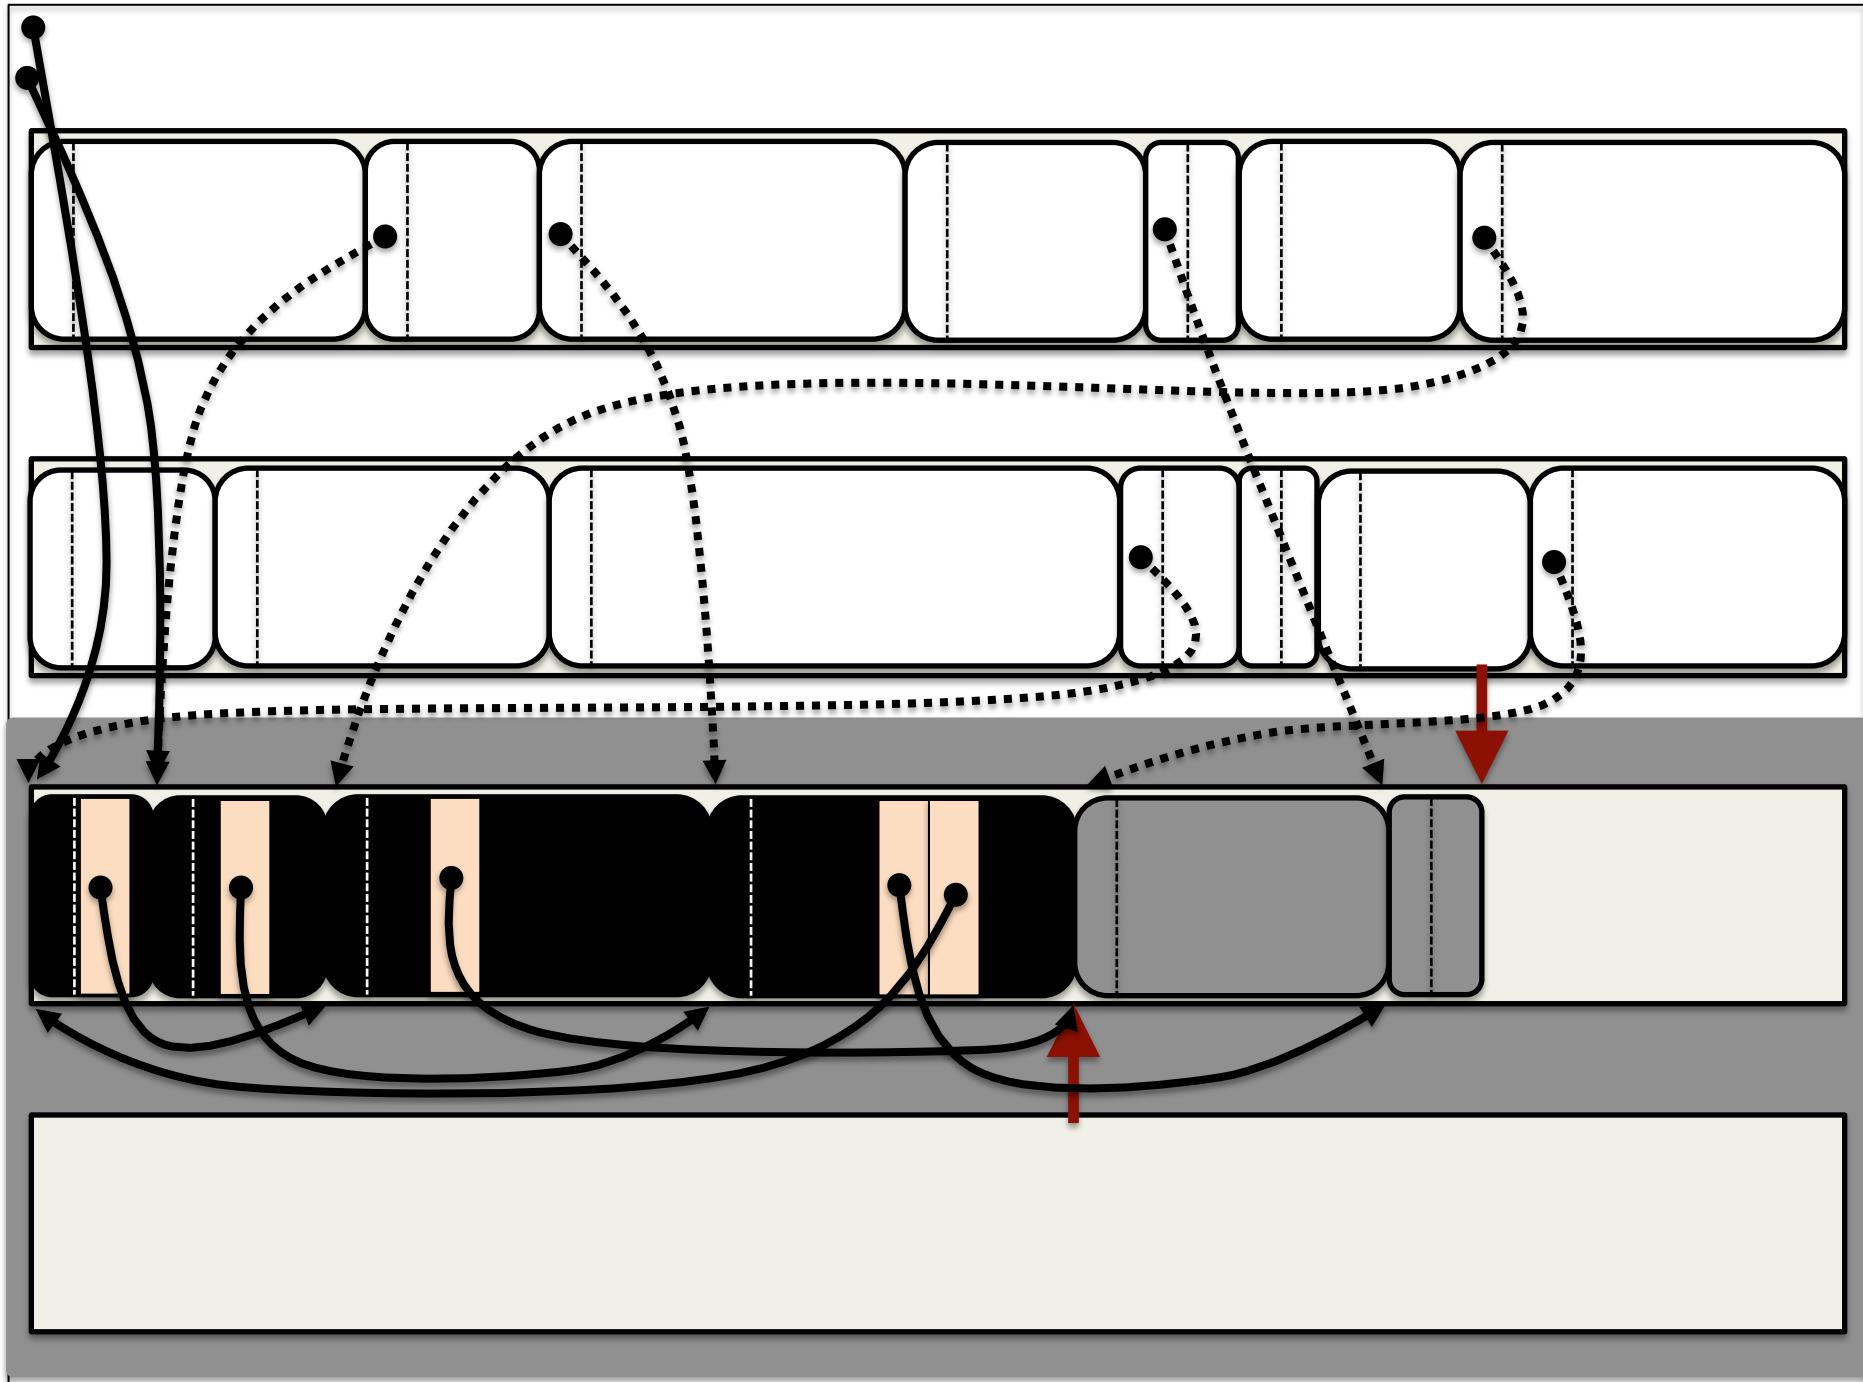
Updating References

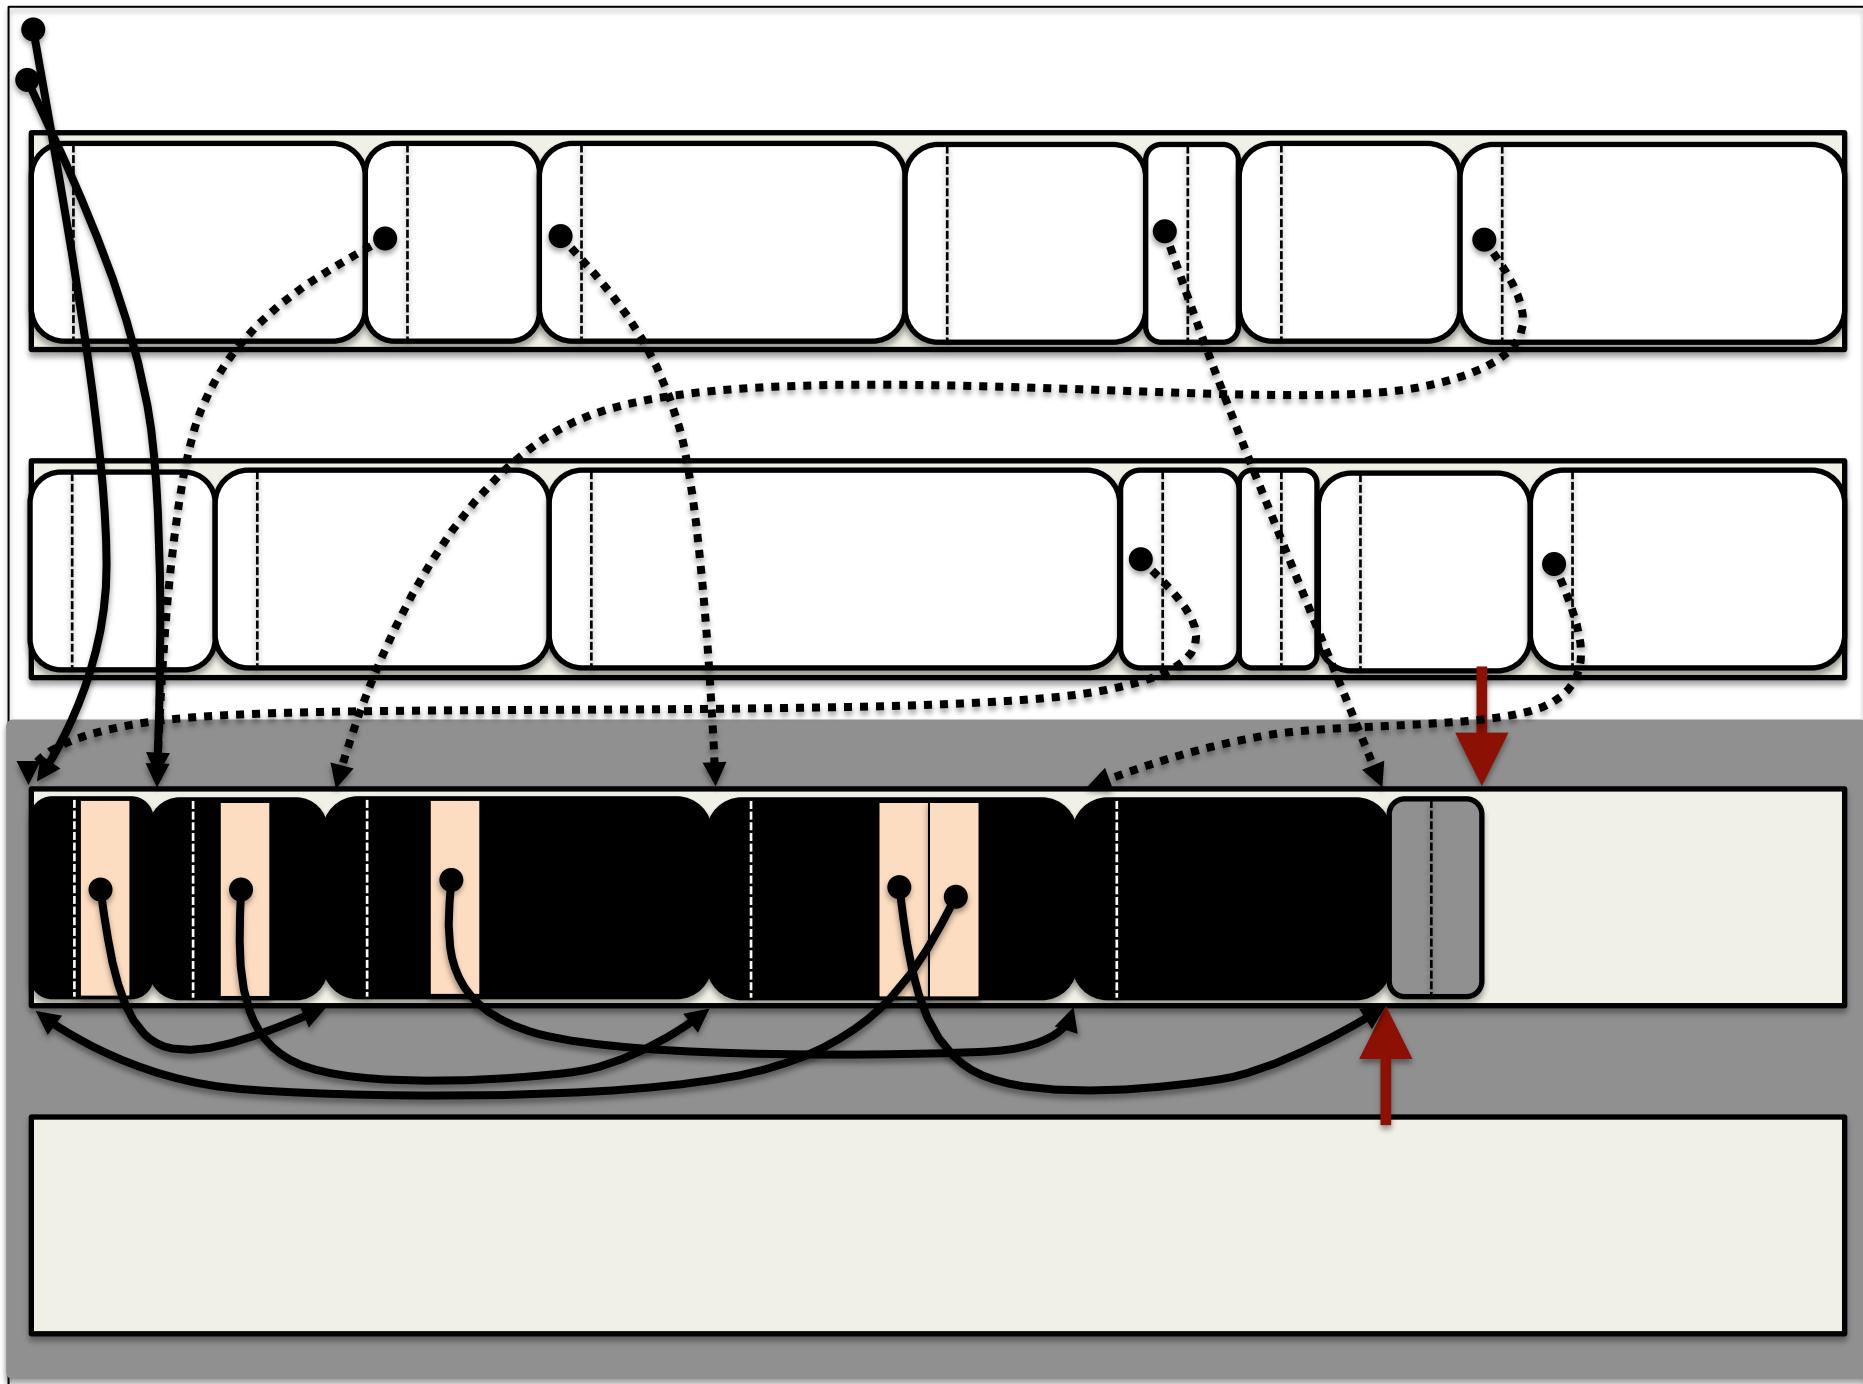
- There can be no references from new to old
 - Part of marking an object black
- Two cases when scanning an object
- The object has not been copied
 - Copy to the new space
 - Update the reference
- The object has already been moved
 - Use the forwarding pointer to update the reference

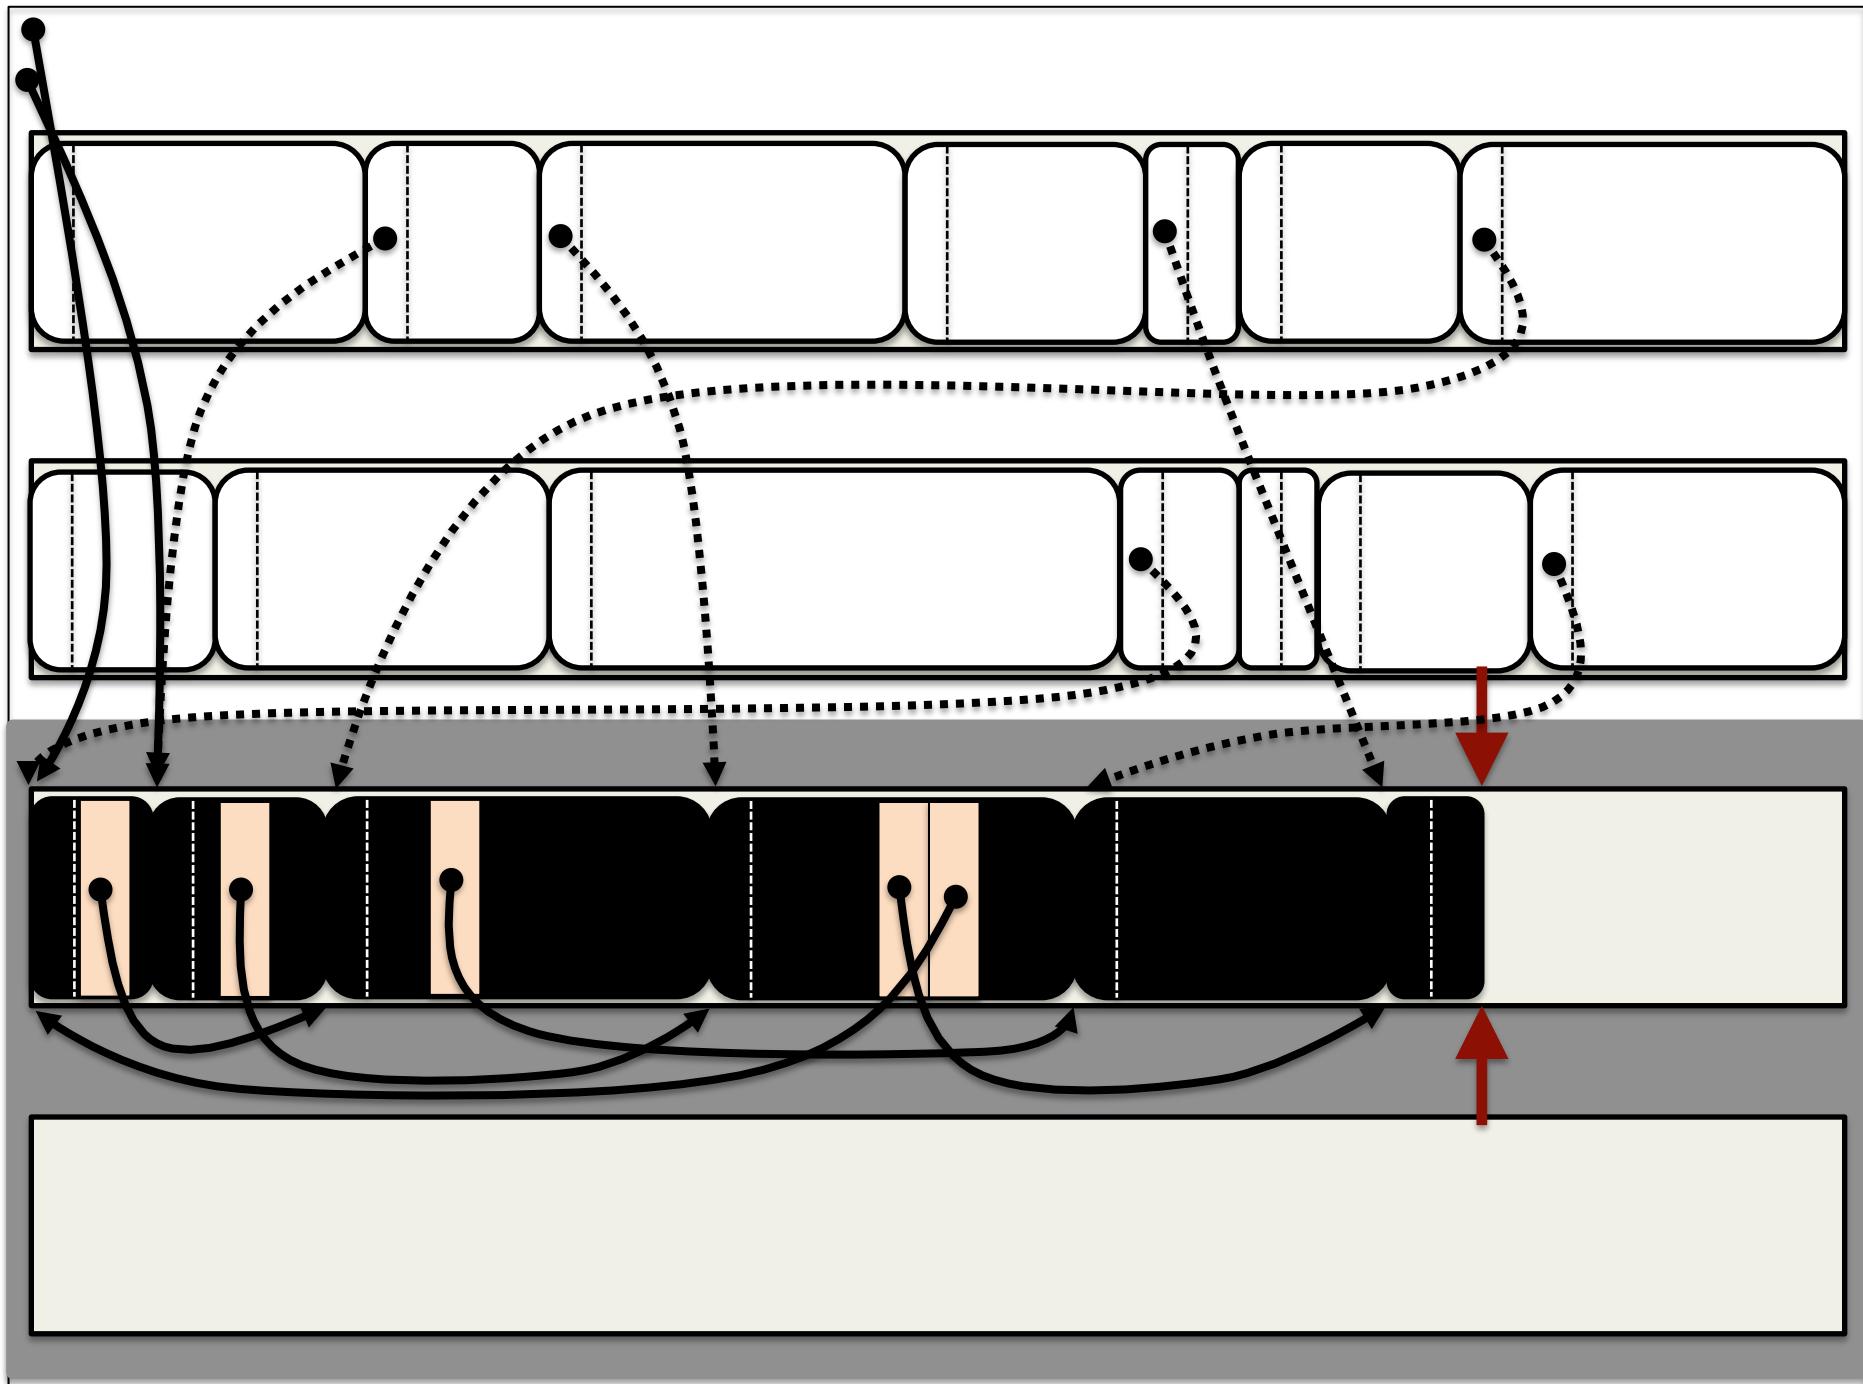












Finishing Up

- Algorithm ends when the finger pointers meet
 - All gray objects have been scanned
- No per-object finishing processes
 - New-space objects are not marked
 - All references point to the new space
 - Forwarding pointers only exist in garbage objects
- Continue allocating in the new space

Assignment 3

- Implement Cheney's algorithm
 - Bump-pointer allocation
 - Semi-Space heap layout
 - Stop the world thread management
 - Single pass mark and copy
- Implement the **MemoryManager** interface
 - Implementation details are up to you
 - Following the algorithm correctly matters
 - Grading will look at the state of the VM after GC

SimpleJava Memory Management

- Interfaces under `vm/memory`
 - `Heap`
 - `MemoryManager`
 - `Region`
- Implementations in `impl/memory`
 - Implementations of heap tracking
 - Two possible memory managers

Allocation

- Allocation handled by the `Region`
 - Currently supports `BumpPointerAllocator`
- Look at the existing allocation code
 - `InfiniteMemoryManager.java`
- Look at `BumpPointerAllocator.java`
 - You shouldn't need to change the implementation
 - Could help while debugging

Heap Class

- Java doesn't give direct pointer access
 - You now know why
- A memory manager needs to access pointers
- In SimpleJava we simulate the heap
 - In production we could replace with raw memory access
 - Remember JIT intrinsics
- The heap is a logical construct
 - Maintains a mapping of pointers to objects

Heap Layout

- The heap is divided into regions
 - Areas of memory to manage independently
- At one extreme there is a single region
 - `InfiniteMemoryManager`
- At the other there can be many
 - We'll talk about this next week
- For your work there will be two

Heap Operations

- Carve off a region from the heap
 - Current implementation assumes bump pointer
 - This may change for the final project
- `memcpy`
 - Copy an object from one location to another
 - Makes an exact duplicate of the object

Object Builder

- Creates correctly structured objects
 - Look at the code in `NewInst` and `NewArray`
- Handles the heap management
- Only way that you should create new objects
 - You may copy objects with `Heap.memcpy`

Roots

- Thread stack and local variables
 - Accessed via `VmThread` interface
- Static fields
 - Accessed via `VmClassLoader`
- String intern table
 - Accessed via `ObjectBuilder`

Pointers to Pointers

- Need to update root locations
- Two options
 - Trace through stack, local, static and intern table
 - Use pointers to pointers
- Look at the `ReferenceLocation` interface

OutOfMemoryError

- A GC-based system can still run out of memory
 - All objects survive the collection
 - Not enough space freed up for an allocation request
- The VM spec calls for an **OutOfMemoryError**
 - SimpleJava doesn't yet throw errors internally
 - Exception handling will be in Assignment 4
- Throw a **VmInternalError**
 - Terminates the VM

Writing a Garbage Collector

- Notoriously tricky to debug
 - Errors show up long after the bug
 - Bugs can be non-deterministic
- Better to focus on getting it right than fixing later
 - Think of the GC as a graph algorithm
 - Make sure that you modify the graph correctly
- Test and verify

Heap Verification

- Unit tests let you examine the heap closely
 - Don't need to set up mutators to do the right thing
 - It's how I will grade the assignment
- Consider adding a verification phase
- Check at the end of GC for invariants
 - No pointers to the old space
 - Re-do trace and count reachable objects
 - All roots properly updated