

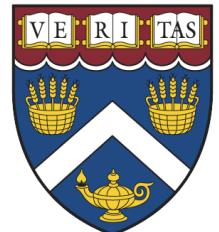
In the last week

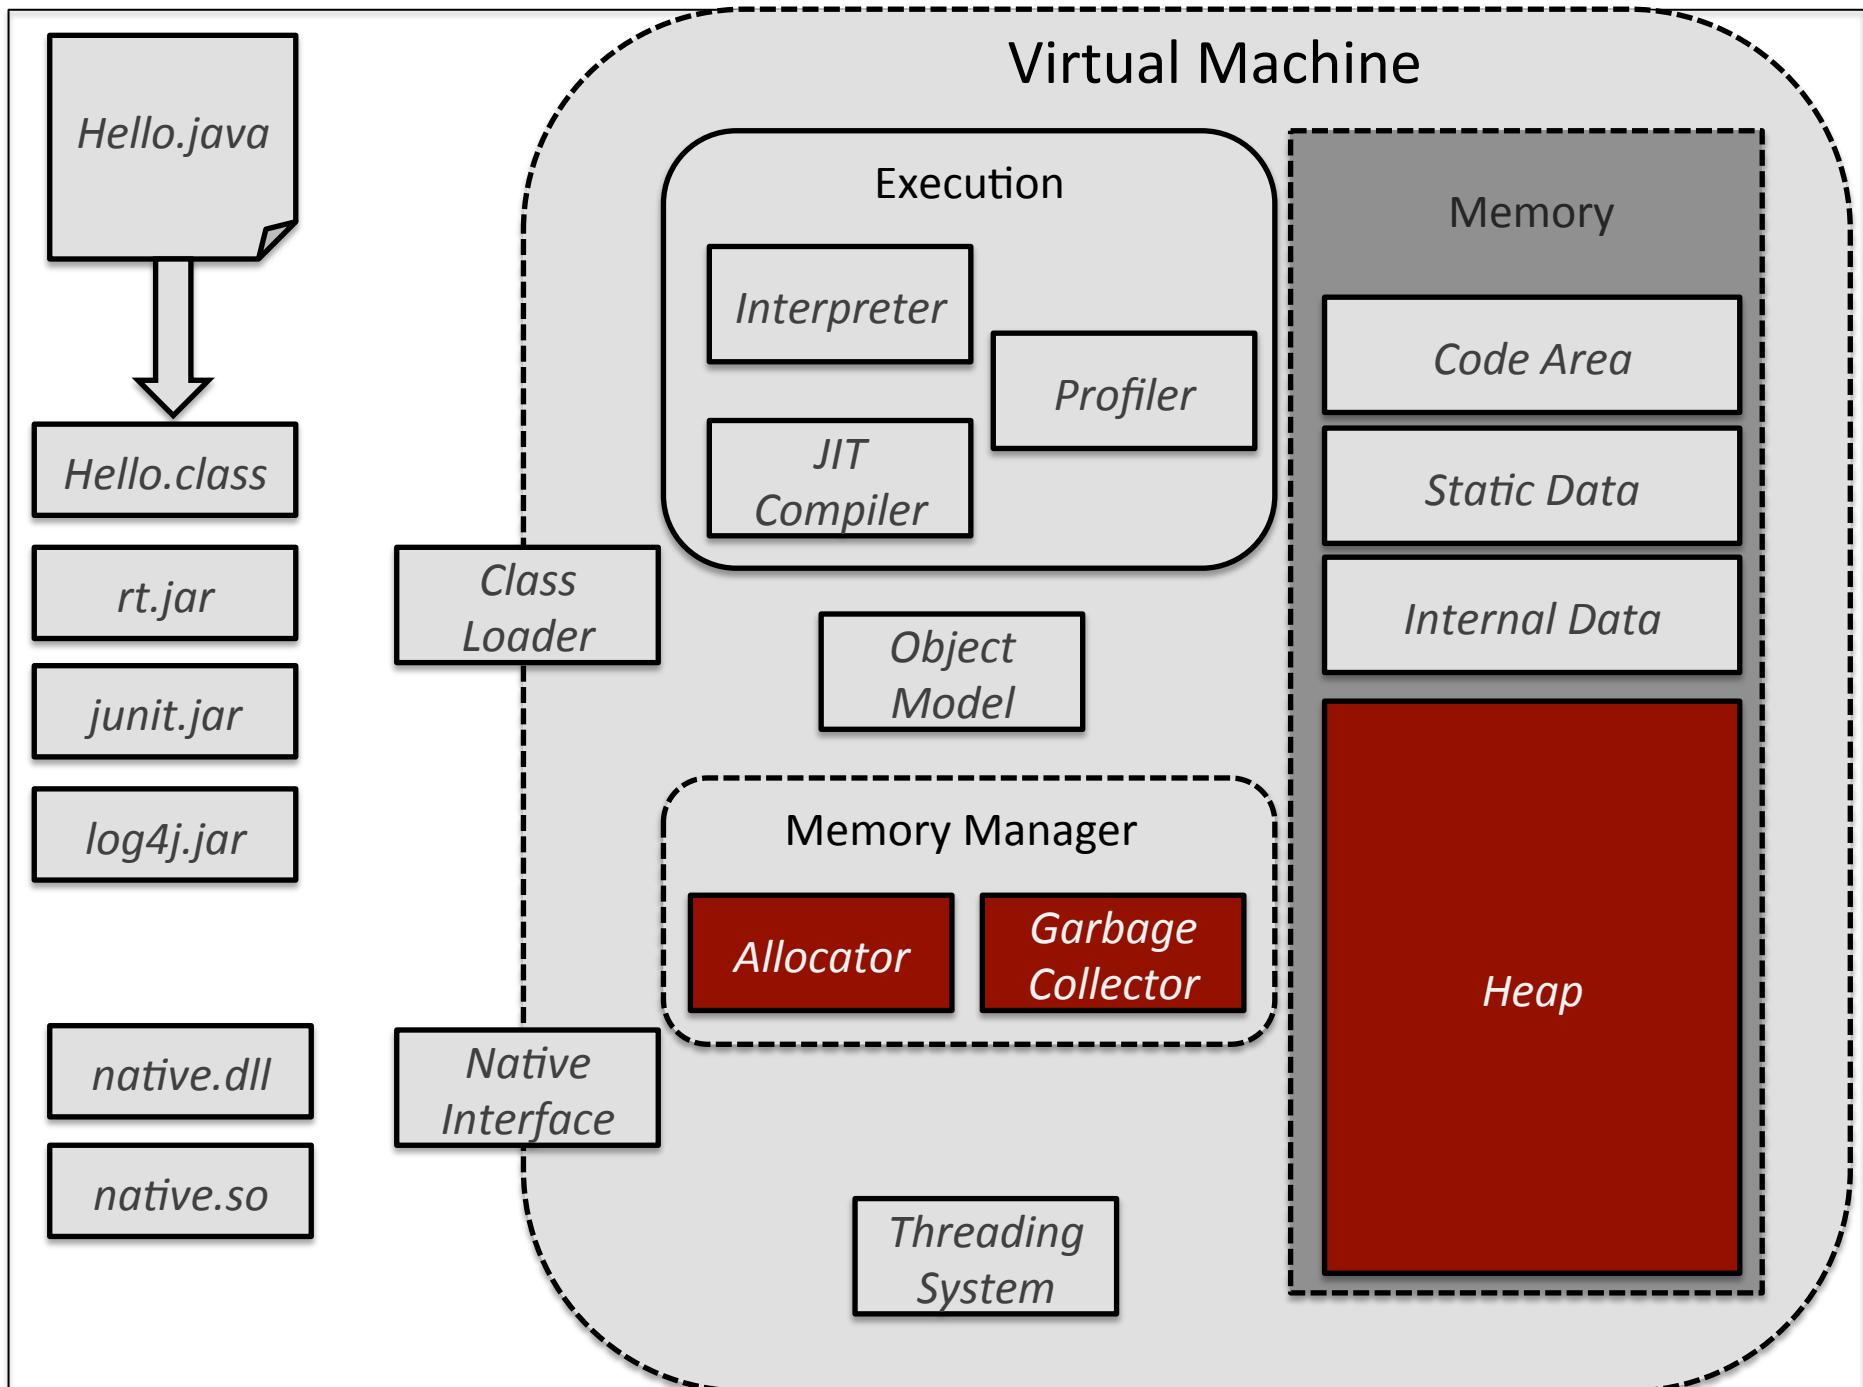
- Assignment 2 grades released
- Assignment 3 underway
 - Question about the heap and region sizes
 - Why isn't there a garbageCollect() method?

Forwarding Pointers

- Part of Assignment 3 involves relocating objects
 - Which means installing forwarding pointers
- Some confusion as to where to store the pointer
 - In C/C++ you would overwrite a header word
 - Set a bit to indicate that you'd installed the pointer
- Since we're using Java we can't set bits
 - Instead just store a `HeapPointer` to the object header
 - Use the `TypeDescriptor` word
 - Use `instanceof` to check if the word contains a pointer

Generational Garbage Collection





Principle of Locality

- Temporal locality
 - Once you use a value, you will likely use it again soon
 - Afterwards it may be a long time before using it again
- Spatial locality
 - Nearby objects tend to be used together
 - Extends to objects that reference one another

Copying Collection

- Take advantage of dynamic spatial locality
 - Breadth-first scan of the heap
 - Move objects based on their connectivity
- Costs vary based on live, not garbage, objects
 - Copy live objects out of the old space
 - Collect everything else en-masse

Copying Overheads

- Copy time overhead
 - Every live object must be copied at every GC
 - Long-lived objects carry the biggest overhead
 - Can also be the largest objects
- Copy reserve overhead
 - Semi space collector requires half the heap
 - Often that space is completely unused

Room For Improvement

- Copy objects as few times as possible
- Copy smaller areas of the heap
 - Copy reserve proportional to data being copied
- GC spends a lot of time tracing older objects
- Target objects that are more likely to be garbage
 - More garbage means fewer objects to copy

Generational Hypotheses

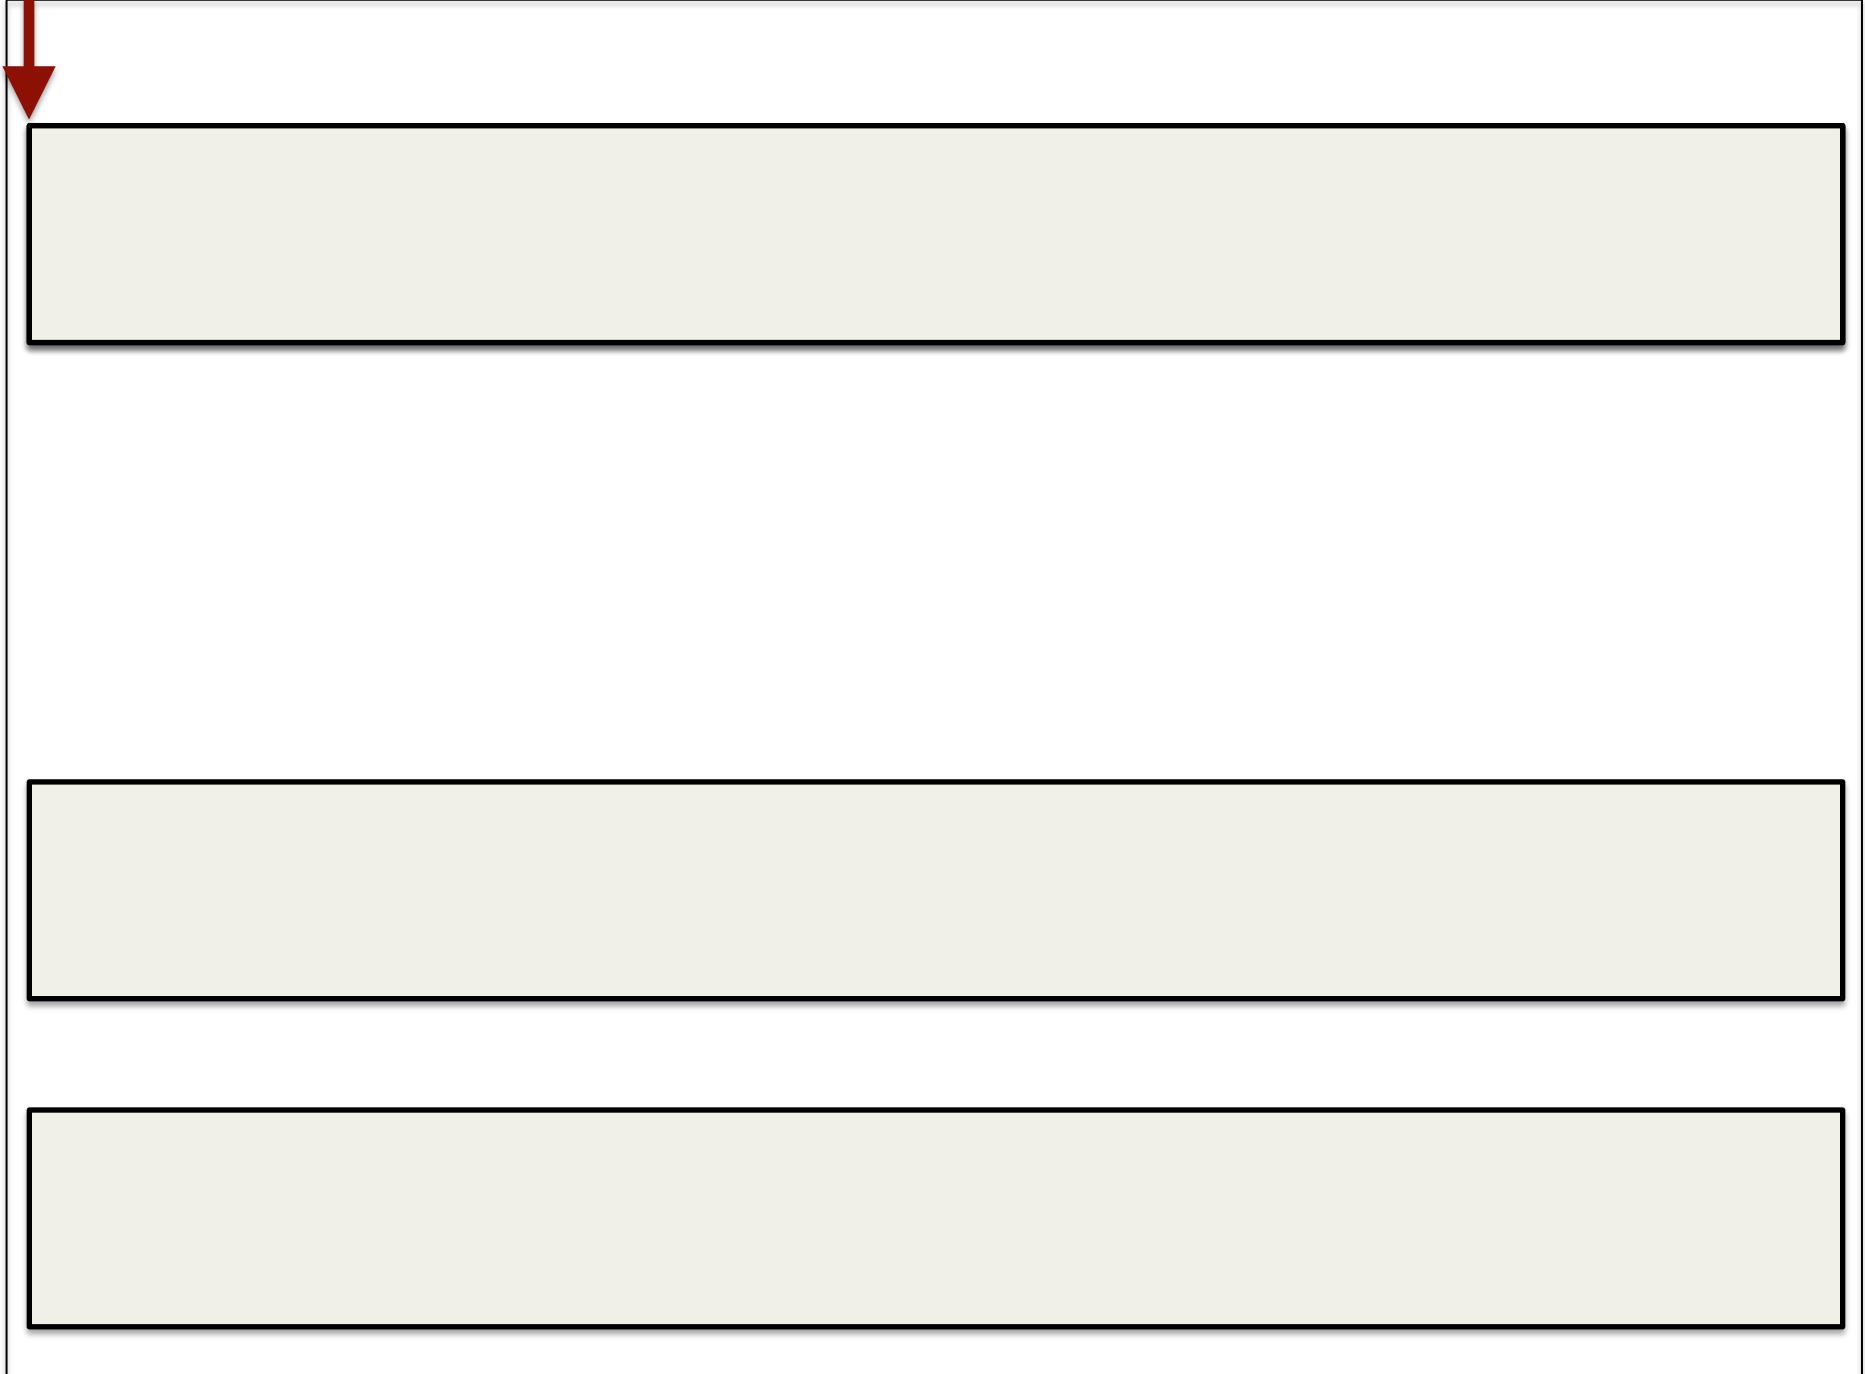
- Heuristics that suggest what is likely to be garbage
- Weak Generational Hypothesis
 - Most objects become garbage shortly after allocation
 - Also known as high infant mortality
- Strong Generational Hypothesis
 - Older objects are more likely to survive a long time
 - Less well supported

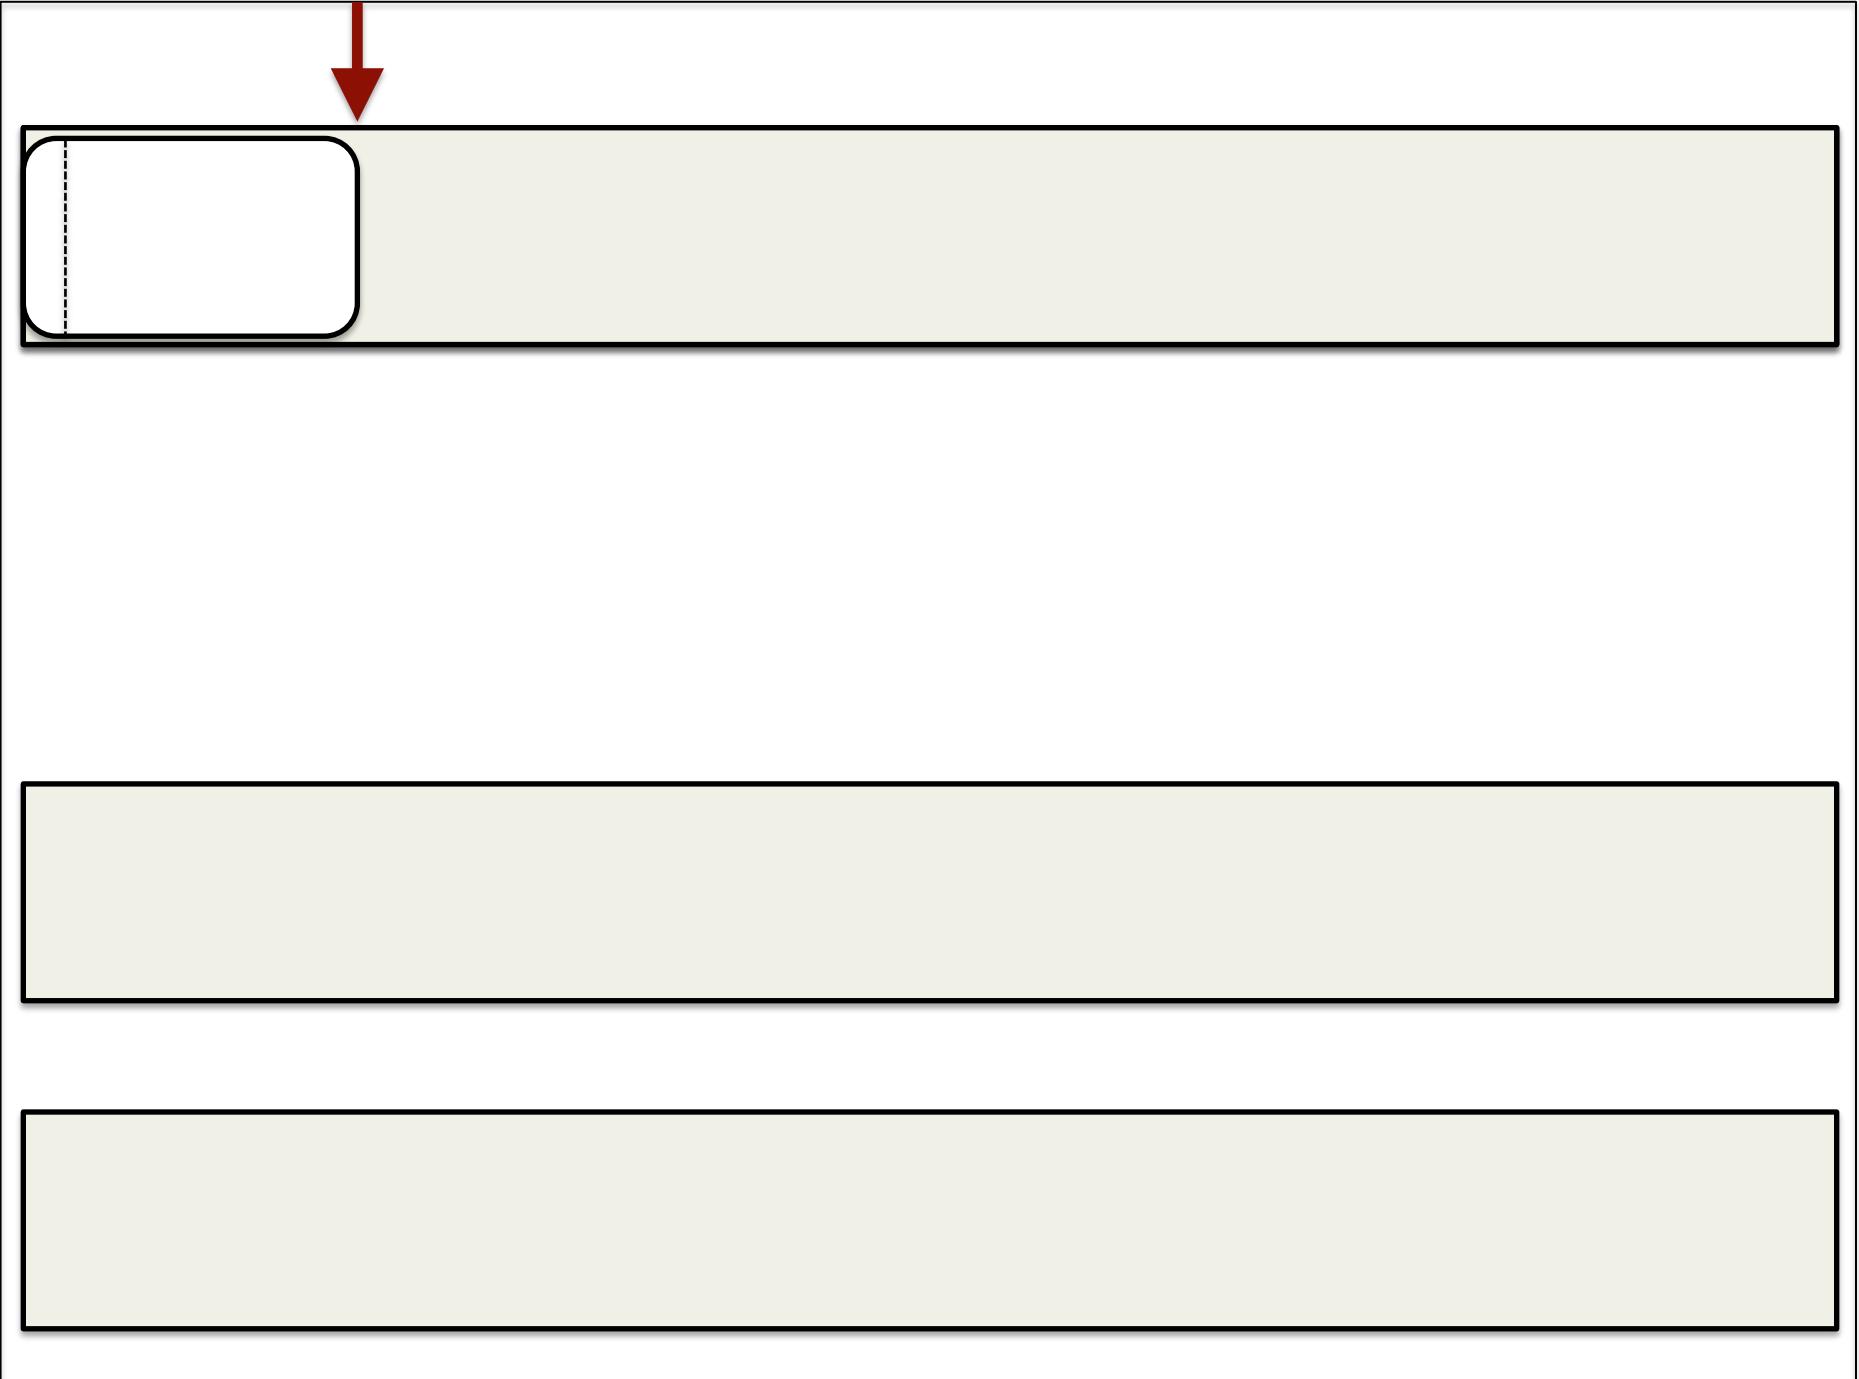
Generational Hypotheses

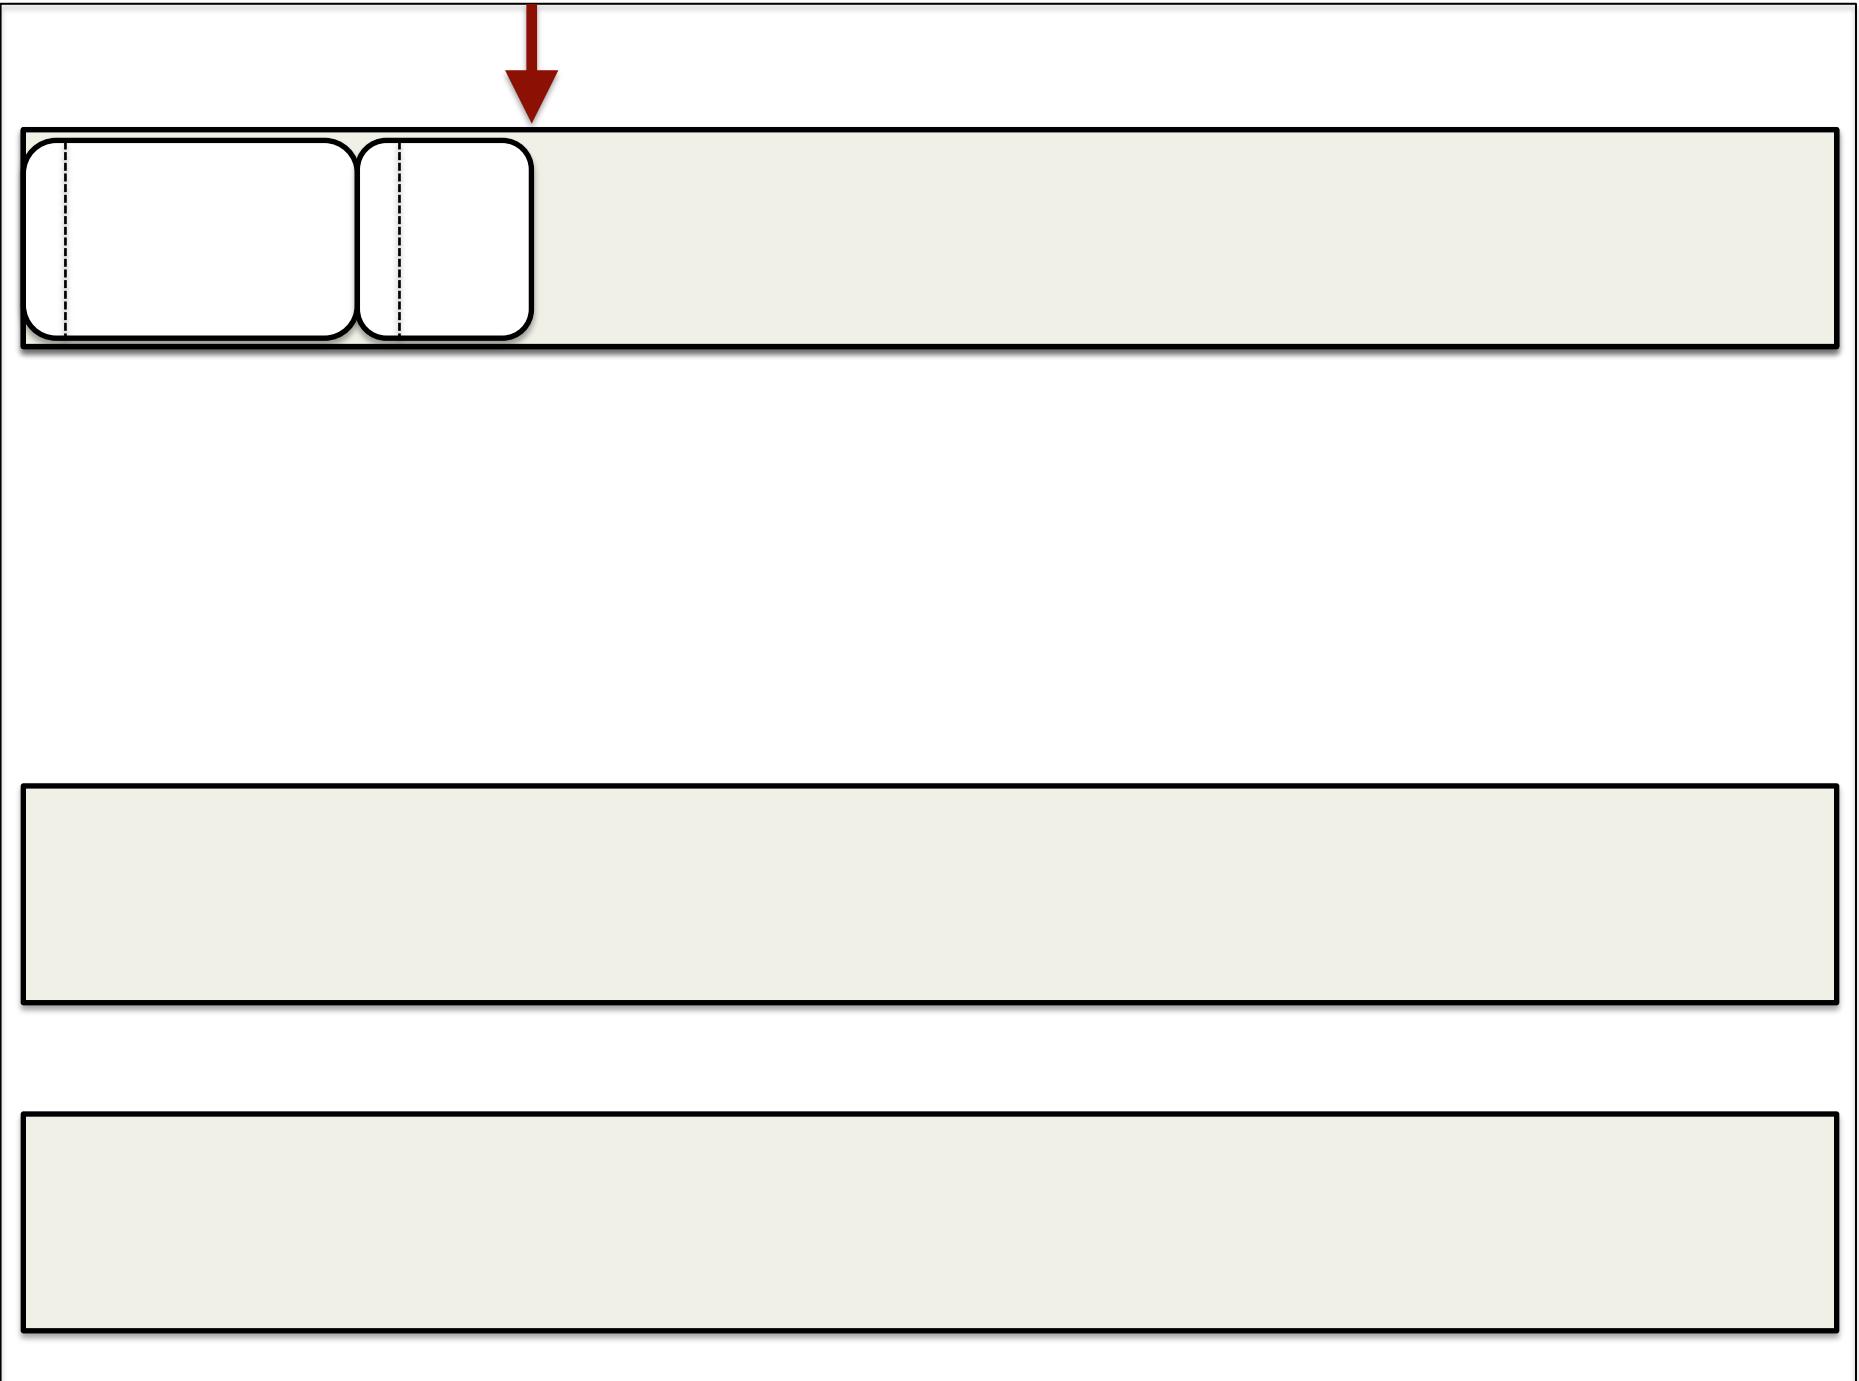
- First proposed in 1984
 - Consistently seen in practice
- Typically around 10% of objects survive
 - Varies by application
 - Frequently far fewer

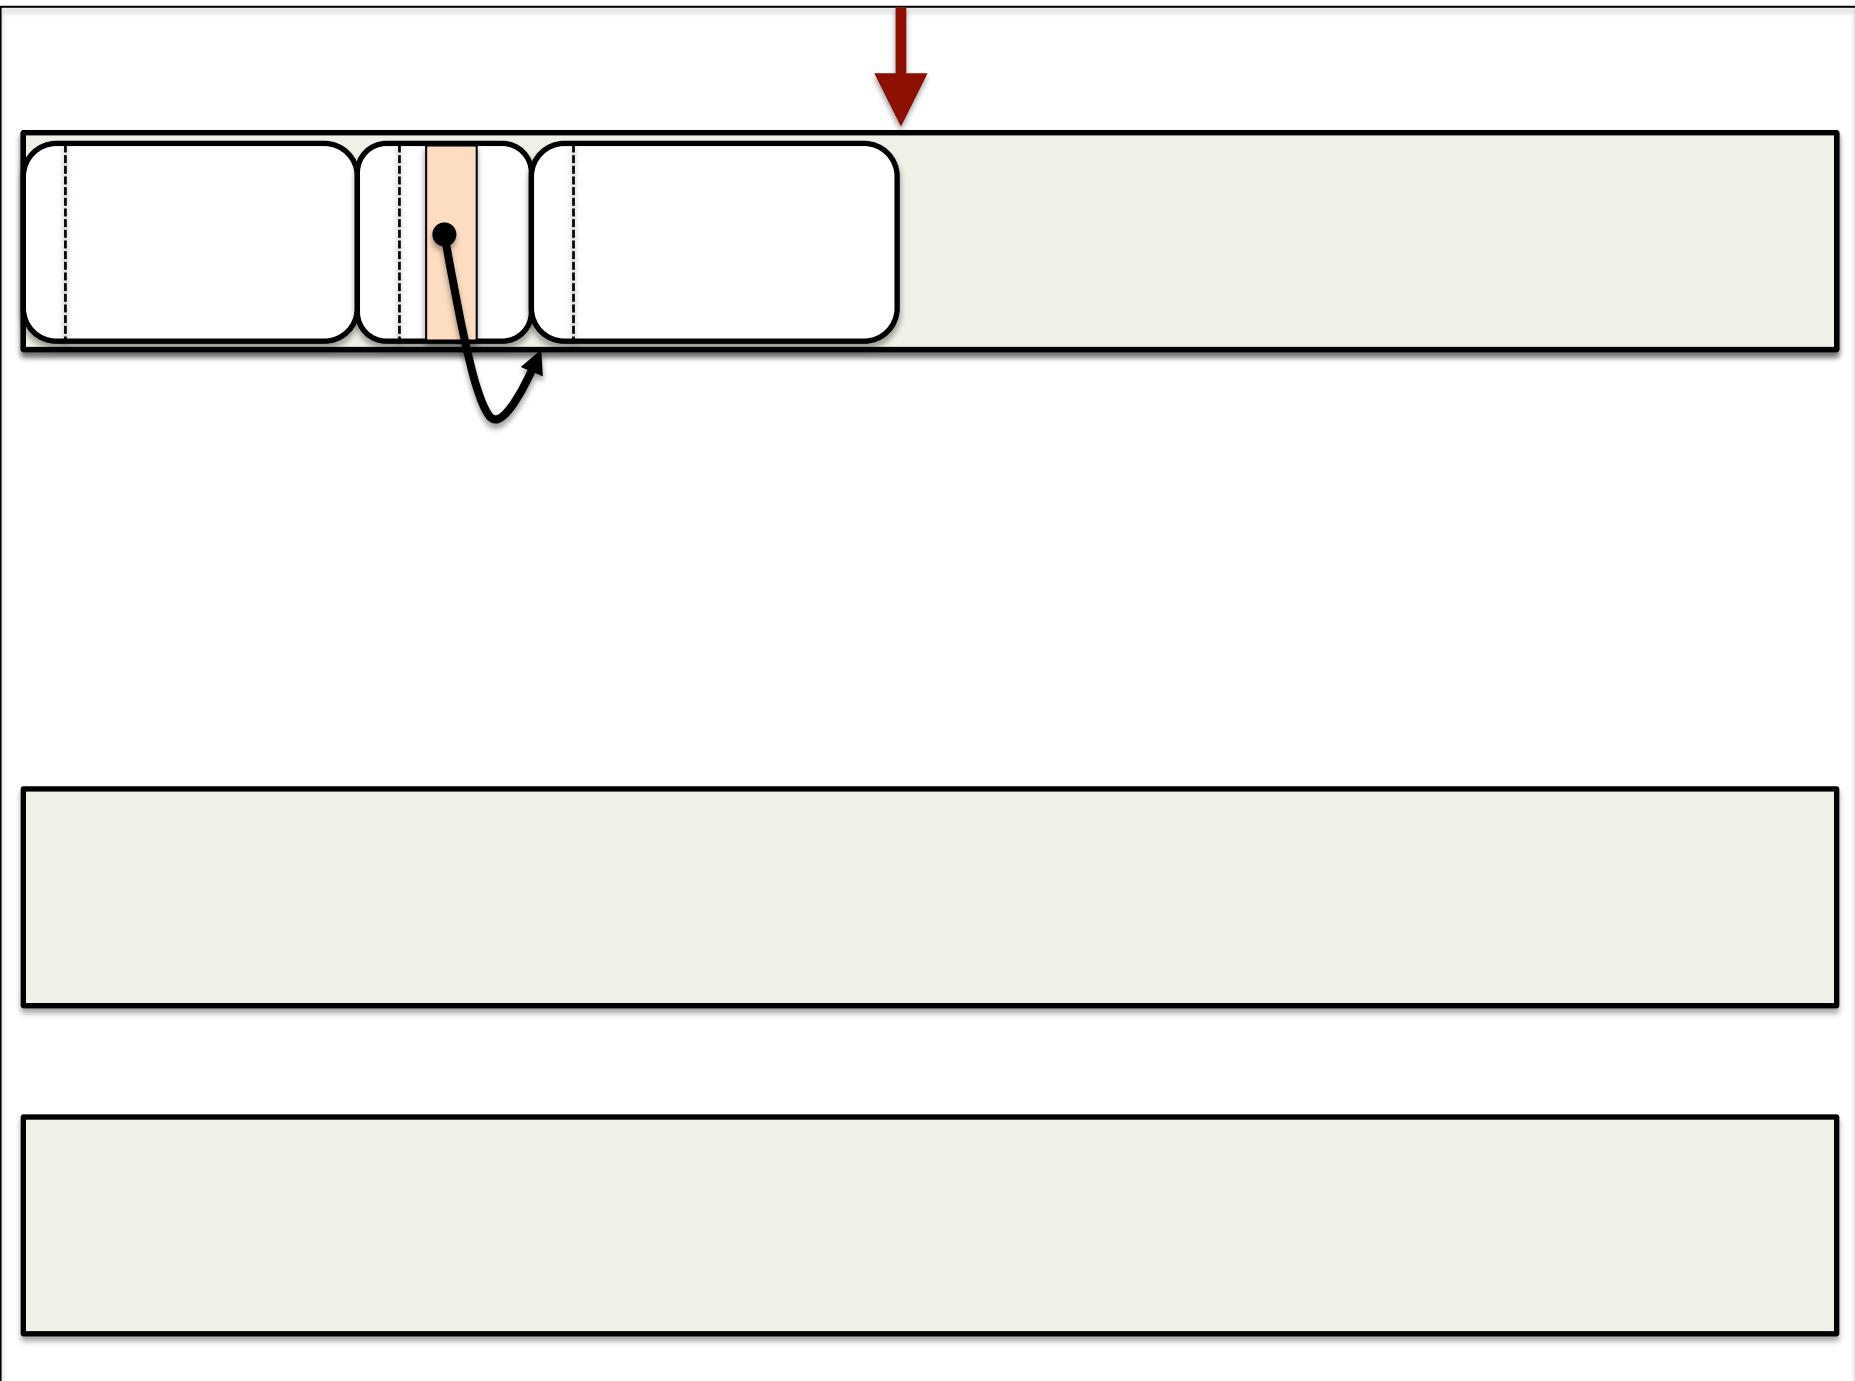
Generational Collectors

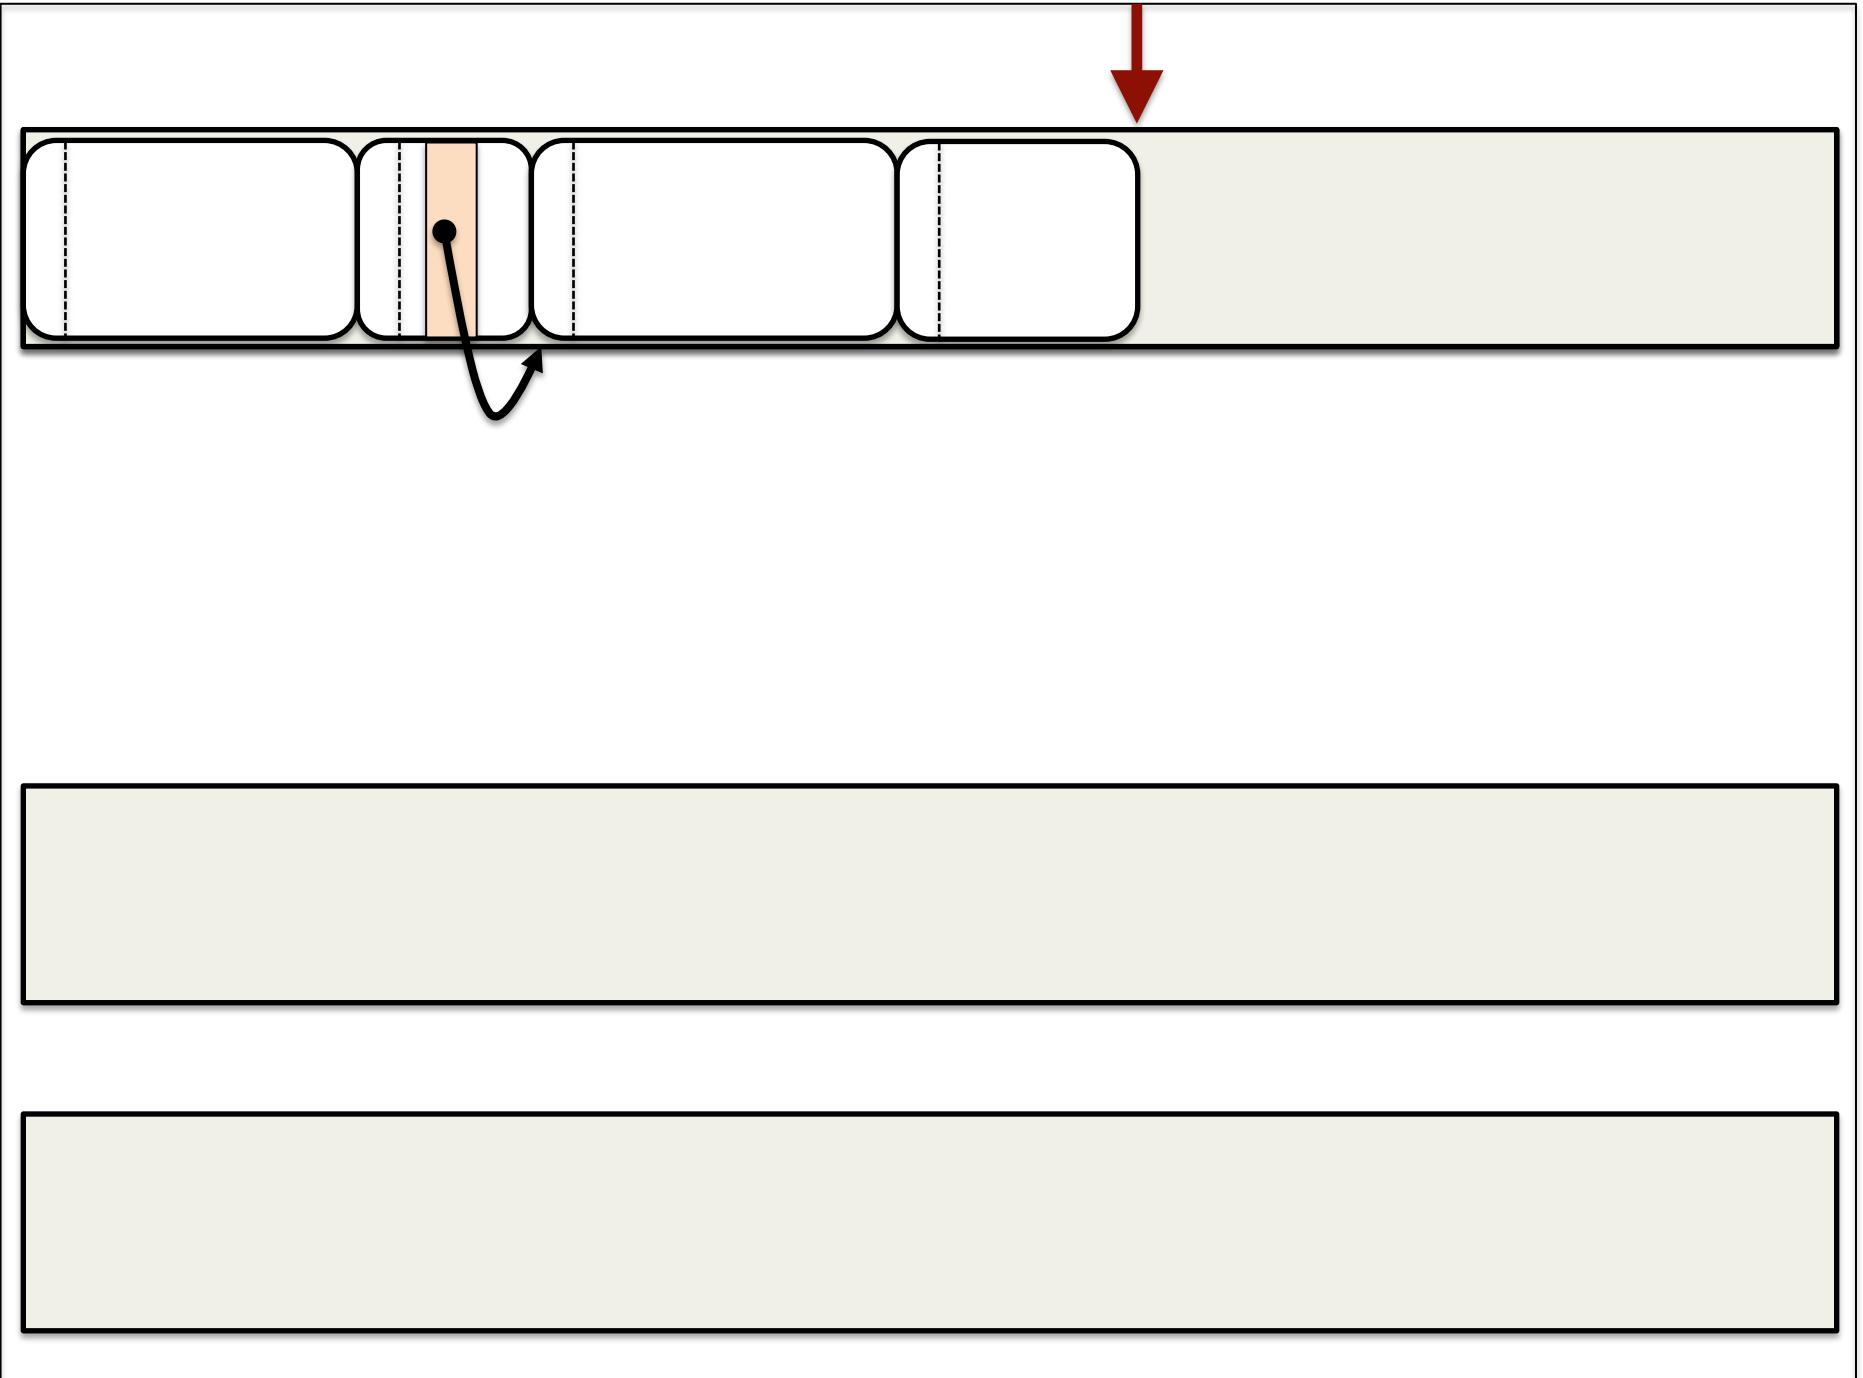
- Partition the heap into two regions
 - Regions may not be the same size
- Allocate to the first region
 - No flipping of semi-spaces
 - This area is the nursery – all allocation happens here
- During GC, move live objects to the second region
 - All objects there have survived a GC
 - Known as the mature space

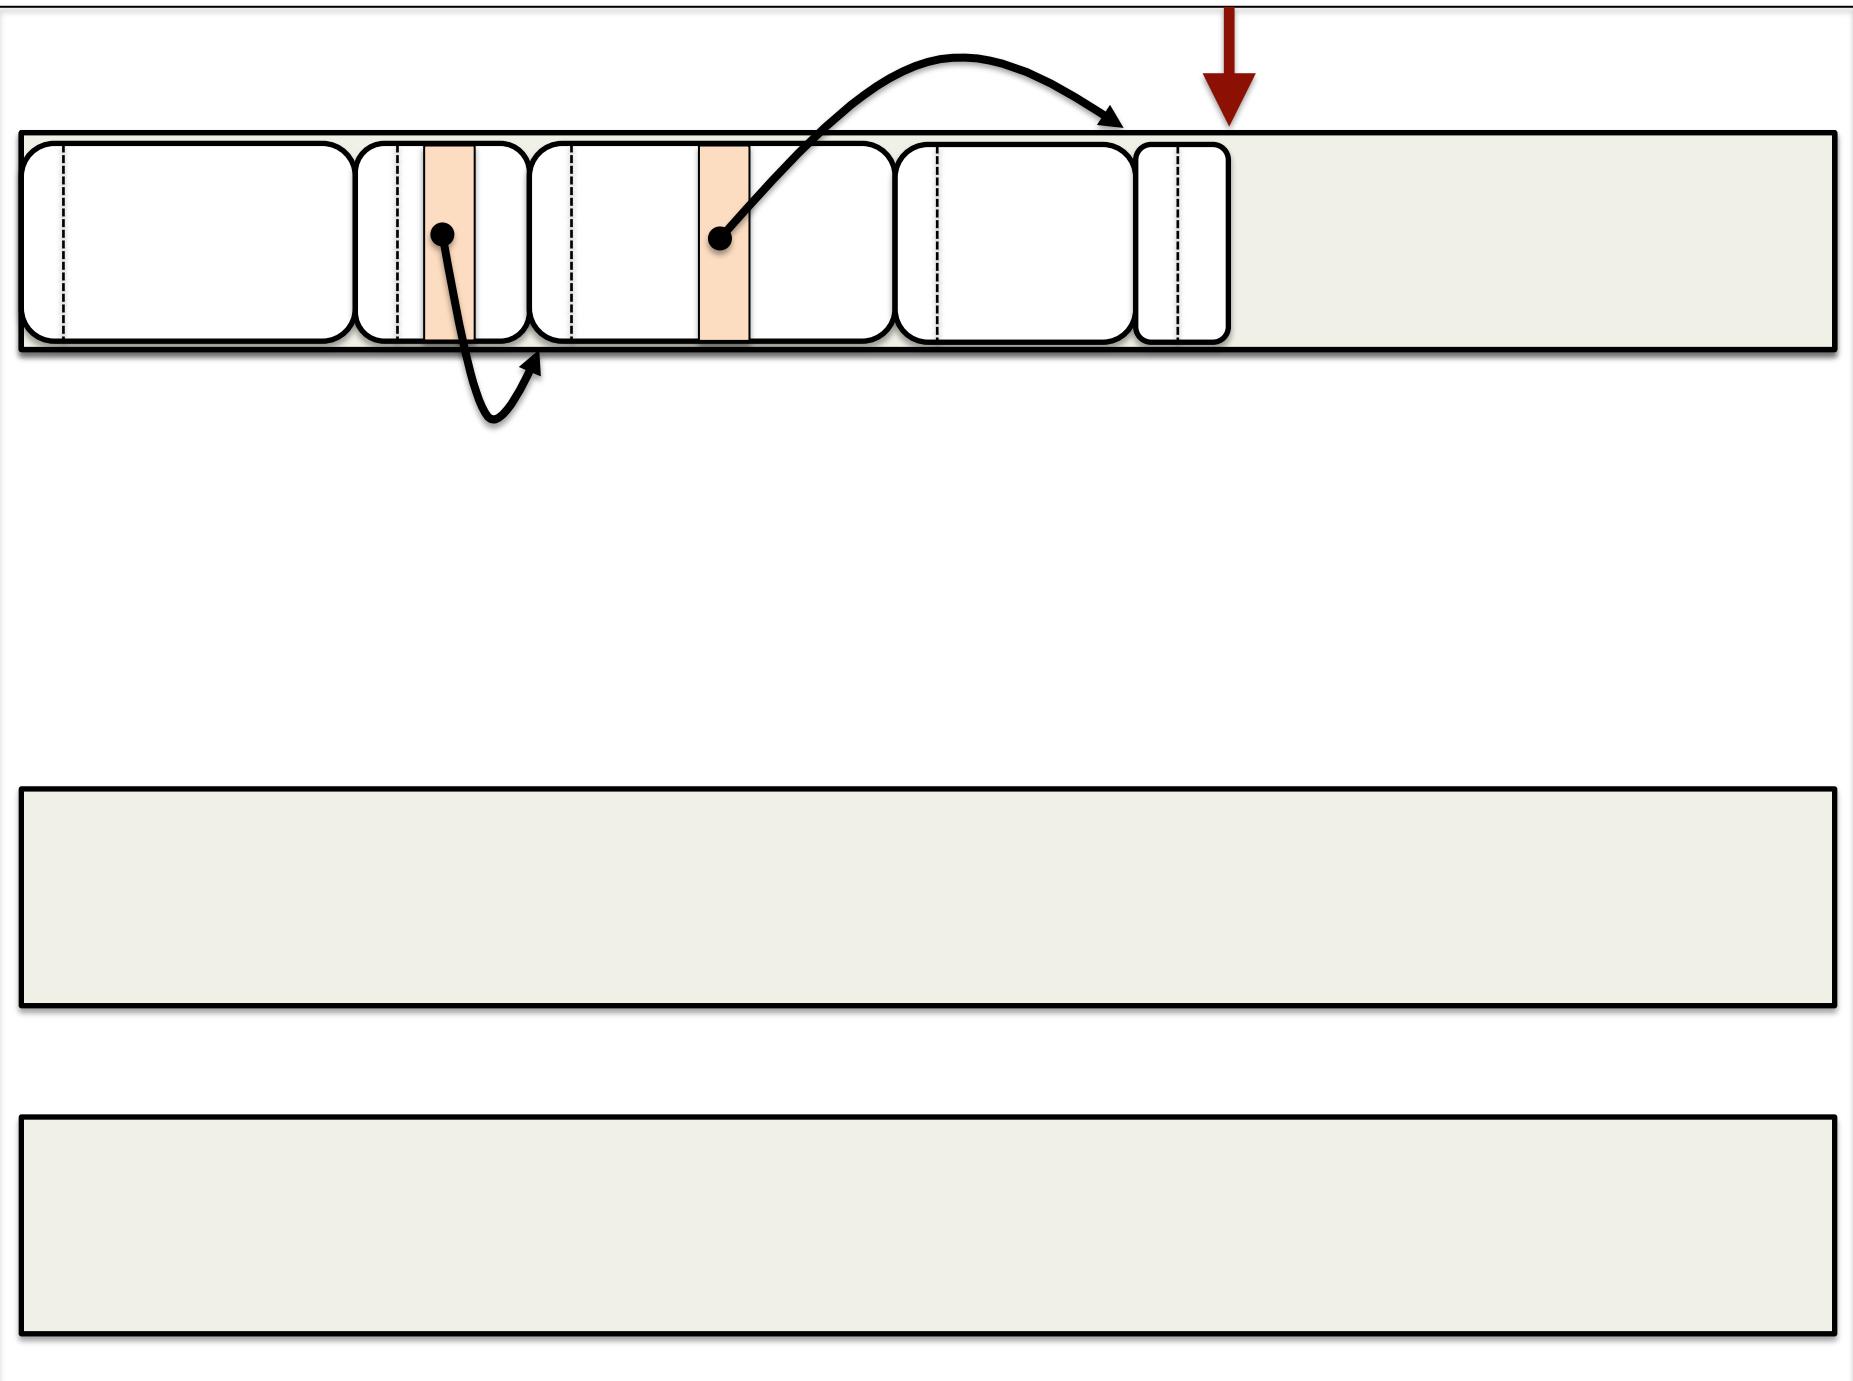


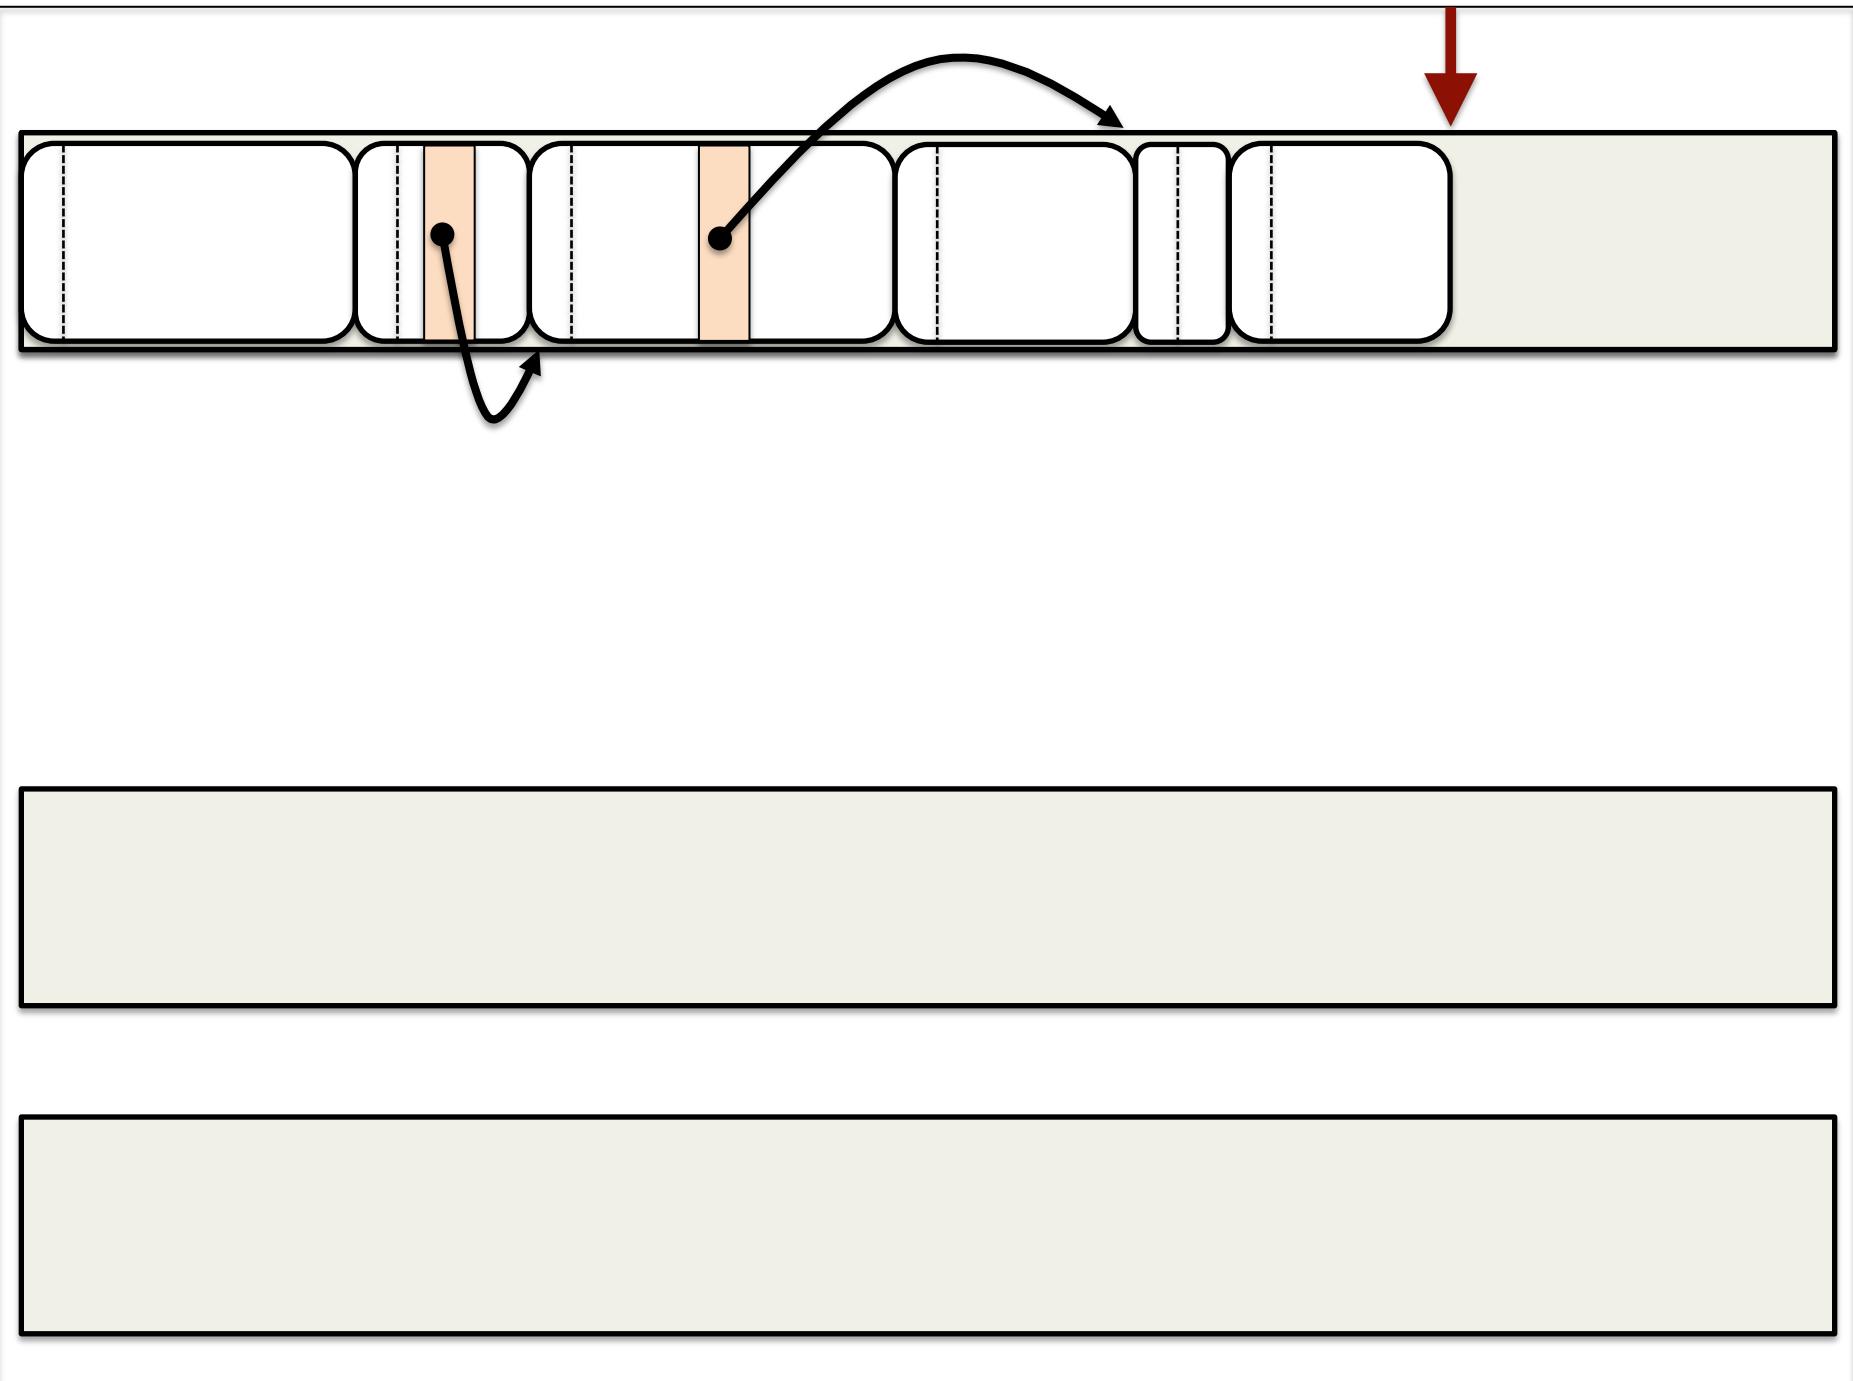


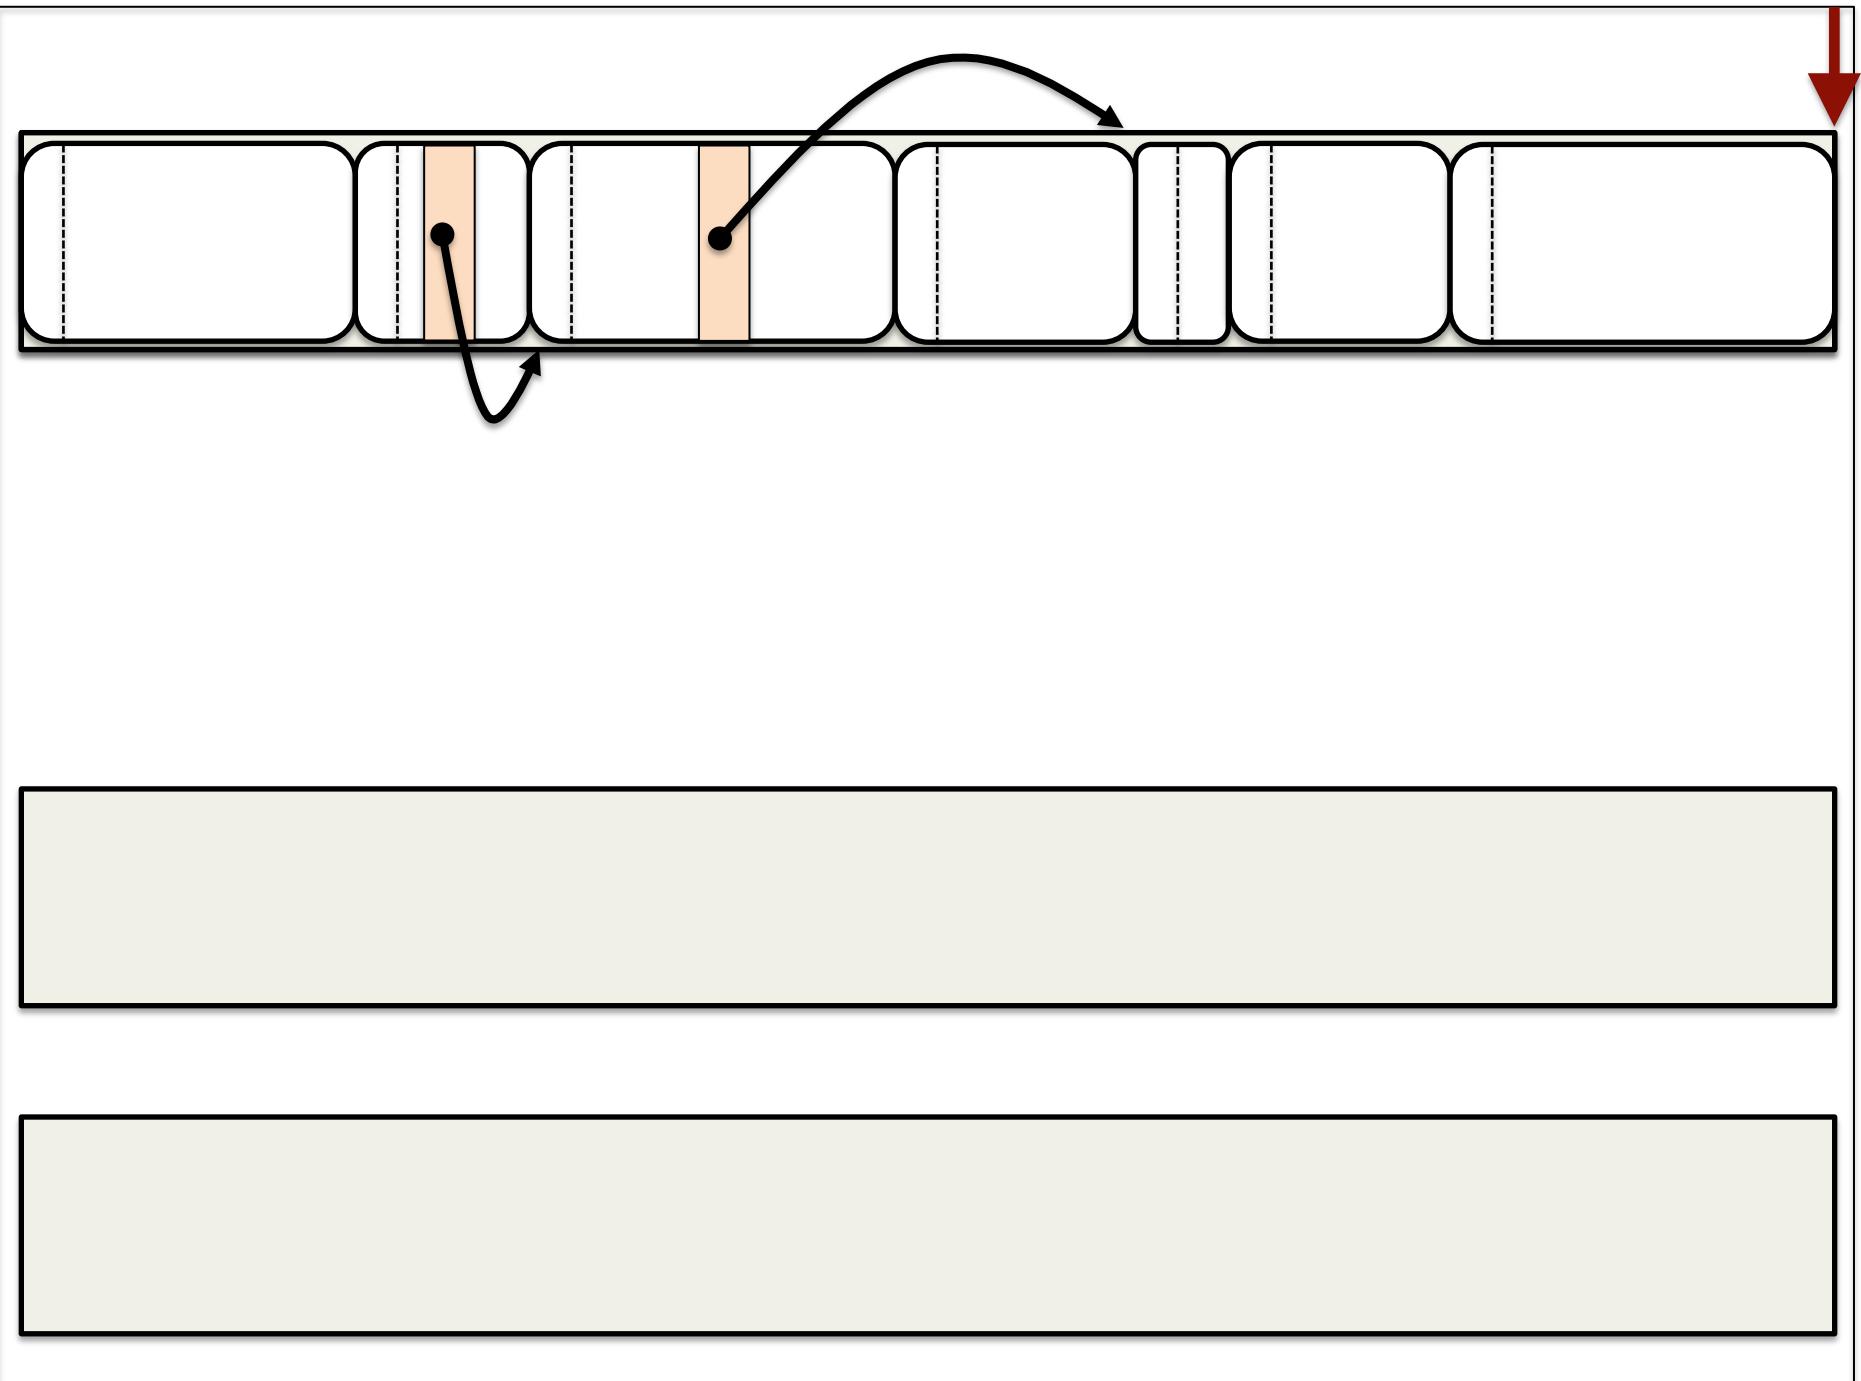


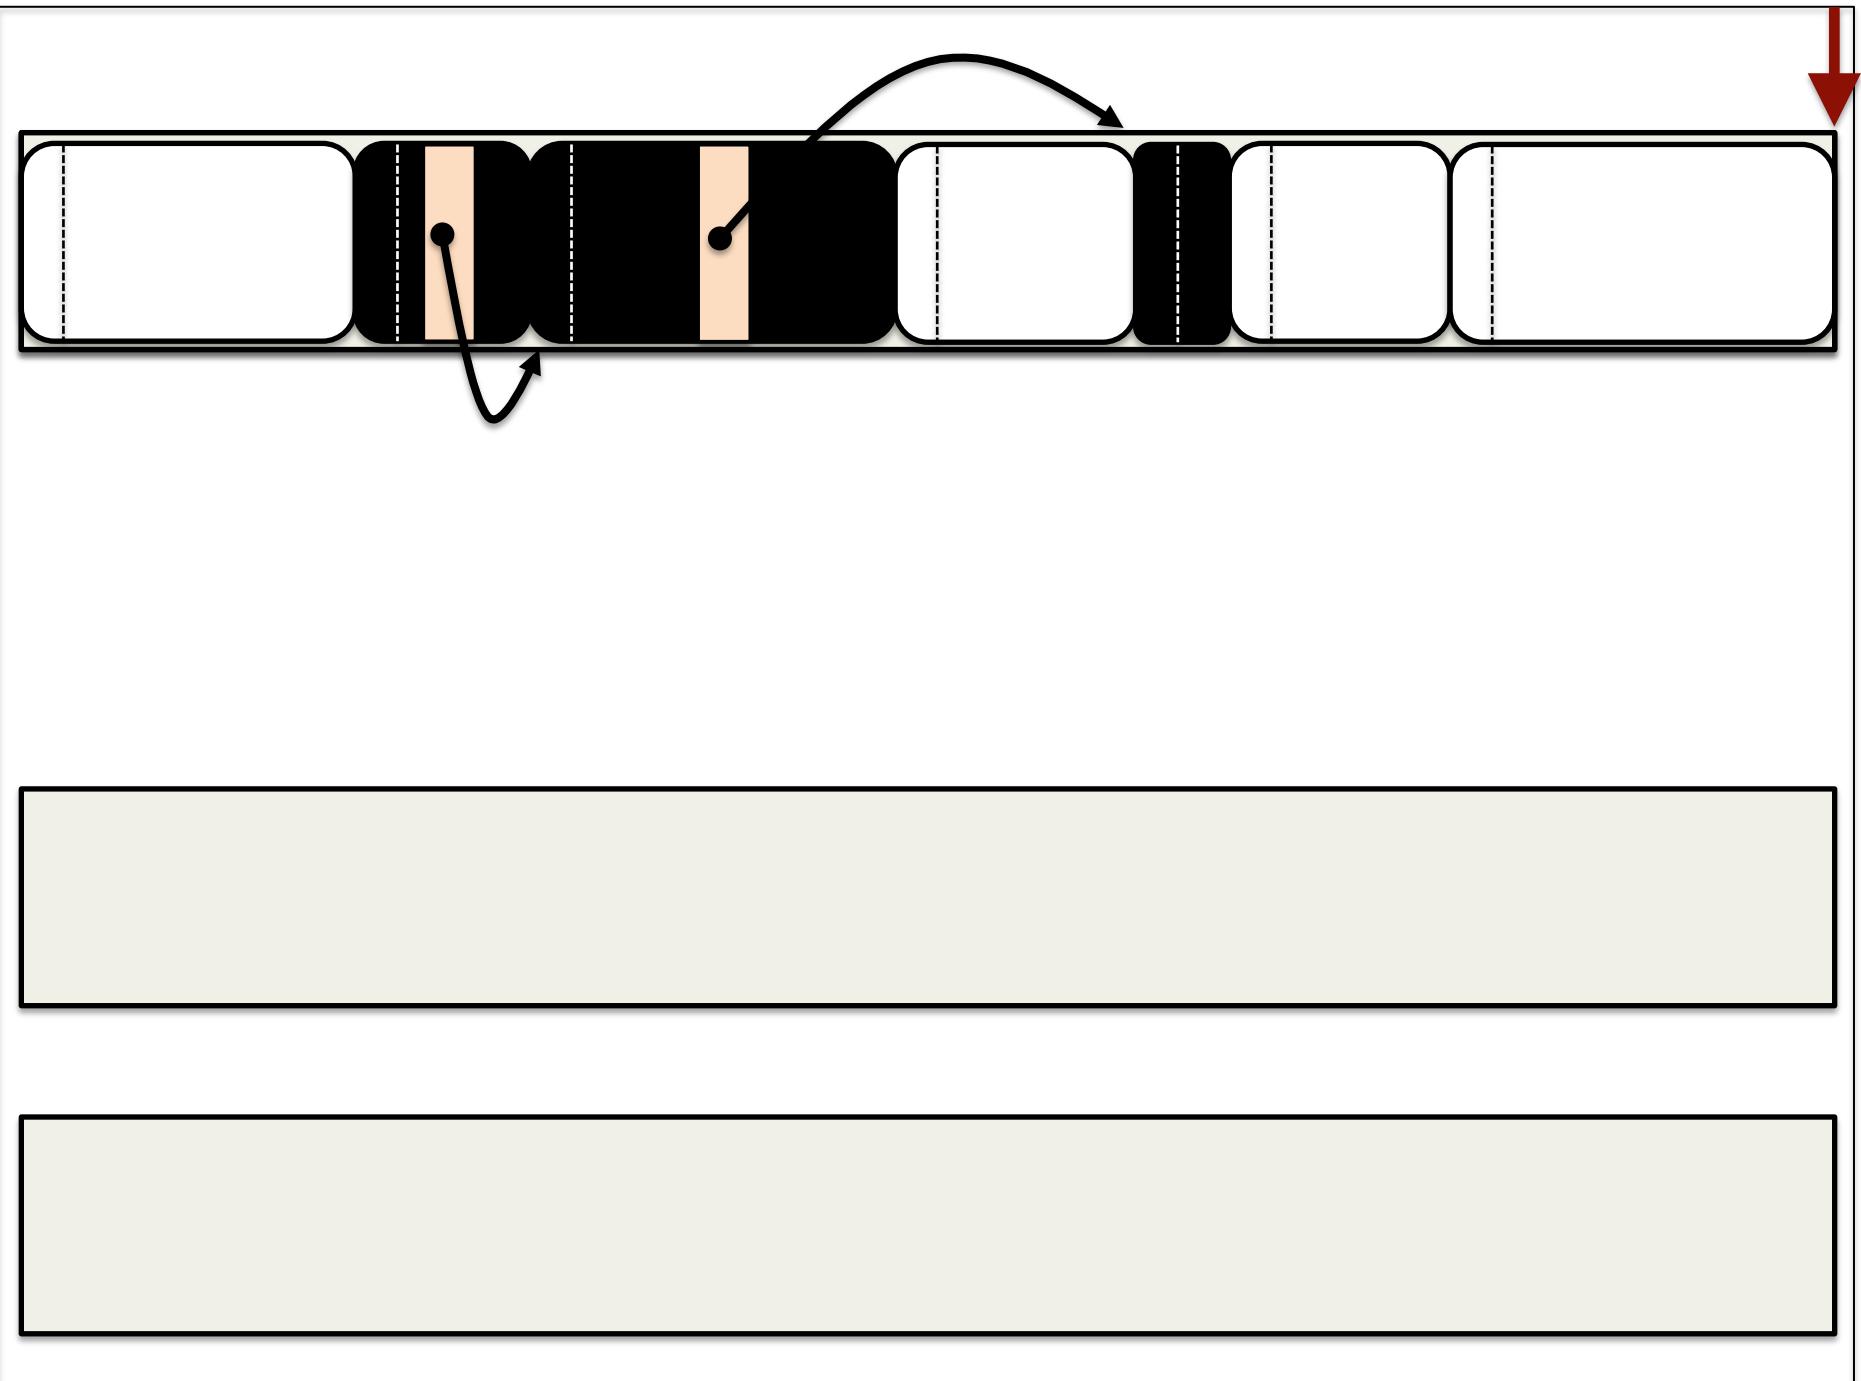


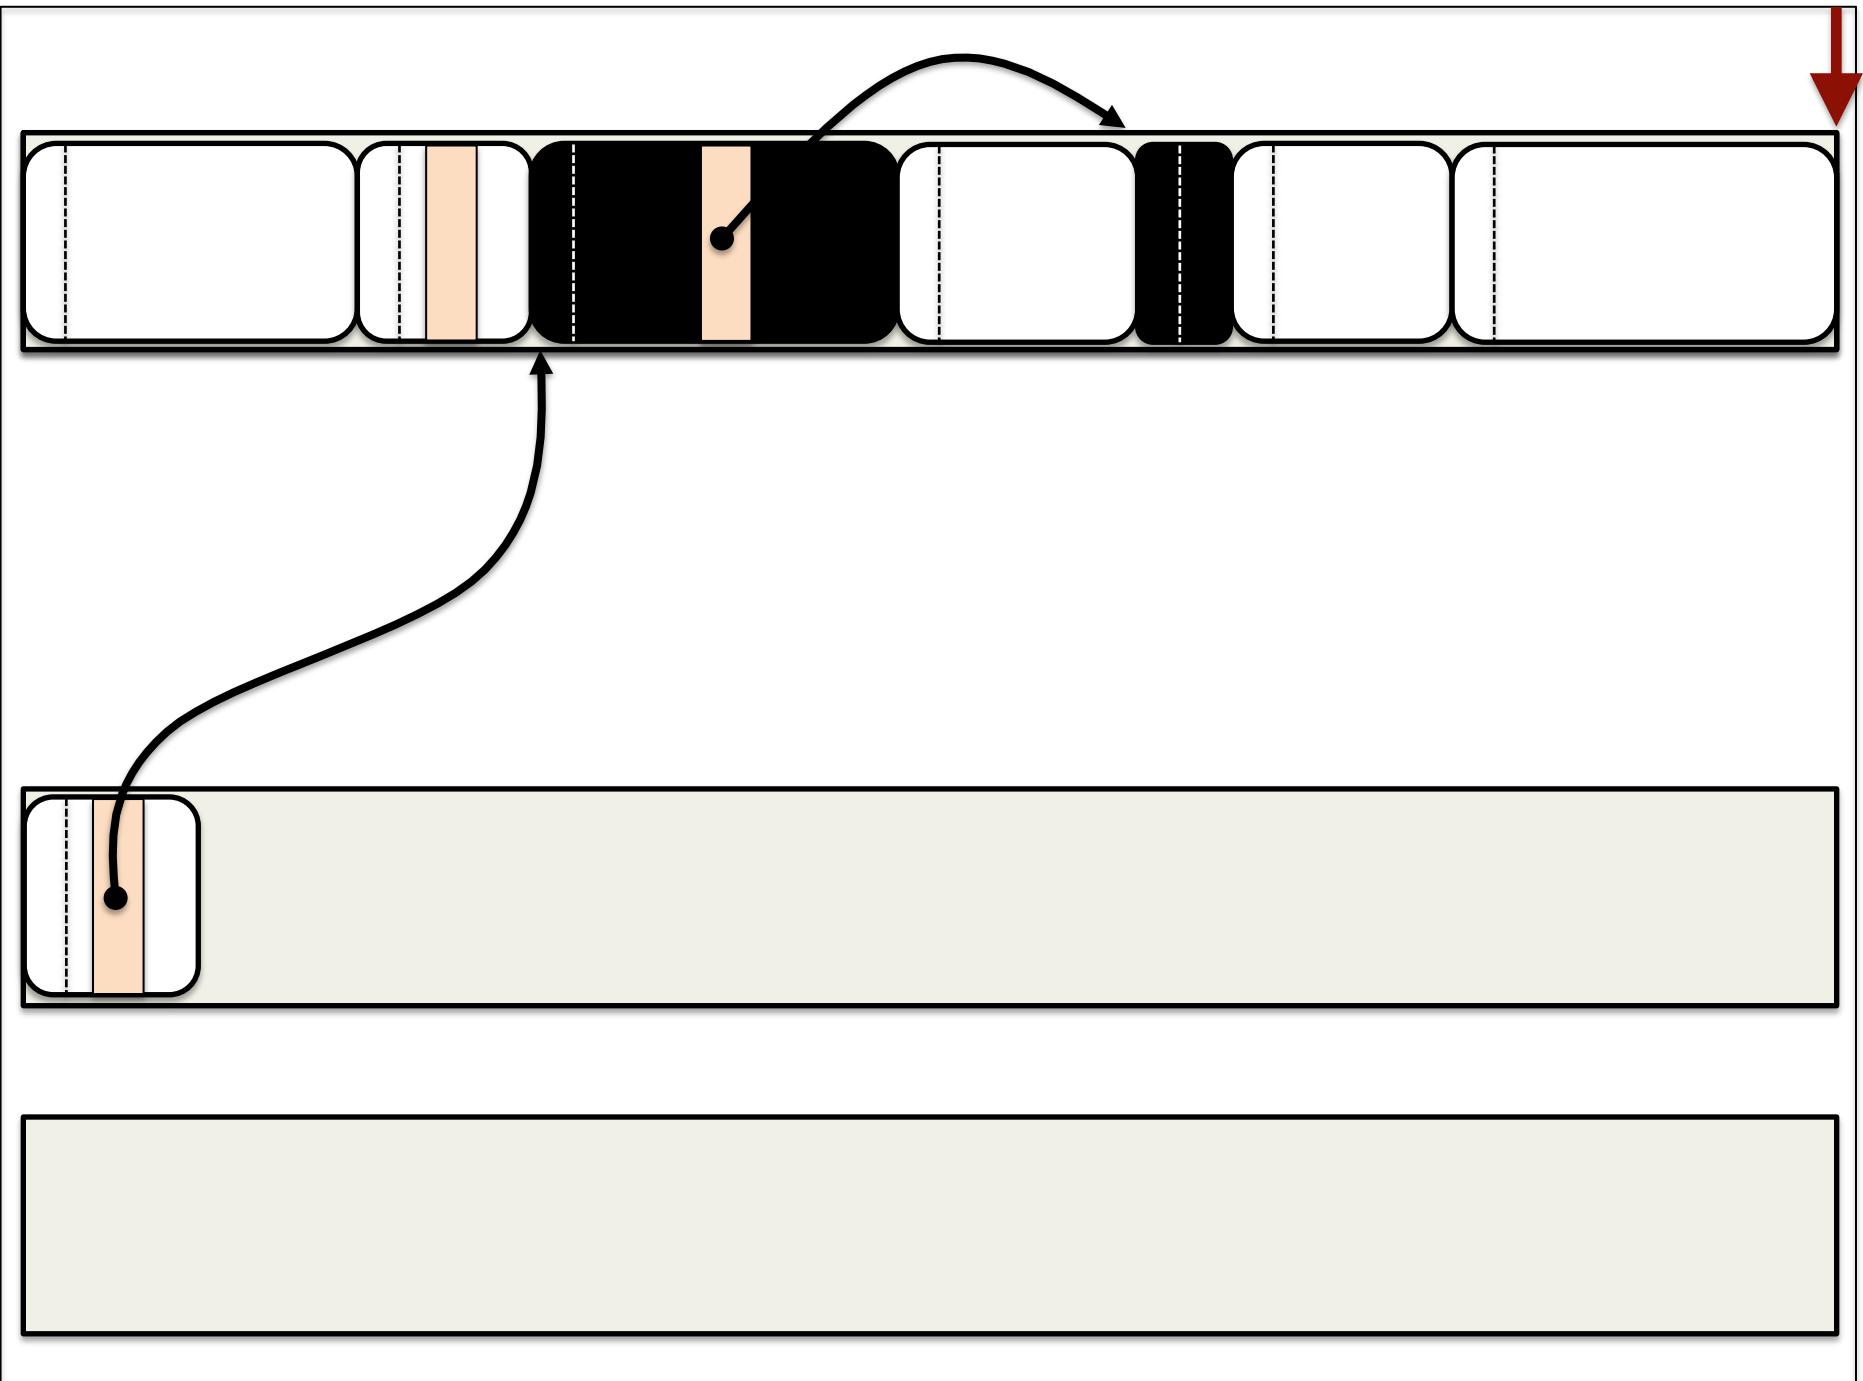


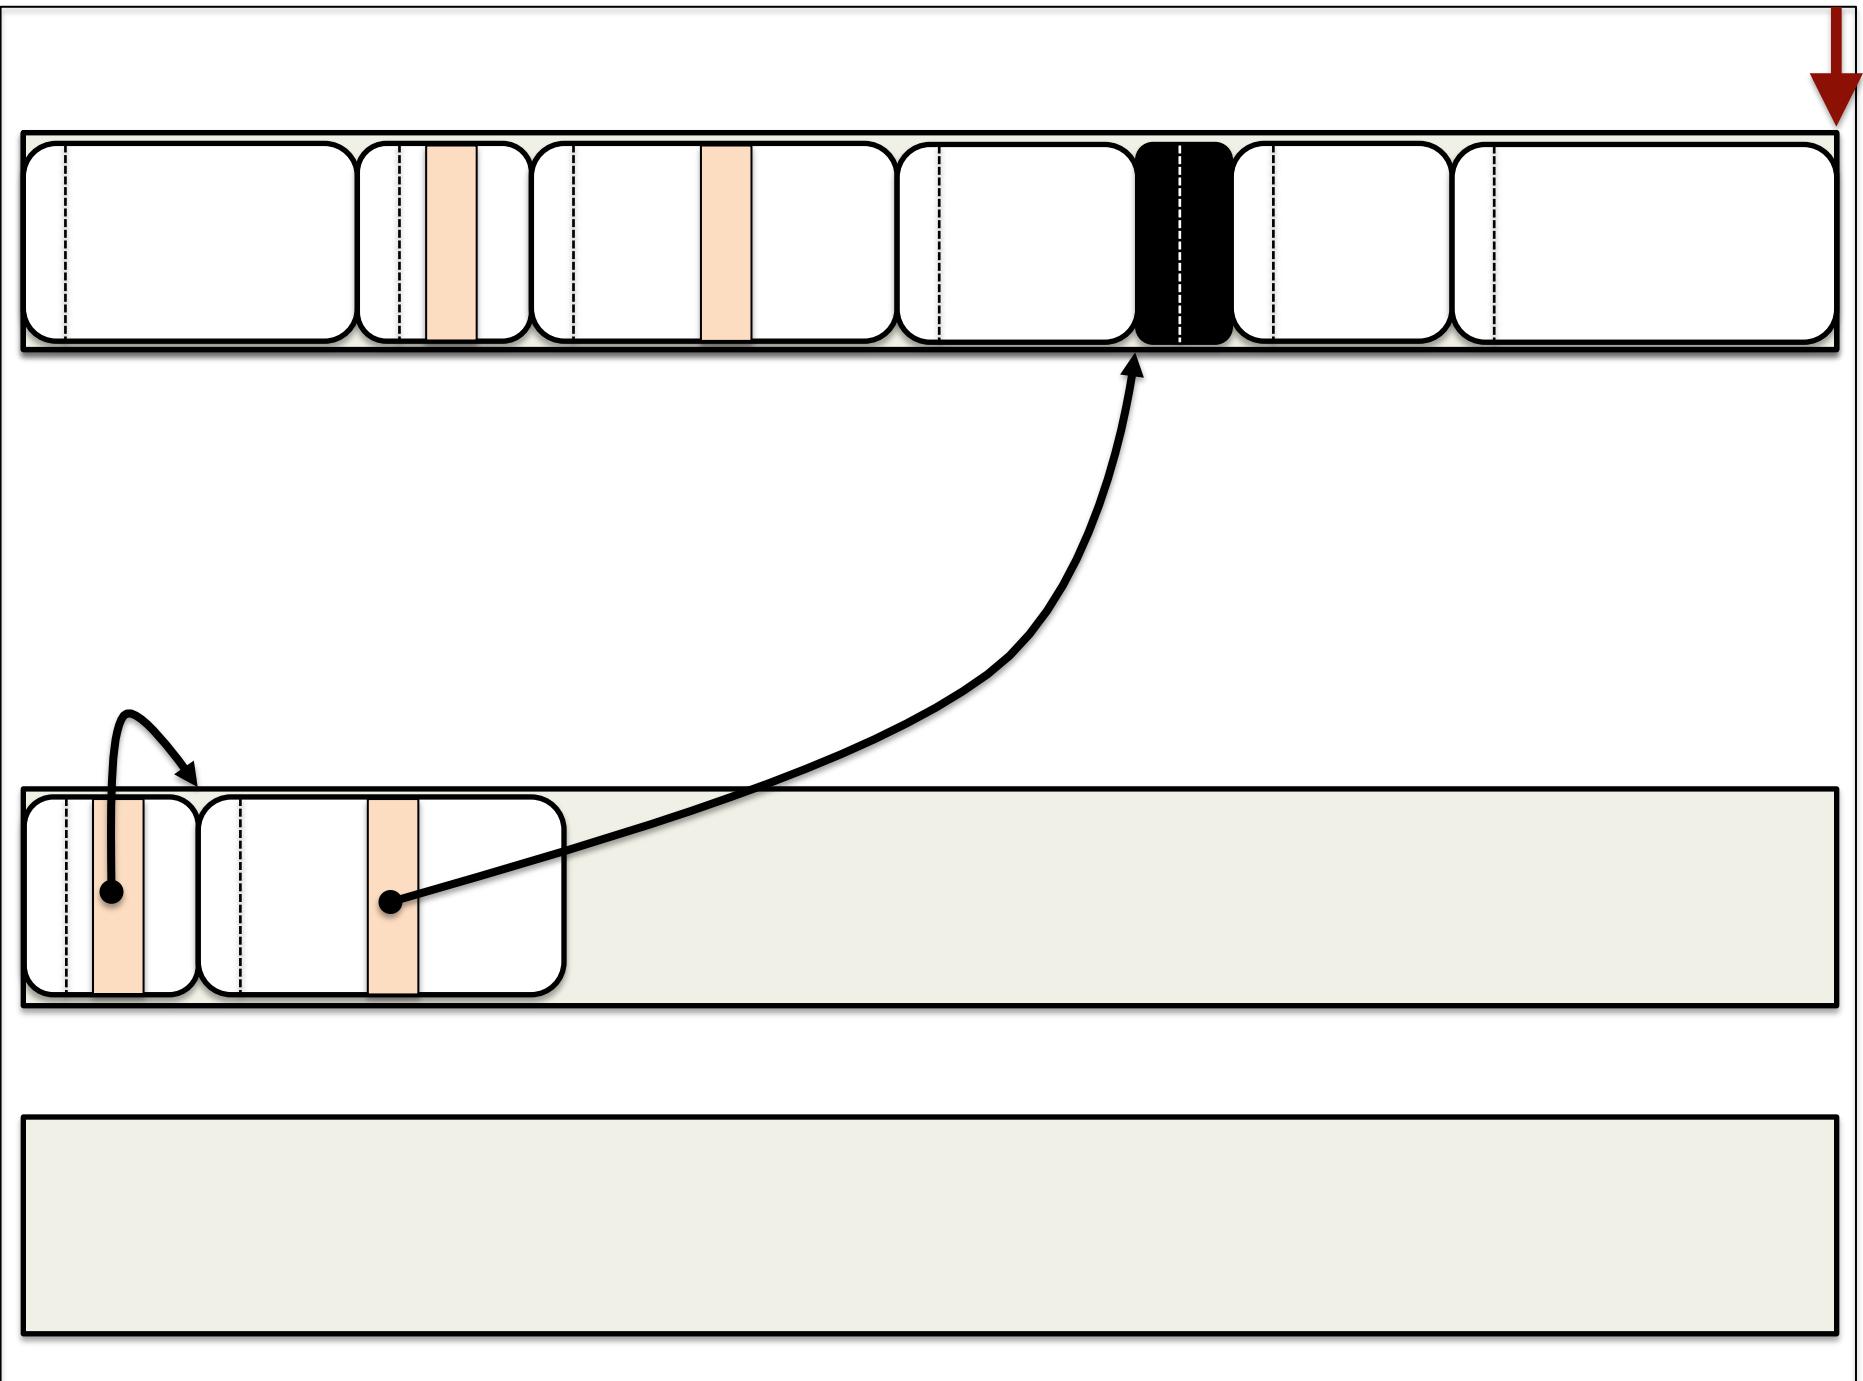


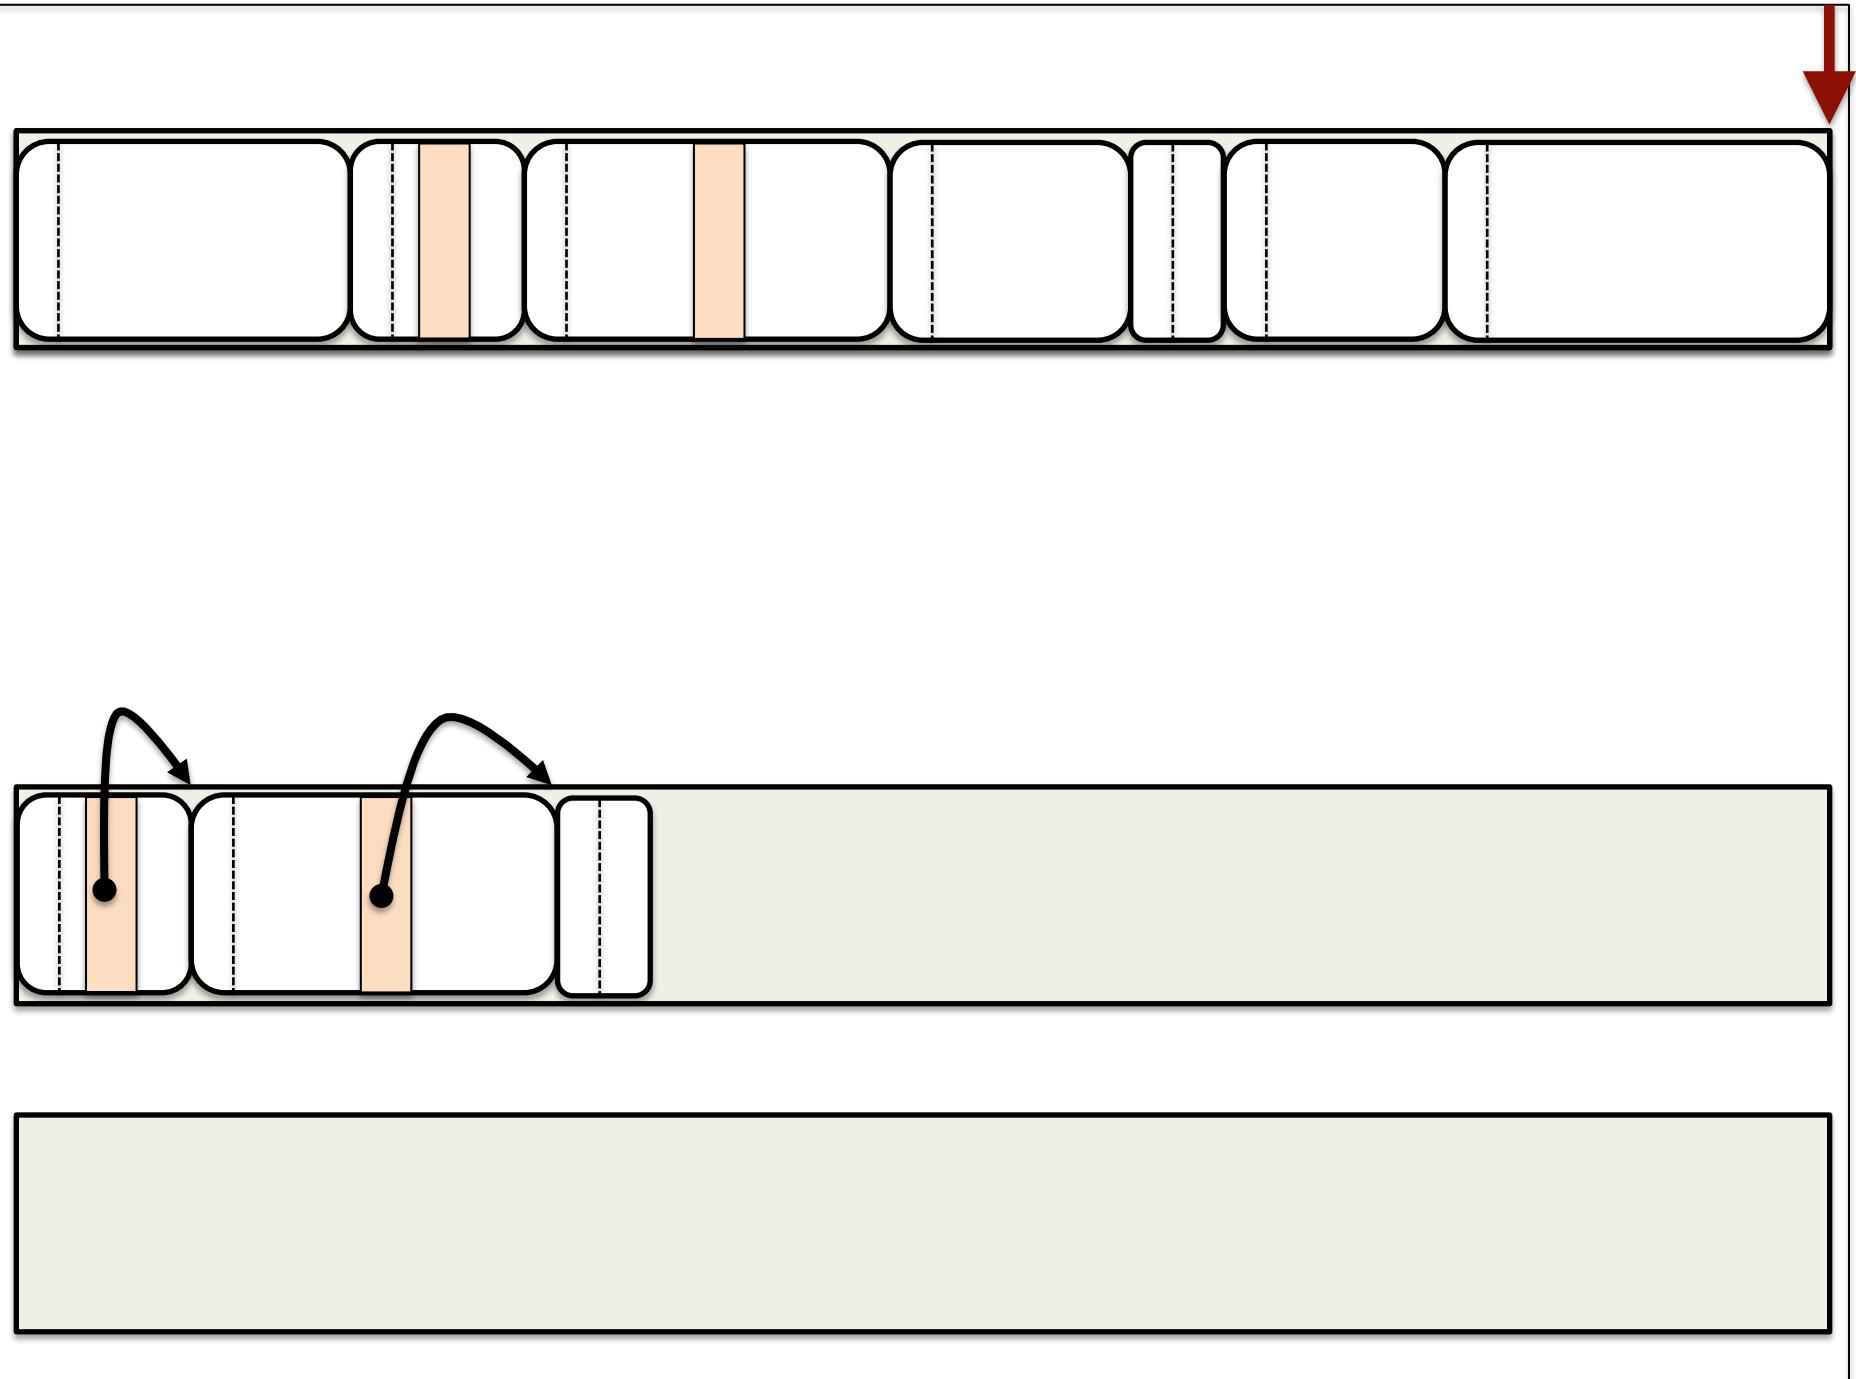


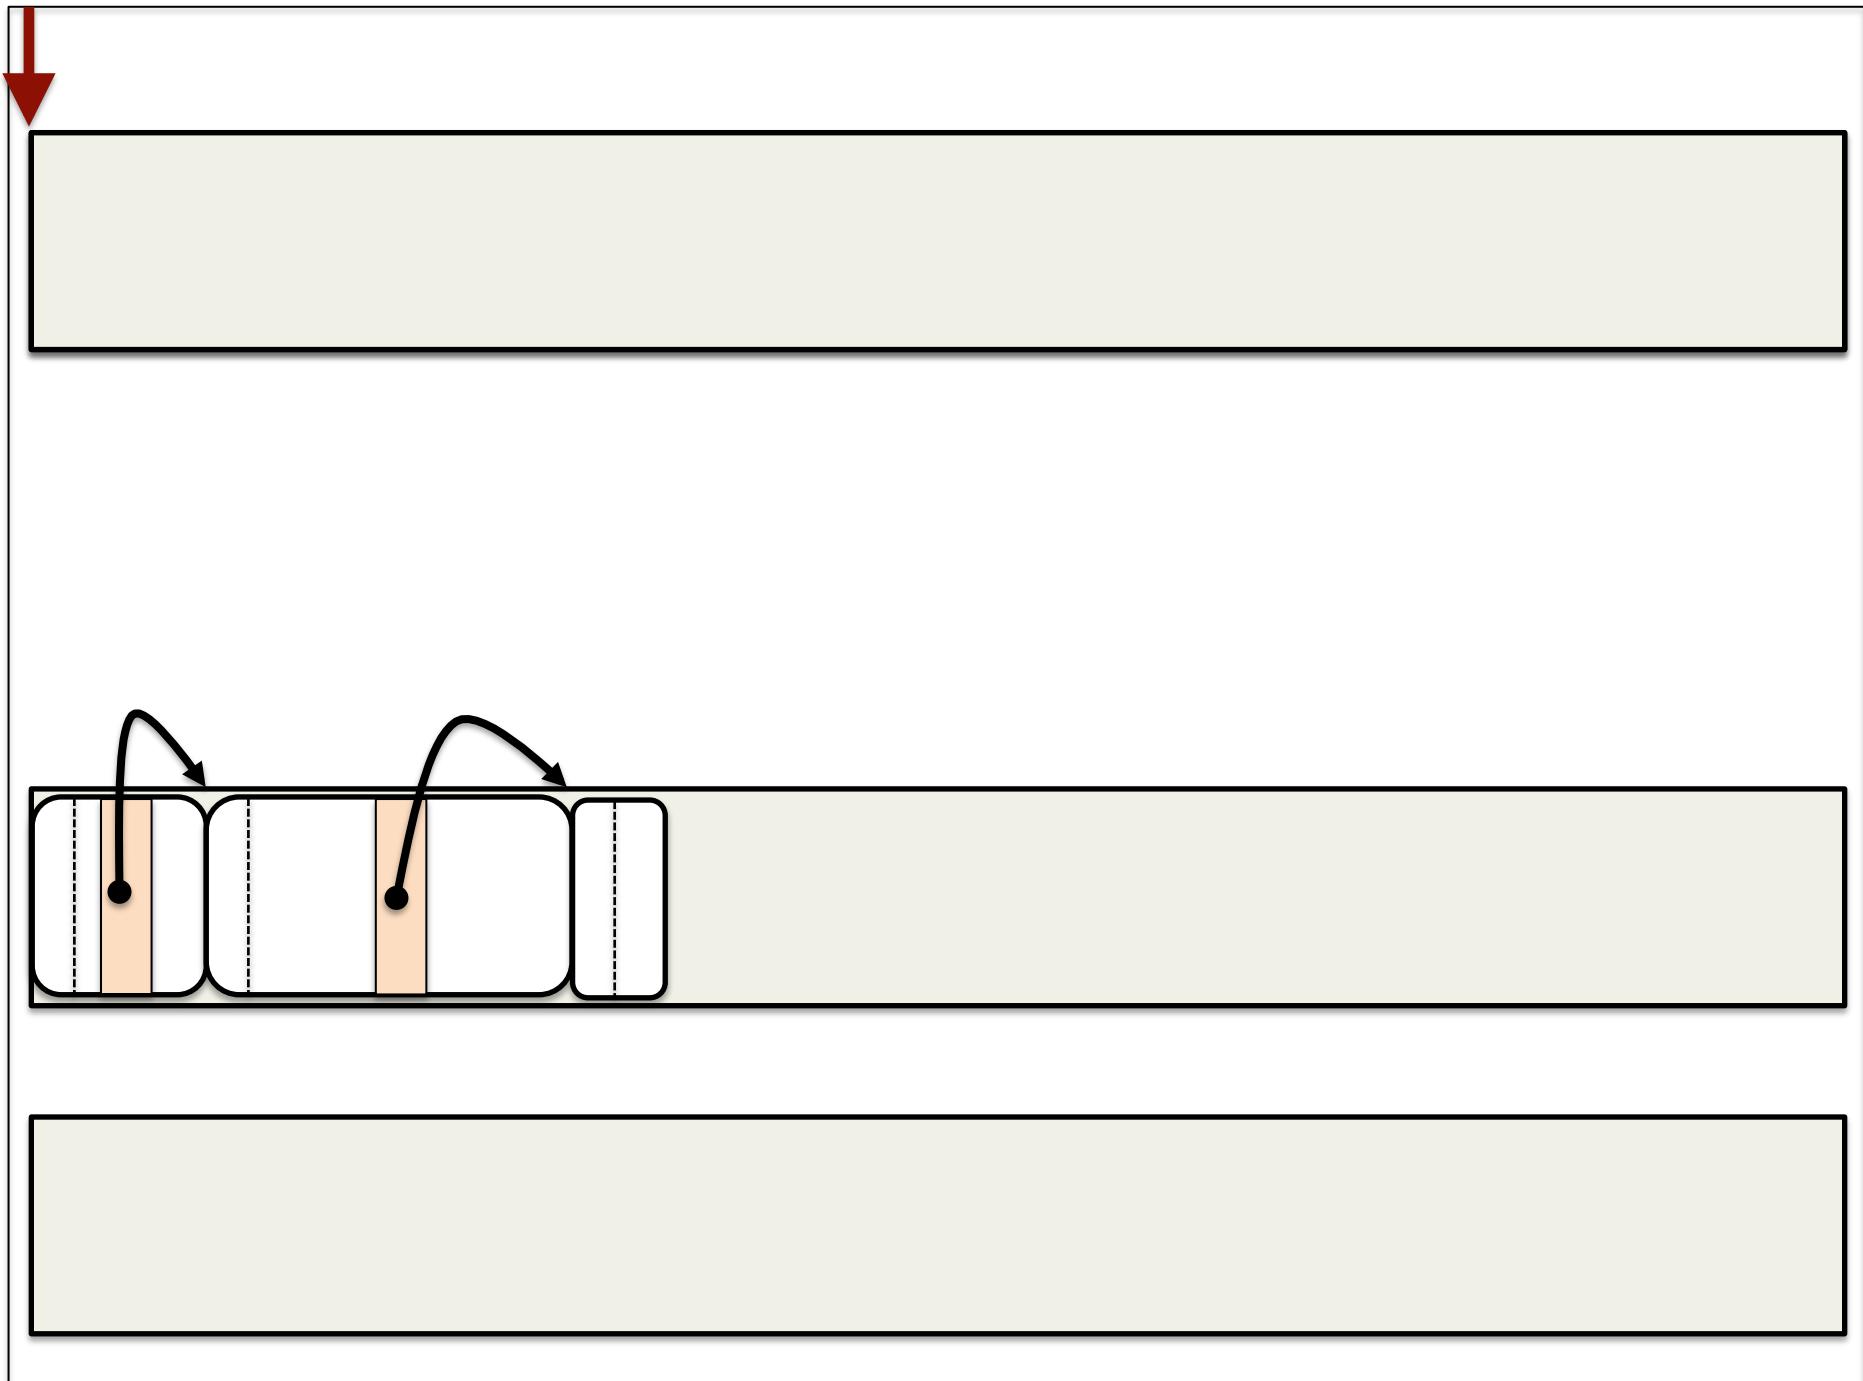






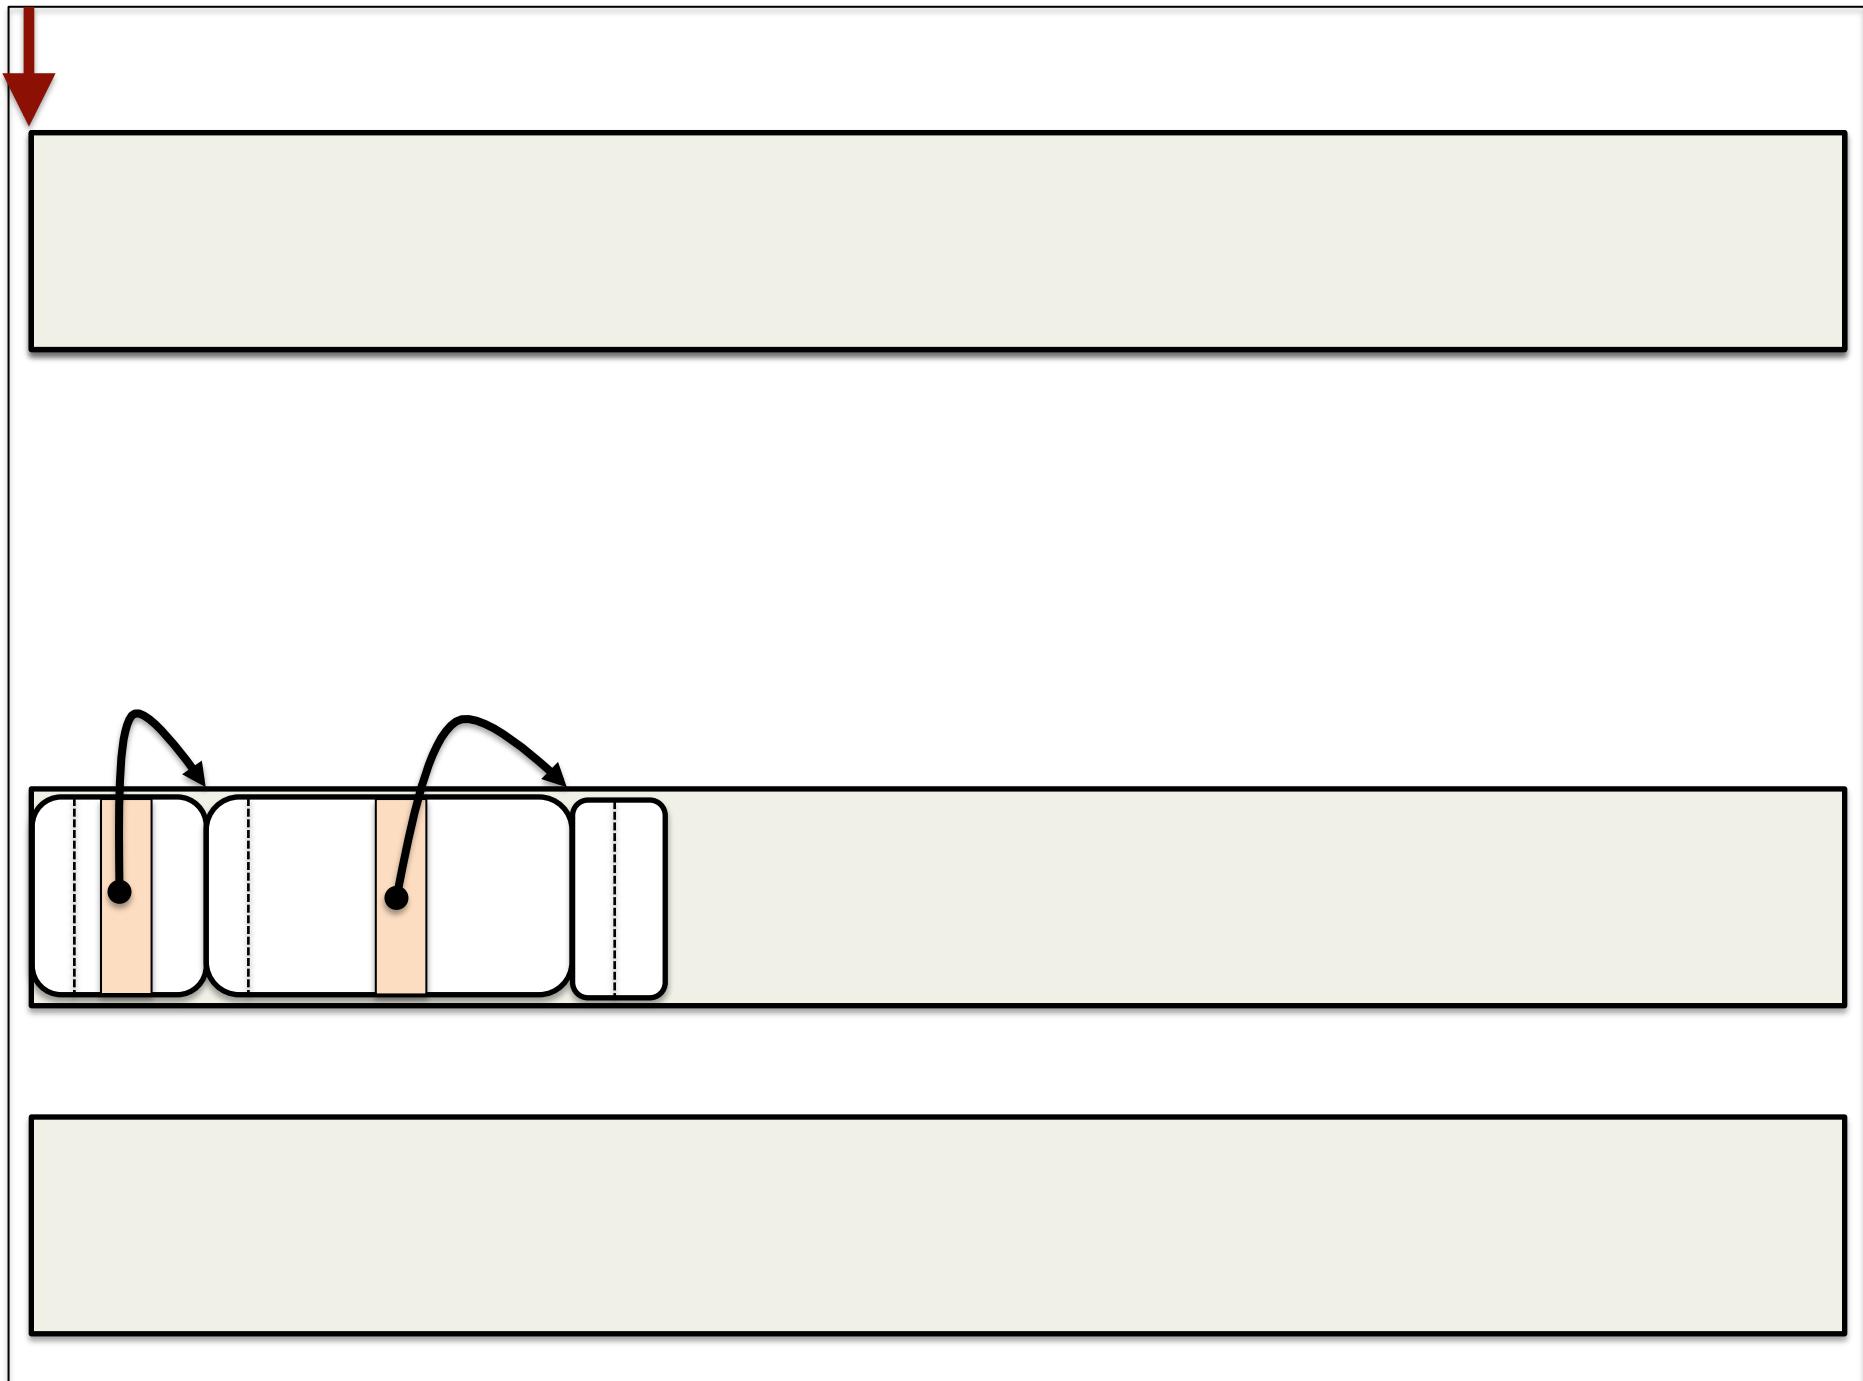


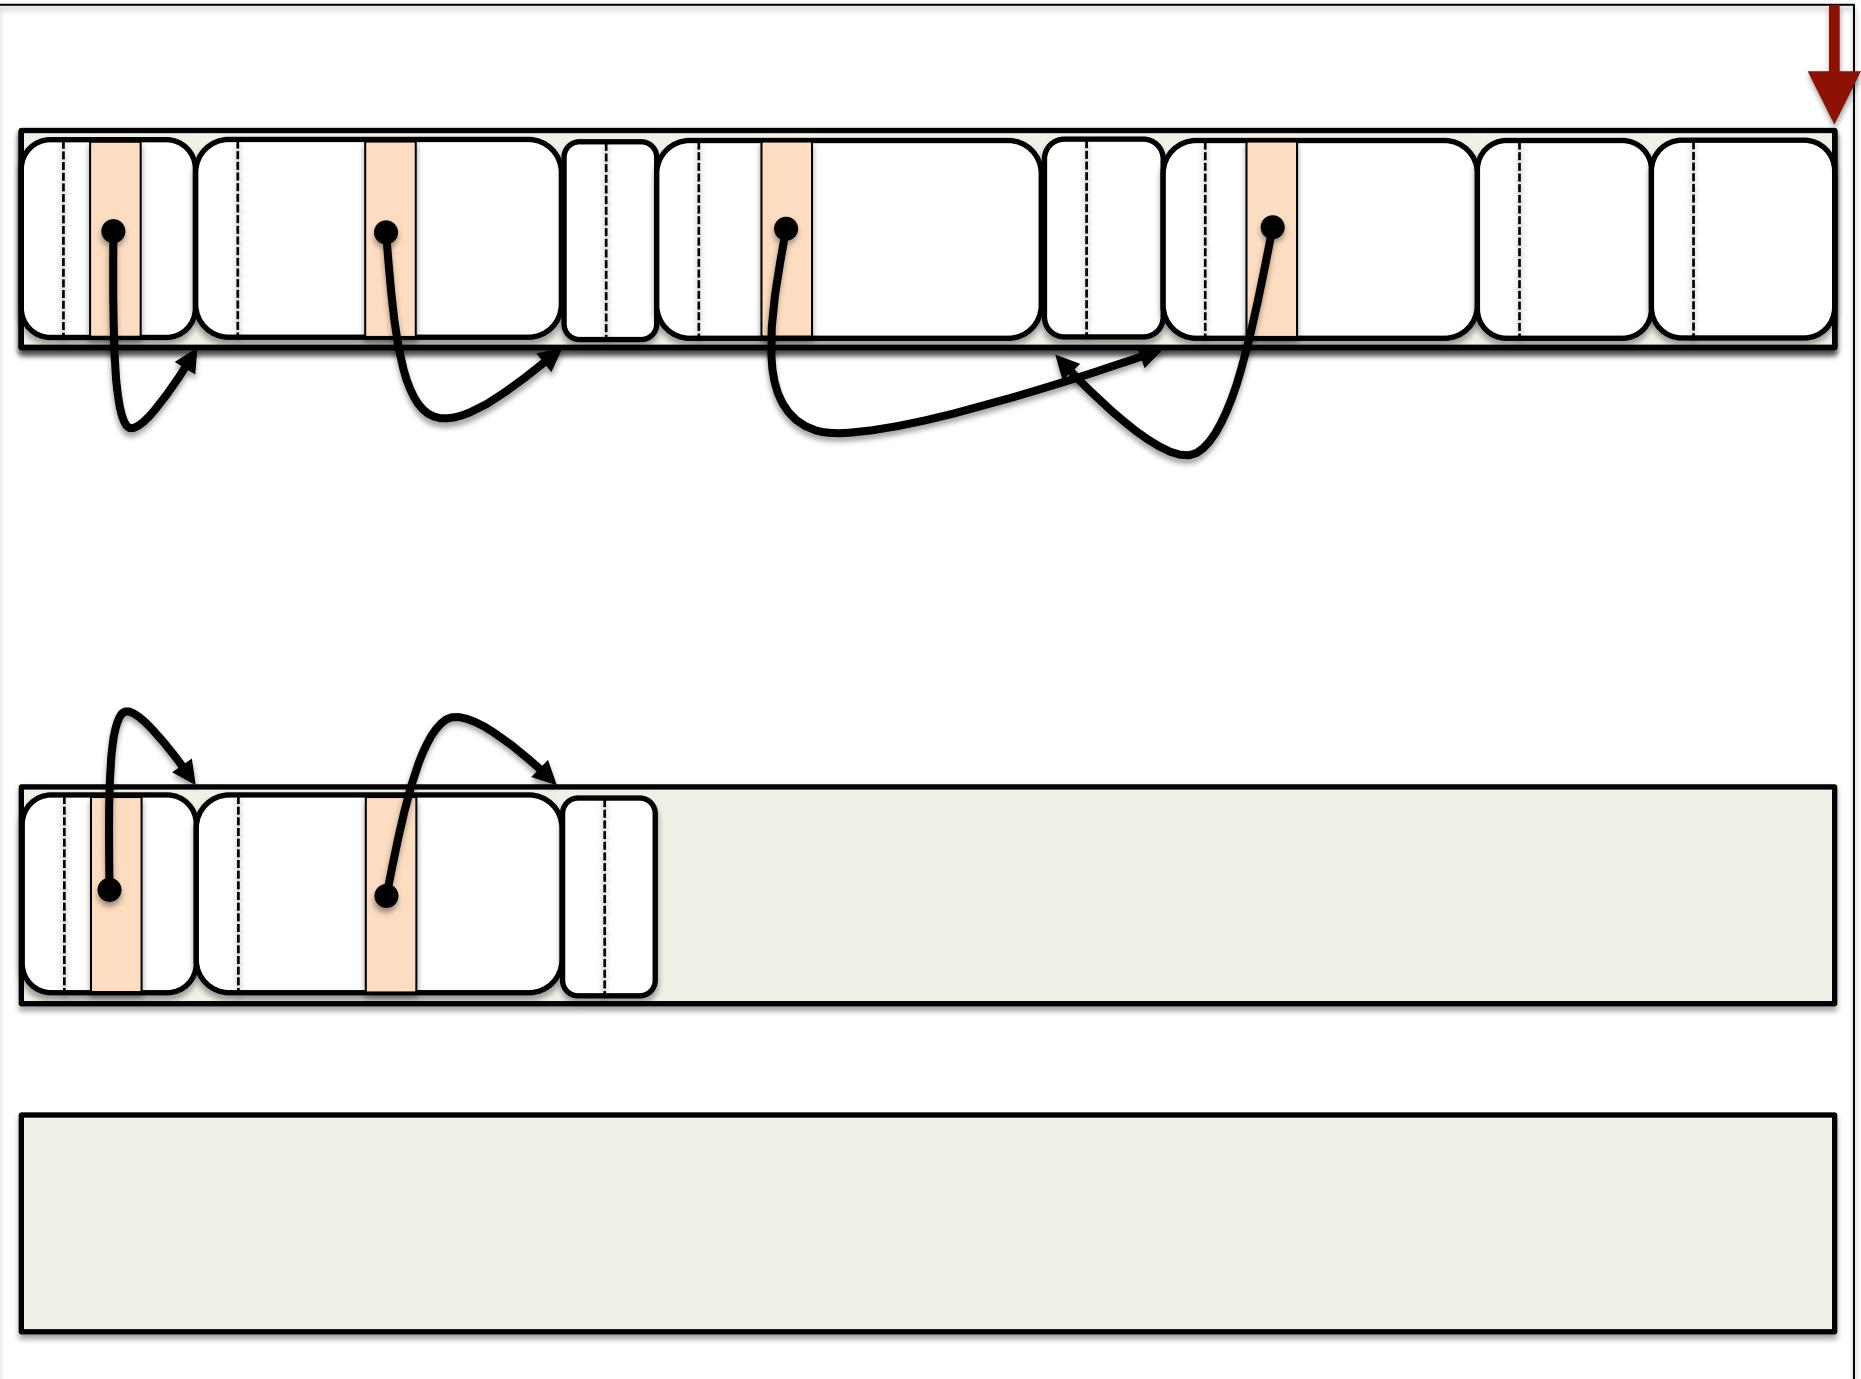


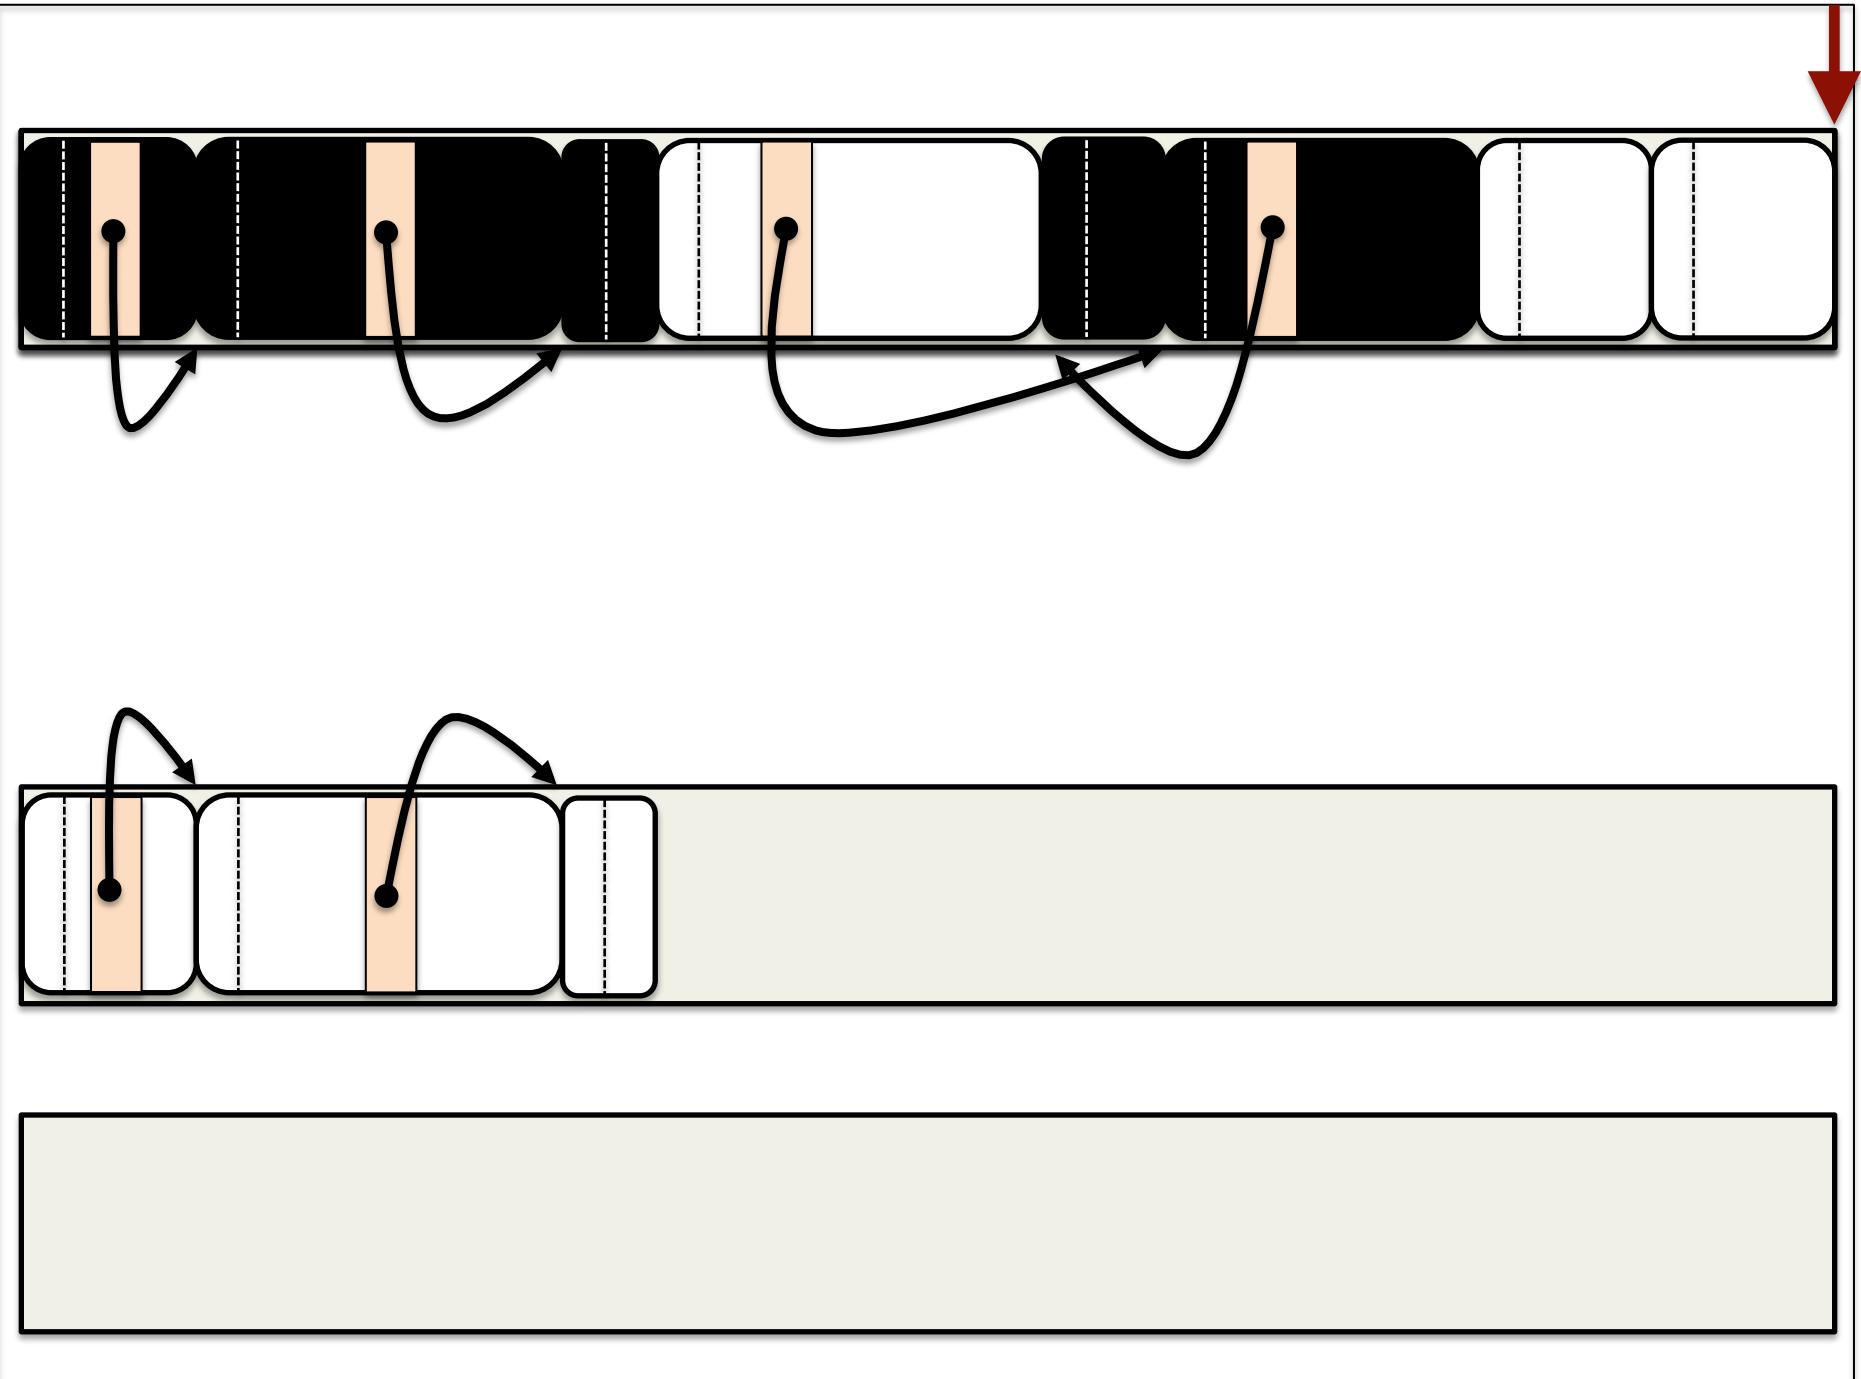


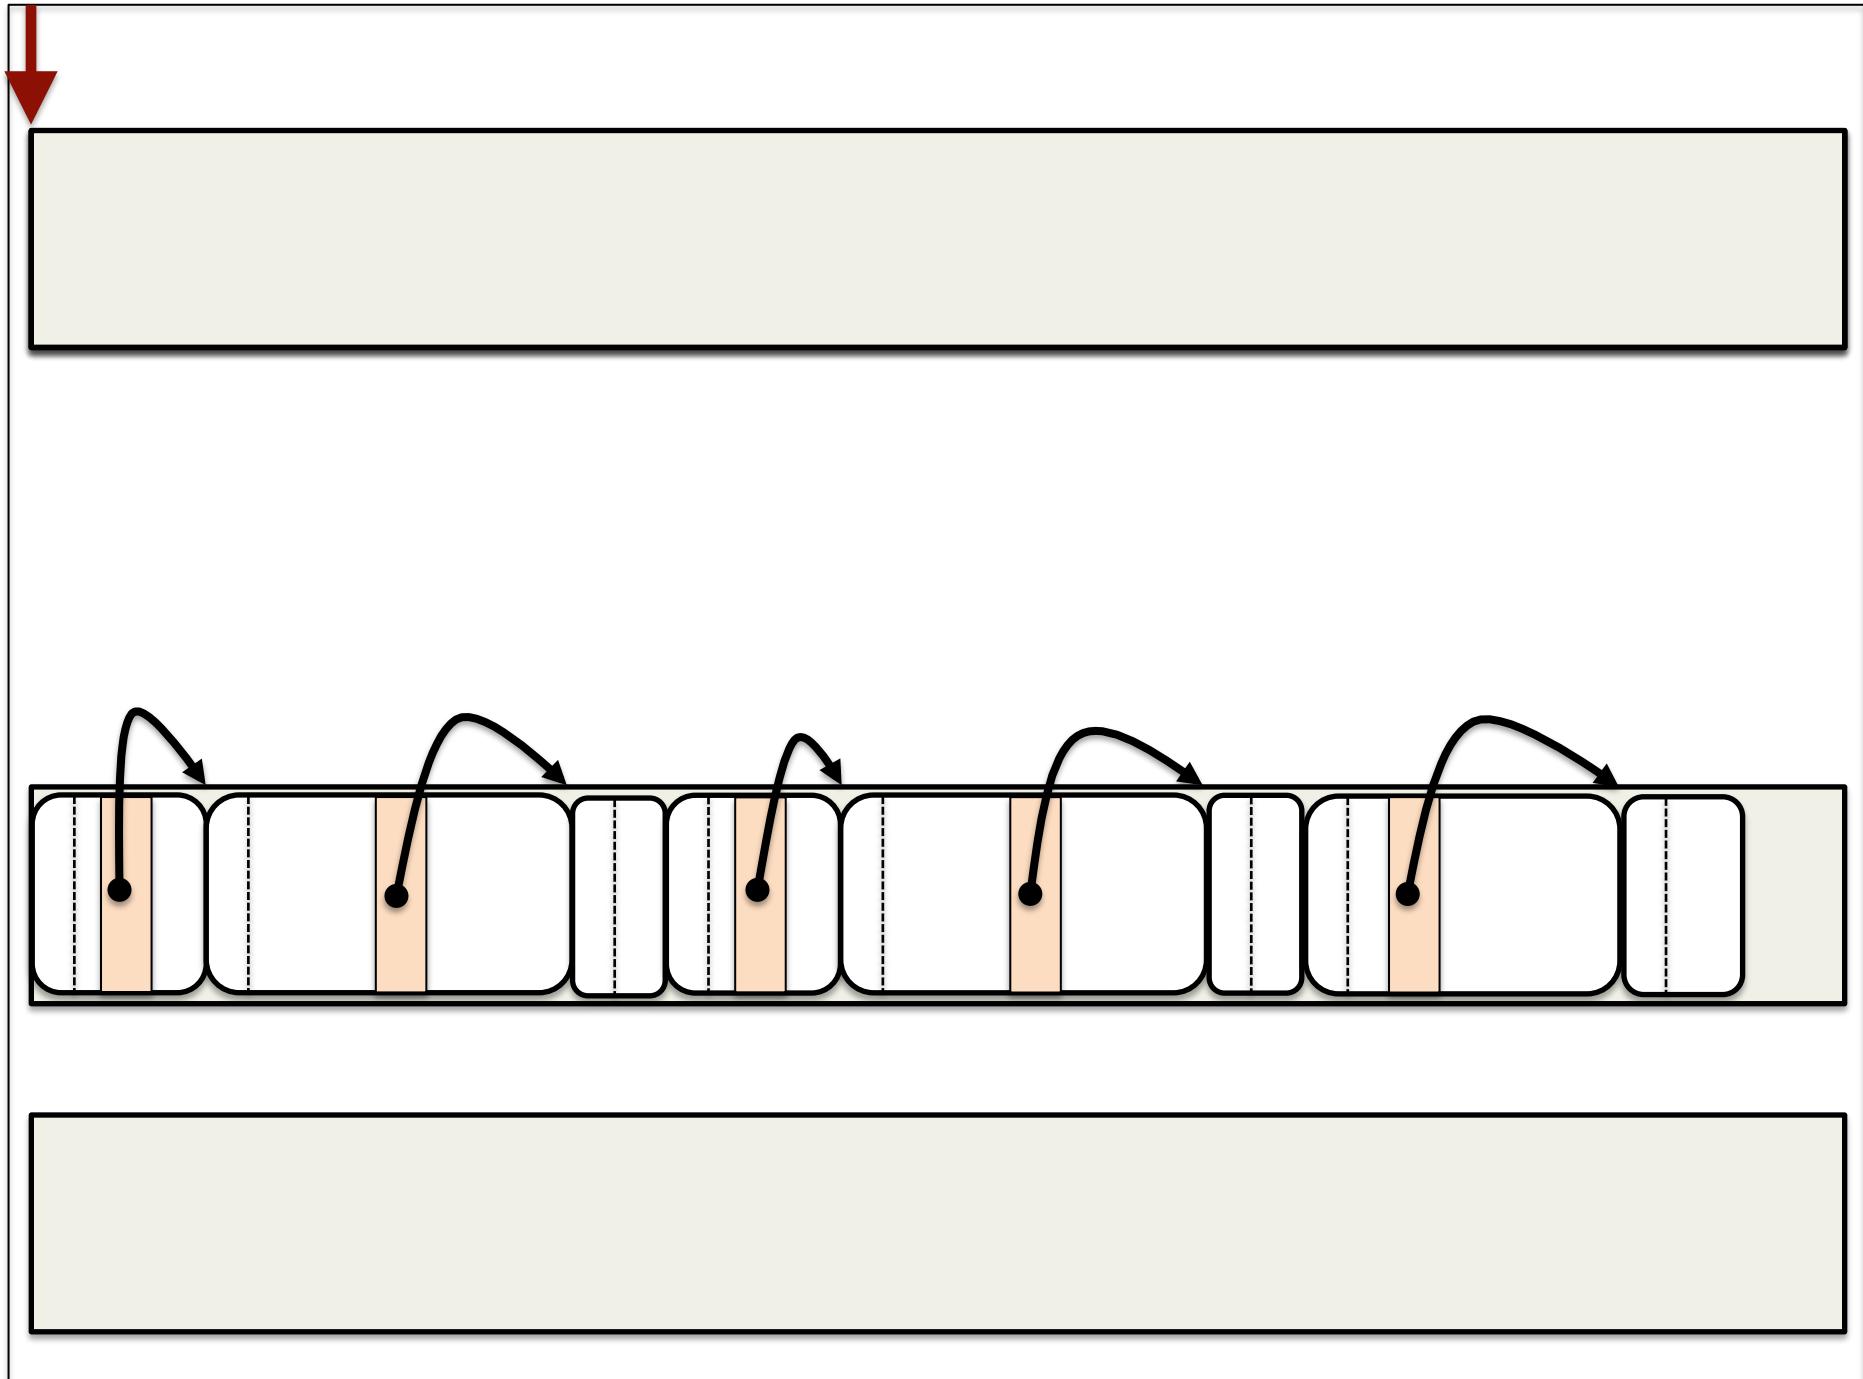
Minor Collections

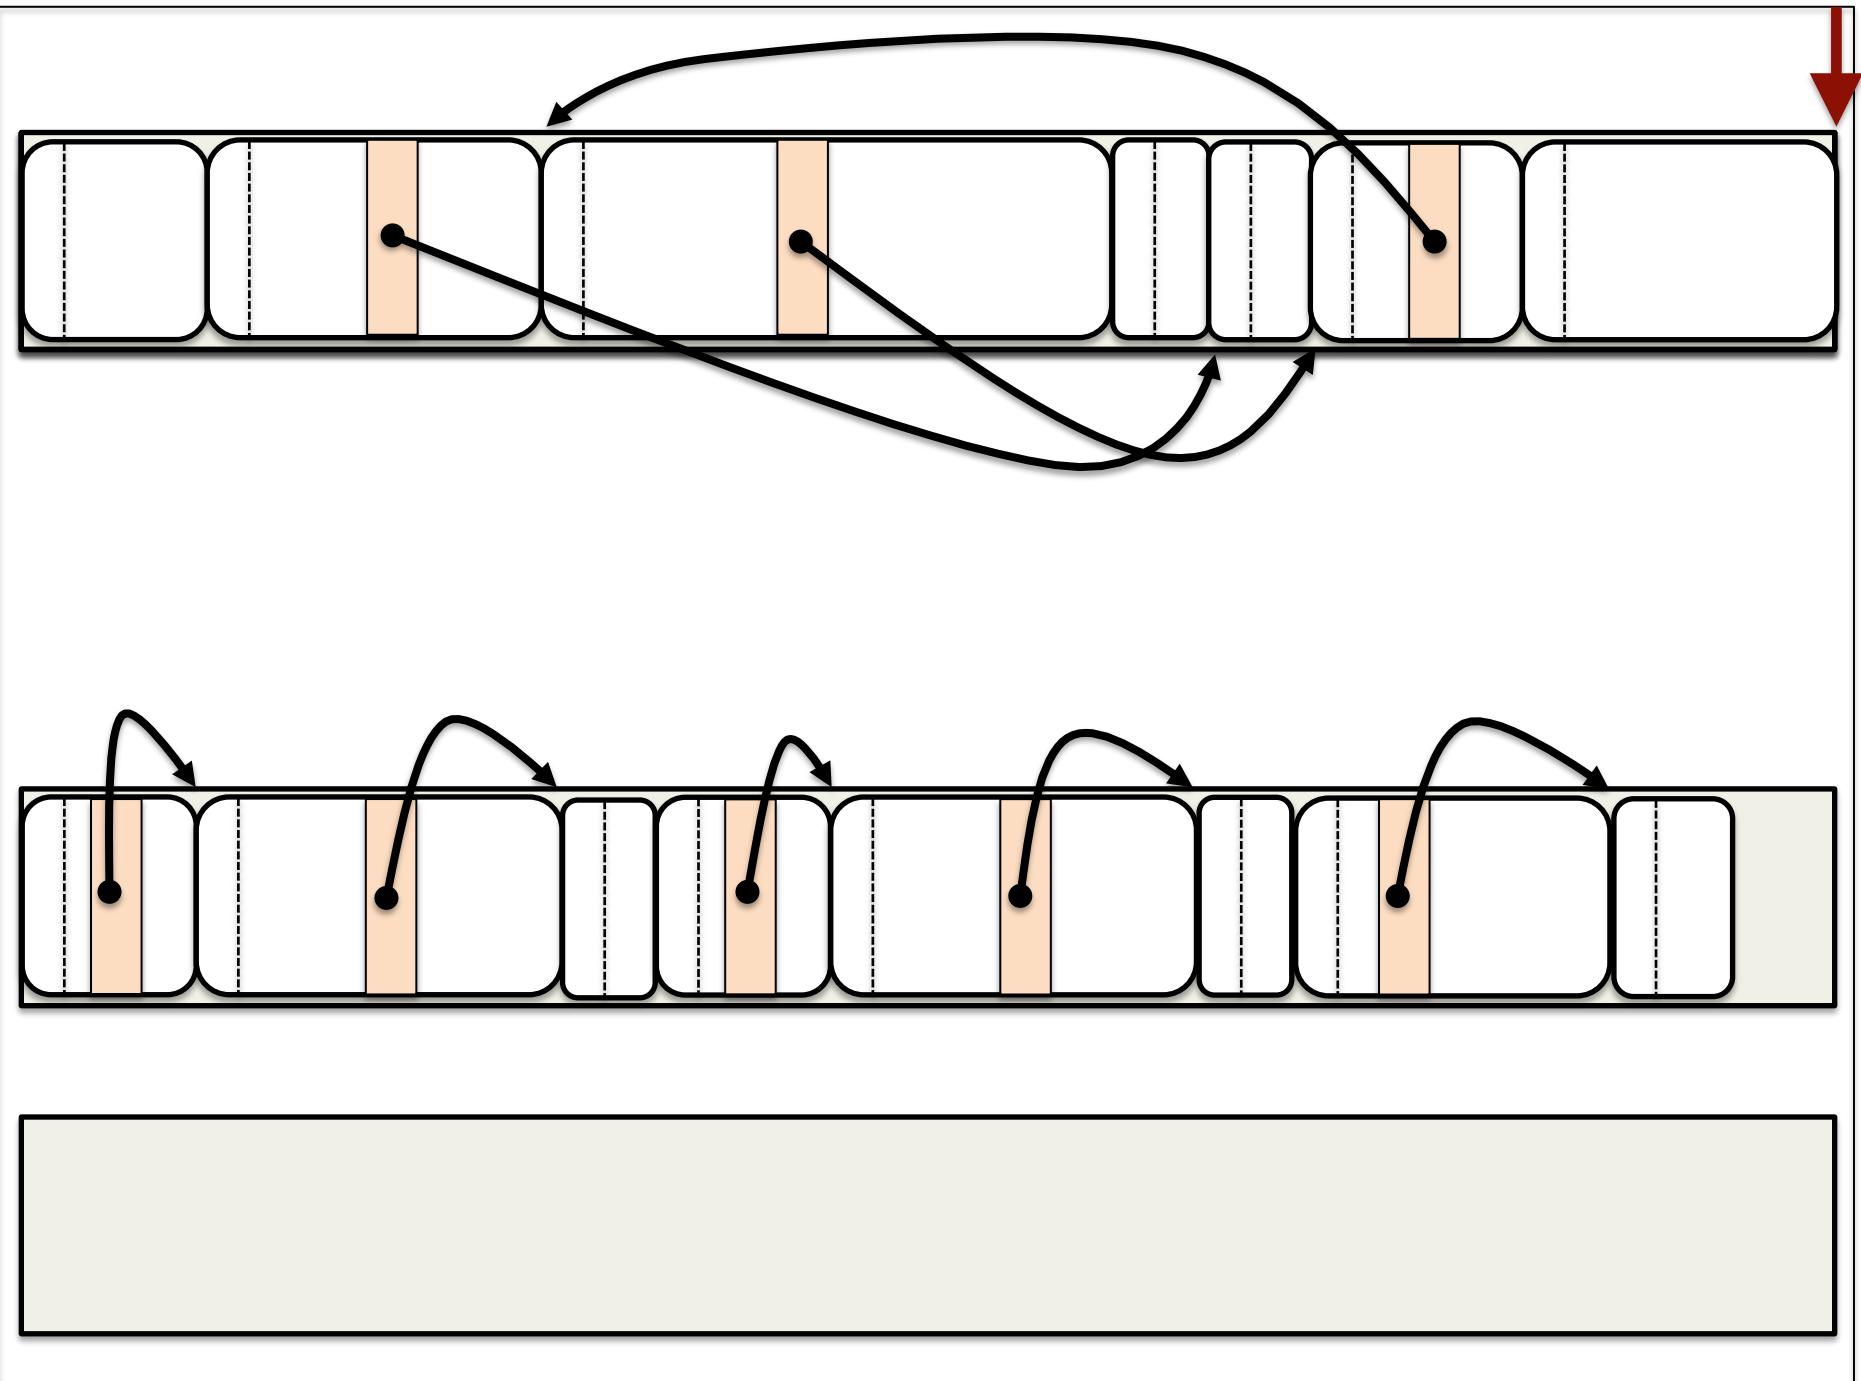
- Minor collections focus only on the nursery
 - Mature space is only used as a destination
- Promote all live objects to the mature space
- Reset bump pointer in the nursery
 - New allocation starts again in the same region

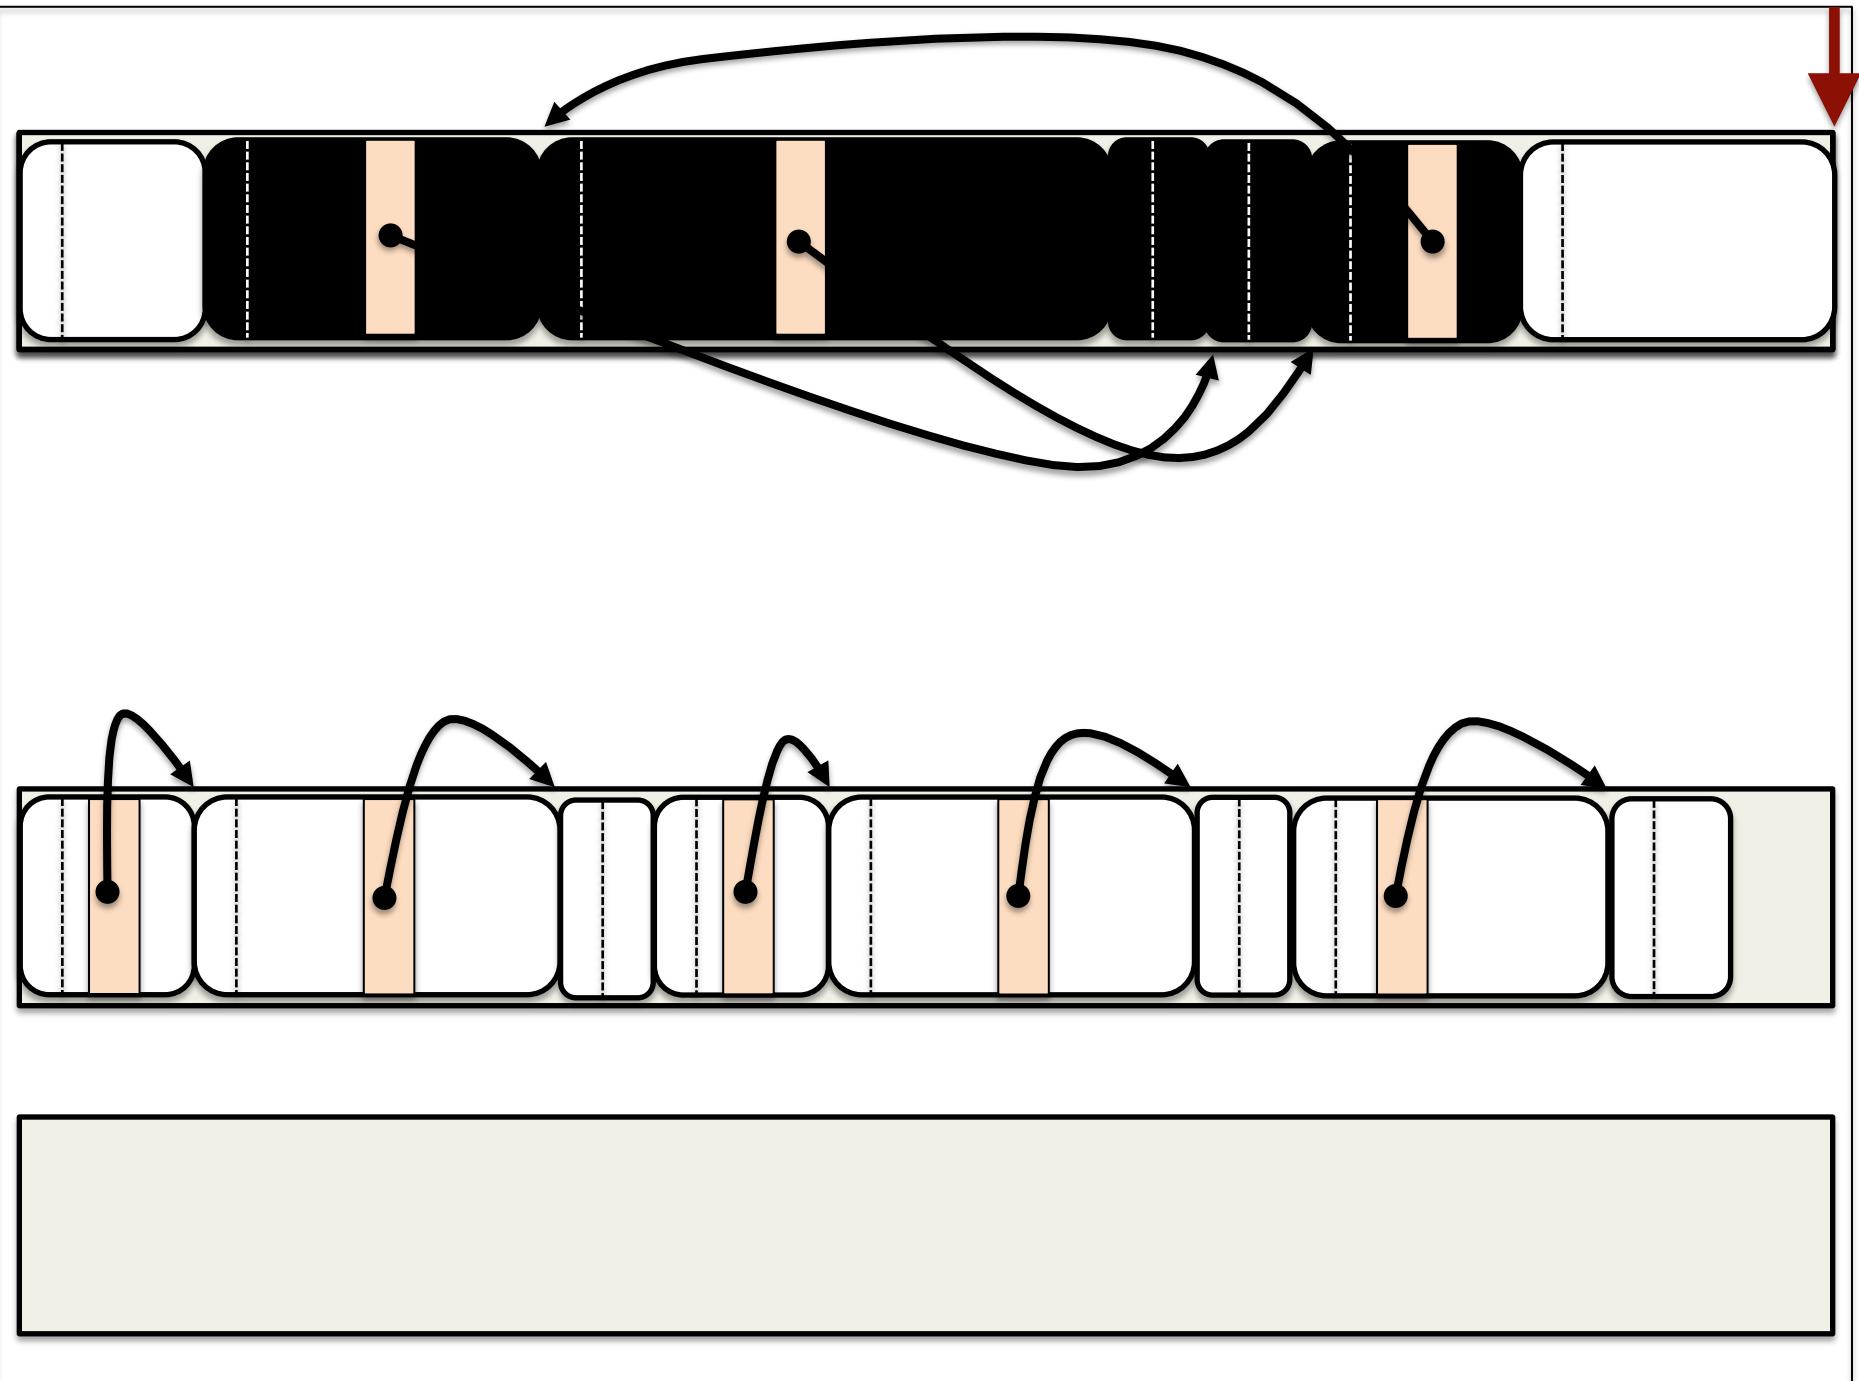


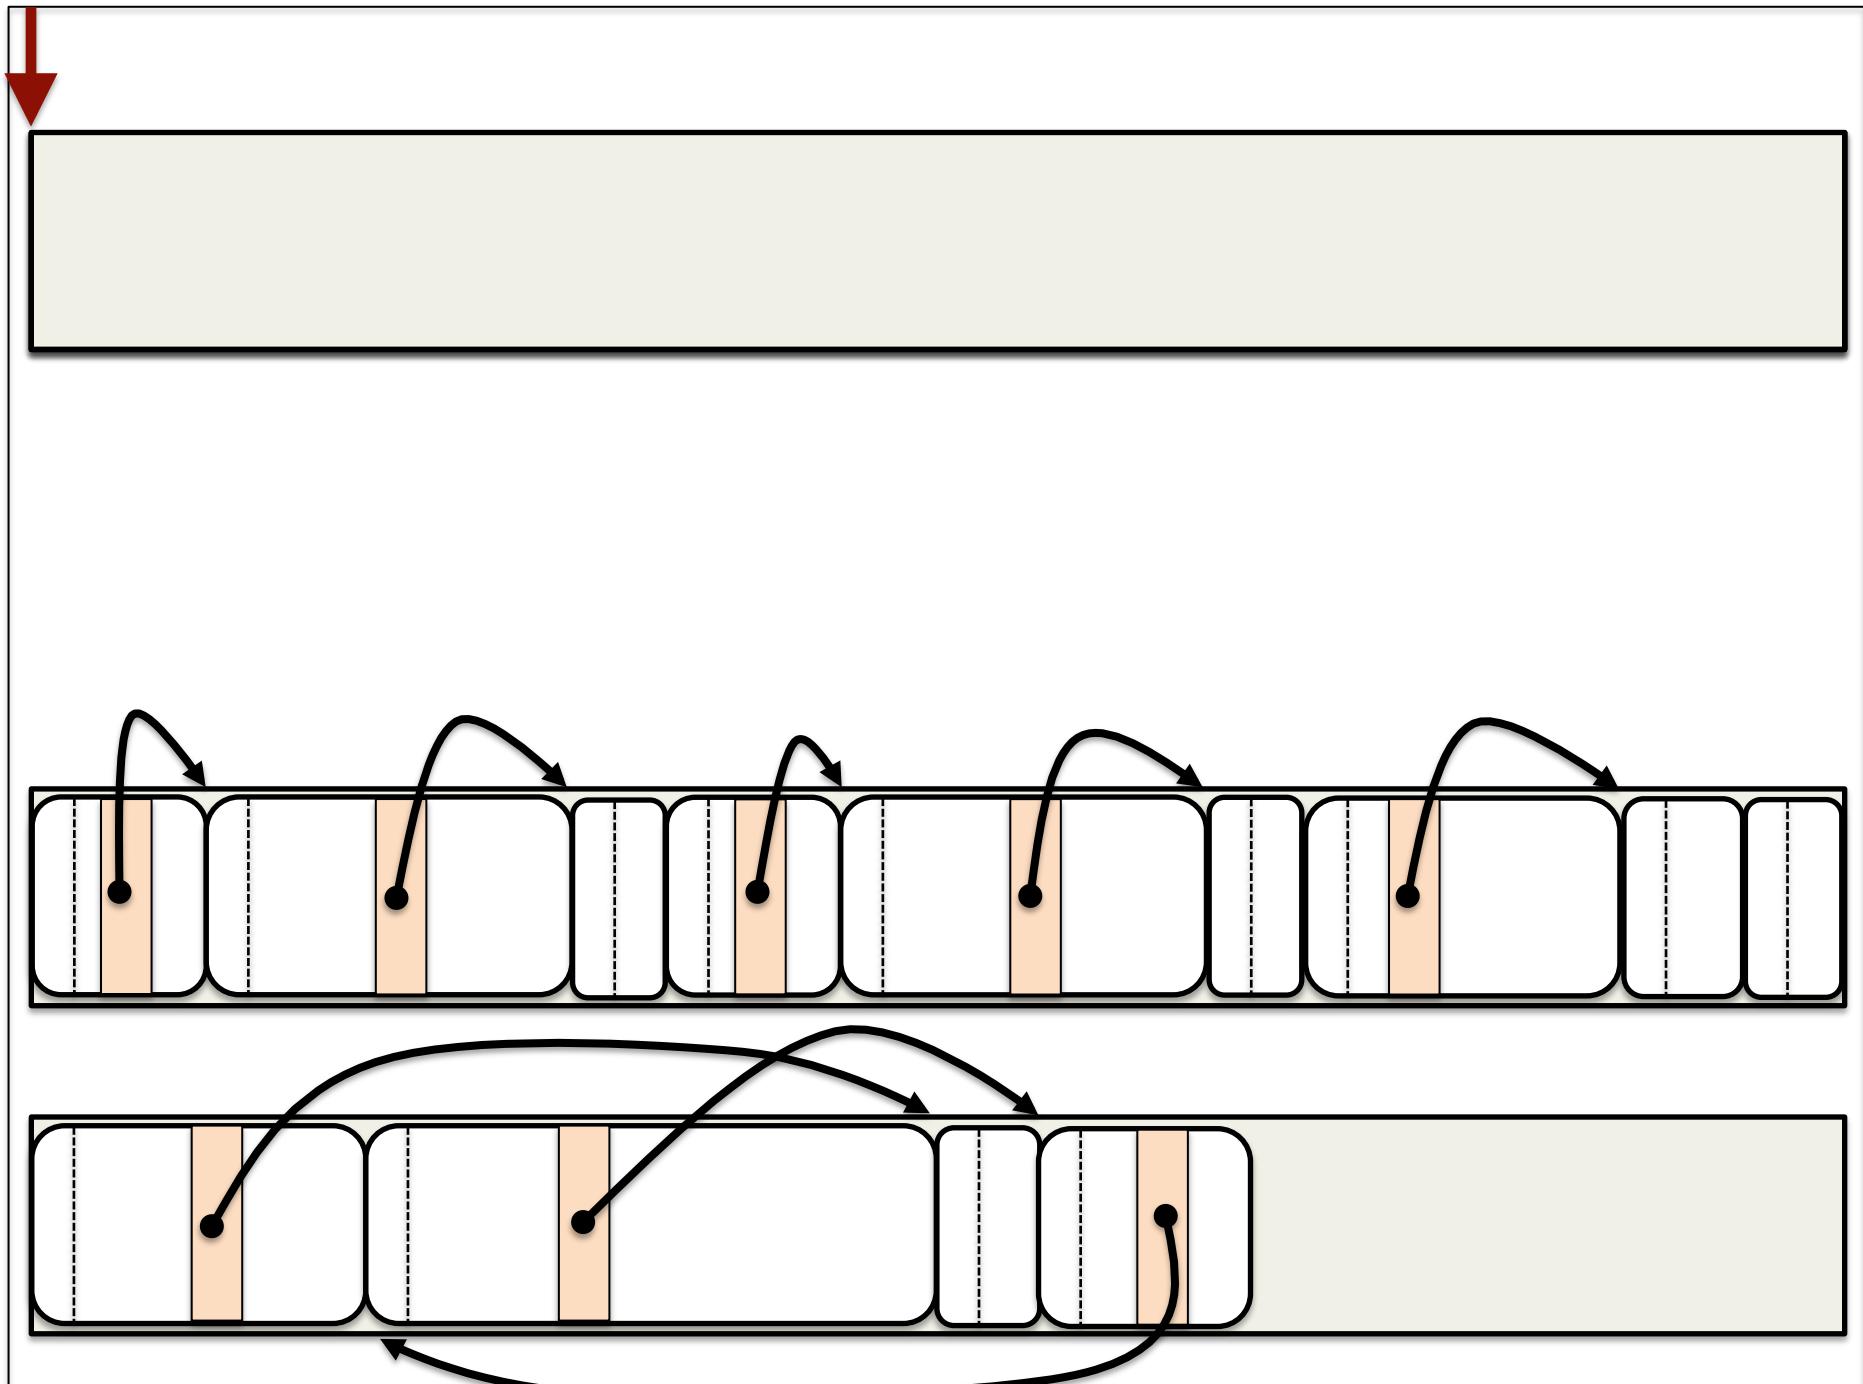






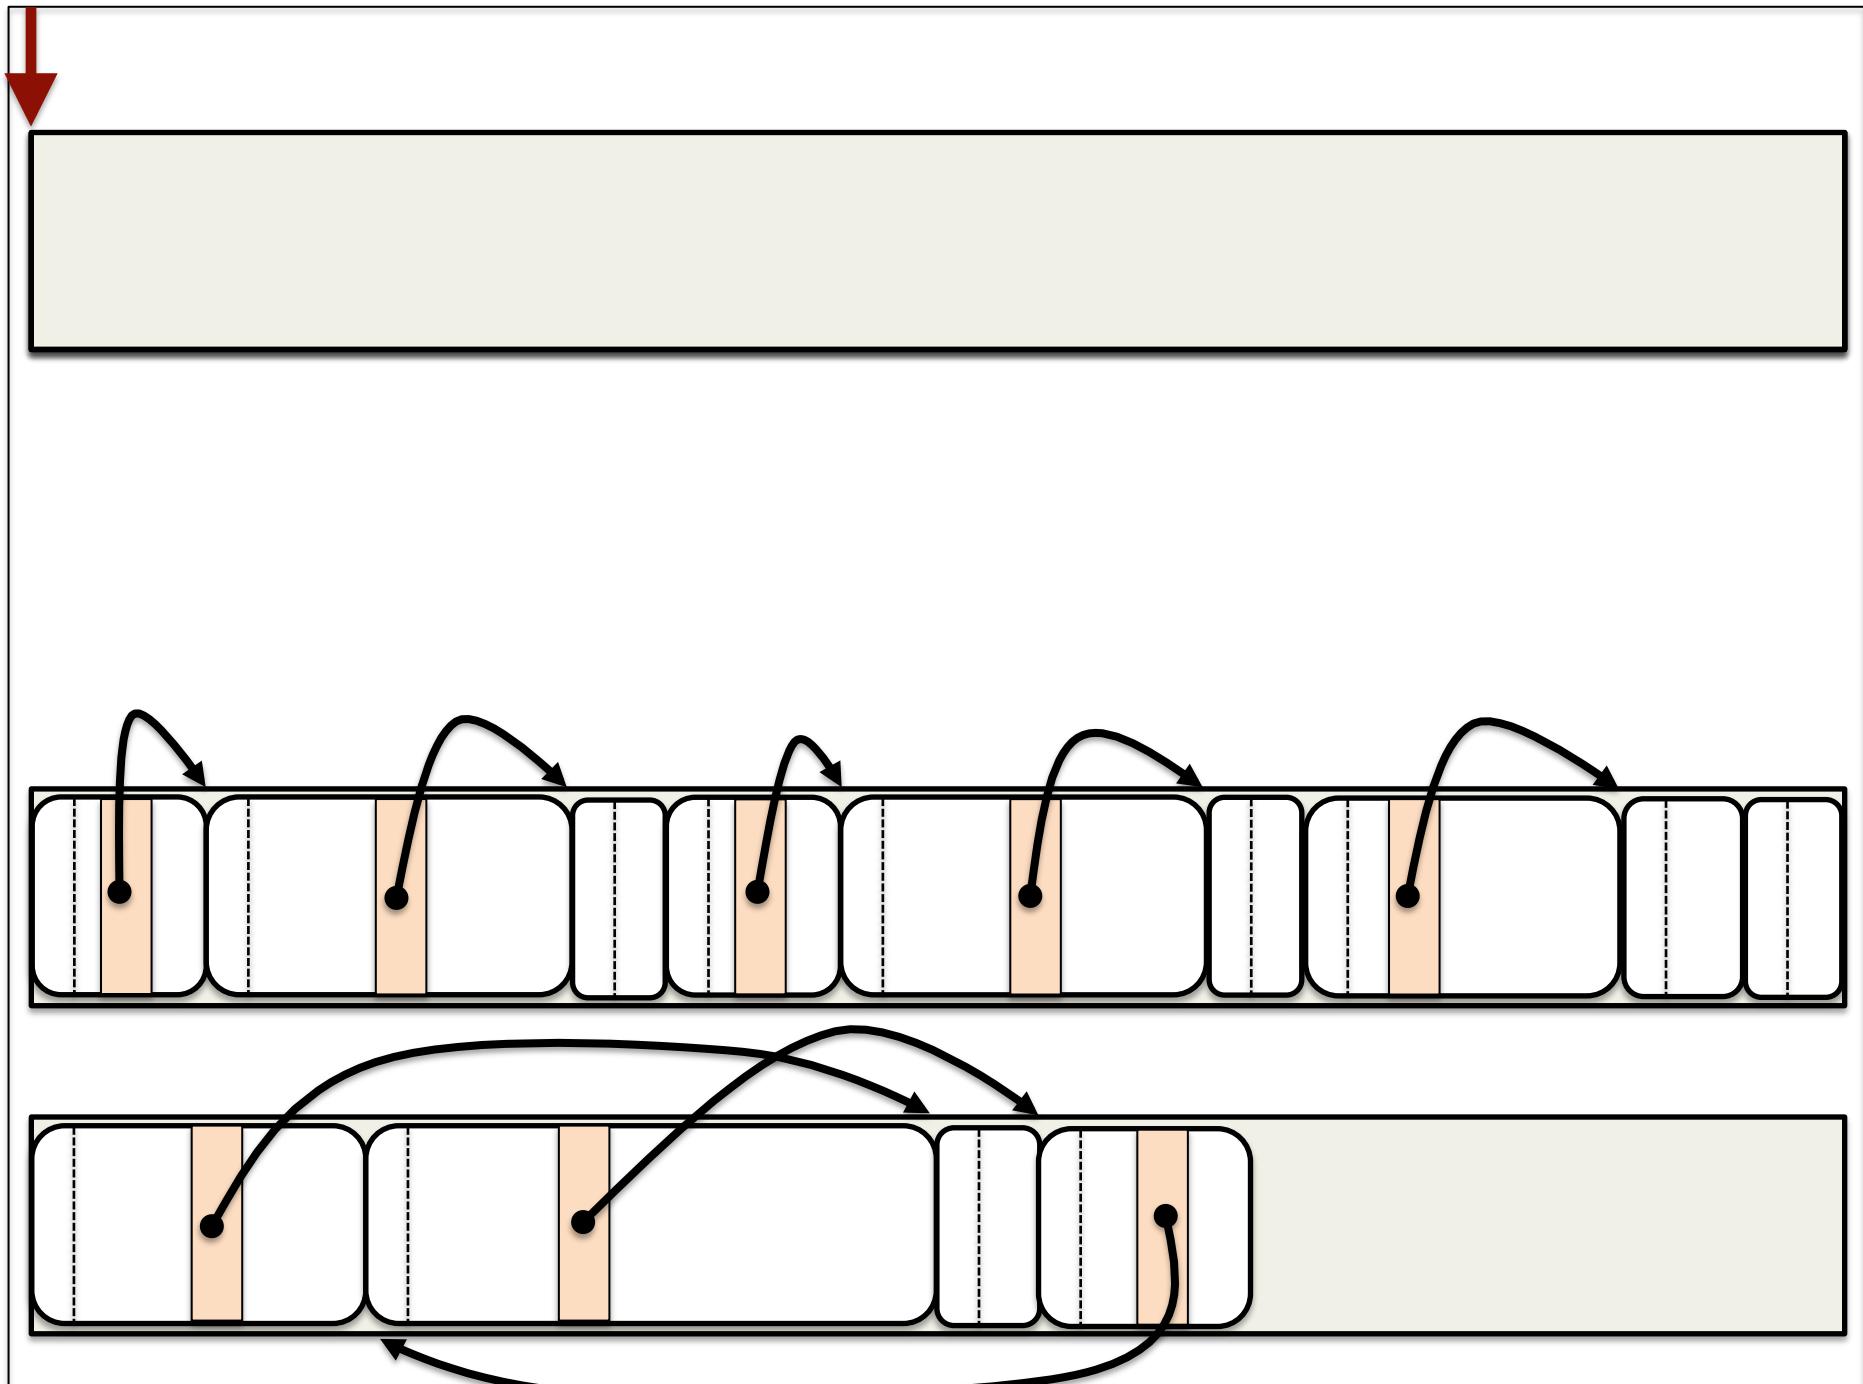


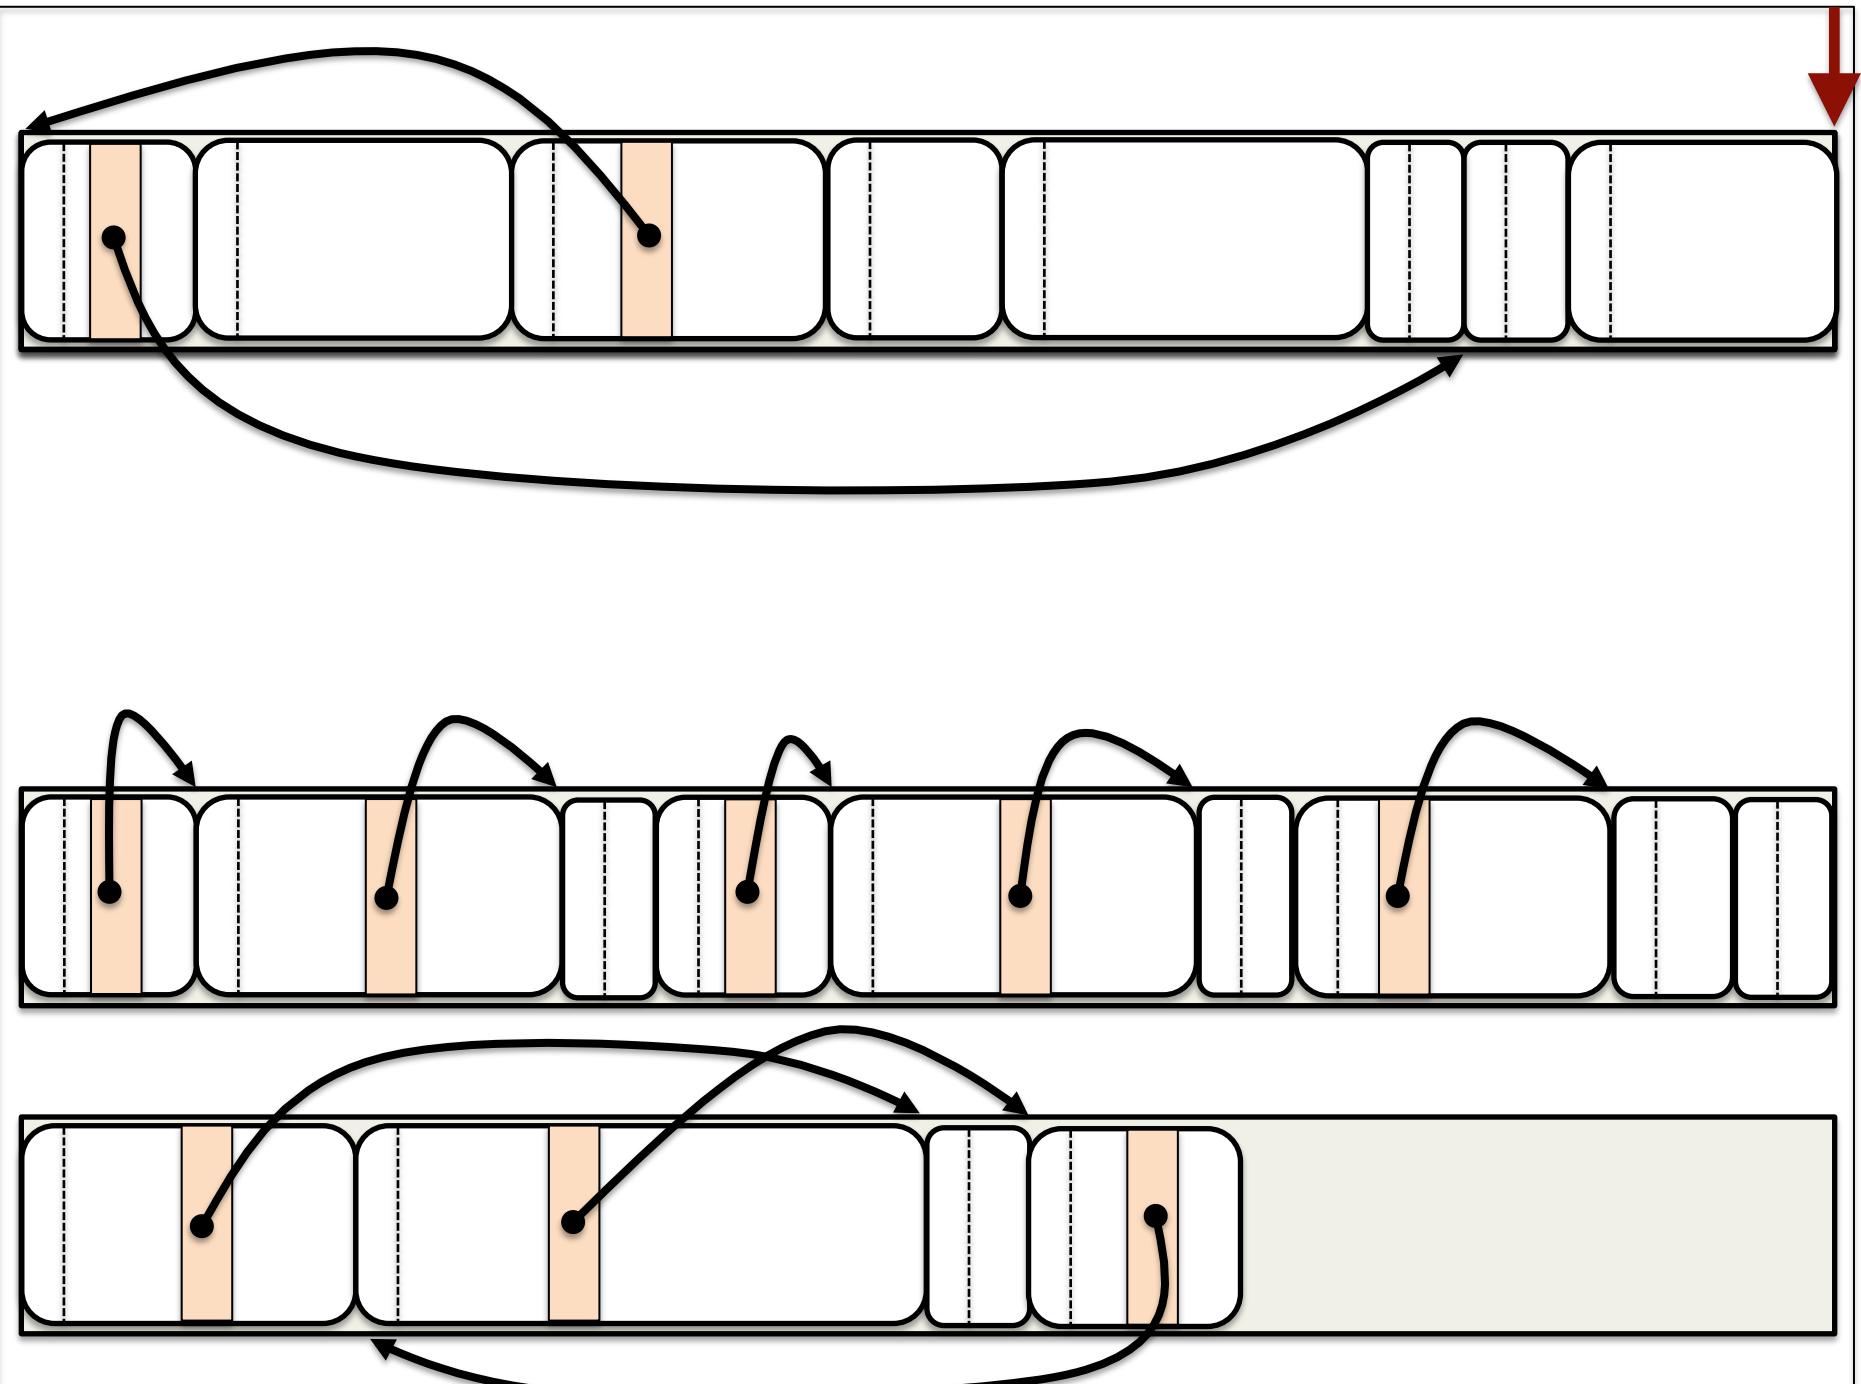


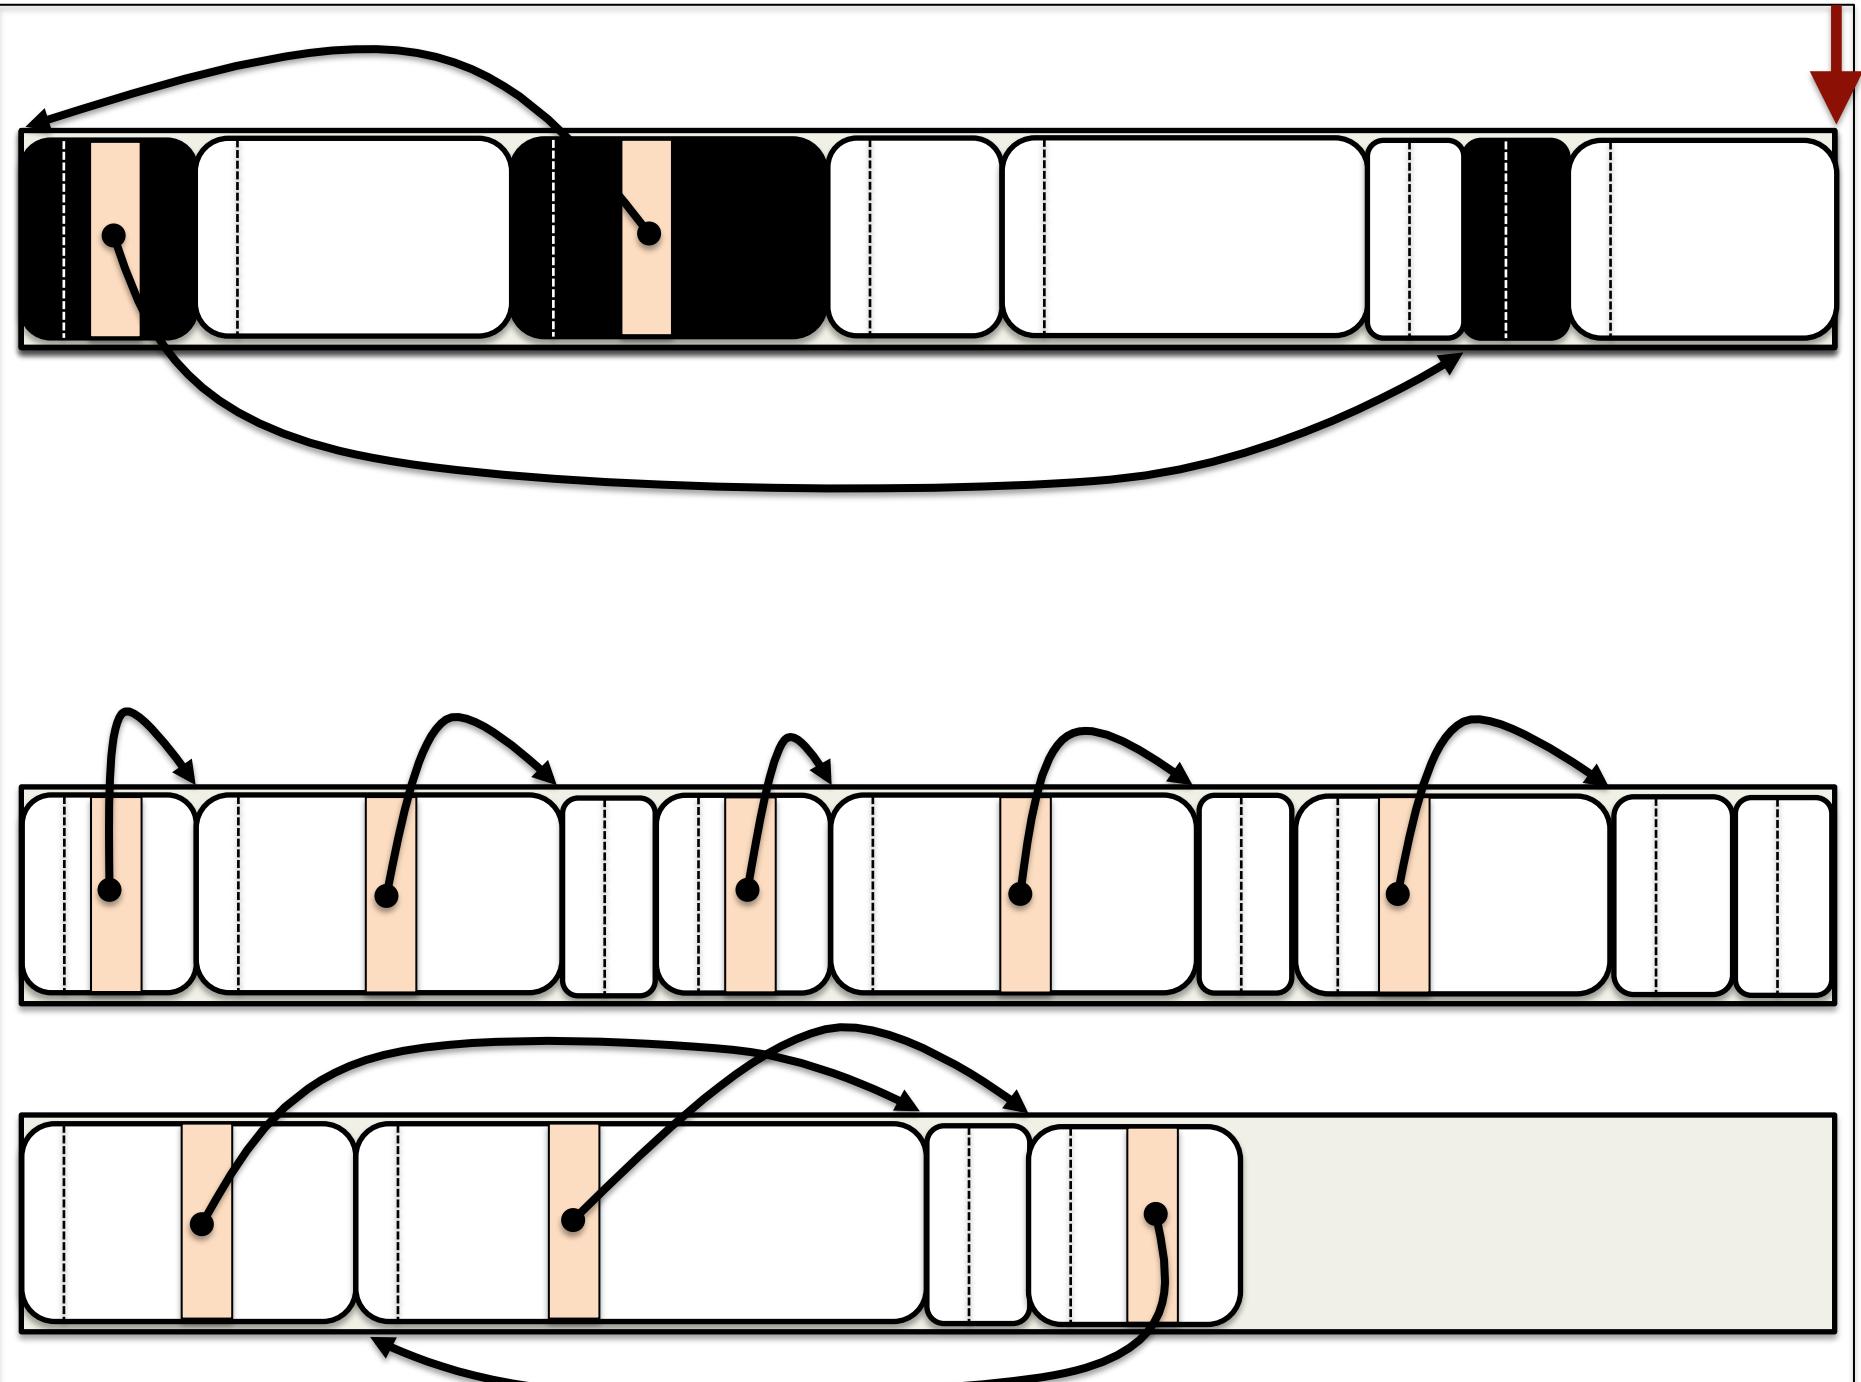


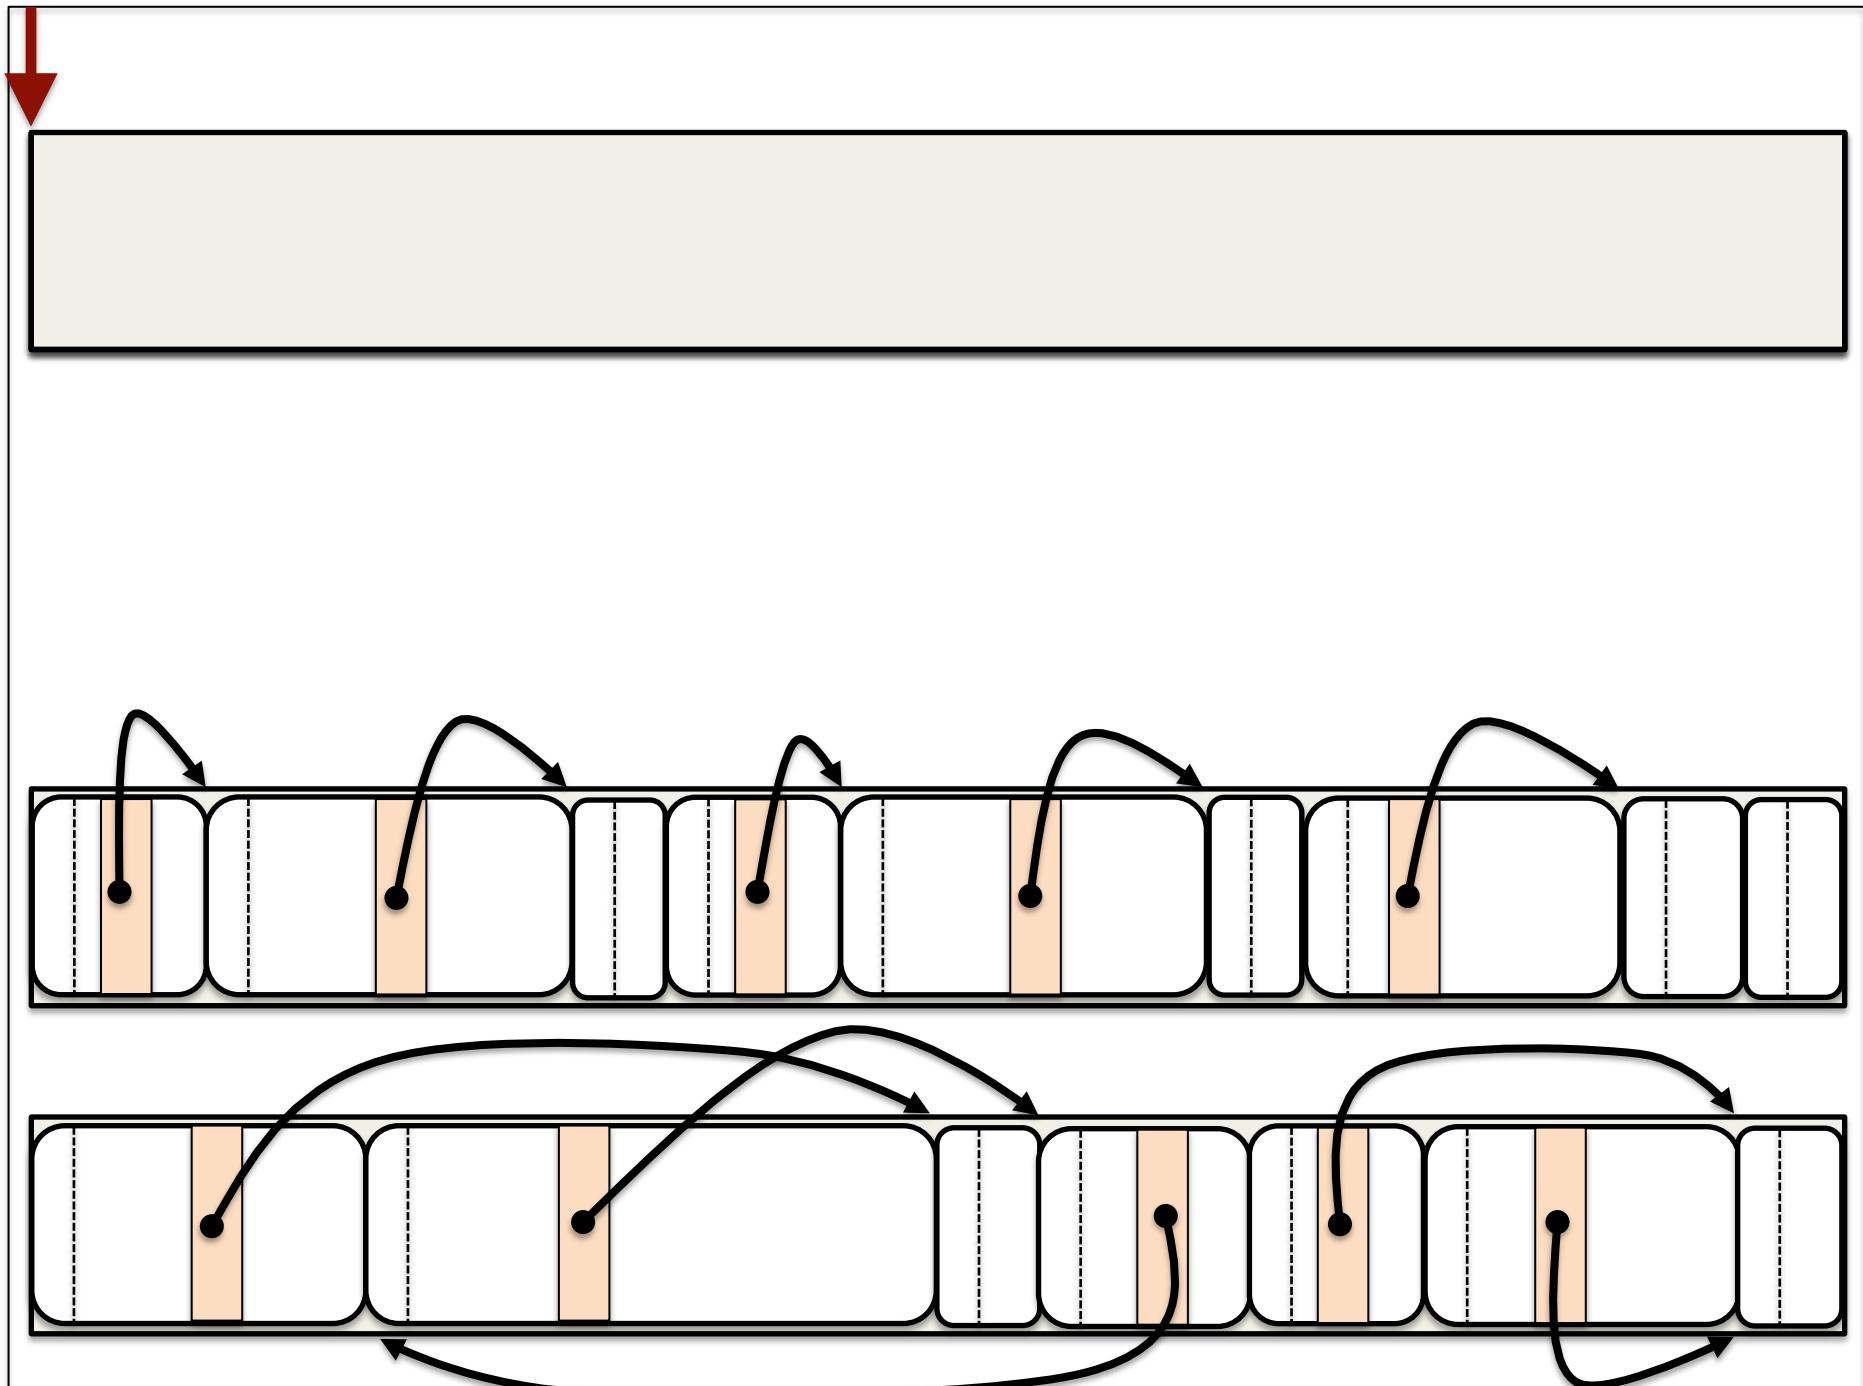
Minor Collections

- Minor collections are very fast
 - Only looking at a small part of the heap
 - Most objects in the nursery are garbage
- Mitigates the drawbacks to copying
 - Maximizes garbage to live ratio
 - Don't copy long-lived objects many times









Major Collection

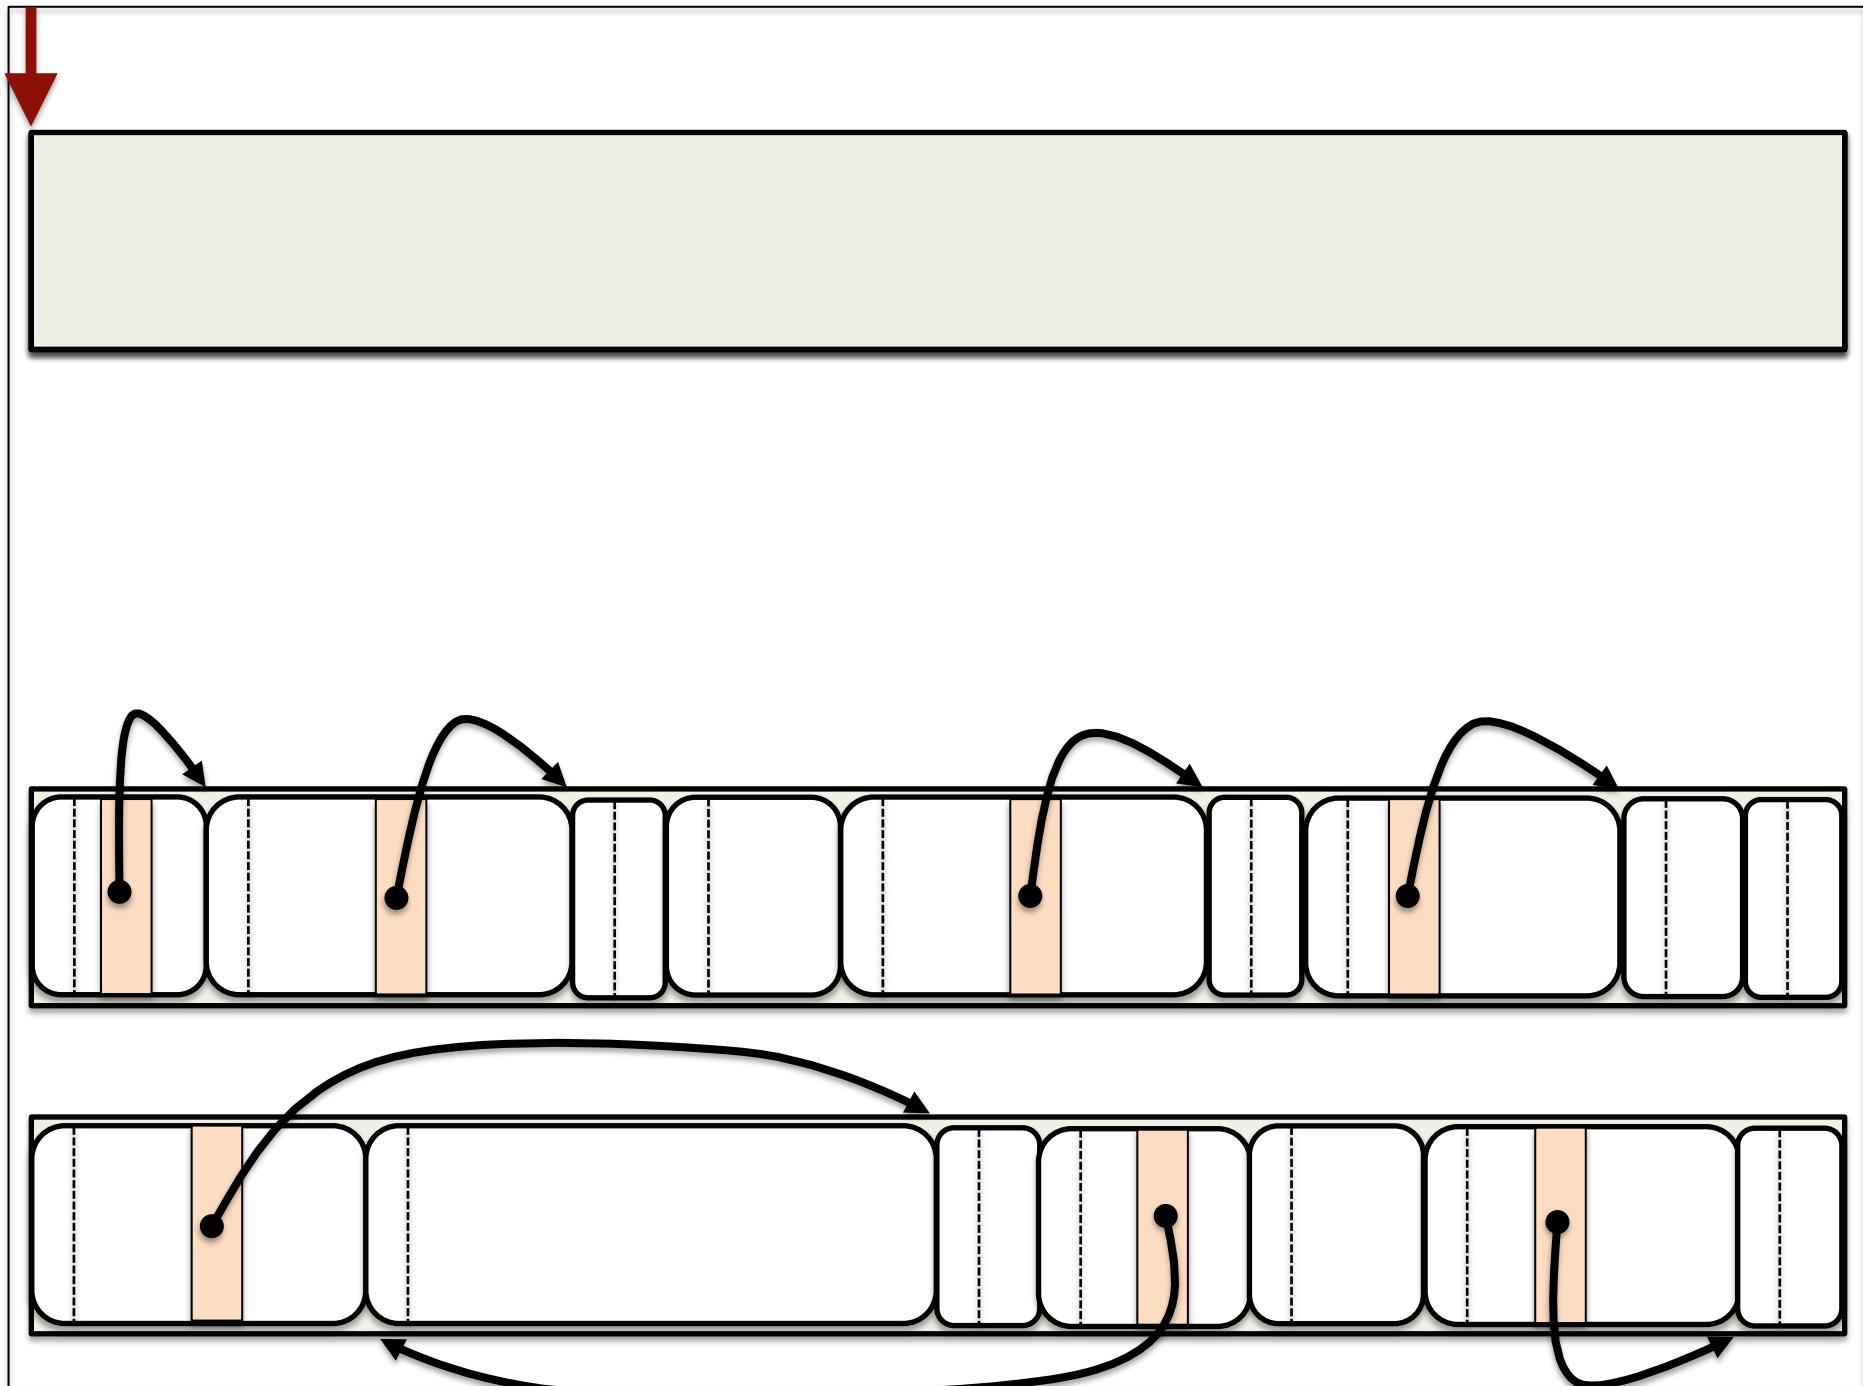
- Eventually the mature space fills up
- We trigger a major collection to free space
 - Collect the whole heap
- Same performance impact as a regular GC
 - Much less frequent than minor collections

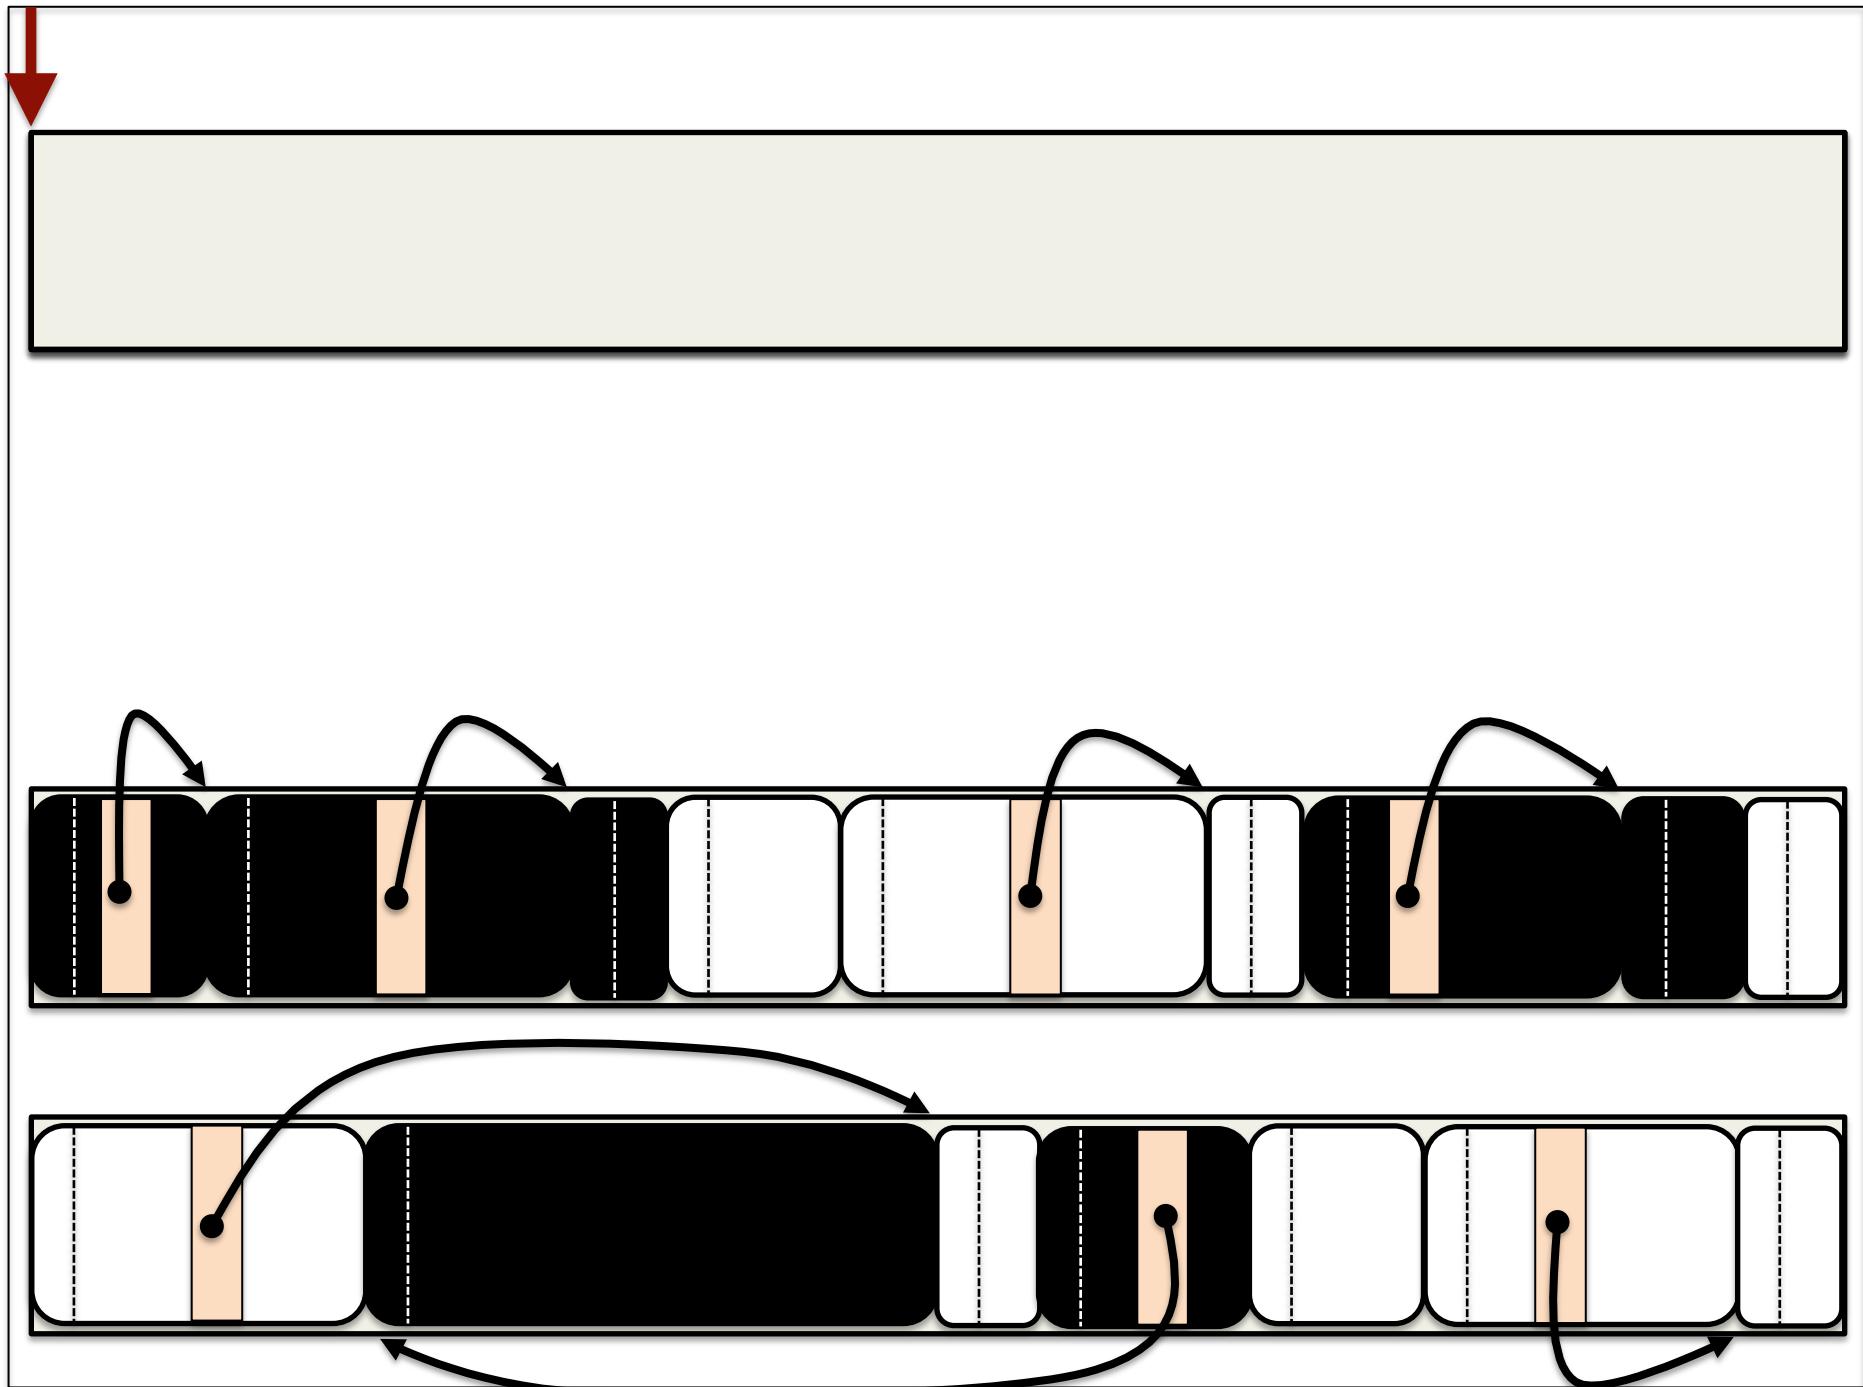
Triggering a Major Collection

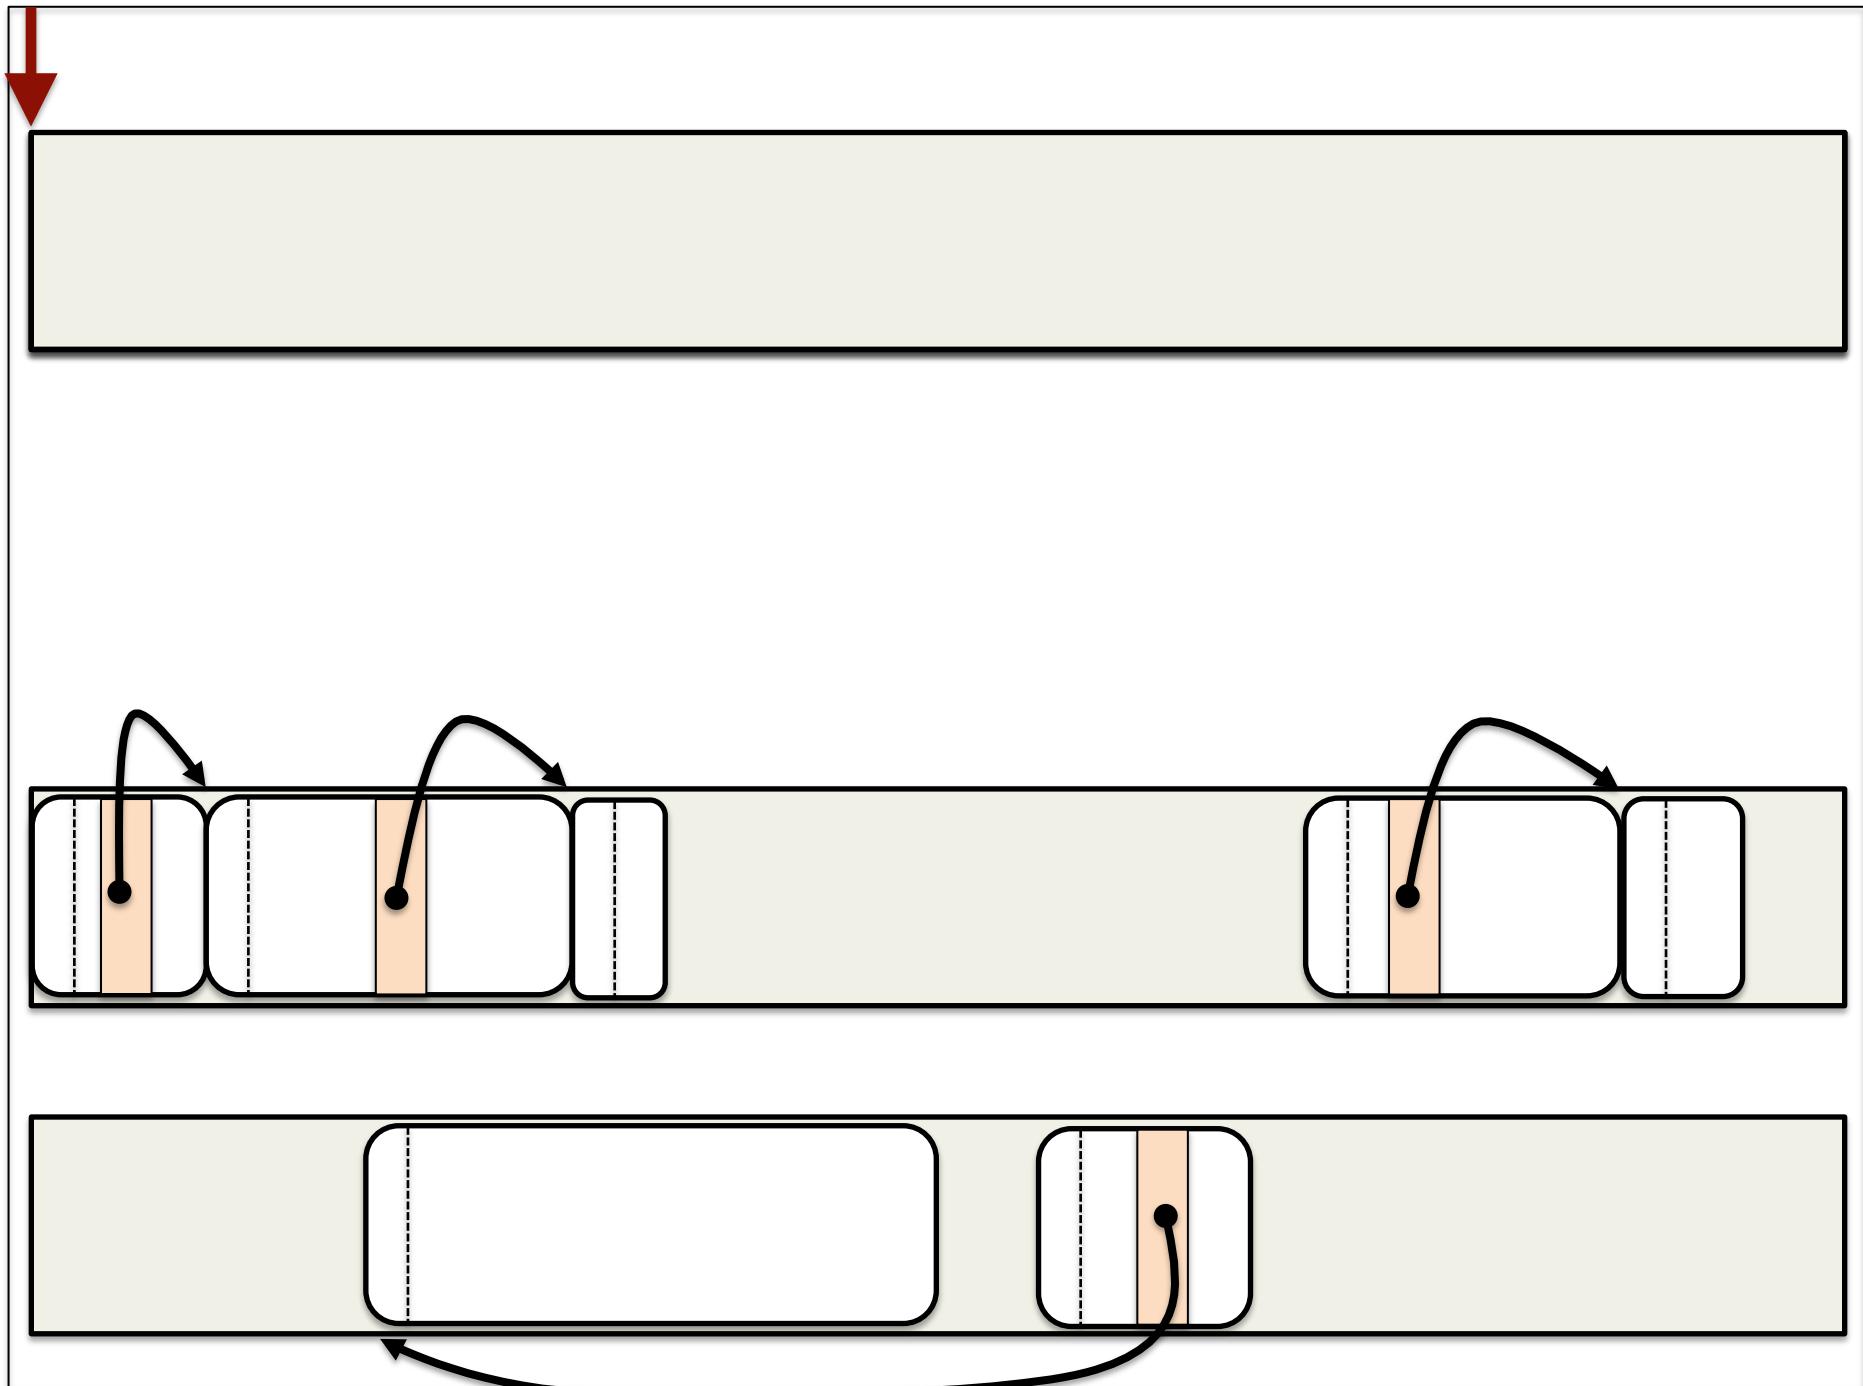
- GC algorithm triggers major collection
 - When no space on nursery promotion
 - Based on heap occupancy
 - When the user calls `System.gc()`
- Trade-off complexity for time
 - Frequent major collections more expensive
 - Switching from minor to major has overhead

Per-Region Management

- The nursery always uses copying collection
 - Need to evacuate objects
 - Want to use bump-pointer allocation
- Mature space can use a different algorithm
 - May not want to use a copy reserve
 - Allocation happens less frequently
 - May have a high live to garbage ratio
 - If strong hypothesis is correct





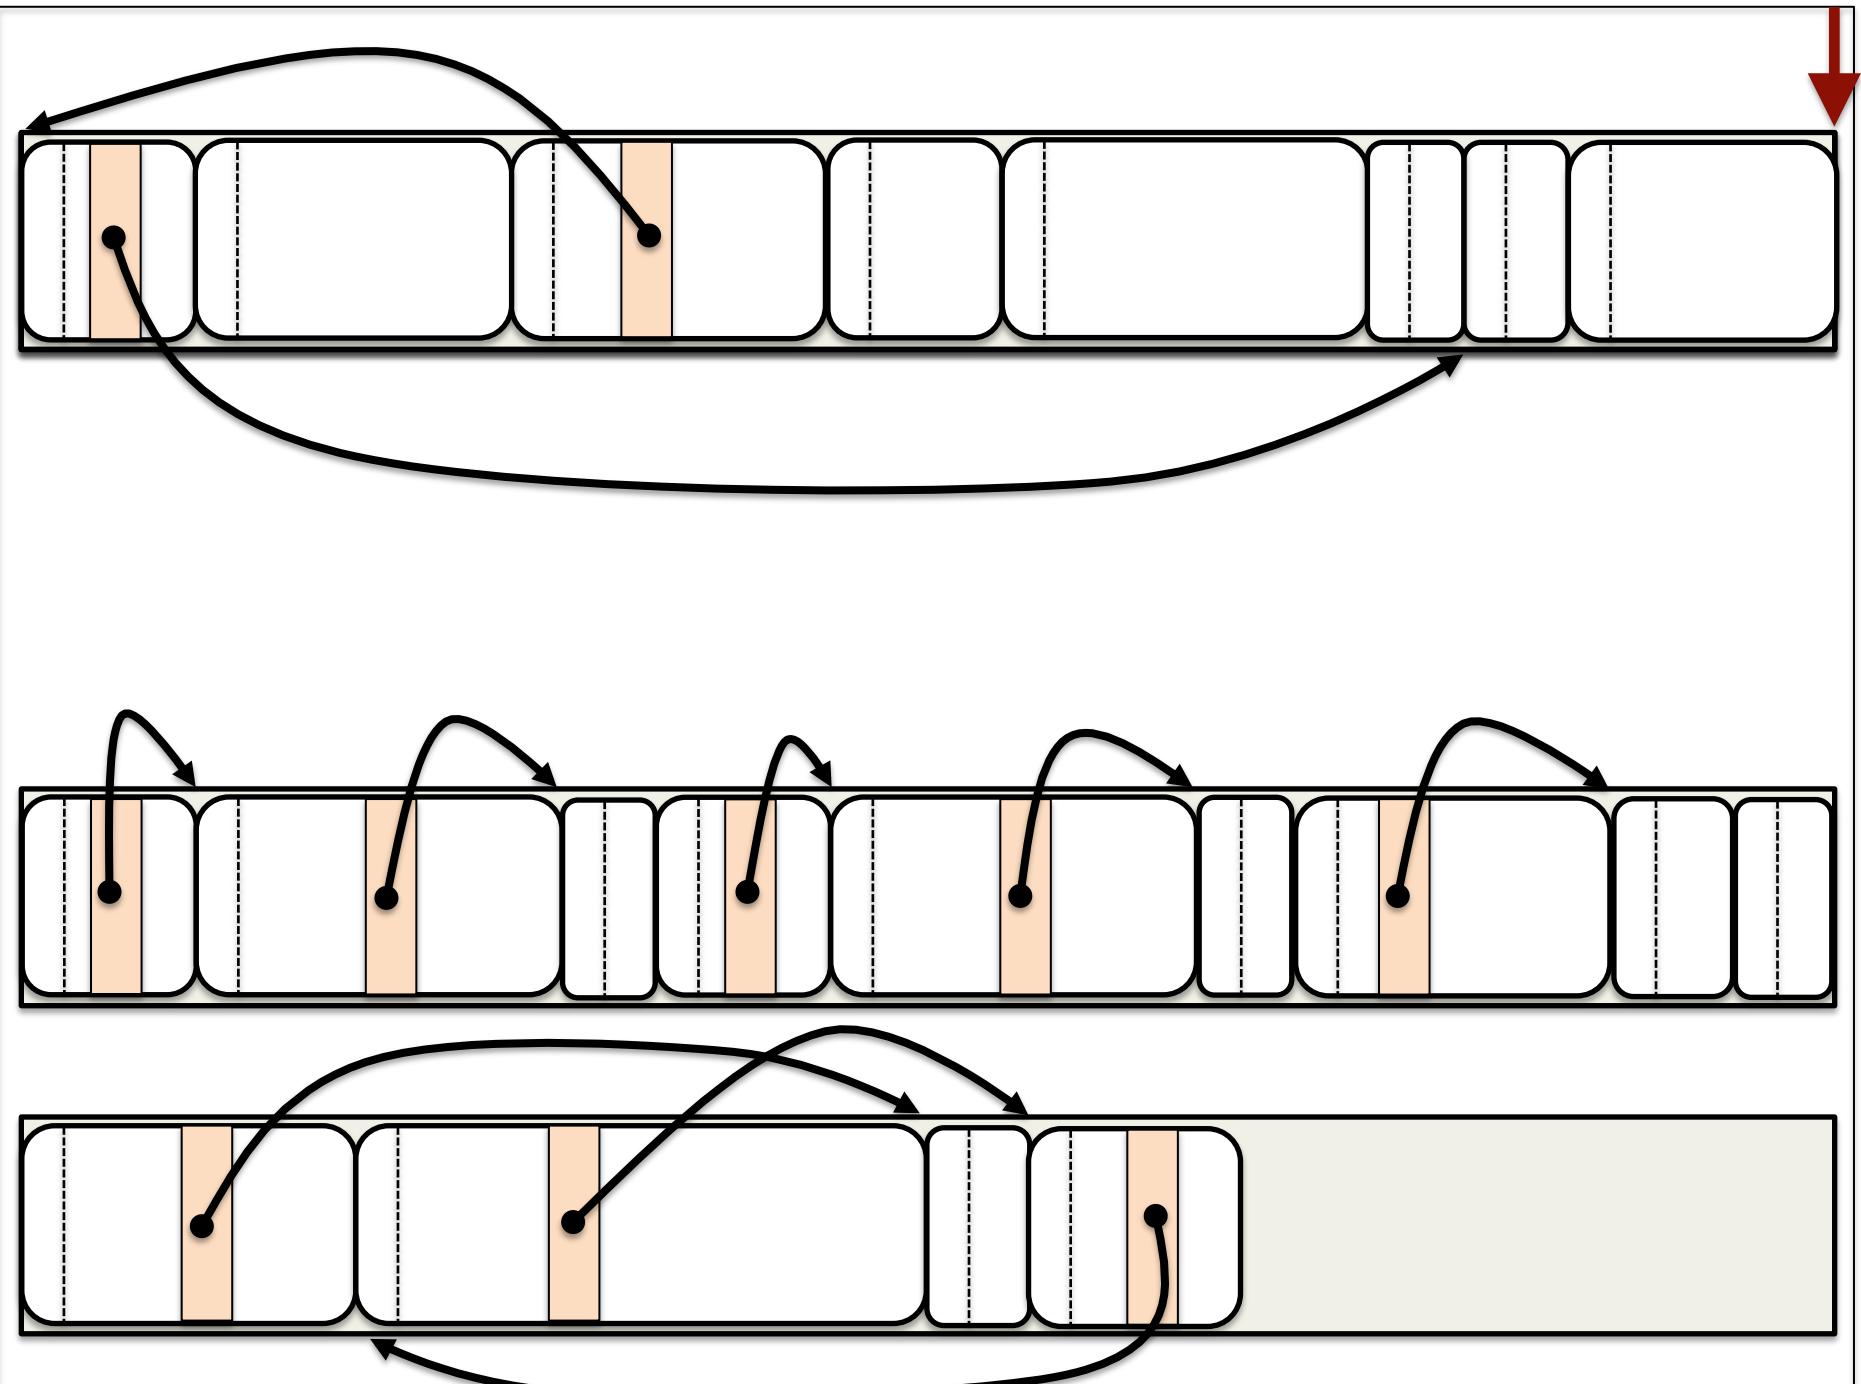


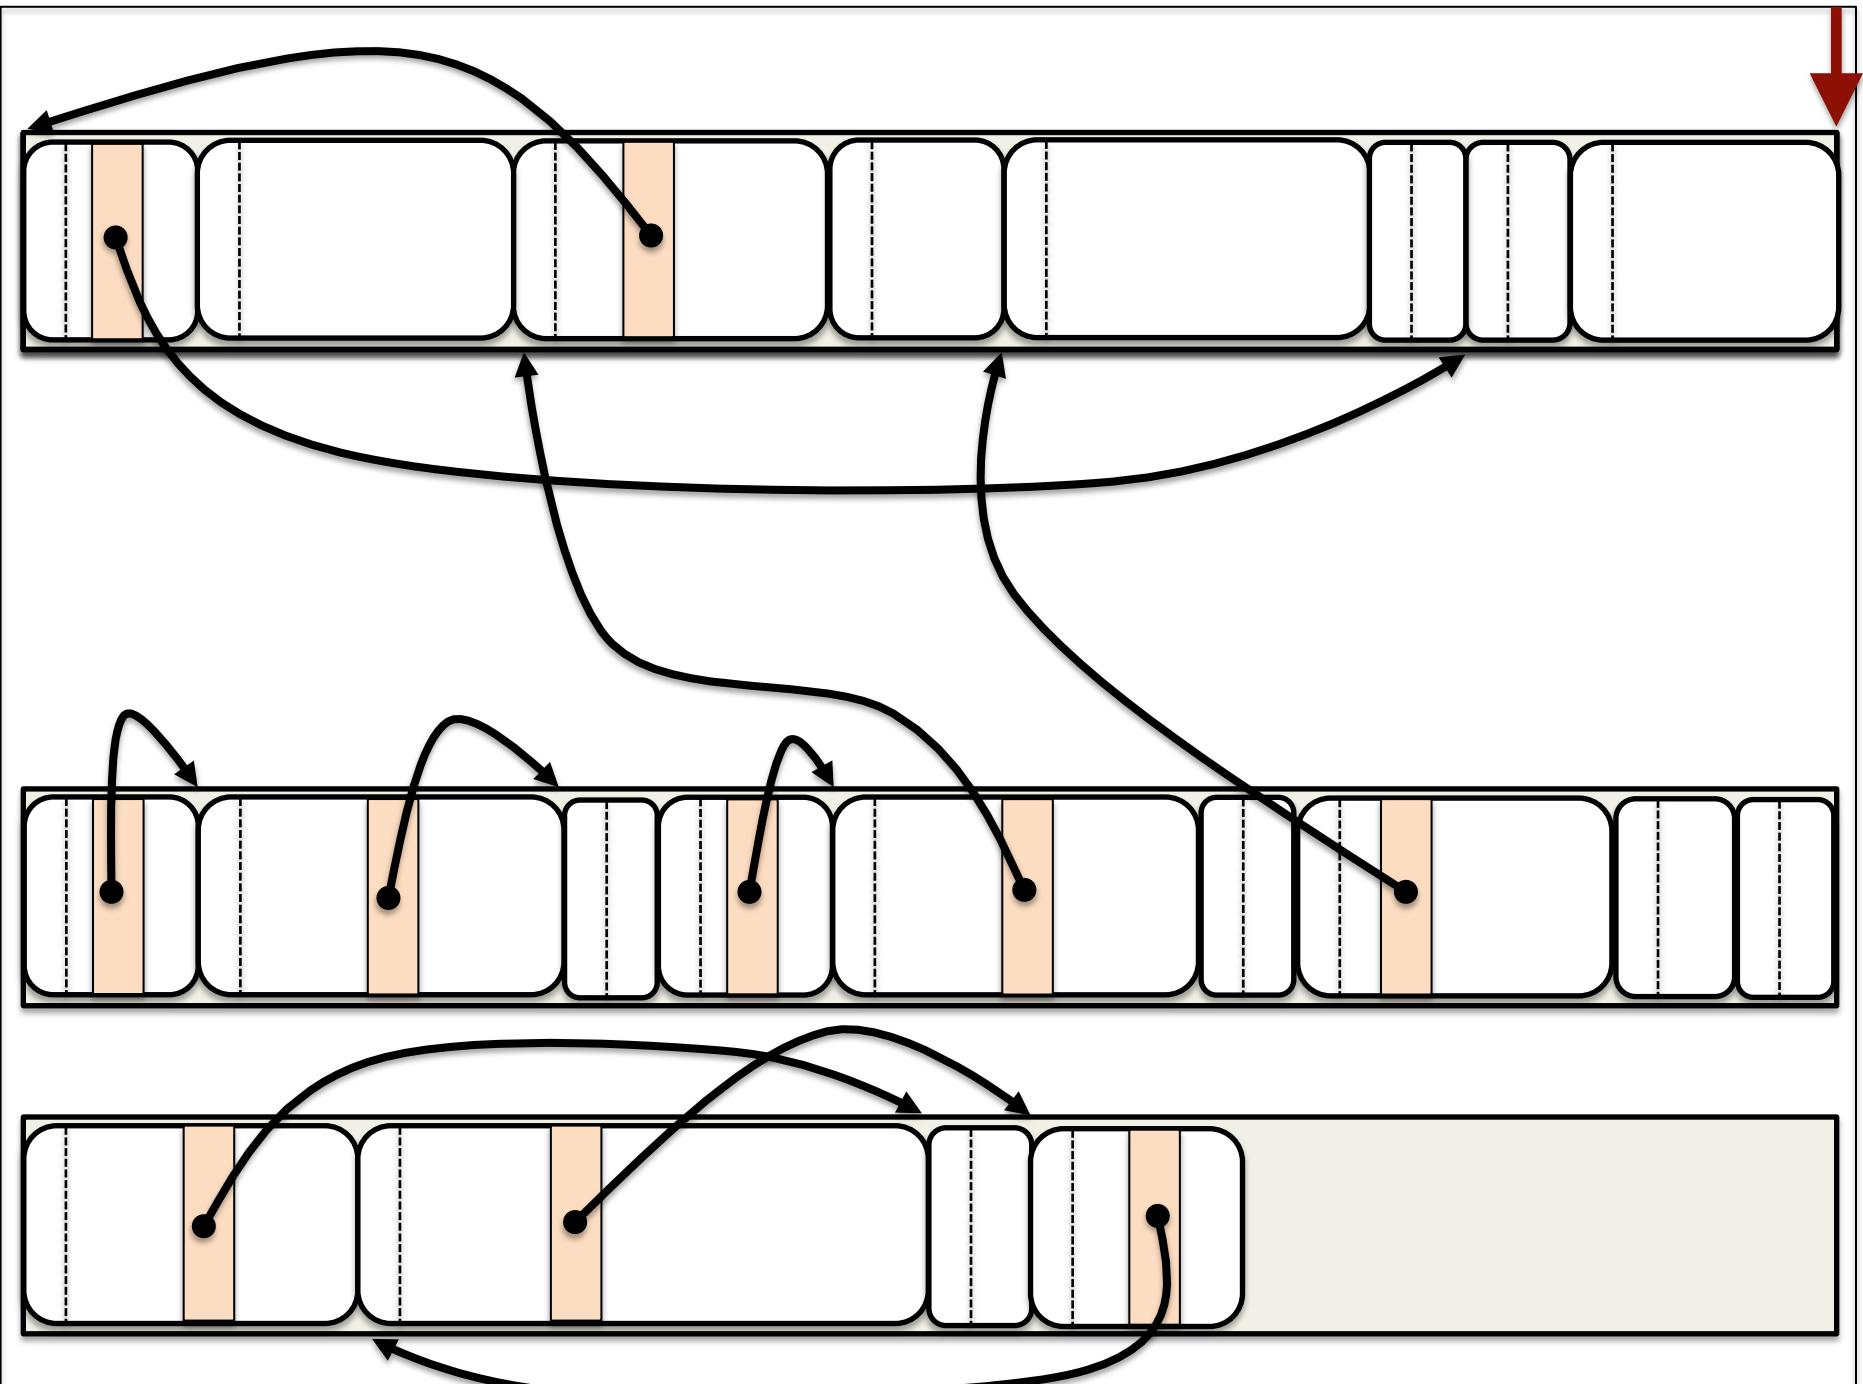
Mature Space Management

- Lots of variation in how to collect mature space
 - Mark/Sweep
 - Semi-Space
 - Mark/Compact
 - Various hybrids
- Tuning can depend on application characteristics
- Concurrent Mark Sweep (CMS) used in Hotspot

Reference Directions

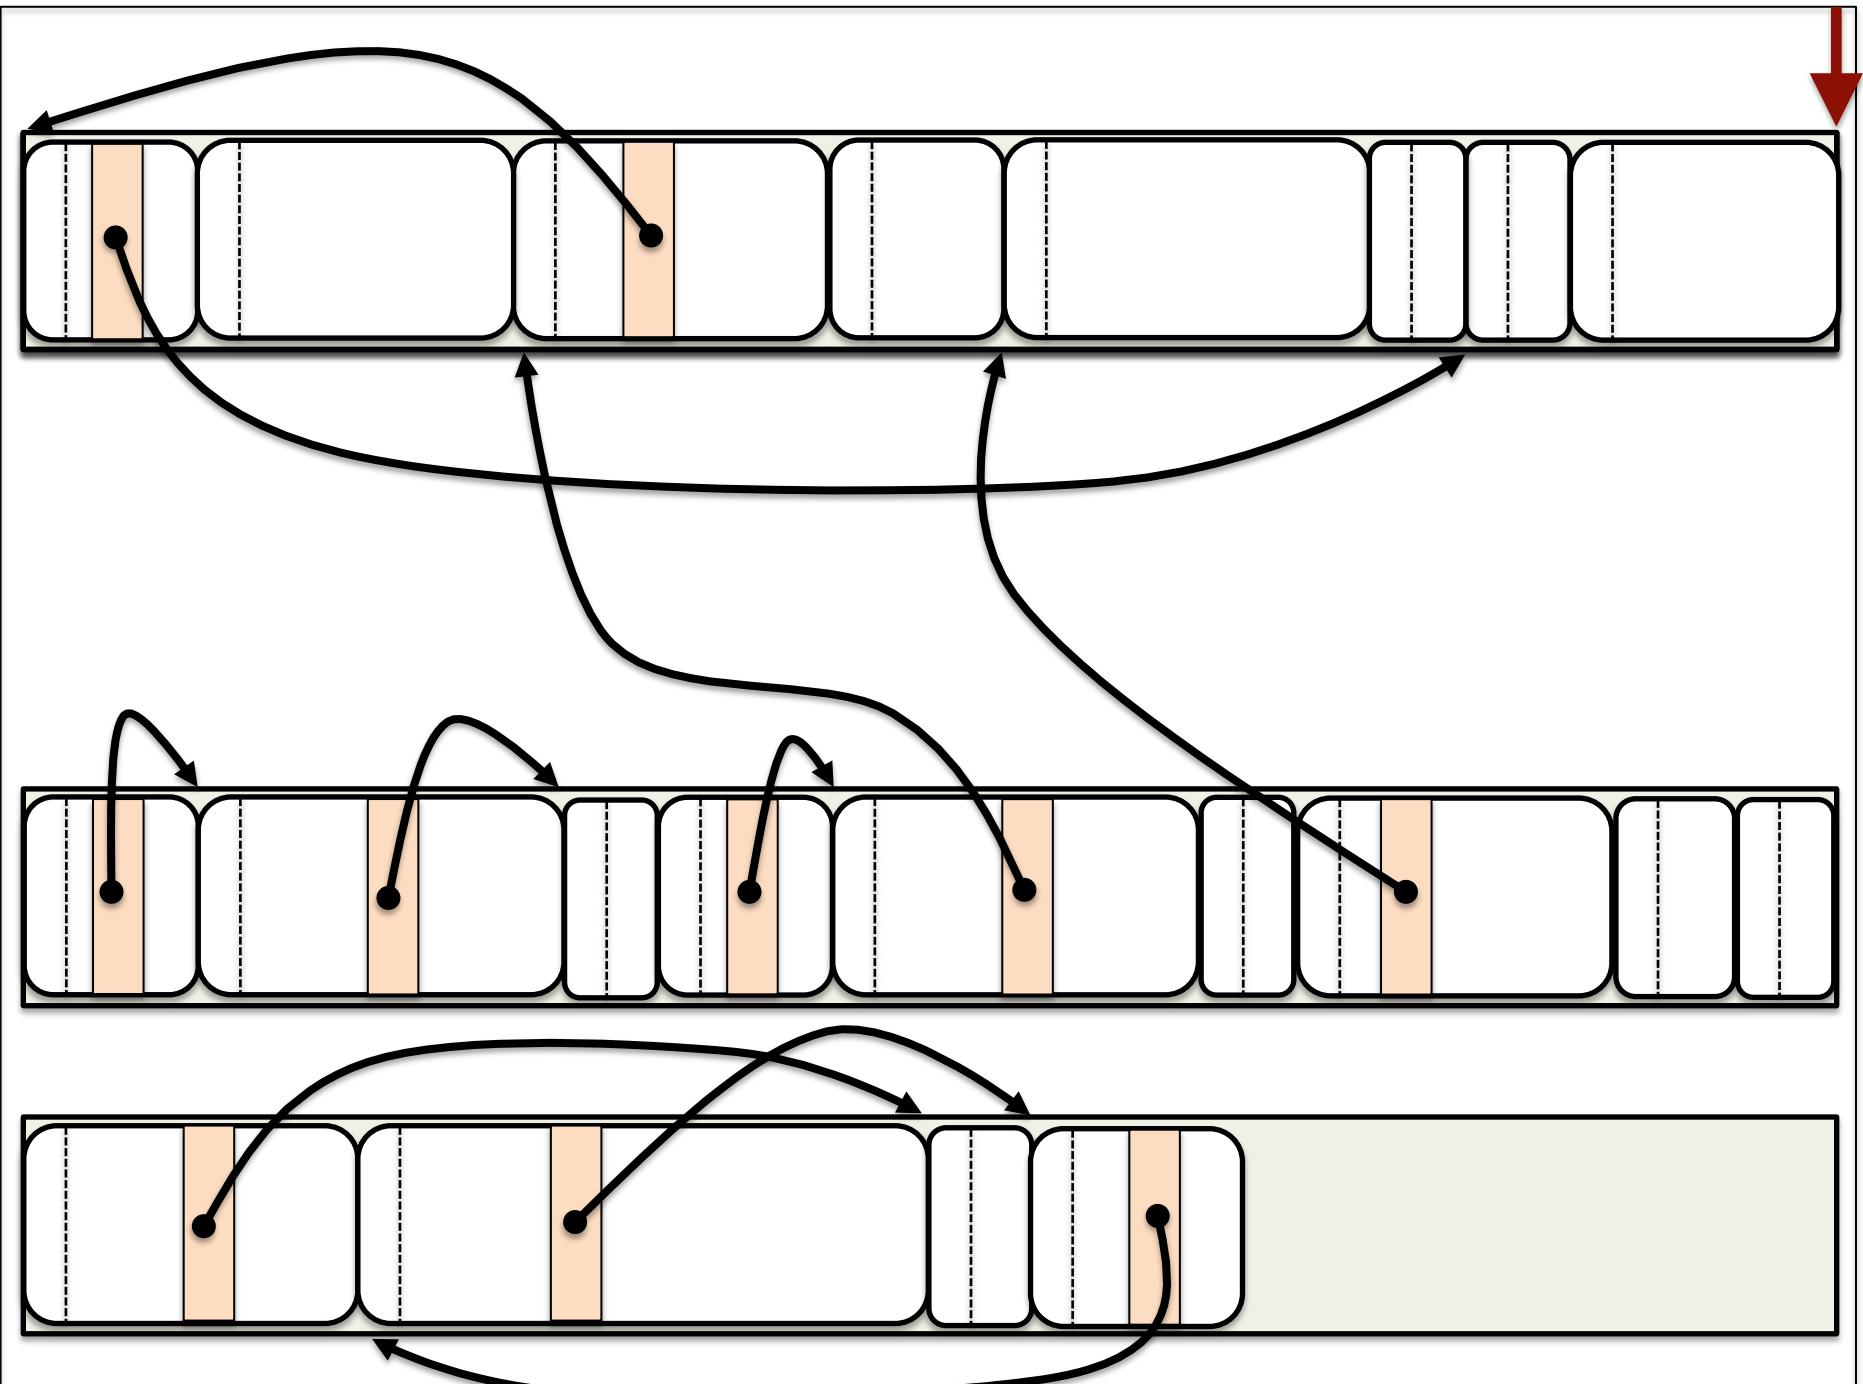
- So far we've assumed immutable objects
 - Can't change an object's contents after allocation
 - Reasonable assumption in some languages
- Mutable objects introduce a new problem
 - Mature objects can point to the nursery
- Mature objects could keep nursery objects live
 - Need to scan mature space for minor collection

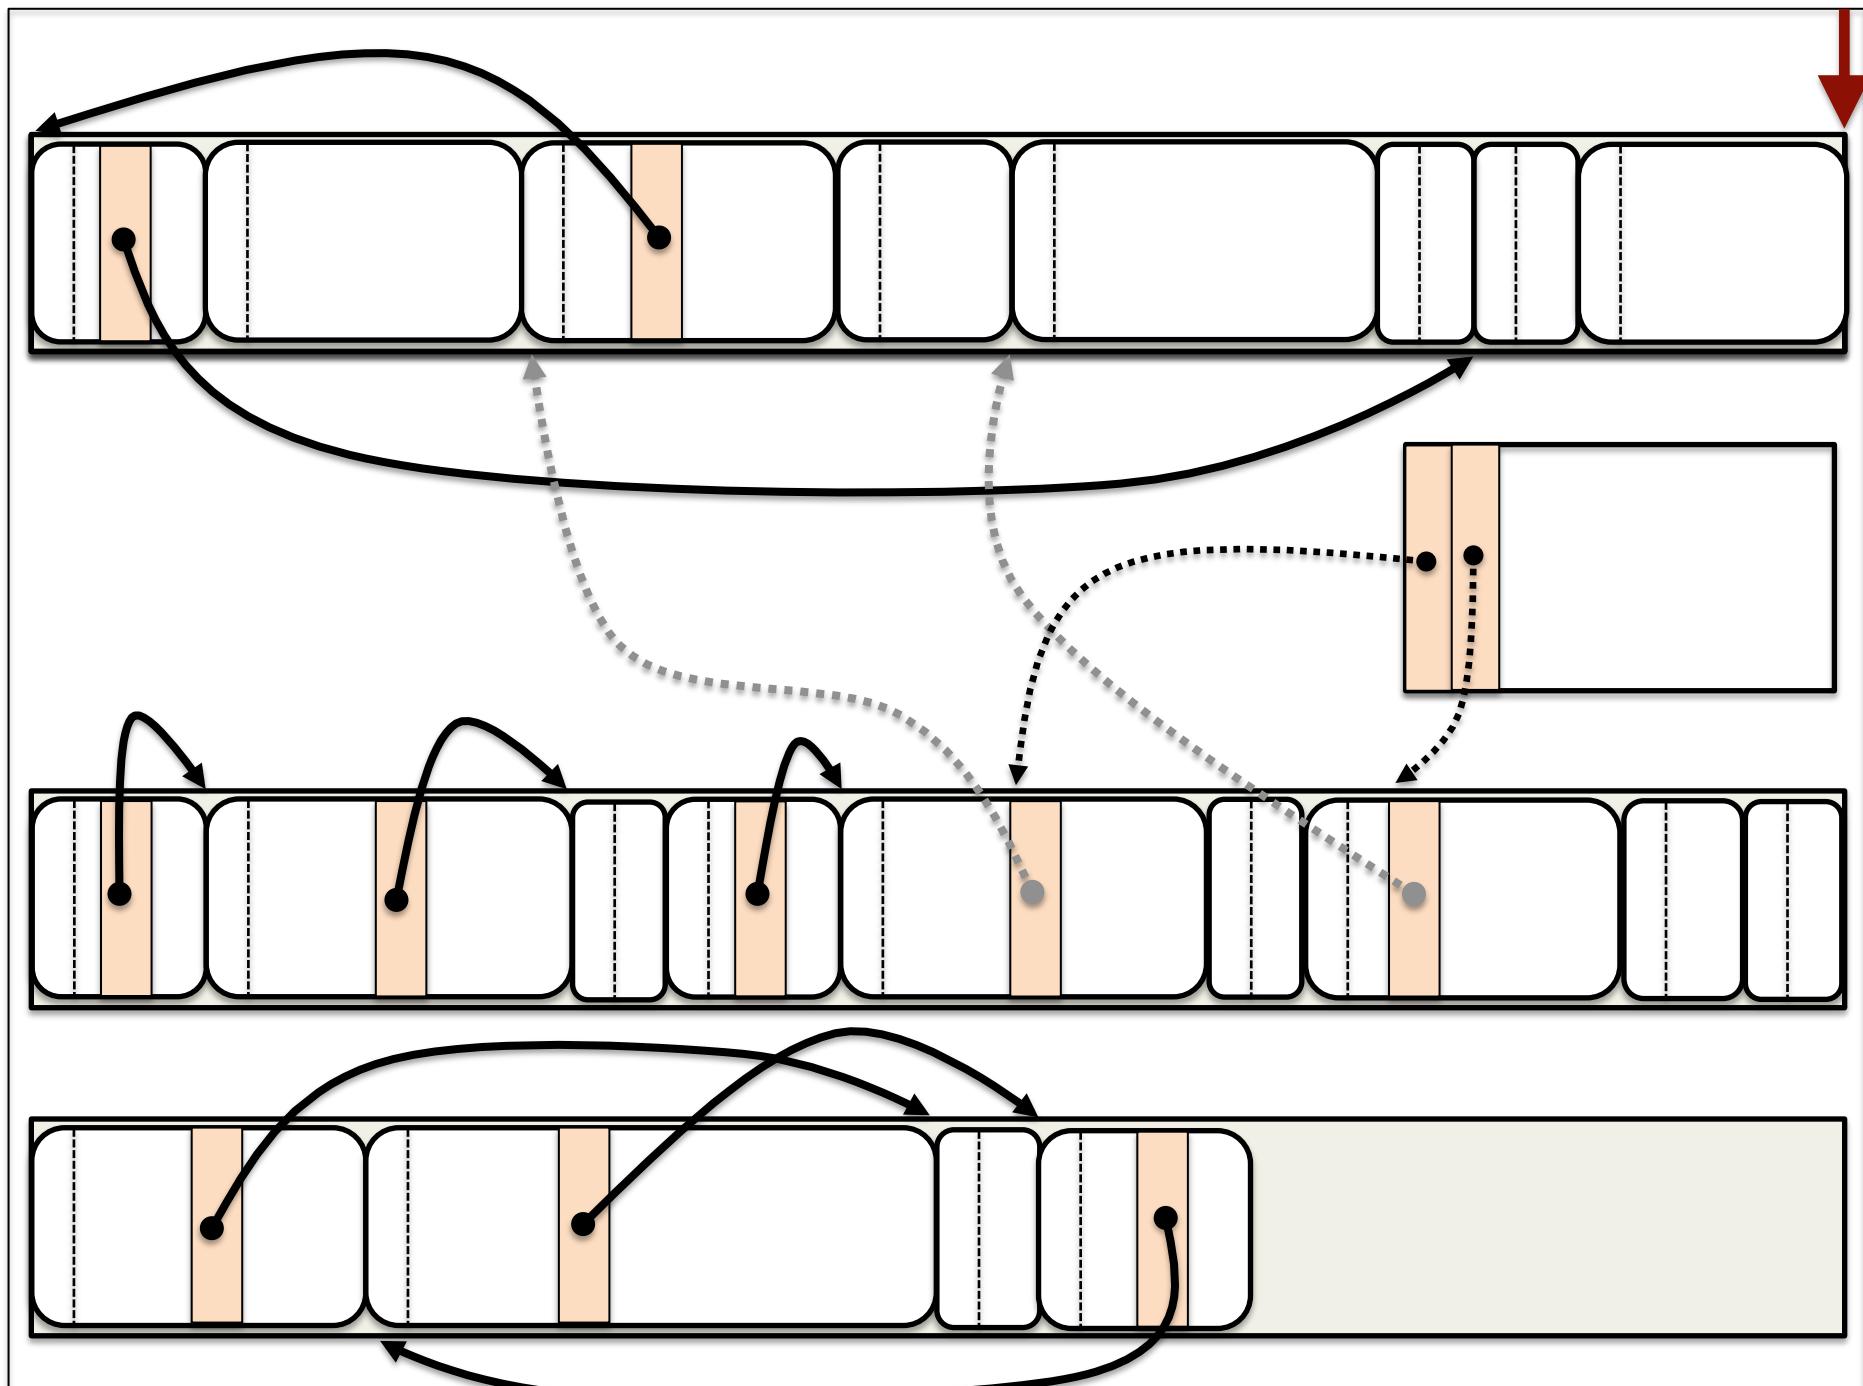




Remembered Sets

- Don't want to scan mature space every time
 - Lose much of the benefit of generational GC
- Track objects that may point to the nursery
 - Generally found to be uncommon
- Use that set as roots for minor GC
- Recall the write barrier concept from RC





Generational Write Barrier

- Insert small piece of code on reference writes
 - Maintain the remembered set
- Need only store references from old to new
 - We'll be scanning the nursery anyway
- Barrier can be heavily optimized
 - Still a source of some overhead
 - Hard to measure – somewhere around 2-3%

Generational Write Barrier

- If the source pointer is in the mature space
 - And the target pointer is in the nursery
 - Update the remembered set

Optimizing the Barrier

- There are two comparisons in the barrier
 - Is the source in the mature space?
 - Is the target in the nursery?
- We know the boundary between the regions
 - Checking for location just a comparison

Heap Layout

- The layout of the heap can help performance
 - Align regions on power of two boundaries
 - Gives a natural tag to regions
 - Find region by masking bits
- Nursery location can speed up allocation
 - Put the nursery at the end of the heap
 - Don't need to do the bump pointer comparison
 - Allocation ends at page fault

Fast and Slow Paths

- We can use heuristics to speed up the barrier
 - Assume that most pointer updates are in the nursery
 - If so, the first comparison normally fails
- Only inline the first comparison
 - If it succeeds, jump to another method
- Performance gain depends on the architecture
 - Size of the instruction cache
 - Effectiveness of branch prediction

Multiple Generations

- Can use more than two generations
 - Maintain some third space to promote from mature
 - Should cluster very old and long-lived objects
- Tends to have diminishing returns
- Limits options for mature space manager
 - Copying overhead eventually dominates

Pretenuring

- We can identify some long-lived objects
 - Static analysis of the code
 - Historical profiling information
 - Heuristics (large arrays)
- We can allocate those directly to mature space
 - Eliminate the copying overhead
 - Can play to mature collector's strengths

Permanent Generation

- Heap region for long-lived structures
 - Class data
 - Pretenured objects
 - Very large objects
- Something of a misnomer
 - Objects may turn out to be garbage
 - May be collected if memory is tight

Nursery Size

- Don't want nursery to be too small
 - Frequent minor collections
 - Objects need time to become garbage
- Don't want nursery to be too large
 - Minor collections become slow
 - More objects copied triggering major collections

Tuning Nursery Size

- Optimal size for throughput is large
 - Approximately half the heap
 - Does the least amount of work overall
- Optimal for responsiveness is much smaller
 - Sacrifice total work for short pauses
 - May need to experiment to find the best size
 - Could tune adaptively