

In the Last Week

- Some good questions around Assignment 4
 - Structure of the line number table
 - Winding the PC back to the invoke instruction
 - How to clear the stack
 - How to rethrow exceptions
 - Using Exception in your code

Facebook Blog Post

- Engineers at Facebook built a new Android app
 - Re-implemented much of their application
- Used a new implementation pattern
 - Reduced method size
 - Greatly increased the number of methods
- Successive problems on old Android versions
 - Overflowing a buffer that tracks methods

Marching Ever Onwards

- Implemented a series of fixes
 - Couldn't break the application into smaller modules
 - Added code directly to the system class loader
 - Realized that the same problem exists at runtime
 - Tried and failed to bring down method count
 - Used JNI to increase the buffer size
 - Realized that vendors had changed the platform
 - Added heuristics to guess where the buffer was
 - Tested across 70 different models of phone

Discussion Boards

- Range of opinions on the problem and solution
 - Should they have hit the limit in the first place?
 - Was it worth targeting the new app to old phones?
 - Could the app be split into services?
 - It would have helped to know what else they tried
- General sympathy for the position they were in
 - Tight release deadline
 - Having to support old versions of the platform
 - Pragmatism when trading off business requirements

Issues With the Solution

- Was there a reason for the smaller buffer size?
 - Was it determined based on the hardware capacity?
 - Could scanning the larger buffer tax the GC?
- Effects on Google
 - Violating the spirit of the interface
 - Later Dalvik patch had to support the hack
- Why was the buffer fixed size to begin with?
 - The dangers of leaky abstractions
- Inherent fragility of the approach

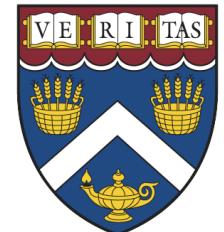
What They Did Right

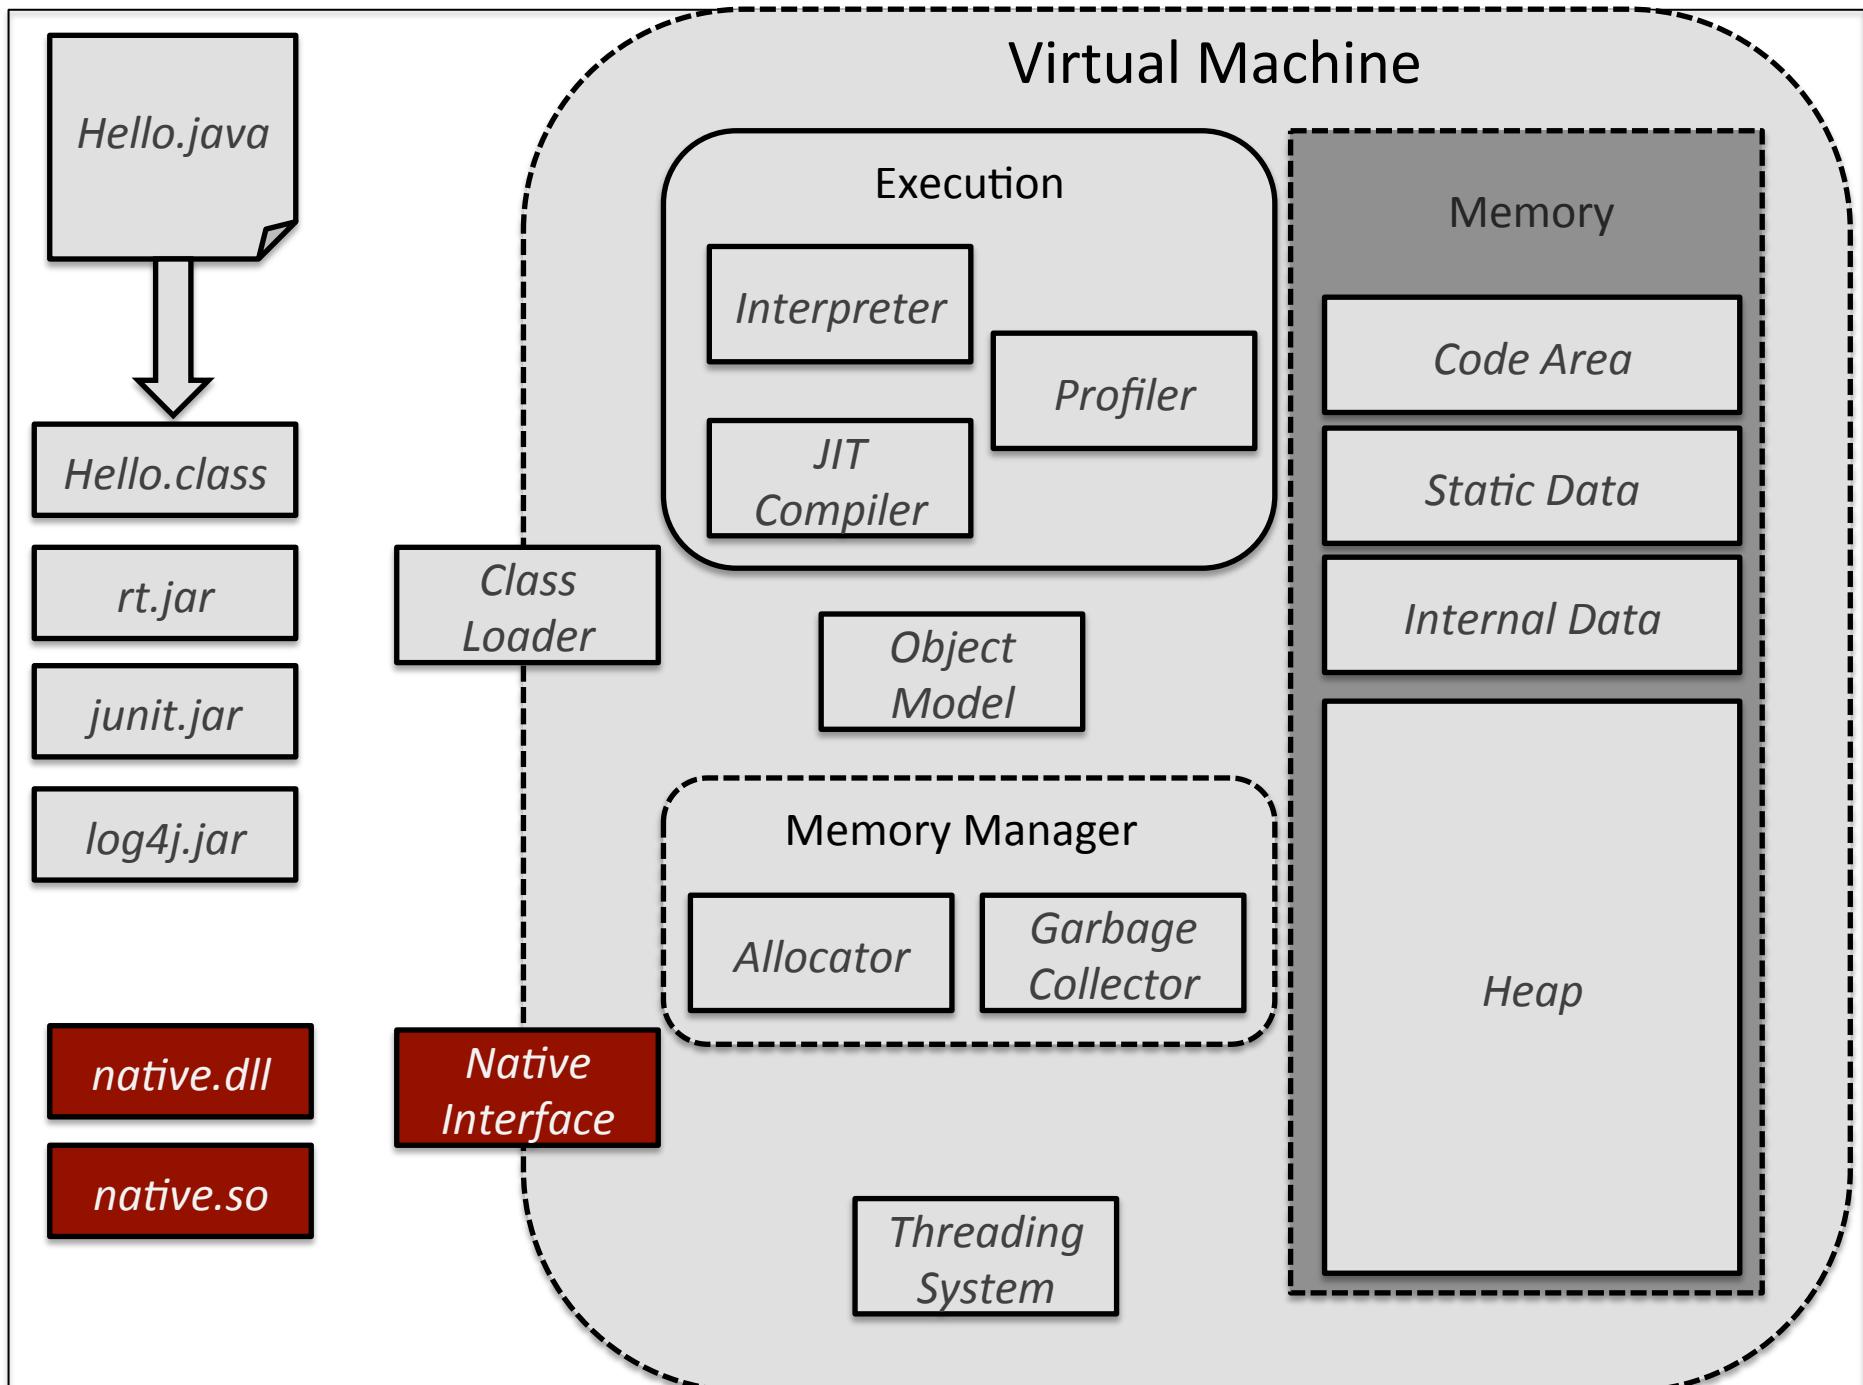
- Lots of testing
 - Gave them some confidence in their approach
- Knew the environment enough to make it work
 - Understood the problem and potential solutions
 - Dived into the source to convince themselves
- Shipped the app

Take-Away Lessons

- Know your platform ahead of time
 - Earlier knowledge of Dalvik internals could help
- Understand the business requirements
 - A temporary fix can be useful at times
 - Getting to market is important
- Manual testing has its place
- If you hack, don't blog about it...

Native Code II



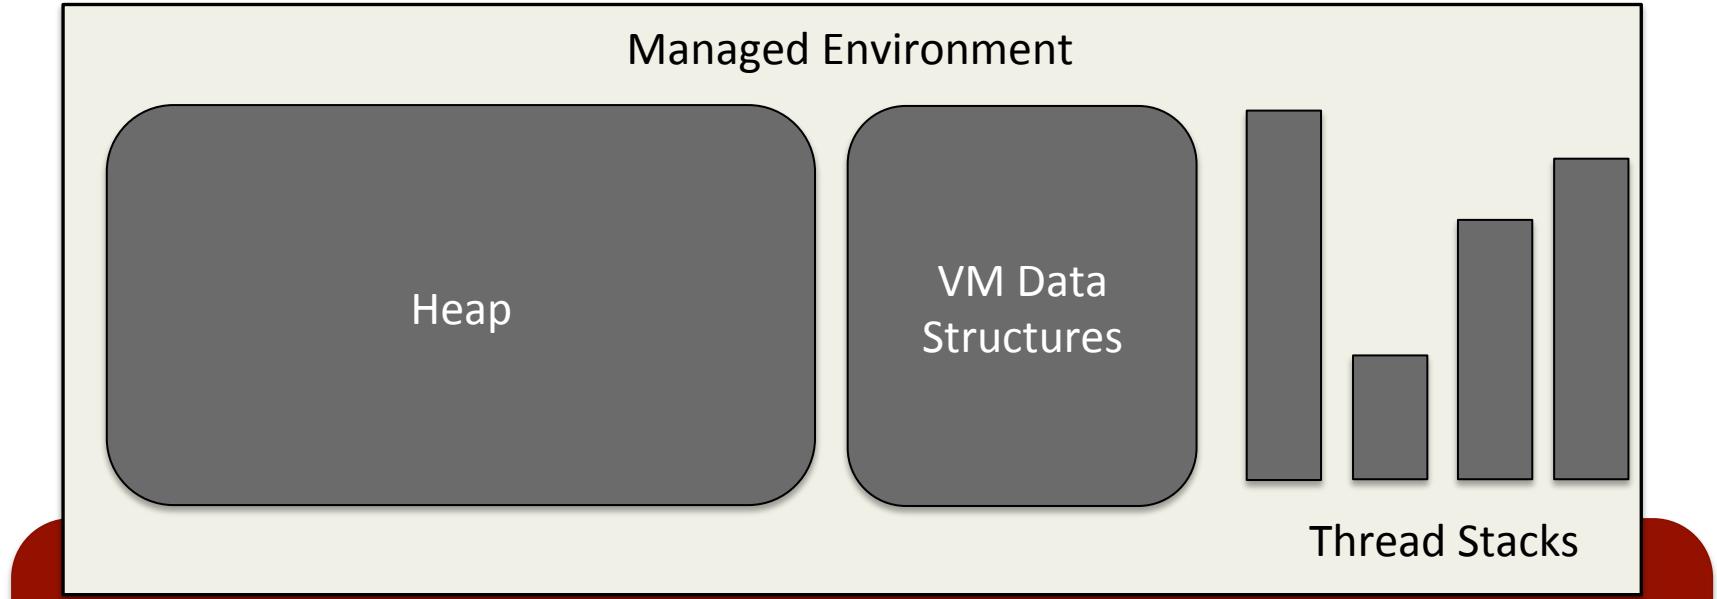


Why Do We Need Native Code

- Need to integrate with existing code
 - Source and build tools may not be available
 - Modifying existing code may be impractical
- Need performance for critical section
 - May want to hand-optimize in assembly code
- Want to do things not allowed by the VM
 - Take advantage of custom platform features

Foreign Function Interface

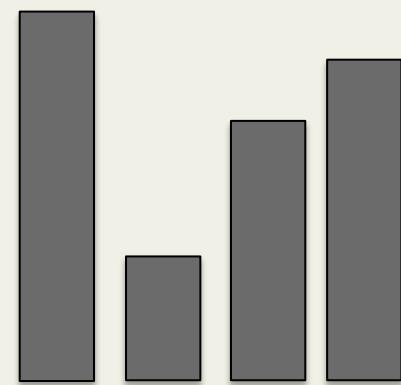
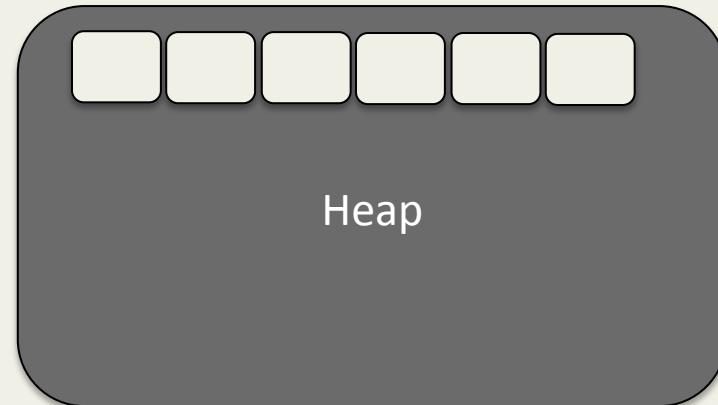
- Integrates managed code with native
 - Imposes a structure on native code
 - Ensures that native code respects VM conventions
- Java Native Interface defines boundary
 - Strict control of interactions between systems
- JNIEnv object provides access to VM
 - Double indirection lets us change implementation



Memory Access

- Native code uses raw pointers
 - Prevents moving garbage collection
 - Assumes object layout

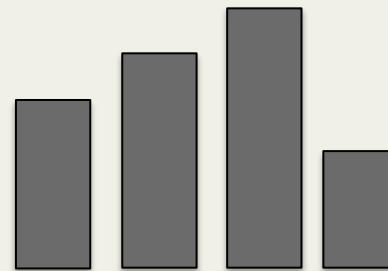
Managed Environment



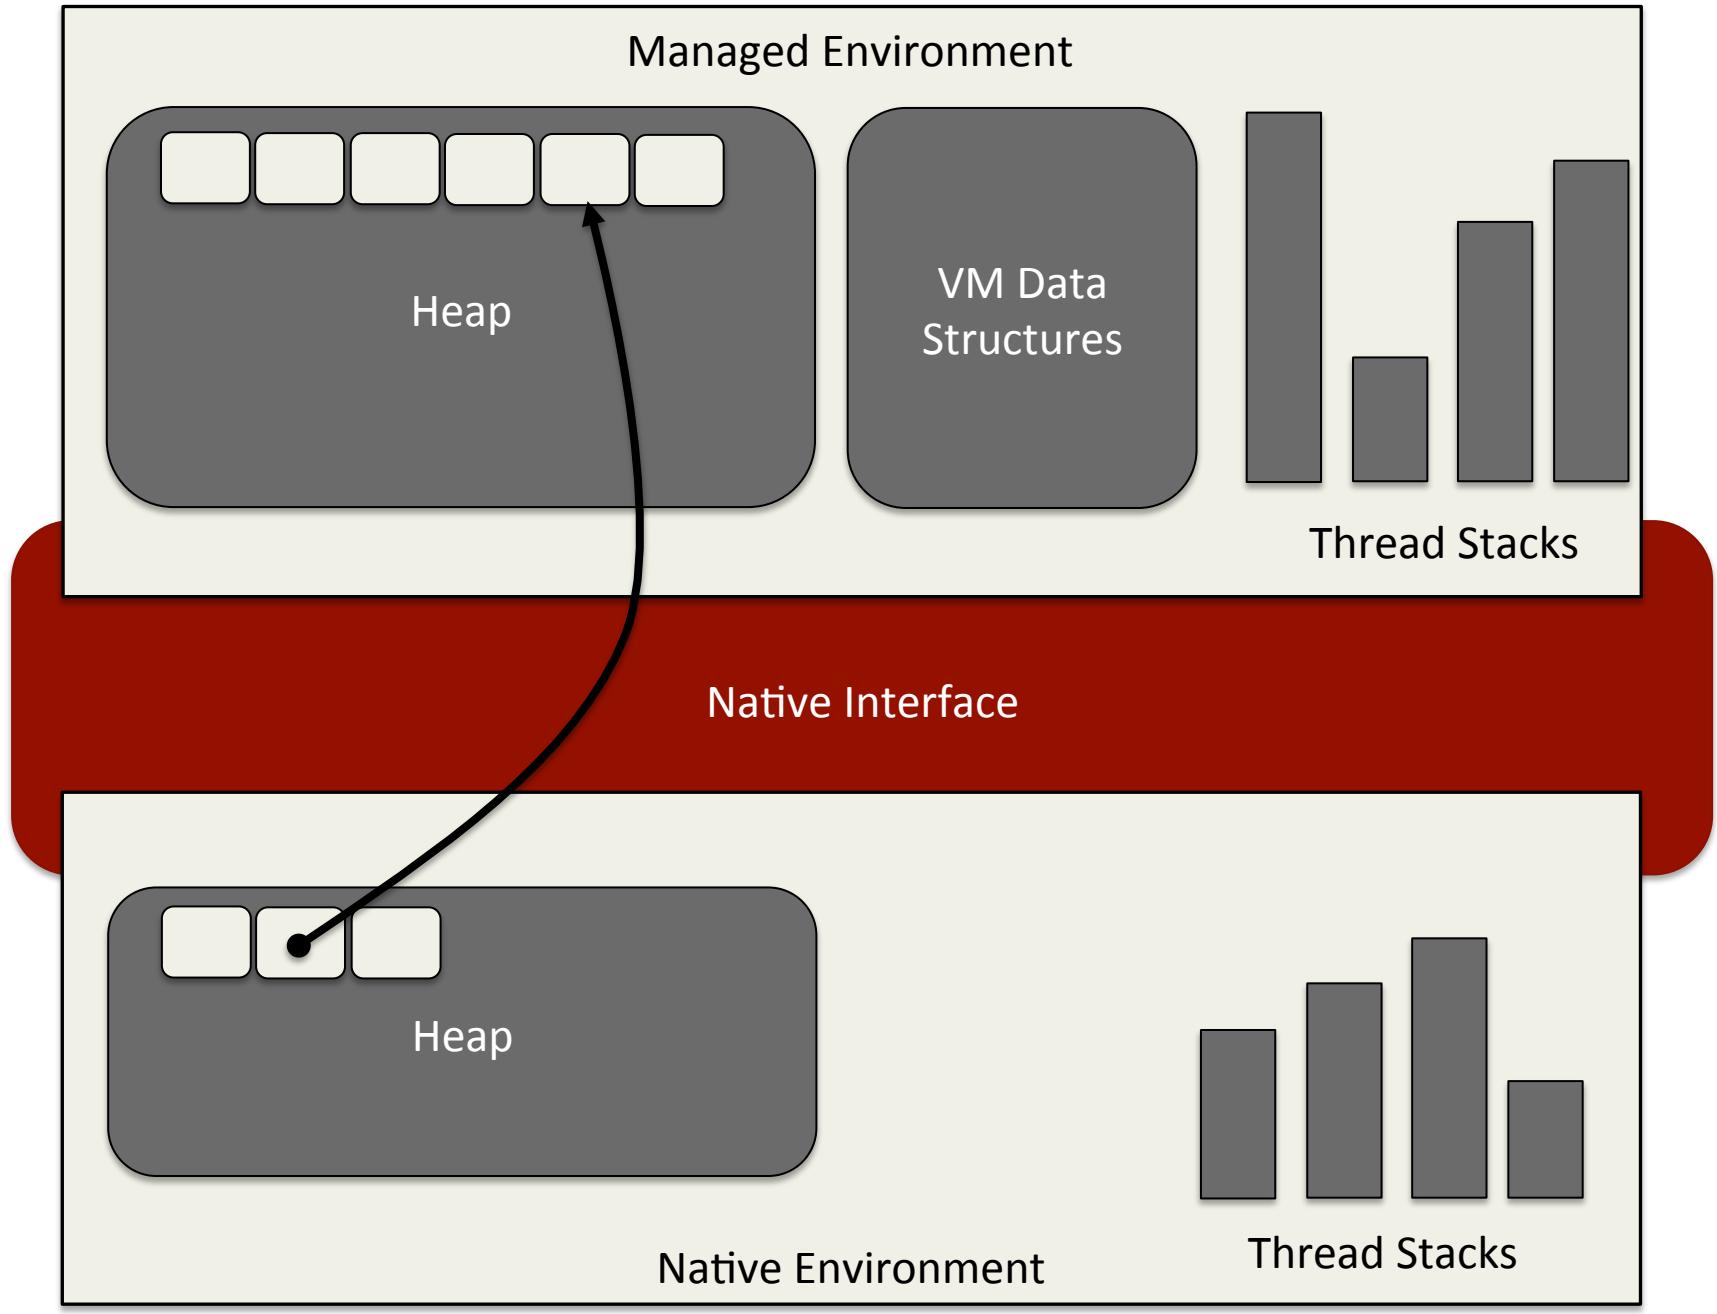
Native Interface

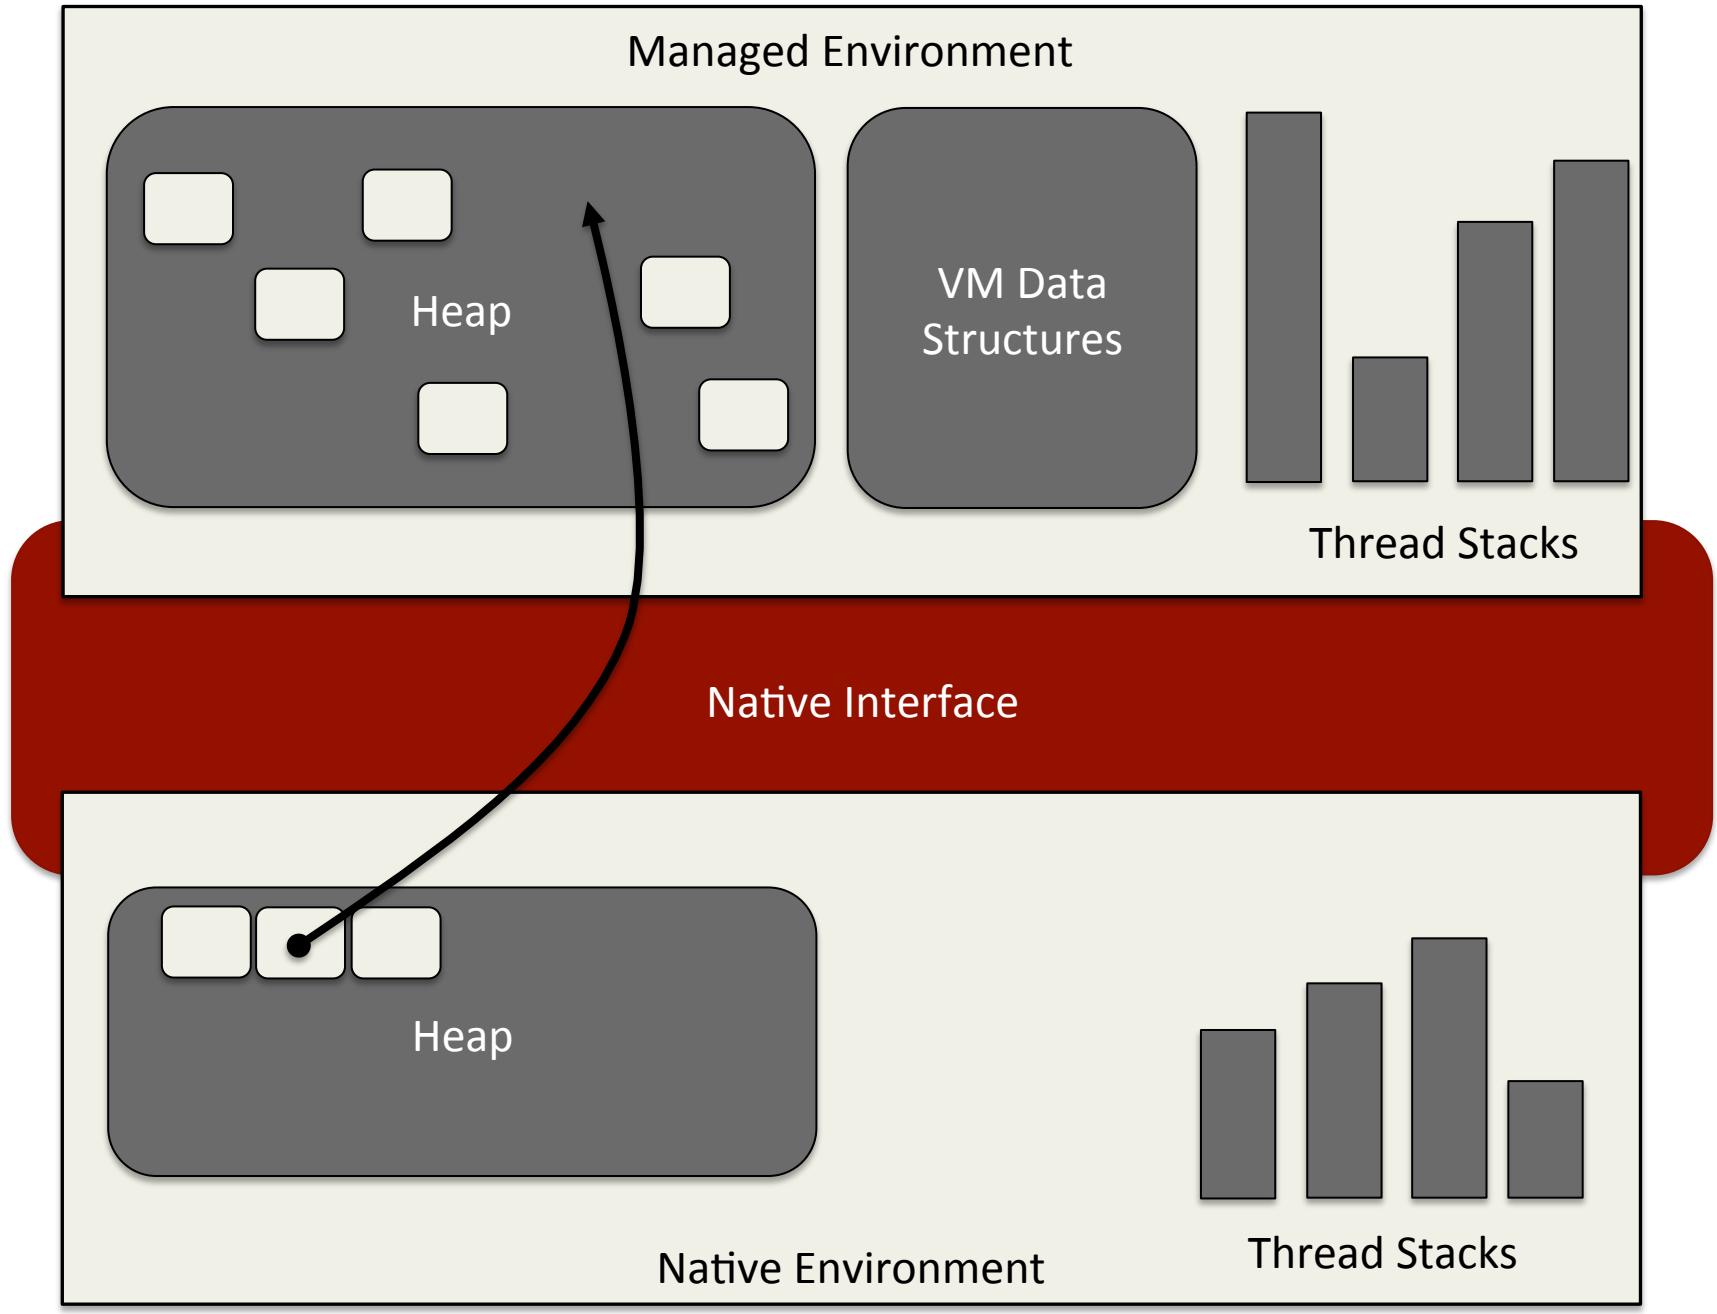


Native Environment



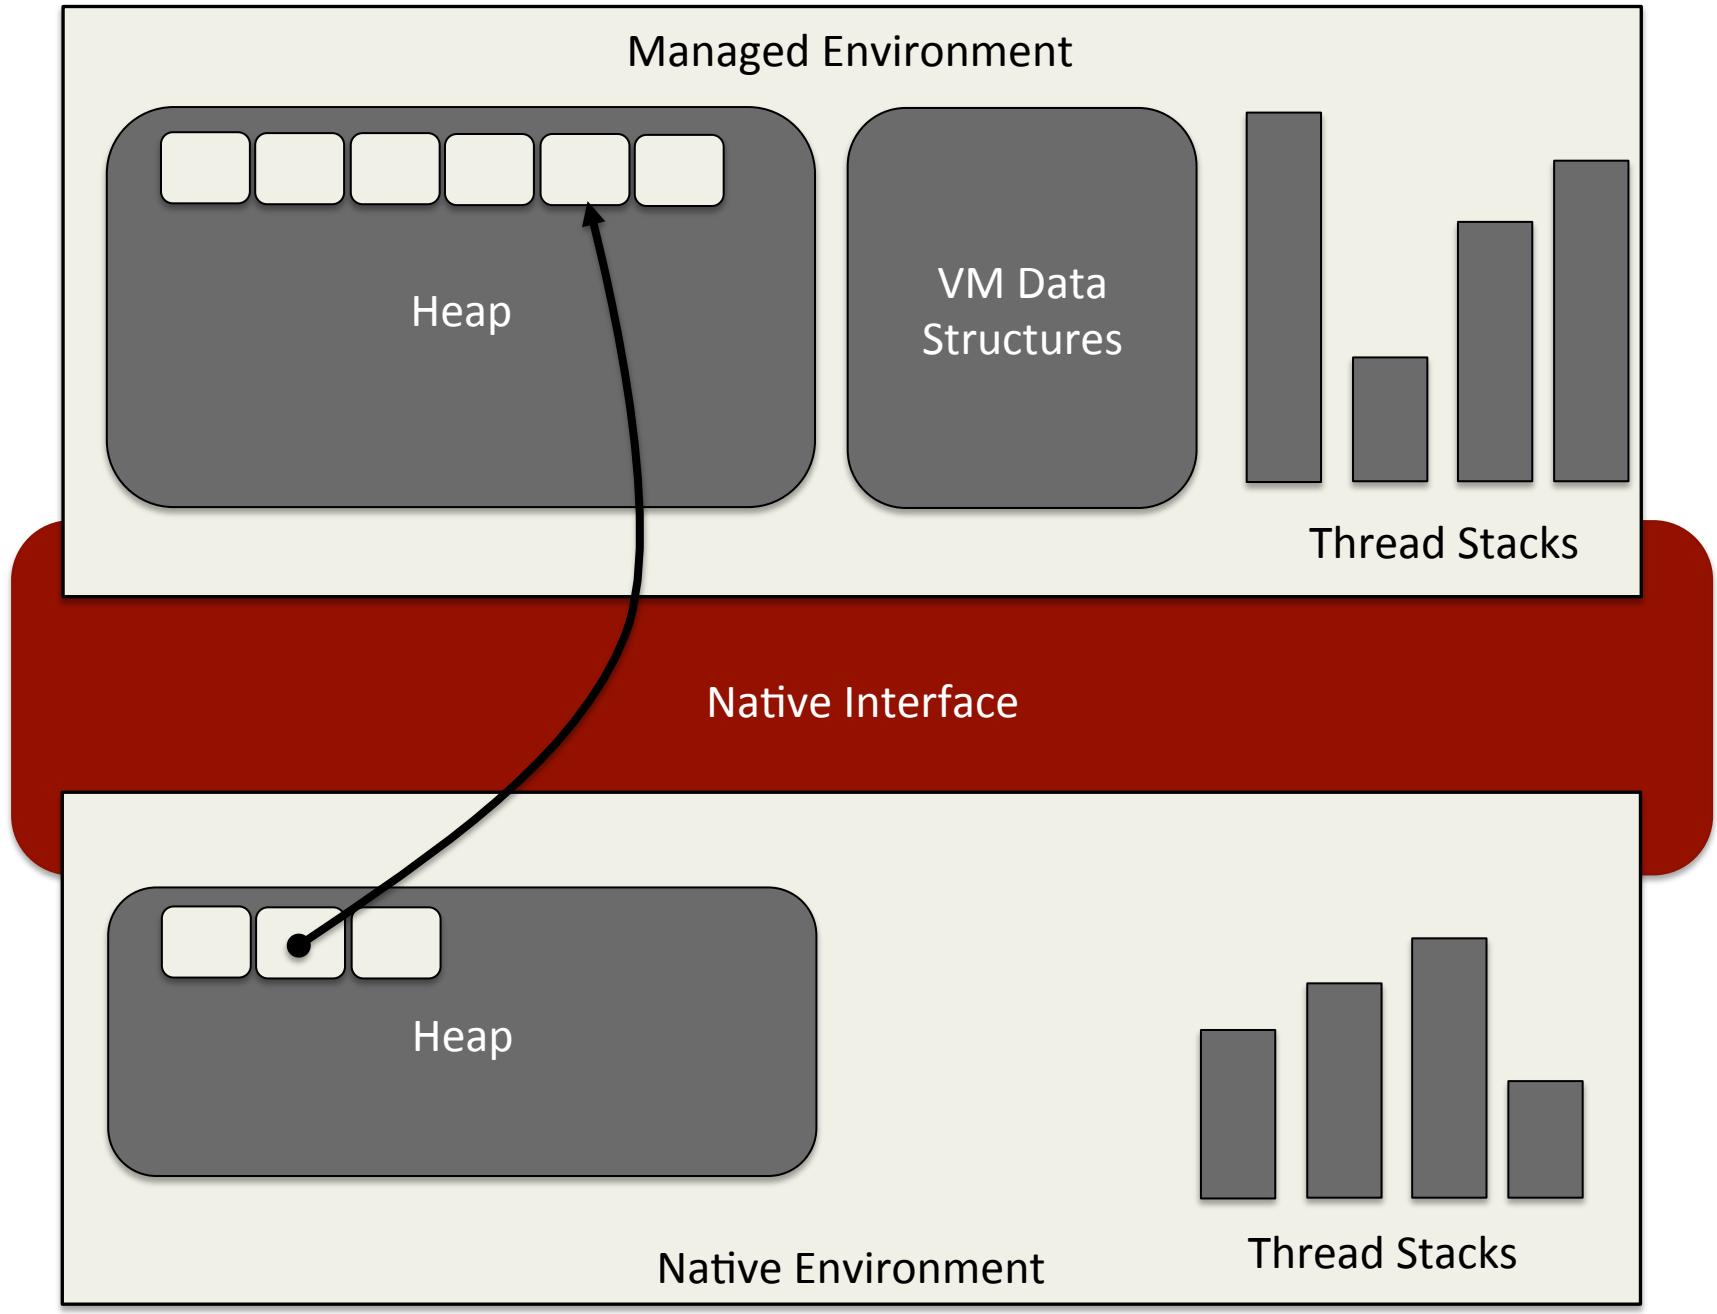
Thread Stacks

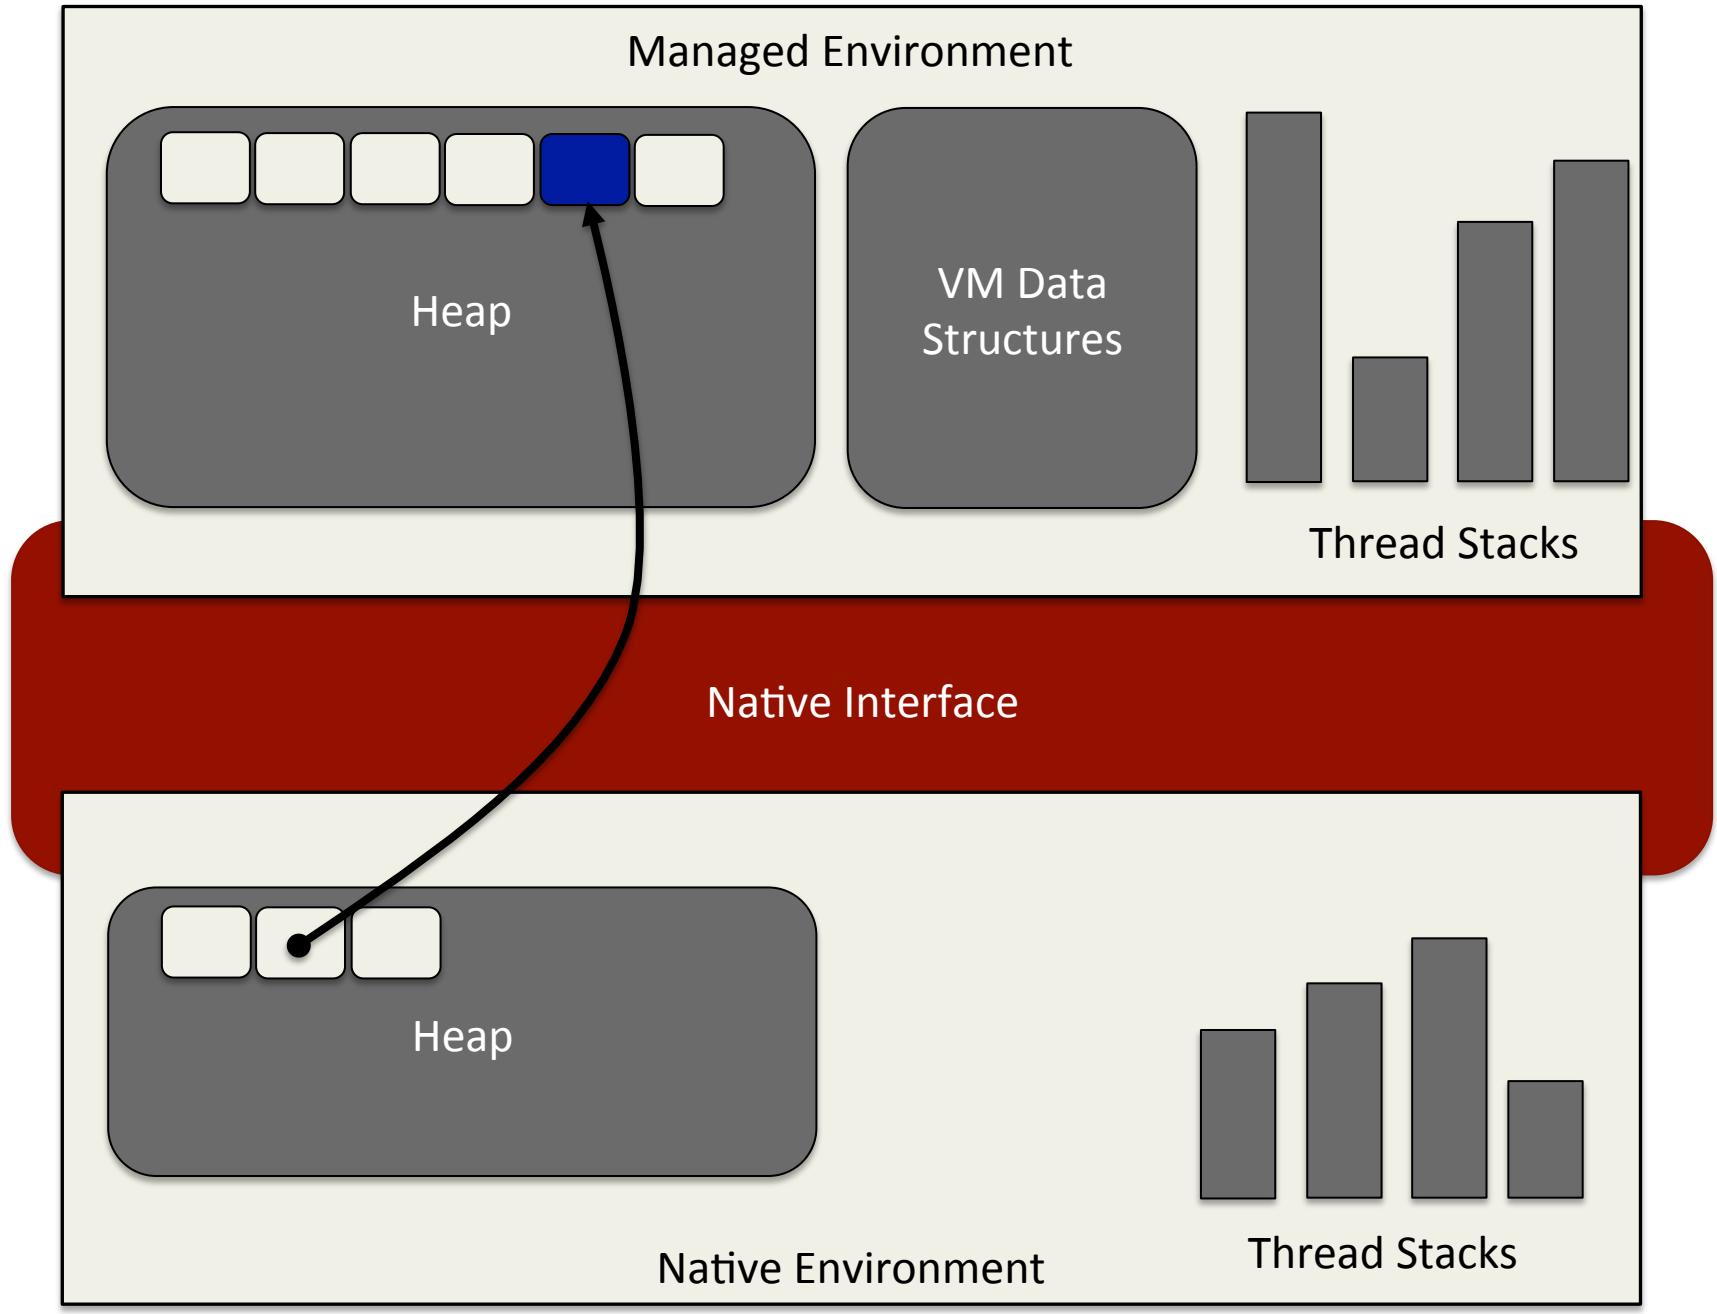


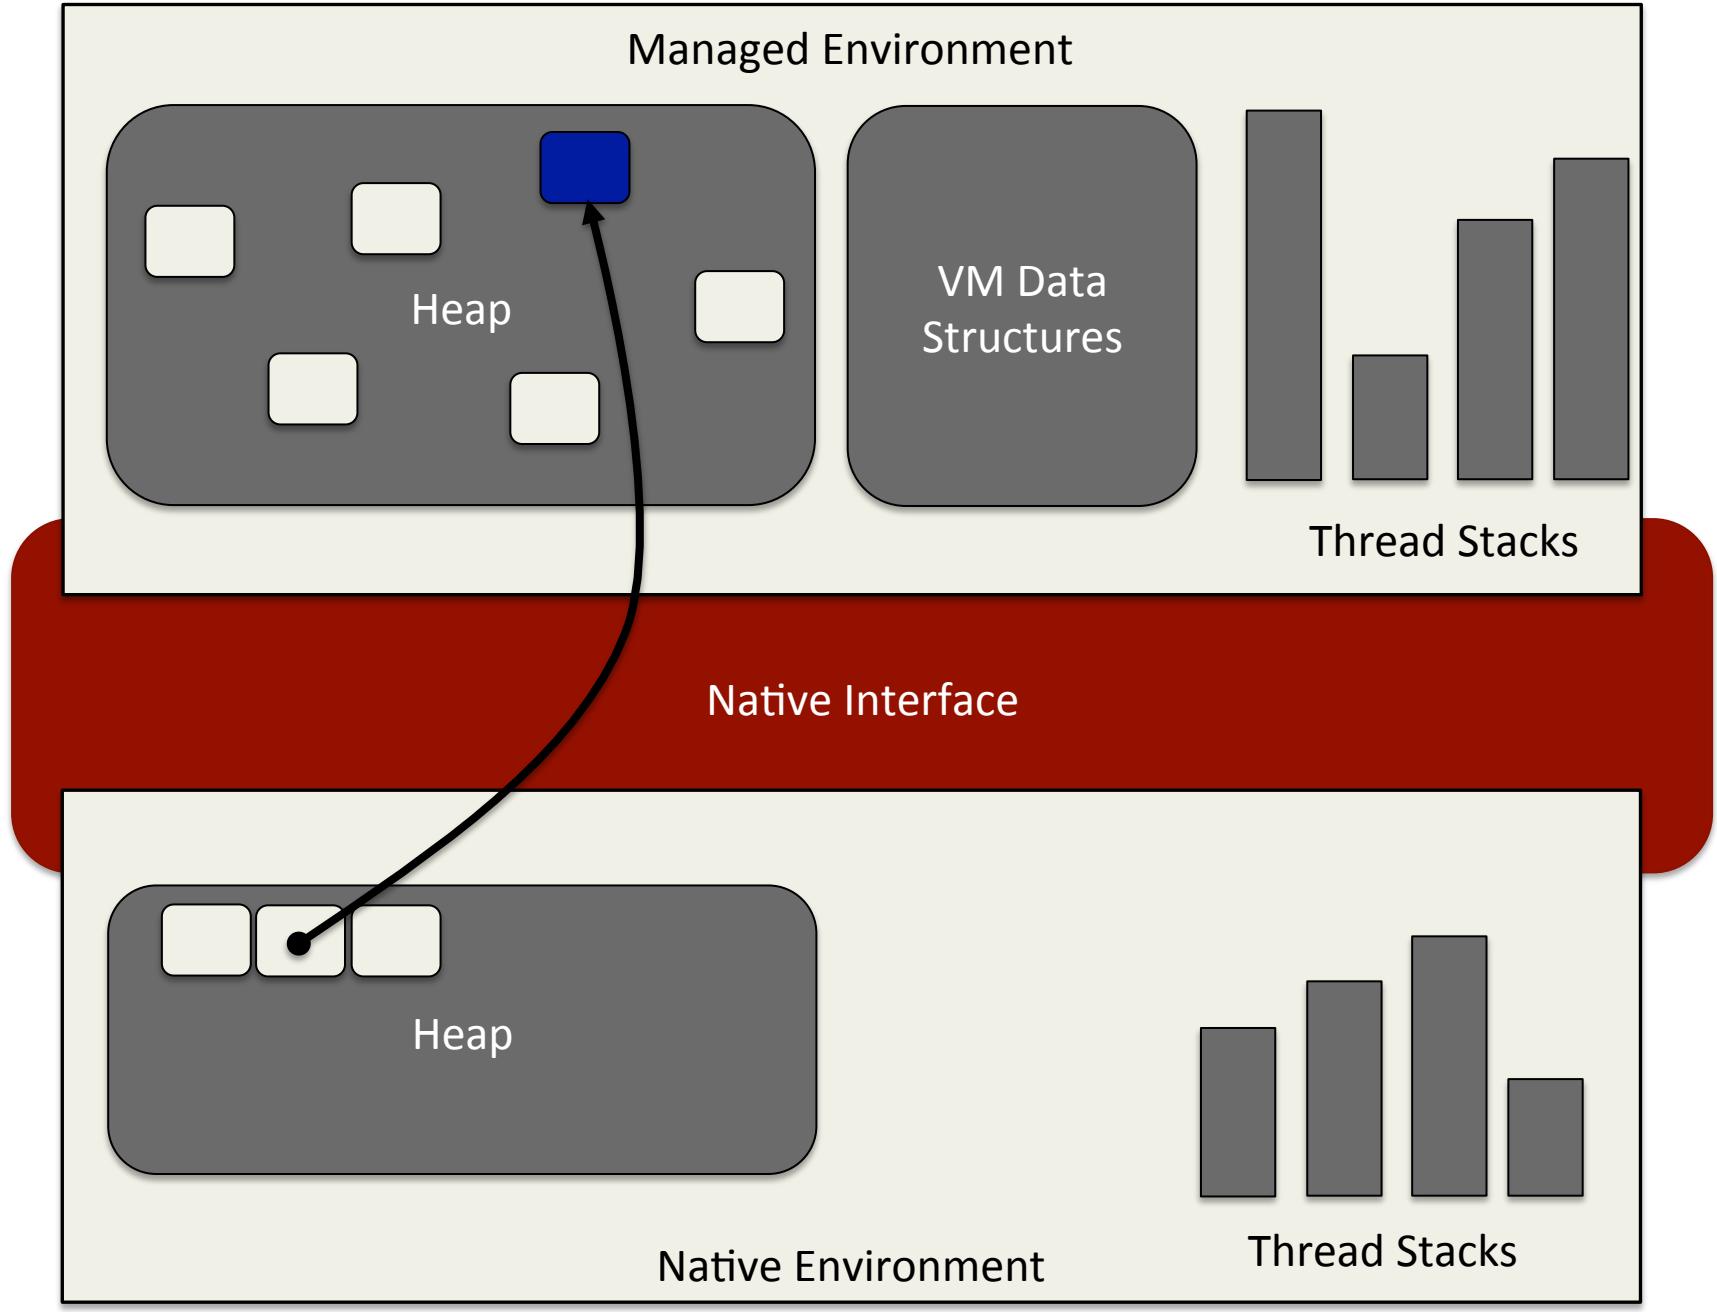


Memory Access

- Native code uses raw pointers
 - Prevents moving garbage collection
 - Assumes object layout
- One possible solution is pinning
 - Designate an object as unmovable
 - Restricts implementation of the VM

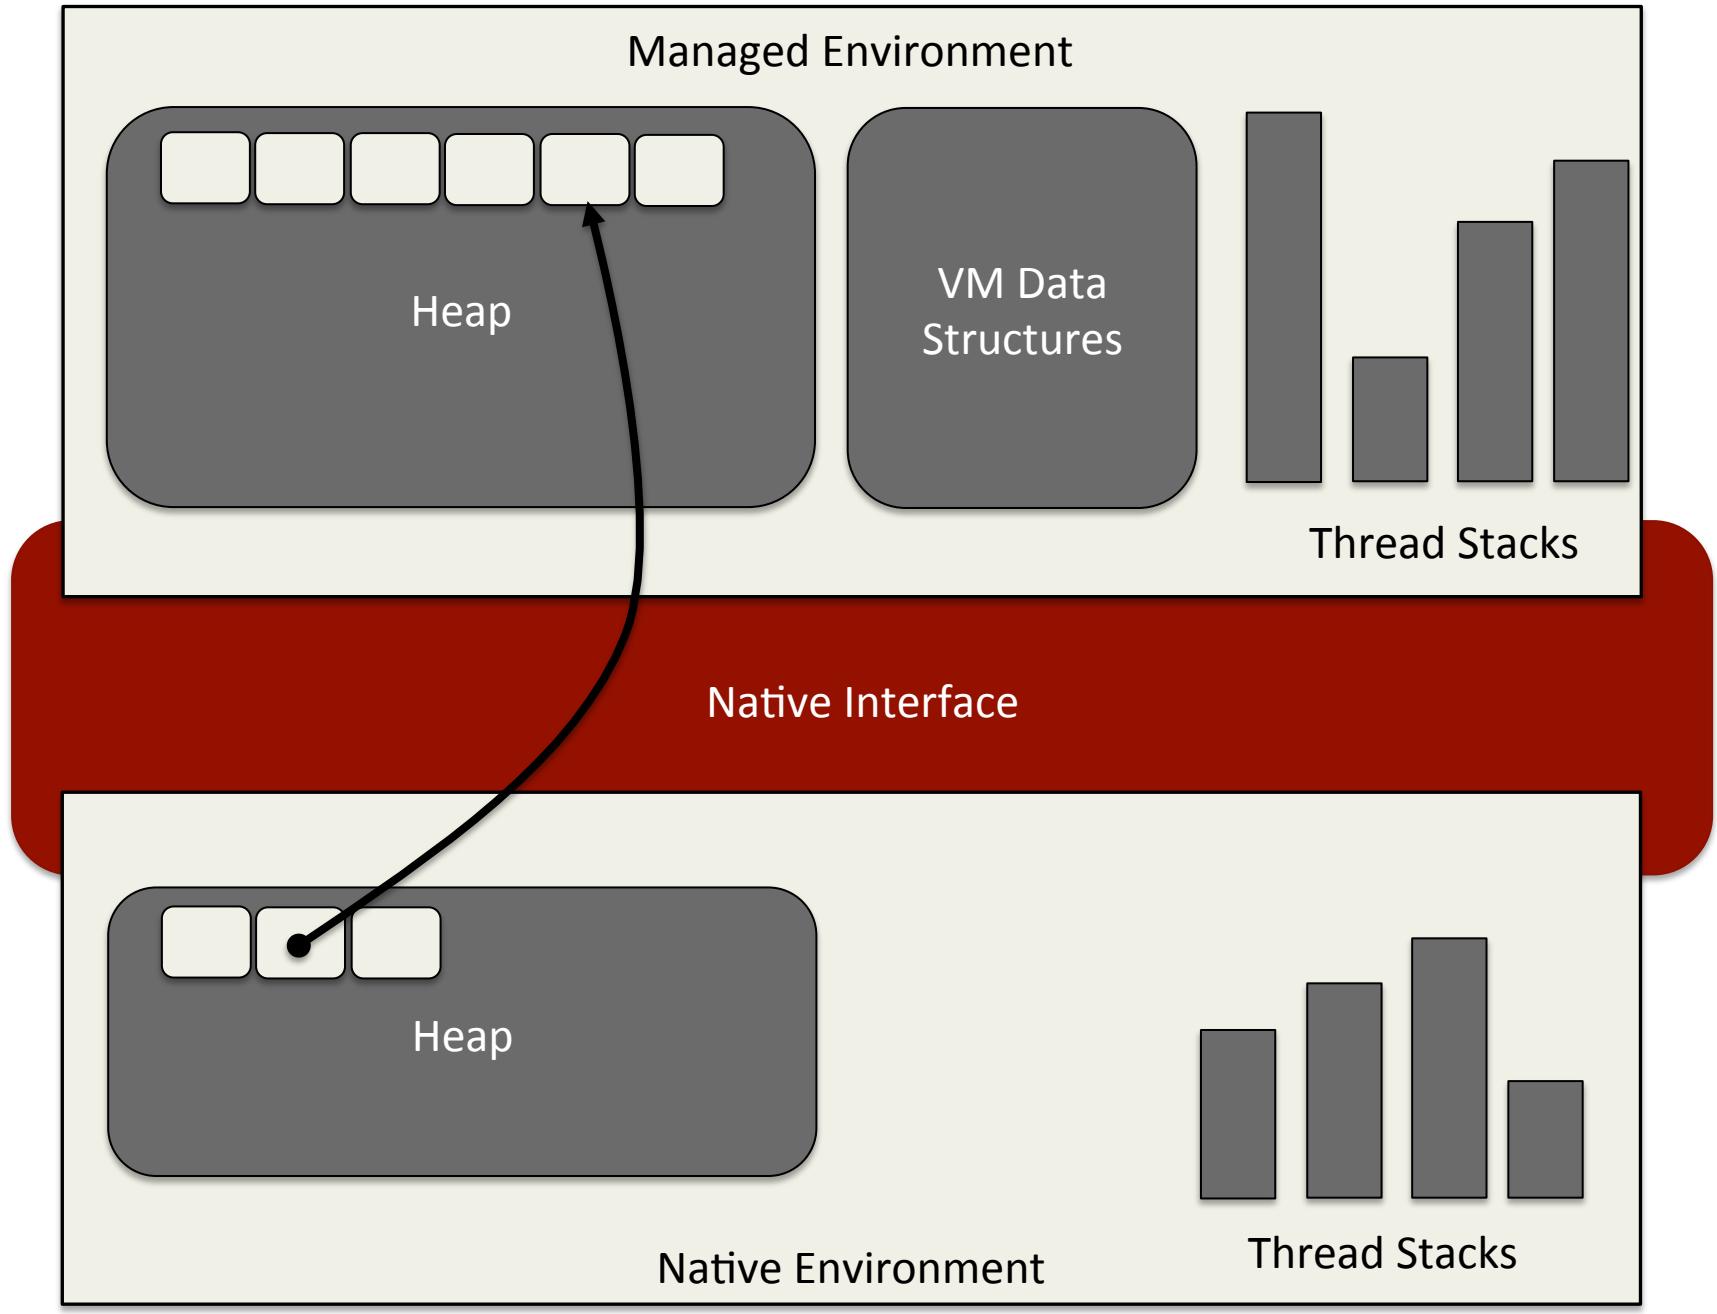


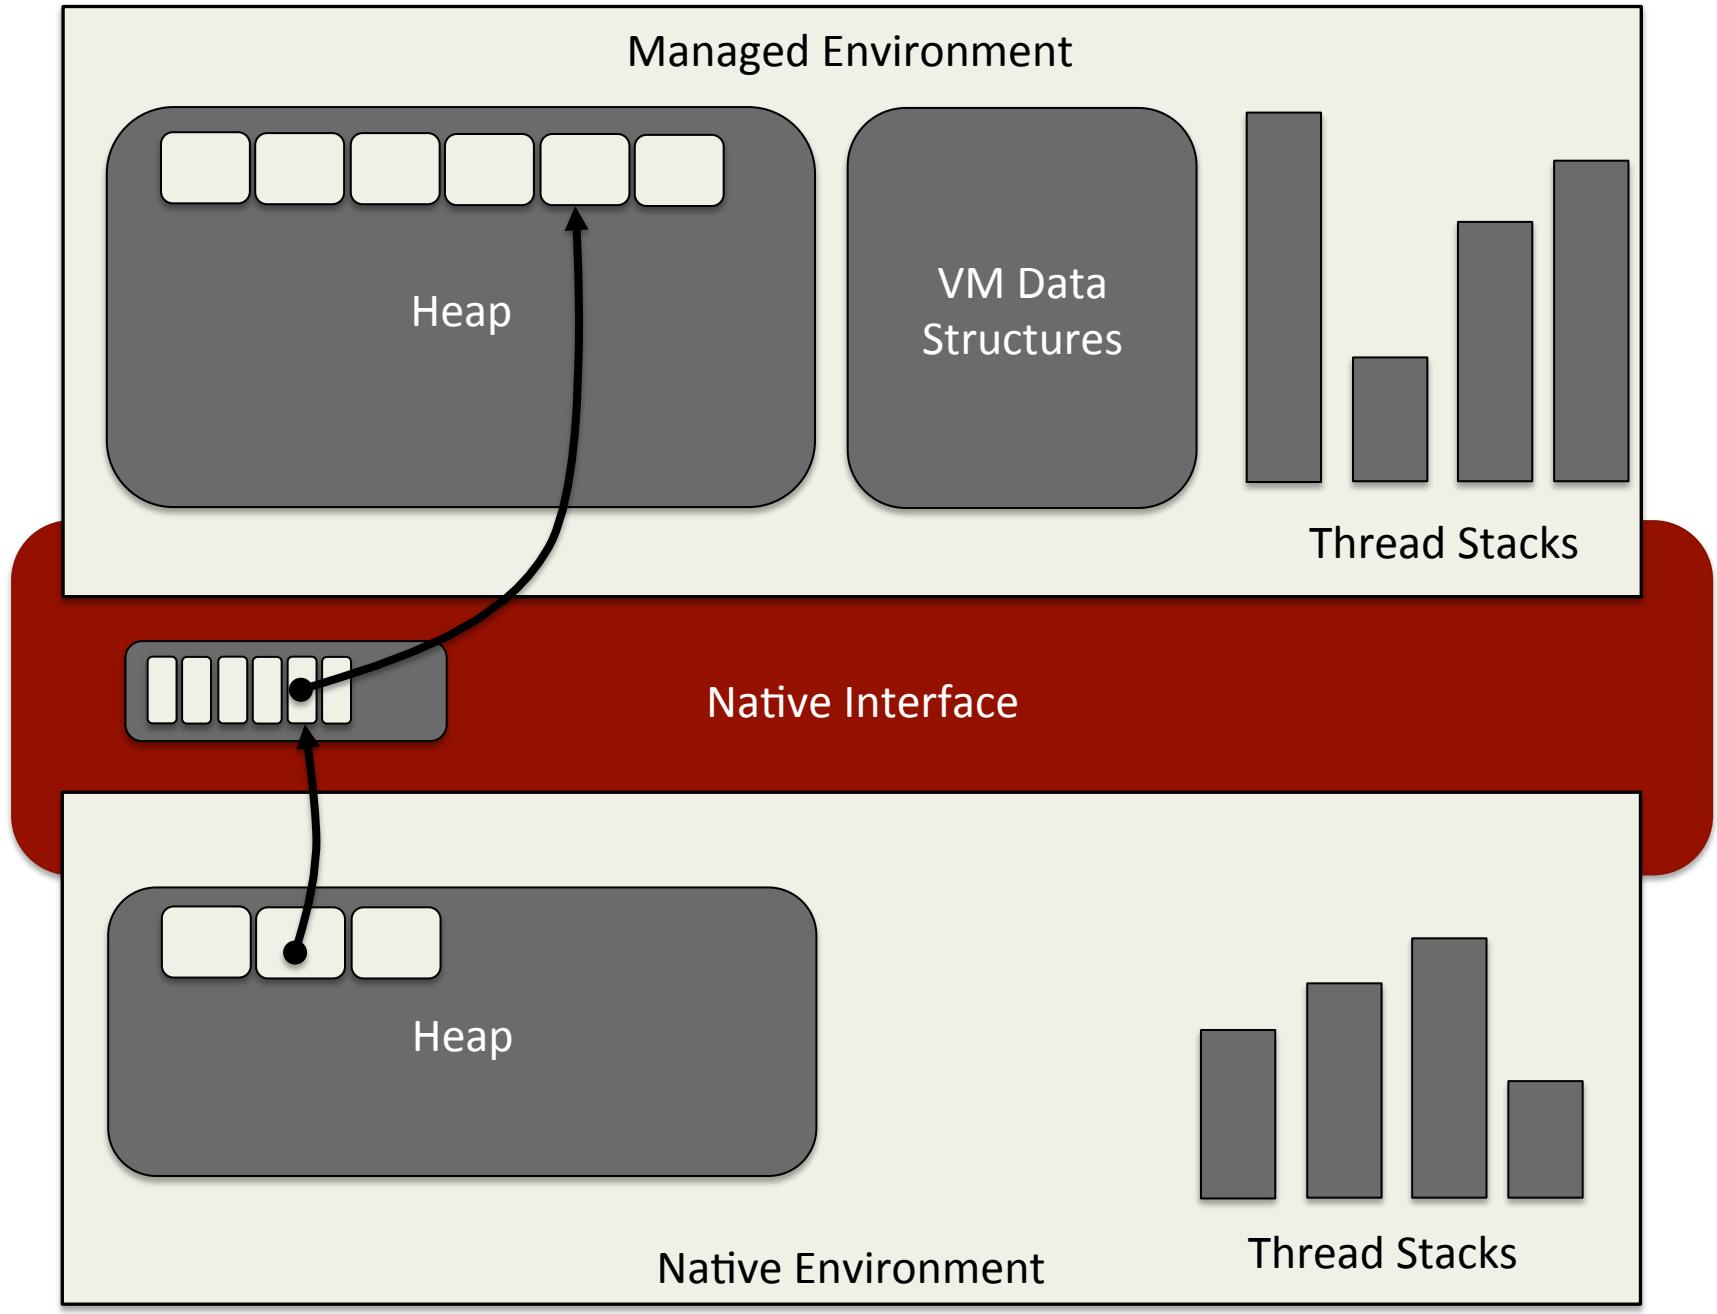


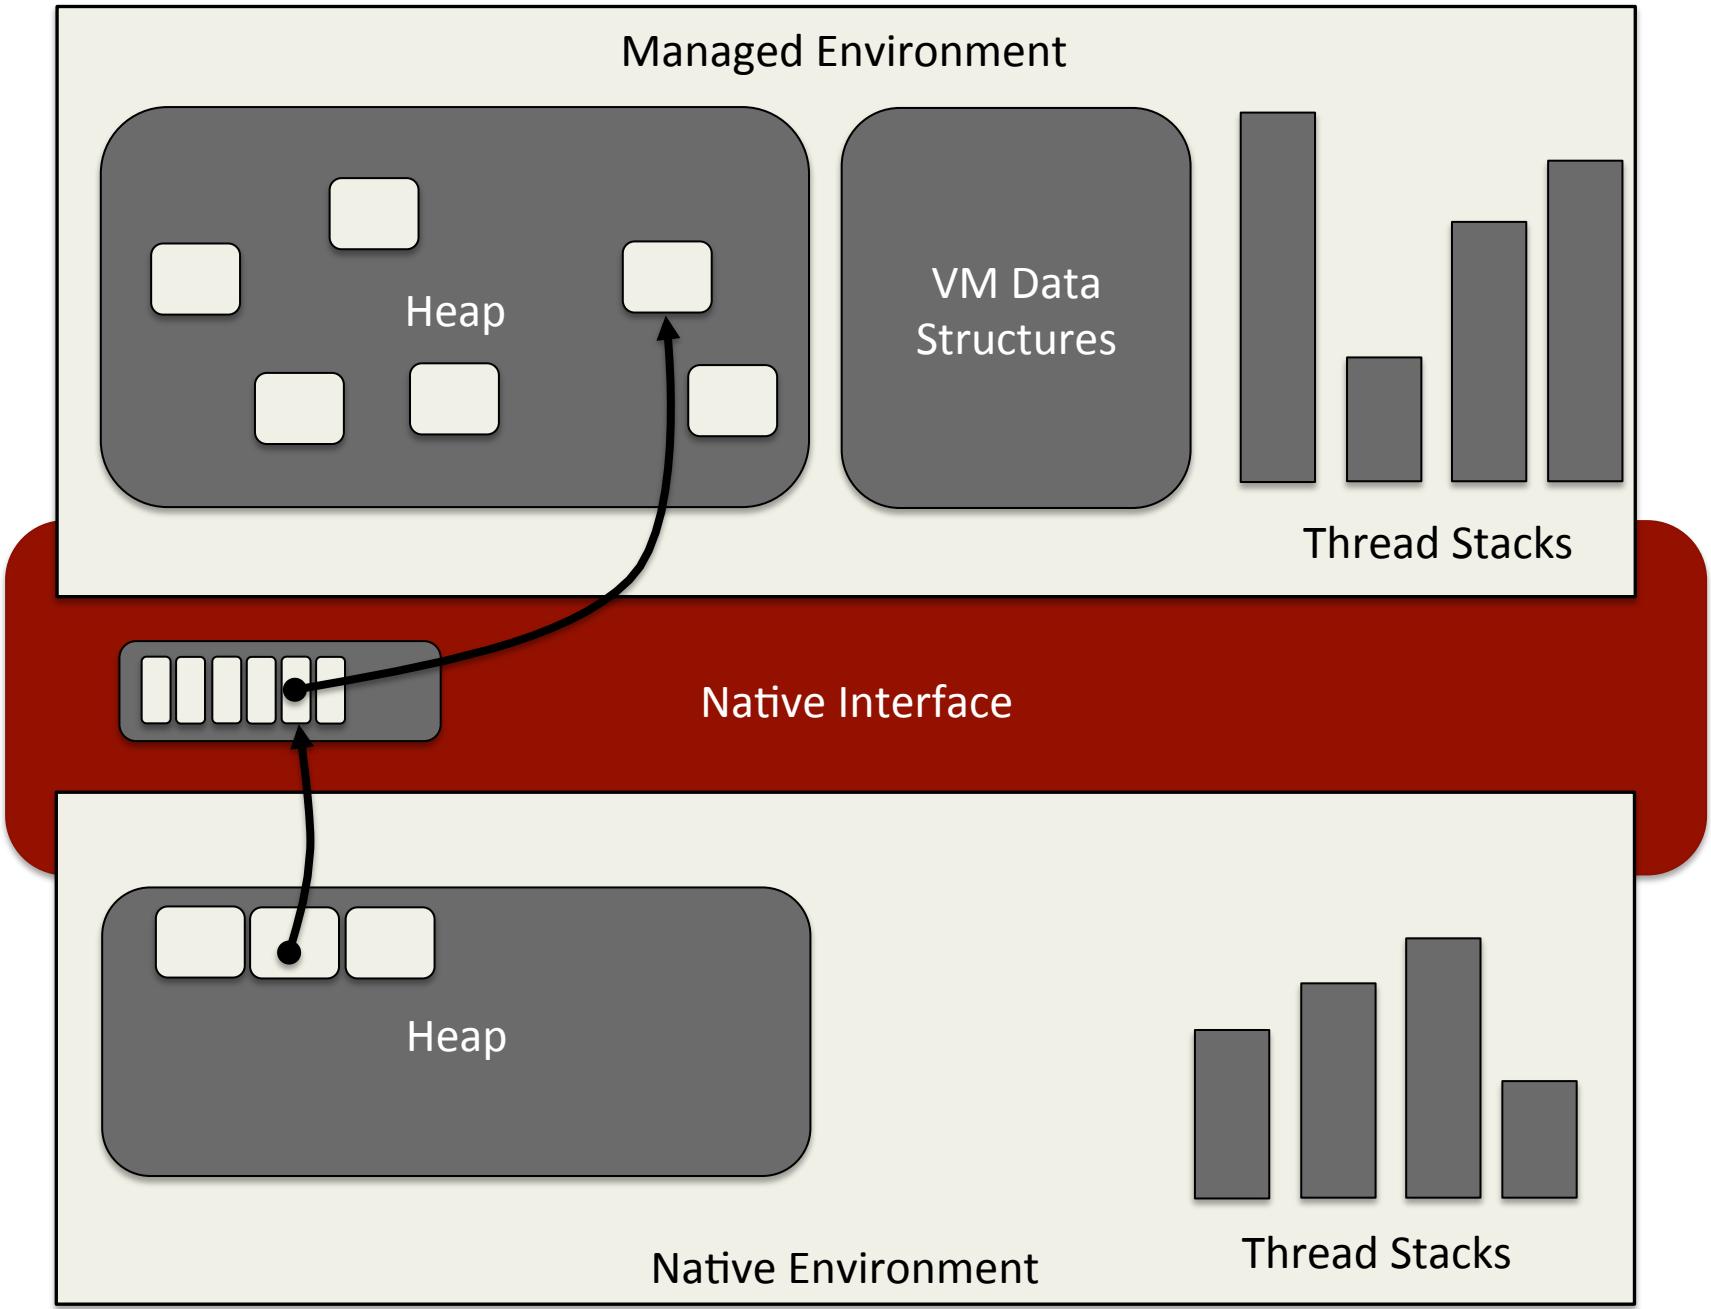


Memory Access

- Native code uses raw pointers
 - Prevents moving garbage collection
 - Assumes object layout
- One possible solution is pinning
 - Designate an object as unmovable
 - Restricts implementation of the VM
- Another approach is to use handles







Handles

- We've talked about handles before
 - Objects can be moved without updating pointers
 - All accesses involve an extra indirection
- In this case the trade-off is better
 - Expect overhead when going from native to VM
 - Keeps object model out of native code
 - Lets us move objects without modifying native

Reserving Objects

- Native needs to tell the VMs what object it uses
 - Don't want a handle for every object
 - Part of the native interface definition
- Objects can be reserved in two ways
 - Local references scoped to the native method
 - Global references need to be explicitly freed
- Global created by promoting local reference

Local References

- Reserved for the execution of a single method
 - Automatically released on method exit
- Local references are thread specific
 - Native code can't share between threads
 - Could lead to dangling reference
- Local references can be released manually
 - Prevent locking of large objects
 - Avoid locking large numbers of objects

Global Reference

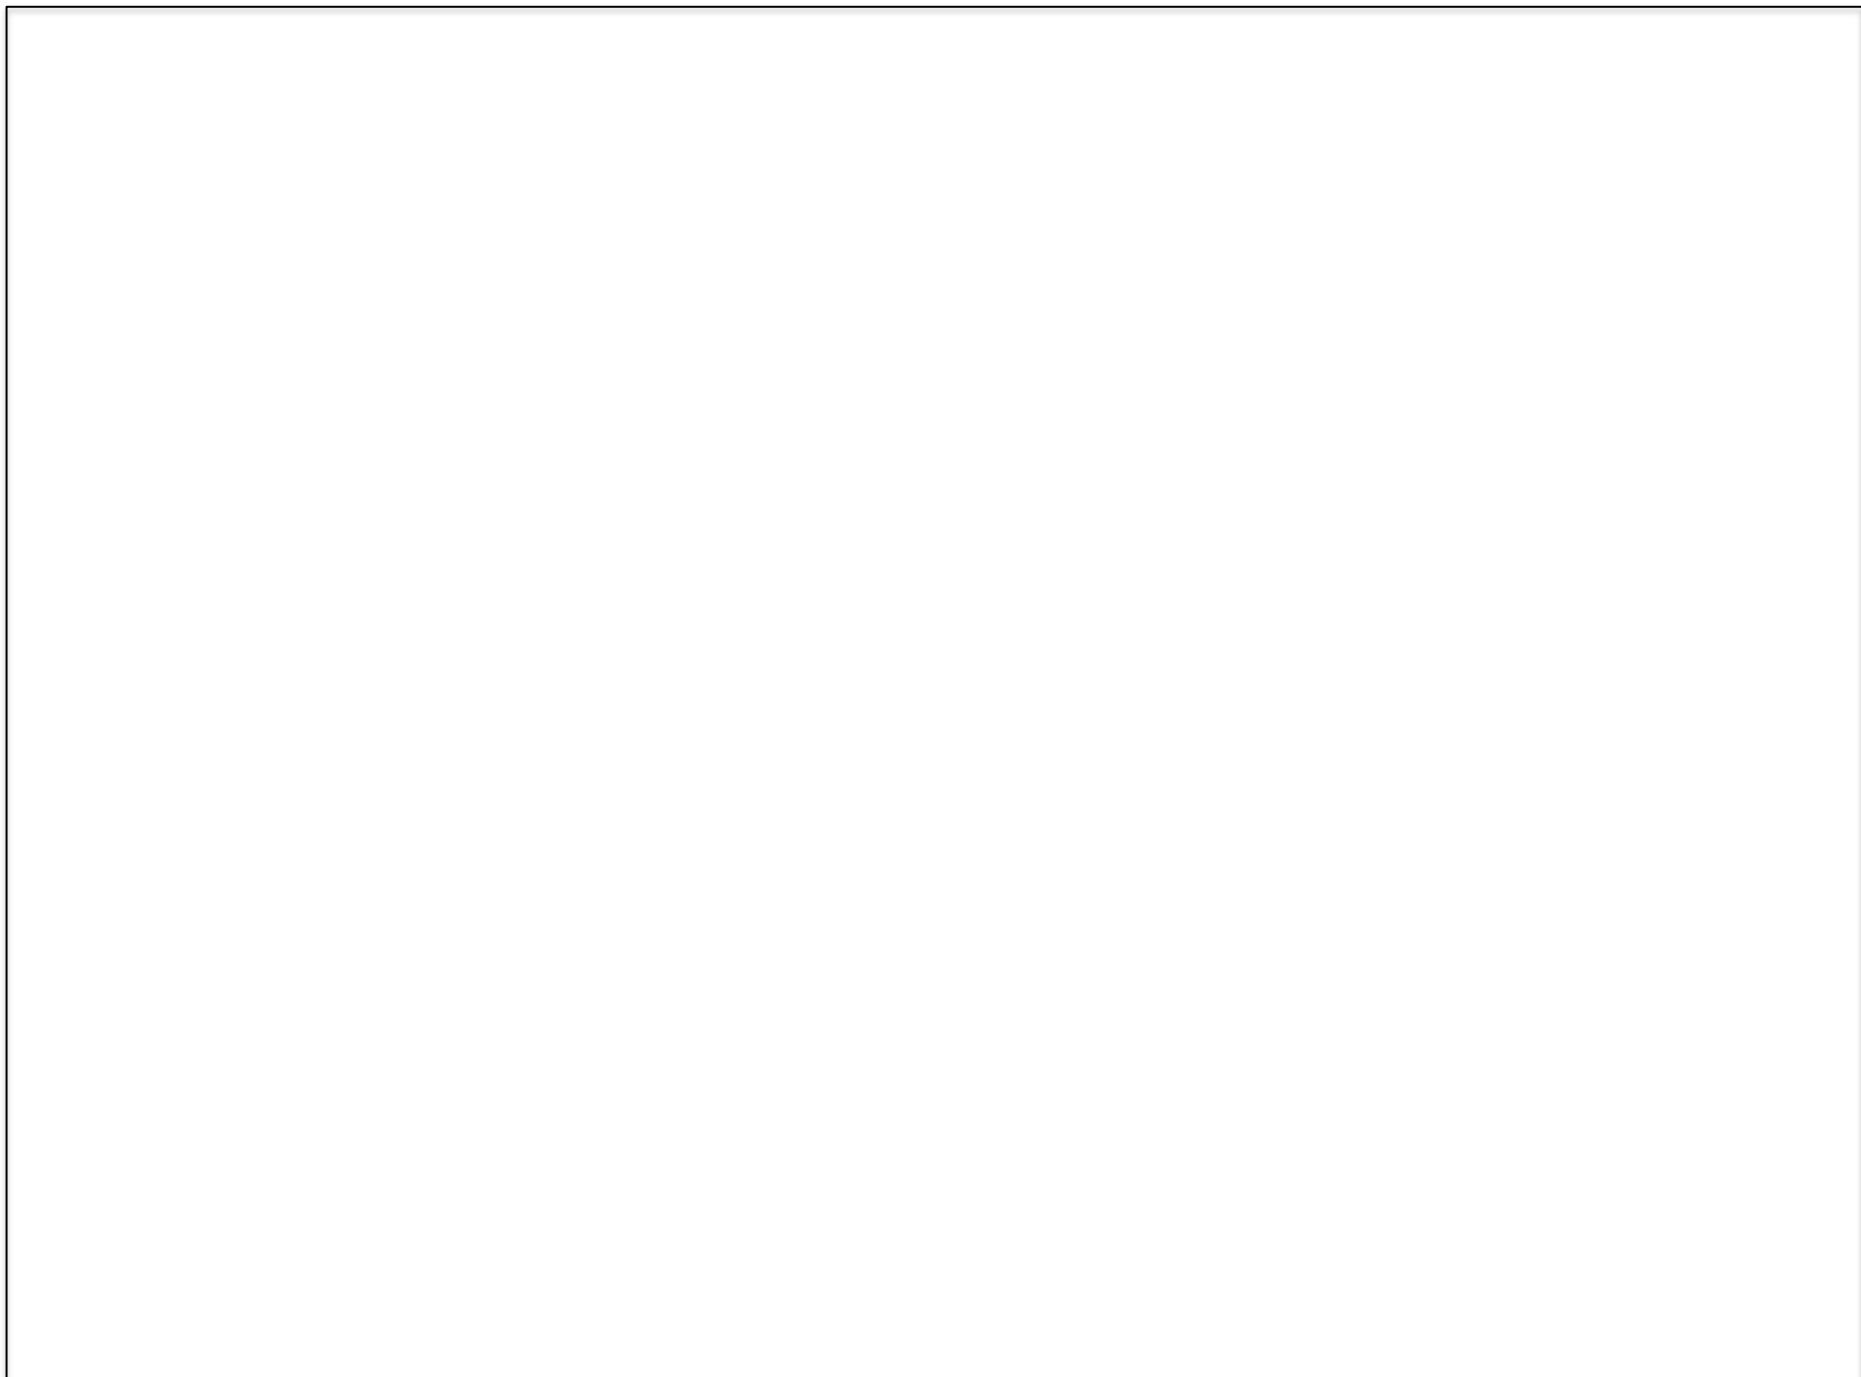
- Created by upgrading a local reference
- Native code takes responsibility for the reference
 - VM will never assume the reference is released
 - Native must explicitly release
- Can be shared between threads
- Potential source of memory leaks

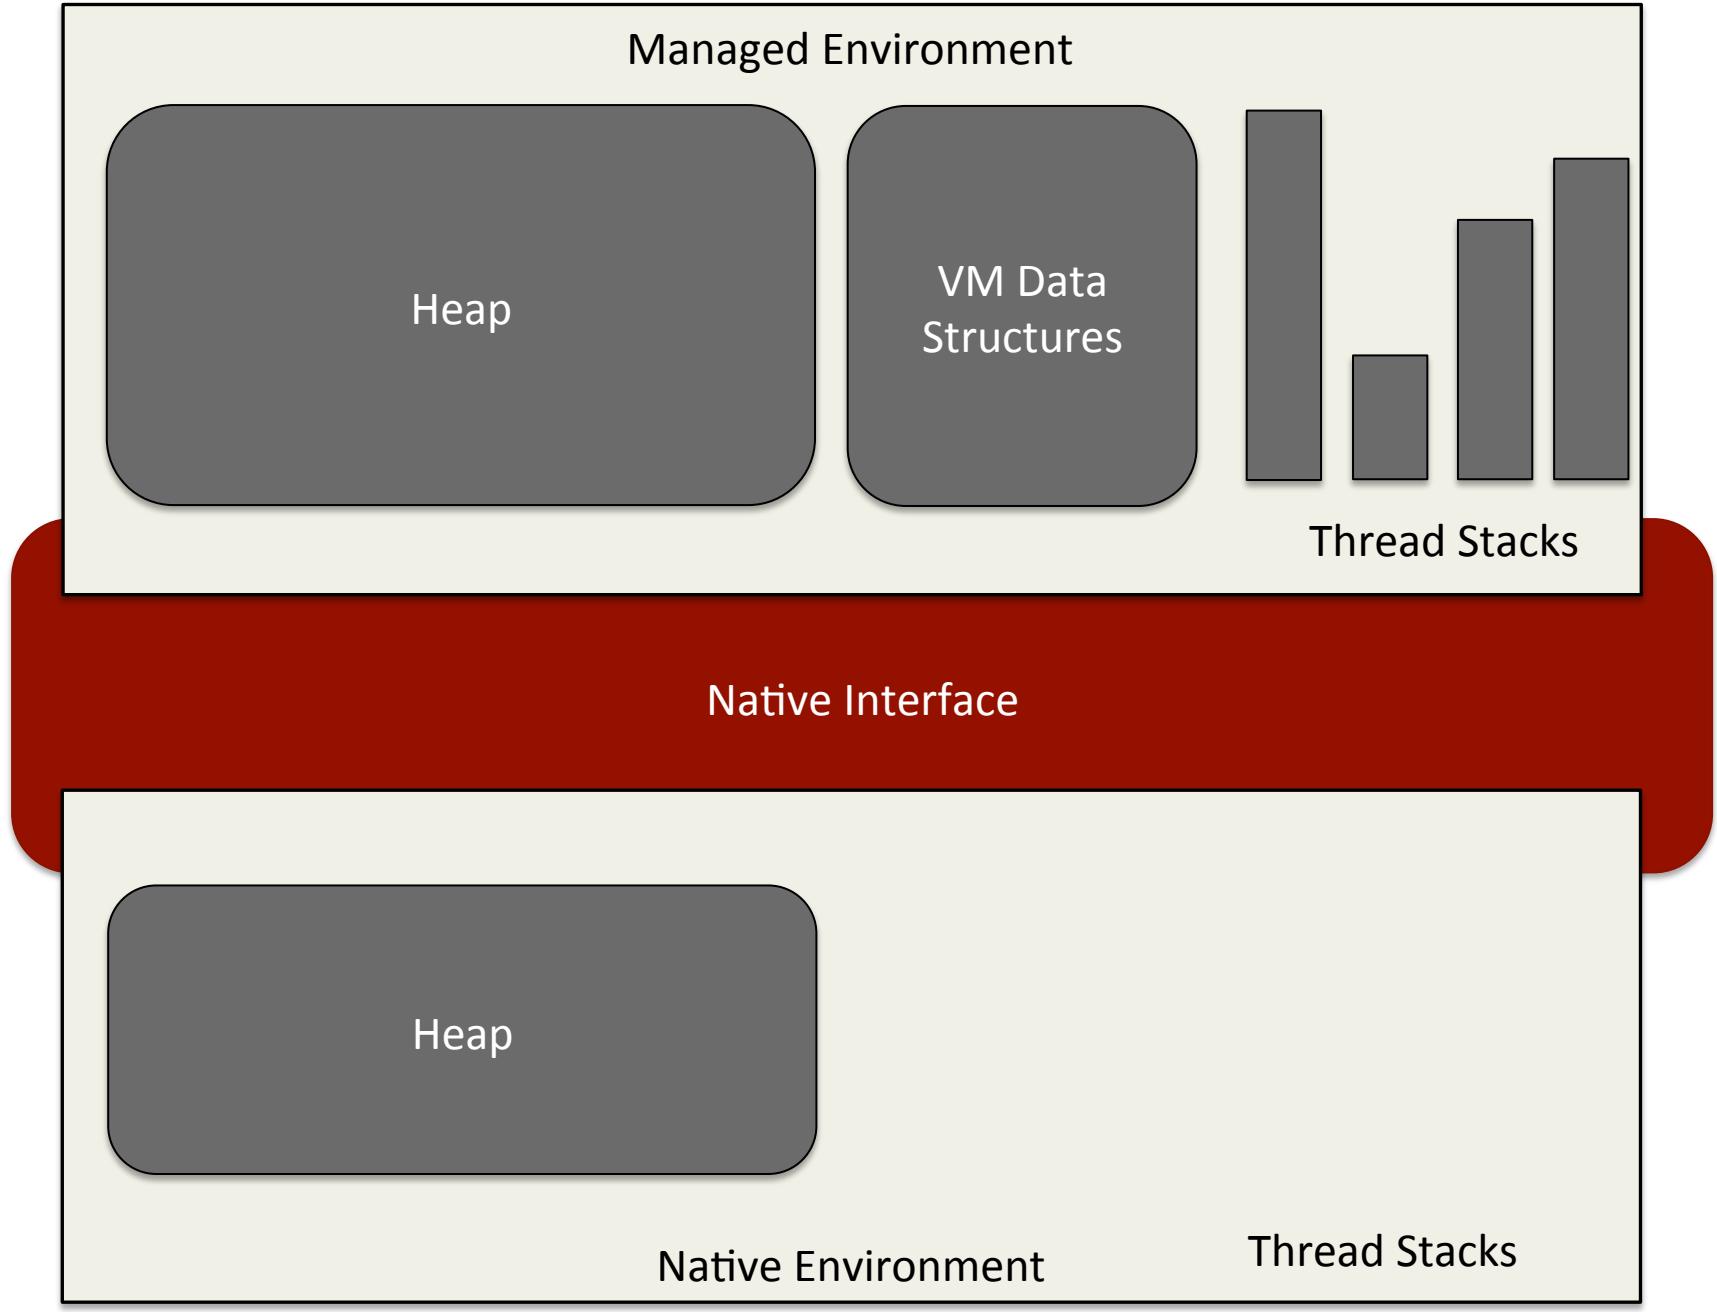
Handle Registry

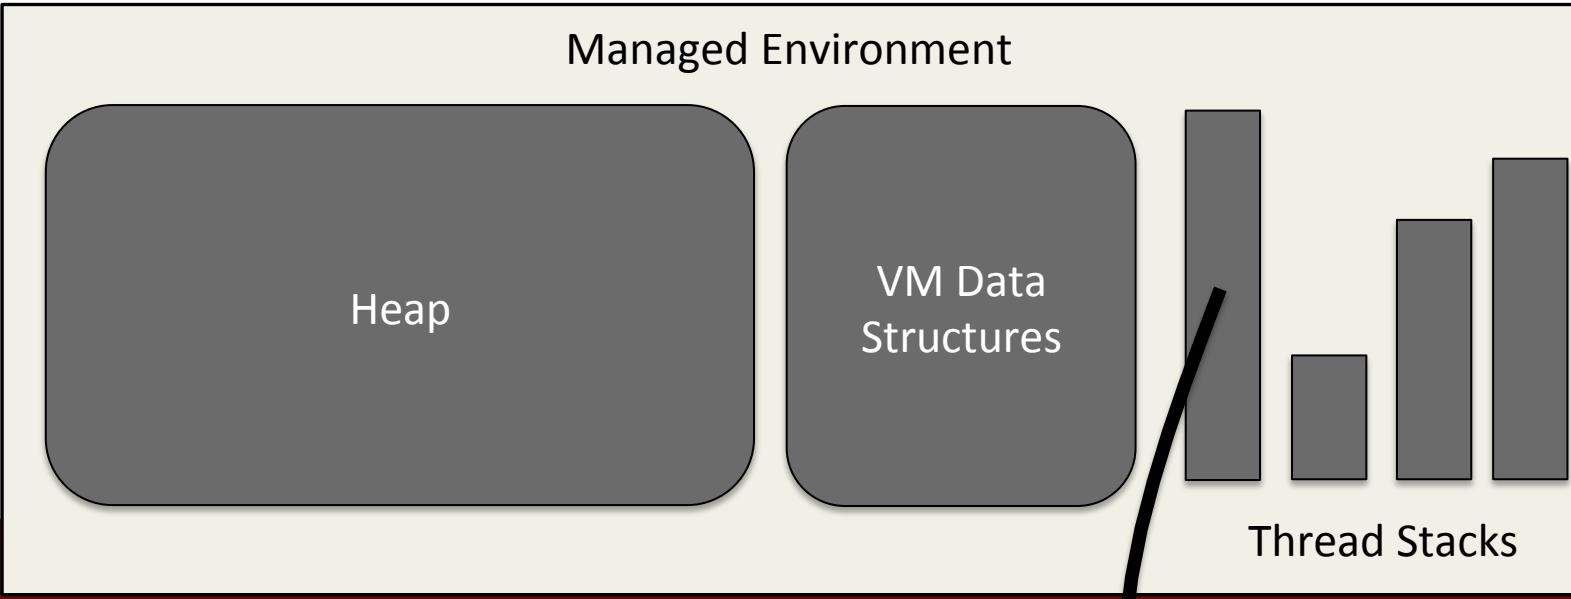
- Mechanism to keep track of references
 - Both local and global
 - Native code holds key index into the registry
 - Registry maps to Java object
- Separate mechanisms for local and global entries
 - Global uses explicit add and remove semantics
 - Local entries are added and removed automatically

Local Registry

- Local references created in two ways
 - Values passed as parameters to native code
 - Return values from VM-level methods
- Explicitly retrieving a reference uses the second
- References added to registry automatically
 - JNI tracks every transition from VM to native
- References deleted when native method returns

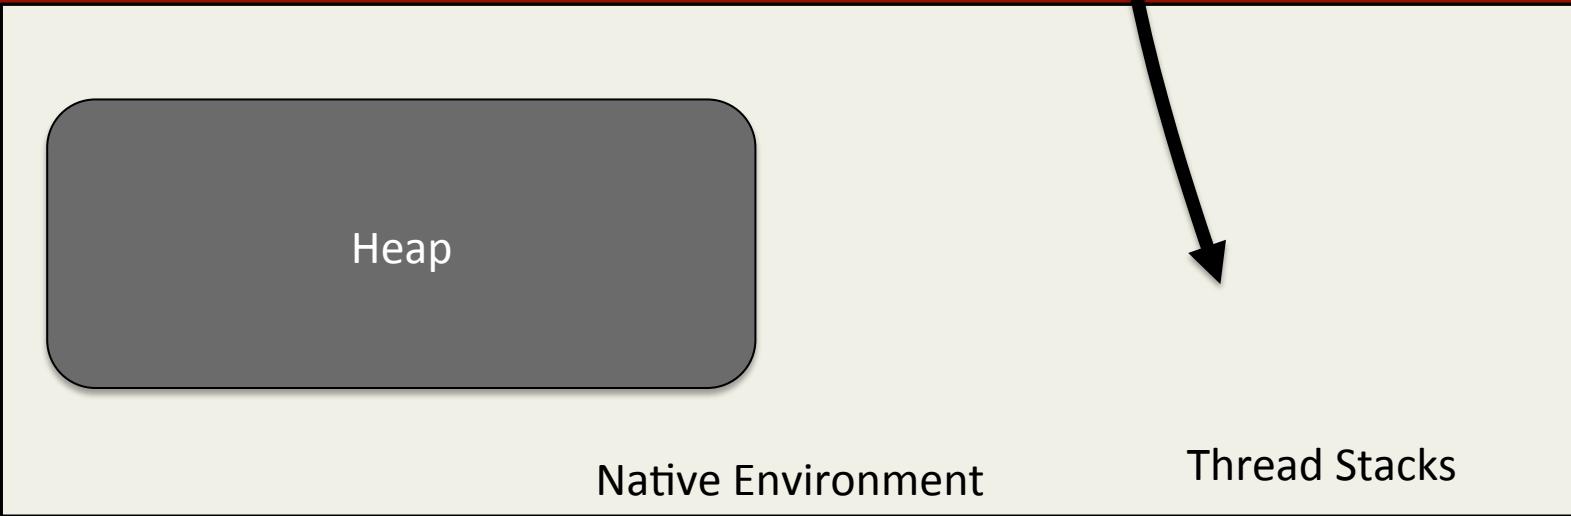


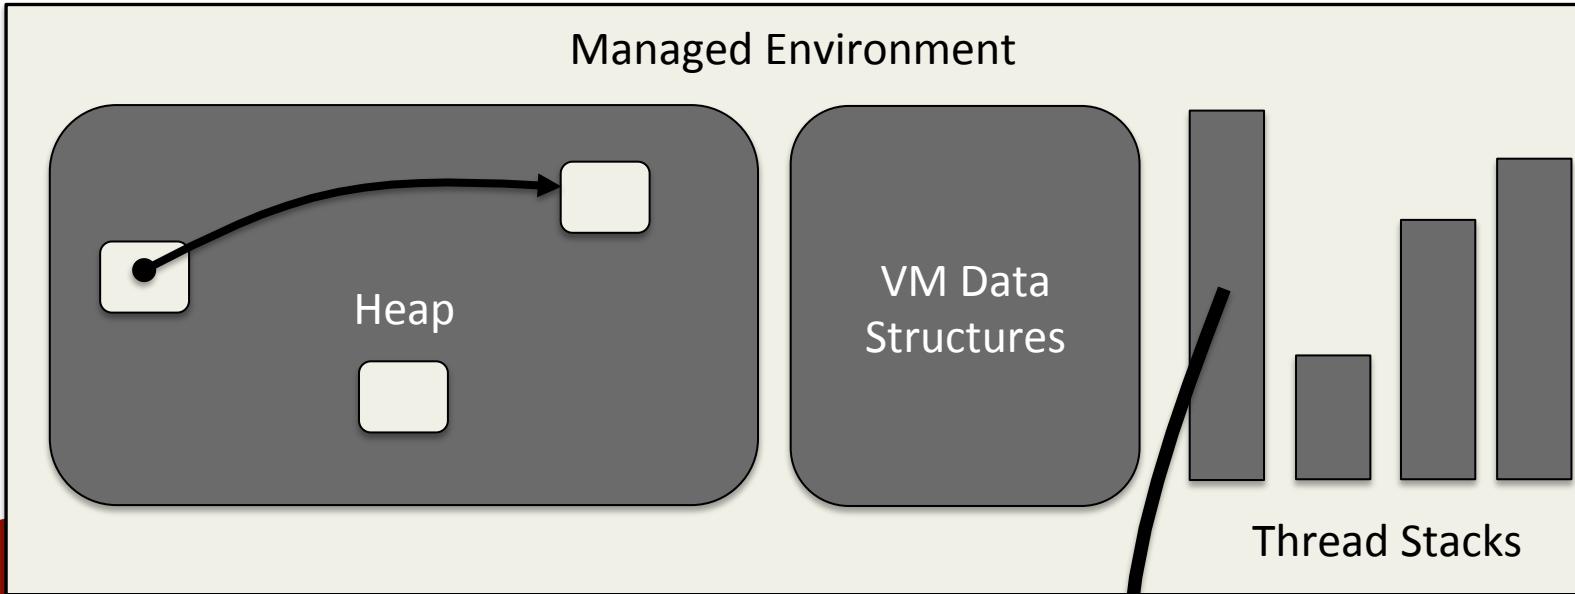




Native Interface

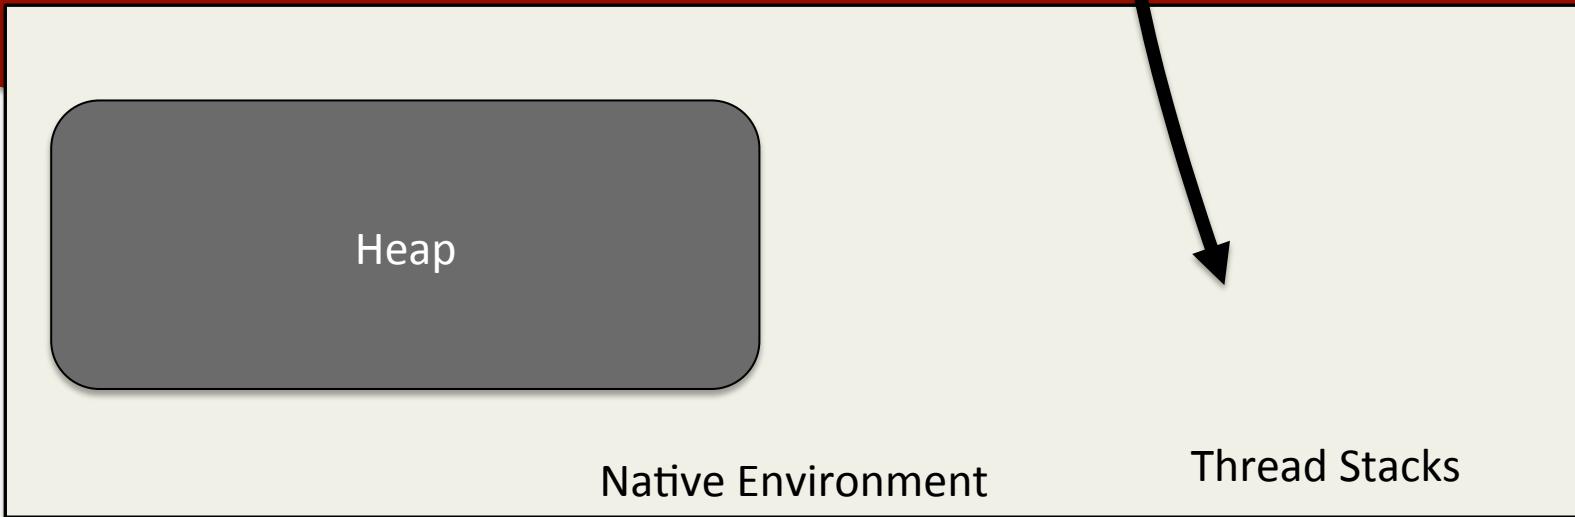
nativeMethod(obj1, obj2)

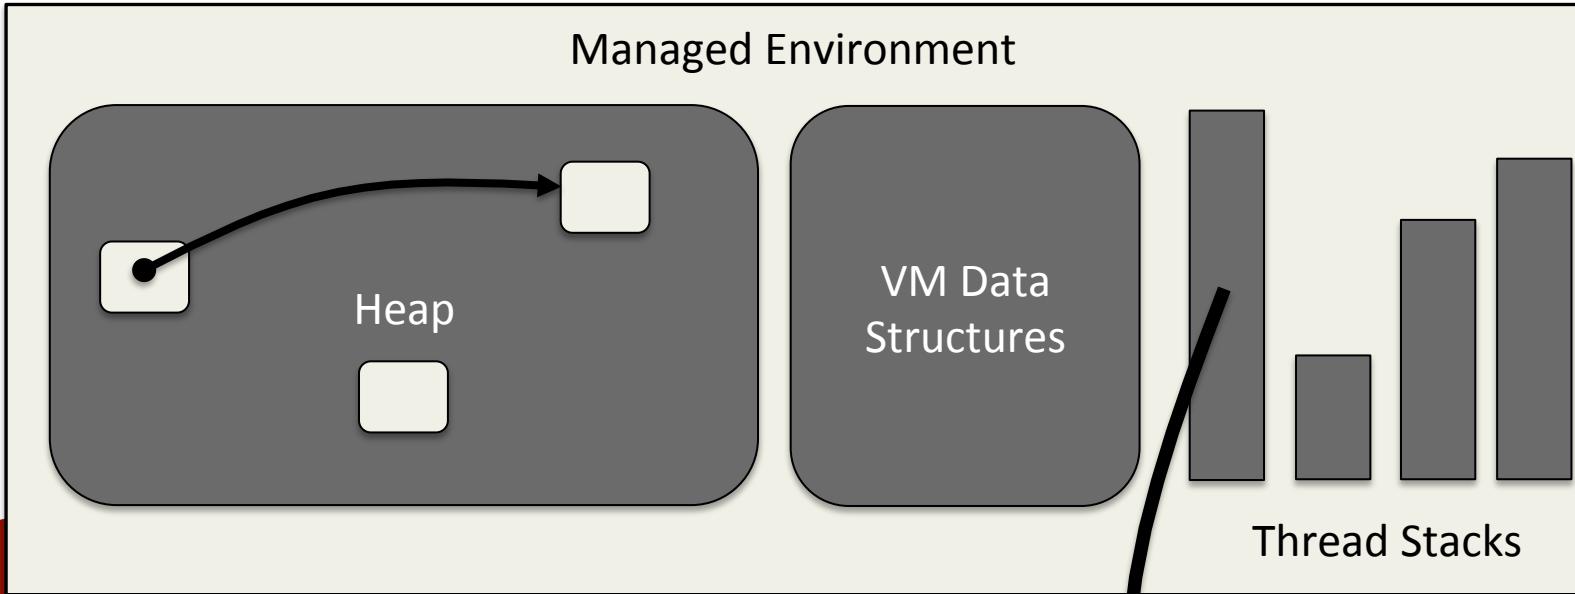




Native Interface

nativeMethod(obj1, obj2)

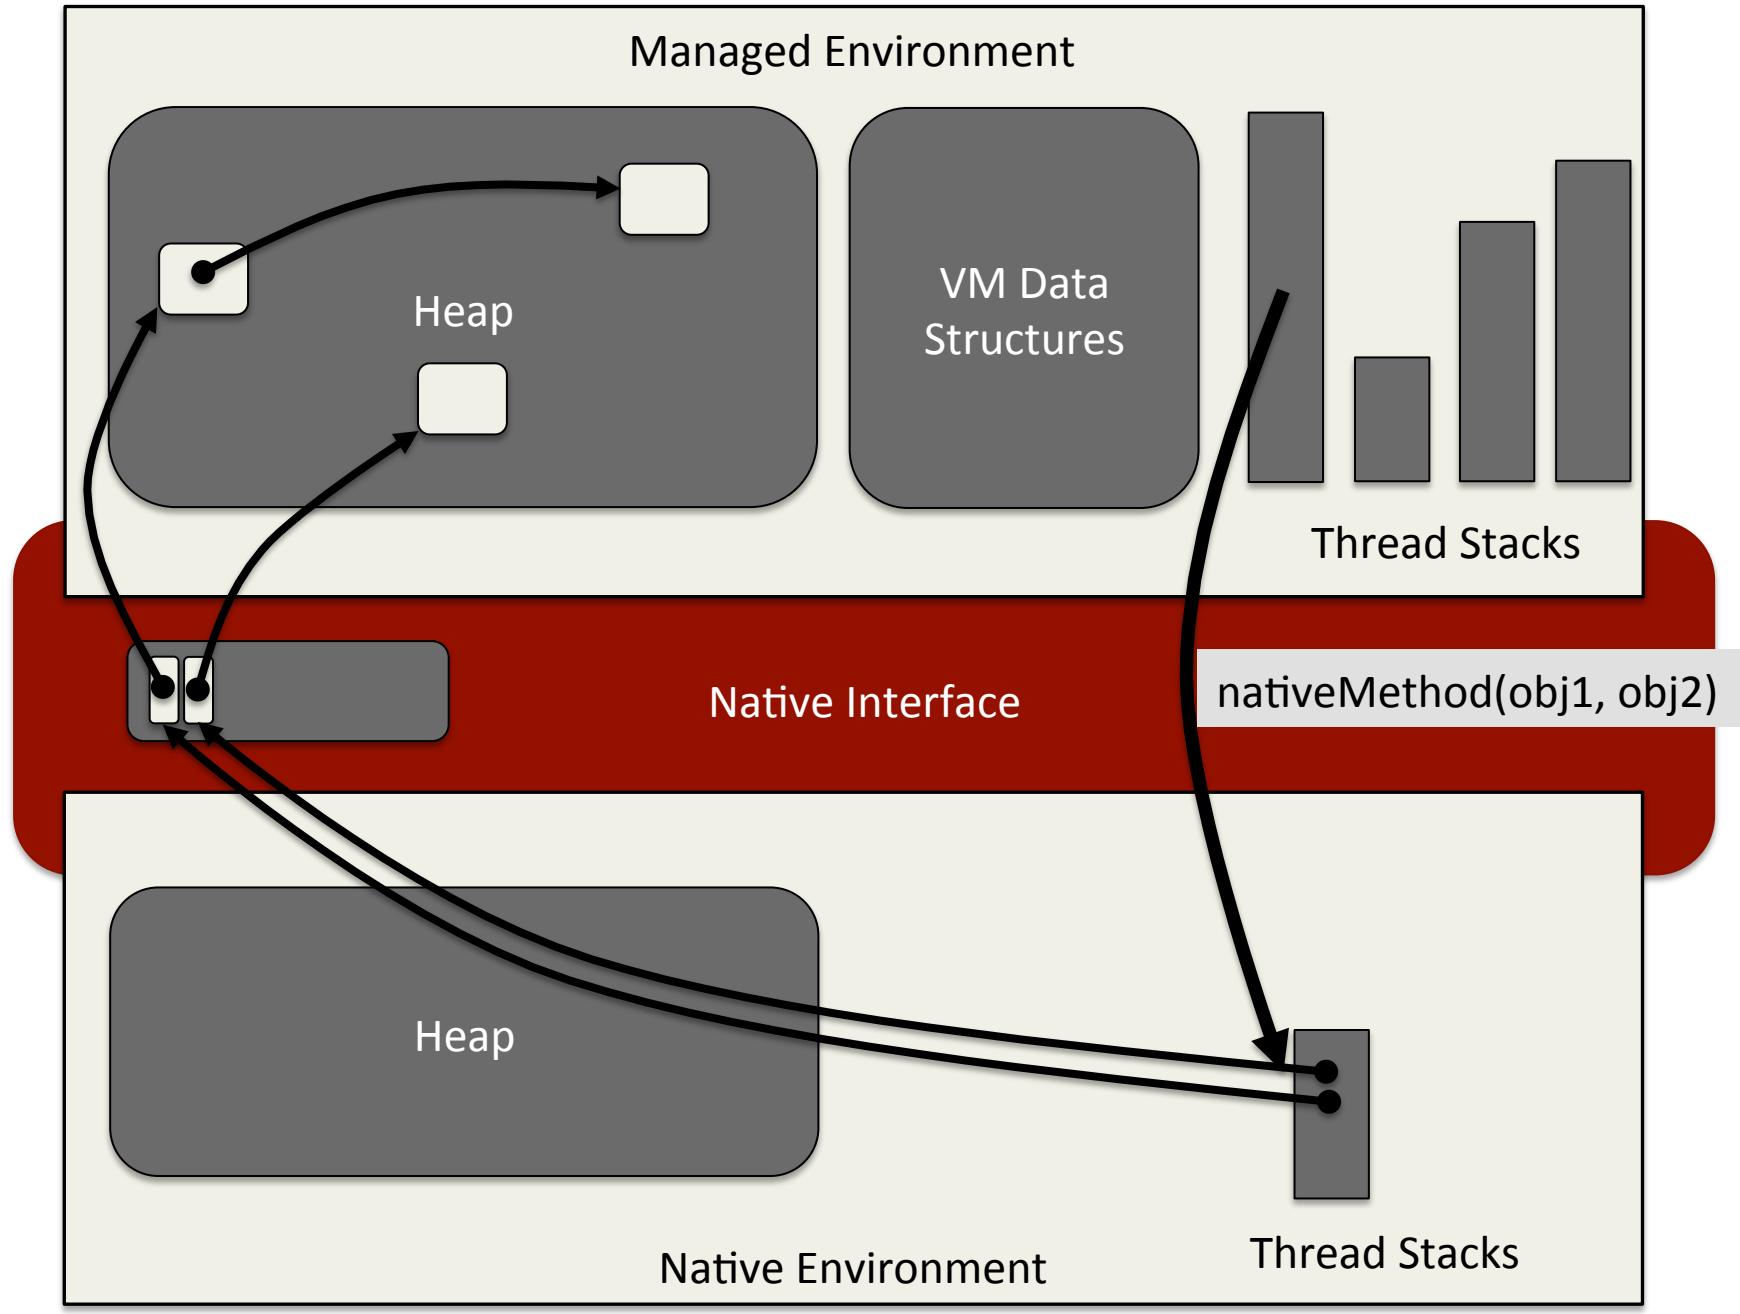


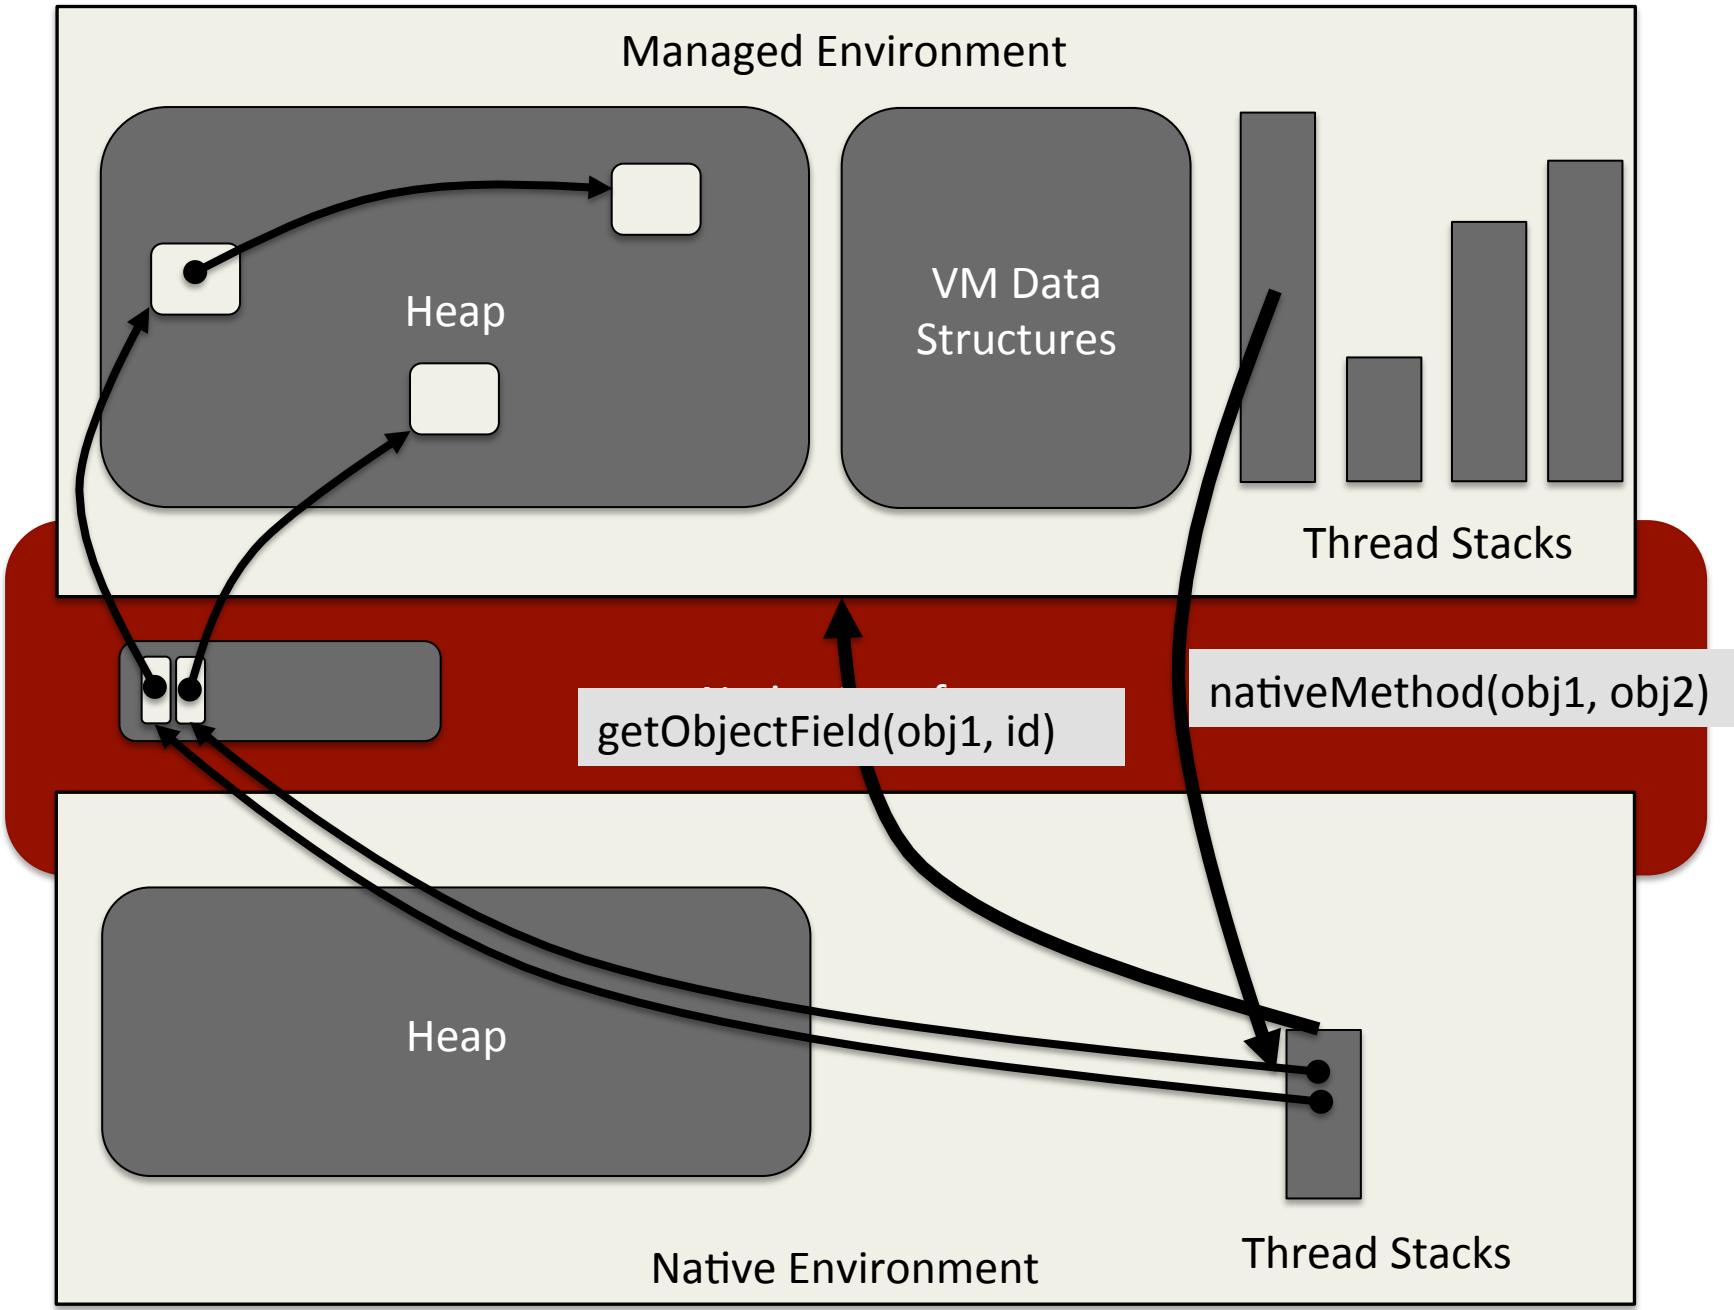


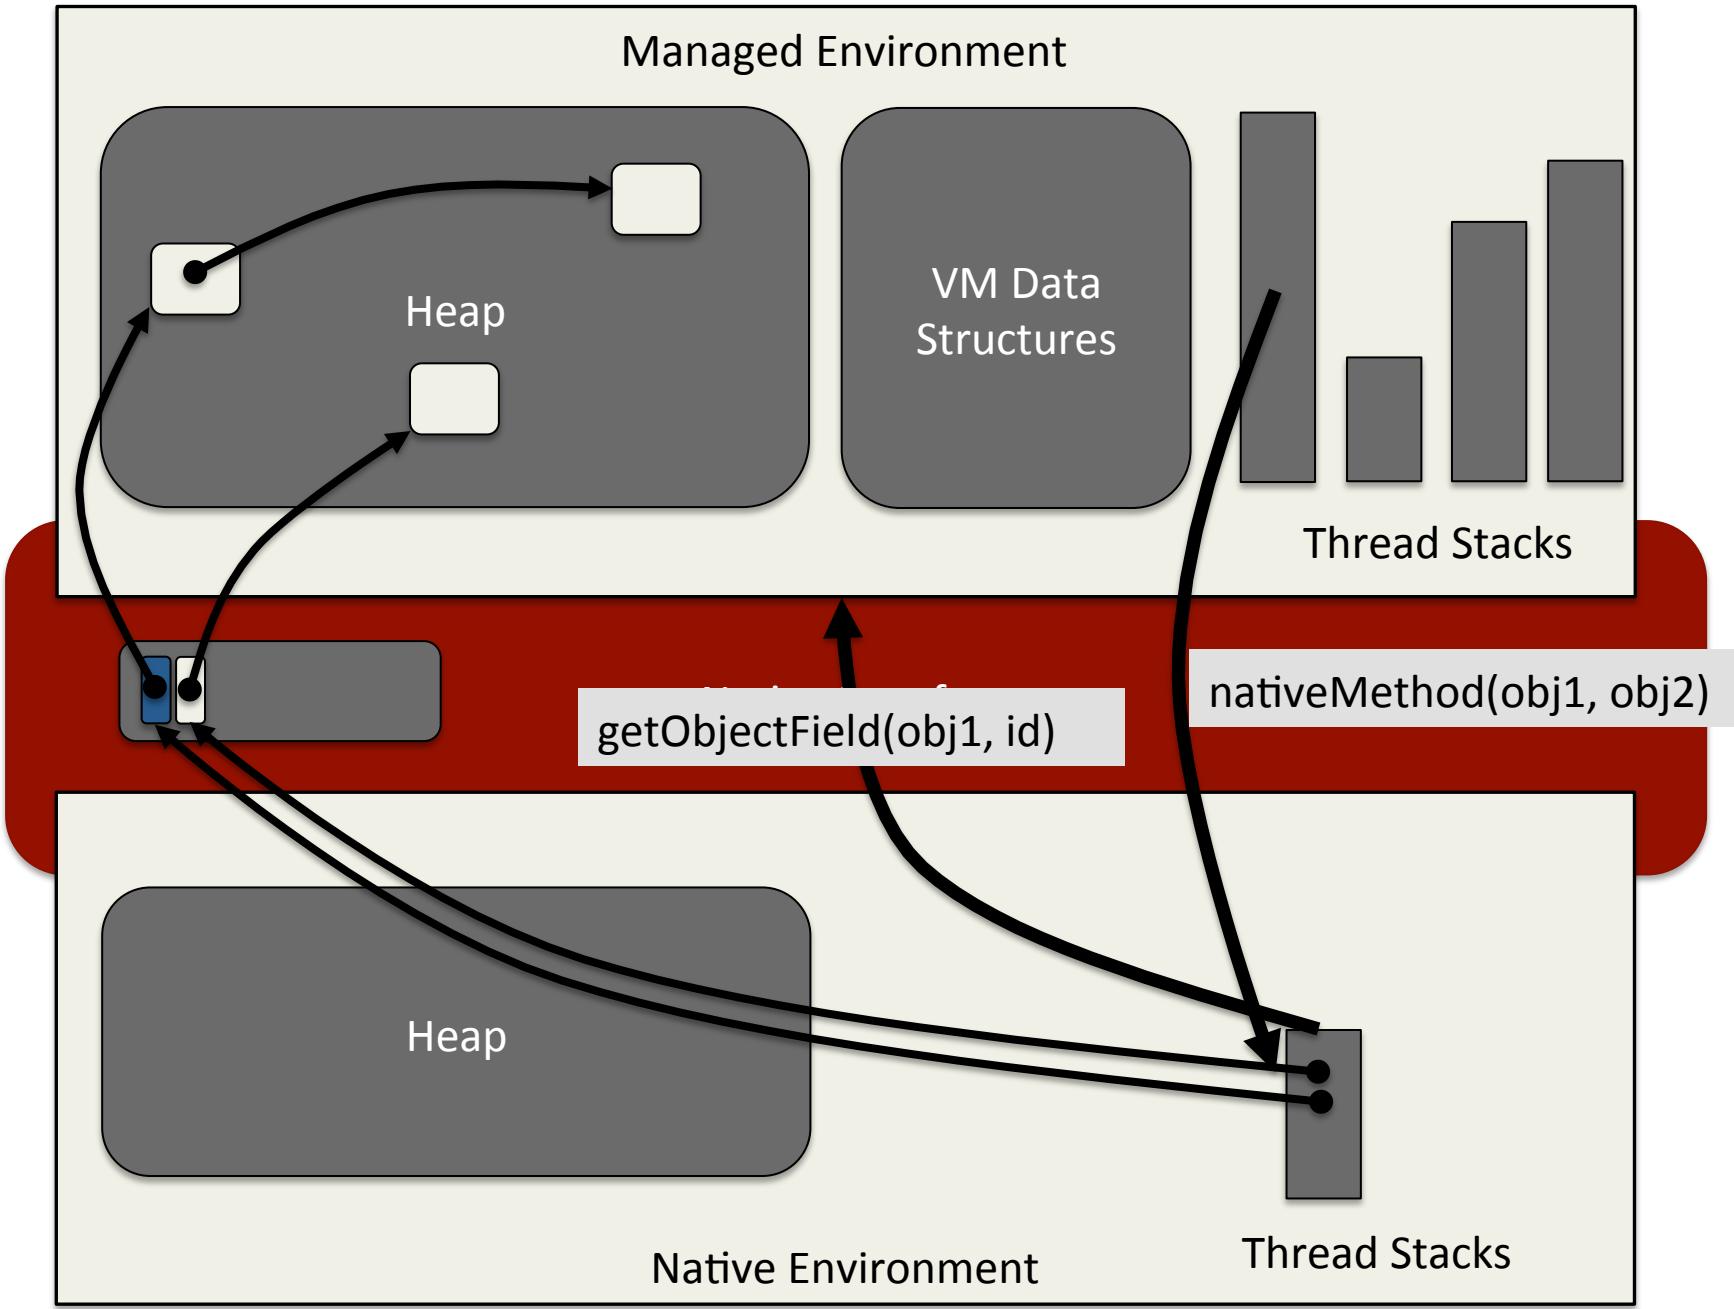
Native Interface

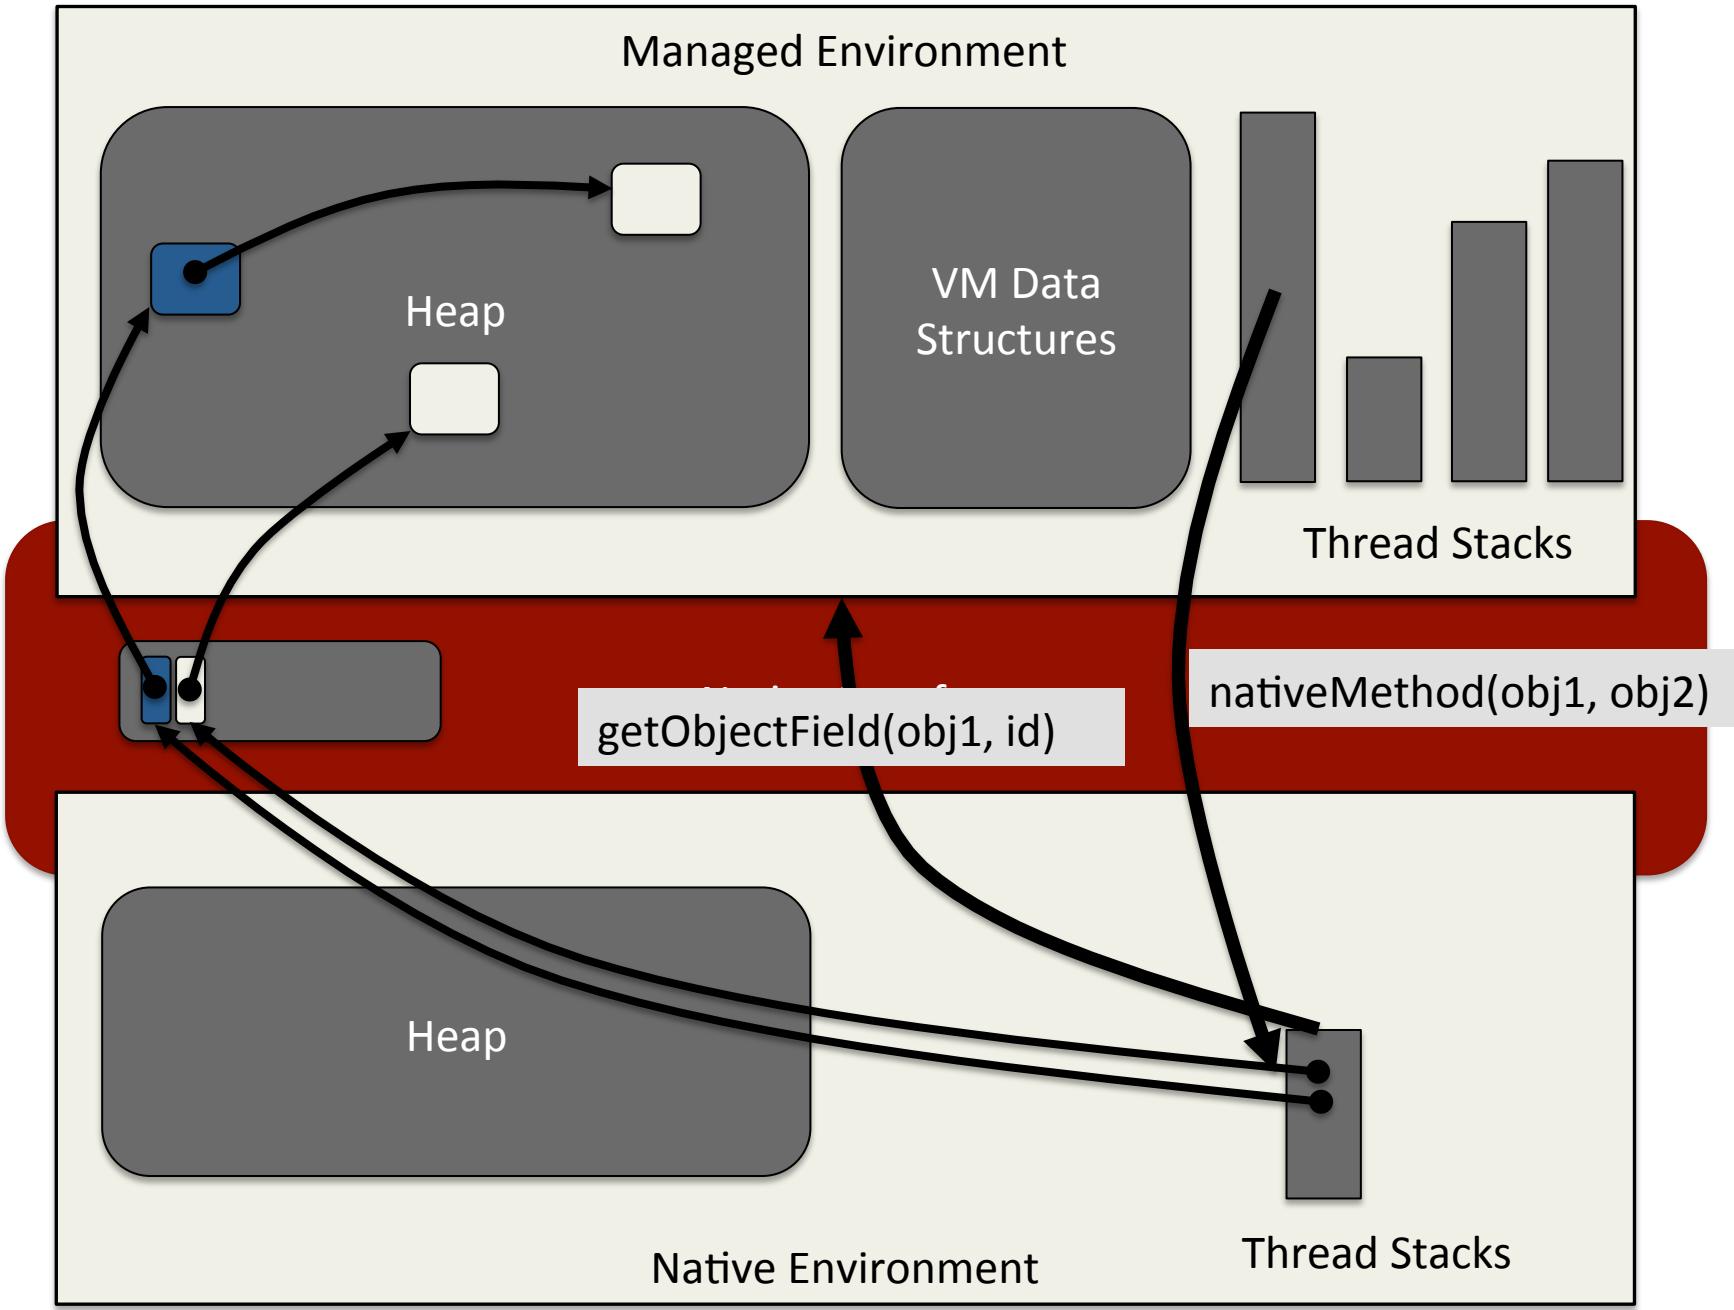
nativeMethod(obj1, obj2)

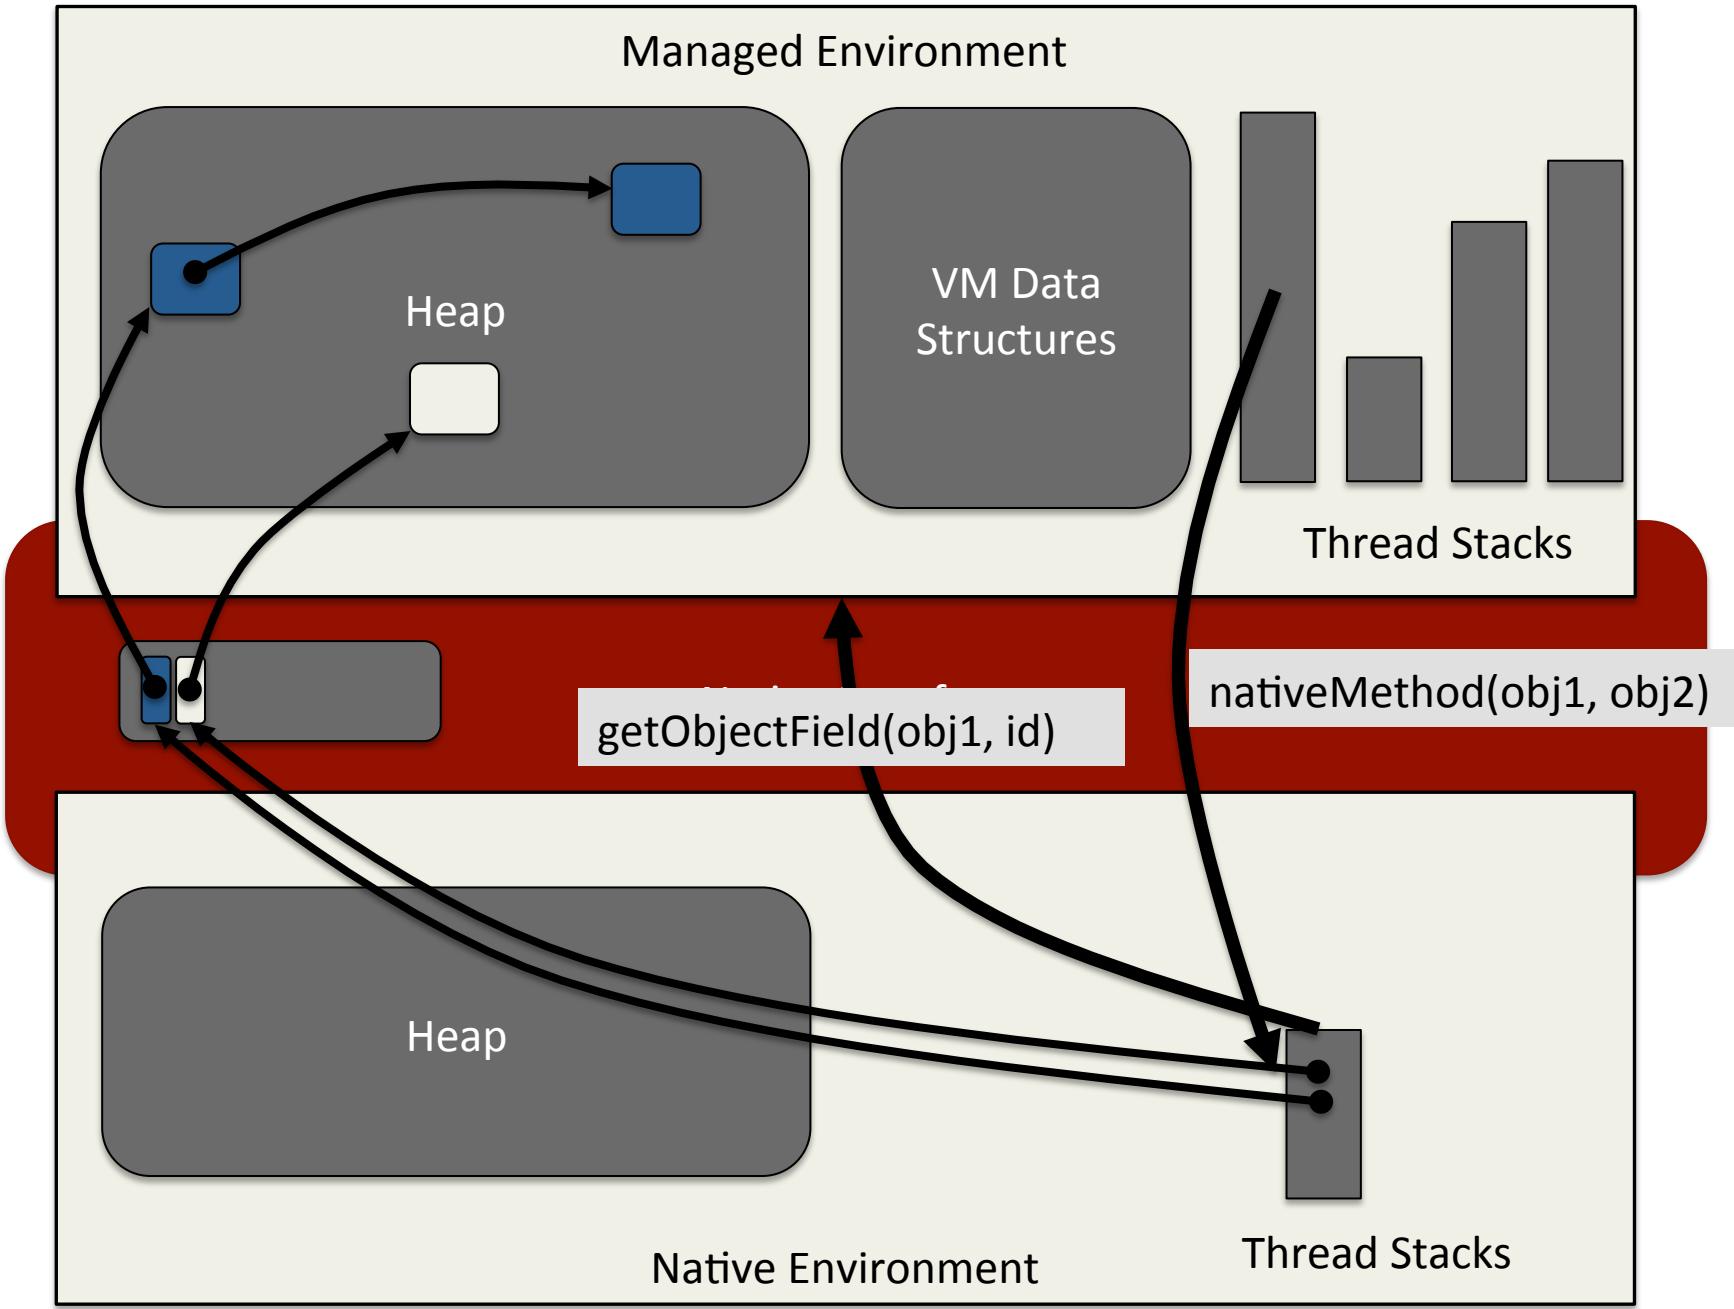


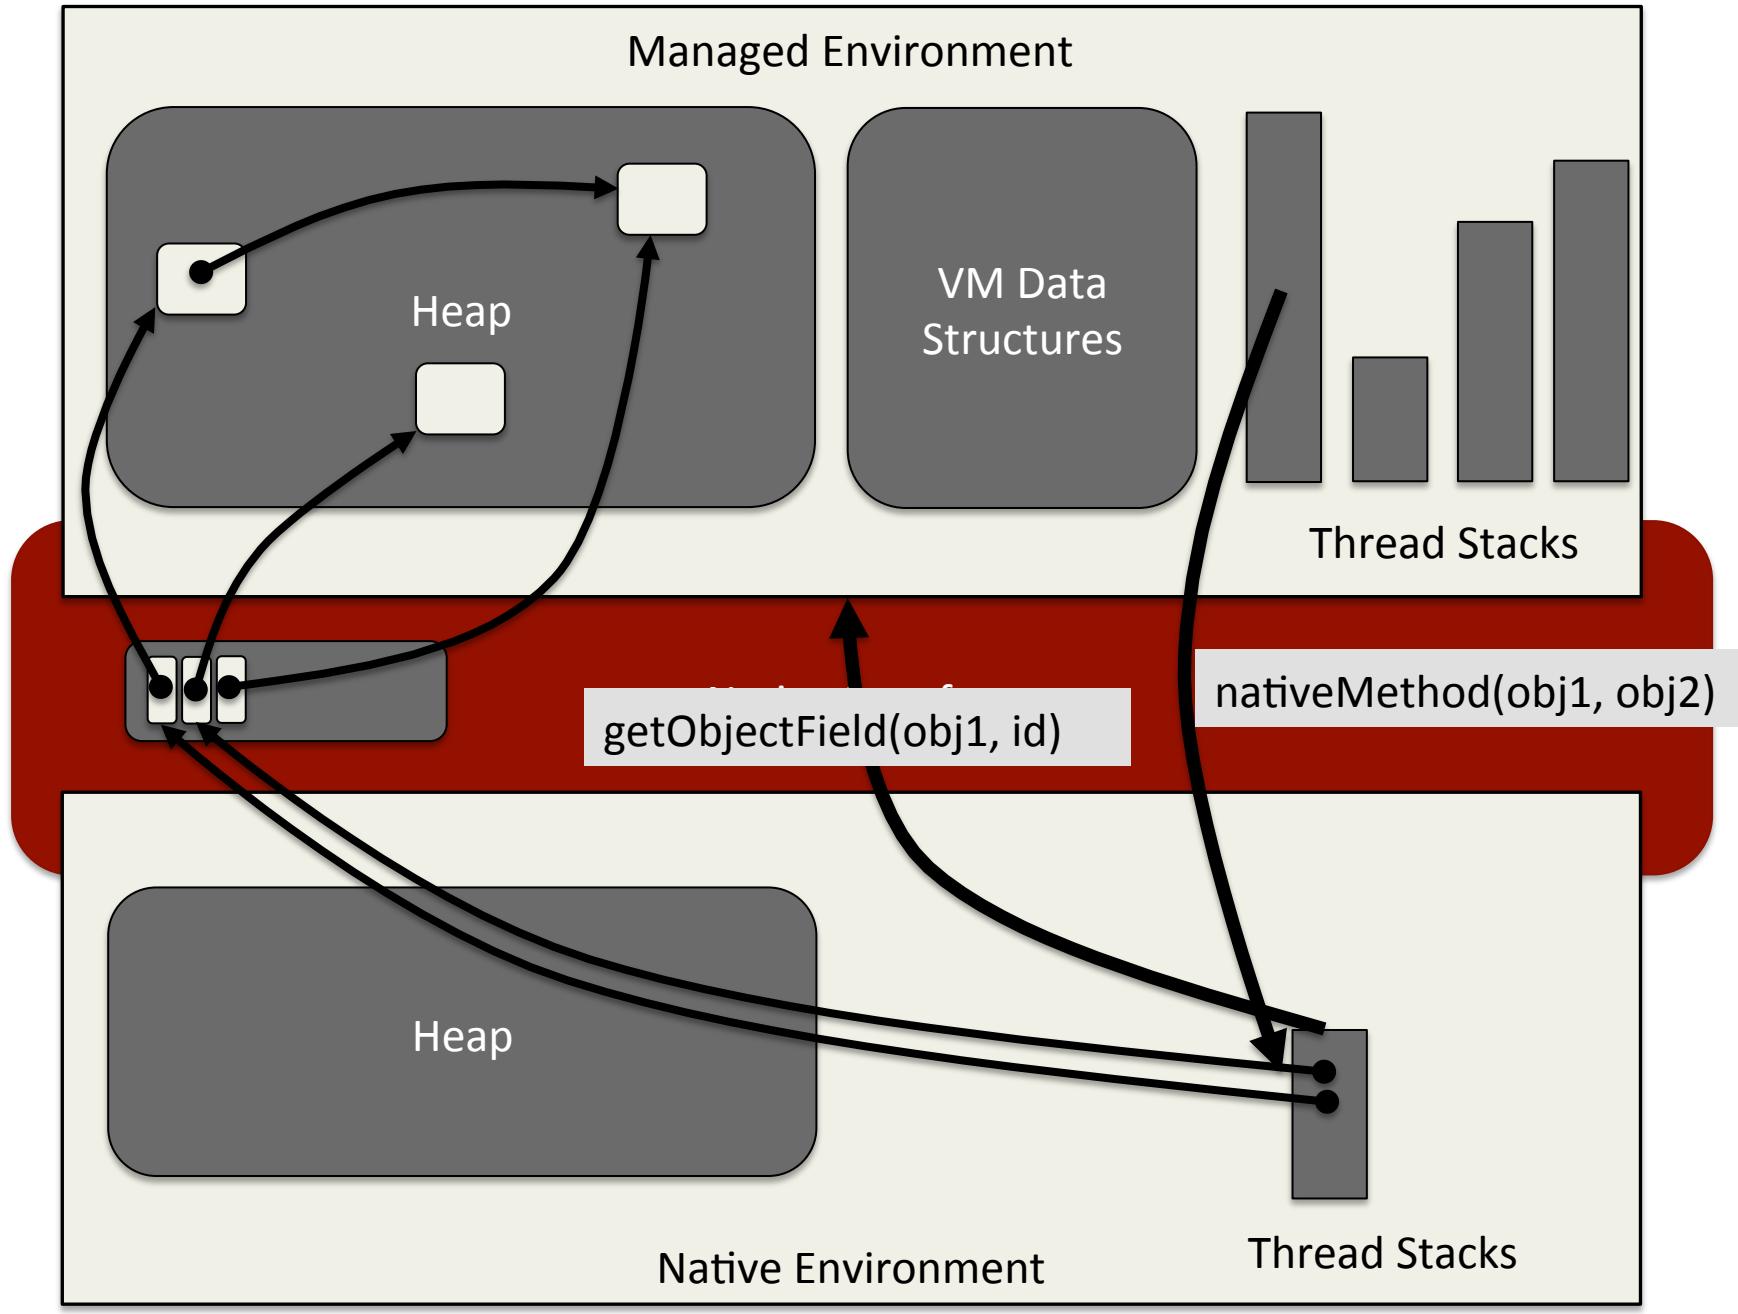


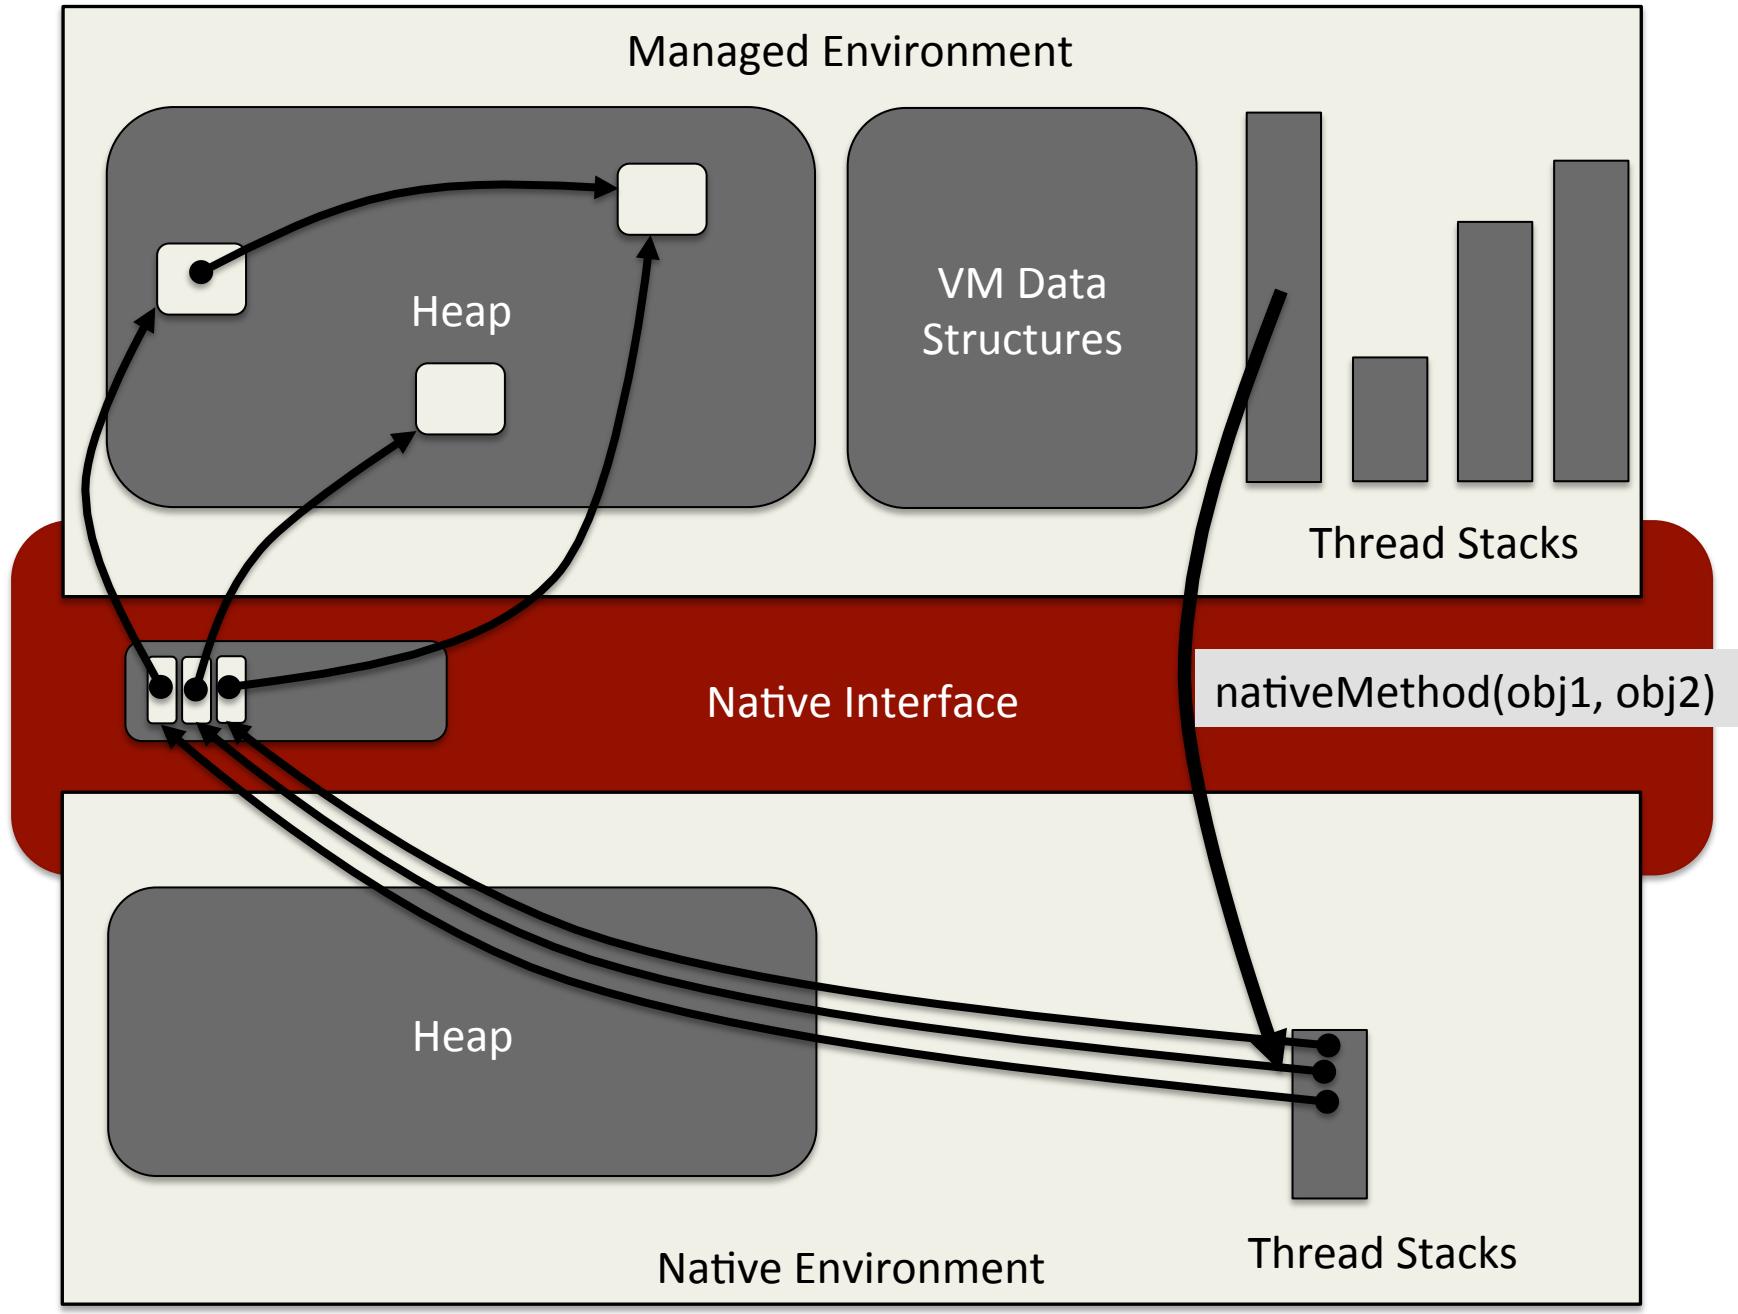


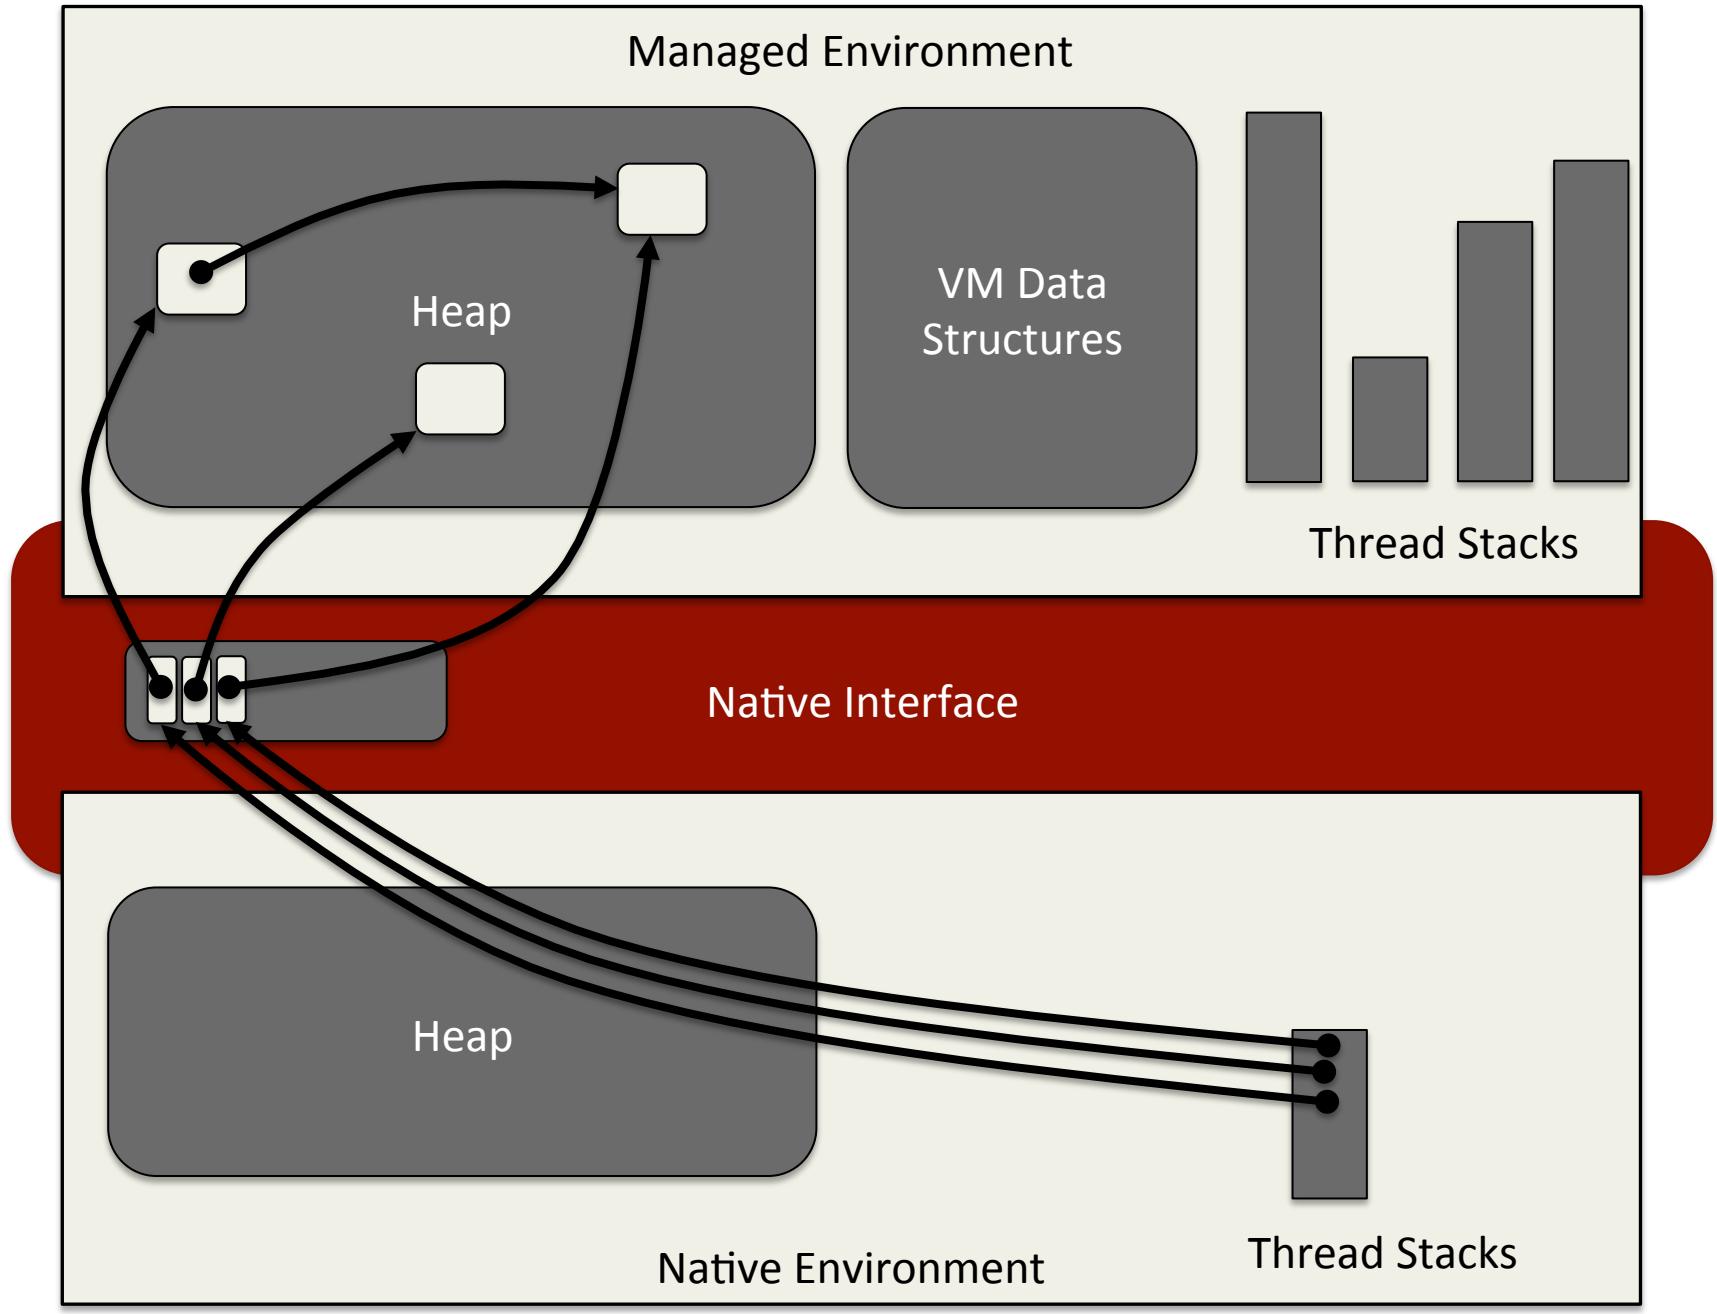


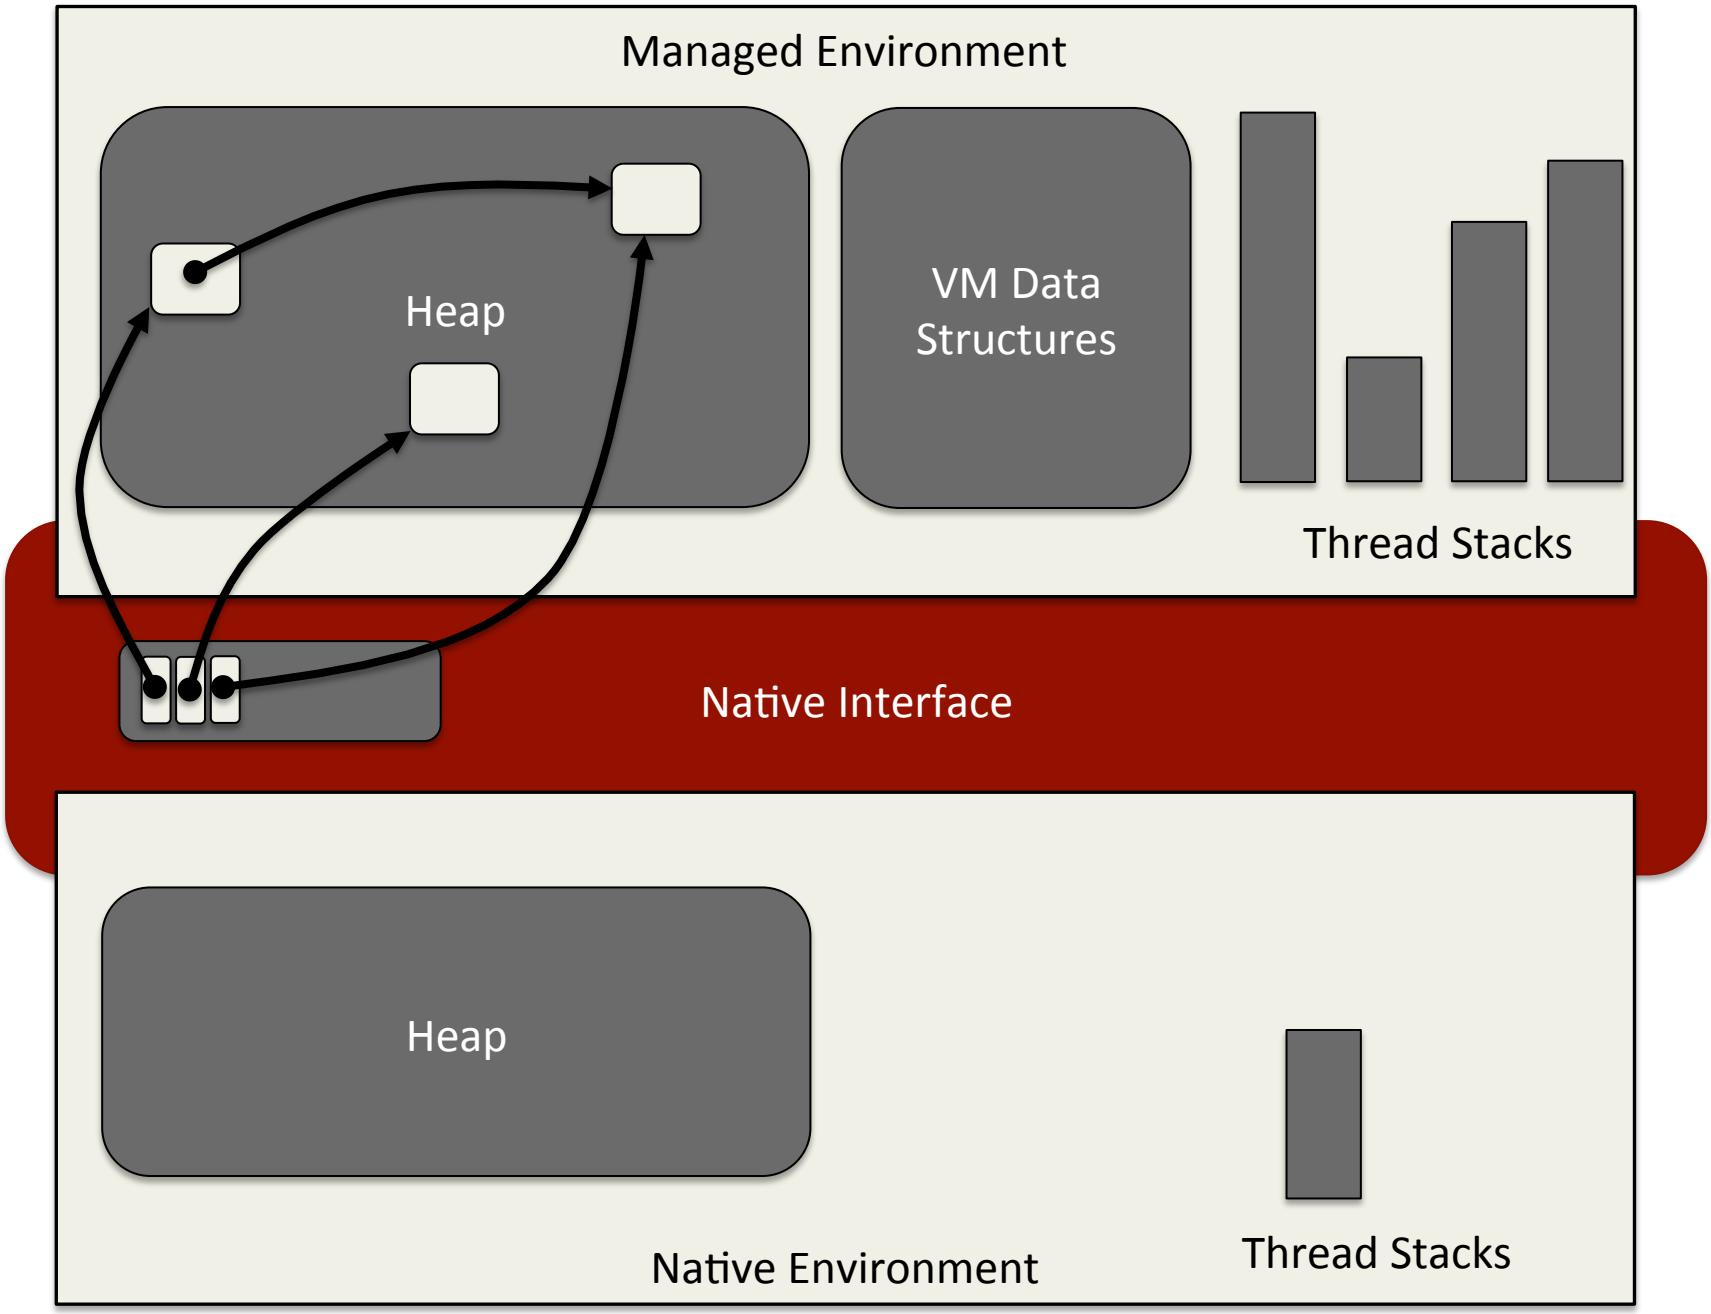


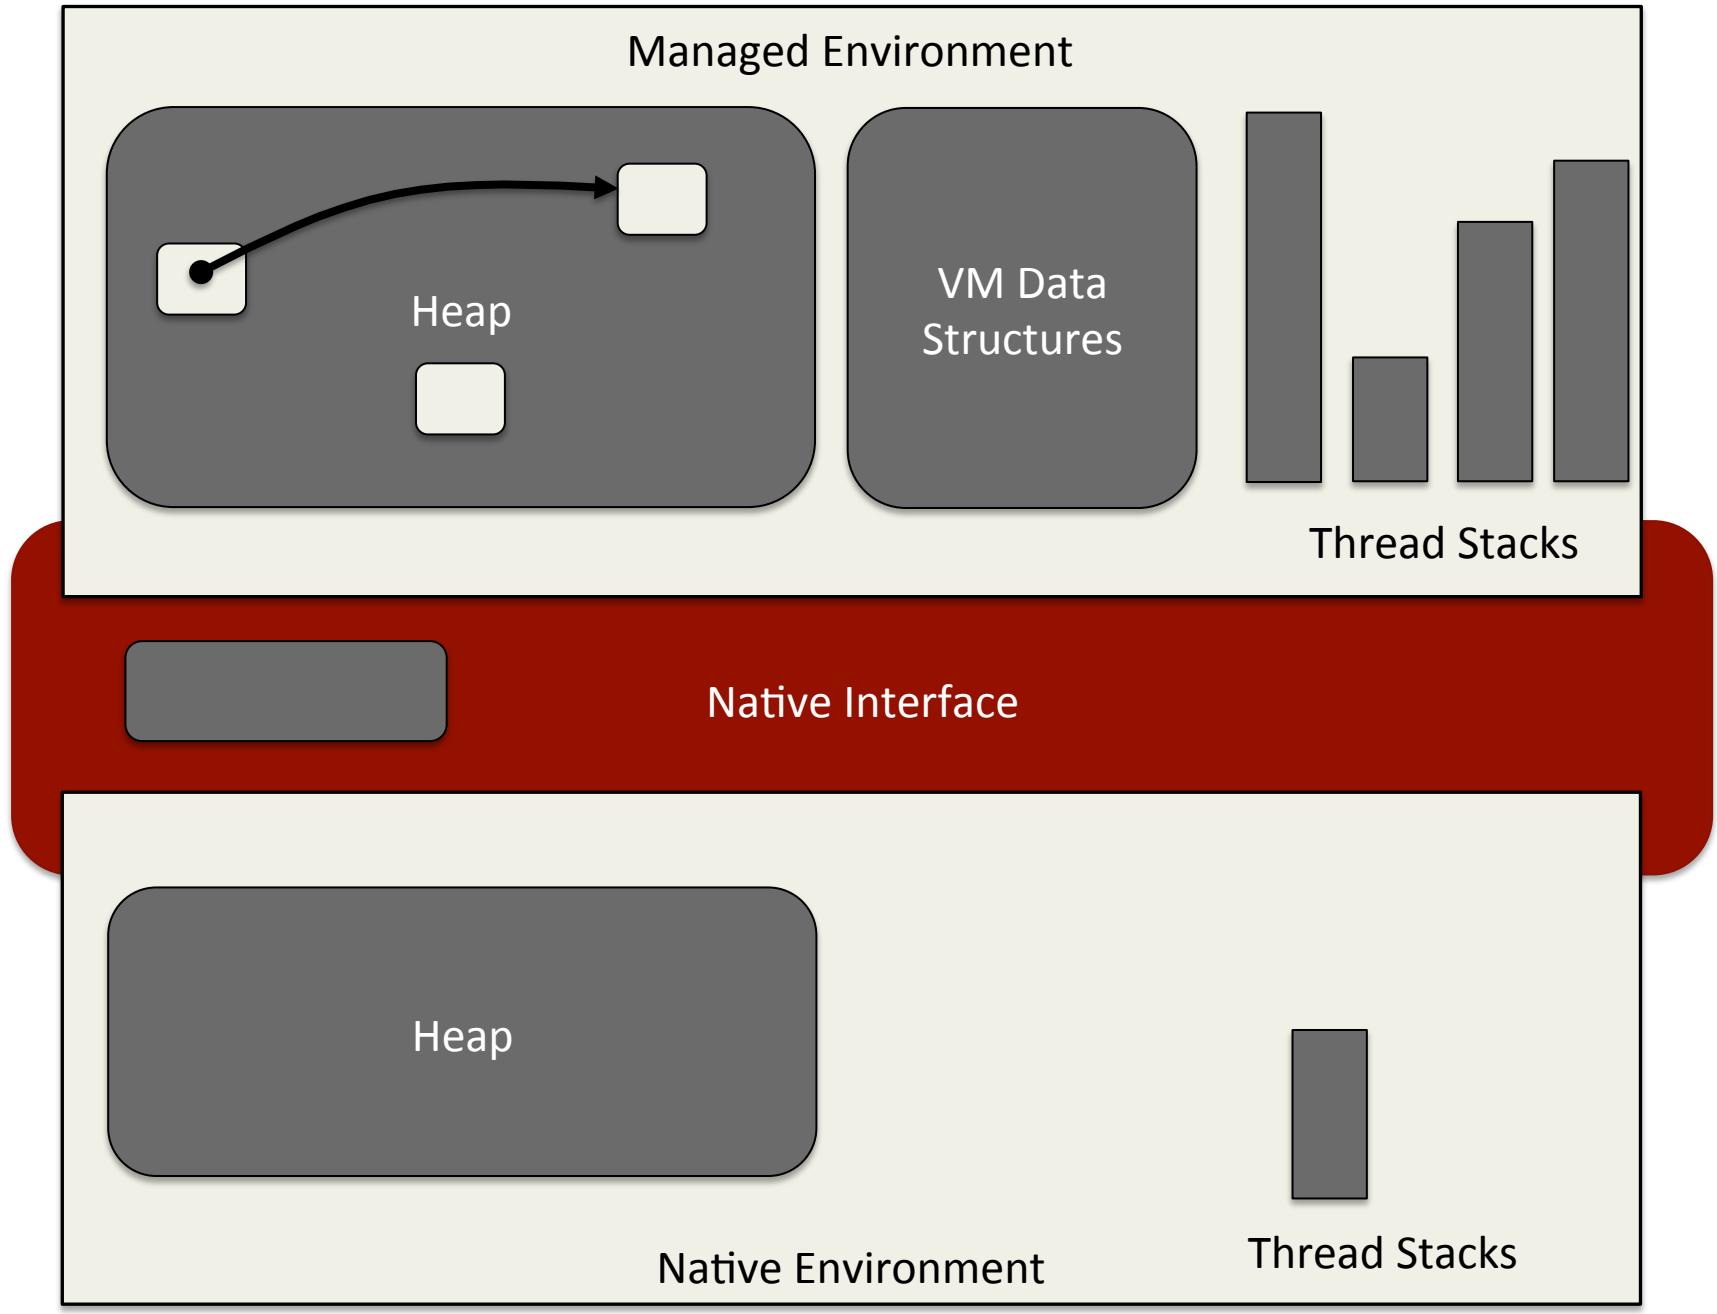












Local Registry Implementation

- Various implementation options
 - Hash map, linked list, etc
- Some similar characteristics to scoped memory
 - Maintain a stack of registry contexts
- Handle space is a limited resource
 - Native code can exhaust available references
 - May compact table by merging multiple references
 - Use a reference count to track reference use

Primitive Arrays

- Every managed object access uses a JNI call
 - Get and set methods track references
 - Allows objects to move behind the scenes
- This adds a significant overhead
- Overhead is particularly bad for primitive arrays
 - No handles to track on get and set
 - Array iteration is a common access pattern

Pinning or Copying

- Primitive arrays can be pinned in memory
 - While pinned, native code can access directly
 - Imposes limitations on the GC
 - Requires that arrays be laid out sequentially
- Arrays can also be copied to native memory
 - Updated contents can then be copied back
 - Most useful for smaller arrays
 - VM can optimize copy operation

Implementing Pinning

- Pinning can be implemented at zero overhead
 - If the garbage collector is non-moving
 - And the VM lays out arrays sequentially
- Neither condition is guaranteed
 - Some VMs use arraylets to break up arrays
 - Others lay arrays out in reverse order
- Can implement pinning semantics by copying
 - Introduces unexpected overhead

Native Safepoints

- Some operations require that threads be paused
- Native threads are a little more flexible
 - They work largely independent of the VM
 - Only interact at well defined points
- Define safepoints as those interactions
 - Accessing managed objects
 - Calling Java methods
 - Returning control to the VM

Exception Handling

- Exception handling in native code is tricky
 - Some native languages support exceptions
 - Others don't, or use non-standard approaches
- JNI doesn't tie in with native mechanisms
- Exceptions thrown in the VM set a flag
 - Native code must check after each call
 - Calling another method before clearing exception is bad
- `JNIEnv` lets native code raise Java exceptions

Invocation API

- Flip side of the native interface
 - Lets the JVM run embedded in a native application
- Use cases similar for Java to native interface
- Environment is similar
 - `JNI_CreateJavaVM` method returns a `JNIEnv`
 - Interactions follow the same rules
- Native code can eventually destroy the VM
 - Individual threads can detach
 - Main thread can unload the VM context