

Eschatology

- Some people have asked about the final project
 - It will be released in two weeks
 - Students can work alone or in a pair
- You will implement a generational GC algorithm
 - Copying collection in the nursery
 - Mark/Sweep collection in the mature space
 - Those working alone will implement simpler algorithms
- I'd like to get an idea of who's working with who
 - Drop me an e-mail if you're planning on forming a team

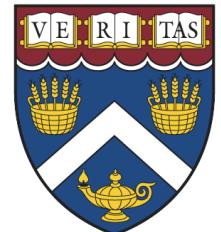
Eschatology II

- Final two classes on December 1st and 8th
 - Shorter special topics
 - Current state of the art
 - Review of the semester
- Suggest topics by e-mail or discussion forum
 - Related areas of interest
 - Things that you'd like to hear more about
 - Anything that wasn't clear during the semester

Eschatology III

- A reminder about the final exam
 - Take-home format
 - Released on the Monday of finals week
 - Due on the Friday
- All material that we've covered is fair game
 - Lecture content
 - Experience when implementing in SimpleJava
 - Assigned readings

Debugging and Profiling



Debugging

- Any mechanism that lets the developer trace code
 - `System.out.println`
 - `new Throwable().printStackTrace()`
- Sometimes stack tracing isn't enough
 - Need deeper insight into code
 - Ability to step through line by line
 - Want to inspect variables or heap objects

Traditional Debugging

- Build application in debug mode
 - Include extra symbol information
 - Turn off optimization
- Use debugger tools to inspect the process
 - Can start with the debugger enabled
 - Or attach the debugger to a running process

Debugging Difficulties

- Debugger can change program characteristics
 - Introduce timing delays
 - Hide race conditions
 - Change cache behavior
- Some bugs don't manifest in debug builds
 - Aggressive optimization changes behavior
 - Operations may have undefined results
 - Developers get away with errors in debug mode

Class File Debug Support

- Java class files ship with debug information
 - File name
 - Line numbers
 - Local variable names
- Bytecode is not optimized
 - Very simple translation from Java to bytecode
 - Makes decompilation simple

Interpreter Debugging

- Well understood process
 - We've run through many examples in class
 - Simple mapping from bytecode to program state
- Breakpoints can be set between bytecodes
 - We've already seen how safepointing works
- Performance is less important
 - Interpreter performance isn't great anyway

Mixed Mode Execution

- Stacks can go between JIT and interpreter
 - Both modes need to support debugging
 - Need to transition smoothly
- May want to hot-swap executing code
 - On-stack replacement
- Threads can also enter native code
 - Native debugging is still a research question

Mixed Mode Execution

- Differences can crop up between modes
 - Less common than in native applications
 - Fewer undefined operations
 - Stricter checking of bytecode
- Some things can still be affected
 - Memory layout changes
 - GC behavior
 - Lots of timing changes

Out of Process Debugging

- Running in debug mode can be slow
 - Artificially low optimization level
 - Overhead of debugger operations
- Want to minimize the overhead
 - Particularly when the debugger's not running
- Launch debugger as separate process
 - Debugger's work doesn't interfere with the program

Virtual Machine

VM Data Structures



Threads

Heap

Virtual Machine

VM Data Structures



Threads

Heap

External Tool

```
public static int add() {  
    final int i = 2;  
    final int j = 3;  
    return i + j;  
}
```

4

Program Counter



Stack

2 3

Local
Variables

public static int add();
flags: ACC_PUBLIC, ACC_STATIC
Code:

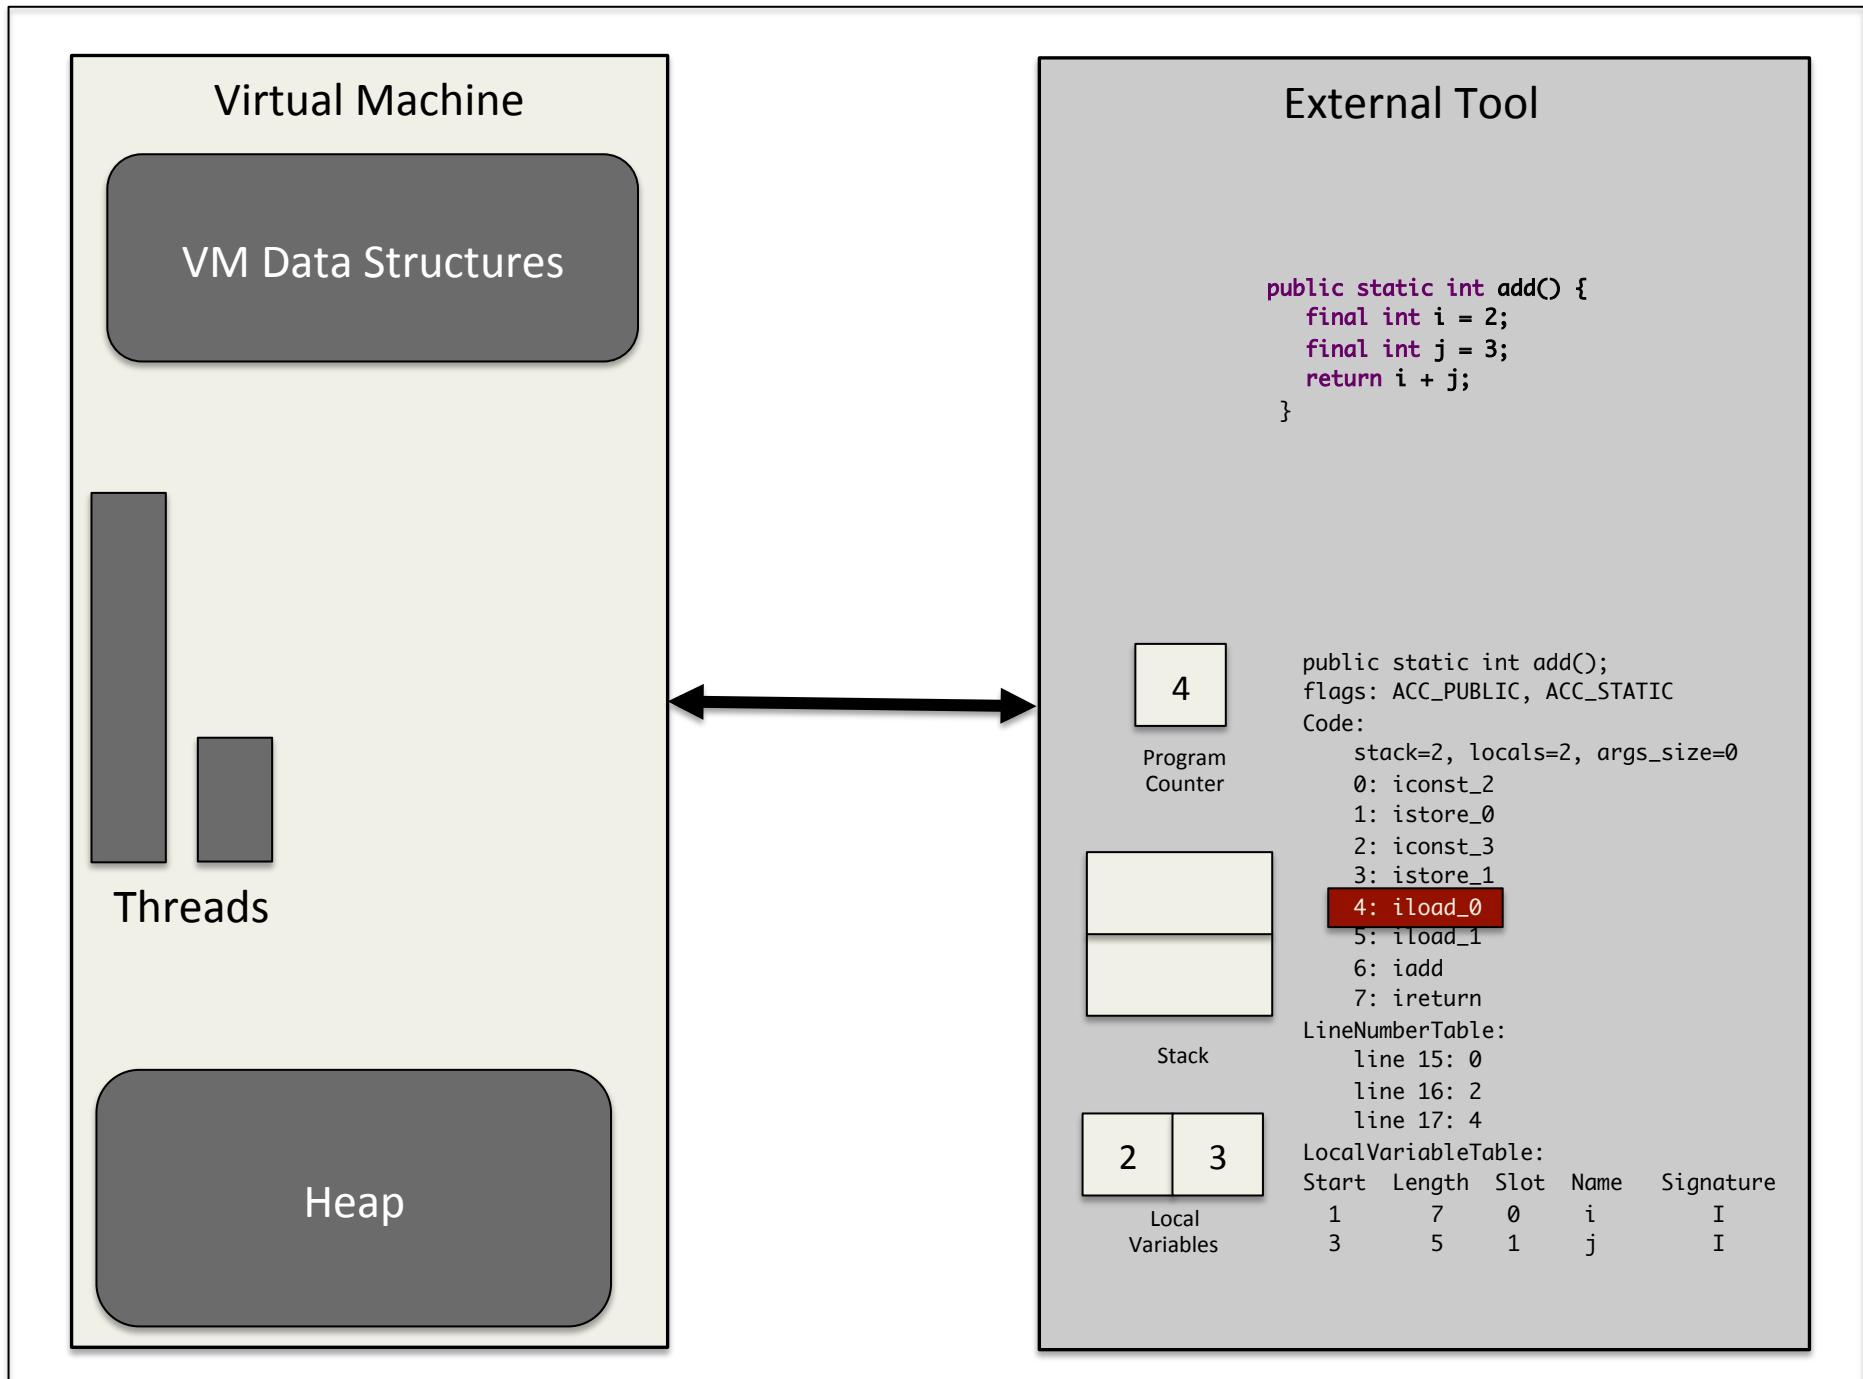
```
stack=2, locals=2, args_size=0  
0:  iconst_2  
1:  istore_0  
2:  iconst_3  
3:  istore_1  
4:  iload_0  
5:  iload_1  
6:  iadd  
7:  ireturn
```

LineNumberTable:

```
line 15: 0  
line 16: 2  
line 17: 4
```

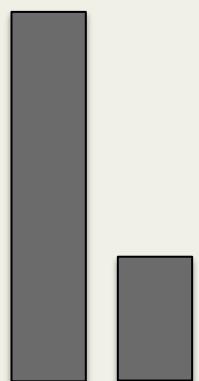
LocalVariableTable:

| Start | Length | Slot | Name | Signature |
|-------|--------|------|------|-----------|
| 1 | 7 | 0 | i | I |
| 3 | 5 | 1 | j | I |



Virtual Machine

VM Data Structures



Debug Agent

Heap

External Tool

```
public static int add() {  
    final int i = 2;  
    final int j = 3;  
    return i + j;  
}
```

4

Program Counter



2 3

Local Variables

public static int add();
flags: ACC_PUBLIC, ACC_STATIC
Code:

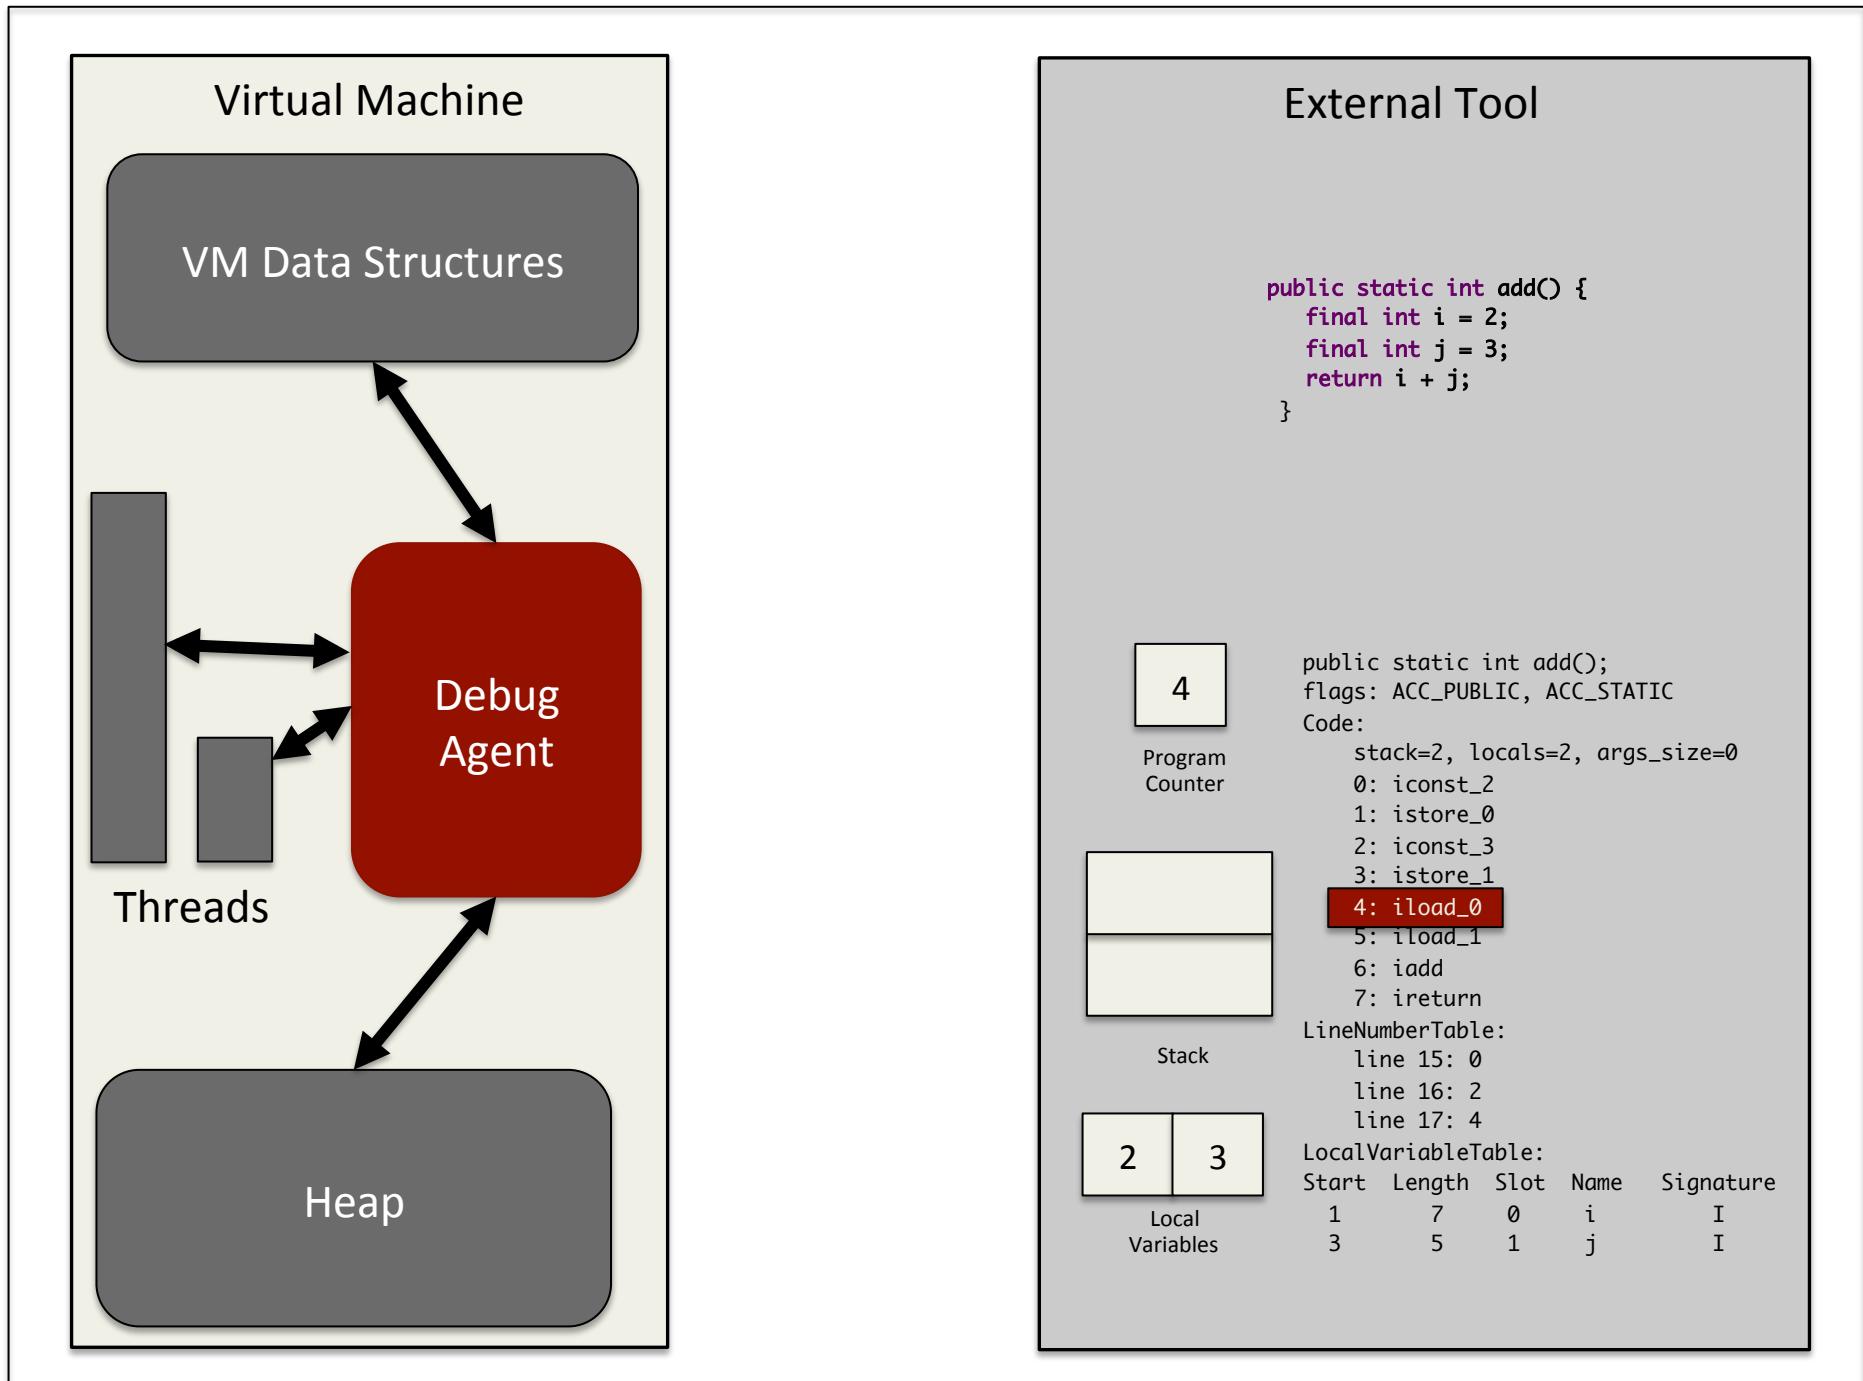
stack=2, locals=2, args_size=0
0: iconst_2
1: istore_0
2: iconst_3
3: istore_1
4: iload_0
5: iload_1
6: iadd
7: ireturn

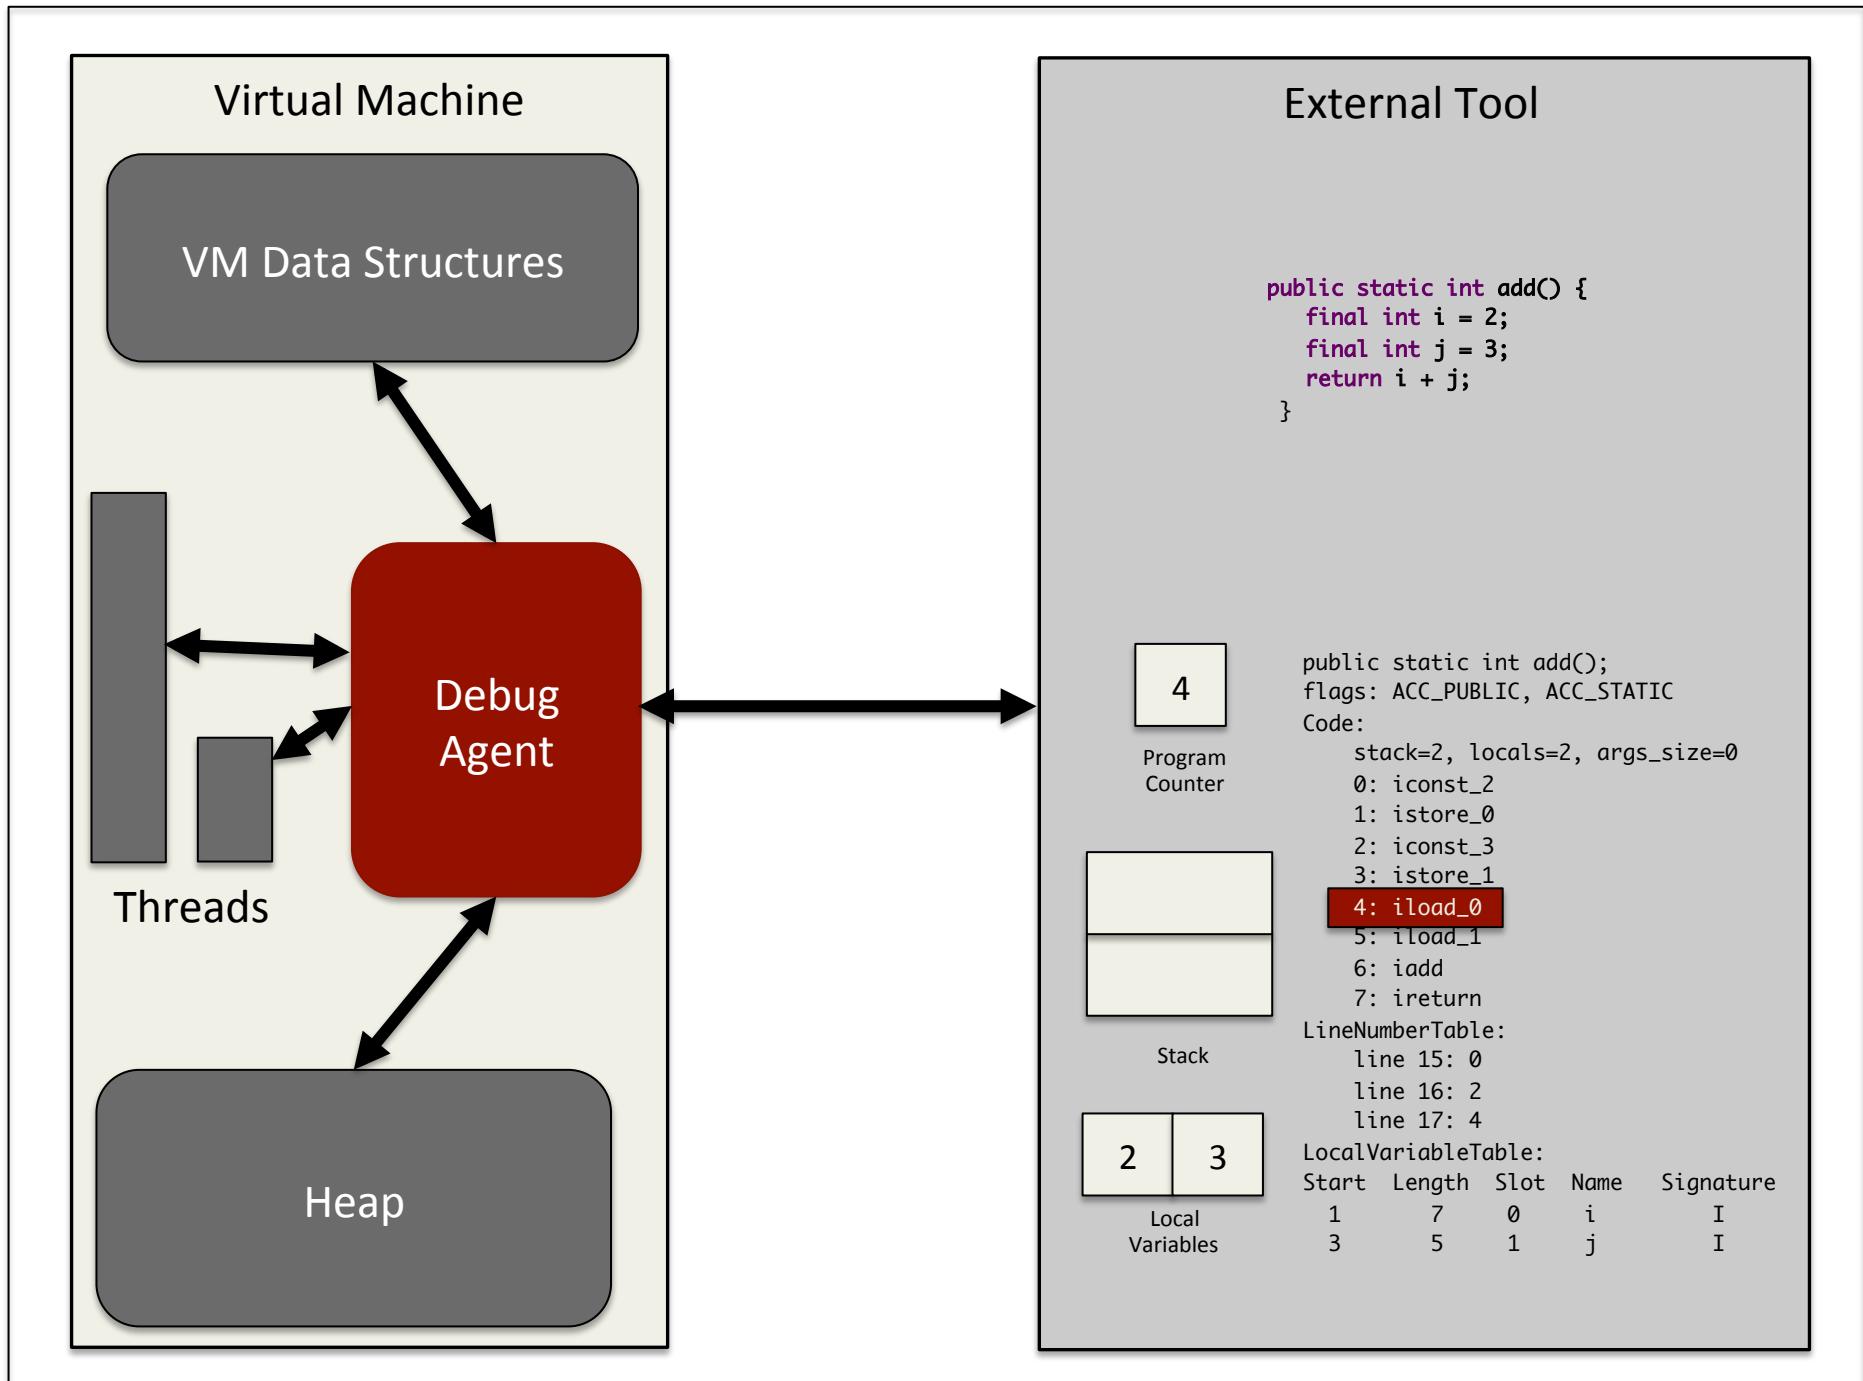
LineNumberTable:

line 15: 0
line 16: 2
line 17: 4

LocalVariableTable:

| Start | Length | Slot | Name | Signature |
|-------|--------|------|------|-----------|
| 1 | 7 | 0 | i | I |
| 3 | 5 | 1 | j | I |





JVMTI

- The JVM Tools Interface
 - Optional part of the VM specification
 - Introduced with version 1.5 of the JDK
 - Supports external tools that interact with the VM
- Disabled by default
 - Load agent on the command line
- Tool determines which features to turn on

JVMTI Agents

- Written in native language
 - Interface definitions in the JDK include directory
 - Generally a small piece of code
- Agent runs in the VM's address space
- Two-way communication with the VM
 - VM notifies agent of interesting events
 - Agent can control the VM
- Agent relays information to external tool

Loading an Agent

- VM loads the agent and calls startup function
 - Agent_OnLoad
- VM only partially created when agent launches
 - System property values have been set
 - No classes loaded
 - No objects created
 - No bytecodes executed
- Failure in agent creation leads to VM shutdown

Attaching an Agent

- Agents can be added once the VM is running
 - Uses `Agent_OnAttach` initialization method
- May have more limited functionality
 - May not be able to instrument running classes
 - Overhead for some features may be higher
- Failure in agent attach ignored by the VM

Timing

- Actions in the `Agent_OnLoad` call restricted
 - Parts of the VM not yet created
- Early loading allows agent to be more powerful
 - Capture VM creation events
 - Instrument system bytecode
 - Influence the newly-created VM

Unloading an Agent

- Agent can create an `Agent_OnUnload` function
 - Free up any resources held by the agent
- Triggers for unloading implementation specific
 - May be manual or automatic
- Called on clean VM shutdown
 - Not called if the VM crashes

Accessing VM Functionality

- The JVMTI agent uses a set of VM functions
 - Environment pointer to thread data and functions
 - Double indirection to allow function replacement
 - Allowable methods limited in some cases
- Control is inverted from JNI
 - Most agent threads not created by Java method call
 - Can't assume that they will eventually return
 - Complicates object reference management

Capabilities

- Set of events and permissions for the agent
 - Specified in `Agent_OnLoad`
- Capabilities are implementation dependent
 - Not all VMs will support all capabilities
 - Capabilities determine agent portability
 - Not all capabilities available after startup
- VM need only provide requested capabilities
 - Can tailor configuration for performance

Event Handling Capabilities

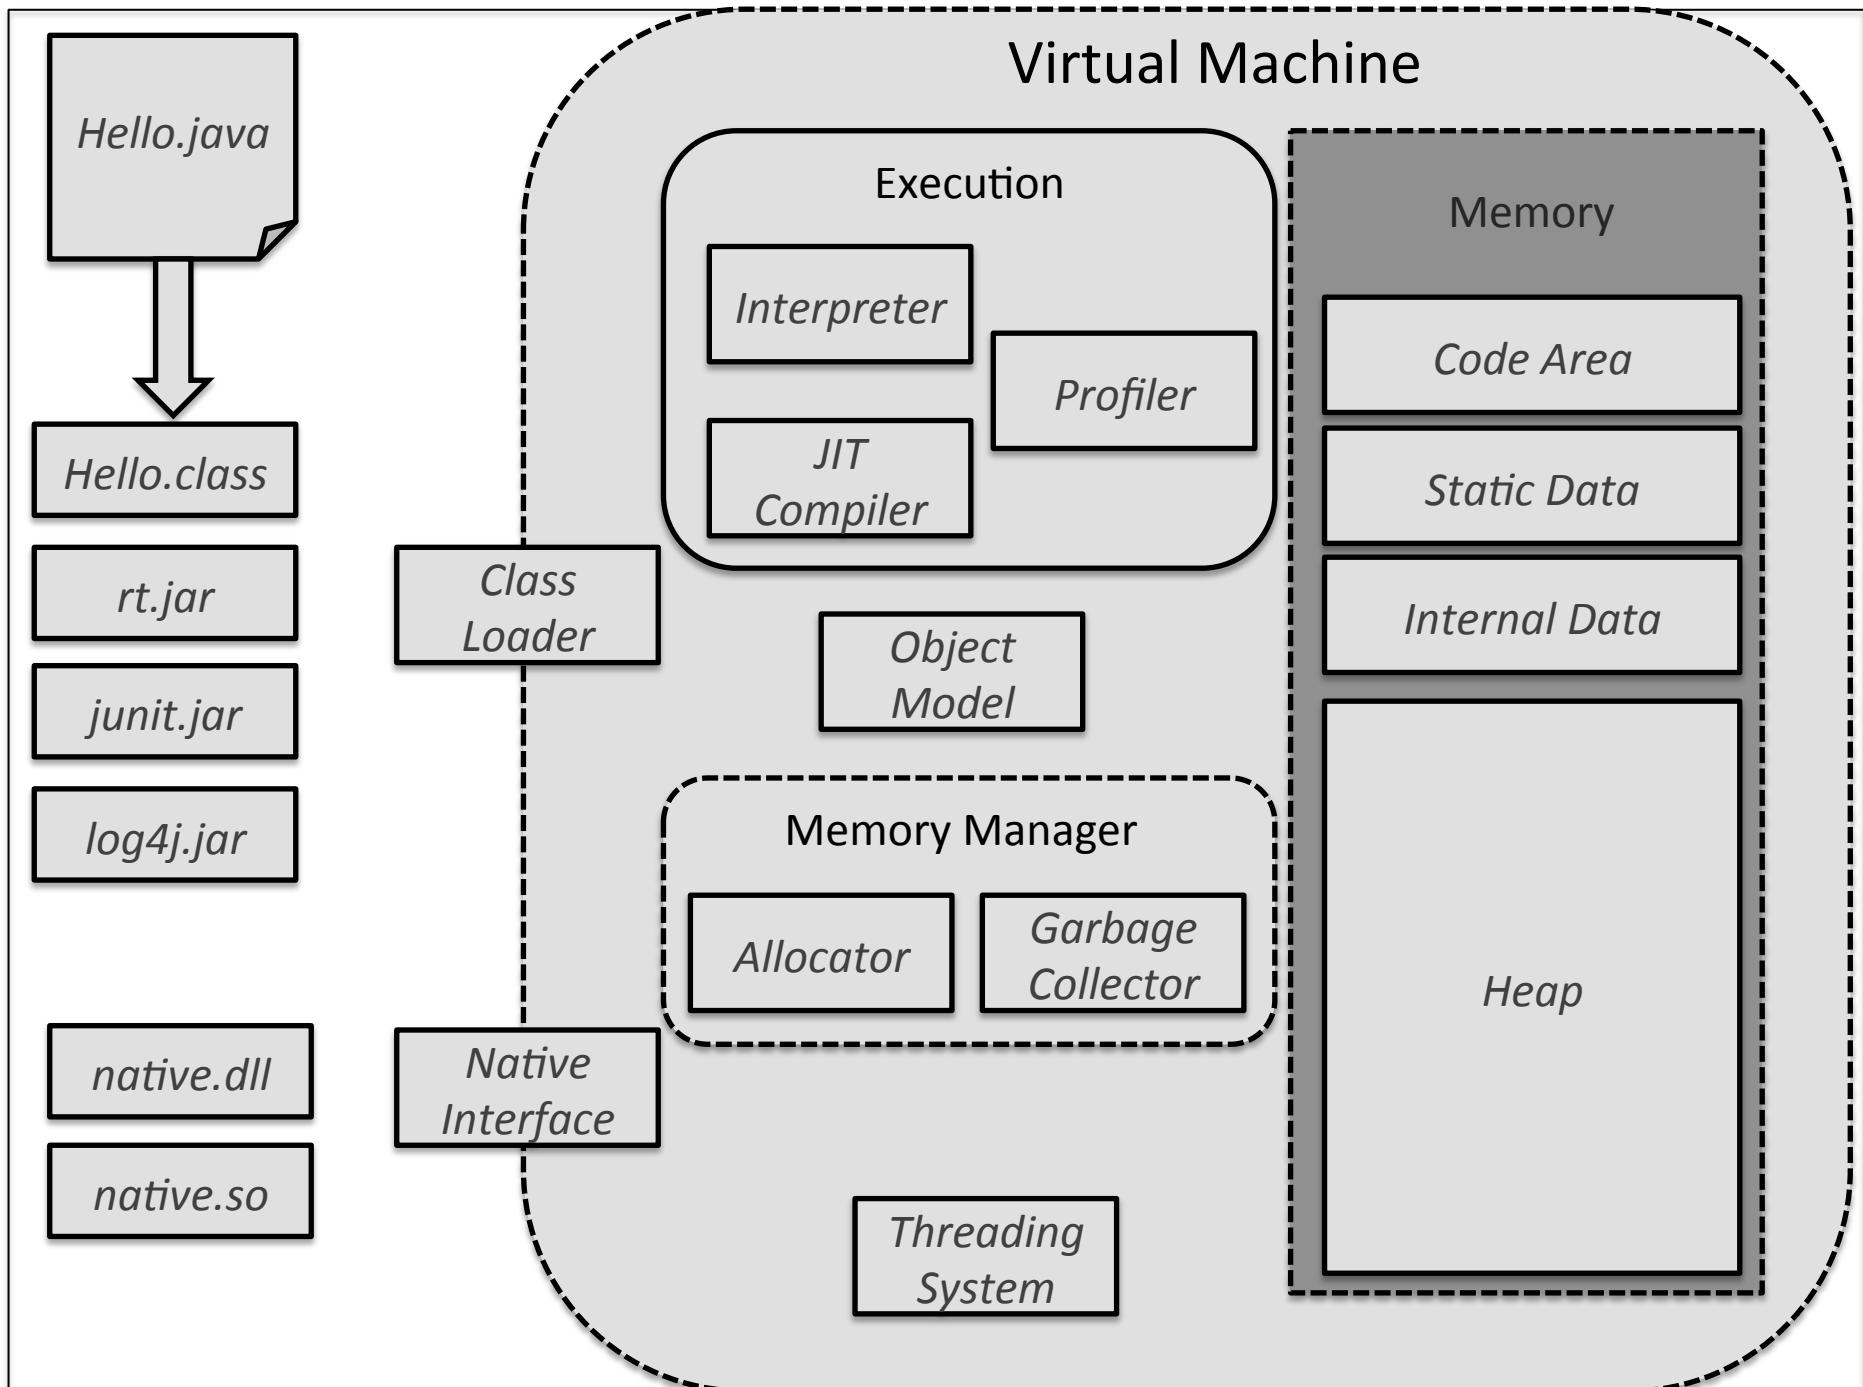
- Some capabilities cause callbacks to event handlers
 - Field modification
 - Field access
 - Method entry/exit
 - Class loading
 - JIT Compilation
 - Native method binding
 - Some allocation
 - Garbage collection
 - Resource exhaustion

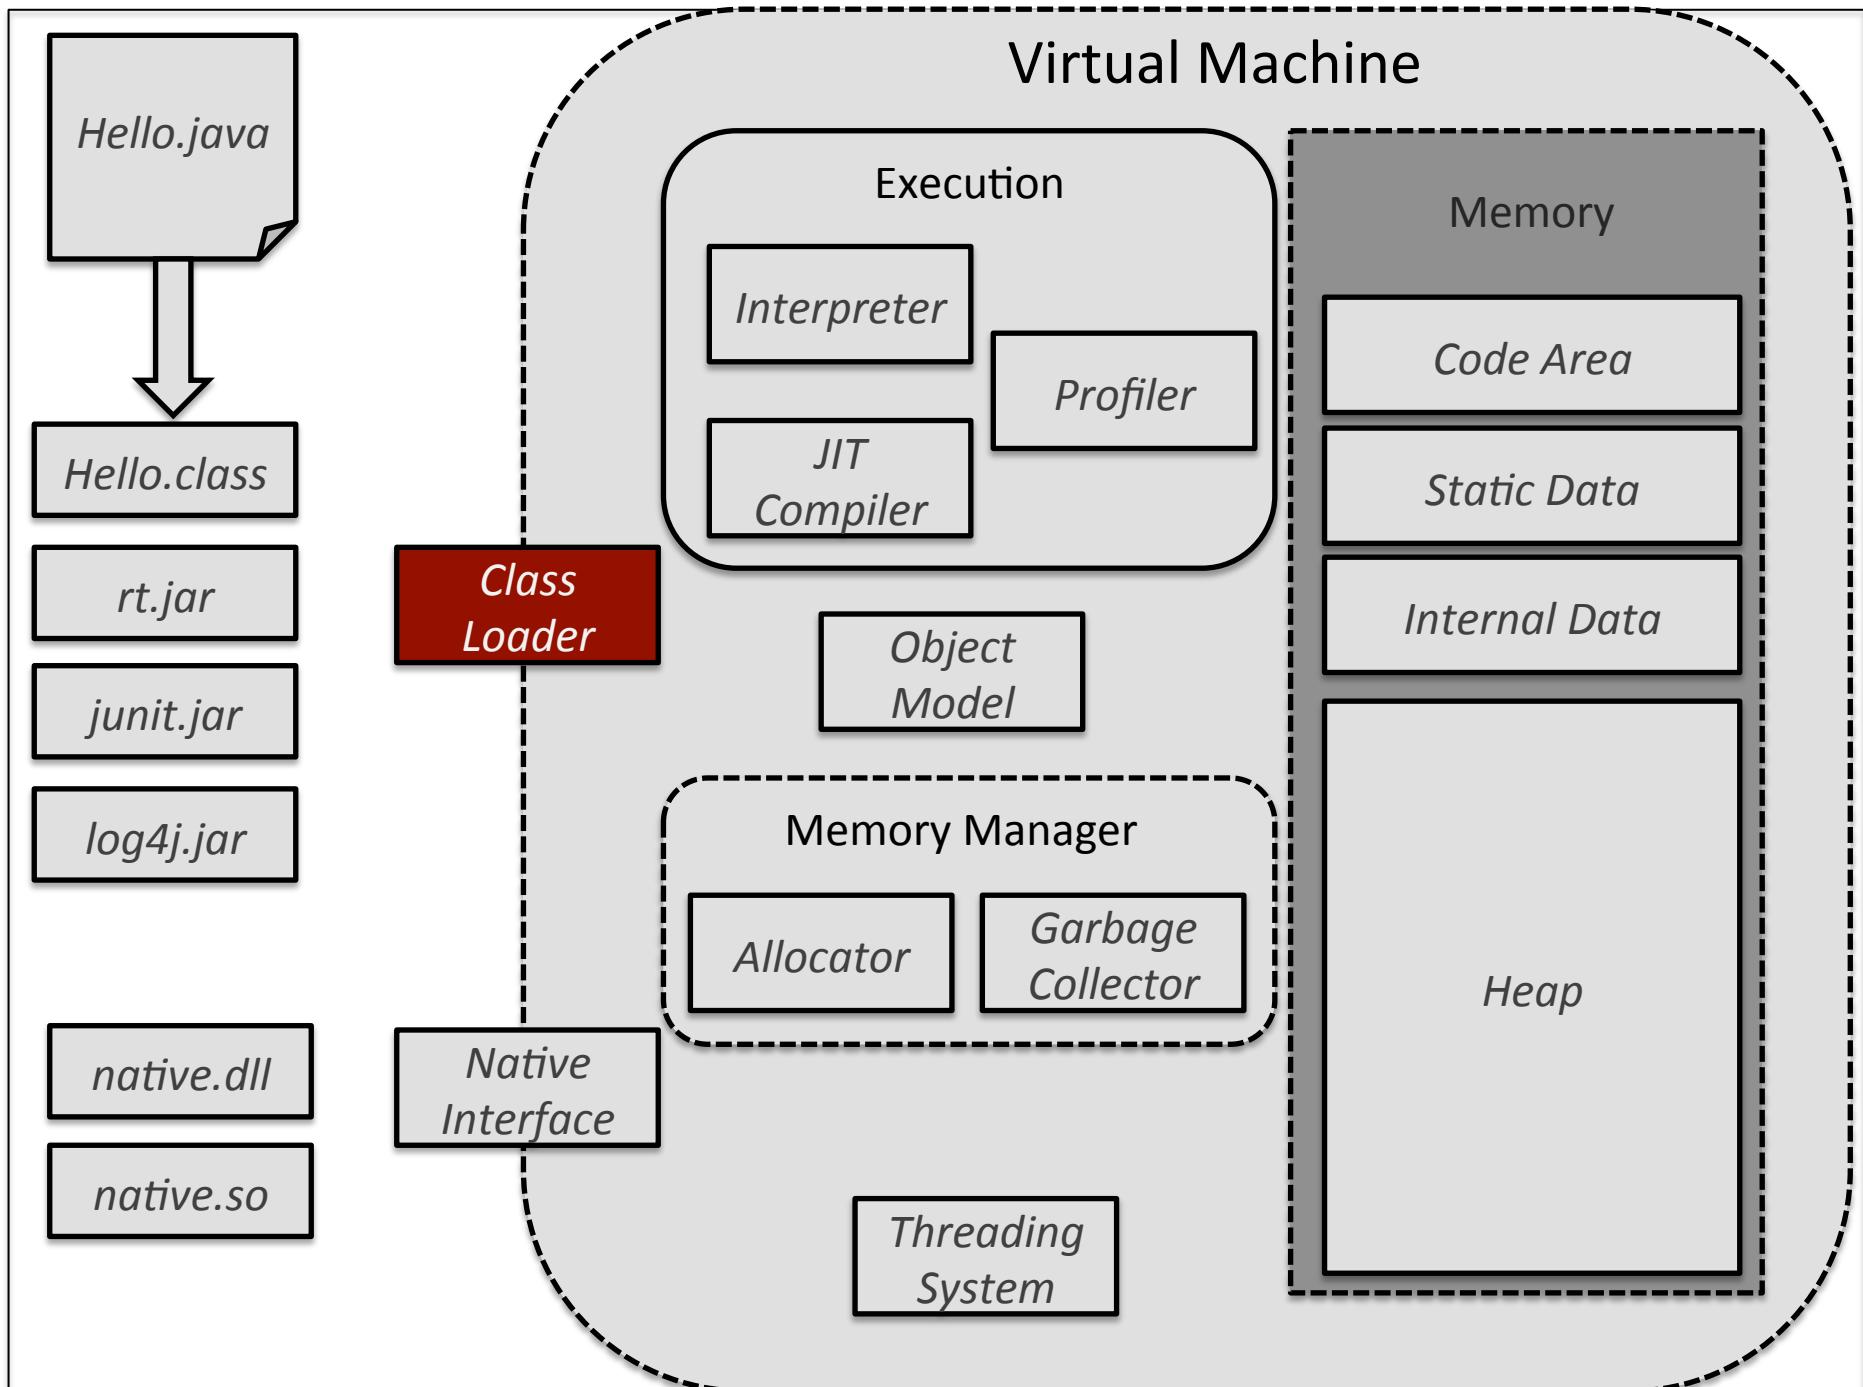
System Monitoring Capabilities

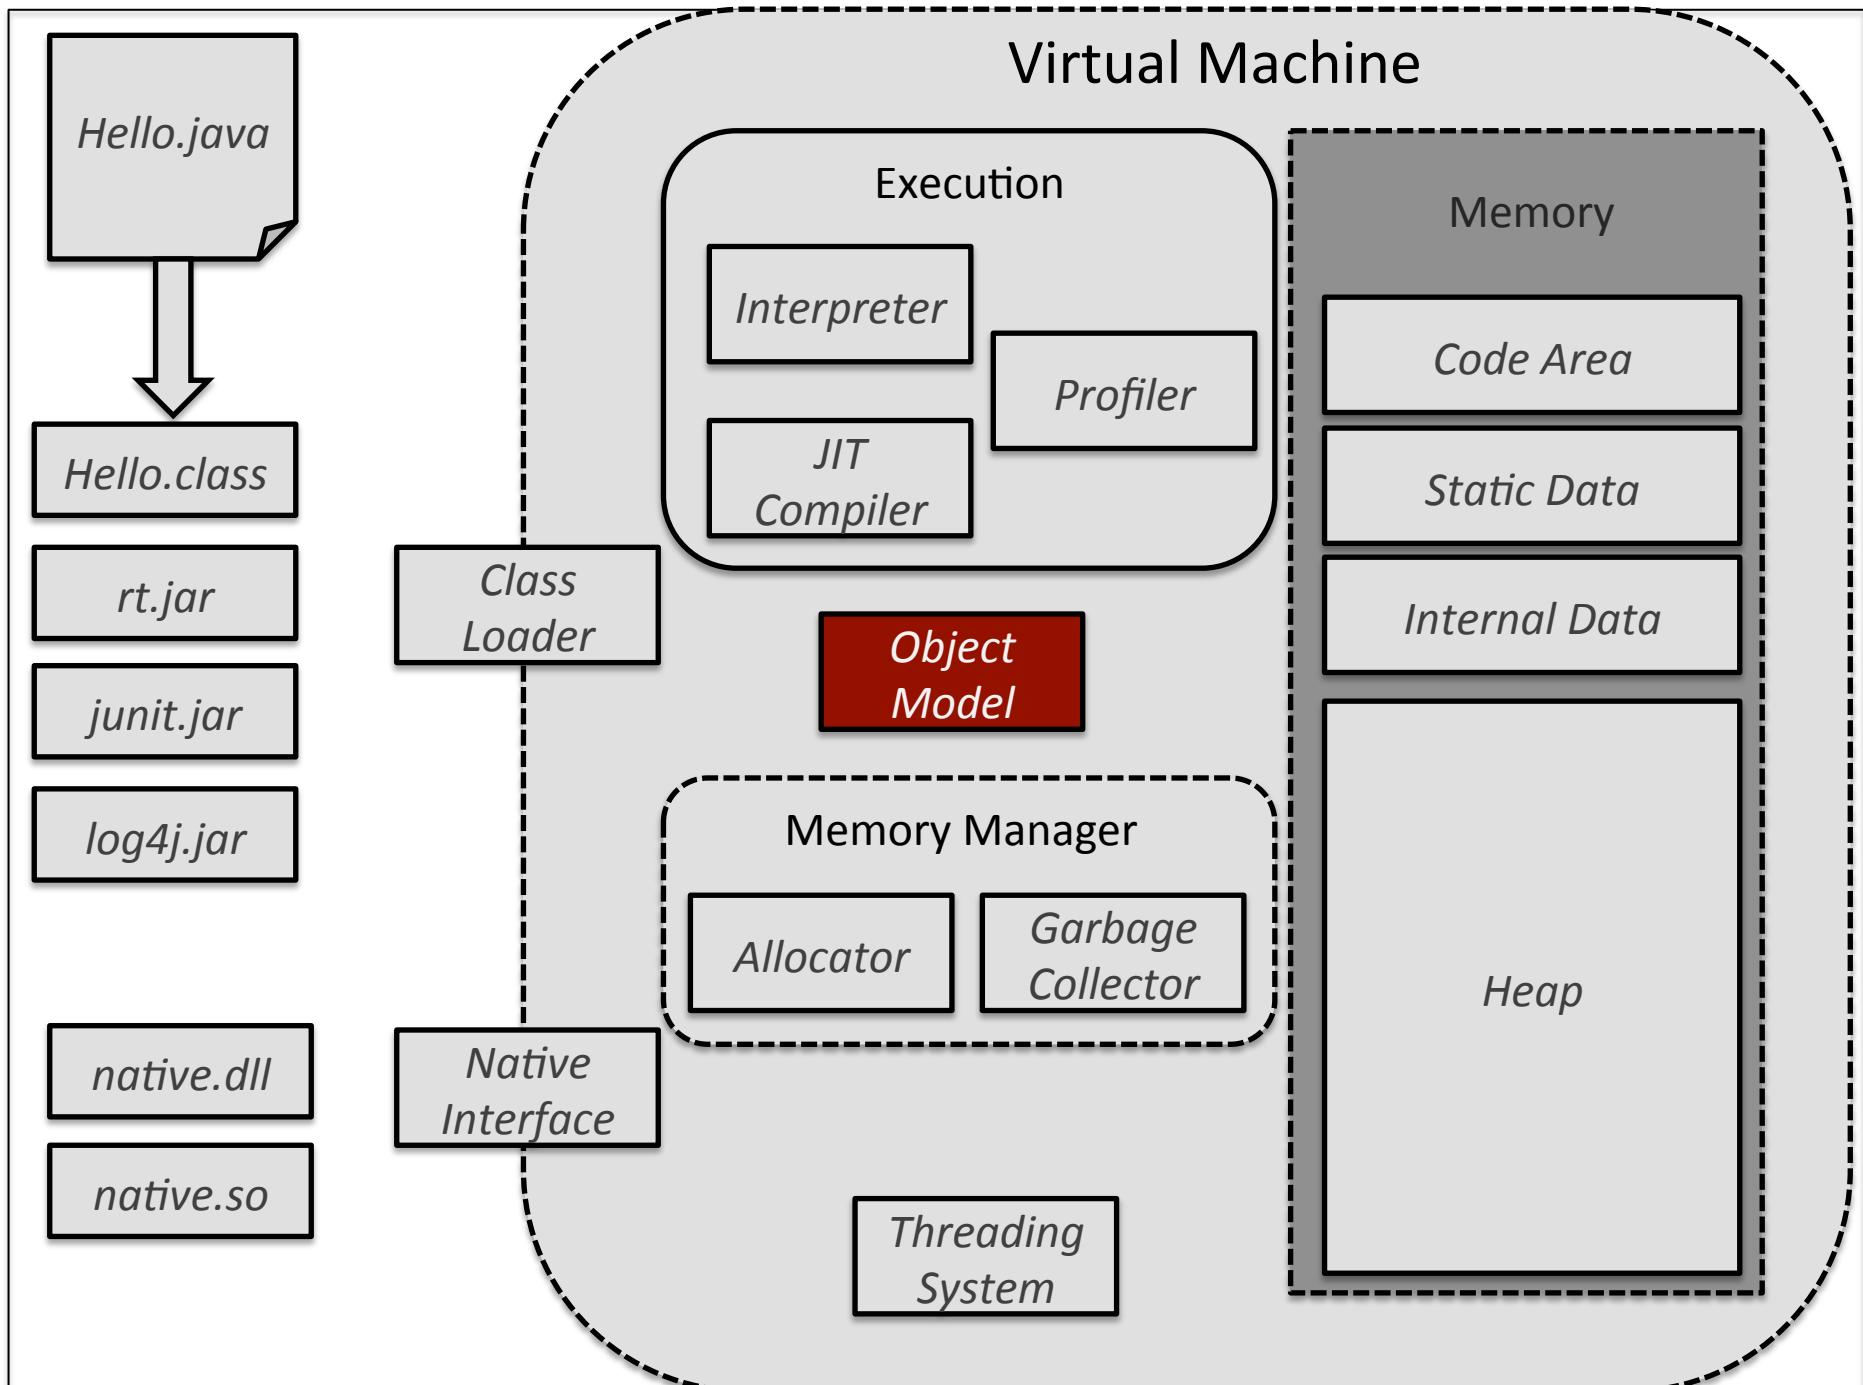
- Others allow tools to access system information
 - Get monitor states
 - Get bytecodes for a method
 - Get source file and line numbers
 - Get class constant pool
 - Get CPU time for a given thread

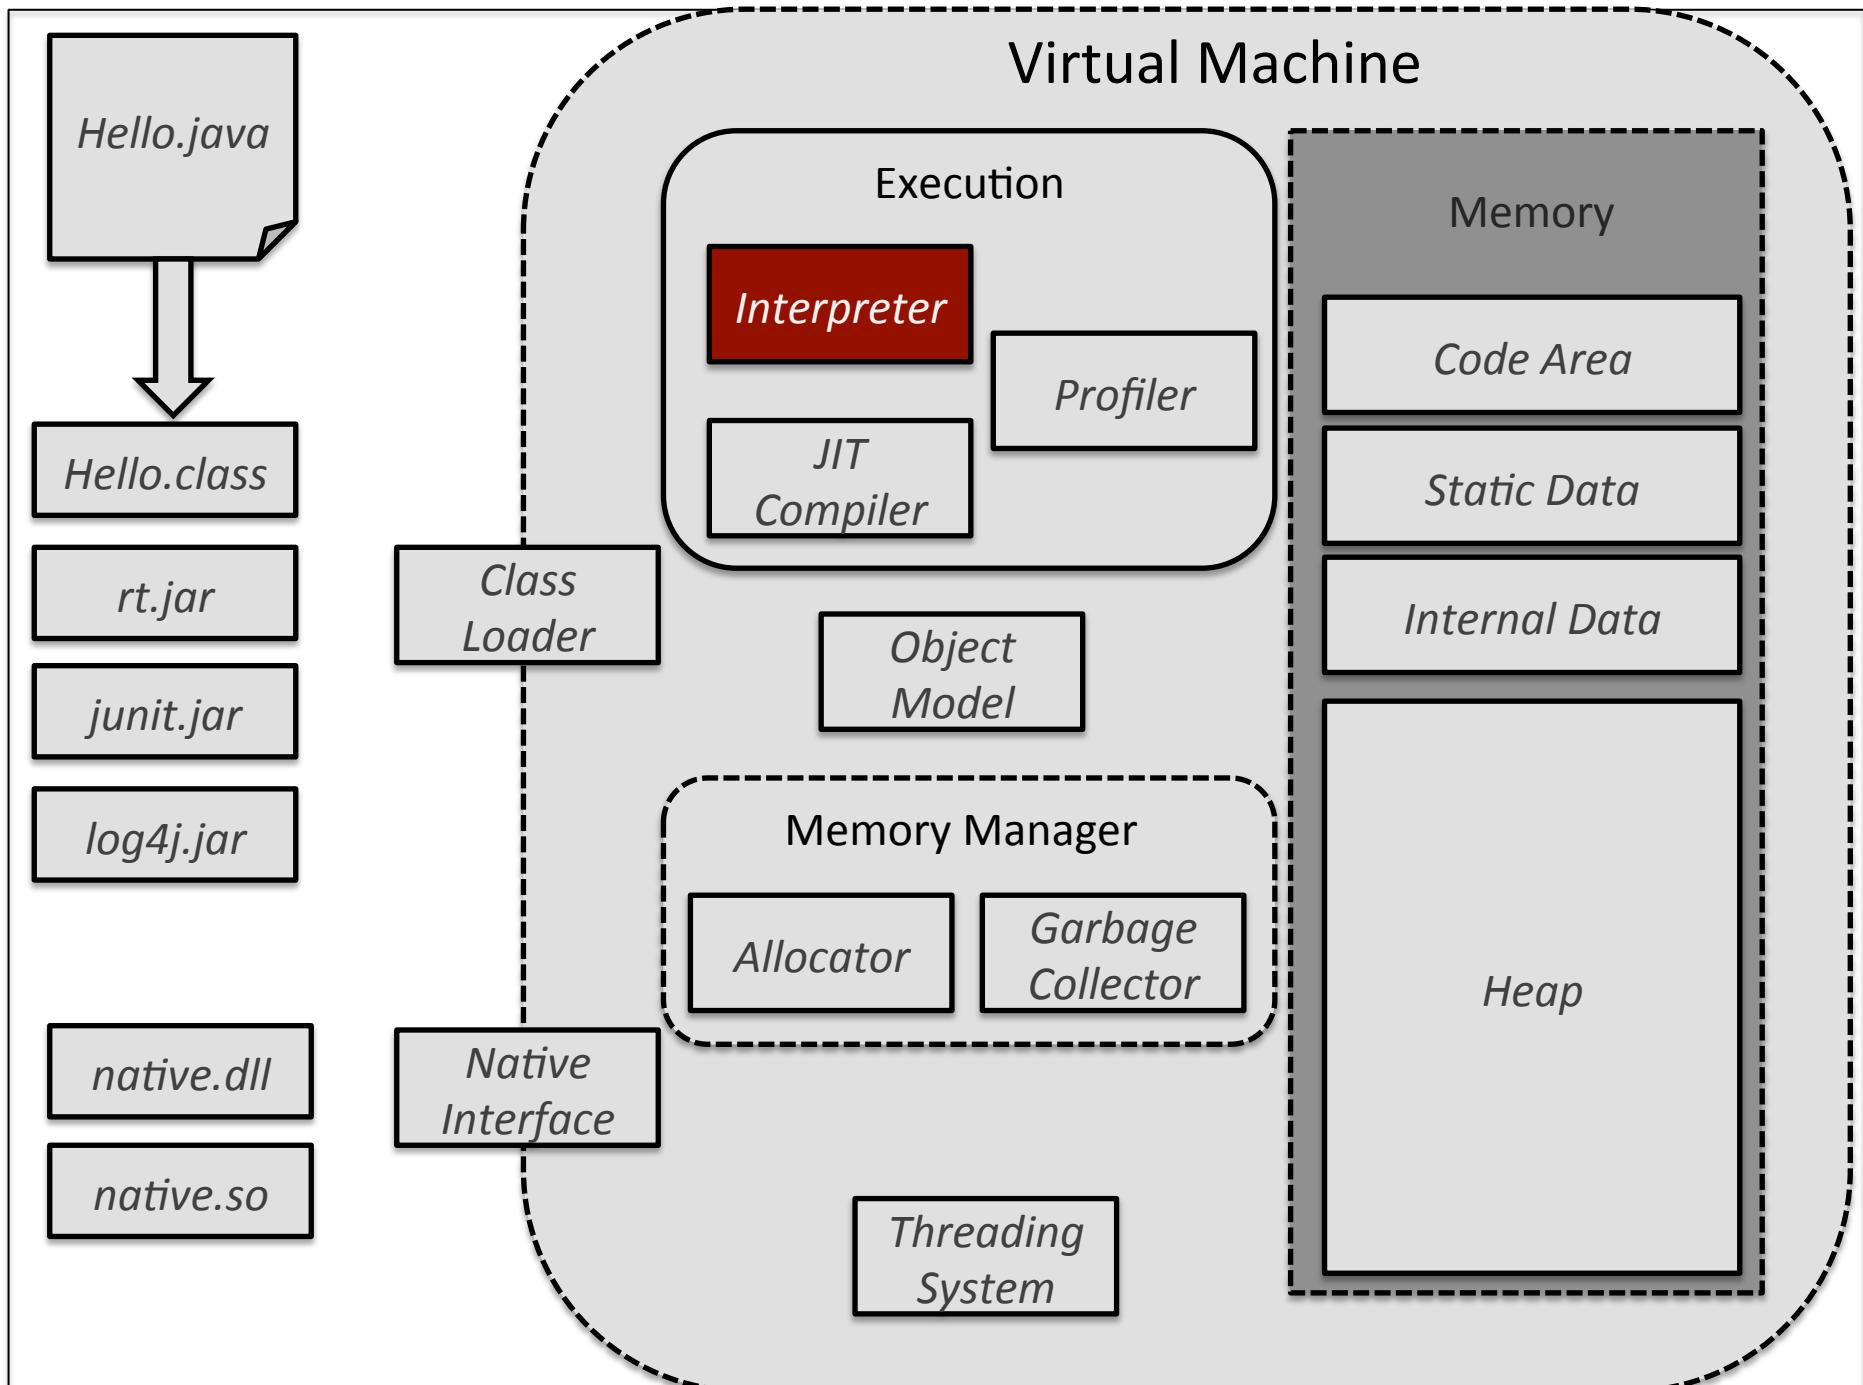
Agent Functionality Capabilities

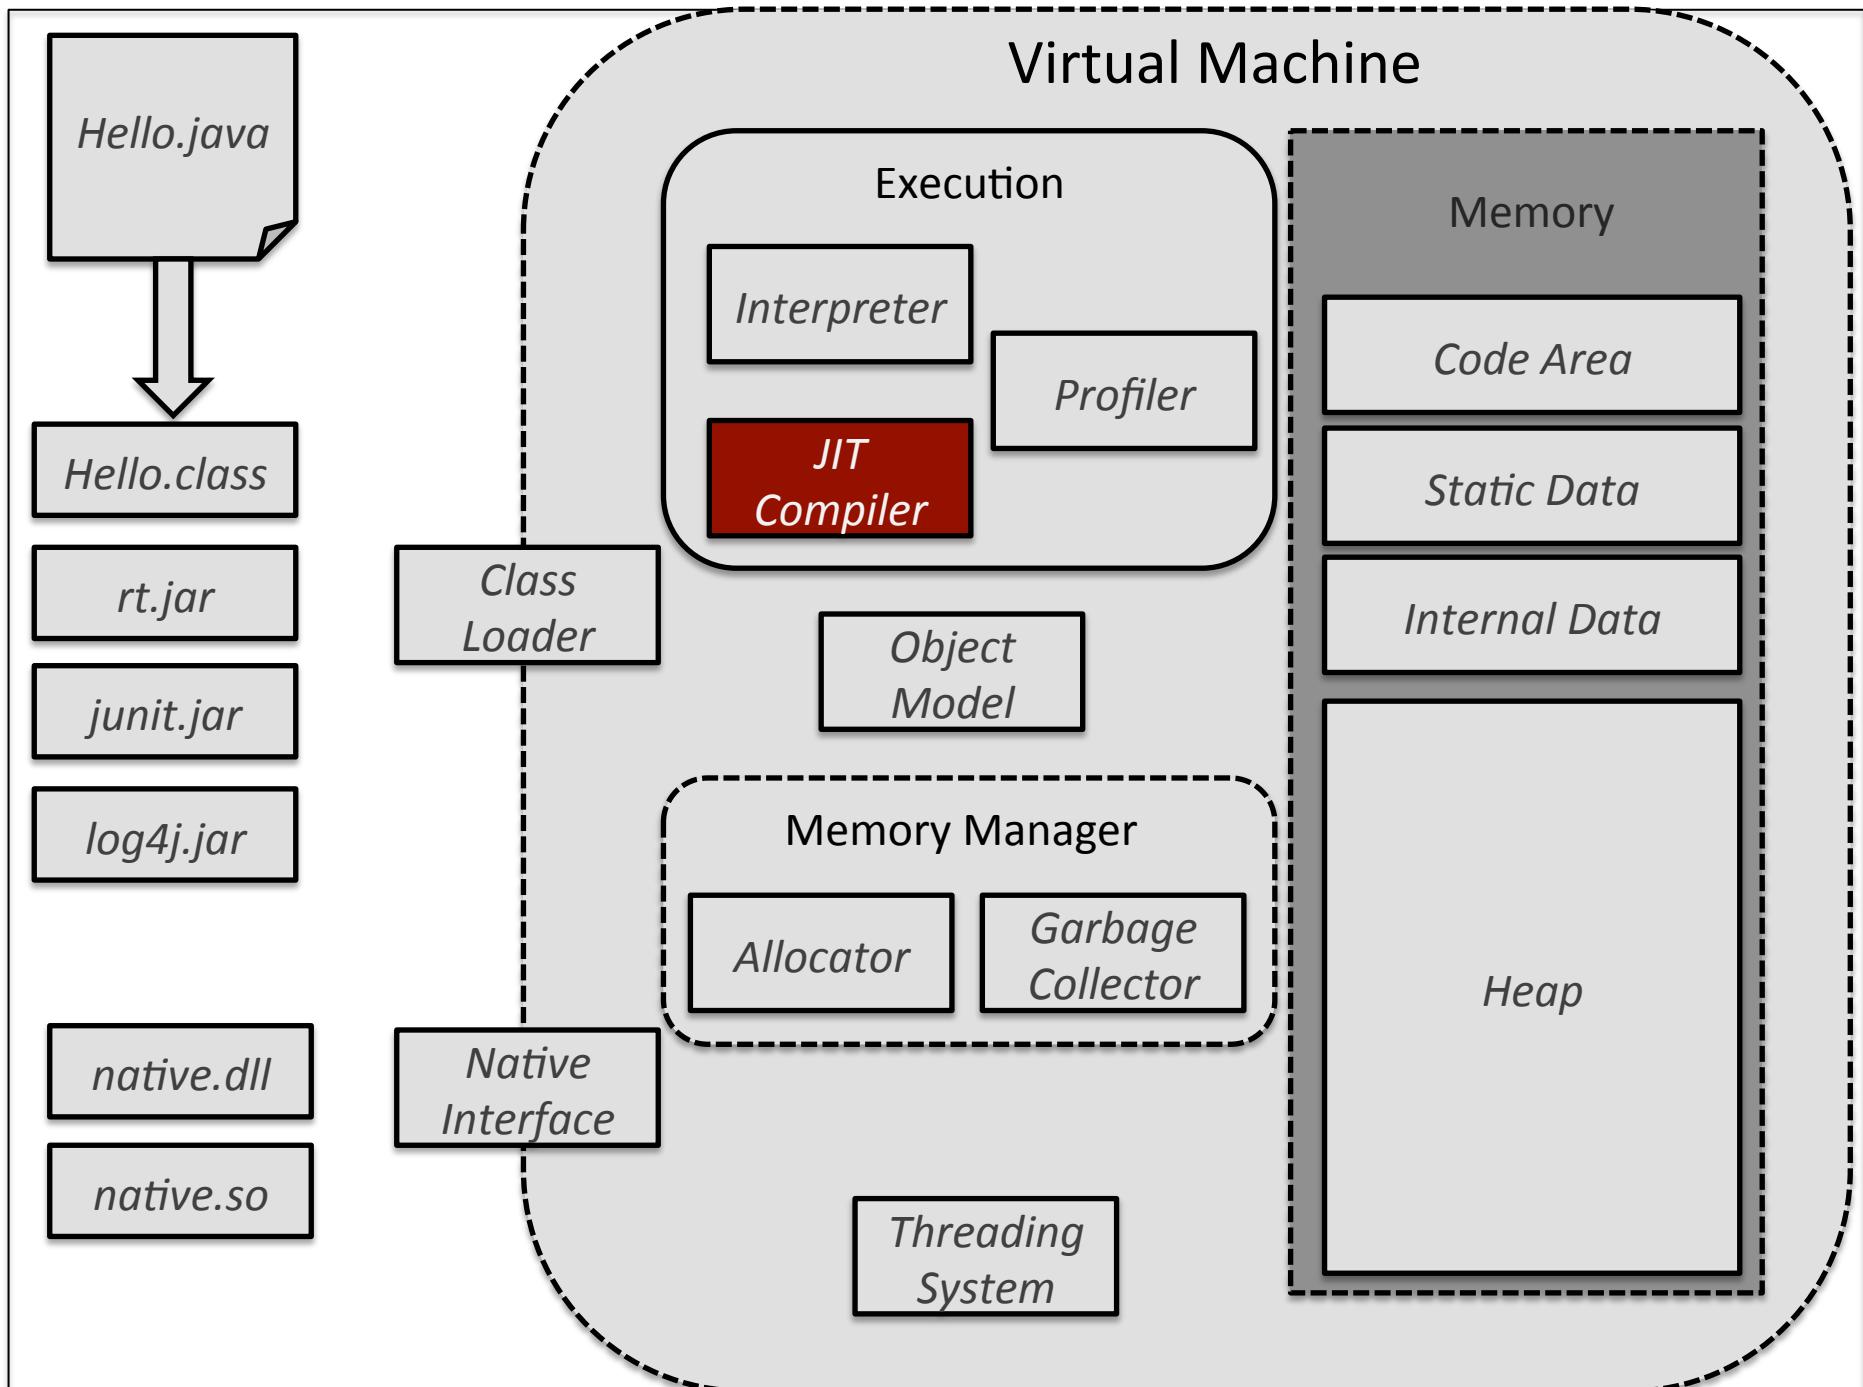
- Other capabilities let the agent perform actions
 - Tag objects on the heap
 - Pop frames from the stack or force method return
 - Redefine or re-transform classes
 - Stop, suspend or interrupt threads
 - Set breakpoints
 - Single-step through code
 - Generate events for other agents

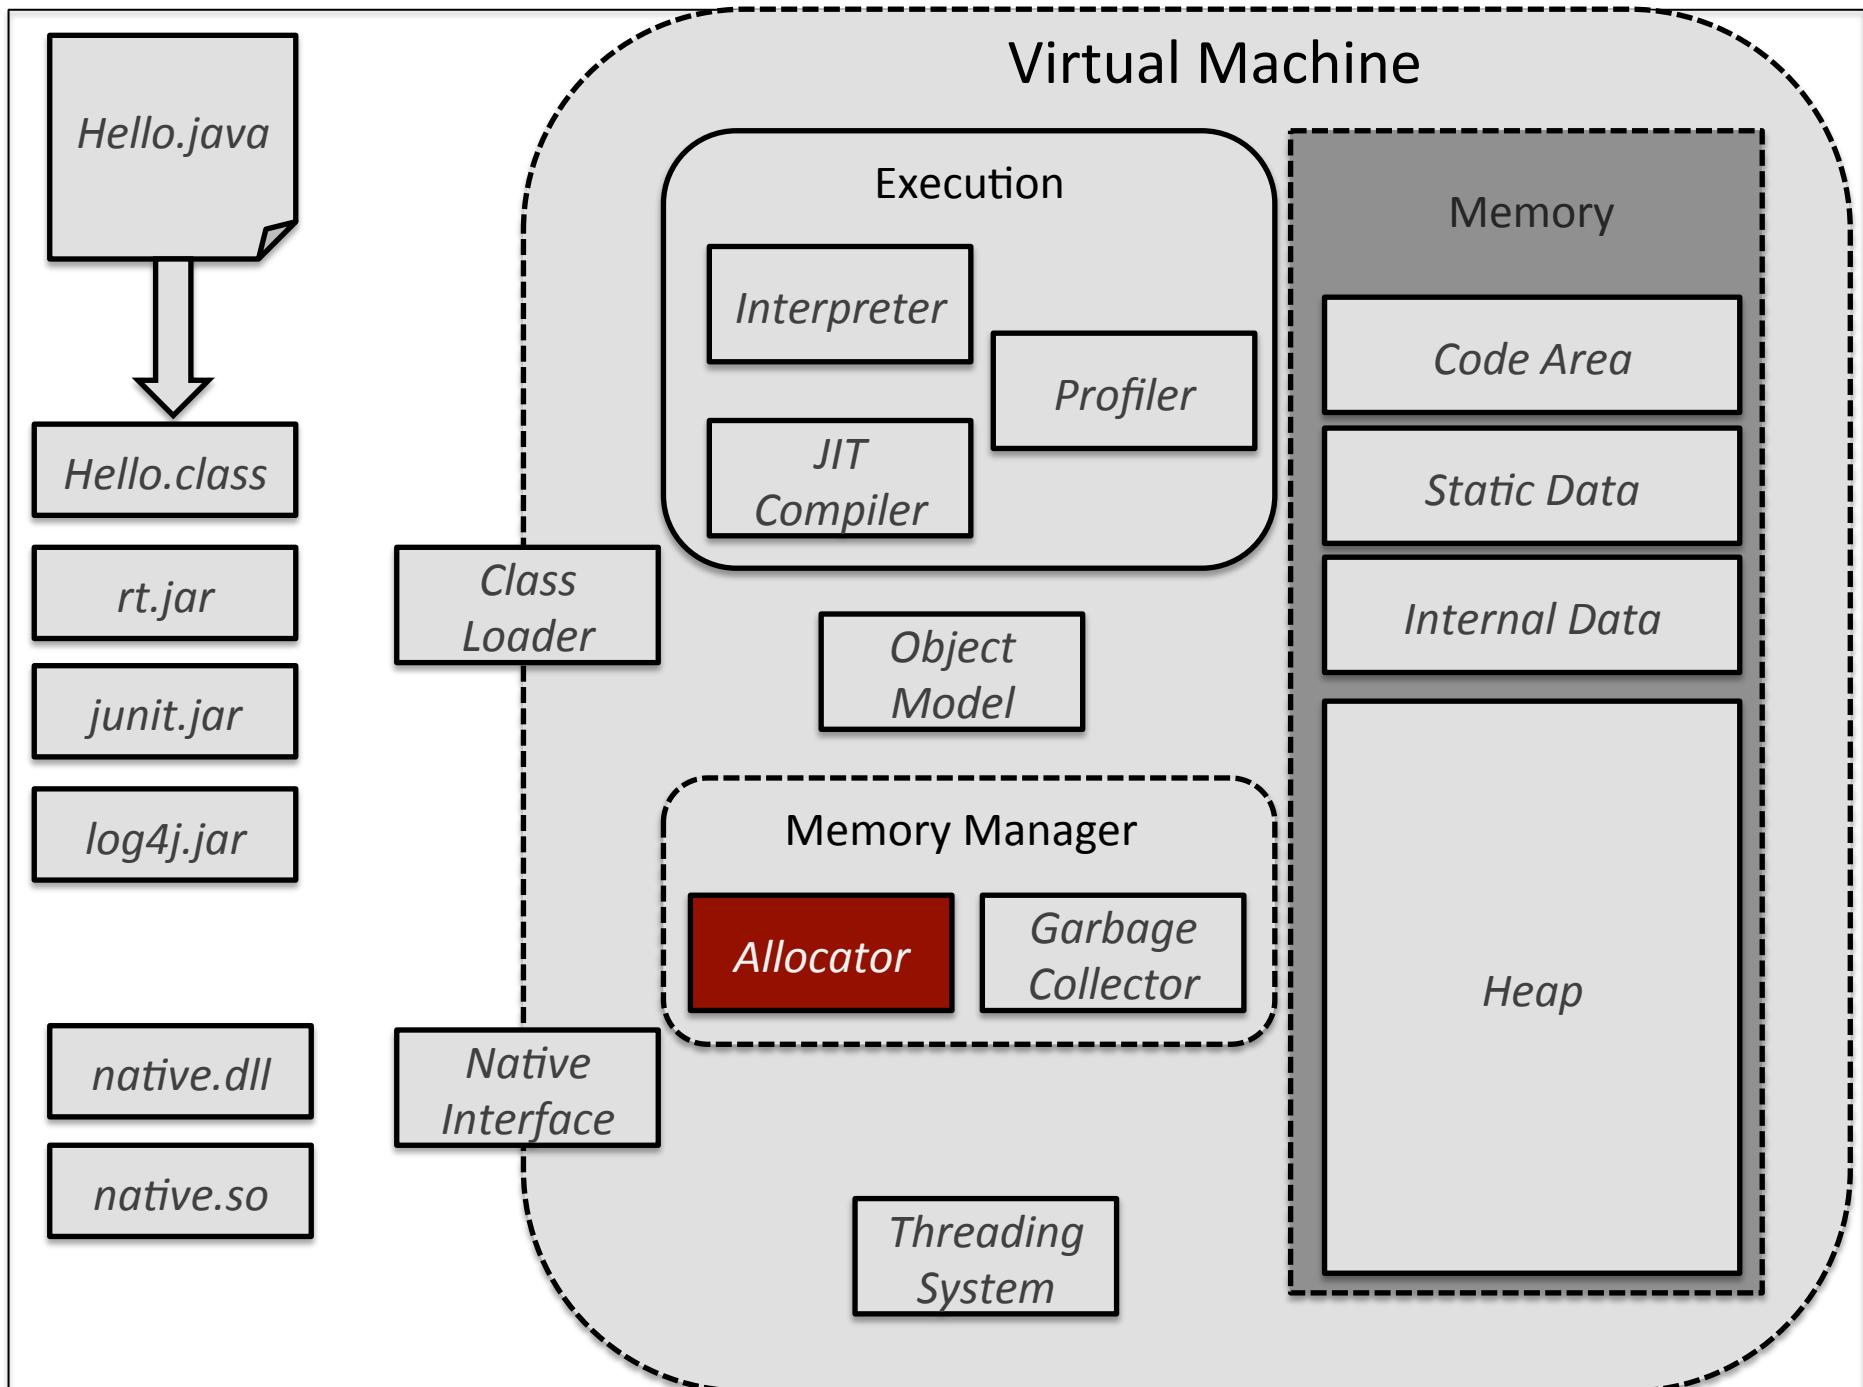


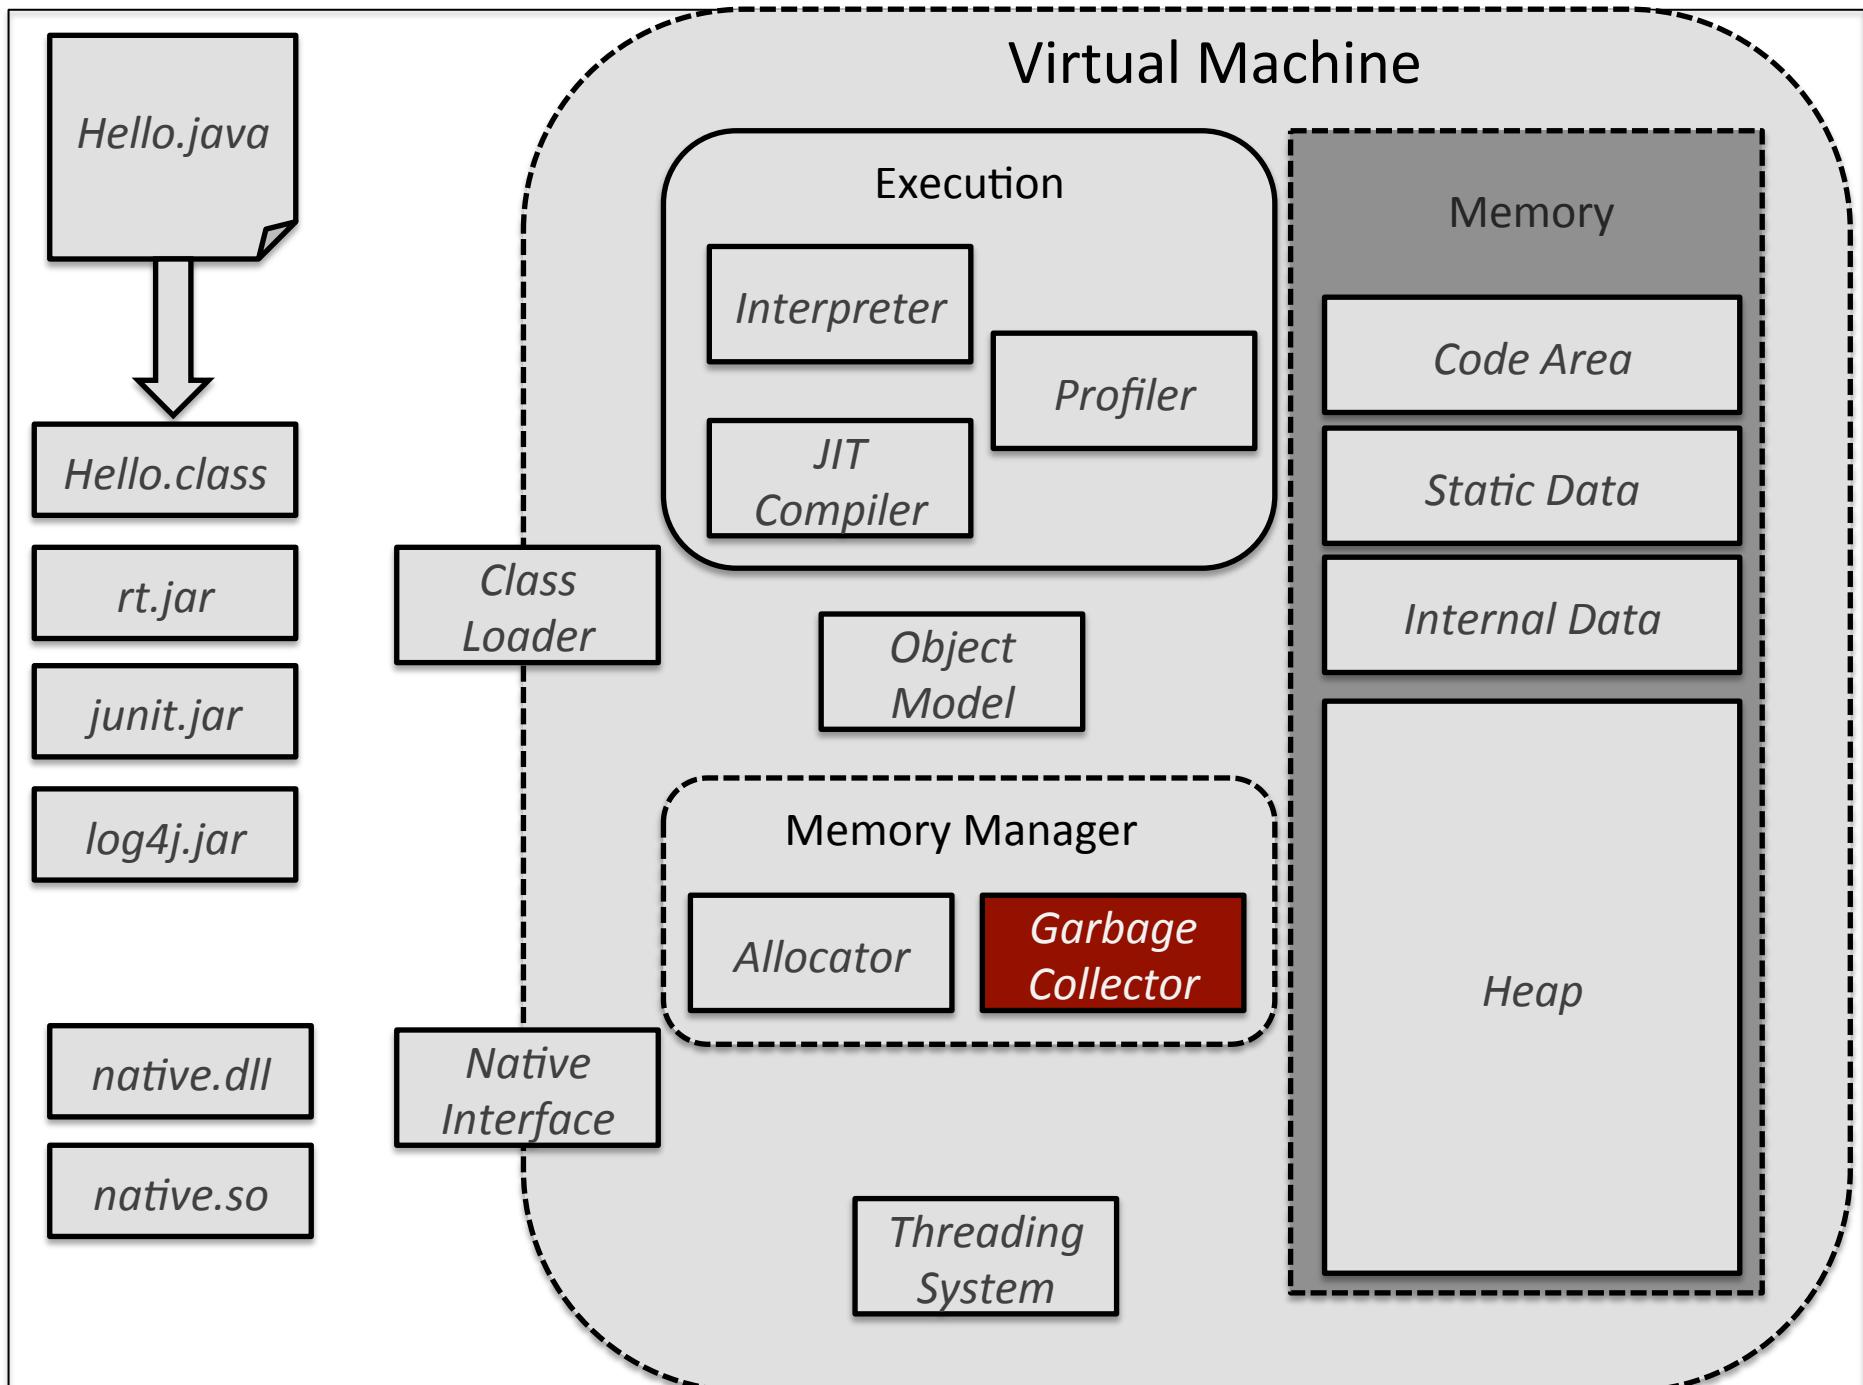


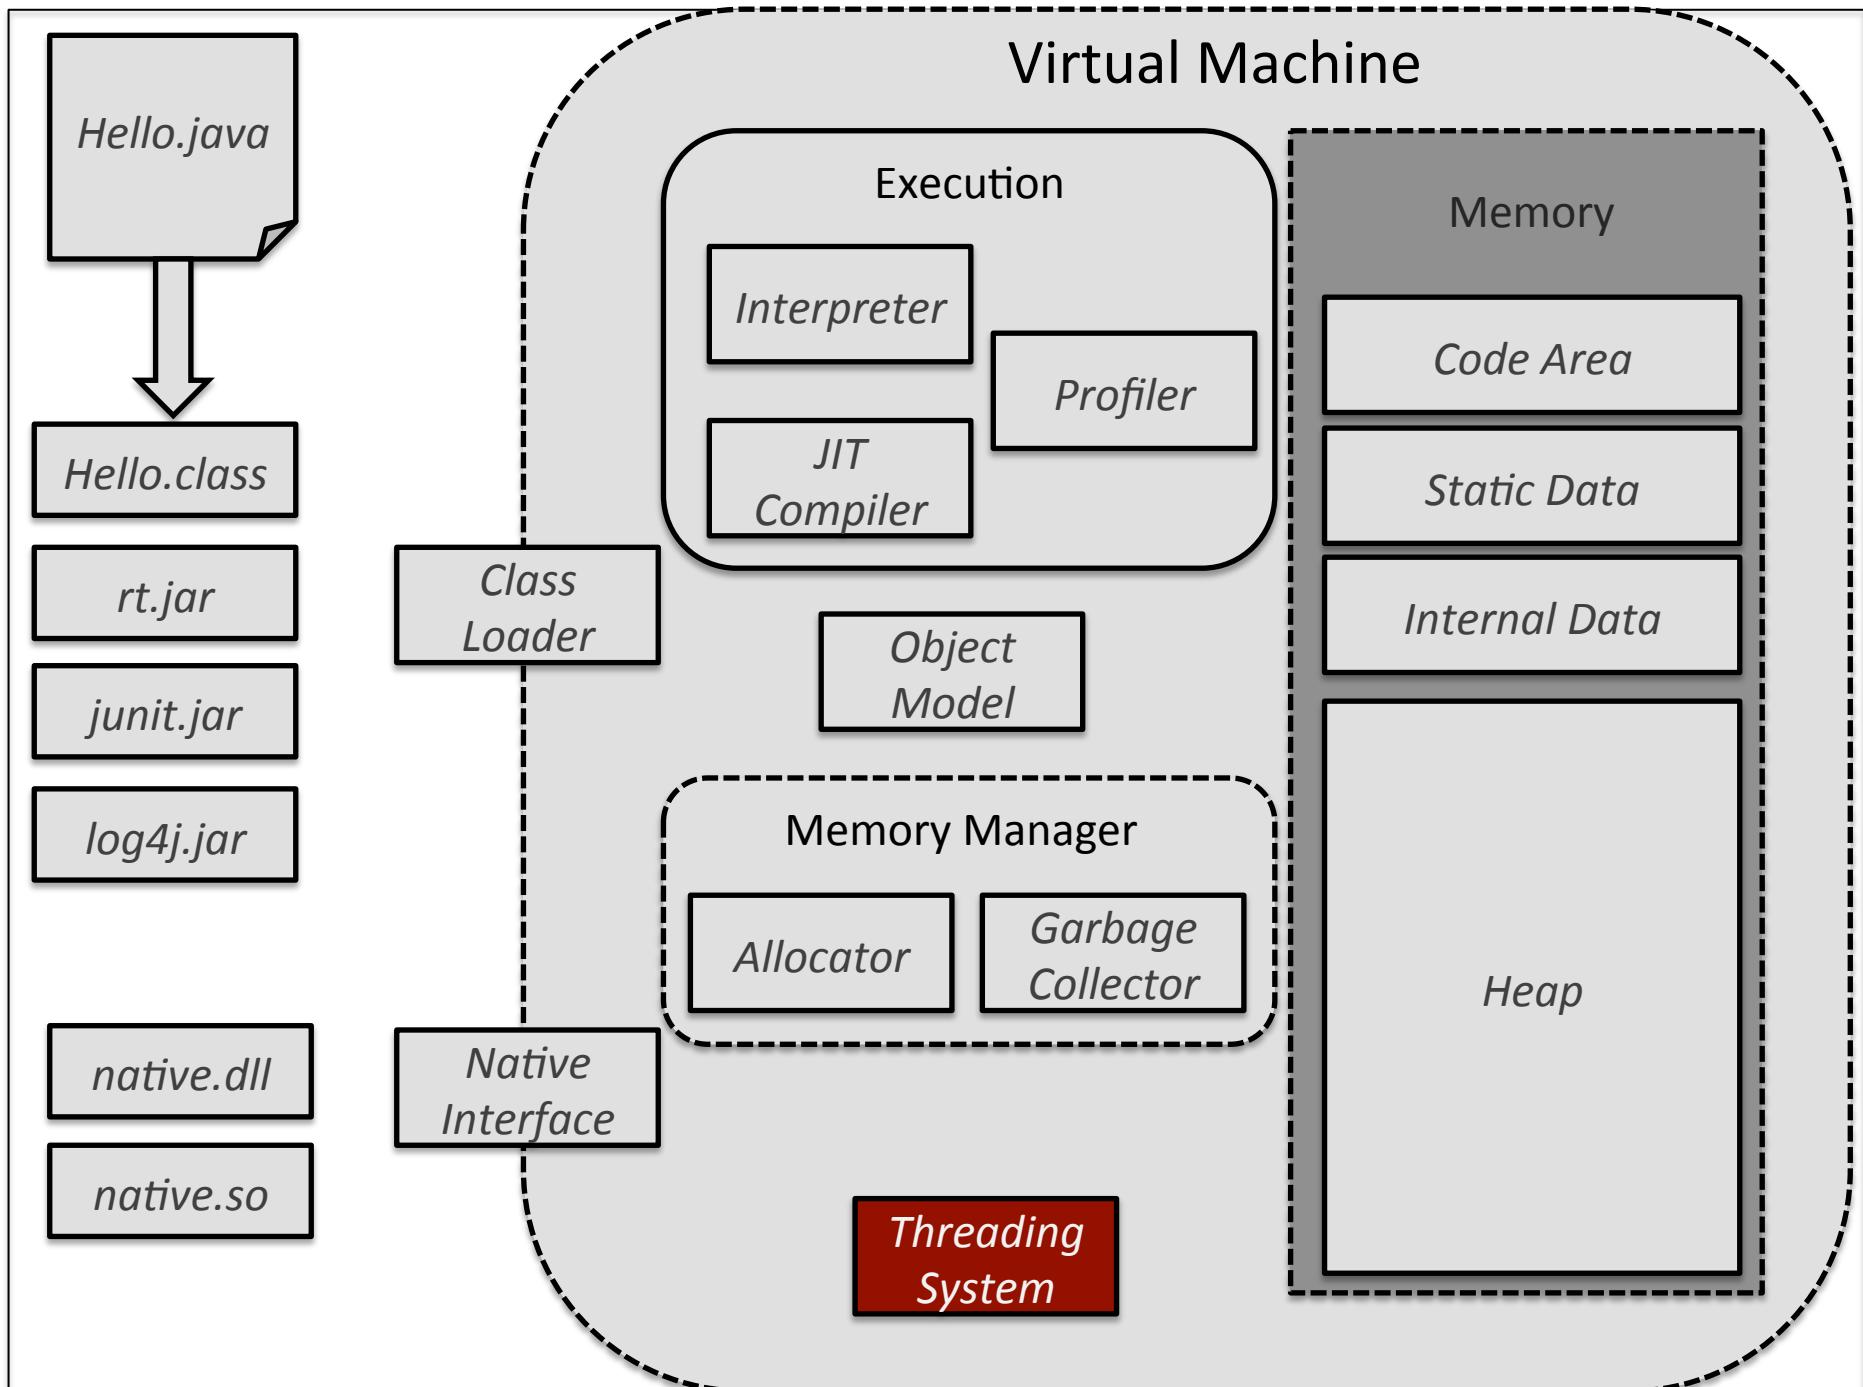


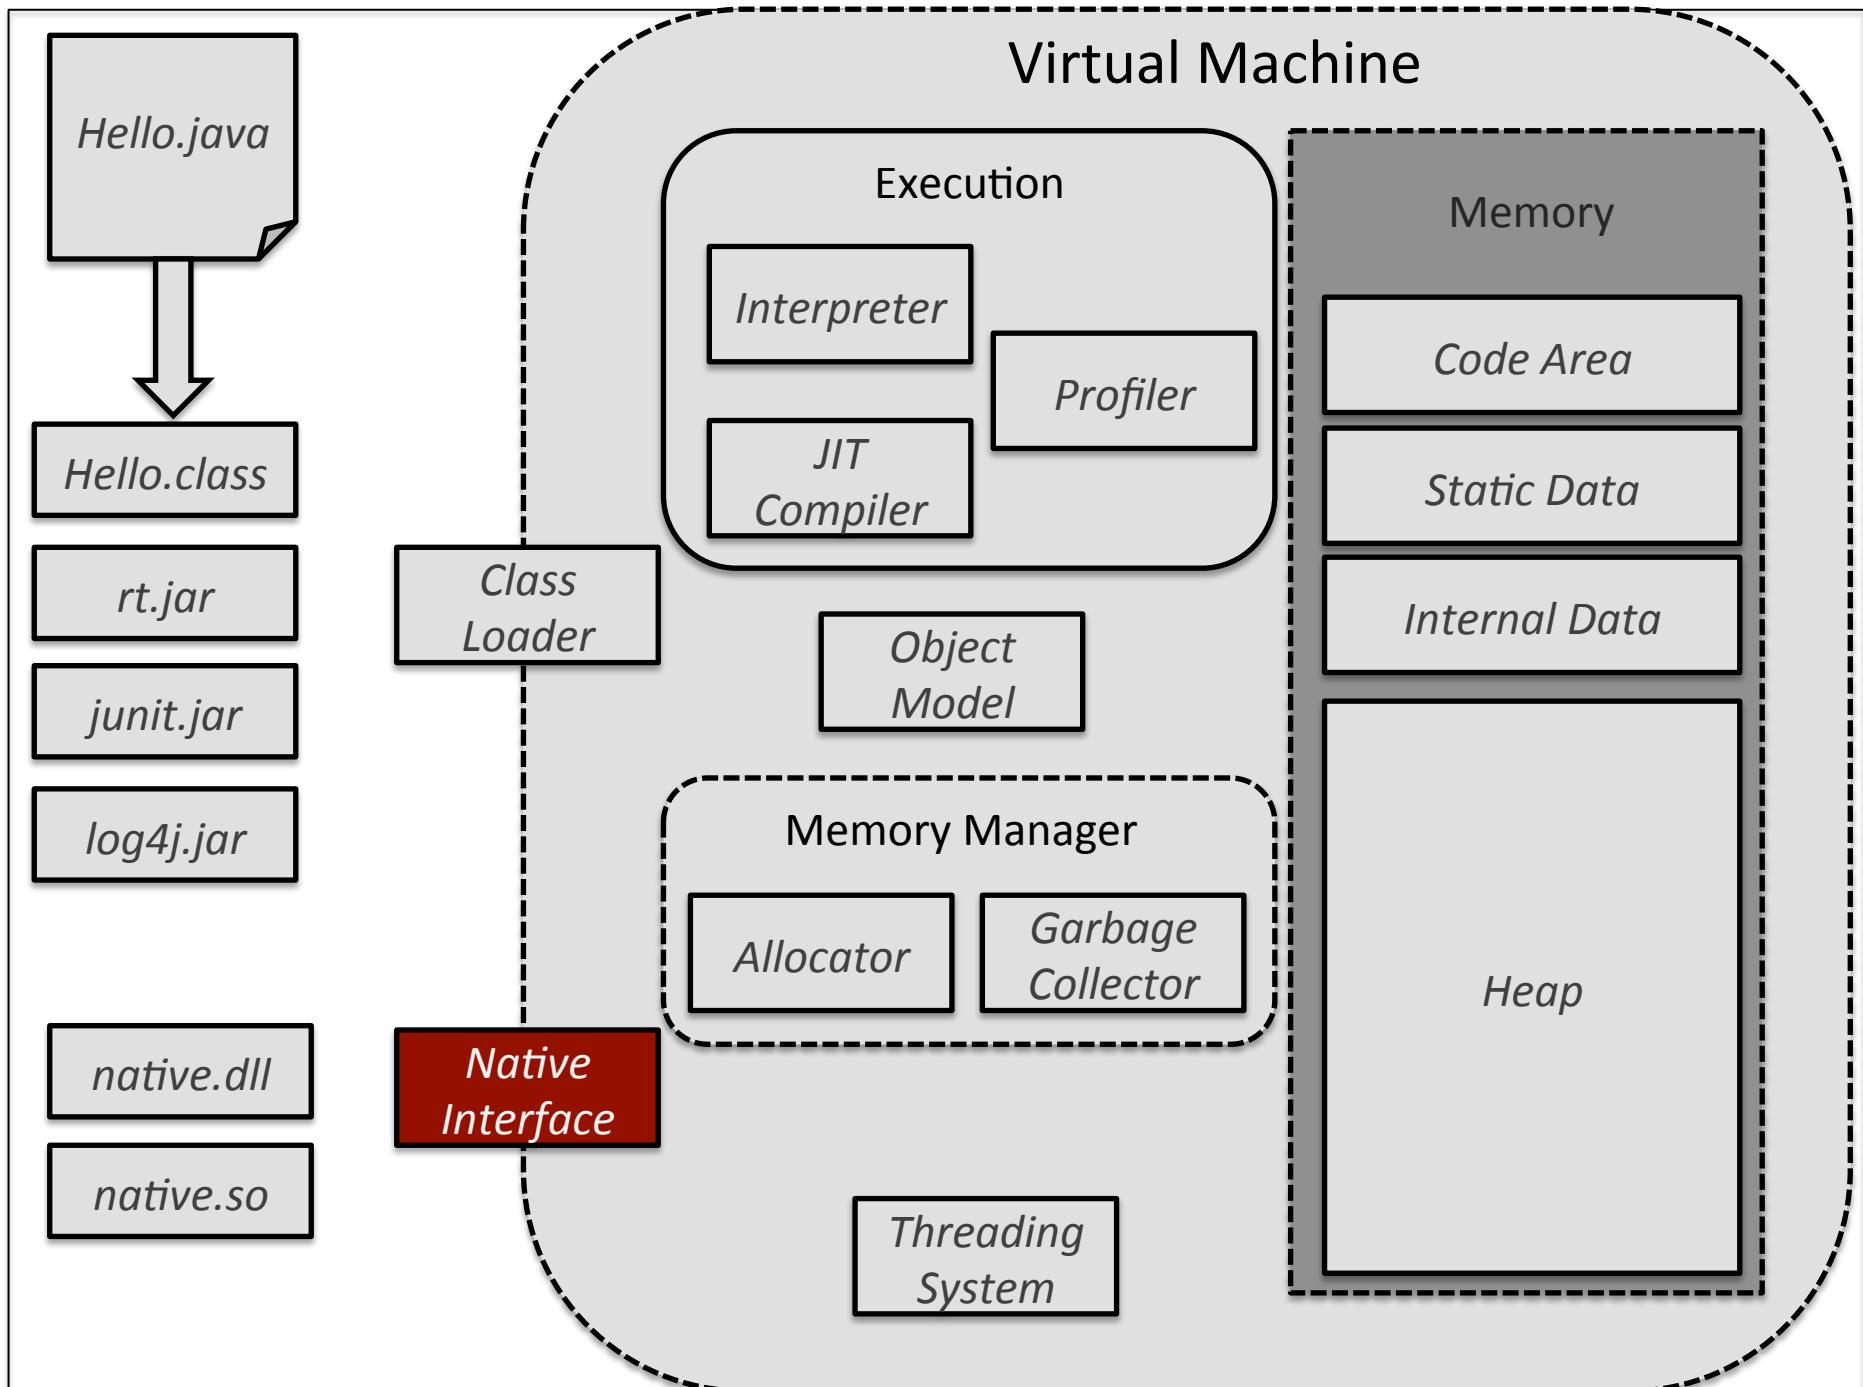












Case Study: Debugging

- Build the external debug tool
- Build the JVMTI agent
 - May need a separate implementation per platform
- Register the agent for relevant capabilities
- Relay events to the external tool

Virtual Machine

VM Data Structures



Threads

Heap

Virtual Machine

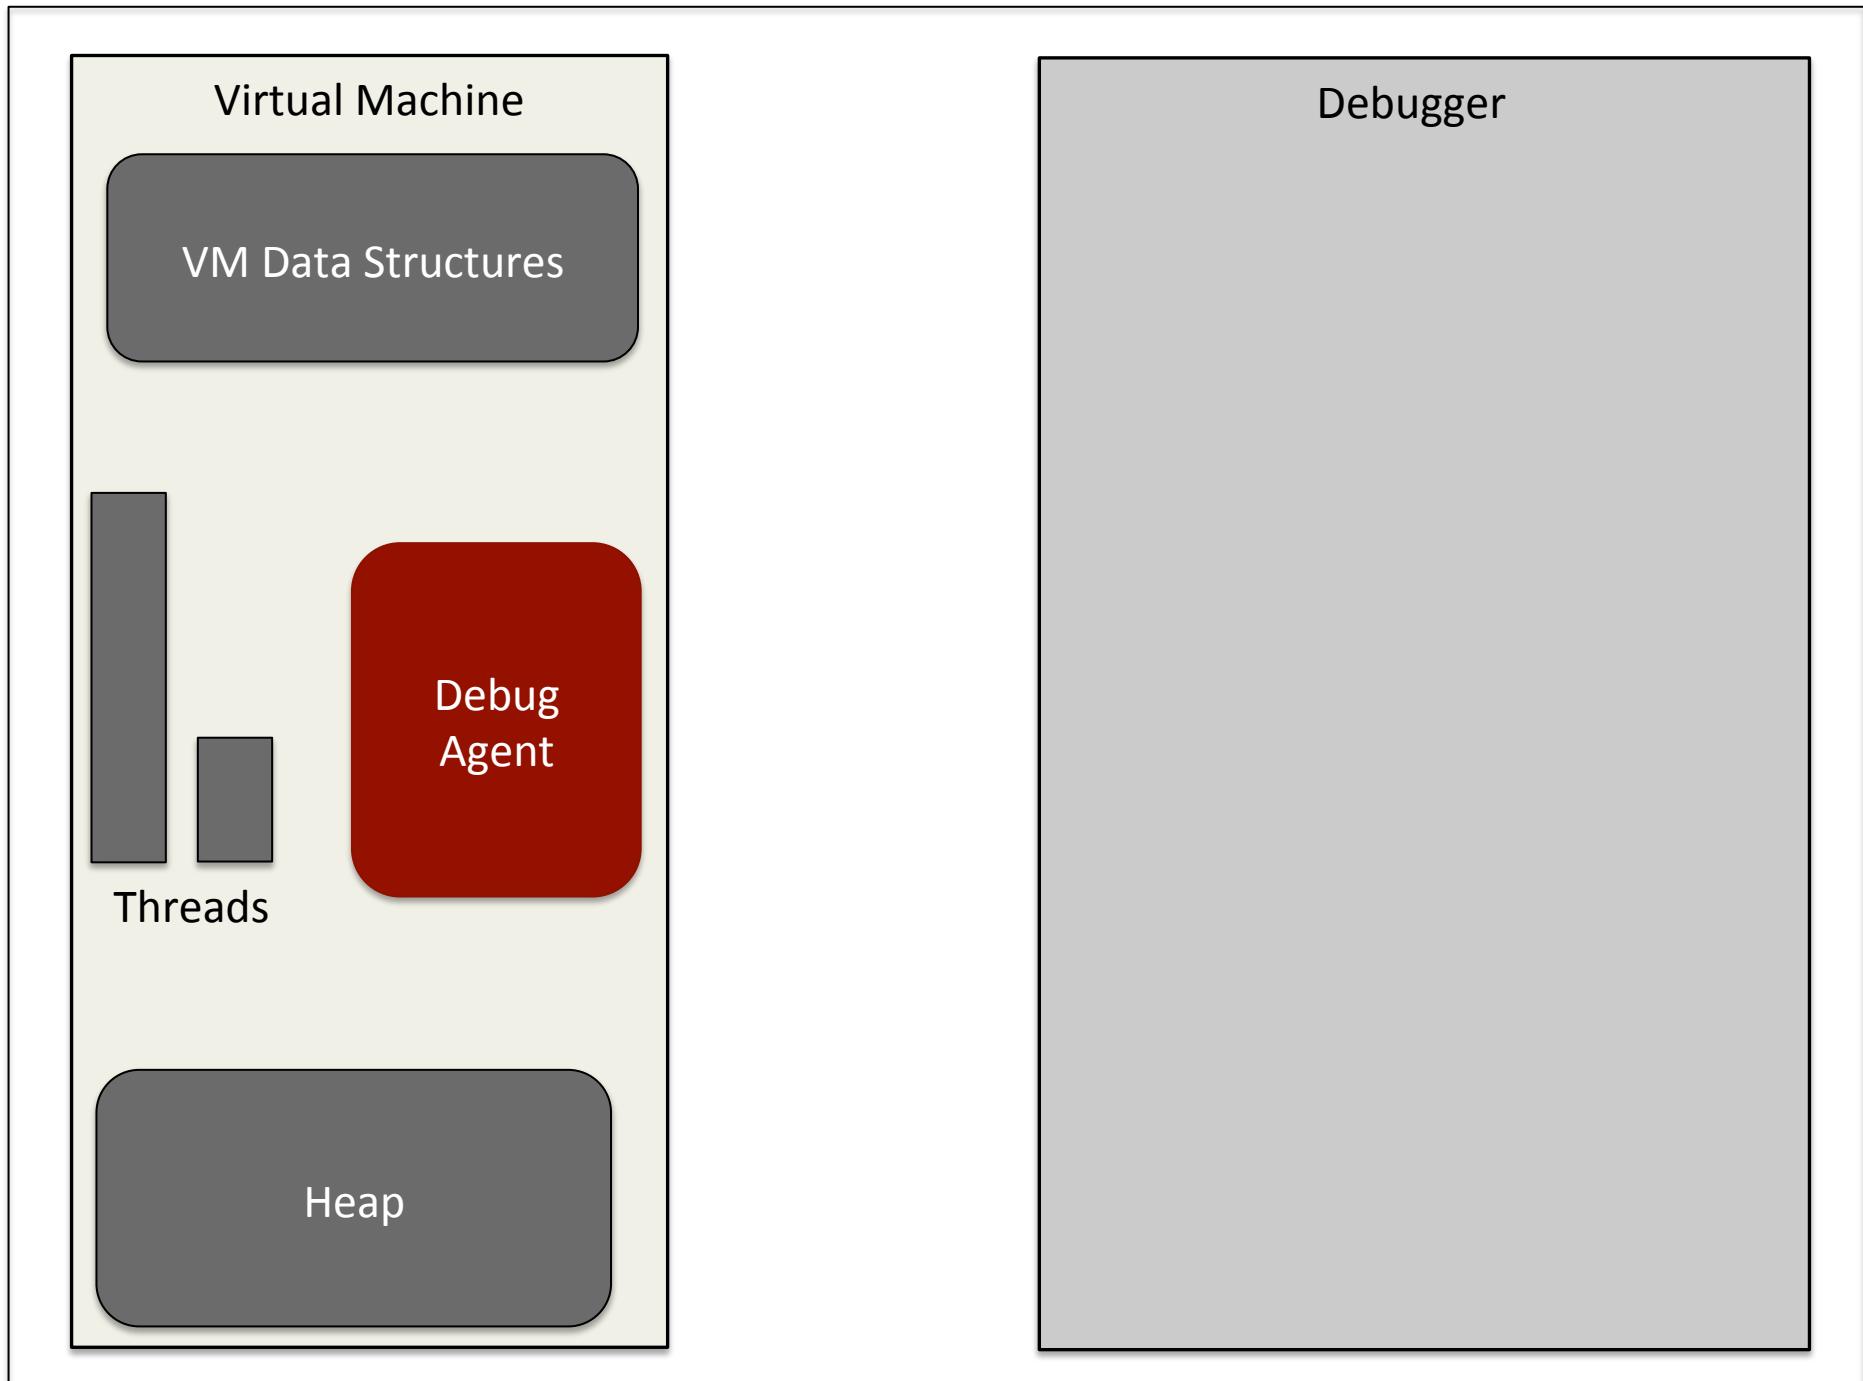
VM Data Structures



Threads

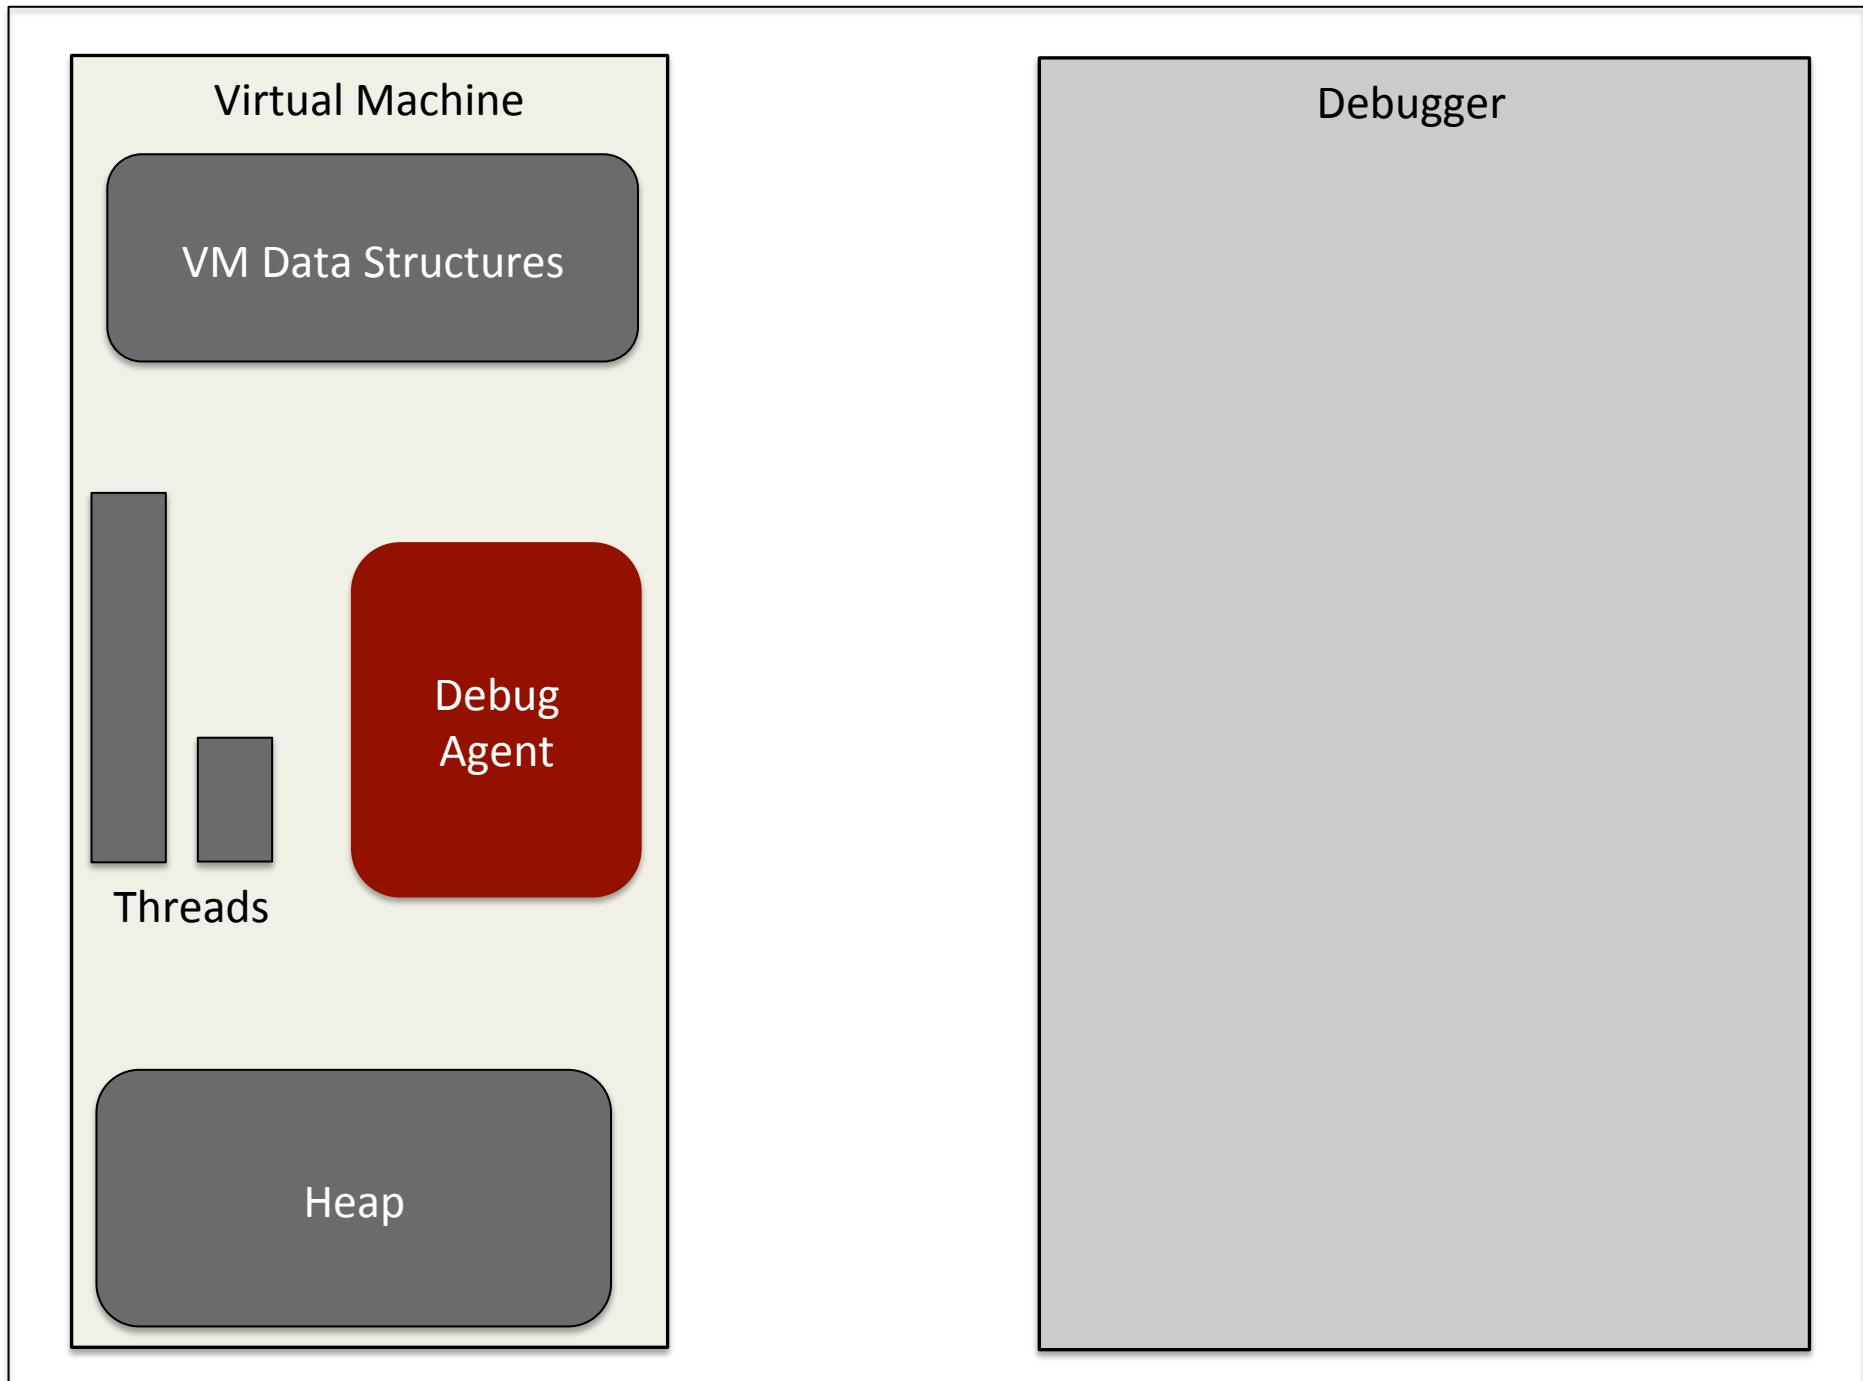
Heap

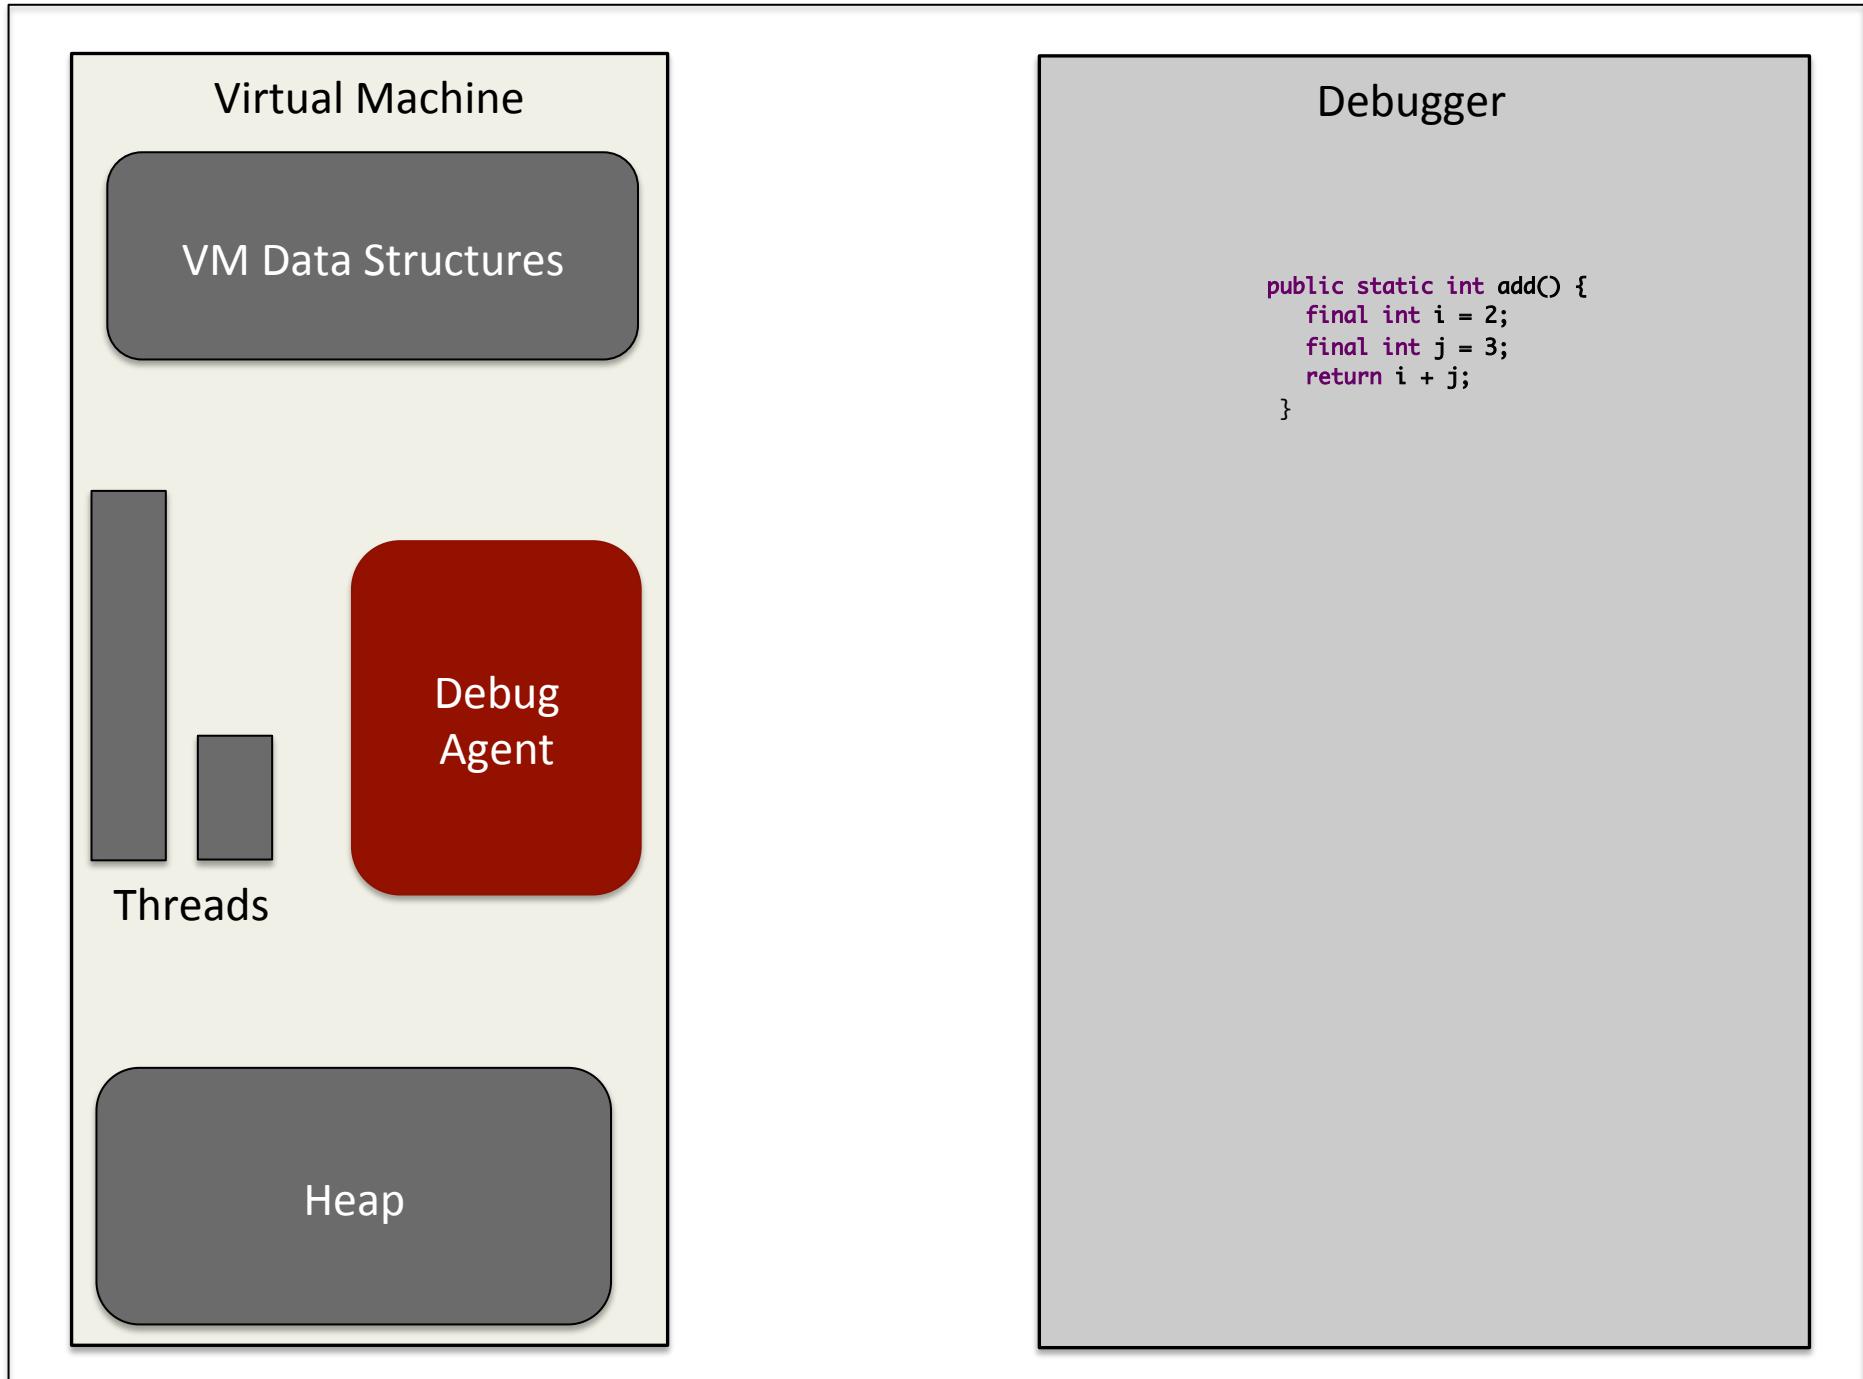
Debugger

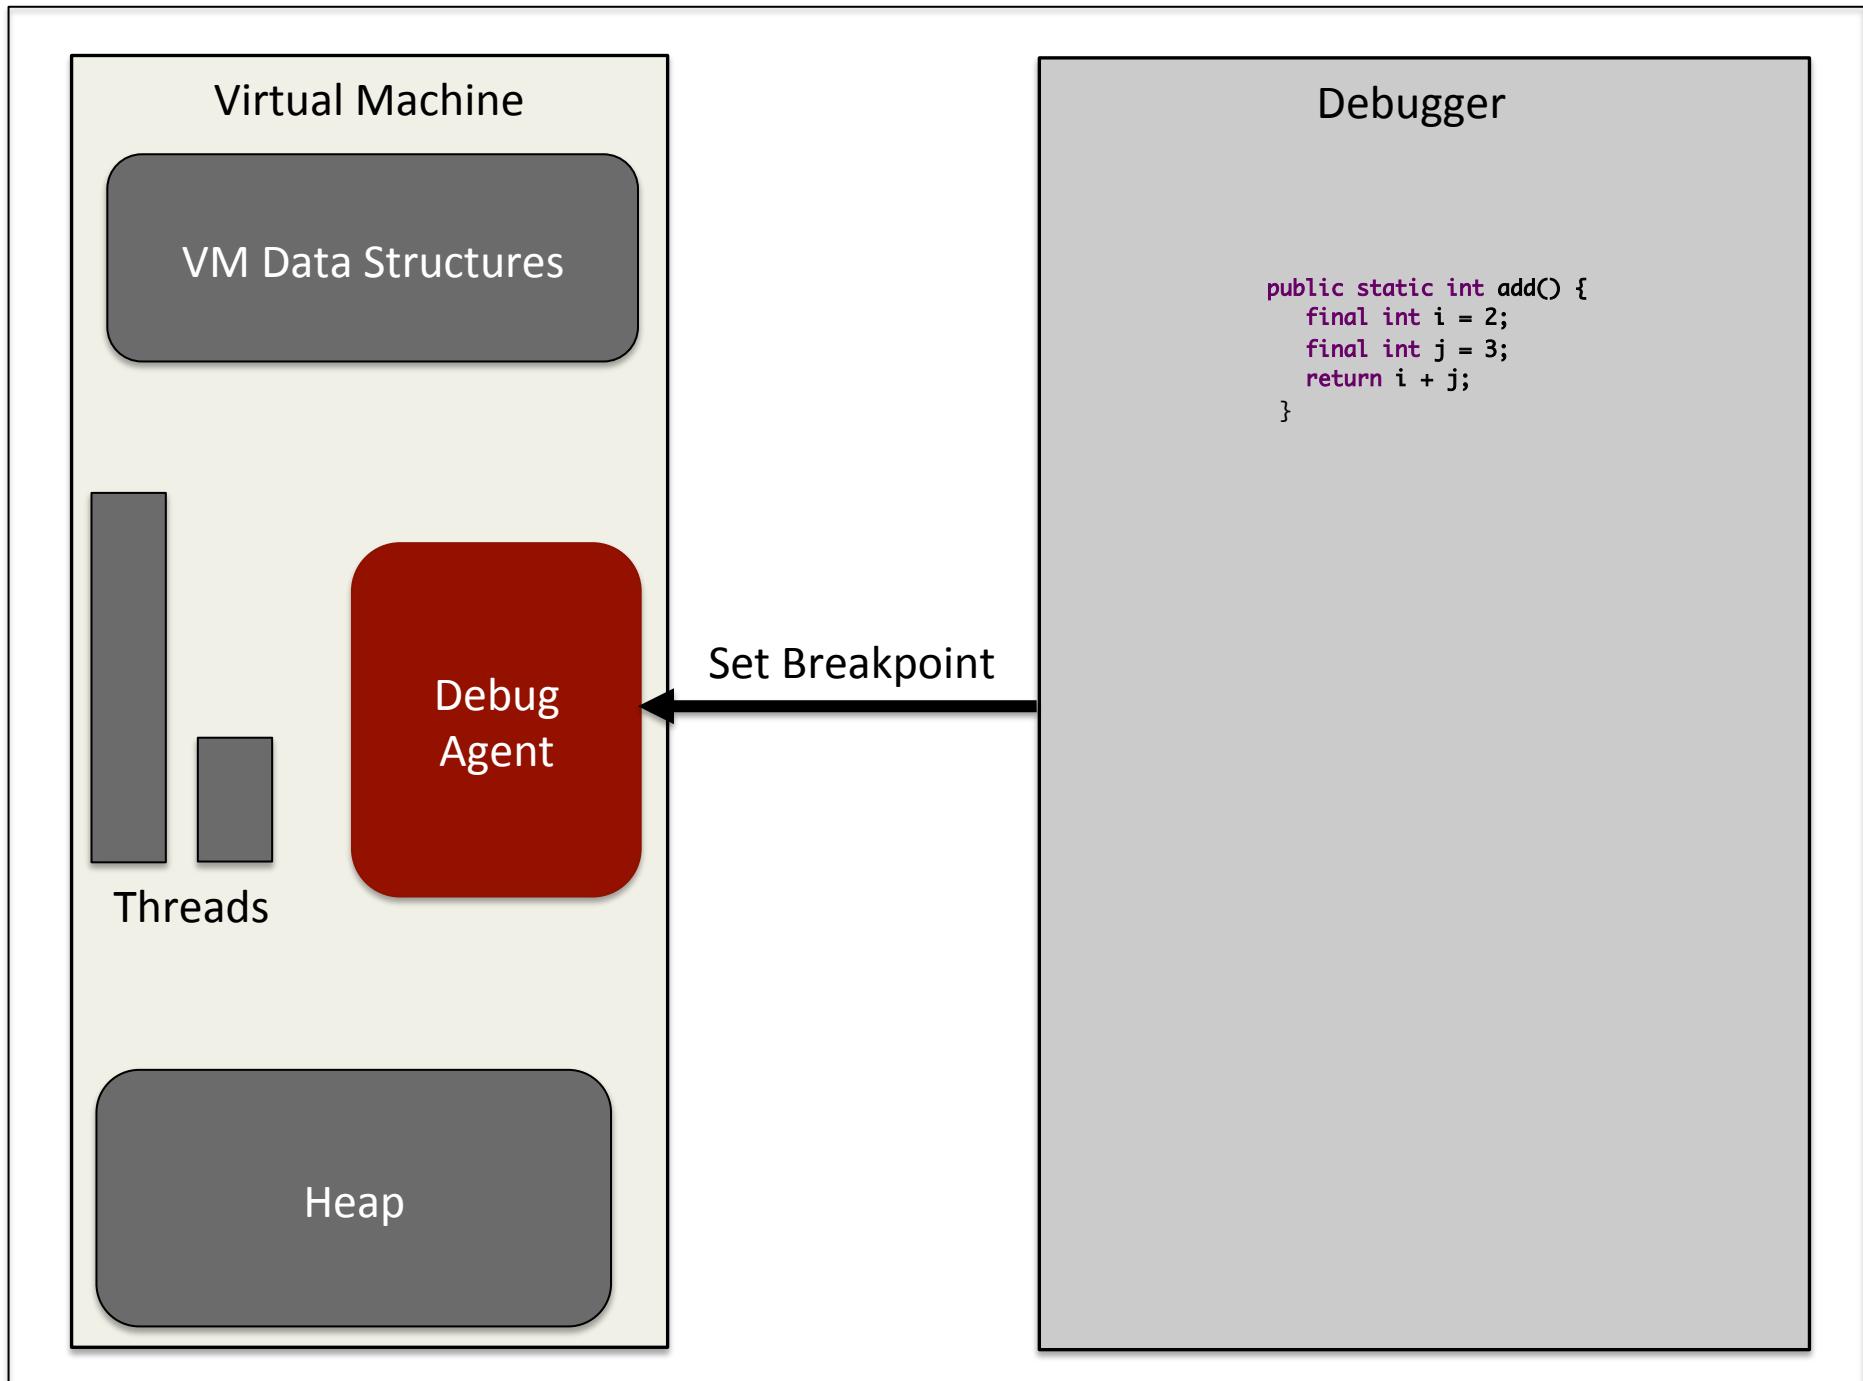


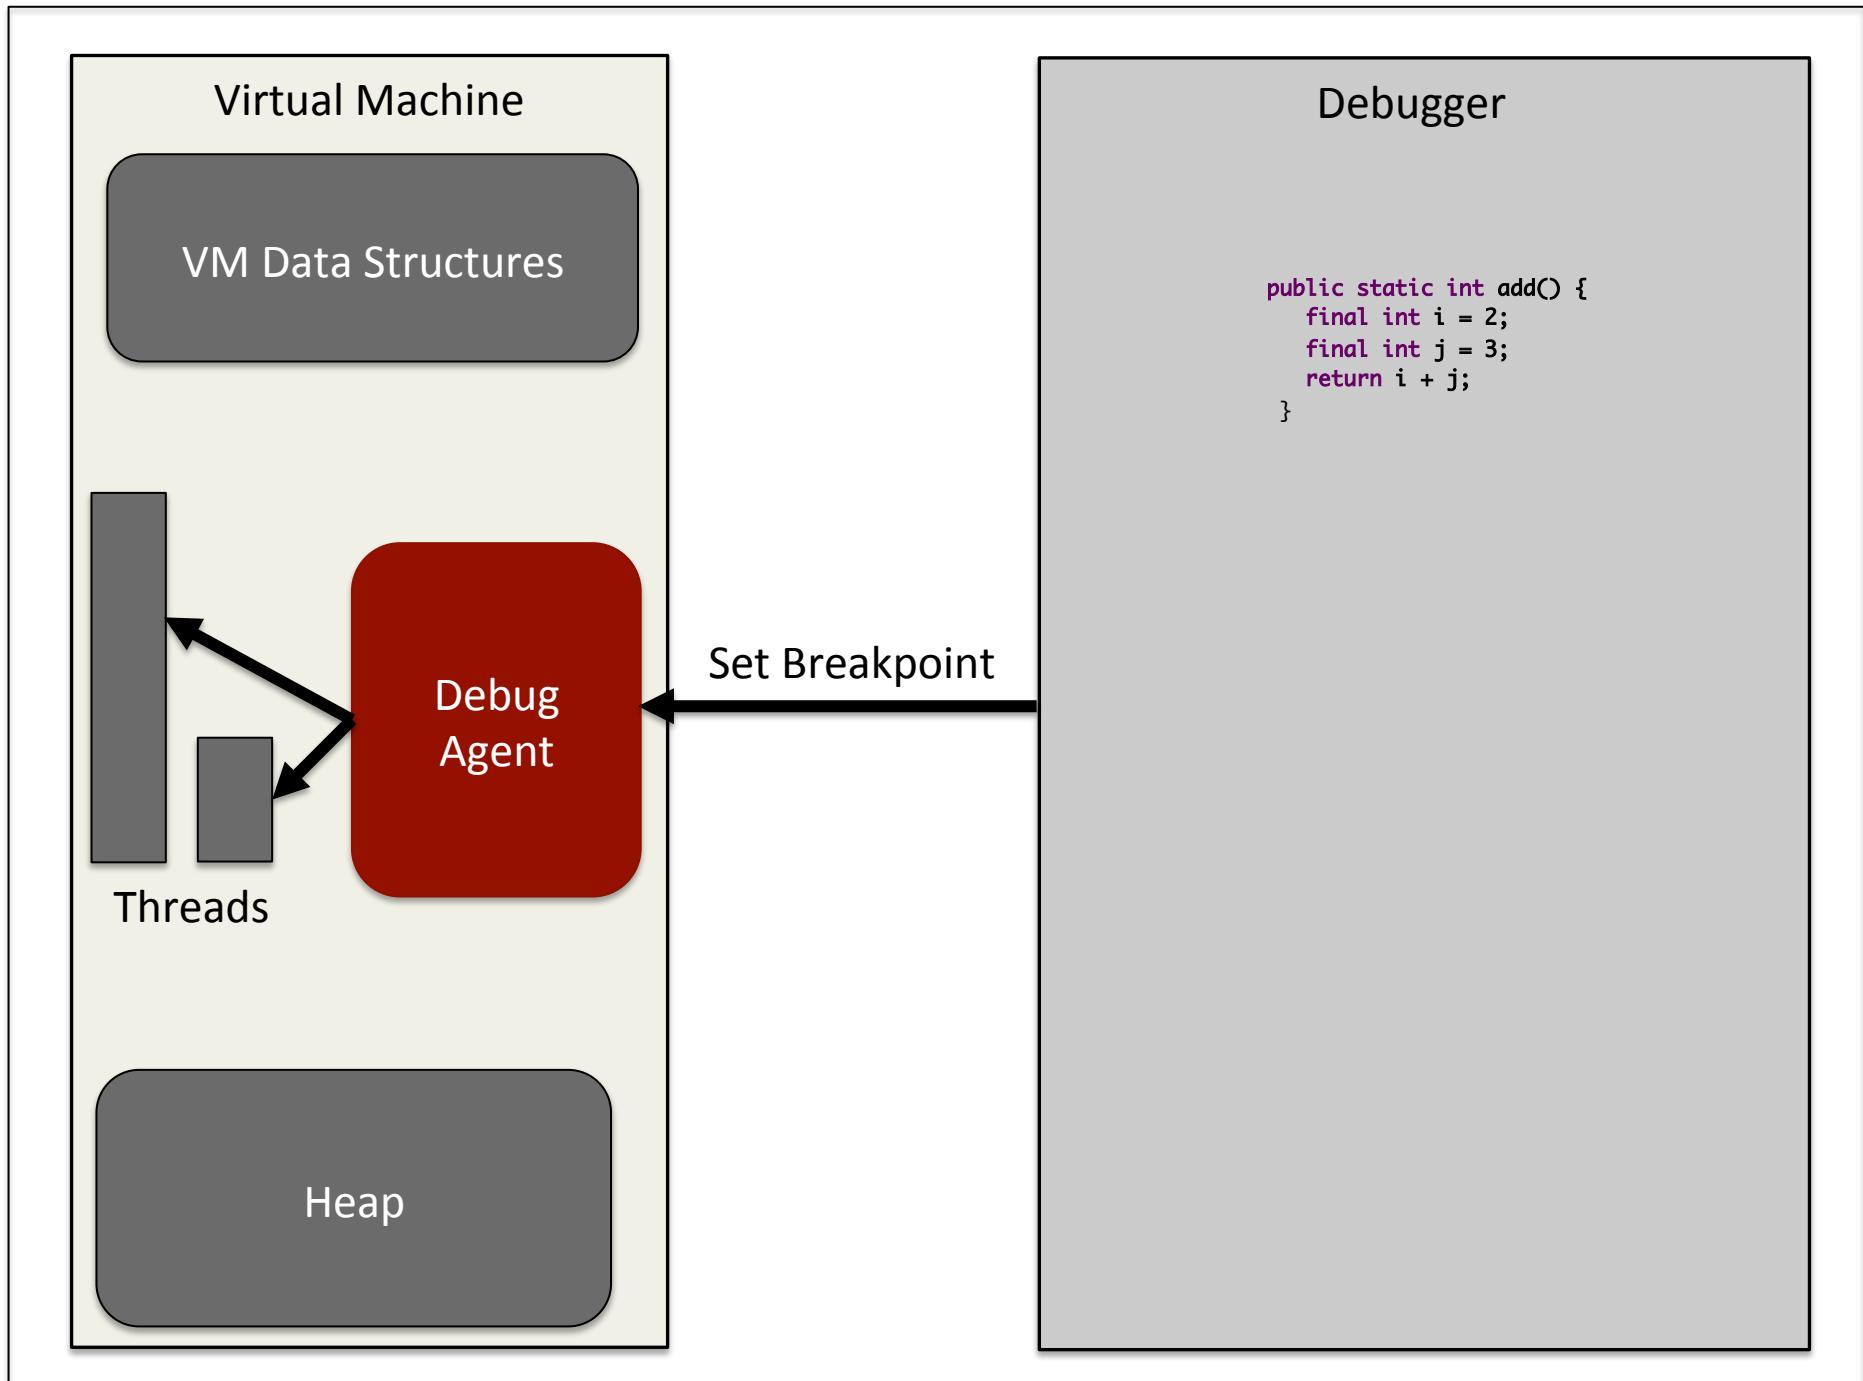
Capabilities

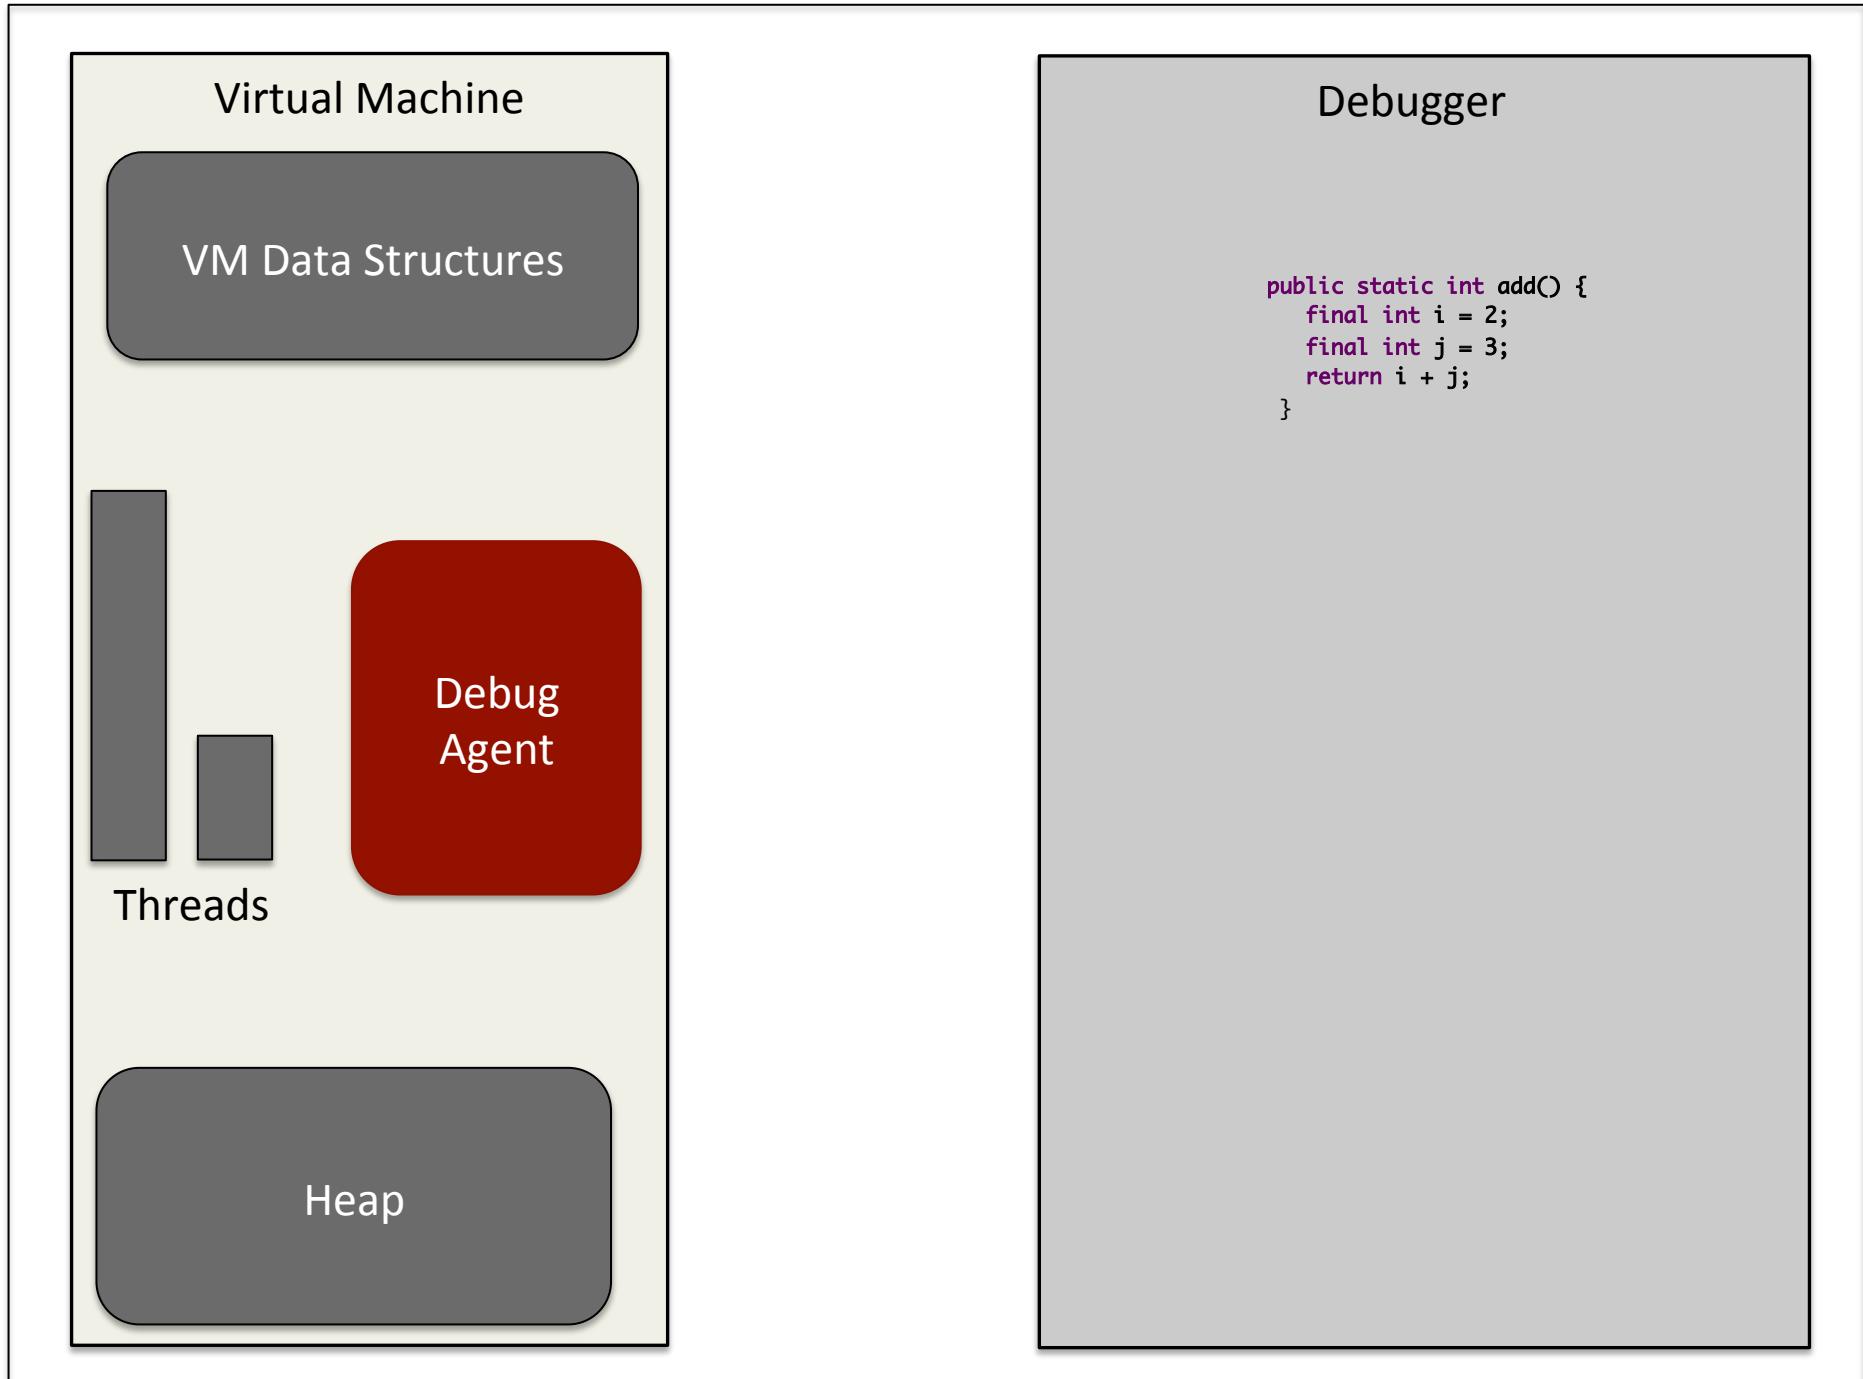
- Select the capabilities for the debugger
 - Receive exception events
 - Ability to tag objects
 - Ability to set breakpoints
 - Ability to signal threads
 - Ability to single-step through code
- Register the capabilities and callbacks at startup

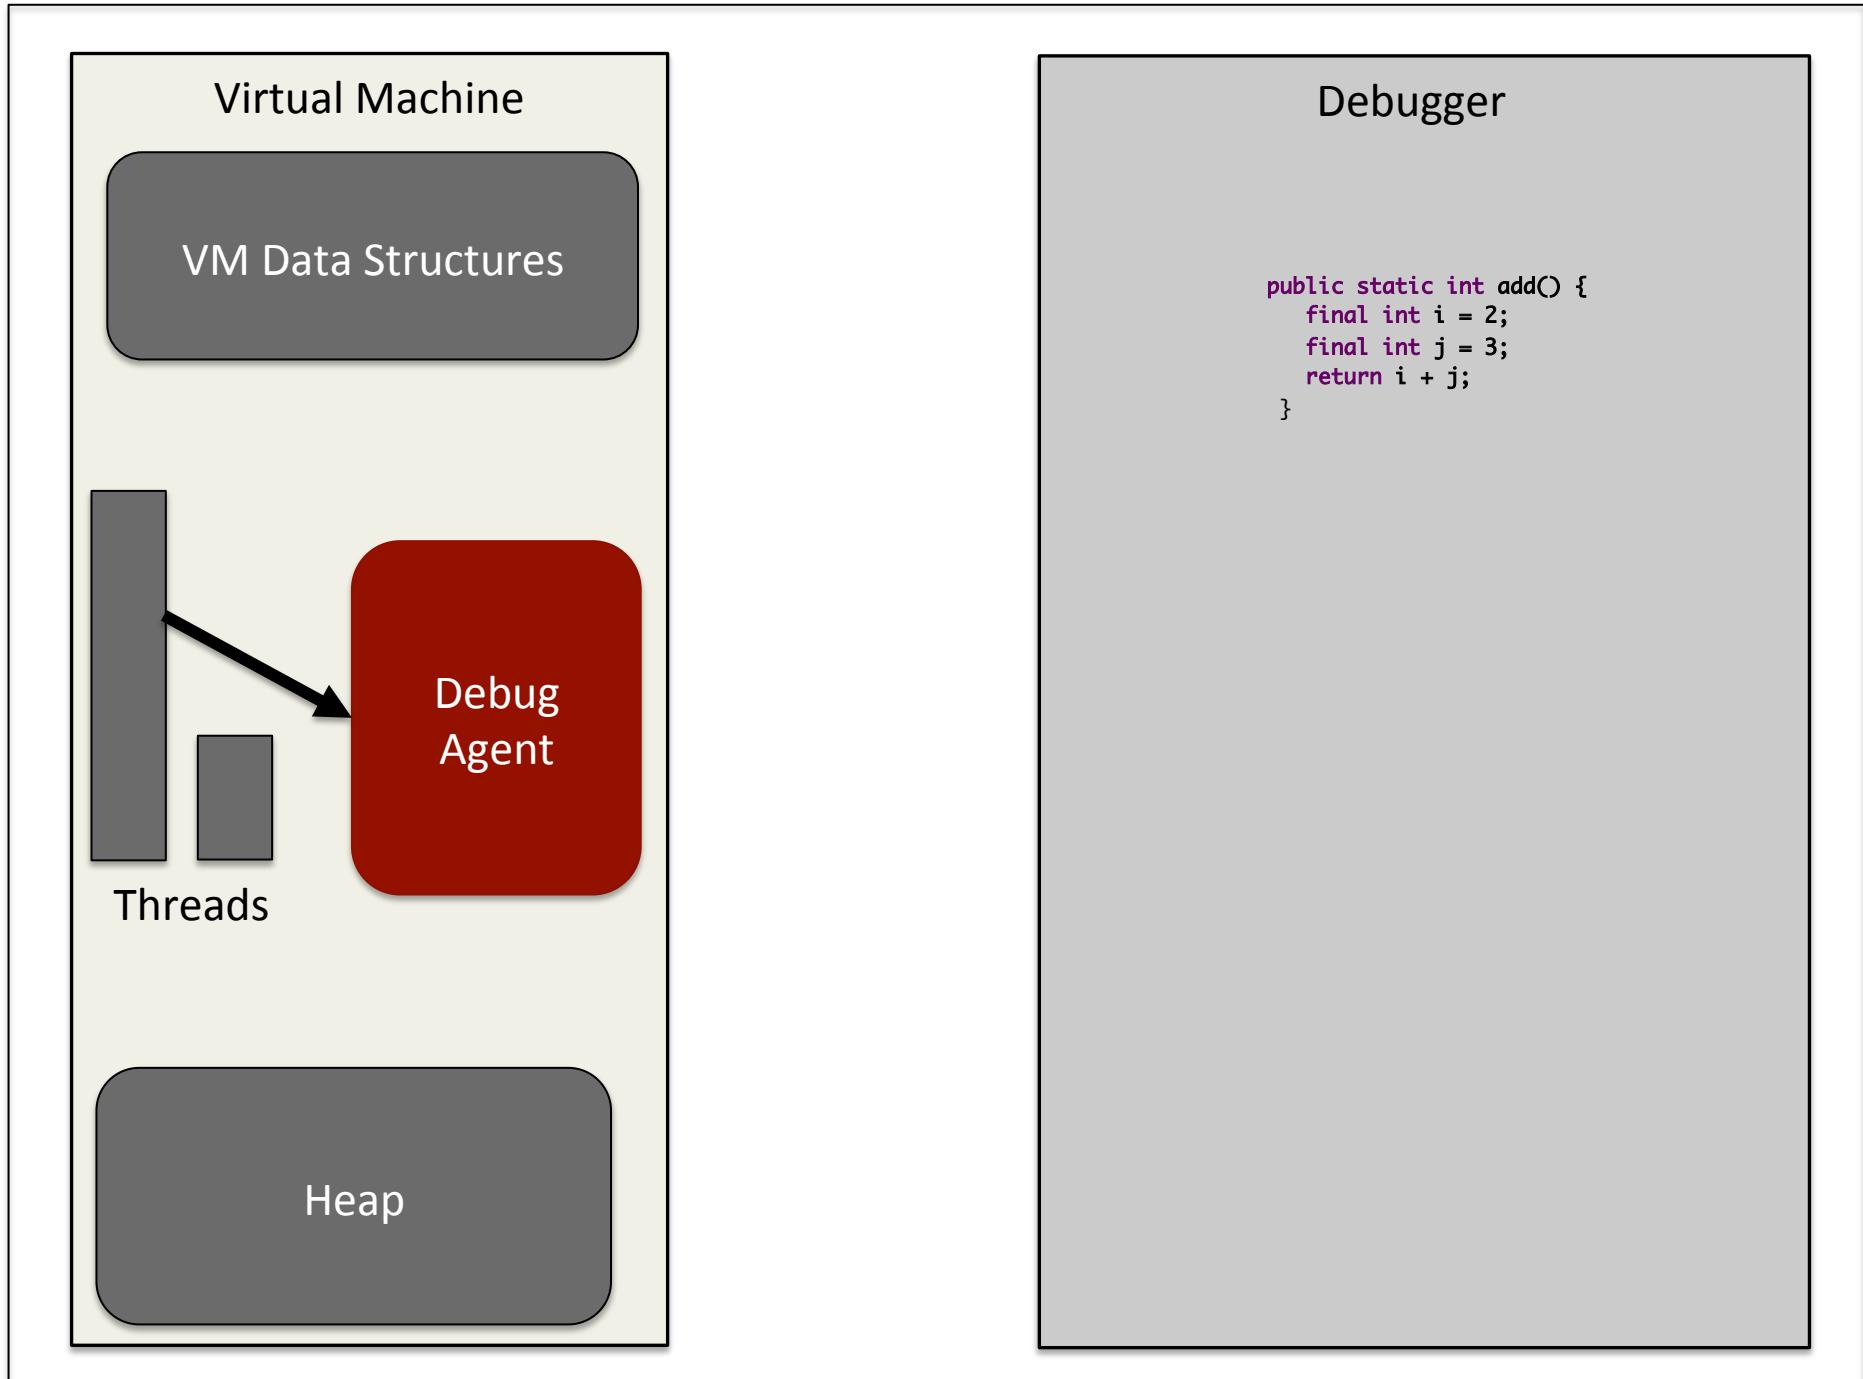


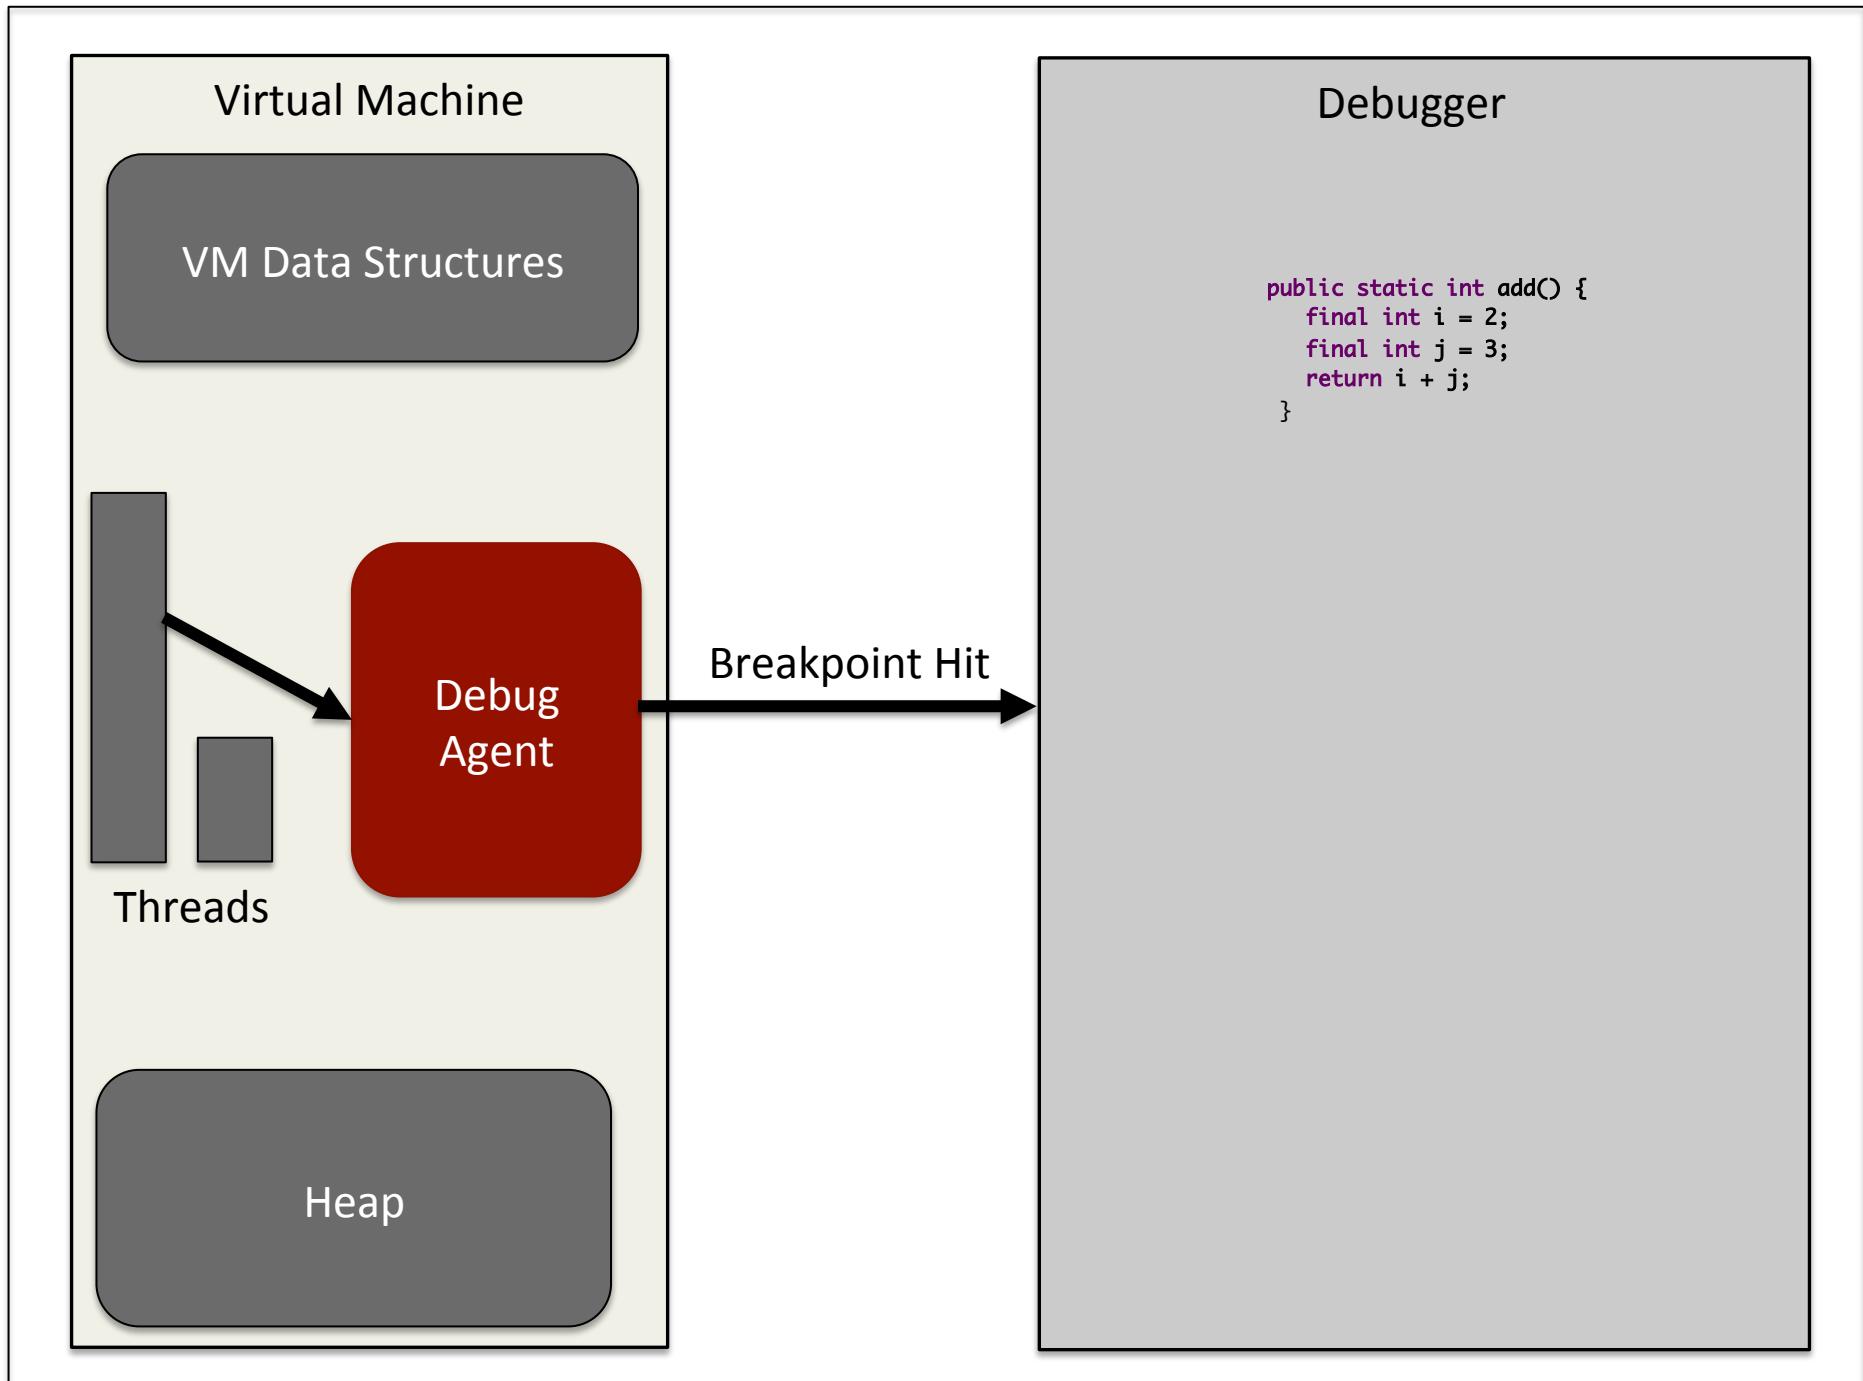






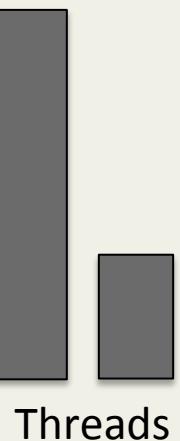






Virtual Machine

VM Data Structures



Debug Agent

Heap

Debugger

```
public static int add() {  
    final int i = 2;  
    final int j = 3;  
    return i + j;  
}
```

4

Program Counter



2 3

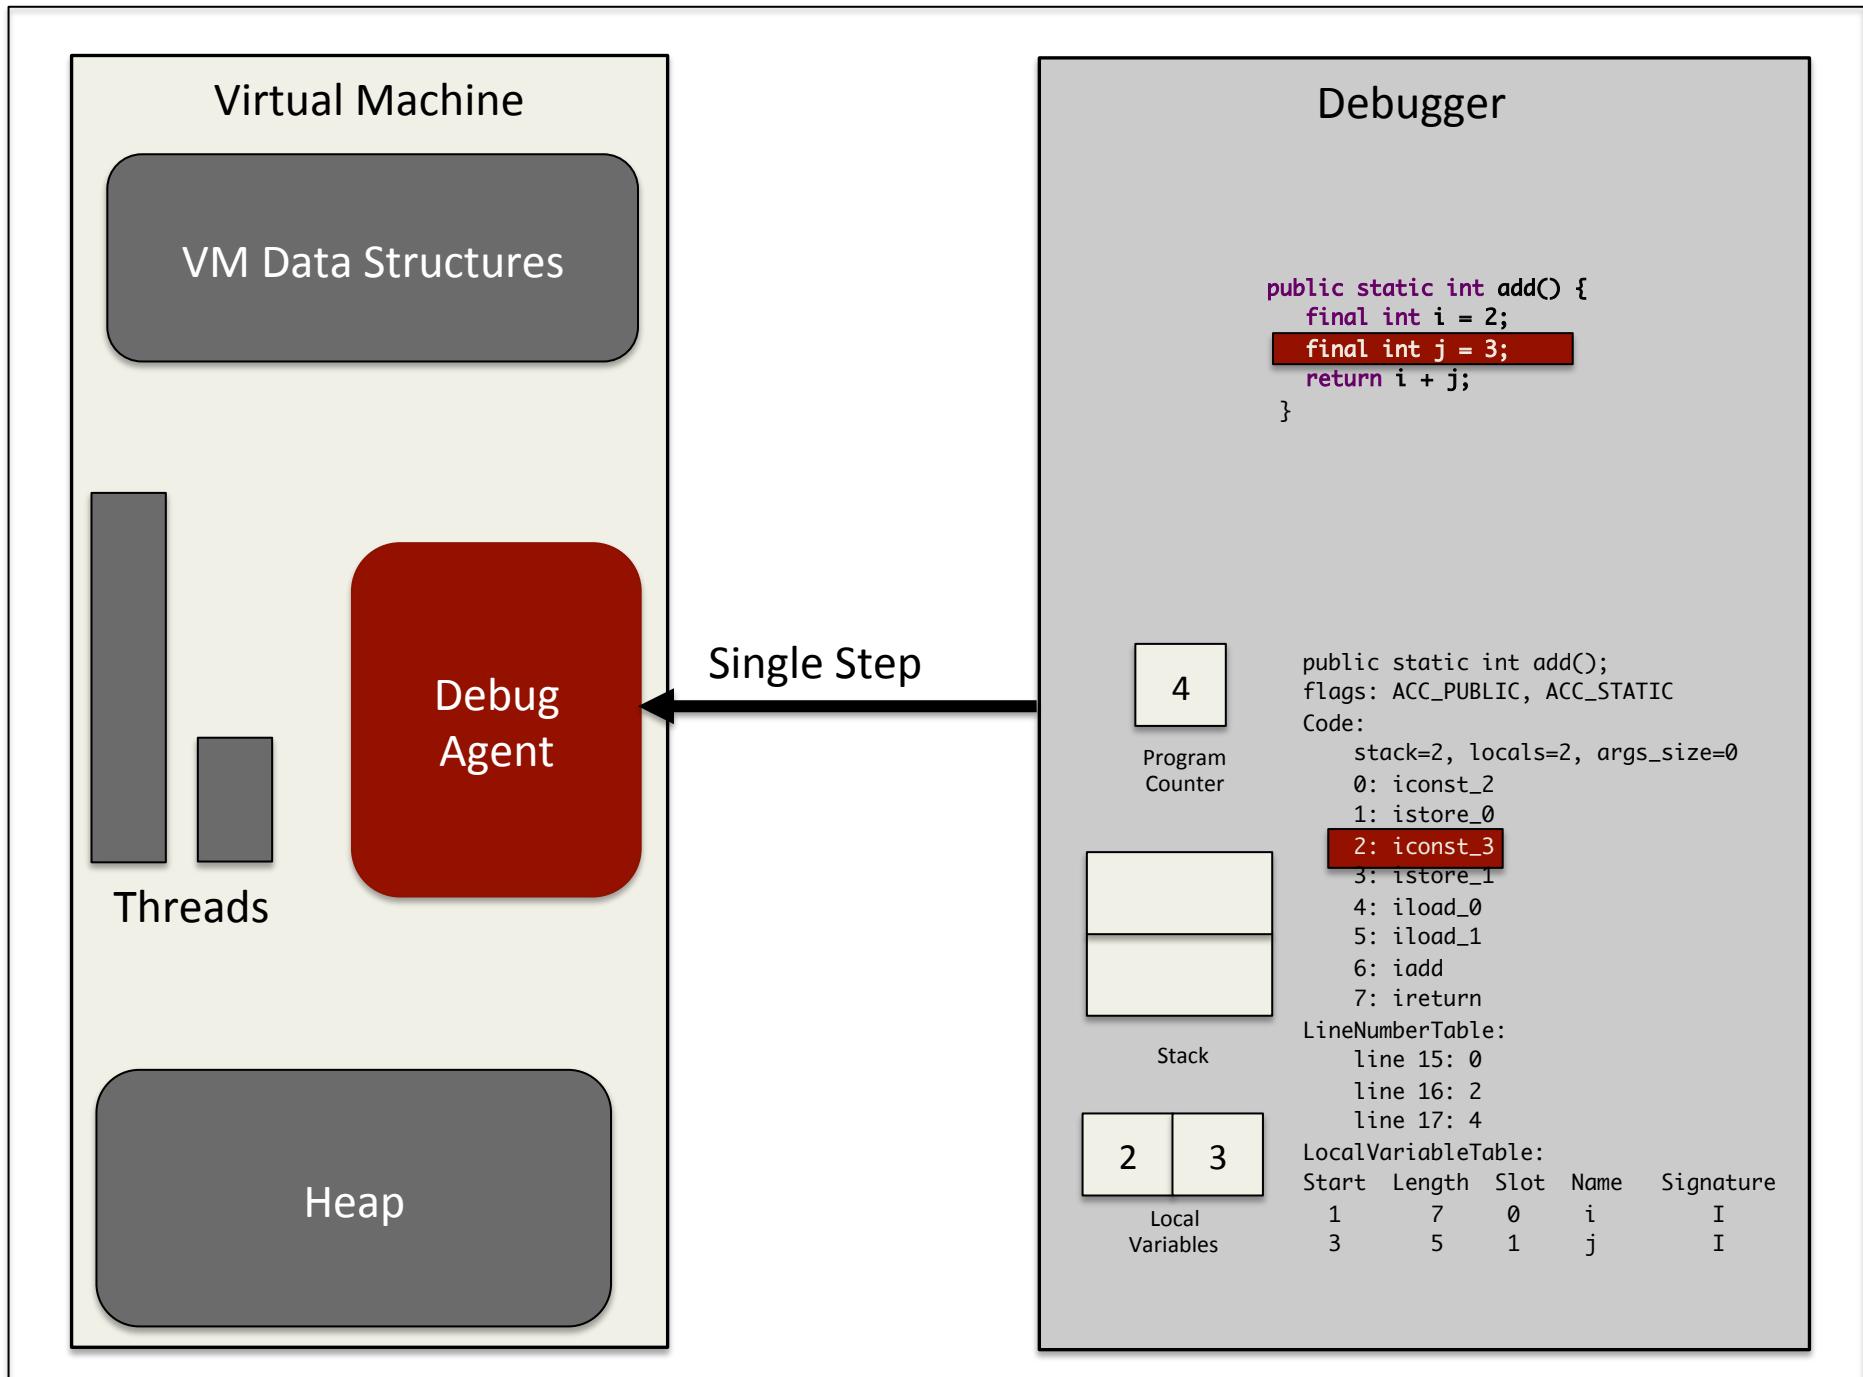
Local Variables

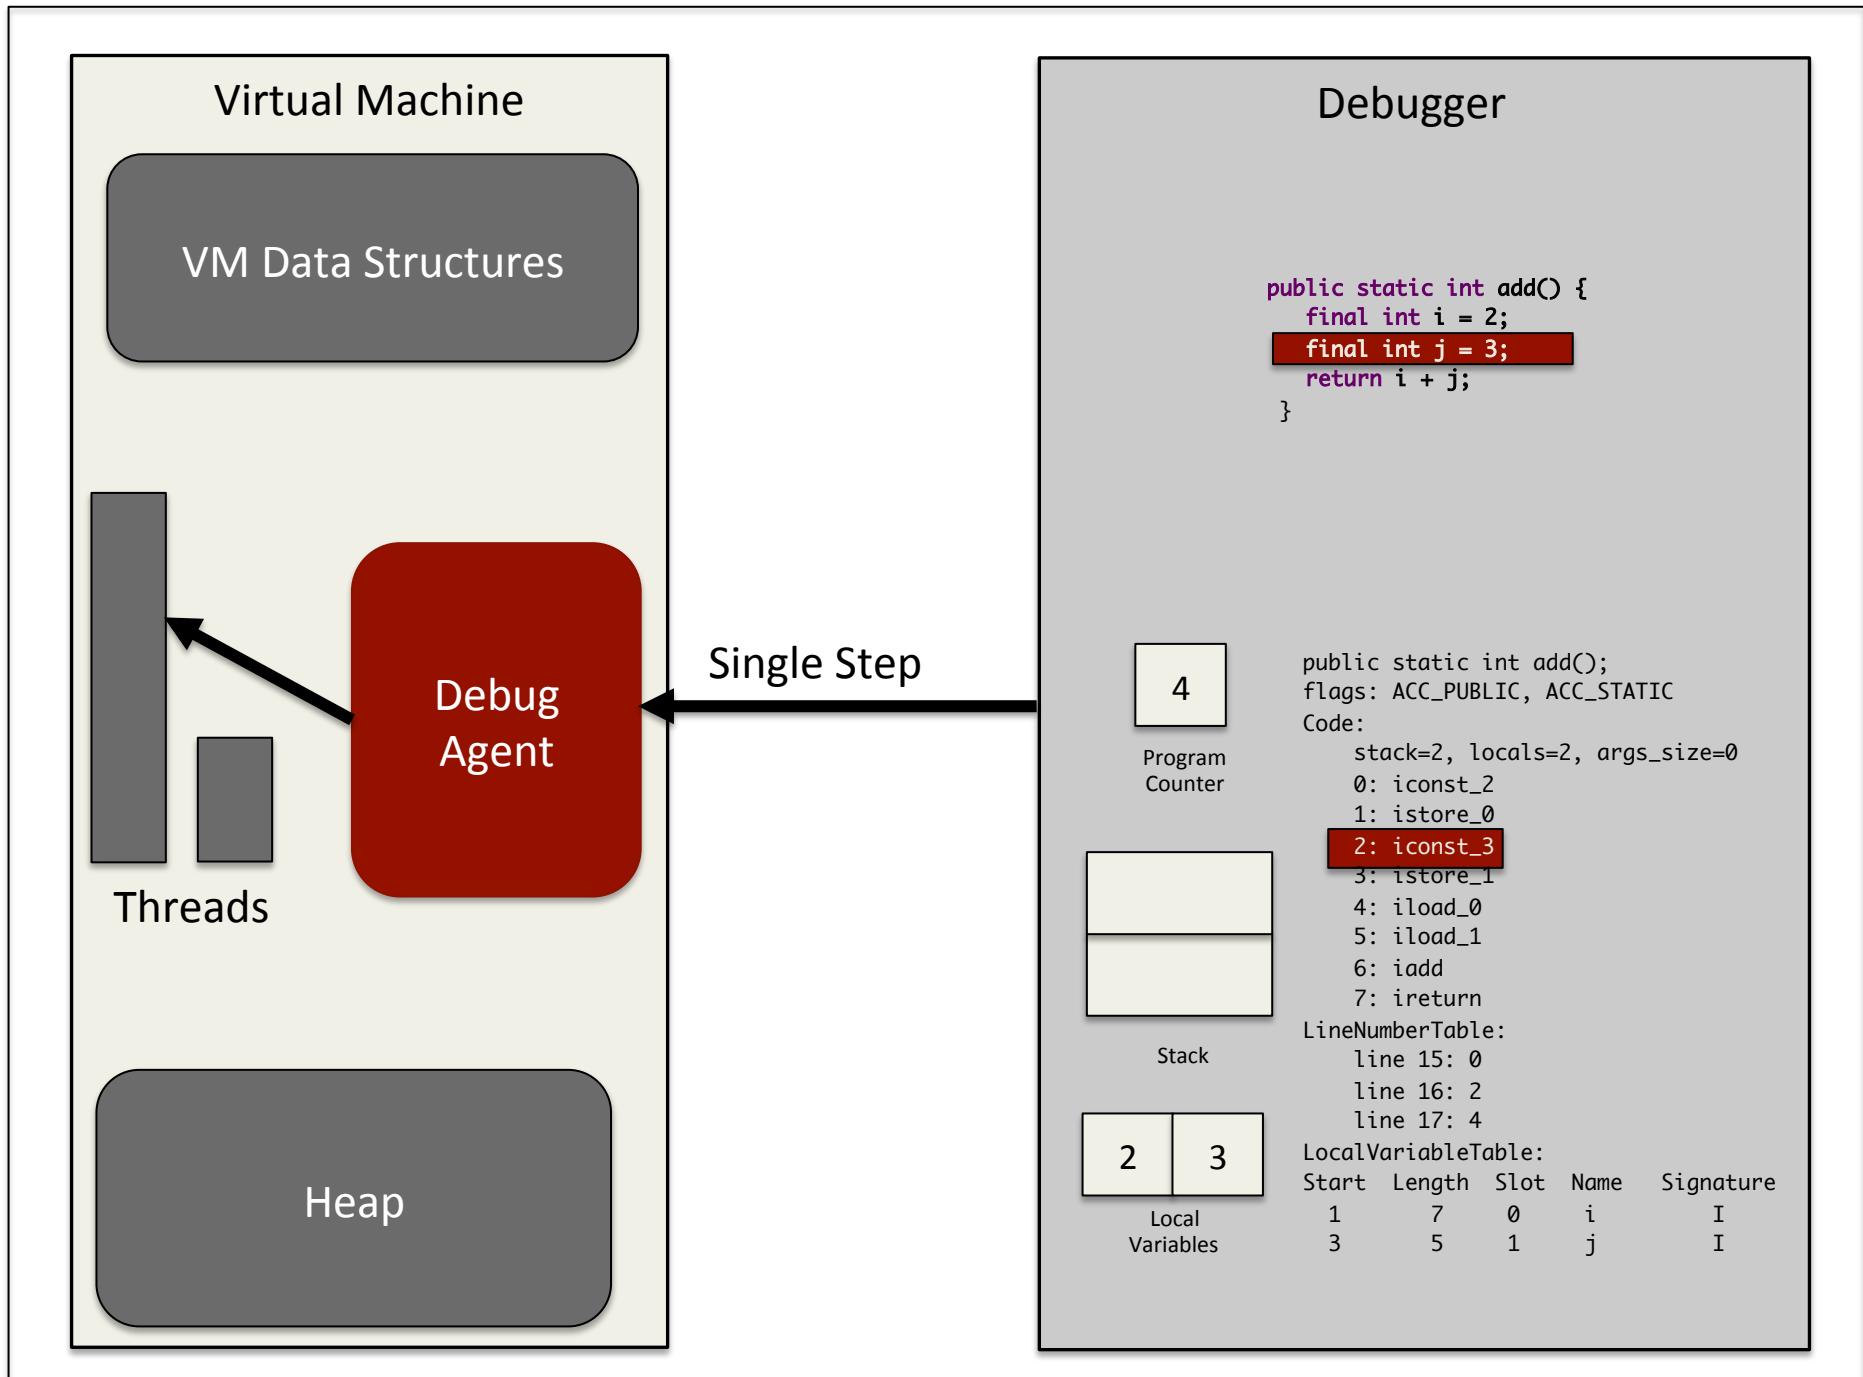
public static int add();
flags: ACC_PUBLIC, ACC_STATIC
Code:

```
stack=2, locals=2, args_size=0  
0:  iconst_2  
1:  istore_0  
2:  iconst_3  
3:  istore_1  
4:  iload_0  
5:  iload_1  
6:  iadd  
7:  ireturn
```

LineNumberTable:
line 15: 0
line 16: 2
line 17: 4

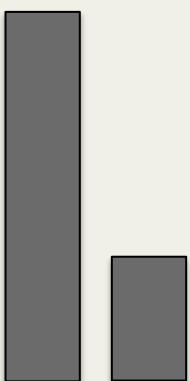
LocalVariableTable:
Start Length Slot Name Signature
1 7 0 i I
3 5 1 j I





Virtual Machine

VM Data Structures



Threads

Debug Agent

Heap

Debugger

```
public static int add() {  
    final int i = 2;  
    final int j = 3;  
    return i + j;  
}
```

4

Program Counter



Stack

2 3

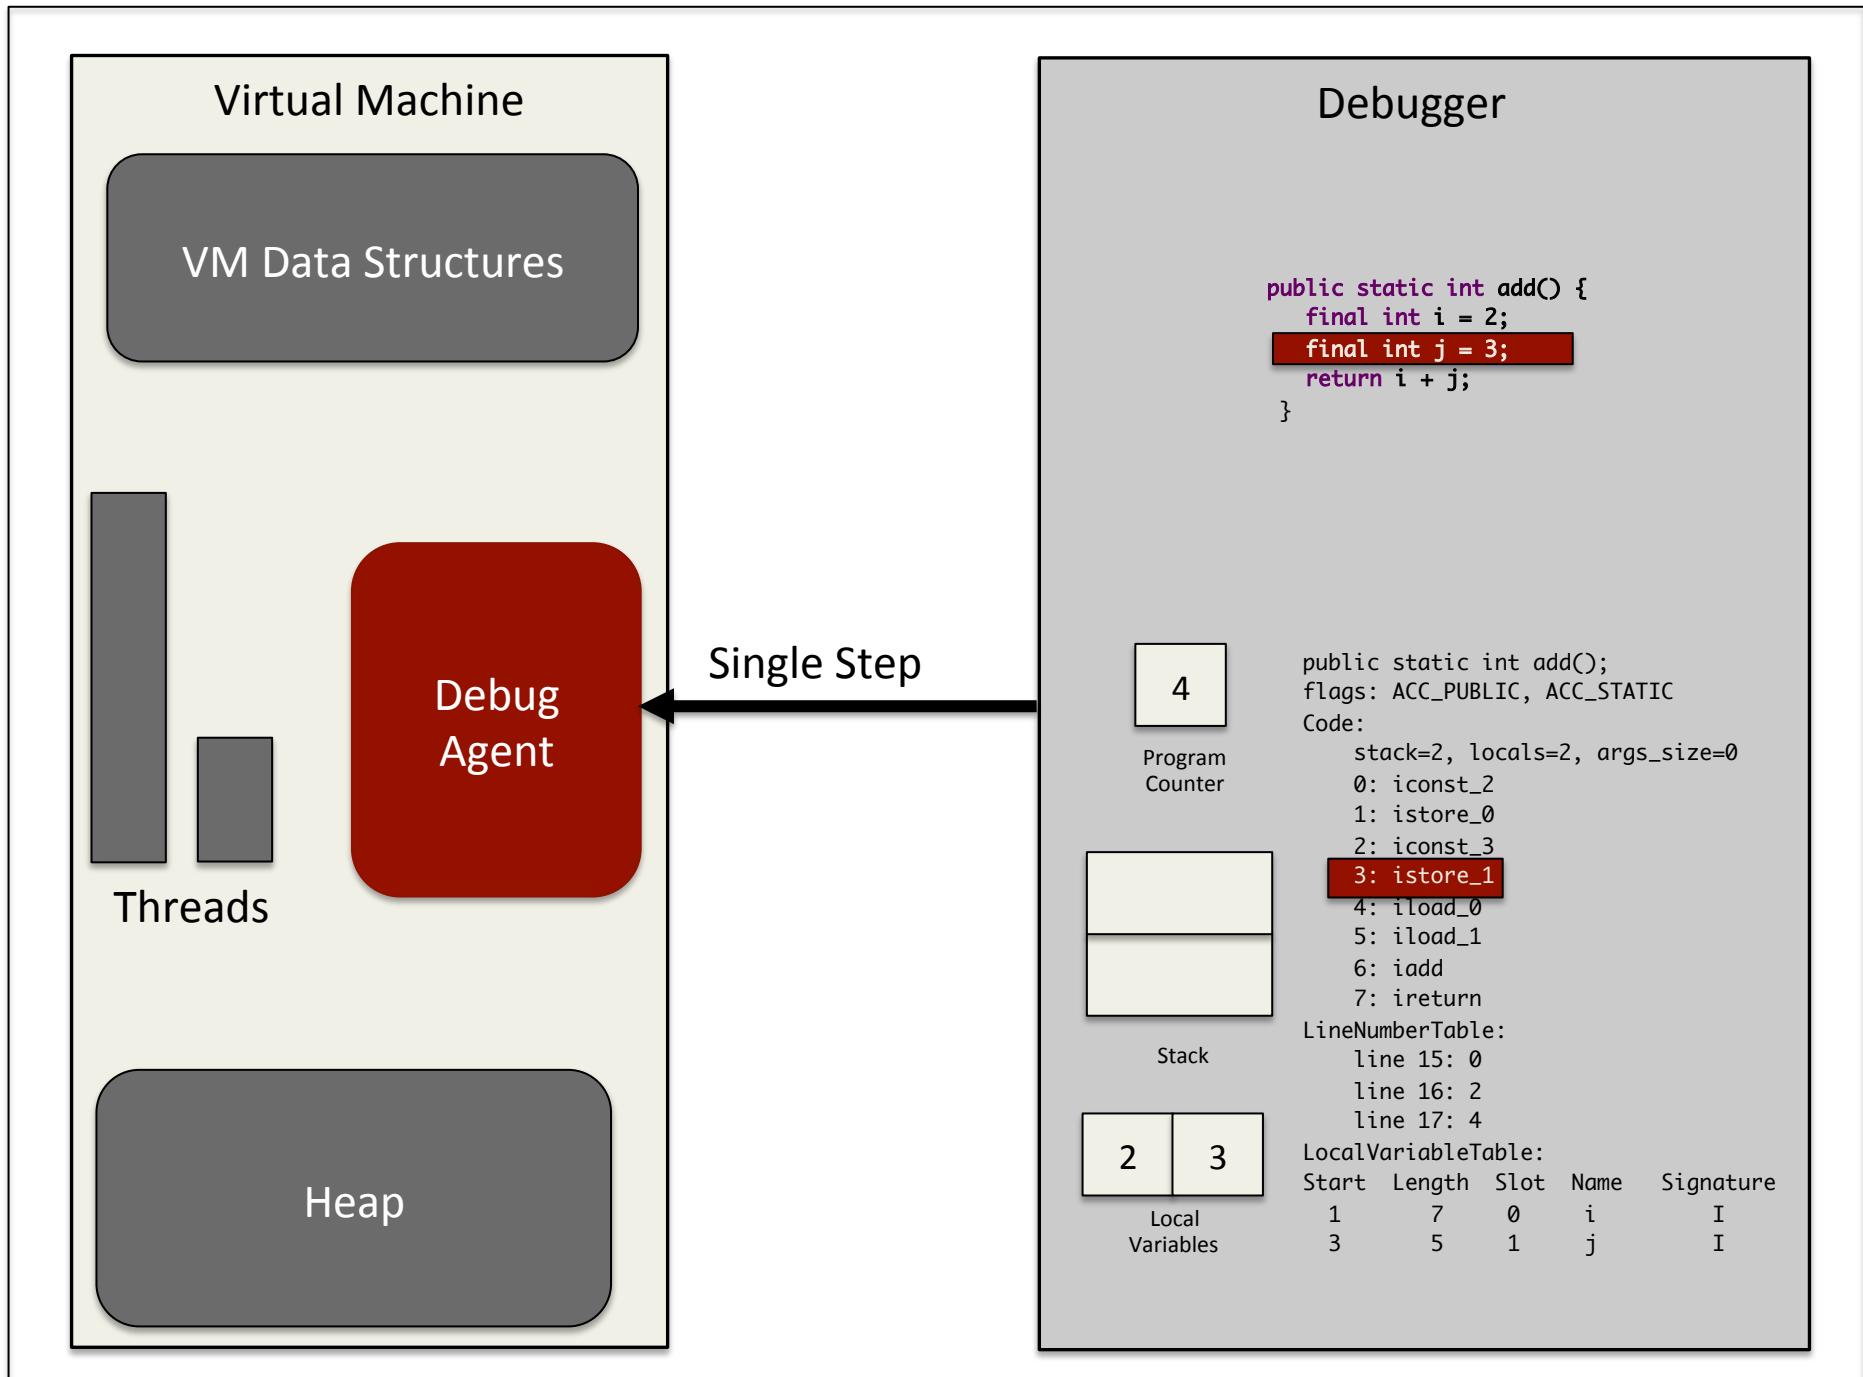
Local Variables

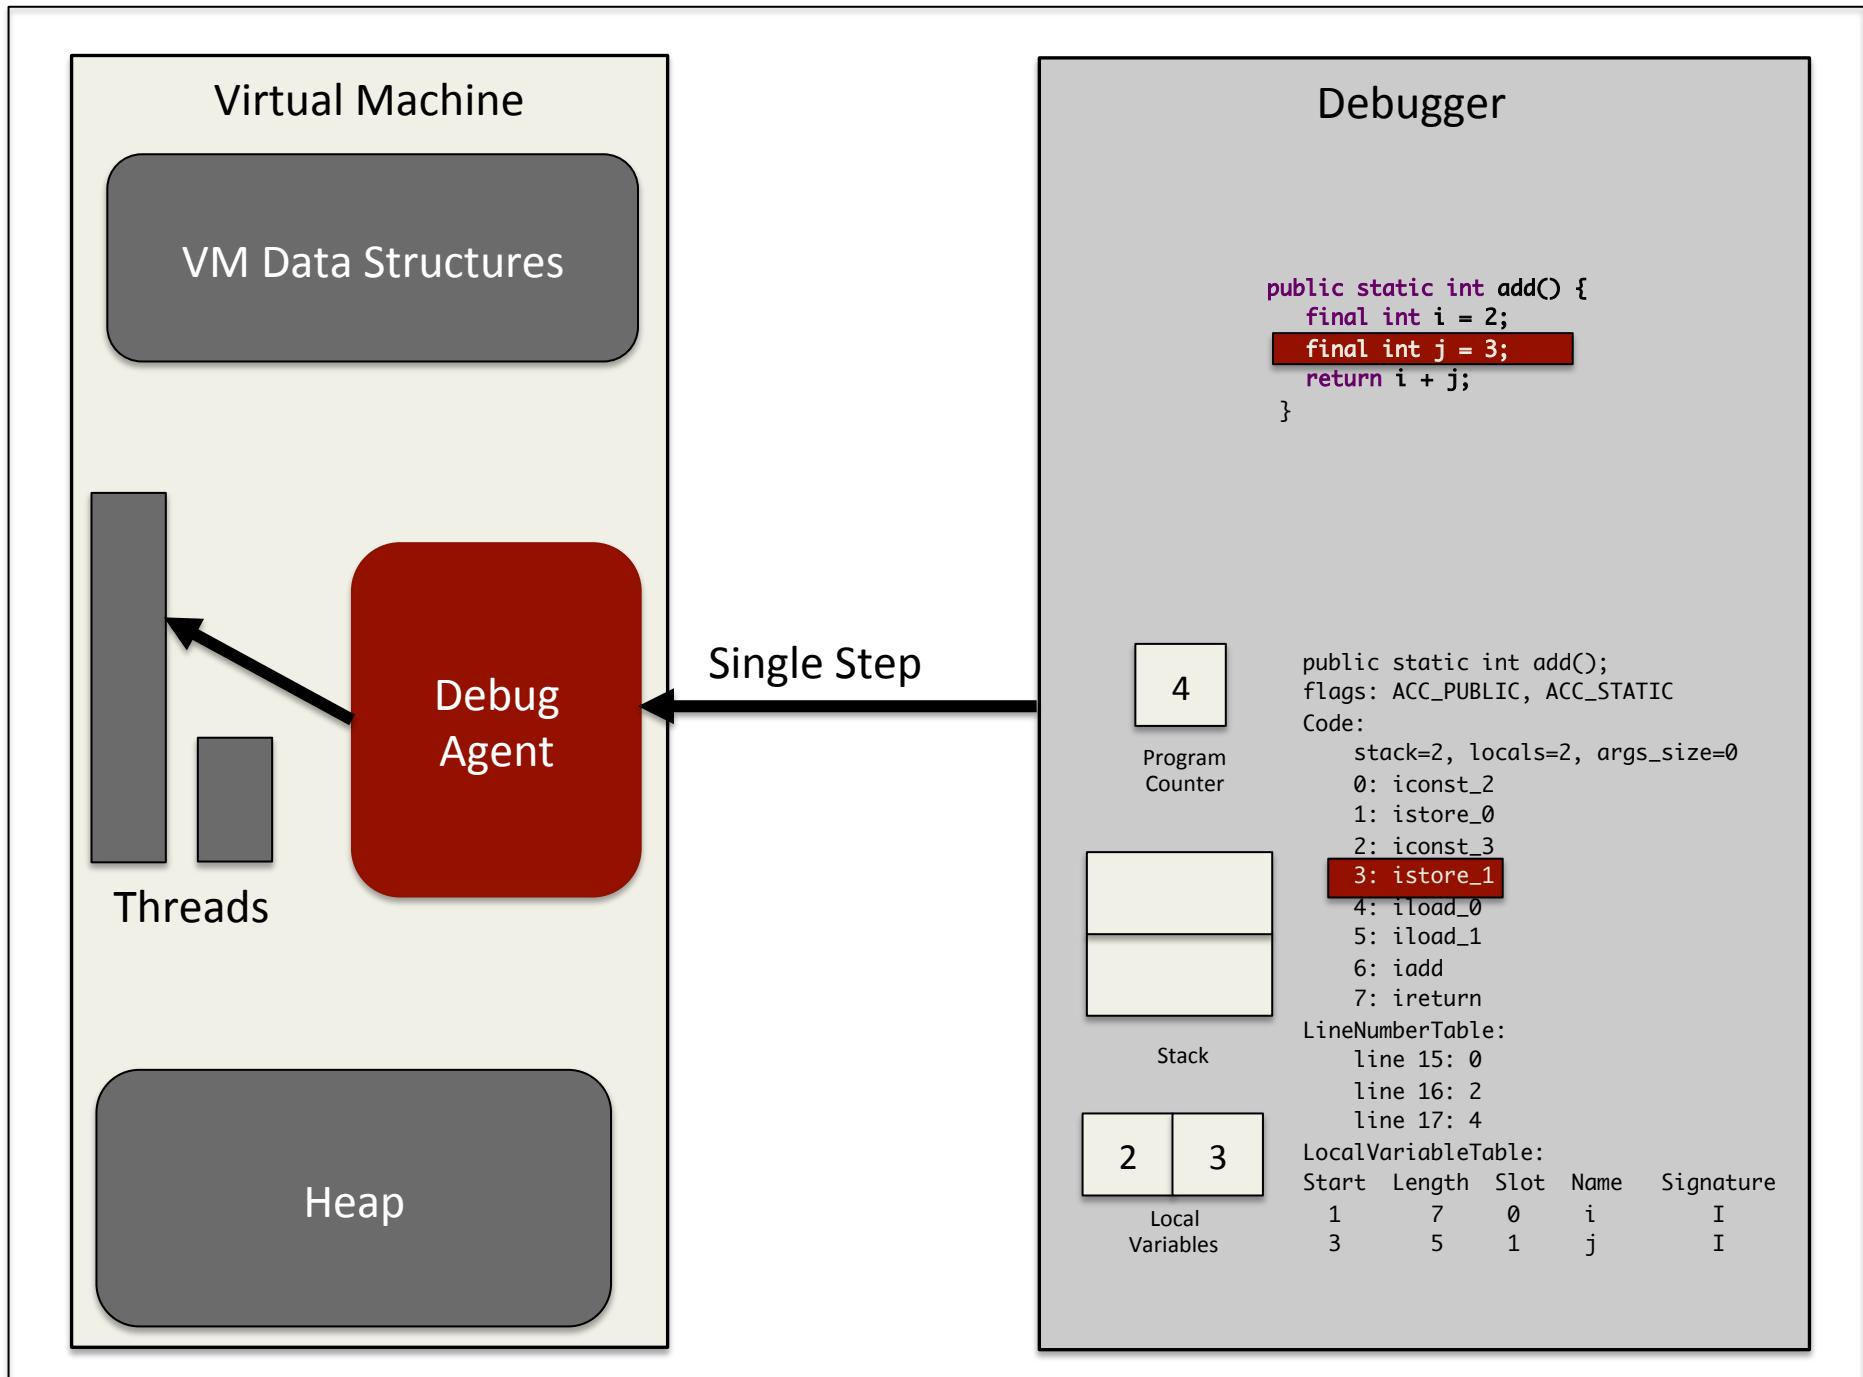
public static int add();
flags: ACC_PUBLIC, ACC_STATIC
Code:

```
stack=2, locals=2, args_size=0  
0:  iconst_2  
1:  istore_0  
2:  iconst_3  
3:  istore_1  
4:  iload_0  
5:  iload_1  
6:  iadd  
7:  ireturn
```

LineNumberTable:
line 15: 0
line 16: 2
line 17: 4

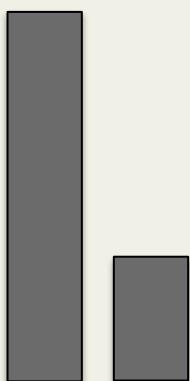
LocalVariableTable:
Start Length Slot Name Signature
1 7 0 i I
3 5 1 j I





Virtual Machine

VM Data Structures



Threads

Debug Agent

Heap

Debugger

```
public static int add() {  
    final int i = 2;  
    final int j = 3;  
    return i + j;  
}
```

4

Program Counter



Stack

2 3

Local Variables

public static int add();
flags: ACC_PUBLIC, ACC_STATIC
Code:

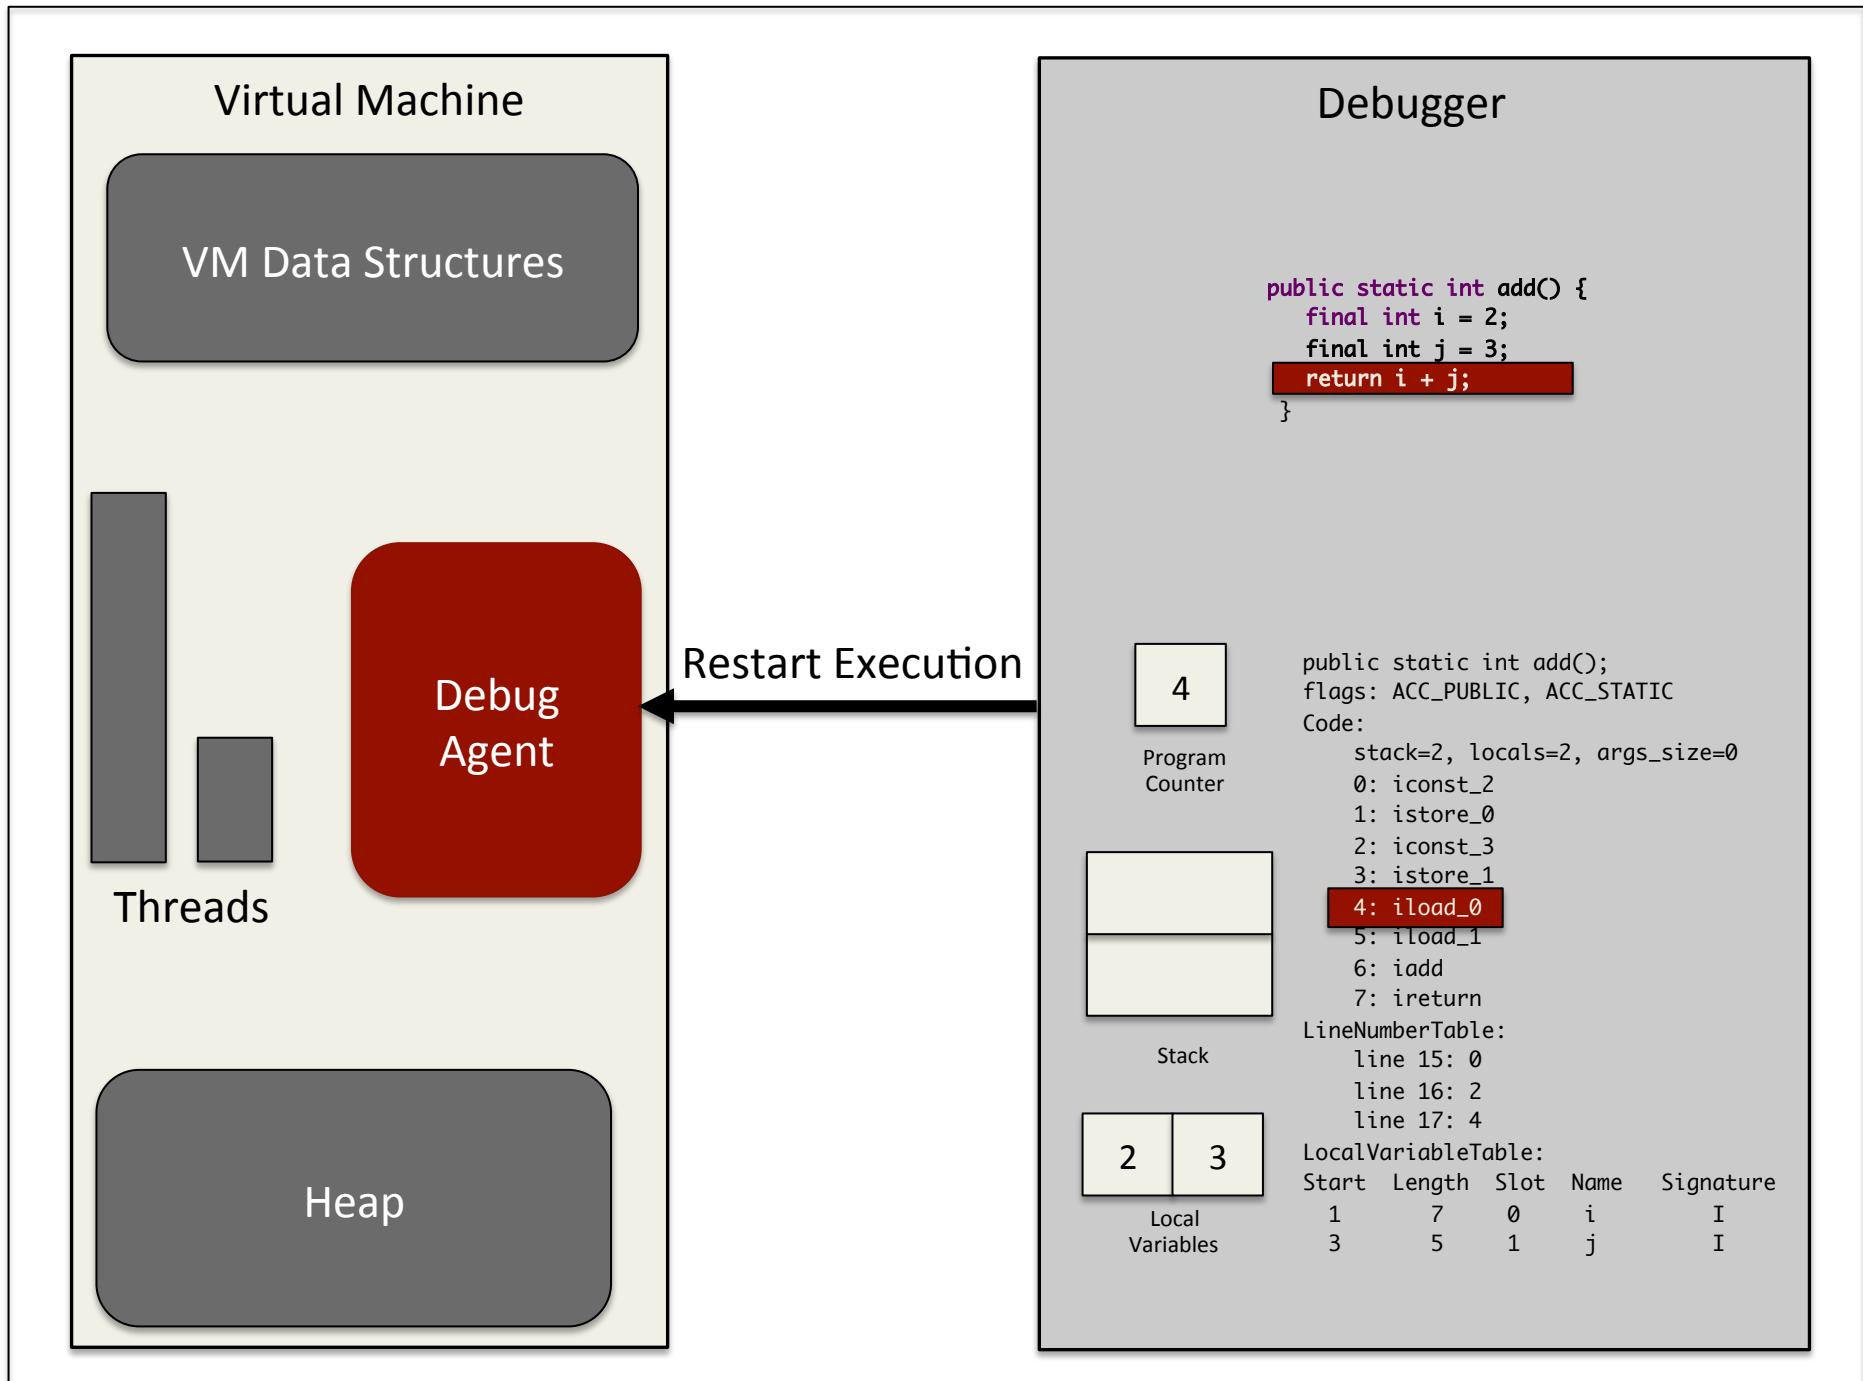
```
stack=2, locals=2, args_size=0  
0:  iconst_2  
1:  istore_0  
2:  iconst_3  
3:  istore_1  
4:  iload_0  
5:  iload_1  
6:  iadd  
7:  ireturn
```

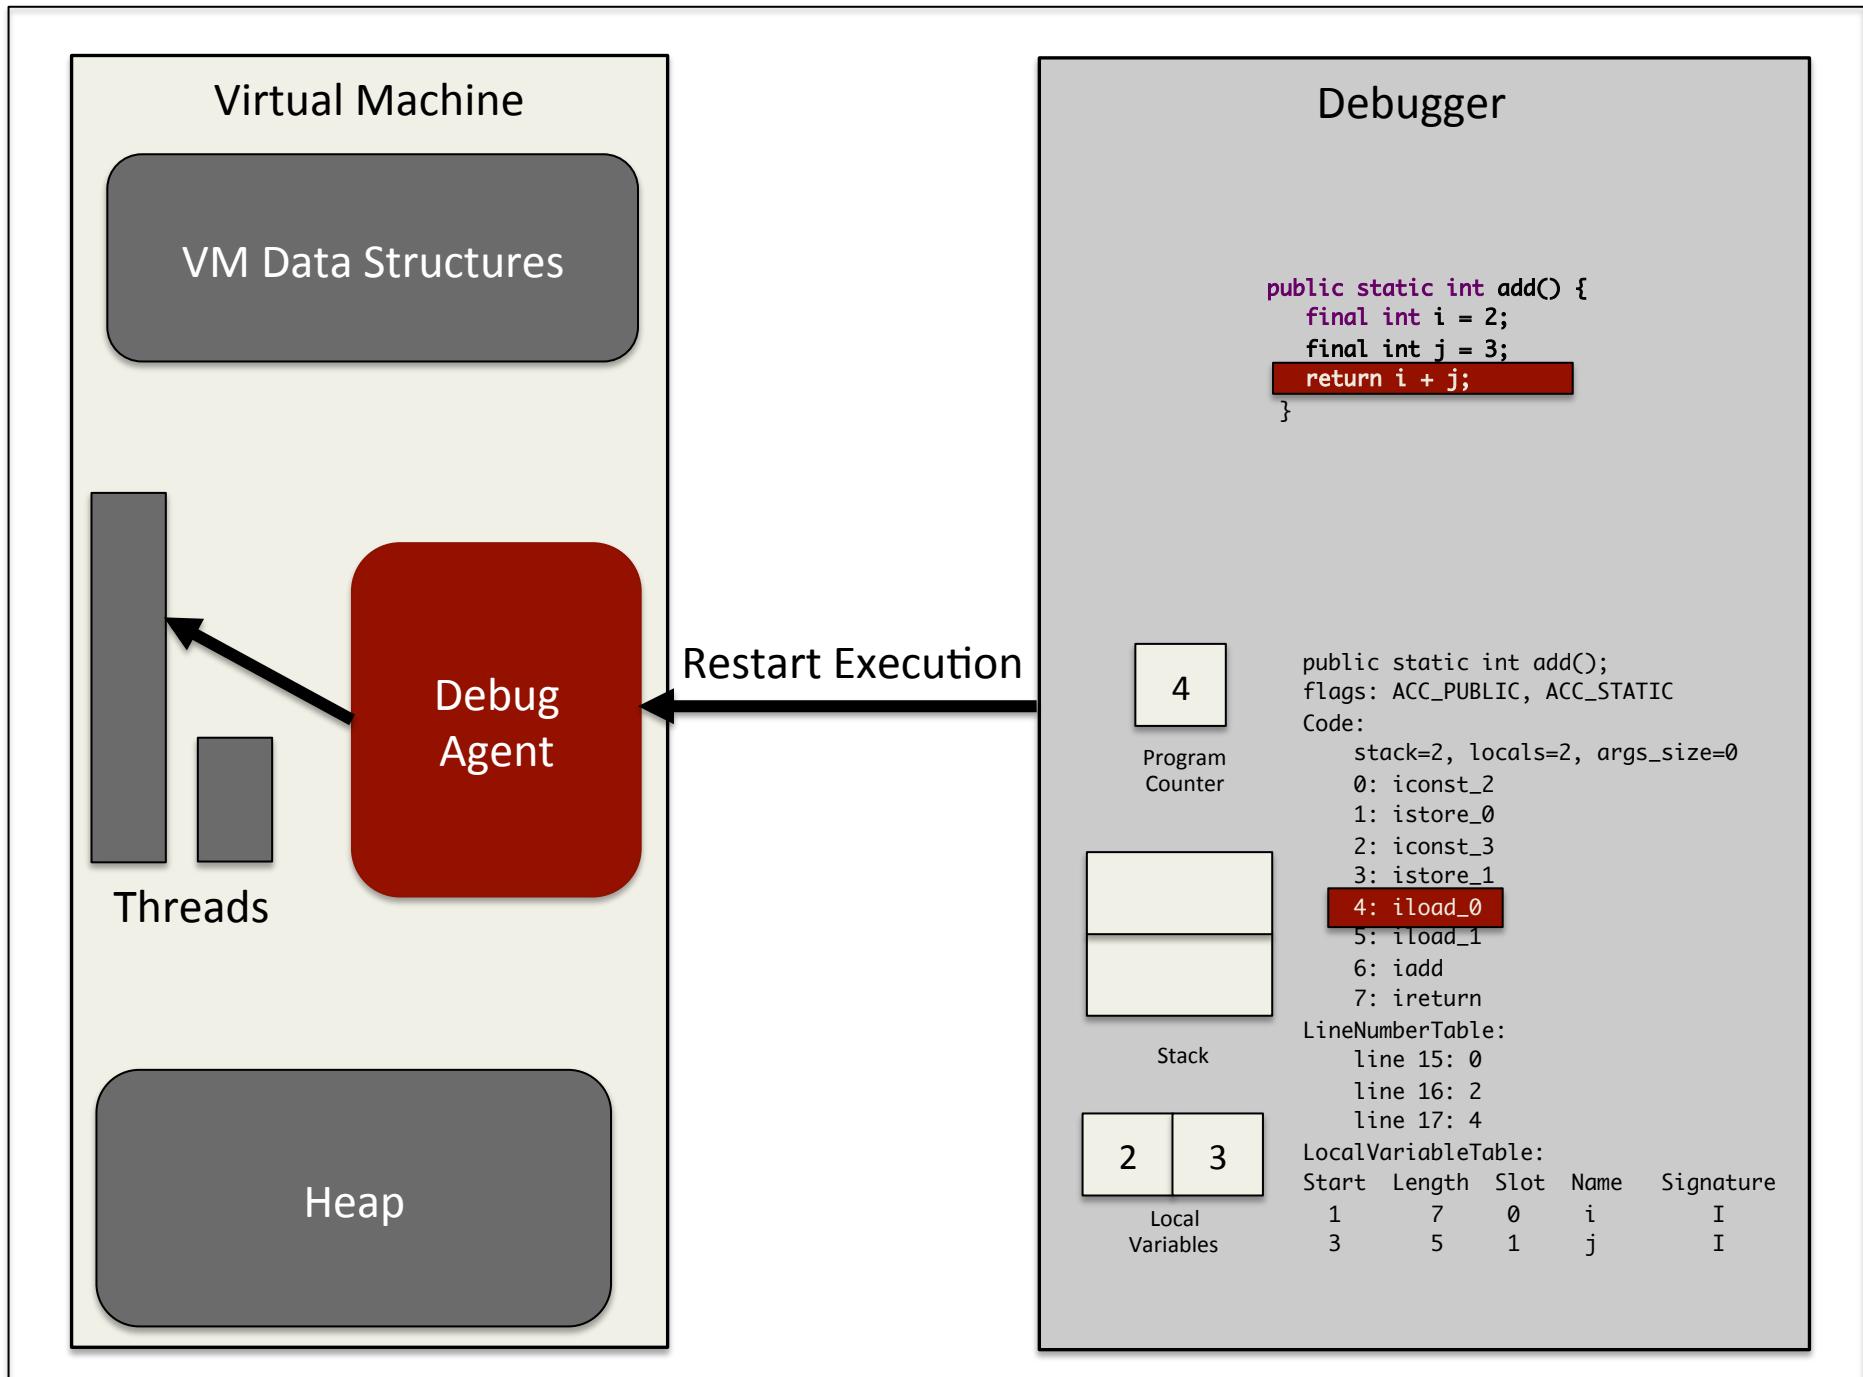
LineNumberTable:

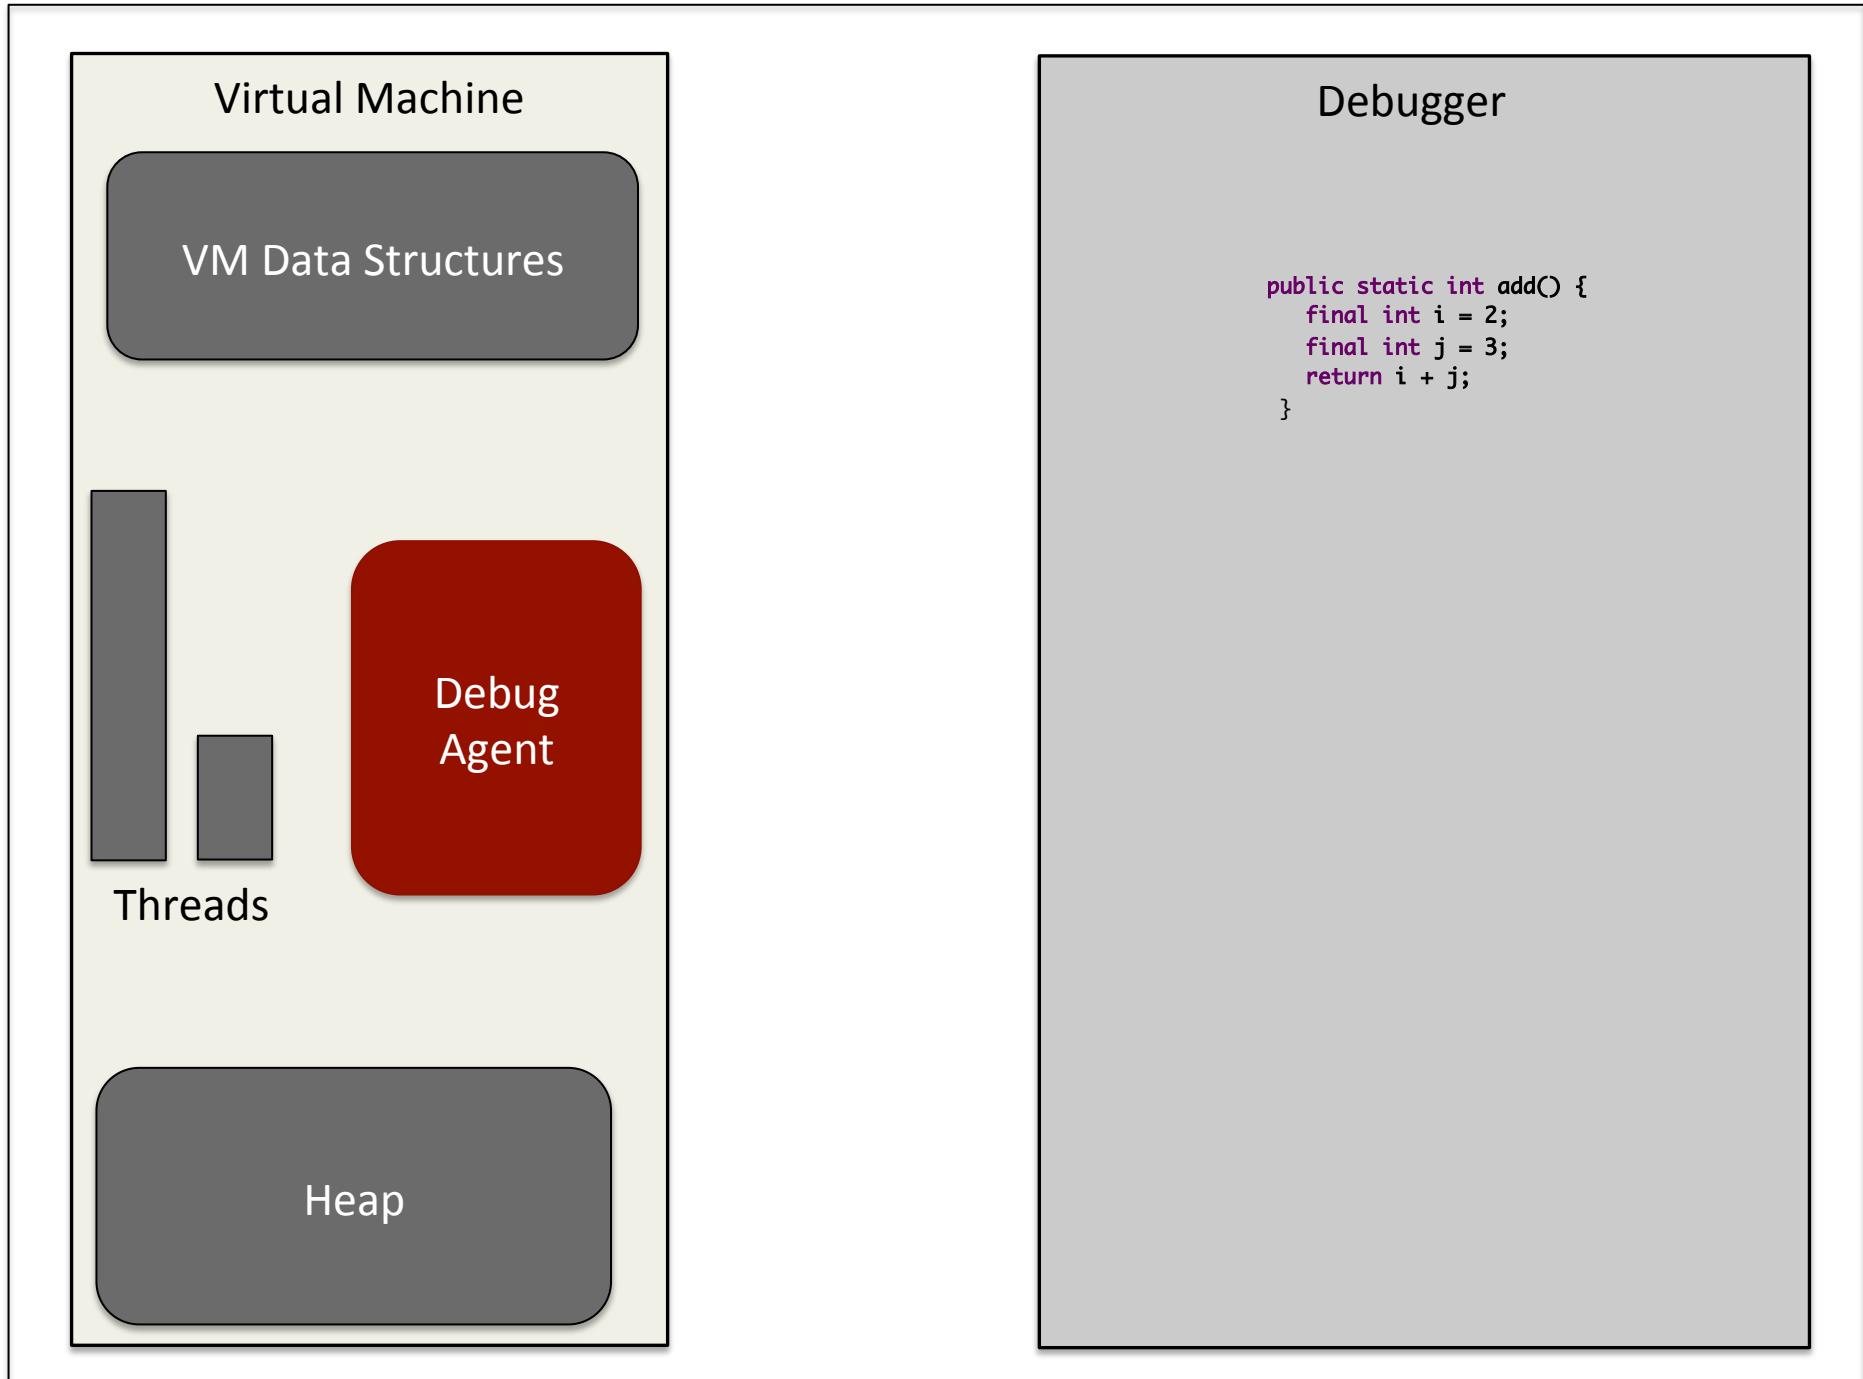
```
line 15: 0  
line 16: 2  
line 17: 4
```

LocalVariableTable:

| Start | Length | Slot | Name | Signature |
|-------|--------|------|------|-----------|
| 1 | 7 | 0 | i | I |
| 3 | 5 | 1 | j | I |







Missing Events

- Not all interesting operations have events
 - Object allocation
 - Operand stack or local variable activity
- Interface allows for bytecode modification
 - Rewrite the class file before the VM loads it
 - Reload a class to insert new code
- Instrumentation happens during class loading
 - After loading but before linking
 - Rewritten bytecode is verified

Bytecode Instrumentation

- Can insert very specific data gathering
 - Generate conditionally events
 - May rely on static analysis
- Profiling code is pure Java bytecode
 - New bytecode becomes part of the program
 - JIT compiler optimizes profiling code
- No native code involved in runtime profiling

Java-Level Agent Data

- May now have two methods of profiling
 - Instrumented bytecode at the Java level
 - JVMTI agent works in native code
- May need communication between the two
 - Can use JNI to access Java-level data
- May output two streams to external tool
 - Less communication inside the VM
 - More complexity in the tool interface

Instrumentation Limitations

- Bytecode instrumentation best for small changes
 - Add a static call to profiling code
 - Log program state at key points during execution
- Hard to make more elaborate changes
 - Can only modify classes, not create them
 - Cross-class dependencies hard to manage
 - Instrumentation code normally written in C
- Custom classloaders are much more flexible

Bytecode Rewriting

- Modifying bytecode is trickier than it seems
 - Need to pass verification
 - Gets harder when changes are not local
- Bytecode indices are important
 - Jump instructions
 - Local variable, line number, exception tables
- Adding local variables can cause unexpected issues
 - Some bytecodes have indices built in

Object Allocation

- Normally access by bytecode instrumentation
 - Instrument new, newarray, anewarray instructions
- Some allocations not captured this way
 - Constant interning
 - Reflection
 - Native allocations
- Special event signaled in these cases

Method Entry and Exit

- Capability exists for tracking method calls
 - Method entry and exit events can be generated
 - Callback method determines if they are interesting
- Tracking all method calls can be expensive
 - Well-factored code should consist of small methods
 - Need to generate events on inlined code
- Can use bytecode instrumentation instead
 - Much more selective about what to track

Heap Monitoring

- Interface provides some heap traversal calls
 - Can opt to find only reachable objects
 - Traversal order is not guaranteed
- Heap objects can be tagged on creation
 - Associate a `jlong` value with an object
 - Lets the agent keep track of that object later

Profiling

- Design of the tool similar to the debugger
- Use a different set of capabilities
 - Get monitor status
 - Access heap data
 - Receive object allocation events
 - Receive method enter and exit events
 - Ability to tag objects
 - Ability to instrument bytecode

Profiler Performance

- Some of the capabilities are heavyweight
 - Tracing the heap
 - Tracking allocations
 - Tracing method enter and exit
- Profilers should minimize disruption
- Can decide up-front what to monitor
 - Separate runs to track heap and execution time
 - Tailor the capabilities that the agent requests

Other uses for Tool Interface

- Code coverage
 - Instrument all paths through a method
- Deadlock or livelock analysis
 - Pause threads and check monitor status
- VM monitoring
 - Correlate VM events to OS or hardware counters
- Extreme bytecode rewriting
 - Modify classes that otherwise can't be rewritten