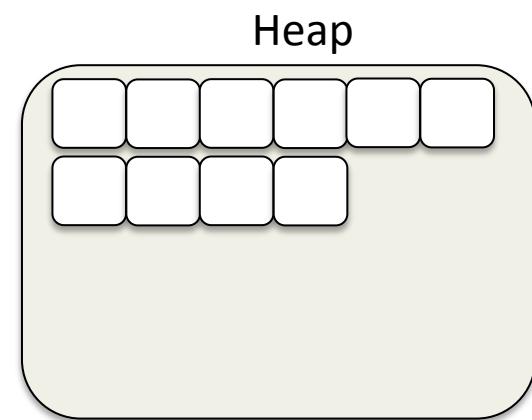


In the Last Week

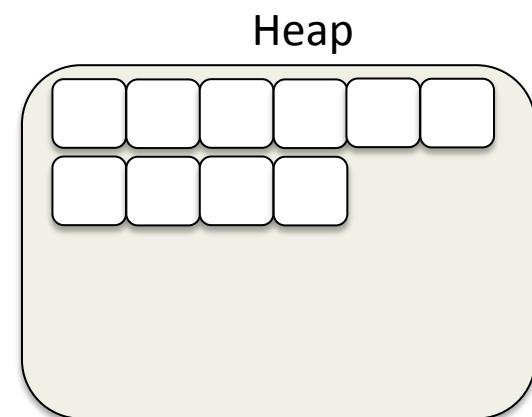
- Assignment 4 due yesterday
- Final project posted now
 - Due on December 7th
- Thanks for special topic and review suggestions
- GC debugging for fun and profit

```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

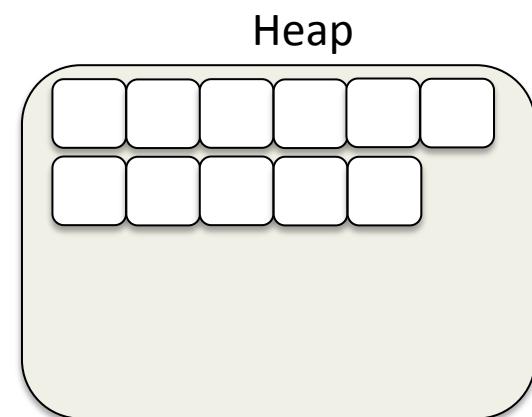
```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```



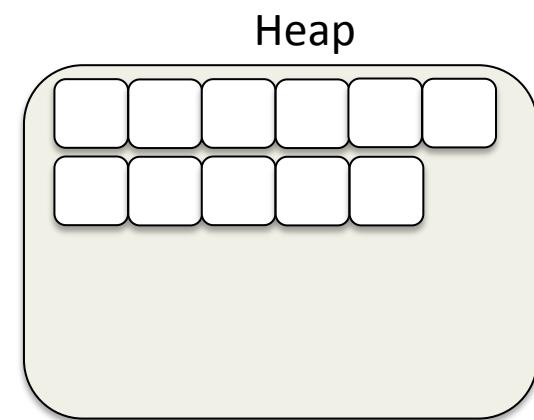
```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```



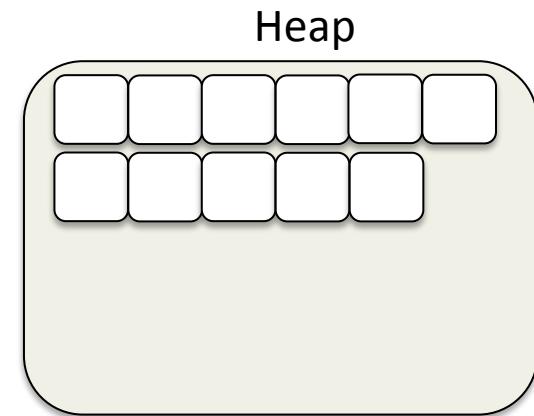
```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```



```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

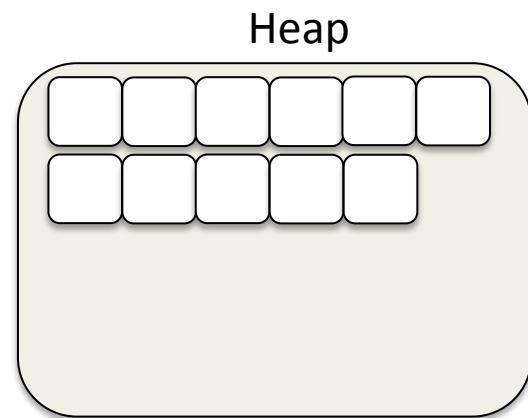


```
constructor() {  
    value = Cls.STATIC;  
}  
  
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```



```
constructor() {  
    value = Cls.STATIC;  
}  
  
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

Class Loader

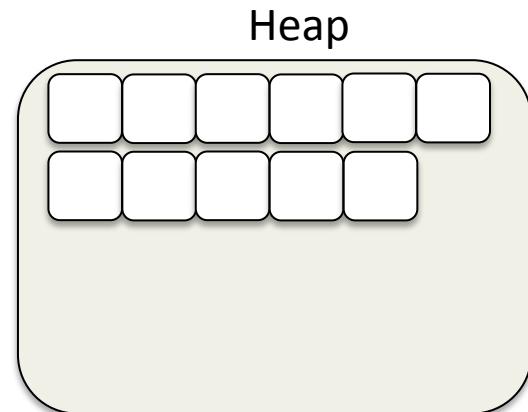


```
static {  
    object = new Object();  
}
```

```
constructor() {  
    value = Cls.STATIC;  
}
```

```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

Class Loader

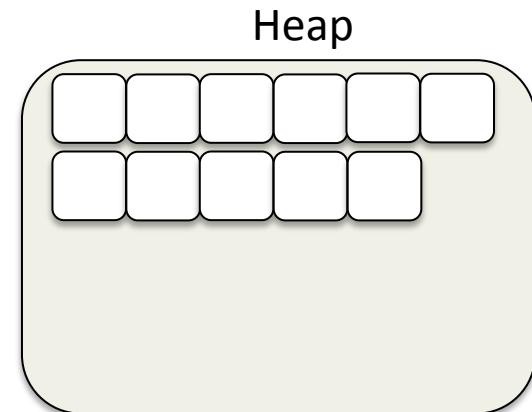


```
static {  
    object = new Object();  
}
```

```
constructor() {  
    value = Cls.STATIC;  
}
```

```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

Class Loader

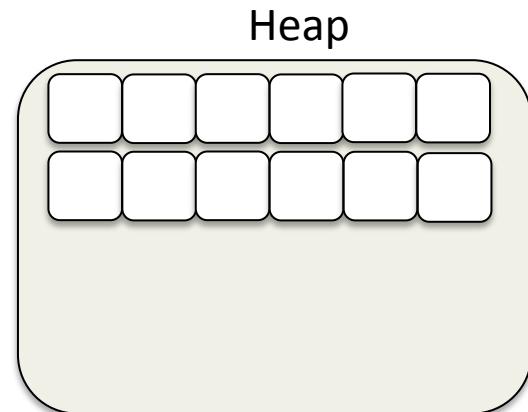


```
static {  
    object = new Object();  
}
```

```
constructor() {  
    value = Cls.STATIC;  
}
```

```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

Class Loader

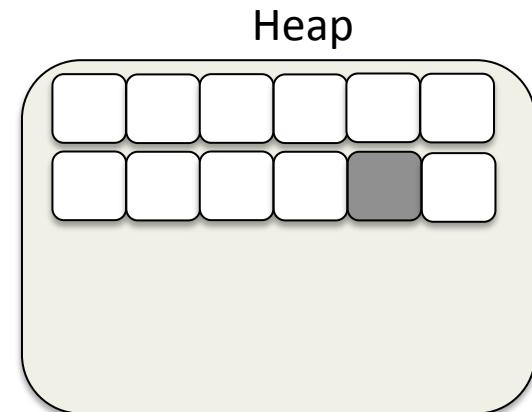


```
static {  
    object = new Object();  
}
```

```
constructor() {  
    value = Cls.STATIC;  
}
```

```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

Class Loader

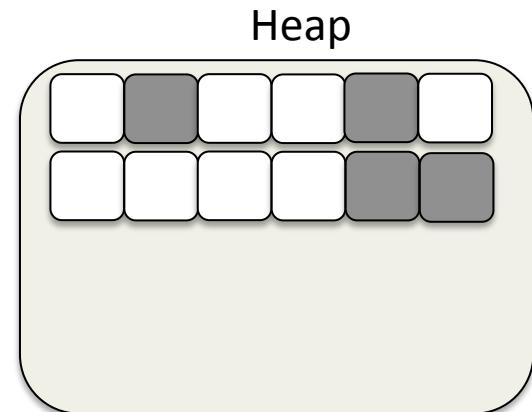


```
static {  
    object = new Object();  
}
```

```
constructor() {  
    value = Cls.STATIC;  
}
```

```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

Class Loader

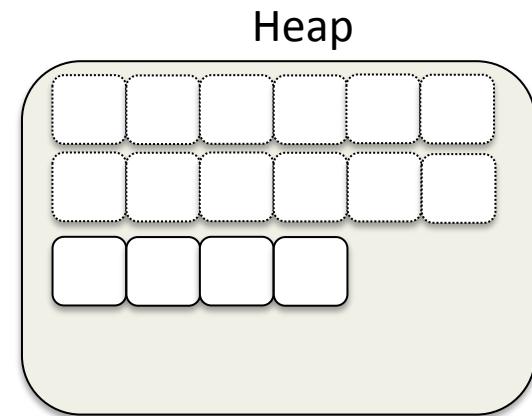


```
static {  
    object = new Object();  
}
```

```
constructor() {  
    value = Cls.STATIC;  
}
```

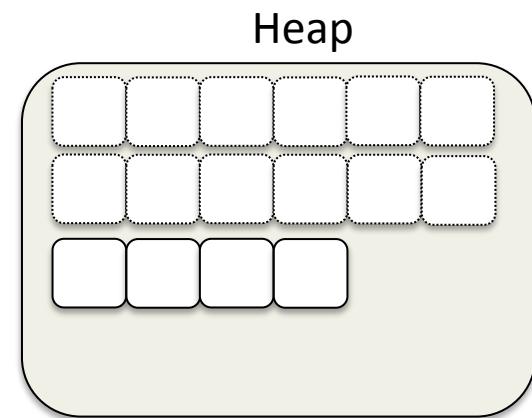
```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

Class Loader

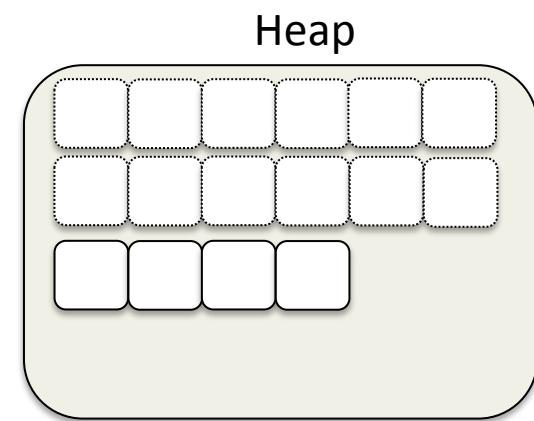


```
constructor() {  
    value = Cls.STATIC;  
}  
  
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```

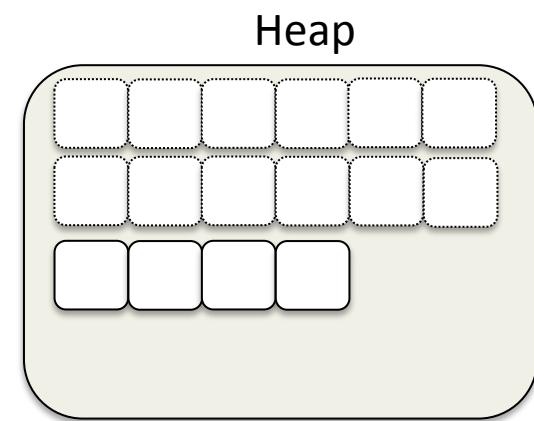
Class Loader



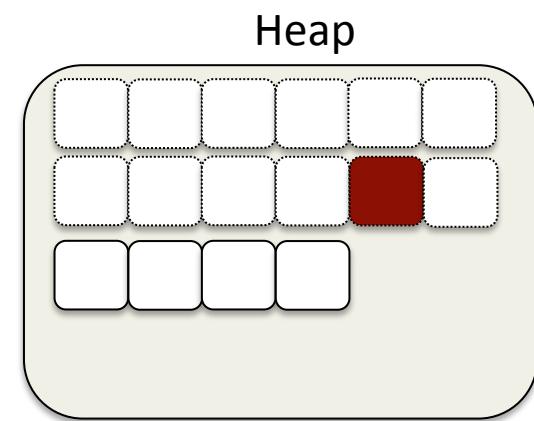
```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```



```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```



```
createObject {  
    ptr = allocate object  
    call constructor(ptr)  
    return ptr;  
}
```



Final Project

- The final project is now posted
 - No takers for the team version
- Implement generational garbage collector
 - See slide deck 14 for details
 - Lecture from Monday October 27th
- Three week timeframe
 - Non-trivial amount of work
 - Instructions assume good understanding of the code

Generational Collector Design

- Fixed-size nursery
 - Bump pointer allocation
 - Evacuate objects to the mature space
- Mark/Sweep mature space management
 - Allocation happens only during minor GC
 - Allocation uses first-fit algorithm
- Write barrier on pointer writes

Part I: Write Barrier

- Don't scan the full heap during minor collections
 - Major performance advantage
 - Need to track references from old to new space
- Use a write barrier to track pointer writes
 - Two bytecode instructions can write pointers
 - Instrument those instructions in the interpreter
- New interface and implementation class

Part I: Write Barrier

- Write barrier maintains list of old-new references
 - Actually, objects that may point from old to new
- Set should only contain relevant objects
 - Don't need to track other references
 - Correctly building the remembered set is important
- Remembered set allocated off to the side
 - Don't need to worry about space overhead

Part II: Non-Contiguous Region

- Mark/Sweep collector causes fragmentation
 - Objects are deleted without moving
- Need an allocator to allocate between objects
 - Current bump pointer region won't work
- New Region class for non-contiguous allocation
 - Method on Heap interface to get an instance

Part II: Non-Contiguous Region

- Allocate to the region using first-fit
 - May be best to track objects and holes
 - Metadata can be stored off to the side
- Support the Region interface's operations
 - `allocate`, `free`, `pointerInRegion`
 - `reset` operation not available
- Will need other operations for mark/sweep GC
 - Some way to scan through all objects
 - Implement on the class, not the interface

Part III: Mark/Sweep Collector

- Mature space managed using mark/sweep
 - Use a single non-contiguous region
 - New memory manager class skeleton
- Root set the same as in Assignment 3
 - Stack and local variables
 - Static fields
 - Interned strings

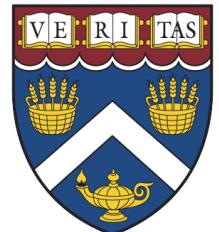
Part III: Mark/Sweep Collector

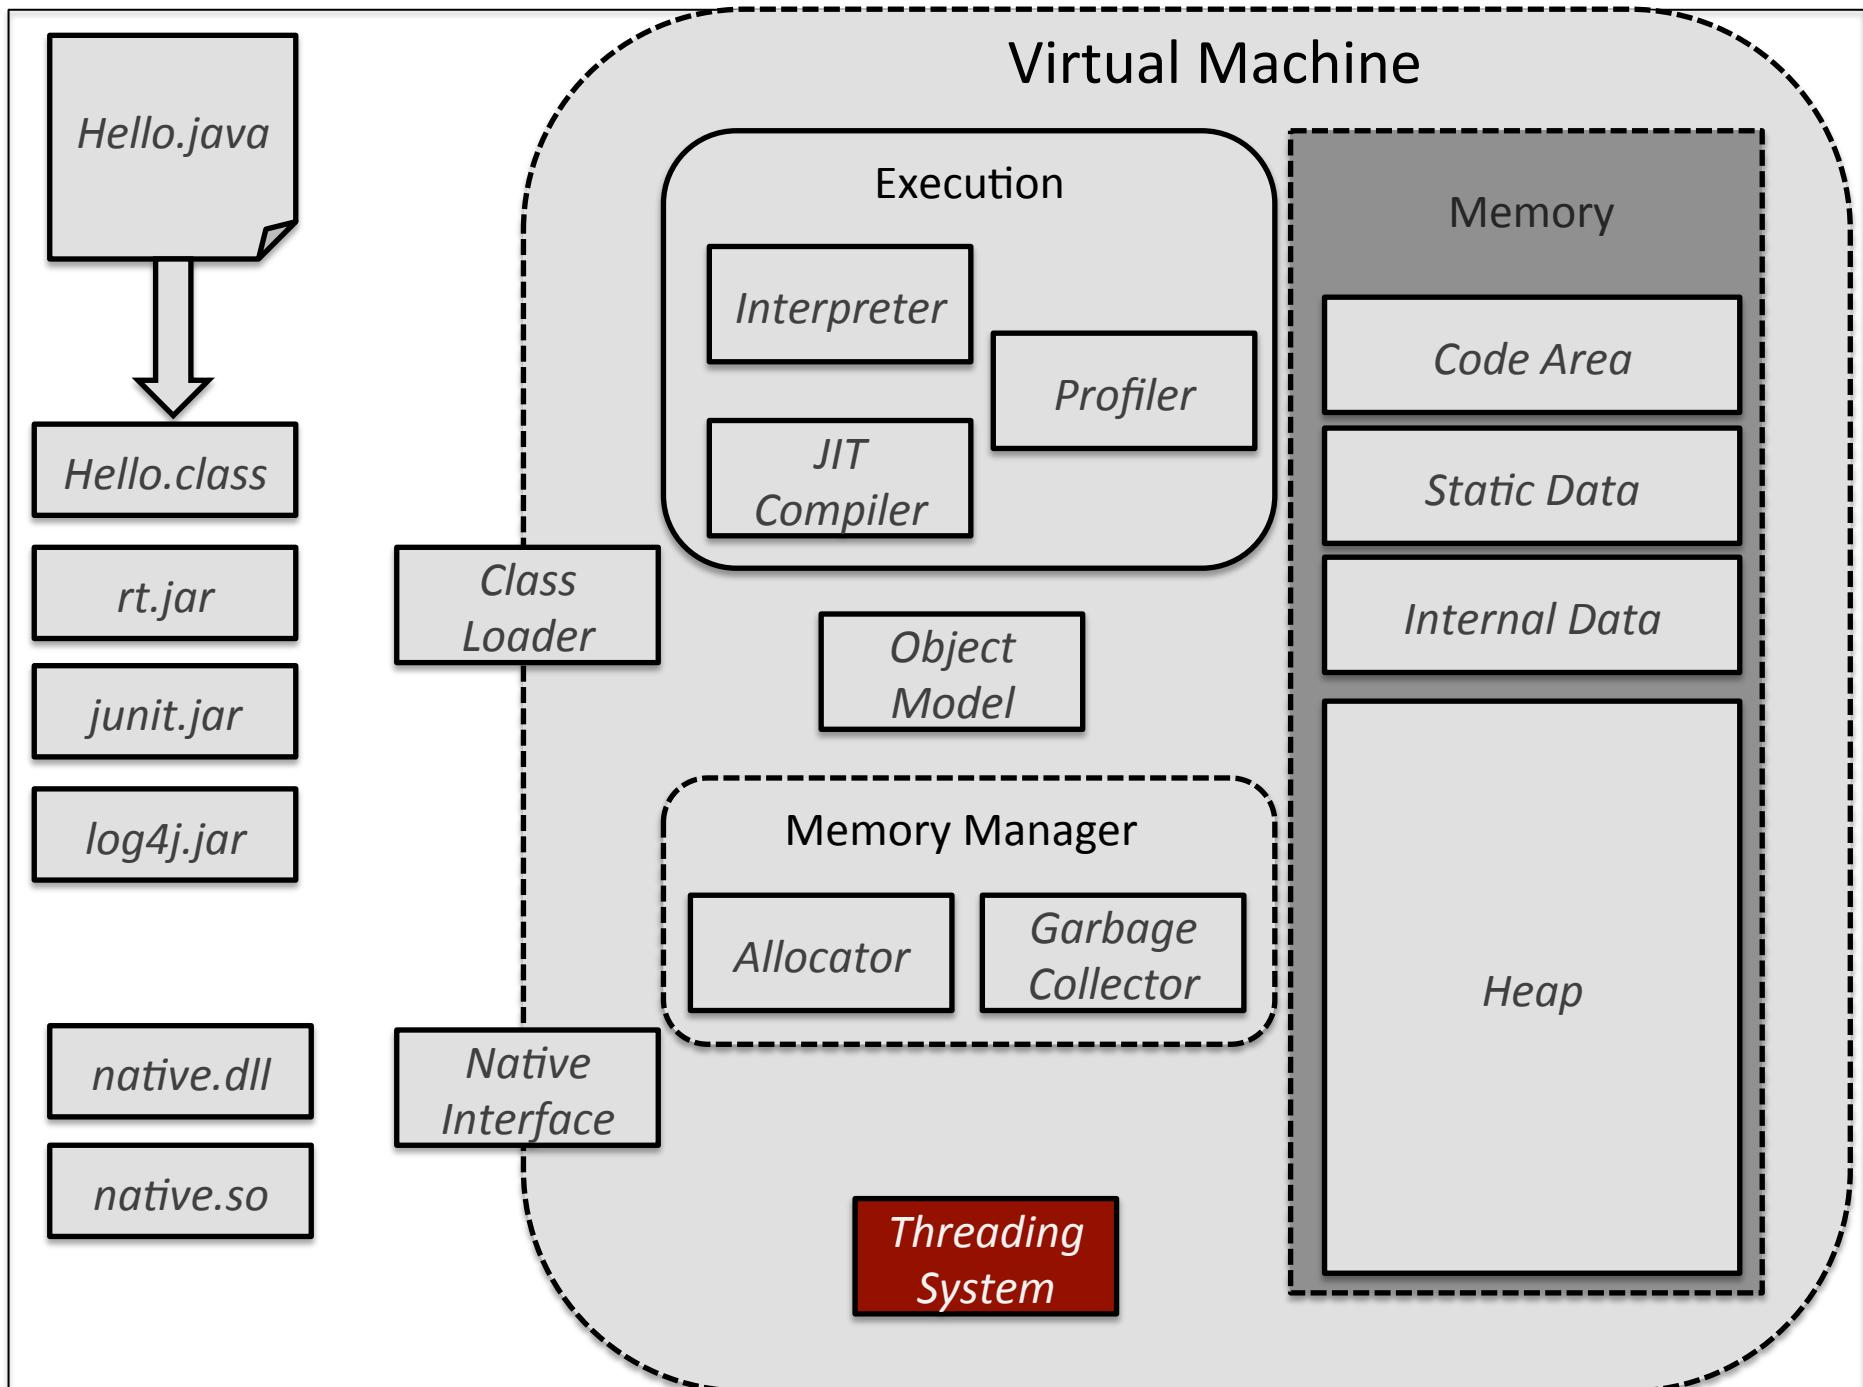
- Trigger garbage collection manually
 - Allocation only happens during minor collection
 - Heap in a strange state
- May want to test this collector separately
 - Add temporary logic to trigger collection
 - Mark/Sweep will be a fully-functional GC
 - See `SimpleJavaVM.java` to switch collectors

Part IV: Nursery Collection

- Final part of the generational collector
- Roughly based on the semi-space collector
 - Root set includes the remembered set
 - Nursery is a single region, not a semi-space
 - Use an explicit gray stack for minor collection
- Trigger major collection after minor GC
 - Call when mature space is $\geq 2 * \text{nursery size}$

Threading and Concurrency





Concurrency

- Multiple operations occurring simultaneously
 - Possibly interacting with one another
- Concurrent systems are now commonplace
 - Offer major performance improvements
 - Make development far more difficult
- VM concurrency falls into two broad groups:
 - Application-level concurrency
 - VM internal concurrency

Processes and Threads

- Primary unit of concurrency is the process
 - Separate memory space
 - Isolated from other processes
- Threads allow concurrency within a process
 - Far more lightweight than processes
 - Shared address space
 - One thread's actions can affect another

Moore's Law

- Frequently misstated
 - The speed of processors doubles every two years
- Actually relates to processor manufacturing
 - Transistor density doubles every two years
- For a long time, the two were equivalent
- Complex circuitry means more heat
 - Modern designs use multiple simpler cores

Hardware Concurrency Support

- Multicore processors allow genuine concurrency
 - Previously required multiple CPUs
 - Single core may also provide hardware multithreading
- CPU optimizes individual threads
 - Out of order execution
- Also provides instructions to handle concurrency
 - Memory fences
 - Atomic CAS instructions

Thread 1



Thread 1



Thread 2



Thread 1

Thread 2

var_1

Ø



Thread 1

int val = var_1

Thread 2

var_1

Ø



Thread 1

val

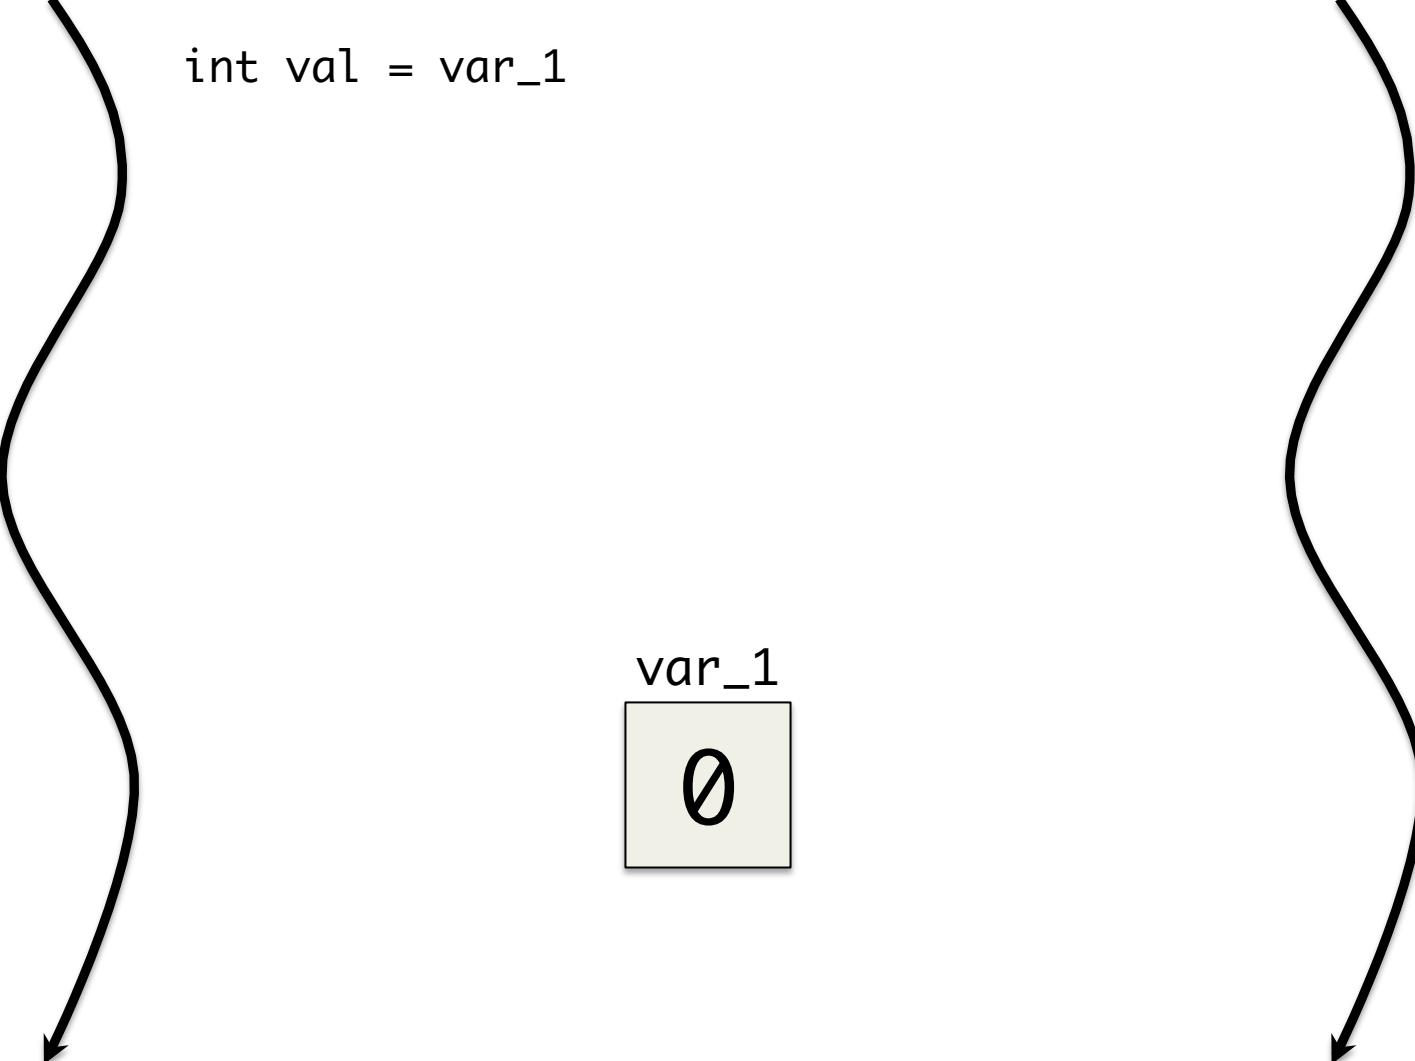
0

int val = var_1

Thread 2

var_1

0



Thread 1

val

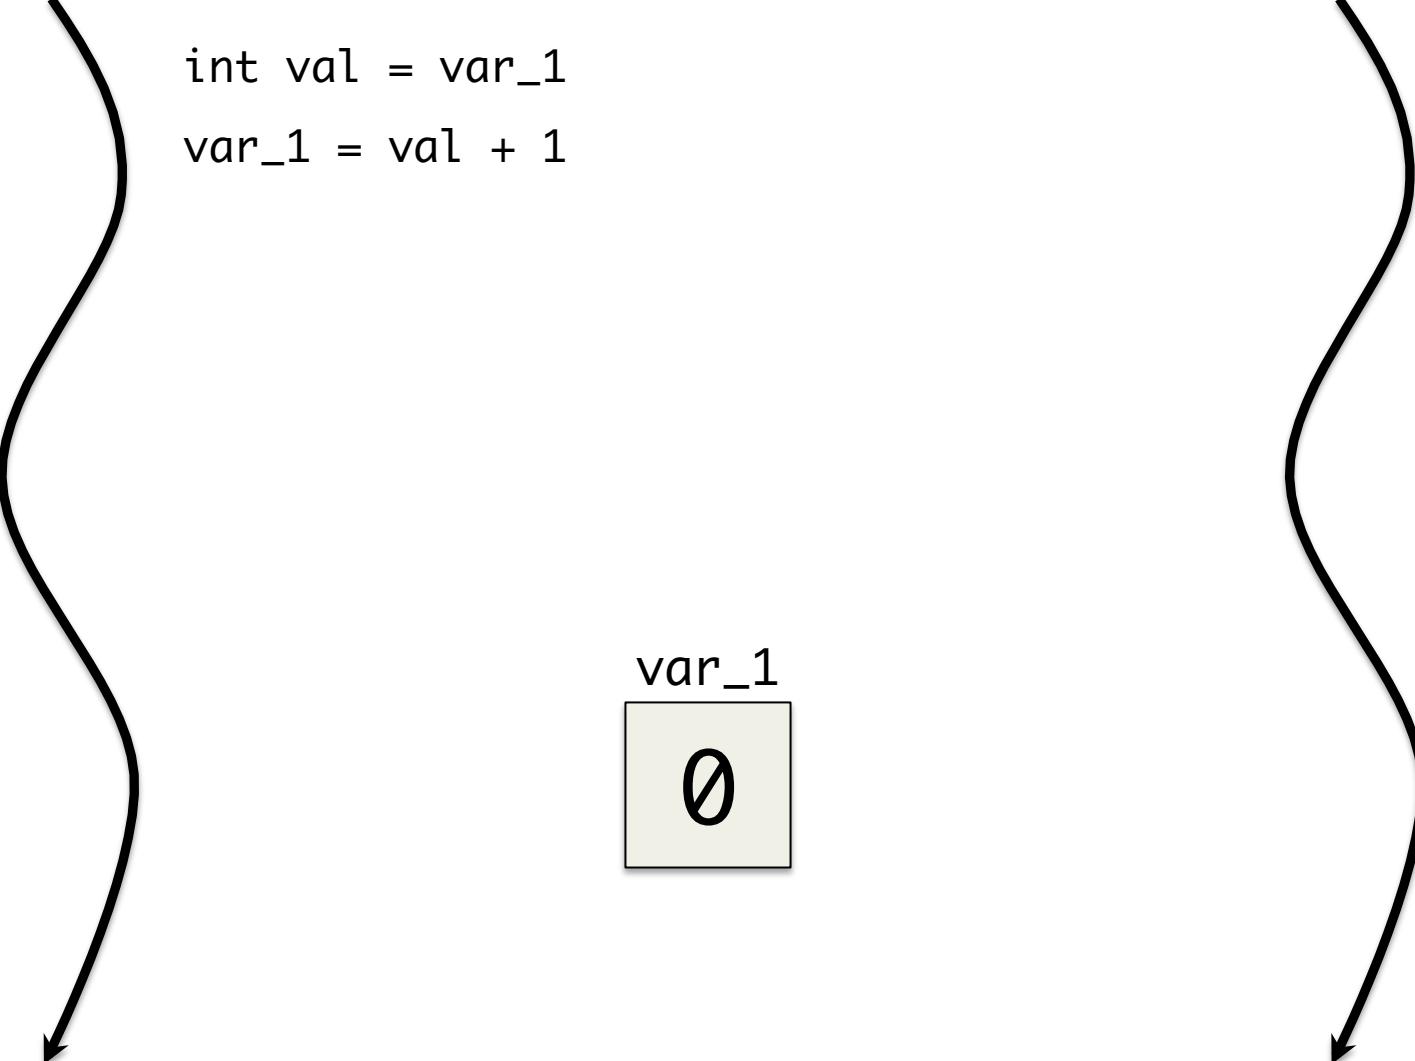
0

```
int val = var_1  
var_1 = val + 1
```

Thread 2

var_1

0



Thread 1

val

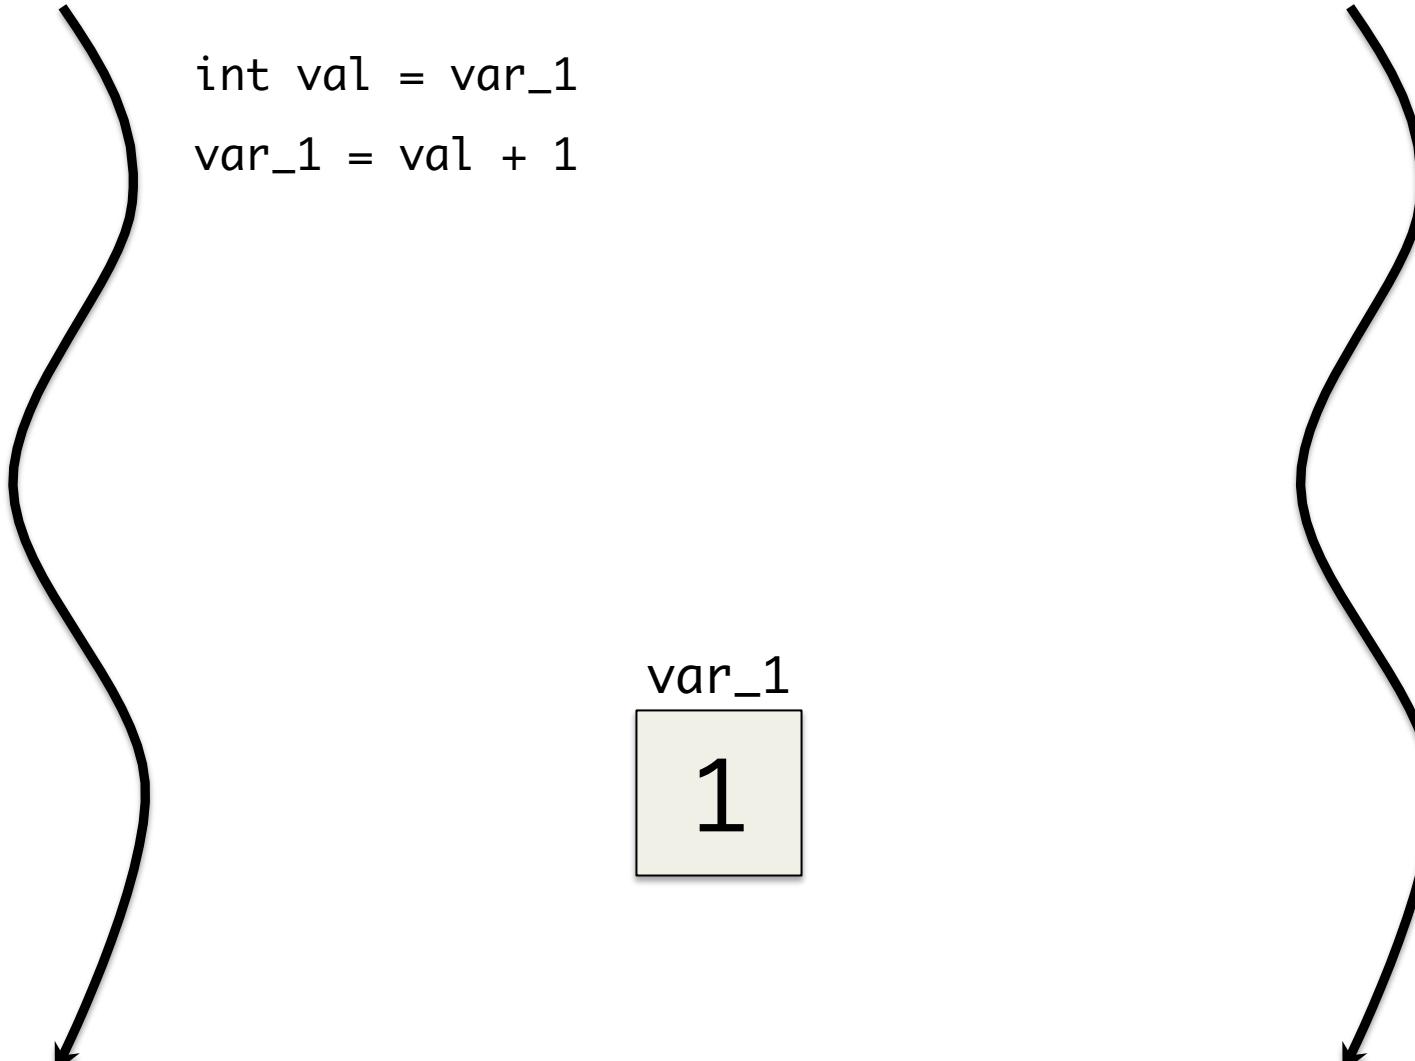
0

```
int val = var_1  
var_1 = val + 1
```

Thread 2

var_1

1



Thread 1

val

0

```
int val = var_1  
var_1 = val + 1
```

Thread 2

```
int val = var_1
```

var_1

1



Thread 1

val

0

```
int val = var_1  
var_1 = val + 1
```

Thread 2

val

1

```
int val = var_1
```

var_1

1

Thread 1

val

0

```
int val = var_1  
var_1 = val + 1
```

Thread 2

val

1

```
int val = var_1  
var_1 = val + 2
```

var_1

1

Thread 1

val

0

```
int val = var_1  
var_1 = val + 1
```

Thread 2

val

1

```
int val = var_1  
var_1 = val + 2
```

var_1

3

Thread 1

Thread 2

var_1

Ø



Thread 1

int val = var_1

Thread 2

var_1



Thread 1

val

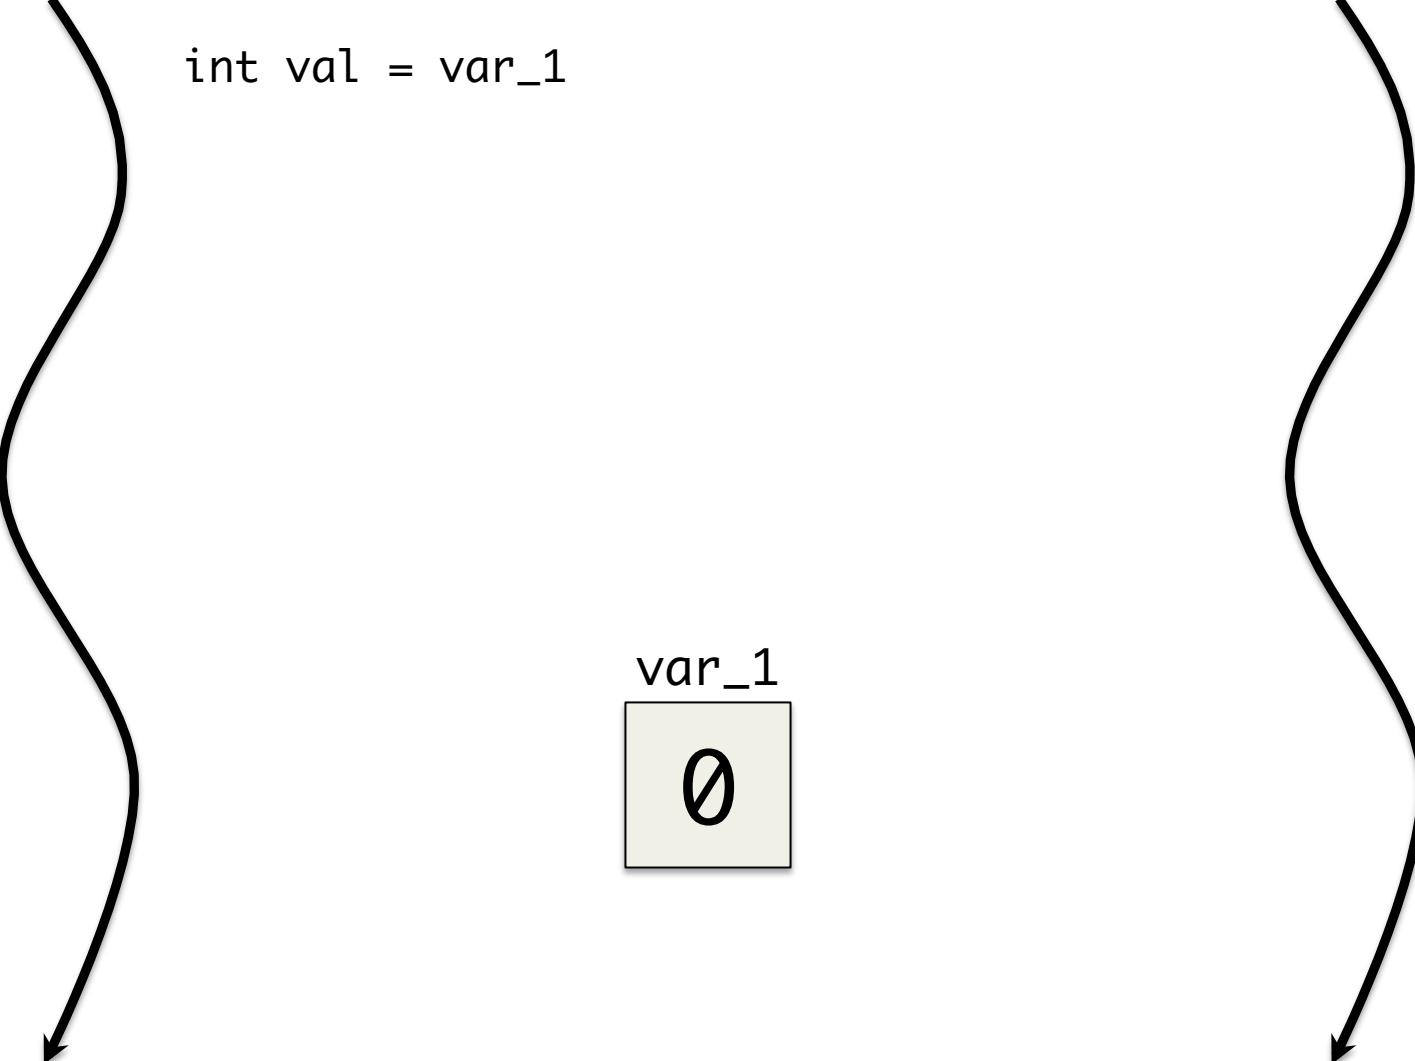
0

int val = var_1

Thread 2

var_1

0



Thread 1

val

0

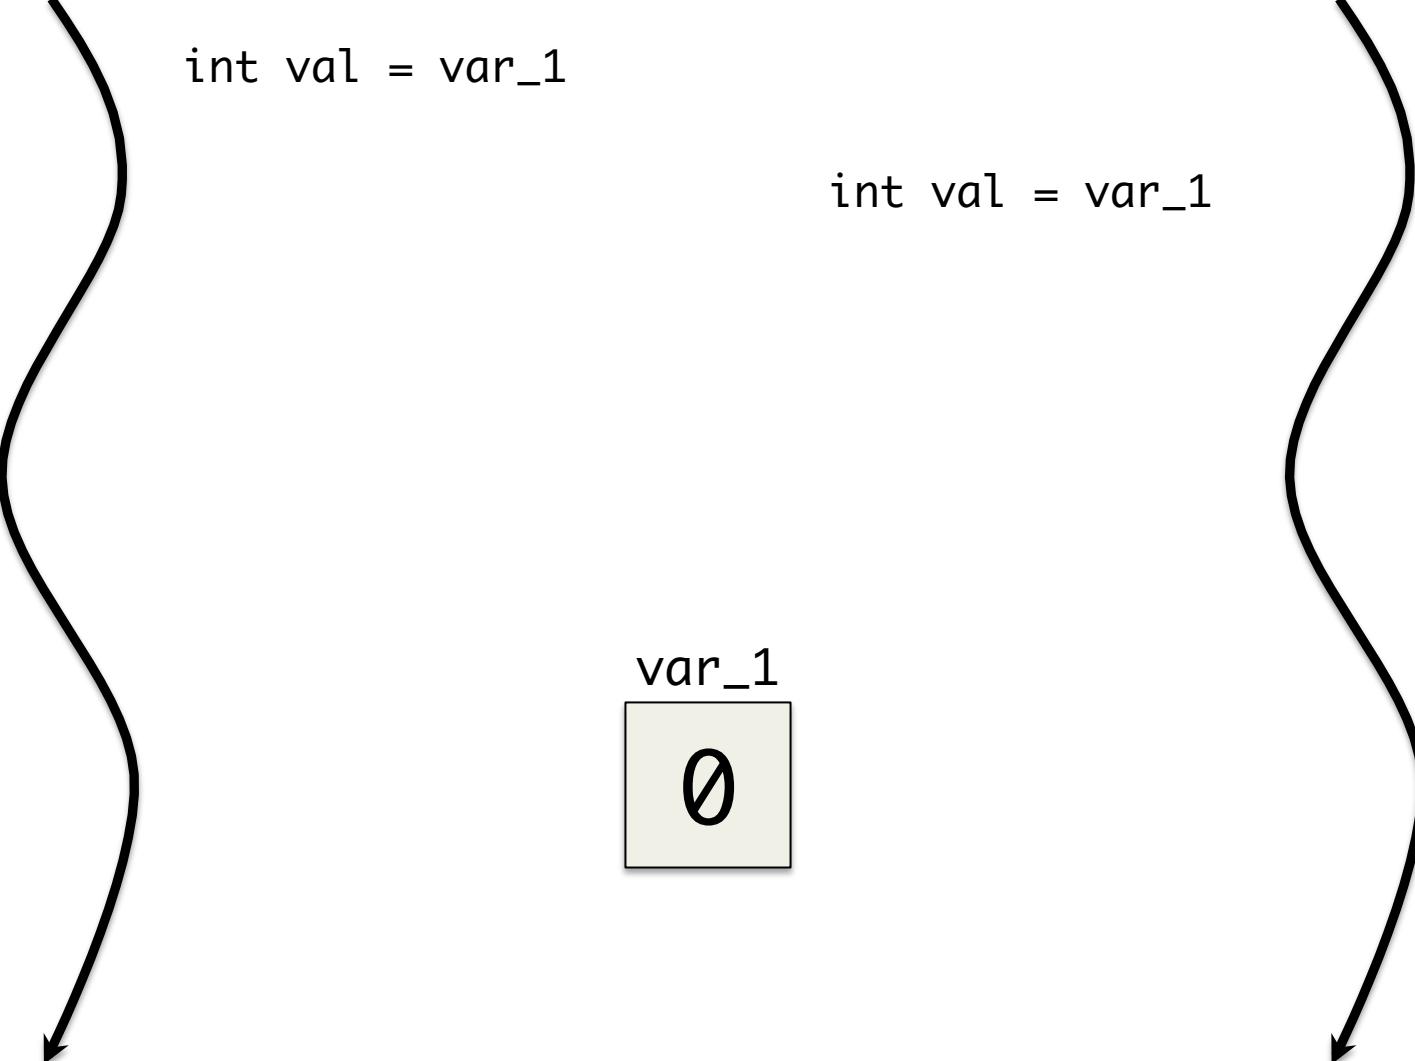
int val = var_1

Thread 2

int val = var_1

var_1

0



Thread 1

val

0

int val = var_1

Thread 2

val

0

int val = var_1

var_1

0



Thread 1

val

0

int val = var_1

Thread 2

int val = var_1
var_1 = val + 2

val

0

var_1

0



Thread 1

val

0

int val = var_1

Thread 2

int val = var_1
var_1 = val + 2

val

0

var_1

2



Thread 1

val

0

int val = var_1

var_1 = val + 1

Thread 2

val

0

int val = var_1

var_1 = val + 2

var_1

2

Thread 1

val

0

int val = var_1

var_1 = val + 1

Thread 2

val

0

int val = var_1

var_1 = val + 2

var_1

1

Thread 1

Thread 2

var_1

Ø



Compare and Swap

- Atomic instruction to update a value
 - Takes an original value
 - Compares with the current contents of memory
- If the test succeeds, update memory immediately
 - No other thread can pre-empt the update
- If the test fails, return the current value
 - Don't perform the update

Thread 1

val

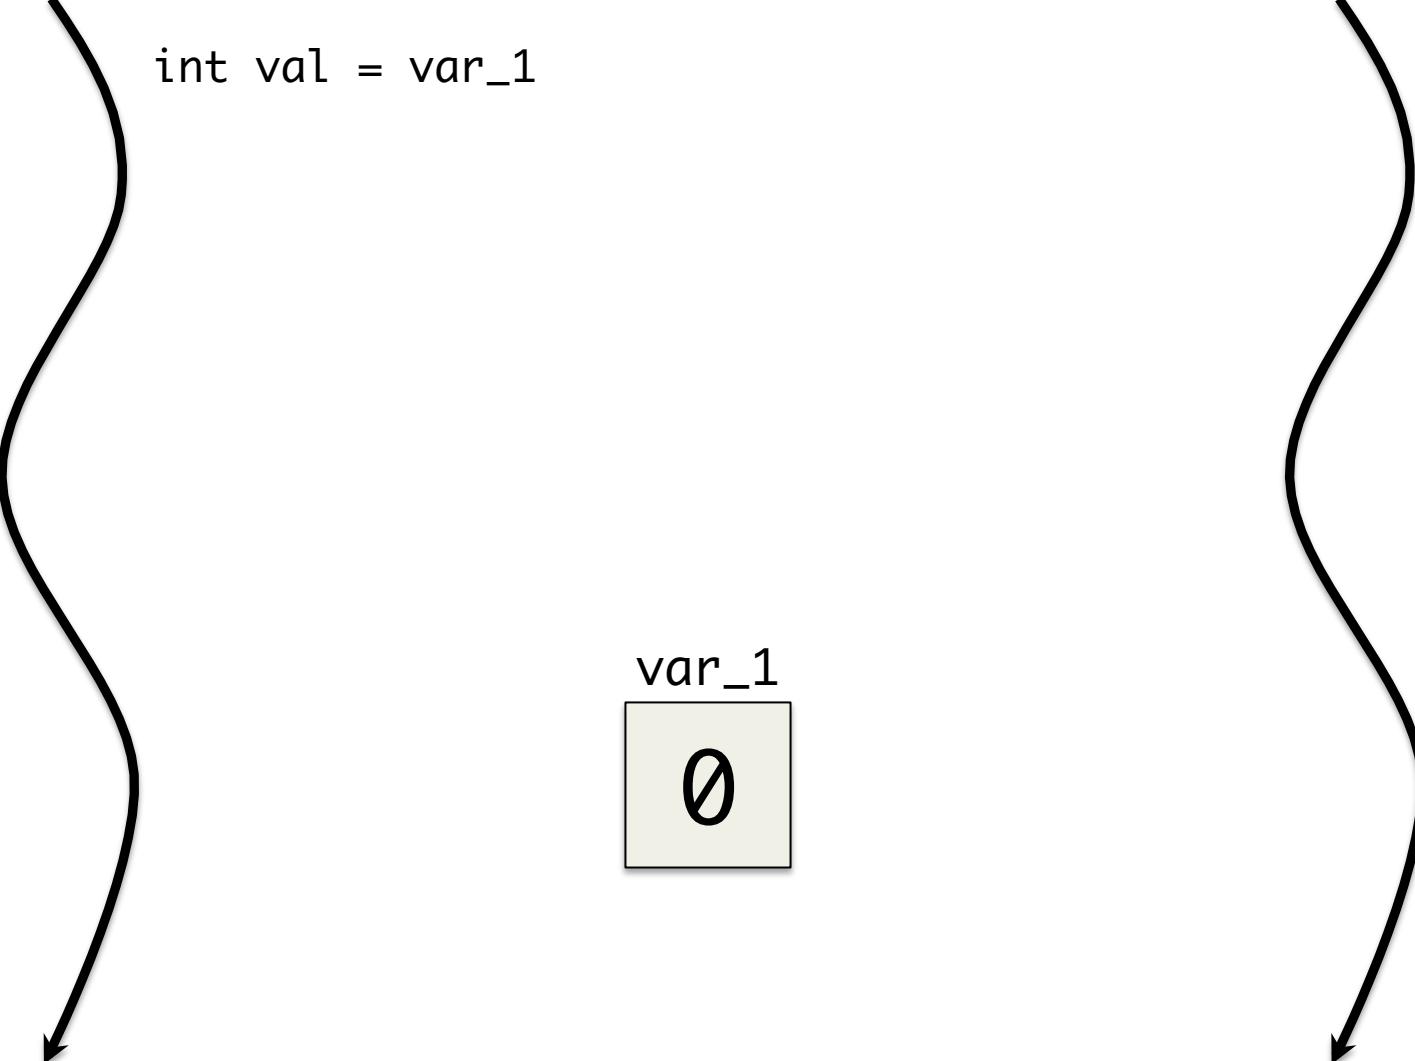
0

int val = var_1

Thread 2

var_1

0



Thread 1

val

0

int val = var_1

Thread 2

val

0

int val = var_1

var_1

0



Thread 1

val

0

int val = var_1

Thread 2

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 2)

var_1

0

Thread 1

val

0

int val = var_1

Thread 2

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 2)

var_1

2

Thread 1

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 1)

Thread 2

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 2)

var_1

2

Thread 1

val

0

int val = var_1

Thread 2

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 2)

var_1 = CAS(var_1, val, val + 1)

var_1

2

Thread 1

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 1)

int val = var_1

var_1

2

Thread 2

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 2)



Thread 1

val
2

int val = var_1

var_1 = CAS(var_1, val, val + 1)

int val = var_1

var_1

2

Thread 2

val
0

int val = var_1
var_1 = CAS(var_1, val, val + 2)

Thread 1

val
2

int val = var_1

var_1 = CAS(var_1, val, val + 2)

var_1 = CAS(var_1, val, val + 1)

int val = var_1

var_1 = CAS(var_1, val, val + 1)

var_1

2

Thread 2

val
0

int val = var_1

var_1 = CAS(var_1, val, val + 2)



Thread 1

val

2

int val = var_1

int val = var_1

var_1 = CAS(var_1, val, val + 1)

var_1

3

Thread 2

val

0

int val = var_1

var_1 = CAS(var_1, val, val + 2)

Operating System Threading

- Various native threading libraries
 - `pthreads` available on most platforms today
 - Windows, Solaris, OS X, iOS have native packages
- Some standard features
 - Atomic operations
 - Mutual exclusion locks
 - Semaphores
 - Thread-local storage

Language-Level Threads

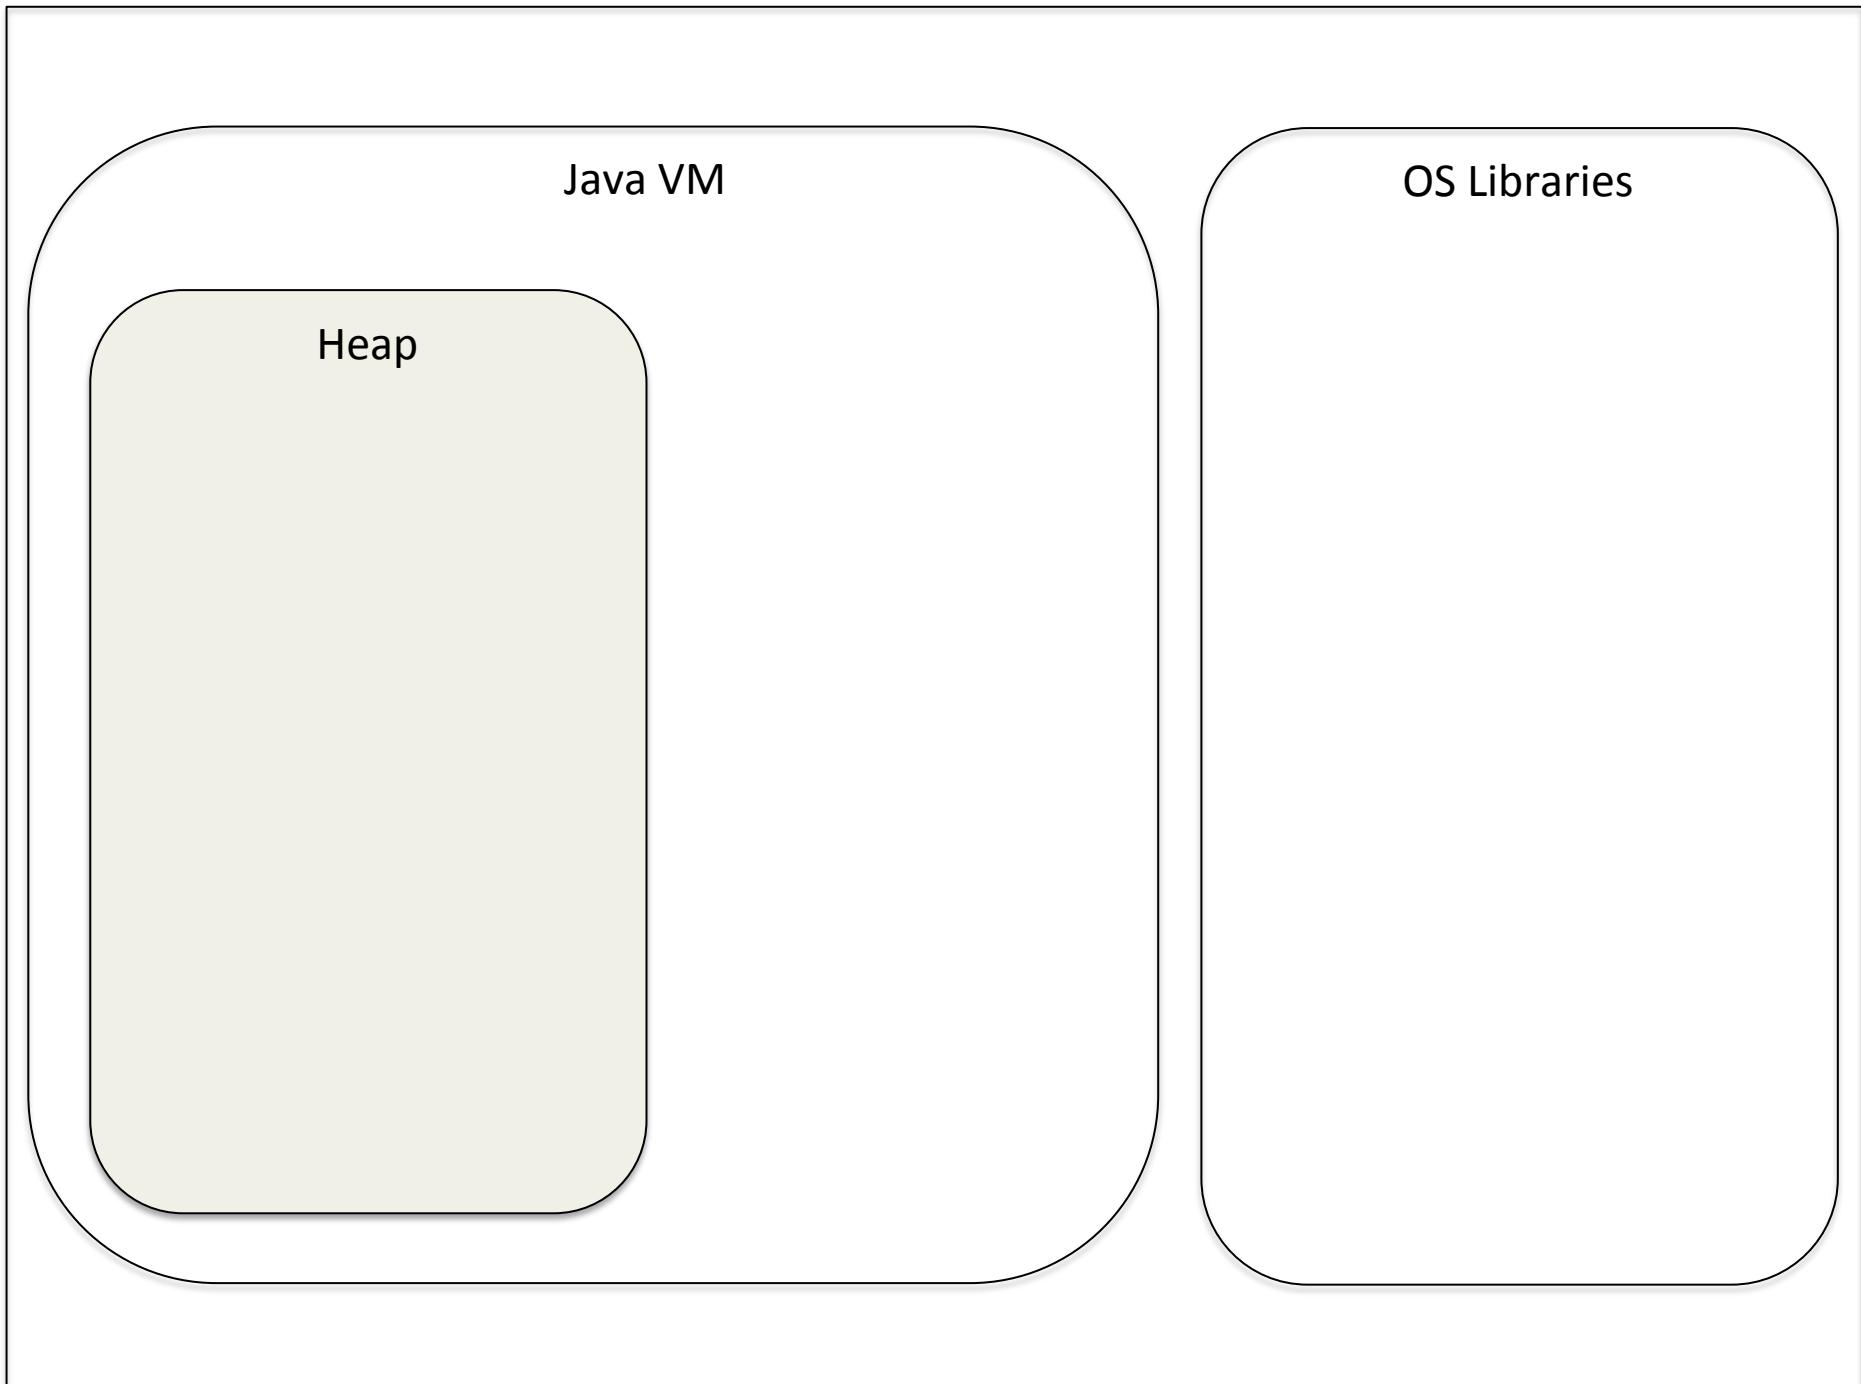
- Higher level abstraction than OS-level
 - Boundary can be fuzzy
- Describes concurrency abstraction
 - More concerned with meaning than mechanism
- Degree of abstraction varies by language
 - Fine control harder to manage but more powerful
 - Very abstract is easy to work with, but limited

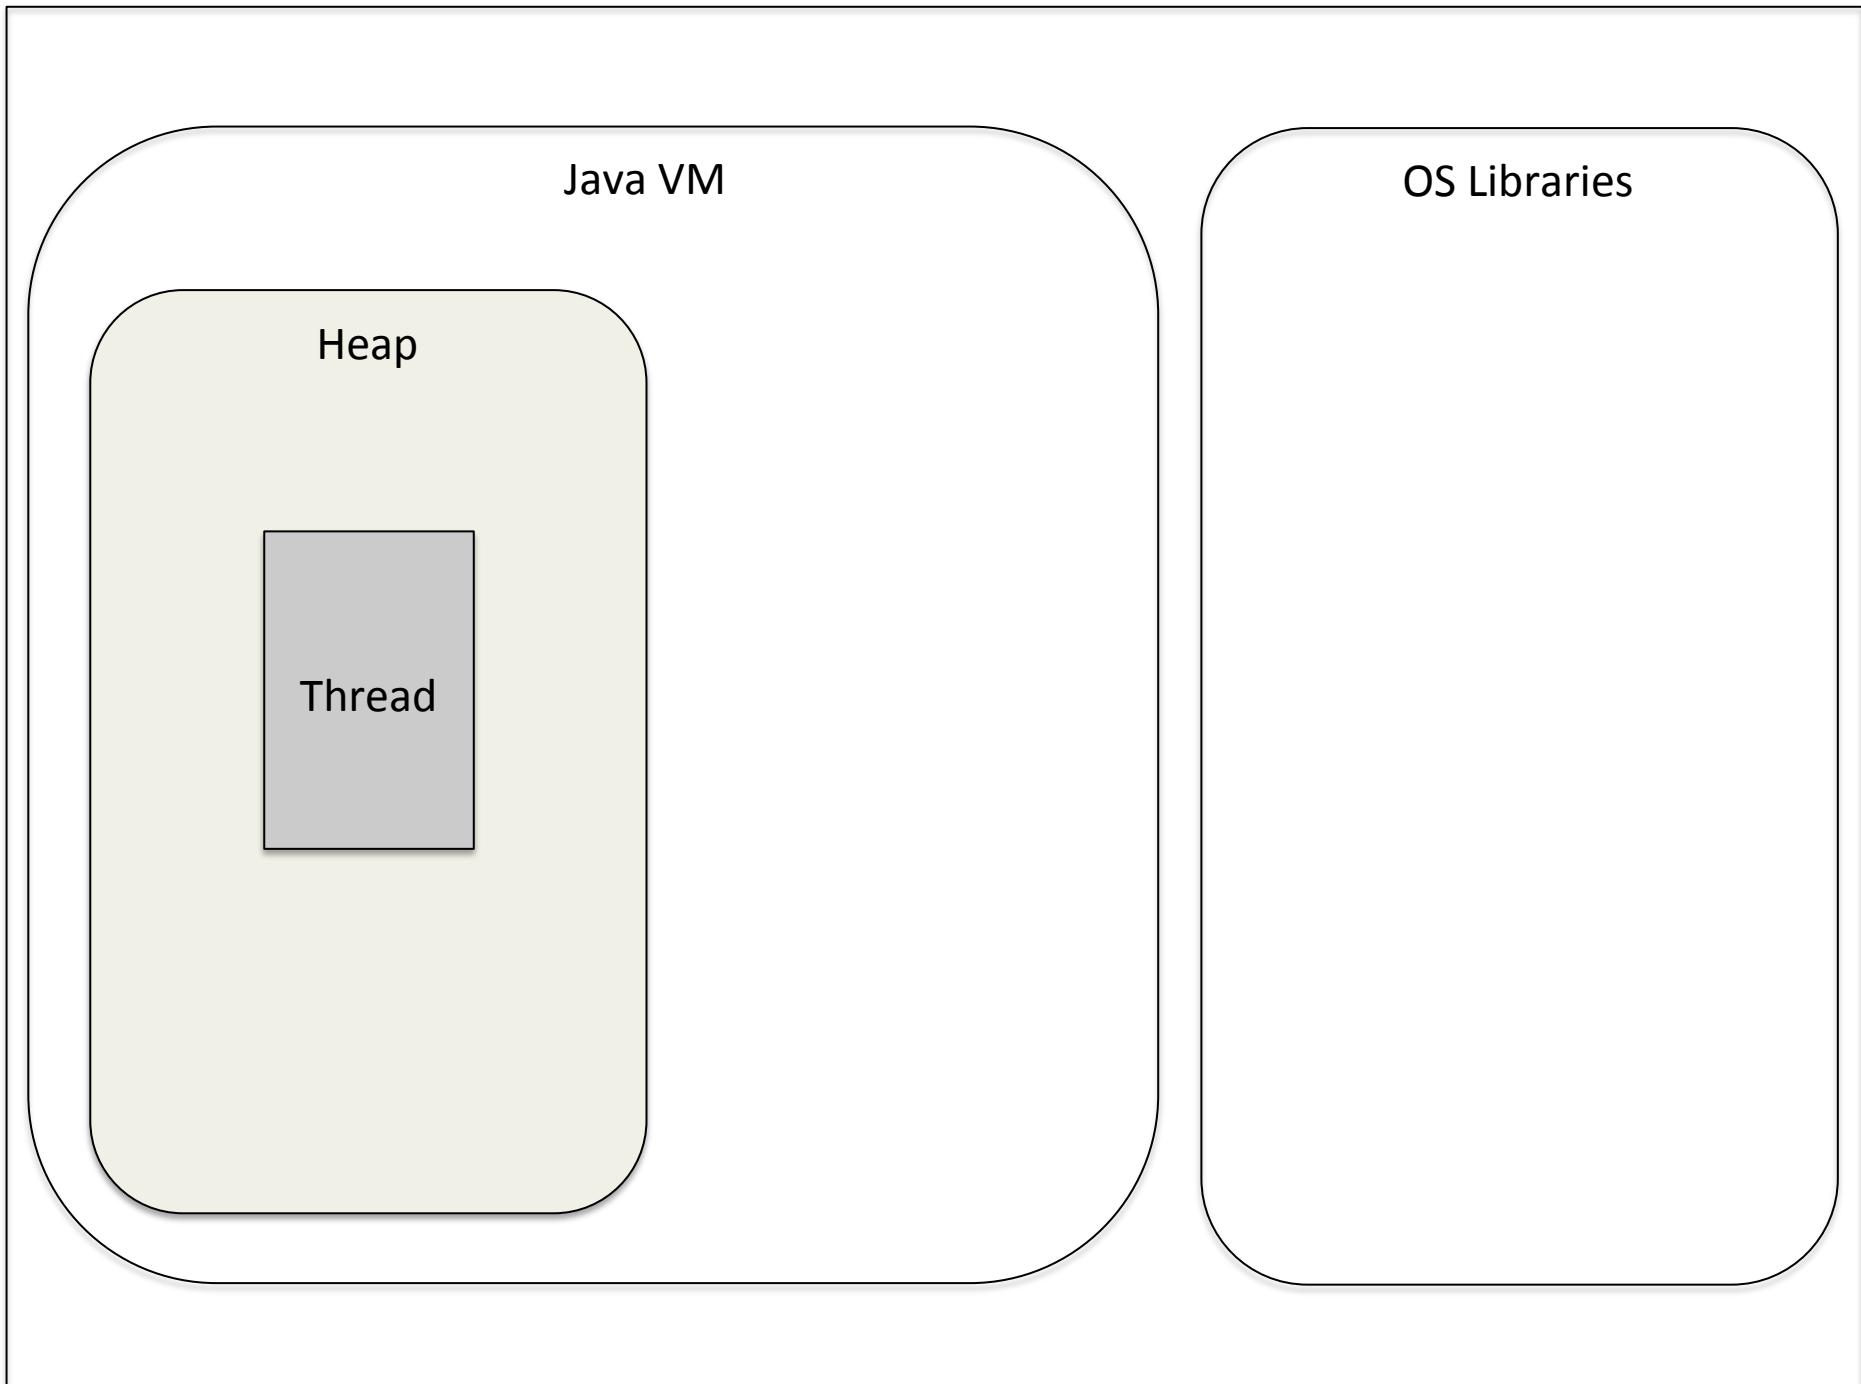
Java Threading

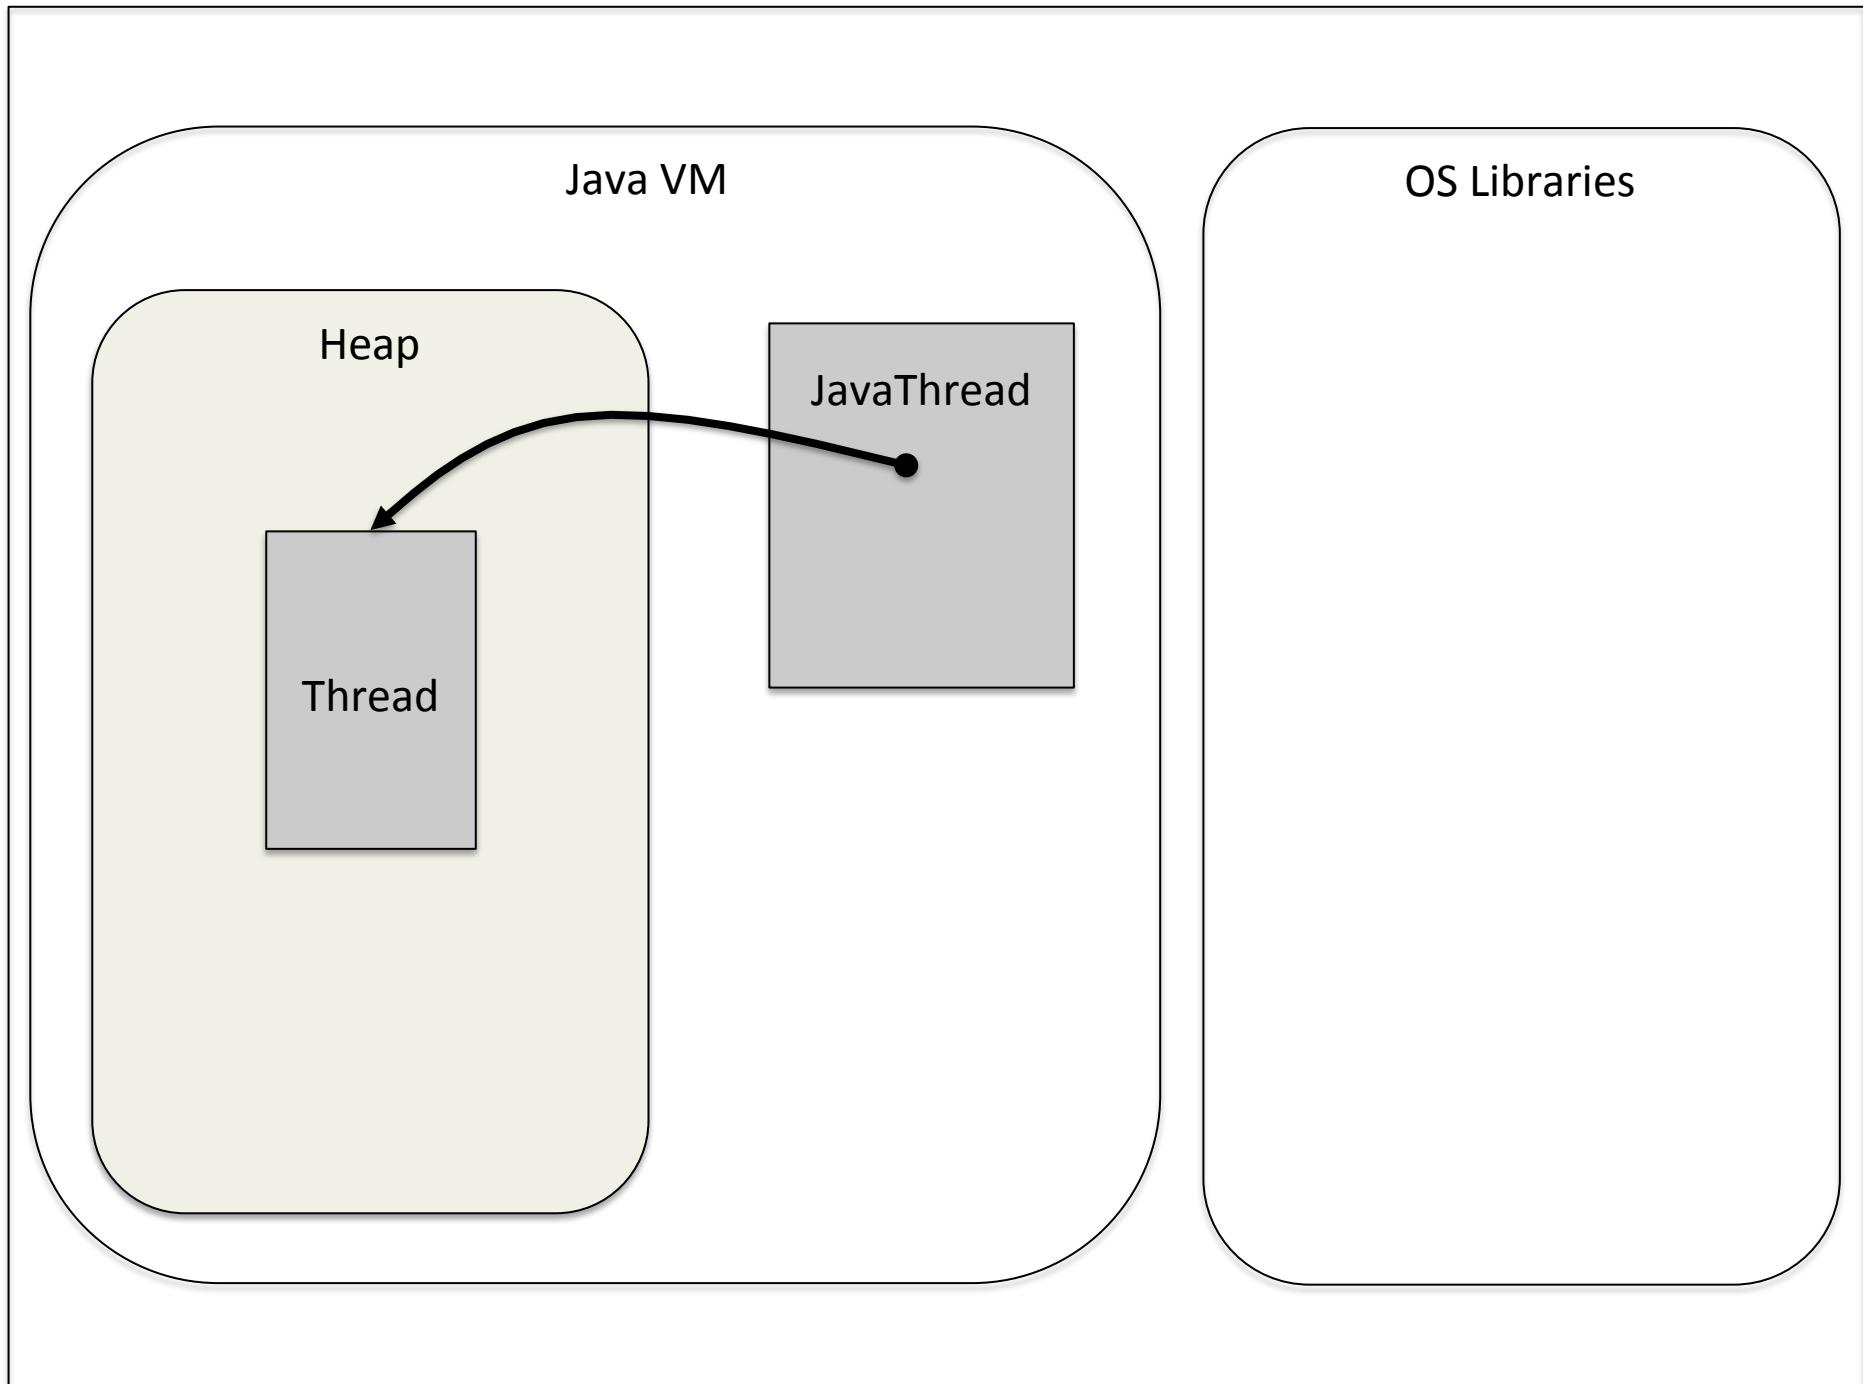
- Basic model
 - Thread class
 - Runnable and Callable interfaces
- Thread created like any other Java object
 - Call start method to launch
- Various other thread operations
 - join, yield, sleep
 - All described in the language spec

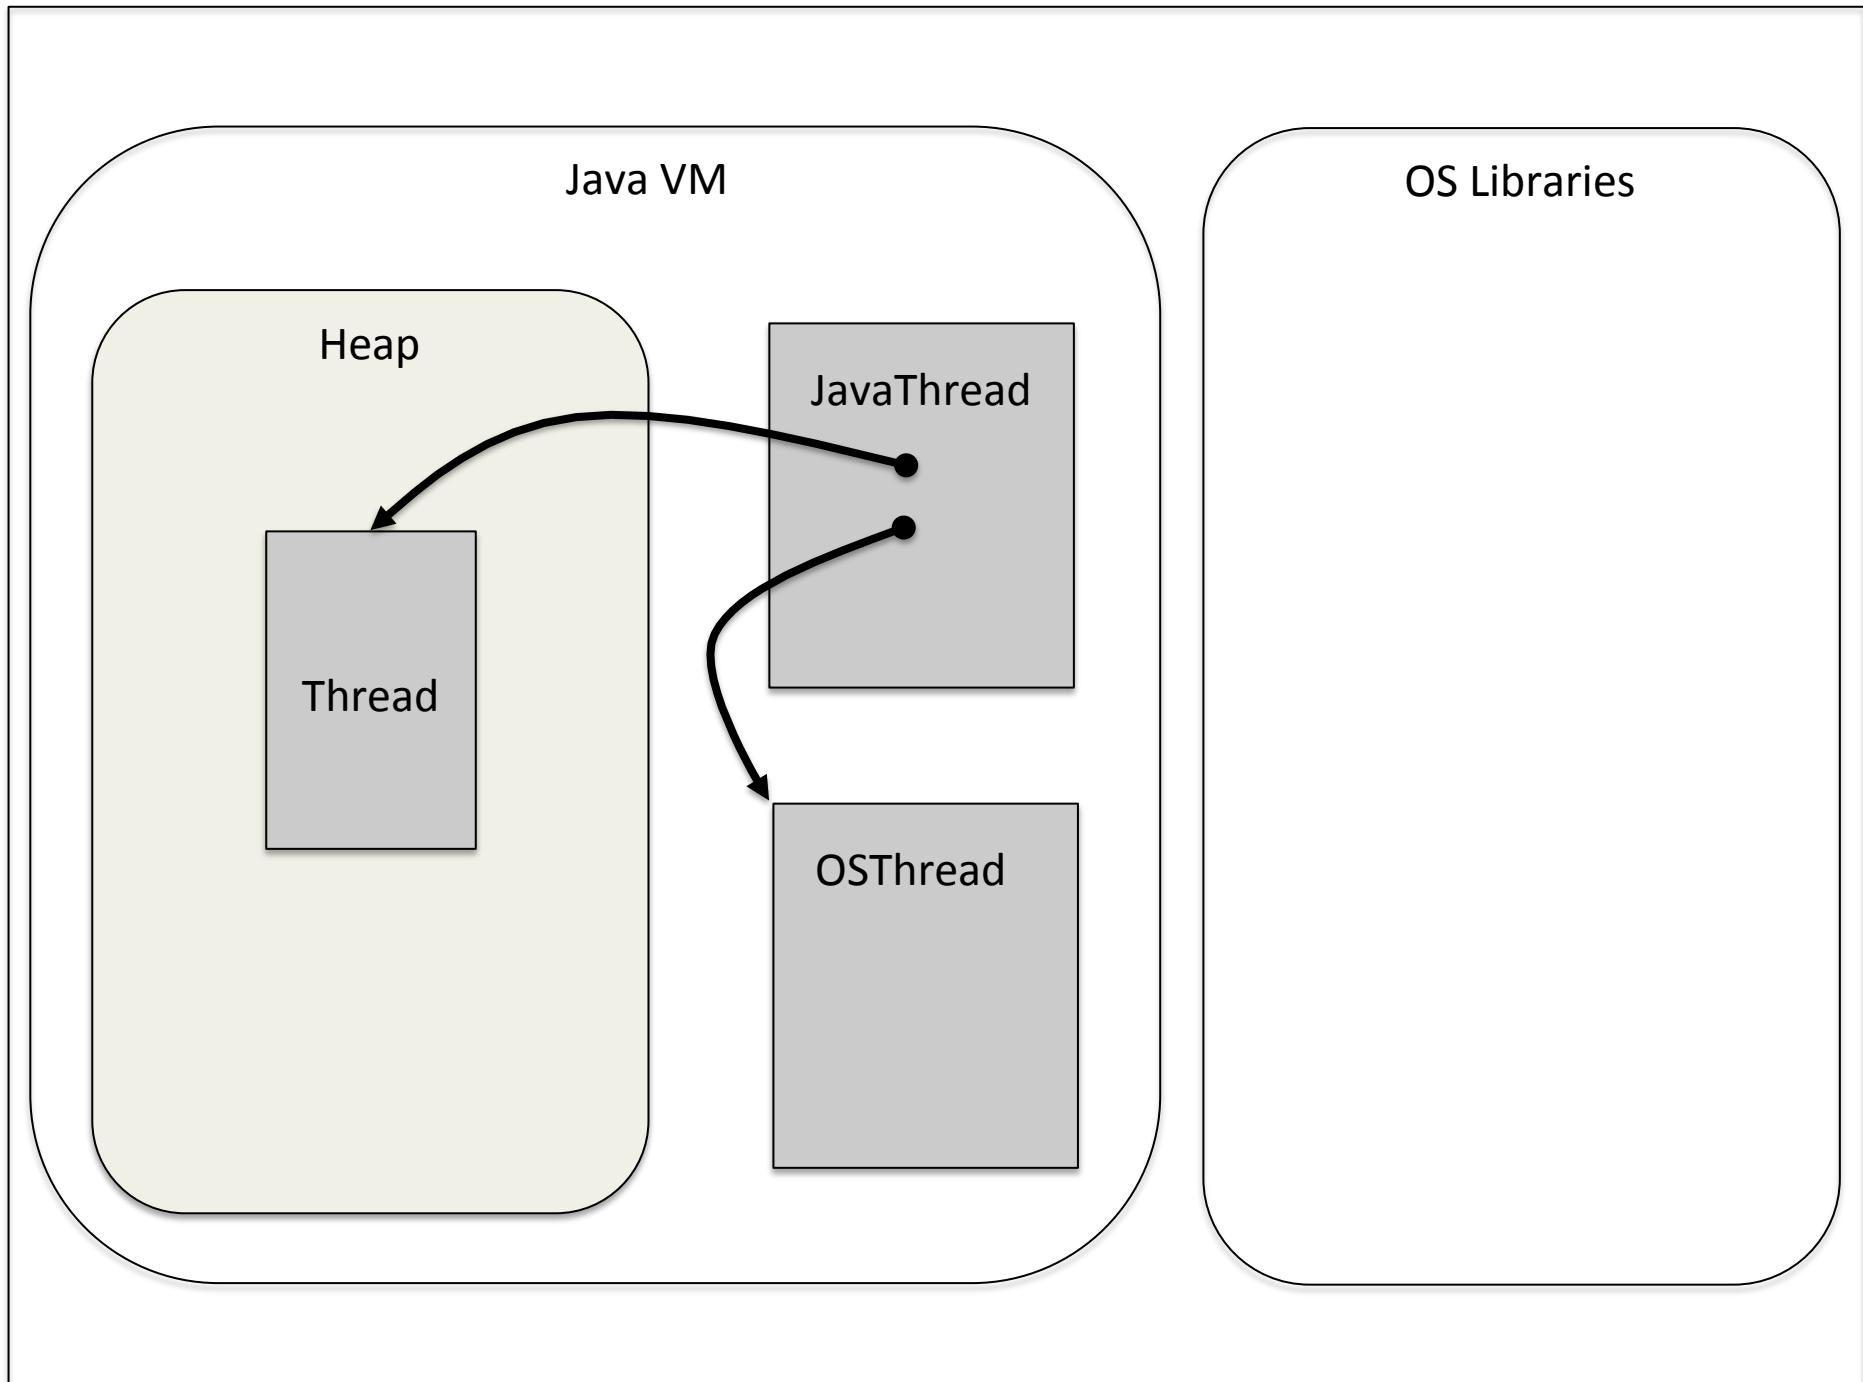
Java VM

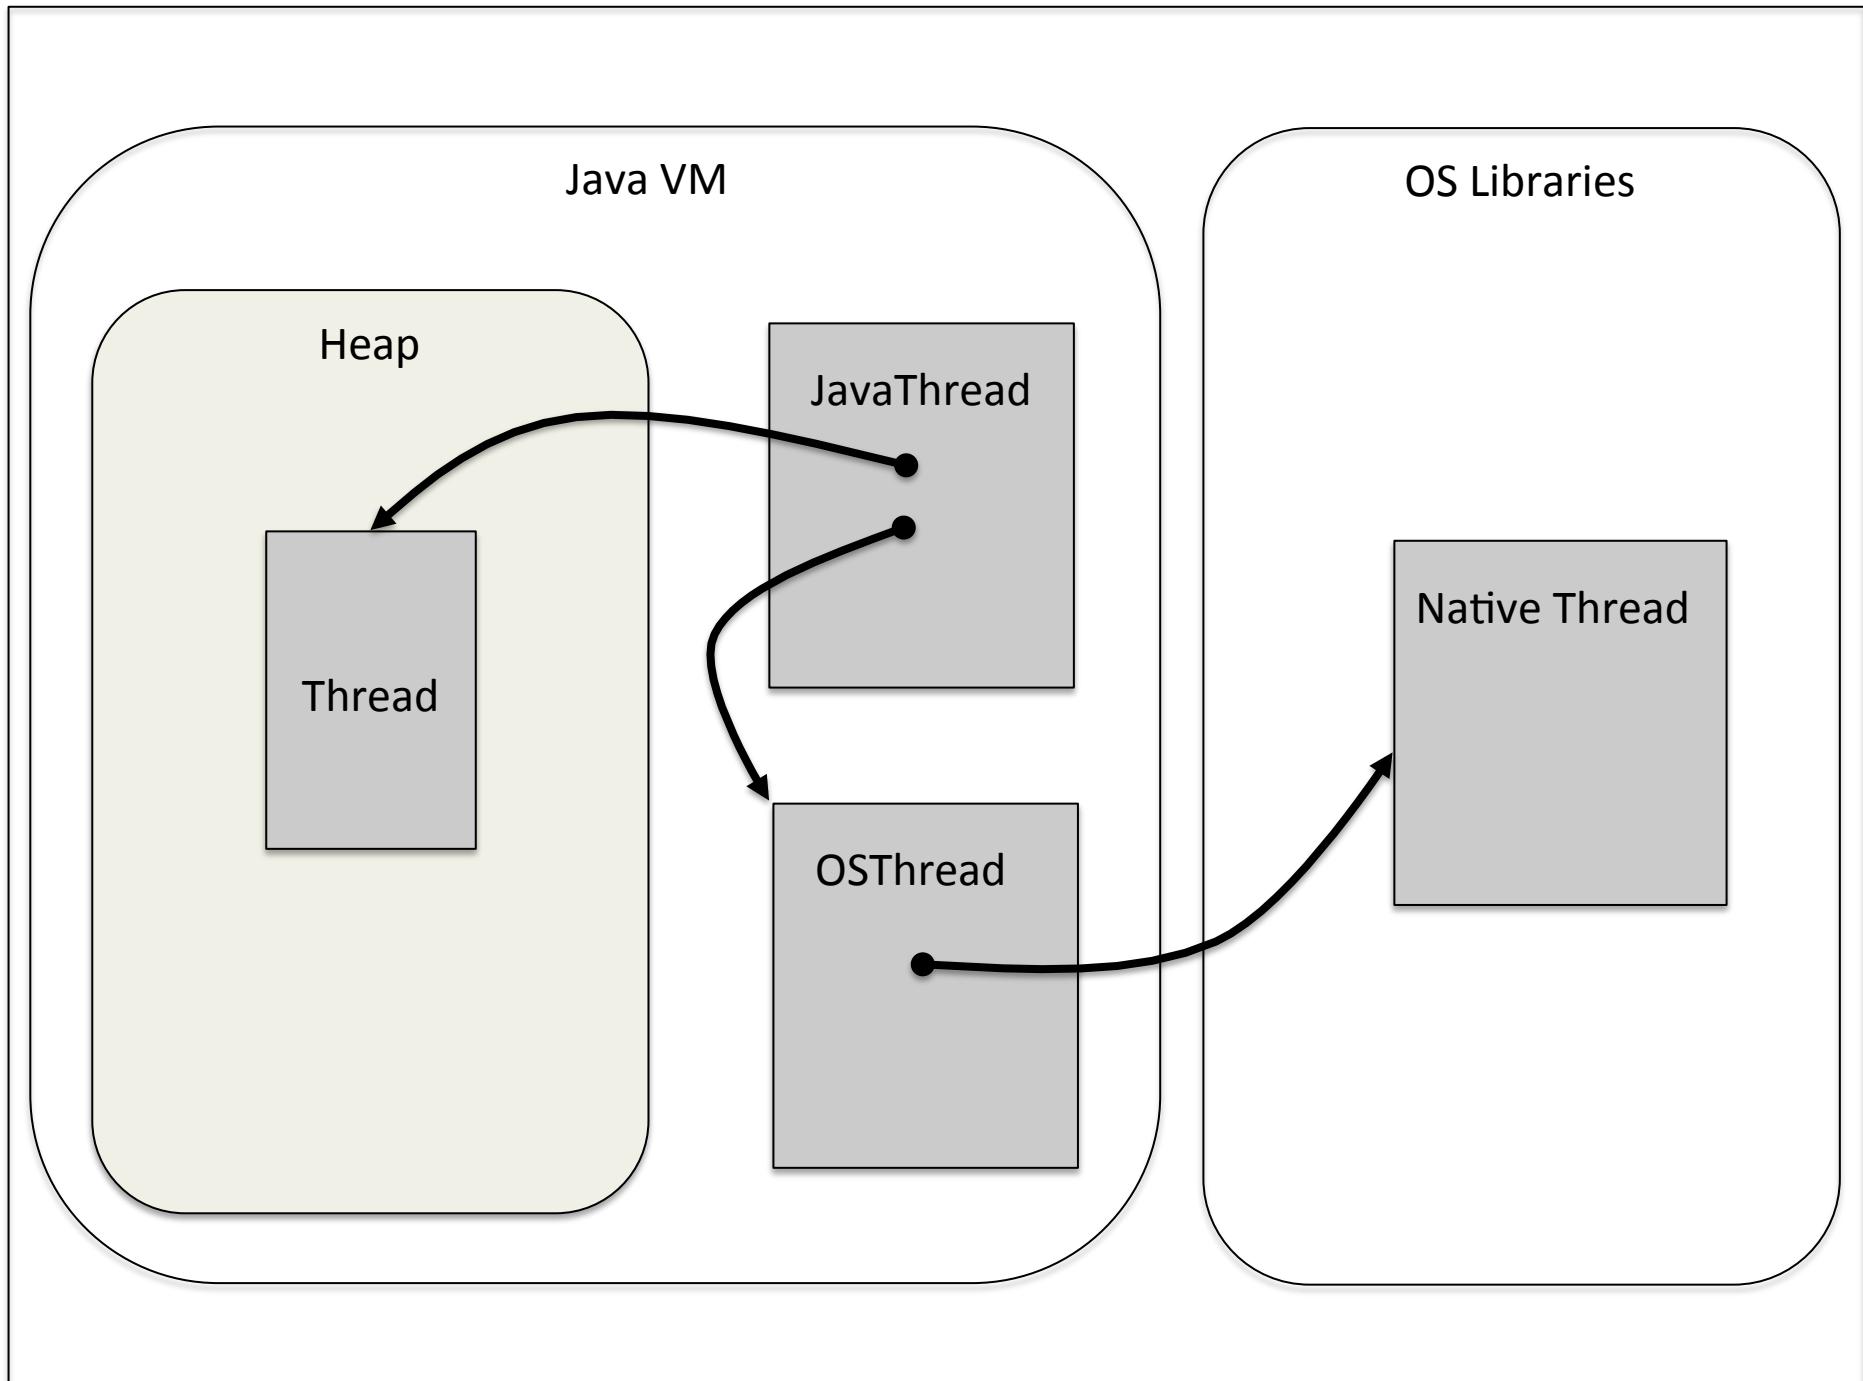
Heap











Scheduling

- Scheduling algorithm is implementation specific
 - Depends greatly on the platform
 - Goal is to present a unified model
- Scheduling can be delegated to the operating system
 - Hotspot maps Java threads directly to OS threads
- VM may alternatively manage threads itself
 - More work for the VM implementer
 - Finer control of scheduling

Green Threads

- Some systems don't support multithreading
 - Mostly embedded devices these days
- VM can use a single OS thread
 - Manually schedule VM threads
- Generally uses straight priority scheduling
 - Highest priority thread runs to completion
 - Thread can be interrupted by higher priority thread
 - Possibility of starvation
- Can't rely on priorities for mutual exclusion

Native Threads

- VM maps `JavaThread` to `OSThread` directly
 - Uses OS-level priority mechanism
- OS has better knowledge of the platform
 - Can optimize scheduling around system events
 - May predict long-latency events
- VM has no control over thread switching

Time Slicing

- OS multiplexes threads to CPU cores
 - There are normally more threads than cores
- Scheduler uses time slicing to share CPU
 - Thread runs for a defined quantum of time
 - Scheduling decision made at the end of that time
 - Handles threads with equal priorities
- More time slicing, worse cache behavior

Hybrid Approach

- VM can map OS to Java threads as one-to-many
 - Run green thread scheduling on each OS thread
- Can group related threads on an OS thread
 - Application thread on one, VM on another
- Green threads can optionally use time slicing
 - Implementation details trickier
 - May have to account for OS time slicing as well

Safepointing

- Time slicing can cause problems with safepointing
 - All non-native threads need to be at a safepoint
 - Can't guarantee when threads were preempted
- Adds an additional step to safepointing
 - Let a thread run to a safepoint
 - Set it as suspended and yield the processor
 - Next highest priority thread takes over
- Need to restore suspended status afterwards
 - Track thread's original status