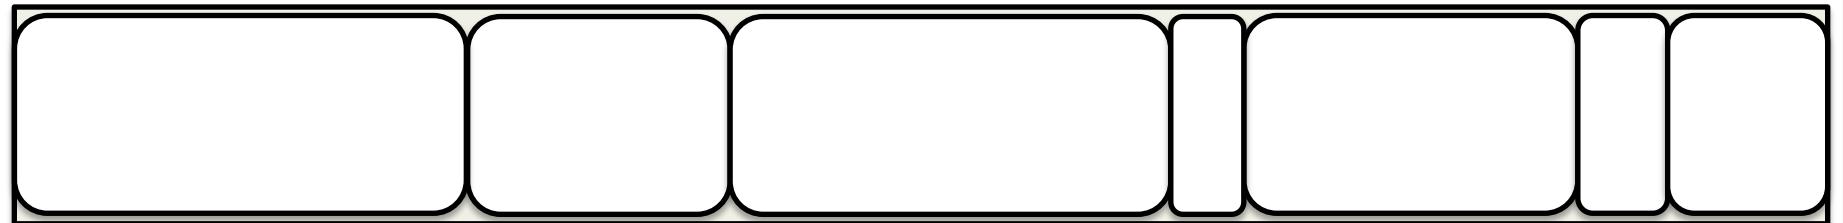
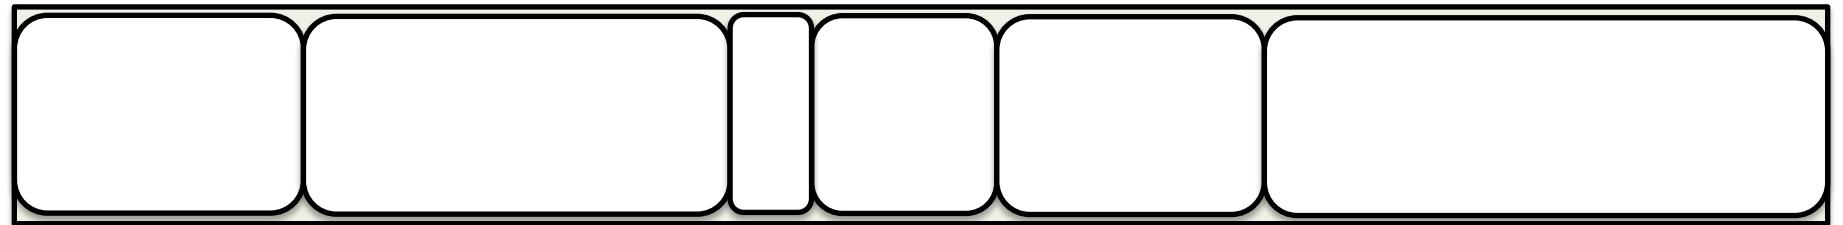
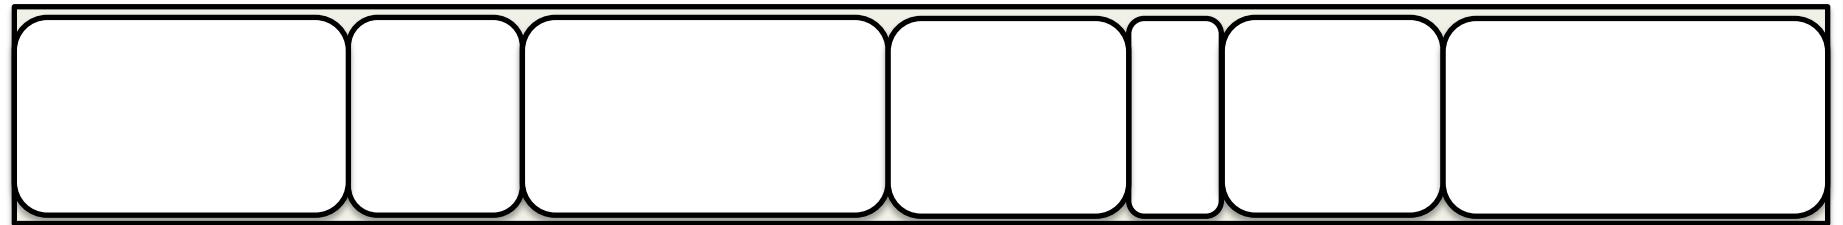


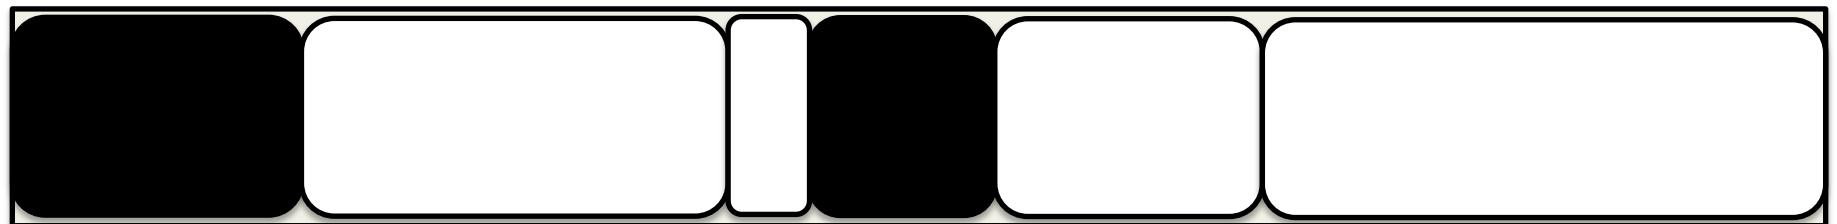
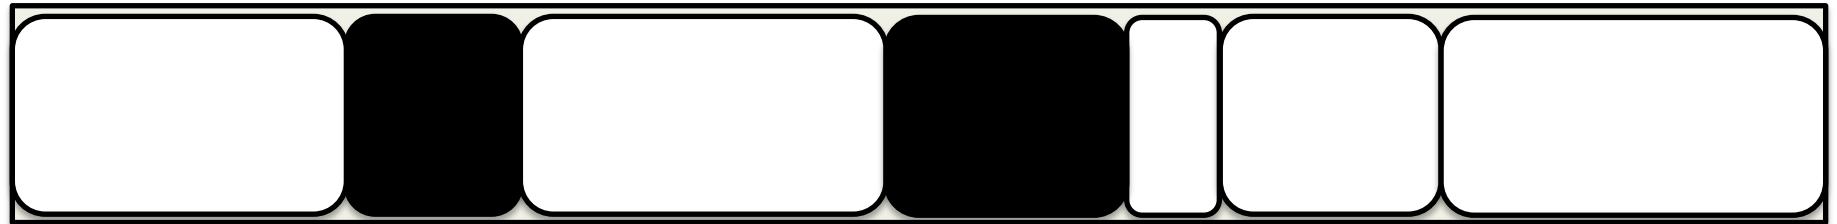
In the Last Week

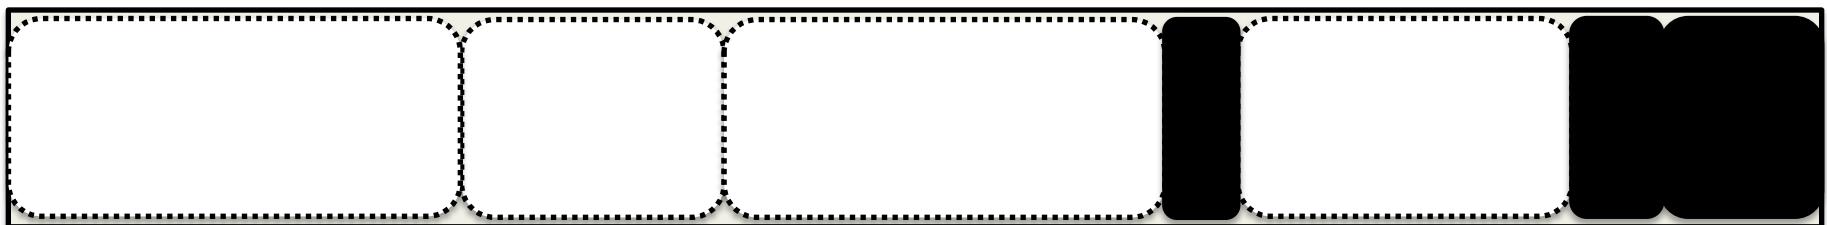
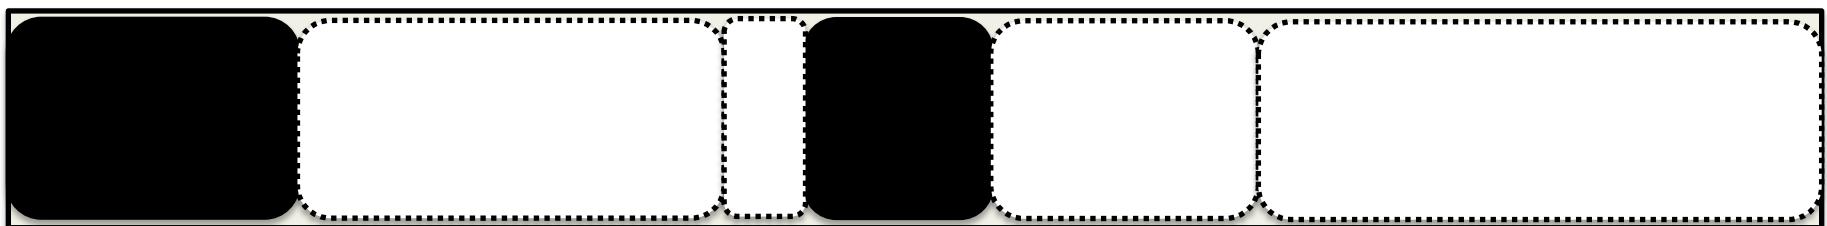
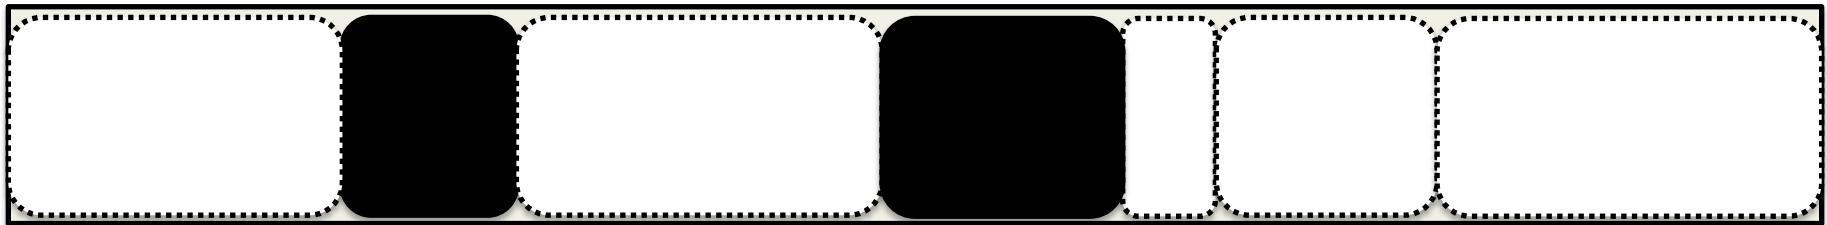
- Fairly quiet week
 - A few more suggestions for special topics
 - Some questions on the final project
- Going beyond the final project specification
- Objects and holes in the non-contiguous allocator

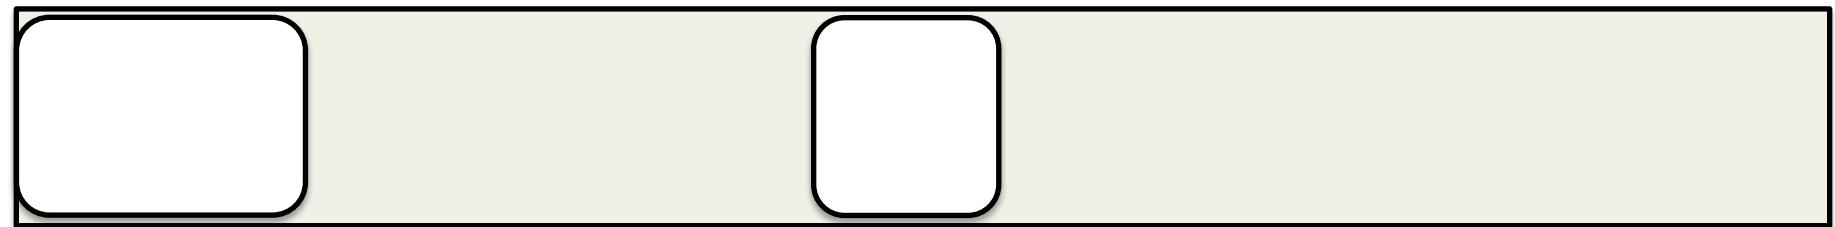
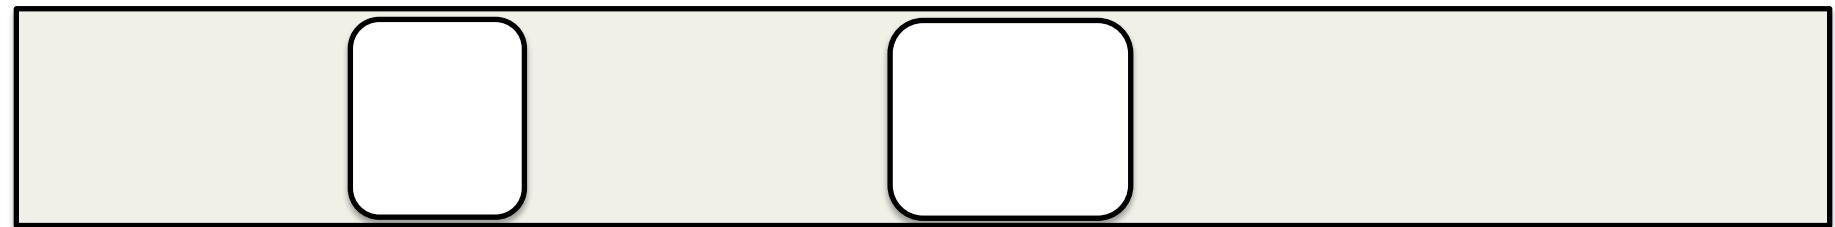
Objects and Holes

- Non-moving collectors create fragmentation
 - Solution is to compact the heap
- Over time fragmentation gets worse
 - Need to be able to fill in the gaps between objects
 - Recall the various allocation algorithms
- It can be useful to think of holes as a separate item
 - List of free locations to allocate objects





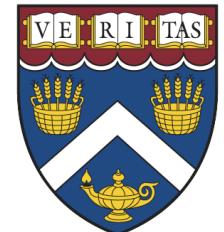


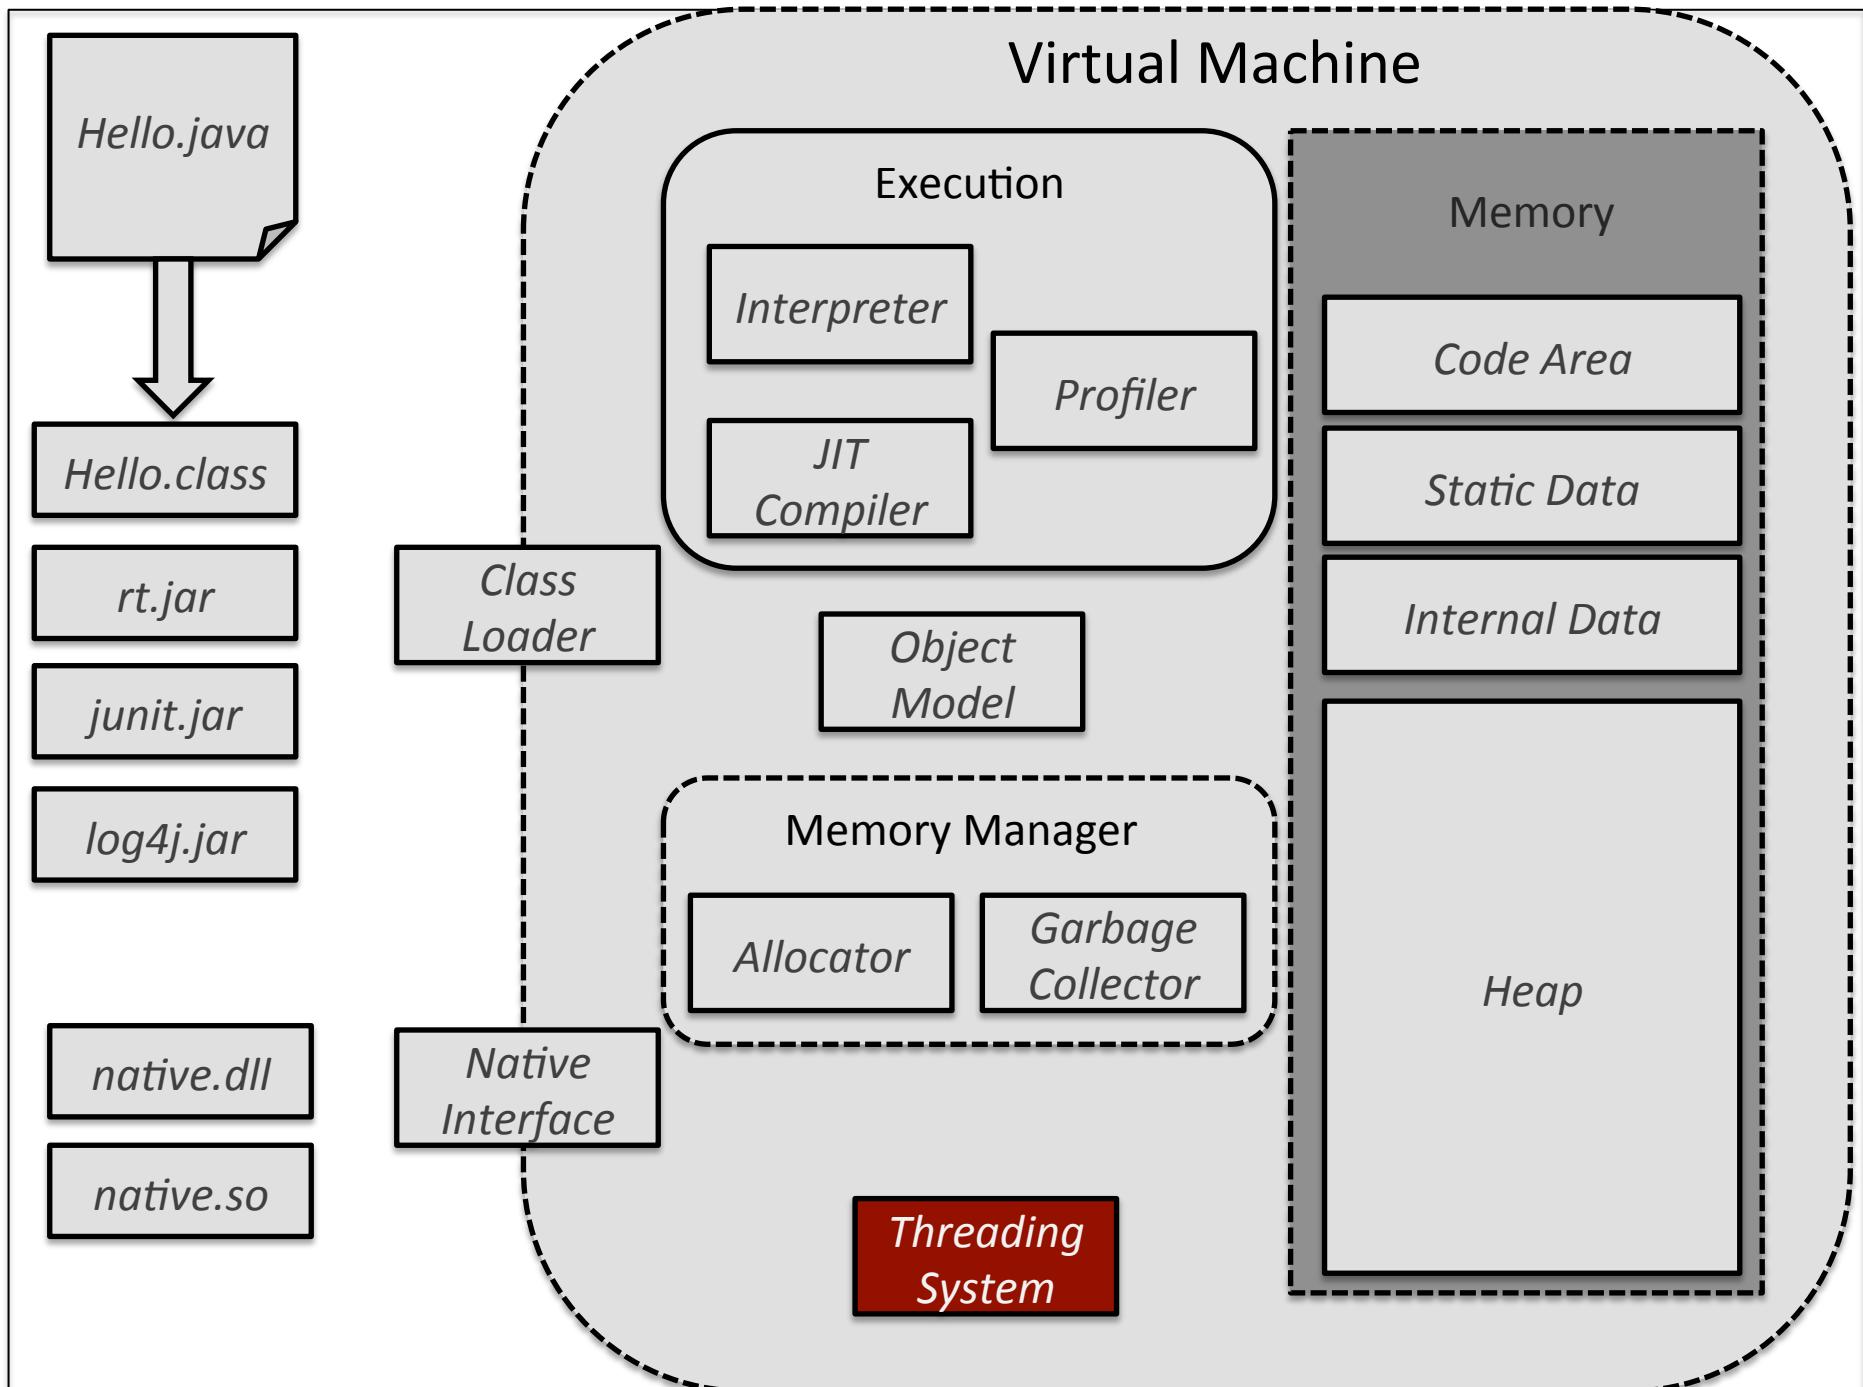


Tracking Space

- It is not necessary to track holes separately
 - You can just track objects allocated on the heap
 - The holes can be derived from the objects
- Allocation is much faster with a free list
 - Just scan through the list to find first fit
 - How real-world allocators are implemented
- Maintaining list of holes adds complication

VM Threads





Memory Fences

- Processor-level synchronization instruction
 - One of the most expensive in the ISA
 - Most processors have finer-grained instructions
- Guarantees memory ordering before and after
 - Use fences to synchronize on monitor operations
- Flushes cached versions of variables

Volatile

- `volatile` variables are never cached locally
 - All writes conceptually go to main memory
 - All threads see the same value immediately
- Useful as lightweight synchronization
 - Fewer features than full monitors
 - May perform better than critical sections
- Some semantic traps for the unwary
 - Incrementing a variable in two threads really two actions

Concurrency and Parallelism

- Two terms often used synonymously
 - Useful to maintain a distinction
- Concurrent threads interact with others
 - Must synchronize between threads
- Parallel threads have limited or no interaction
 - Very little synchronization
 - Best performance advantage

Daemon Threads

- Threads in the VM fall into two groups
 - Daemon threads normally belong to the VM
 - Non-daemon threads belong to the application
 - Applications can turn threads into daemons
- Daemon threads don't normally terminate
 - The VM runs until all non-daemon threads end
 - Daemon threads abandoned at shut-down
 - No cleanup operations happen

Internal Daemon Threads

- Periodic task thread
 - Implements scheduled events and timer interrupts
- Signal dispatcher thread
 - Monitors for OS-level signals
- Compilation and garbage collection
- VMThread
 - Main dispatcher for system operations

VMThread

- Single instance in a VM
 - Runs background tasks separately from Java threads
- VM tasks are added to a synchronized queue
 - VMThread waits for a change to the queue
 - Does safepointing based on the operation
- Single VM thread simplifies implementation
 - Don't need to worry about concurrency

VMThread

- Around fifty possible tasks
 - Some heavyweight VM operations
 - Many of them handlers for rare cases
- Four modes for tasks to run
 - Safepoint/no safepoint
 - Blocking/non blocking

VMThread Operations

- Logging and reporting
- Verification
- Deoptimize code
- Redefine classes
- Compute stack traces
- Revoke biased locks

Compiler Threads

- Compilation can happen in parallel
 - Doesn't block the Java threads
 - Interpreter can continue to execute a method
- Most code replacement doesn't require a safepoint
 - Use the compiled code for the next execution
- On Stack Replacement is a heavier operation
 - Task scheduled for the VMThread

Memory Model and JIT Compiler

- Java concurrency has implications for JIT
 - Limits code reordering
 - Limits memory prefetching
- **final** fields treated specially
 - Initial write in constructor comes before object is shared
 - Could be reordered in inlined code
 - Afterwards, don't have to reload the final field

Profiler Sampling

- Some VMs use a thread for their profiler
 - Periodic thread samples stack state
 - Builds up map of hot methods
- Lower overhead than measuring every call
 - Background thread doesn't block Java thread
- Results can be less accurate
 - More likely to miss short methods

Multi-Thread Memory Management

- Garbage collectors take advantage of threads
 - Can run far faster than single-threaded GC
 - Implementation is more complex
 - Consider trade-off between throughput and pauses
- Allocation introduces contention
- GC can be split across multiple threads
 - Add more phases to the garbage collector algorithm
 - Threads do different work at different points

Parallel and Concurrent GC

- Significant distinction in garbage collection
- Parallel GC splits the collection between threads
 - Still stop all running Java threads
- Concurrent GC interleaves with the mutator
 - Implementation can be very complex
- Modern garbage collectors are both
 - Concurrent phases interact with mutators
 - Parallel phases speed up critical sections

Parallel GC Algorithms

- Most effective when there is little coordination
 - Allow majority of work to happen independently
 - Scalability limited by the number of threads
- Many algorithms use work stealing
 - Each thread maintains a list of tasks
 - When a thread's list is empty, look to other thread's lists
- Work stealing can introduce contention
 - Work queues are a vital place for performance
 - Can minimize by stealing from the back of the queue

Allocation by Concurrent Threads

- So far allocation has happened in the nursery
 - Bump pointer algorithm
- Allocation is fast
 - Load the current allocation pointer
 - Add the size of the object to allocate
 - Compare with the limit pointer
 - Update the allocation pointer

Thread 1

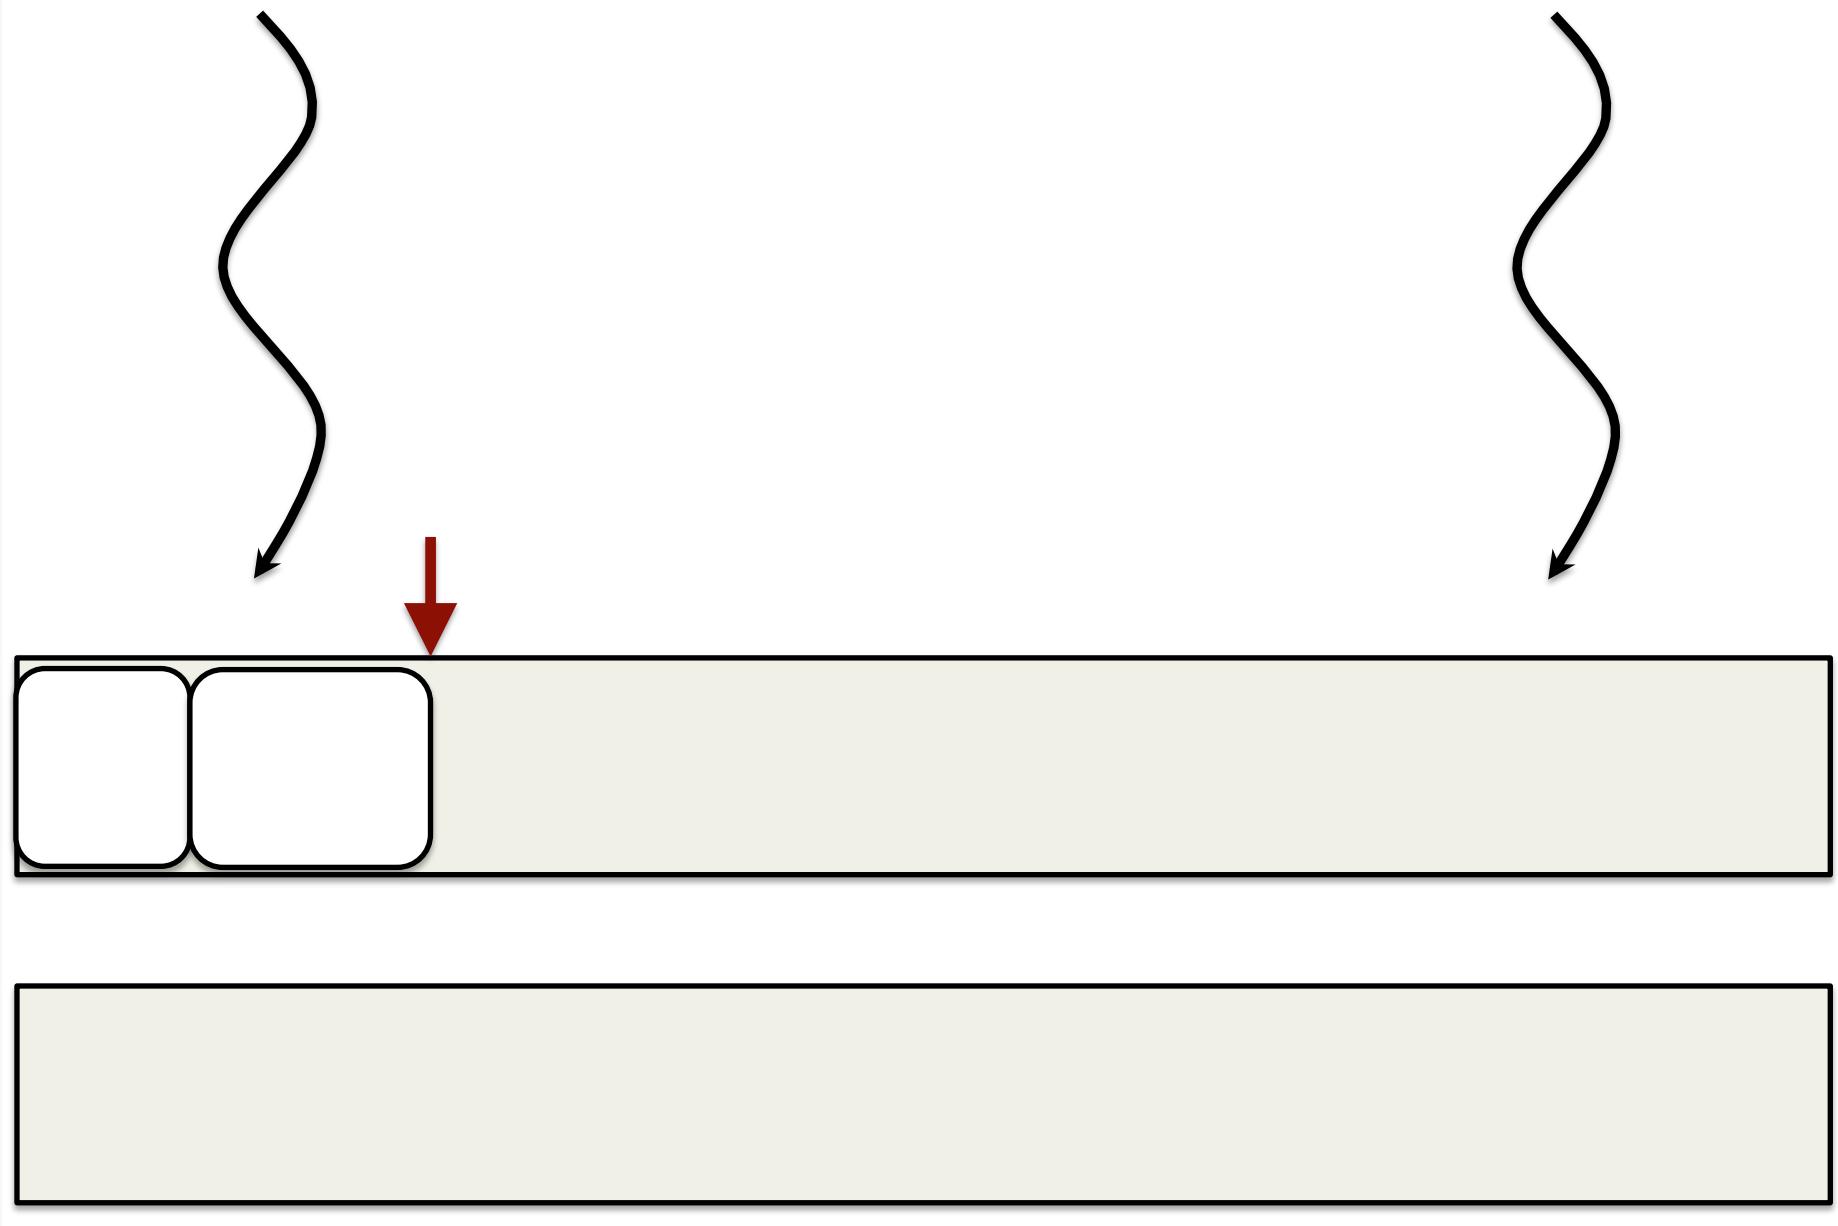


Thread 2



Thread 1

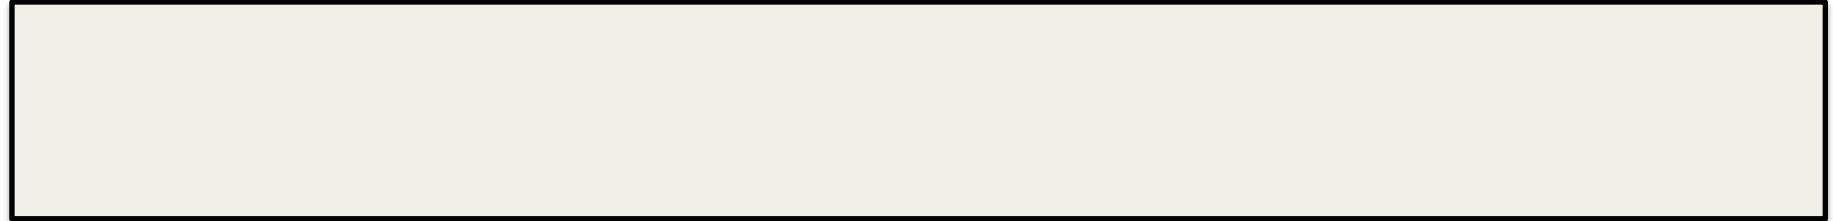
Thread 2



Thread 1

Thread 2

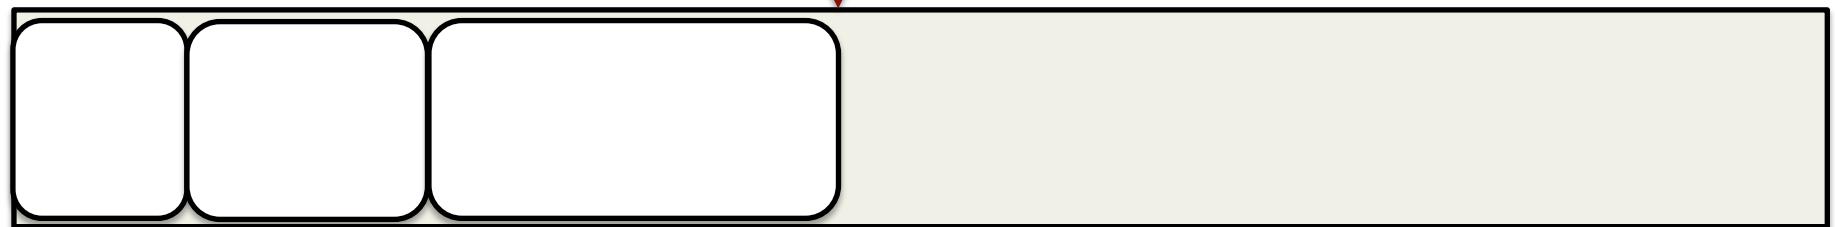
allocate



Thread 1

Thread 2

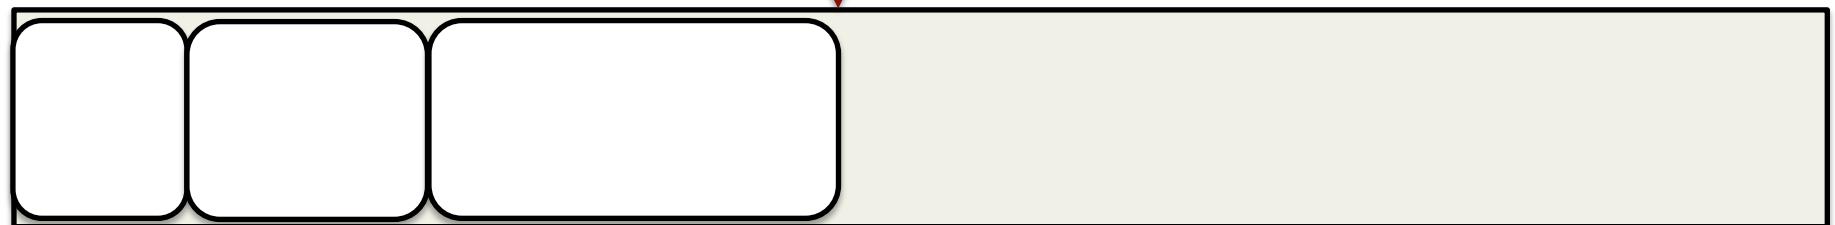
allocate



Thread 1

Thread 2

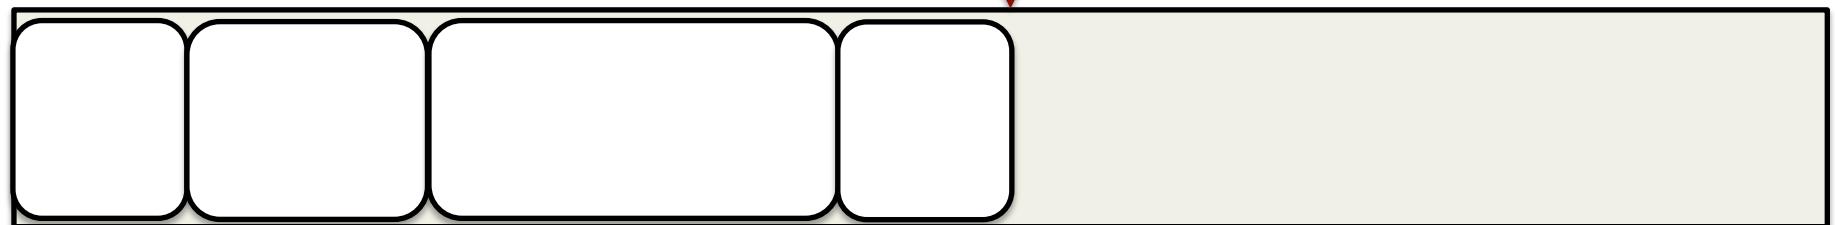
allocate



Thread 1

Thread 2

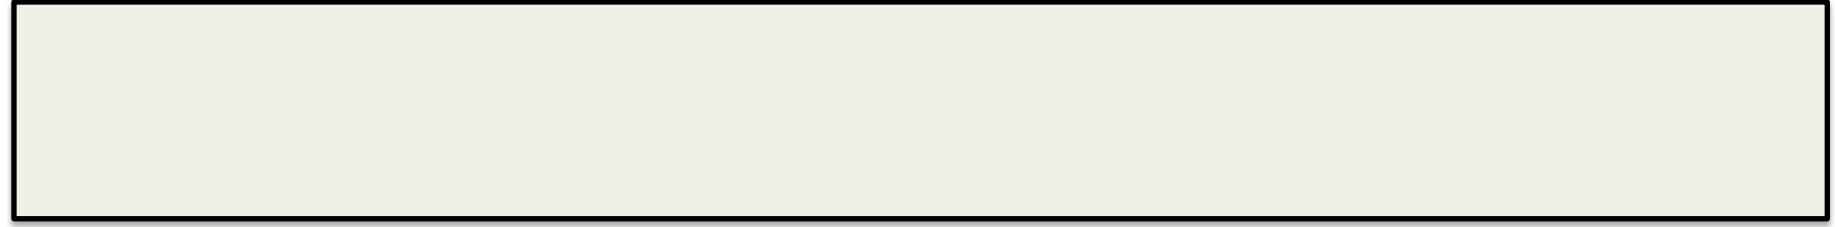
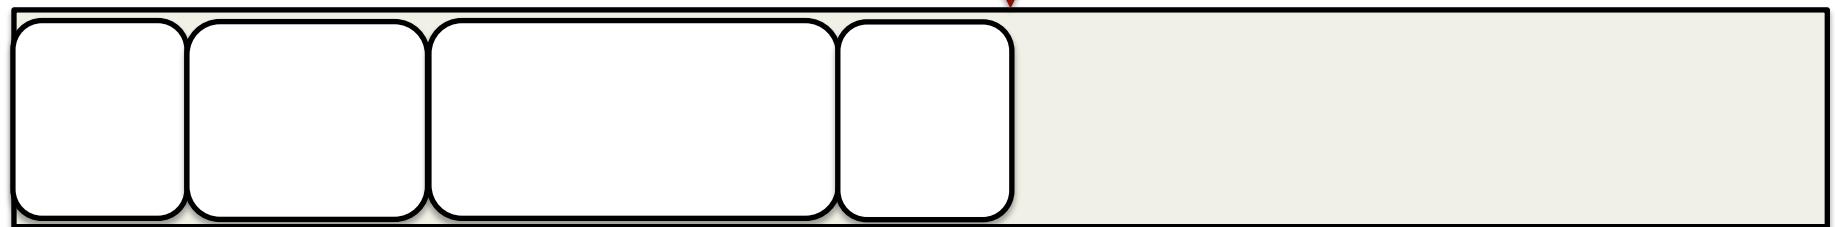
allocate



Thread 1

Thread 2

allocate

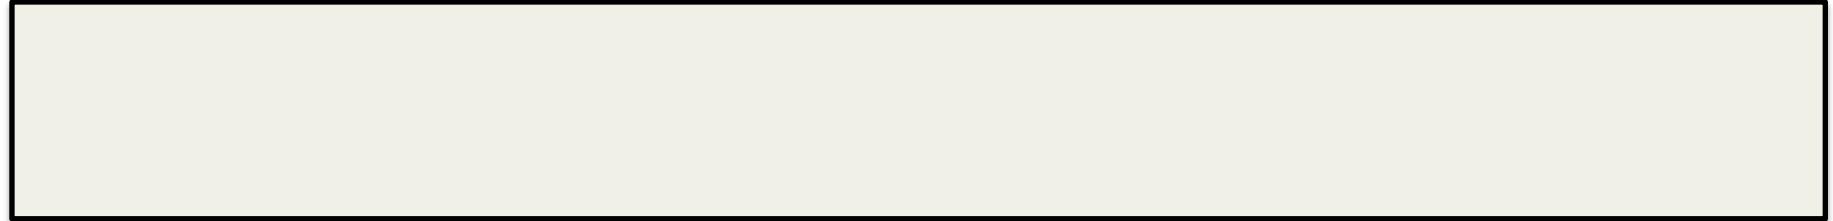
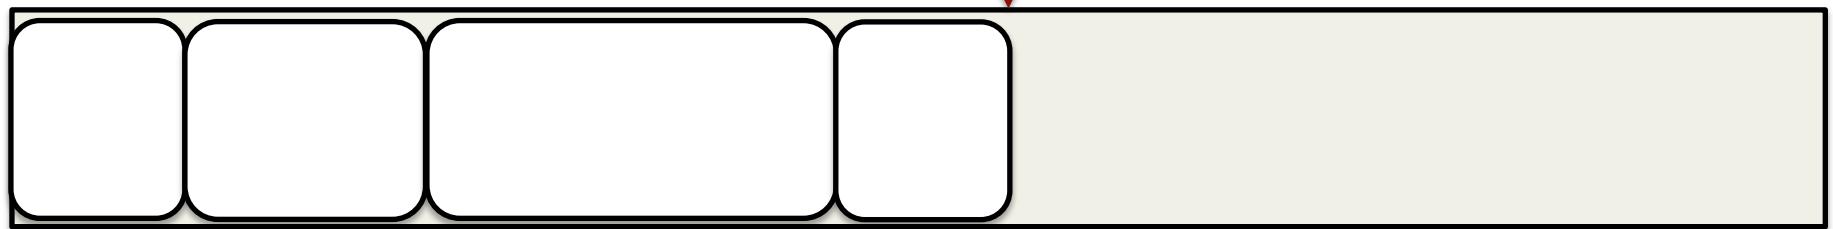


Thread 1



```
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)
```

Thread 2

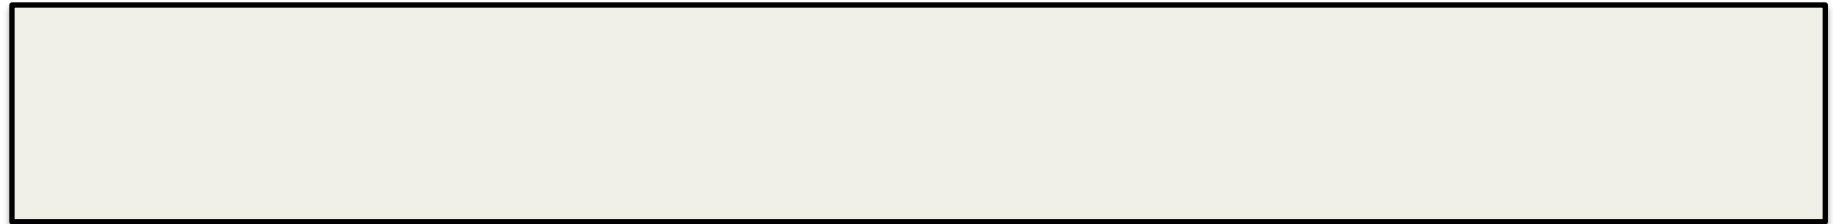
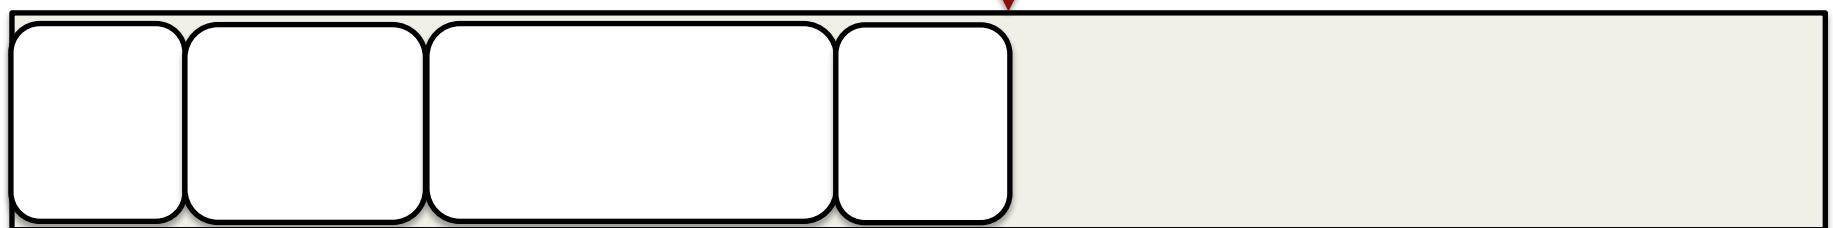


Thread 1

```
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)
```

Thread 2

```
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)
```



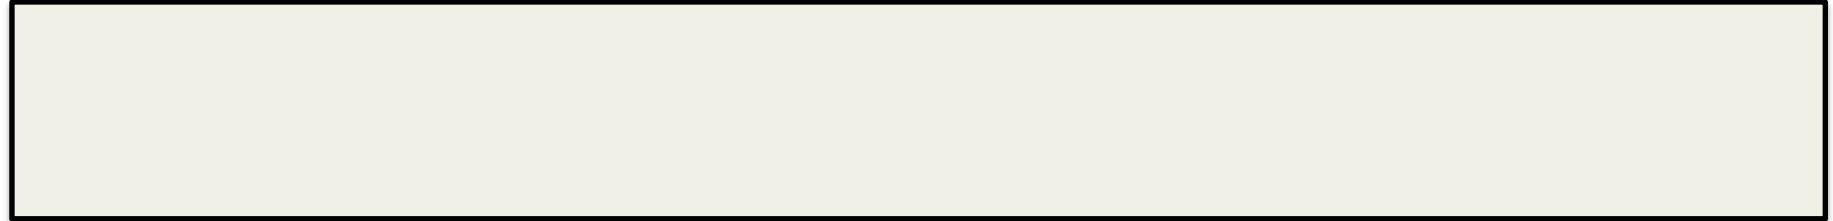
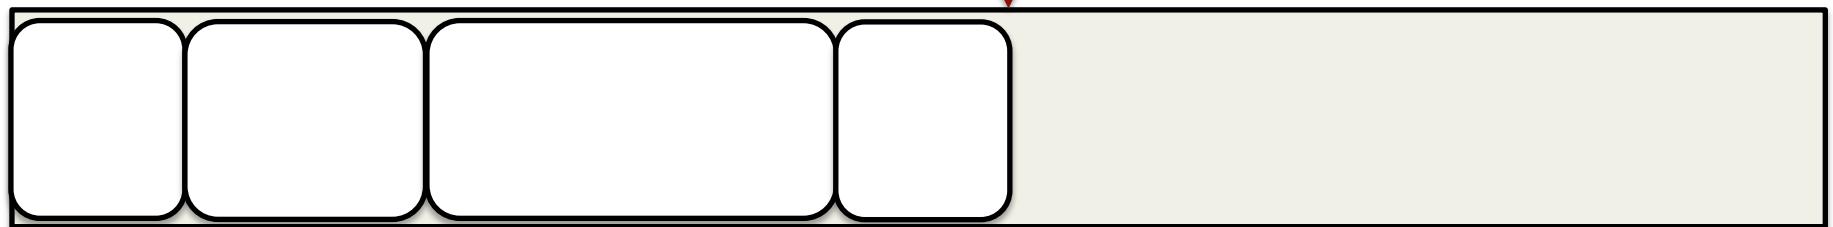
Thread 1

`ptr = getPtr()`

`newPtr = ptr + size`
`setPtr(newPtr)`

Thread 2

`ptr = getPtr()`
`newPtr = ptr + size`
`setPtr(newPtr)`



Thread 1

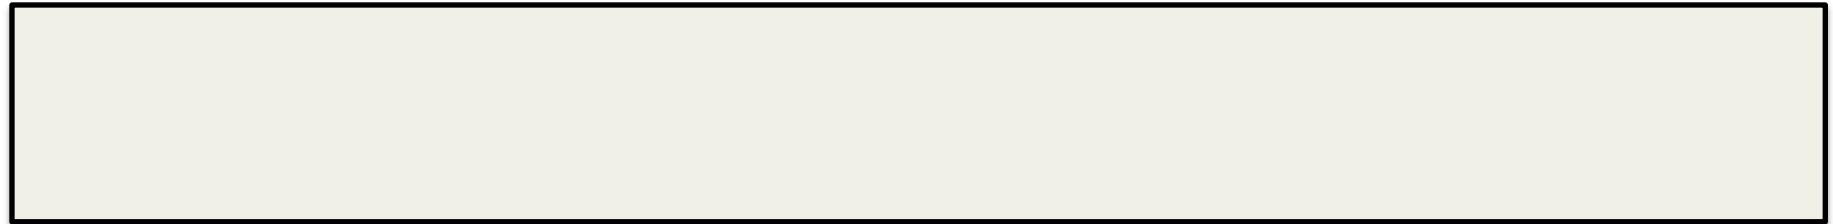
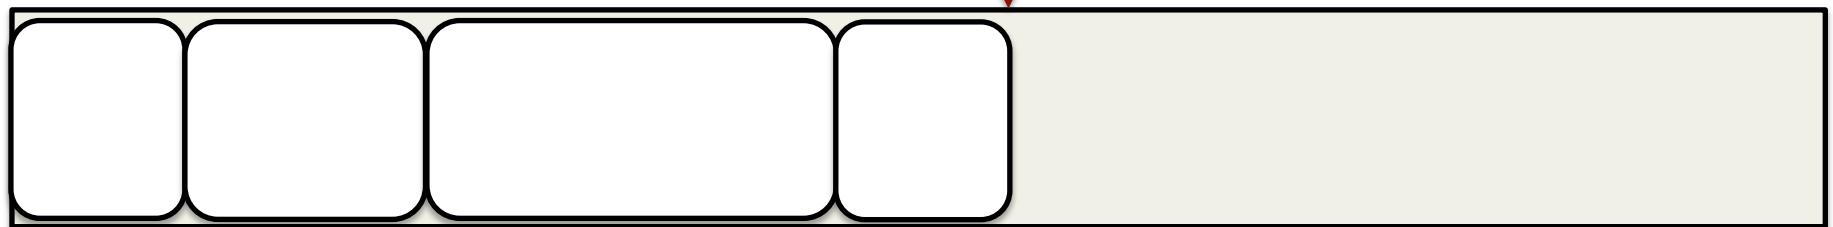
`ptr = getPtr()`



`newPtr = ptr + size
setPtr(newPtr)`

Thread 2

`ptr = getPtr()
newPtr = ptr + size
setPtr(newPtr)`



Thread 1

`ptr = getPtr()`

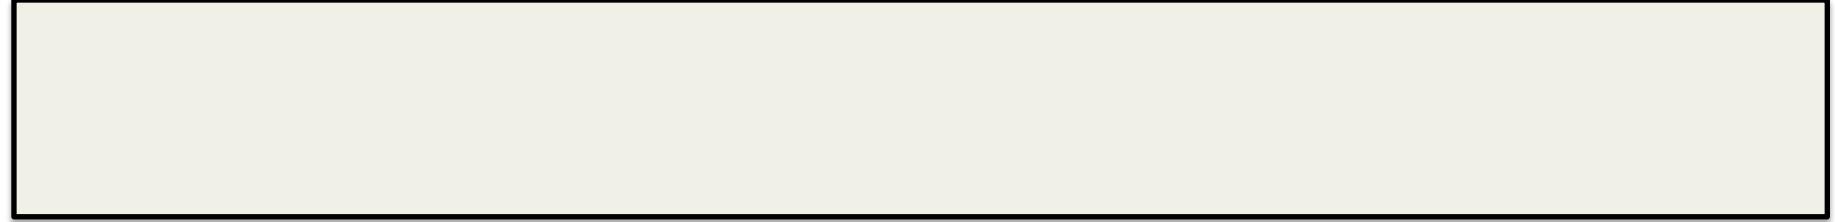
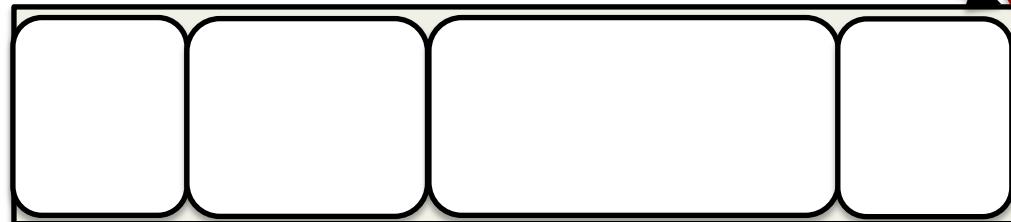
`ptr`



`newPtr = ptr + size`
`setPtr(newPtr)`

Thread 2

`ptr = getPtr()`
`newPtr = ptr + size`
`setPtr(newPtr)`



Thread 1

`ptr = getPtr()`

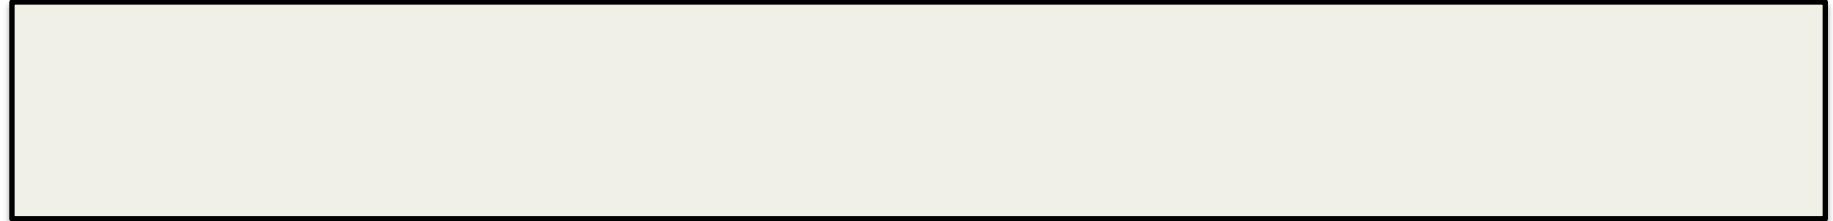
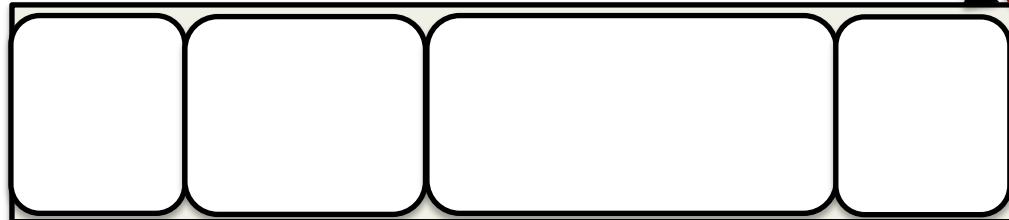
`ptr`



Thread 2

`ptr = getPtr()`
`newPtr = ptr + size`
`setPtr(newPtr)`

`ptr`



Thread 1

$\text{ptr} = \text{getPtr}()$

ptr



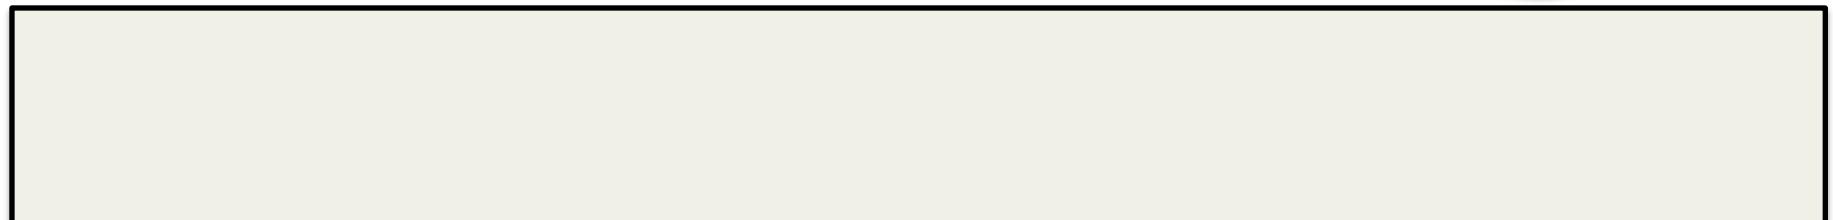
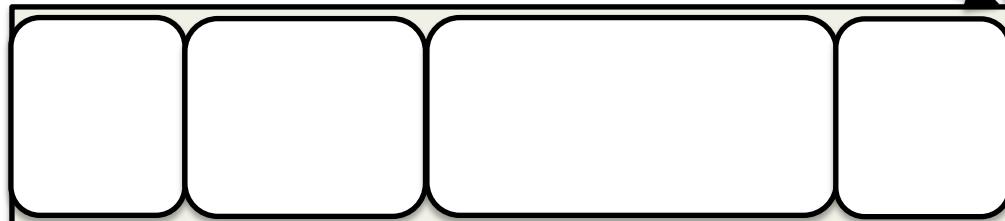
$\text{newPtr} = \text{ptr} + \text{size}$
 $\text{setPtr}(\text{newPtr})$

Thread 2

$\text{ptr} = \text{getPtr}()$
 $\text{newPtr} = \text{ptr} + \text{size}$
 $\text{setPtr}(\text{newPtr})$

newPtr

ptr



Thread 1

`ptr = getPtr()`

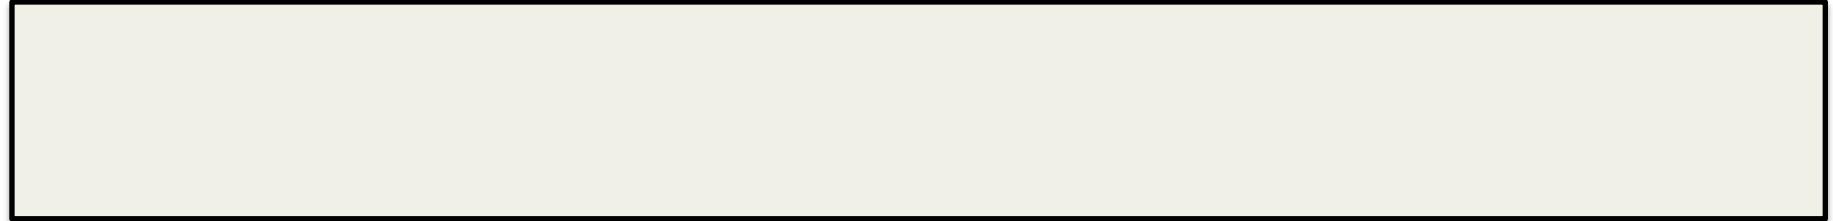
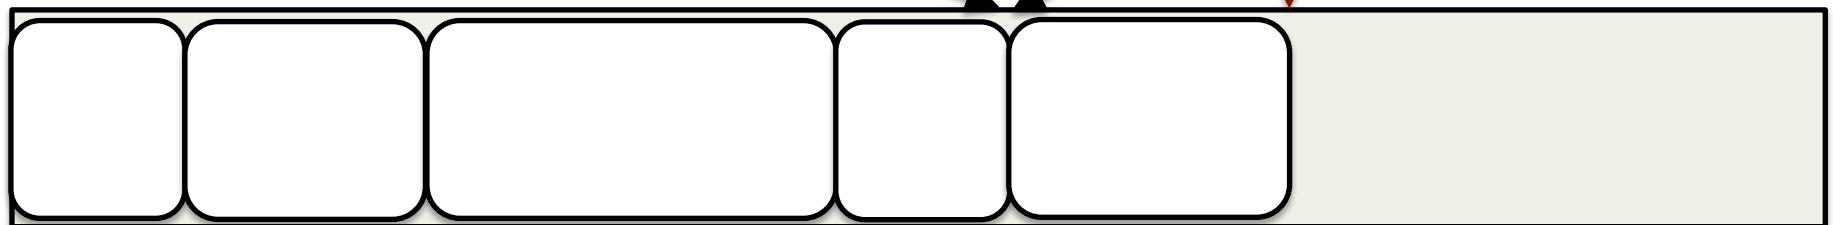
`ptr`



Thread 2

`ptr = getPtr()`
`newPtr = ptr + size`
`setPtr(newPtr)`

`ptr`



Thread 1

`ptr = getPtr()`

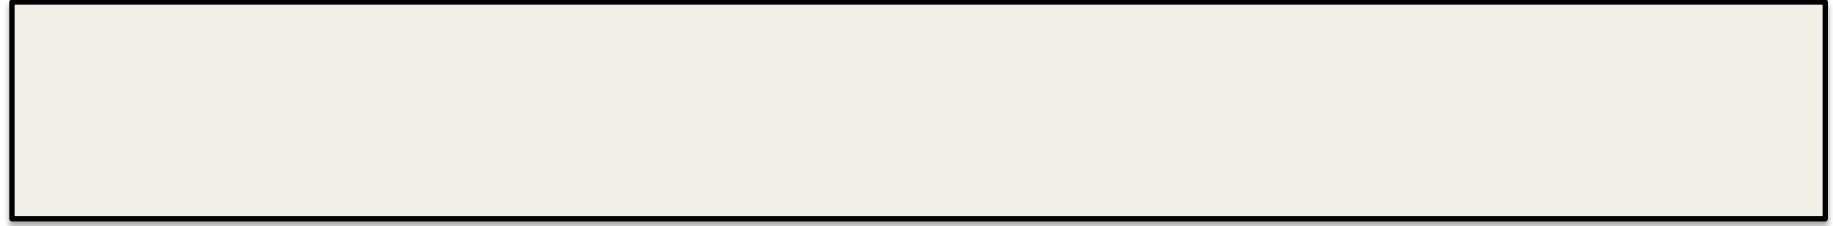
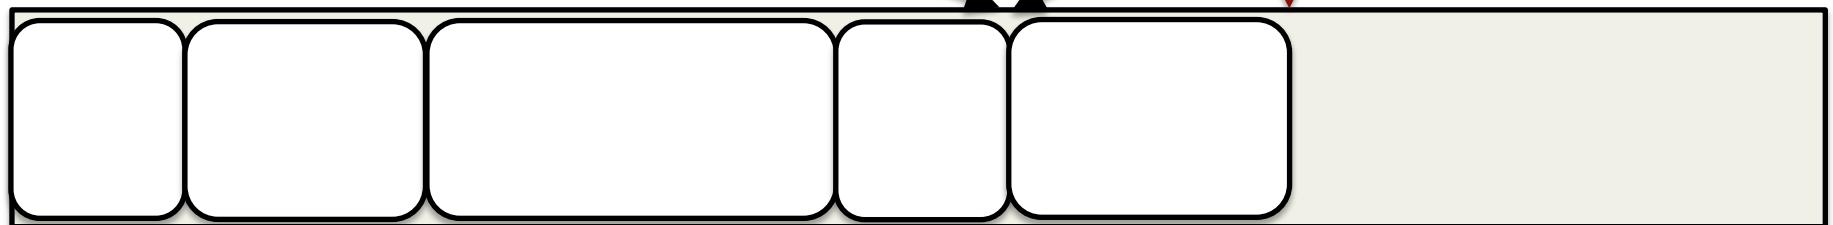
`ptr`



Thread 2

`ptr = getPtr()`
`newPtr = ptr + size`
`setPtr(newPtr)`

`ptr`



Thread 1

`ptr = getPtr()`

`newPtr`



`ptr`

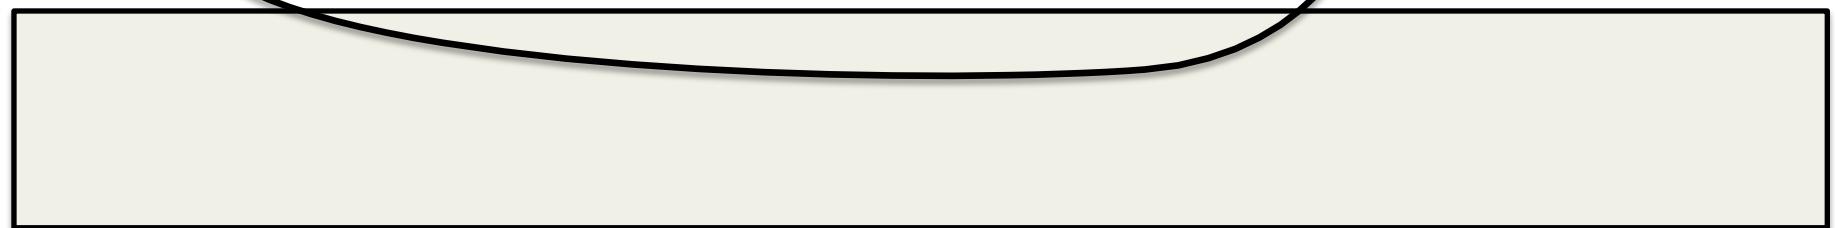
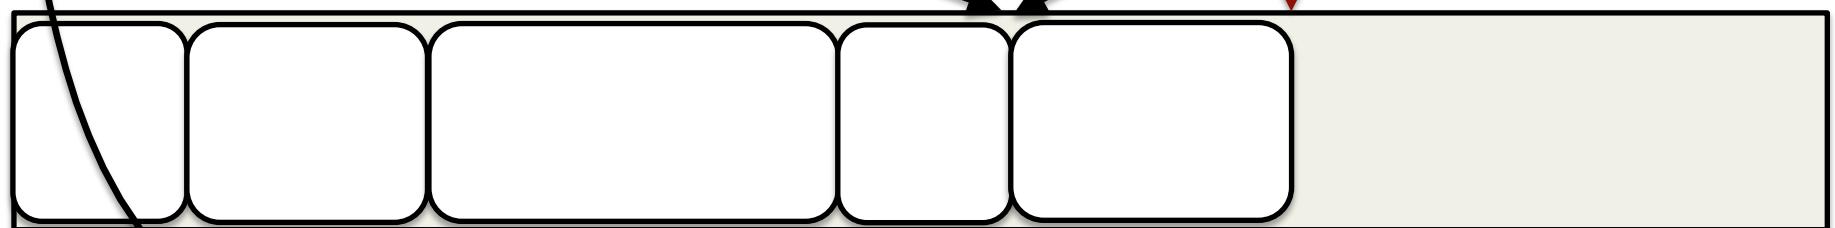


`newPtr = ptr + size`
`setPtr(newPtr)`

Thread 2

`ptr = getPtr()`
`newPtr = ptr + size`
`setPtr(newPtr)`

`ptr`



Thread 1

$\text{ptr} = \text{getPtr}()$

ptr

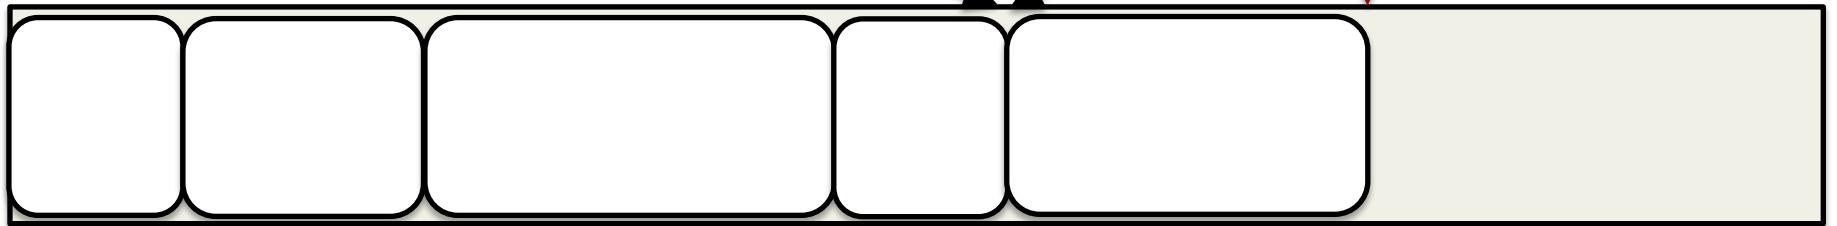


$\text{newPtr} = \text{ptr} + \text{size}$
 $\text{setPtr}(\text{newPtr})$

Thread 2

$\text{ptr} = \text{getPtr}()$
 $\text{newPtr} = \text{ptr} + \text{size}$
 $\text{setPtr}(\text{newPtr})$

ptr



Solving Allocation Races

- We saw last week that this is a data race
 - Two threads accessing the same memory location
 - No synchronization in between
- Standard solution would be to synchronize
 - Acquire some shared lock before allocation
 - Use atomic instructions

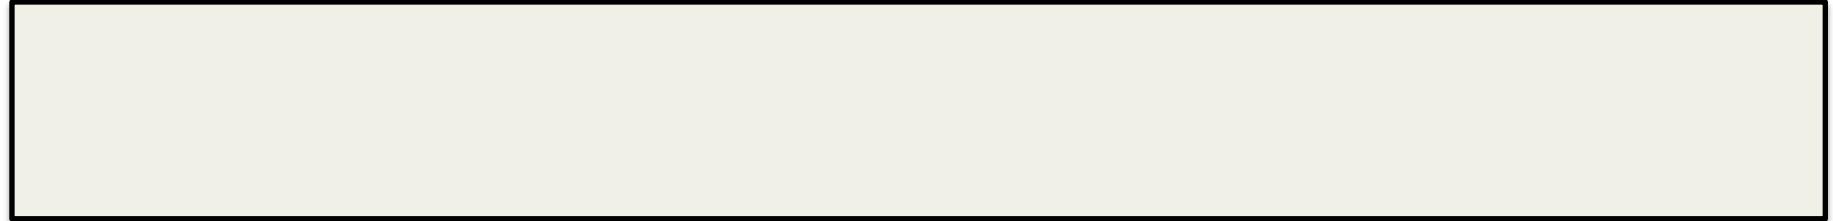
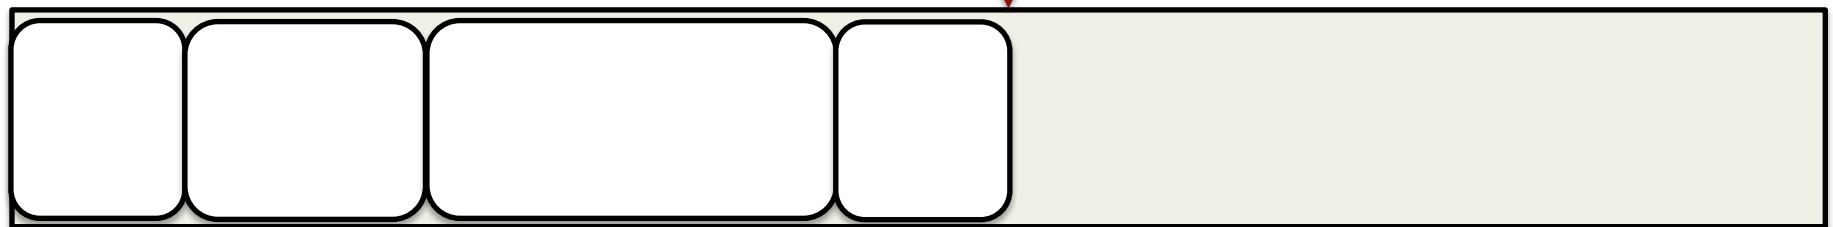
Thread 1

`ptr = getPtr()`

`newPtr = ptr + size`
`setPtr(newPtr)`

Thread 2

`ptr = getPtr()`
`newPtr = ptr + size`
`setPtr(newPtr)`



Thread 1

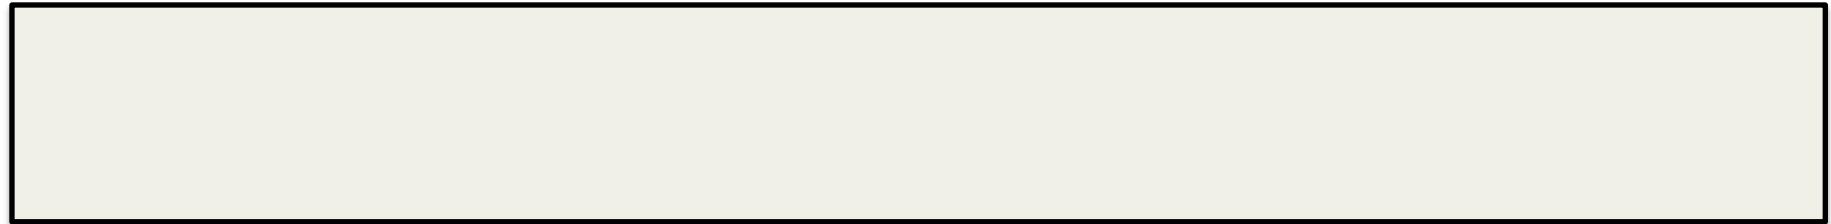
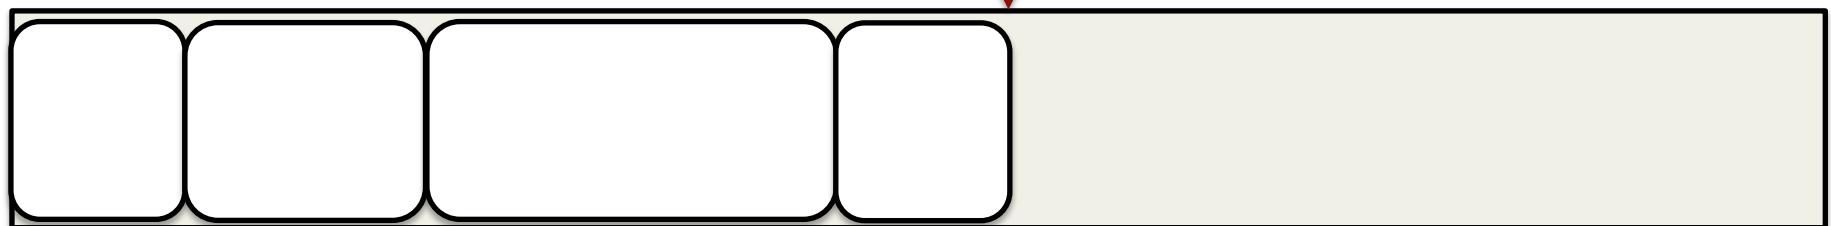
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

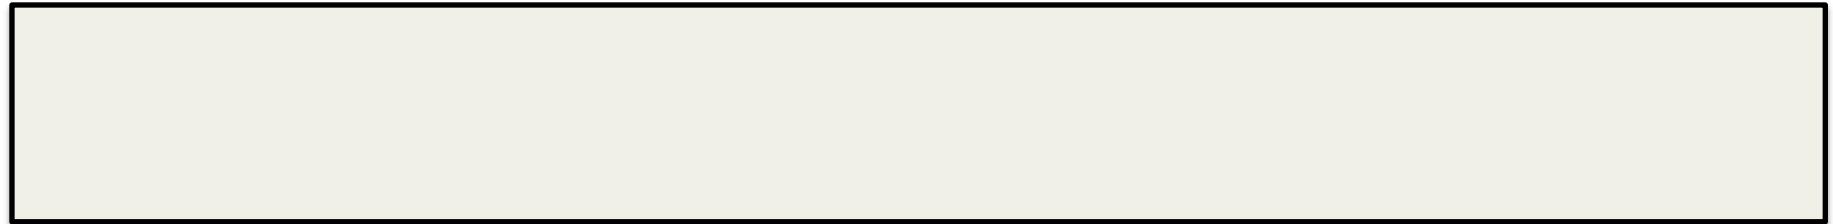
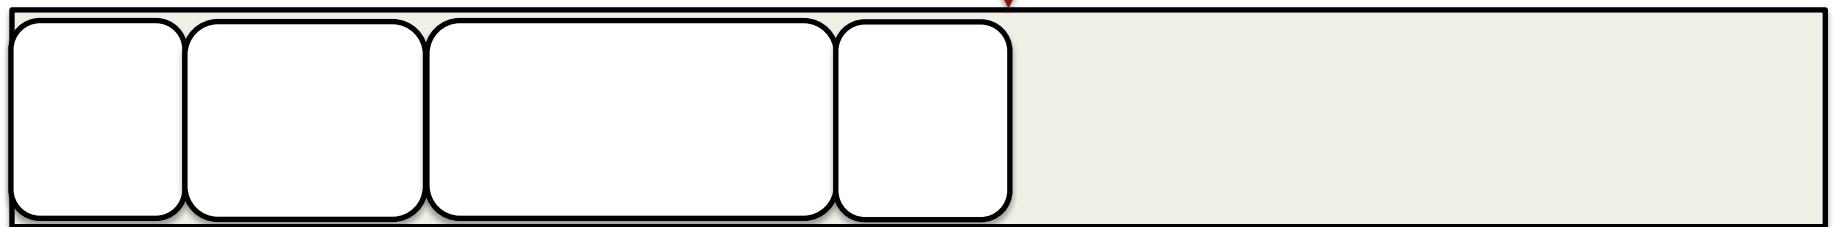
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

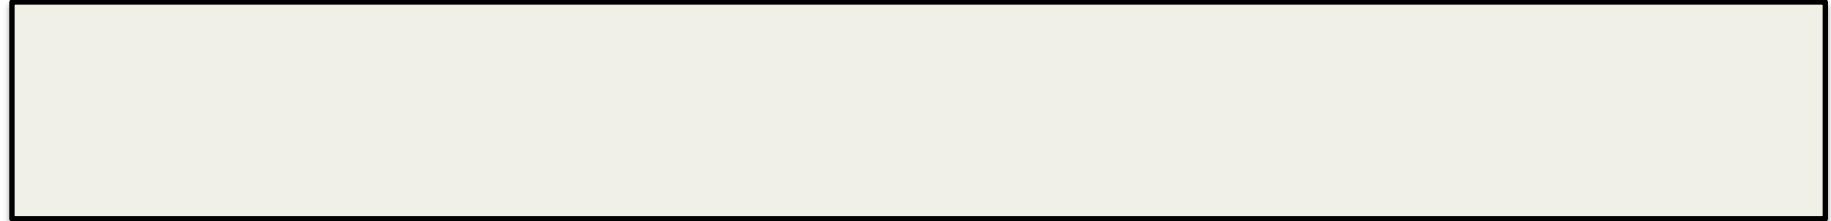
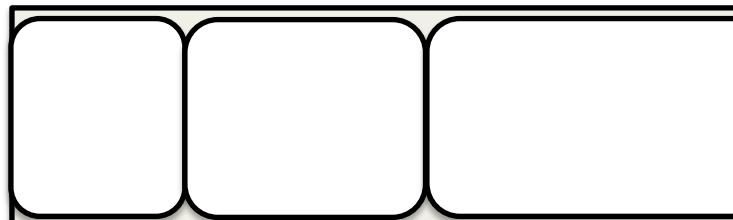
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

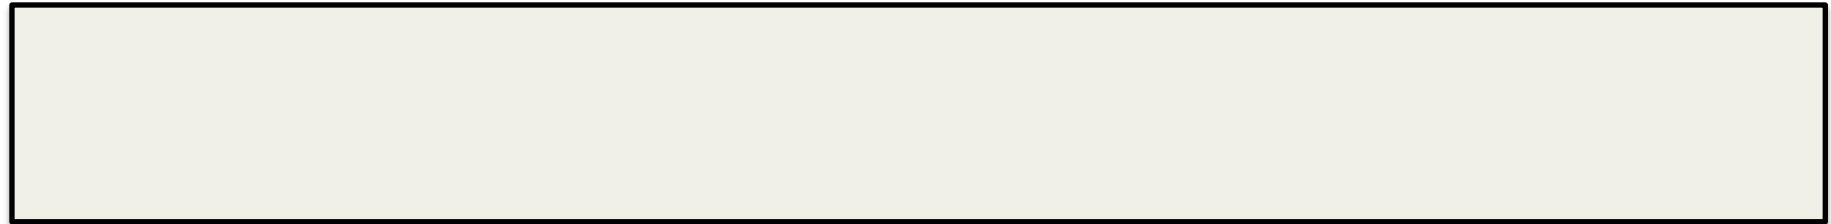
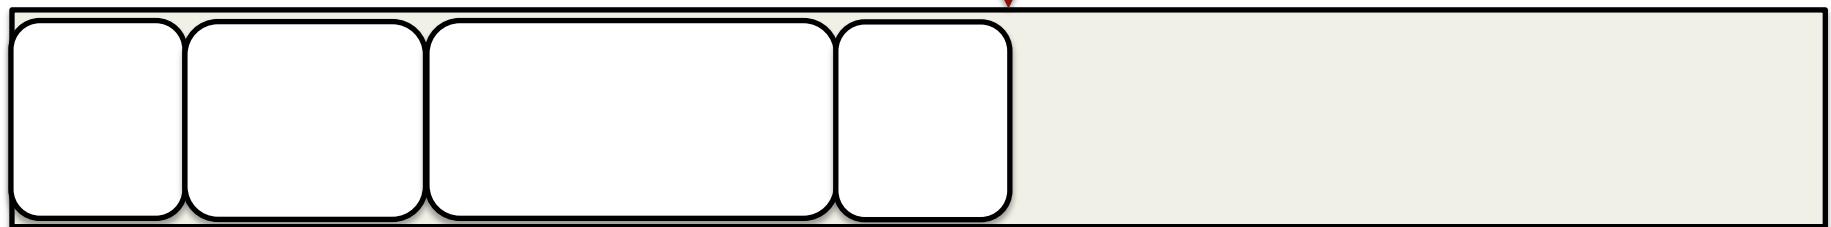
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

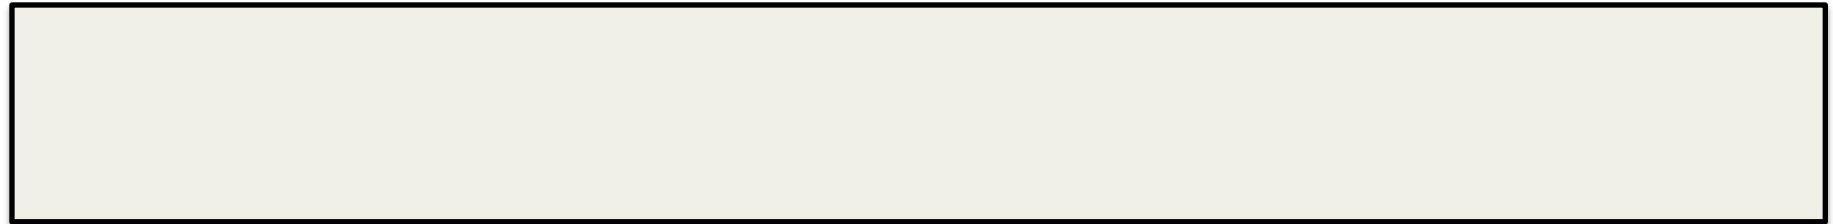
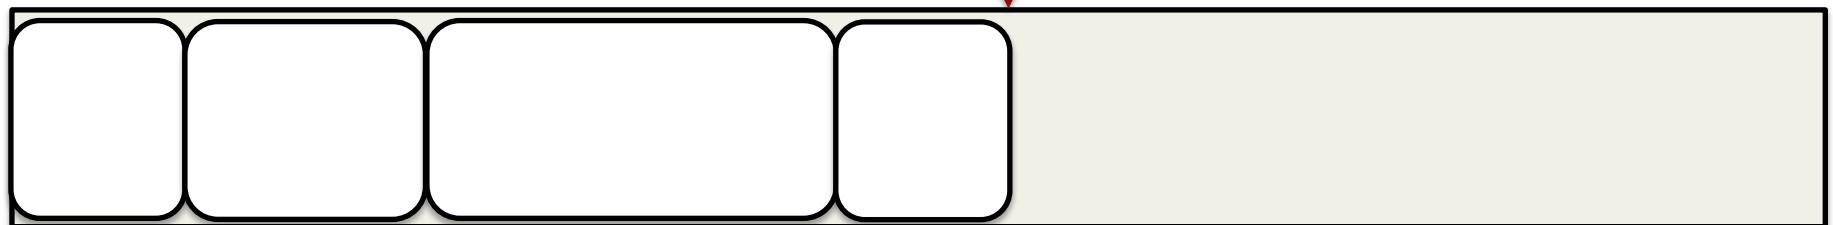
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



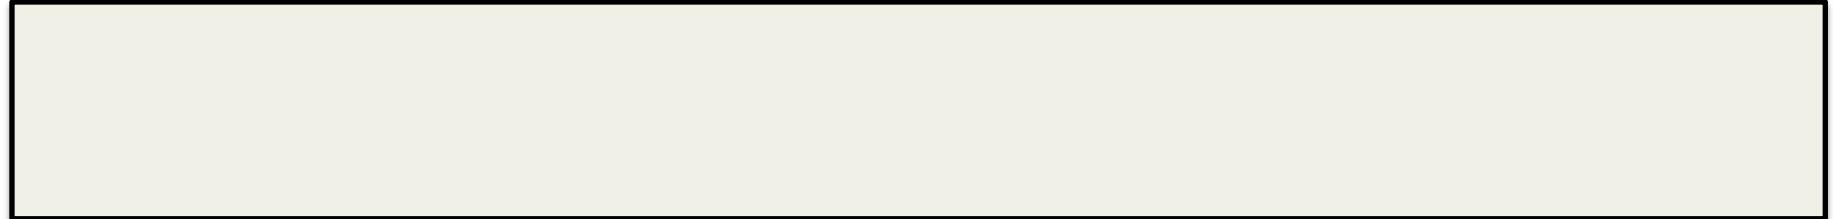
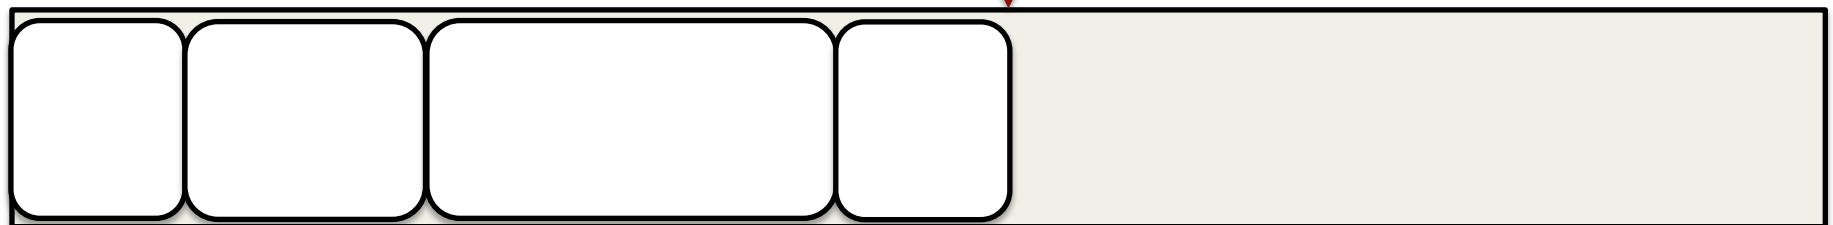
Thread 1

```
lock(heap)  
ptr = getPtr()
```

```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

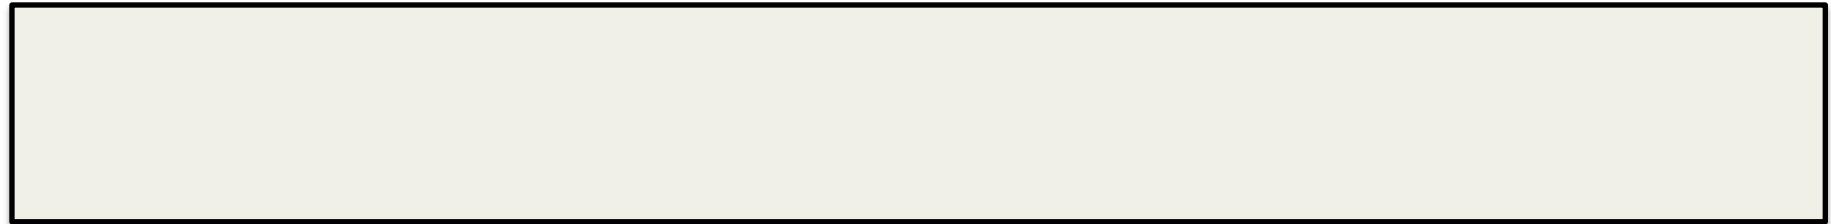
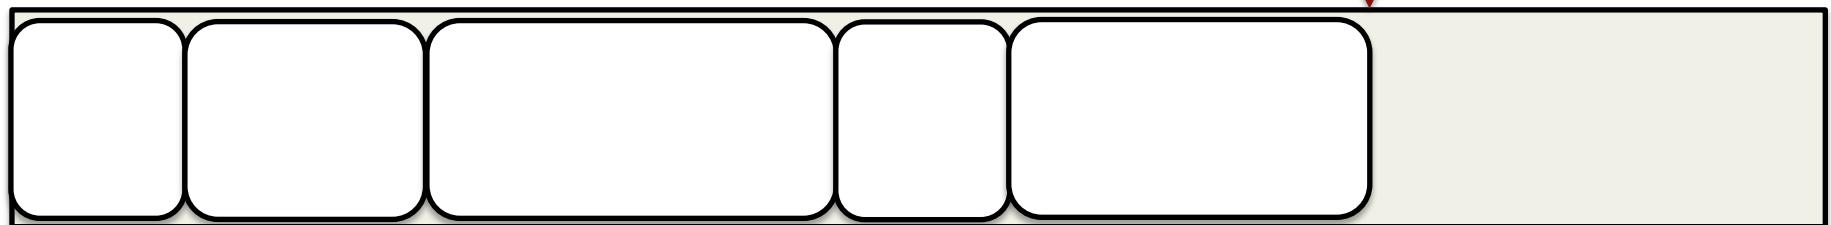
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

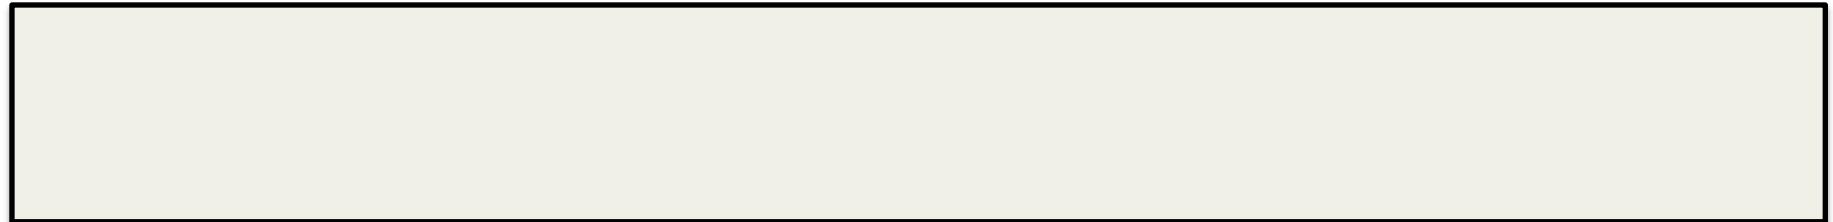
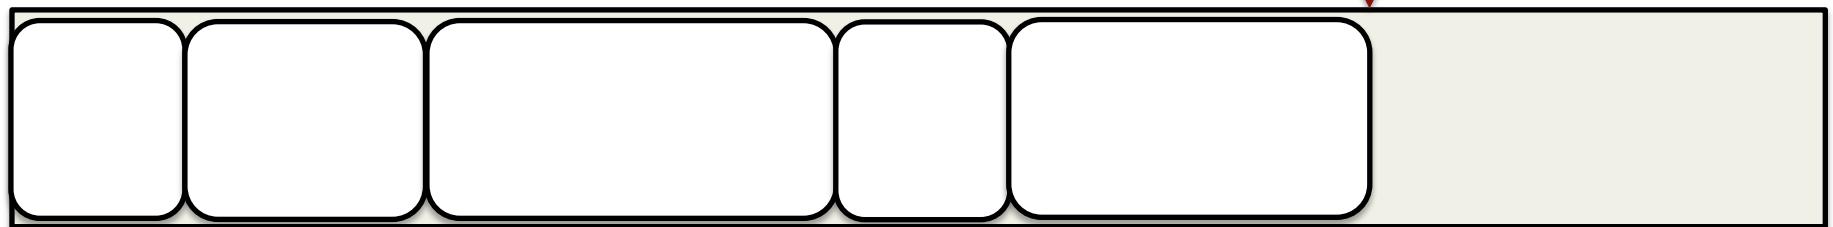
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

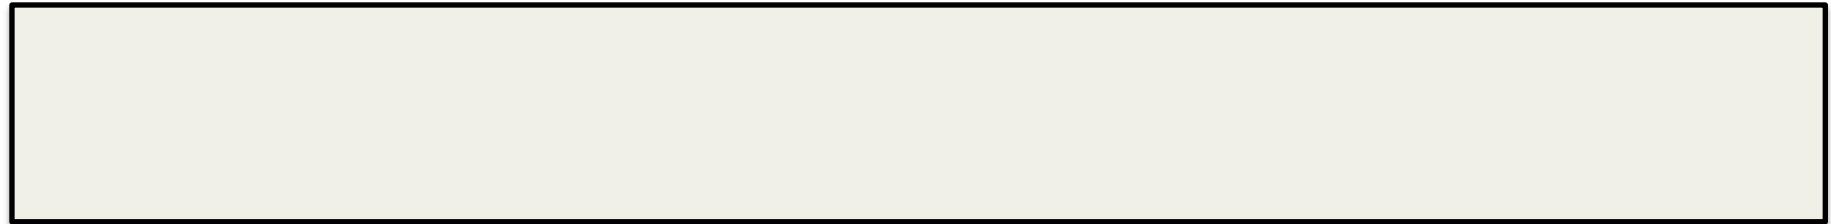
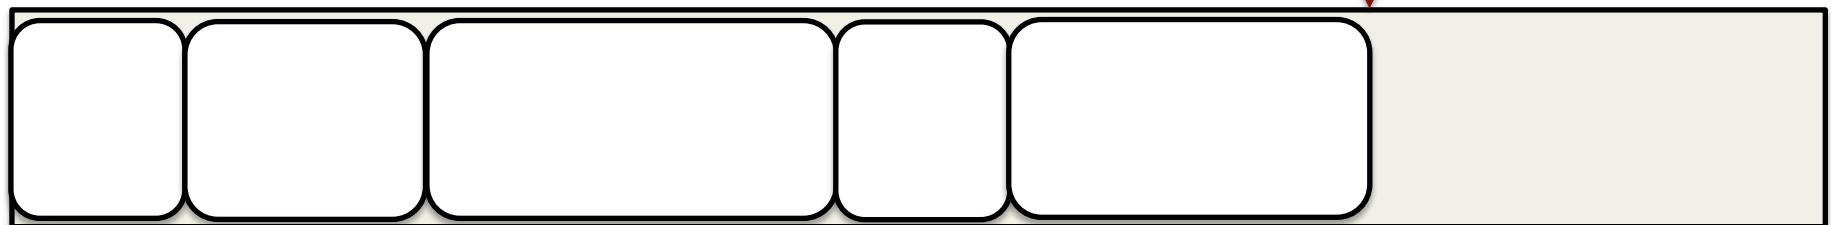
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



Thread 1

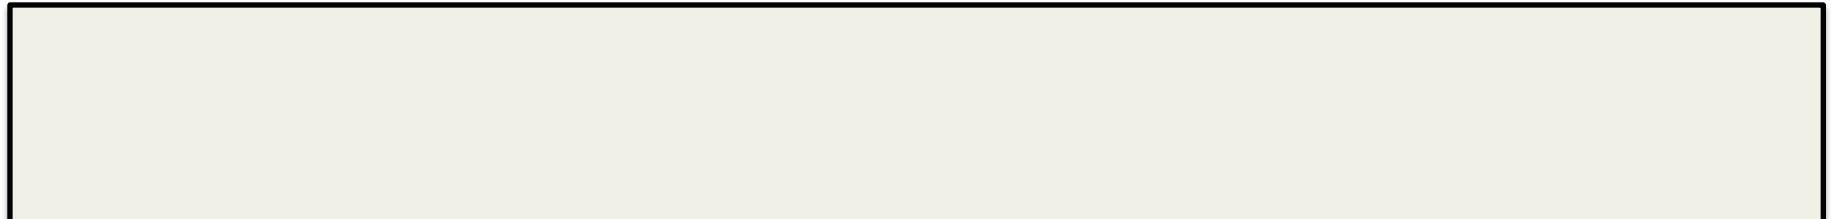
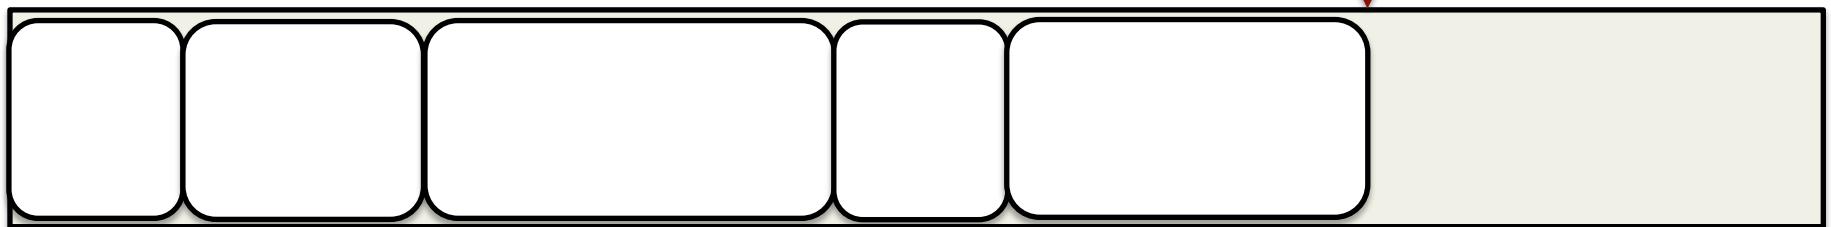
```
lock(heap)  
ptr = getPtr()
```



```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



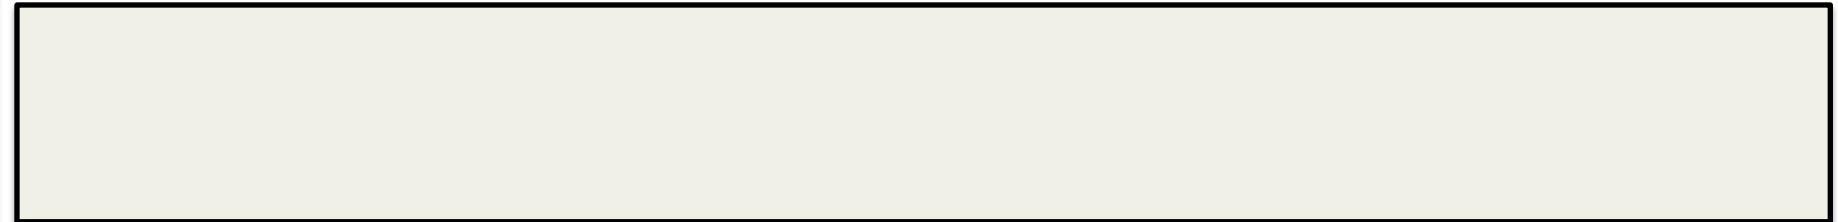
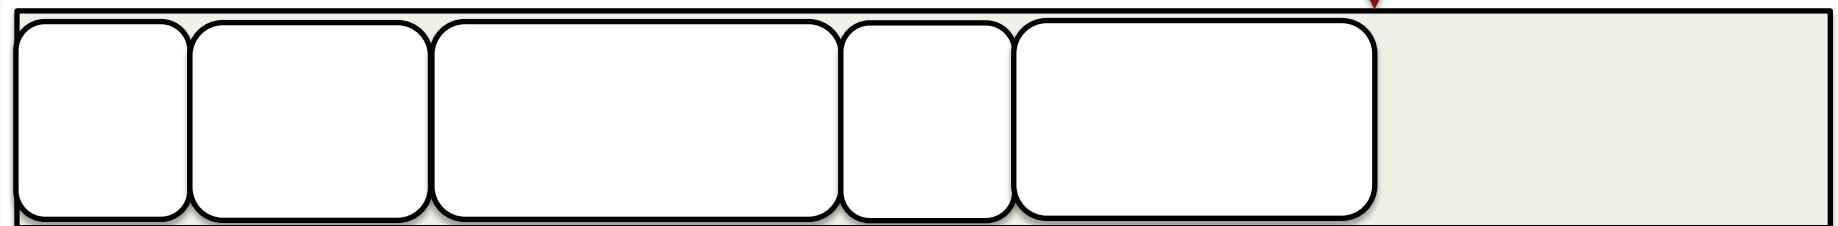
Thread 1

```
lock(heap)  
ptr = getPtr()
```

```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```



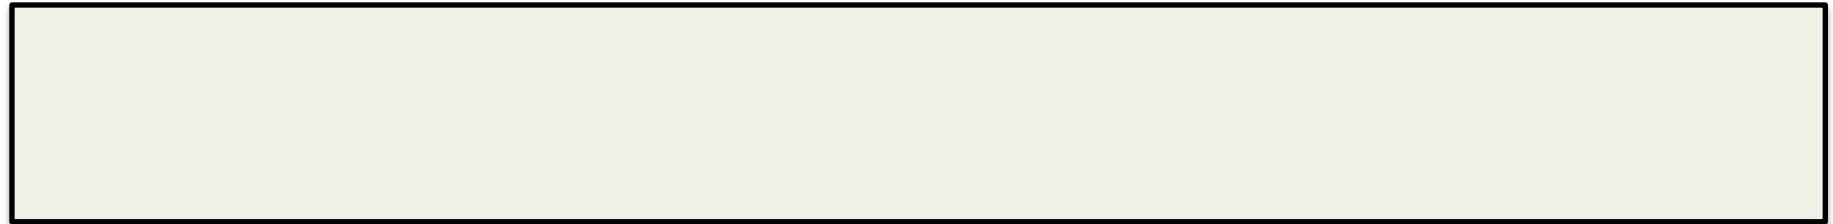
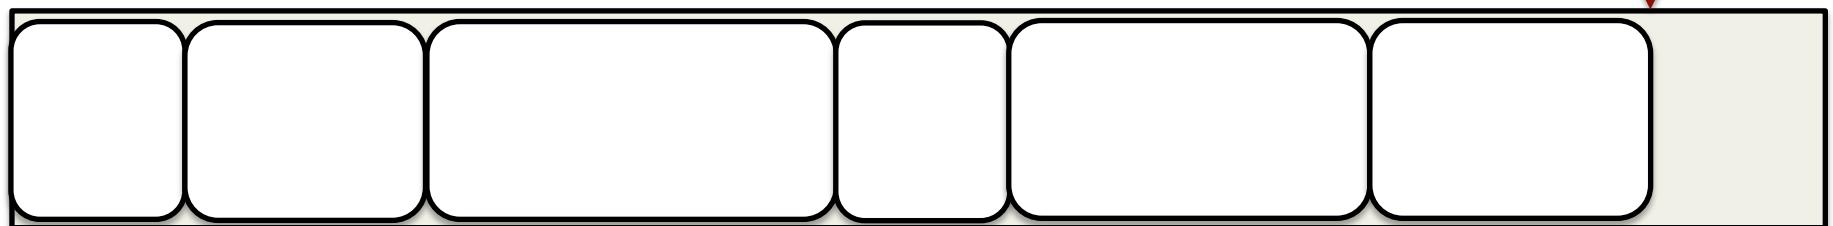
Thread 1

```
lock(heap)  
ptr = getPtr()
```

```
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

Thread 2

```
lock(heap)  
ptr = getPtr()  
newPtr = ptr + size  
setPtr(newPtr)  
unlock(heap)
```

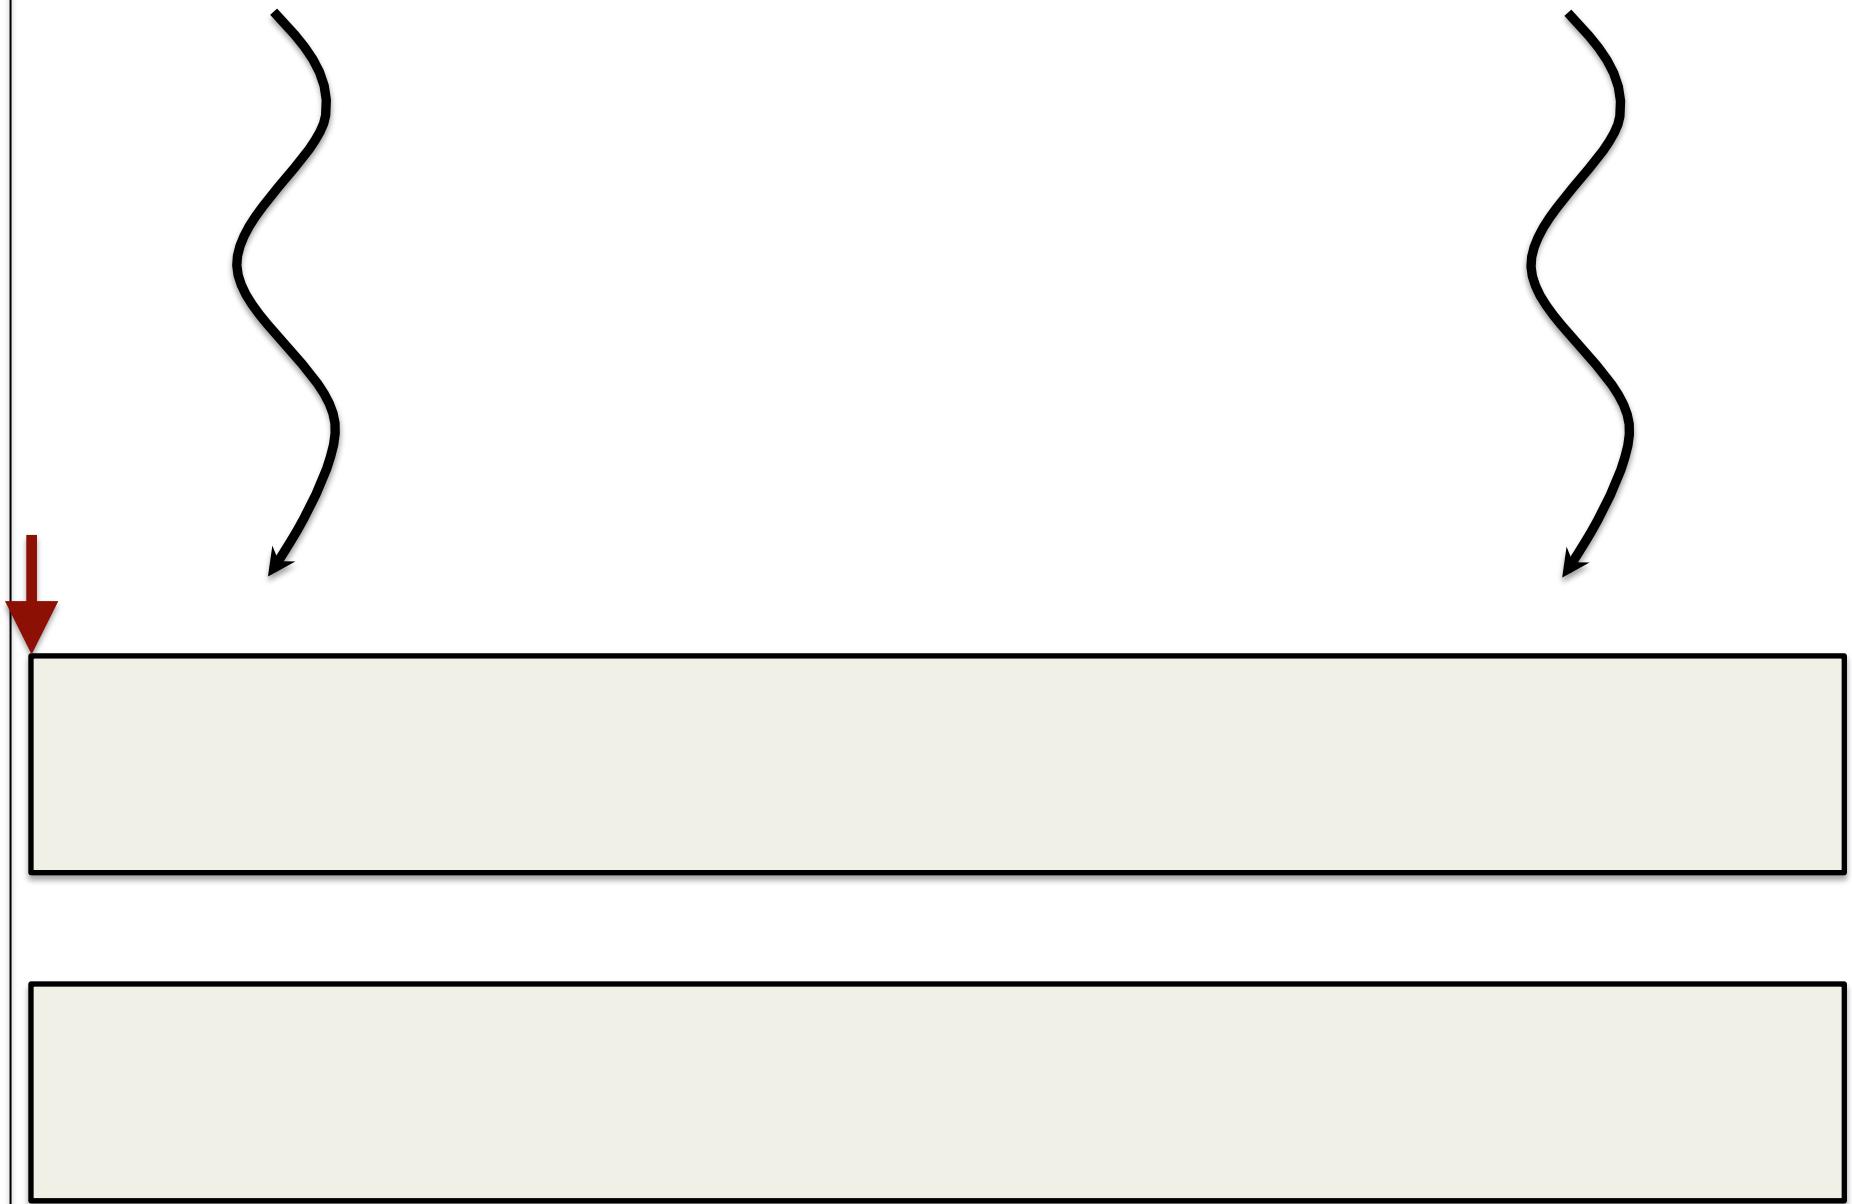


Contention

- Allocation is a very common operation
 - Particularly in managed environments
 - Developers assume allocation is easy
- Contention on the allocator is bad
 - Particularly as the number of threads increases
 - Single bottleneck for the whole system
- Better to allocate without synchronization

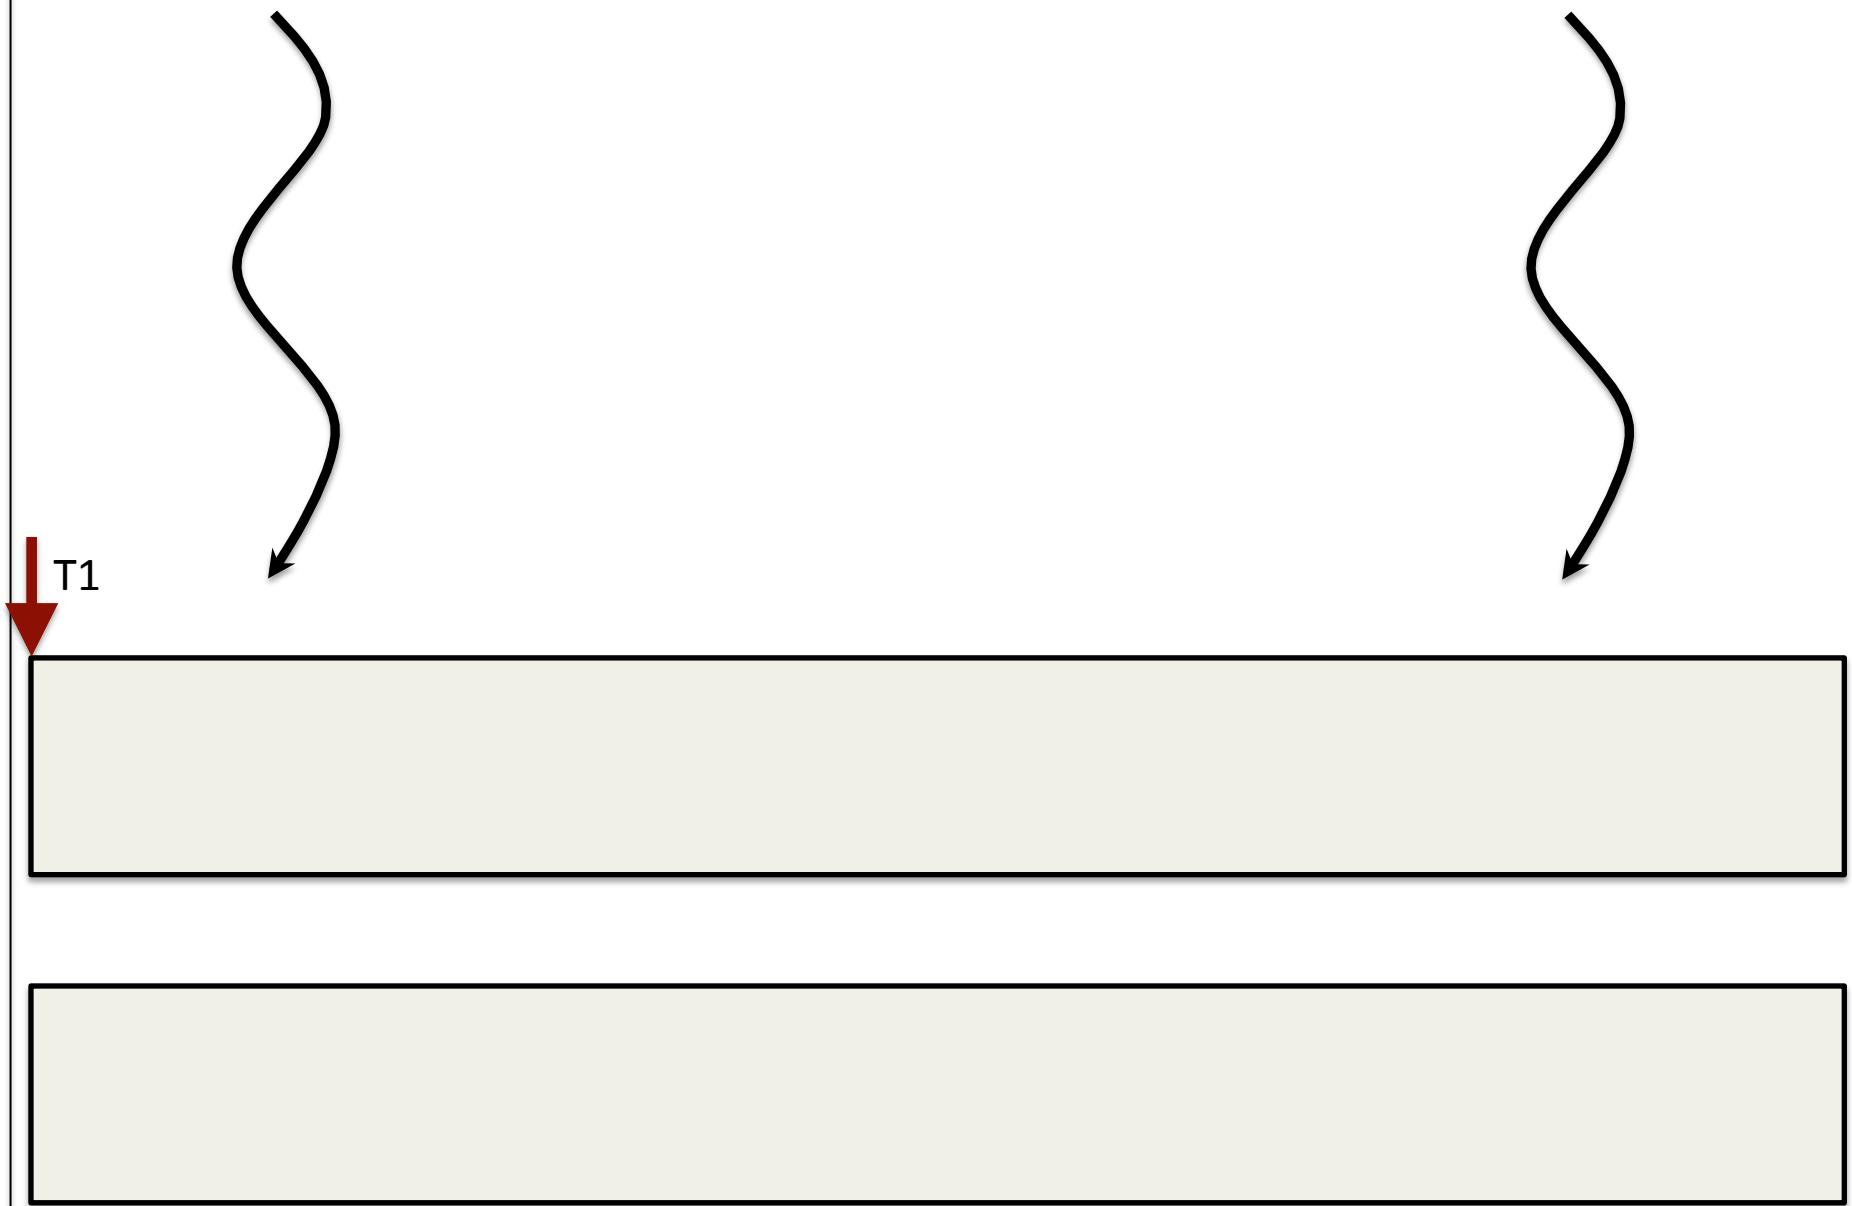
Thread 1

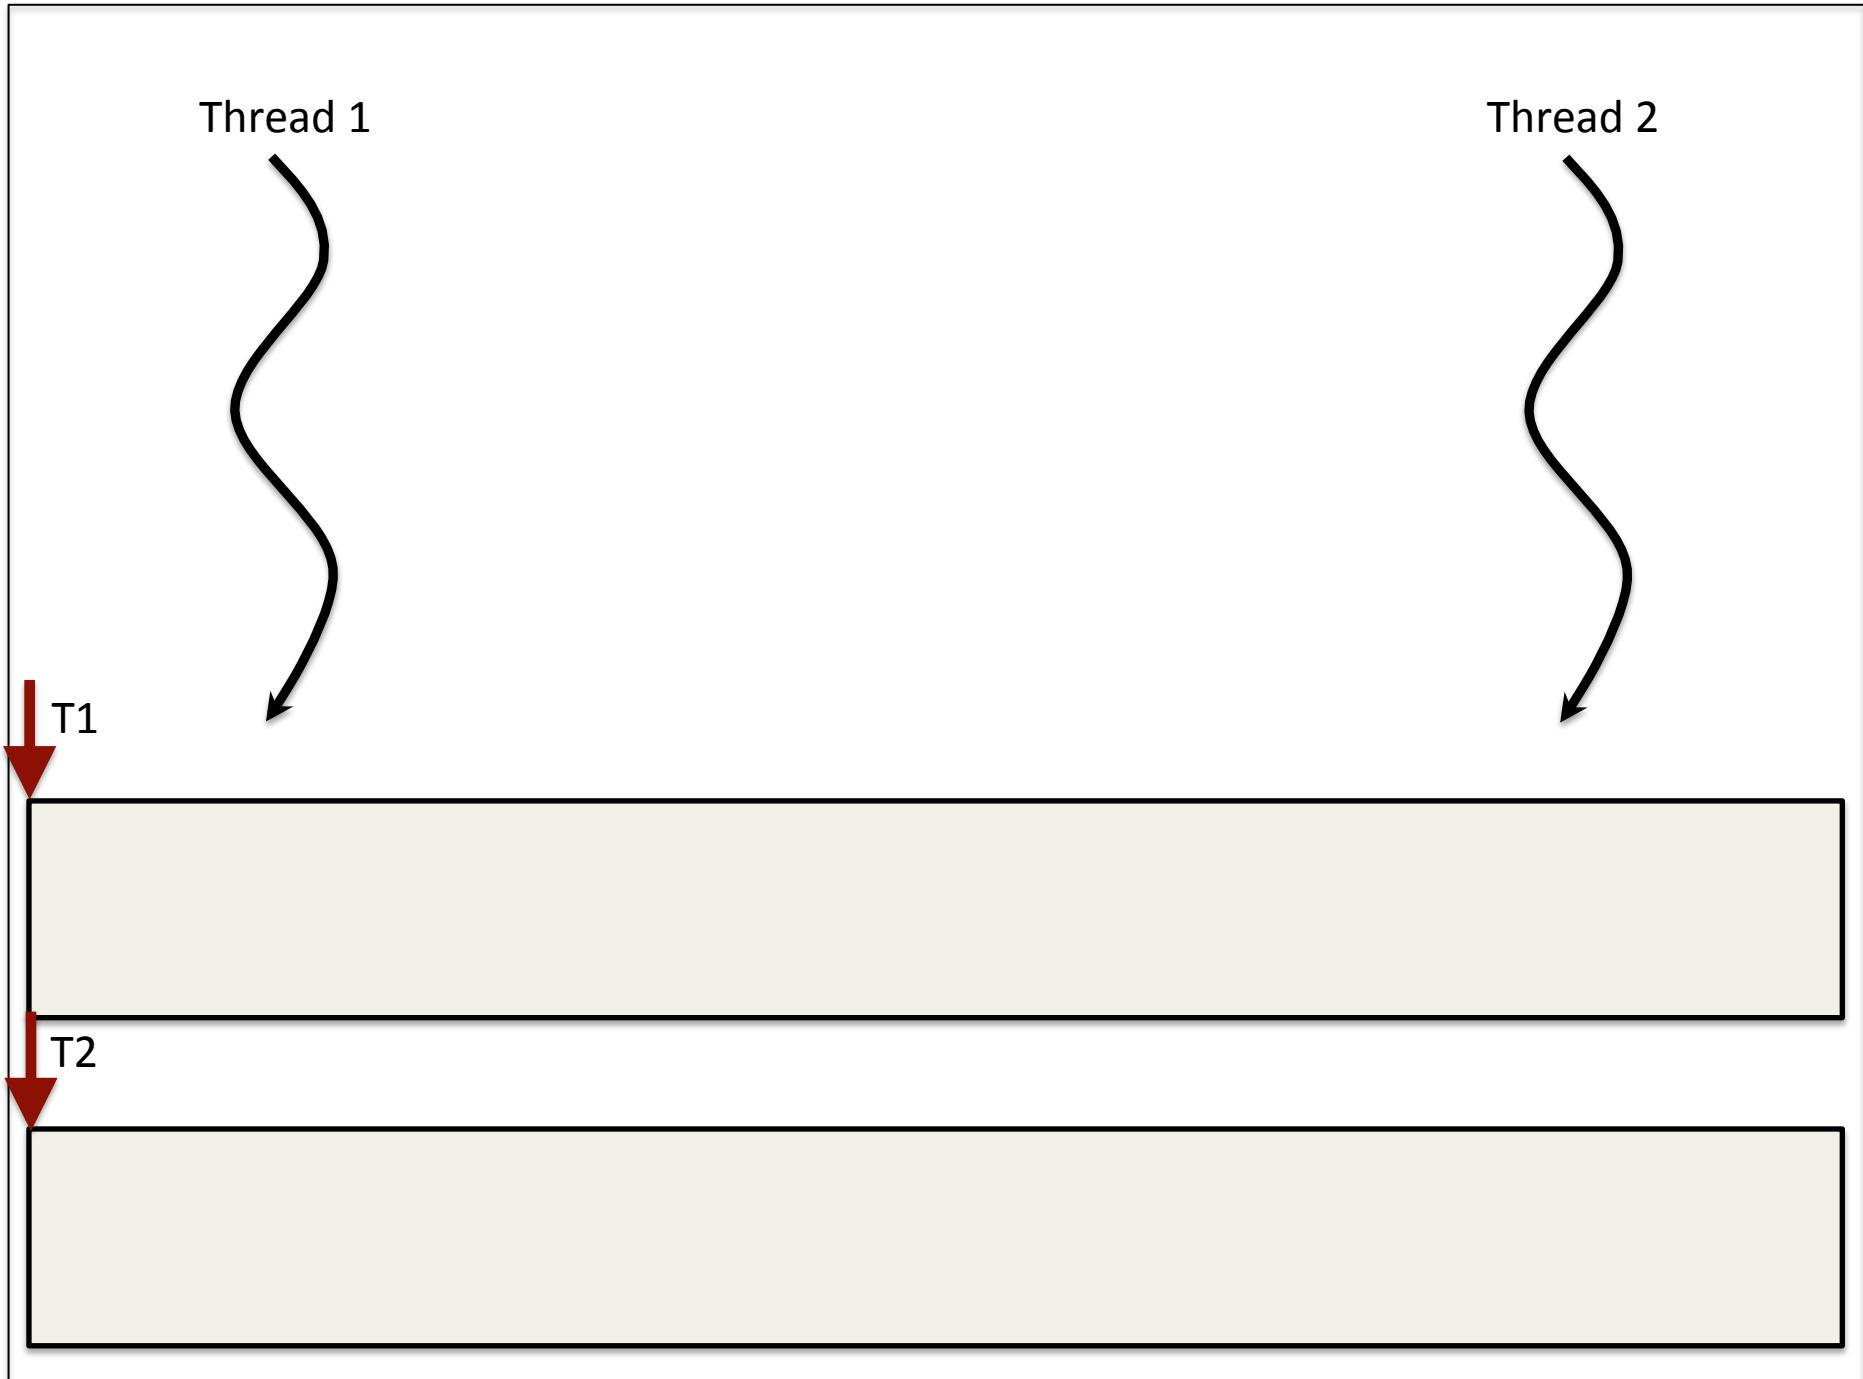
Thread 2



Thread 1

Thread 2





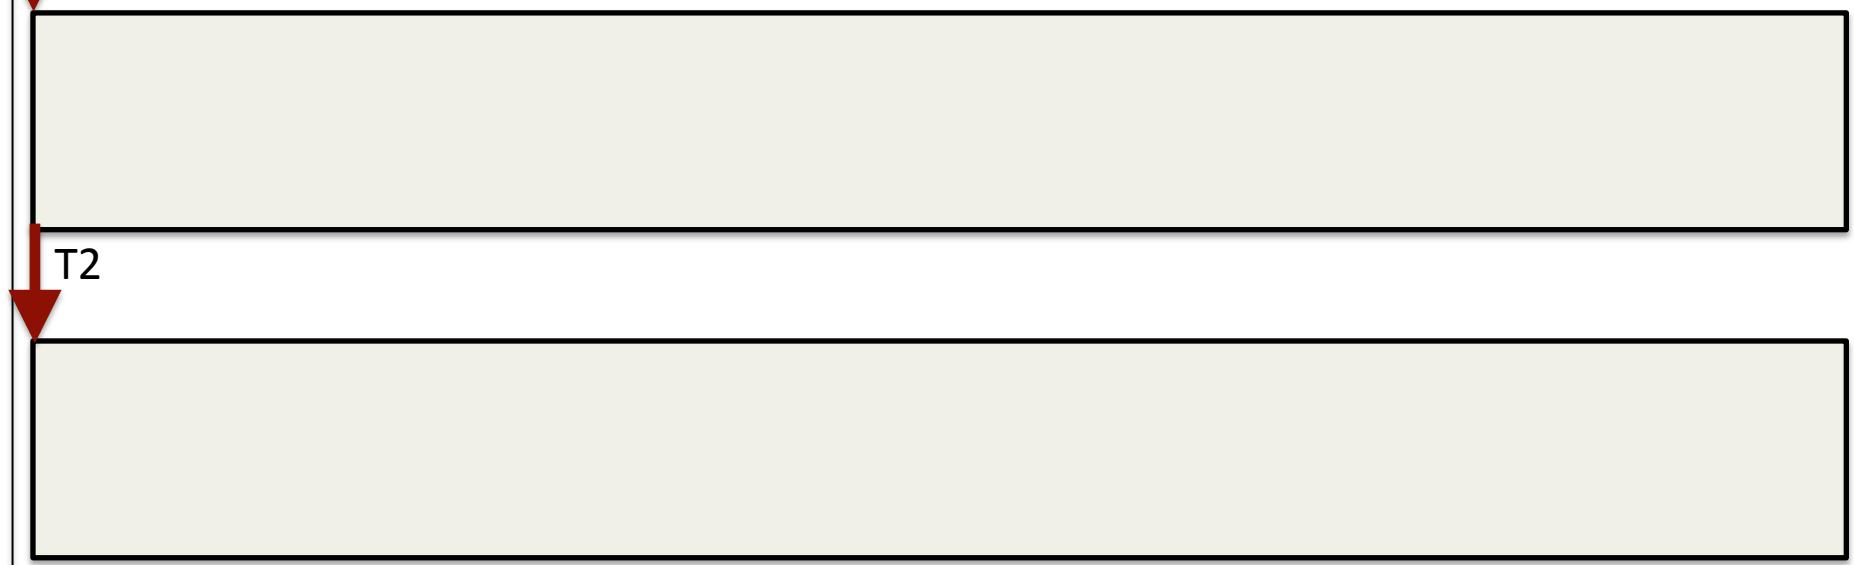
Thread 1

Thread 2

allocate

T1

T2

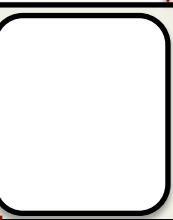


Thread 1

Thread 2



T1



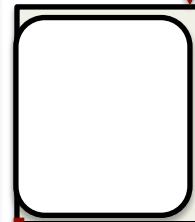
T2



Thread 1

Thread 2

allocate

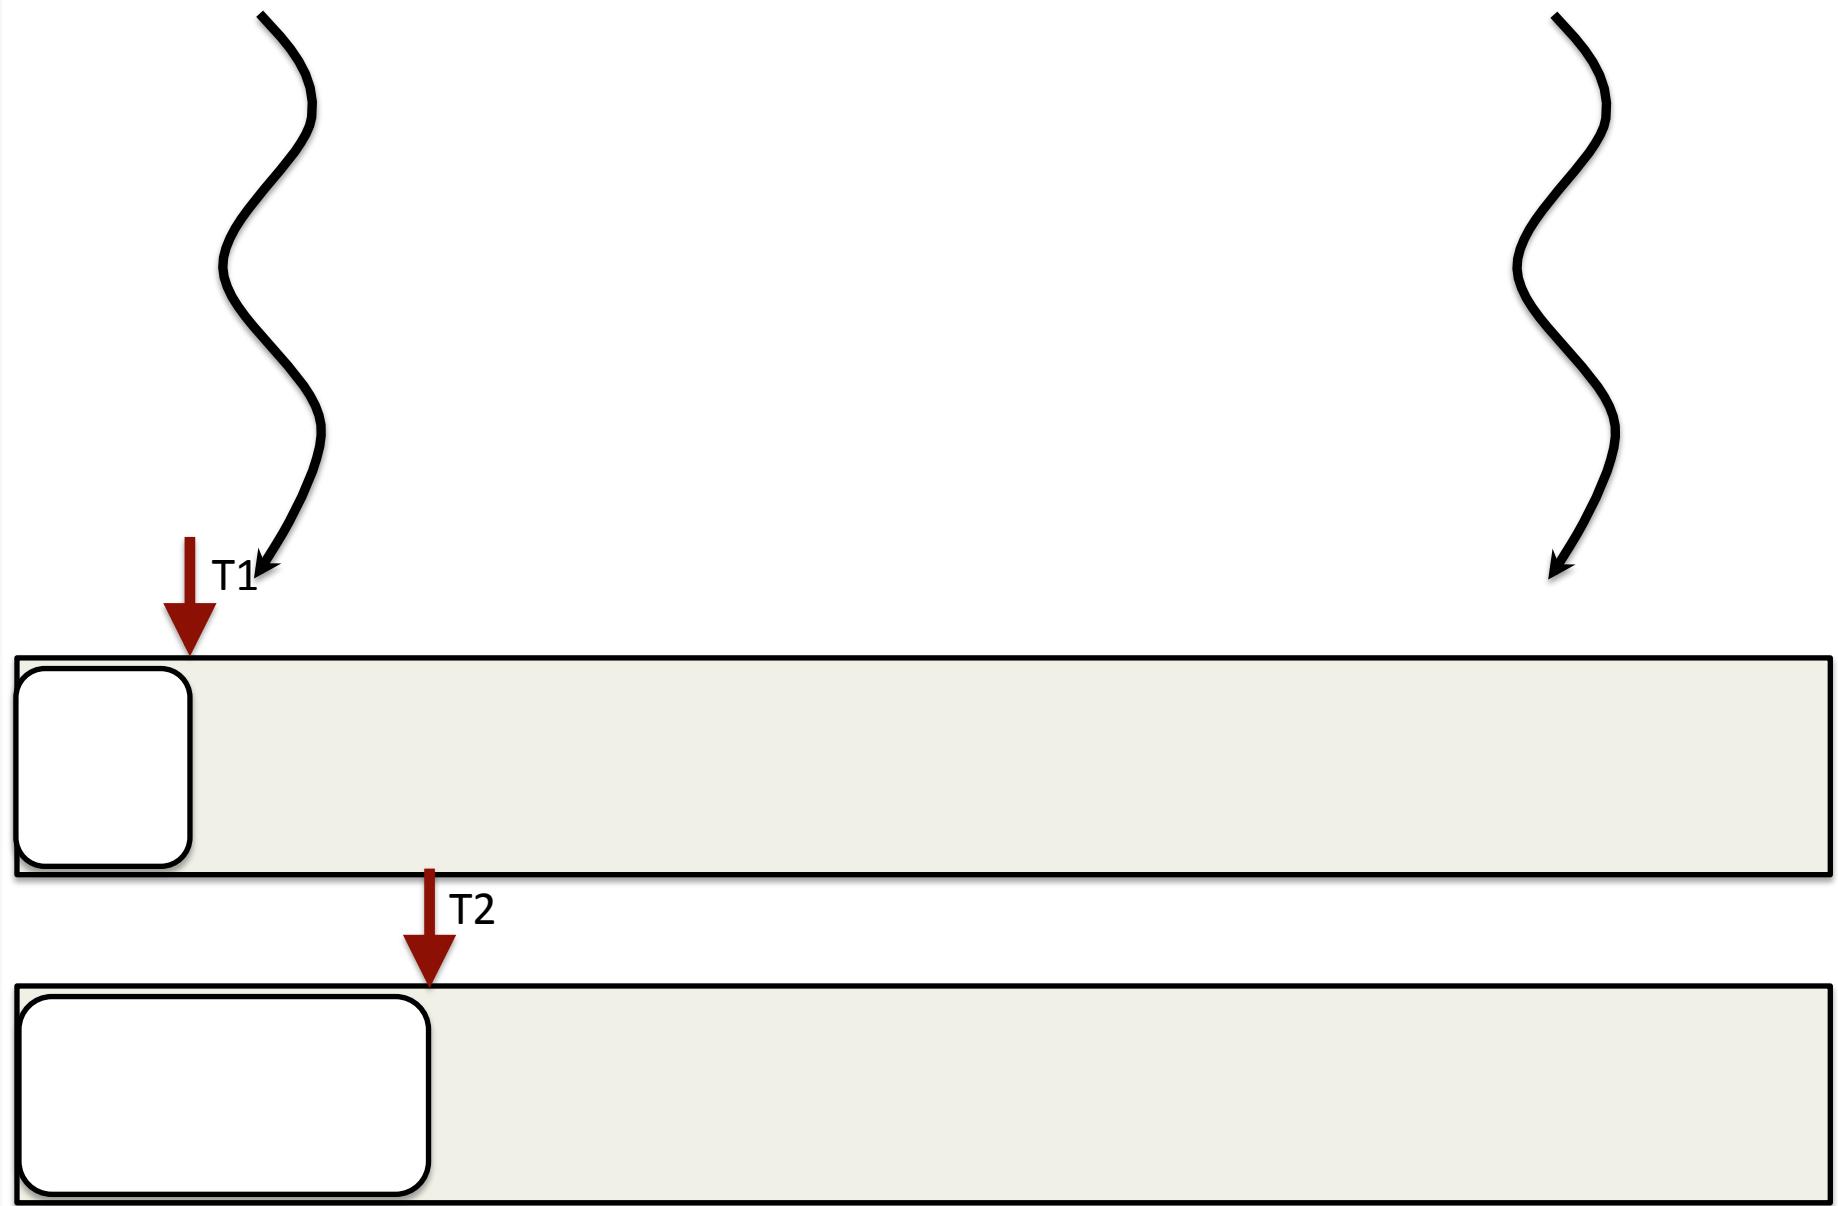


T2



Thread 1

Thread 2



Thread 1

Thread 2

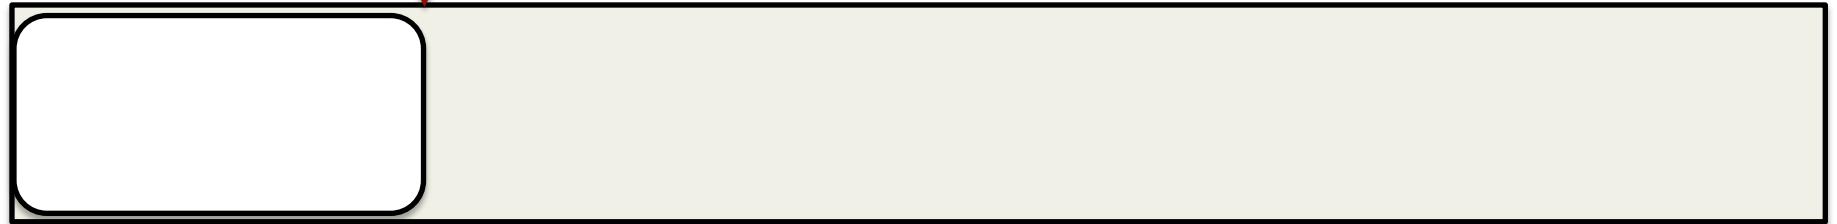
allocate



T1

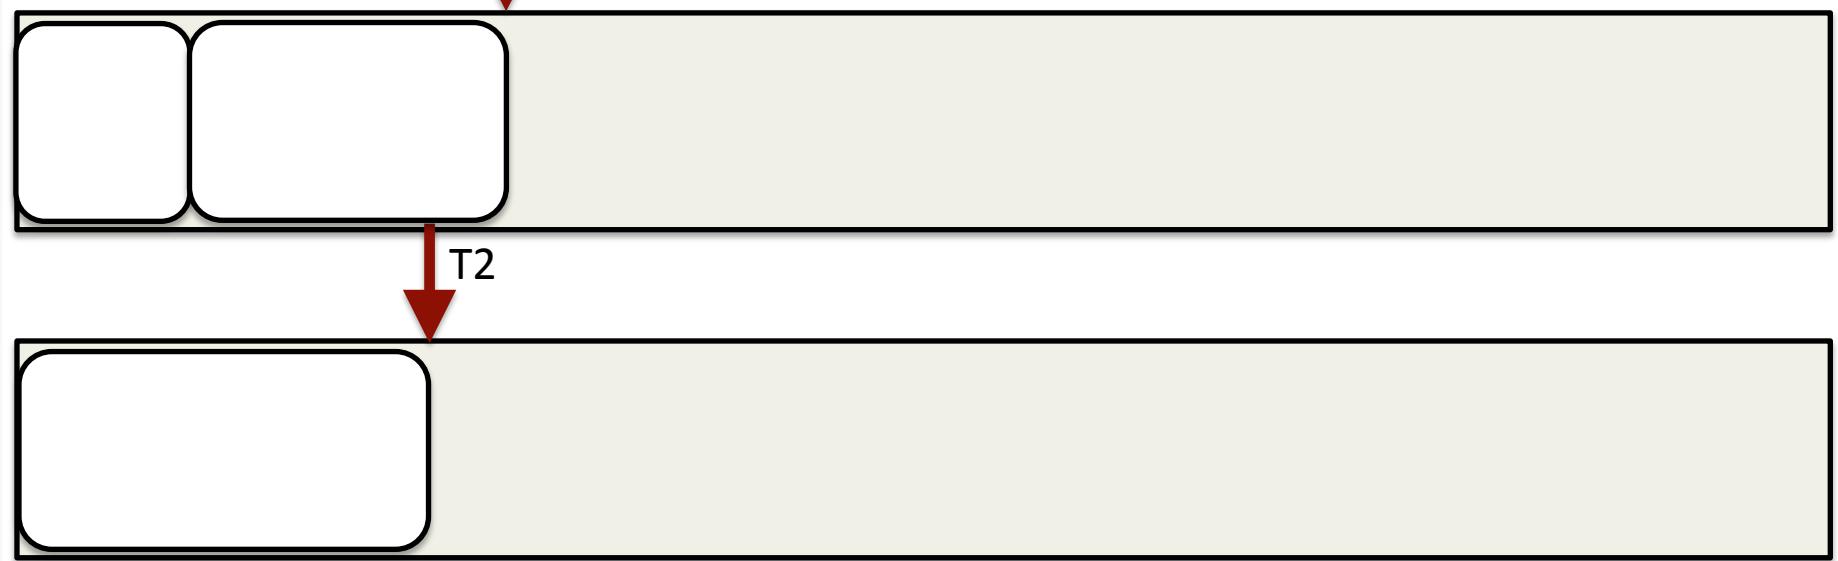


T2



Thread 1

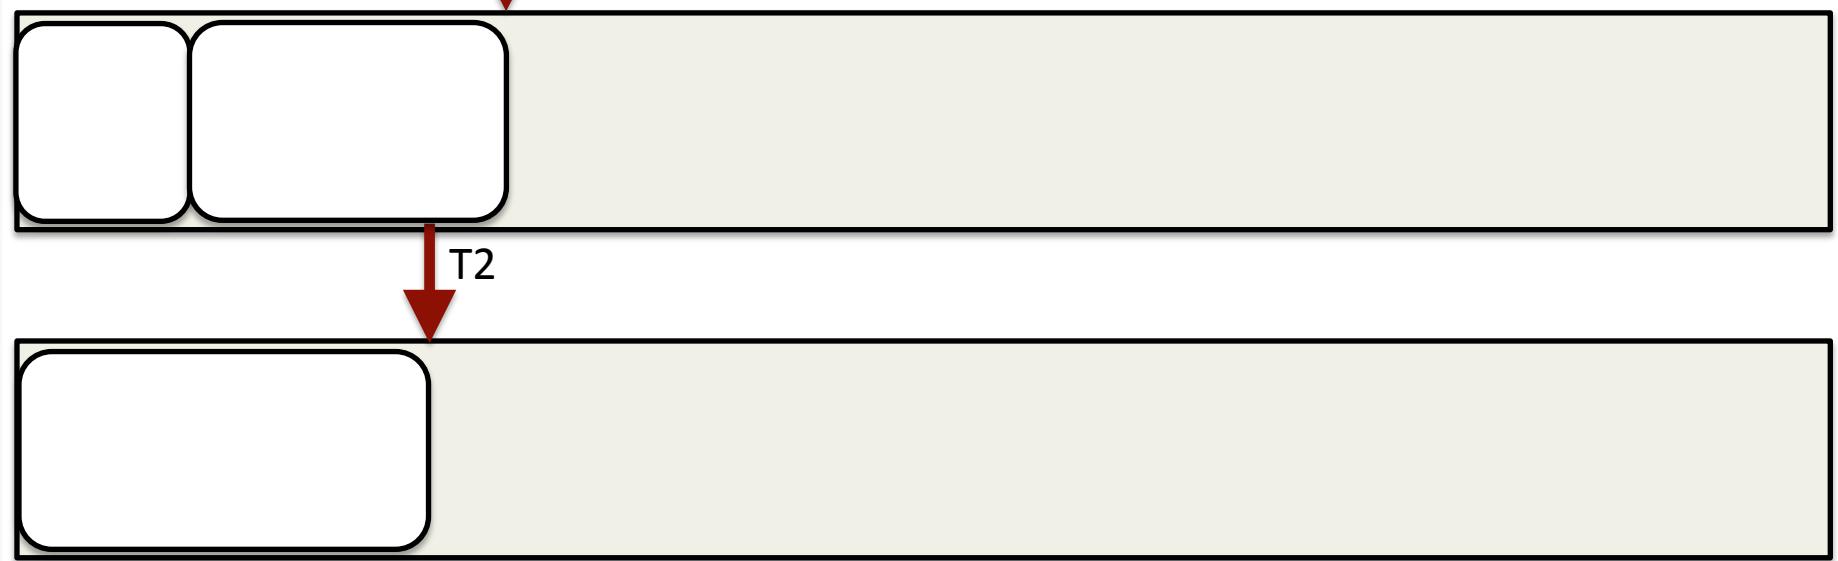
Thread 2



Thread 1

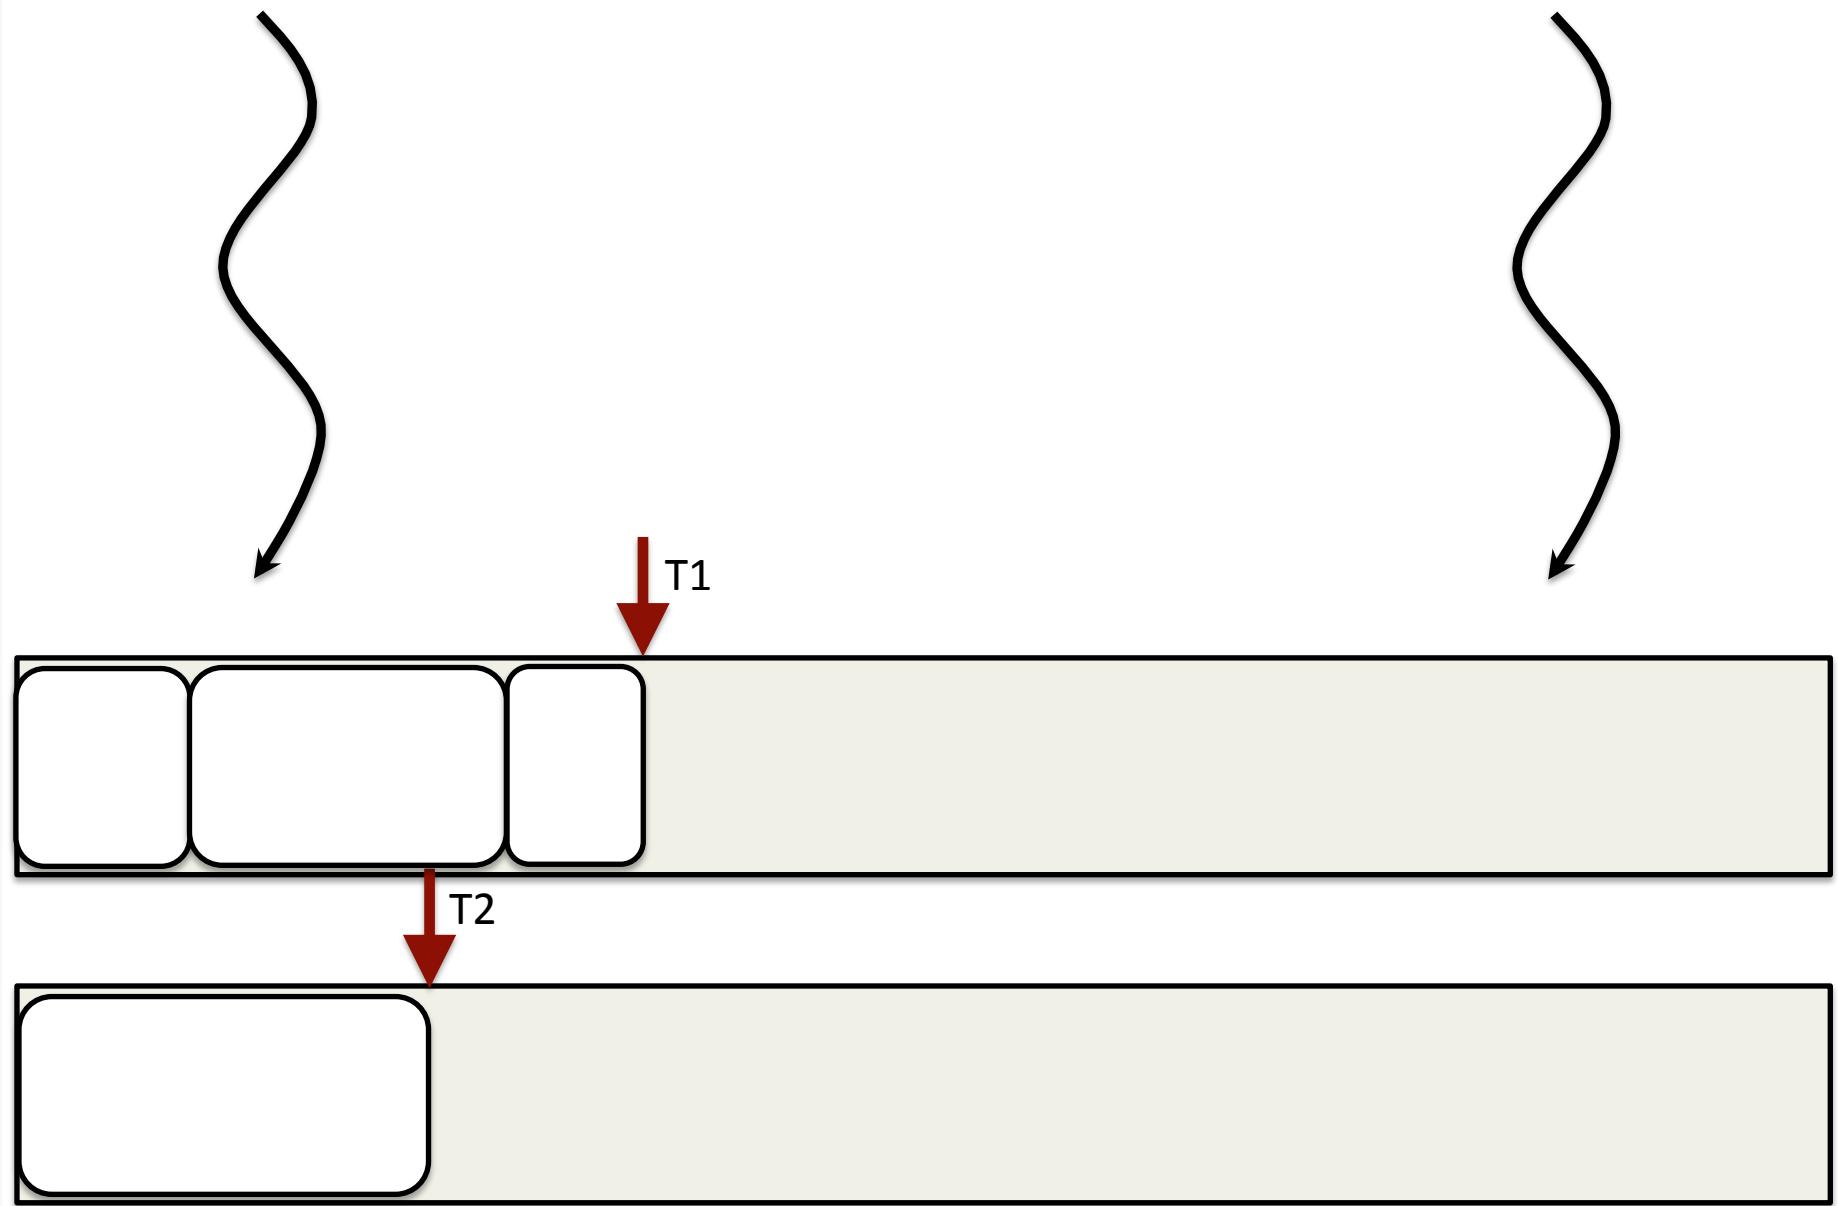
Thread 2

allocate



Thread 1

Thread 2



Thread-Local Allocation

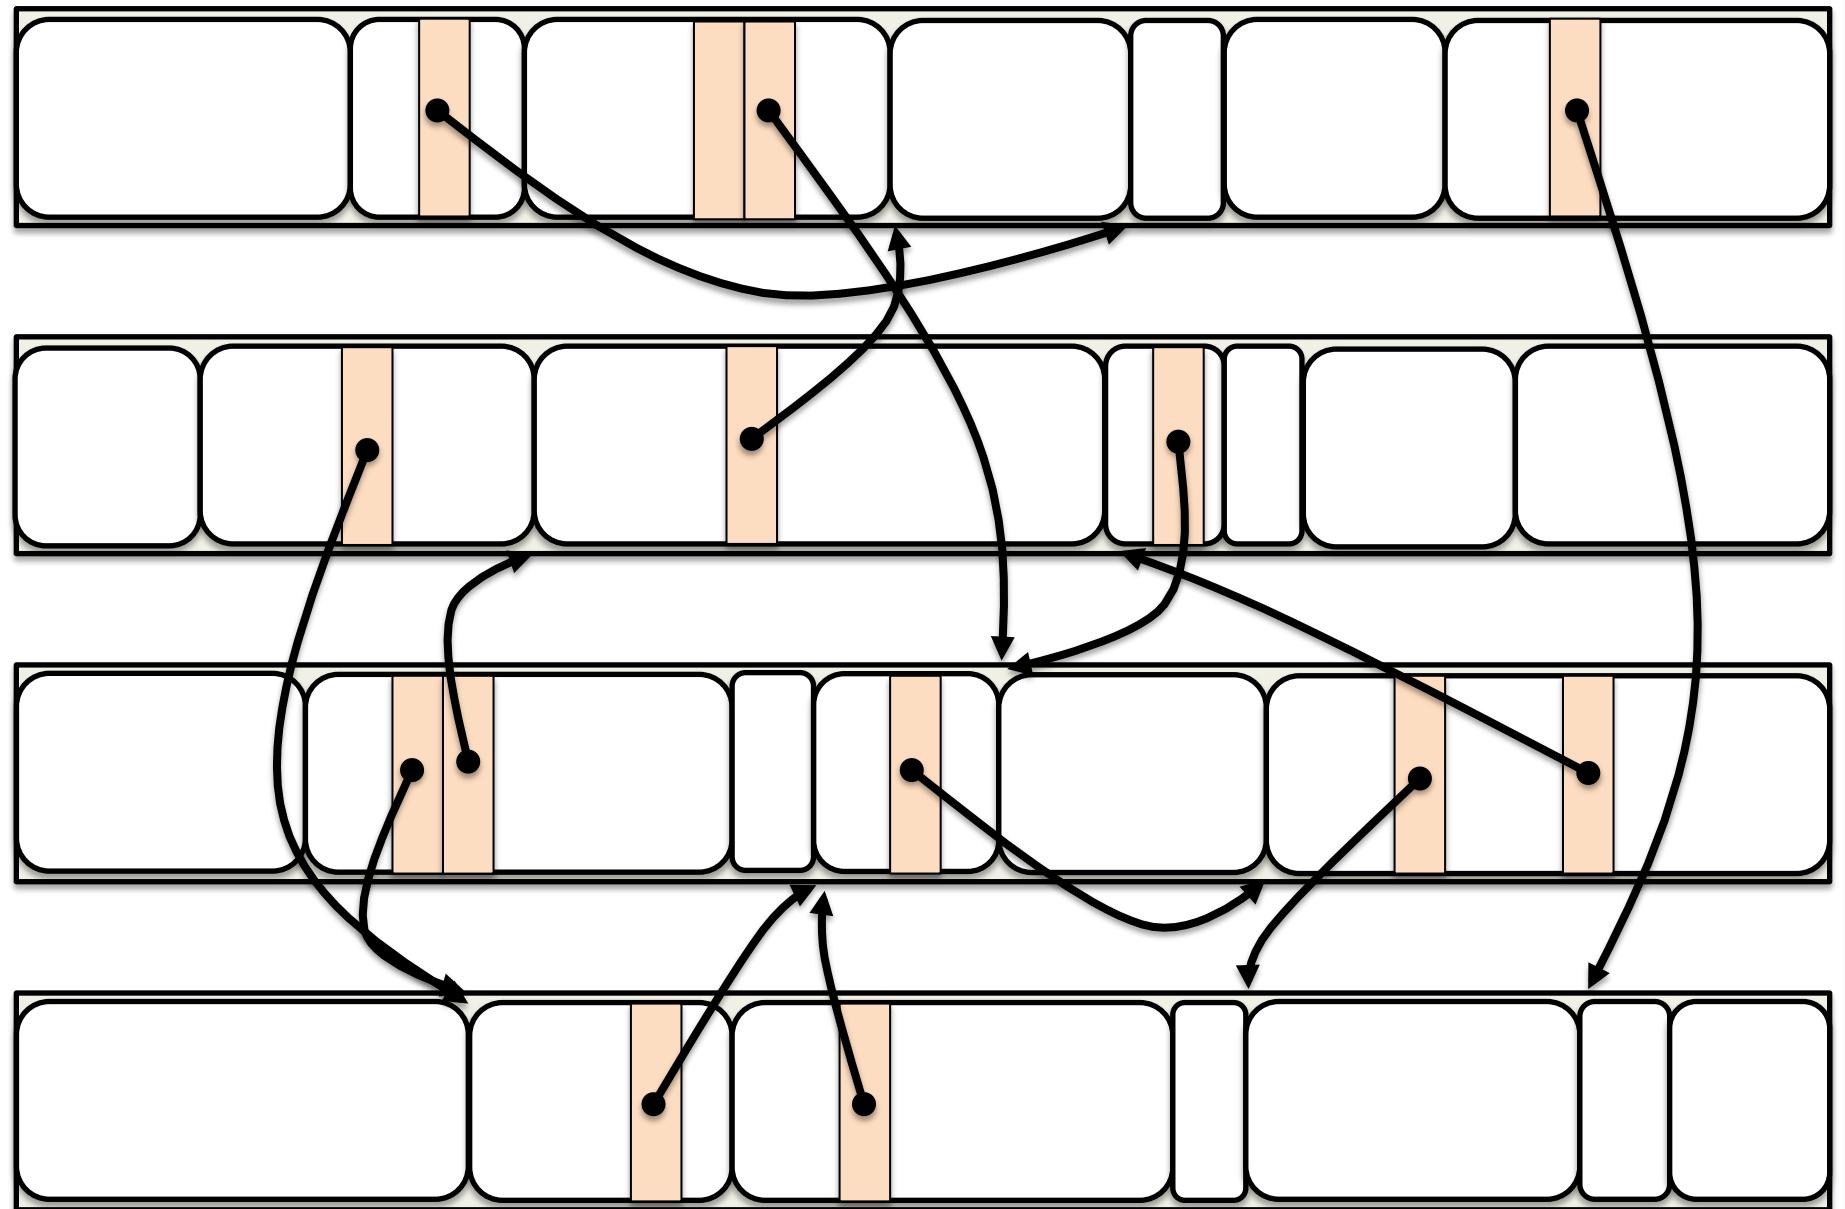
- The nursery is divided into smaller regions
 - Thread Local Allocation Buffer (TLAB)
- Each thread has its own TLAB for allocation
 - Store the buffer details in thread-local storage
 - No contention for the allocator
- Nursery maintains a set of available TLABs
 - Threads grab a new one when needed
 - Contention on the TLAB store
 - Garbage collect when there are none left

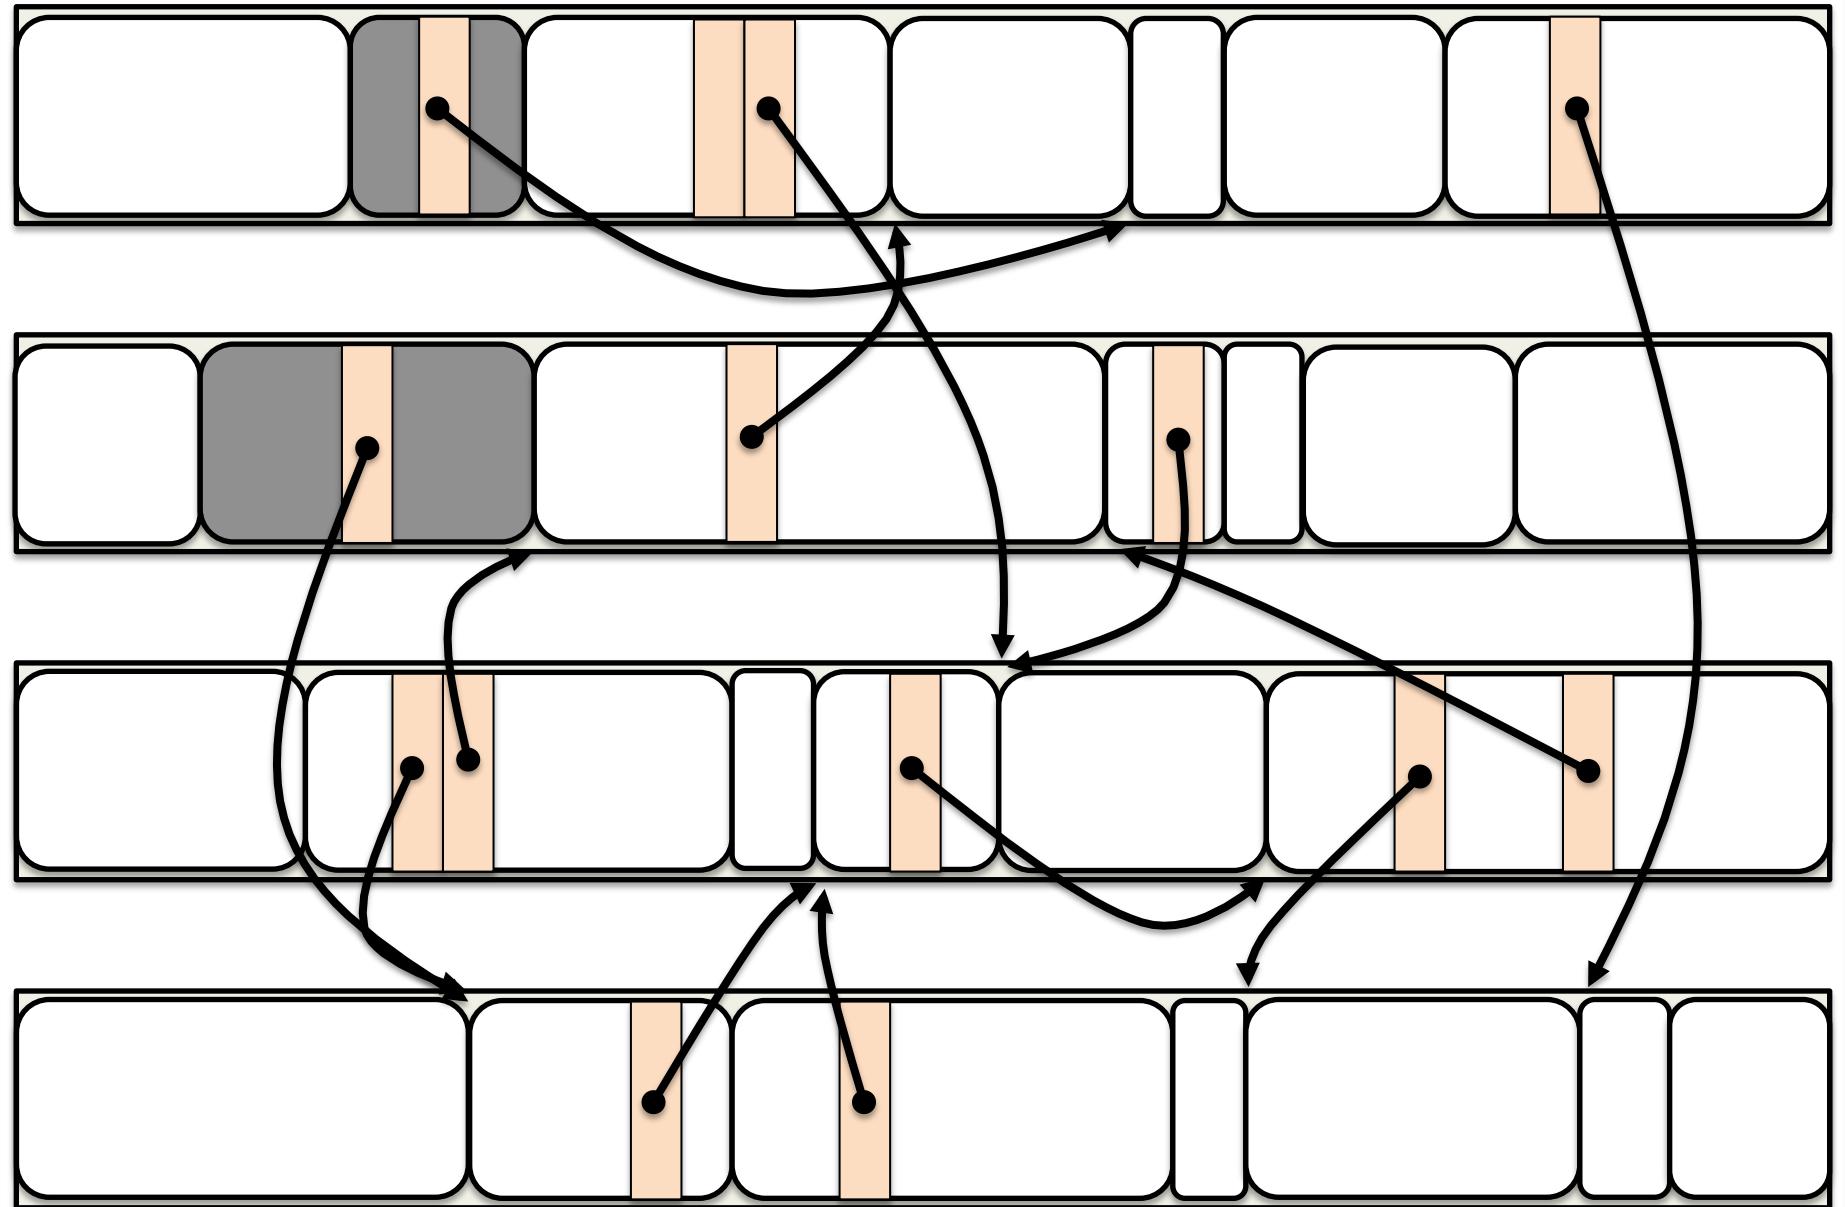
Concurrent Marking

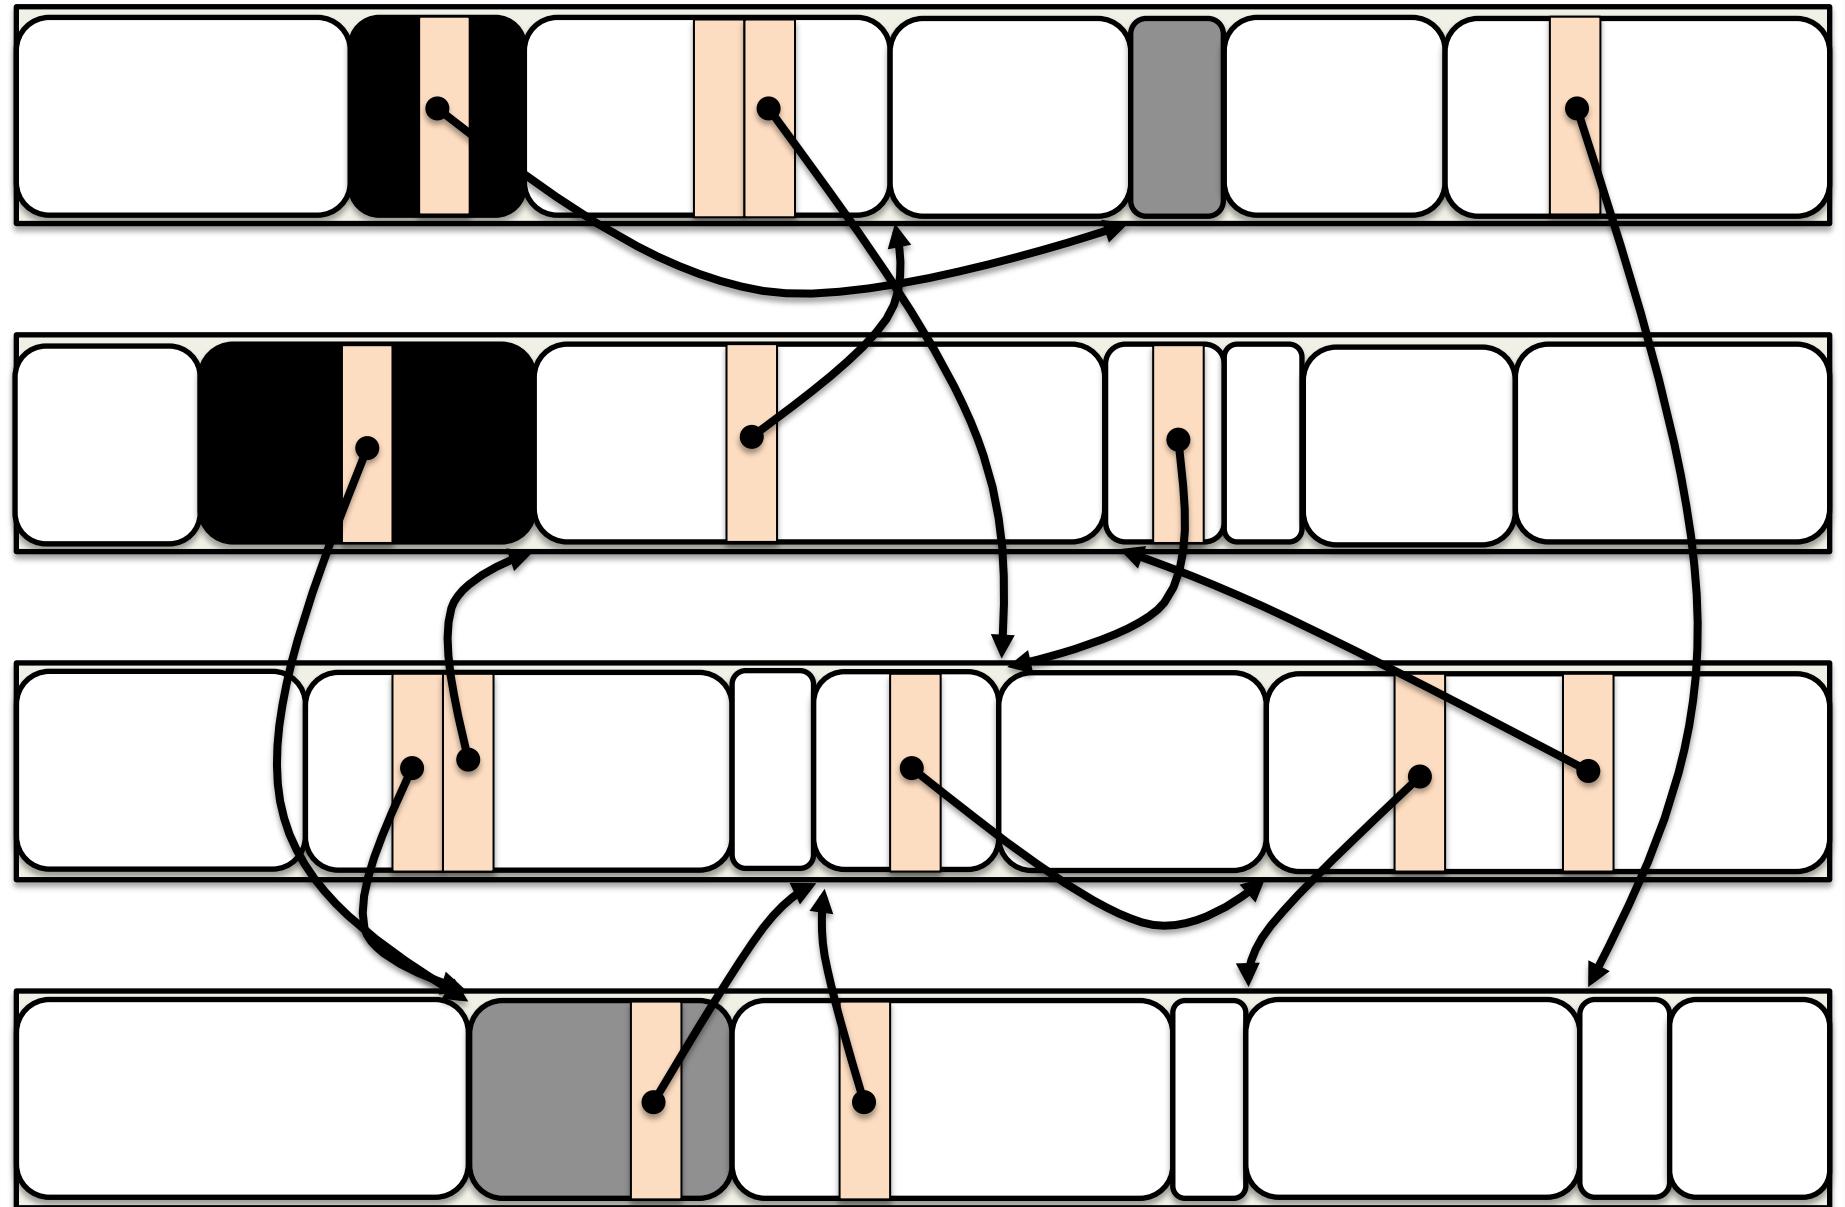
- Marking a large region can be expensive
 - Even with remembered sets
 - Complex write barriers come with other overheads
- Marking can largely be performed concurrently
 - Background threads mark as the mutators run
 - Garbage objects don't become live again
- What happens if the object graph changes?

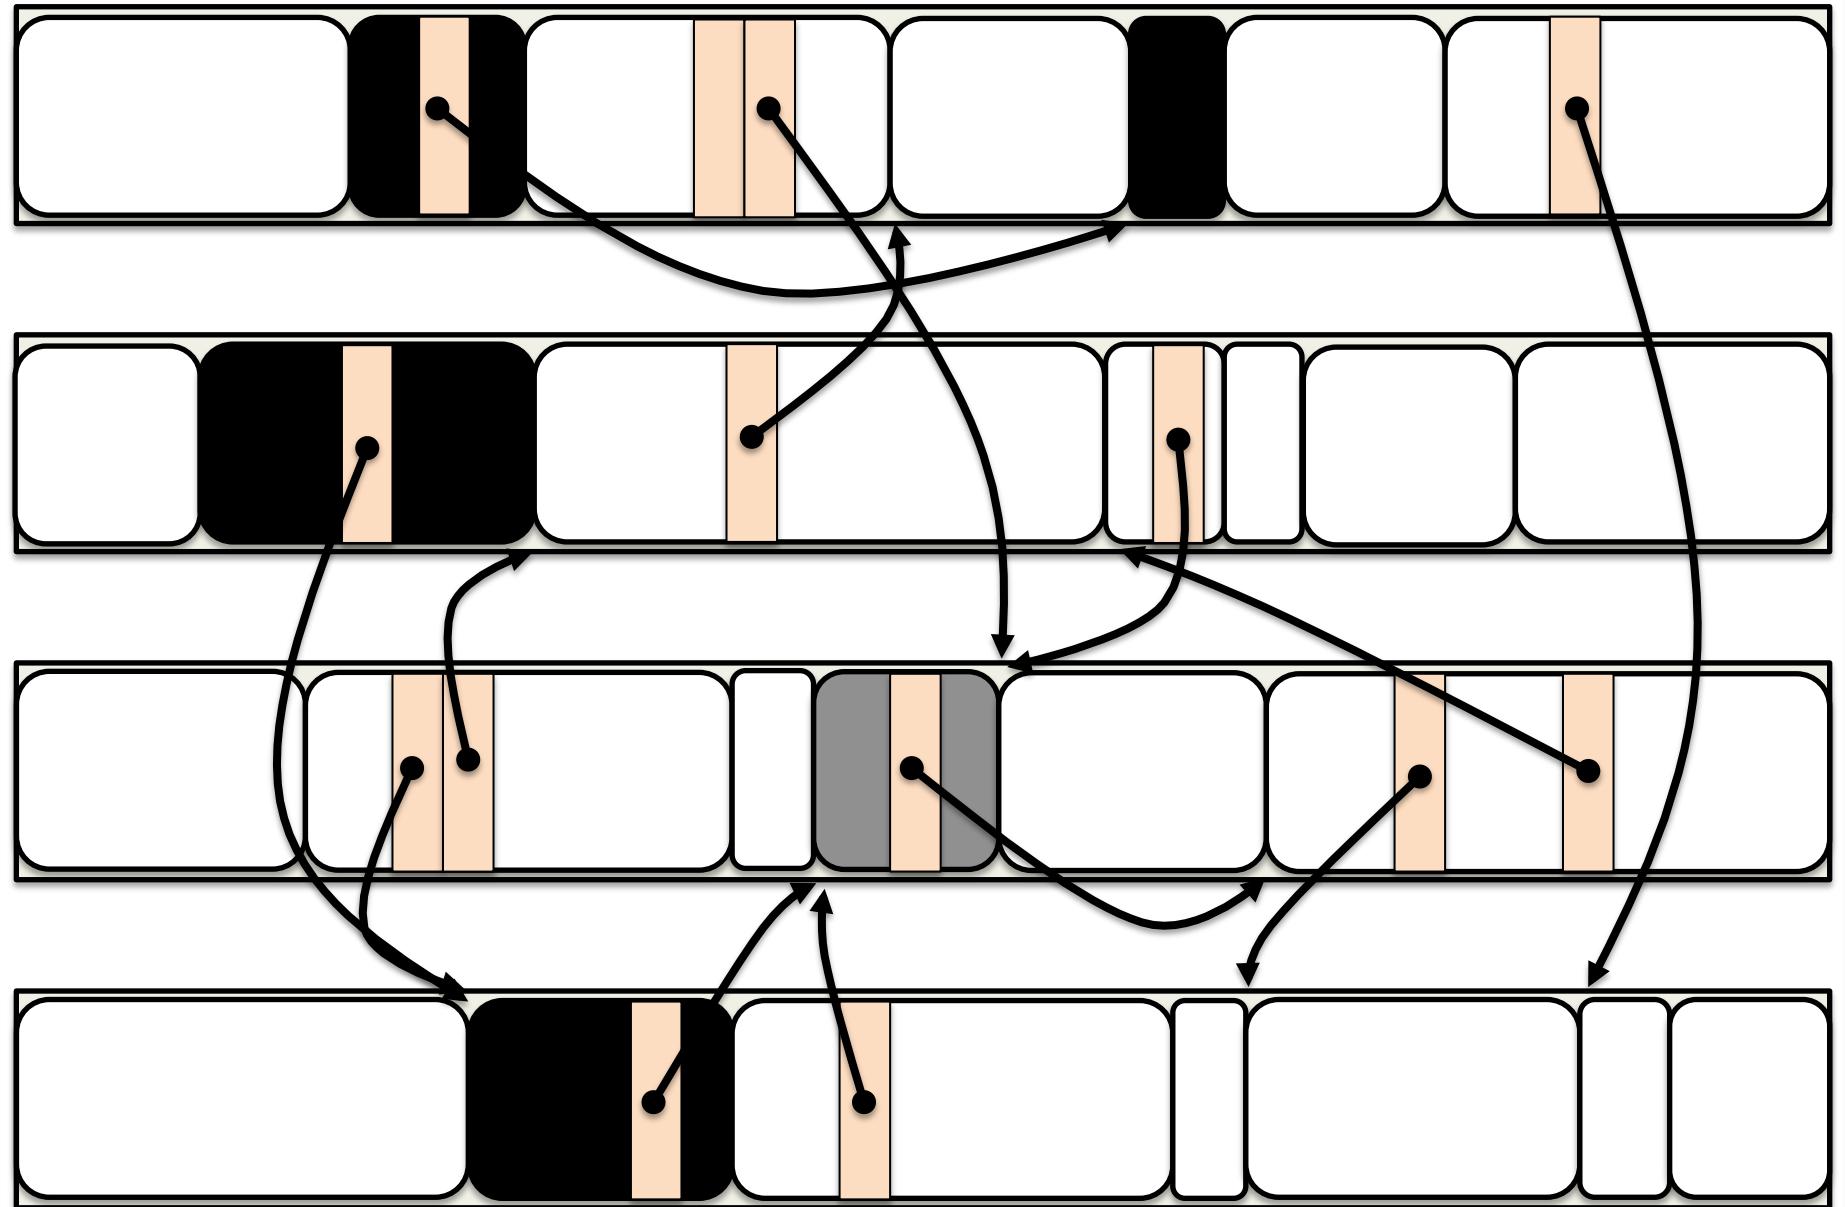
Tricolor Marking

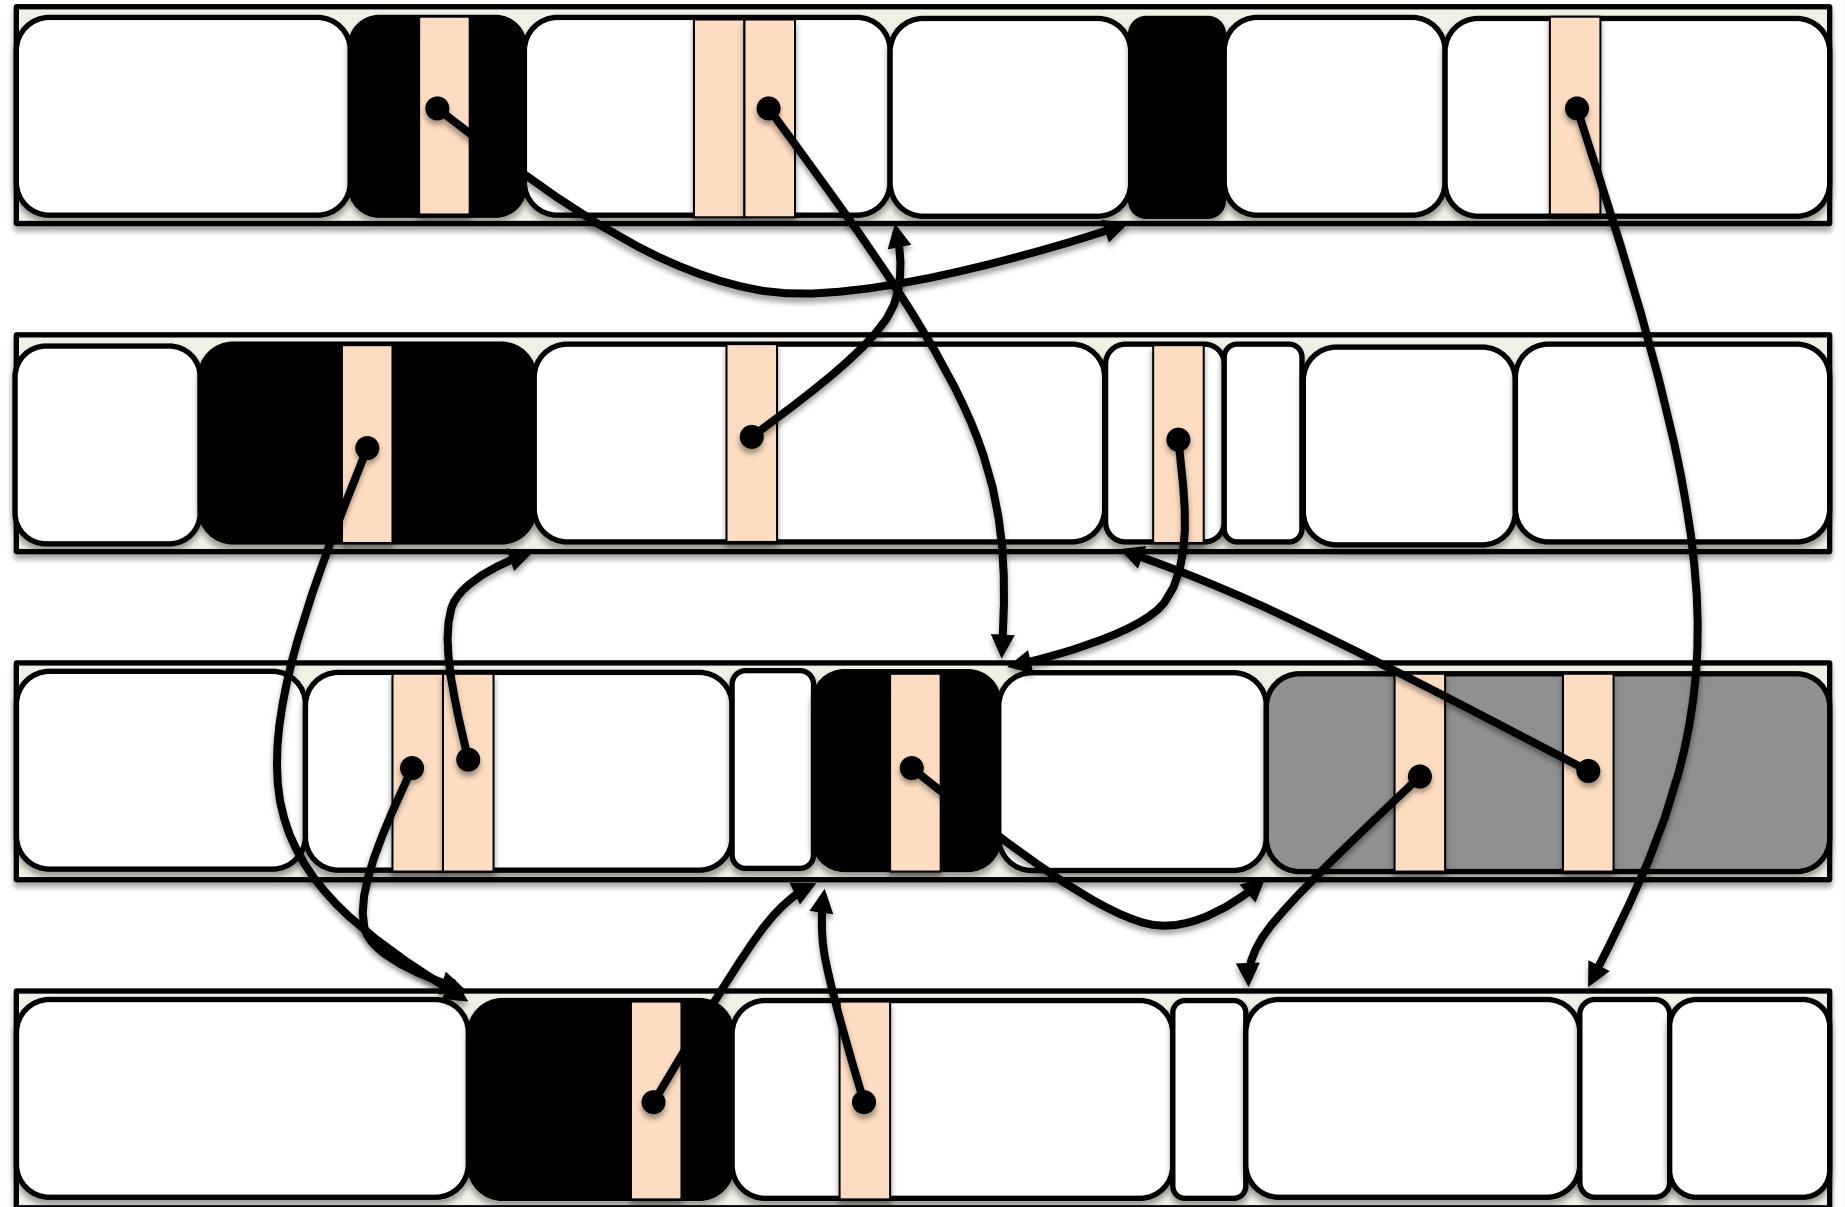
- Convenient way to think about marking
- Objects exist in three possible states
 - White
 - Gray
 - Black
- Maintain invariants between colors
 - Black objects can't point to white
 - Objects cannot go from black or gray to white

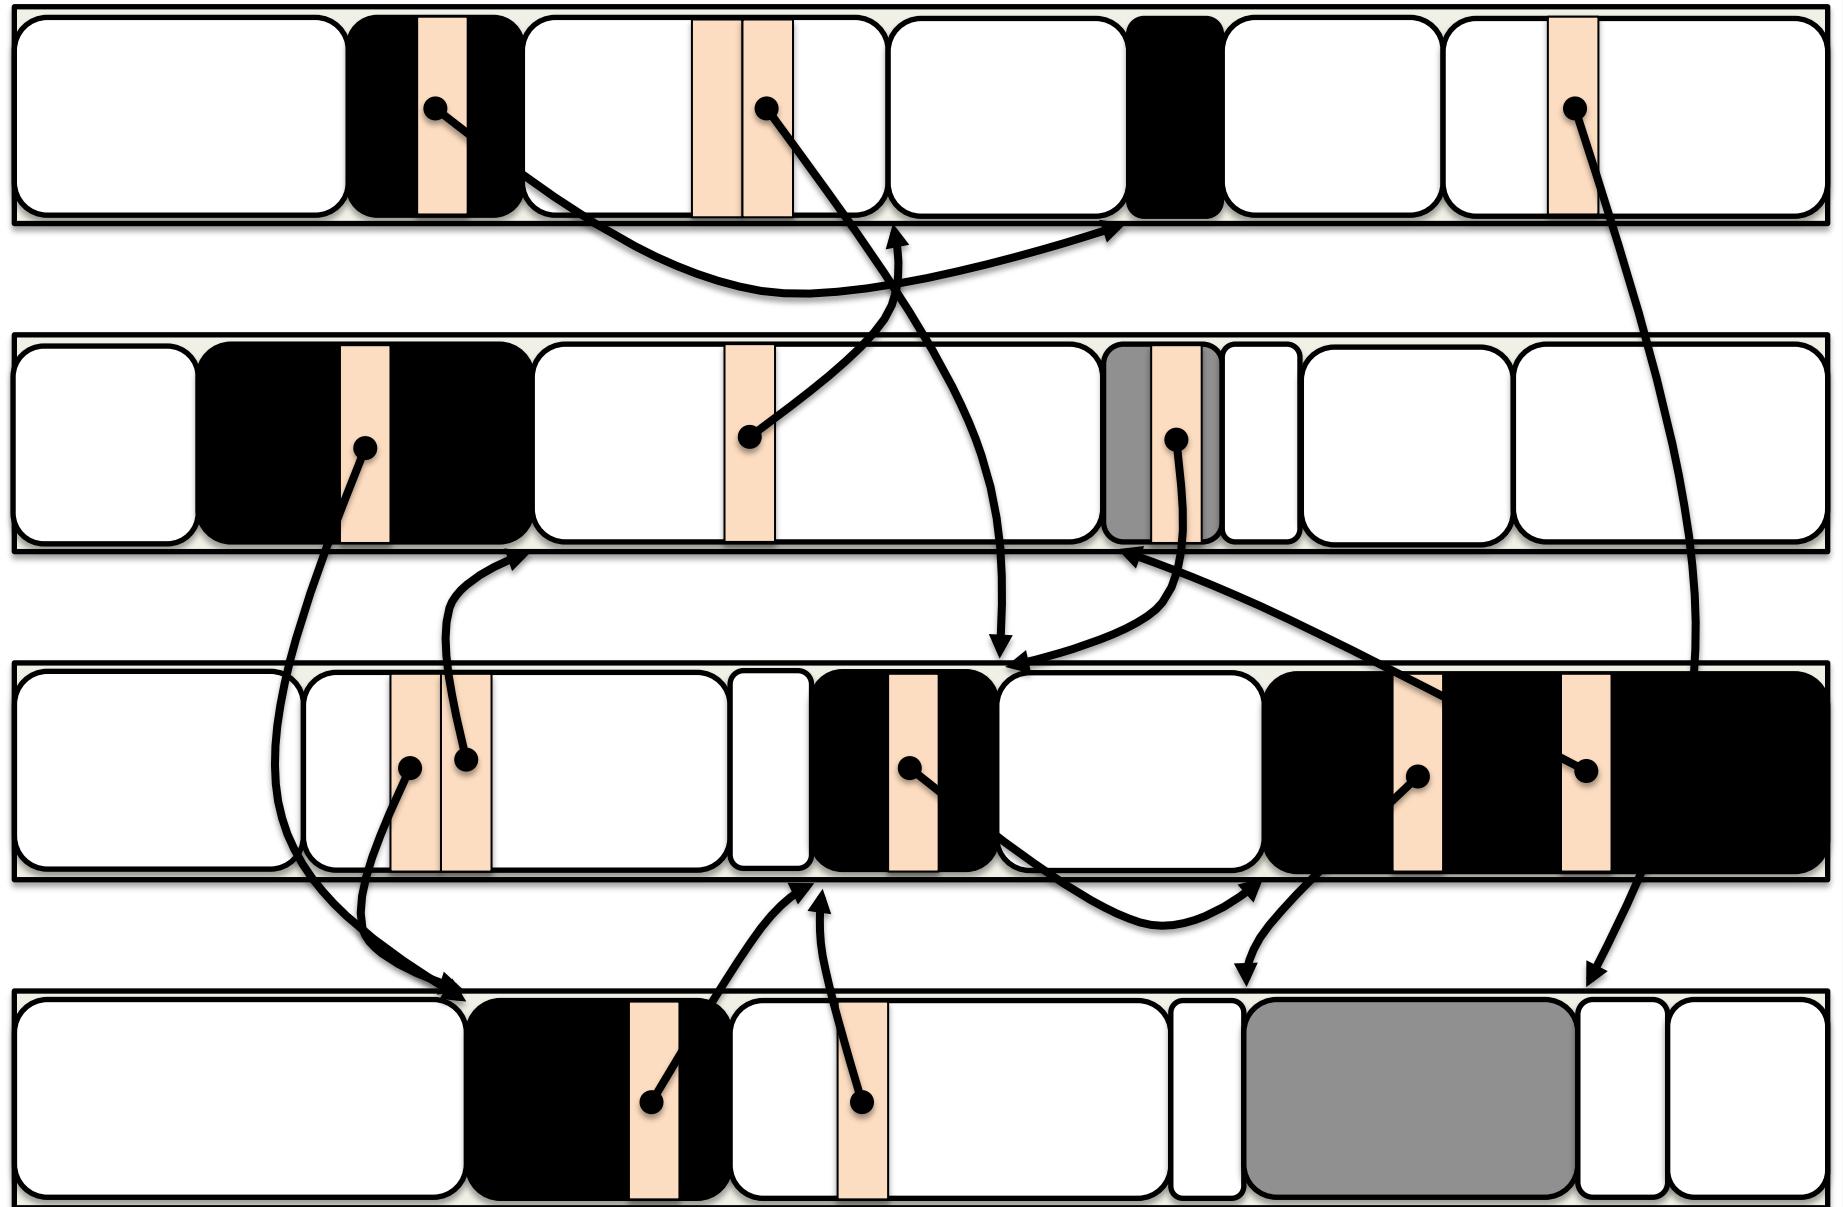


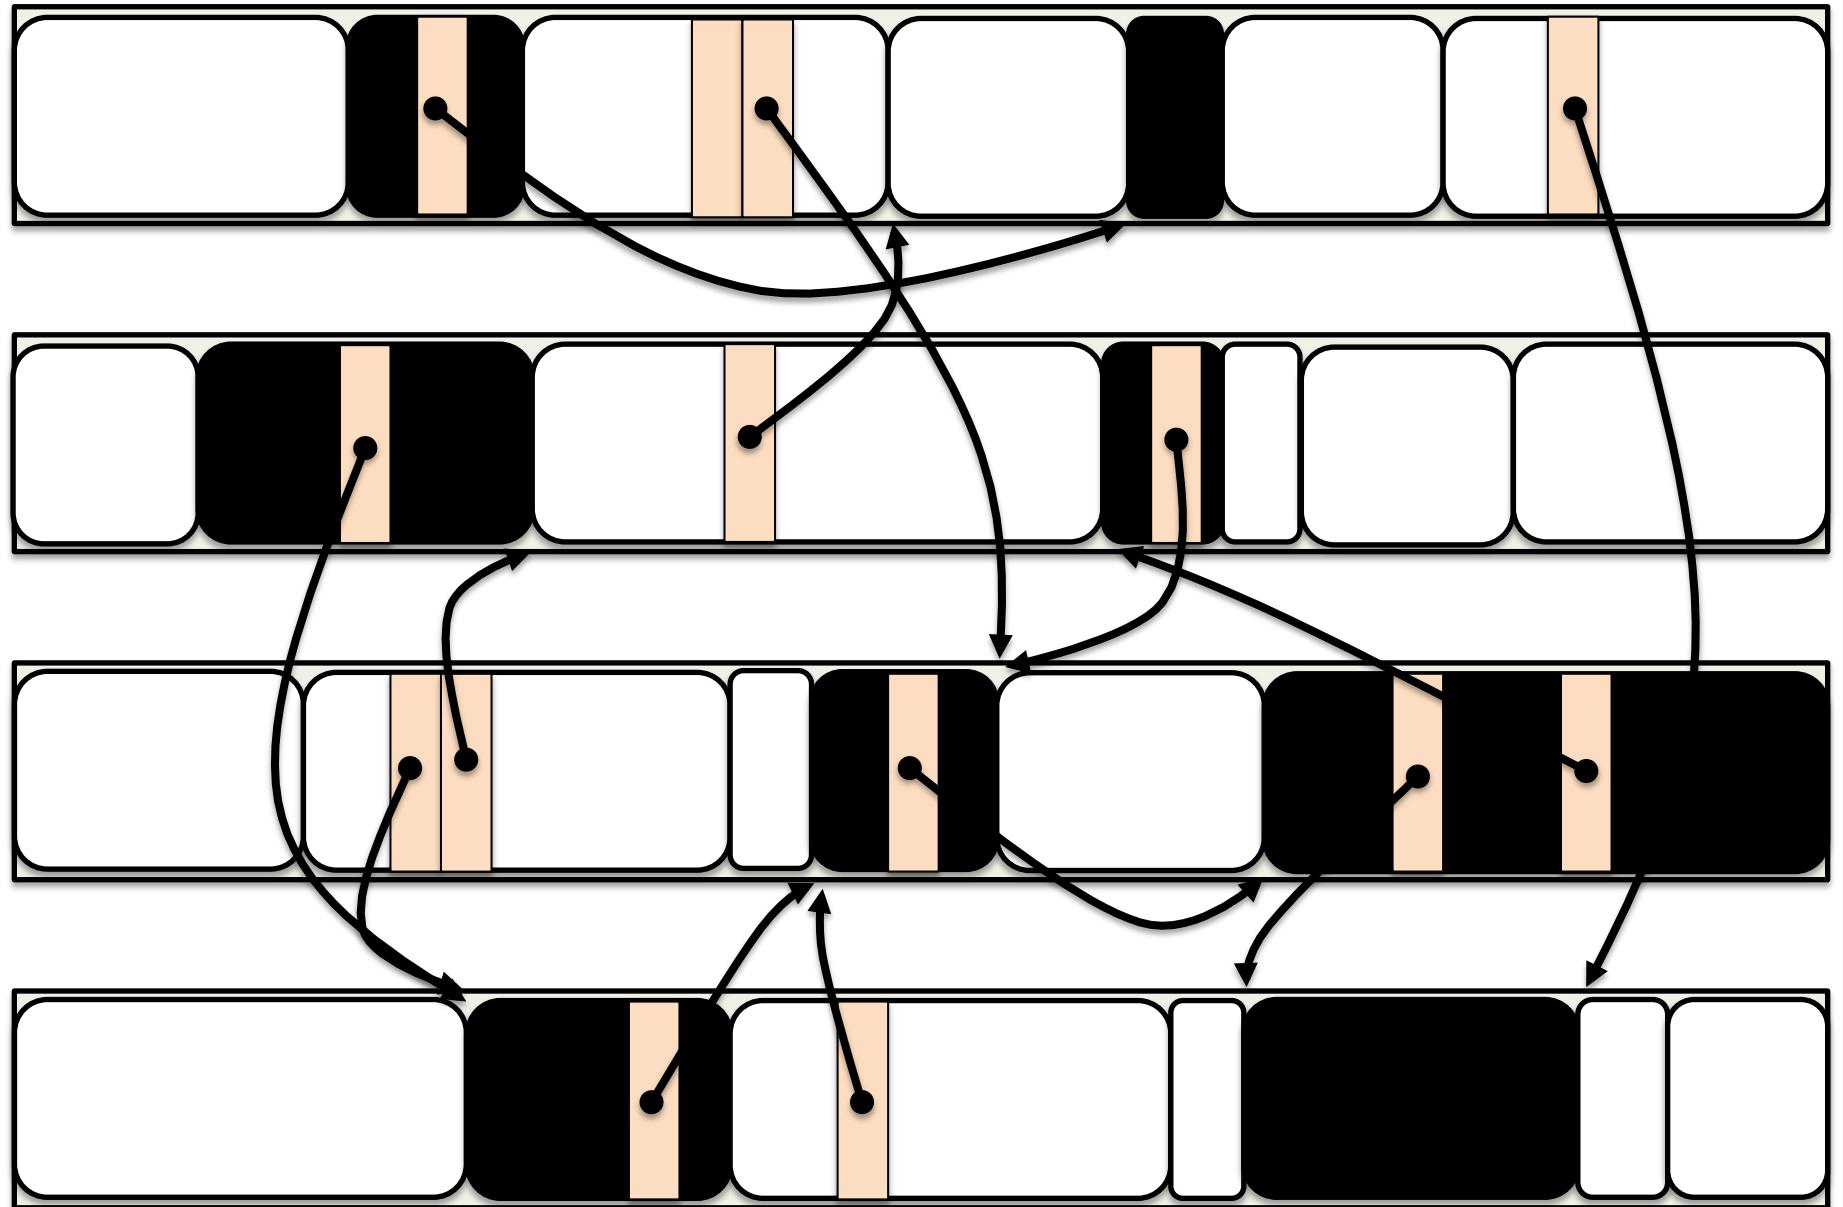






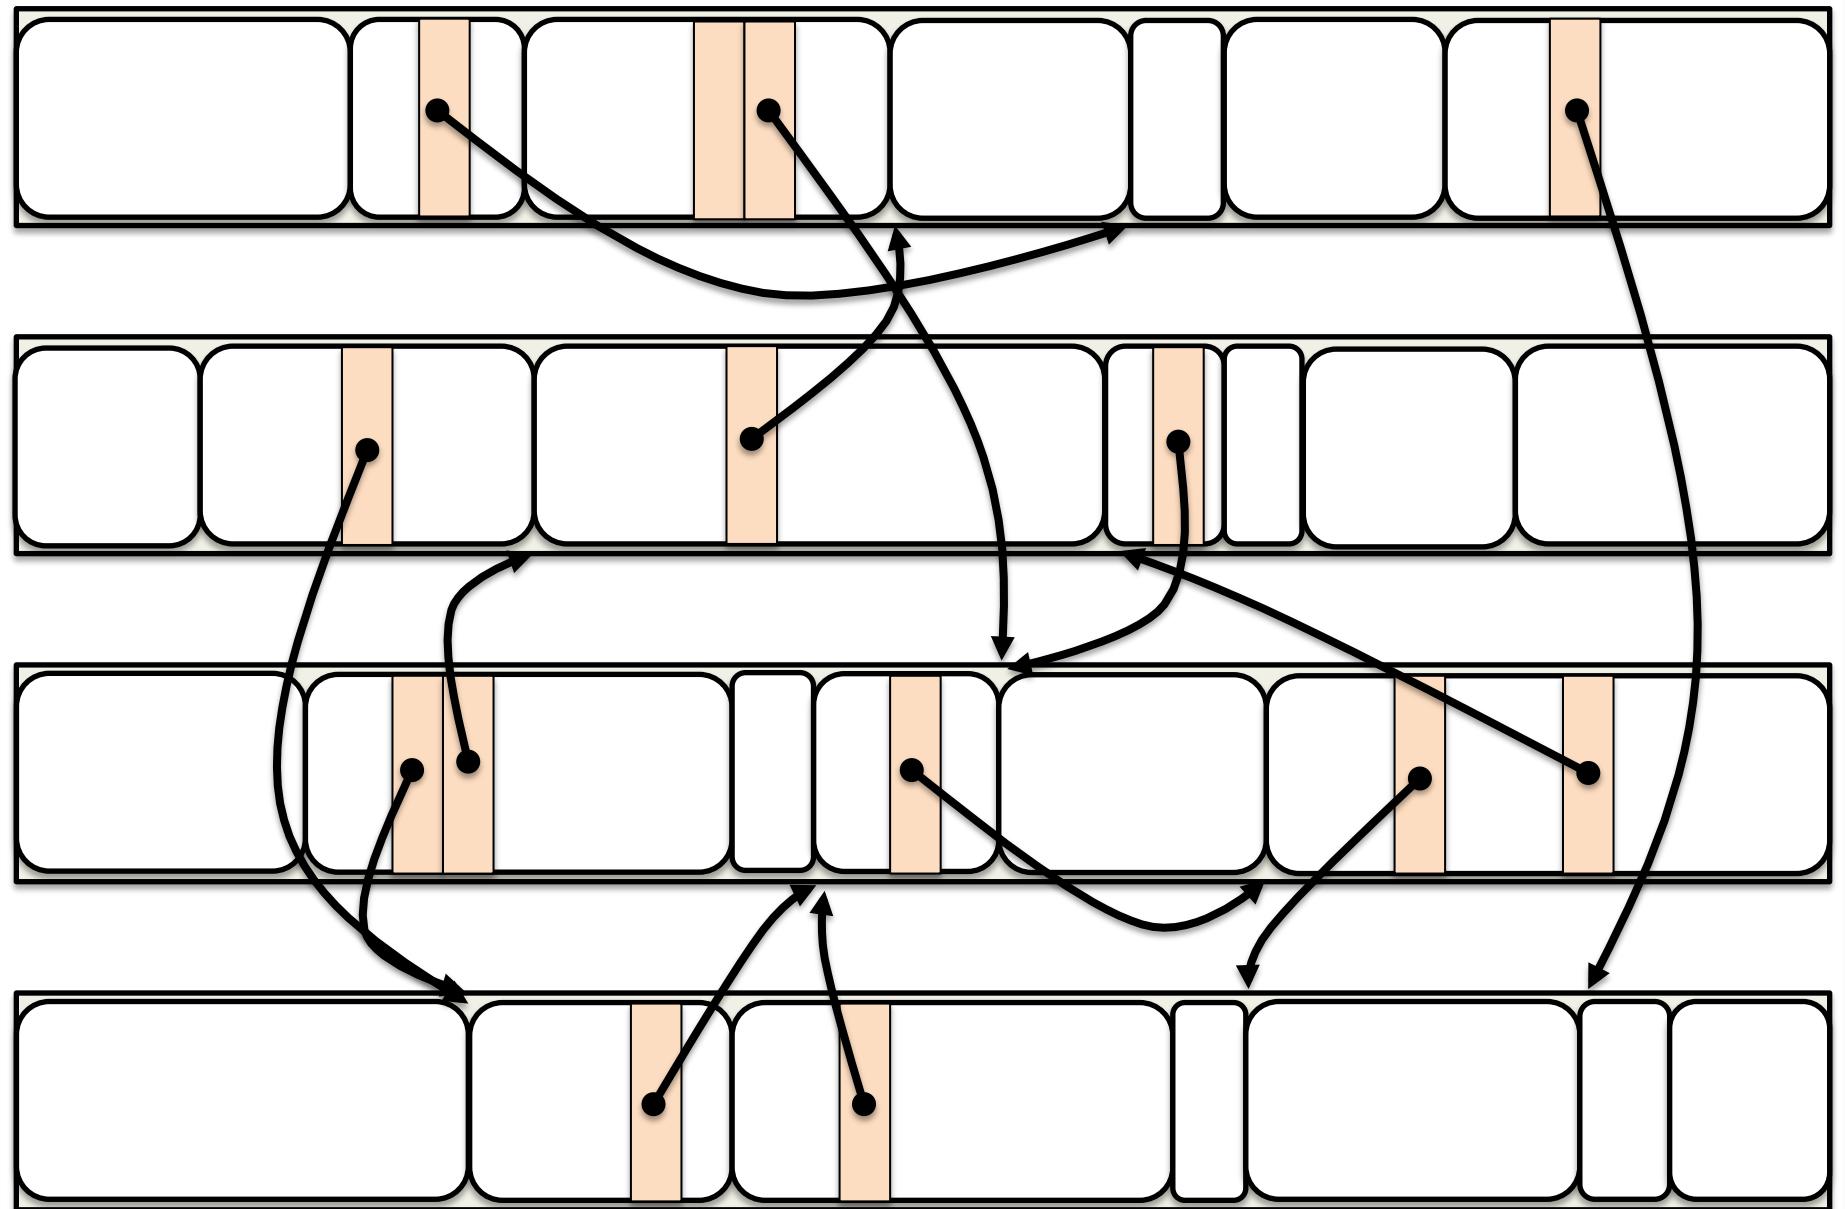


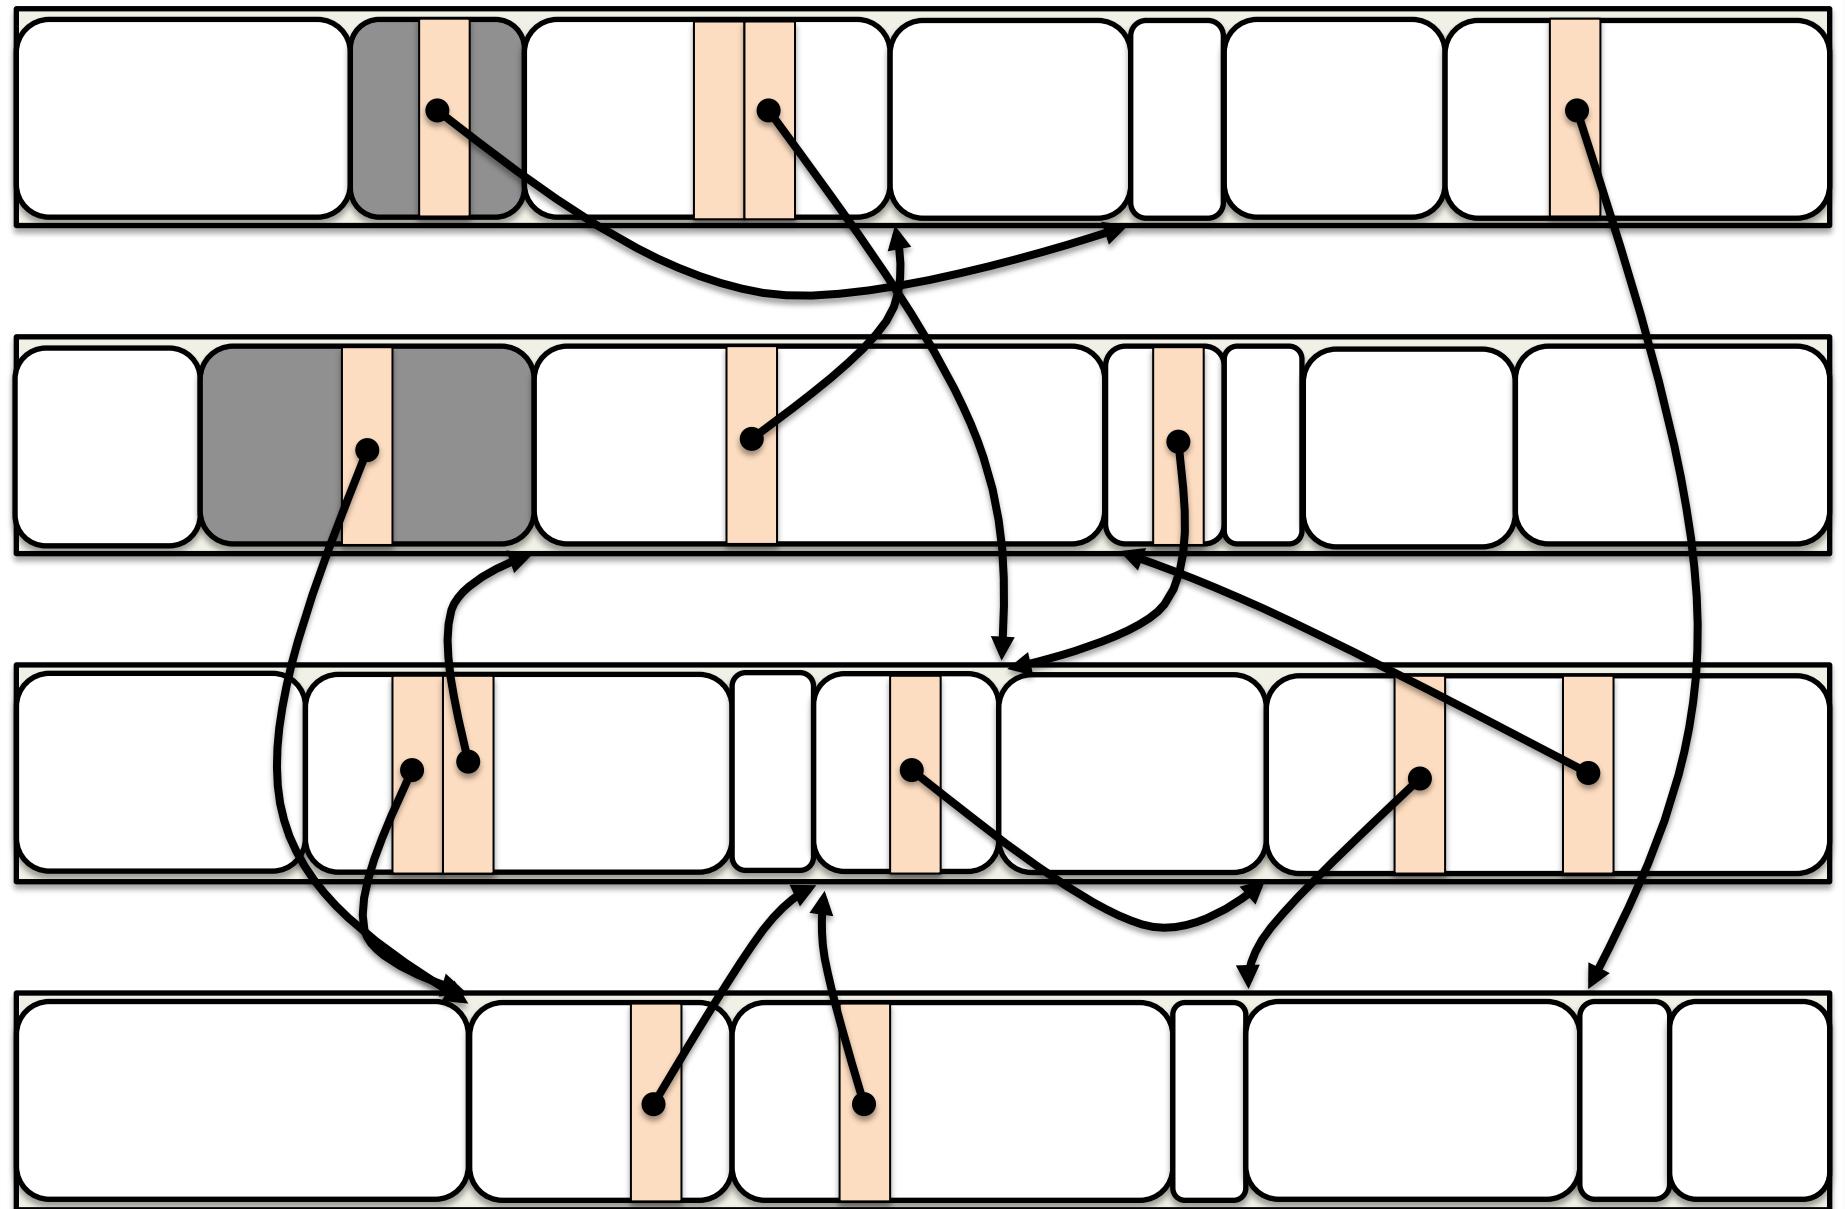


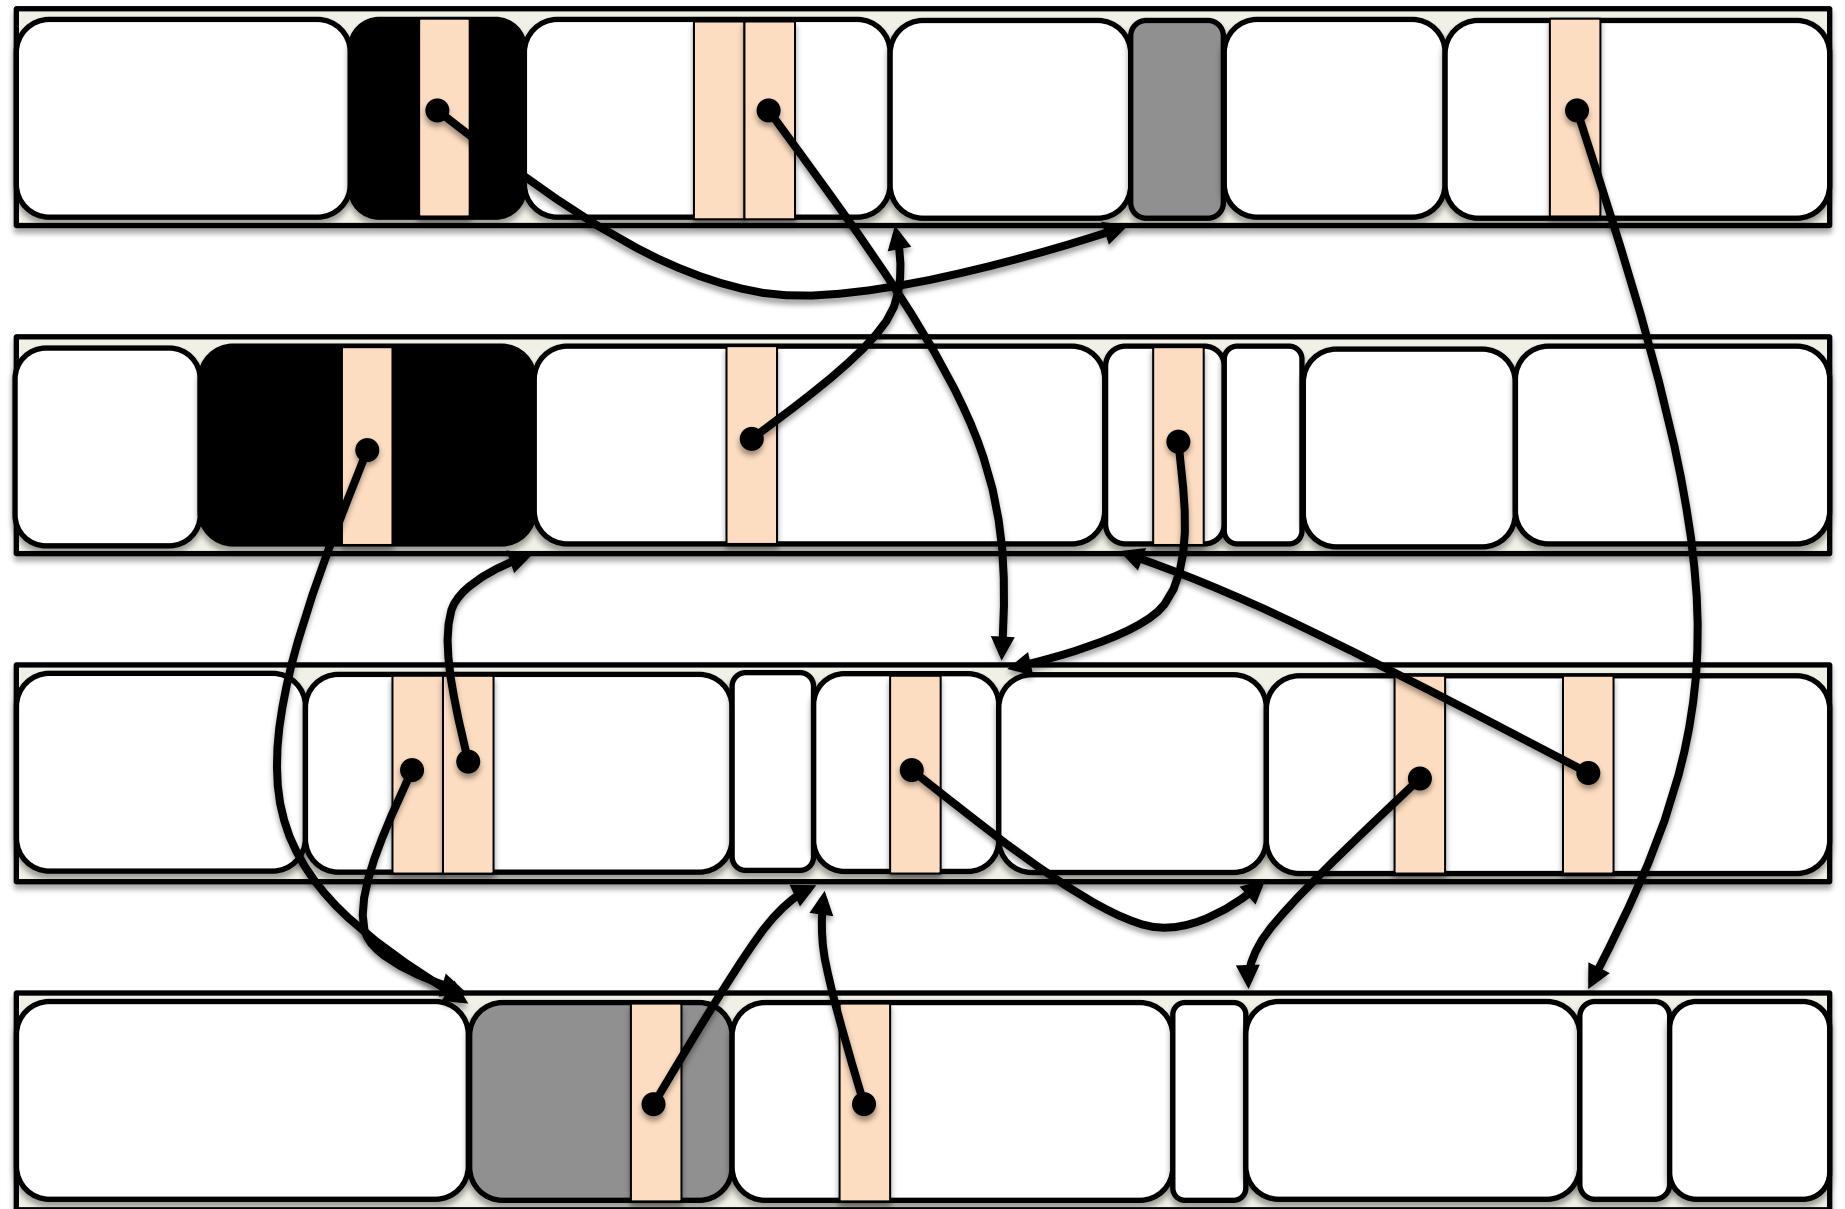


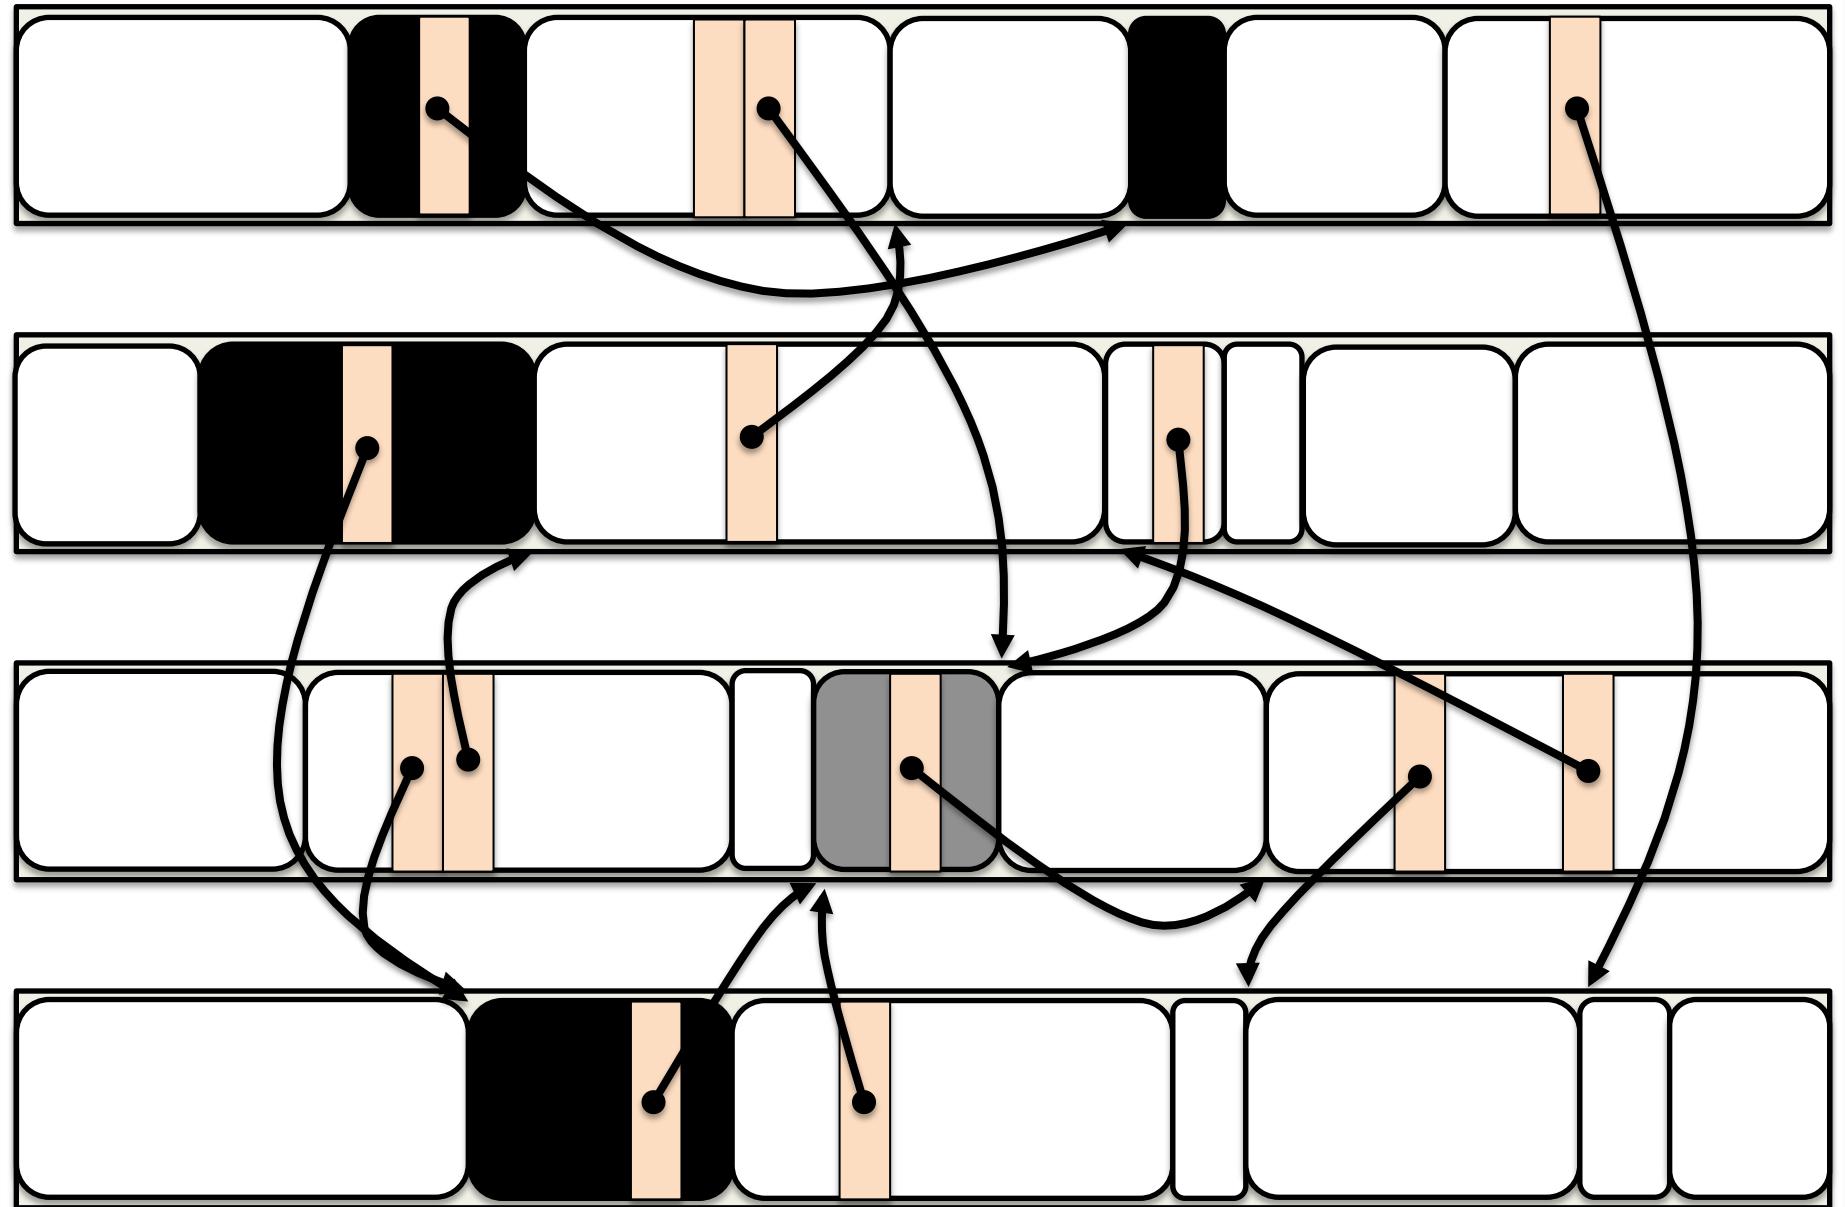
Snapshot At The Beginning

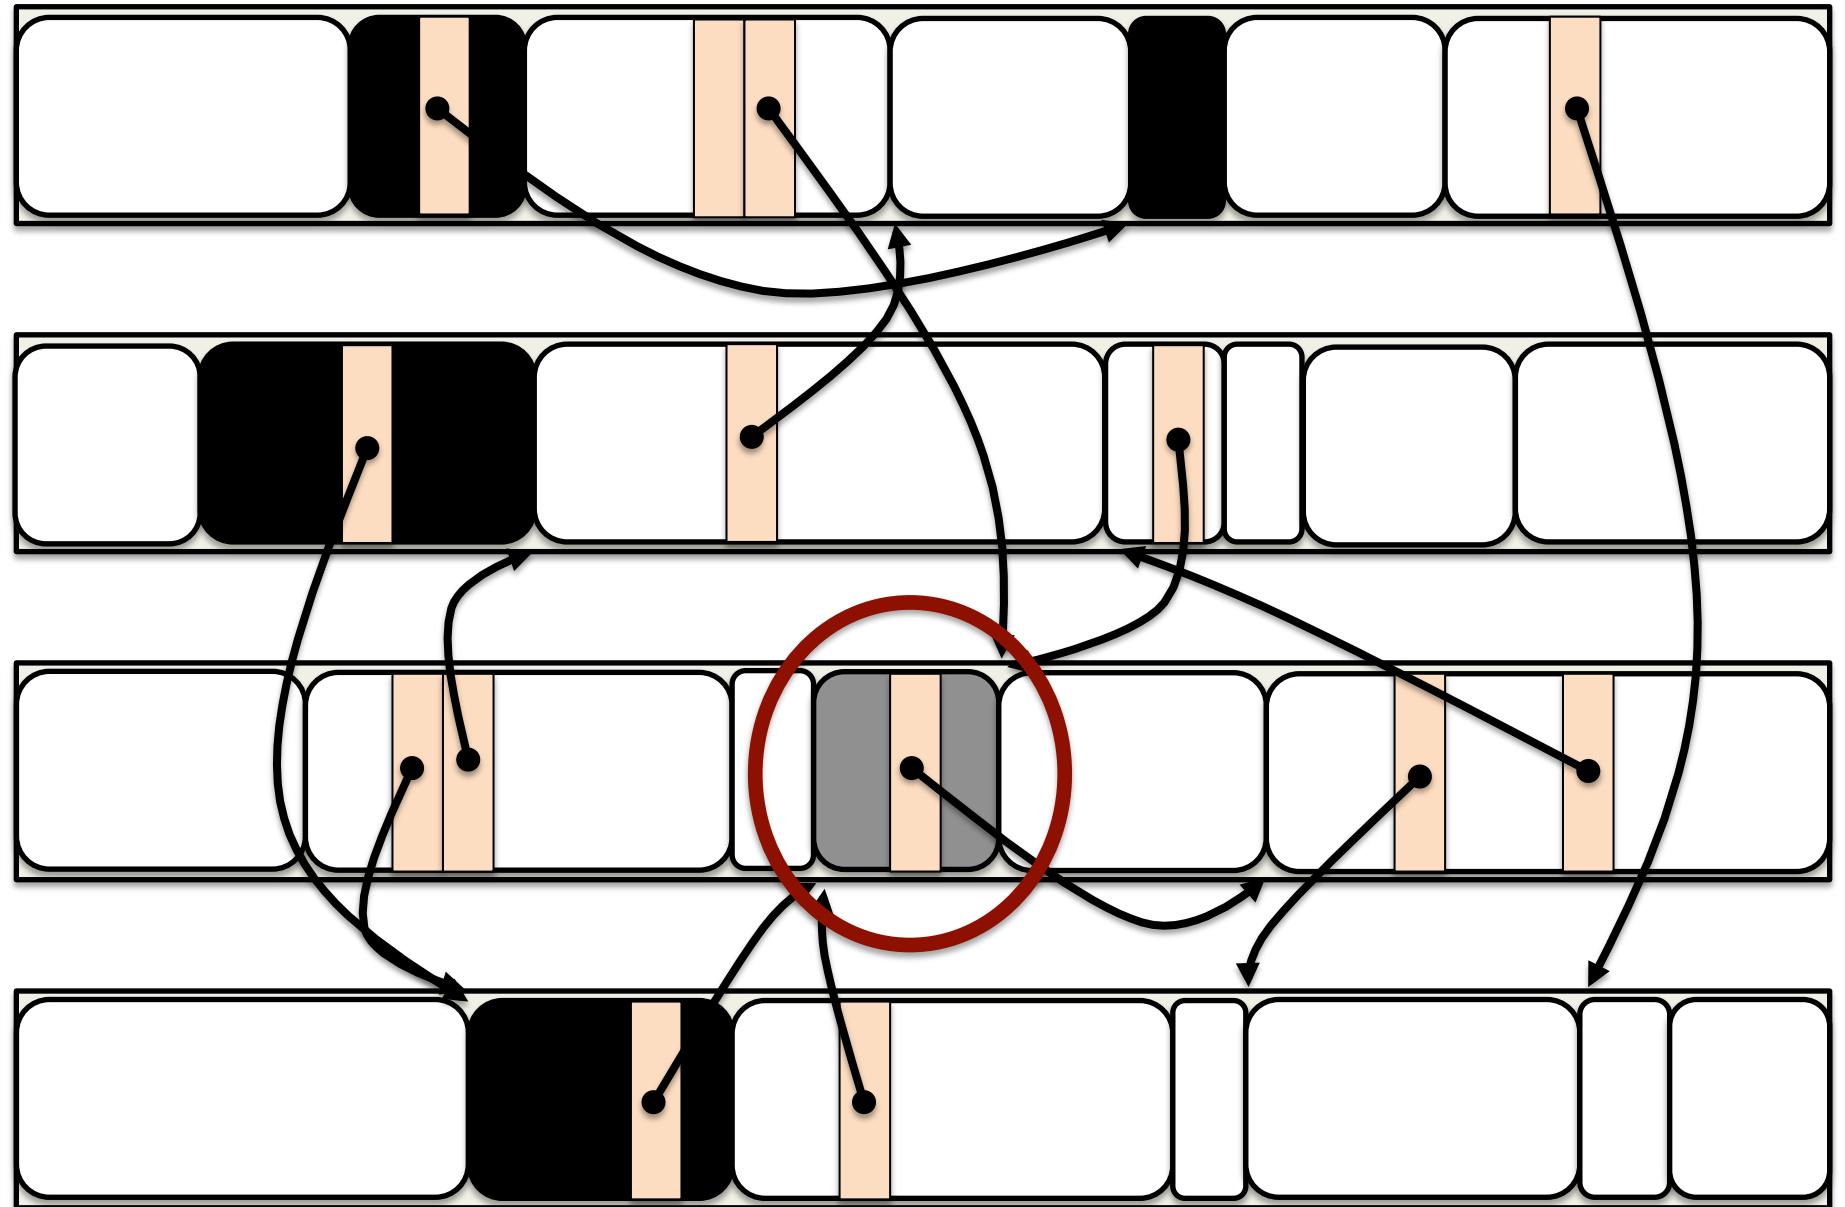
- GC requires that all live objects are marked
 - Does not require that all garbage is unmarked
- Concurrent algorithm works from a point in time
 - Mark all objects that were live at that point
 - Plus all newly-allocated objects
- Implies that all live objects are now marked
 - As well as newly-allocated garbage

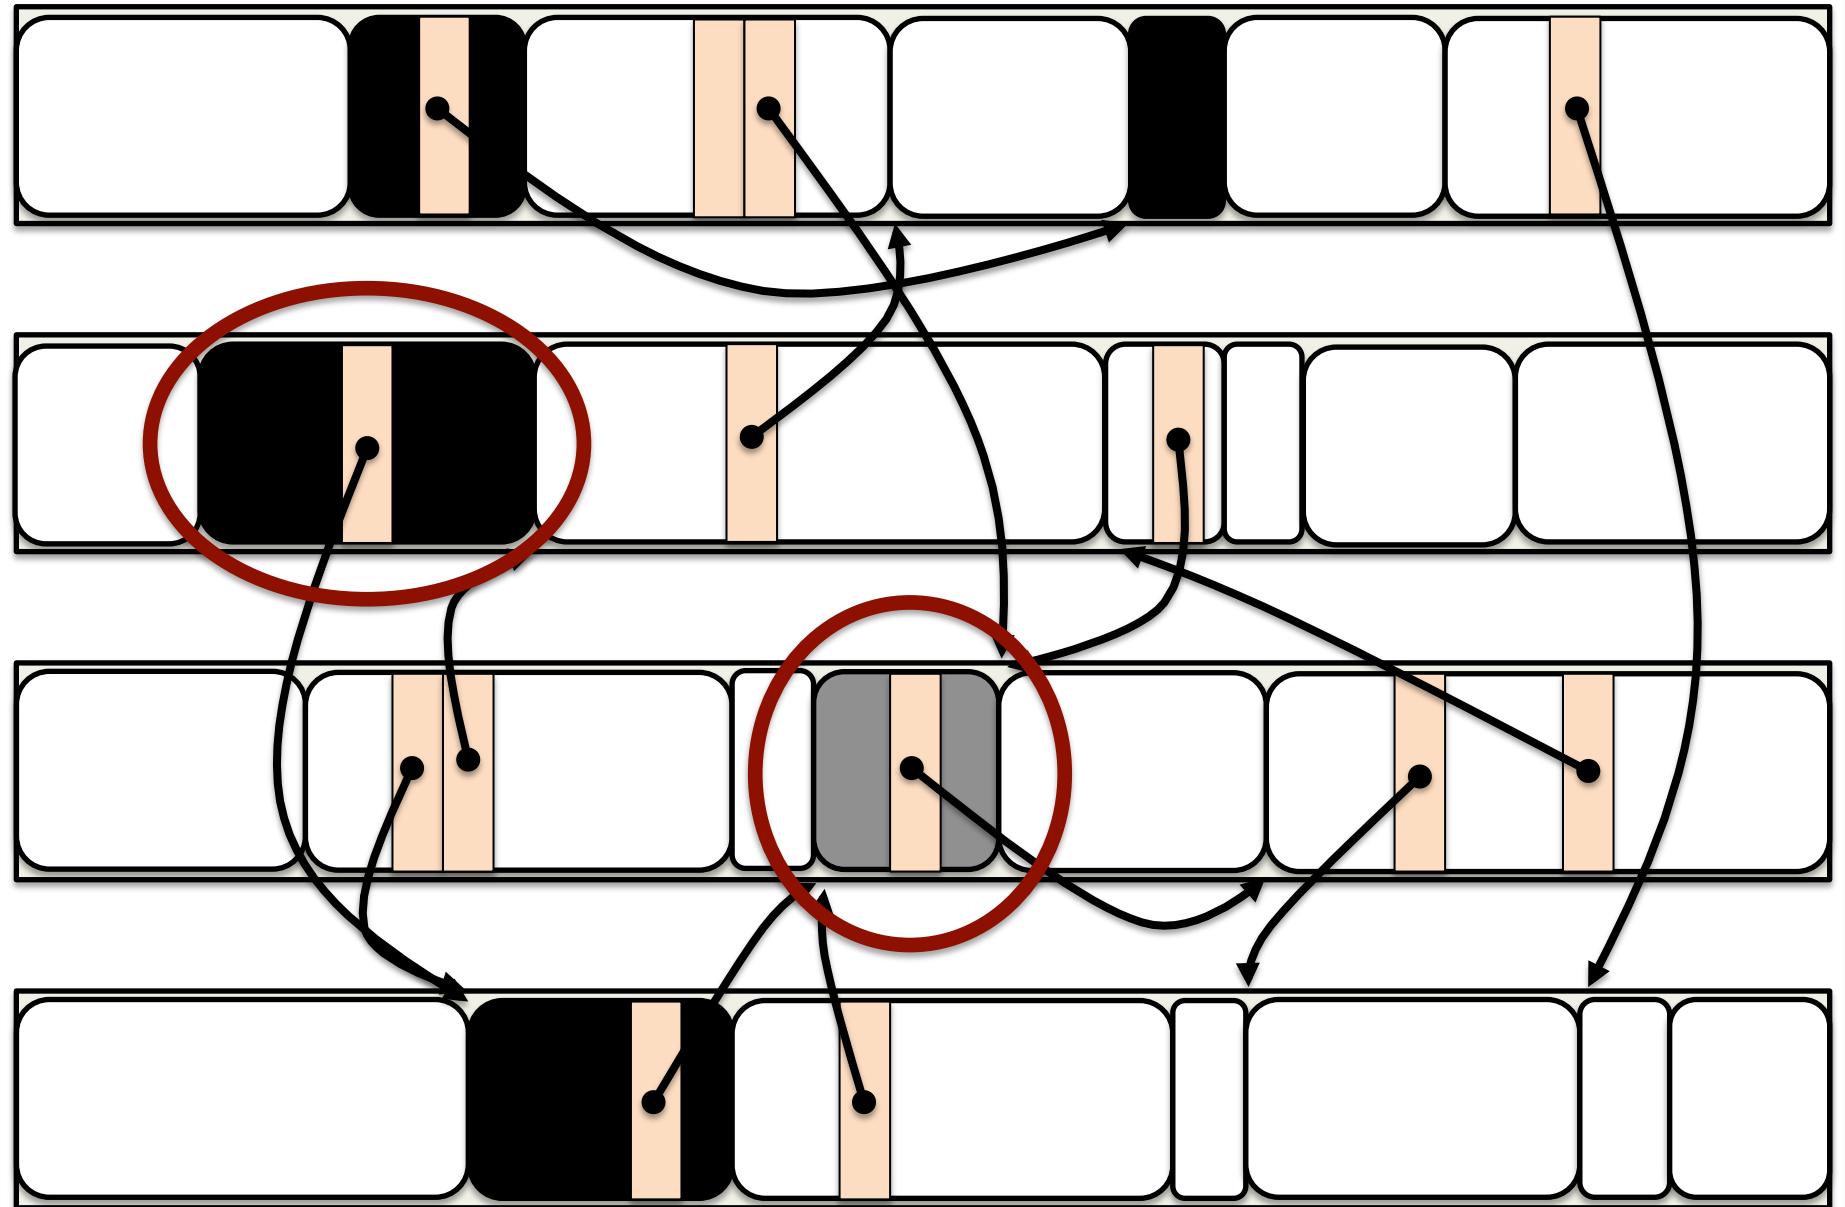


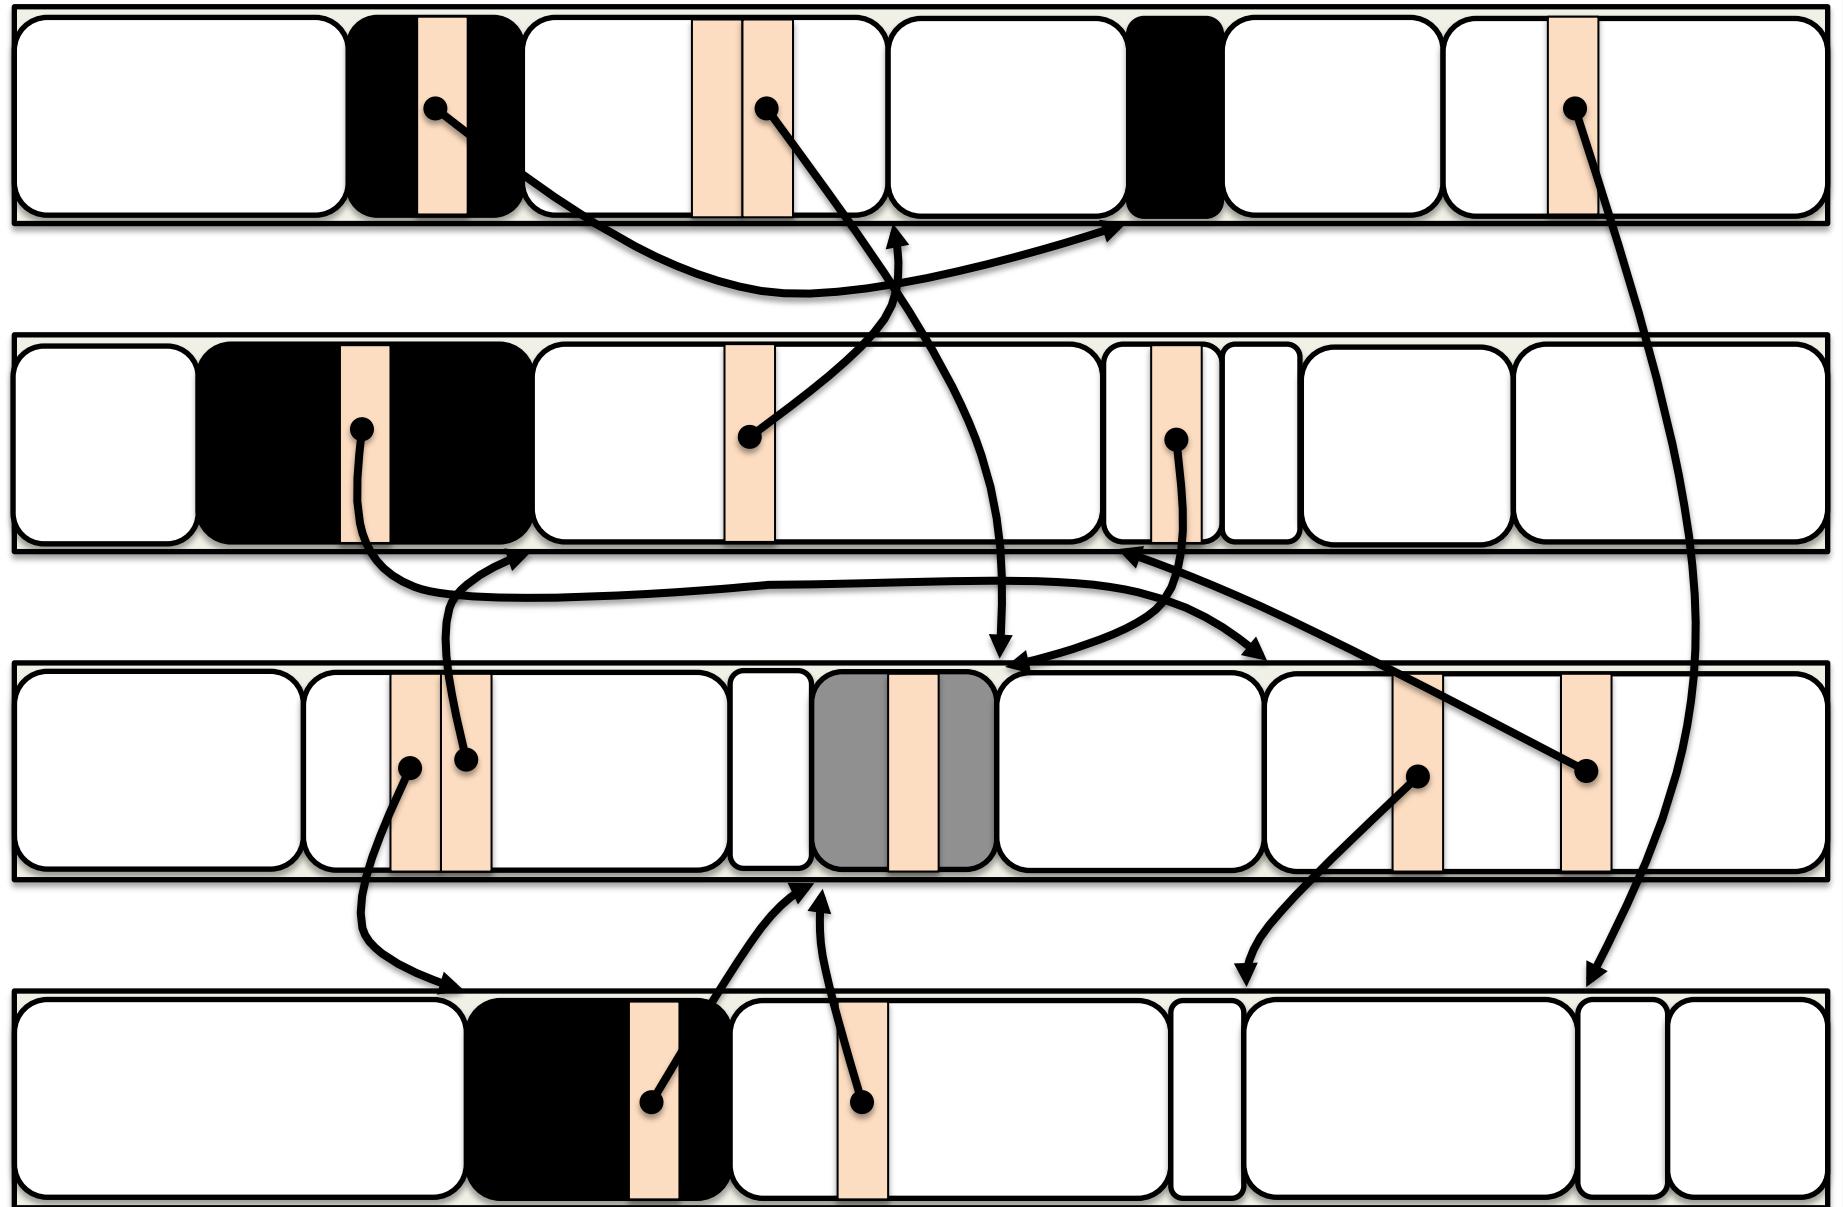


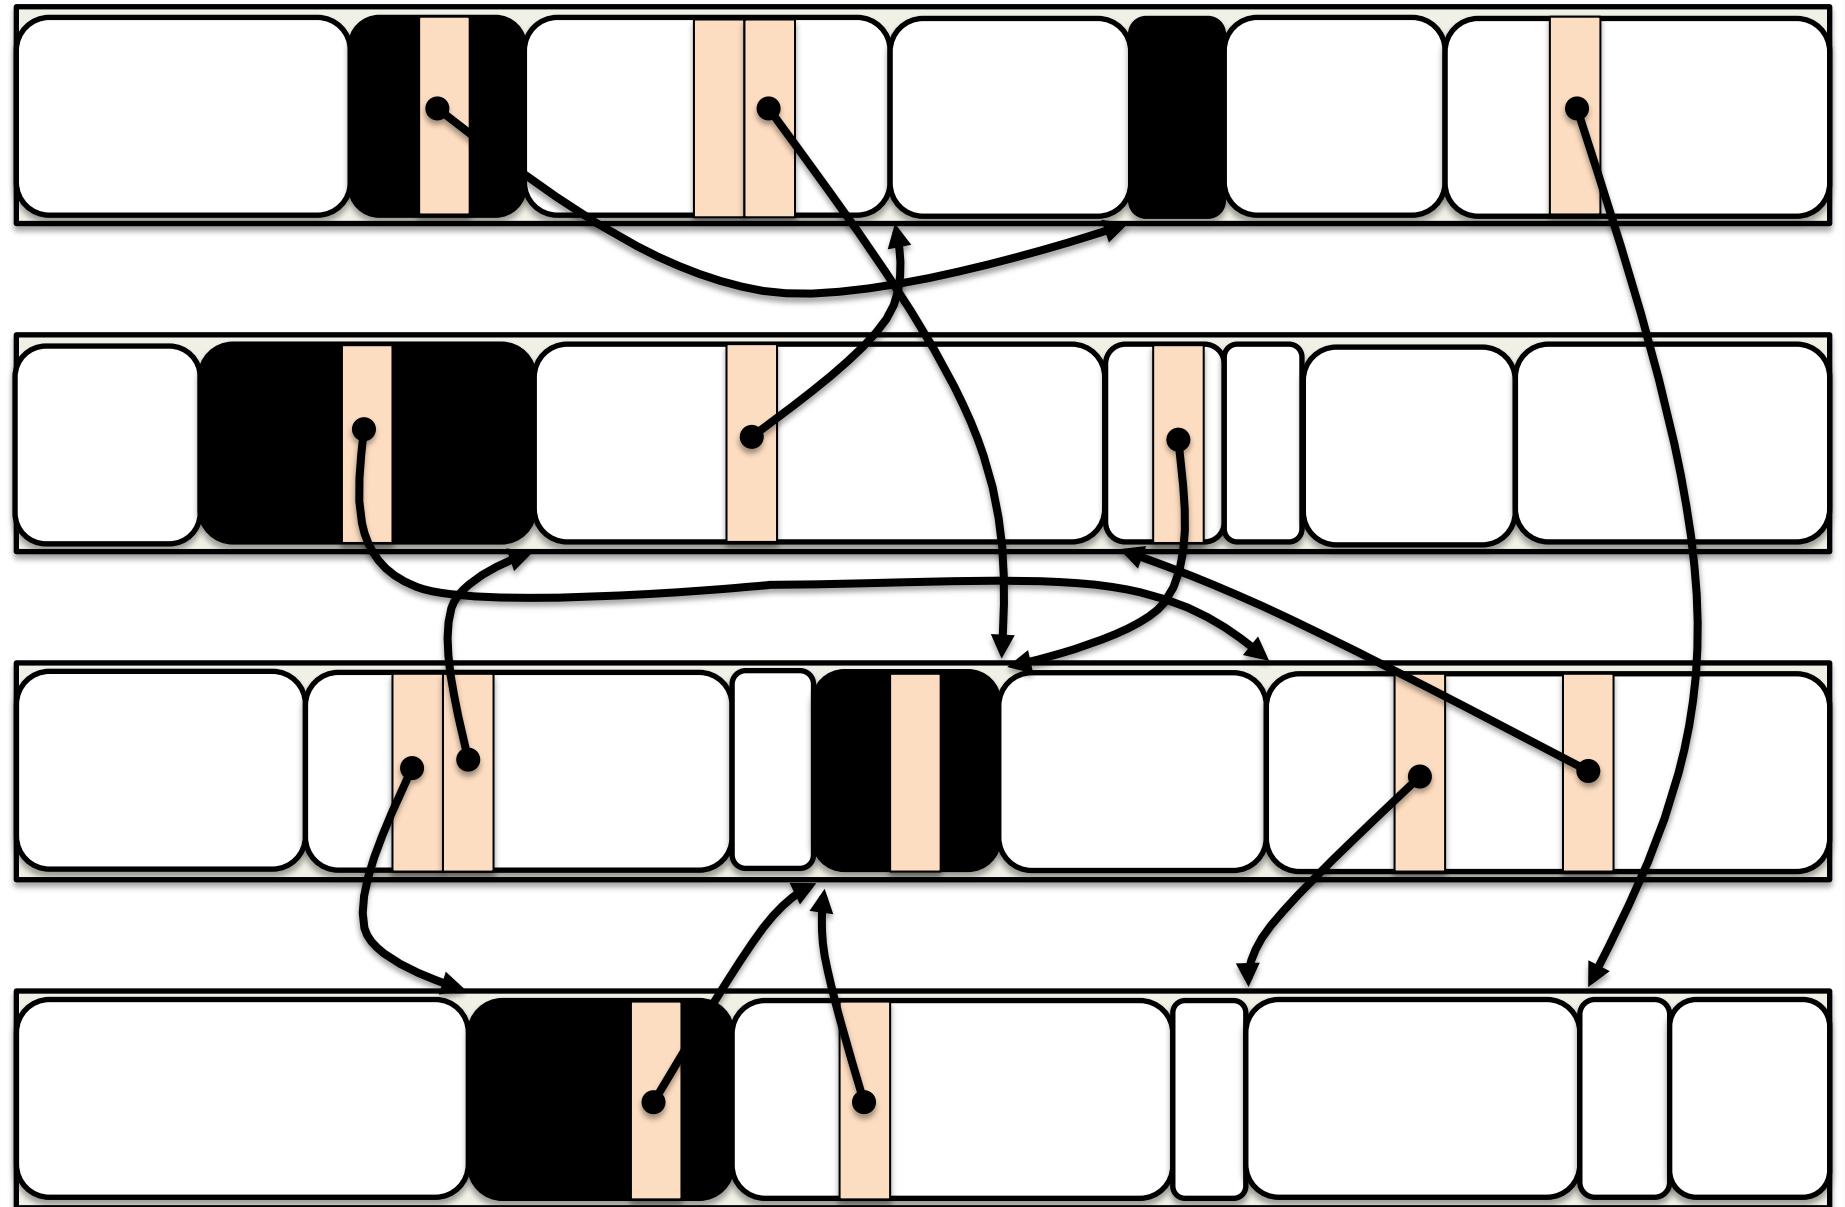


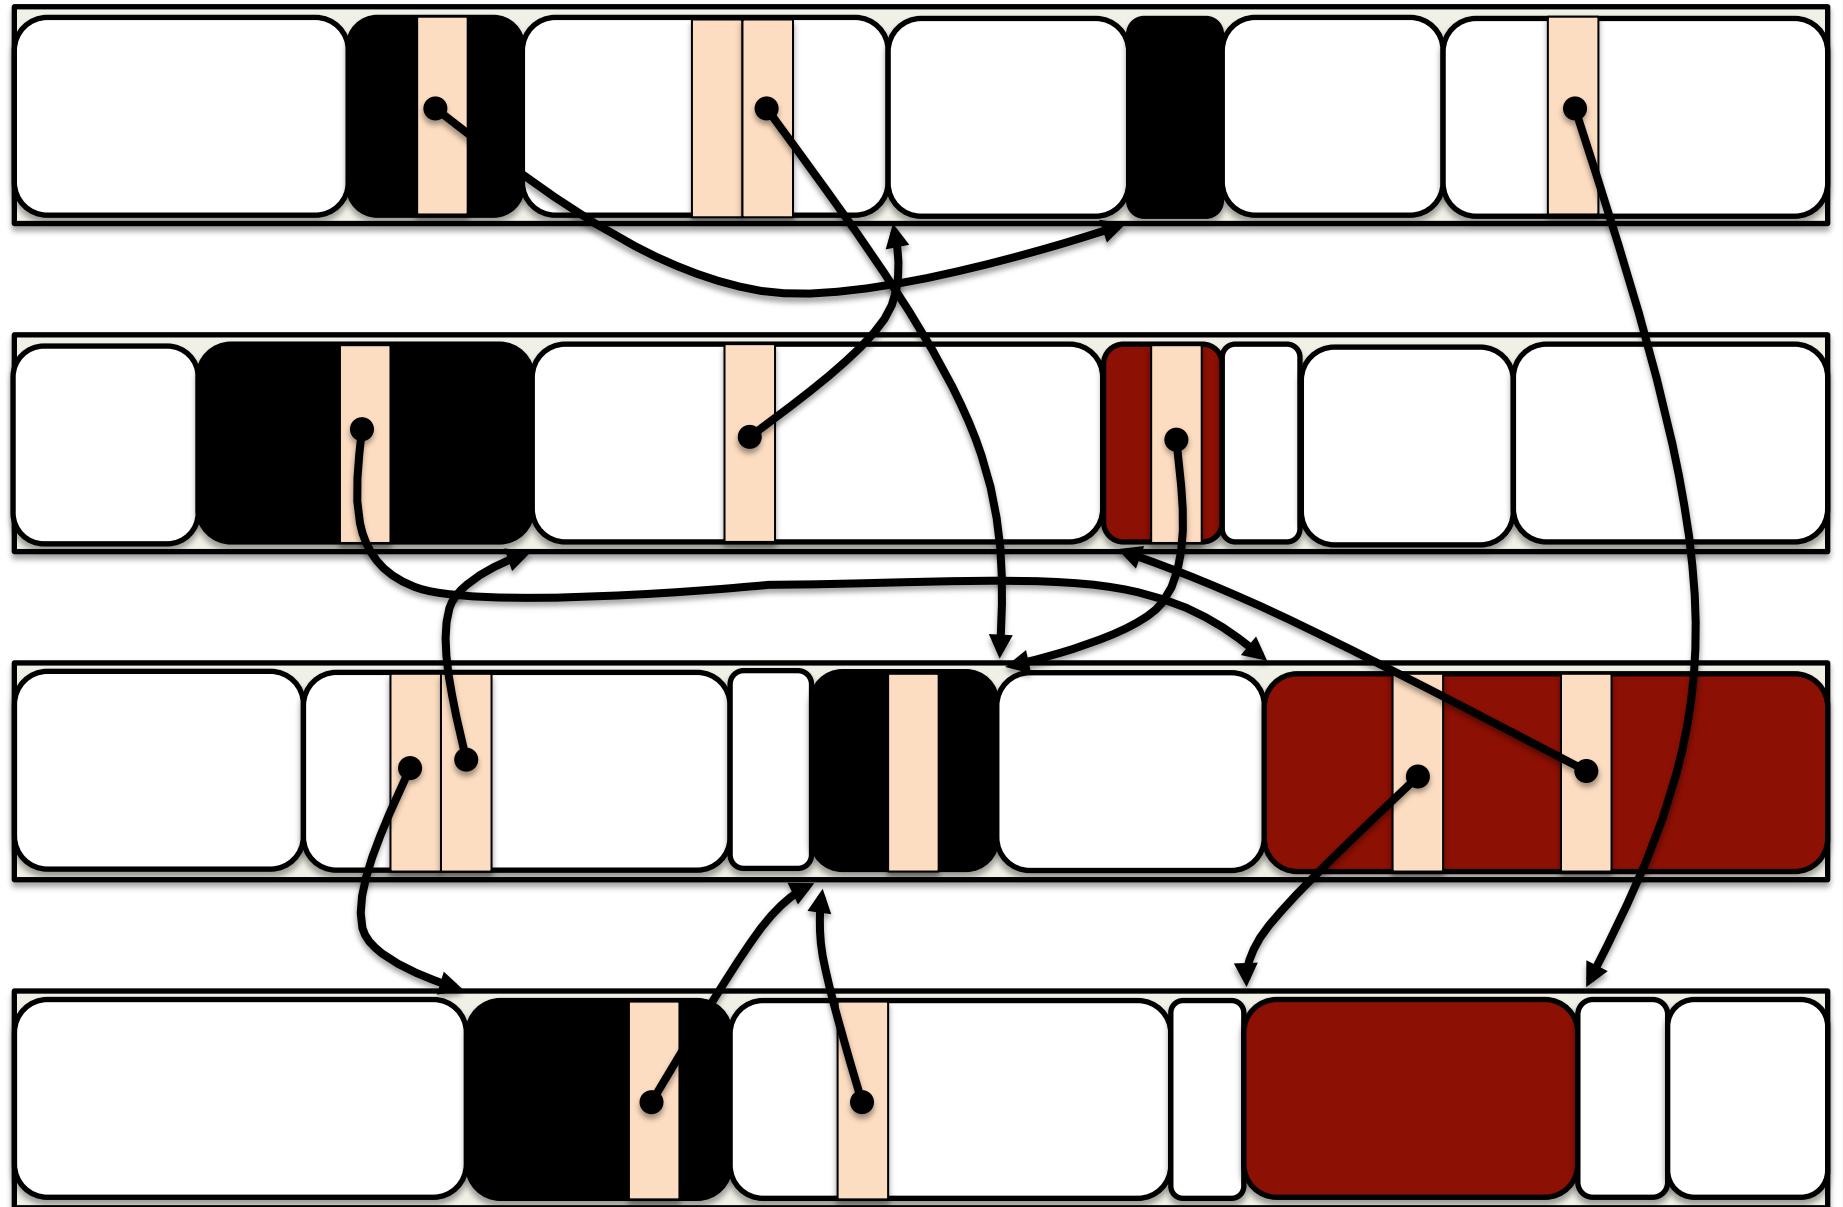






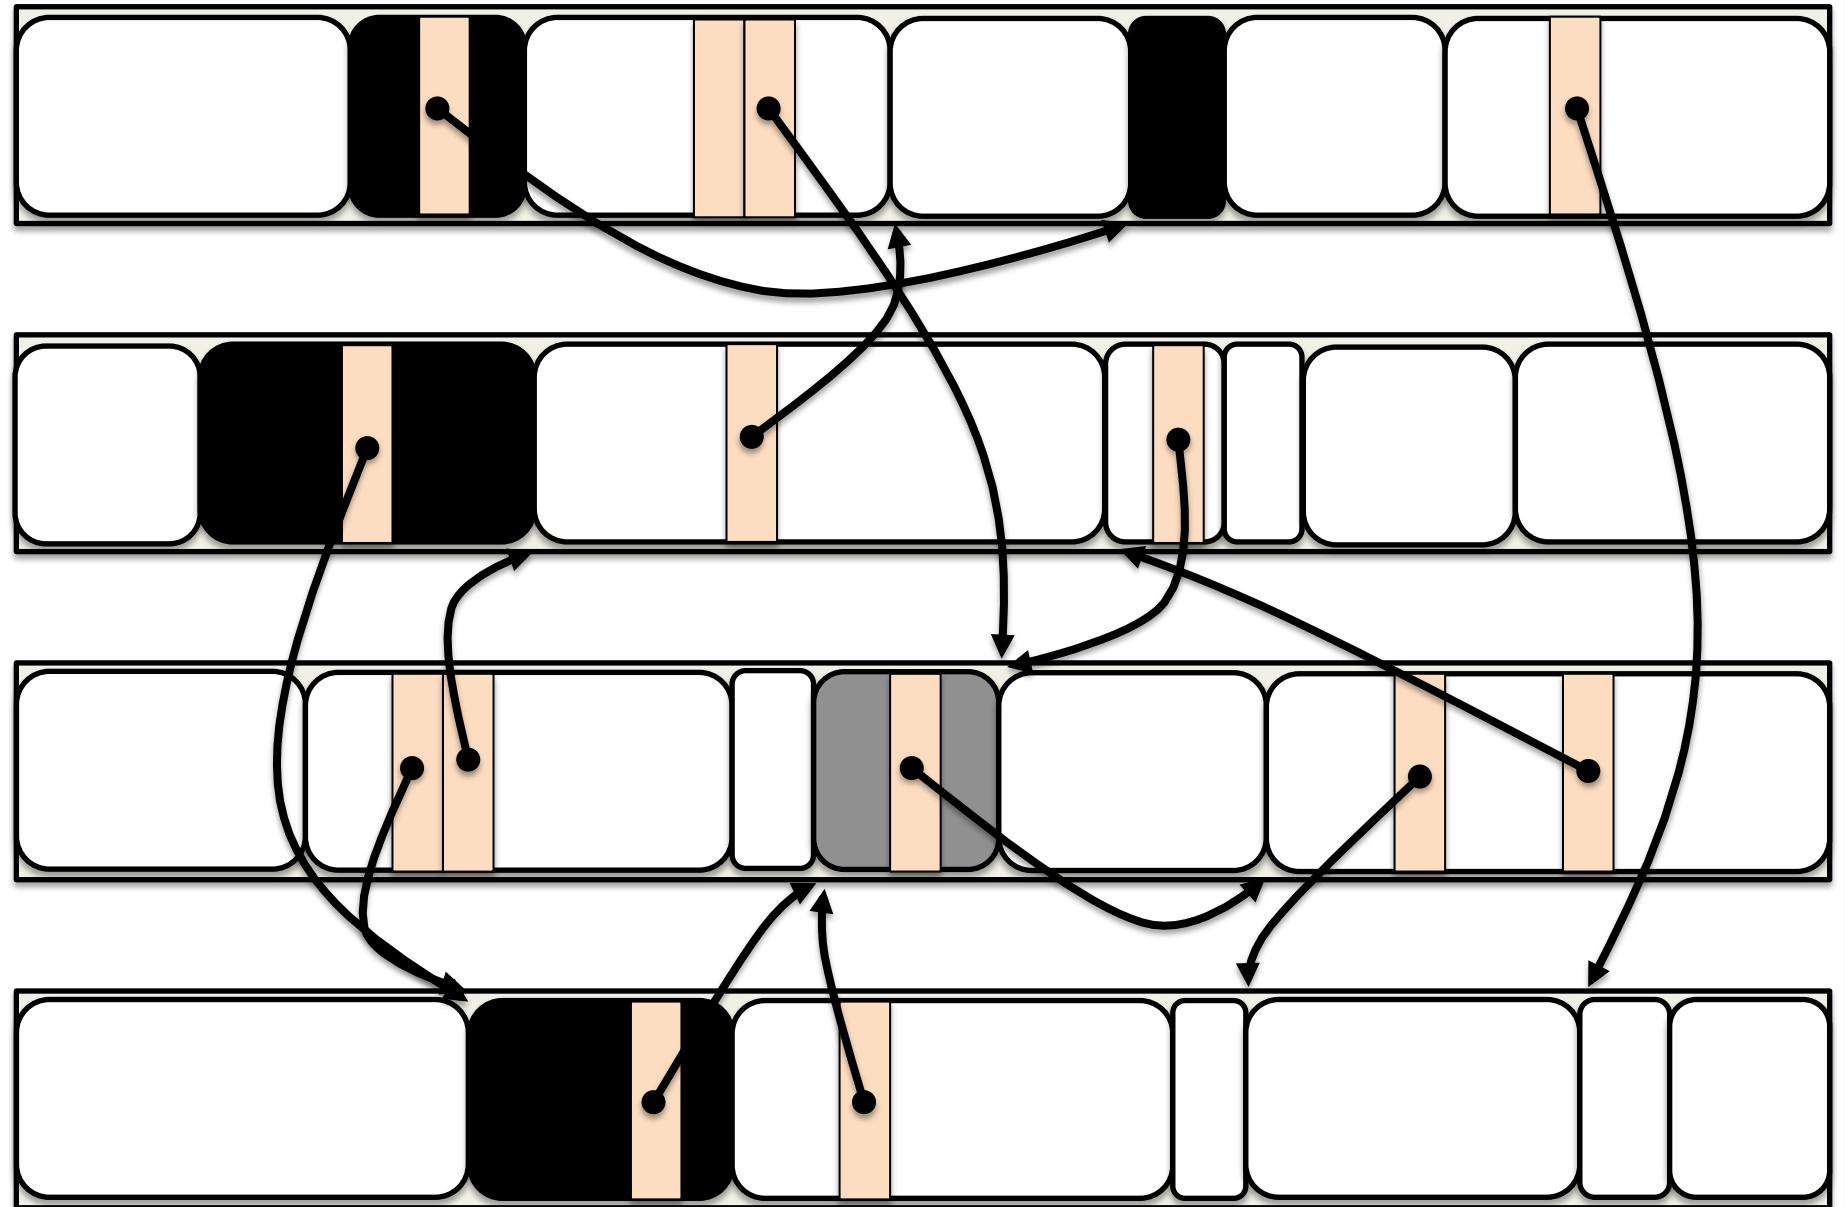


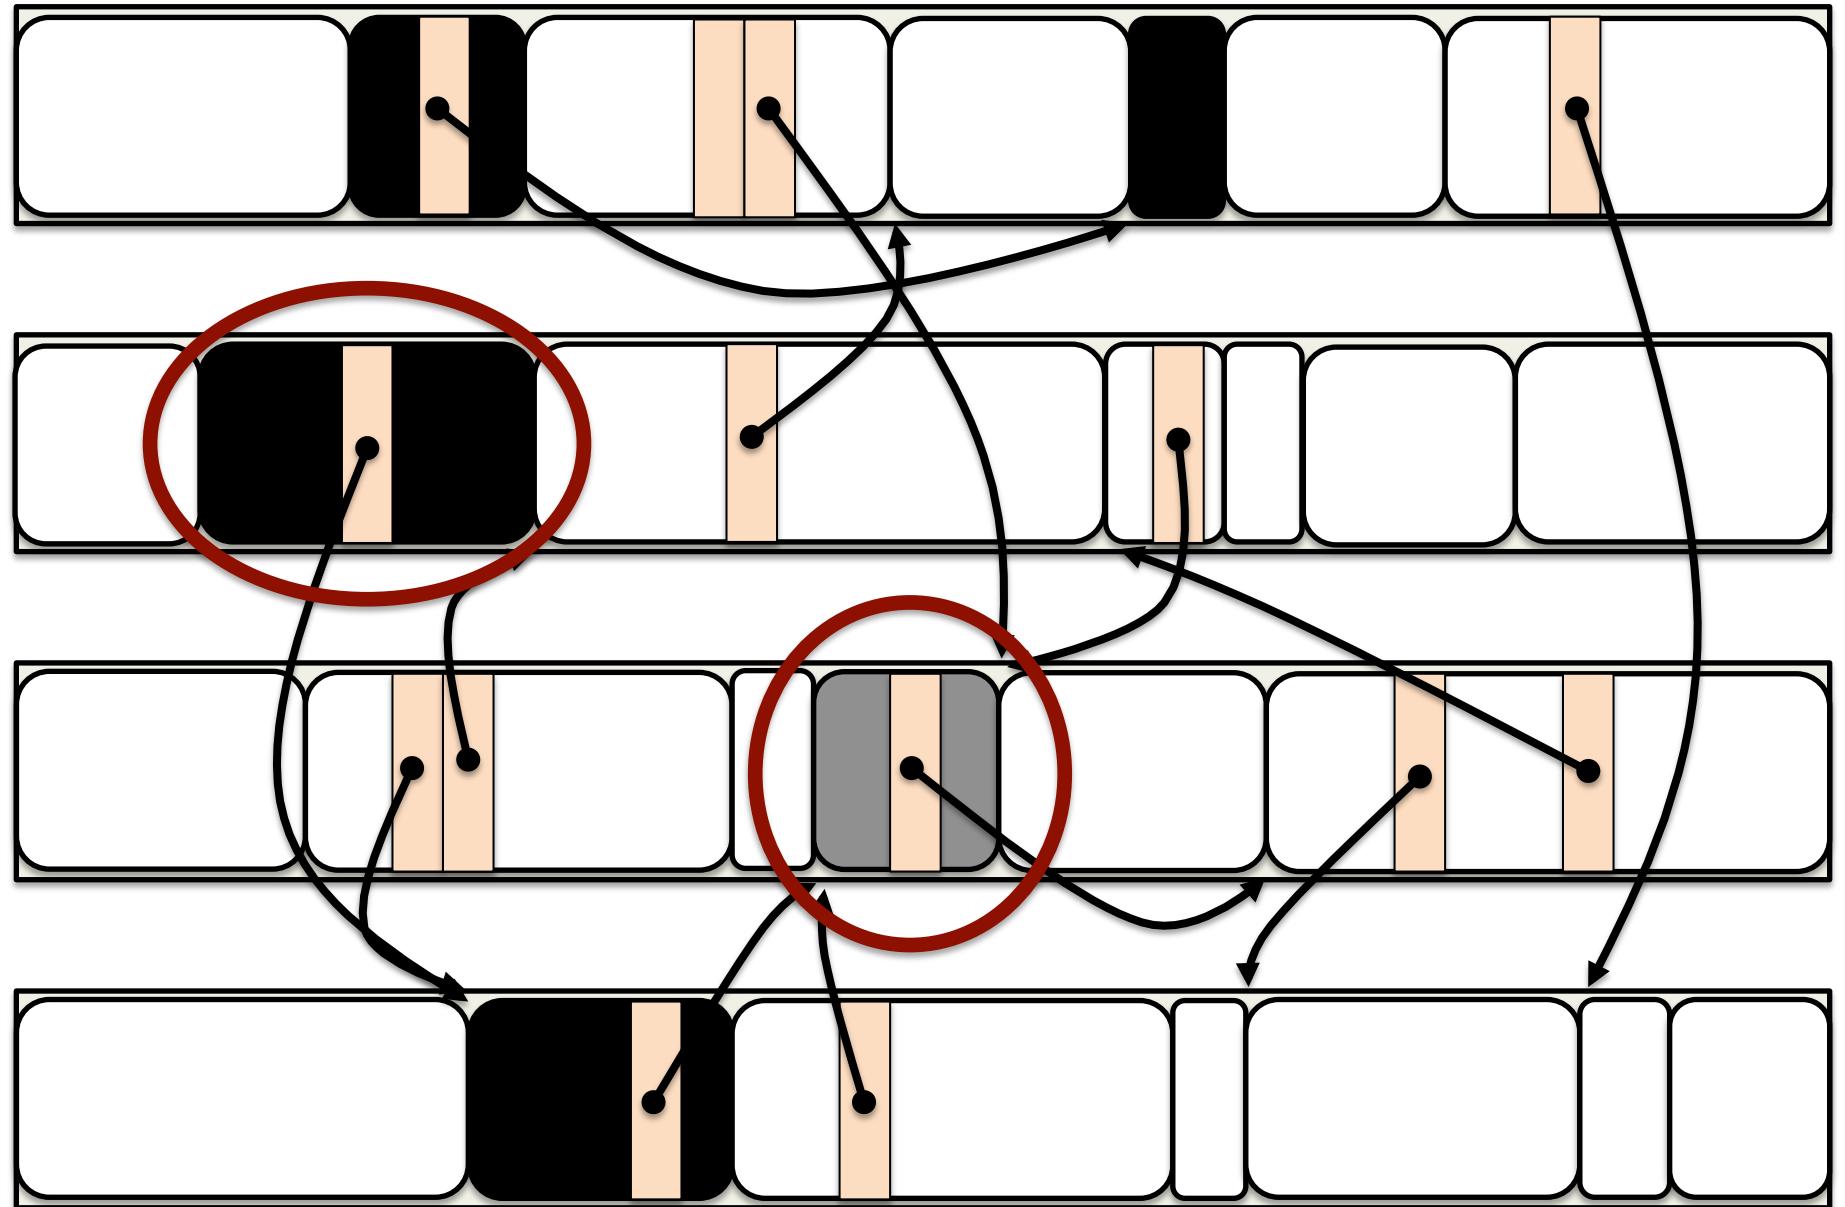


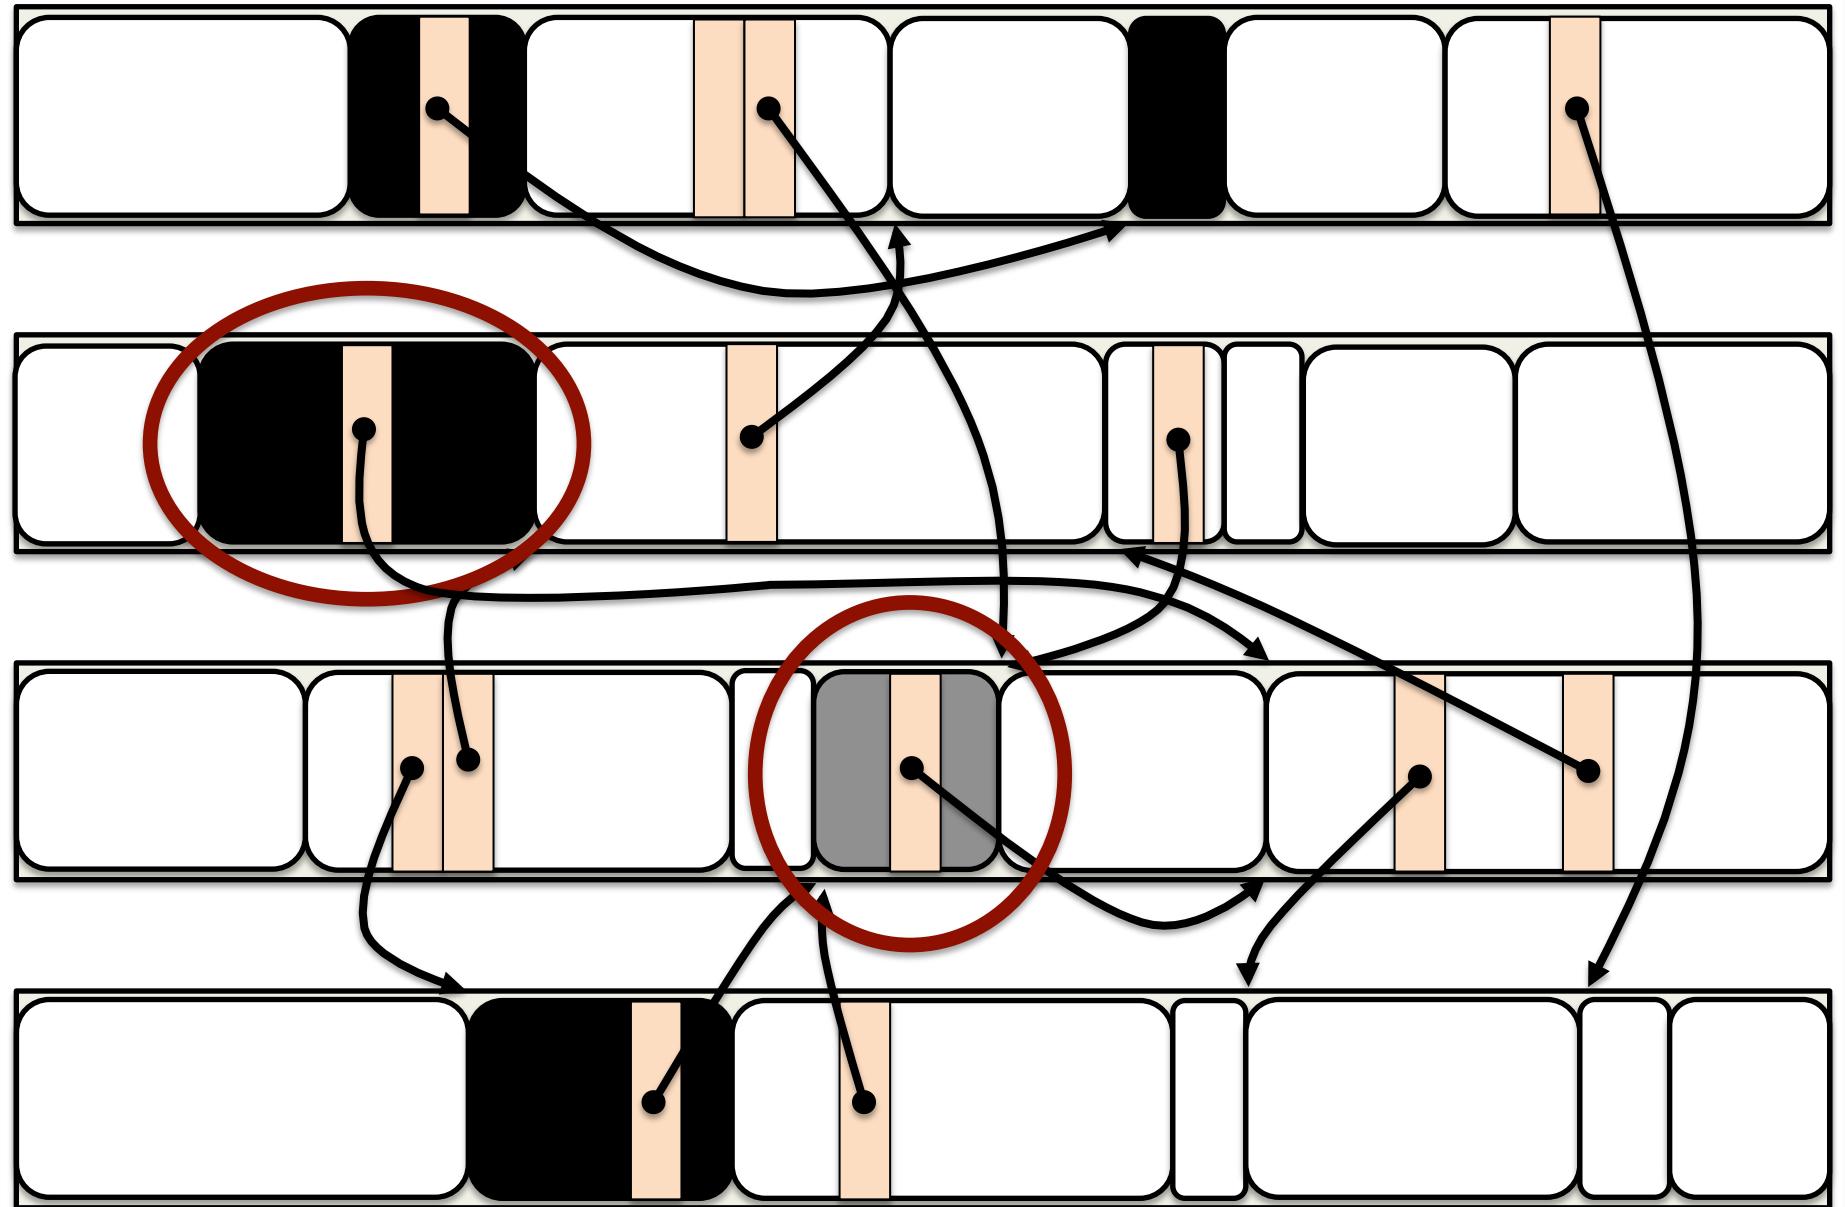


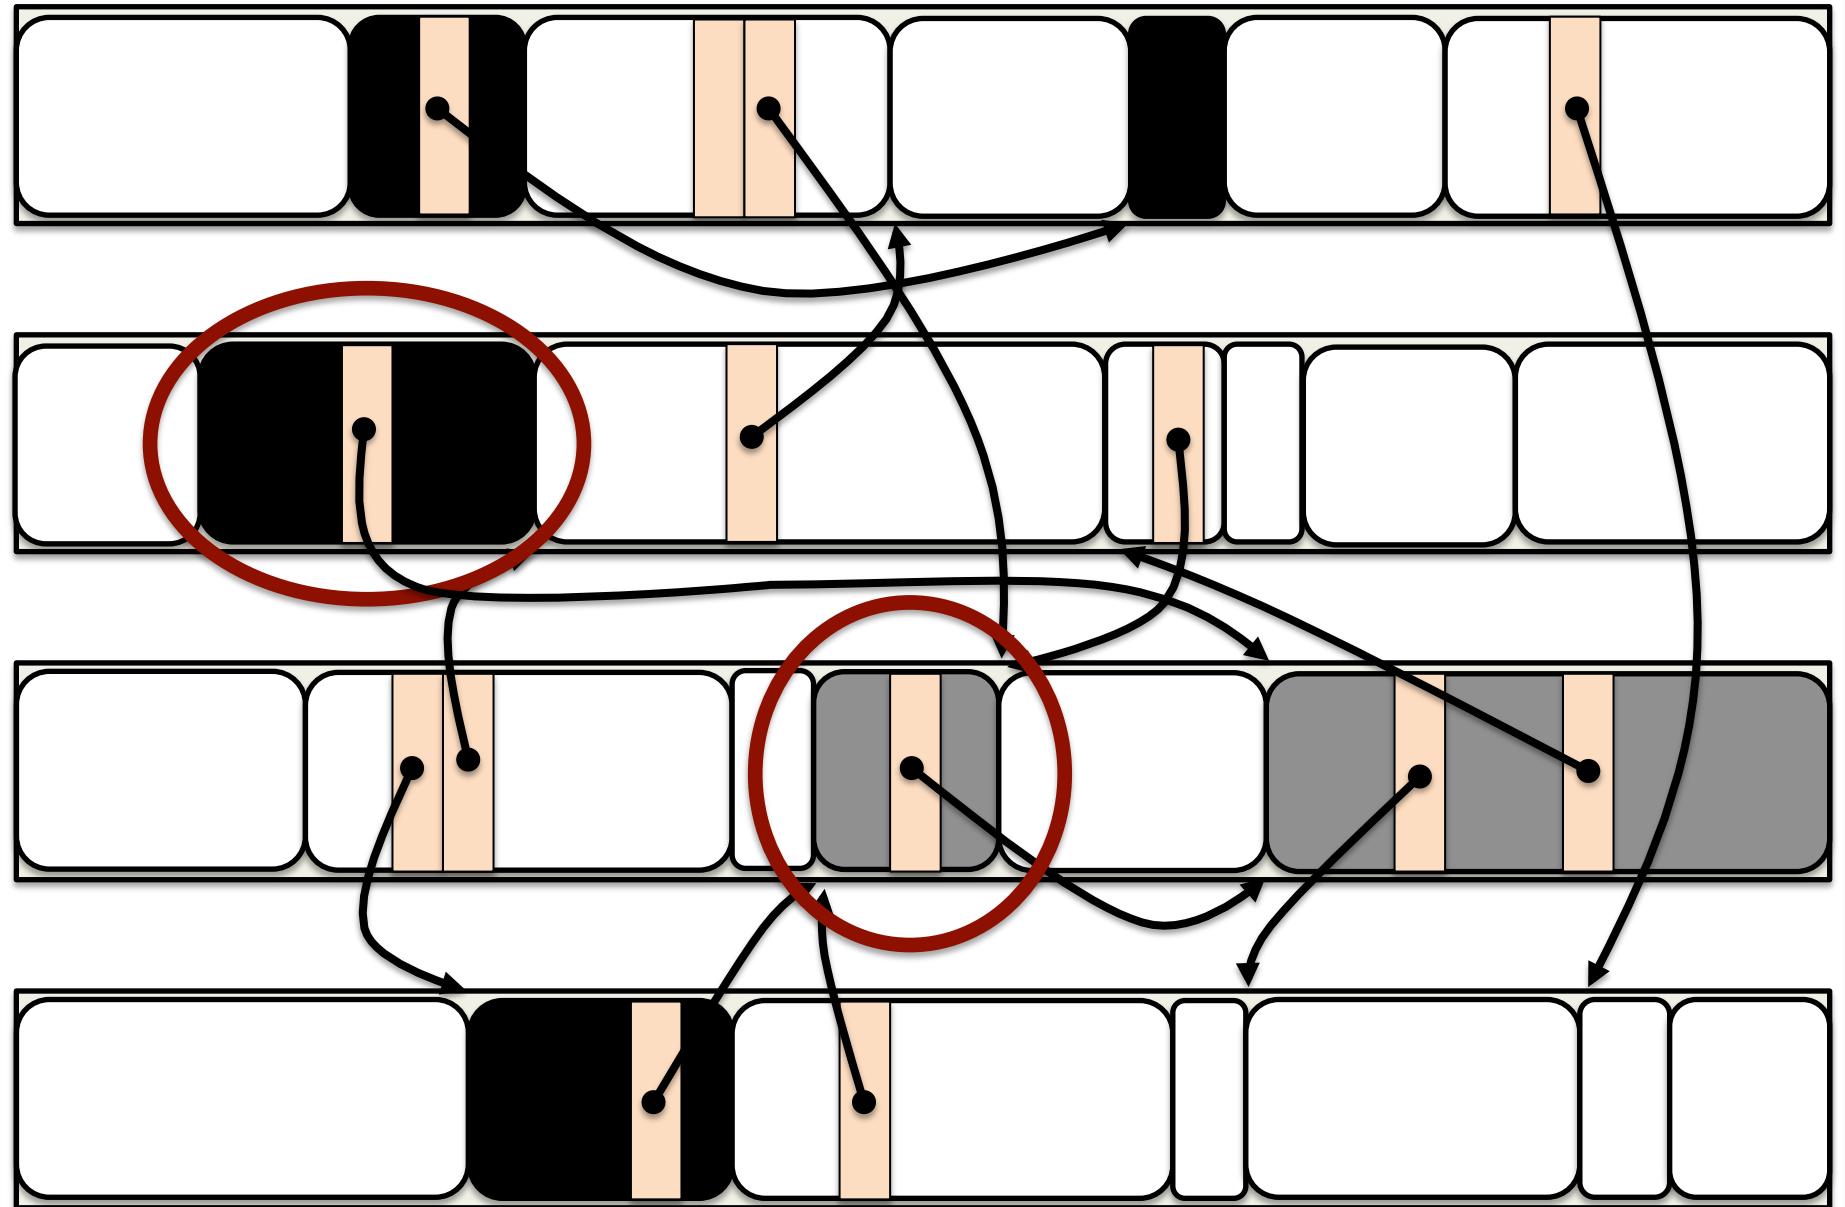
Tracking All Objects

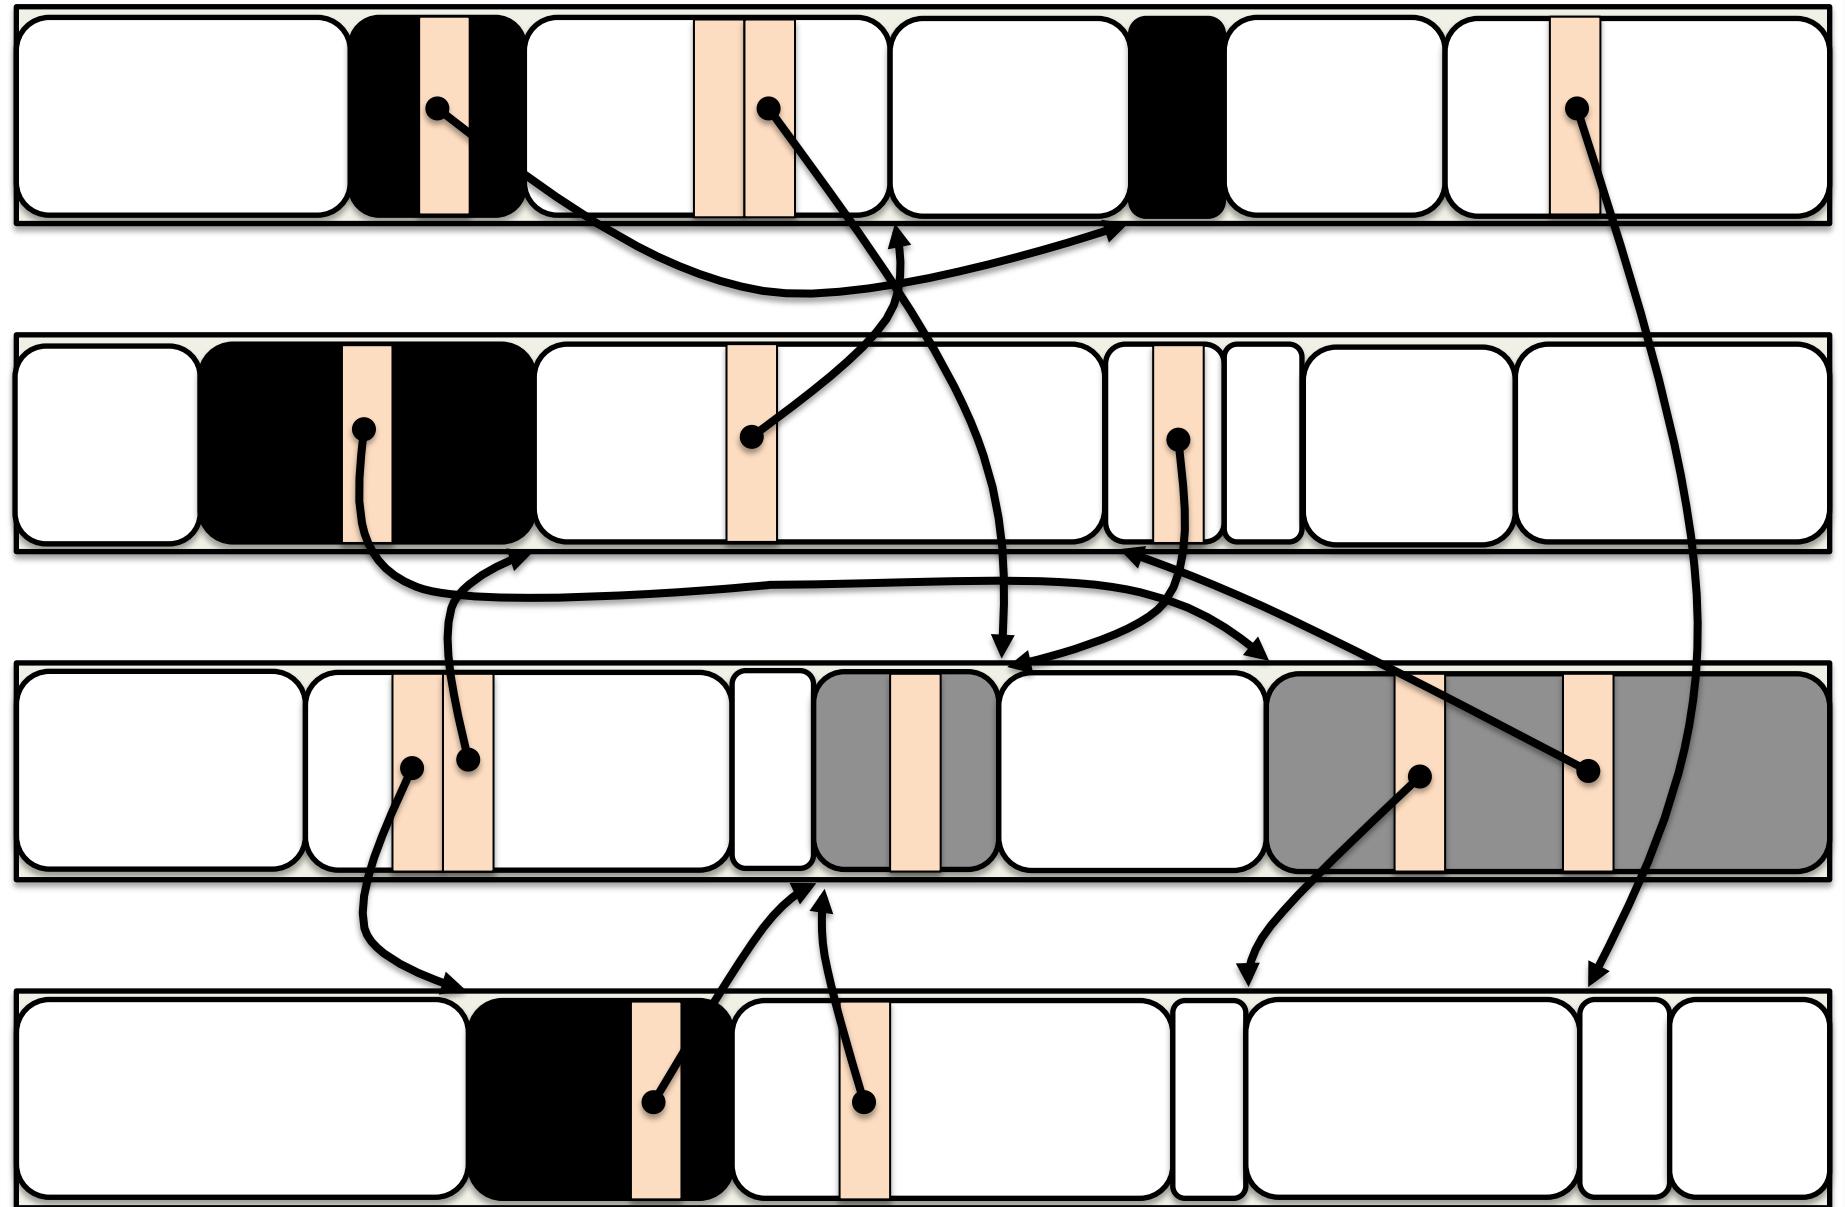
- We can lose an object if:
 - A white object reference is stored in a black object
 - All paths from gray objects to white are removed
- We can consider stacks to be black in this case
- Fix this by intercepting writes to gray objects
 - Check if the target is white
 - If so, set it to be gray
- Implement using a write barrier

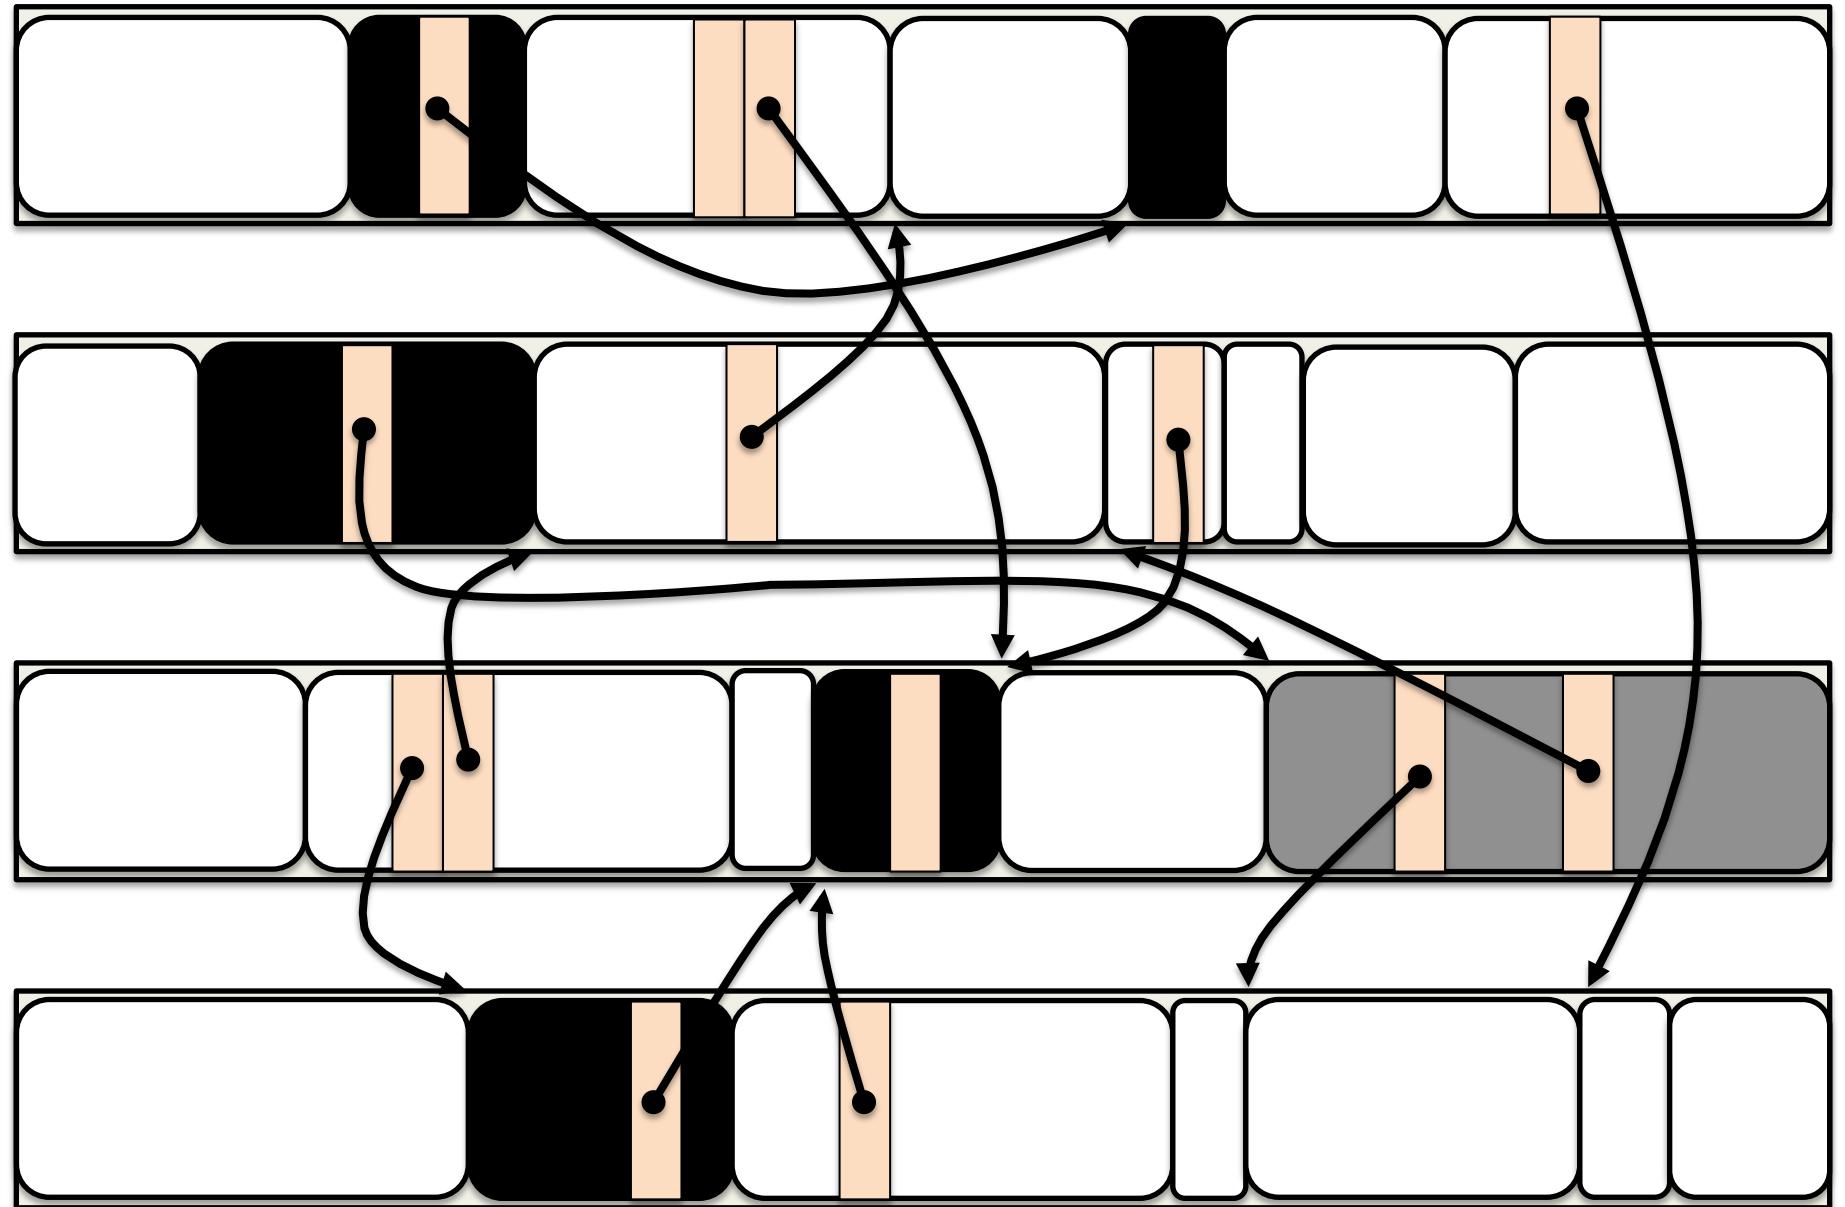


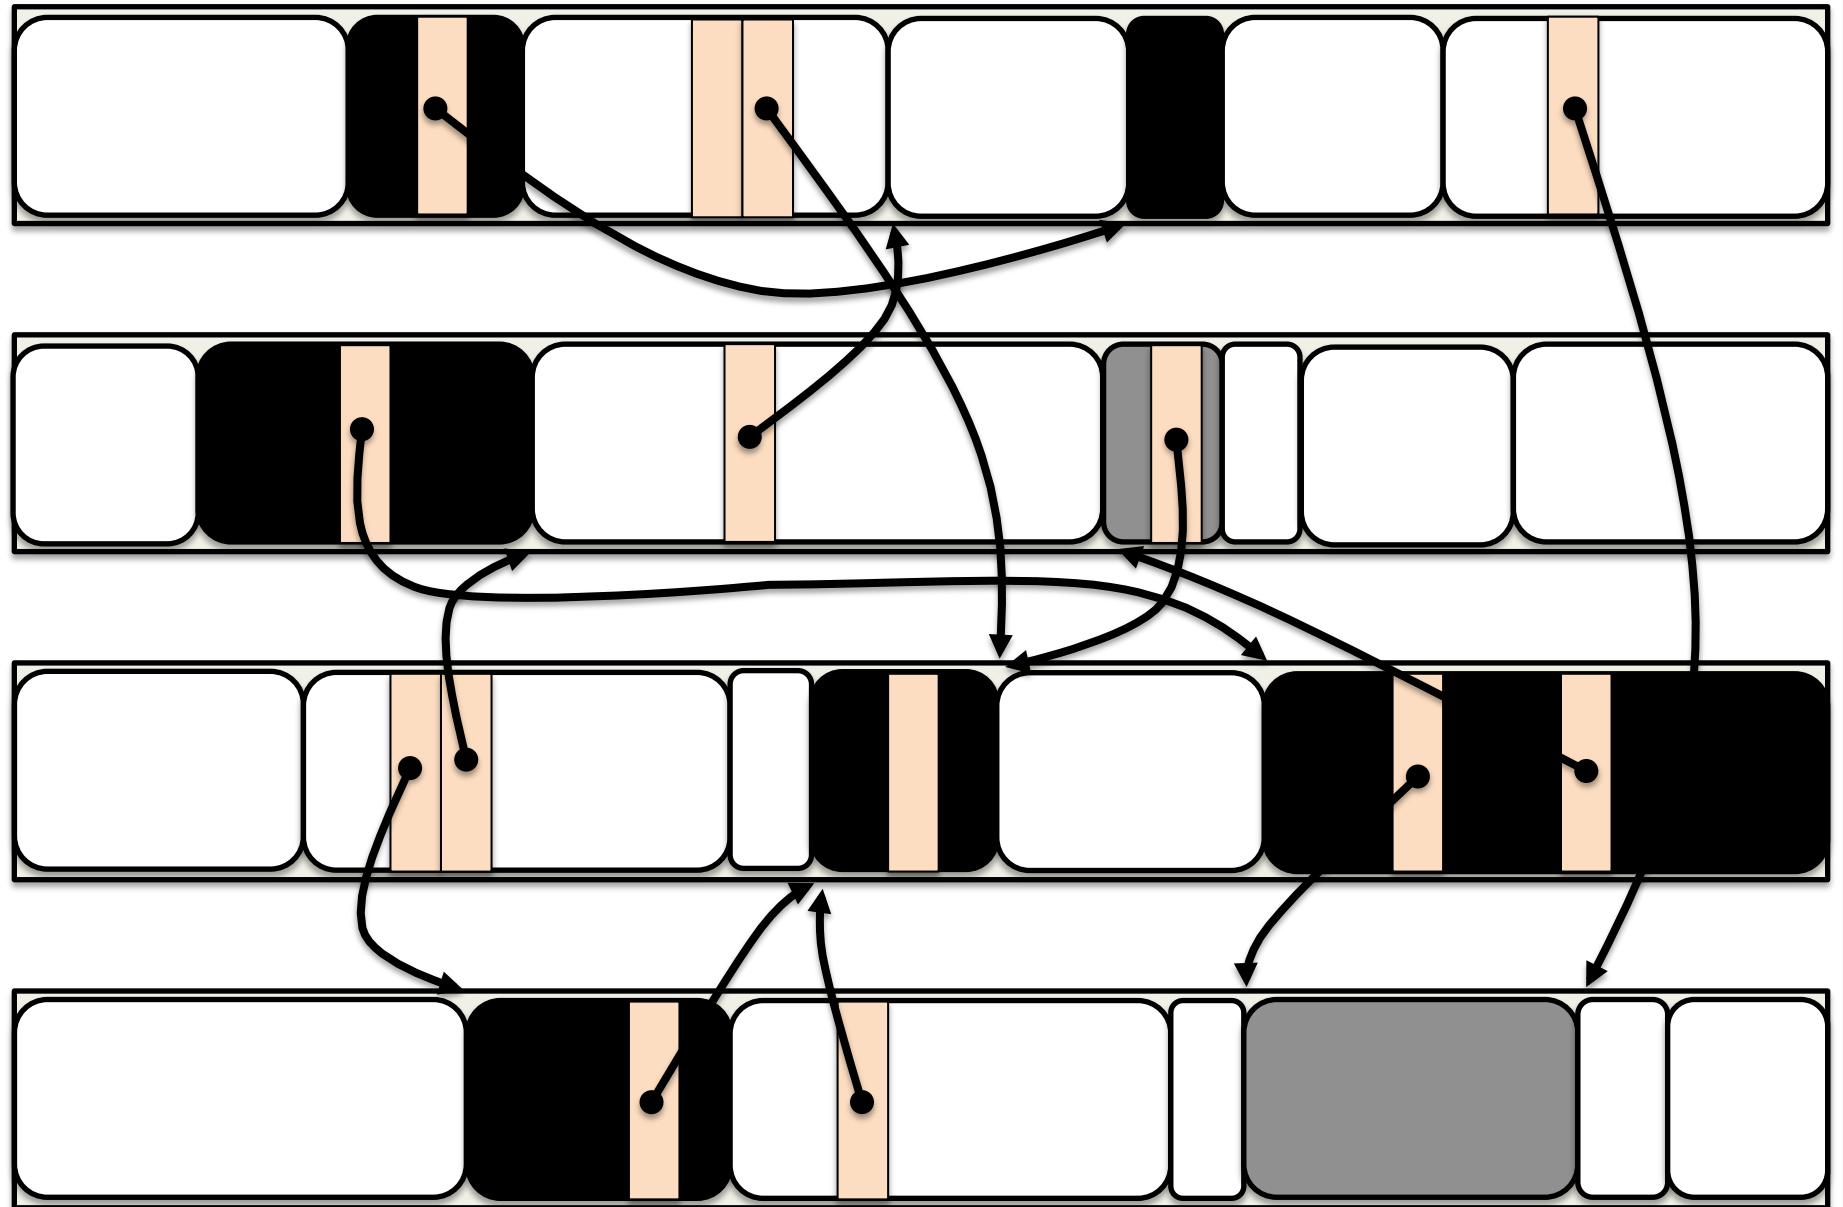


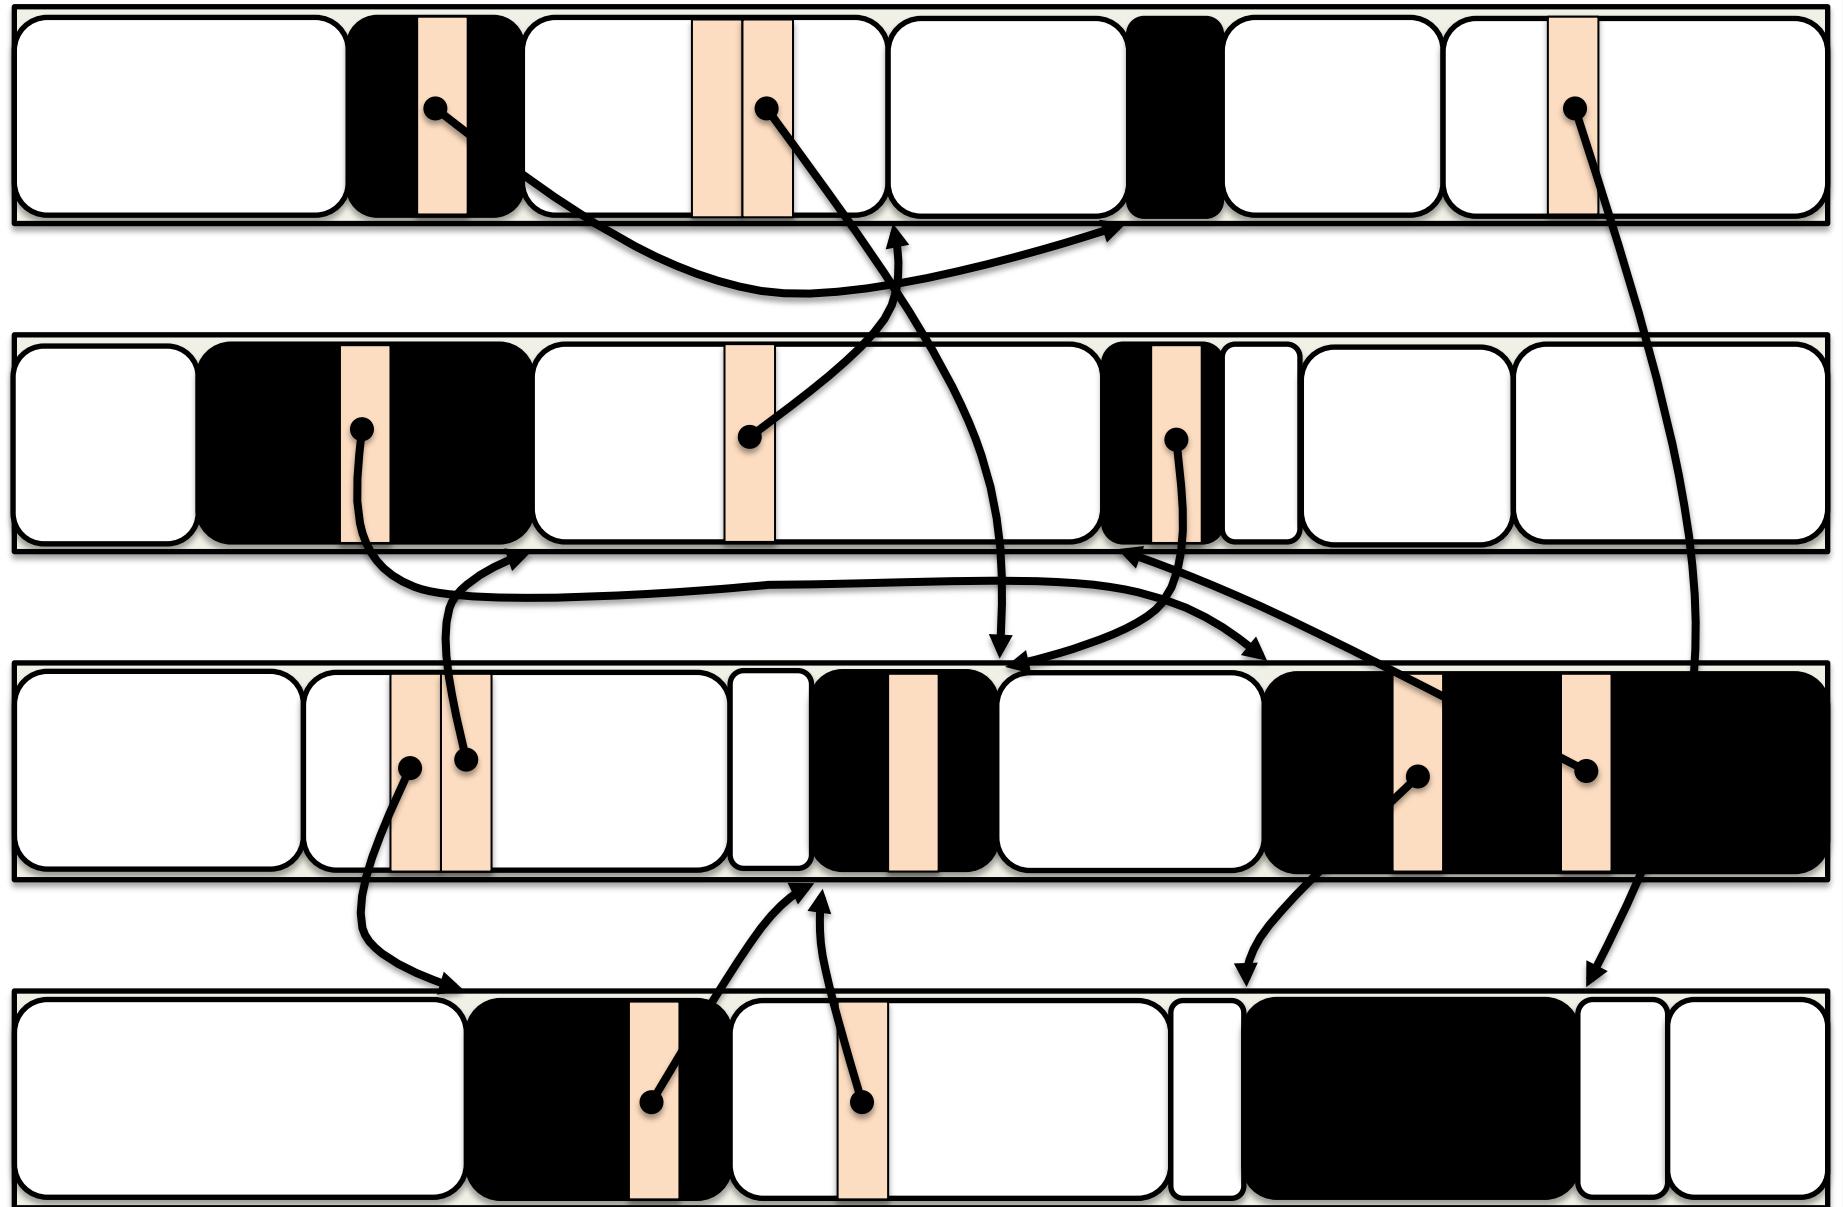












SATB Write Barrier

- If the source object is gray
 - And the target object is white
 - Add the target object to the gray stack

Allocating Objects

- We have three choices for newly-allocated objects
 - Allocate white
 - Allocate gray
 - Allocate black
- New objects contain no references
 - Recall the default values for fields
 - They are already scanned
- For this algorithm, we allocate objects as black

Snapshot At The Beginning

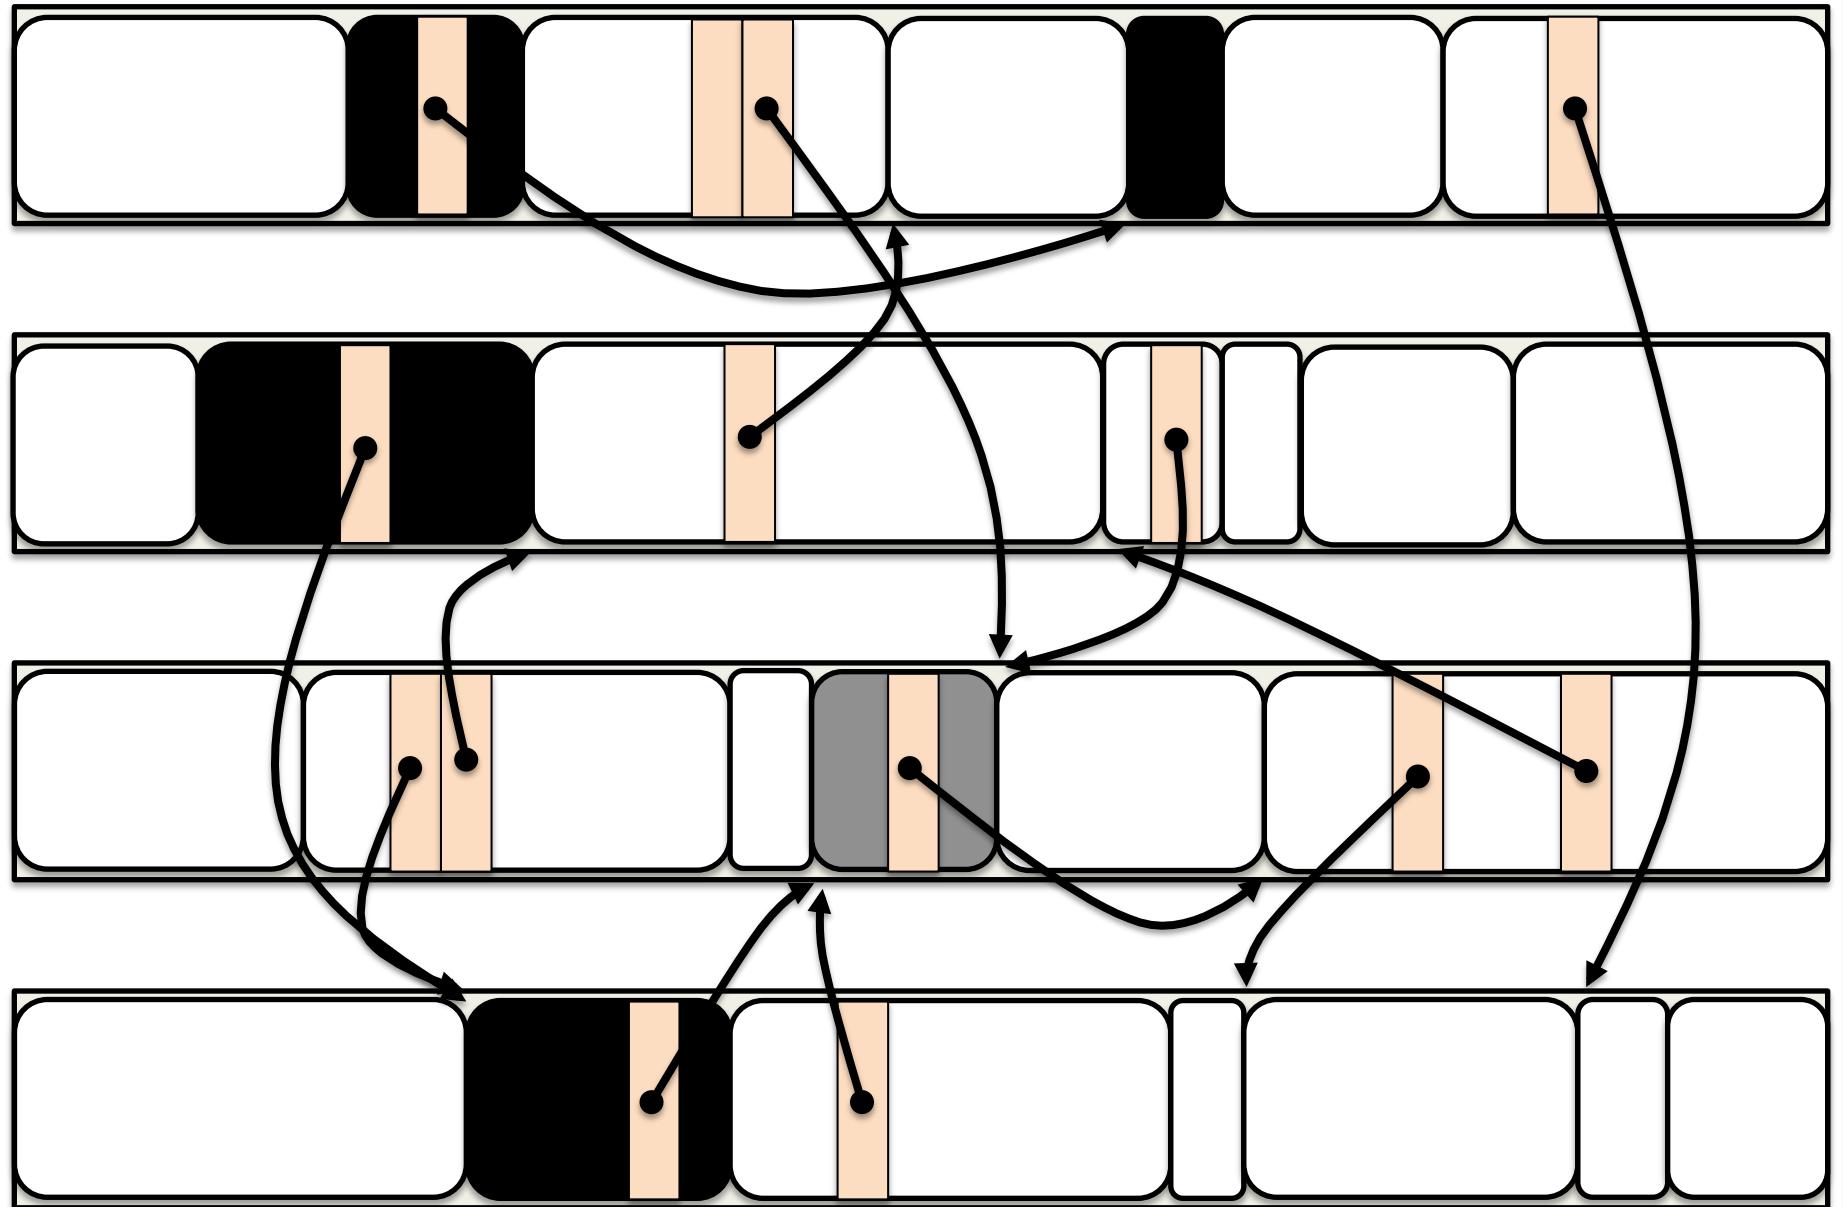
- Concurrent algorithm
 - Very short initial pause to scan the root set
 - All other work happens concurrently
- Limited amount of scanning
 - Once an object is black, it never changes
- Conservative
 - Retains all newly-allocated objects
 - Reduces benefit of generational collection

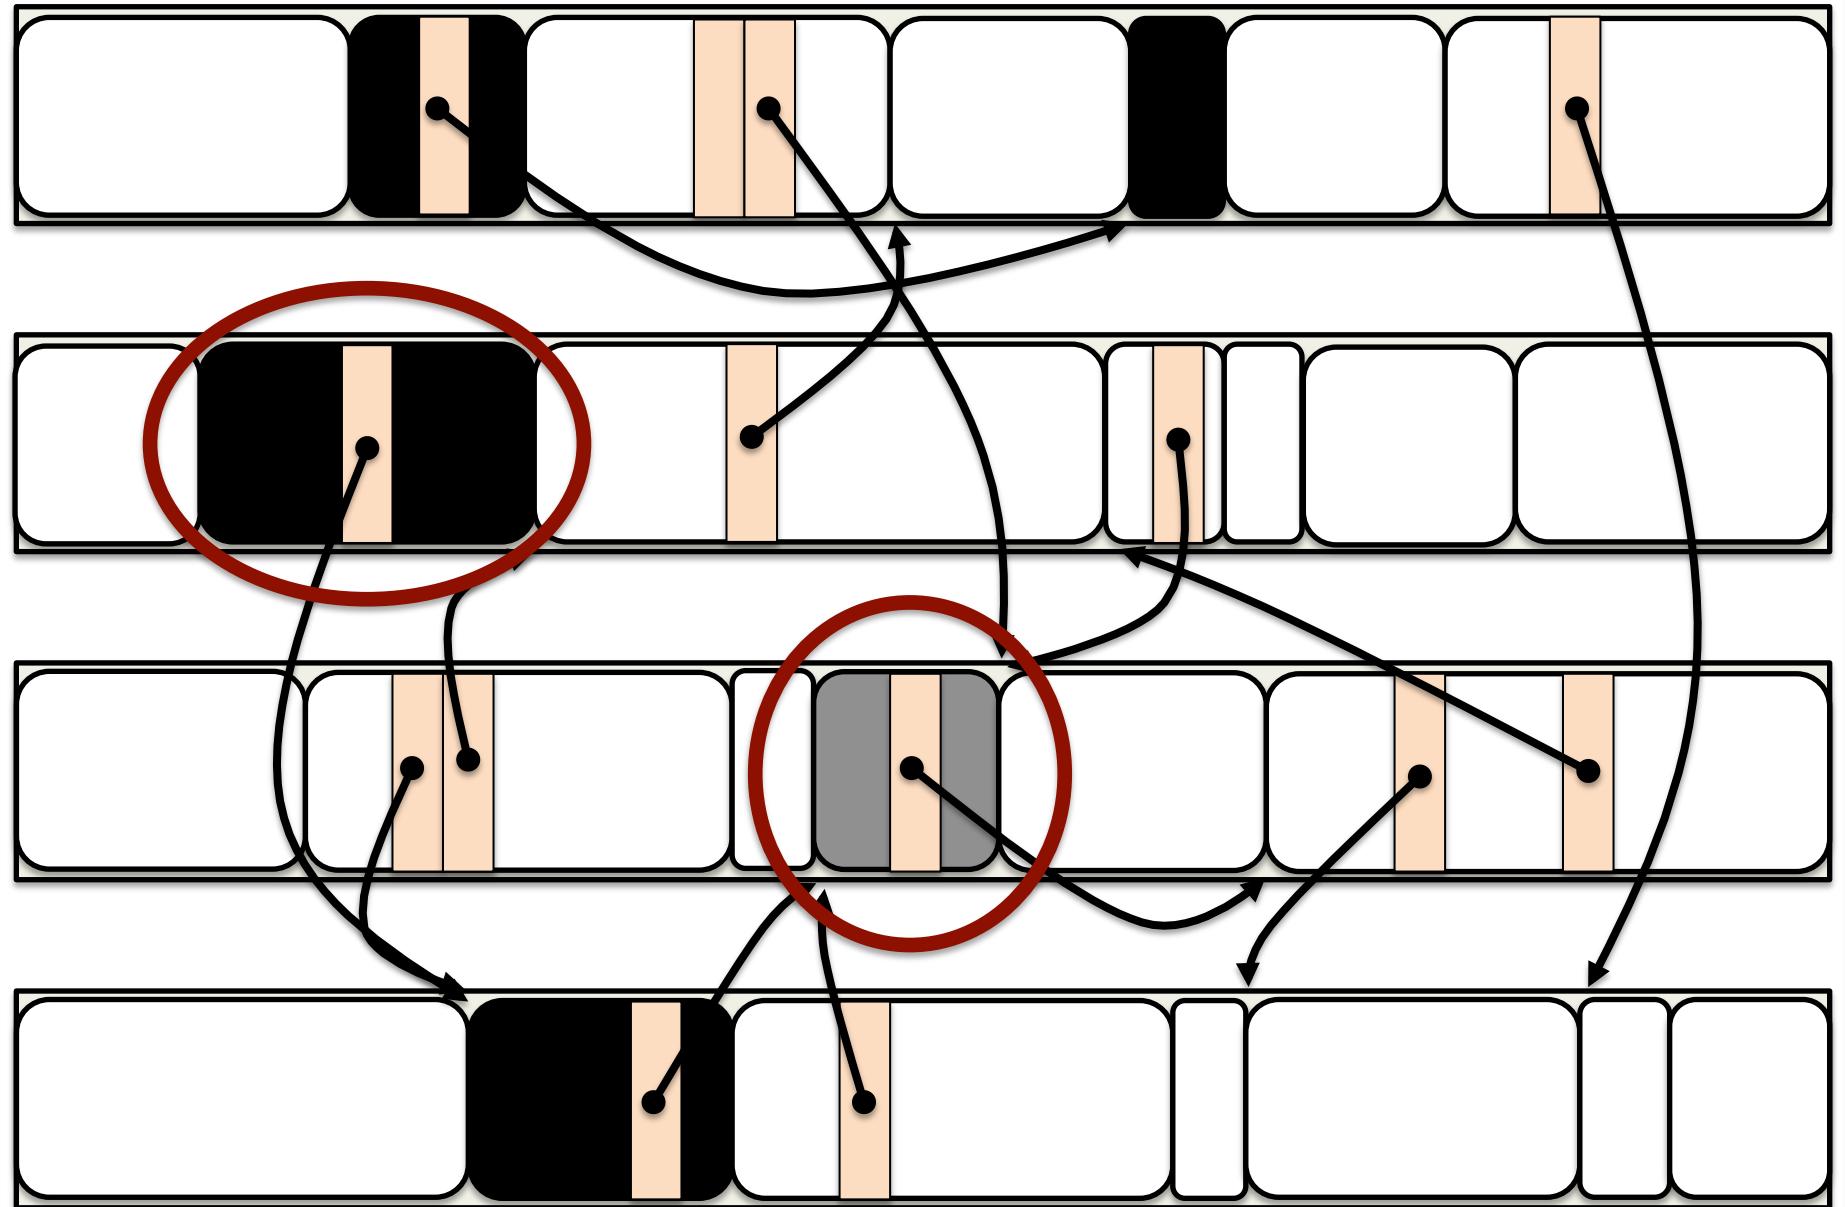
Incremental Update

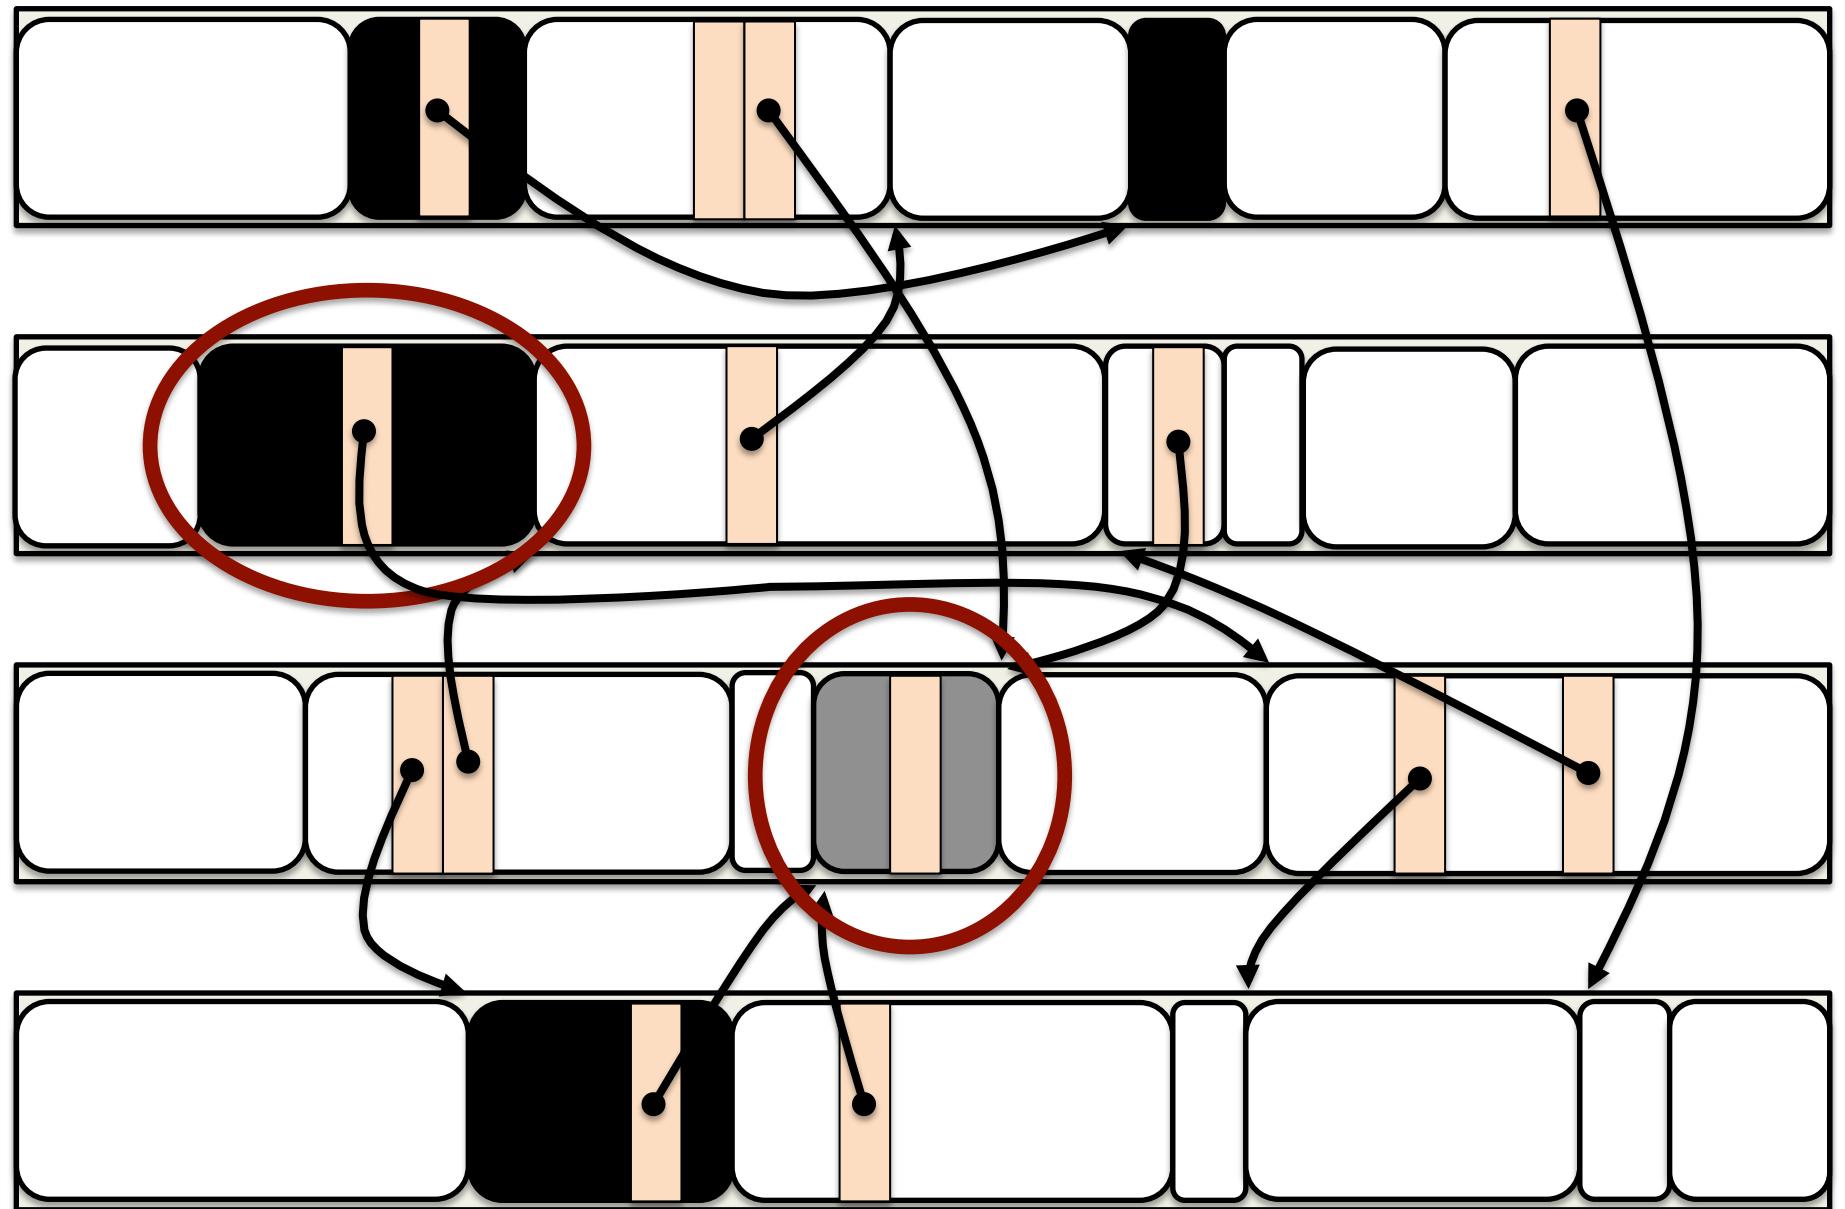
- Reduces the conservatism of SATB
- Takes the opposite approach
 - Tries to mark objects alive at the end of the phase
 - New objects that are unreachable should not be marked
- Heuristic approach – may mark some garbage
 - Generally far less than SATB

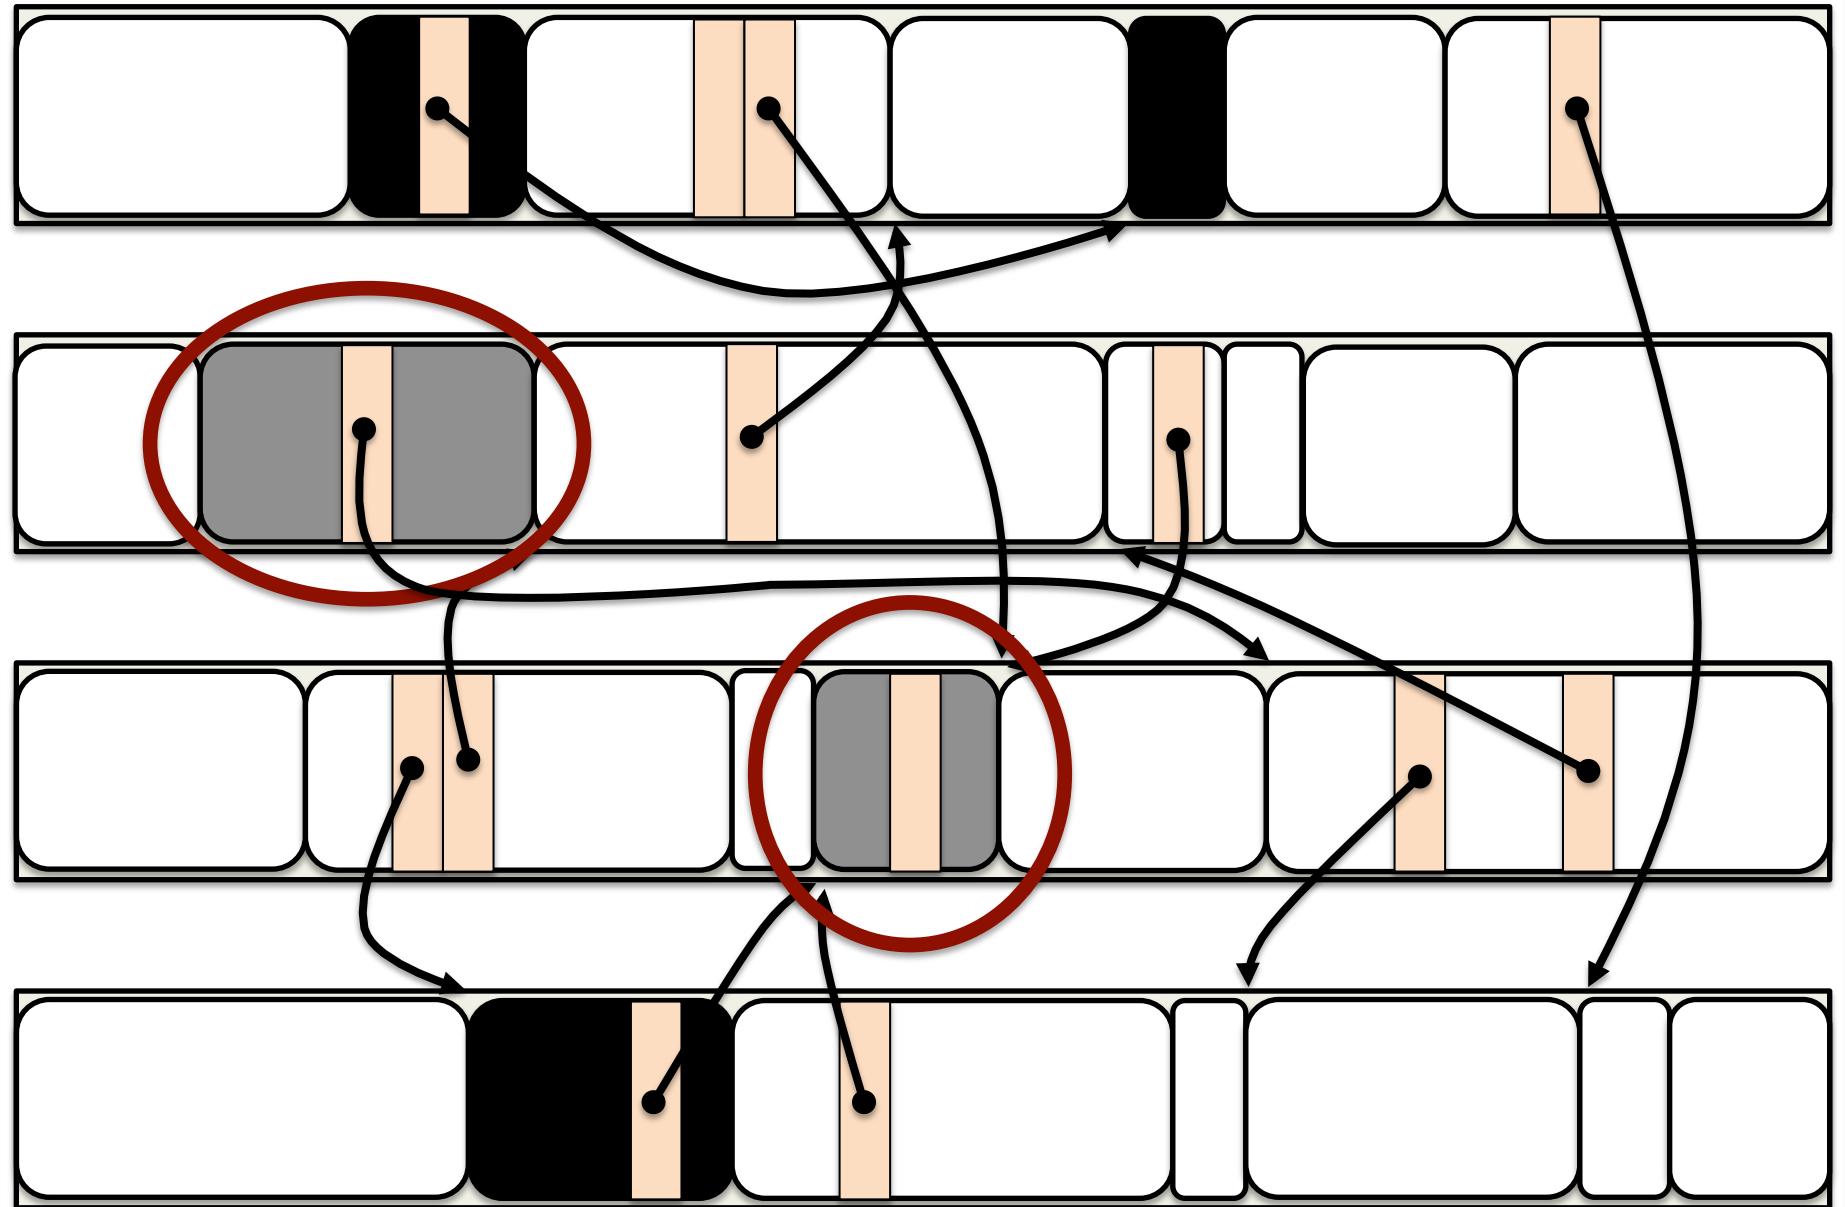
Incremental Update

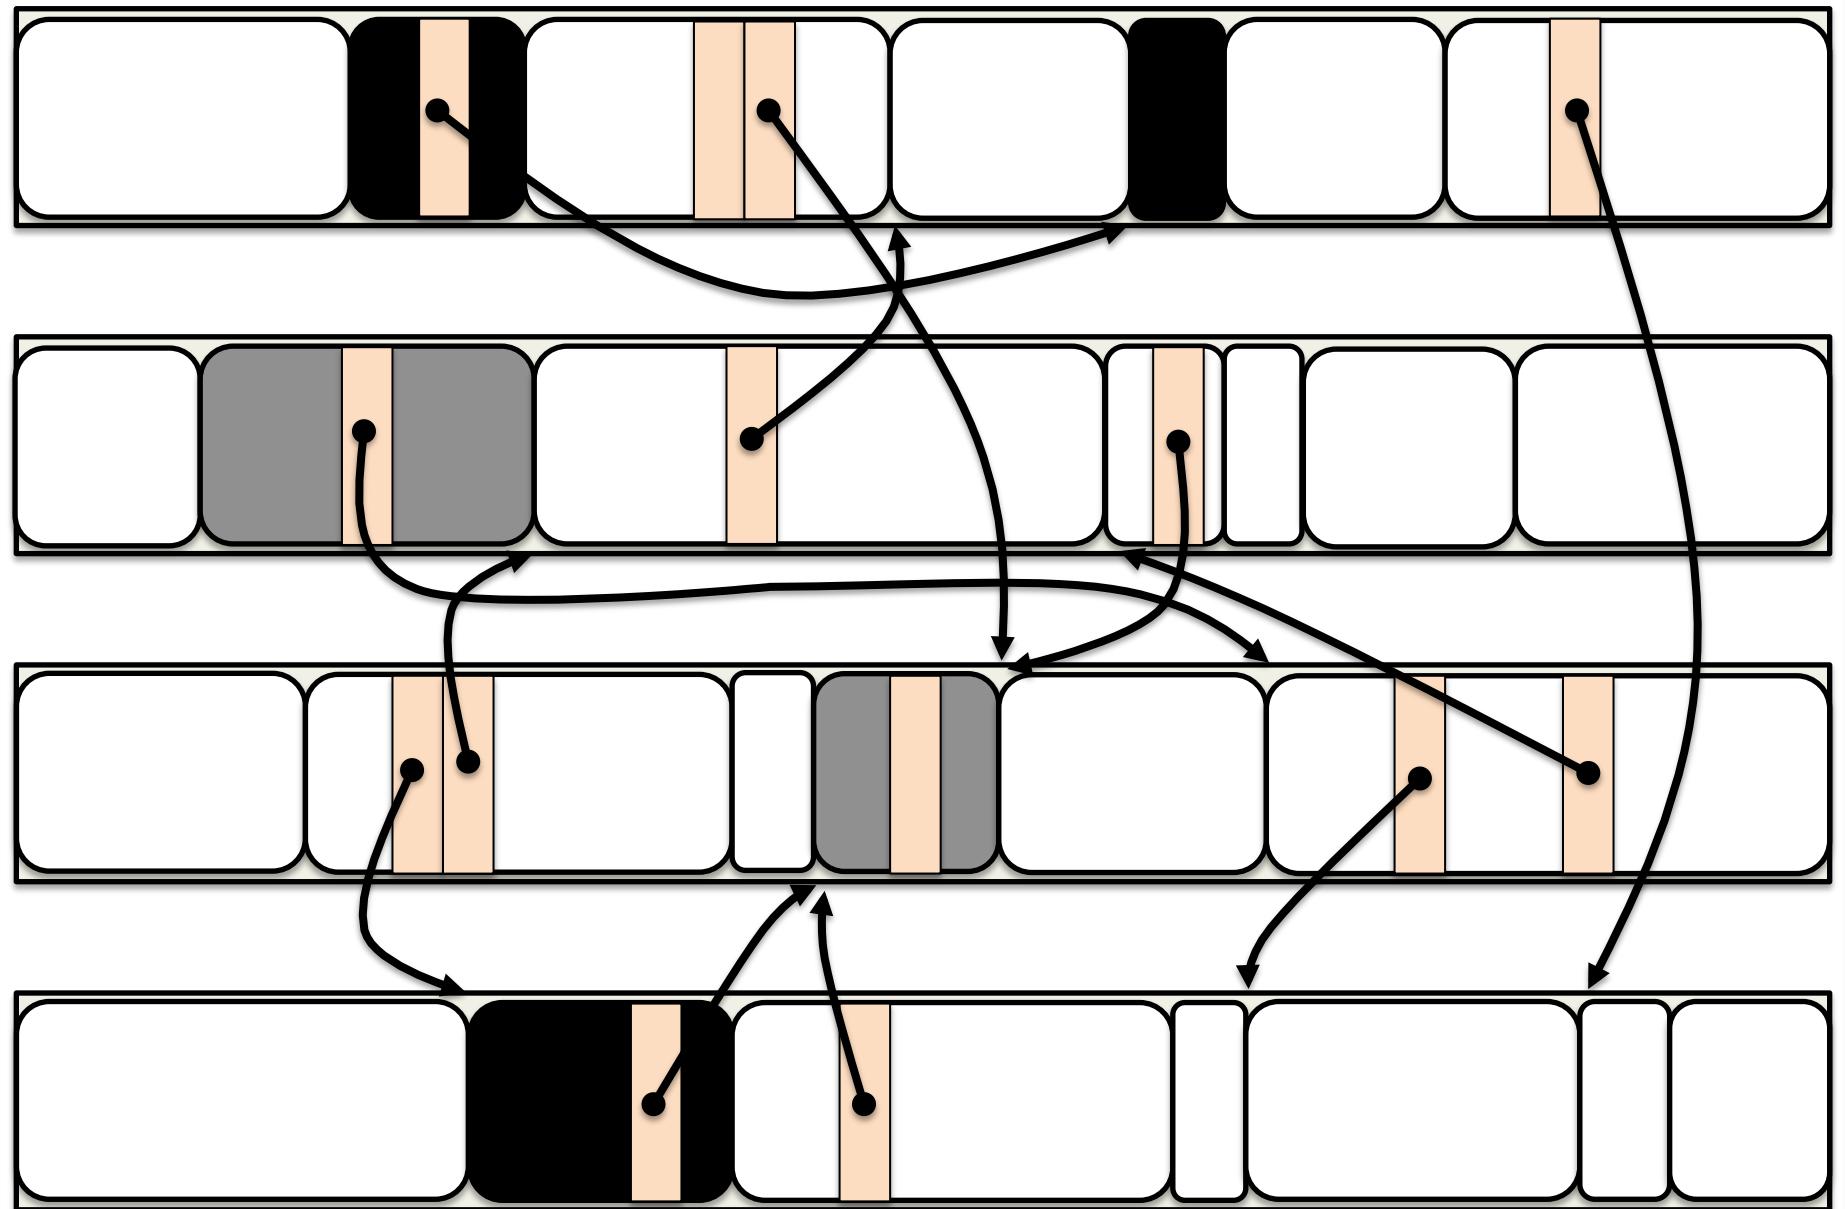
- Maintains the same invariant
 - Black objects must not point to white objects
- Captures writes to black objects
 - If the target is gray, do nothing
 - If the target is white, turn the black object gray
- Eventually the black object is re-scanned
 - The reference to the white object is located
 - If the reference has been overwritten, it is not marked

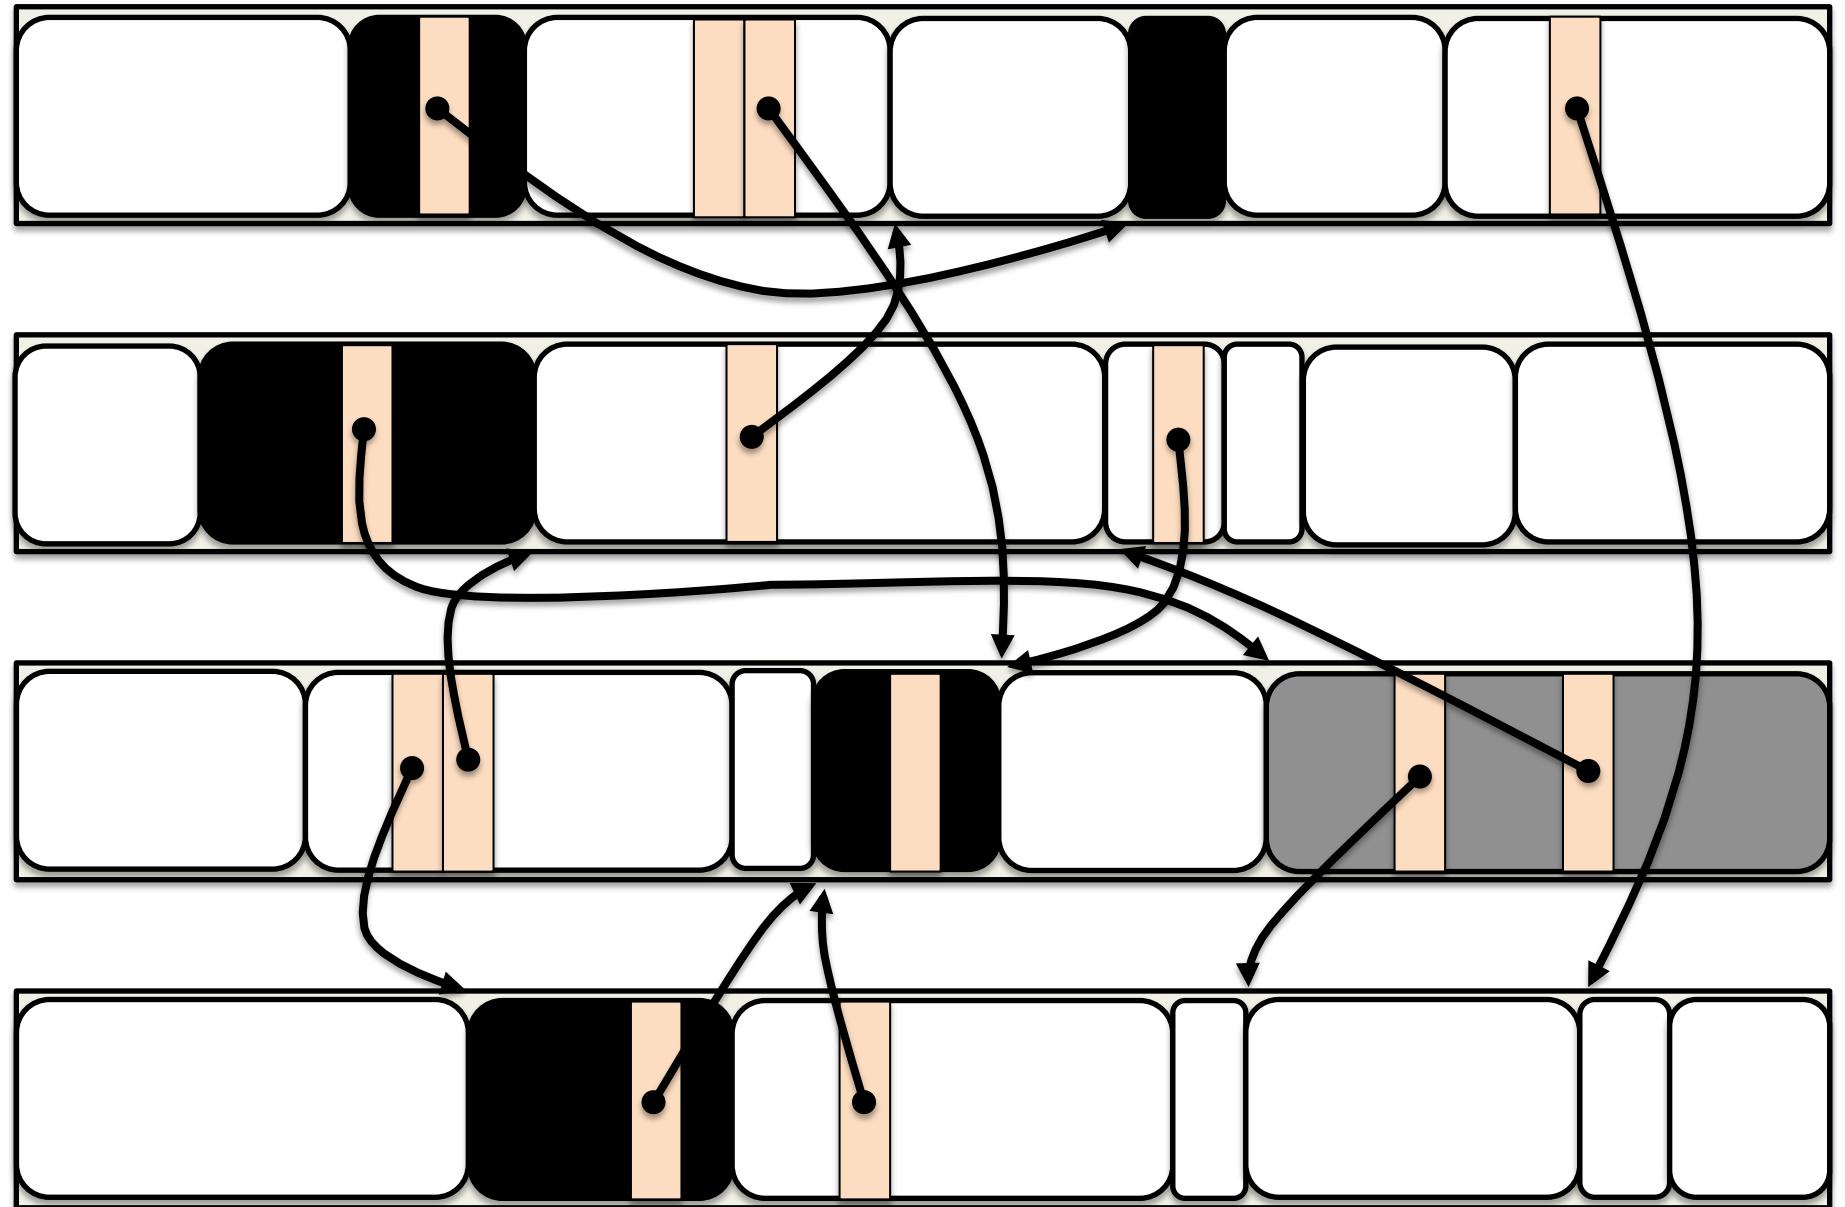


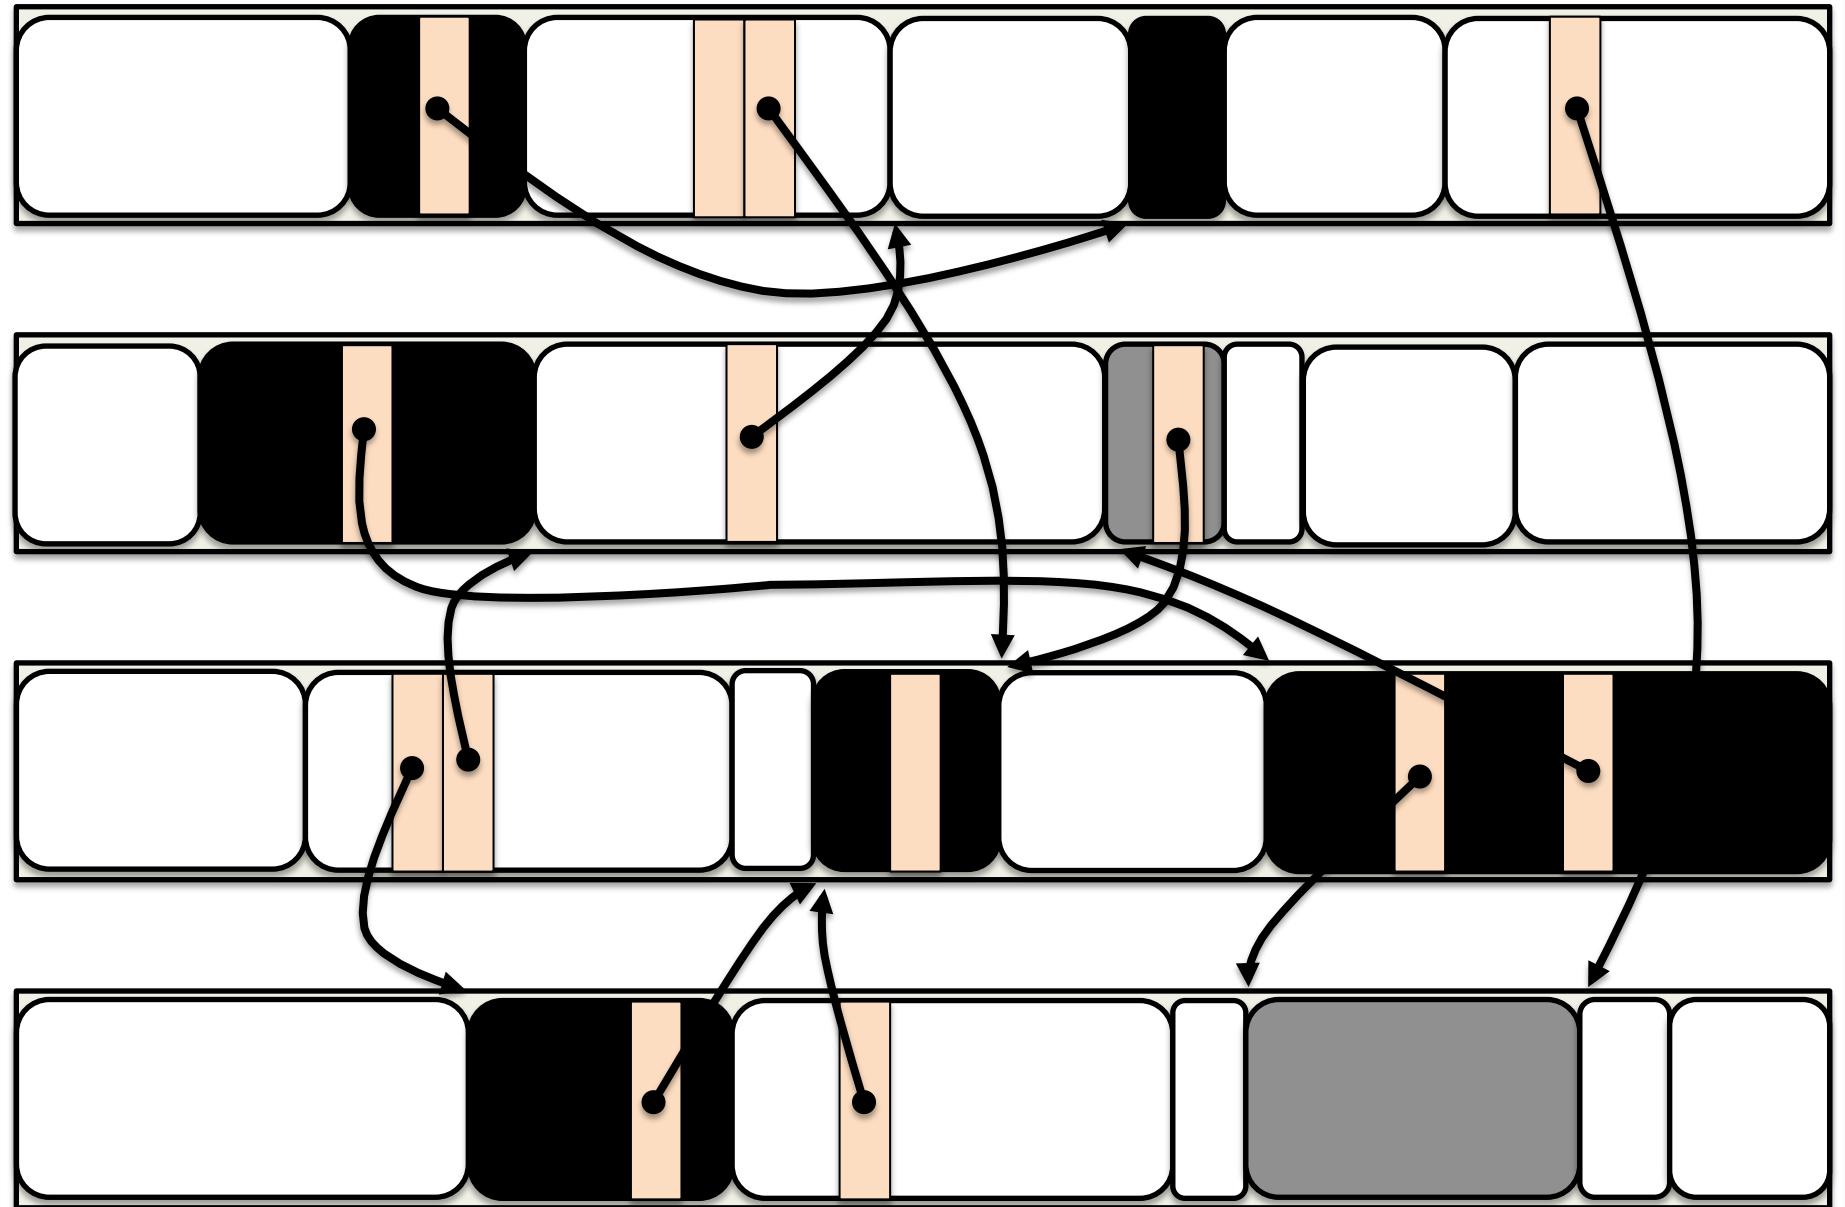


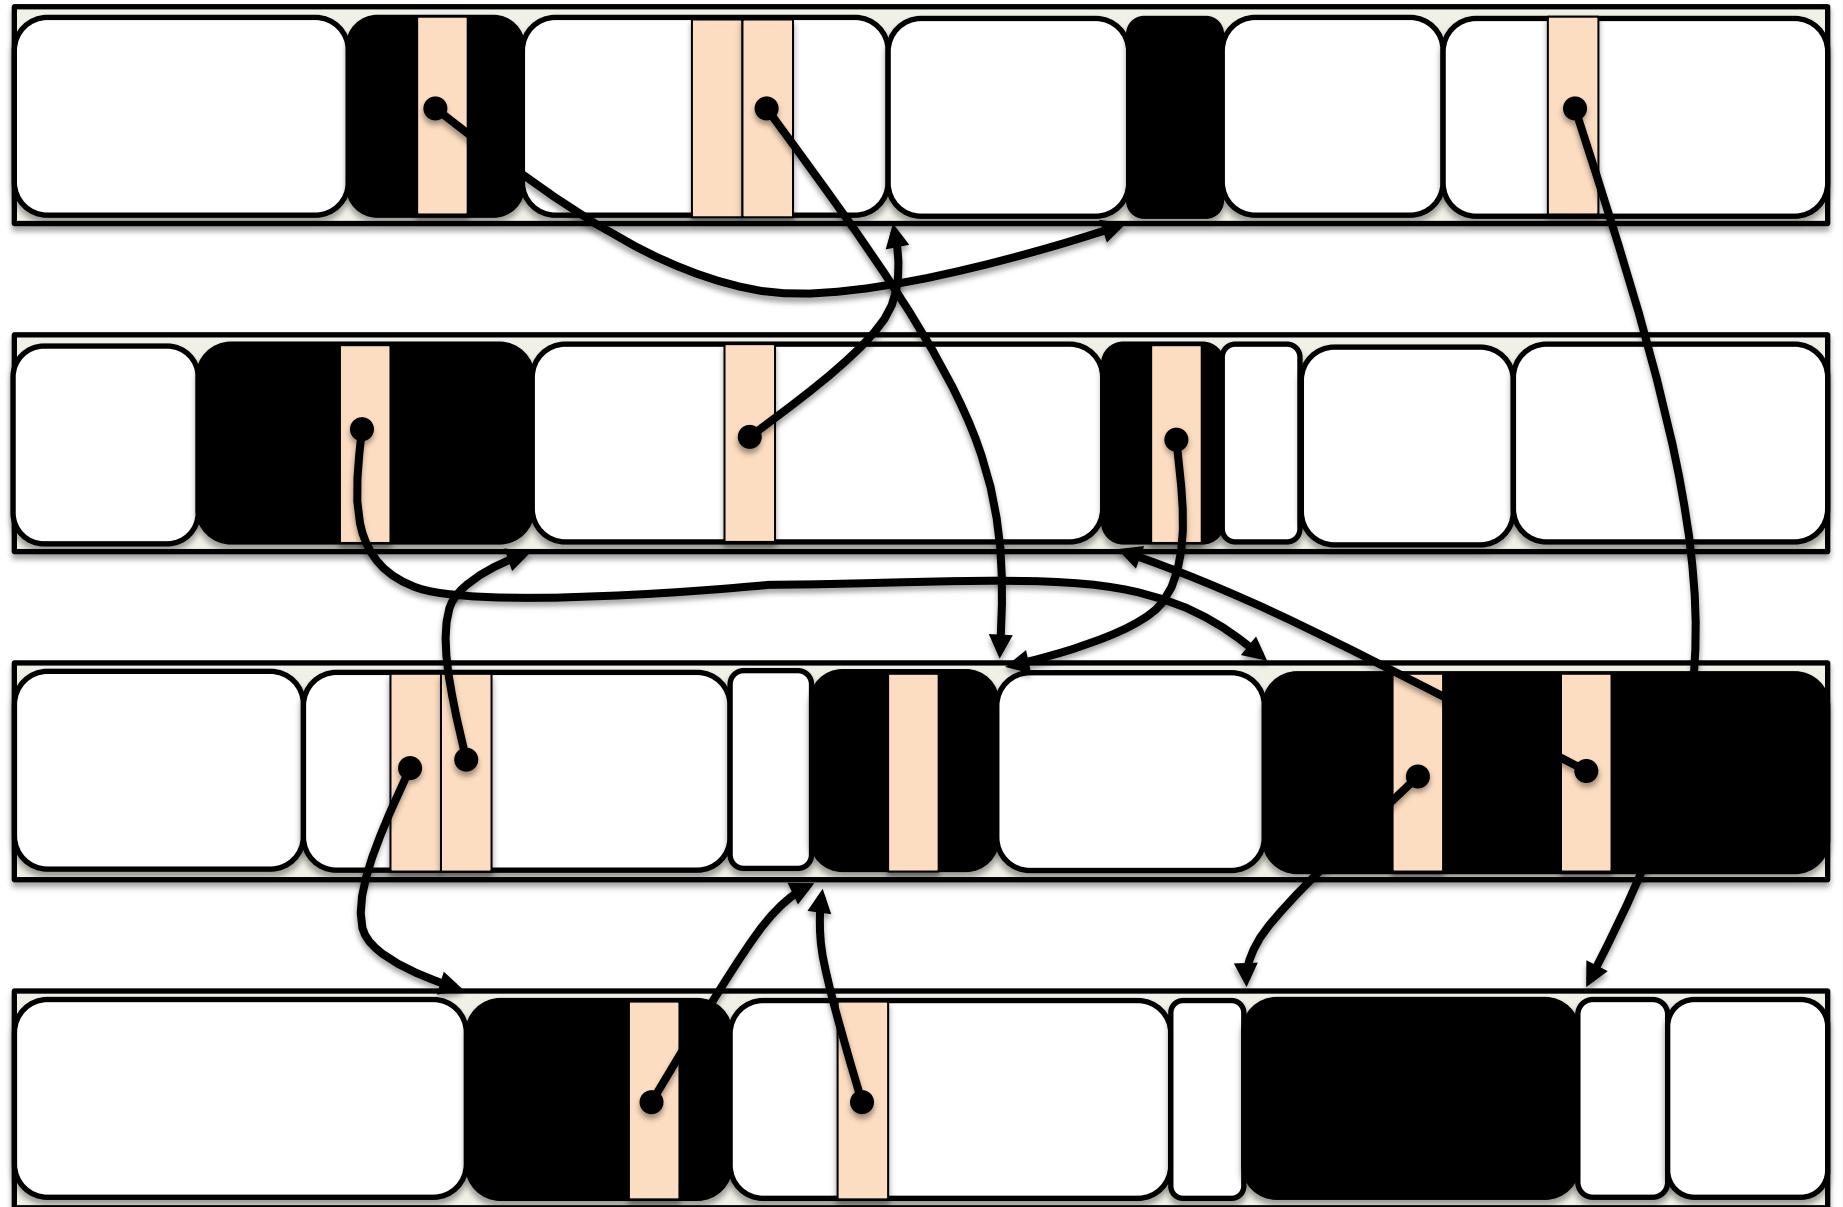












Incremental Update Barrier

- If the target object is white
 - And the destination object is black
 - Add the destination to the gray stack

Allocation

- Objects in Incremental Update allocated white
 - Object marked if it is stored in a black object
 - If object is never reachable, it is never marked
- Has potential to skip a lot of garbage
 - Most objects are never referred to from the heap
- Concurrent mark maintain generational behavior

Remark Phase

- Must re-scan roots when the mark is over
 - Write barrier tracks updates to the heap
 - No write barrier on stack or local writes
 - Only reference to a white object may be in a root
- Algorithm is iterative
 - A single object may turn gray and black many times
 - Remark phase may trigger more changes
- Algorithm is guaranteed to terminate
 - Citation available on request

Incremental Update

- Less conservative than Snapshot At The Beginning
 - Allocating white retains less garbage
 - White pointers may be overwritten before scanning
- Implementation is considerably trickier
 - Need to iterate over root scanning
- Mark phase may last longer
 - Write barriers must be active for longer

Concurrent Write Barriers

- Concurrent marking barriers have overhead
 - Need to check the mark state of source and target
- More expensive than basic generational barrier
 - Runs at the same time as the mutator
 - Must be active for all threads, not just GC
- Want to be able to turn off barriers at times
 - Use the expensive barrier only during marking

Phased Write Barrier

- If the source object is gray
 - And the target object is white
 - Add the target object to the gray stack

Phased Write Barrier

- If the system is currently marking
 - And the source object is gray
 - And the target object is white
 - Add the target object to the gray stack

Changing Write Barrier State

- We know from last week that is can be a race
 - The GC state is changed by one thread
 - Another thread doesn't see the update
- State change requires a safepoint
 - Safepoint ensures that there is a memory fence
 - Each thread gets the latest value of the state variable
- Safepoint operation is very lightweight
 - Don't need to synchronize with all other threads

Incremental Marking

- SATB lets us mark an evolving object graph
 - Nothing requires that the algorithm be concurrent
- We can easily change concurrent to incremental
 - Do a little bit of work every so often
 - Ideal time is during allocation
- Spread the work out without adding threads

Concurrent Mark Sweep

- Previous iteration of Hotspot's GC
 - Still in wide production use
- Composed from the primitives that we've now seen
 - Generational algorithm with Mark/Sweep mature space
 - Incremental update mark phase
 - Parallel sweep phase
- Sweeping can be completely parallel
 - By the time an object is swept, there are no references
 - No way for the mutator to be affected

Concurrent Copying

- So far all copying has happened during GC
 - All mutator threads are stopped
 - We know that the heap is in a bad state
 - We don't really care
- Copy phase can be slow
 - Minimize by copying smaller areas
 - Leads to more pauses, albeit shorter ones
- Ideally we could copy objects concurrently

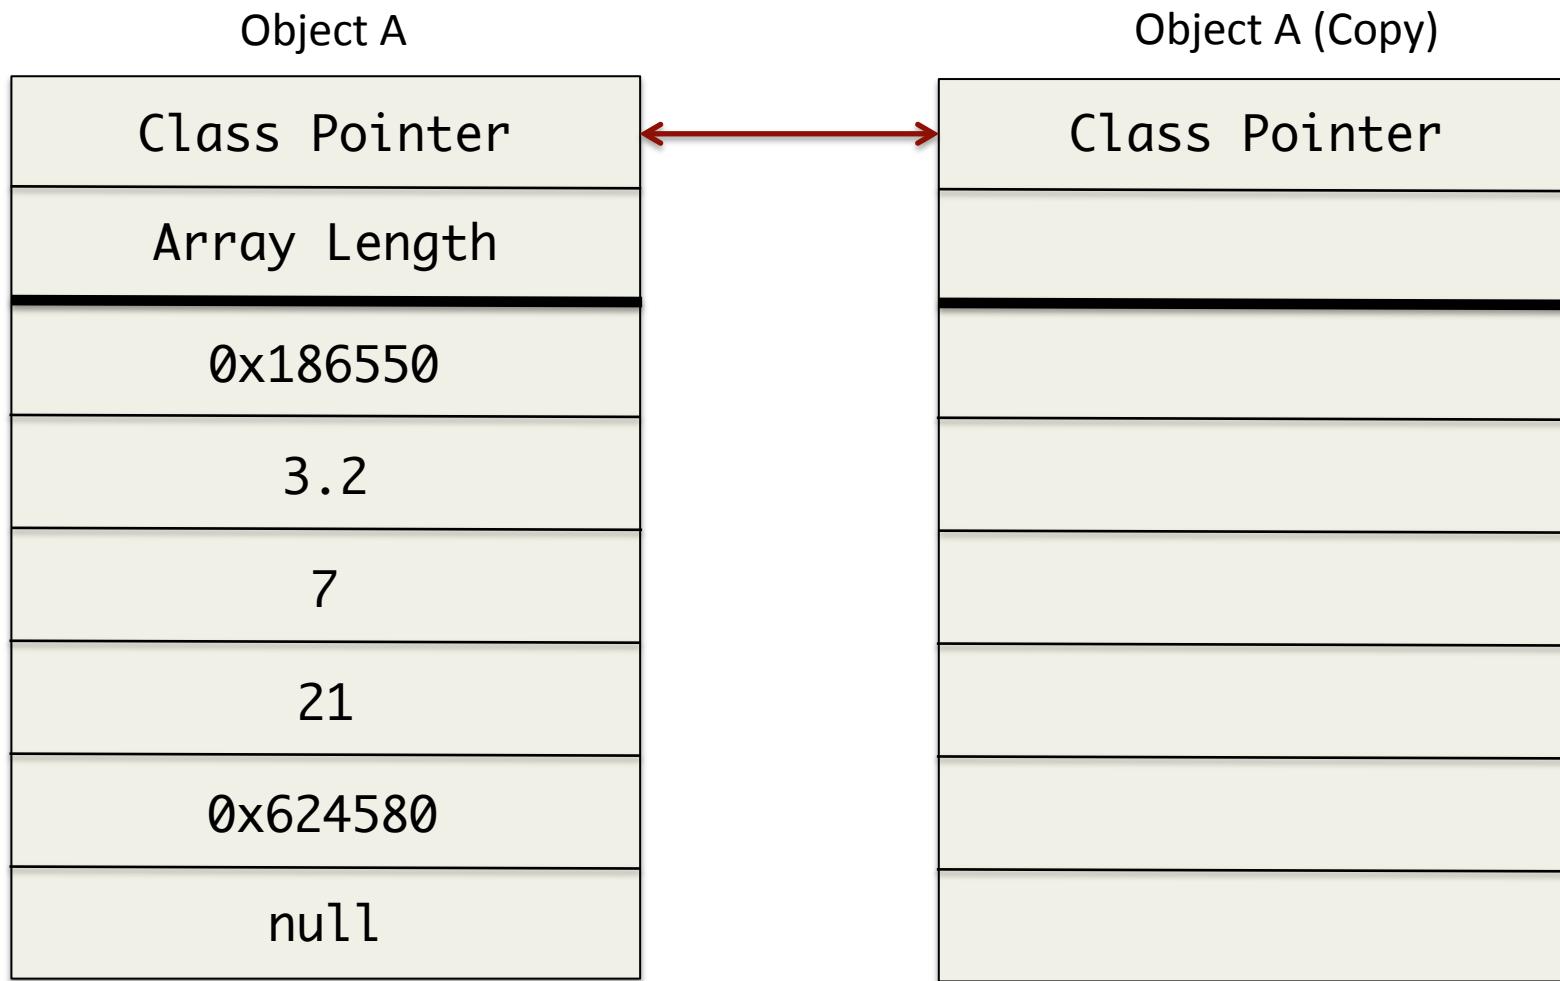
Object A

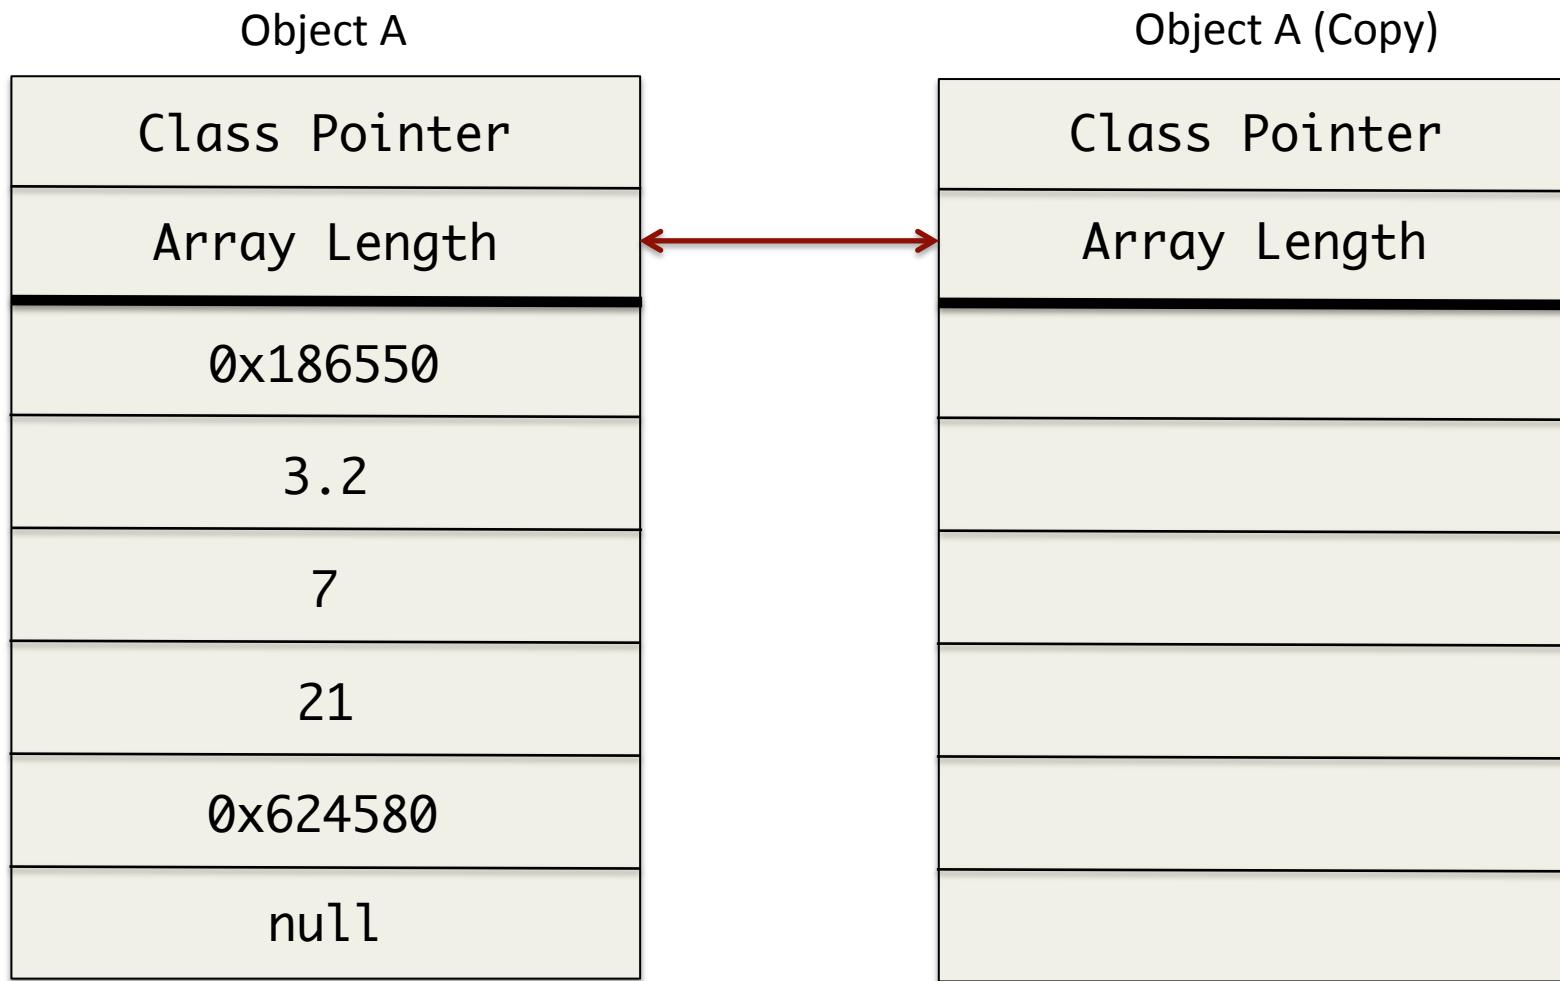
Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

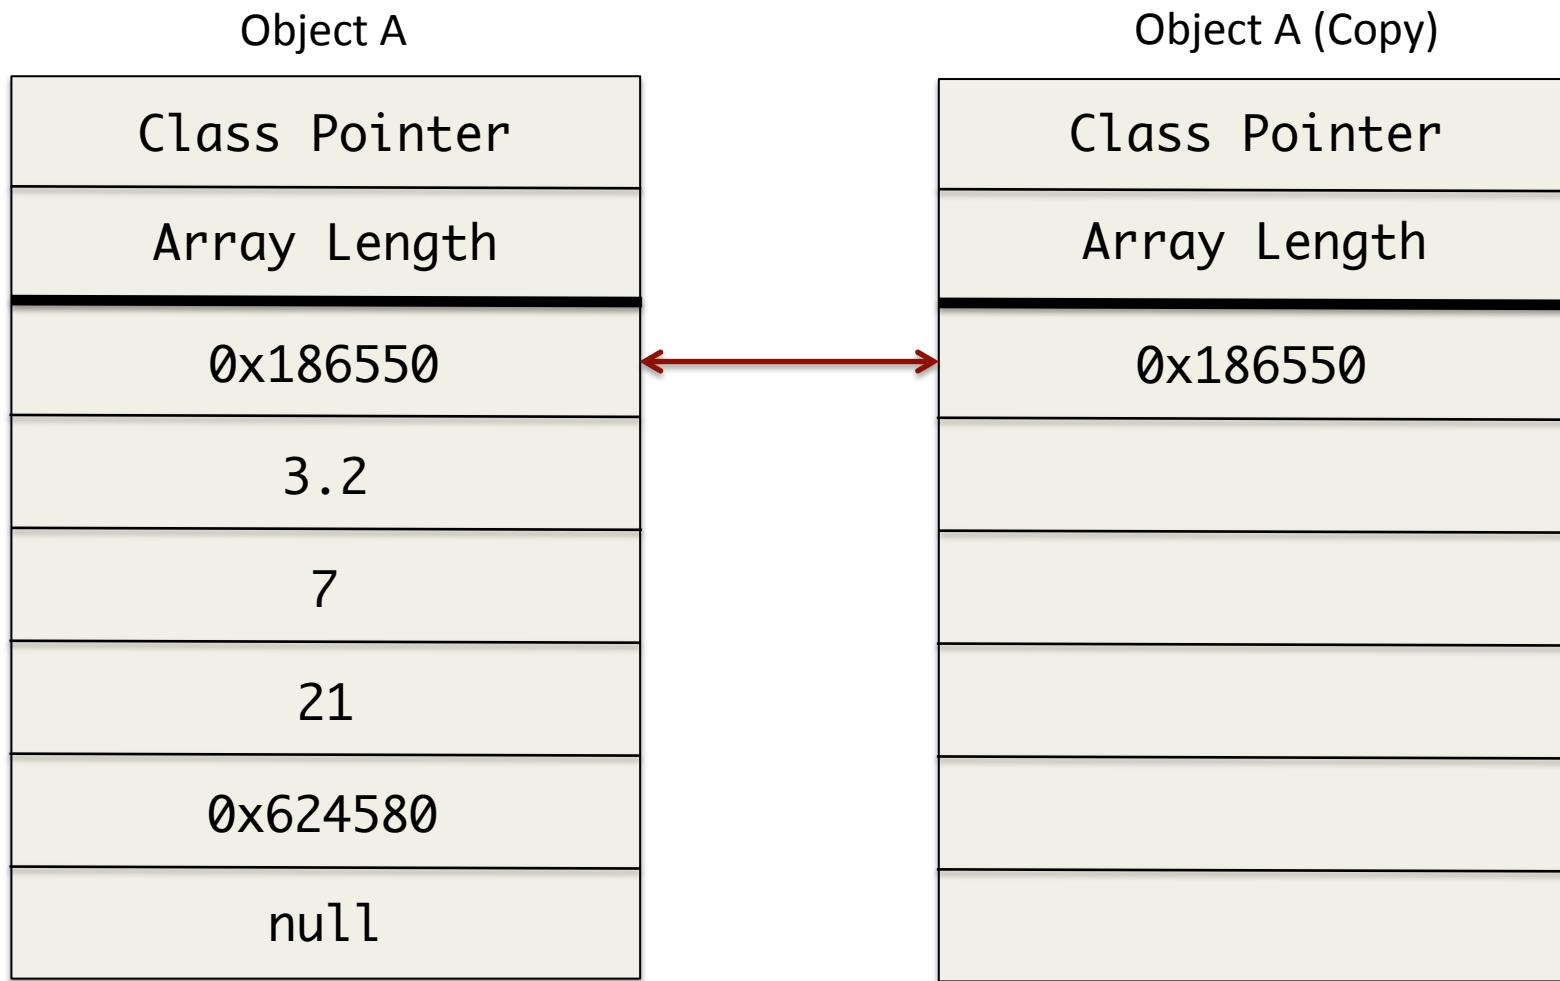
Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)







Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null



Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7



Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21



Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580



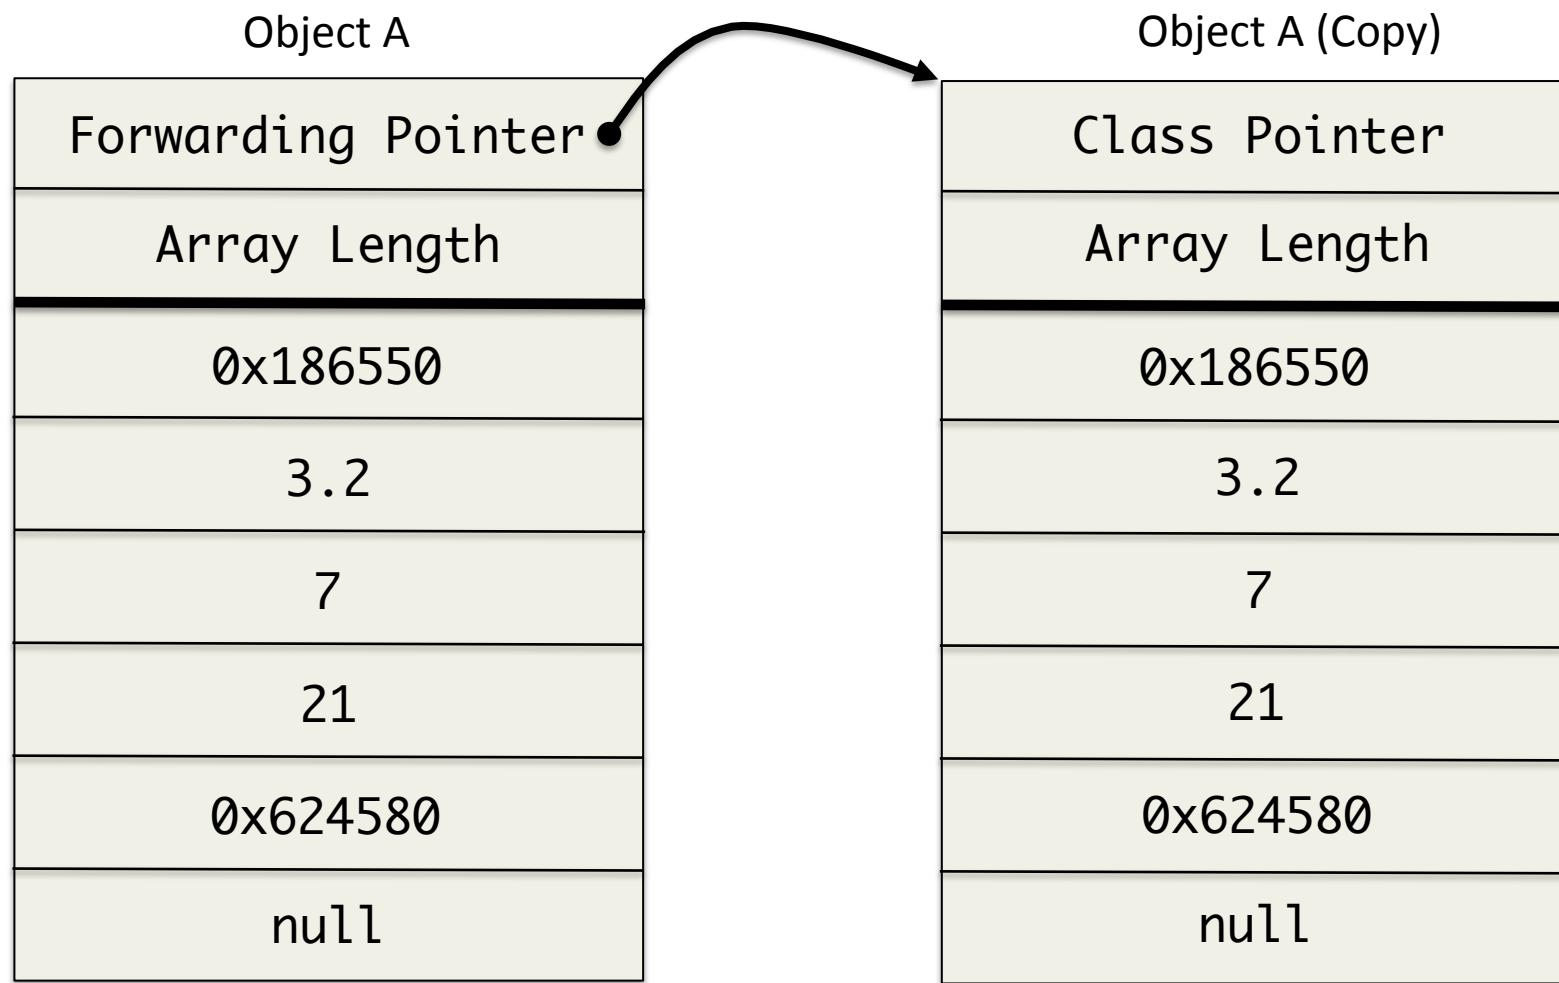
Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

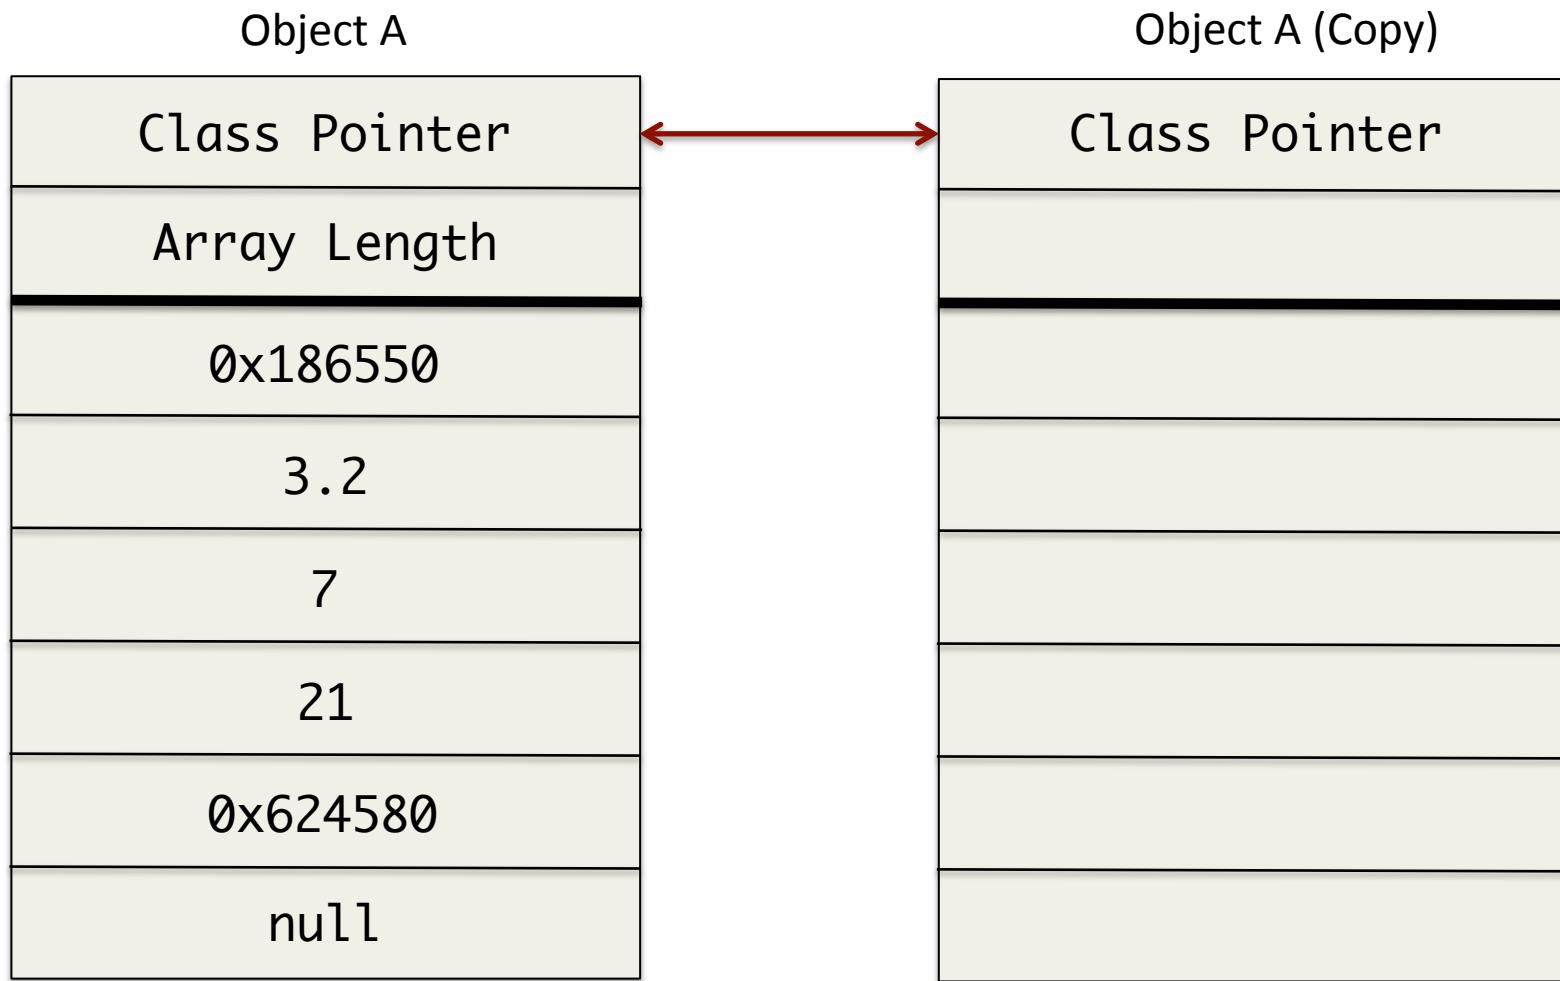


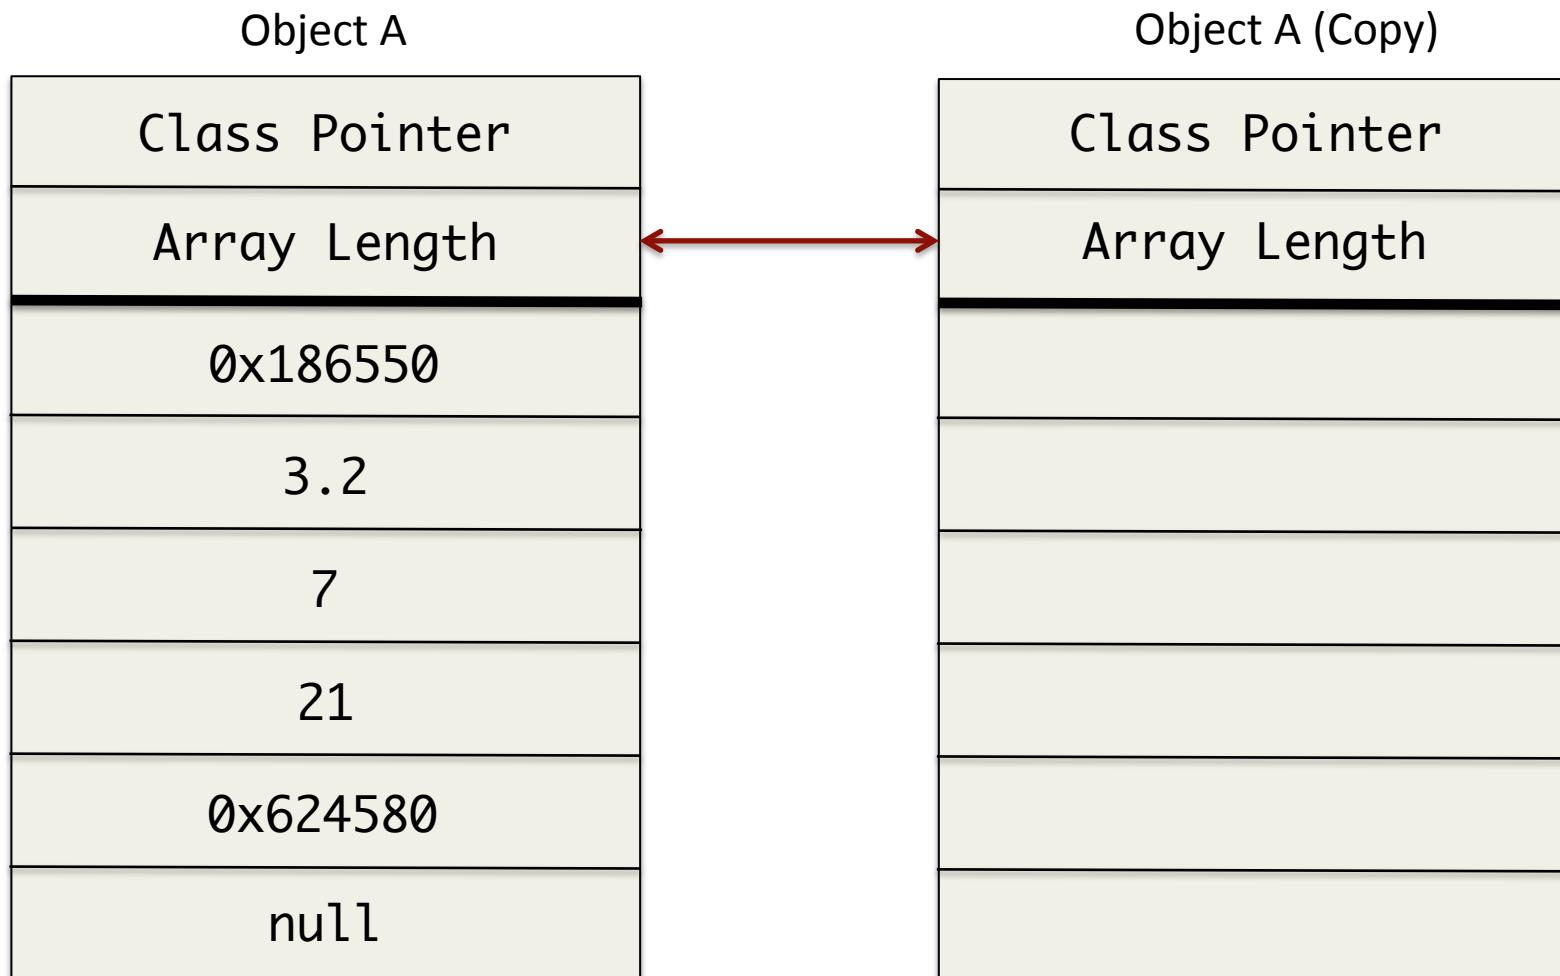


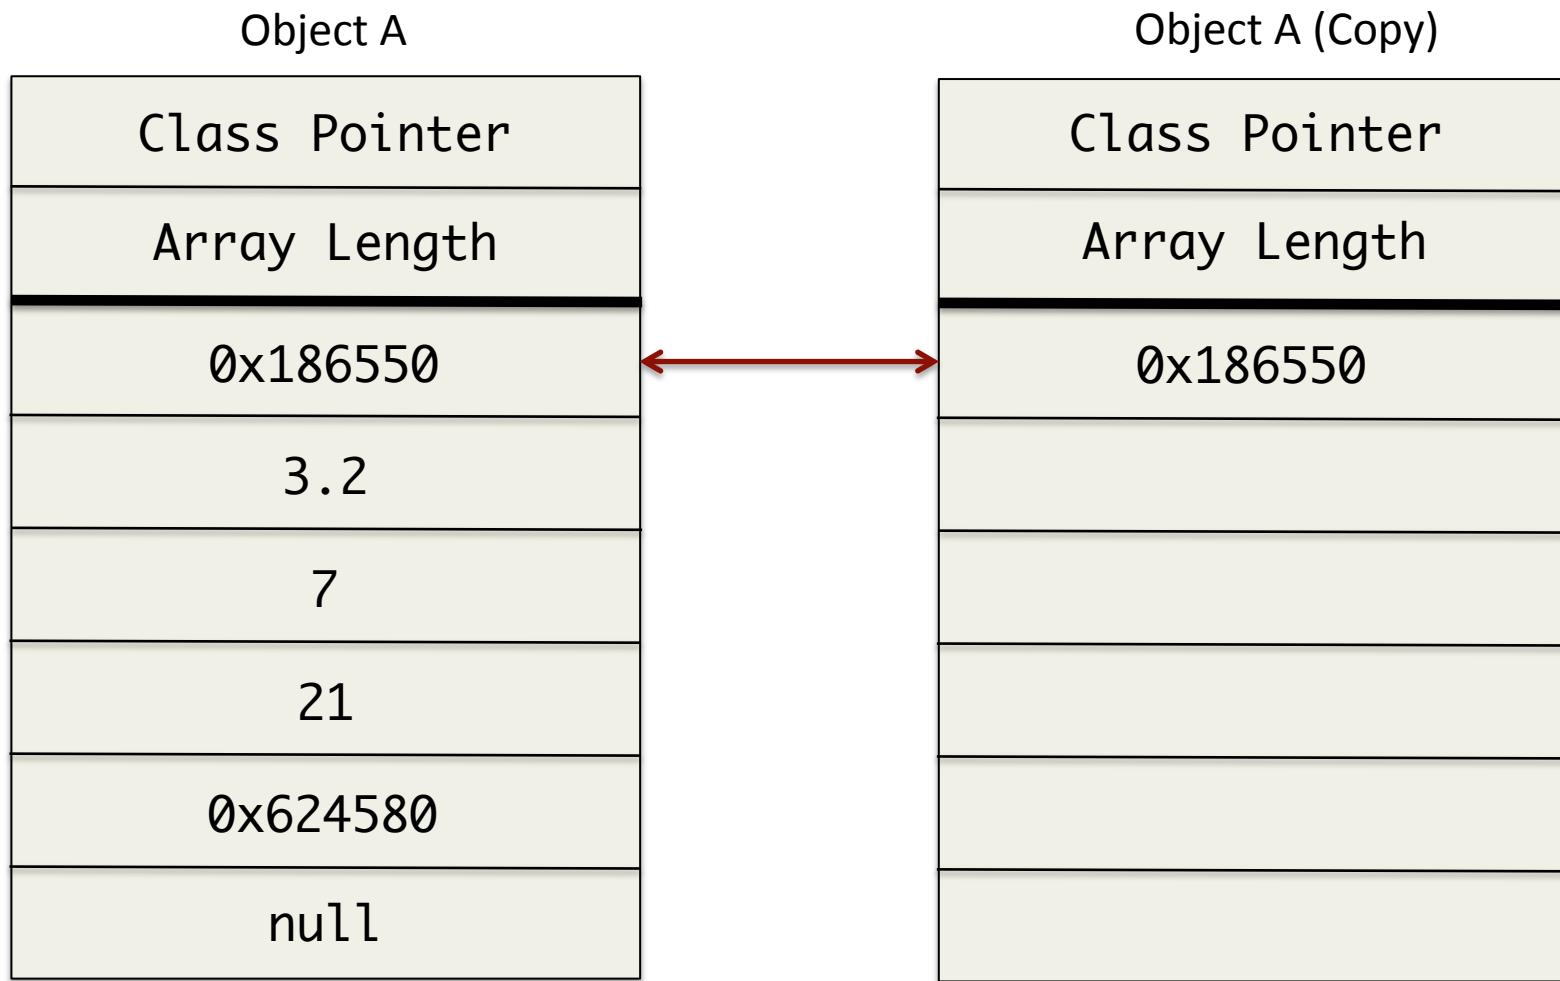
Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)







Object A

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null



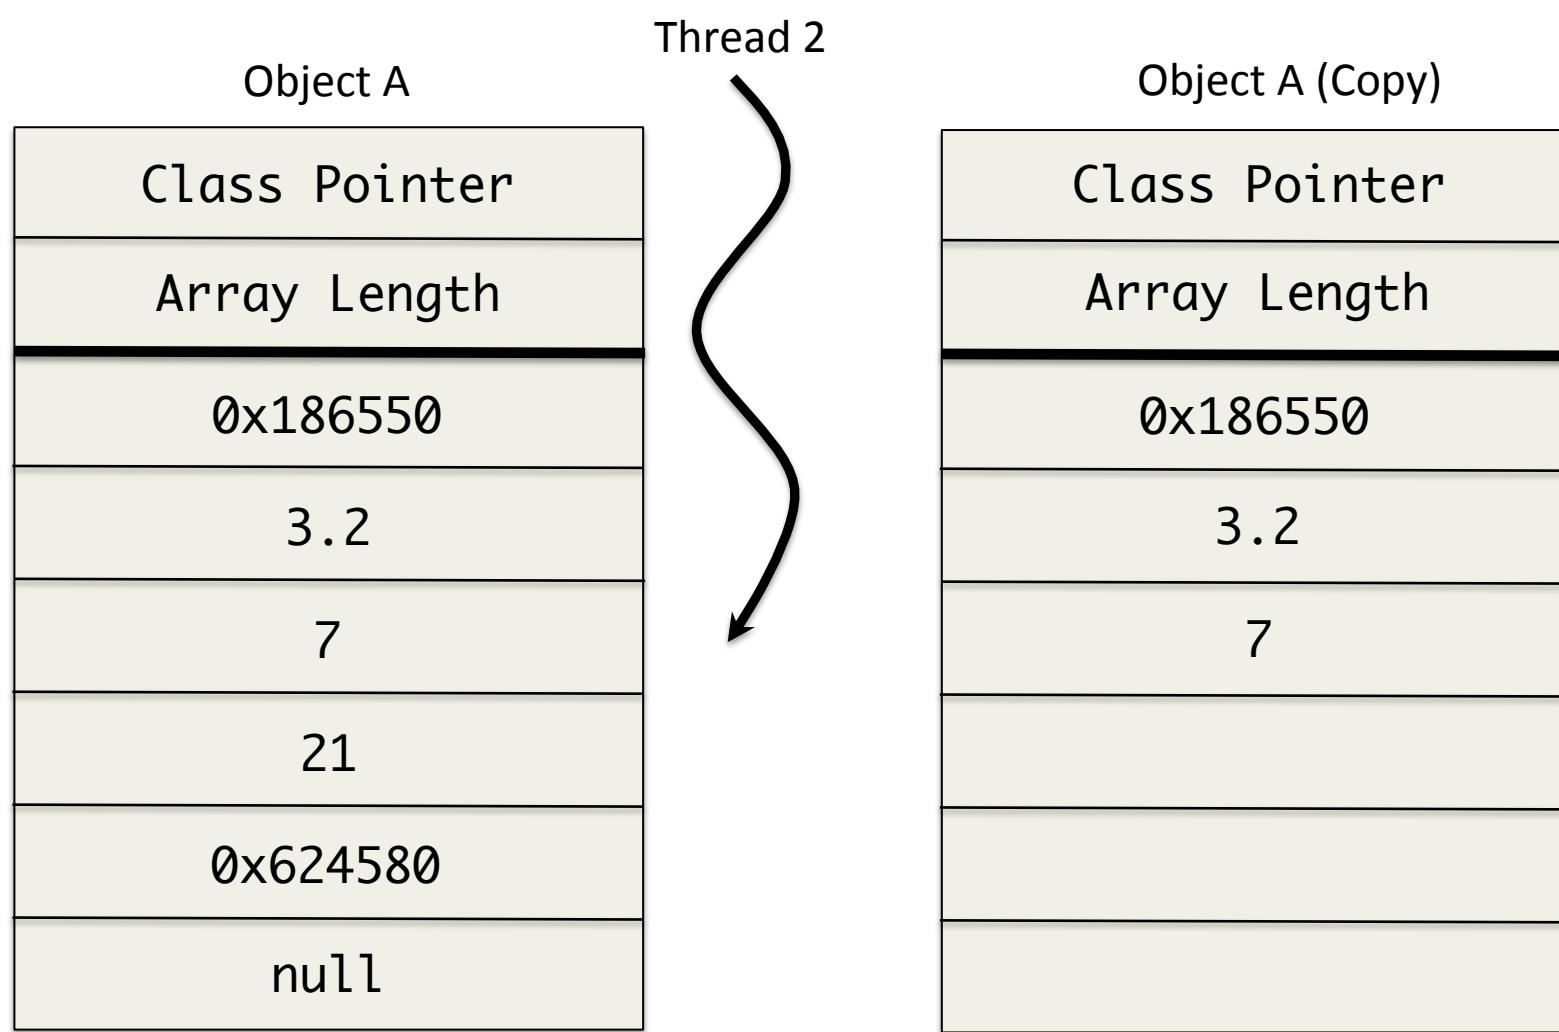
Object A

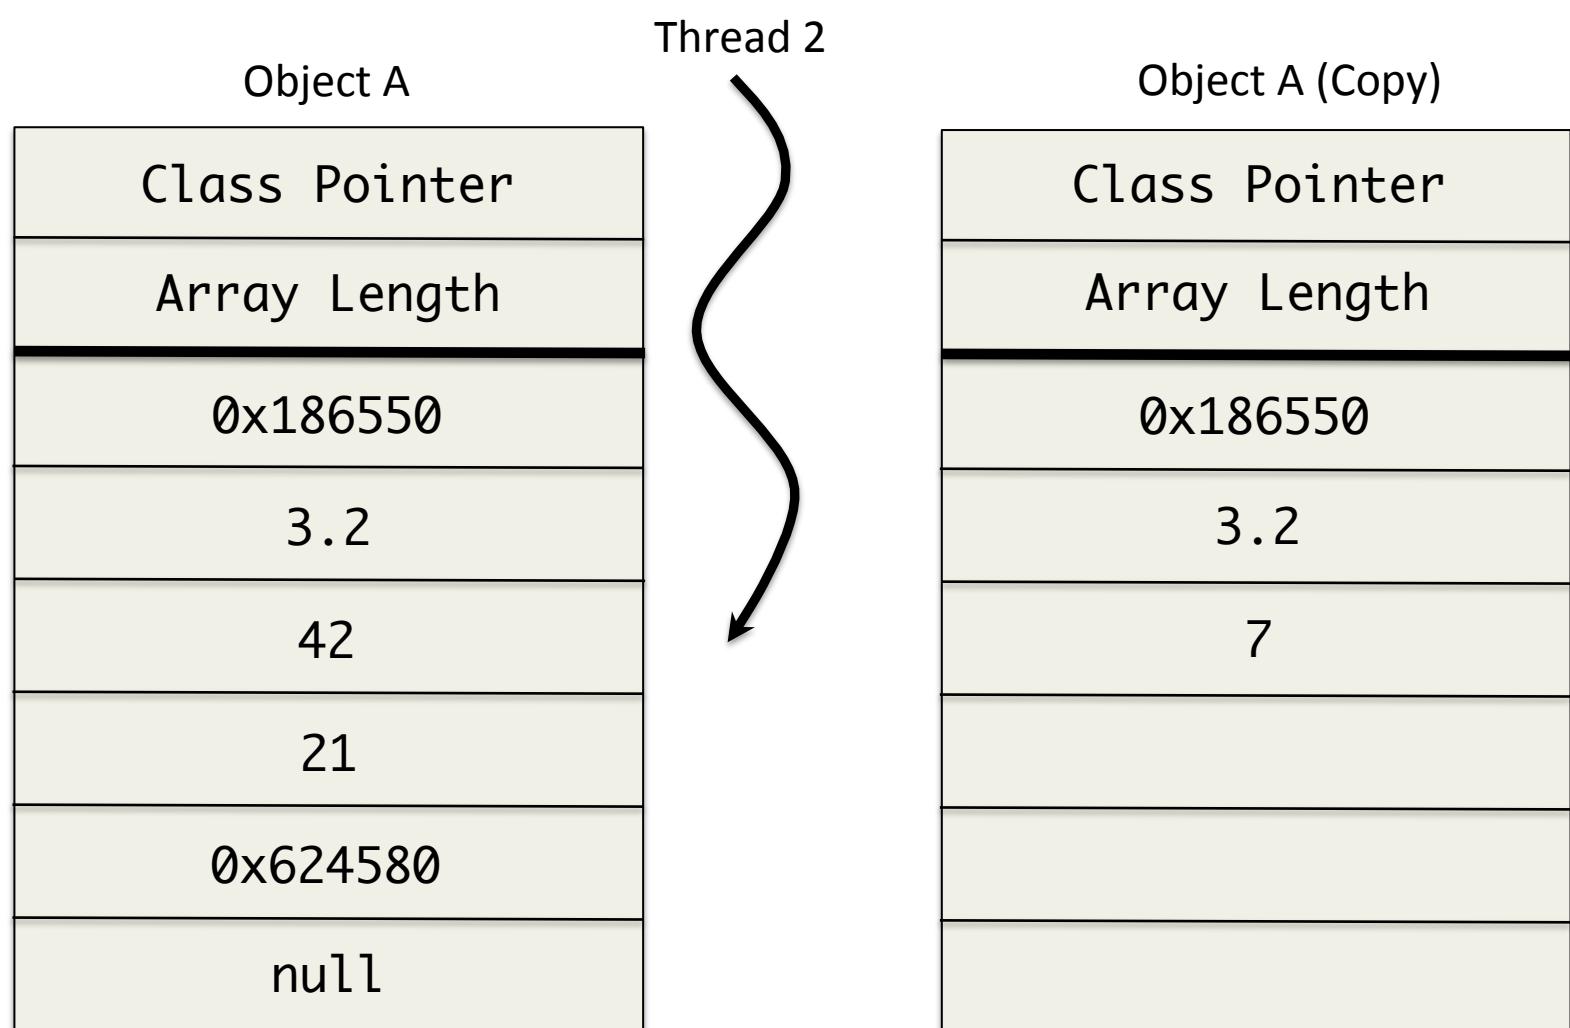
Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7







Object A

Class Pointer
Array Length
0x186550
3.2
42
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21



Object A

Class Pointer
Array Length
0x186550
3.2
42
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580



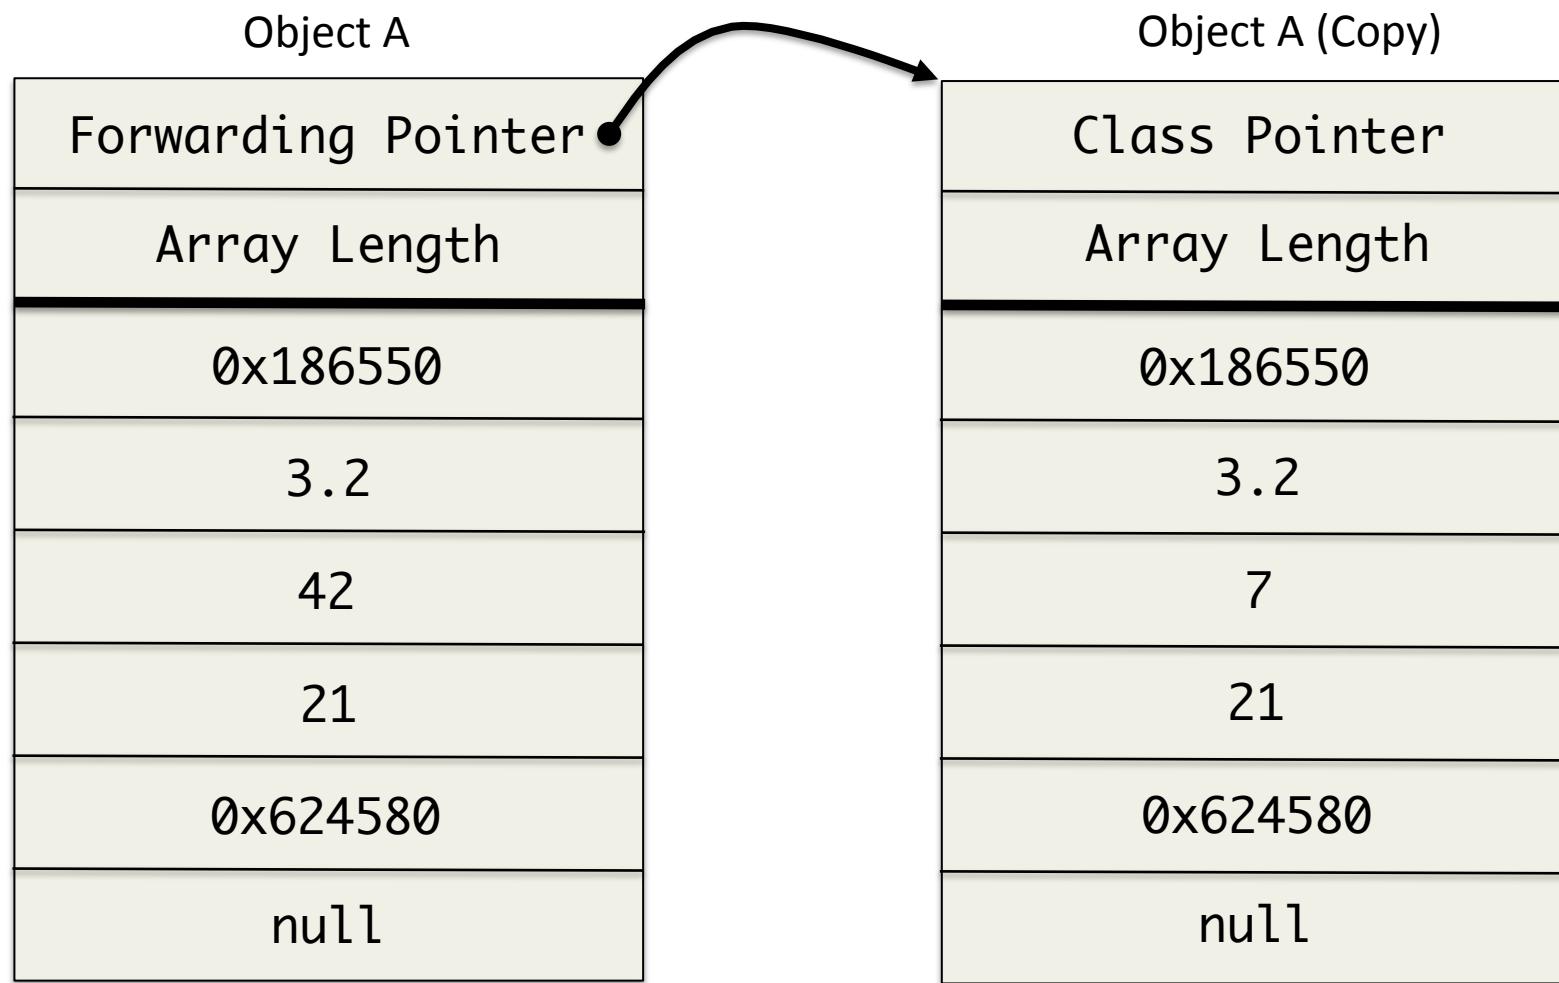
Object A

Class Pointer
Array Length
0x186550
3.2
42
21
0x624580
null

Object A (Copy)

Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null





Concurrent Copying

- Lost update problem makes concurrent copy hard
 - Particularly in light of the Java memory model
- One option would be to lock each object
 - GC and mutator threads contend for access
 - Set up contention manager to favor the mutator
- Some other options available

Sapphire

- The Sapphire collector is split into three phases
 - Mark, copy, flip
- Creates a shell for each live object prior to copy
 - Installs forwarding pointer to the shell
- Installs write barrier during copy phase
 - Barrier installed on all writes (not just pointers)
 - Mutator writes to both old and new versions

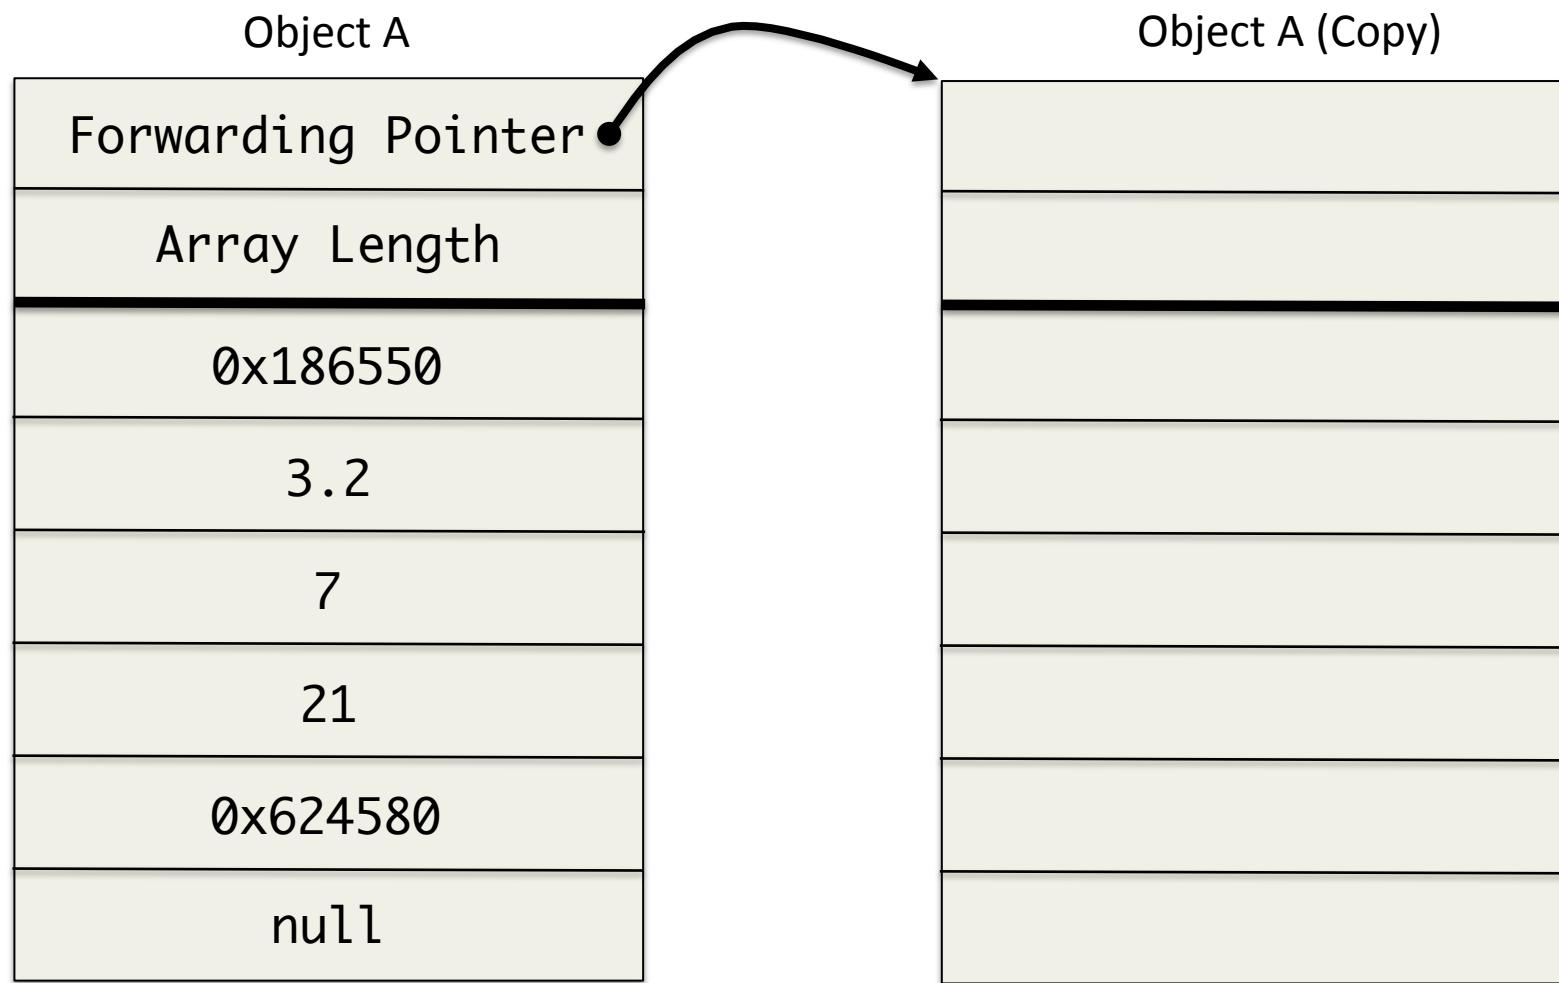
Object A

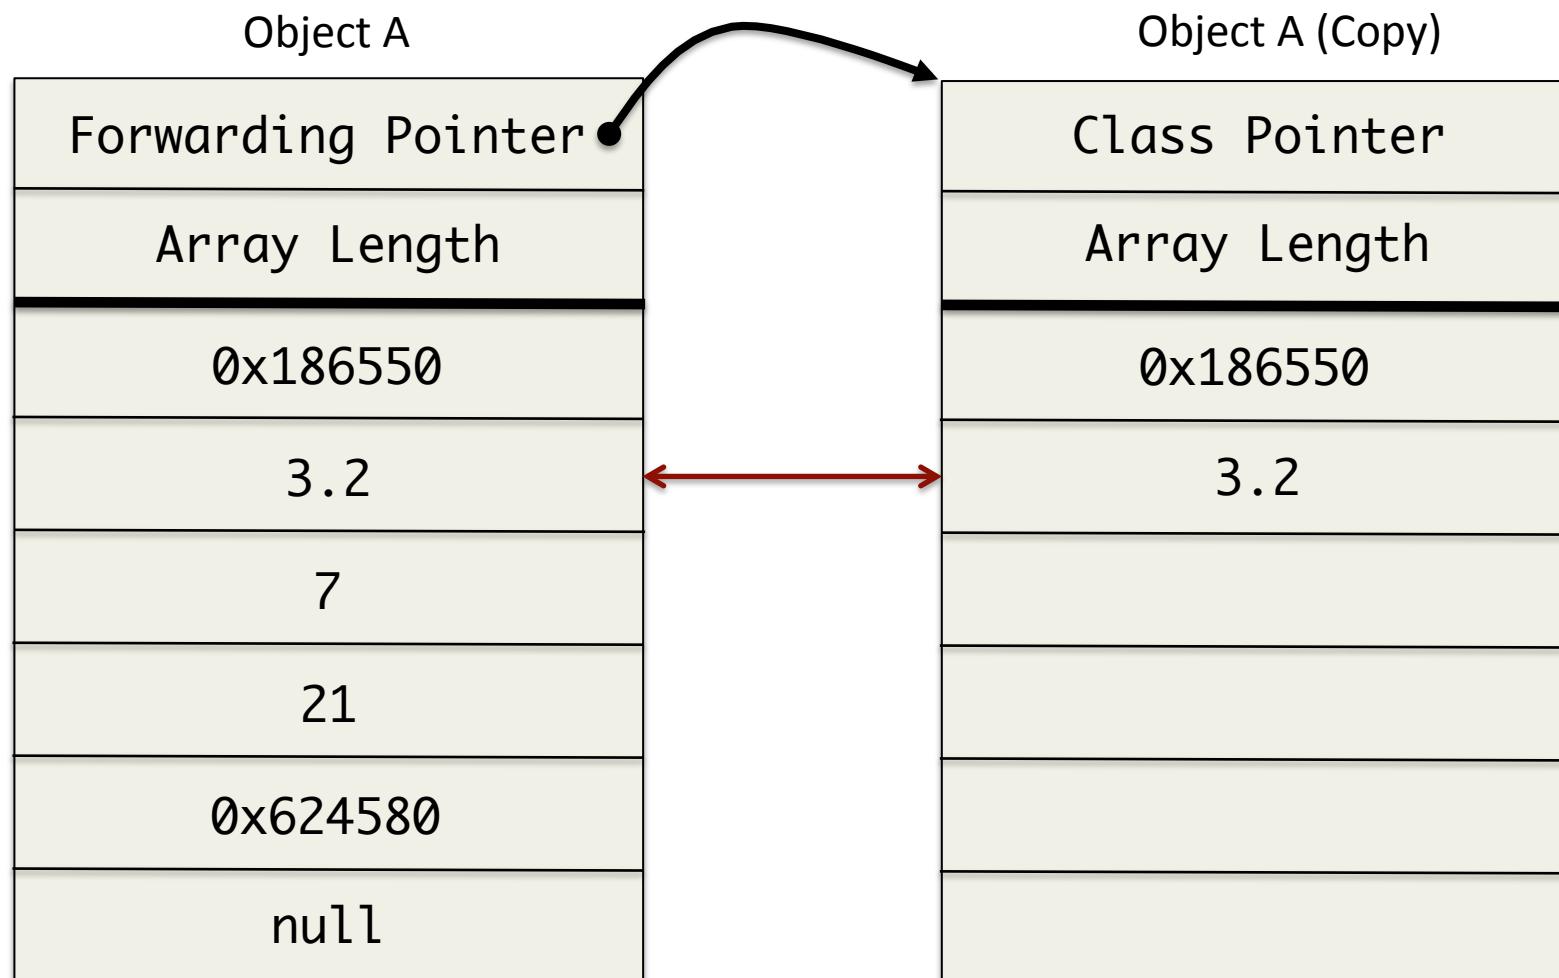
Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

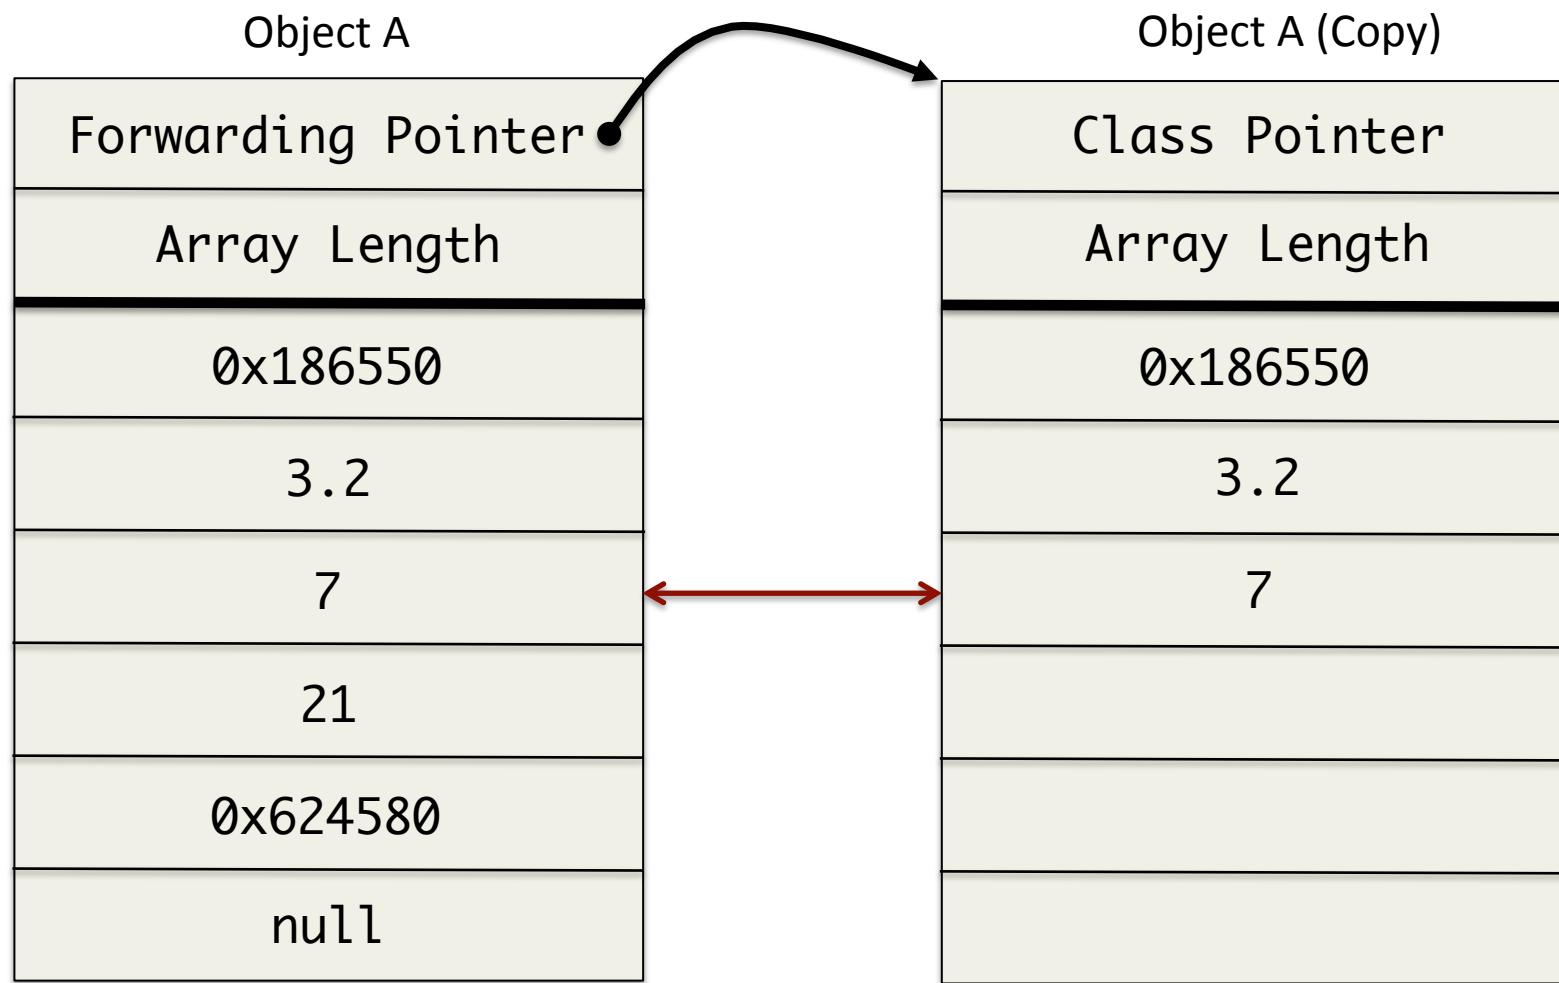
Object A

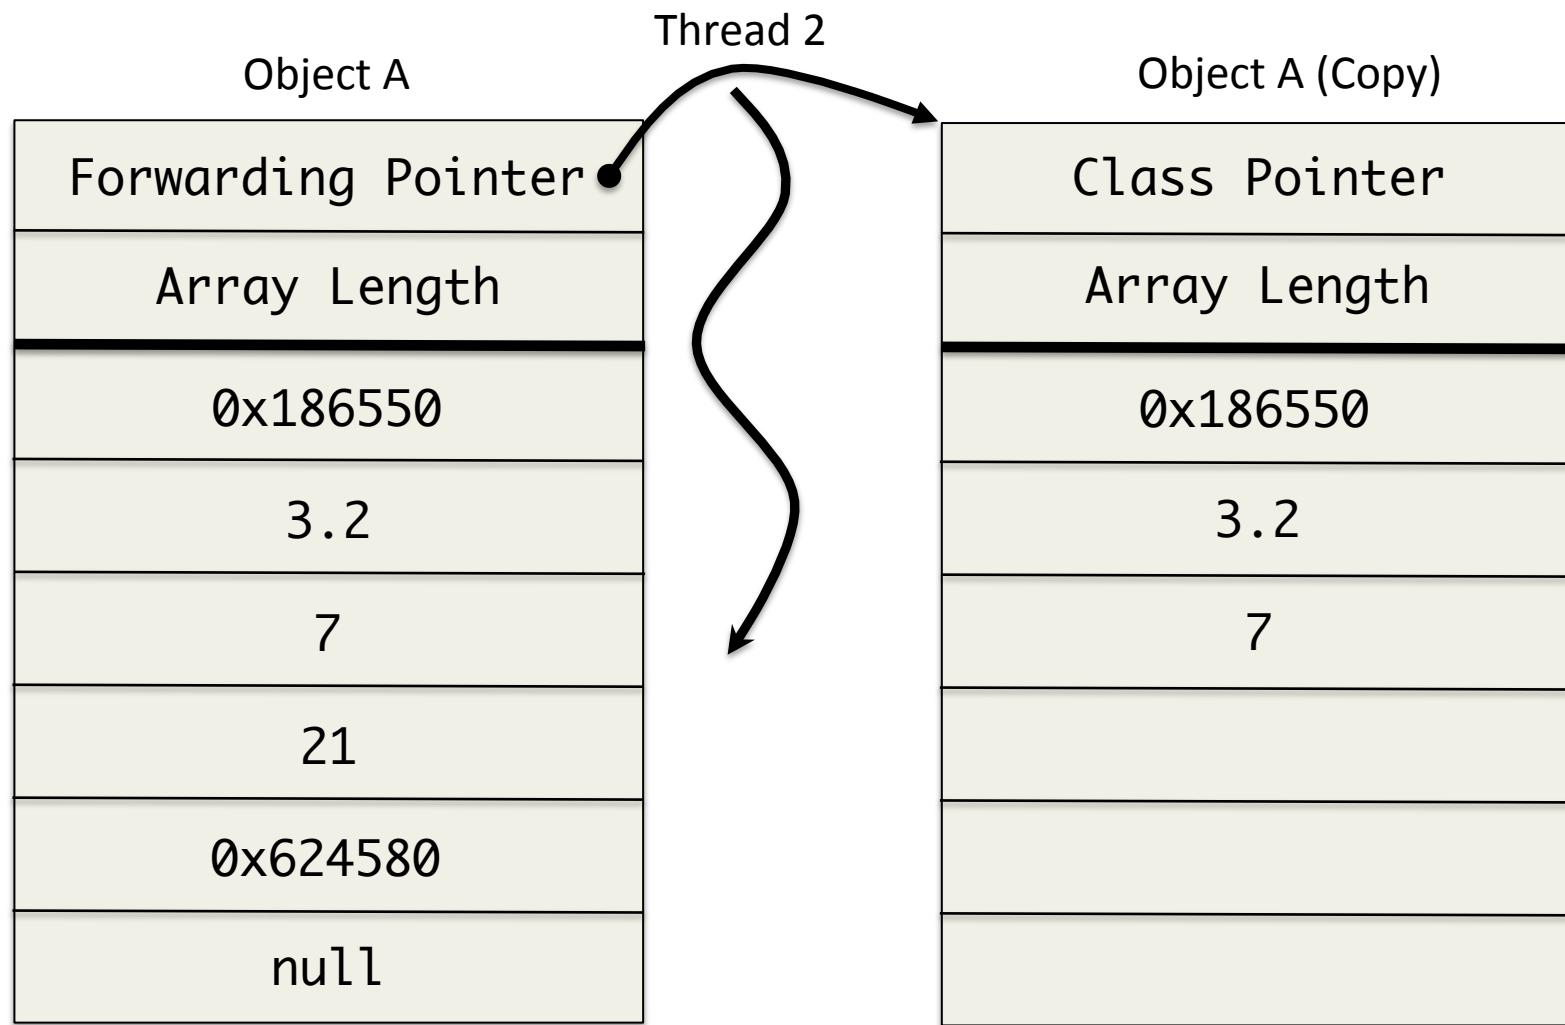
Class Pointer
Array Length
0x186550
3.2
7
21
0x624580
null

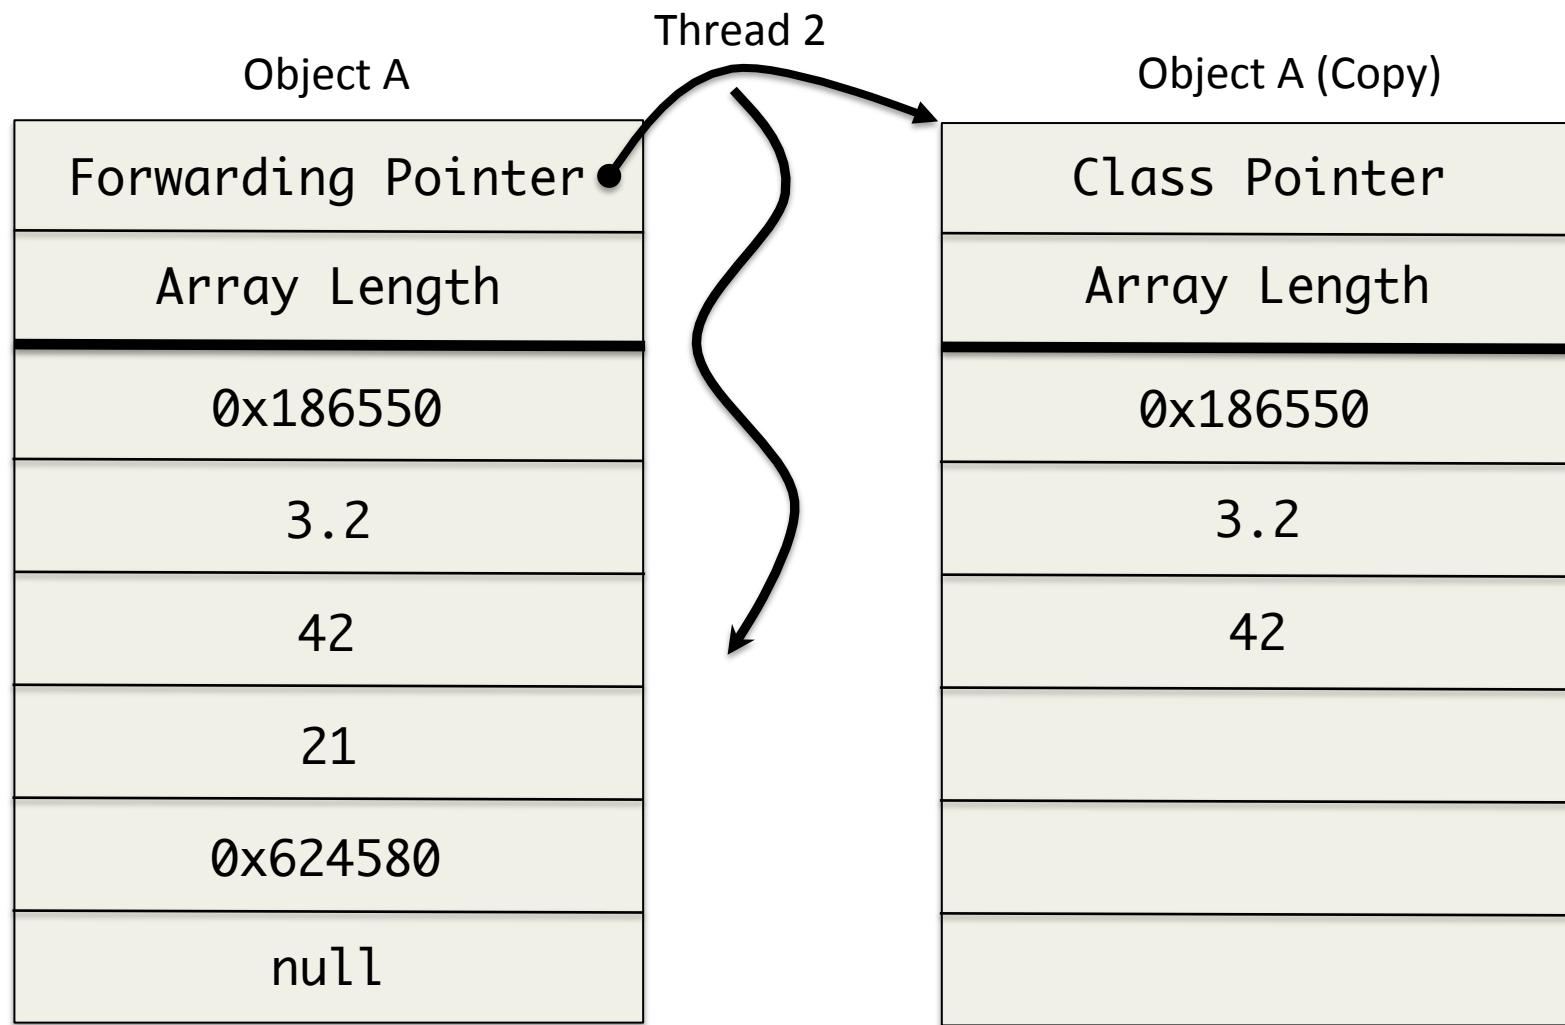
Object A (Copy)

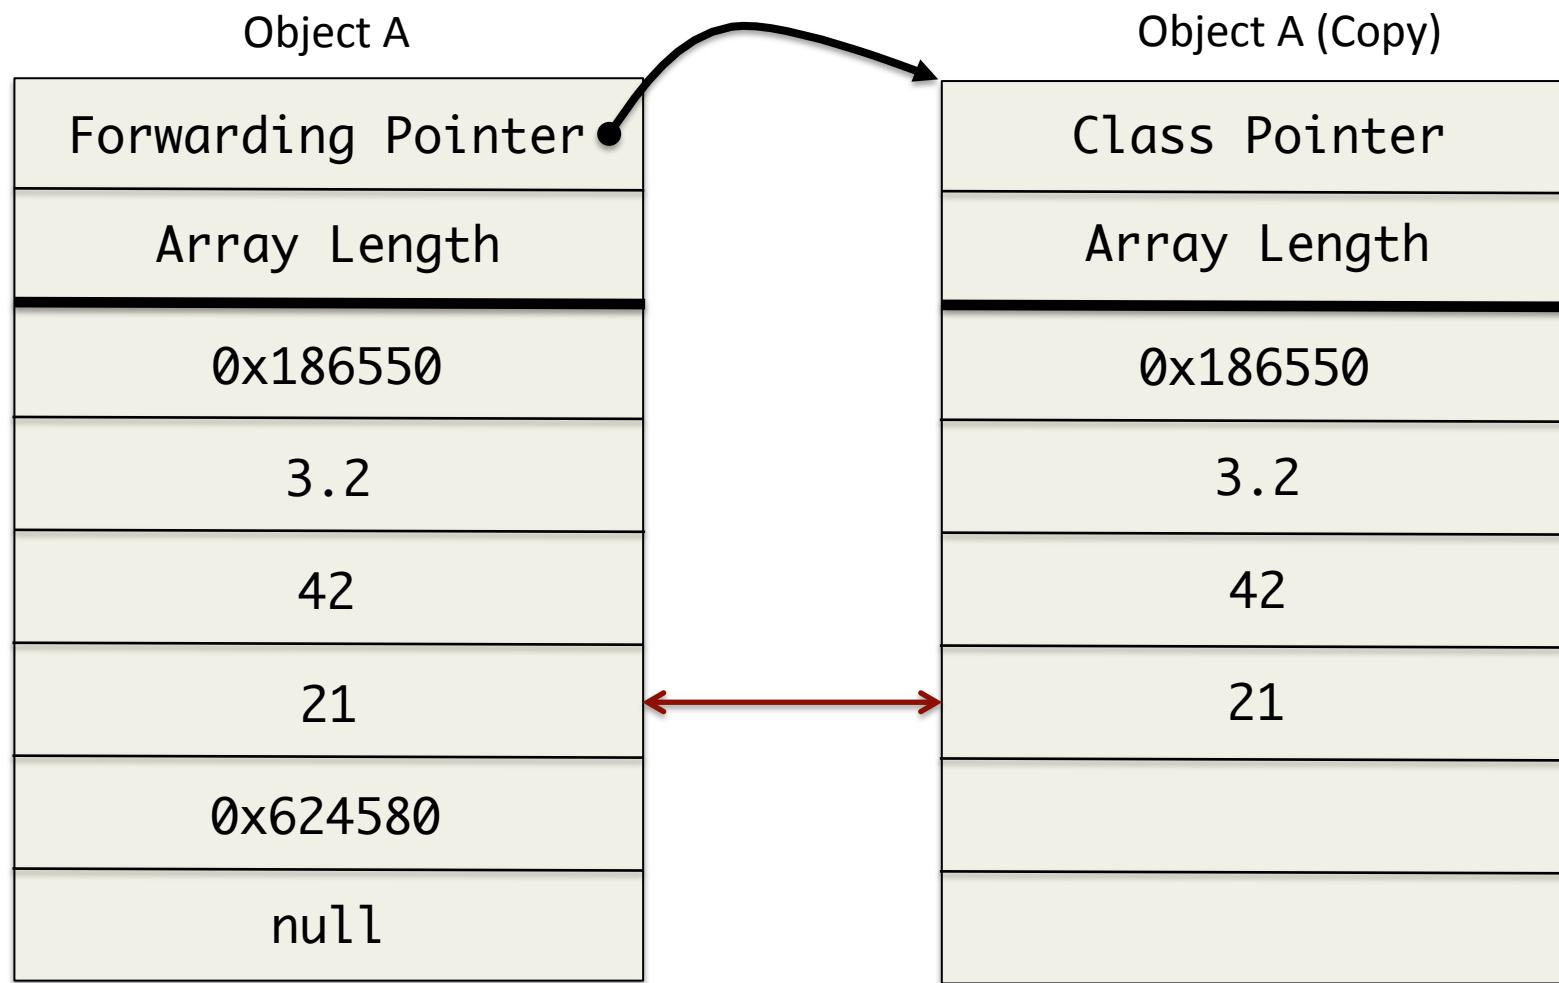


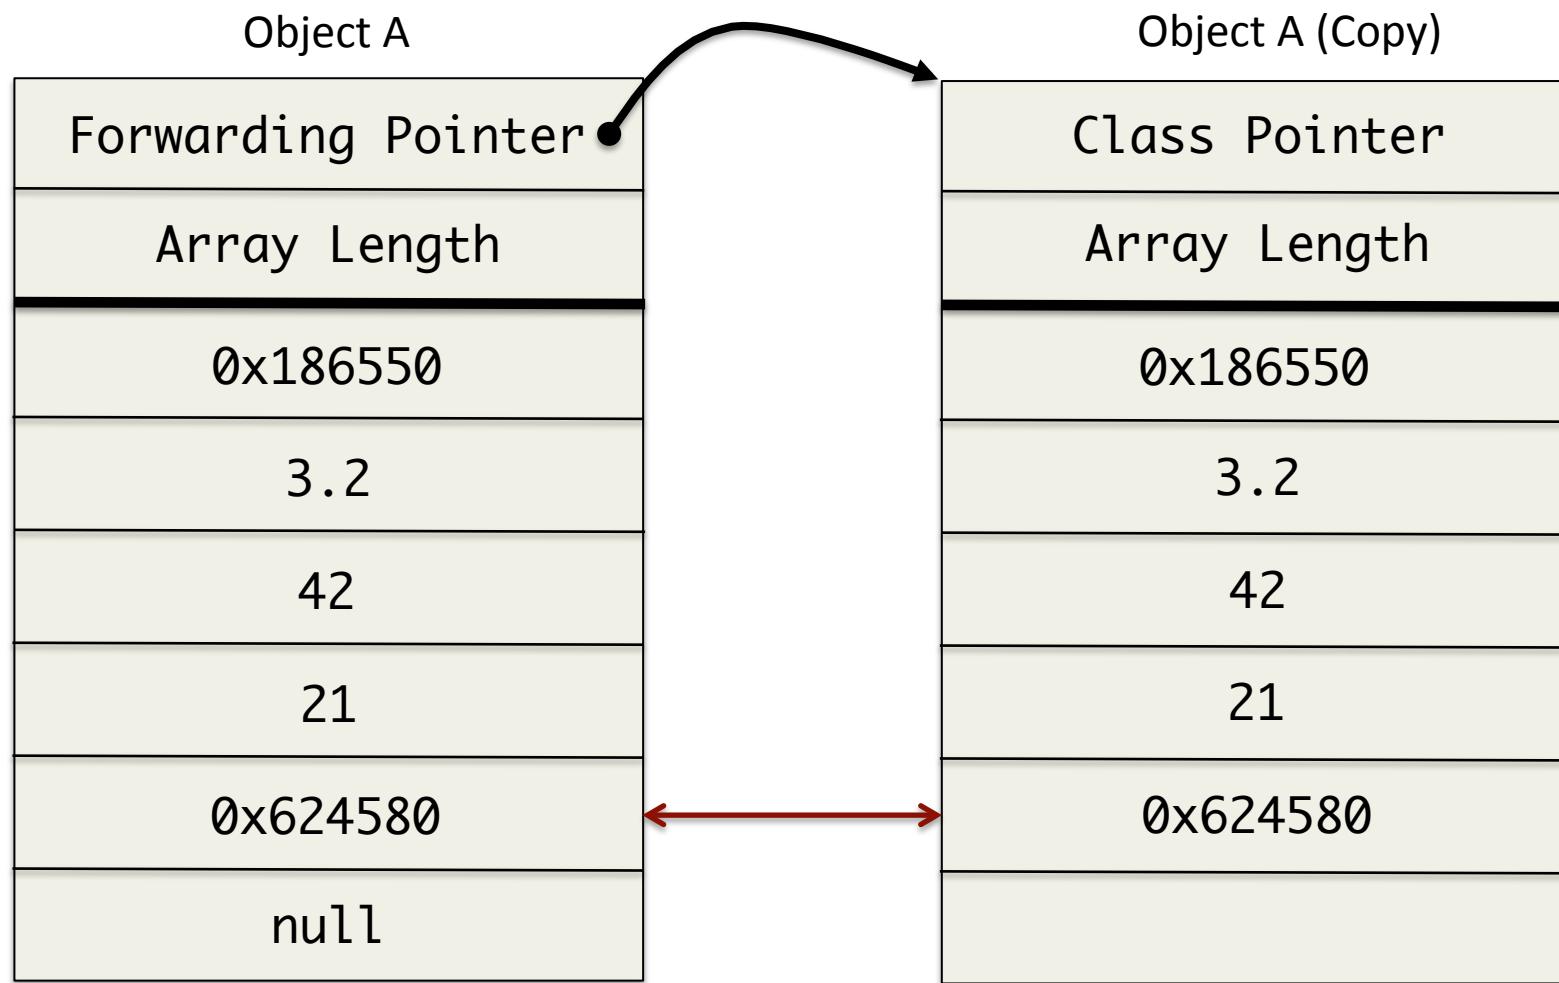


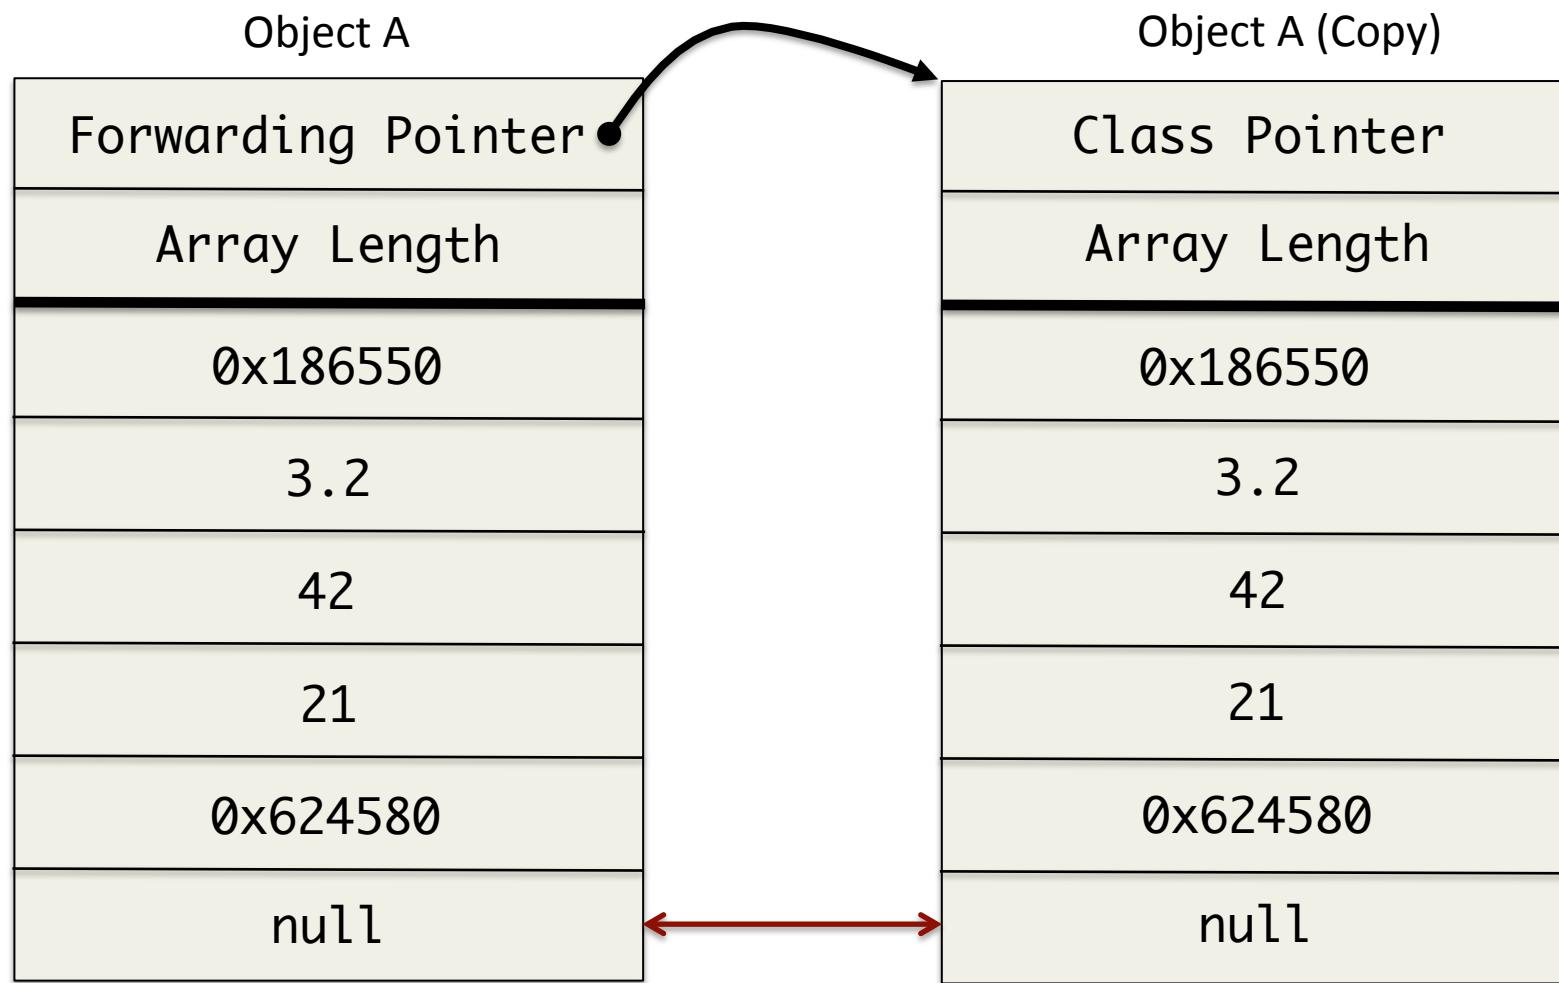












Sapphire Pros and Cons

- Enables lock-free concurrent copying
 - Synchronization between threads use atomics
 - Much lower overhead than locks
- Mutator ends up doing a lot more work
 - Follows forwarding pointers
 - Double writes
 - Complex write barrier

Concurrent Copying

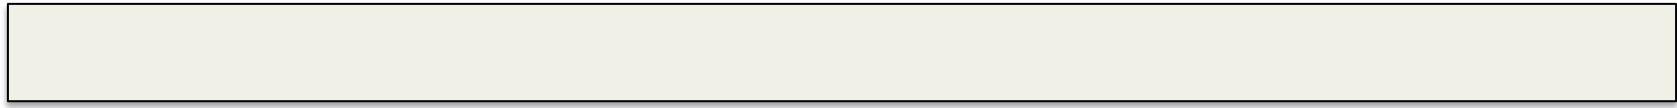
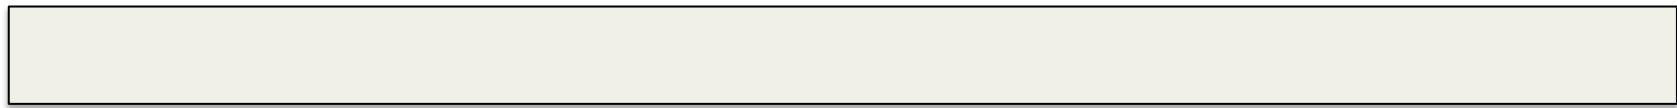
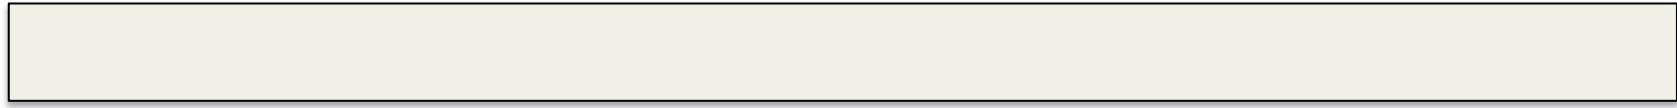
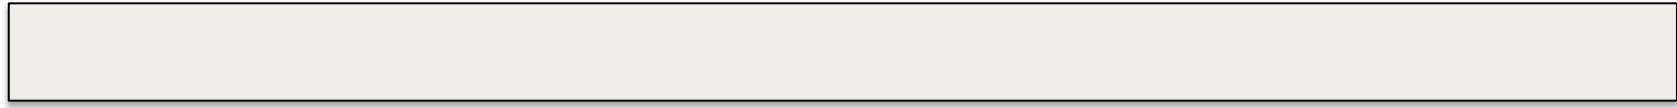
- Generally, the overheads aren't worth it
 - Increased write barrier complexity
 - More synchronization operations bad for the cache
- Most algorithms stop the world to copy
 - Limit each pause to a small region of the heap
 - Target regions likely to have more garbage
- Copying can be highly parallelized

Garbage First Algorithm

- Latest iteration of the Hotspot collector
 - Not enabled by default
 - Targets pause times over throughput
 - Many server applications want the inverse
- GC is parallel and concurrent
 - Copies objects to compact the heap
 - Pauses are predictable and tunable
 - Offers soft real-time without guarantees

G1 Regions

- Heap is split into equally-sized regions
 - Region sizes range from 1-32 Mb
 - Aim to have no more than 2048 regions
 - More tuning to enable more than 64 Gb heap
- Young generation made up of two block types
 - Nursery (or Eden)
 - Survivor set
- Very large objects given a region to themselves
 - Multiple regions can coalesce if necessary



--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

Nursery	Nursery						
---------	---------	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

							Nursery
--	--	--	--	--	--	--	---------

	Nursery		Nursery				
--	---------	--	---------	--	--	--	--

Nursery							
---------	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

Nursery	Nursery		Survivor				
---------	---------	--	----------	--	--	--	--

--	--	--	--	--	--	--	--

				Survivor			
--	--	--	--	----------	--	--	--

		Survivor					Nursery
--	--	----------	--	--	--	--	---------

	Nursery		Nursery				
--	---------	--	---------	--	--	--	--

Nursery							
---------	--	--	--	--	--	--	--

		Survivor			Survivor		
--	--	----------	--	--	----------	--	--

							Survivor
--	--	--	--	--	--	--	----------

Nursery	Nursery		Survivor				
---------	---------	--	----------	--	--	--	--

	Mature	Mature			Mature		
--	--------	--------	--	--	--------	--	--

Mature	Mature			Survivor		Mature	Mature
--------	--------	--	--	----------	--	--------	--------

		Survivor		Mature		Nursery	
--	--	----------	--	--------	--	---------	--

Mature	Nursery	Mature	Nursery				
--------	---------	--------	---------	--	--	--	--

Nursery	Mature			Mature		Mature	
---------	--------	--	--	--------	--	--------	--

		Survivor			Survivor		
--	--	----------	--	--	----------	--	--

	Mature						Survivor
--	--------	--	--	--	--	--	----------

Nursery	Nursery		Survivor	Humongous		
---------	---------	--	----------	-----------	--	--

	Mature	Mature			Mature	
--	--------	--------	--	--	--------	--

Mature	Mature			Survivor		Mature	Mature
--------	--------	--	--	----------	--	--------	--------

		Survivor		Mature		Nursery	
--	--	----------	--	--------	--	---------	--

Mature	Nursery	Mature	Nursery				
--------	---------	--------	---------	--	--	--	--

Nursery	Mature			Mature		Mature	
---------	--------	--	--	--------	--	--------	--

		Survivor			Survivor		
--	--	----------	--	--	----------	--	--

	Mature		Humongous			Survivor	
--	--------	--	-----------	--	--	----------	--

Remembered Sets

- Maintain remembered sets for each region
 - Track references to that region from outside
 - Don't track internal references
- Maintain set using a write barrier
 - Use block alignment to generate region ID
 - Simple shift operation
- Remembered sets give estimate of live data

Minor GC

- Parallel stop-the-world copying collection
- Promote all live young objects
 - Move objects from nursery to the survivor regions
 - Move objects from the survivor to the mature
- Don't evacuate all young objects at once
 - Determine nursery regions by pause time
 - Determine survivor regions by garbage ratio estimates

Concurrent Marking

- Initial marking to scan the root set
 - Requires that threads be stopped
 - Piggy-backs on a minor GC pause
- Concurrent marking across the heap
 - Uses incremental update
- Remark phase to complete incremental update
 - Stops the world

Cleanup Phase

- After mark phase we know all live objects
 - Remembered sets aggregate the information
- We can explicitly target regions with more garbage
 - Minimize copy overhead
 - Reclaim more garbage in a bounded time
- In many cases a region is entirely garbage
 - Remember the train algorithm

Evacuation

- Final part of major collection
- Iterate over regions selected for copying
 - Copy all live objects to an empty region
 - Highly parallel
 - Uses a form of TLAB to minimize contention
- Number of regions copied can be tuned
 - Trade-off throughput for pause time