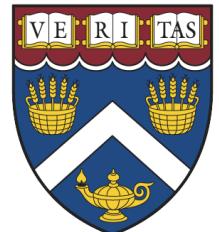


In the Last Week

- Some questions around the final project
 - What's going on with the GC footprint?
 - What are the parameters to `onPointerCreate`?
- Final exam logistics
 - Class tonight and next Monday
 - Final exam released early morning on the 15th
 - Exam due at midnight on the 19th
 - No late days for the final exam

Advanced Topics

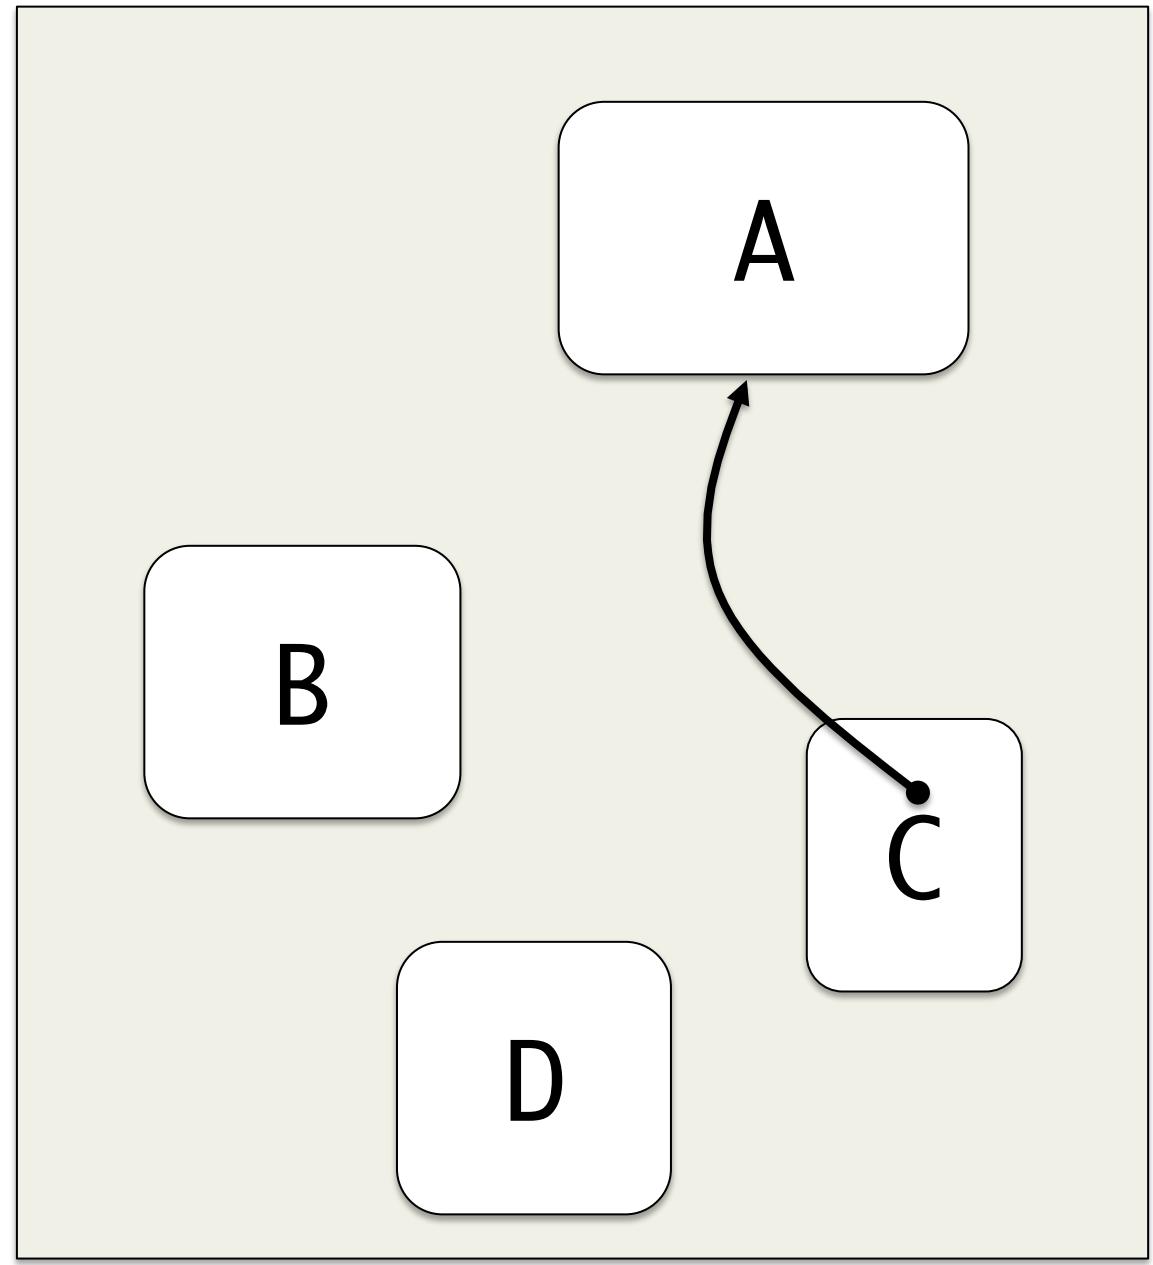


WEAK AND SOFT REFERENCES

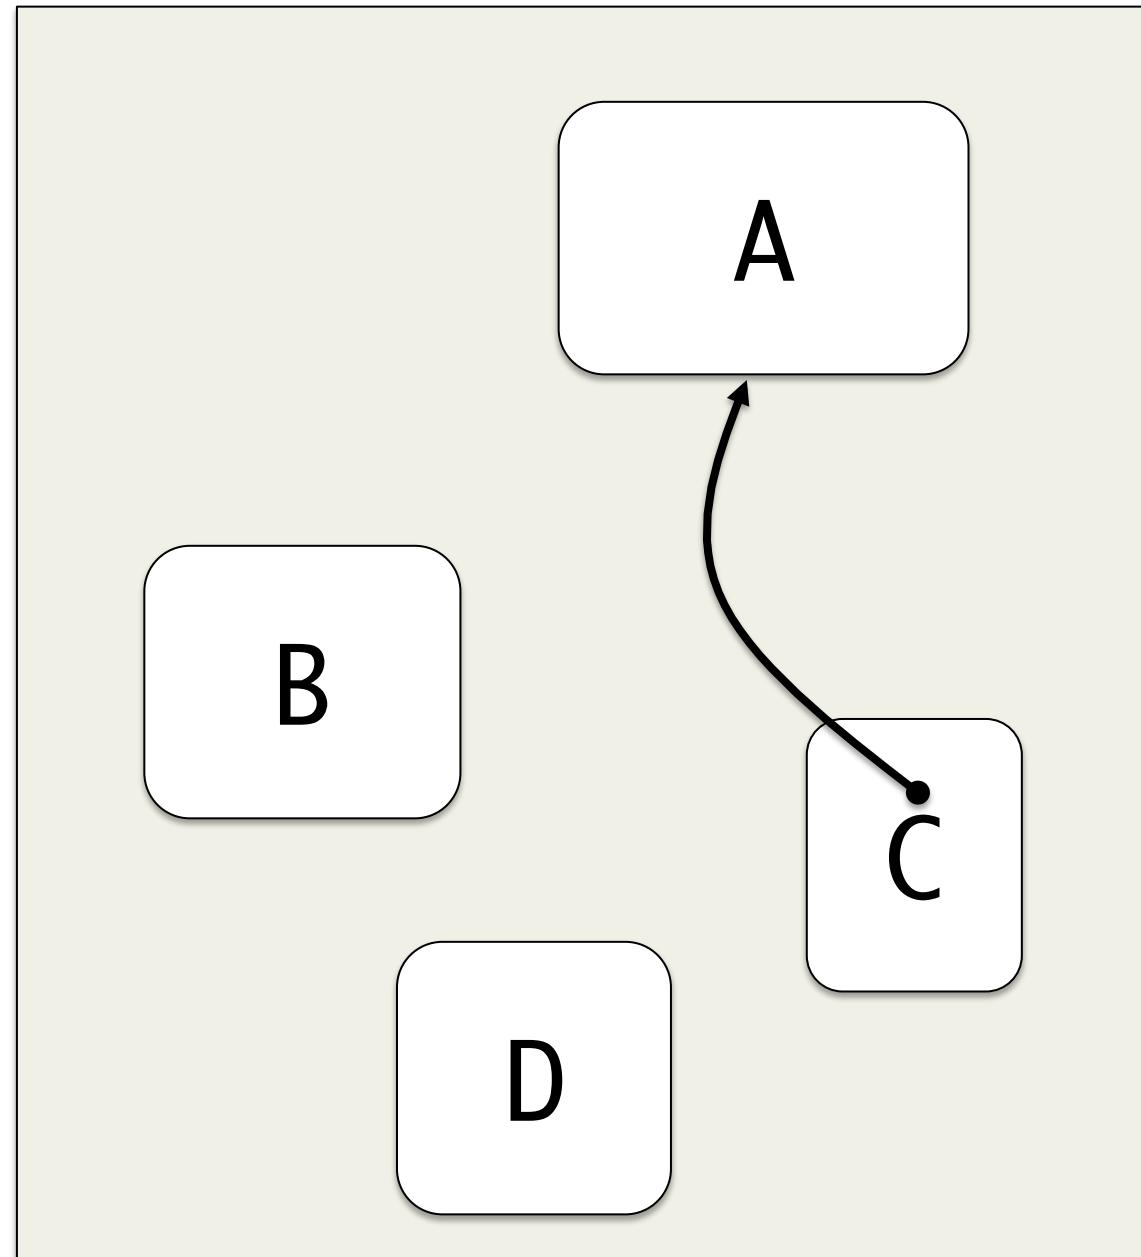
GC Reachability

- Basic tracing model is simple
 - Roots define an entry point to the object graph
 - Live objects are reachable from the roots
- This model can sometimes be inconvenient
 - Not all references are equal
- Weak references allow reachability without liveness
 - Objects referred to only by weak references are garbage

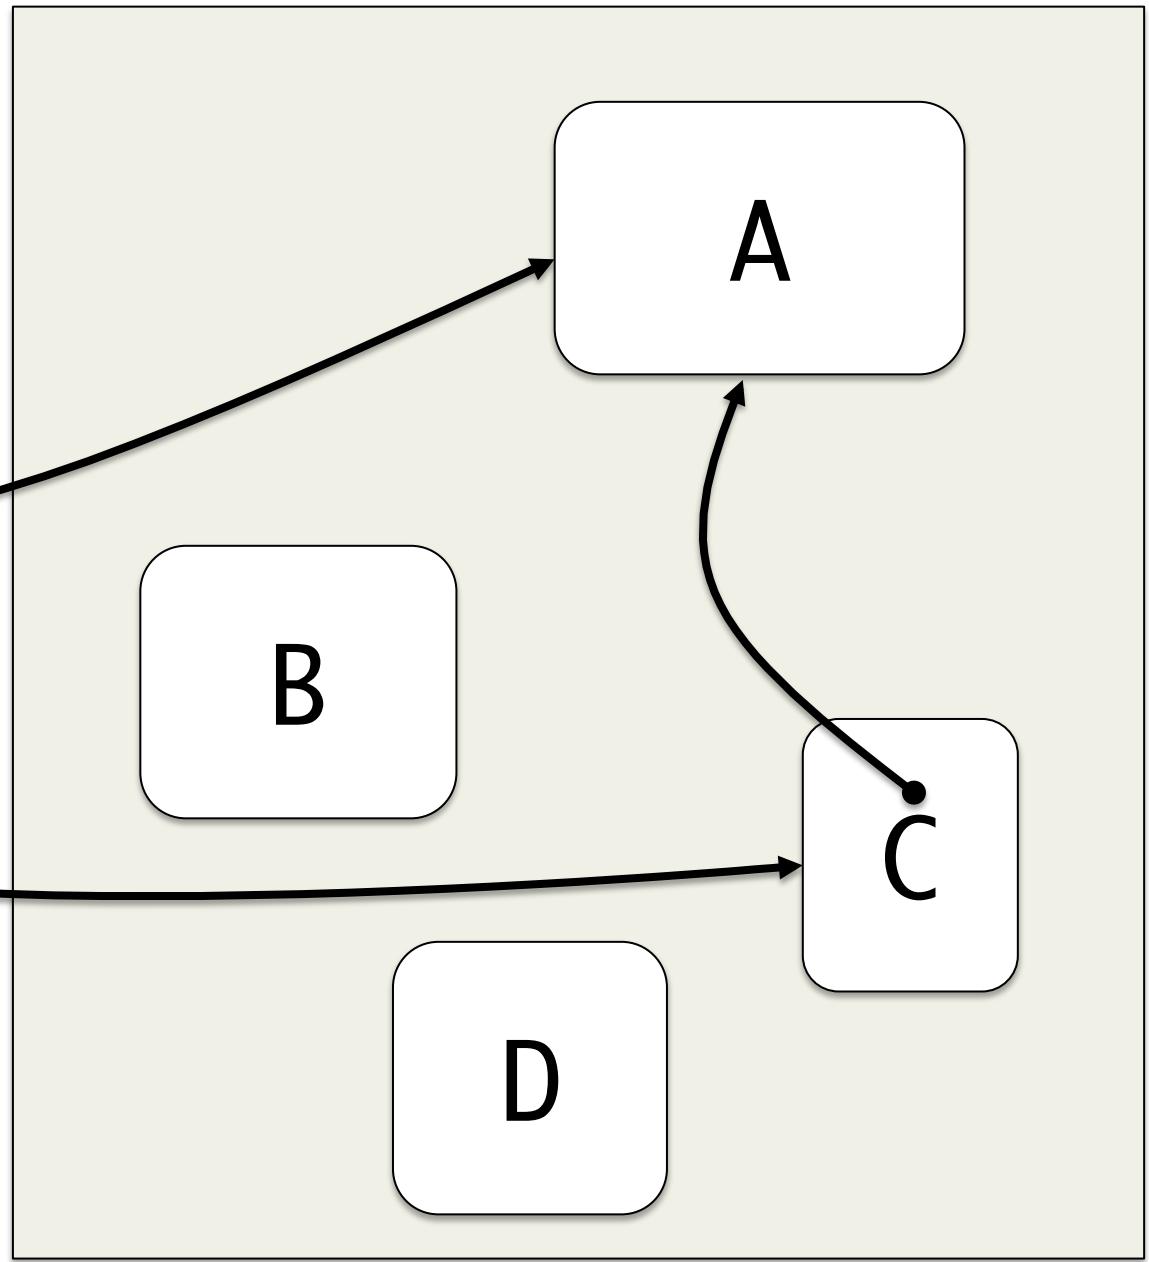
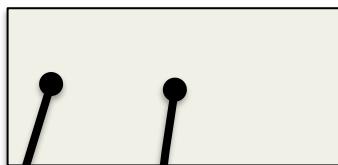


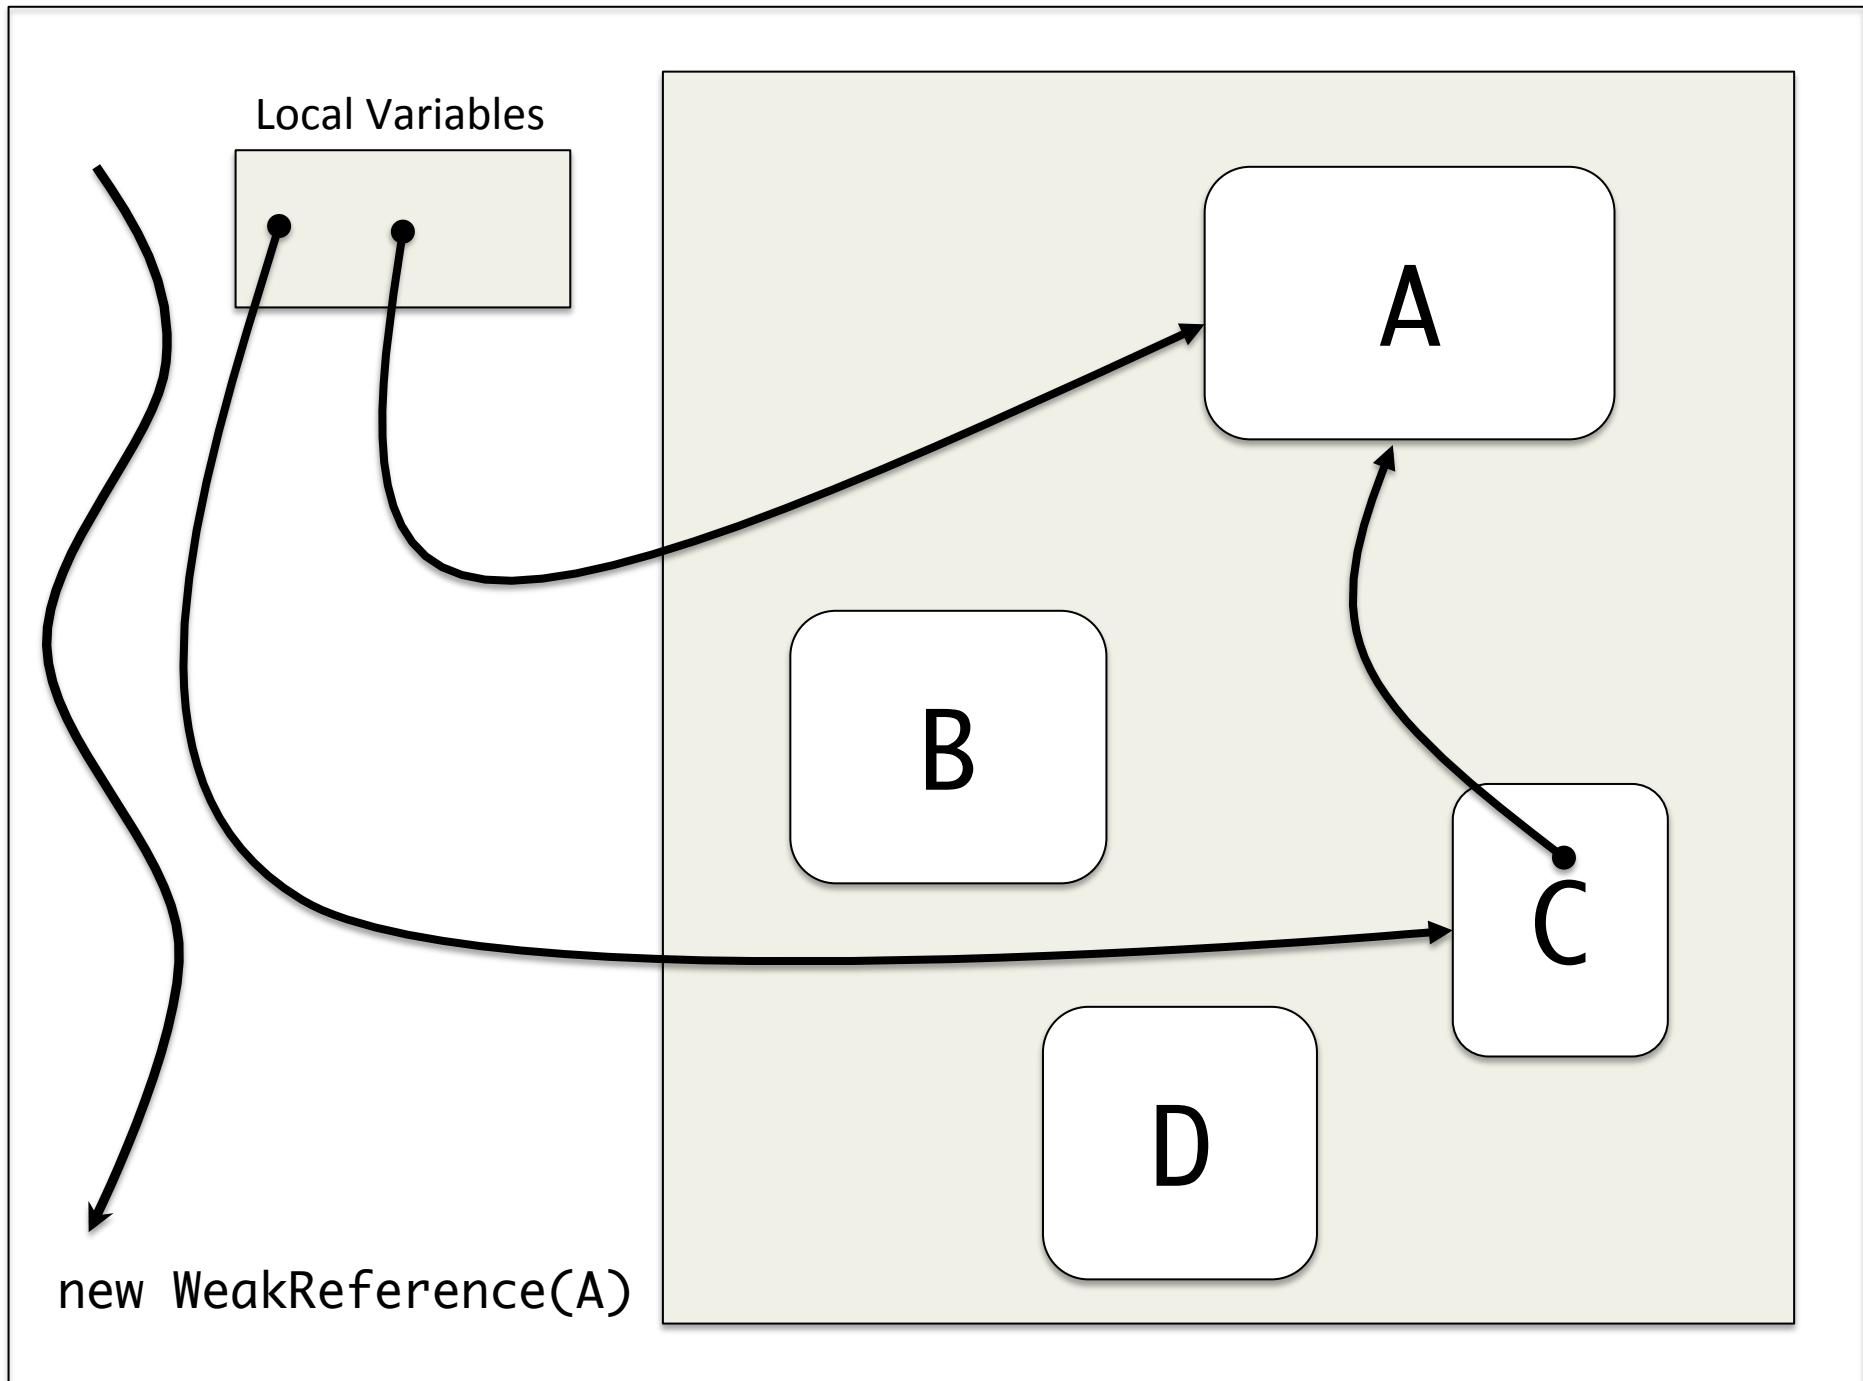


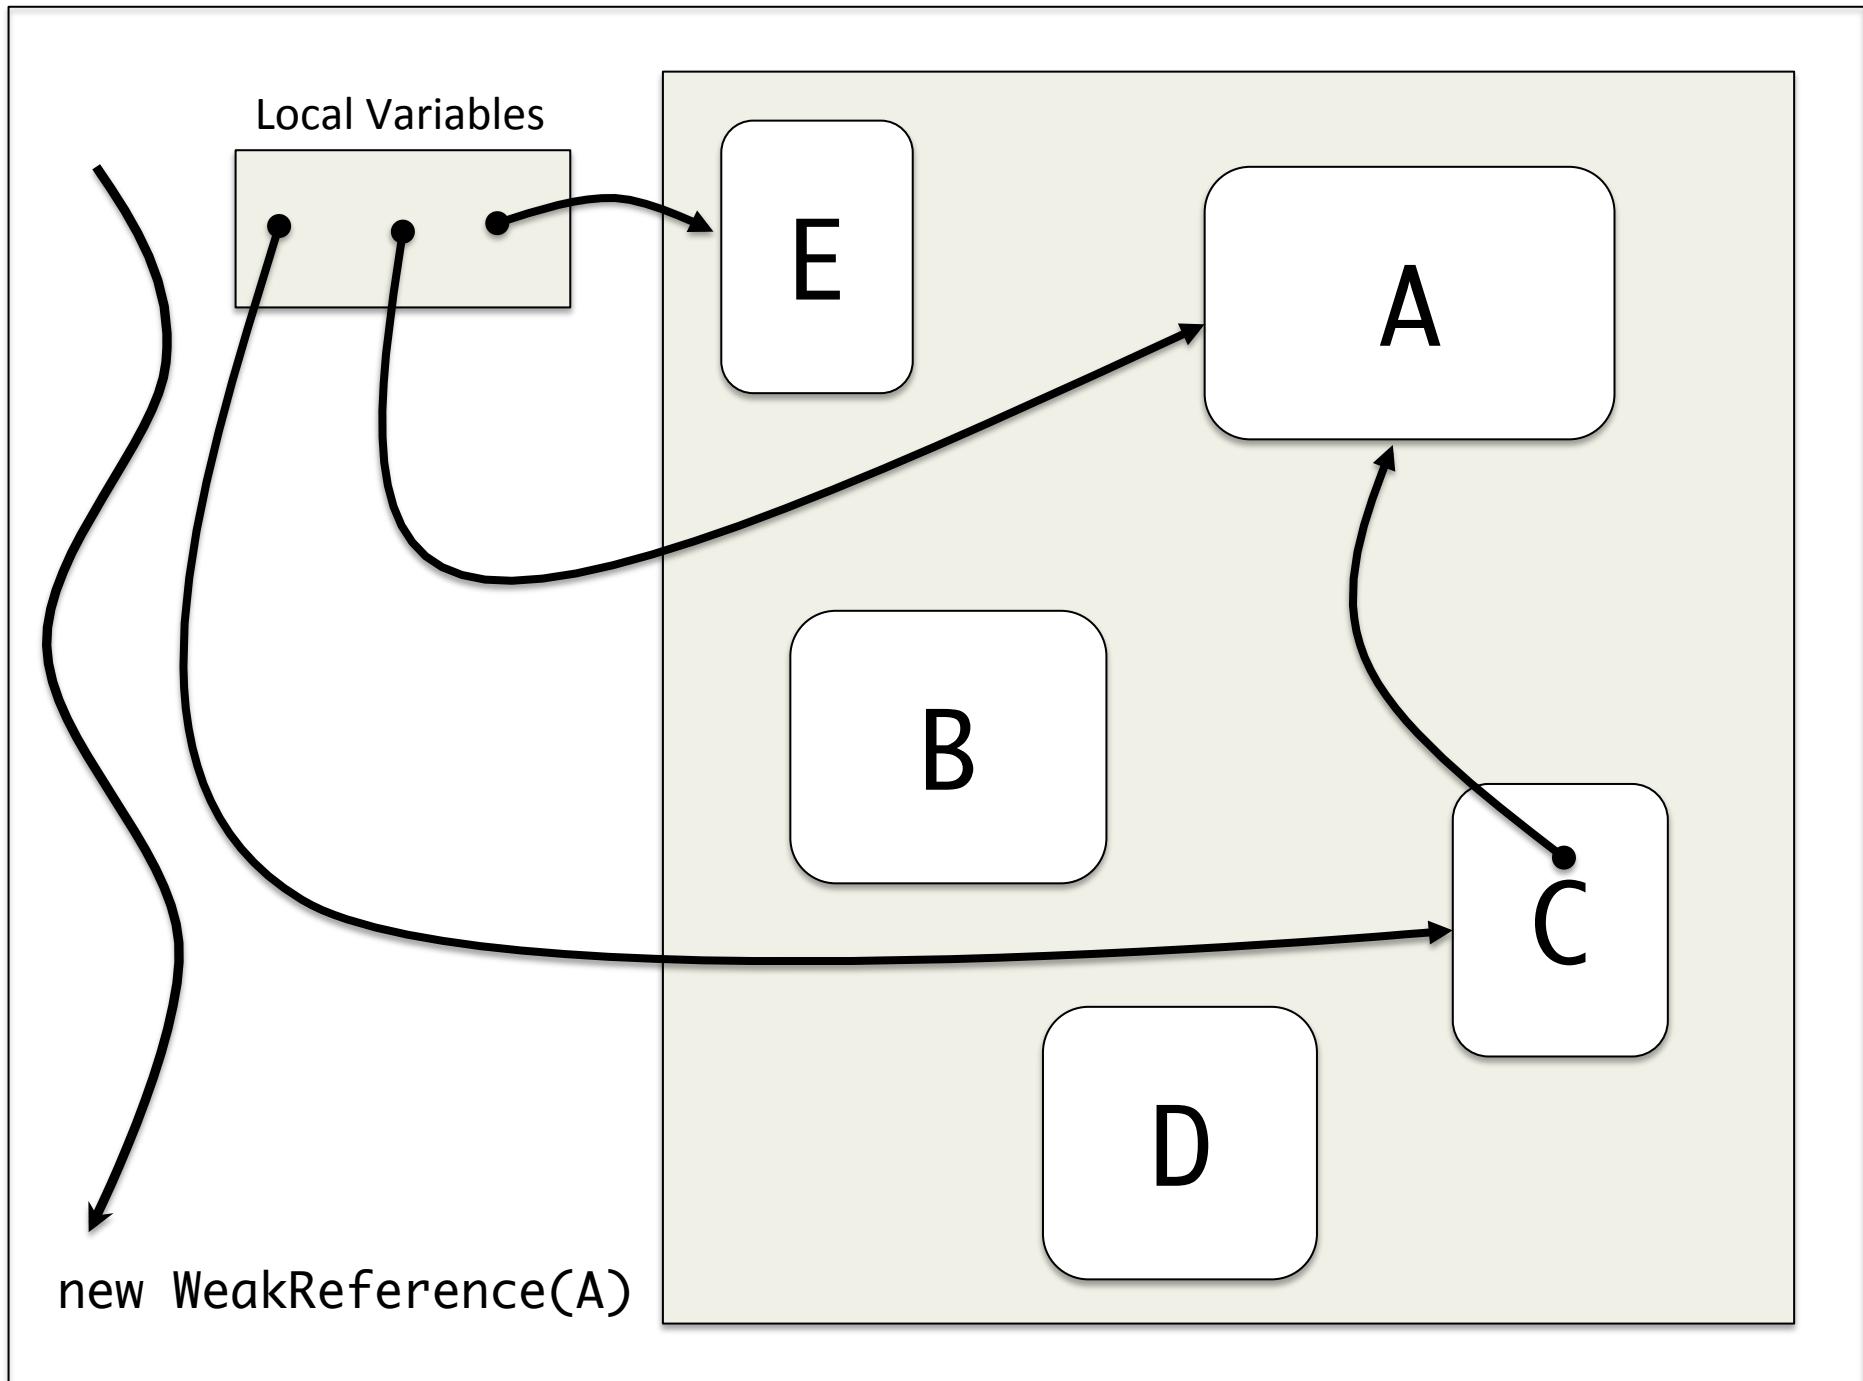
Local Variables

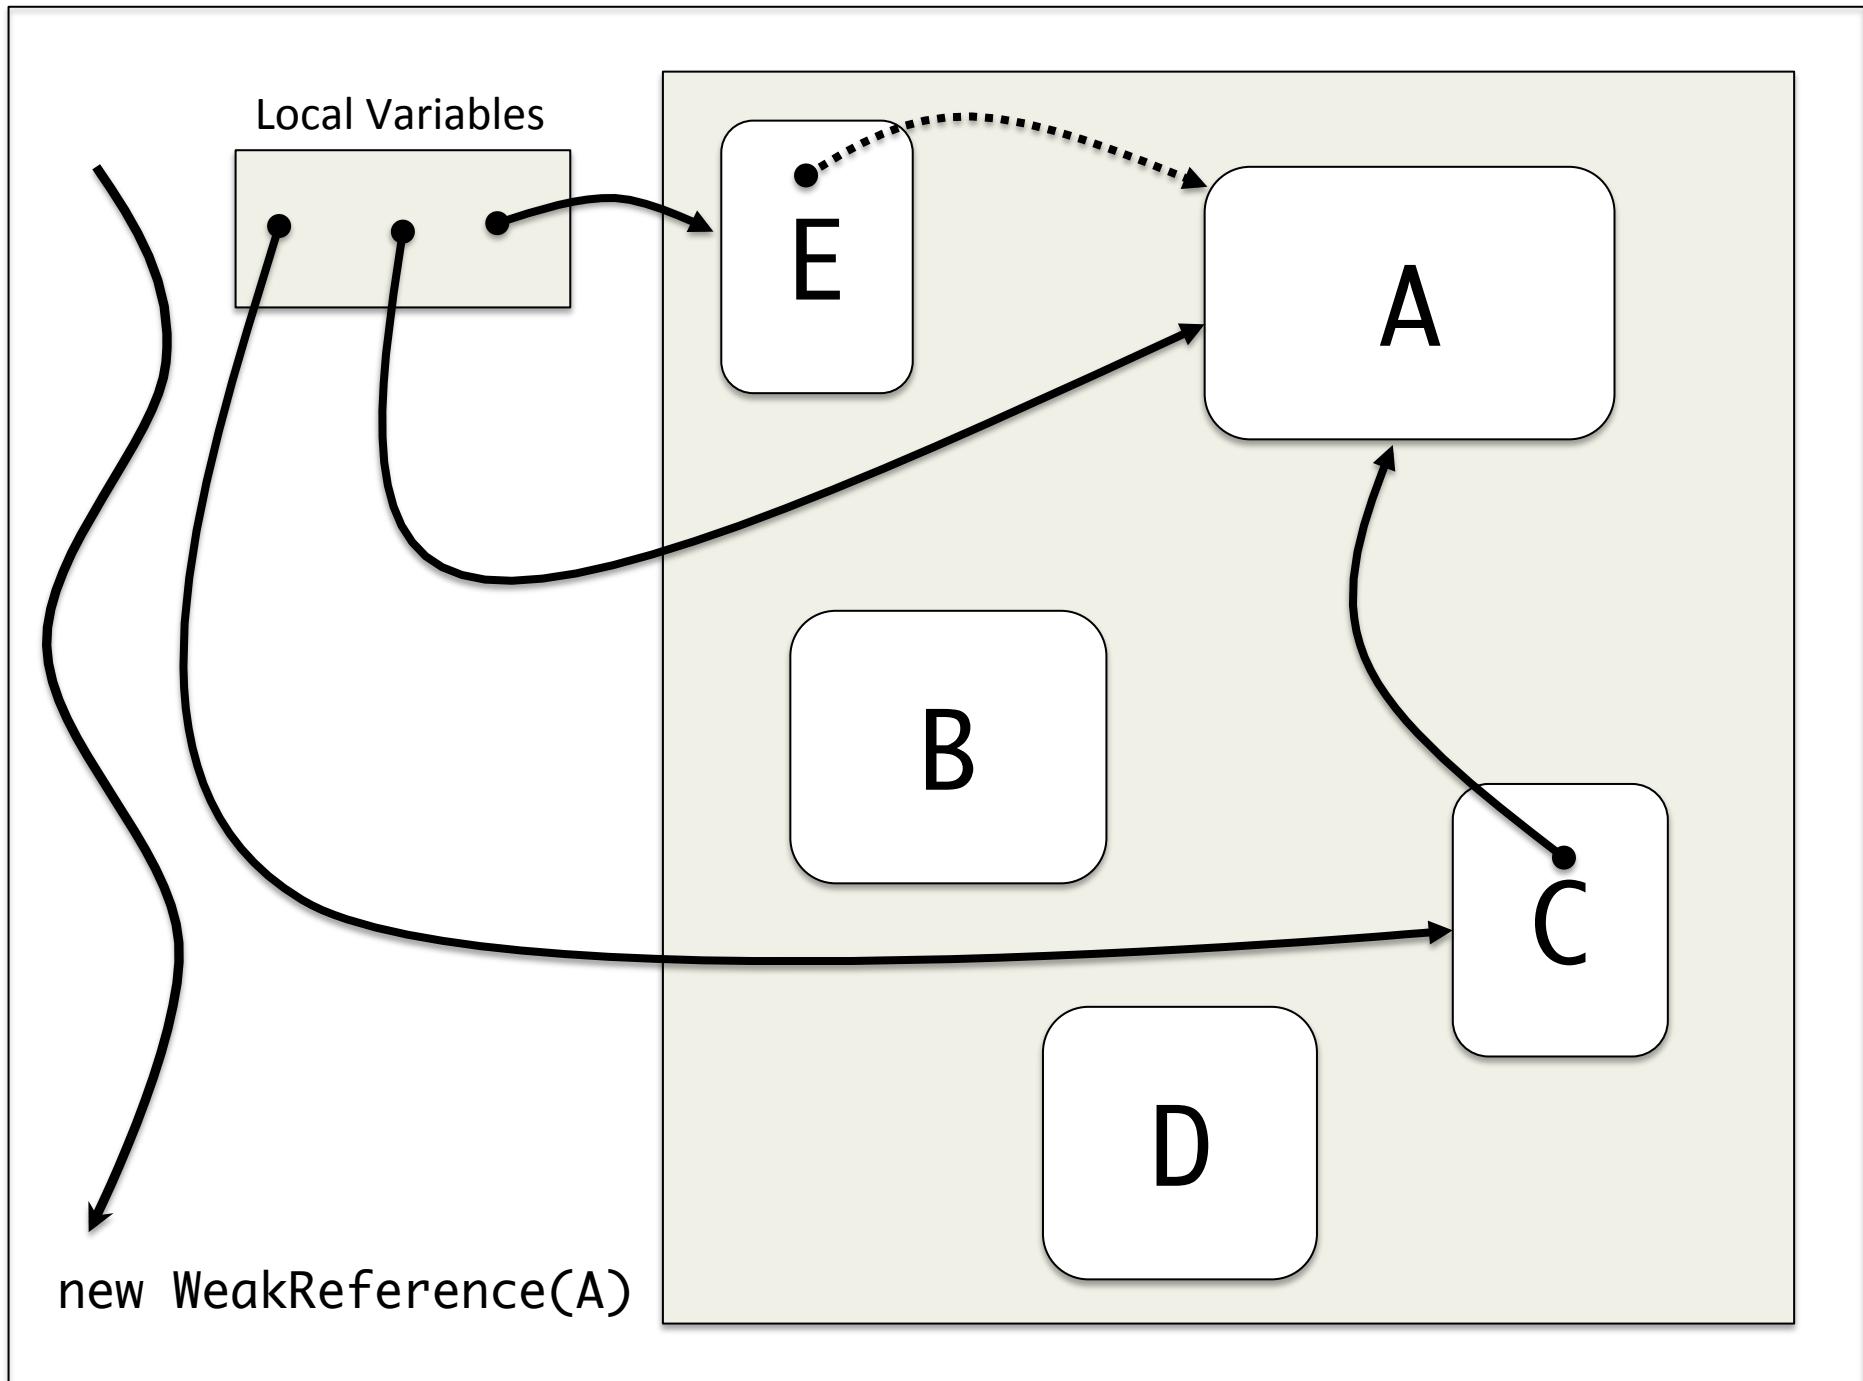


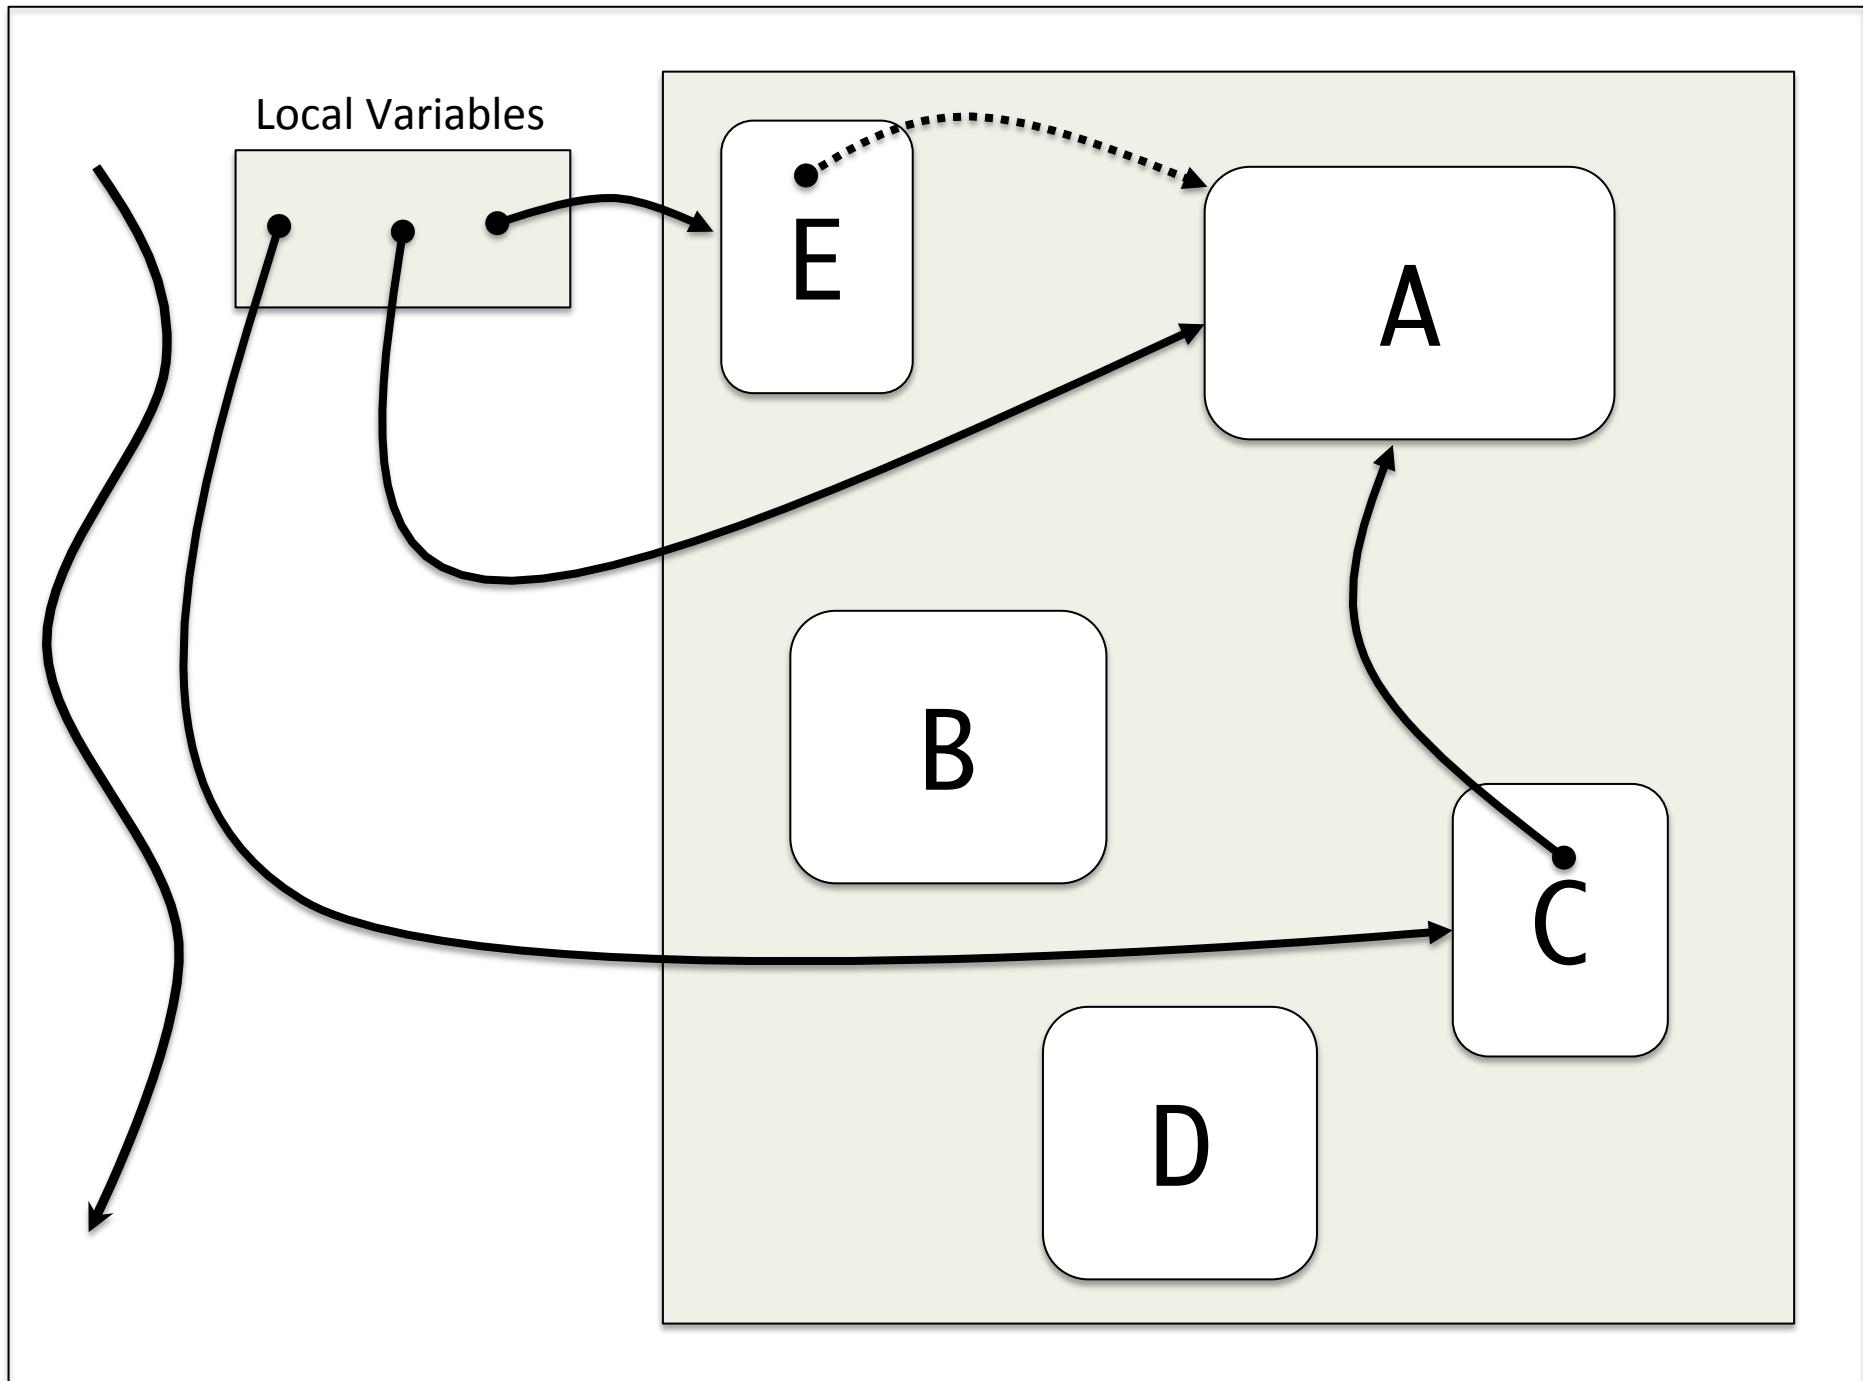
Local Variables

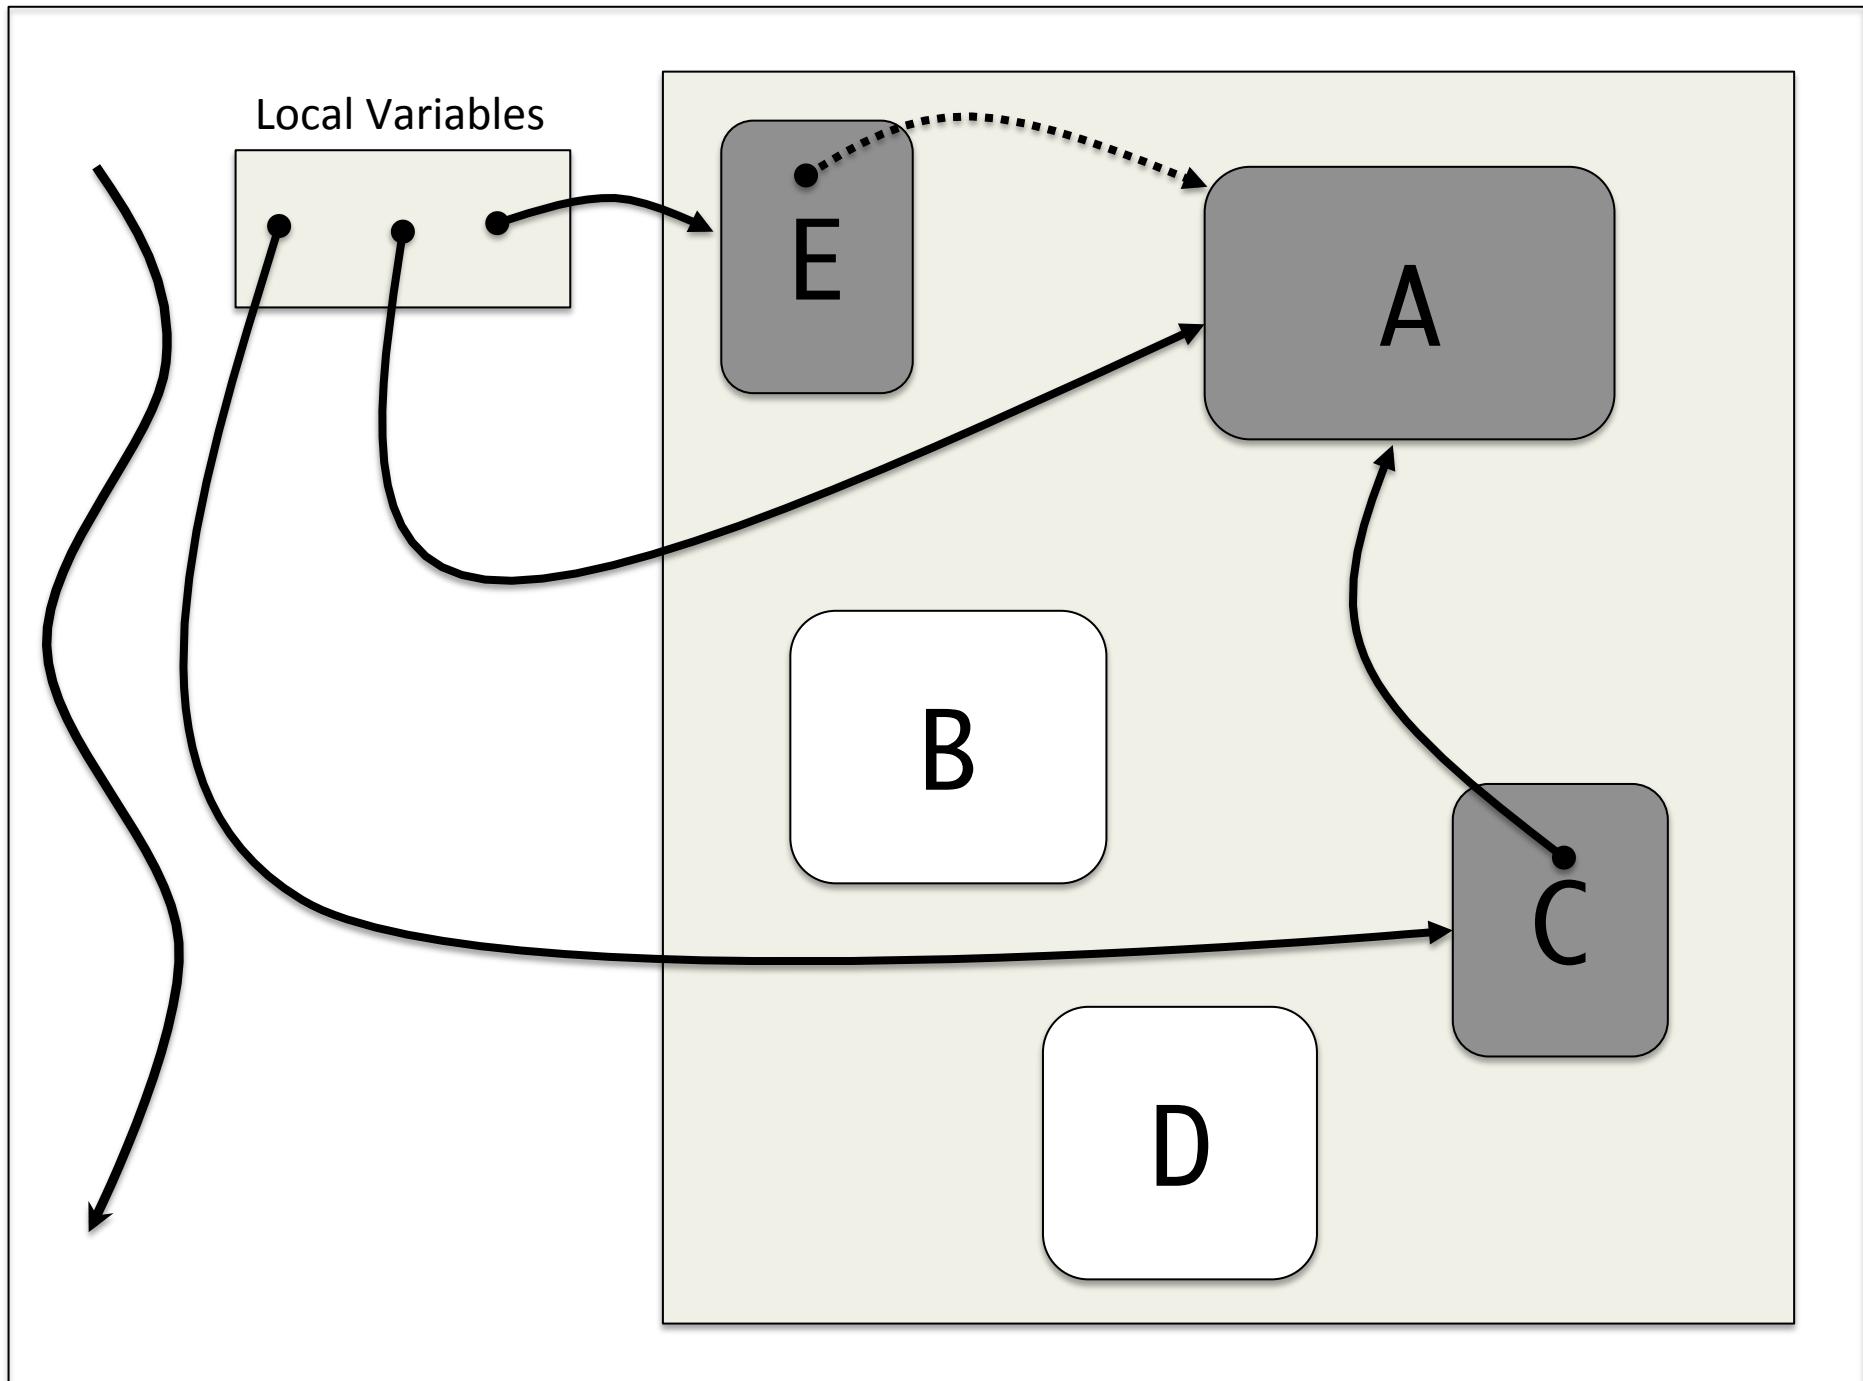


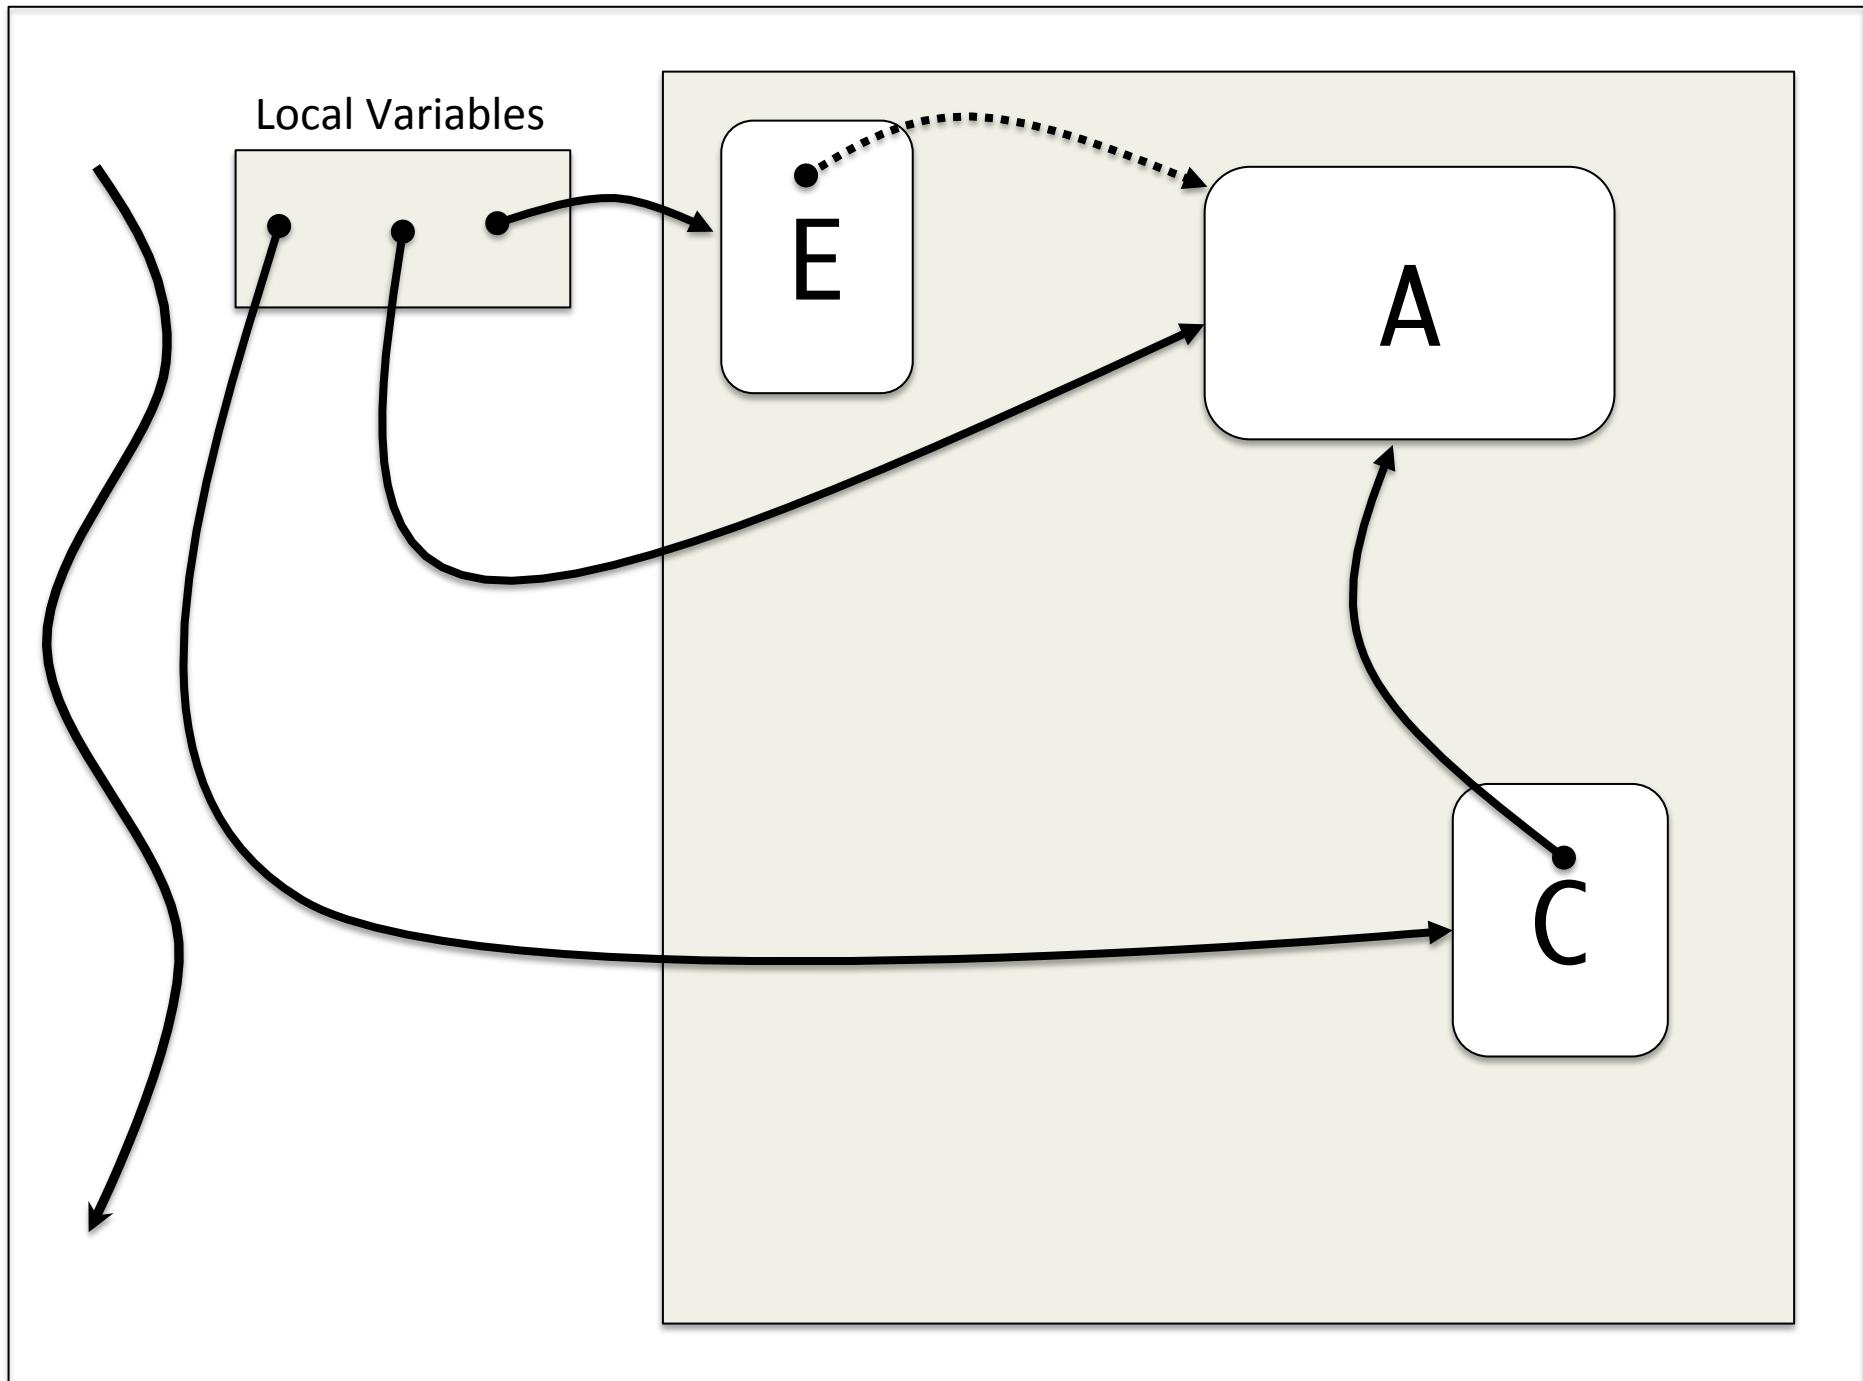


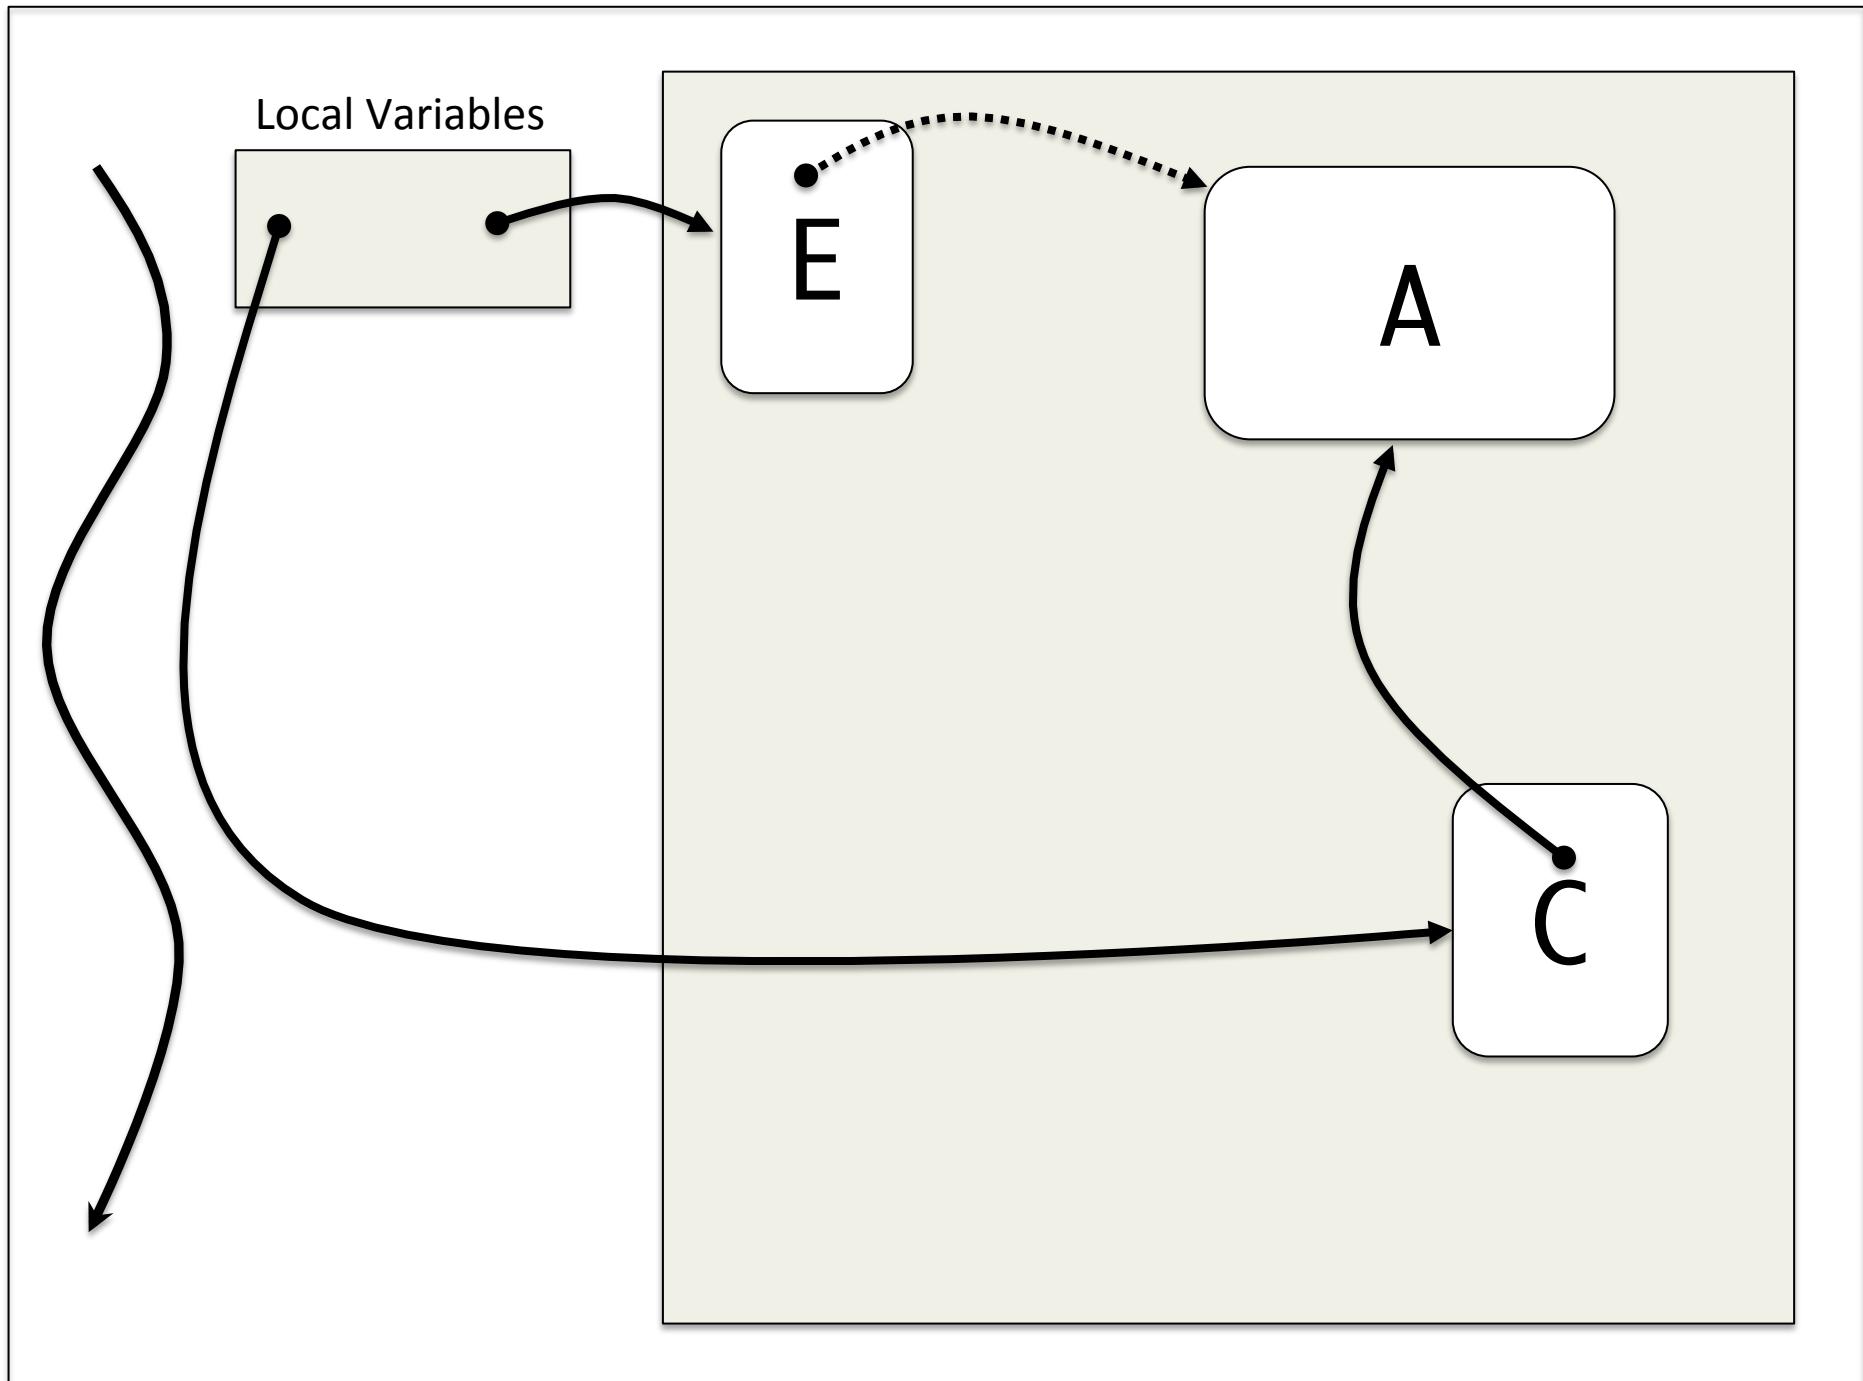


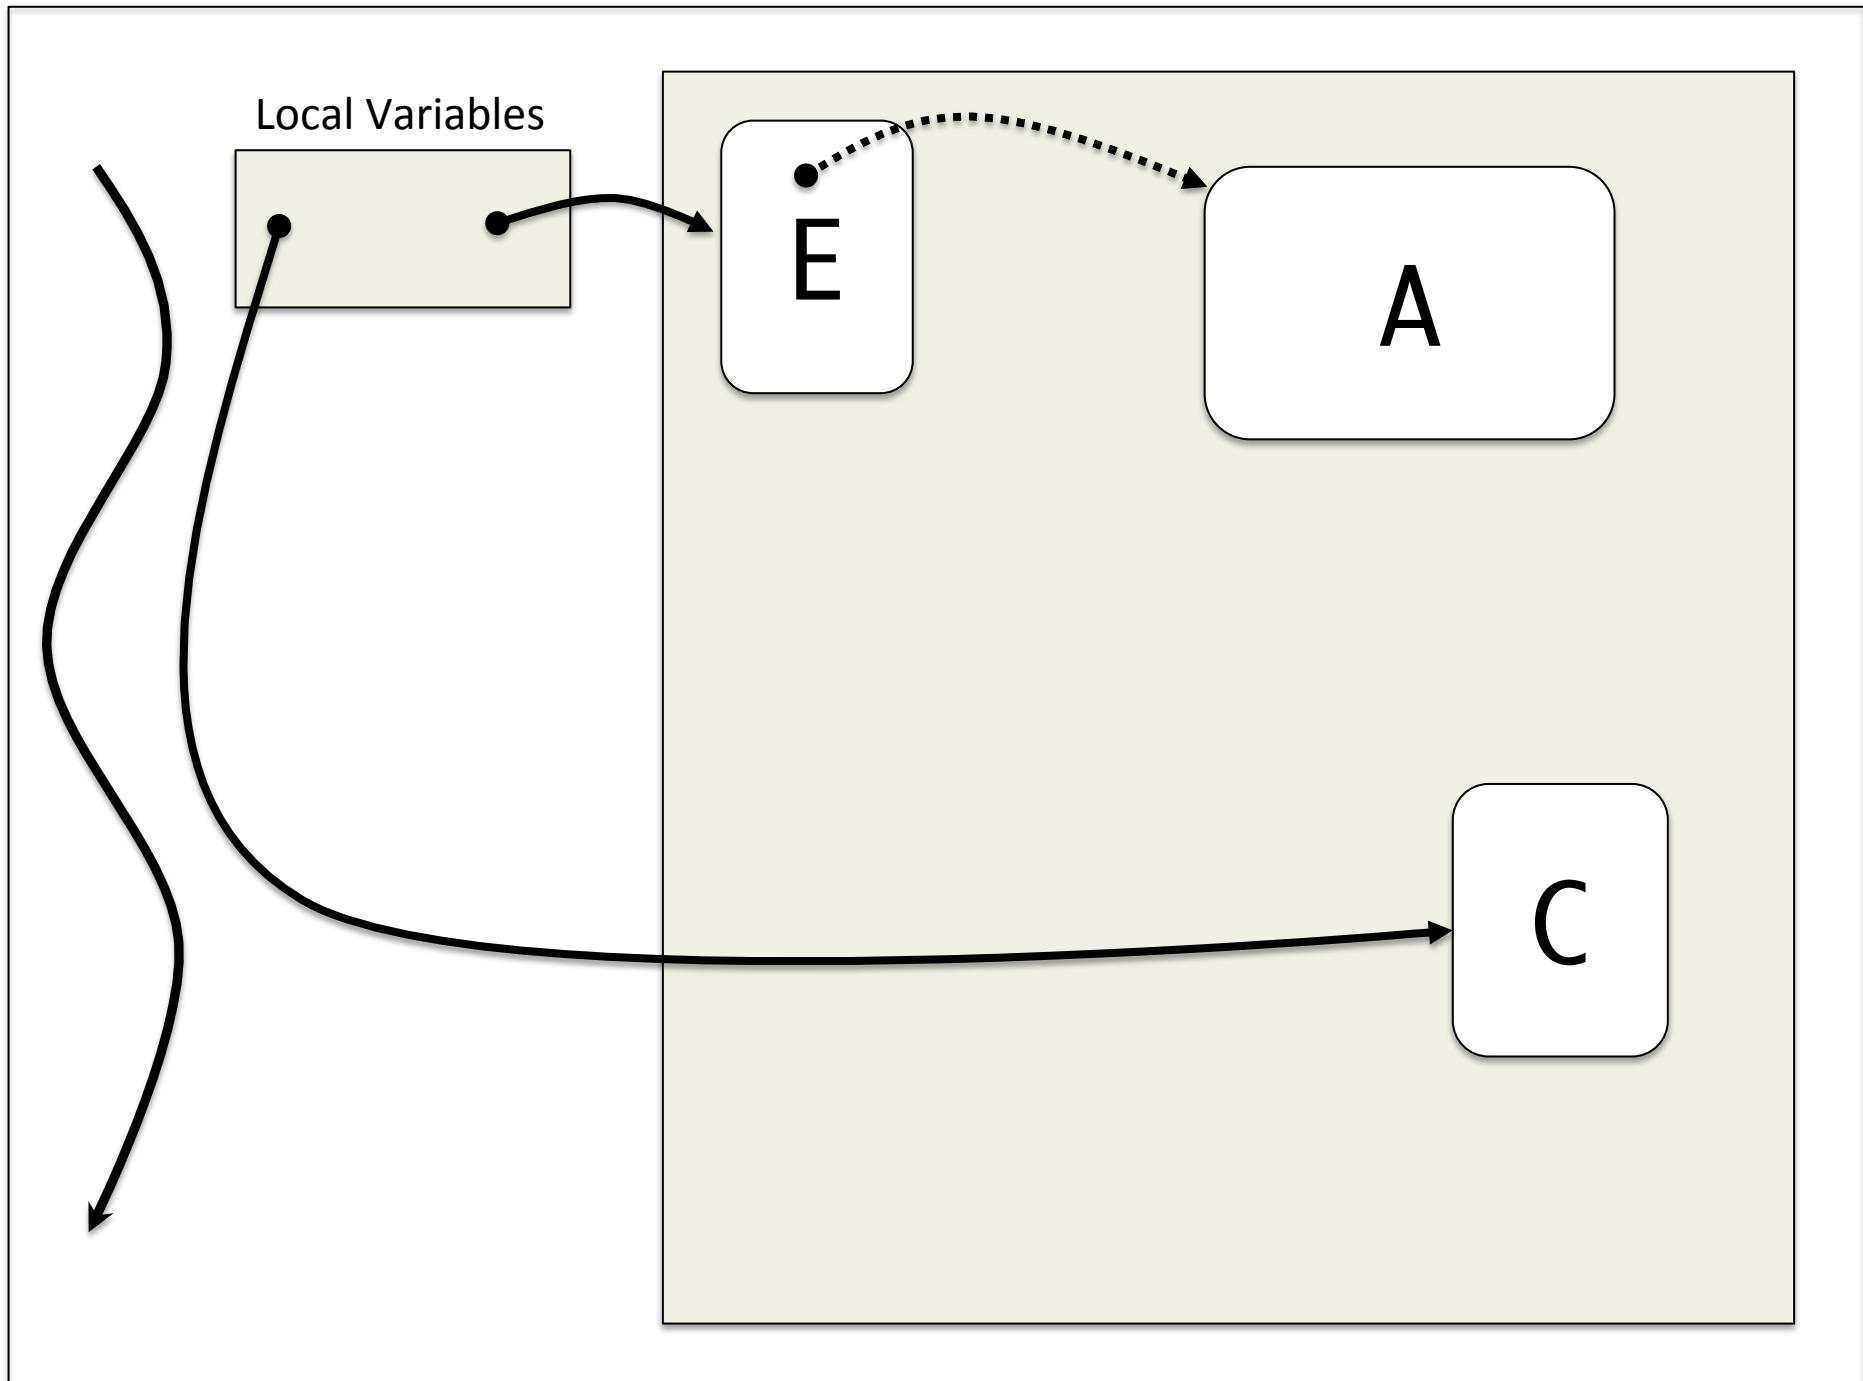


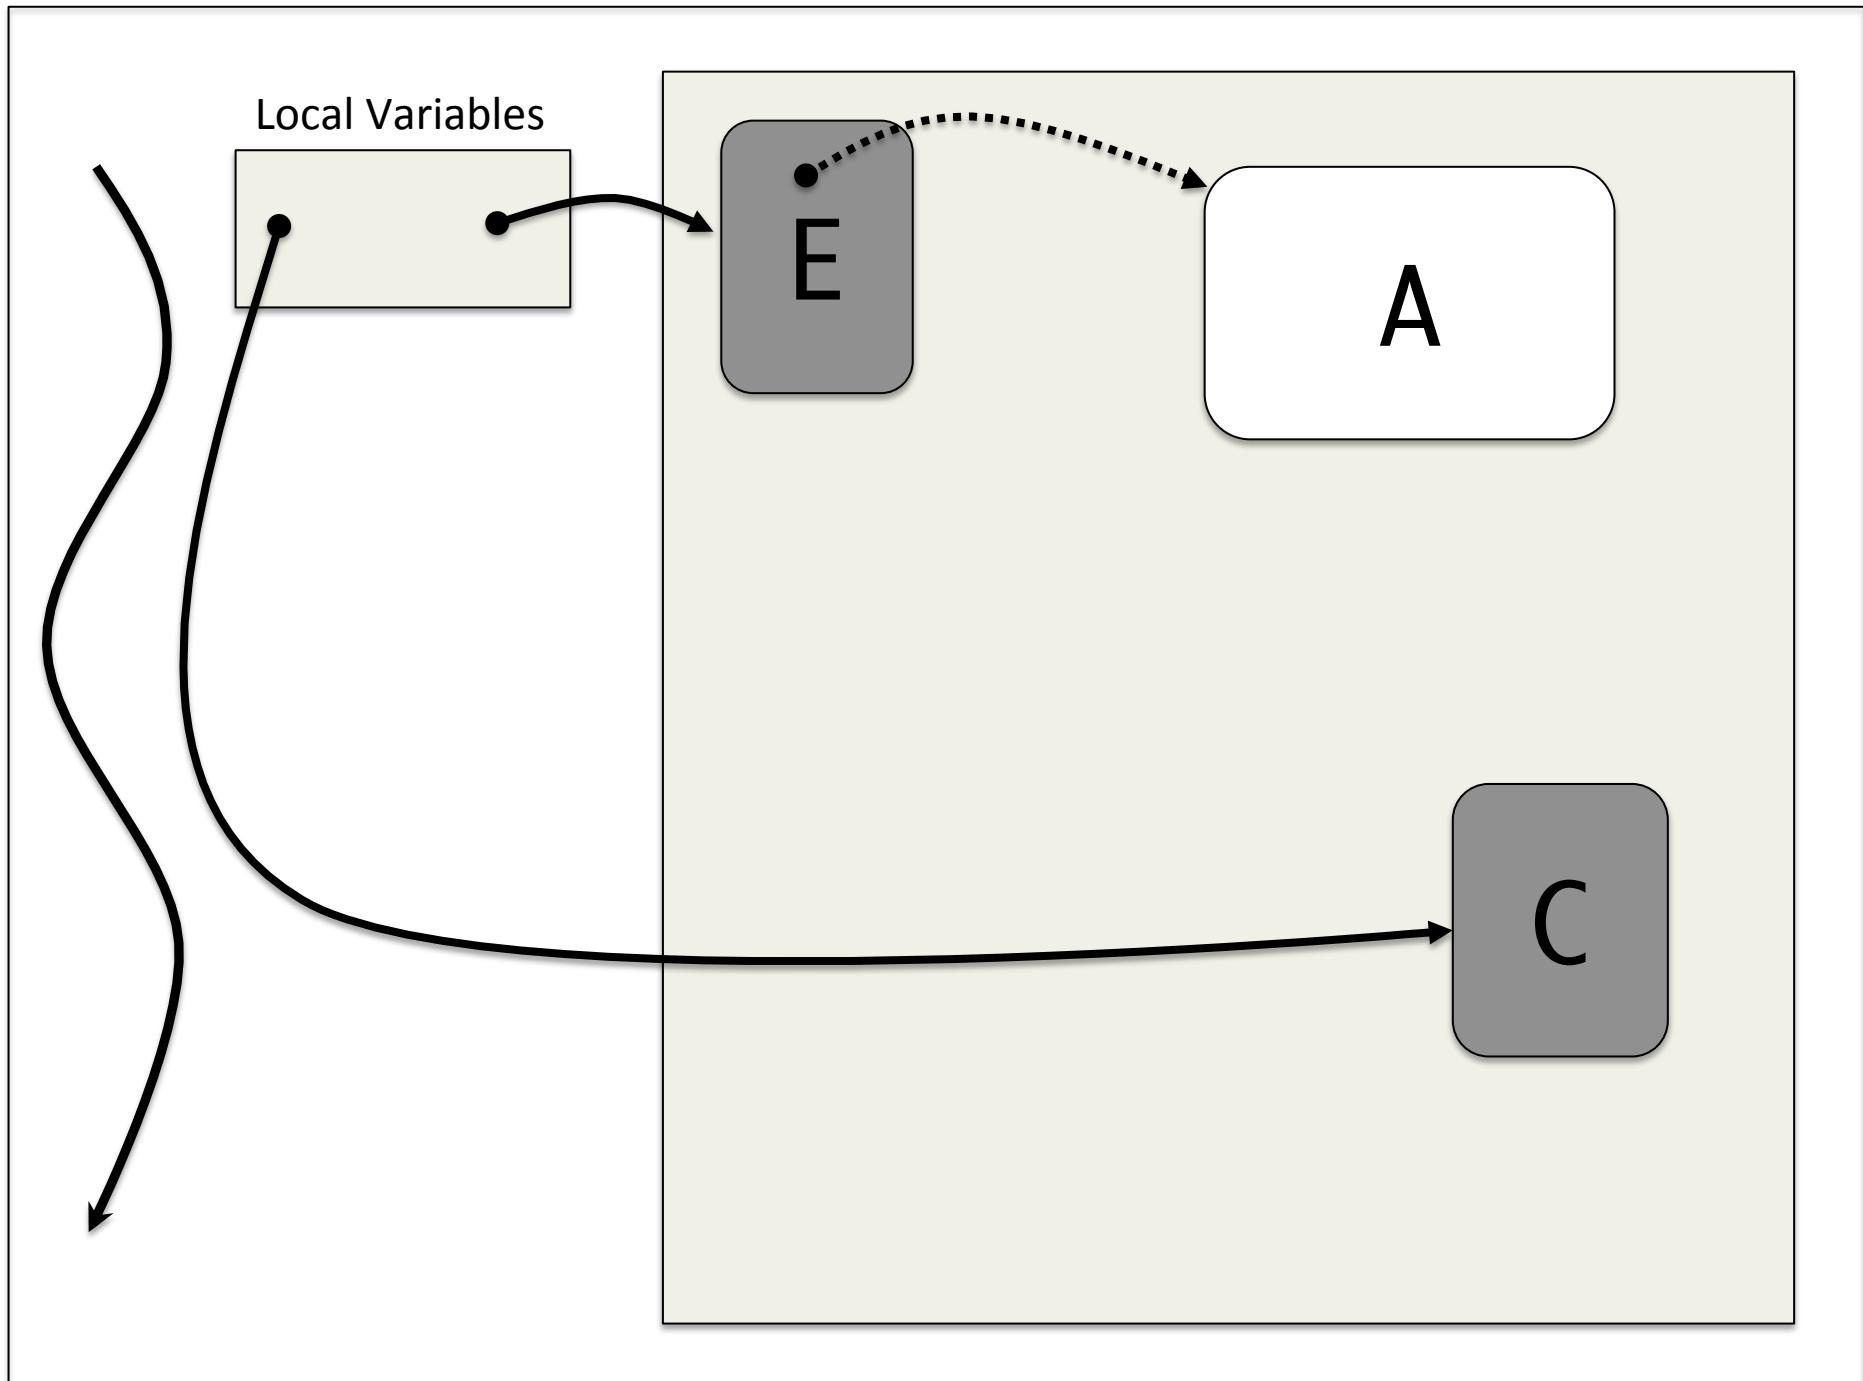


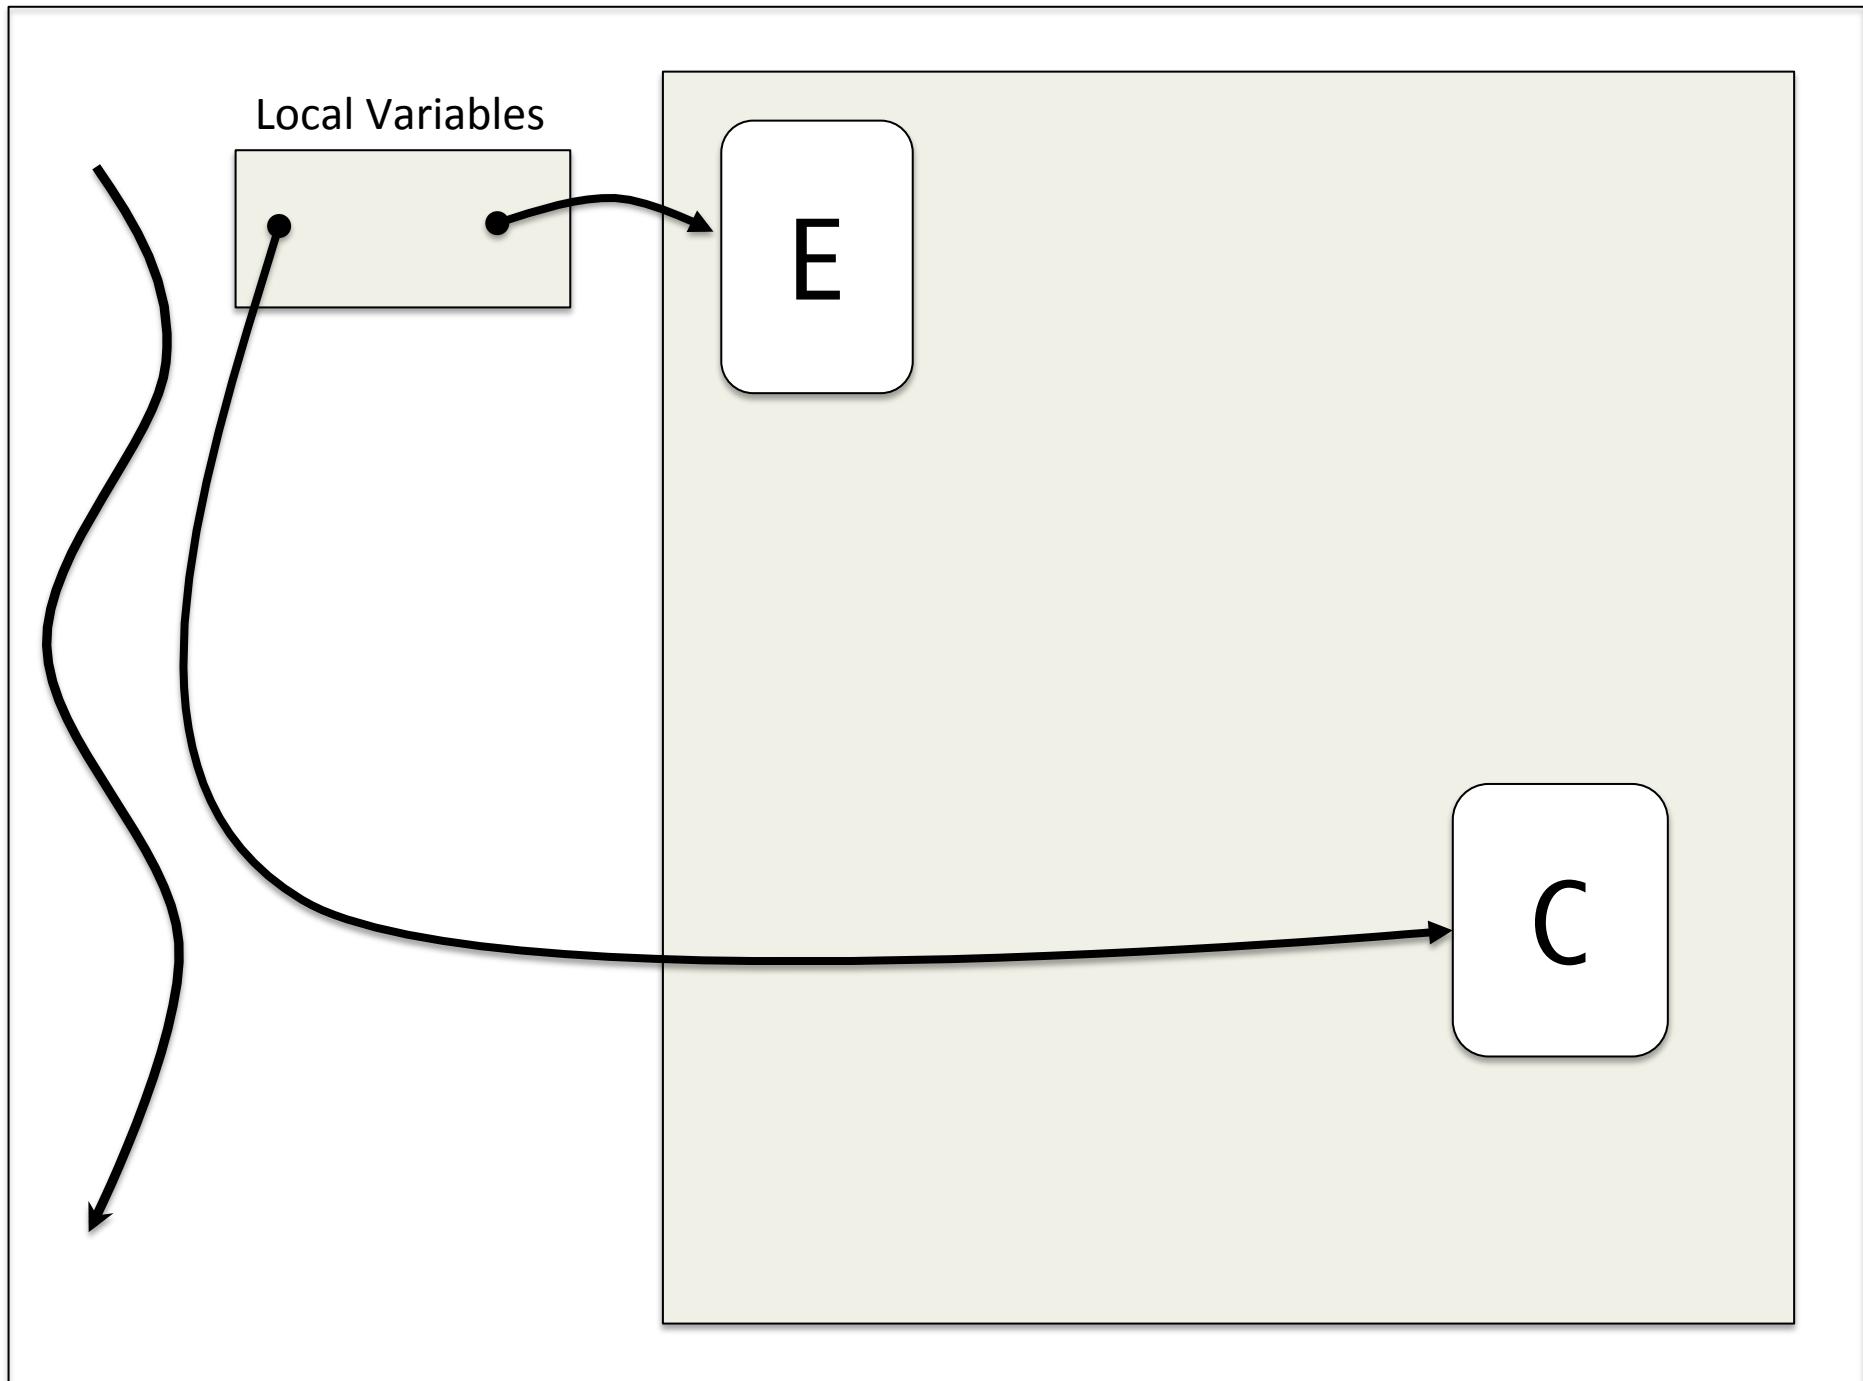












Weak Reference Uses

- Associative caches
 - Object is the key, value holds extra information
 - Key shouldn't keep an object alive
- Reference cycles
 - Long-standing issue with reference counting
 - Make one reference is a structure weak
- Observer pattern
 - Easily disassociate from event managers

Implementing Weak References

- WeakReference class has special semantics
 - Tracing algorithms treat it differently
 - Value may be changed by the garbage collector
- WeakReference is a primordial class
 - VM doesn't allow redefinition
 - Referent field is volatile
- GC must keep track of weak references
 - Need to nullify them if the object is collected

Reference Queues

- Gives visibility into object reachability
 - Weak reference is registered with a queue
- GC adds reference to the queue when nullifying
 - WeakReference object is placed on the queue
 - Referent is freed
- Developer can check for objects on queue
 - Knows which objects have been freed
 - Can take action on other related objects

Soft References

- Similar semantics to weak references
 - Soft reference does not keep object alive
 - Represented as a separate object in the Java heap
- Soft references are not cleared immediately
 - Softly-reachable object may survive multiple GCs
 - Free depends on memory pressure
- Guaranteed to be cleared before OutOfMemory

Using Soft References

- Large data structures that can be re-created
 - Result of complex computation
 - File, network or database resources
- Store structures in memory for fast access
 - Refer to cached structure using soft reference
 - Freed when the GC needs to reclaim the memory
- Program performance tied to GC algorithm

Freeing Soft References

- Free schedule is implementation specific
 - Can happen at any time from developer perspective
 - Spec guarantees that it is before memory runs out
- May want to release objects incrementally
 - GC tends to run slowly when memory is tight
 - Massive free operation leads to large pause
- VMs implement caching algorithms

Soft References in Hotspot

- SoftReference object has a timestamp
 - Set to the time of the last GC
 - Shows roughly when the reference was last accessed
- GC knows how much live data is in the heap
 - Figures out based on the mark phase
- Pre-set ratio of survival time to garbage ratio
 - Each GC has a dynamic limit on time since last use
 - Least recently used objects are freed first

Phantom References

- Third type of reference object in Java
 - Not commonly used
- Serves only as a notification of collection
 - Associate a PhantomReference with a ReferenceQueue
 - Test for the reference to be added to the queue
- Allows for complex clean-up behavior
 - Unpredictable

Reachability Hierarchy

- Strongly reachable
 - Can be reached by following normal references
- Softly reachable
 - Reachable only through soft references
- Weakly reachable
 - Reachable only through weak references
- Phantom reachable
 - Reachable only through phantom reference
- Unreachable
 - No references to the object

FINALIZATION

Destructors

- Destructors are a common concept
 - Code that runs when an object is freed
 - Useful for cleaning up resources
- Implementation in C++ is intuitive
 - Two types of storage – stack and heap
 - Stack objects freed when they leave scope
 - Heap objects have explicit free call
- Destructors run at well-understood times

Finalizers

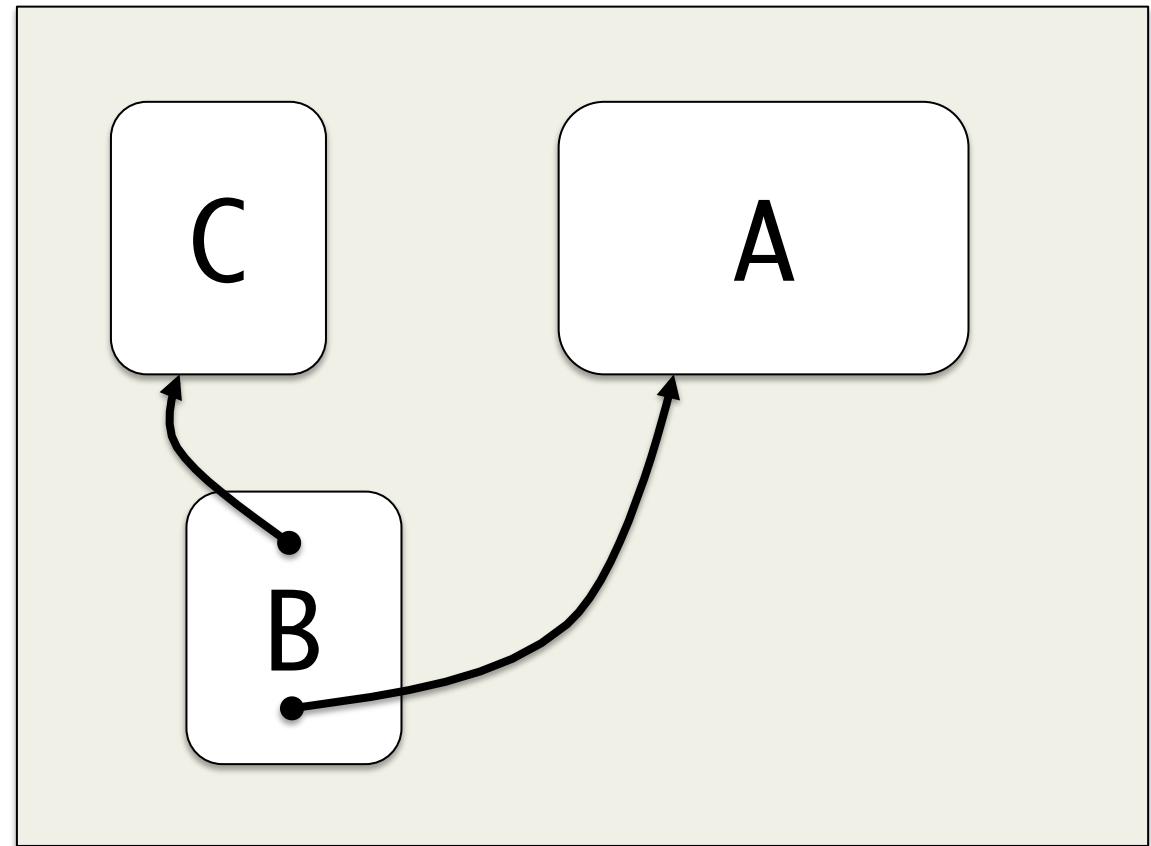
- Similar concept in managed languages
 - Any Java class can override `Object.finalize()` method
 - Called when an object is reclaimed by the GC
- Two key differences
 - Timing of finalizer is unpredictable
 - Finalization done in a separate thread or threads
- Finalize methods are harder to get right

Finalization Semantics

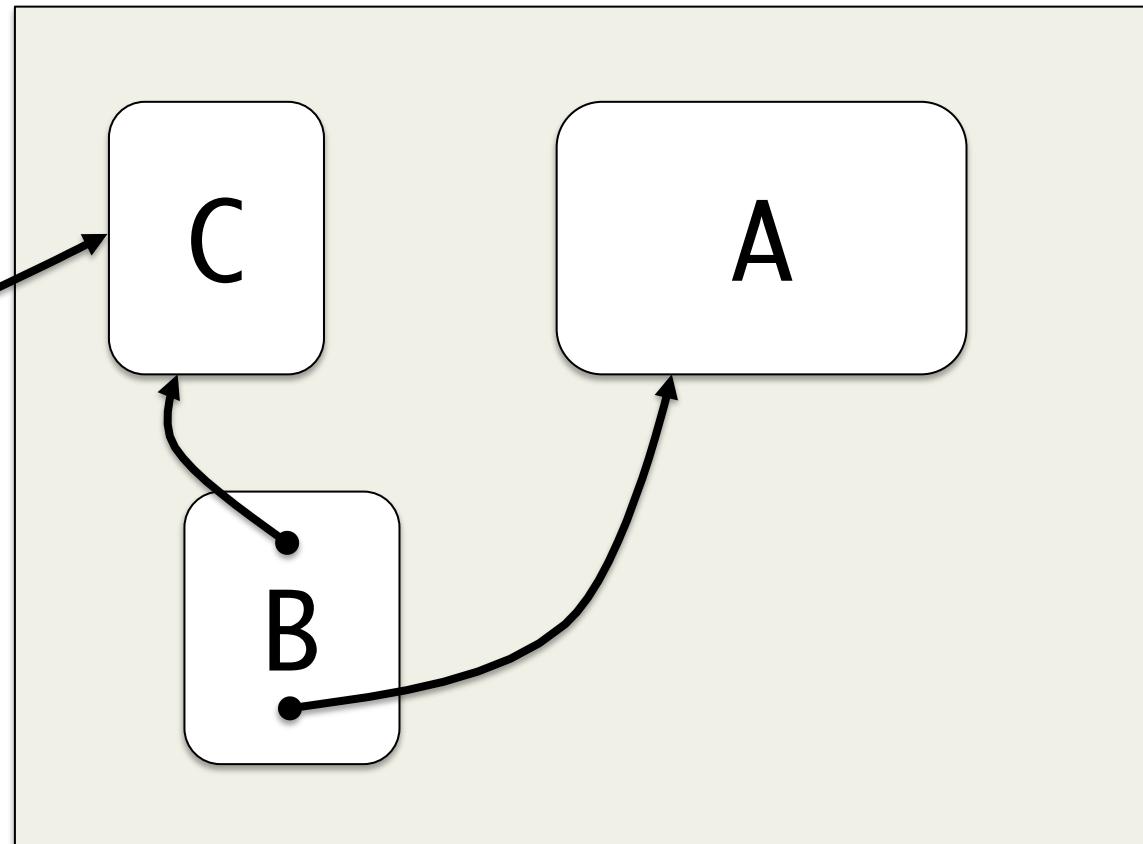
- Finalize method can contain any Java code
 - Recommended use is to clean up resources
- Finalization is handled by a daemon thread
 - Thread guaranteed not to currently hold any locks
- Exceptions thrown by a finalizer are ignored

Implementing Finalization

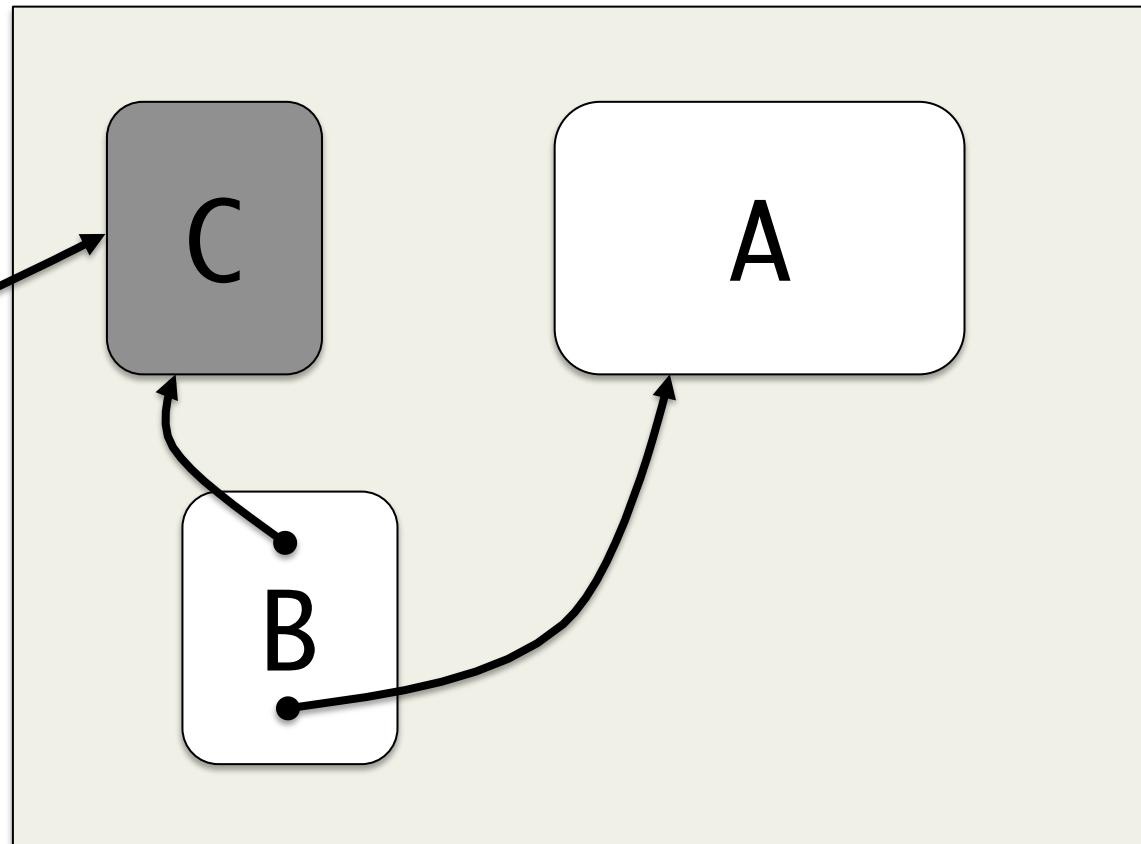
- Maintain a finalization queue
 - Check for a finalizer when freeing an object
 - May also flag finalizable objects at allocation time
- Add to the queue only before freeing the memory
 - Finalizers run before the memory is re-used
 - Some algorithms don't reuse all parts of the heap
- Run finalizers incrementally or concurrently
 - Finalizers can interact with other threads



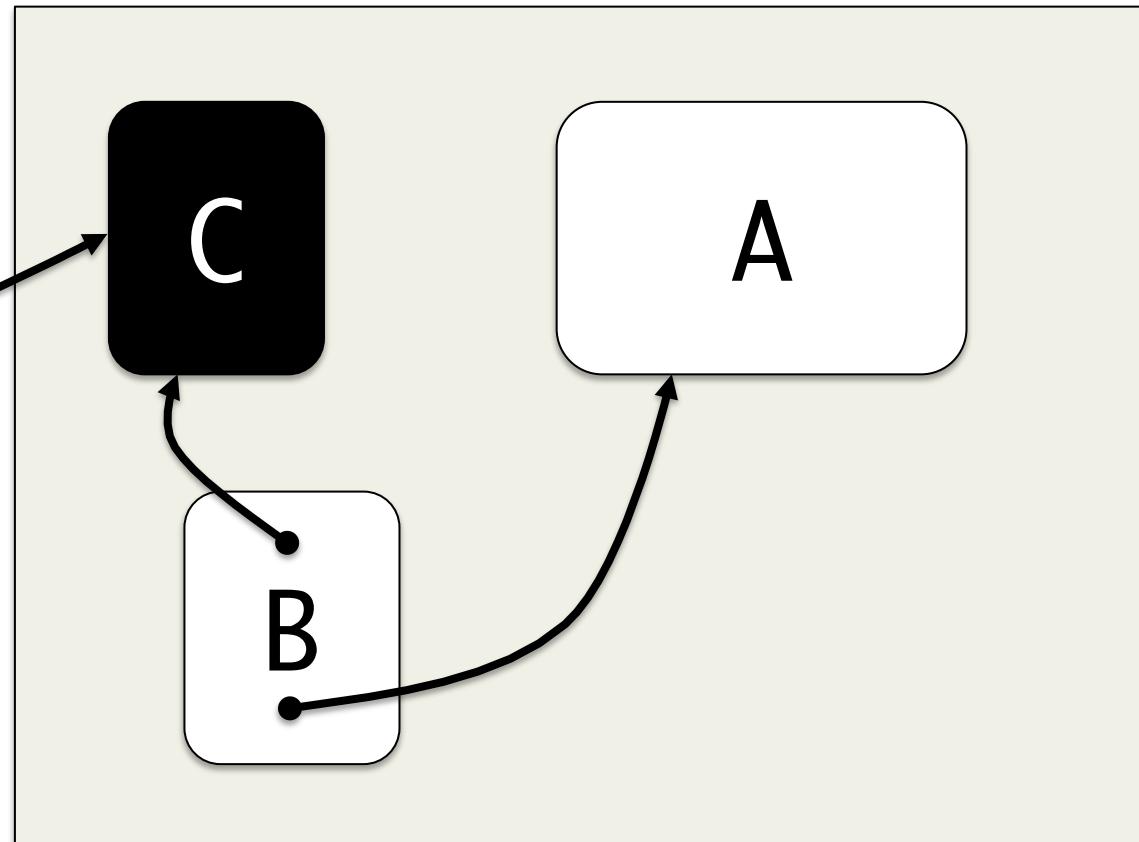
Local Variables

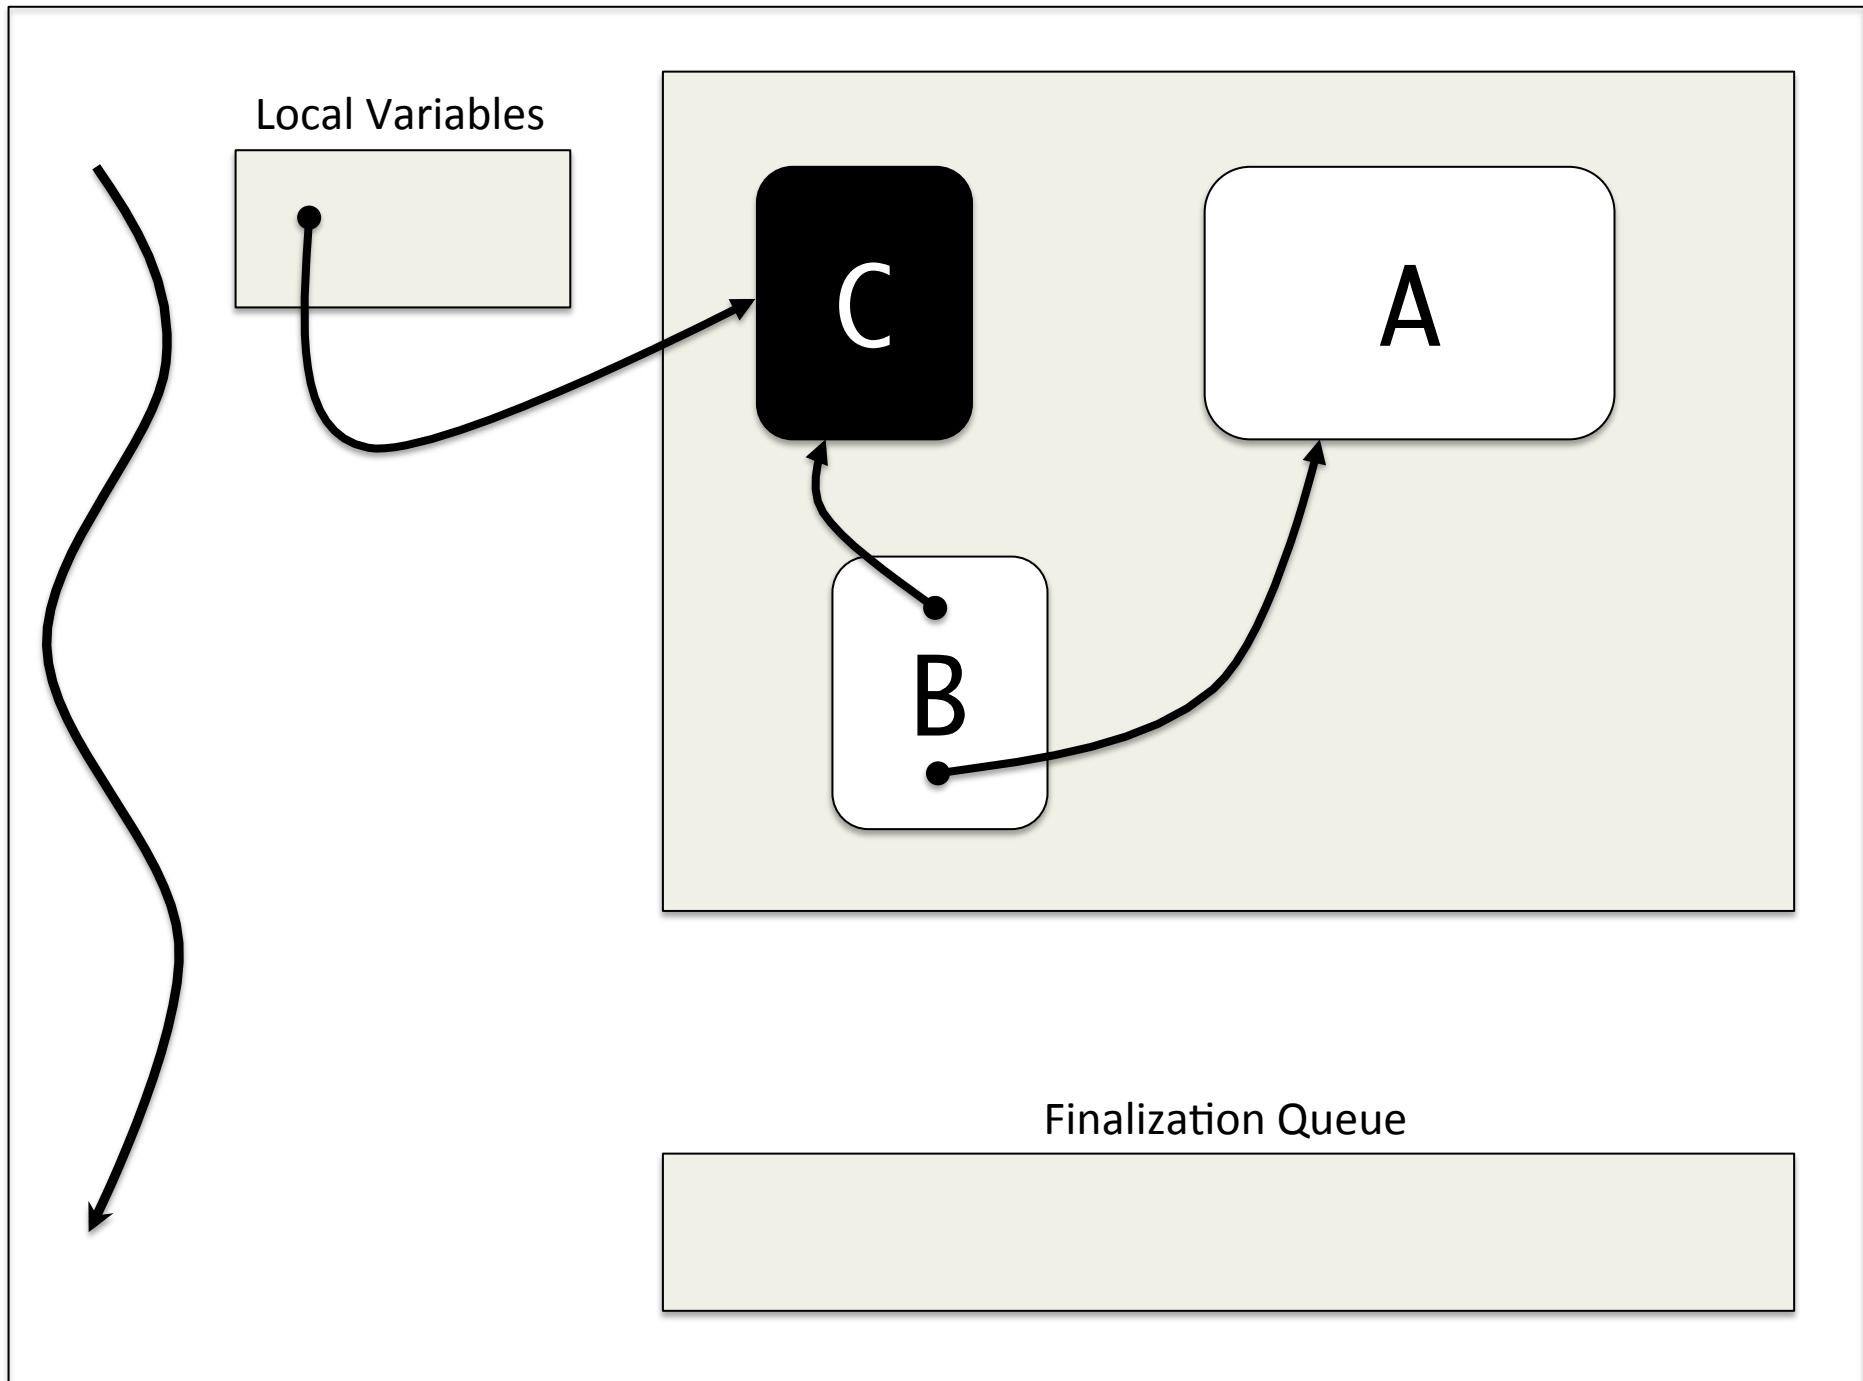


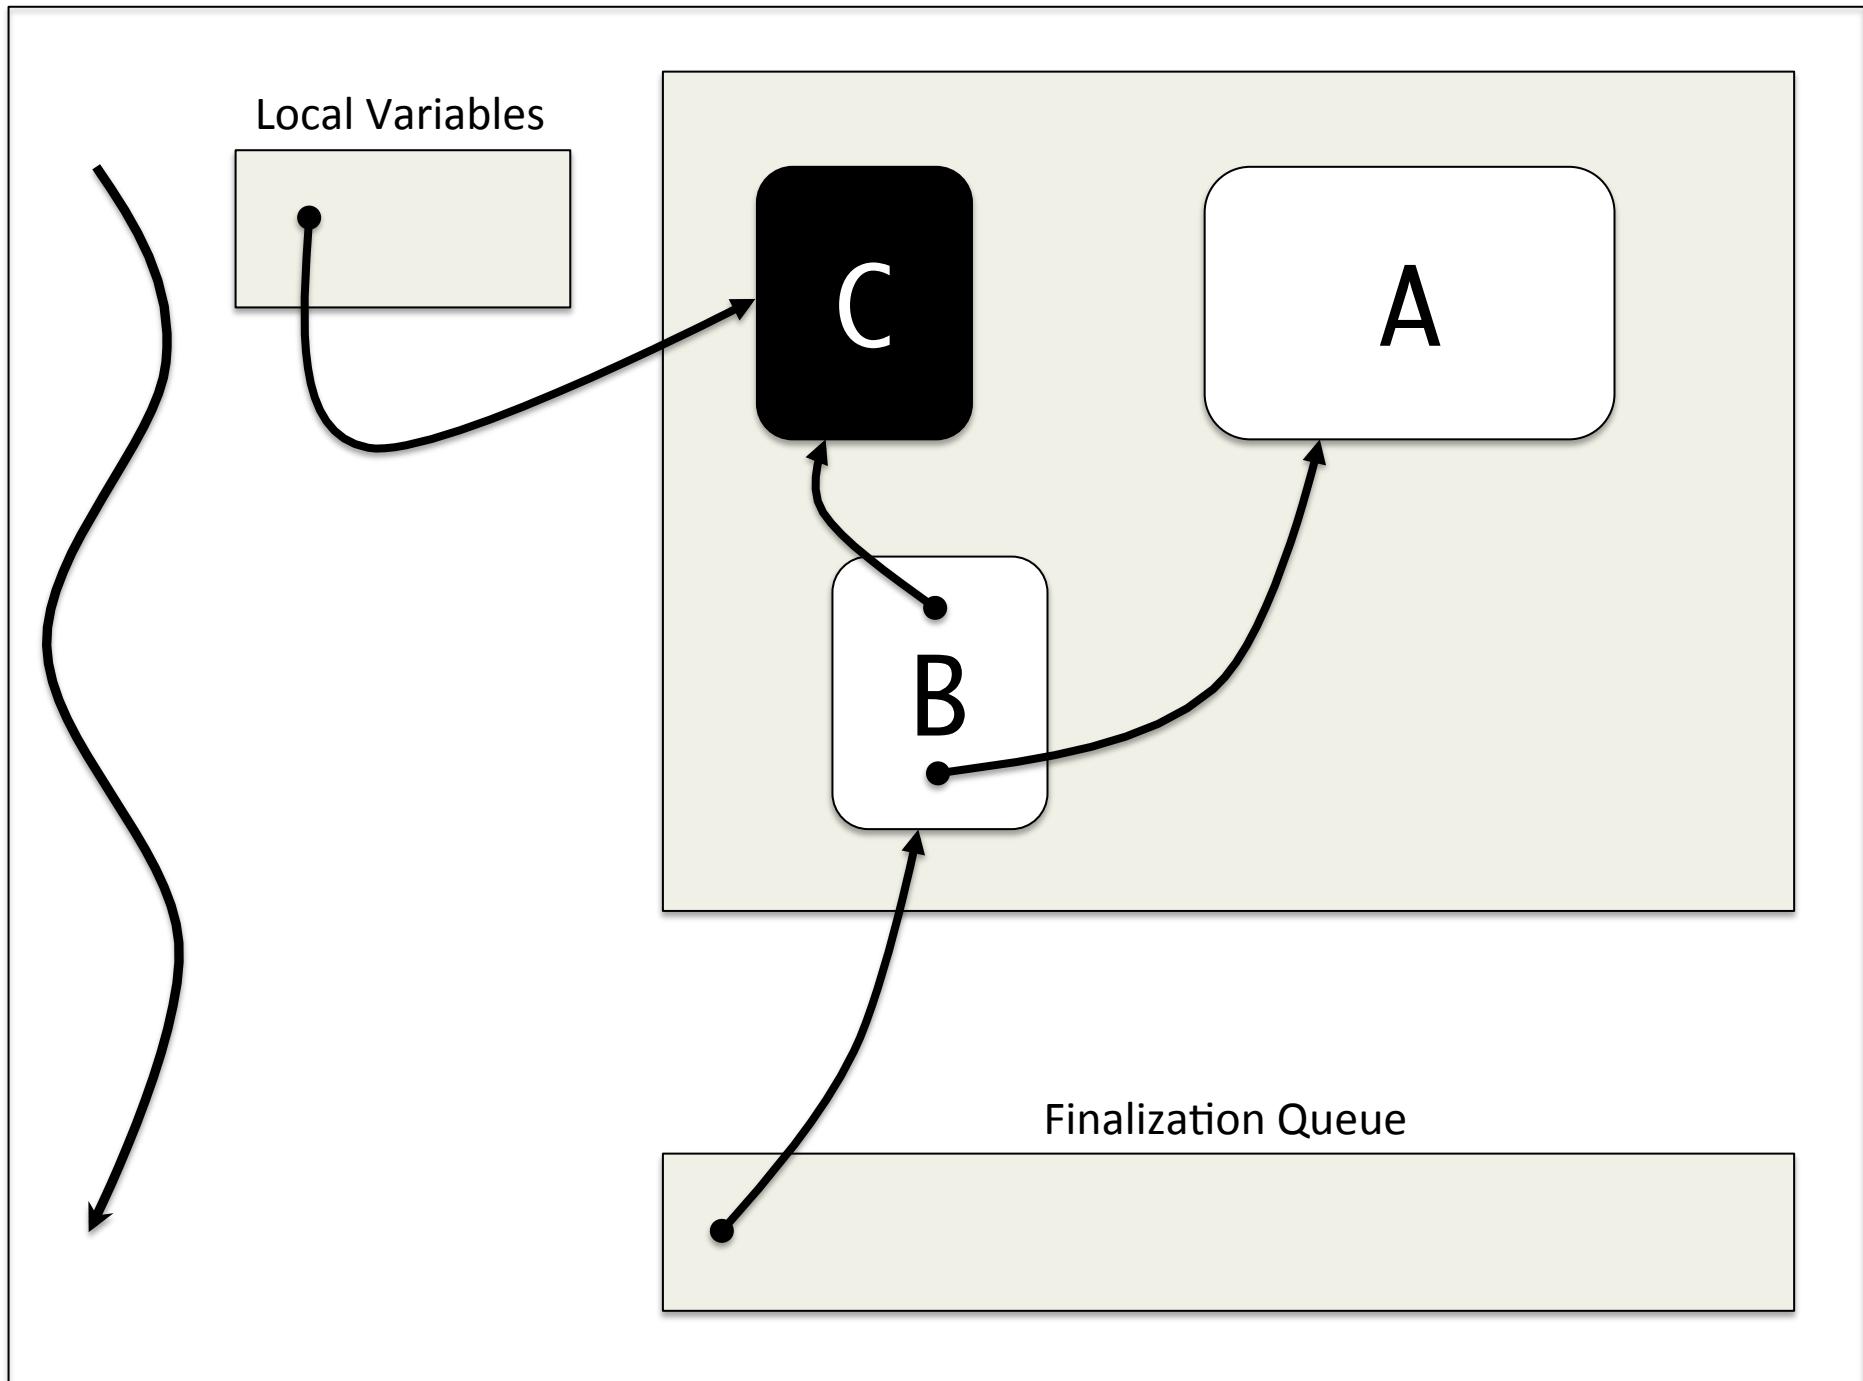
Local Variables

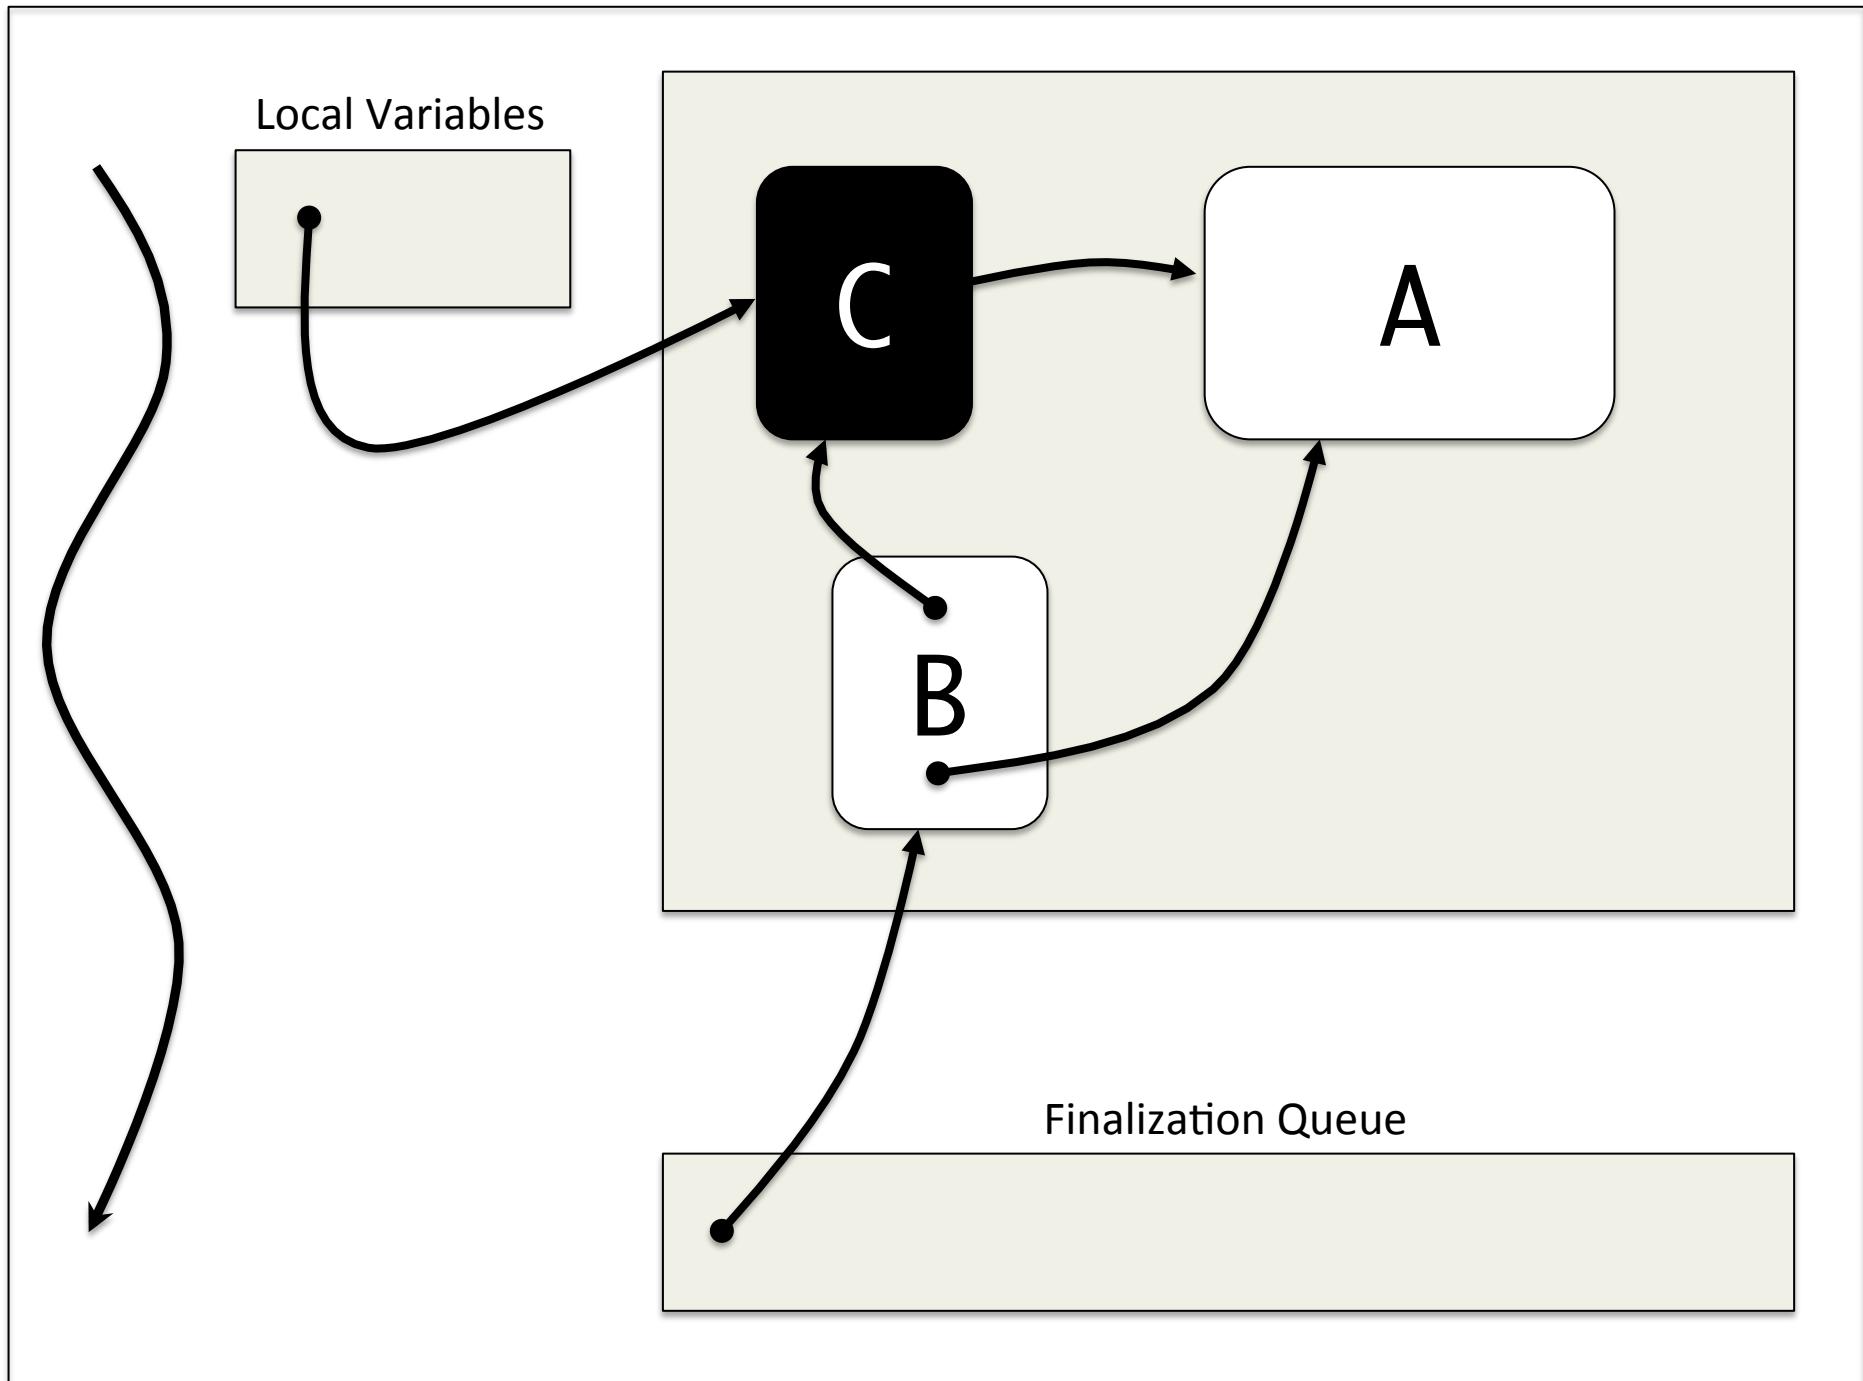


Local Variables

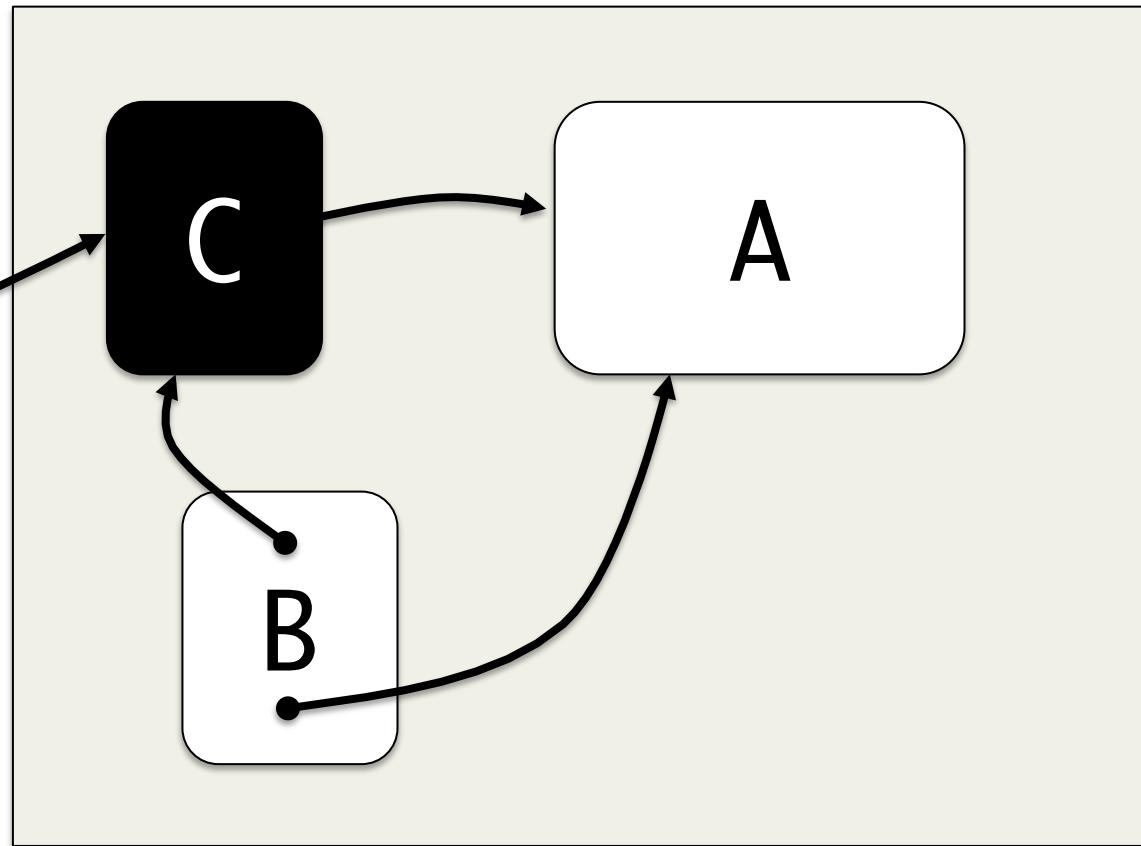






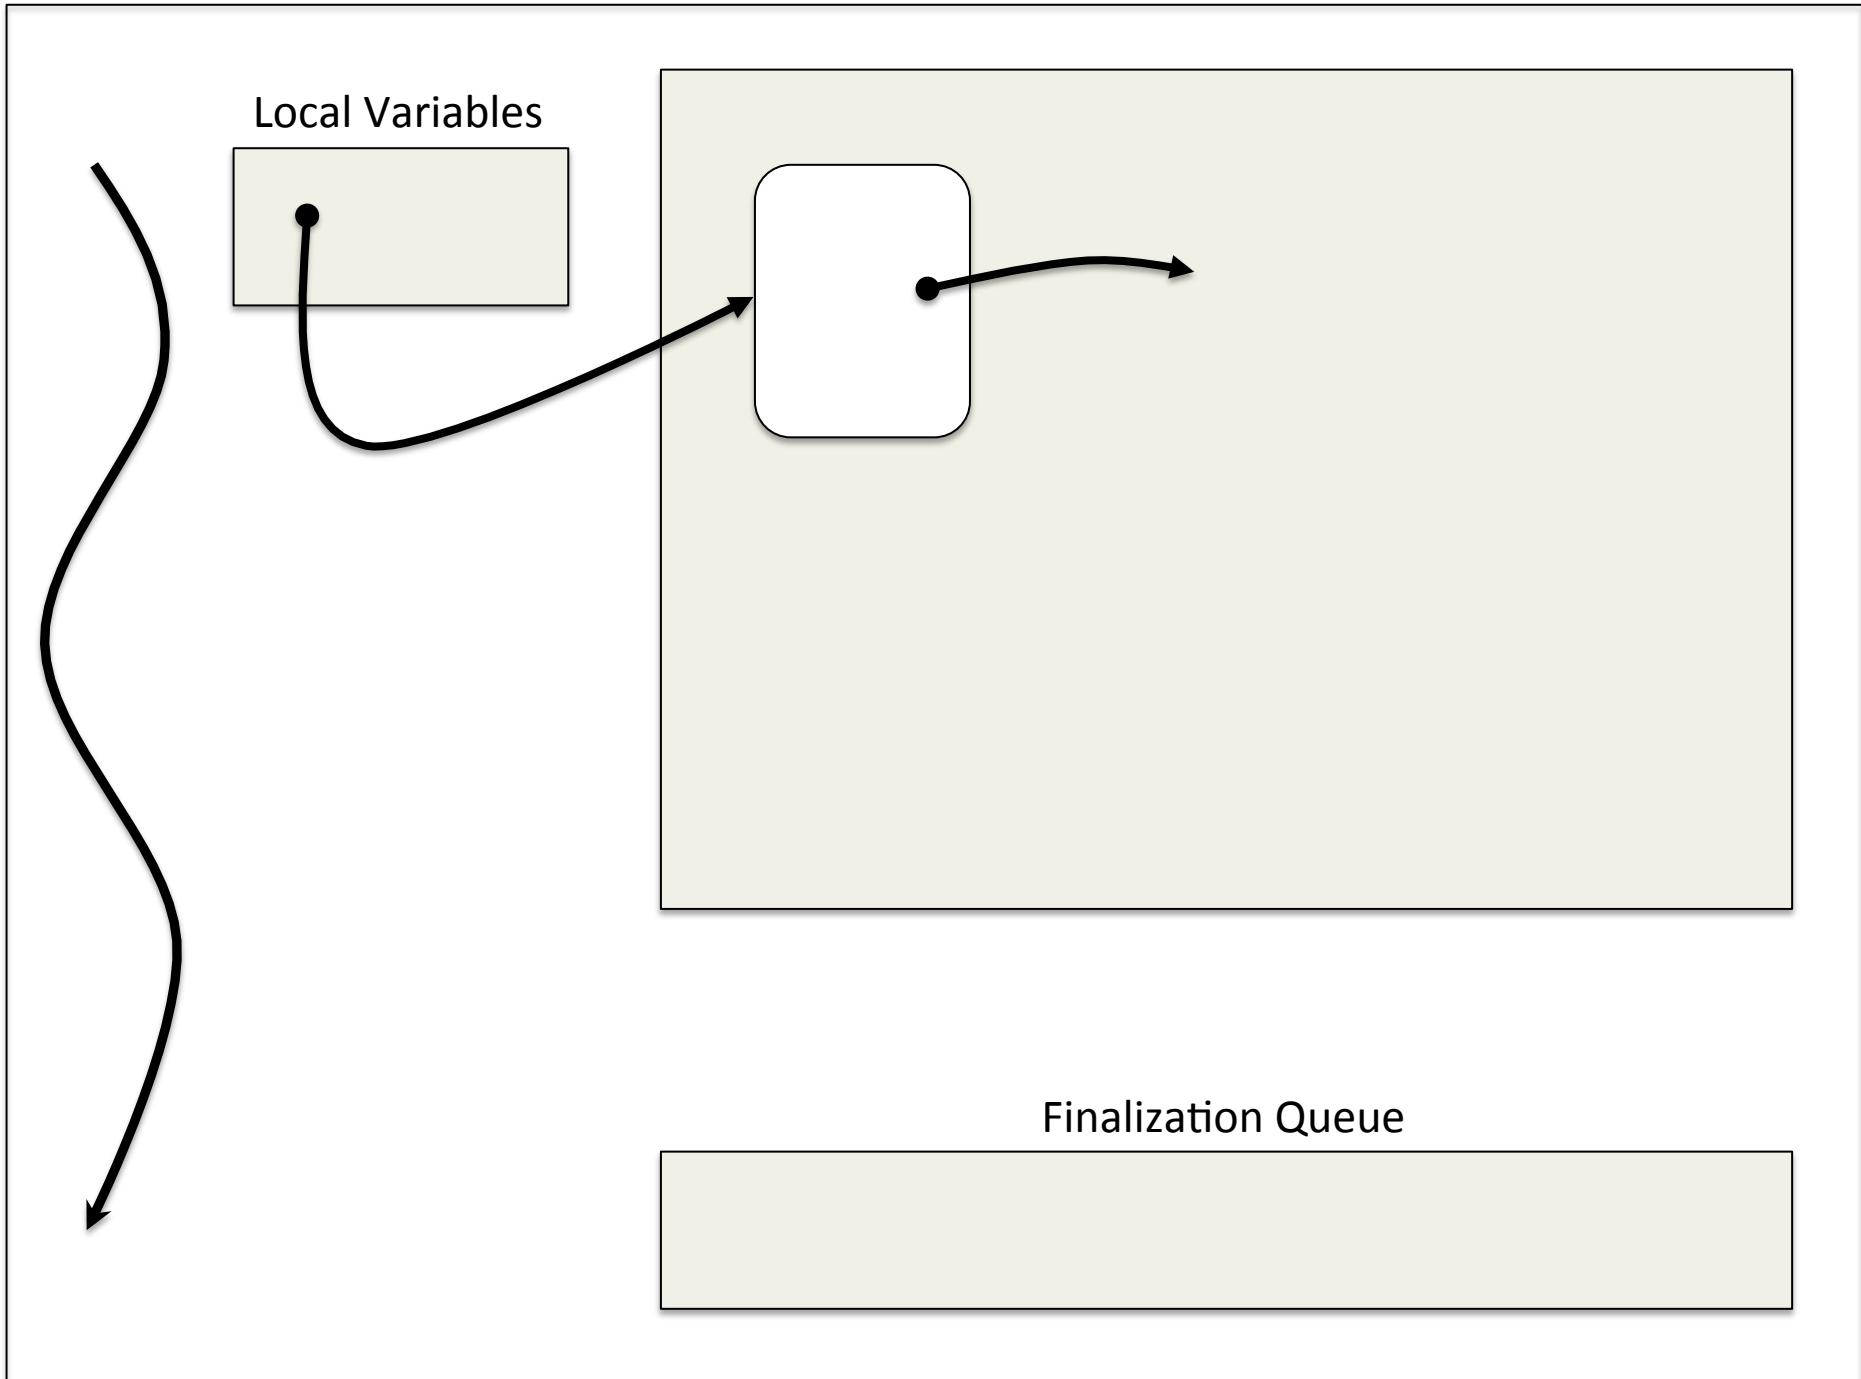


Local Variables



Finalization Queue





Object Resurrection

- A finalizer can make a garbage object reachable
 - May be the this pointer
 - May be another object entirely
- Resurrection greatly complicates GC
 - Can't delete any object that may be referred to
 - Need to wait until all finalizers run before deleting
- Can use a two-phase approach

Phased Finalization

- Trace the heap to find objects to be freed
 - Applies to all tracing algorithms
- Run all finalizers in the finalization queue
 - Remember that any object can only be finalized once
 - This phase may resurrect objects
- Trace the heap again
 - Identify objects that are definitely garbage

Concurrent Finalization

- Some algorithms allow concurrent finalization
 - Extension of the phased finalization algorithm
- Works if garbage can be retained for a GC cycle
 - Ensure that all finalizers run before the next GC
 - Tracing phase of next GC finds unreachable objects
- Track phantom objects separately
 - Those found unreachable and not resurrected

Double Finalization

- A given object's finalizer is called at most once
- Resurrected objects may have different semantics
 - An object may be resurrected after finalization
 - Objects may resurrect themselves
- The object will not be finalized again

Multi-Threaded Finalization

- Objects are finalized by a finalizer thread
 - May be multiple parallel threads
 - Could be a specialized daemon
- The VM can borrow an application thread
 - Must ensure that no locks are held
 - Incremental finalization on constrained systems
- Finalizers can cause deadlock

Why Use Finalizers

- Free up native memory
 - Natural extension to GC semantics
- Last-resort resource clean up
 - Log errors in explicit resource tracking
- Can be useful if you understand the semantics

AHEAD OF TIME COMPILATION

Ahead of Time Compilation

- AOT compilation is a familiar model
 - Write your code
 - Run the compiler
 - Ship the resulting executable
- Java AOT step is generally very minimal
 - Convert the Java code to bytecode
- AOT compilers can build executable from Java

AOT Advantages

- Shift compilation cost from the runtime
 - No interpretation and profiling phase
 - No JIT compilation
- Produce a static binary
 - Single image contains all required code
 - Can do whole-program analysis
- Well adapted to constrained environments
 - Low power prohibits JIT compilation
 - Uncooperative environments don't allow dynamic code

AOT Limitations

- Up-front compilation cost
 - Resulting executable is platform specific
- Can't use dynamic classloading
 - All classes must be declared ahead of time
 - Harder to adapt the application to the environment
- Limited profiling and debugging
 - JVMTI instruments bytecode
 - Requires interpreter or JIT compilation

Ahead of Time Compilers

- Various compiler implementations
 - GCJ
 - Excelsior
 - LLVM
- Compiler requires a class library implementation
 - Licensing issues with the standard JDK library
 - GNU Classpath offers partial compatibility

LLVM

- Open source compilation framework
- Various language-specific front-ends
 - Common intermediate representation
- Basis of various compilers in everyday use
 - Clang used as a replacement for GCC
- VMKit was a JVM leveraging LLVM
 - LLVM + MMTk + pthreads + glue

GC in LLVM

- Limited support for garbage collection
 - Doesn't provide an actual GC implementation
 - Provides hooks for runtime developers
- Binary interface in LLVM support intrinsics
 - Creating GC safepoints
 - Computing stack maps
 - Read and write barriers
- Intermediate representation supports other aspects
 - Type maps
 - Root locations

LLVM as a JIT

- LLVM was thought to be static only
 - Optimizations required large analysis
 - Doesn't take advantage of profiling data
 - Support for mixed mode operation limited
 - Compilation is much slower than required
- Apple released LLVM-based JIT in Webkit
 - Dynamically compiling Javascript in Safari
 - Significant engineering work

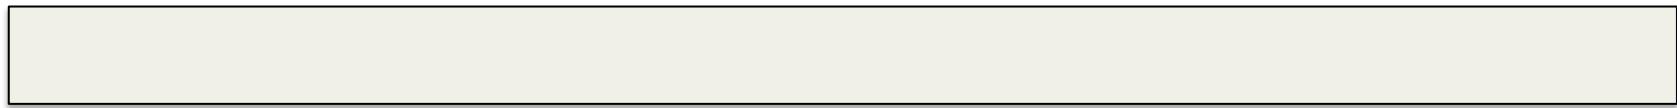
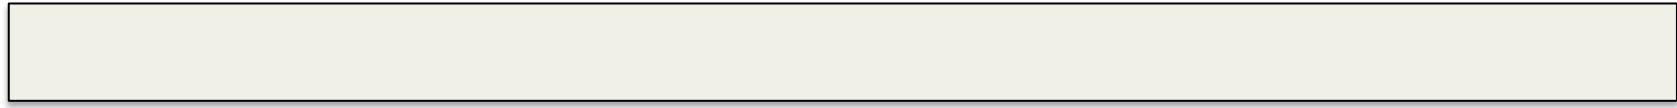
GARBAGE FIRST ALGORITHM

Garbage First Algorithm

- Latest iteration of the Hotspot collector
 - Not enabled by default
 - Targets pause times over throughput
 - Many server applications want the inverse
- GC is parallel and concurrent
 - Copies objects to compact the heap
 - Pauses are predictable and tunable
 - Offers soft real-time without guarantees

G1 Regions

- Heap is split into equally-sized regions
 - Region sizes range from 1-32 Mb
 - Aim to have no more than 2048 regions
 - More tuning to enable more than 64 Gb heap
- Young generation made up of two block types
 - Nursery (or Eden)
 - Survivor set
- Very large objects given a region to themselves
 - Multiple regions can coalesce if necessary



--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

Nursery	Nursery						
---------	---------	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

							Nursery
--	--	--	--	--	--	--	---------

	Nursery		Nursery				
--	---------	--	---------	--	--	--	--

Nursery							
---------	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--

Nursery	Nursery		Survivor				
---------	---------	--	----------	--	--	--	--

--	--	--	--	--	--	--	--

				Survivor			
--	--	--	--	----------	--	--	--

		Survivor				Nursery	
--	--	----------	--	--	--	---------	--

	Nursery		Nursery				
--	---------	--	---------	--	--	--	--

Nursery							
---------	--	--	--	--	--	--	--

		Survivor			Survivor		
--	--	----------	--	--	----------	--	--

							Survivor
--	--	--	--	--	--	--	----------

Nursery	Nursery		Survivor				
---------	---------	--	----------	--	--	--	--

	Mature	Mature			Mature		
--	--------	--------	--	--	--------	--	--

Mature	Mature			Survivor		Mature	Mature
--------	--------	--	--	----------	--	--------	--------

		Survivor		Mature		Nursery	
--	--	----------	--	--------	--	---------	--

Mature	Nursery	Mature	Nursery				
--------	---------	--------	---------	--	--	--	--

Nursery	Mature			Mature		Mature	
---------	--------	--	--	--------	--	--------	--

		Survivor			Survivor		
--	--	----------	--	--	----------	--	--

	Mature						Survivor
--	--------	--	--	--	--	--	----------

Nursery	Nursery		Survivor	Humongous		
---------	---------	--	----------	-----------	--	--

	Mature	Mature			Mature	
--	--------	--------	--	--	--------	--

Mature	Mature			Survivor		Mature	Mature
--------	--------	--	--	----------	--	--------	--------

		Survivor		Mature		Nursery	
--	--	----------	--	--------	--	---------	--

Mature	Nursery	Mature	Nursery				
--------	---------	--------	---------	--	--	--	--

Nursery	Mature			Mature		Mature	
---------	--------	--	--	--------	--	--------	--

		Survivor			Survivor		
--	--	----------	--	--	----------	--	--

	Mature		Humongous			Survivor	
--	--------	--	-----------	--	--	----------	--

Remembered Sets

- Maintain remembered sets for each region
 - Track references to that region from outside
 - Don't track internal references
- Maintain set using a write barrier
 - Use block alignment to generate region ID
 - Simple shift operation
- Remembered sets give estimate of live data

Minor GC

- Parallel stop-the-world copying collection
- Promote all live young objects
 - Move objects from nursery to the survivor regions
 - Move objects from the survivor to the mature
- Don't evacuate all young objects at once
 - Determine nursery regions by pause time
 - Determine survivor regions by garbage ratio estimates

Concurrent Marking

- Initial marking to scan the root set
 - Requires that threads be stopped
 - Piggy-backs on a minor GC pause
- Concurrent marking across the heap
 - Uses incremental update
- Remark phase to complete incremental update
 - Stops the world

Cleanup Phase

- After mark phase we know all live objects
 - Remembered sets aggregate the information
- We can explicitly target regions with more garbage
 - Minimize copy overhead
 - Reclaim more garbage in a bounded time
- In many cases a region is entirely garbage
 - Remember the train algorithm

Evacuation

- Final part of major collection
- Iterate over regions selected for copying
 - Copy all live objects to an empty region
 - Highly parallel
 - Uses a form of TLAB to minimize contention
- Number of regions copied can be tuned
 - Trade-off throughput for pause time

BENCHMARKING

Benchmarking

- Benchmarks are big business
 - Drive adoption of a given platform
 - Encourage investment in technology
- They provide great ammunition for flame wars
 - What would Slashdot be without benchmarks?
- Benchmarks are almost all worthless

Apples to Apples Comparison

- An ideal benchmark compares like for like
 - Single application running on different JVMs
 - Identical hardware and software platforms
 - Running in single-user mode
 - Networking and other interrupts turned off
- A good experiment uses many runs
 - Average over many ($n > 30$) iterations
 - Take standard deviation, confidence intervals, etc

Micro Benchmarks

- Small kernels of code that test a given operation
- Generally implemented as a tight loop
 - Allocations per second
 - Monitor acquisitions and releases
 - Floating point operations per second
- Very hard to make a micro benchmark interesting

```
public class SquareRoot {  
  
    public static void main(final String[] args) {  
        int iterations = Integer.parseInt(args[0]);  
        long start = System.currentTimeMillis();  
  
        for (int i = 0; i < iterations; i++) {  
            Math.sqrt((double) i);  
        }  
  
        long end = System.currentTimeMillis();  
  
        System.out.println("Time: " + (end - start));  
    }  
}
```

JIT Compilation

- The JIT is designed to defeat microbenchmarks
 - Dead loop elimination
 - Constant folding
 - Escape analysis
- Many microbenchmarks measure startup costs
 - Code runs in interpreted mode for a time
 - JIT compiler is triggered
 - Benchmark behavior changes completely

Phase Behavior

- Native code tends to start at full speed
 - Compilation happens ahead of time
- Managed runtimes take time to get started
 - Startup time required to initialize the VM
 - Initially run in interpreted mode
 - JIT compilation kicks in after a while
 - Runs at full speed after compilation period

Macro Benchmarks

- Aim to emulate real-world applications
- Various SPEC benchmark suites
 - SPECjbb2000, 2008, 2013
 - SPECjvm98, 2008
- DaCapo benchmark suite
 - Targets interesting memory management behavior
- Java Linpack
 - Java version of popular floating point benchmark

Macro Benchmark Benefits

- More likely to reflect real applications
 - Run for a reasonably long time
 - Workloads can be analyzed and considered
 - More often designed to account for phase behavior
- More varied operations
 - Don't obsess over a single part of the VM
 - Use a mix of custom and library code

Macro Benchmark Issues

- May not test what you think
 - Applications sometimes spend time in strange places
- Repetitive workloads
 - Increase run time by doing the same thing again
- Macro benchmark suites become obsolete
 - SpecJVM98 now runs in seconds, not minutes
- Limited number of benchmark suites
 - VM vendors can target and optimize for them

GC Benchmarking

- GC is notoriously hard to benchmark
 - Synthetic applications behave differently from real
 - Benchmark times too short for interesting behavior
- Micro benchmarks focus on the nursery
 - Generate lots of small, similar objects
 - Create occasional large, long-lived objects
- Many macro benchmarks never trigger full GC
 - Well-tuned system can run for surprisingly long

Java Benchmarking Advice

- Include a sufficient warm-up phase
 - Get past the startup overhead
 - Run the workload for a while to let the JIT work
- Measure steady-state behavior
 - Assume that Java applications run for a long time
- Avoids benchmarks targeting Java's weaknesses
 - By focusing on Java's strengths?

General Benchmarking Thoughts

- The only important benchmark is your application
 - Make sure that your own code runs well
 - Target the right platform for the job
- Benchmarking shows that any application can be tuned
 - Use that to your advantage
 - Profile your realistic workloads
 - Optimize the important parts
- Try out your application with different VM settings
 - Or even running on different JVMs

MULTITENANT VM

VM Overhead

- The Java VM is a heavyweight system
 - VM internal data structures and mechanisms
 - Full compiler implementation
- Many applications use the same libraries
 - Collections, Spring, XML processing, etc
- Every Java application incurs the same overhead

Multitenancy

- Ideally, multiple programs could share the VM
 - Single VM startup pause
 - Compile common classes once
- Classloading provides limited sharing
 - Doesn't give the illusion of a private VM
 - Class loading semantics decrease potential sharing
 - Difficulty in terminating the program
 - Semantics for JVMTI

Multi-Tasking Virtual Machine

- Project Barcelona in Sun labs (pre-Oracle)
- Introduced isolates to Java
 - Defined an isolate as a sub-process within a VM
 - Isolates share no state with other isolates in the system
 - Each isolate appears to be its own VM
- Allowed for aggressive optimization across isolates
 - Transparent to the developer

MVMs Today

- Specification is available to implementers
 - JSR 121: Application Isolation API
- Two reference implementations from Sun
 - MVM, MVM2
- Interest died down after the initial work
 - Machine-level virtualization made it less interesting
- Interest appears to be picking back up now
 - Vendors interested in PaaS cloud offerings

VM SHUTDOWN

Shutdown Modes

- Ideally, the JVM always shuts down gracefully
 - May happen if the VM encounters an internal bug
- VM can abruptly halt
 - Receives a SIGKILL or TerminateProcess interrupt
 - The host machine crashes
- Graceful shutdown happens in several cases
 - Final non-daemon thread exits
 - VM receives an external interrupt signal
- Shutdown involves several phases

Shutdown Hooks

- Program designates work during termination
 - Flush buffers
 - Close connections
- Hook is specified as an un-started Thread
 - Recommendation for defensive programming
 - Expectation that work should not take long
- Hooks run in any order at shutdown time

Forcing Finalization

- VM option forces all finalizers to run on exit
 - Gets around the limitation on whether finalizers run
- Finalizers can contain arbitrary code
 - Developers may not have considered shutdown
 - Code could deadlock or slow shutdown
- May be many objects to finalize
 - Large heap sizes imply more objects to finalize

JVM Exit Preparation Step

- Any JVM shutdown hooks run after Java code
 - Delete-on-exit files
- Clean up most VM structures and threads
 - Profiler, signal thread, GC threads
- Notify JVMTI agents
 - VM death events sent to all registered agents
- Native code may still be running
 - Set a flag in the JNI thread environment

Final Shutdown

- Terminate any remaining Java threads
 - Daemon or shutdown hooks
- Bookkeeping tasks for any Java thread
 - Release JNI reference handles
 - Mark thread as terminated
- Terminate the VMThread
 - Bring any threads to a VM safepoint
 - Stop JIT compiler threads
- Return control to the OS

METACIRCULARITY

Self-Hosting

- A compiler is just a regular computer program
 - Takes source files as inputs
 - Produces executable code as output
- A self-hosting is one that can compile itself
 - Compiler written in the target language
- Chicken-and-egg problem
 - Interrupt the recursion with a bootstrap compiler

Metacircularity

- Special case of a self-hosting runtime
 - The language's own features optimize the runtime
- Both VM and application are optimized together
 - VM data structures managed by the GC
 - VM code passes through the JIT
- VM can be written in a higher-level language
 - Use Java synchronization rather than pthreads

JIT Compilation

- Much of the VM implementation relies on the JIT
 - Compiles both the VM and the application
- The JIT is an metacircular JVM is written in Java
 - Which must be JIT compiled
- Same problem as self-hosted compilers
 - Solve in the same way
 - Create a small bootstrap VM
 - Bootstrap VM is used only to get the system running

Intrinsics

- Some required language features are unavailable
 - Low-level pointer access is the most obvious
- Remember JIT compiler intrinsics
 - JIT has a better implementation than the class file
 - Class file is ignored
- We can use intrinsics to implement new features
 - Intercept all calls to a given class or interface
 - JikesRVM uses the org.vmmagic package

Magic

- Magic package split into two parts
- Pragma annotations
 - Tell the compiler to change the method
 - Remove safepoints, control compilation, etc
- Intrinsic annotation
 - Tell the compiler to replace the entire method

Magic Intrinsics

- `org.jikesrvm.runtime.Magic` has method stubs
 - Directly access registers
 - Manipulate stack frames
 - Cast arbitrary bits to Java types
 - Pointer arithmetic
 - Atomic operations such as CAS and memory fences
- Method stubs just throw runtime exceptions
 - Really there for the javac compiler
 - JIT compiler recognizes calls and rewrites