



CS 165

Data Systems

Have fun learning to design and build modern data systems

class 20

updates 2.0

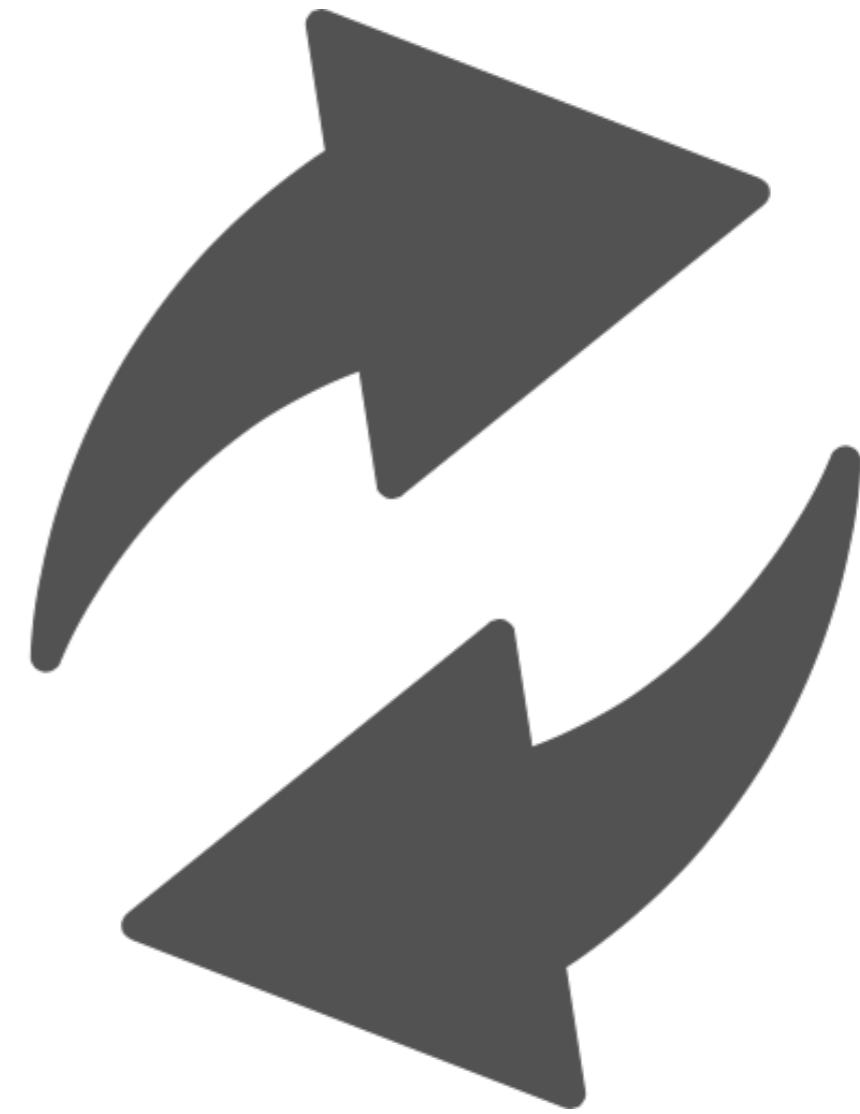
prof. Stratos Idreos

[HTTP://DASLAB.SEAS.HARVARD.EDU/CLASSES/CS165/](http://DASLAB.SEAS.HARVARD.EDU/CLASSES/CS165/)



UPDATE table_name
SET column1=value1,column2=value2,...
WHERE some_column=some_value

INSERT INTO table_name
VALUES (value1,value2,value3,...)

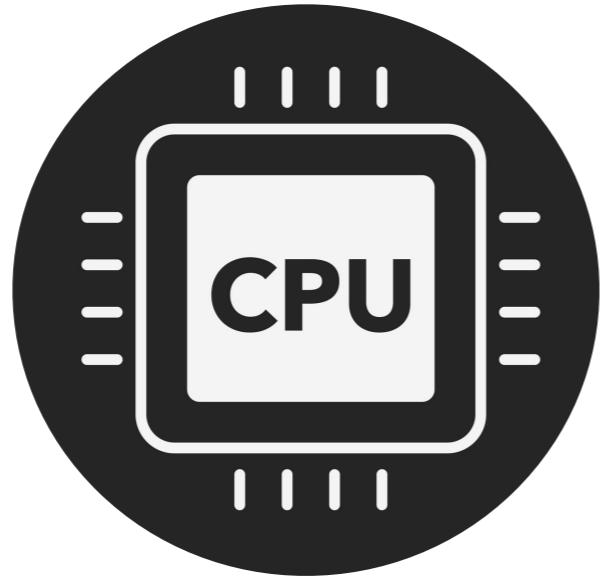


updates

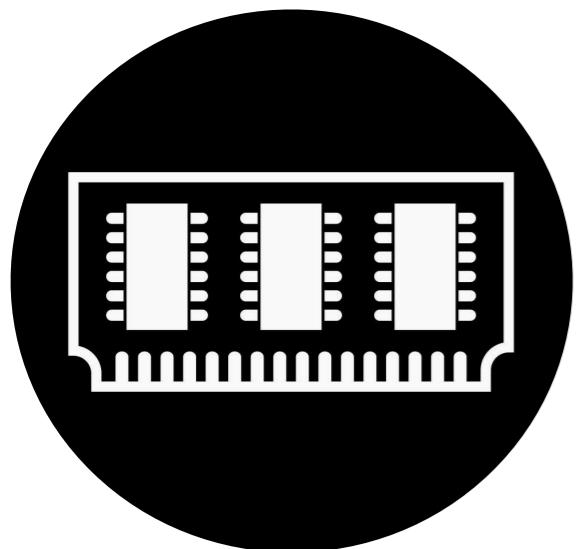
the world has changed
a little bit by now...

updates





monitor CPU utilization



monitor memory
hierarchy utilization



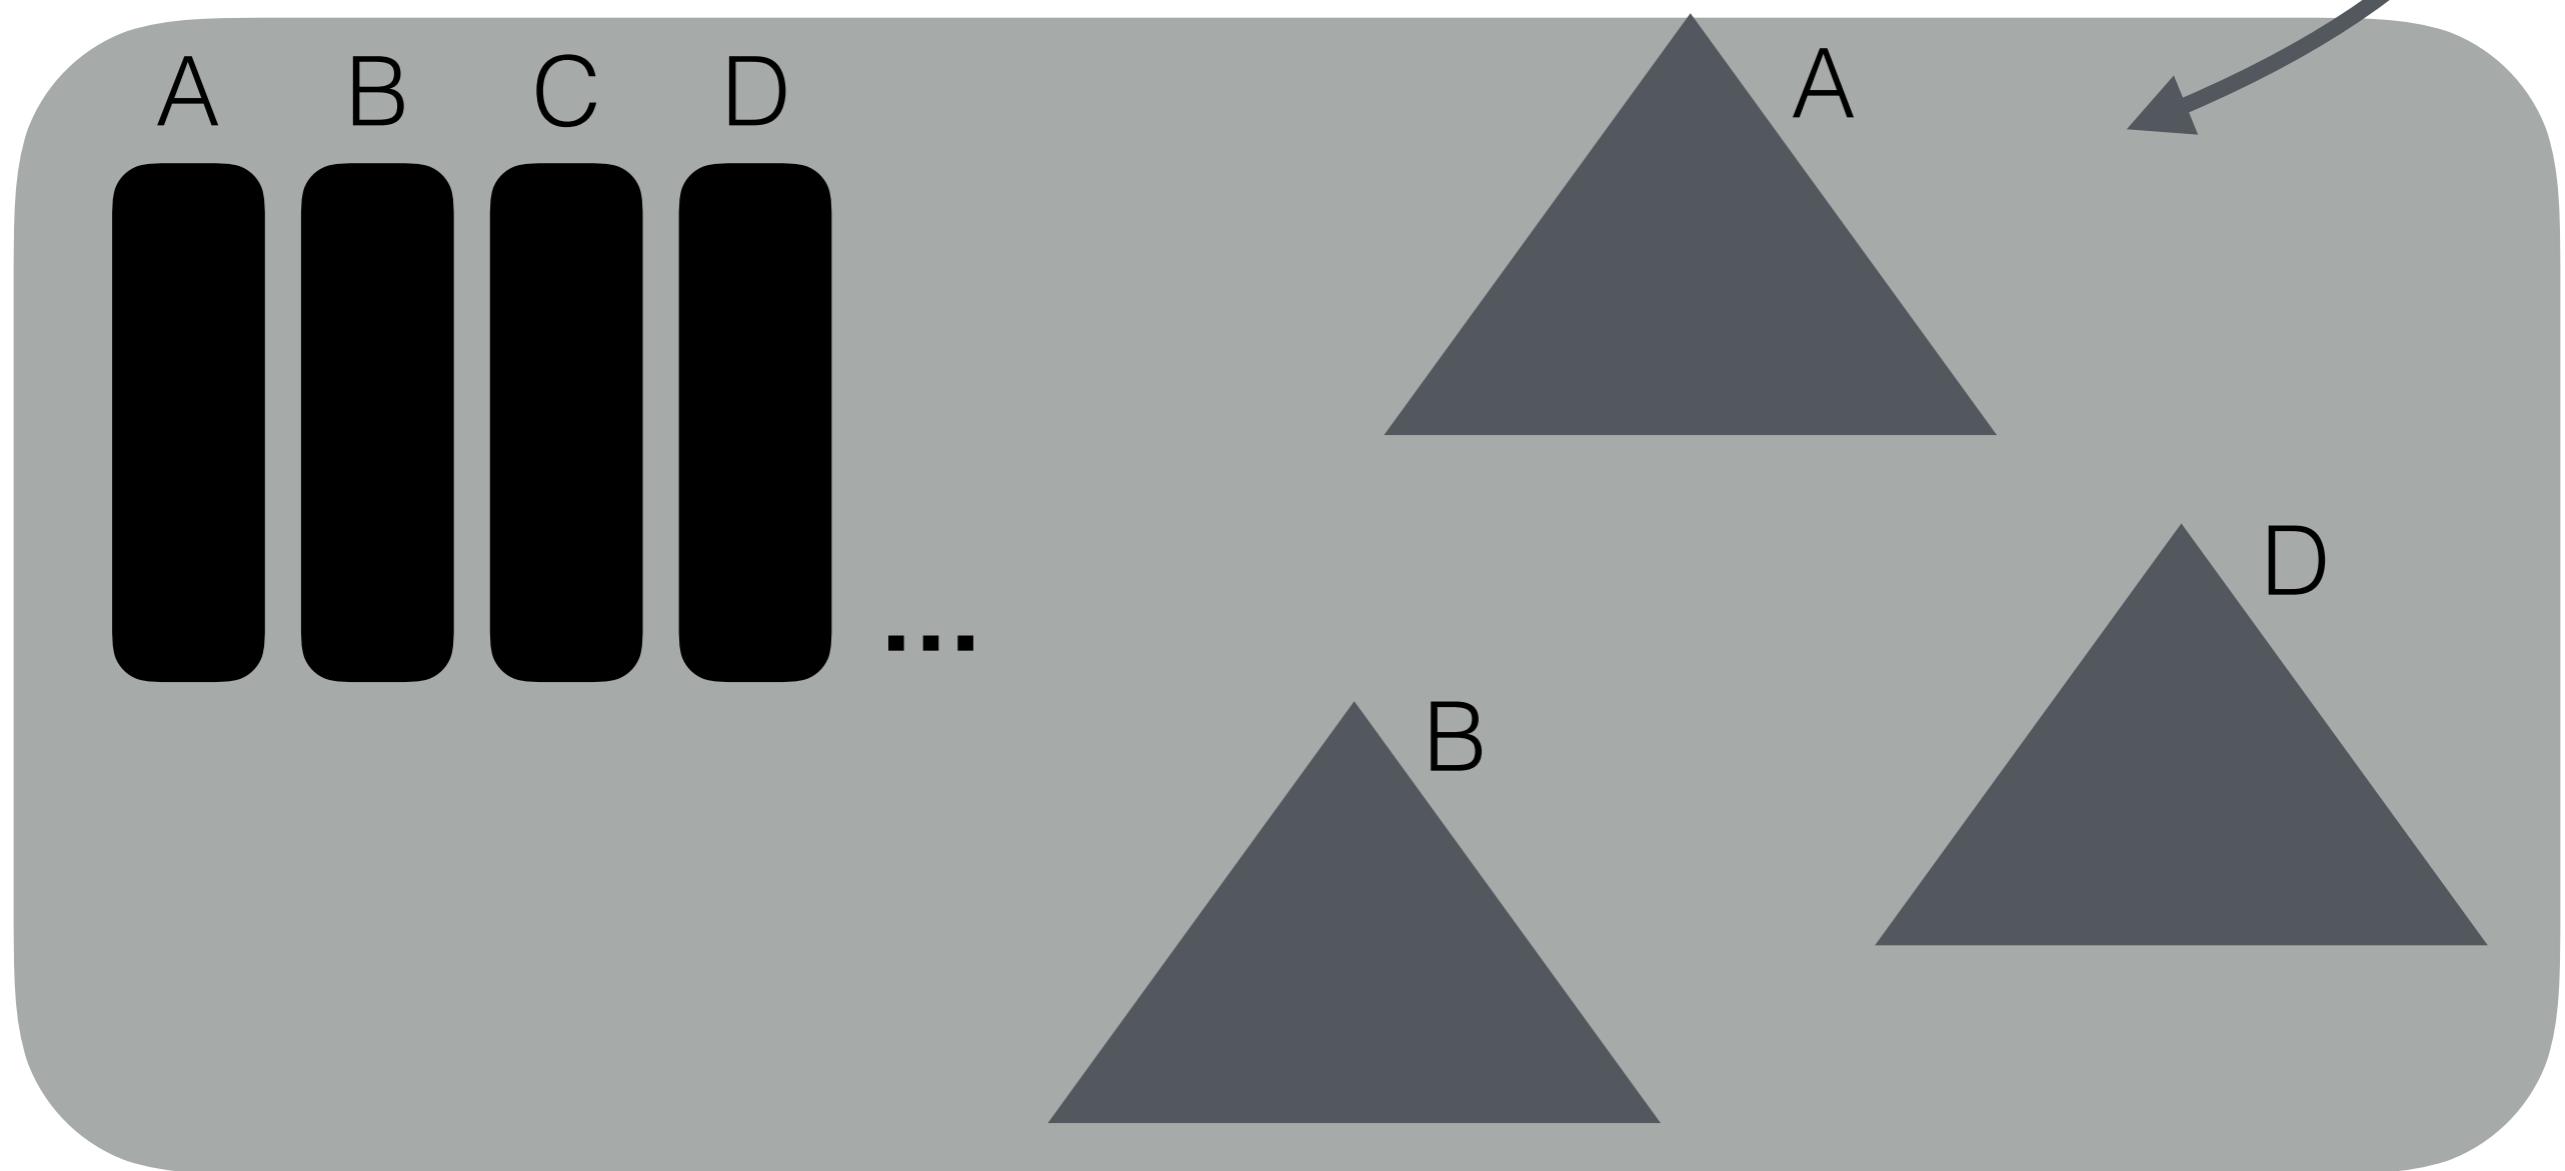
monitor clicks
(frequency, locations,
specific links, sequences)



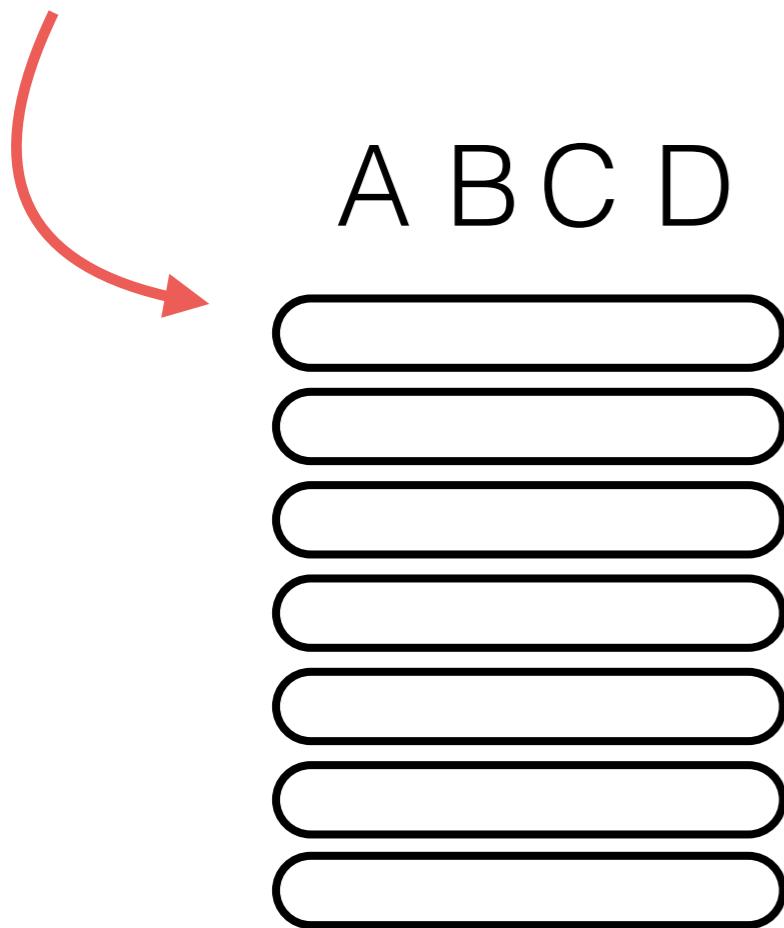
insert new entry (a,b,c,d,...) on table x

update N columns, K trees, statistics, ...

table x

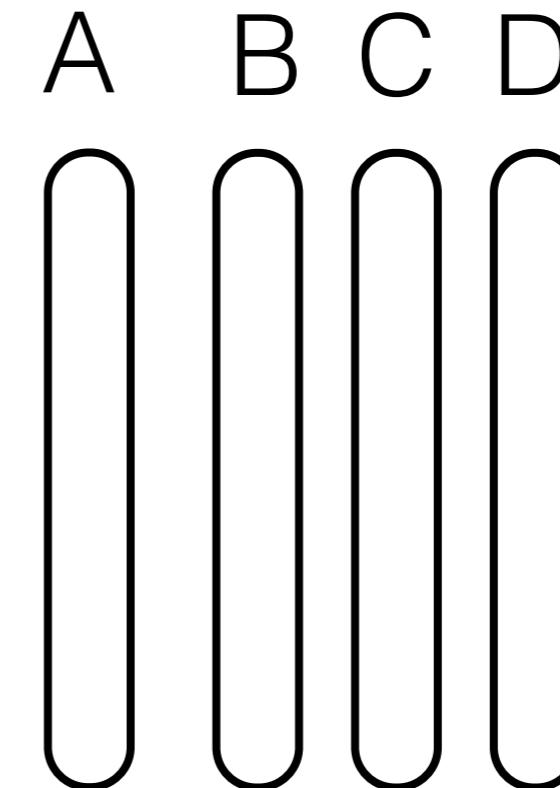


updates



write optimized-store

periodic
merge
and/or
on-the-fly
merge



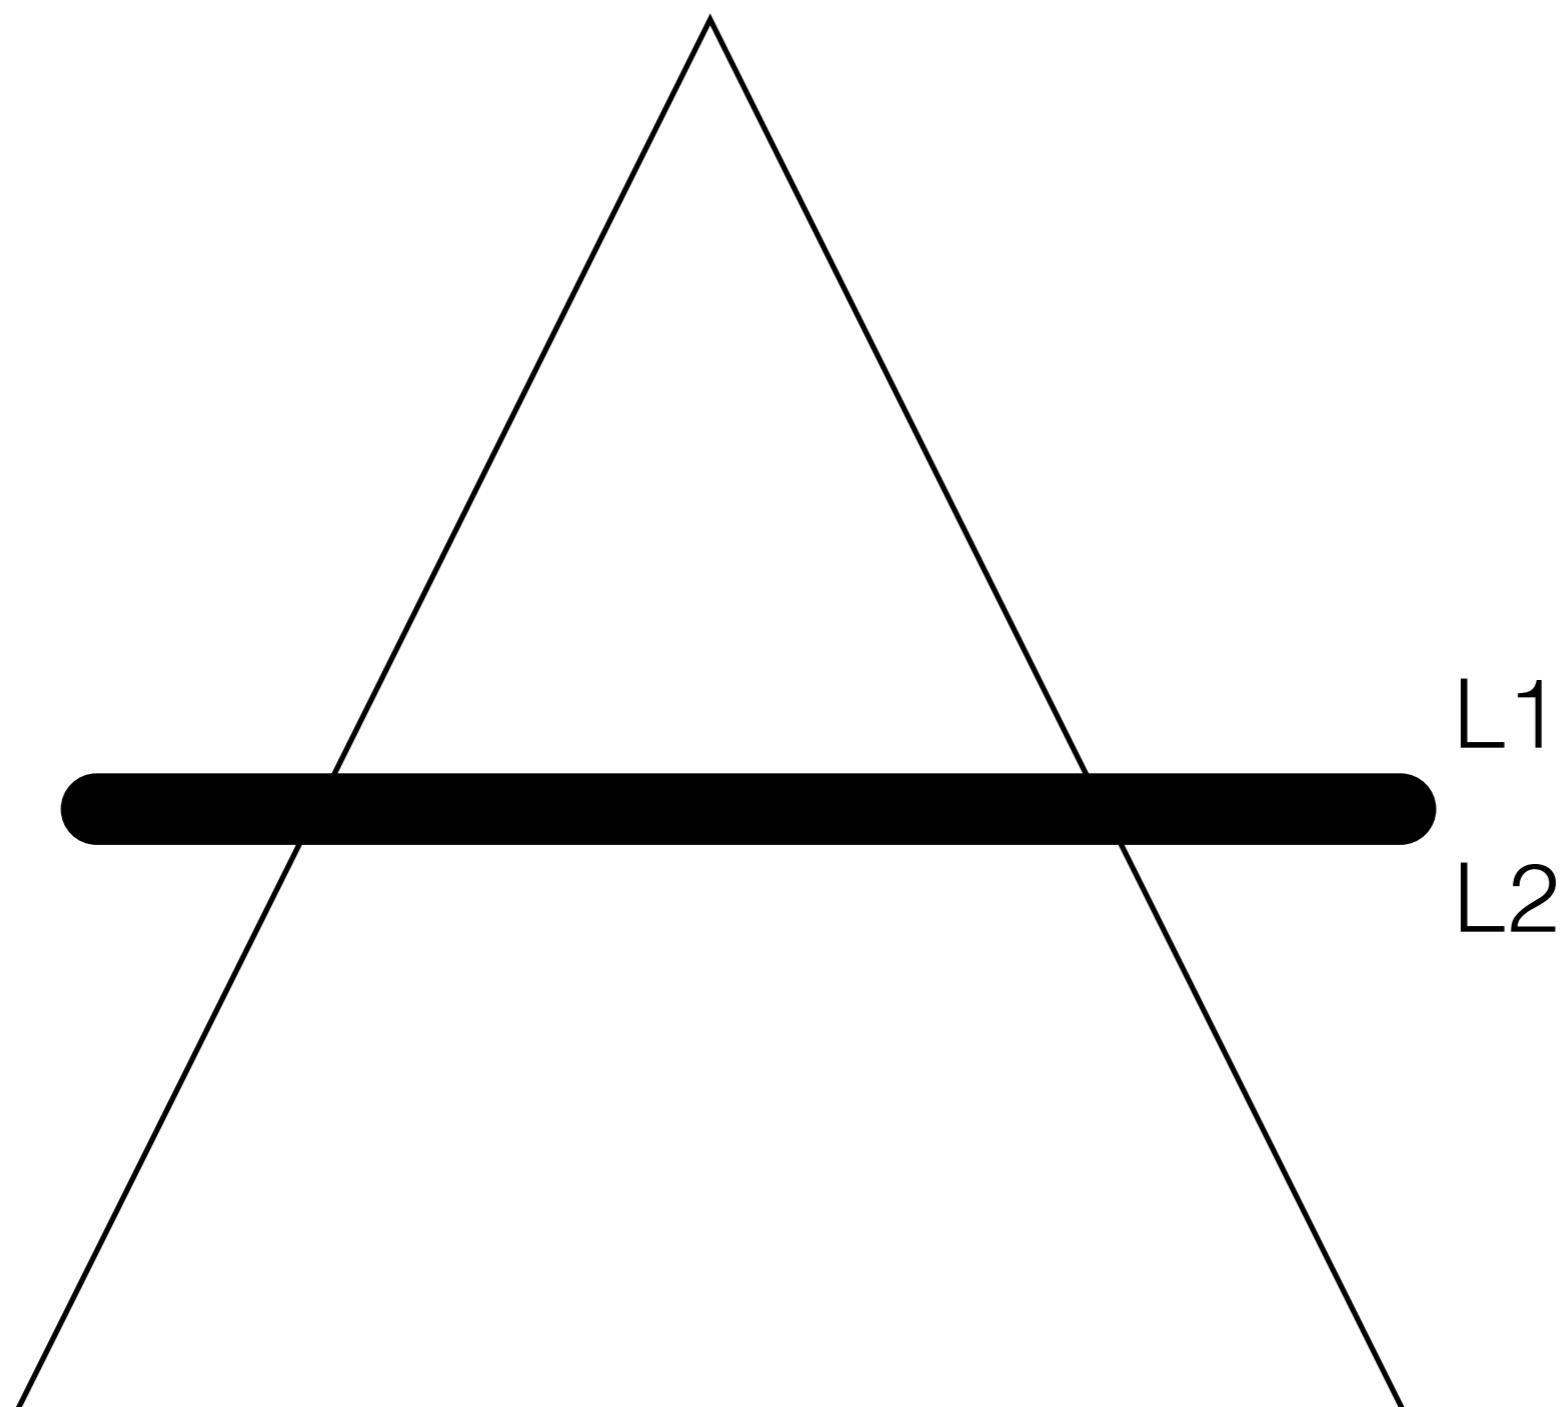
read optimized-store

reads

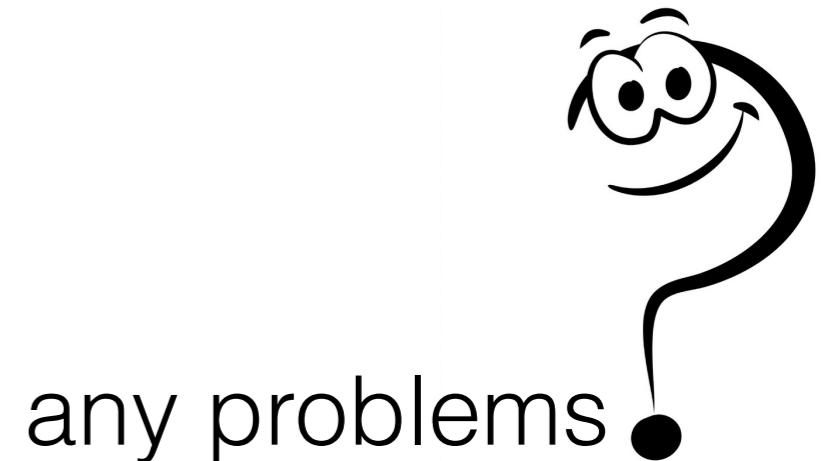
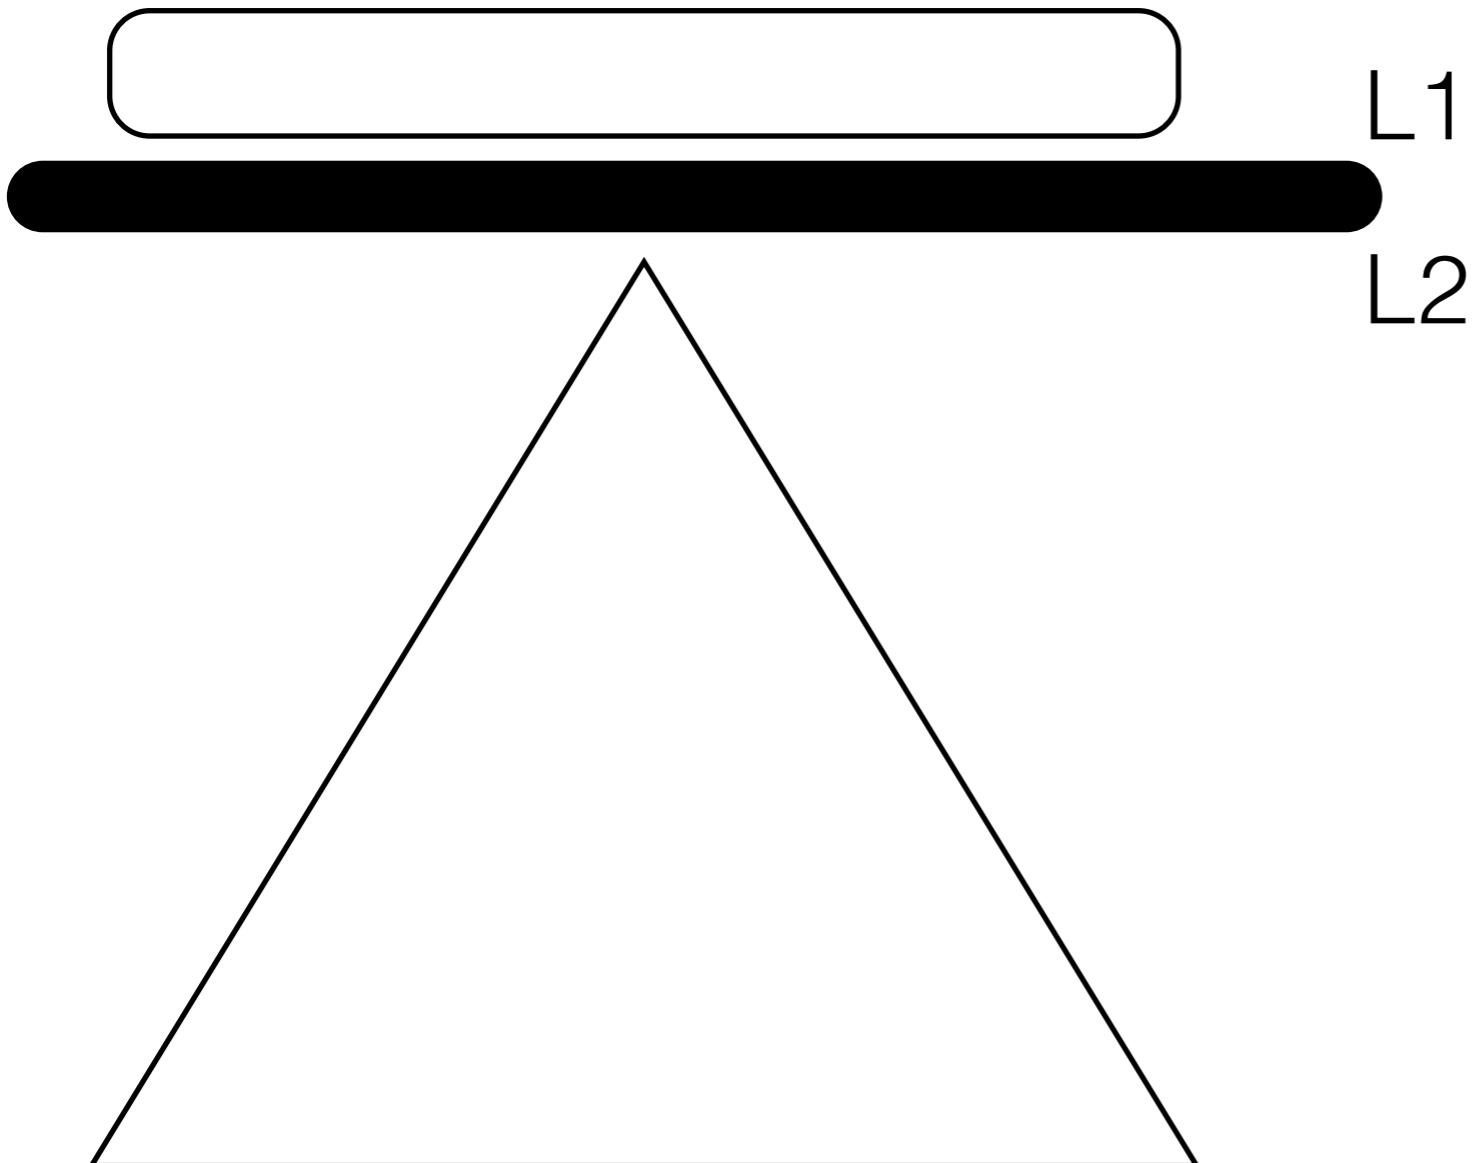
A case for fractured mirrors

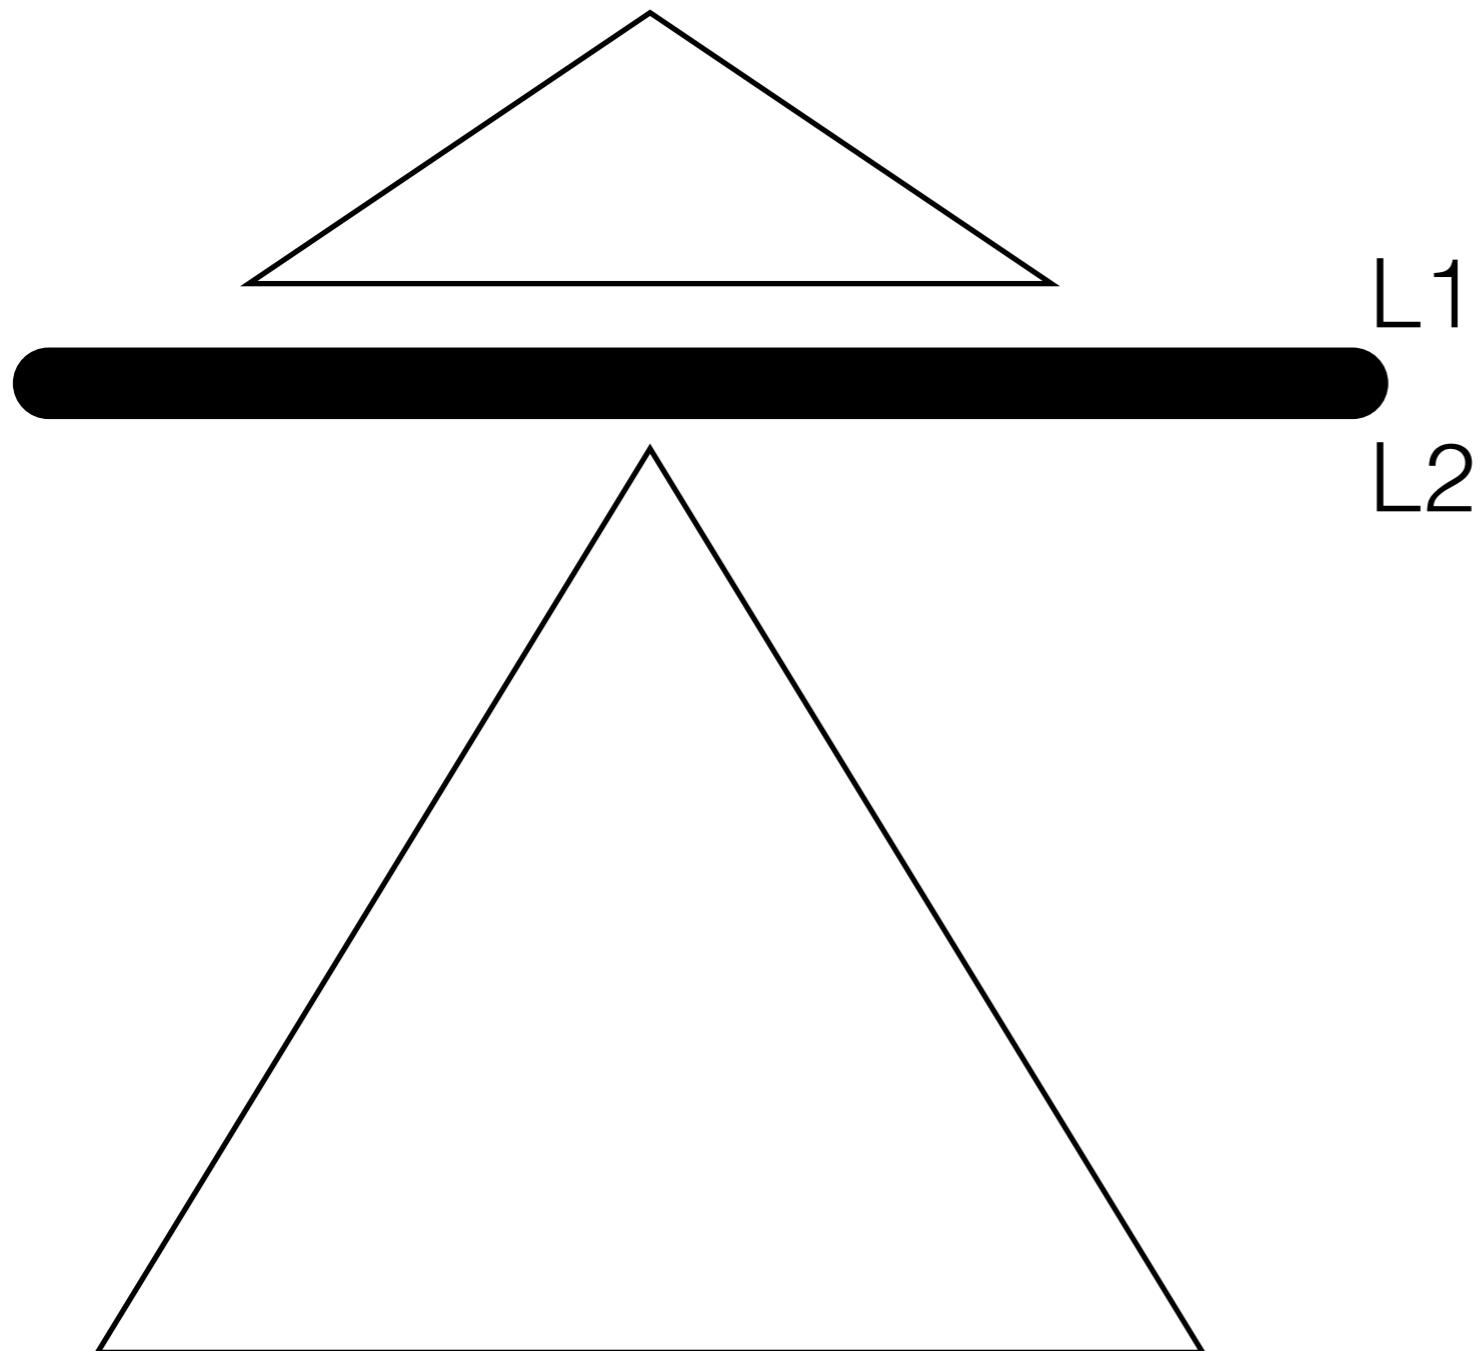
Ravishankar Ramamurthy, David J. DeWitt, Qi Su

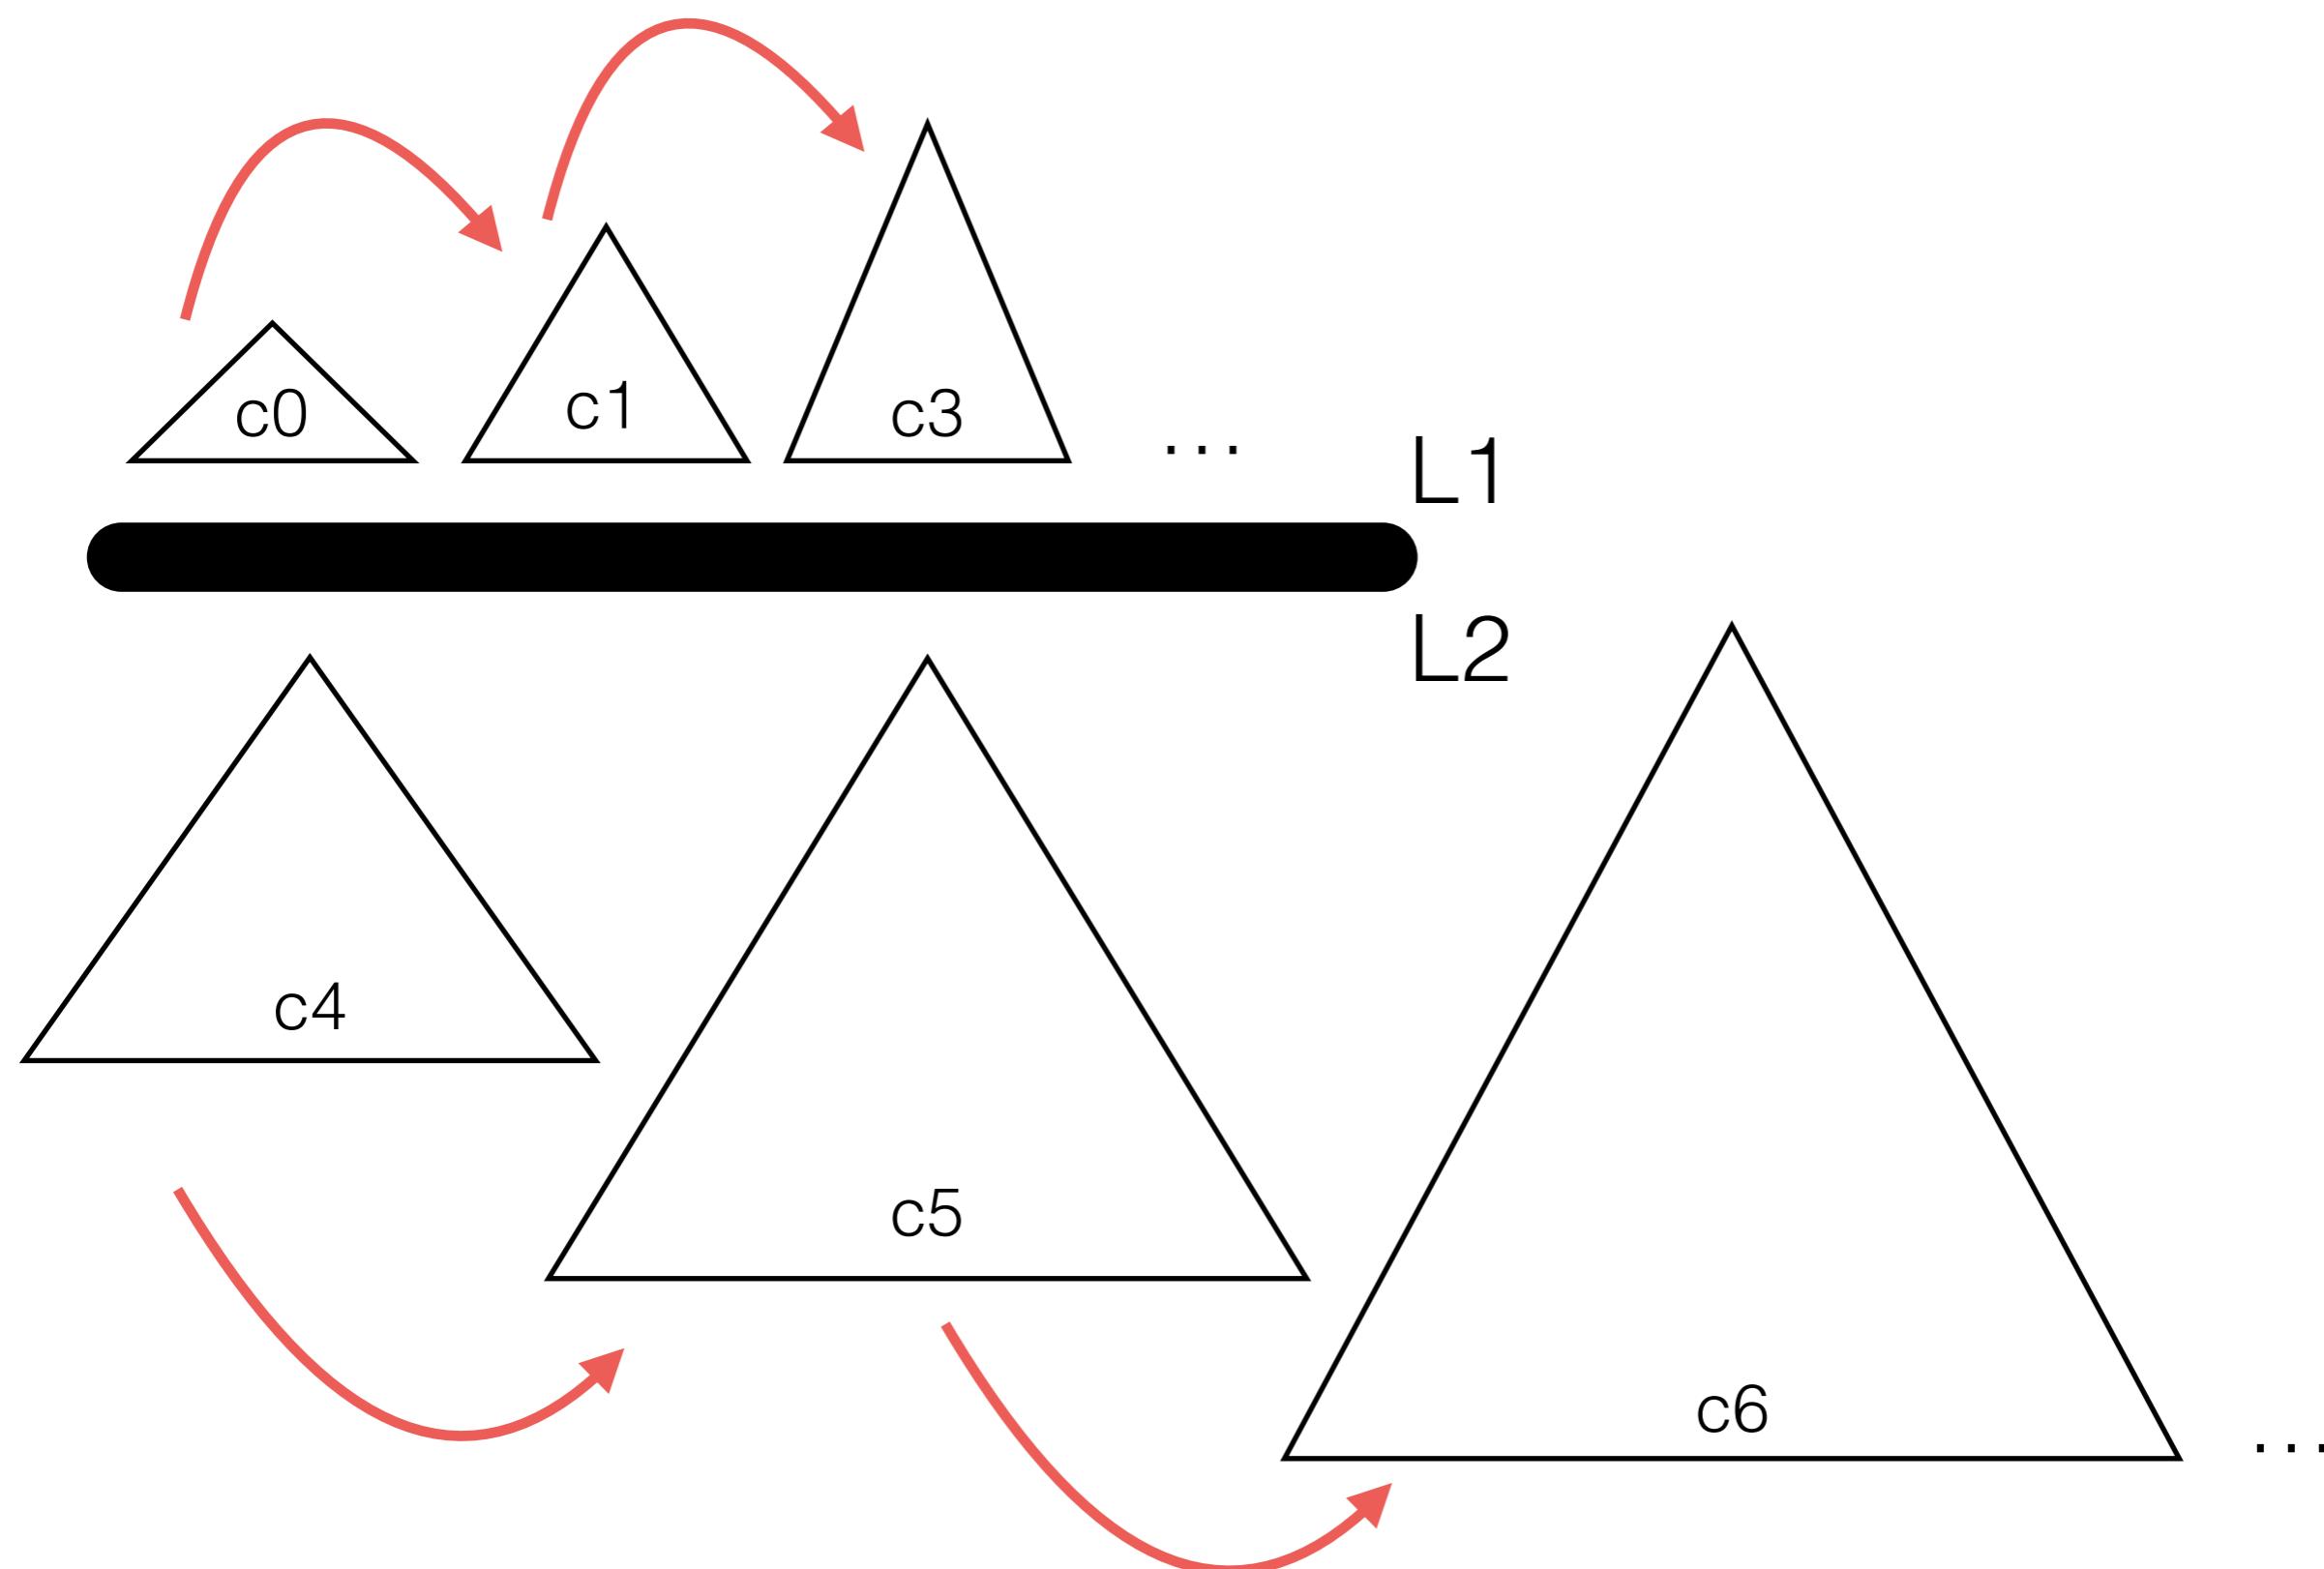
Very Large Databases Journal (**VLDBJ**), 2003

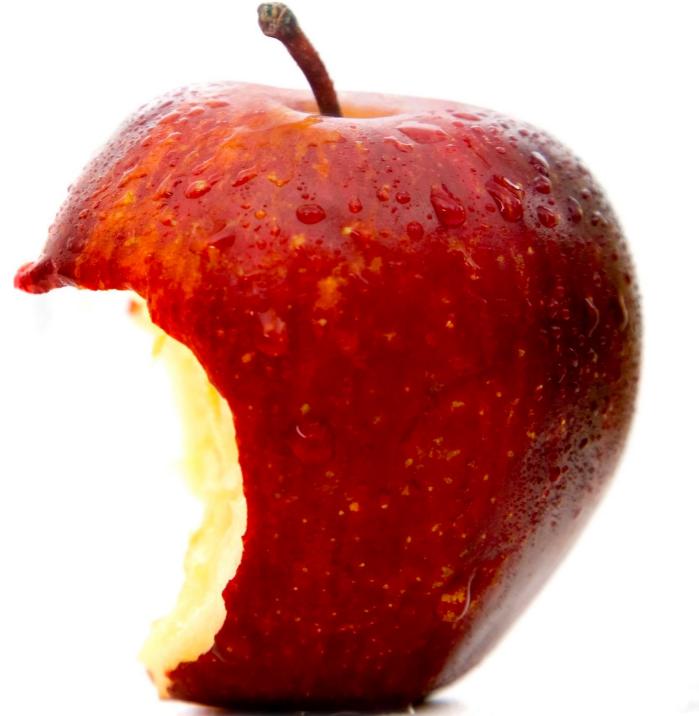


buffer





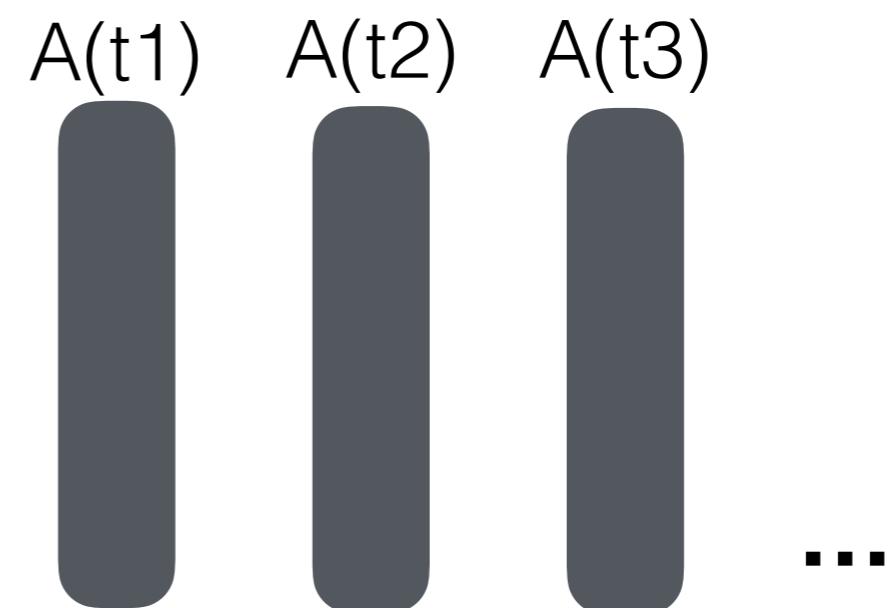




in-place updates: the cardinal sin

The Transaction Concept: Virtues and Limitations

Jim Gray, Tandem TR 81.3, 1981

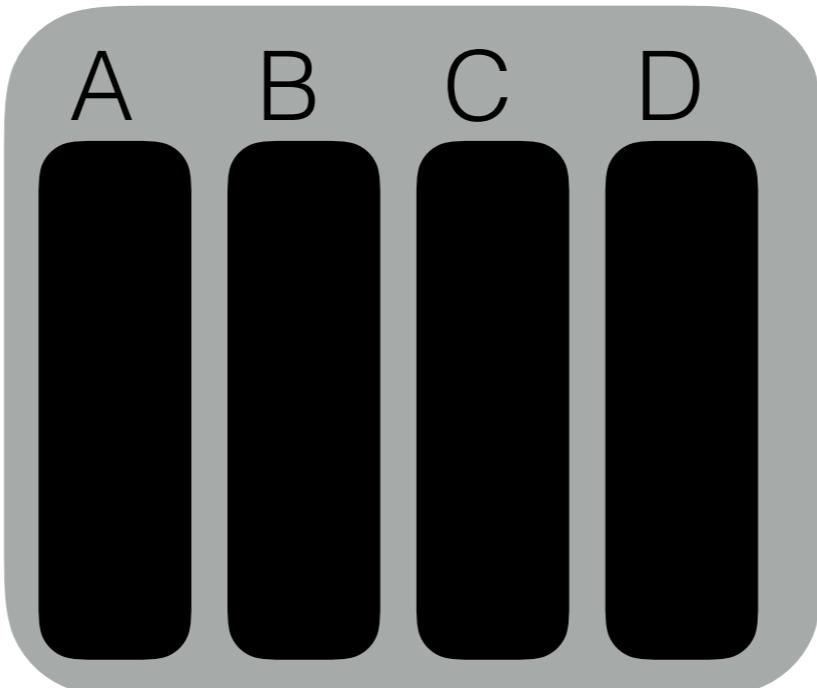


what can go wrong?

failures & >1 queries

update all rows
where A=v1 & B=v2
to (a=a/2,b=b/4,c=c-3,d=d+2)

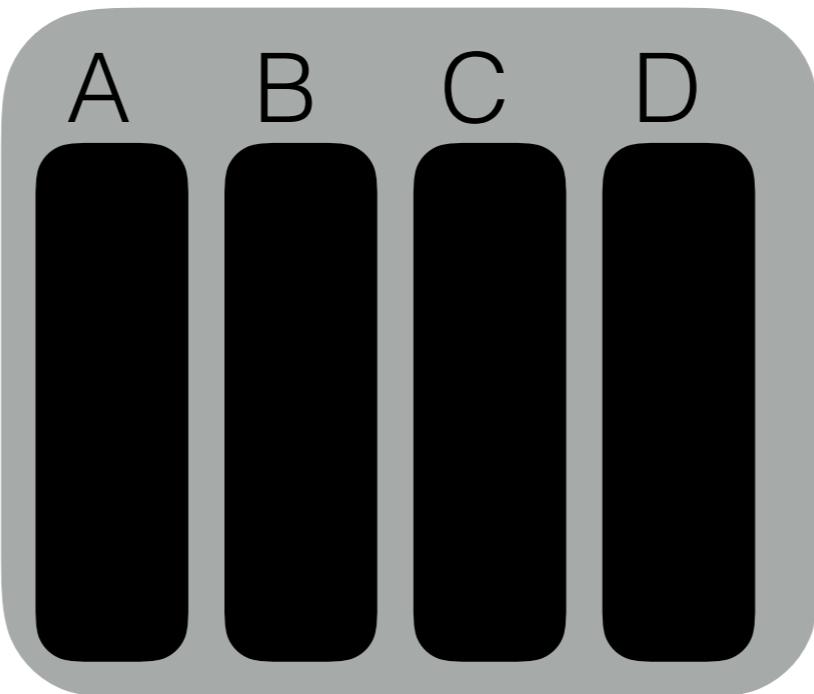
update all rows
where A=v1 & B=v2
to (a=a/2,b=b/4,c=c-3,d=d+2)



search (scan/index)
to find row to update

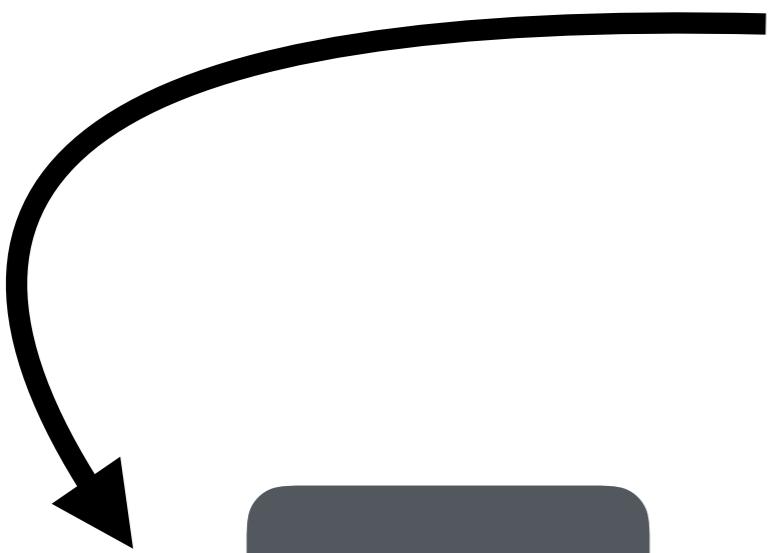
select+project actions

update all rows
where A=v1 & B=v2
to (a=a/2,b=b/4,c=c-3,d=d+2)

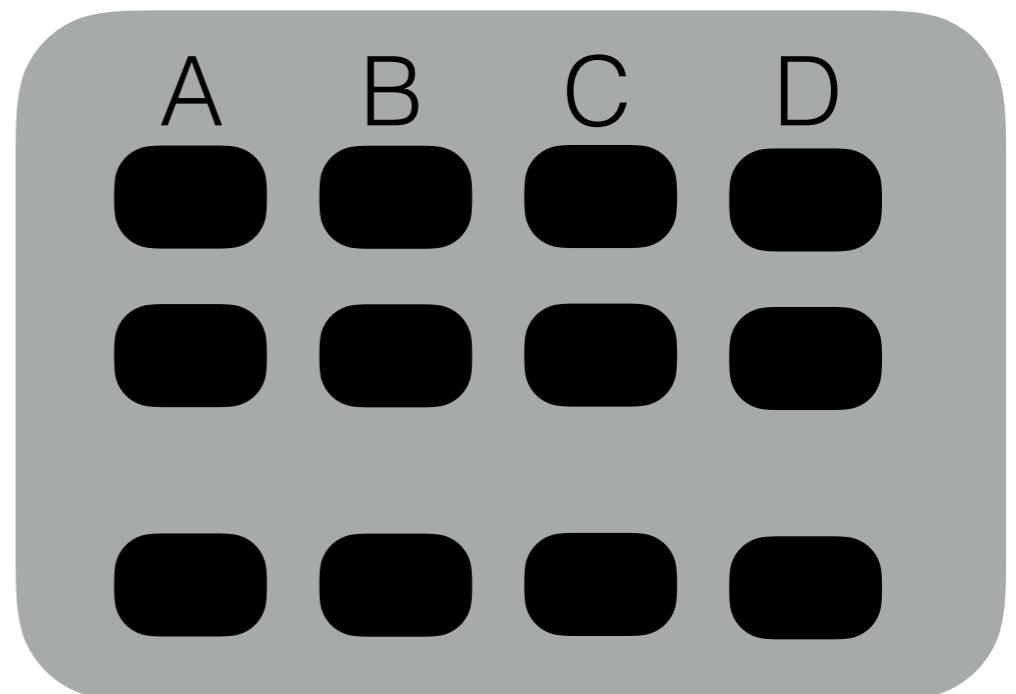
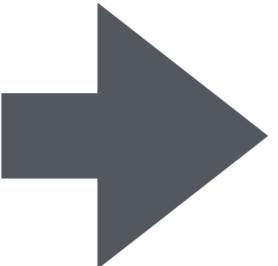


search (scan/index)
to find row to update

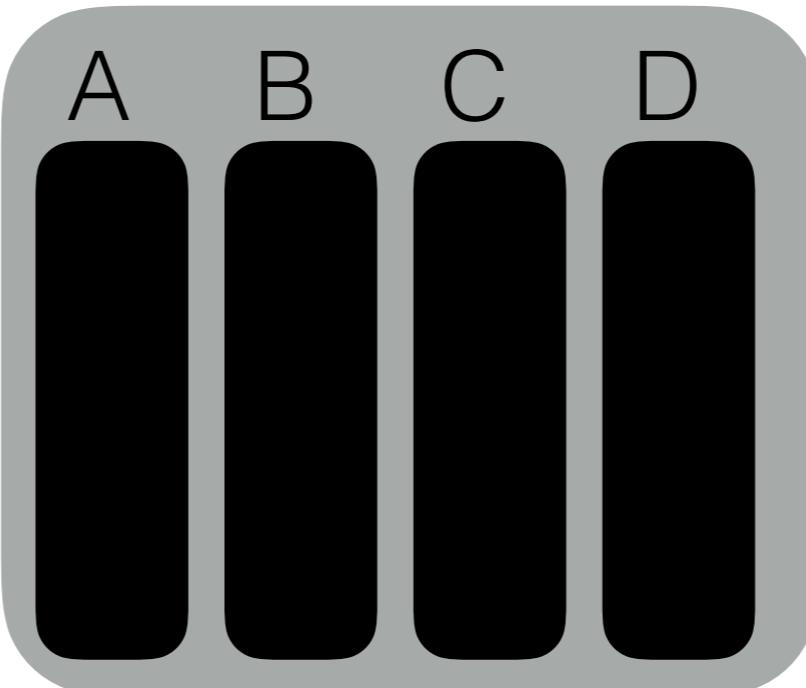
select+project actions



list of rowIDs (positions)
(sort)



update all rows
where A=v1 & B=v2
to (a=a/2,b=b/4,c=c-3,d=d+2)

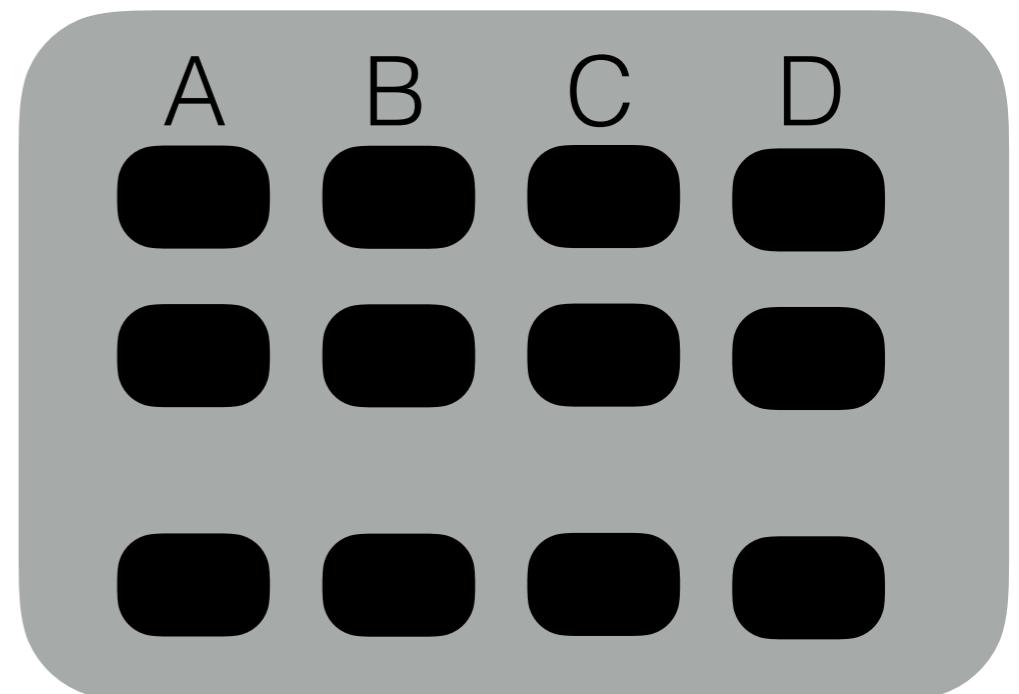
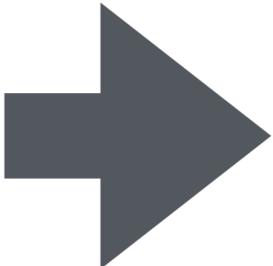


search (scan/index)
to find row to update

select+project actions

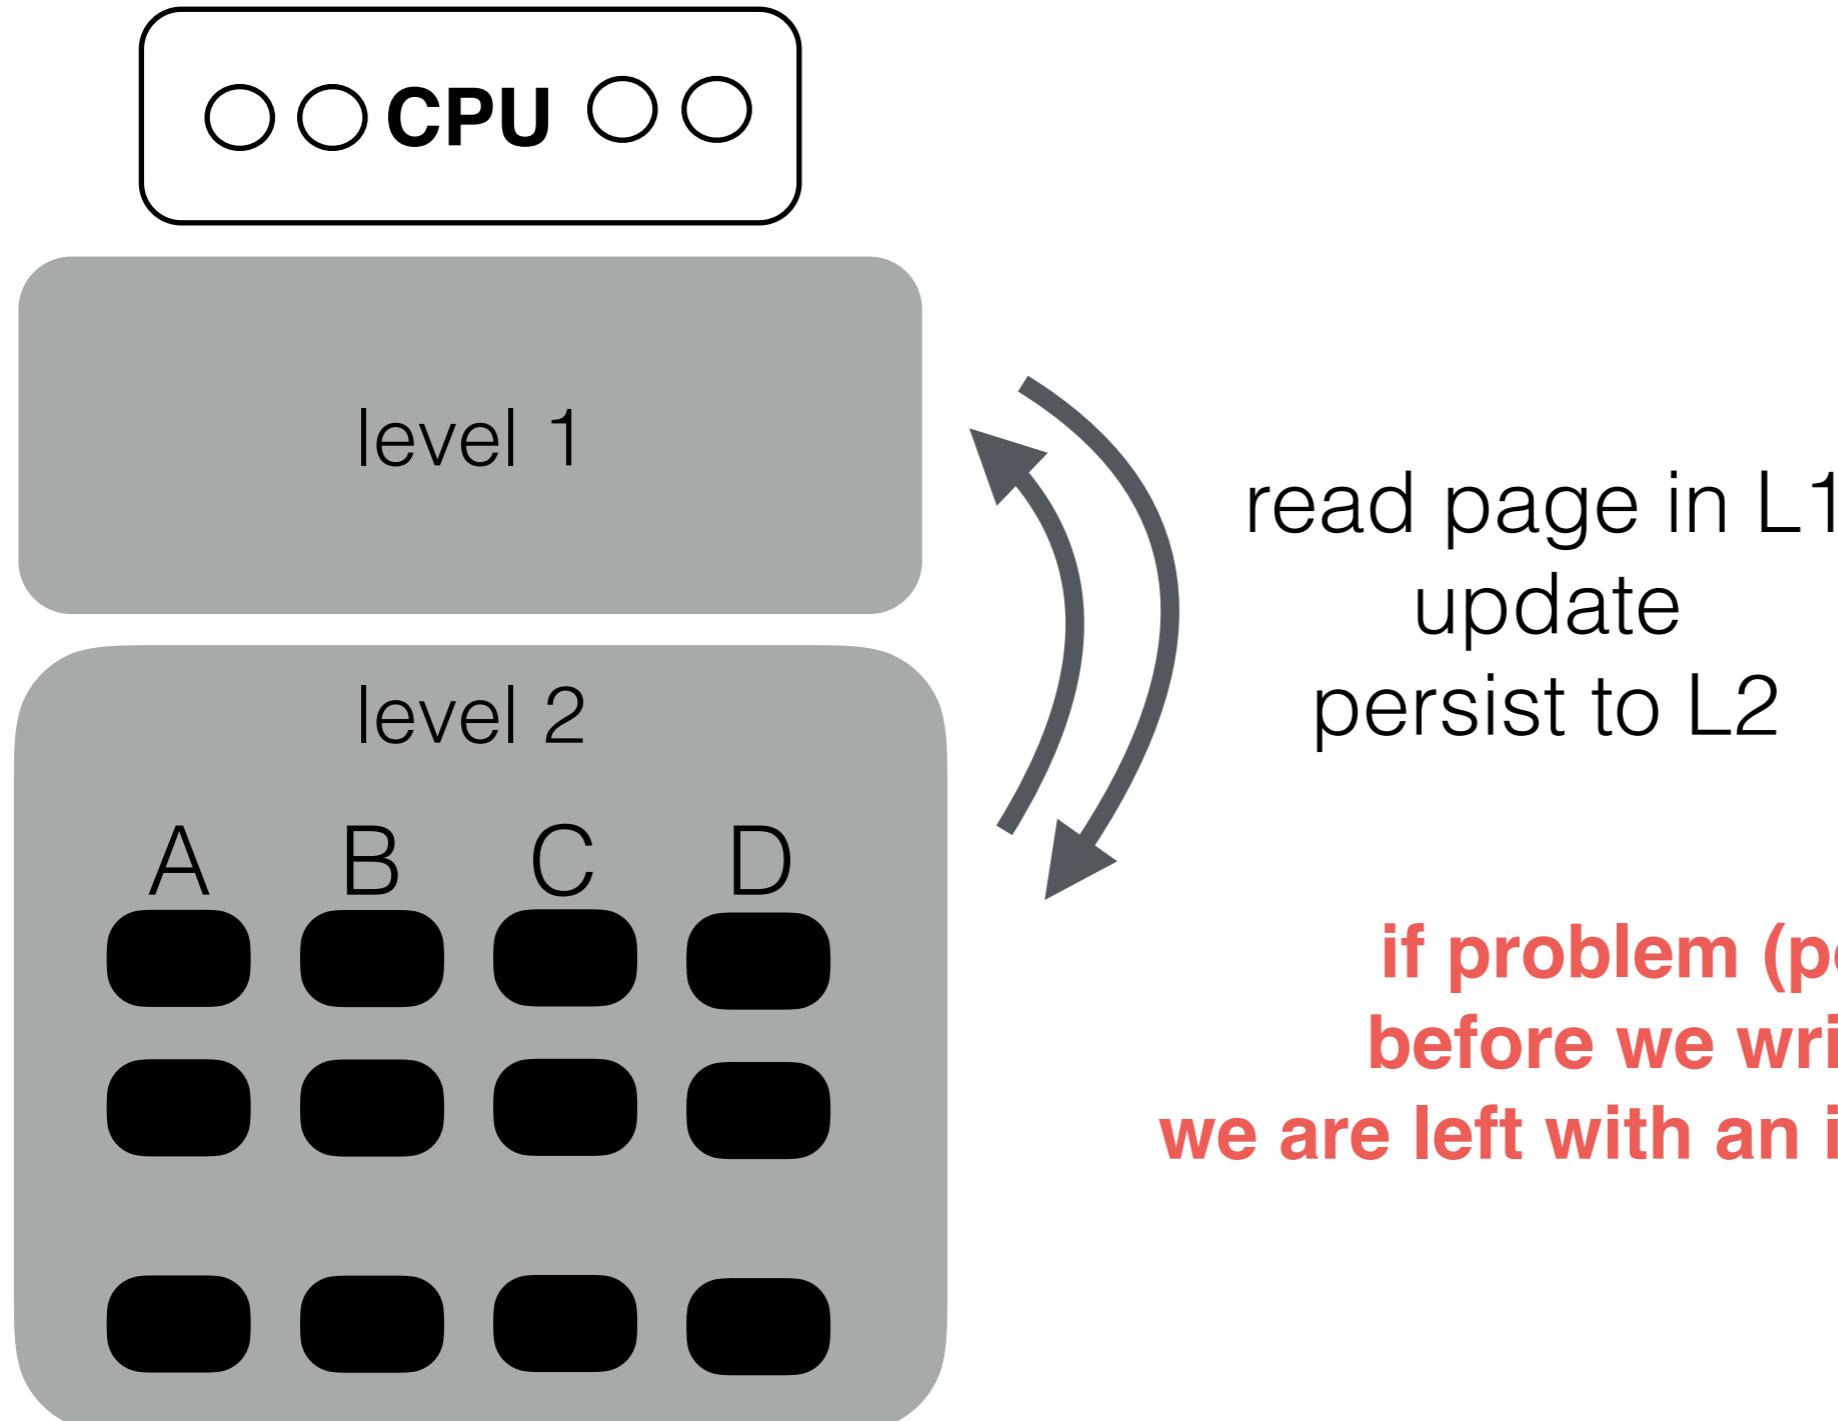


list of rowIDs (positions)
(sort)



we know what to update but nothing happened yet

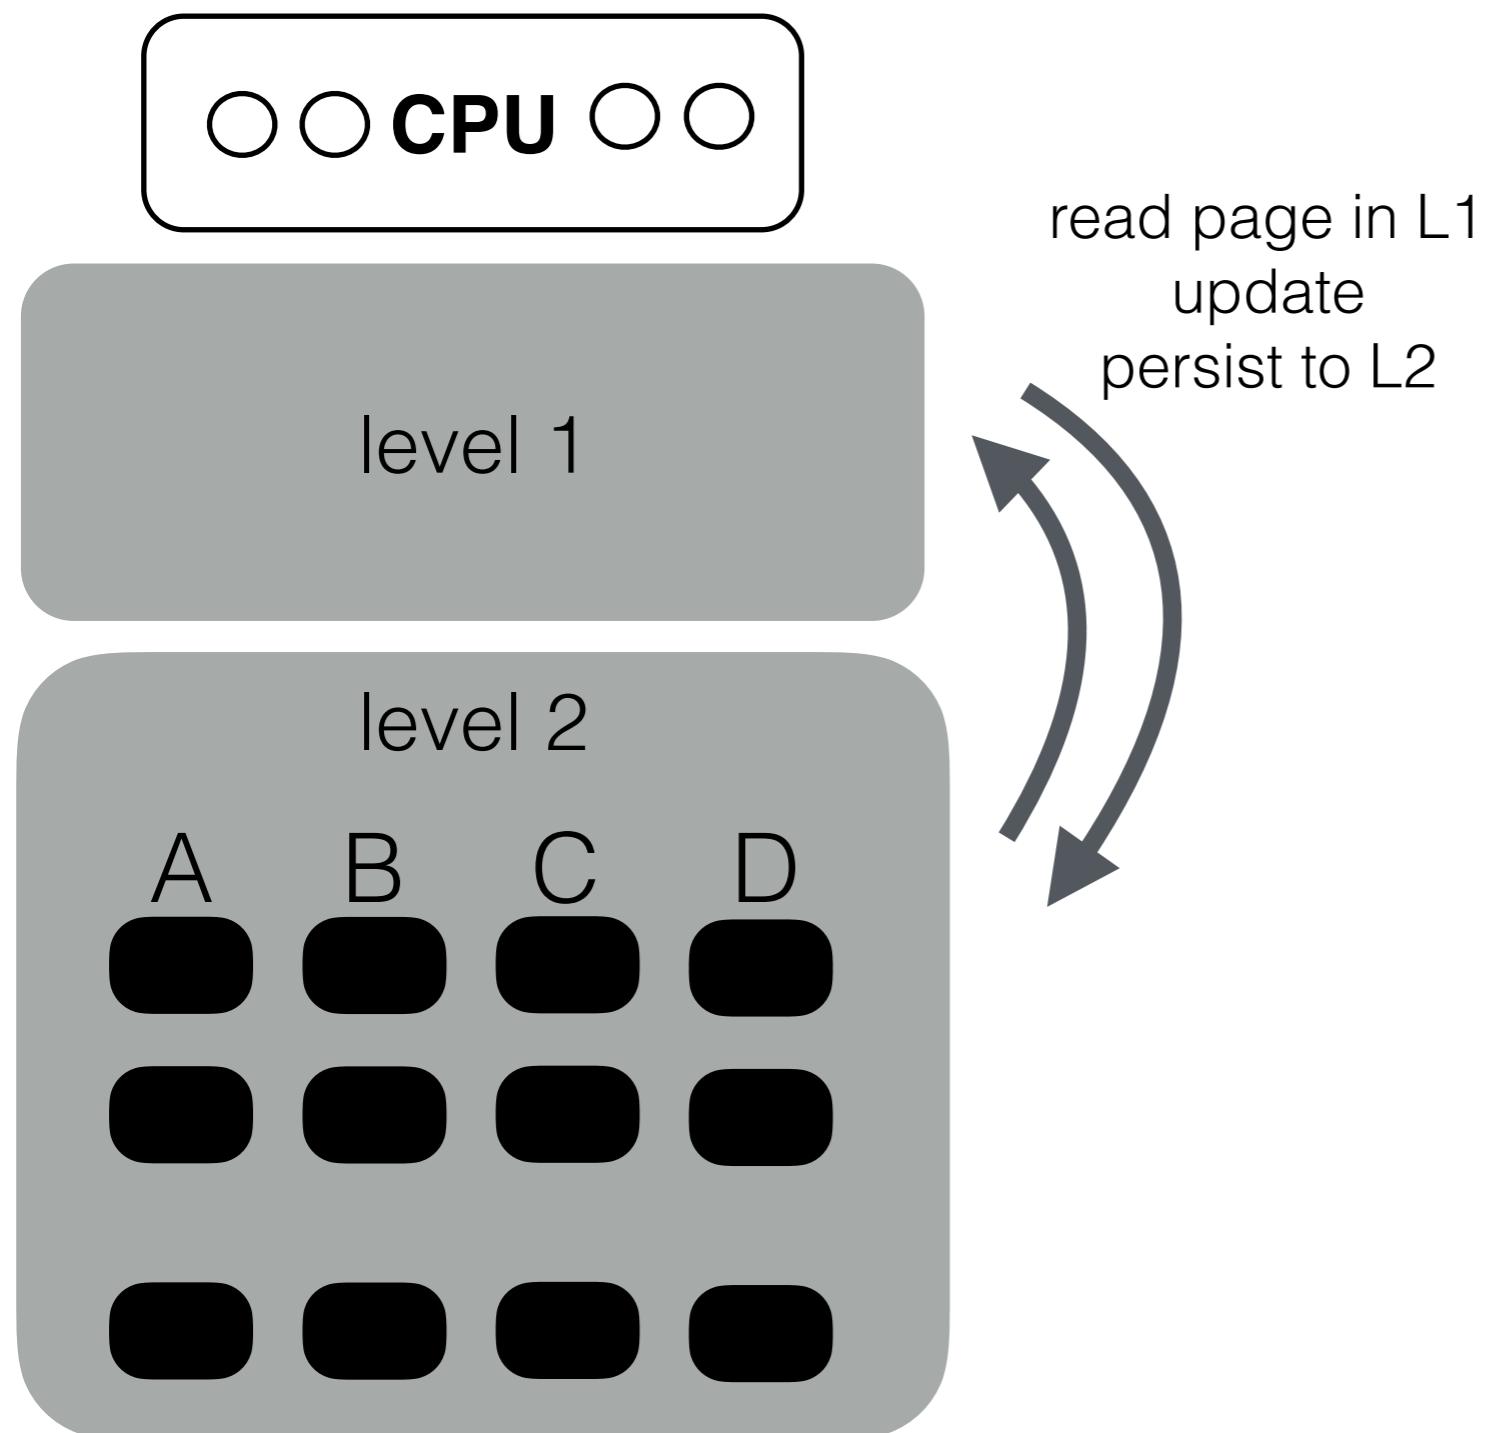
update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)



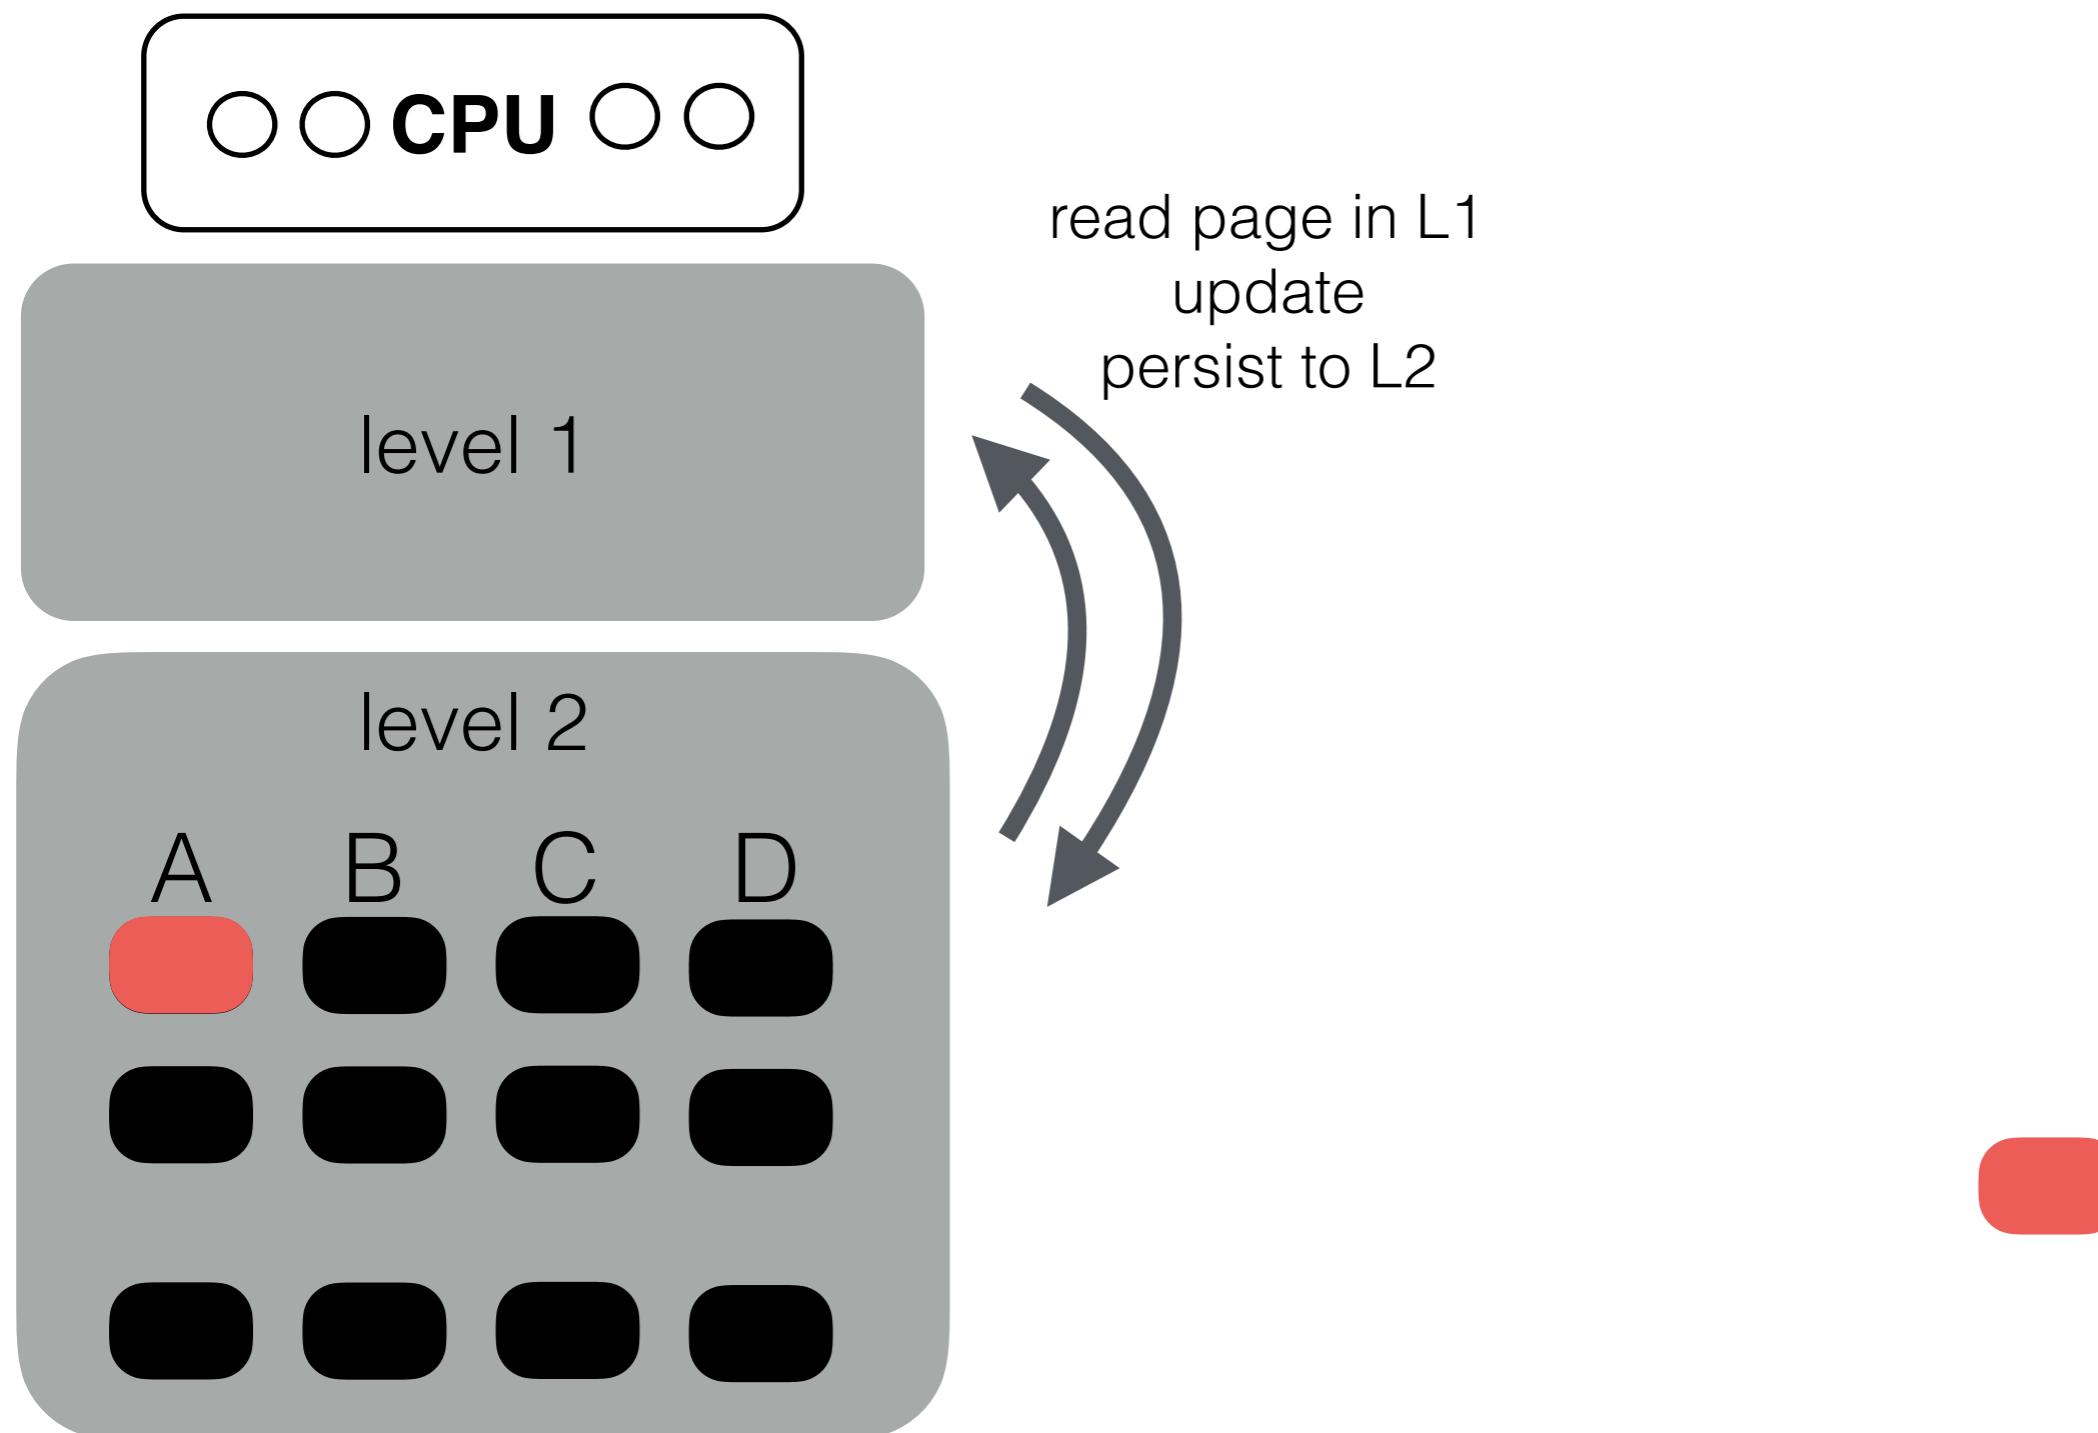
WAL: keep persistent notes as we go so we can resume or undo



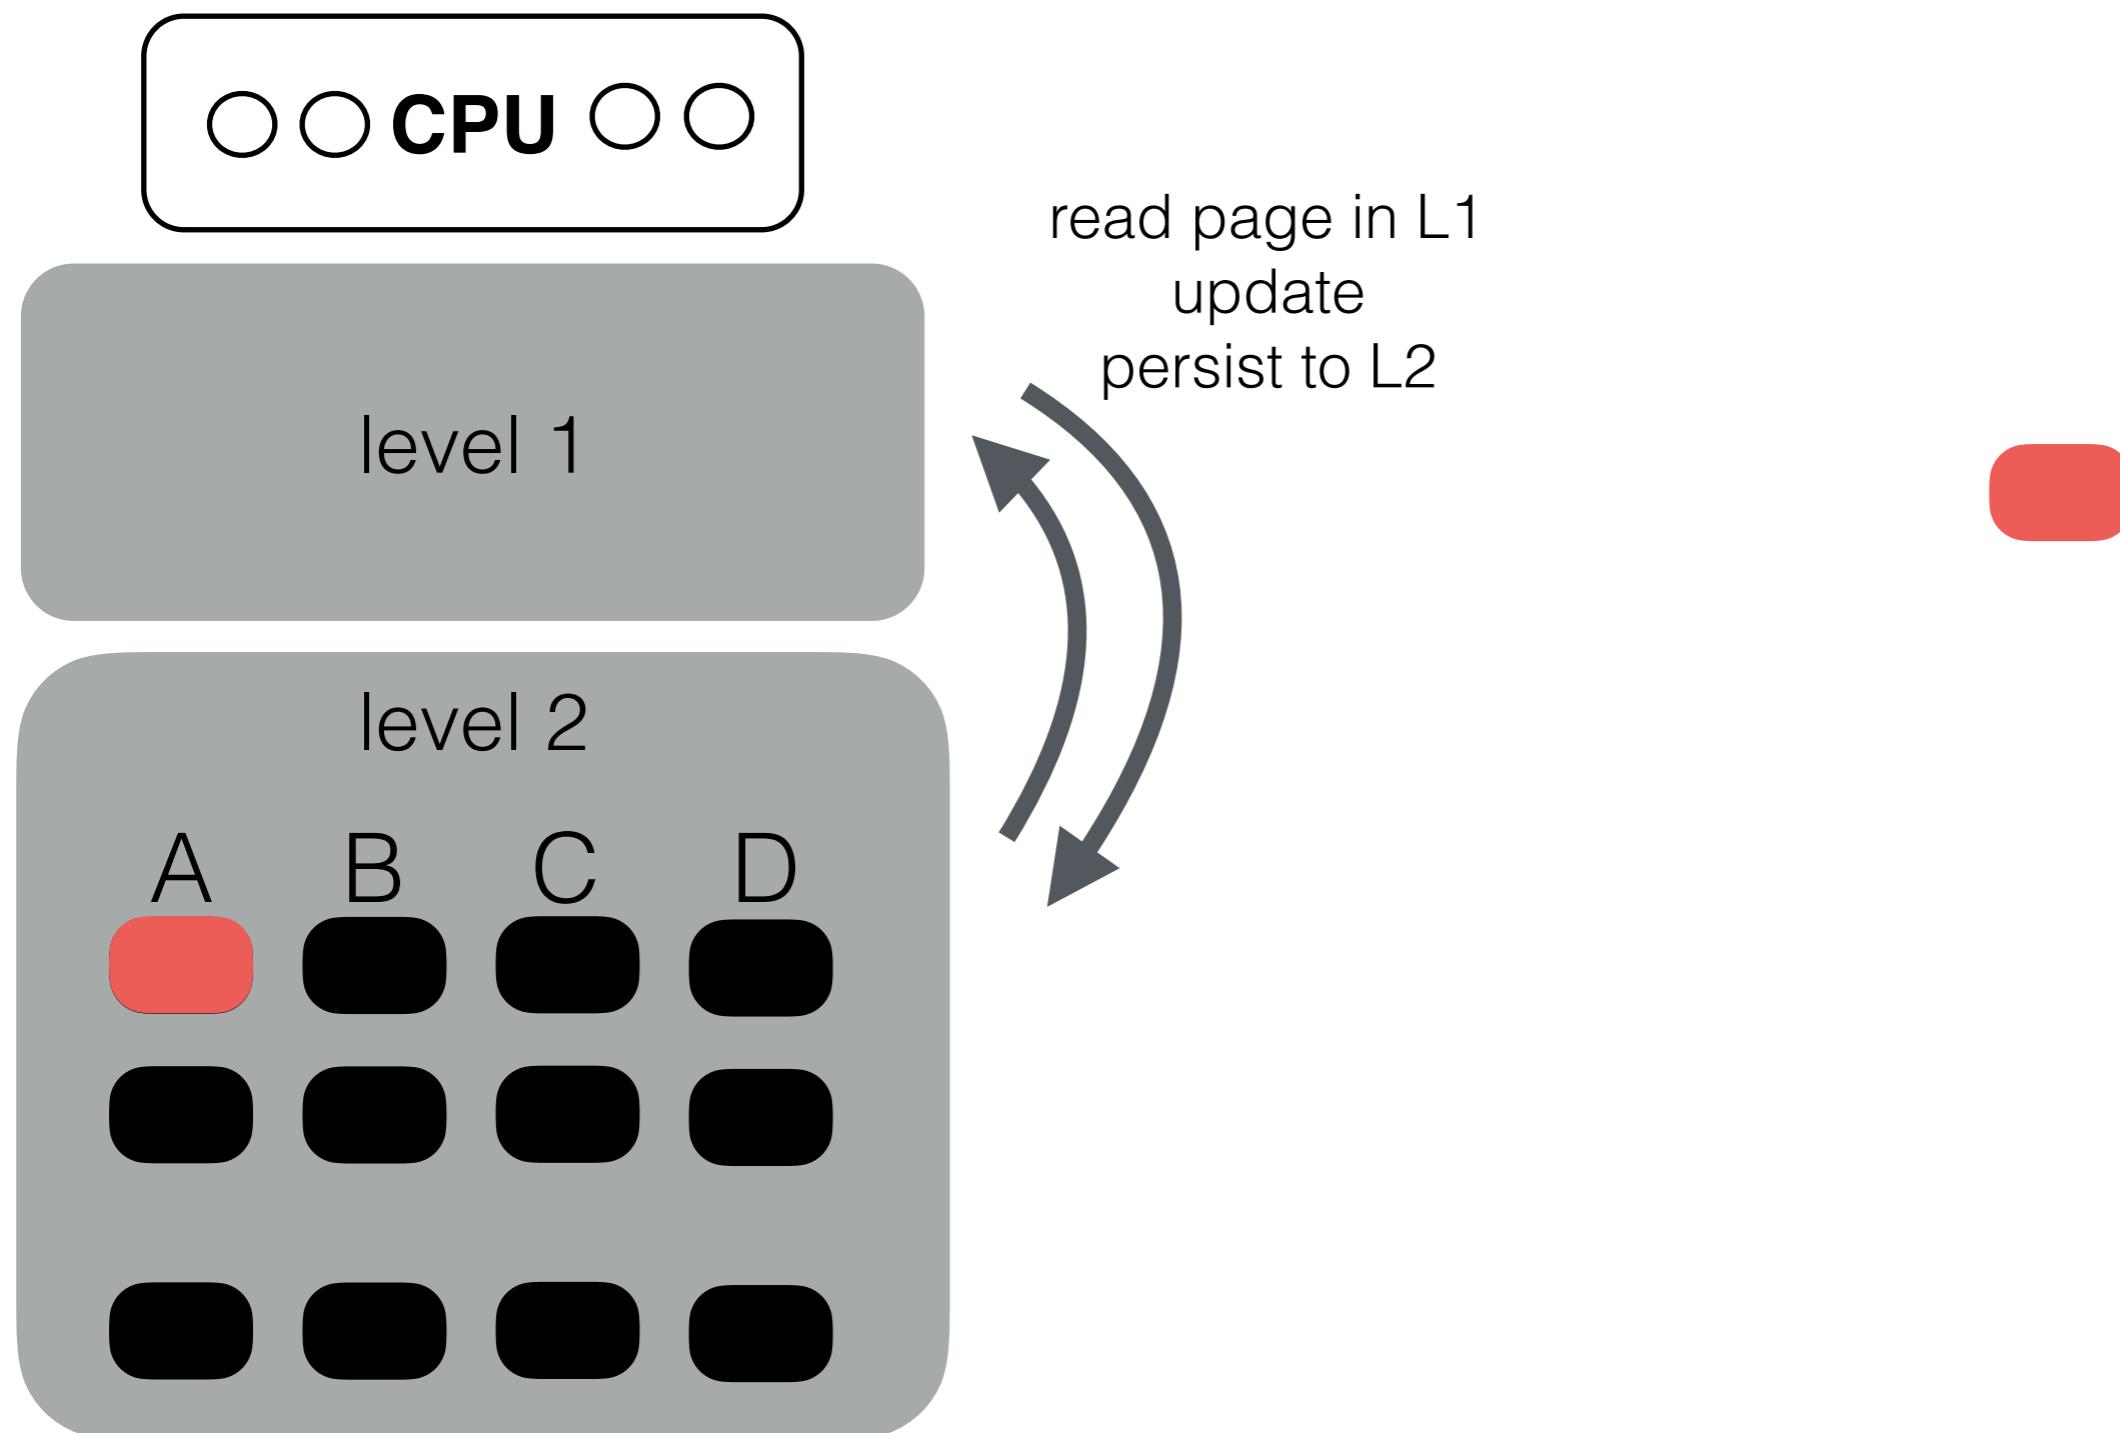
update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)



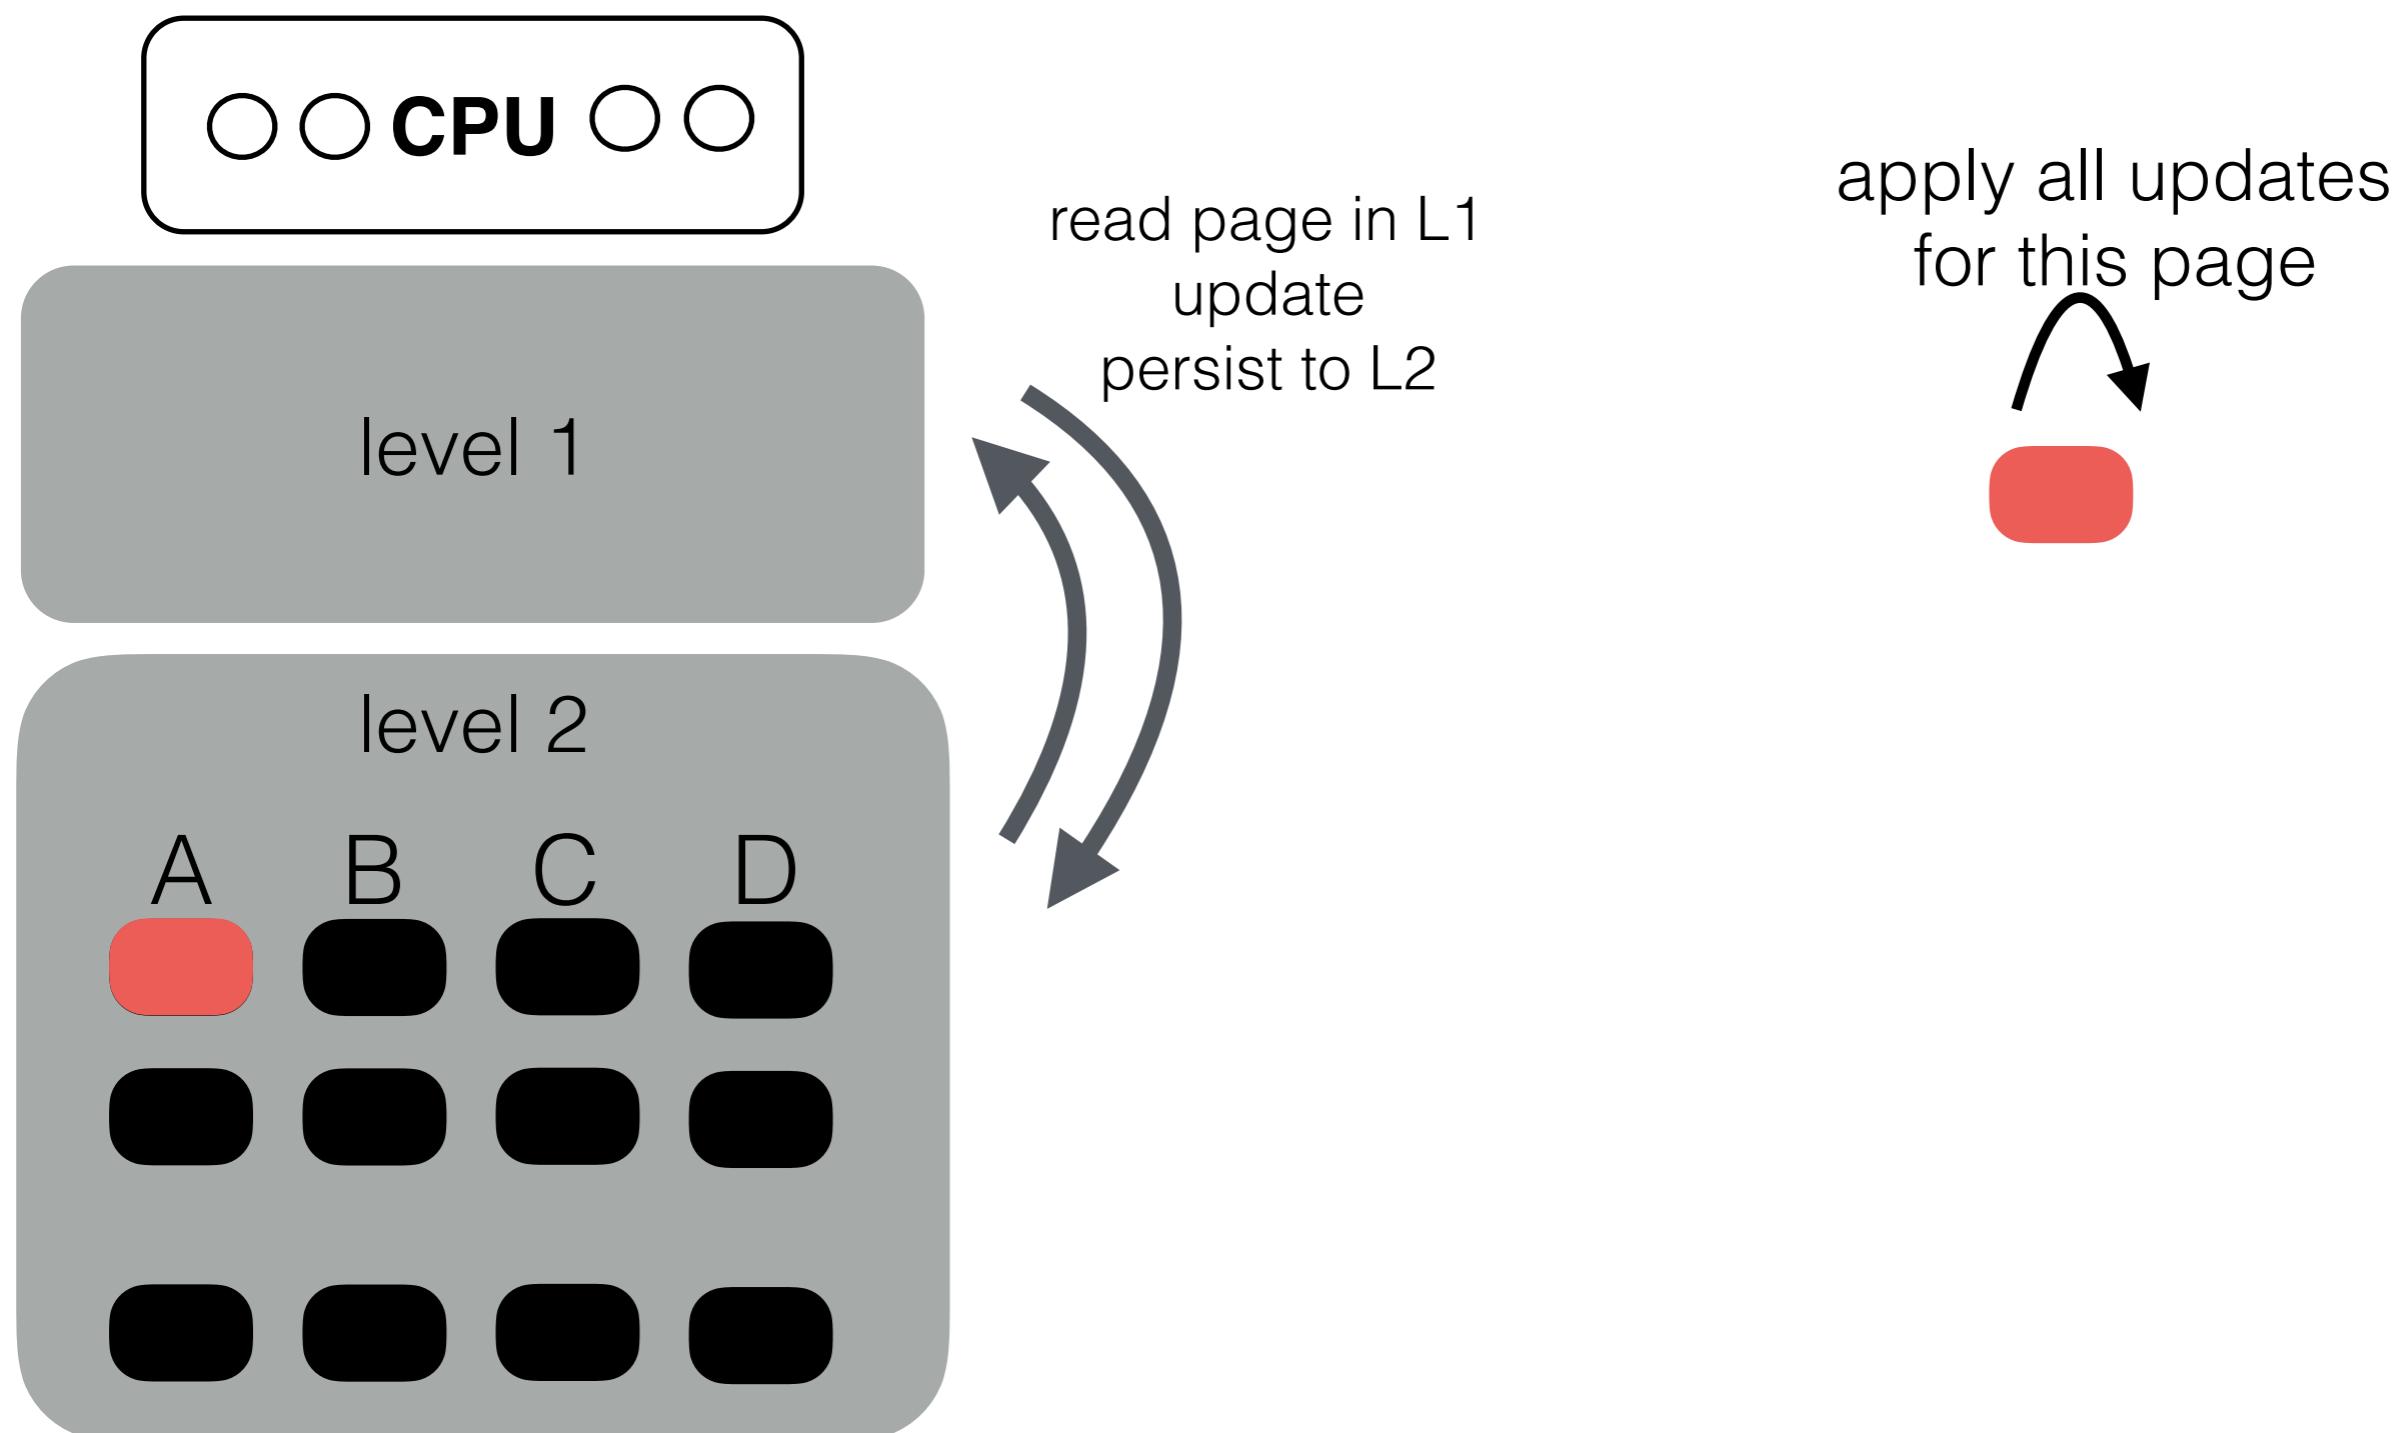
update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)



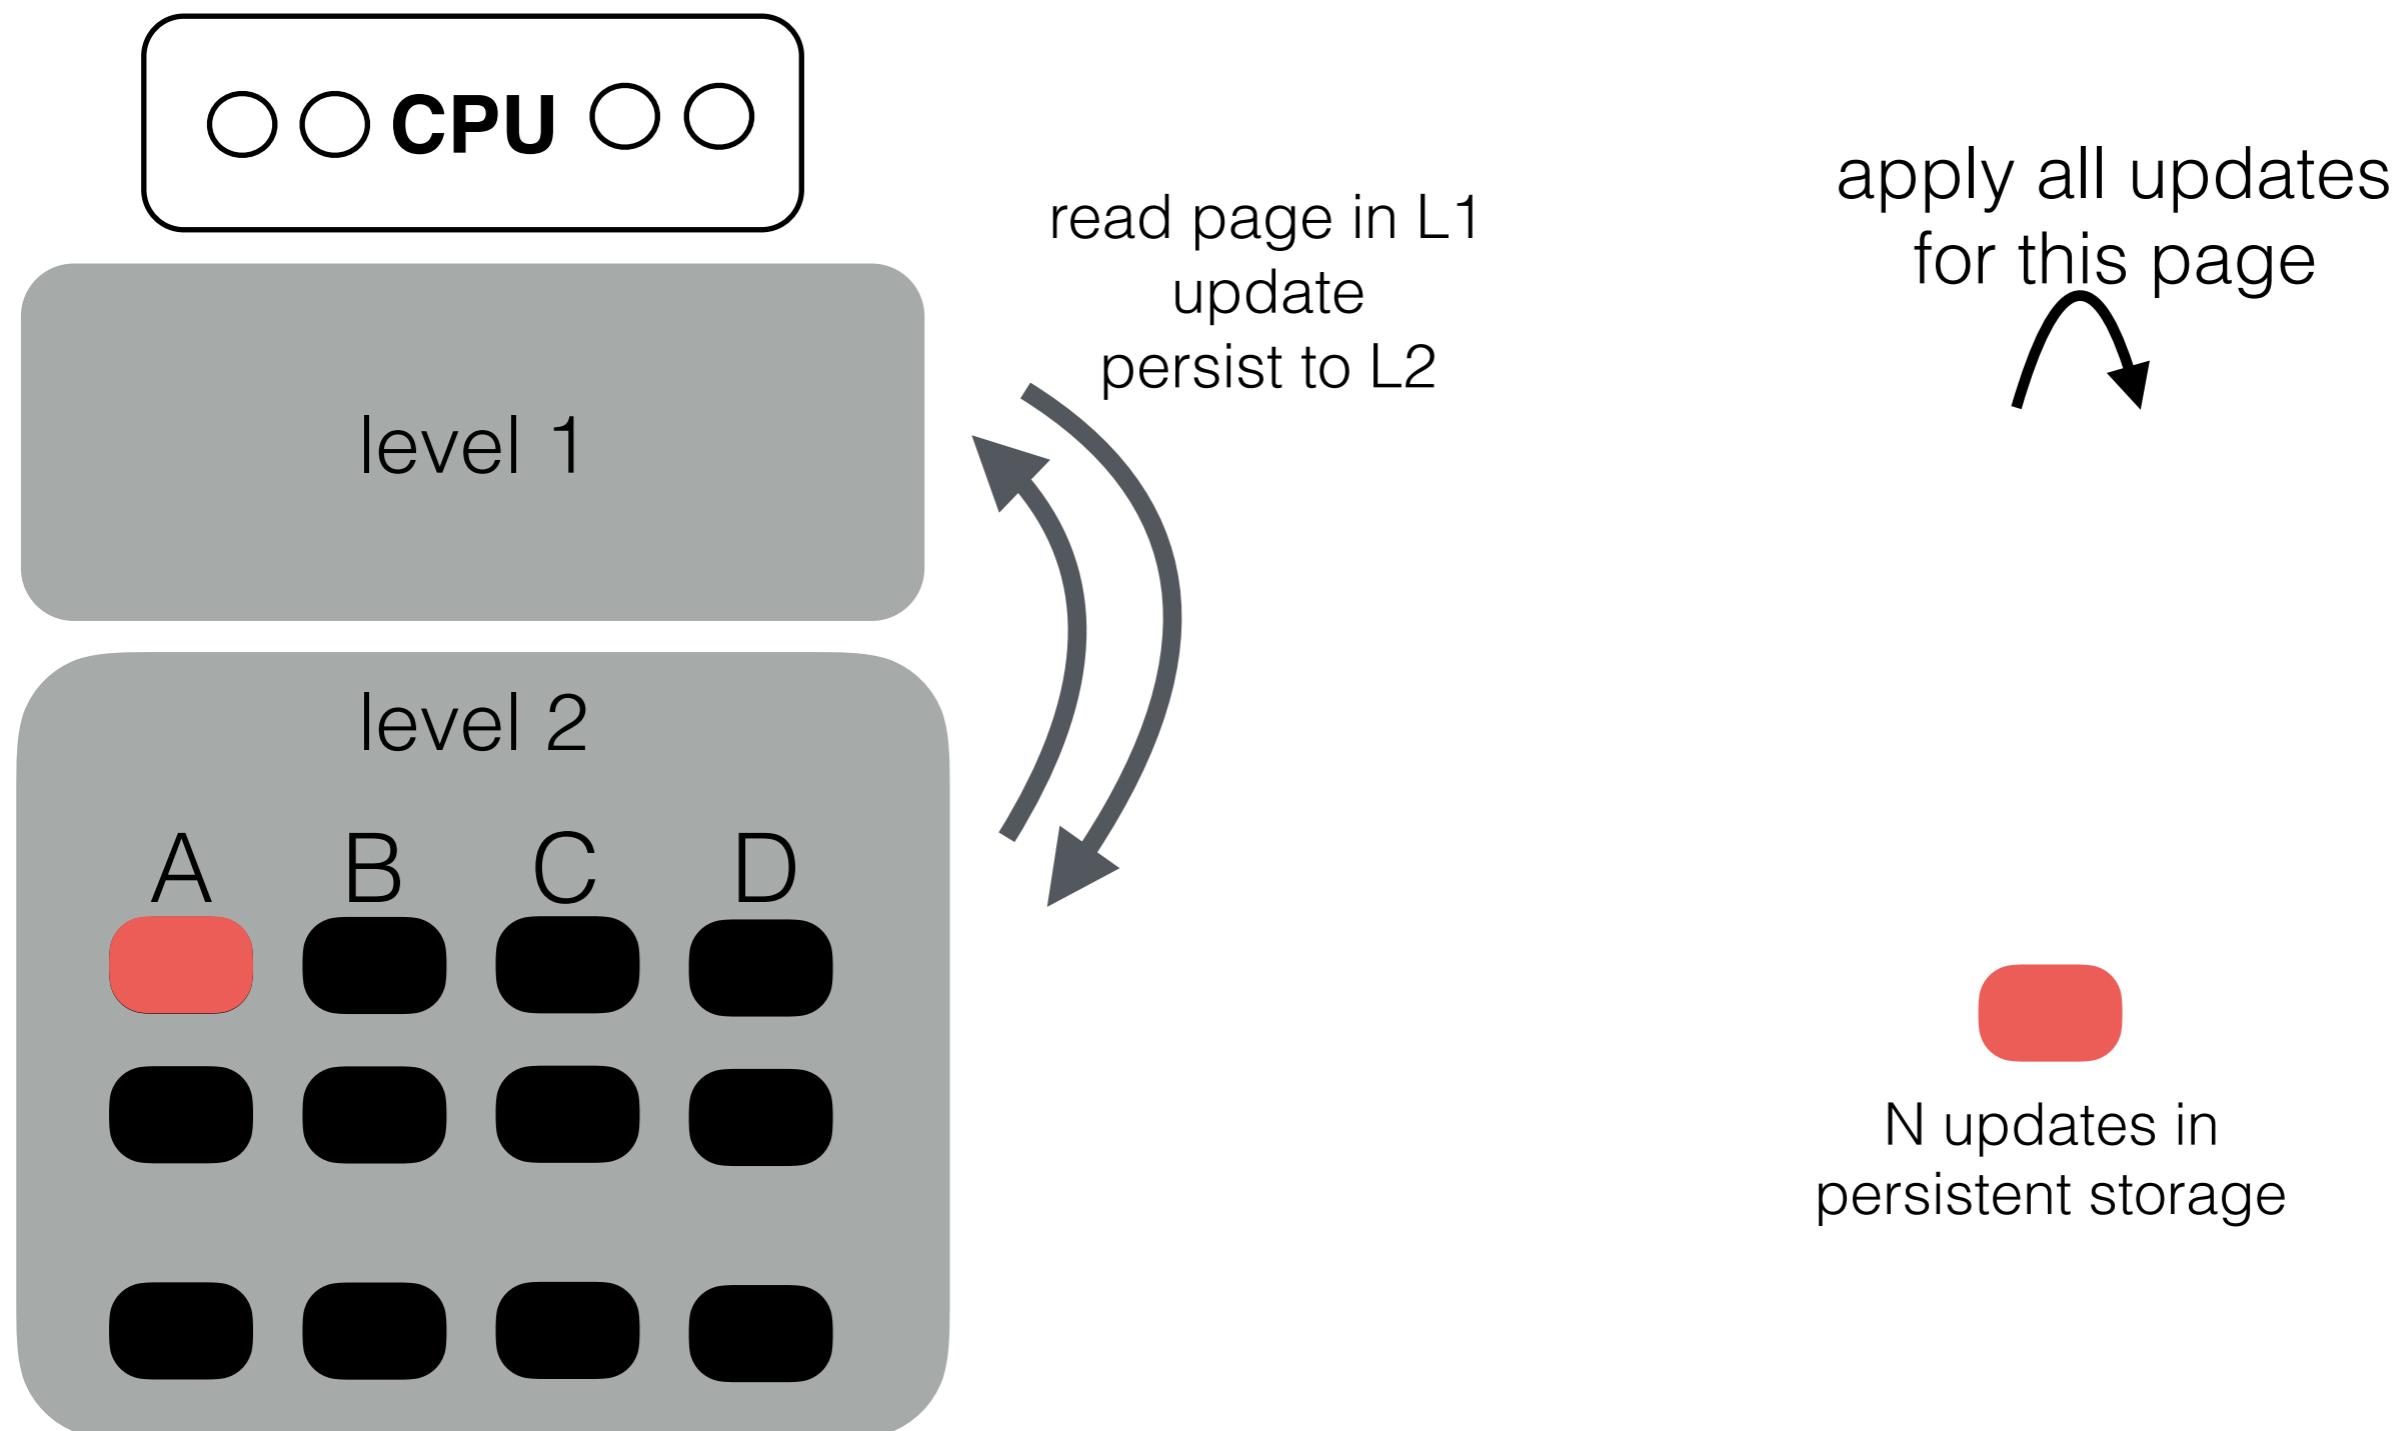
update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)



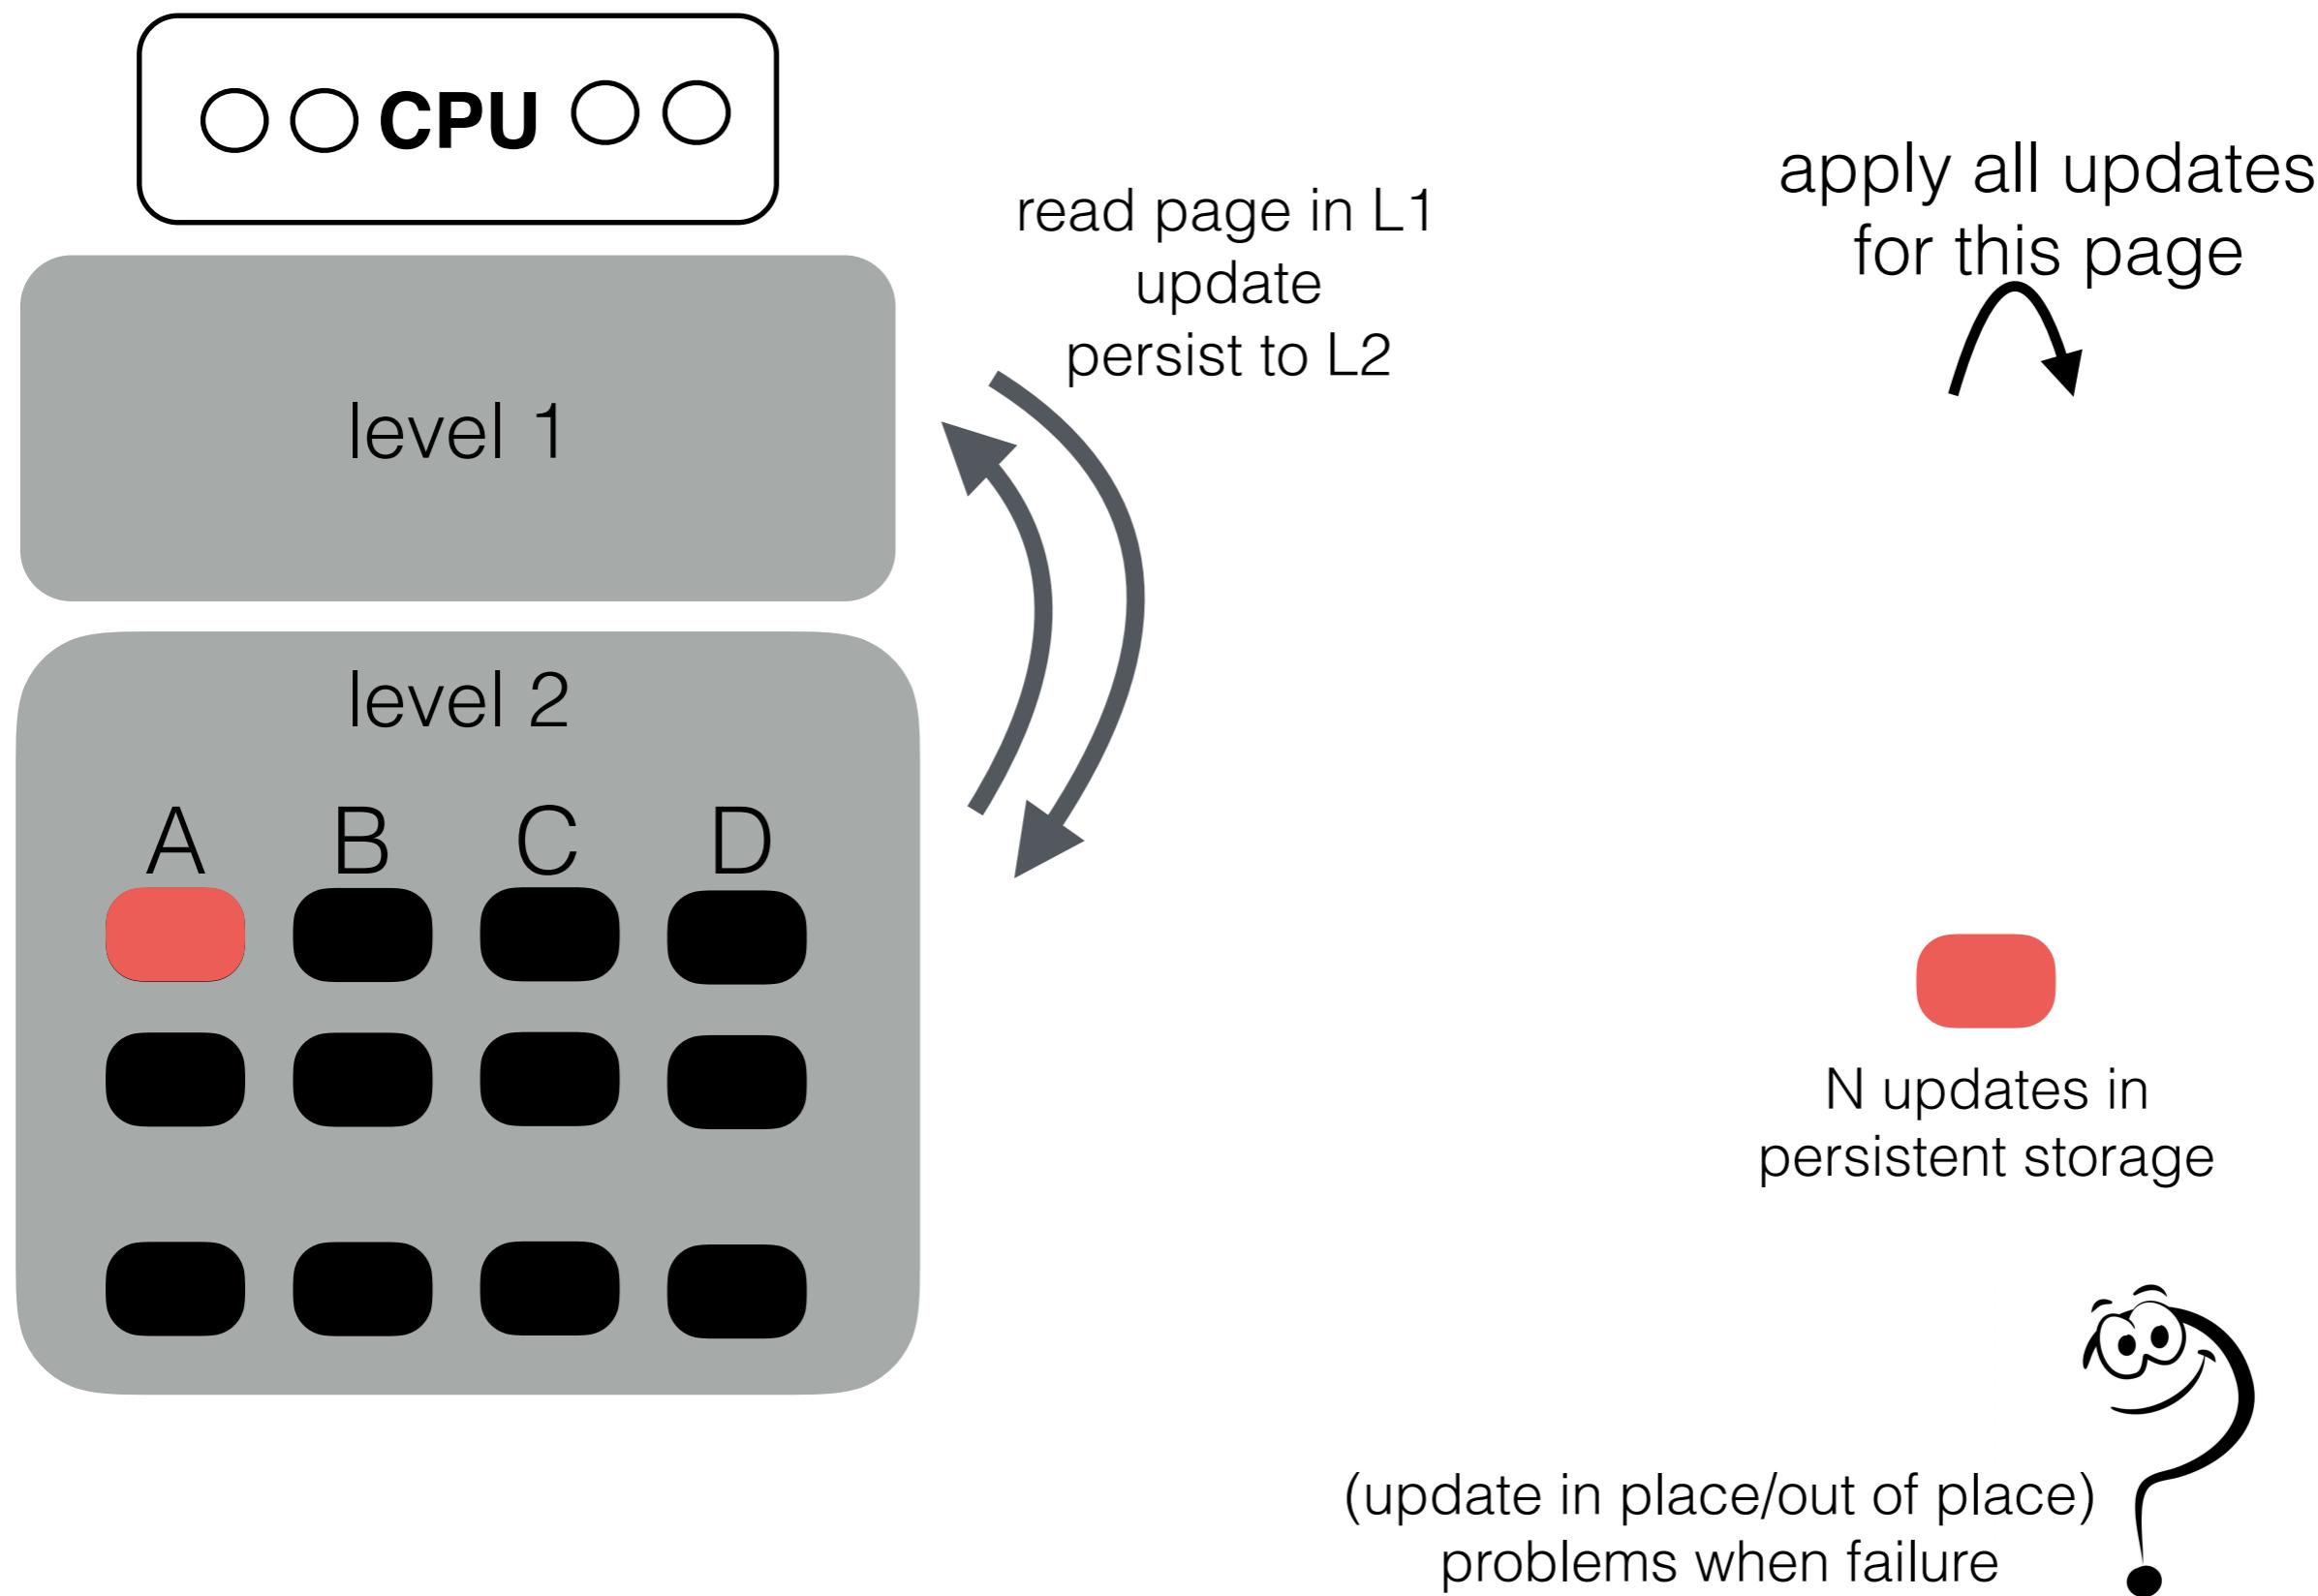
update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)



update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)



update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)

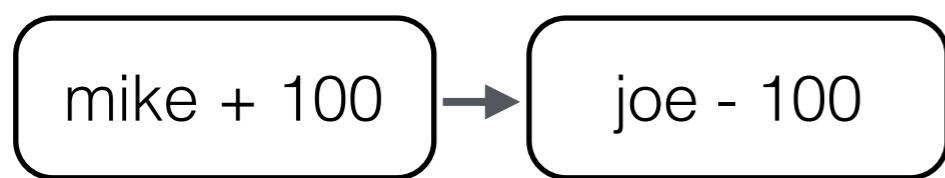
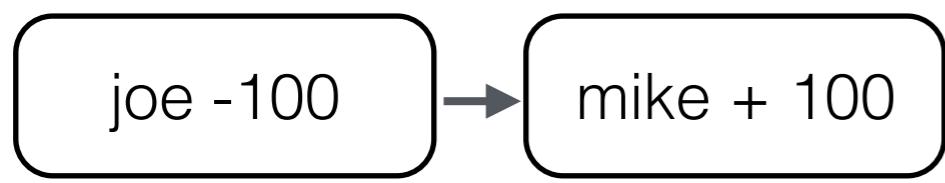


classic example

joe owes mike 100\$

both joe and mike have a Bank of Bla account

possible actions

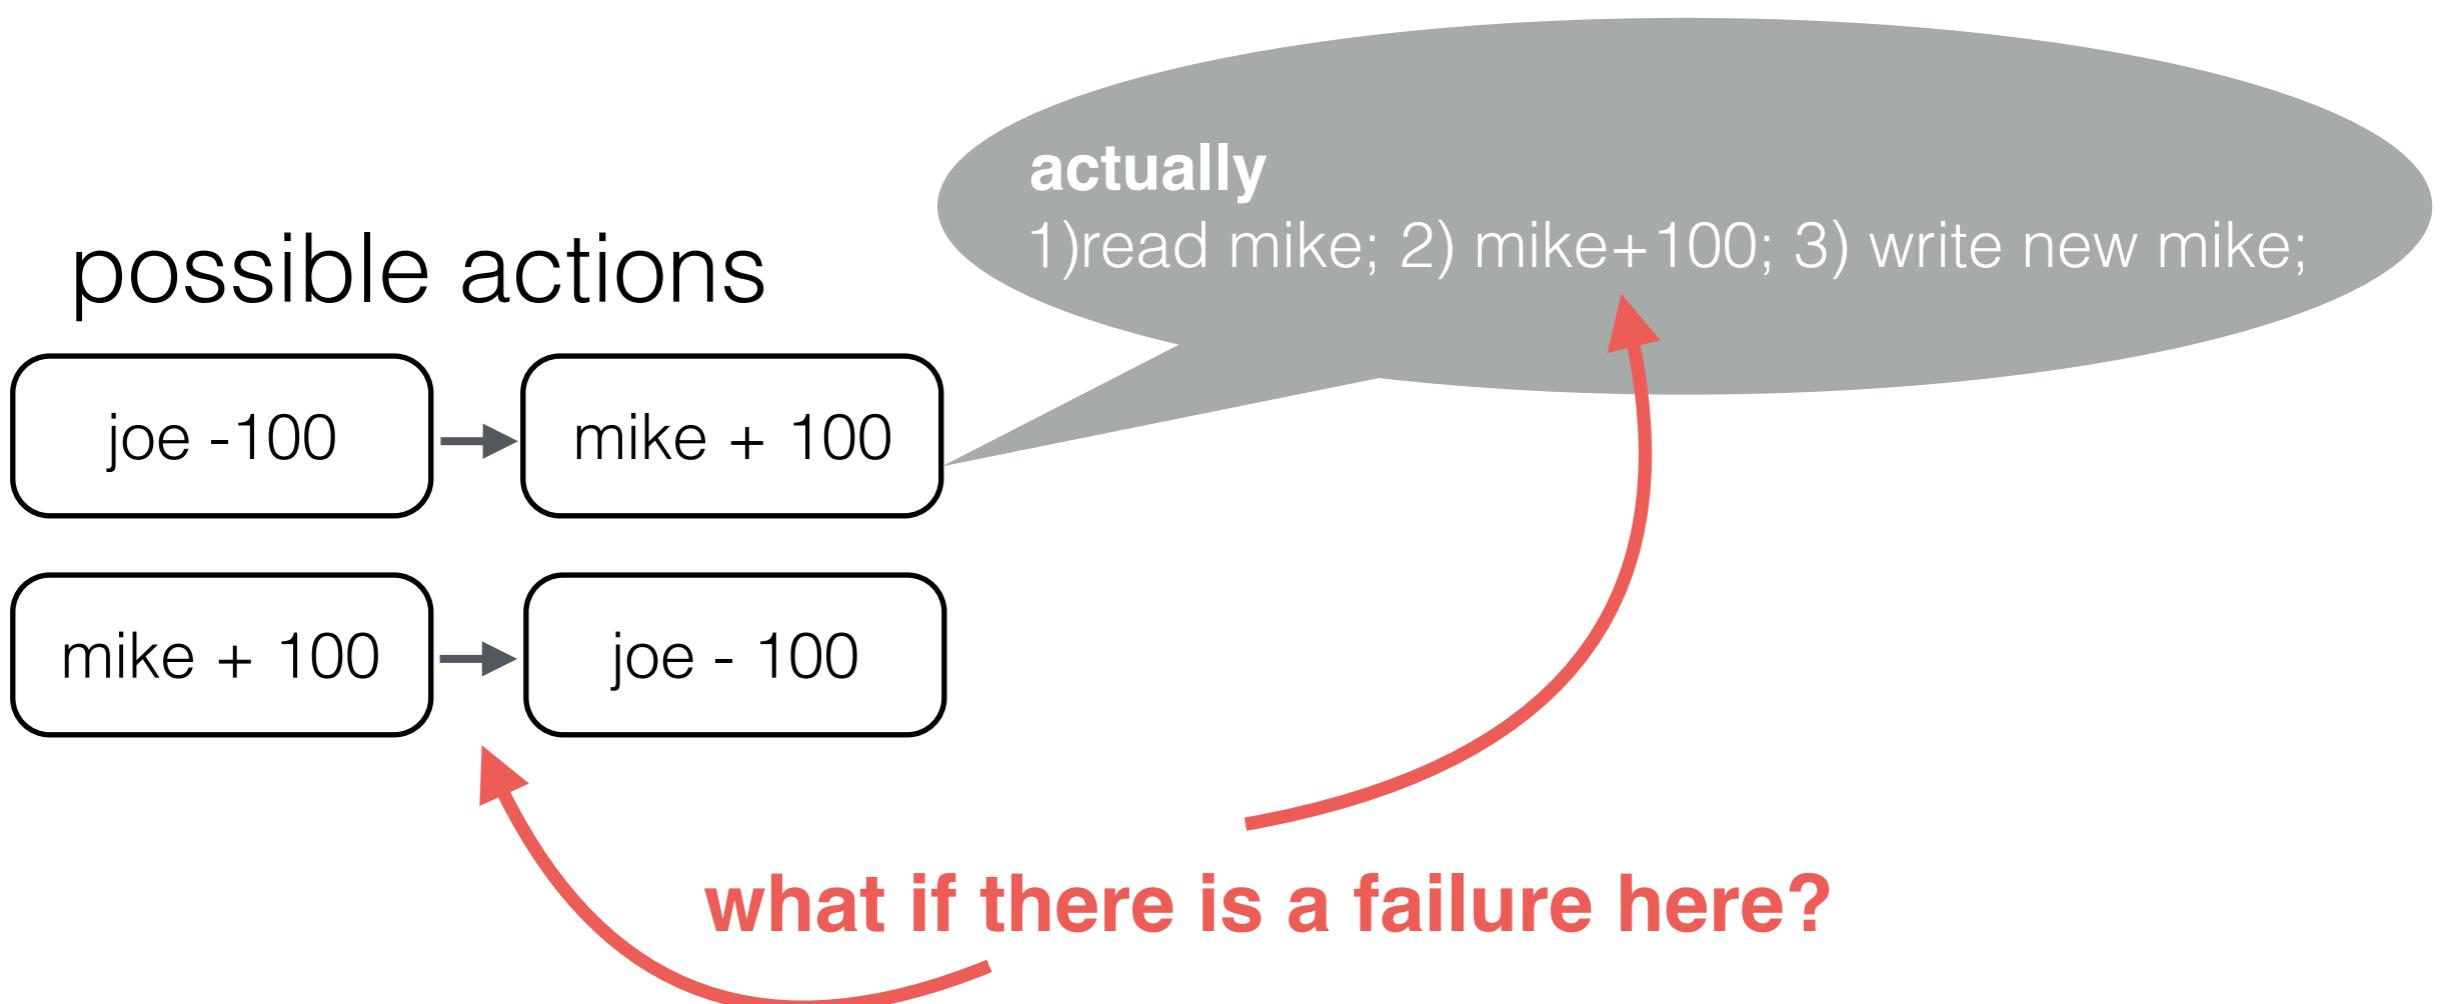


what if there is a failure here?

classic example

joe owes mike 100\$

both joe and mike have a Bank of Bla account



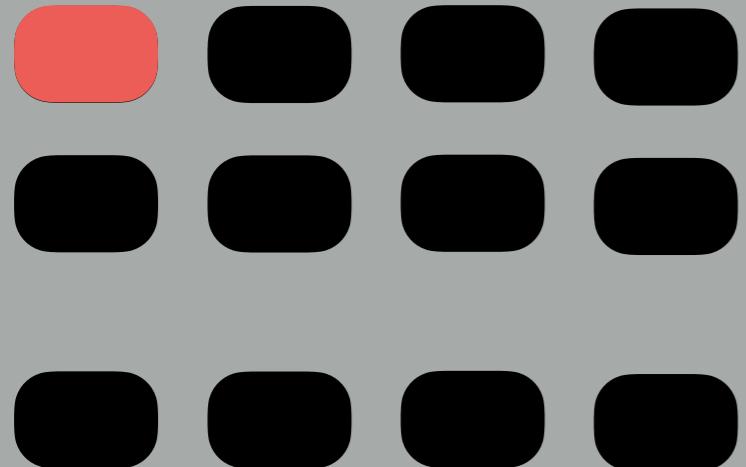
update all rows where A=v1 & B=v2 to (a=a/2,b=b/4,c=c-3,d=d+2)

○ ○ **CPU** ○ ○

level 1

level 2

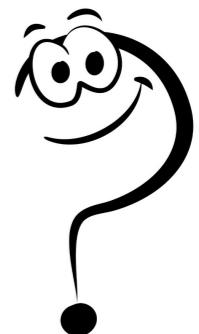
A B C D



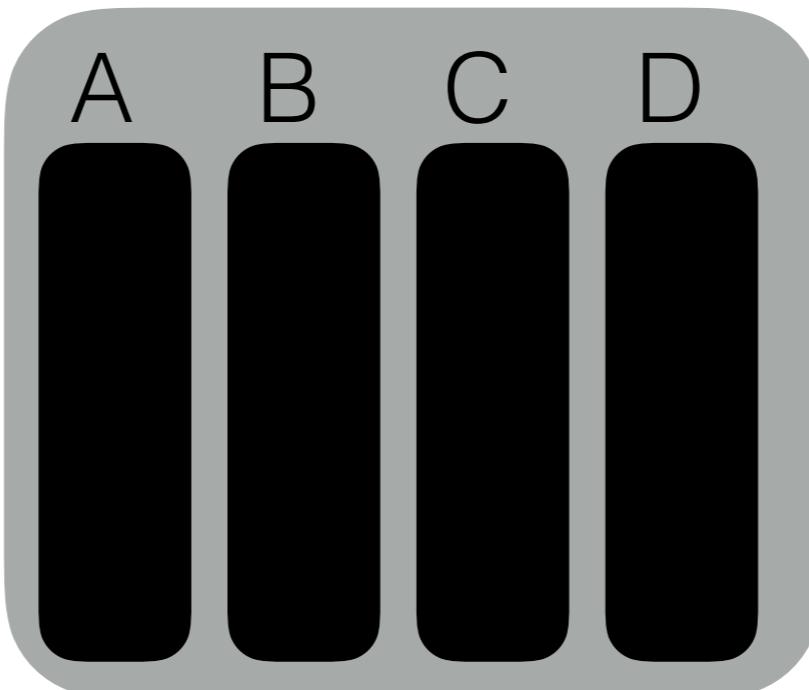
read page in L1
update
persist to L2



what do we need to remember (log)

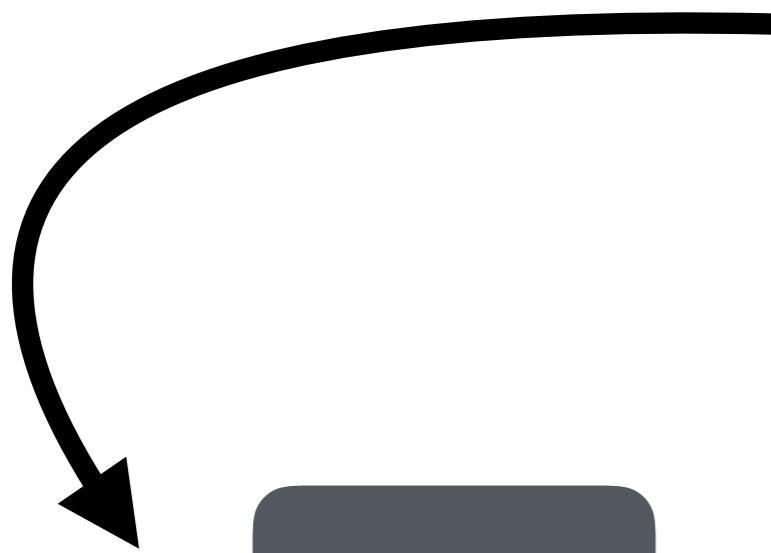


update all rows
where A=v1 & B=v2
to (a=a/2,b=b/4,c=c-3,d=d+2)

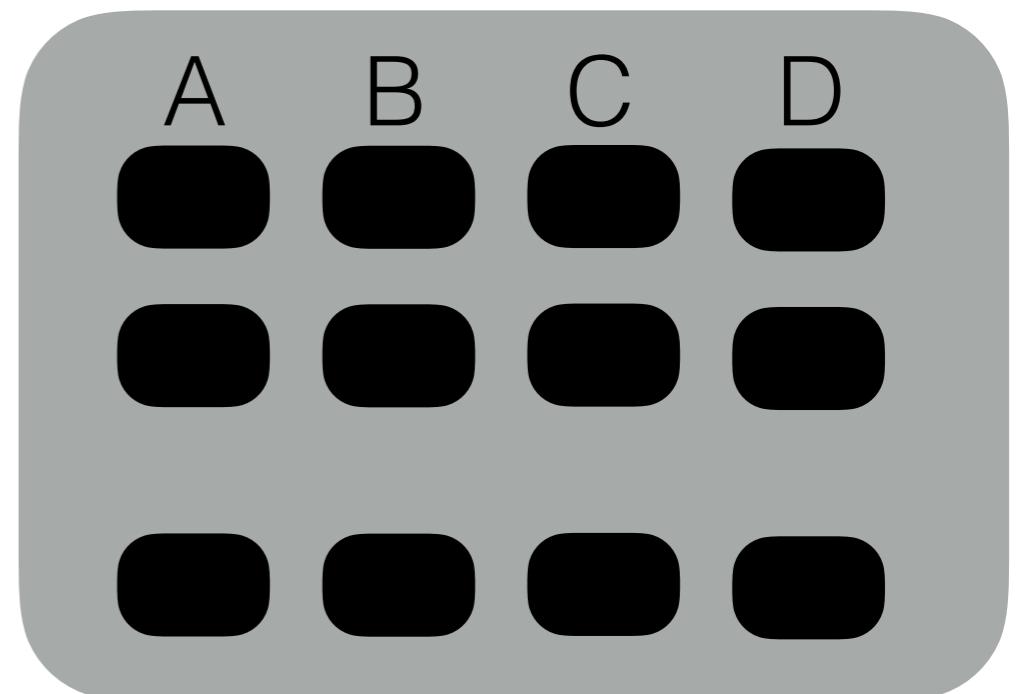
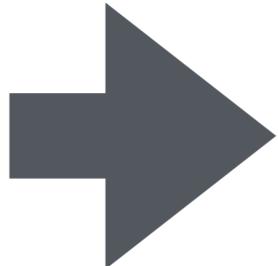


search (scan/index)
to find row to update

select+project actions

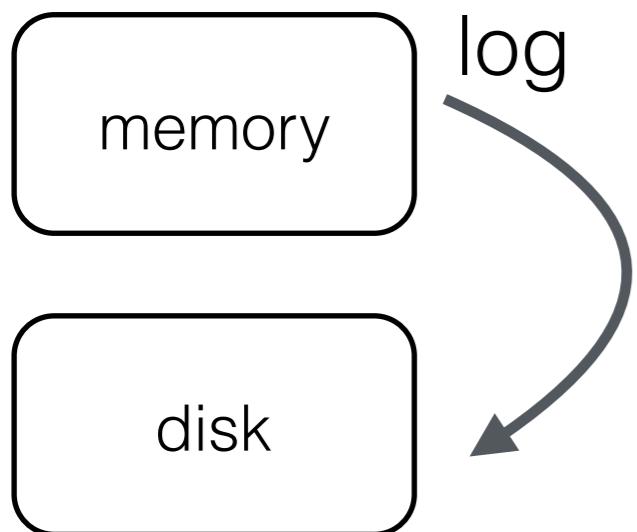


list of rowIDs (positions)
(sort)

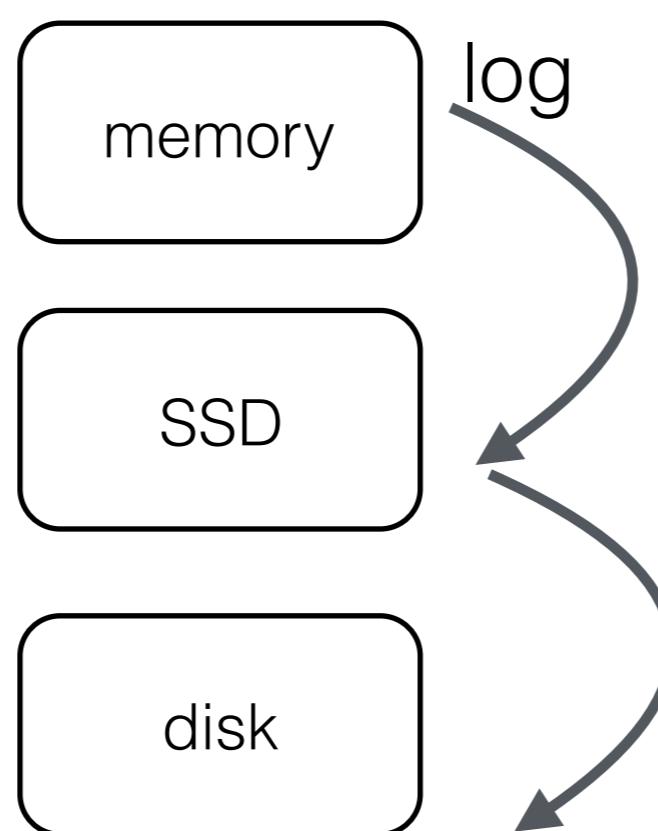


restart: get set of positions again, and either resume from last written page
or undo all previously written pages

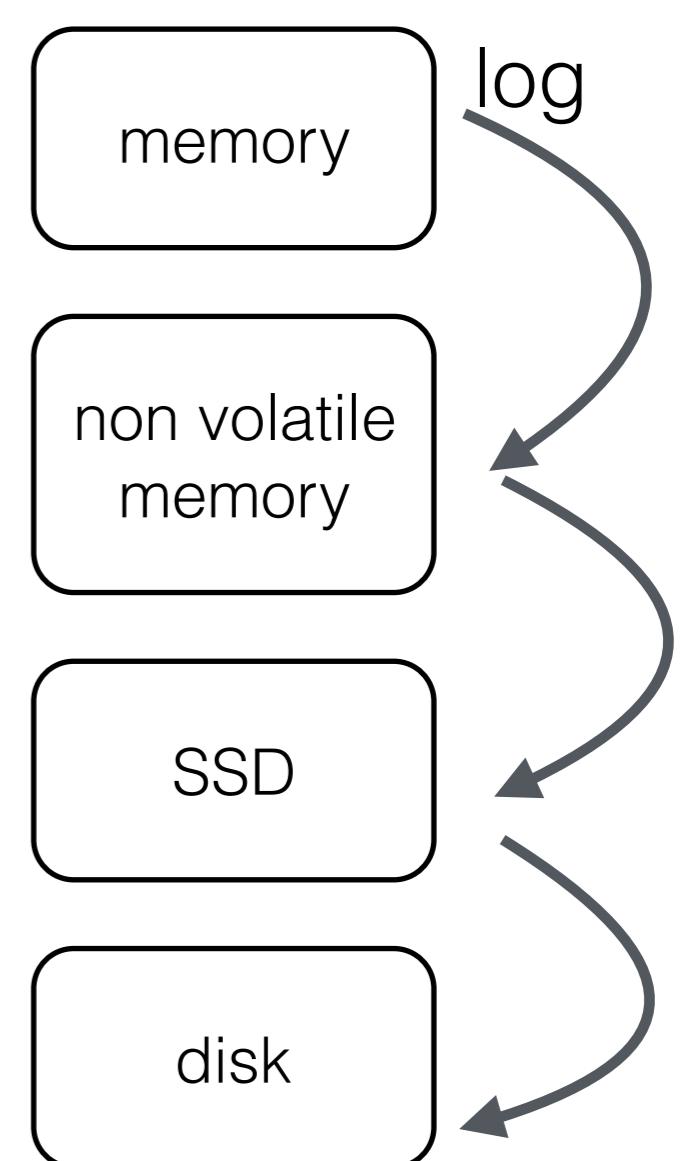
buffer updates



buffer updates



buffer updates



flash disks

read/write asymmetry

need to erase a block to write it

finite number of erase cycles

MaSM: efficient online updates in data warehouses

Manos Athanassoulis, Shimin Chen, Anastasia Ailamaki, Phillip Gibbons, Radu Stoica
ACM SIGMOD Conference, 2011



C Mohan, IBM Research
ACM SIGMOD Edgar F. Codd Innovations award 1993

(rumor has it he has his own Facebook server)
(also check his noSQL lecture)

ARIES: A Transaction Recovery Method Supporting Fine-Granularity Locking and Partial Rollbacks Using Write-Ahead Logging

C. Mohan, Donald J. Haderle, Bruce G. Lindsay, Hamid Pirahesh, Peter M. Schwarz
ACM Transactions on Database Systems, 1992

classic example

joe owes mike 100\$

both joe and mike have a Bank of Bla account



what about logging and recovery during e-shopping
e.g., shopping cart, wish list, check out

Quantifying eventual consistency with PBS

Peter Bailis, Shivaram Venkataraman, Michael J. Franklin,
Joseph M. Hellerstein, Ion Stoica
Communications of the ACM, 2014

db

complex
legacy
tuning
expensive
...

noSQL

simple
clean
just enough
...

as apps become
more **complex**

as apps need to be
more **scalable**

newSQL

why noSQL

**Mesa: Geo-Replicated, Near Real-Time,
Scalable Data Warehousing**

Ashish Gupta et al.
International Conference on
Very Large Databases (**VLDB**), 2014

F1: A Distributed SQL Database That Scales

Jeff Shute et al.
International Conference on
Very Large Databases (**VLDB**), 2013

(more in 265)

all or nothing

remember what we changed so we can undo or resume
(depends on applications semantics)

transaction

any database query that
can/should be seen as a single task

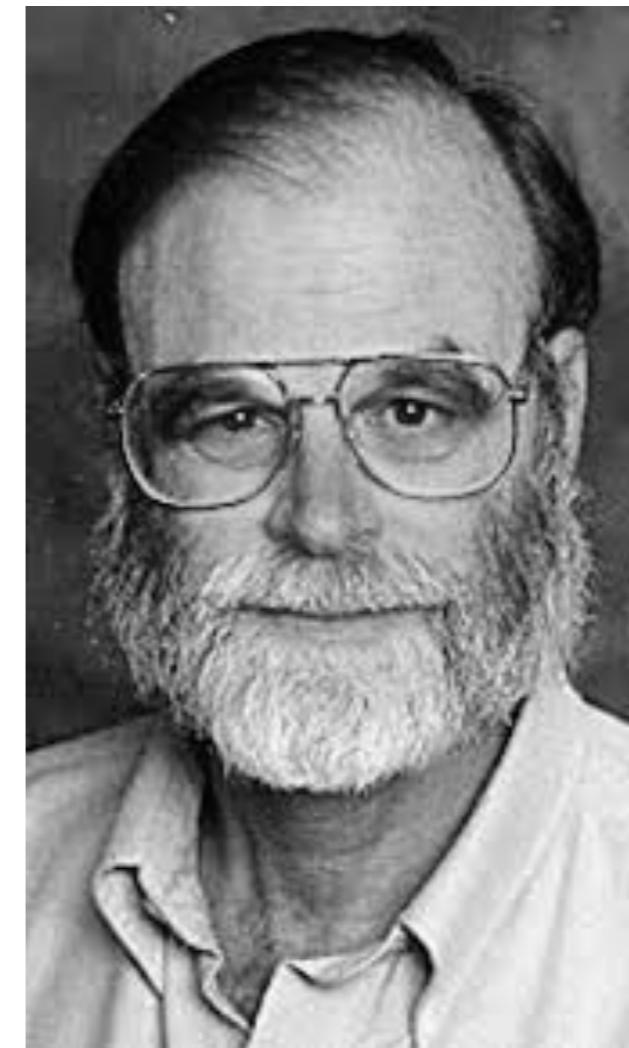
parse
optimize

plan(
read X
write Z

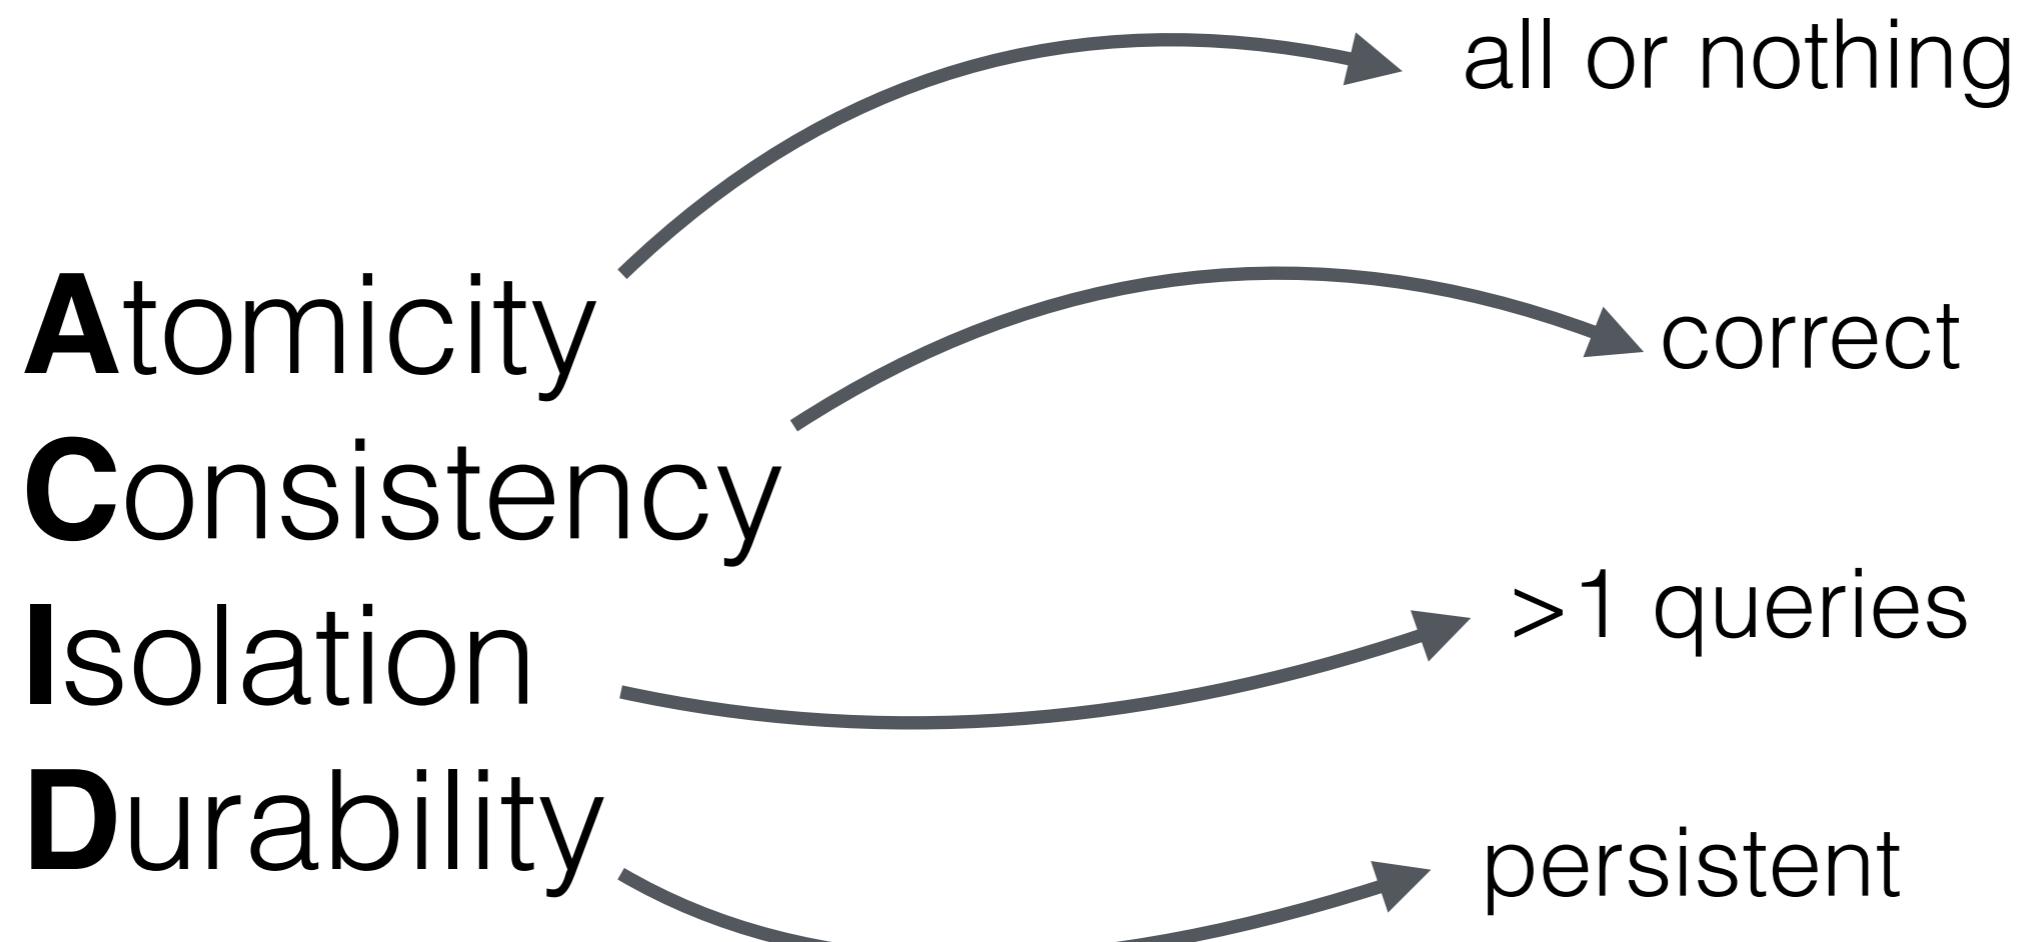
...
)

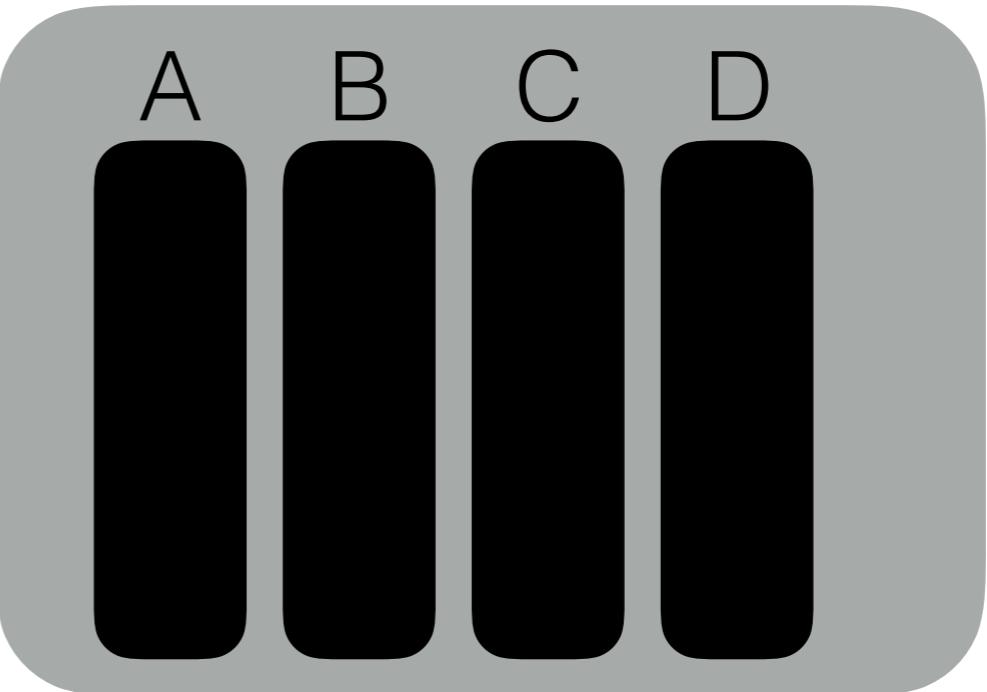


Atomicity **C**onsistency **I**solation **D**urability



Jim Gray, IBM, Tandem, DEC, Microsoft
ACM Turing award
ACM SIGMOD Edgar F. Codd Innovations award





queries arrive concurrently

- q1) select $A > v_1$, $B = B + b$, $C = C - c$
- q2) select $B > v_2$, $\text{avg}(C)$
- q3) select $C < v_3 \& B > v_4$, $A = C$
- q4) write more queries

Goal:

process queries in parallel
leave db in a consistent state
return correct results

to be released after class



textbook: chapters 16, 17, 18



updates 2.0

DATA SYSTEMS

prof. Stratos Idreos



HARVARD
School of Engineering
and Applied Sciences